



Cloud Manager documentation

Cloud Manager

NetApp
October 11, 2021

This PDF was generated from <https://docs.netapp.com/us-en/occm/index.html> on October 11, 2021.
Always check docs.netapp.com for the latest.

Table of Contents

Cloud Manager documentation	1
Release notes	2
Cloud Manager.....	2
Important changes in Cloud Manager	45
Cloud Volumes ONTAP AMI change	45
SaaS changes	45
Machine type changes	45
Account settings	45
New permissions	46
New endpoints	50
Get started with Cloud Manager.....	51
Learn about Cloud Manager	51
Getting started checklist	52
Signing up to NetApp Cloud Central	54
Logging in to Cloud Manager	55
Set up a Cloud Central account	56
Set up a Connector	65
Where to go next	87
Manage Cloud Volumes ONTAP	88
Learn	88
Get started in AWS	125
Get started in Azure	177
Get started in GCP	199
Provision and manage storage	230
Manage storage VMs	259
Monitor performance	272
Improving protection against ransomware	279
Administer	281
Provision volumes using a file service	307
Azure NetApp Files	307
Amazon FSx for ONTAP	319
Cloud Volumes Service for AWS	335
Cloud Volumes Service for GCP	357
Manage on-prem ONTAP clusters	373
Discovering ONTAP clusters	373
Using Active IQ data to manage ONTAP clusters	376
Managing storage for ONTAP clusters	380
Remove an on-prem ONTAP working environment	384
Replicate ONTAP data	385
Learn about the Replication service	385
Replicating data between systems	386
Managing data replication schedules and relationships	389
Learn about replication policies	390

Use templates to standardize resource creation	394
Learn about Application Templates	394
Template building blocks	396
Build application templates for your organization.	399
Check resources for template compliance	420
Create resources using templates	424
Back up data to the cloud	429
Learn about Cloud Backup	429
Get started	436
Set up licensing for Cloud Backup	490
Managing backups for Cloud Volumes ONTAP and on-premises ONTAP systems	493
Restoring data from backup files	503
Cross-account and cross-region configurations	510
Copy and synchronize data	526
Cloud Sync overview	526
Get started	529
Tutorials	567
Managing sync relationships	578
Manage data brokers	583
Creating and viewing reports to tune your configuration	588
Uninstalling the data broker	590
Cloud Sync APIs	590
Cloud Sync technical FAQ	595
Gain insight into data privacy	603
Learn about Cloud Data Sense	603
Get started	609
Viewing governance details about the data stored in your organization	656
Viewing compliance details about the data stored in your organization	659
Managing your private data	671
Adding personal data identifiers using Data Fusion	688
Viewing compliance reports	691
Responding to a Data Subject Access Request	697
Categories of private data	699
Removing data sources from Cloud Data Sense	704
Frequently asked questions about Cloud Data Sense	707
Tier on-premises data to the cloud	713
Learn about Cloud Tiering	713
Get started	717
Set up licensing for Cloud Tiering	746
Measure network latency and throughput performance	751
Managing data tiering from your clusters	752
Get an overview of data tiering from your clusters	760
Cloud Tiering technical FAQ	761
Reference	764
Organize your resources using tags	768

Learn about tagging	768
Manage tags for your resources	771
Enable real-time global file sharing	775
Learn about Global File Cache	775
Before you begin to deploy Global File Cache	780
Getting started	784
Before you begin to deploy Global File Cache Edge instances	795
Deploy Global File Cache Edge instances	801
End-user training	804
Additional information	804
Optimize cloud compute costs	806
Learn about the Compute service	806
Start optimizing your cloud compute costs	807
Viewing your Amazon S3 buckets	810
Administer Cloud Manager	812
Manage Connectors	812
Manage cloud provider credentials	827
Adding NetApp Support Site accounts to Cloud Manager	850
Managing your account	850
Monitoring operations in your account	859
Managing an HTTPS certificate for secure access	862
Removing Cloud Volumes ONTAP working environments	863
Configuring a Connector to use a proxy server	864
Reference	866
Automate with the REST API	878
Automation using the Cloud Manager REST API platform	878
Automation resources for infrastructure as code	878
Where to get help and find more information	879
Earlier versions of Cloud Manager documentation	881
Legal notices	882
Copyright	882
Trademarks	882
Patents	882
Privacy policy	882
Open source	882

Cloud Manager documentation

Release notes

Cloud Manager

What's new in Cloud Manager 3.9

Cloud Manager typically introduces a new release every month to bring you new features, enhancements, and bug fixes.



Looking for a previous release?

[What's new in 3.8](#)

[What's new in 3.7](#)

Cloud Manager 3.9.11 (4 Oct 2021)

- [Amazon FSx for ONTAP enhancements](#)
- [Cloud Volumes ONTAP enhancements](#)
- [Cloud Backup enhancements](#)
- [Cloud Sync enhancements](#)
- [Application Templates enhancements \(11 Oct\)](#)
- [Resource Tagging enhancements \(11 Oct\)](#)

Amazon FSx for ONTAP enhancements

- Now you can [create CIFS volumes in FSx for ONTAP using Cloud Manager](#).
- Now you can [edit FSx for ONTAP volumes using Cloud Manager](#).

Cloud Volumes ONTAP enhancements

- Cloud Manager can now deploy and manage Cloud Volumes ONTAP 9.10.0.

[Learn about the new features included in this release of Cloud Volumes ONTAP](#).

- We reduced the amount of time that it takes to deploy a Cloud Volumes ONTAP working environment in Microsoft Azure or in Google Cloud when normal write speed is enabled. The deployment time is now 3-4 minutes shorter on average.

Cloud Backup enhancements

- The size of each backup file is now displayed in the Backup List page and when performing a volume or file restore operation.

This is useful if you want to delete large backup files that are unnecessary, or so you can compare backup file sizes to identify any abnormal backup files that could be the result of a malicious software attack.

- A new TCO calculator is available to help you understand the total cost of ownership for Cloud Backup, and to compare these costs to traditional backup solutions and estimate potential savings. Check it out [here](#).
- Now you can easily [unregister Cloud Backup for a working environment](#) if you no longer want to use

backup functionality (or be charged) for that working environment.

Cloud Sync enhancements

- Cloud Sync now supports additional sync relationships for [Box](#) when using the Cloud Sync API:
 - Amazon S3 to Box
 - IBM Cloud Object Storage to Box
 - StorageGRID to Box
 - Box to an NFS server
 - Box to an SMB server

[Learn how to set up a sync relationship using the API.](#)

- You can now [create a report](#) for SFTP paths.

Application Templates enhancements (11 Oct)

- Now you can create a duplicate of an existing template. This can save a lot of time in case you want to create a new template that is very similar to an existing template. Just make the duplicate with a new name and change the couple items that make the template unique. [See more details about creating templates.](#)

Resource Tagging enhancements (11 Oct)

- Support has been added so you can label certain GCP resources.

Now you can manage labels on your GCP resources using the Cloud Manager Tagging service. You can view GCP labels and label values that have been applied to resources, and you can apply those labels to other GCP resources that you are managing. [See the GCP resources that you can label.](#)

Cloud Manager 3.9.10 build 2 (12 September 2021)

We fixed a Cloud Backup bug where a volume restore operation failed when the language code for the volume was different than the language code for the storage VM.

Cloud Manager 3.9.10 (2 September 2021)

- [Support for Amazon FSx for ONTAP](#)
- [Cloud Volumes ONTAP enhancement](#)
- [Cloud Data Sense enhancements](#)
- [Cloud Sync enhancements](#)
- [Cloud Tiering enhancements](#)
- [Replication enhancements](#)
- [Cloud Backup enhancements](#)
- [Resource Tagging enhancements](#)
- [A new Notification Service has been added](#)

Support for Amazon FSx for ONTAP

- [Amazon FSx for ONTAP](#) is a fully managed service allowing customers to launch and run file systems powered by NetApp's ONTAP storage operating system. FSx for ONTAP provides the same features, performance, and administrative capabilities NetApp customers use on premises, with the simplicity, agility, security, and scalability of a native AWS service.

[Learn about Amazon FSx for ONTAP.](#)

- You can configure an FSx for ONTAP working environment in Cloud Manager.

[Create an Amazon FSx for ONTAP working environment.](#)

- Using a Connector in AWS and Cloud Manager, you can create and manage volumes, replicate data, and integrate FSx for ONTAP with NetApp cloud services, such as Data Sense and Cloud Sync.

[Get started with Cloud Data Sense for Amazon FSx for ONTAP.](#)

Cloud Volumes ONTAP enhancement

Data is automatically encrypted on Cloud Volumes ONTAP in Azure using [Azure Storage Service Encryption](#) with a Microsoft-managed key. But you can now use your own customer-managed encryption key instead by completing the following steps:

1. From Azure, create a key vault and then generate a key in that vault.
2. From Cloud Manager, use the API to create a Cloud Volumes ONTAP working environment that uses the key.

[Learn more about these steps.](#)

Cloud Data Sense enhancements

- Added support for scanning data on NFS volumes on Amazon FSx for ONTAP systems. [See how to configure scanning for your FSx for ONTAP systems.](#)
- The capability to add "Status" information to your files using Data Sense has changed terminology to use "Tags". These are file level tags - not to be confused with resource level tagging that can be applied to volumes, EC2 instances, virtual machines, etc.

Cloud Sync enhancements

You can now sync data to or from an Amazon FSx for ONTAP file system.

- [Learn about Amazon FSx for ONTAP](#)
- [View supported sync relationships](#)
- [Learn how to create a sync relationship for Amazon FSx for ONTAP](#)

Cloud Tiering enhancements

- New Cloud Tiering BYOL license replaces FabricPool license

A new **Cloud Tiering** license is now available for tiering configurations that are supported within Cloud Manager using the Cloud Tiering service. It is a floating license that you can use across multiple on-premises ONTAP clusters. The **FabricPool** license that you may have used in the past is retained only for configurations that aren't supported within Cloud Manager.

[Learn more about the new Cloud Tiering license.](#)

- Now you can tier inactive data from your on-prem ONTAP clusters to any Object Storage service which uses the Simple Storage Service (S3) protocol.

[See how to tier data to S3-compatible object storage.](#)

Replication enhancements

You can now replicate data from a Cloud Volumes ONTAP system or an on-premises ONTAP cluster to an Amazon FSx for ONTAP file system.

[Learn about the Replication service.](#)

Cloud Backup enhancements

- Now you can create an on-demand backup at any time to capture the current state of a volume. This is useful if important changes have been made to a volume and you don't want to wait for the next scheduled backup to protect that data.

[See how to create an on-demand backup.](#)

- When configuring backups to Amazon S3 from an on-premises ONTAP system, now you can define a connection to a Private Interface Endpoint in the activation wizard. This allows you to use a network interface that connects your on-prem system privately and securely to a service powered by AWS PrivateLink. [See details about this option.](#)
- For additional security and control, you can choose your own customer-managed keys for data encryption in the activation wizard instead of using the default Amazon S3 encryption keys.

This is available when configuring backups from an on-premises ONTAP system or from a Cloud Volumes ONTAP system in AWS.

- The restriction of being able to browse a single directory with flat files up to a maximum of 30,000 files has been removed. Now you can restore files from directories that have a larger number of files.

Resource Tagging enhancements

- Now you can manage tags on your Azure resources using the Cloud Manager Tagging service. You can view Azure tags and tag values that have been applied to resources, and you can apply those tags to other Azure resources that you are managing.

[See the Azure resources that you can tag.](#)

The Cloud Manager Operator IAM role associated with the Connector needs the latest permissions to tag Azure resources. [See the new permissions that are required.](#)

- Some additional AWS EC2 resources can now be tagged. [See the AWS resources that you can tag.](#)

The Cloud Manager IAM role associated with the Connector needs the latest permissions to tag AWS EC2 instances. [See the new permissions that are required.](#)

A new Notification Service has been added

The Notification service has been introduced so you can view the status of Cloud Manager operations that you

have initiated during your current login session. You can verify whether the operation was successful, or if it failed. [See how to monitor operations in your account.](#)

Cloud Manager 3.9.9 build 2 (18 Aug 2021)

This patch update includes bug fixes.

Cloud Manager 3.9.9 update (8 Aug 2021)

- [Application Template enhancements](#)
- [Resource Tagging enhancements](#)

Application Template enhancements

- Now you can create a Cloud Volumes ONTAP working environment, and your first volume, using templates. This support is provided only for AWS environments at this time, and only for single-node clusters. See how to [create a template for a Cloud Volumes ONTAP working environment](#).
- A new feature enables you to add tags to a volume in a volume template. Tagging enables you to group different resources to identify applications, regions, or departments in order to perform automation or to allocate costs to certain departments or regions.

Resource Tagging enhancements

Now you can manage tags from your AWS EC2 Instances in Cloud Manager. You can view AWS tags and tag values that have been applied to EC2 Instances, and you can apply those tags to other EC2 Instances that you are managing. [Learn more about tagging](#).

Cloud Manager 3.9.9 (1 Aug 2021)

- [Cloud Backup enhancements](#)
- [Cloud Sync enhancements](#)
- [Cloud Data Sense enhancements](#)
- [Monitoring enhancement](#)
- [Connector enhancement](#)

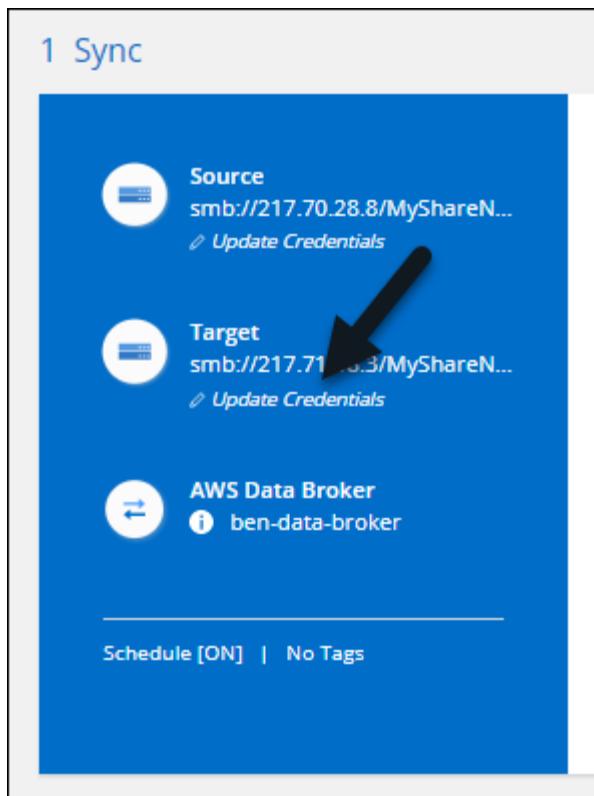
Cloud Backup enhancements

- When configuring backups to Azure Blob from an on-premises ONTAP system, you can define a connection to an Azure Private Endpoint in the activation wizard. This allows you to use a network interface that connects you privately and securely to a service powered by Azure Private Link.
- An Hourly backup policy is now supported in addition to the existing Daily, Weekly, and Monthly policies. The Hourly backup policy provides a minimal Recovery Point Objective (RPO).

Cloud Sync enhancements

- Cloud Sync now enables you to update the data broker with the latest credentials of the source or target in an existing sync relationship.

This enhancement can help if your security policies require you to update credentials on a periodic basis. [Learn how to update credentials](#).



- When creating a sync relationship, you can now add tags to the object storage target in a sync relationship.

Adding tags is supported with Amazon S3, Azure Blob, Google Cloud Storage, IBM Cloud Object Storage, and StorageGRID.

Relationship Tags

Cloud Sync assigns the relationship tags to all of the files transferred to the S3 bucket.

This enables you to search for the transferred files by using the tag values.

Save on Object's Tags Save On Object's Metadata

Tag Key Up to 128 characters	Tag Value Up to 256 characters
---------------------------------	-----------------------------------

+ Add Relationship Tag Optional Field | [Up to 5]

- Cloud Sync now supports [Box](#) as the source in a sync relationship to Amazon S3, StorageGRID, and IBM Cloud Object Storage when using the Cloud Sync API.

[Learn how to set up a sync relationship using the API.](#)

- When you deploy a data broker in Google Cloud, you can now choose whether to enable or disable a public IP address for the virtual machine instance.

[Learn how to deploy a data broker in Google Cloud.](#)

- When you choose the source or target volume for Azure NetApp Files, Cloud Sync now displays a dual-protocol volume no matter which protocol you chose for the sync relationship.

Cloud Data Sense enhancements

- Ability to manage file settings for multiple files at a time.

In earlier versions of Cloud Data Sense you could perform the following actions on one file at a time: add a status tag, assign a user, and add an AIP label. Now you can select multiple files from the Data Investigation page and perform these actions on multiple files.

- When viewing the Age of Data graph in the Governance dashboard, in addition to viewing data based on the last time it was modified, now you can view the data by when it was created or by when it was last accessed (when it was read).

This information is provided in the Data Mapping Report as well.

- When deploying Data Sense on-premises, now you can install scanning software on additional on-prem hosts when you plan to scan configurations that include petabytes of data. These additional *scanner nodes* provide increased processing power when scanning very large configurations.

See how to [deploy Data Sense software on multiple hosts](#).

Monitoring enhancement

We changed the default name of the Acquisition Unit instance to CloudInsights-AU-*UUID* so that the name is more descriptive (the UUID is a generated hash).

Cloud Manager deploys this instance when you enable the Monitoring service on a Cloud Volumes ONTAP working environment.

[Learn more about the Monitoring service.](#)

Connector enhancement

The Connector is now supported on a host that's running Red Hat Enterprise Linux 7.9.

[View system requirements for the Connector.](#)

Cloud Manager 3.9.8 build 3 (25 July 2021)

This patch update includes bug fixes and security improvements.

Cloud Manager 3.9.8 update (13 July 2021)

- [Application Template enhancements](#)
- [New Resource Tagging feature](#)

Application Template enhancements

- Support has been added so you can add in the template that you want to [replicate the data in the volume you are creating](#) to another volume using the [Replication service](#). When you replicate data to other NetApp storage systems and continually update the secondary data, your data is kept current and remains available whenever you need it.

- Now you can download a report that includes all the volumes that have "drifted" from your template settings instead of just viewing this information for a single volume in the Dashboard. In this manner you can identify these volumes and assign someone to bring the volumes back into compliance. See [how to download your drift report](#).

New Resource Tagging feature

- A new Cloud Manager feature enables you to apply tags to your existing ONTAP resources to help organize and manage those resources. Tags are metadata that you can use to group resources to identify applications, environments, regions, billing codes, cloud providers, and more.

[Learn more about tagging.](#)

Cloud Manager 3.9.8 (7 July 2021)

- [Cloud Volumes ONTAP enhancements](#)
- [Cloud Backup enhancements](#)
- [Cloud Tiering enhancements](#)
- [Cloud Data Sense enhancements](#)
- [Cloud Sync enhancements](#)
- [Connector enhancement](#)
- [Support Dashboard enhancement](#)
- [Digital Wallet enhancement](#)

Cloud Volumes ONTAP enhancements

This release of Cloud Manager includes enhancements to the management of Cloud Volumes ONTAP.

Enhancements available in all cloud providers

- New charging methods are available for Cloud Volumes ONTAP.
 - Capacity-based BYOL:** A capacity-based license enables you to pay for Cloud Volumes ONTAP per TiB of capacity. The license is associated with your NetApp account and enables you to create as many Cloud Volumes ONTAP systems, as long as enough capacity is available through your license. Capacity-based licensing is available in the form of a package, either *Essentials* or *Professional*.
 - Freemium offering:** Freemium enables you to use all Cloud Volumes ONTAP features free of charge from NetApp (cloud provider charges still apply). You're limited to 500 GiB of provisioned capacity per system and there's no support contract. You can have up to 10 Freemium systems.

[Learn more about these licensing options.](#)

Here's an example of the charging methods that you can choose from when deploying a new Cloud Volumes ONTAP system in Azure:

Cloud Volumes ONTAP Charging Methods

[Learn more about our charging methods](#)



Pay-As-You-Go by the hour



Bring your own license

Bring your own license type

Capacity-Based

Package

Professional



Freemium (Up to 500GB)

- Write once, read many (WORM) storage is no longer in Preview and is now available for general use with Cloud Volumes ONTAP. [Learn more about WORM storage](#).

Enhancements available in AWS

Starting with the 9.9.1 release, Cloud Volumes ONTAP now supports the m5dn.24xlarge instance type with the following charging methods: PAYGO Premium, bring your own license (BYOL), and Freemium.

[View supported configurations for Cloud Volumes ONTAP in AWS.](#)

Enhancements available in Azure

- When creating a Cloud Volumes ONTAP system in Azure, you now have the option to select an existing resource group for the VM and its associated resources.

Location & Connectivity

Location <p>Azure Region <input style="width: 100%; height: 30px; border: 1px solid #ccc; padding: 5px; margin-bottom: 10px;" type="text" value="WEST US"/></p> <p>Availability Zone <i>(Optional)</i> <input style="width: 100%; height: 30px; border: 1px solid #ccc; padding: 5px;" type="text" value="Select an Availability Zone"/></p>	Connectivity <p>Resource Group <input checked="" style="margin-right: 10px;" type="radio" value="existing"/> Create a new group <input checked="" style="margin-left: 10px;" type="radio" value="existing"/> Use an existing group</p> <p>Resource Group Name <input style="width: 100%; height: 30px; border: 1px solid #ccc; padding: 5px;" type="text" value="RG1"/></p>
---	---

The following permissions enable Cloud Manager to remove Cloud Volumes ONTAP resources from a resource group, in case of deployment failure or deletion:

```
"Microsoft.Network/privateEndpoints/delete",
"Microsoft.Compute/availabilitySets/delete",
```

Be sure to provide these permissions to each set of Azure credentials that you've added to Cloud Manager. You can find the latest list of permissions on the [Cloud Manager policies page](#).

- As a security enhancement, Cloud Manager now disables **Blob public access** when creating a storage account for Cloud Volumes ONTAP.
- By default, Cloud Manager now enables an Azure Private Link connection on the boot diagnostics storage account for new Cloud Volumes ONTAP systems.

This means *all* storage accounts for Cloud Volumes ONTAP will now use a private link.

[Learn more about using an Azure Private Link with Cloud Volumes ONTAP.](#)

Enhancements available in Google Cloud

- Starting with the 9.9.1 release, Cloud Volumes ONTAP now supports Balanced persistent disks (pd-balanced).

These SSDs balance performance and cost by providing lower IOPS per GiB.

- The custom-4-16384 machine type is no longer supported with new Cloud Volumes ONTAP systems.

If you have an existing system running on this machine type, you can keep using it, but we recommend switching to the n2-standard-4 machine type.

[View supported configurations for Cloud Volumes ONTAP in GCP.](#)

Cloud Backup enhancements

- Cloud Backup now allows you to create backups using a different account/subscription than the one you

are using for your Cloud Volumes ONTAP system. You can also create backup files in a different region than the one in which your Cloud Volumes ONTAP system is deployed.

This capability is available when using AWS or Azure, and only when enabling backup on an existing working environment - it is not available when creating a new Cloud Volumes ONTAP working environment.

- For additional security and control, you can choose your own customer-managed keys for data encryption in the activation wizard instead of using the default Microsoft-managed encryption keys.

This is available when configuring backups from an on-premises ONTAP system or from a Cloud Volumes ONTAP system in Azure.

- The restriction of restoring only 8 files at a time using single-file restore from your backup files has been removed. Now you can restore up to 100 files at a time.

Cloud Tiering enhancements

When tiering data to Azure Blob storage, now your Connector can be running on your premises. You are no longer required to use a Connector installed in an Azure VNet.

Cloud Data Sense enhancements

- A new feature enables you to [move any source files that Data Sense is scanning to any NFS share](#). This allows you to move sensitive or security-related files to a special area so you can do more analysis.
- You can now choose to quickly map data into categories instead of doing a full classification scan. This enables you to [view the Data Mapping report](#) from the Governance Dashboard to get an overview of your data when there are certain data sources that you do not need to run a complete scan on.
- Now you can [assign a file to a specific Cloud Manager user](#) so that person can be responsible for any follow-up actions that need to be done on the file. This capability can be used with the existing feature to add custom Status tags to a file.

A new Filter in the Investigation page enables you to easily view all files that have the same person in the "Assigned To" field.

- Some users with smaller scanning requirements have asked to be able to use a smaller Cloud Data Sense instance. Now you can. There are some limitations when using these smaller instances, so [see what these restrictions are first](#).
- Data scans have a negligible impact on your storage systems and on your data. However, if you are concerned with even a very small impact, you can configure Data Sense to perform "slow" scans now.
- The Last Accessed Time value has been added to the File Details page and to reports that you output in CSV format so you can see when users have last accessed the file.

Cloud Sync enhancements

- Cloud Sync now supports sync relationships between ONTAP S3 Storage and a Google Cloud Storage bucket from the user interface.

[View supported sync relationships](#).

- Cloud Sync can now copy object metadata and tags between object-based storage when you create a sync relationship and enable a setting.

[Learn more about the Copy for Objects setting](#).

- You can now set up the data broker to access credentials from an external HashiCorp Vault by authenticating with a Google Cloud service account.

[Learn more about using a HashiCorp Vault with a data broker.](#)

- When setting up a sync relationship to an AWS S3 bucket, the Sync Relationship wizard now enables you to define the tags or metadata that you want to save on the objects in the target S3 bucket.

The tagging option was previously part of the sync relationship's settings.

Connector enhancement

We redesigned the **Add Connector** wizard to add new options and to make it easier to use. You can now add tags, specify a role (for AWS or Azure), upload a root certificate for a proxy server, view code for Terraform automation, view progress details, and more.

- [Create a Connector in AWS](#)
- [Create a Connector in Azure](#)
- [Create a Connector in GCP](#)

Support Dashboard enhancement

NetApp Support Site (NSS) accounts are now managed from the Support Dashboard, rather than from the Settings menu. This change makes it easier to find and manage all support-related information from a single location.

NSS User Name	NSS User ID	Attached Working Environments
testcloud2	61e6b48b-371e-4681-a...	—

Digital Wallet enhancement

The Digital Wallet page is now located in its own tab under the Resources section. [Learn how to manage licenses from your Digital Wallet.](#)

The screenshot shows the Cloud Manager dashboard. At the top, there are account and workspace dropdowns. Below the header, a navigation bar includes 'Canvas' (selected), 'Replication', 'Backup & Restore', 'K8s', 'Data Sense', 'File Cache', 'Compute', 'Sync', and 'All Services (+8)'. A red arrow points to the 'Digital Wallet' section, which has a blue icon of a gear and a lock, with the text 'View & Manage Digital Wallet'. Other sections visible include 'Canvas', 'Timeline', and 'File Cache'.

Cloud Manager 3.9.7 update (7 June 2021)

- Cloud Sync support for storage classes in Google Cloud
- Cloud Tiering enhancements
- Cloud Backup enhancements
- Application Templates enhancements
- Cloud Data Sense (Cloud Compliance) enhancements
- Global File Cache enhancements

Cloud Sync support for storage classes in Google Cloud

When a Google Cloud Storage bucket is the target in a sync relationship, you can now choose the storage class that you want to use. Cloud Sync supports the following storage classes:

- Standard
- Nearline
- Coldline
- Archive

Cloud Tiering enhancements

- In earlier releases you could transition tiered data from the Standard storage class to another storage class after 30 days when using Amazon S3 or Google Cloud object storage. This release adds some new functionality:
 - You can choose the number of days when tiered data will move to a more cost-effective tier. This is called data "life cycle management".
 - Azure Blob now supports the *Cool* access tier.
 - Google Cloud Storage supports moving tiered data to multiple storage tiers over time. For example, you can move tiered data from the *Standard* class to the *Nearline* class after 45 days, and then to the *Coldline* class after 75 days, and then to the *Archive* class after 270 days.
- Now you can proactively move data back to the performance tier from the cloud tier if you want to stop using tiering on a volume, or if you decide to keep all user data on the performance tier, but keep Snapshot copies on the cloud tier. This capability is available when using ONTAP 9.8 and greater.

See how to [migrate data back to the performance tier](#).

Cloud Backup enhancements

- Two known limitations for backing up data protection (DP) volumes have been resolved. Your systems must have ONTAP 9.8 or greater installed:
 - Before, cascaded backup worked only if the SnapMirror relationship type was Mirror-Vault or Vault. Now you can make backups if the relationship type is MirrorAllSnapshots.
 - Cloud Backup now can use any label for the backup as long as it is configured in the SnapMirror policy. The restriction of requiring labels with the names daily, weekly, or monthly is gone.

Application Templates enhancements

- A new feature enables you to conditionally enable certain actions when the user is running the template. For example, if a Cloud Volumes ONTAP volume is created with NetApp storage efficiency enabled, then Cloud Backup is also enabled on that volume. If storage efficiency is not enabled, then Cloud Backup is not enabled.
- You can now create a volume on an on-premises ONTAP system using templates.
- New functionality called "drift" has been added as an option when creating your templates.

This feature enables Cloud Manager to monitor the hard-coded values you entered for a parameter in a template. After a storage admin has created a volume using that template, if Cloud Manager later sees that the parameter value has been changed so that it no longer aligns with the template definition, you can see all the volumes that have "drifted" from the designed template. In this manner you can identify these volumes and make changes to bring them back into compliance.

- Now you can run a template from the Template Dashboard instead of having to open a working environment to launch the template there.

Cloud Data Sense (Cloud Compliance) enhancements

- Cloud Compliance has been renamed as **Cloud Data Sense** as of this release. With all the new Governance and other capabilities that have been included in the product, the Compliance name was not promoting the full set of capabilities.
- A new *Full Data Mapping* report is available from the Governance Dashboard to provide an overview of the data being stored in your corporate data sources to assist you with decisions of migration, back up, security, and compliance processes.

The report provides overview pages that summarize all your working environments and data sources, and then provides a breakdown for each working environment. Go [here](#) for more details.

- A new filter in the Data Investigation page enables you to view a list of all files that are duplicated across your storage systems.

This is helpful to identify areas where you can save storage space, or identify files that have specific permissions or sensitive information that you do not want duplicated across your storage. [See how to viewing all duplicated files.](#)

- You can add a custom Status tags to files that Data Sense is scanning. The Status is not added to the file in the same way as AIP Labels are added. The Status is just seen by Cloud Manager users so you can indicate if a file needs to be deleted, or checked for some reason. [See how to apply and view Status tags in your files.](#)

A new Filter in the Investigation page enables you to easily view all files that have a Status assigned.

- Cloud Data Sense can scan for Personal Identifiable Information (PII) in two additional types of files: .DCM and .DICOM.
- The File Size, Created Date, and Last Modified Date values have been added to reports that you output in CSV format. Created Date is also a new filter you can use to narrow down Investigation page search results.

Global File Cache enhancements

Global File Cache software version 1.1.0 has been released. A new "Edge Synchronization" feature is available that keeps multiple Edges at a remote office synchronized. When a file is fetched at one Edge, then the same file on all Edges participating in Edge Sync is updated and cached. See the [new and fixed issues in this release](#).

Cloud Manager 3.9.7 (30 May 2021)

- [Cloud Volumes ONTAP enhancements](#)
- [Digital Wallet](#)

Cloud Volumes ONTAP enhancements

This release of Cloud Manager includes enhancements to the management of Cloud Volumes ONTAP.

Enhancements available in AWS

- A new Professional Package enables you to bundle Cloud Volumes ONTAP and Cloud Backup Service by using an annual contract from the AWS Marketplace. Payment is per TiB. This subscription doesn't enable you to back up on-prem data.

If you choose this payment option, you can provision up to 2 PB per Cloud Volumes ONTAP system through EBS disks and tiering to S3 object storage (single node or HA).

Go to the [AWS Marketplace page](#) to view pricing details and go to the [Cloud Volumes ONTAP Release Notes](#) to learn more about this licensing option.

- Cloud Manager now adds tags to EBS volumes when it creates a new Cloud Volumes ONTAP working environment. The tags were previously created after Cloud Volumes ONTAP was deployed.

This change can help if your organization uses service control policies (SCPs) to manage permissions.

Enhancements available in all cloud providers

- If you enabled data tiering on a volume using the *auto* tiering policy, you can now adjust the minimum cooling period using the API.

[Learn how to adjust the minimum cooling period.](#)

- When you create a new NFS volume, Cloud Manager now displays custom export policies in ascending order, making it easier for you to find the export policy that you need.
- Cloud Manager now deletes older cloud snapshots of root and boot disks that are created when a Cloud Volumes ONTAP system is deployed and every time its powered down. Only the two most recent snapshots are retained for both the root and boot volumes.

This enhancement helps reduce cloud provider costs by removing snapshots that are no longer needed.

Note that a Connector requires a new permission to delete Azure snapshots. View the latest Cloud Manager policy for Azure.

"Microsoft.Compute/snapshots/delete"

Digital Wallet

A new **Digital Wallet** feature enables you to more easily view and manage your Cloud Volumes ONTAP licenses and Cloud Backup licenses from a single location.

[Learn more about Digital Wallet.](#)

The screenshot shows the Digital Wallet interface. At the top, there are two tabs: "Cloud Volumes ONTAP Licenses" (selected) and "Cloud Backup". Below the tabs is a circular progress bar labeled "24 Total Licenses". To the right of the bar, there's a "License Distribution" section with four categories: "BYOL License" (14), "PAYGO License" (6), "Free Trial" (2), and "Eval" (2). Below this, there are three buttons: "BYOL (14)", "PAYGO (6)", and "Eval (2)". A search icon is located to the right of the PAYGO button. The main area displays a table of Cloud Volumes ONTAP instances:

Cloud Volumes ONTAP Name	Type	System License	Extra Capacity License Status
Cloud Volumes ONTAP Number 1	Single Node	Active	None
Cloud Volumes ONTAP Number 2	Single Node	Active	3 Active
Cloud Volumes ONTAP Number 3	Single Node	Active	3 Active

Cloud Manager 3.9.6 update (24 May 2021)

Cloud Manager was updated to include the latest version of Cloud Volumes ONTAP.

Cloud Volumes ONTAP 9.9.1

Cloud Manager can now deploy and manage Cloud Volumes ONTAP 9.9.1.

[Learn about the new features included in this release of Cloud Volumes ONTAP.](#)

Cloud Manager 3.9.6 build 2 (11 May 2021)

We fixed a bug that caused failures when creating a Cloud Volumes ONTAP working environment in Azure.

Cloud Manager 3.9.6 update (5 May 2021)

- [Cloud Backup enhancements](#)
- [Monitoring enhancements](#)

- [Replication enhancement](#)
- [Account enhancement](#)
- [Cloud Compliance enhancements](#)

Cloud Backup enhancements

- The separate Backup and Restore Dashboards have been combined under a new **Backup & Restore** tab to make it easier for you to manage all your backup and restore operations from a single location. See [the Backup & Restore Dashboard](#) for details.
- Now you can create backups from your on-premises ONTAP systems to Google Cloud Storage or to your NetApp StorageGRID systems. See [Backing up to Google Cloud Storage](#) and [Backing up to StorageGRID](#) for details.
- A new feature in ONTAP 9.9.1 enables you to use System Manager to send backups of your on-premises ONTAP volumes to object storage you've set up through Cloud Backup. [See how to use System Manager to back up your volumes to the cloud using Cloud Backup](#).
- Backup policies have been improved with the following enhancements:
 - Now you create a custom policy that includes a combination of daily, weekly, and monthly backups.
 - When you change a backup policy, the change applies to all new backups **and** to all volumes using the original backup policy. In the past the change only applied to new volume backups.
- Some smaller improvements have also been made:
 - When configuring the cloud destination for your backup files, now you can select a different region than the region in which the Cloud Volumes ONTAP system resides.
 - The number of backup files you can create for a single volume has been increased from 1,019 to 4,000.
 - In addition to the earlier ability to delete all backup files for a single volume, now you can delete just a single backup file for a volume, or you can delete all backup files for an entire working environment, if needed.

Monitoring enhancements

- You can now enable the Monitoring service on a Cloud Volumes ONTAP working environment even if you have an existing Cloud Insights tenant.
- When you enable the Monitoring service, Cloud Manager sets up a free trial of Cloud Insights. On the 29th day, your plan now automatically transitions from the Trial Version to the [Basic Edition](#).

[Learn more about using the Monitoring service with Cloud Volumes ONTAP](#).

Replication enhancement

We redesigned the Replication tab for ease of use and to match the current look and feel of the Cloud Manager user interface.

The screenshot shows the Cloud Manager interface with the following key statistics at the top:

- 5 Volumes Relationships
- 32 TB Replicated Capacity
- 0 Currently Transferring
- 5 Healthy
- 0 Failed

Below these stats is a table titled "5 volumes relationships" with the following columns:

Health Status	Source Volume	Target Volume	Lag Duration	Status	Mirror State	Last Successful Transfer
Green checkmark	Volume1 WtName1	Volume1_copy WtName2	A day	idle	uninitialized	jan 1, 1970 02:00:10 am 10 GB
Green checkmark	Volume1 (WE Name) Source Volume	Volume1_copy (WE Name) Target Volume	Green checkmark Healthy	Replication Health		jan 1, 1970 02:00:10 am 10 GB
Green checkmark						jan 1, 1970 02:00:10 am 10 GB
Green checkmark						jan 1, 1970 02:00:10 am 10 GB
Green checkmark						jan 1, 1970 02:00:10 am 10 GB

Below the table, there's a section titled "Transfer info" with the following details:

Idle Transfer Status	check Transfer Type	50 TB Total Transfer Size	Jan 1 hour Total Transfer Time	low Transfer Priority
100 MB/s Max Transfer Rate	Log Duration	Snapshot copied Mirror State	500K / 2.0B Used Size in Used in Cloud	1.0 Network Compression Ratio

Account enhancement

The Timeline in Cloud Manager now shows actions and events related to account management. The actions include things like associating users, creating workspaces, and creating Connectors. Checking the Timeline can be helpful if you need to identify who performed a specific action, or if you need to identify the status of an action.

[Learn how to filter the Timeline to the Tenancy service.](#)

Cloud Compliance enhancements

- Cloud Compliance has been renamed to “Governance & Compliance”. There are two tabs in Cloud Manager: “Governance” and “Compliance”. The “Governance” tab brings you to the [Governance Dashboard](#) within the “Governance & Compliance” service, and the “Compliance” tab brings you to the [Compliance Dashboard](#).
- Scanning of data stored on Azure Blob is now supported when using the [MinIO service](#). See [Scanning object storage that uses S3 protocol](#) for details.
- New personal data type. Cloud Compliance can now find Austrian SSNs in files.

Cloud Manager 3.9.6 (2 May 2021)

- [Cloud Tiering enhancements](#)
- [Application Template enhancements](#)
- [Cloud Sync enhancements](#)

Cloud Tiering enhancements

- When selecting the volumes that you want to tier from an ONTAP system, now there's a checkbox in the Tier Volumes page to select **all** volumes to make it easier to apply the same policy to all volumes. [See how to select all volumes in the cluster](#).
- If you need to change the number of "cooling days" that determine how long data in a volume must remain

inactive before it is moved to object storage, now you can specify up to 183 days (up from 63 days) when using ONTAP 9.8 or greater.

Application Template enhancements

- The user interface has been enhanced in the AppTemplates service so that it is easier for template designers to move between actions and to see which action they are currently defining.
- Now you can integrate Cloud Compliance when creating a volume template for either Cloud Volumes ONTAP or Azure NetApp Files. So you can enable Compliance for each newly created volume, or enable Cloud Backup for each newly created volume... or create a template that enables both Backup and Compliance on the created volume.

Cloud Sync enhancements

- You can now view the errors found in reports and you can delete the last report or all reports.

[Learn more about creating and viewing reports to tune your configuration.](#)

- A new **Compare by** setting is now available for each sync relationship.

This advanced setting enables you to choose whether Cloud Sync should compare certain attributes when determining whether a file or directory has changed and should be synced again.

[Learn more about changing the settings for a sync relationship.](#)

Cloud Manager 3.9.5 (11 Apr 2021)

- [Cloud Volumes ONTAP enhancements](#)
- [Cloud Sync enhancements](#)
- [Cloud Compliance enhancements](#)
- [New Application Templates feature](#)
- [Connector enhancement](#)
- [Account enhancements](#)

Cloud Volumes ONTAP enhancements

This release of Cloud Manager includes enhancements to the management of Cloud Volumes ONTAP.

Enhancement available in all cloud providers

Cloud Manager now enables logical space reporting on the initial storage VM that it creates for Cloud Volumes ONTAP.

When space is reported logically, ONTAP reports the volume space such that all the physical space saved by the storage efficiency features are also reported as used.

Enhancements available in AWS

- Cloud Volumes ONTAP now supports *General Purpose SSD (gp3)* disks, starting with the 9.7 release. gp3 disks are the lowest-cost SSDs that balance cost and performance for a broad range of workloads.

[Learn more about using gp3 disks with Cloud Volumes ONTAP.](#)

- Cloud Volumes ONTAP no longer supports Cold HDD (sc1) disks.

Enhancement available in Azure

When Cloud Manager creates storage accounts in Azure for Cloud Volumes ONTAP, the TLS version for the storage account is now version 1.2.

Cloud Sync enhancements

- The standalone Cloud Sync service has been retired. You should now access Cloud Sync directly from Cloud Manager where all of the same features and functionality are available.

After logging in to Cloud Manager, you can switch to the Sync tab at the top and view your relationships, just like before.

- When setting up a sync relationship, you can choose from Google Cloud buckets in different projects, if you provide the required permissions to the data broker's service account.

[Learn how to set up the service account.](#)

- Cloud Sync now copies metadata between Google Cloud Storage and S3 providers (AWS S3, StorageGRID, and IBM Cloud Object Storage).
- You can now restart a data broker from Cloud Sync.

The screenshot shows the 'Data Broker Group' section of the Cloud Manager interface. It lists one data broker, 'data-broker1', which is connected to 'AWS'. The broker status is 'Active'. A context menu is open over the broker entry, with the 'Restart Data Broker' option highlighted.

- Cloud Sync now identifies when a data broker isn't running the latest software release. This message can help to ensure that you're getting the latest features and functionalities.

The screenshot shows a data broker entry for 'My First Azure Data Broker'. The broker is listed as 'Active' and connected to 'Azure' with a transfer rate of 30 mbps. A note at the bottom of the card states: 'Notice: This Data Broker isn't running the latest software release (1.8). If you experience performance problems please contact us.'

Cloud Compliance enhancements

- Added support for scanning NFS or CIFS file shares that reside on-premises or in the cloud.

Now you can scan file shares that reside on non-NetApp storage systems. See [scanning file shares](#) for details.

- Added support for scanning object storage that uses the S3 protocol.

In addition to scanning Amazon S3 buckets, now you can scan data from any Object Storage service which uses the S3 protocol. This includes NetApp StorageGRID, IBM Cloud Object Store, and more. See [scanning object storage](#) for details.

- The feature called "Highlights" has been renamed to "Policies". See [how to use Policies](#) to help in your compliance and governance efforts.
- Now you can see if there are duplicates of certain files in your storage systems. This is useful to identify areas where you can save storage space. It can also help to ensure that files containing sensitive information are not unnecessarily duplicated in your storage systems.

Learn how to [search for duplicate files](#).

- The [Governance dashboard](#) has added charts to show [top data repositories listed by data sensitivity](#) and [data listed by types of Open Permissions](#).

New Application Templates feature

Templates enable you to standardize resource creation in your working environments. For example, you can hard-code required parameters in a "volume template" that are later applied when a storage admin creates a volume. This can include required disk type, size, protocol, cloud provider, and more. You can also turn on certain services, like Cloud Backup, for every created volume.

This makes it easy for your storage admins to create volumes that are optimized for specialized workload requirements; such as databases or streaming services. And it makes life easier for your storage architects knowing that each volume is created optimally for each application. Learn about [Application Templates](#) and how you can use them in your environment.

Connector enhancement

If you configured a proxy server, you can now enable an option to send API calls directly to Cloud Manager without going through the proxy. This option is supported with Connectors that are running in AWS or in Google Cloud.

[Learn more about this setting](#).

Account enhancements

- You can now create a service account user.

A service account acts as a "user" that can make authorized API calls to Cloud Manager for automation purposes. This makes it easier to manage automation because you don't need to build automation scripts based on a real person's user account who can leave the company at any time. And if you're using federation, you can create a token without generating a refresh token from the cloud.

[Learn more about using service accounts](#).

- You can now allow private previews in your account to get access to new NetApp cloud services as they are made available as a preview in Cloud Manager.
- You can also allow third-party services in your account to get access to third-party services that are

available in Cloud Manager.

[Learn more about these options.](#)

Cloud Manager 3.9.4 update (8 Apr 2021)

Active IQ enhancements

- If Active IQ discovers unused Cloud Volumes ONTAP licenses in your account, you can click a button to create a new Cloud Volumes ONTAP system using the license. Or you can apply the license to an existing Cloud Volumes ONTAP system to extend the capacity of that license by 368 TB.

[See how to use your available licenses.](#)

Cloud Manager 3.9.4 update (15 Mar 2021)

Cloud Compliance enhancements

- A new [Governance dashboard](#) is now available so that you can increase the efficiency and control the costs related to the data on your organizations' storage resources.

For example, the dashboard identifies the amount of stale data, non-business data, and very large files in your systems so you can decide whether you want to move, delete, or tier some files to less expensive object storage.

- You can view a list of [all users or groups who have access to a file](#).
- Cloud Compliance is now supported in Government regions in AWS.

Cloud Manager 3.9.4 (8 Mar 2021)

- [Cloud Volumes ONTAP enhancements](#)
- [Connector enhancements](#)
- [Cloud Sync enhancements](#)
- [Cloud Tiering enhancements](#)
- [Active IQ enhancements](#)
- [ANF enhancements](#)

Cloud Volumes ONTAP enhancements

This release of Cloud Manager includes enhancements to the management of Cloud Volumes ONTAP.

Enhancement available in all cloud providers

Cloud Manager can now deploy and manage Cloud Volumes ONTAP 9.9.0.

[Learn about the new features included in this release of Cloud Volumes ONTAP.](#)

Enhancements available in AWS

- You can now deploy Cloud Volumes ONTAP 9.8 in the AWS Commercial Cloud Services (C2S) environment.

[Learn how to get started in C2S.](#)

- Cloud Manager has always enabled you to encrypt Cloud Volumes ONTAP data using the AWS Key Management Service (KMS). Starting with Cloud Volumes ONTAP 9.9.0, data on EBS disks and data tiered to S3 are encrypted if you select a customer-managed CMK. Previously, only EBS data would be encrypted.

Note that you'll need to provide the Cloud Volumes ONTAP IAM role with access to use the CMK.

[Learn more about setting up the AWS KMS with Cloud Volumes ONTAP.](#)

Enhancement available in Azure

You can now deploy Cloud Volumes ONTAP 9.8 in the Azure Department of Defense (DoD) Impact Level 6 (IL6).

Enhancements available in Google Cloud

- We've reduced the number of IP addresses that are required for Cloud Volumes ONTAP 9.8 and later in Google Cloud. By default, one less IP address is required (we unified the intercluster LIF with the node management LIF). You also have the option to skip the creation of the SVM management LIF when using the API, which would reduce the need for an additional IP address.

[Learn more about IP address requirements in Google Cloud.](#)

- When you deploy a Cloud Volumes ONTAP HA pair in Google Cloud, you can now choose shared VPCs for VPC-1, VPC-2, and VPC-3. Previously, only VPC-0 could be a shared VPC. This change is supported with Cloud Volumes ONTAP 9.8 and later.

[Learn more about Google Cloud networking requirements.](#)

Connector enhancements

- Cloud Manager now notifies Admin users through an email when a Connector isn't running.

Keeping your Connectors up and running helps to ensure the best management of Cloud Volumes ONTAP and other NetApp Cloud Services.

- Cloud Manager now displays a notification if you need to change the instance type for your Connector.

Changing the instance type ensures that you can use the new features and capabilities that you're currently missing. [Learn more about machine type changes.](#)

Cloud Sync enhancements

- Cloud Sync now supports sync relationships between ONTAP S3 Storage and SMB servers:
 - ONTAP S3 Storage to an SMB server
 - An SMB server to ONTAP S3 Storage

[View supported sync relationships.](#)

- Cloud Sync now enables you to unify a data broker group's configuration directly from the user interface.

We don't recommend changing the configuration on your own. You should consult with NetApp to understand when to change the configuration and how to change it.

[Learn more about defining a unified configuration.](#)

Cloud Tiering enhancements

- When tiering to Google Cloud Storage, you can apply a lifecycle rule so that the tiered data transitions from the Standard storage class to lower-cost Nearline, Coldline, or Archive storage after 30 days.
- Cloud Tiering now displays if you have any undiscovered on-prem ONTAP clusters so that you can add them to Cloud Manager to enable tiering or other services on those clusters.

[Learn how to discover these additional clusters.](#)

Active IQ enhancements

- When Active IQ displays the list of your on-prem clusters (based on your NSS account), you can click a button to [discover the cluster](#) and add it to the Cloud Manager Canvas. This makes it easier to manage all your storage systems from Cloud Manager.
- When Active IQ determines that one or more clusters require firmware updates, you can click a button to [download the Ansible playbook and upgrade the cluster firmware](#).
- A new [Cloud-Ready Workloads tab](#) provides a list of the workloads or volumes that we have identified as ideal to move to the cloud from your on-prem ONTAP clusters. Moving some of these volumes could reduce your costs and improve performance and resiliency.

See [What is Lift and Shift?](#)

ANF enhancements

- Now you can dynamically change the service level for a volume to meet workload needs and optimize your costs. The volume is moved to the other capacity pool with no impact to the volume. [Learn more](#).

Cloud Manager 3.9.3 update (16 Feb 2021)

Cloud Backup Service enhancements

- Now you can restore volumes to on-premises ONTAP systems from backup files that reside in Amazon S3, Azure Blob, and Google Cloud Storage.
- A new Restore Dashboard has been added that provides details about all the volumes and files you have restored.

The Dashboard is also the starting place to perform all volume and file restore operations. See [the Restore Dashboard](#) for details. In previous releases the restore volumes option was included in the Backup Dashboard.

- Cloud Backup is now supported on Cloud Volumes ONTAP HA systems in Google Cloud.

Cloud Manager 3.9.3 update (14 Feb 2021)

Cloud Compliance enhancements

- View and manage Azure Information Protection (AIP) labels in files you are scanning.

- After you integrate the AIP label functionality into Cloud Compliance, you can view the labels that are assigned to files, add labels to files, and change labels. See [how to integrate AIP labels](#) in your workspace.
- Assign labels individually to files, or use the Policies functionality to [add labels to all files that match the Policy criteria](#). With Policies, labels are updated continuously as Cloud Compliance finds matches in your files.
- Filter data in the Investigation page by AIP label to view all files that match the label.
- Send email alerts to Cloud Manager users (daily, weekly, or monthly) when any of your Policies return results so you can get notifications to protect your data.

Select this option when [creating or editing any Policy](#).

- View File Owner and Permission information when [viewing individual file details](#).

You can also use this criteria to further filter your data in the Investigation page.

- Delete files directly from Cloud Compliance.

You can [permanently remove files](#) that seem insecure or risky to leave in your storage system.

Cloud Manager 3.9.3 update (10 Feb 2021)

- [Cloud Tiering enhancements](#)
- [Cloud Sync enhancements](#)

Cloud Tiering enhancements

- Cloud Tiering now activates write-back prevention on a cluster when an aggregate is at >90% capacity (70% for ONTAP 9.6 and earlier). By preventing cold data write-backs on heavily utilized local tiers, Cloud Tiering preserves the local tier for active data.

When this happens, an indication appears in the Manage Aggregates table.

6 Aggregates ⓘ

Aggregate Name	Total Size	Used Size	Cold Data	Tiered Size
agg10fg1	0 B	0 B	0 B (Estimated)	0 B
agg11fg1	0 B	0 B	0 B (Estimated)	0 B
agg12fg2soft	3.52 GB	3.28 GB	2.96 GB	5.62 GB
	0 B	0 B	0 B (Estimated)	0 B
	2.64 GB	2.52 GB	1.76 GB (Estimated)	0 B
	2.64 GB	1.52 GB	1.06 GB (Estimated)	0 B

Notice: If the local tier is at >90% capacity, cold data is read directly from the cloud tier without being written back to the local tier. By preventing cold data write-backs on heavily utilized local tiers, FabricPool preserves the local tier for active data. To resolve that, you can change the value (0-99%), add disks to the aggregate, or move volumes.

- You can now add on-prem ONTAP clusters more easily from the Cloud Tiering service.

When you click **Add cluster** from the Cloud Tiering page, you're now sent directly to the **Add Working Environment** wizard.

- You can now filter the Timeline to show actions specific to the Cloud Tiering service.

The screenshot shows the 'Timeline' page with a 'Filters' section. A dropdown menu is open under the 'Service' filter, showing two options: 'Cloud Manager' (unchecked) and 'Cloud Tiering' (checked). Below the dropdown are 'Clear' and 'Apply' buttons. The main timeline table is partially visible at the bottom.

Cloud Sync enhancements

- We've simplified the process for syncing data to or from Cloud Volumes ONTAP. You can now select a Cloud Volumes ONTAP working environment and choose an option to sync data to or from this working environment.

The screenshot shows the NetApp Cloud Manager interface. On the left, there's a 'Canvas' section with a 'Add Working Environment' button and a diagram of a cloud icon labeled 'SINGLE' containing 'CloudVolumesONTAP' and 'Cloud Volumes ONTAP' with a '51 GiB Capacity'. An 'aws' logo is also present. On the right, a detailed view of the 'CloudVolumesONTAP' environment is shown. It includes a summary card with 'Compliance Off', 'Monitoring On', 'File Cache Off', and 'Sync On'. Below this are buttons for 'Enable', 'Sync data from this location', 'Sync data to this location', and 'View Dashboard'. A data sync summary shows '559.16TiB Data Synced'. There are also three circular icons with dots and a hand cursor icon.

- In the last release, we introduced a new Reports feature that provides information that you can use with the help of NetApp personnel to tune a data broker's configuration and improve performance. These reports are now supported with object storage.

Cloud Manager 3.9.3 (9 Feb 2021)

- [Monitoring enhancements](#)
- [Support improvements](#)

Monitoring enhancements

- The Monitoring service is now supported with Cloud Volumes ONTAP for Azure.
- The Monitoring service is also supported in Government regions in AWS and Azure.

The Monitoring service gives you complete visibility into your Cloud Volumes ONTAP infrastructure. Enable the service to monitor, troubleshoot, and optimize your Cloud Volumes ONTAP resources.

[Learn more about the Monitoring service.](#)

Support improvements

We've updated the Support Dashboard by enabling you to add your NetApp Support Site credentials, which registers you for support. You can also initiate a NetApp Support case directly from the dashboard. Just click the Help icon and then **Support**.

The screenshot shows the NetApp Support Dashboard. At the top, there's a navigation bar with 'Support Dashboard' (with a question mark icon), 'Account' (which is the active tab), and 'Connector'. Below the navigation, there's a section for 'Support Registration' with a green checkmark icon and the text 'Registered for Support'. It also shows the account serial number '960002998664...' and a blue square icon. To the right of this are buttons for 'View Active NSS Credentials', 'Learn More', and 'Add NSS Credentials'. The main content area is divided into several sections: 'Documentation' (with a book icon), 'Knowledge Base' (with a microphone icon), 'Communities' (with a people icon), 'Feedback' (with a smiley face icon), 'API' (with a gear icon), and 'Support' (with a person icon). Each section contains a brief description and a link to more information.

Cloud Manager 3.9.2 update (11 Jan 2021)

- [Cloud Compliance enhancements](#)
- [Cloud Backup enhancements](#)

Cloud Compliance enhancements

- Added support for scanning Microsoft OneDrive accounts.

Now you can add your corporate OneDrive accounts to Cloud Compliance in order to scan folders and files from all your OneDrive users. See [scanning OneDrive accounts](#) for details.

- The "Policies" feature now allows you can create your own custom Policies that provide results for searches specific to your organization.

In the last release, Cloud Compliance provided a set predefined Policy filters that all users could use. Now you can create your own Policies to return specific scan results in the Investigation page. See how to [create your own custom policies](#).

- Ability to scan backup files from on-premises ONTAP systems for free.

If you don't want Cloud Compliance to scan volumes directly on your on-prem ONTAP systems, a new Beta feature released this month allows you to run compliance scans on backup files created from your on-prem ONTAP volumes. So if you're already creating backups of your on-prem ONTAP volumes using [Cloud Backup](#), you can use this new feature to run compliance scans on those backup files - for **FREE**.

See how to [back up on-prem ONTAP volumes to object storage](#) and how you can [scan those backup files](#).

- Cloud Compliance can now find the personal data type "IP Address" in files. See the list of all [personal data types](#) that Cloud Compliance finds in scans.

Cloud Backup enhancements

You can restore individual files to additional destination working environments:

- Backup files in Azure Blob can be used to restore individual files to Cloud Volumes ONTAP systems installed on Azure, and to on-premises ONTAP systems.
- Backup files in Amazon S3 can be used to restore individual files to on-premises ONTAP systems (restoring files to Cloud Volumes ONTAP systems installed on AWS was already supported).

View the [backup and restore matrix](#) to see which working environments are supported for creating backups, restoring volumes, and restoring files.

Cloud Manager 3.9.2 (4 Jan 2021)

- [Cloud Volumes ONTAP enhancements](#)
- [Cloud Tiering enhancements](#)
- [General enhancements](#)

Cloud Volumes ONTAP enhancements

This release of Cloud Manager introduces the following enhancements for Cloud Volumes ONTAP.

Support for AWS Outposts

A few months ago, we announced that Cloud Volumes ONTAP had achieved the Amazon Web Services (AWS) Outposts Ready designation. Today, we're pleased to announce that we've validated Cloud Manager and Cloud Volumes ONTAP with AWS Outposts.

If you have an AWS Outpost, you can deploy Cloud Volumes ONTAP in that Outpost by selecting the Outpost VPC in the Working Environment wizard. The experience is the same as any other VPC that resides in AWS. Note that you will need to first deploy a Connector in your AWS Outpost.

There are a few limitations to point out:

- Only single node Cloud Volumes ONTAP systems are supported at this time
- The EC2 instances that you can use with Cloud Volumes ONTAP are limited to what's available in your Outpost
- Only General Purpose SSDs (gp2) are supported at this time

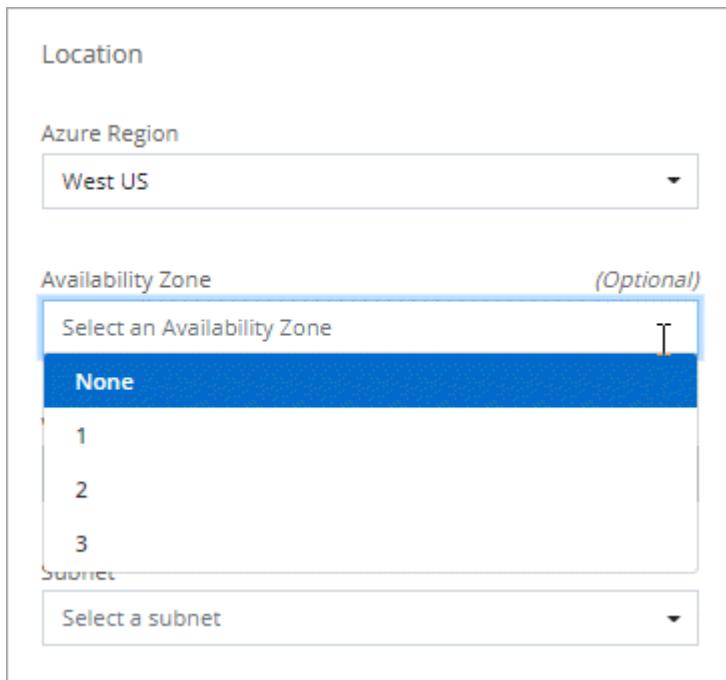
Support for Ultra SSD VNVVRAM in all supported Azure regions

Cloud Volumes ONTAP can now use an Ultra SSD as VNVVRAM when you use the E32s_v3 VM type with a single node system [in any supported Azure region](#).

VNVVRAM provides better write performance.

Ability to choose an Availability Zone in Azure

You can now choose the Availability Zone in which you'd like to deploy a single node Cloud Volumes ONTAP system. If you don't select an AZ, Cloud Manager will select one for you.



Support for bigger disks and new instances in GCP

- Cloud Volumes ONTAP now supports 64 TB disks in GCP.



The maximum system capacity with disks alone remains at 256 TB due to GCP limits.

- Cloud Volumes ONTAP now supports the following machine types:
 - n2-standard-4 with the Explore license and with BYOL
 - n2-standard-8 with the Standard license and with BYOL
 - n2-standard-32 with the Premium license and with BYOL

Cloud Tiering enhancements

- A new Cloud Performance Test gives you the ability to measure network latency and throughput performance from an ONTAP cluster to an object store before and after setting up data tiering.

Your cluster performance results			
Node: Node A	Last Check: 11/14/2020 08:54 am		 Recheck Performance
Operation	Size	Avg. Latency (ms)	Throughput
Write	4 KB	17687	337.5 MB
Read	8 KB	5348	14.33 MB
Read	8 KB	5568	27.98 MB
Read	32 KB	5668	108.5 MB
Read	256 KB	5809	706.8 MB

Notice: We recommend that you run this check when the cluster is under 50% CPU utilization.

- The Tiering Setup wizards were redesigned for ease of use.

Additional enhancements

- New Support Dashboard

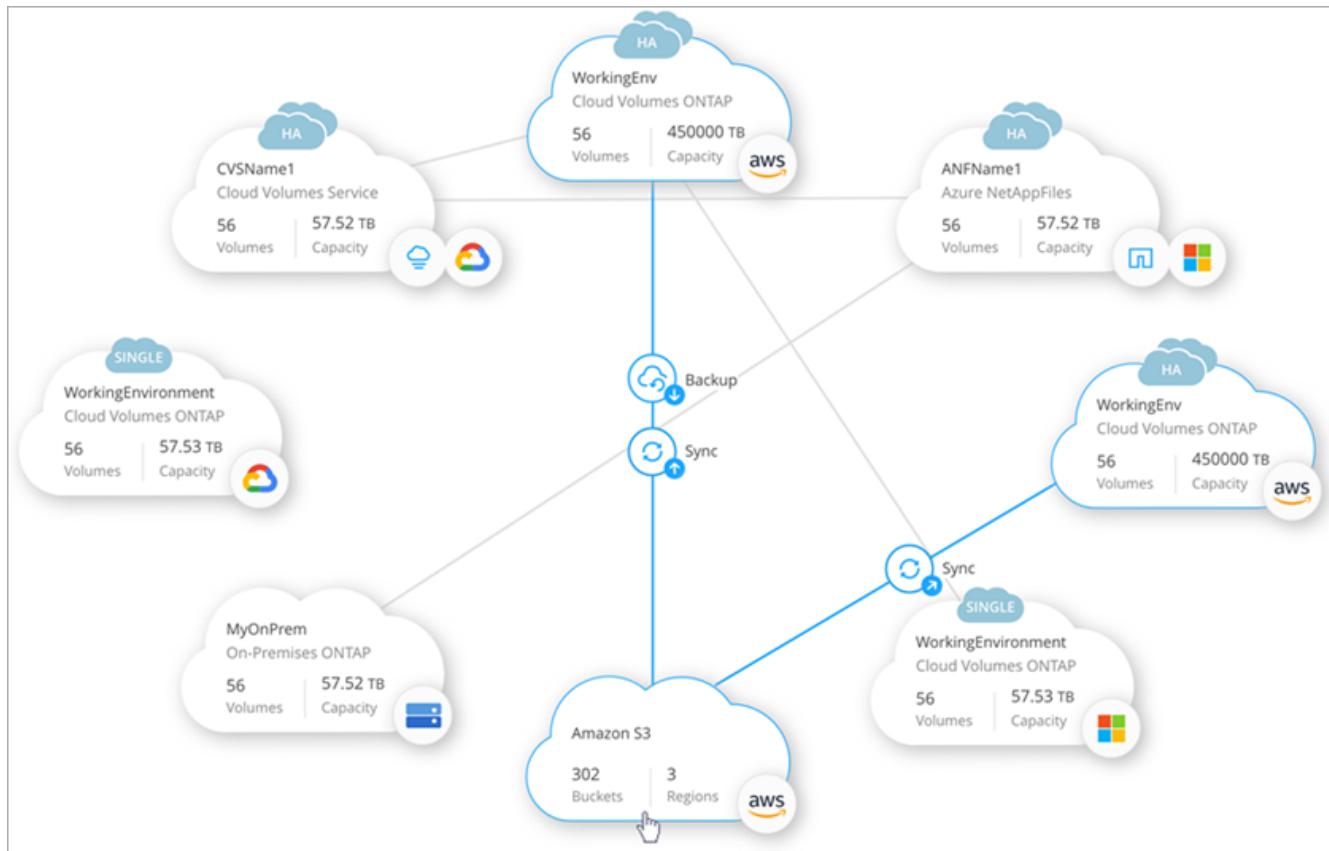
In the Help menu, a new Support Dashboard includes links to resources that can enable you to get help, submit feedback, and contact NetApp Support. You can also send and download AutoSupport messages from the **Connector AutoSupport** tab.

The screenshot shows the NetApp Support Dashboard. At the top, there's a navigation bar with 'Support Dashboard' (highlighted with a blue circle), 'Account' (selected tab), and 'Connector'. Below the navigation, there's a section for 'Account serial number' (960002998664...) and a green checkmark indicating 'Registered for Support' with a 'Support Registration' link. To the right are buttons for 'View Active NSS Credentials', 'Learn More', and 'Add NSS Credentials'. The main content area is divided into several sections: 'Documentation' (with a 'Cloud Manager Documentation' link), 'Knowledge Base' (with a 'Knowledge Base' link), 'Communities' (with a 'Communities' link), 'Feedback' (with a 'Feedback' link), 'API' (with an 'API Documentation' link), and 'Support' (with an 'Open an Issue' button). On the right side of the 'Support' section, there are links for 'Contact NetApp Support' and 'Call Us'.

- Visual representation between working environments

Cloud Manager makes it easier to view the relationships between the services enabled on your working environments.

For example, the following image shows an example of two working environments where data is backed up from Cloud Volumes ONTAP to Amazon S3, and where data is synced between Amazon S3 and two Cloud Volumes ONTAP systems.



Cloud Manager 3.9.1 (7 Dec 2020)

- General enhancements
- Cloud Volumes ONTAP AMI change
- Cloud Backup enhancements
- Cloud Compliance enhancements
- Cloud Tiering enhancements
- Cloud Sync enhancements

General enhancements

- We've renamed the **Working Environments** tab to **Canvas**.

This tab starts as a blank canvas and enables you to add your working environments by deploying, allocating, and discovering storage across your hybrid cloud.

The screenshot shows the Cloud Manager interface with the 'Canvas' tab selected. On the left, there's a 'Cloud Environment' section with icons for Azure, AWS, and On-Prem. Below this, a 'Working Environments' section lists several items:

- 1 On-Premises
1 TB Allocated Capacity
- 1 Cloud Volume ONTAP (High-Availability)
0.59 TB Allocated Capacity
- 1 Cloud Volumes ONTAP
0.29 TB Allocated Capacity
- 1 Azure NetApp Files
0.00 Byte Allocated Capacity
- 1 Amazon S3
89 Buckets

At the bottom left, it says 'Cloud Manager 3.9.0 Build: 1 Nov 1, 2020 09:22:53 am UTC Demo'. On the right, there's a 'Visual View' button.

- It's now easier to navigate between Cloud Manager and Spot.

A new **Storage Operations** section in Spot enables you to navigate directly to Cloud Manager. After you're done, you can get back to Spot from the **Compute** tab in Cloud Manager.

Cloud Volumes ONTAP AMI change

Starting with the 9.8 release, the Cloud Volumes ONTAP PAYGO AMI is no longer available in the AWS Marketplace. If you use the Cloud Manager API to deploy Cloud Volumes ONTAP PAYGO, you'll need to [subscribe to the Cloud Manager subscription in the AWS Marketplace](#) before deploying a 9.8 system.

Cloud Backup enhancements

- You now have the ability to restore individual files from a backup file.
 - If you need to restore a few files from a certain point in time, now you can just restore those files instead of having to restore the whole volume.
 - You can restore the files to a volume in the same working environment, or to a volume in a different working environment that's using the same cloud account.
 - This single file restore option relies on a new Cloud Restore instance that is deployed in your environment. [Go here for details about this new functionality.](#)
- You can configure Cloud Backup in a Google Cloud environment now while deploying a new Cloud Volumes ONTAP system. In the past you could only configure Cloud Backup on existing Cloud Volumes ONTAP systems.
- Now you can restore volumes that you had backed up from on-prem ONTAP systems to Cloud Volumes ONTAP systems deployed in AWS or Azure.

Cloud Compliance enhancements

- Ability to scan data directly from your on-premises ONTAP clusters

If you have discovered your on-prem clusters in Cloud Manager, now you can run Compliance scans directly on those volumes. No longer do you have to copy those volumes to a Cloud Volumes ONTAP system before you can run a Compliance scan.

- Ability to install Cloud Compliance in your on-premises location

If you plan to scan on-premises ONTAP cluster data, now you can install Cloud Compliance on-premises as well. It is still integrated in the Cloud Manager UI and it can still be used to scan other working environments, including cloud based volumes, buckets, and databases.

[See the prerequisites and installation steps here.](#)

- Ability to easily scan CIFS data protection volumes

In the past you have been able to scan NFS DP volumes. This release allows you to easily scan CIFS DP volumes directly within Cloud Compliance. [Learn how.](#)

- A new "Policies" feature provides a predefined selection of combination filters that return results in the Investigation page

Ten Policies are available with this release. For example, the "HIPAA – Stale data over 30 days" Policy identifies files that contain Health information that is over 30 days old. [See the full list of predefined policies.](#)

You can select Policies from a tab in the Compliance Dashboard and as a filter in the Investigation page.

- Cloud Compliance can now find the sensitive personal data type "Political Opinions Reference" in files. See the list of all [sensitive personal data types](#) that Cloud Compliance finds in scans.
- A new filter for "file size" is available from the Investigation page to refine your search results for files of a certain size

Note that the list of required endpoints for Cloud Compliance deployments has been revised based on cloud provider. [Review this list for AWS, Azure, and on-prem requirements.](#)

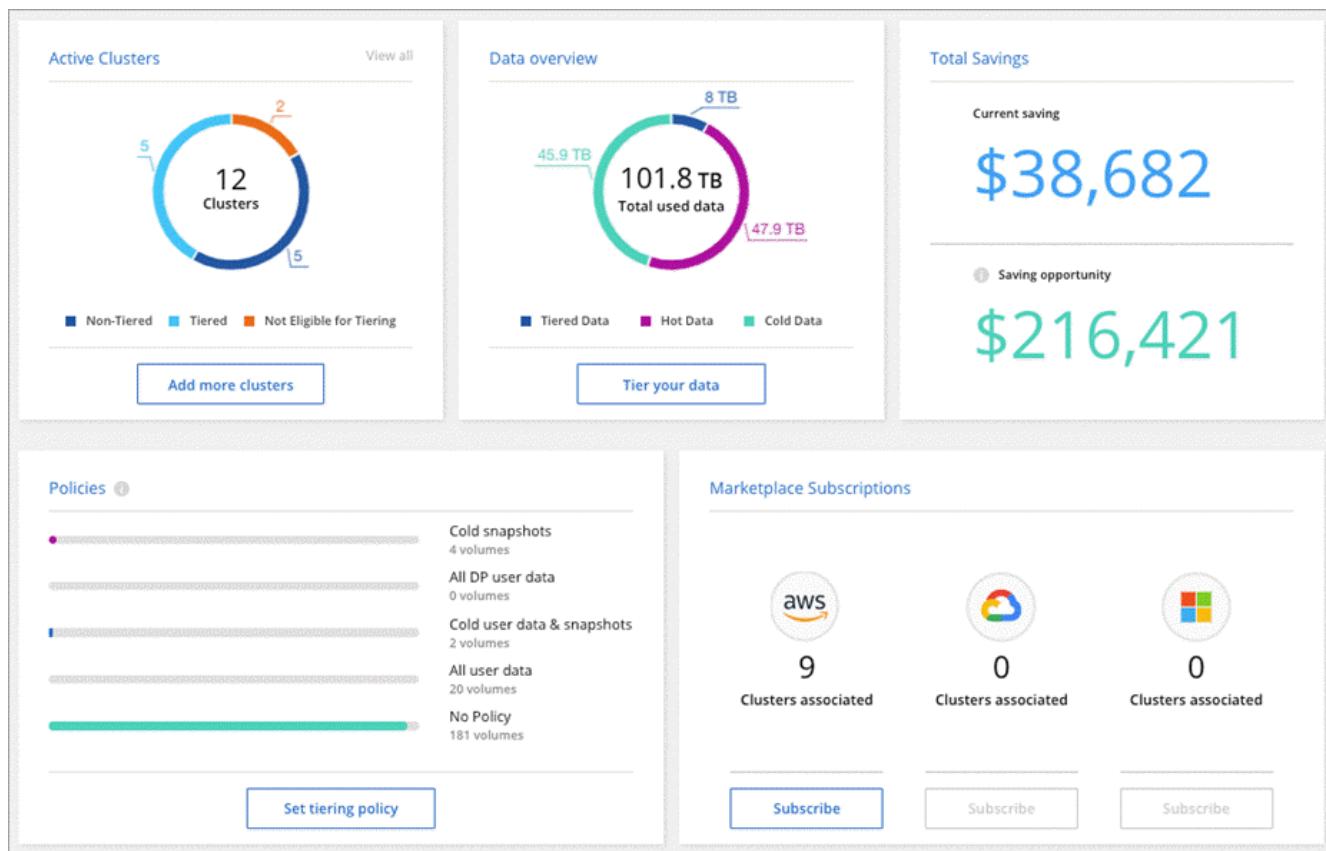
Cloud Tiering enhancements

- You can now change the tiering policy and minimum cooling days for multiple volumes at the same time.

The screenshot shows a 'Configure volumes' dialog overlaid on a list of volumes. The dialog has three numbered callouts: 1 points to the 'Configure volumes' button, 2 points to the 'Selected' checkbox, and 3 points to the 'Tiering Policy' dropdown. The list below shows three volumes: vol2, vol1, and vol3, each with its tier status set to 'Tiered Volume' and 'Cold snapshots'.

Volume Name	SVM Name	Volume Size	Used Size	Snapshot Used Size	Cold Data	Tier Status	Tiering Policy
vol2	svm_AFF1	200 GB	4.11 MB	424 KB	2.88 MB	70 %	Tiered Volume Cold snapshots
vol1	svm_AFF1	200 GB	1 MB	0 B	716.8 KB	70 %	Tiered Volume Cold snapshots
vol3	svm_AFF1	200 GB	1 MB	0 B	716.8 KB	70 %	Tiered Volume Cold snapshots

- Cloud Tiering now provides an aggregated view of data tiering from each of your on-premises clusters. This overview provides a clear picture of your environment and enables you to take proper actions. [Learn more about this page.](#)



Cloud Sync enhancements

- You can now manage data broker groups.

Grouping data brokers together can help improve the performance of sync relationships. Manage groups by adding a new data broker to a group, viewing information about data brokers, and more.

[Learn how to manage data brokers.](#)

- Cloud Sync now supports an ONTAP S3 Storage to ONTAP S3 Storage sync relationship.

[View the entire list of supported sources and targets.](#)

Cloud Manager 3.9 Update (18 Nov 2020)

Cloud Backup is now supported on Cloud Volumes ONTAP in Google Cloud. Click [here](#) for details.

Note: Only single-node systems are currently supported.

Cloud Volumes ONTAP 9.8 (16 Nov 2020)

Cloud Volumes ONTAP 9.8 is available in AWS, Azure, and Google Cloud Platform. This release includes support for [HA pairs in GCP](#).



The GCP service account associated with the Connector [needs the latest permissions](#) to deploy an HA pair in GCP.

[Learn what else is new in Cloud Volumes ONTAP 9.8.](#)

Cloud Manager 3.9 update (8 Nov 2020)

We released an enhancement to Cloud Manager 3.9.

Cloud Compliance enhancements

- Now you can create custom personal data identifiers from your databases. This gives you the full picture about where potentially sensitive data resides in **all** your files.

A feature we call "Data Fusion" allows you to scan your files to identify whether unique identifiers from your databases are found in those files—basically making your own list of "personal data" that is identified in Cloud Compliance scans.

[Learn how to create custom personal identifiers from your databases.](#)

- Added support for scanning MySQL database schemas.

Go to [scanning database schemas](#) for the list of all supported databases and for instructions.

Cloud Manager 3.9 (3 Nov 2020)

- [Azure Private Link for Cloud Volumes ONTAP](#)
- [Active IQ cluster insights](#)
- [Cloud Tiering enhancements](#)

Azure Private Link for Cloud Volumes ONTAP

By default, Cloud Manager now enables an Azure Private Link connection between Cloud Volumes ONTAP and its associated storage accounts. A Private Link secures connections between endpoints in Azure.

- [Learn more about Azure Private Links](#)
- [Learn more about using an Azure Private Link with Cloud Volumes ONTAP](#)

Active IQ cluster insights

Active IQ cluster insights are now available within Cloud Manager. This initial release provides the following functionality:

- Shows a list of your on-prem clusters based on your NetApp Support Site (NSS) credentials.
- Identifies which of those clusters have been discovered within Cloud Manager, and those that have not been discovered.
- Enables you to view unused Cloud Volumes ONTAP licenses.
- Identifies if any of your discovered ONTAP clusters need to have their shelf or disk firmware updated.

Go to [Monitoring ONTAP clusters](#) for details. This information is provided to Cloud Manager from the [Active IQ Digital Advisor](#).

Cloud Tiering enhancements

- When you set up data tiering from your volumes, Cloud Tiering now identifies the Snapshot used size for each volume. This information can help you decide which type of data to tier to the cloud.

The screenshot shows a table with columns: Volume Name, SVM Name, Volume Size, Used Size, Snapshot used size, Cold Data (Estimated), Tier Status, and TieringPolicy. The 'Snapshot used size' column is highlighted with a red box. For the first row, the values are: Volume 1, SVMNameB..., 462 TB, 100 TB, 50 TB, 70 TB | 70%, Available for Tiering, and Cold User Data. There is also a small edit icon next to the last column.

Volume Name	SVM Name	Volume Size	Used Size	Snapshot used size	Cold Data (Estimated)	Tier Status	TieringPolicy
Volume 1	SVMNameB...	462 TB	100 TB	50 TB	70 TB 70%	Available for Tiering	Cold User Data

- Cloud Tiering now enables inactive data reporting on HDD aggregates, if the cluster is running ONTAP 9.6 or later.

This enhancement makes it easier for Cloud Tiering to show you the potential savings from tiering cold data.

- Cloud Tiering now prompts you to change thick-provisioned volumes to thin-provisioned volumes, if that's required to enable data tiering on the volumes in an aggregate.

Cloud Manager transition to SaaS

We've introduced a software-as-a-service experience for Cloud Manager. This new experience makes it easier for you to use Cloud Manager and enables us to provide additional features to manage your hybrid cloud infrastructure.

The previous Cloud Manager experience

Cloud Manager software was previously comprised of a user interface and a management layer that sent requests to cloud providers. To get started, you would deploy Cloud Manager in your cloud network or on-premises network and then access the user interface that runs on that instance.

That experience has changed.

The new SaaS experience

The Cloud Manager interface is now accessible through a SaaS-based user interface that you log in to from NetApp Cloud Central. You no longer need to access a user interface from software that runs in your network.

In most cases, you need to deploy a *Connector* in your cloud or on-premises network. The Connector is software that's needed to manage Cloud Volumes ONTAP and other cloud data services. (The Connector is actually the same as the existing Cloud Manager software that you have installed.)

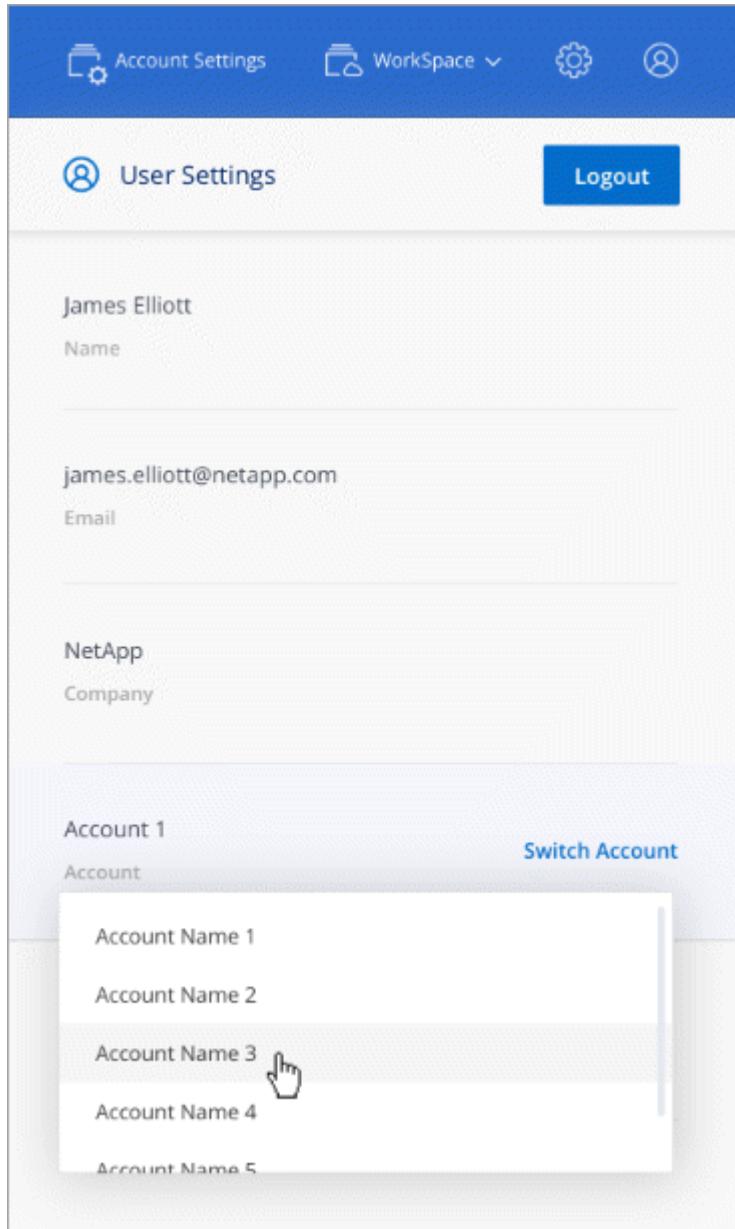
Benefits

This SaaS-based approach provides several benefits:

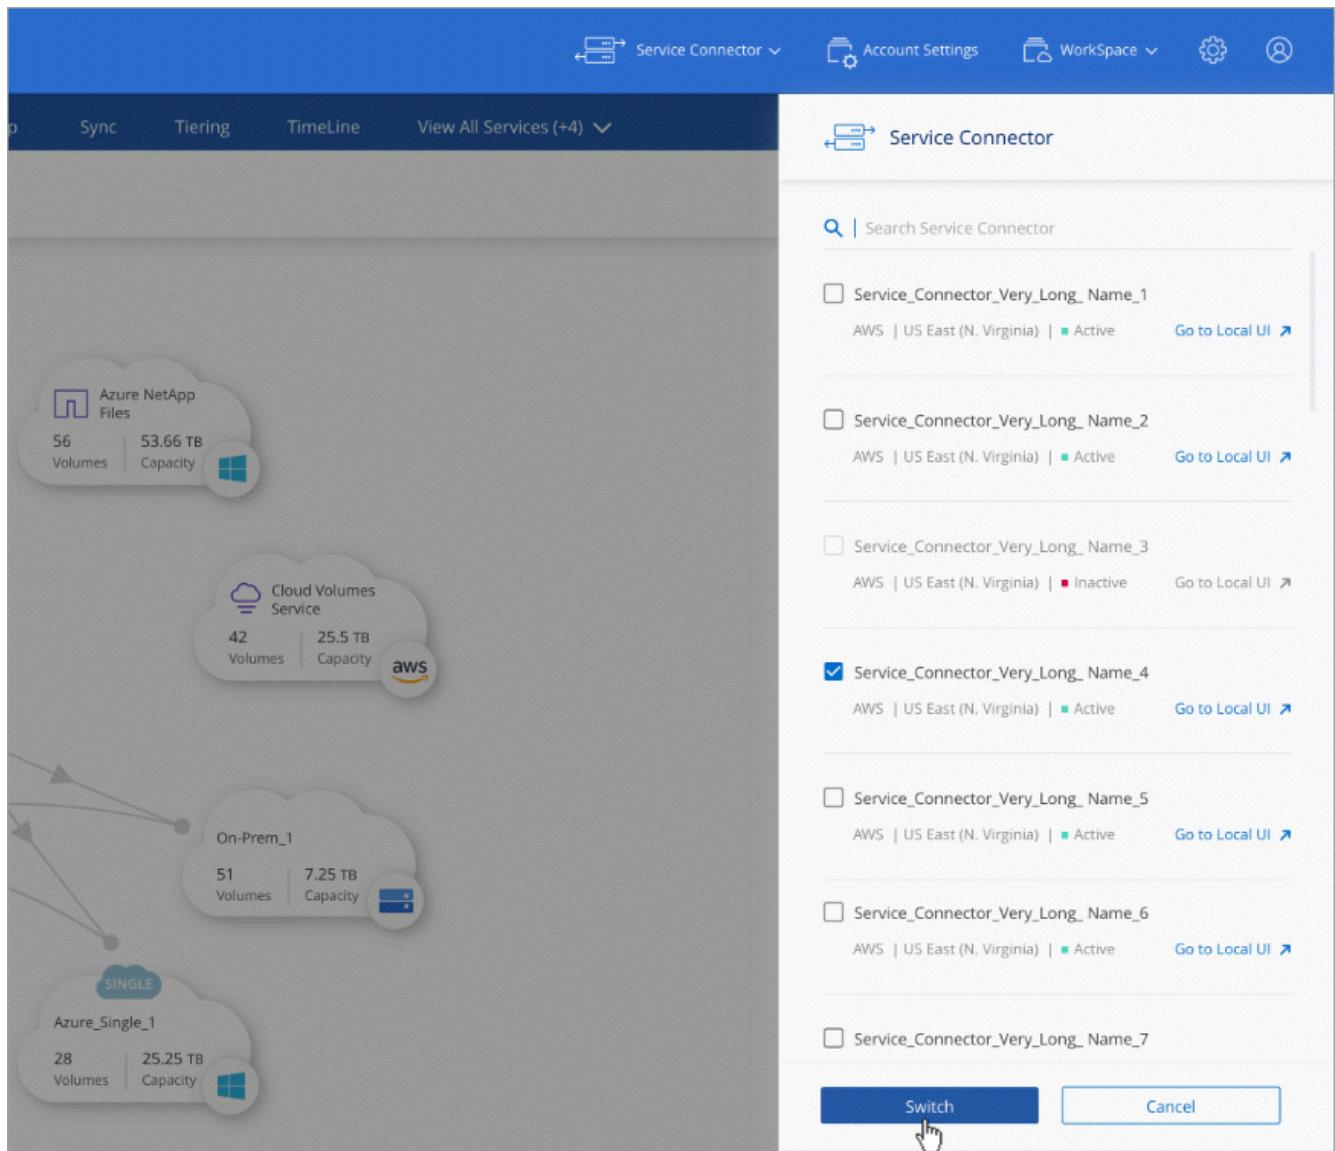
- It enables us to offer additional management capabilities for Azure NetApp Files and Cloud Volumes Service without needing to deploy software in your environment.
- You can easily switch between your Cloud Central accounts.

If a user is associated with multiple Cloud Central accounts, they can change to a different account at any

time from the User Settings menu. They can then see the Connectors and working environments that are associated with that account.



- You can easily switch between Connectors (what you know today as the Cloud Manager software) that are installed in different networks or different cloud providers.

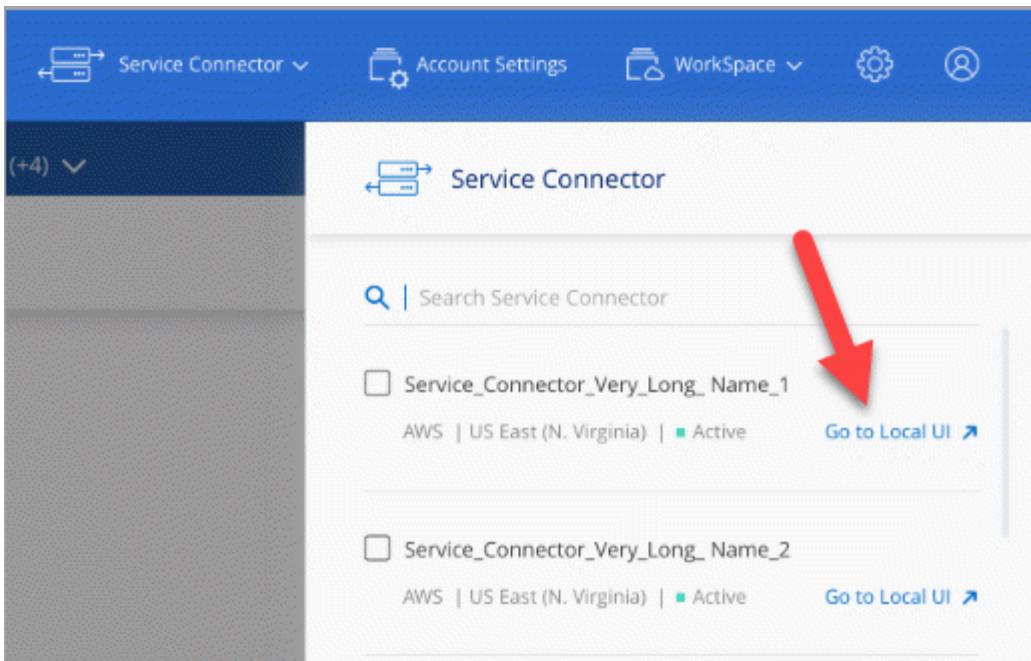


The local user interface

While you should perform almost all tasks from the SaaS user interface, a local user interface is still available on the Connector. This interface is needed for a few tasks that need to be performed from the Connector itself:

- Setting a proxy server
- Installing a patch
- Downloading AutoSupport messages

You can access the local user interface directly from the SaaS user interface:



Instance, VM, and machine type changes

To ensure that adequate resources are available for new and upcoming features in Cloud Manager, we've changed the minimum required instance, VM, and machine type as follows:

- AWS: t3.xlarge
- Azure: DS3 v2
- GCP: n1-standard-4

When you upgrade the machine type, you'll get access to features like a new Kubernetes experience, Global File Cache, Monitoring, and more.

These default sizes are the minimum supported [based on CPU and RAM requirements](#).

Cloud Manager will prompt you with instructions to change the machine type of the Connector.

Known issues

Known issues identify problems that might prevent you from using this release of the product successfully.

You can find known issues for Cloud Volumes ONTAP in the [Cloud Volumes ONTAP Release Notes](#) and for ONTAP software in general in the [ONTAP Release Notes](#).

Possible conflict with IP addresses in the 172 range

Cloud Manager deploys the Connector with two interfaces that have IP addresses in the 172.17.0.0/16 and 172.18.0.0/16 ranges.

If your network has a subnet configured with either of these ranges, then you might experience connectivity failures from Cloud Manager. For example, discovering on-prem ONTAP clusters in Cloud Manager might fail.

The workaround is to change the IP addresses of the Connector's interfaces. Contact NetApp Support for help.

Known limitations

Known limitations identify platforms, devices, or functions that are not supported by this release of the product, or that do not interoperate correctly with it. Review these limitations carefully.

SSL decryption isn't supported

Cloud Manager doesn't support firewall configurations that have SSL decryption enabled. If SSL decryption is enabled, error messages appear in Cloud Manager and the Connector instance displays as inactive.

For enhanced security, you have the option to [install an HTTPS certificate signed by a certificate authority \(CA\)](#).

Blank page when loading the local UI

If you load the local user interface for a Connector, the UI might fail to display sometimes, and you just get a blank page.

This issue is related to a caching problem. The workaround is to use an incognito or private web browser session.

Pay-as-you-go not available for Azure and Google Cloud partners

If you are a Microsoft Cloud Solution Provider (CSP) partner or a Google Cloud partner, NetApp pay-as-you-go subscriptions are not available. You must purchase a license and deploy NetApp cloud solutions with a BYOL license.

Pay-as-you-go subscriptions are not available for the following NetApp cloud services:

- Cloud Volumes ONTAP Explore, Standard, or Premium
- Cloud Tiering Service
- Cloud Backup Service
- Cloud Data Sense

Only one HA pair in an Azure resource group

You must use a dedicated resource group for each Cloud Volumes ONTAP HA pair that you deploy in Azure. Only one HA pair is supported in a resource group.

Cloud Manager experiences connection issues if you try to deploy a second Cloud Volumes ONTAP HA pair in an Azure resource group.

Connectors should remain running

A Connector should remain running at all times. It's important for the continued health and operation of the services that you enable.

For example, a Connector is a key component in the health and operation of Cloud Volumes ONTAP PAYGO systems. If a Connector is powered down, Cloud Volumes ONTAP PAYGO systems will shut down after losing communication with a Connector for longer than 14 days.

SaaS platform is disabled for Government regions

If you deploy a Connector in an AWS GovCloud region, an Azure Gov region, or an Azure DoD region, access to Cloud Manager is available only through a Connector's host IP address. Access to the SaaS platform is disabled for the entire account.

This means that only privileged users who can access the end-user internal VPC/VNet can use Cloud Manager's UI or API.

It also means that the following services aren't available from Cloud Manager:

- Cloud Data Sense
- Kubernetes
- Cloud Tiering
- Global File Cache

The SaaS platform is required to use these services.



The Monitoring service is supported and available in Government regions.

Cloud Manager doesn't support downgrades of Cloud Volumes ONTAP

Cloud Manager doesn't support downgrading Cloud Volumes ONTAP to a previous version. Contact NetApp technical support for help with downgrades.

Shared Linux hosts are not supported

The Connector isn't supported on a host that is shared with other applications. The host must be a dedicated host.

Cloud Manager doesn't support FlexGroup volumes

While Cloud Volumes ONTAP supports FlexGroup volumes, Cloud Manager does not. If you create a FlexGroup volume from System Manager or from the CLI, then you should set Cloud Manager's Capacity Management mode to Manual. Automatic mode might not work properly with FlexGroup volumes.

Cloud Manager doesn't support ONTAP S3

While Cloud Volumes ONTAP in Azure supports objects using S3, Cloud Manager doesn't provide any management capabilities for this feature. Using the CLI is the best practice to configure S3 client access from Cloud Volumes ONTAP. For details, refer to the [S3 Configuration Power Guide](#).

Important changes in Cloud Manager

This page highlights important changes in Cloud Manager that can help you use the service as we introduce new enhancements. You should continue to read the [What's new](#) page to learn about all new features and enhancements.

Cloud Volumes ONTAP AMI change

Starting with the 9.8 release, the Cloud Volumes ONTAP PAYGO AMI is no longer available in the AWS Marketplace. If you use the Cloud Manager API to deploy Cloud Volumes ONTAP PAYGO, you'll need to [subscribe to the Cloud Manager subscription in the AWS Marketplace](#) before deploying a 9.8 system.

SaaS changes

We have introduced a software-as-a-service experience for Cloud Manager. This new experience makes it easier for you to use Cloud Manager and enables us to provide additional features to manage your hybrid cloud infrastructure.

- [Cloud Manager transition to SaaS](#)
- [Learn how Cloud Manager works](#)

Machine type changes

To ensure that adequate resources are available for new and upcoming features in Cloud Manager, we've changed the minimum required instance, VM, and machine type as follows:

- AWS: t3.xlarge
- Azure: DS3 v2
- GCP: n1-standard-4

When you upgrade the machine type, you'll get access to features like a new Kubernetes experience, Global File Cache, Monitoring, and more.

These default sizes are the minimum supported [based on CPU and RAM requirements](#).

Cloud Manager will prompt you with instructions to change the machine type of the Connector.

Account settings

We introduced Cloud Central accounts to provide multi-tenancy, to help you organize users and resources in isolated workspaces, and to manage access to Connectors and subscriptions.

- [Learn about Cloud Central accounts: users, workspaces, Connectors, and subscriptions](#)
- [Learn how to get started with your account](#)
- [Learn how to manage your account after you set it up](#)

New permissions

Cloud Manager occasionally requires additional cloud provider permissions as we introduce new features and enhancements. This section identifies new permissions that are now required.

You can find the latest list of permissions on the [Cloud Manager policies page](#).

AWS

- Starting with the 3.9.9 release, the following permissions are required for the tagging service:

```
{  
    "Action": [  
        "ec2:CreateTags",  
        "ec2>DeleteTags",  
        "ec2:DescribeTags",  
        "tag:getResources",  
        "tag:getTagKeys",  
        "tag:getTagValues",  
        "tag:TagResources",  
        "tag:UntagResources"  
    ],  
    "Resource": "*",  
    "Effect": "Allow",  
    "Sid": "tagServicePolicy"  
},
```

- Starting with the 3.8.1 release, the following permissions are required to use Cloud Backup with Cloud Volumes ONTAP. [Learn more](#).

```
{
    "Sid": "backupPolicy",
    "Effect": "Allow",
    "Action": [
        "s3:DeleteBucket",
        "s3:GetLifecycleConfiguration",
        "s3:PutLifecycleConfiguration",
        "s3:PutBucketTagging",
        "s3>ListBucketVersions",
        "s3:GetObject",
        "s3>ListBucket",
        "s3>ListAllMyBuckets",
        "s3:GetBucketTagging",
        "s3:GetBucketLocation",
        "s3:GetBucketPolicyStatus",
        "s3:GetBucketPublicAccessBlock",
        "s3:GetBucketAcl",
        "s3:GetBucketPolicy",
        "s3:PutBucketPublicAccessBlock"
    ],
    "Resource": [
        "arn:aws:s3:::netapp-backup-*"
    ]
},
}
```

Azure

- Starting with the 3.9.10 release, you can use customer-managed encryption keys with Cloud Volumes ONTAP. A Connector requires these new permissions to set up an encryption key with a single node Cloud Volumes ONTAP system:

```
"Microsoft.Compute/diskEncryptionSets/write",
"Microsoft.KeyVault/vaults/deploy/action",
"Microsoft.Compute/diskEncryptionSets/delete"
```

- Starting with the 3.9.10 release, you can manage tags on your Azure resources using the Cloud Manager Tagging service. A Connector requires these new permissions for this service:

```
"Microsoft.Resources/tags/read",
"Microsoft.Resources/tags/write",
"Microsoft.Resources/tags/delete"
```

- Starting with the 3.9.8 release, Cloud Manager can remove Cloud Volumes ONTAP resources from a

resource group, in case of deployment failure or deletion. Be sure to provide these permissions to each set of Azure credentials that you've added to Cloud Manager:

```
"Microsoft.Network/privateEndpoints/delete",  
"Microsoft.Compute/availabilitySets/delete",
```

- Starting with the 3.9.7 release, Cloud Manager can now delete older cloud snapshots of root and boot disks that are created when a Cloud Volumes ONTAP system is deployed and every time its powered down. A Connector requires a new permission to delete Azure snapshots:

```
"Microsoft.Compute/snapshots/delete"
```

- To avoid Azure deployment failures, make sure that your Cloud Manager policy in Azure includes the following permission:

```
"Microsoft.Resources/deployments/operationStatuses/read"
```

- Starting with the 3.8.7 release, the following permission is required to encrypt Azure managed disks on single node Cloud Volumes ONTAP systems using external keys from another account. [Learn more](#).

```
"Microsoft.Compute/diskEncryptionSets/read"
```

- The following permissions are required to enable Global File Cache on Cloud Volumes ONTAP. [Learn more](#).

```
"Microsoft.Resources/deployments/operationStatuses/read",  
"Microsoft.Insights/Metrics/Read",  
"Microsoft.Compute/virtualMachines/extensions/write",  
"Microsoft.Compute/virtualMachines/extensions/read",  
"Microsoft.Compute/virtualMachines/extensions/delete",  
"Microsoft.Compute/virtualMachines/delete",  
"Microsoft.Network/networkInterfaces/delete",  
"Microsoft.Network/networkSecurityGroups/delete",  
"Microsoft.Resources/deployments/delete",
```

GCP

New permissions to deploy Cloud Data Sense in Google Cloud

Starting with the 3.9.10 release, the following permissions are required to deploy Cloud Data Sense in Google Cloud:

- `compute.subnetworks.use`
- `compute.subnetworks.useExternalIp`
- `compute.instances.addAccessConfig`

New permission for changing machine type

We recently discovered that the following permission is required for Cloud Volumes ONTAP machine type changes when switching between machine type families.

- `compute.instances.setMinCpuPlatform`

New permissions for HA pairs

Starting with the 3.9 release, the service account for a Connector requires additional permissions to deploy a Cloud Volumes ONTAP HA pair in GCP:

- `compute.addresses.list`
- `compute.backendServices.create`
- `compute.networks.updatePolicy`
- `compute.regionBackendServices.create`
- `compute.regionBackendServices.get`
- `compute.regionBackendServices.list`

New permissions for data tiering

Starting with the 3.9 release, additional permissions are required to set a service account on the Cloud Volumes ONTAP instance. This service account provides permissions for data tiering to a Google Cloud Storage bucket.

- `iam.serviceAccounts.actAs`
- `storage.objects.get`
- `storage.objects.list`

New permissions for Kubernetes management

Starting with the 3.8.8 release, the service account for a Connector requires additional permissions to discover and manage Kubernetes clusters running in Google Kubernetes Engine (GKE):

- `container.*`

New permissions for data tiering

Starting with the 3.8 release, the following permissions are now required to use a service account for data tiering. [Learn more about this change](#).

- `storage.buckets.update`
- `compute.instances.setServiceAccount`
- `iam.serviceAccounts.getIamPolicy`
- `iam.serviceAccounts.list`

New endpoints

The Connector requires outbound internet access to manage resources and processes within your public cloud environment. This section identifies new endpoints that are now required.

You can find the [full list of endpoints accessed from your web browser here](#) and the [full list of endpoints accessed by the Connector here](#).

- Users need to access Cloud Manager from a web browser by contacting the following endpoint:

<https://cloudmanager.netapp.com>

- Connectors require access to the following endpoint to obtain software images of container components for a Docker infrastructure:

<https://cloudmanagerinfraprod.azurecr.io>

Ensure that your firewall enables access to this endpoint from the Connector.

Get started with Cloud Manager

Learn about Cloud Manager

Cloud Manager enables IT experts and cloud architects to centrally manage their hybrid multi-cloud infrastructure using NetApp's cloud solutions.

Features

Cloud Manager is an enterprise-class, SaaS-based management platform that keeps you in control of your data no matter where it is.

- Set up and use [Cloud Volumes ONTAP](#) for efficient, multi-protocol data management across clouds.
- Set up and use file-storage services: [Azure NetApp Files](#), [Amazon FSx for ONTAP](#), [Cloud Volumes Service for AWS](#), and [Cloud Volumes Service for Google Cloud](#).
- Discover and manage your on-prem ONTAP clusters by creating volumes, backing up to the cloud, replicating data across your hybrid cloud, and tiering cold data to the cloud.
- Enable integrated cloud services and software like [Cloud Data Sense](#), [Cloud Insights](#), [Cloud Backup service](#), [Trident](#), and more.

[Learn more about Cloud Manager](#).

Supported object storage providers

Cloud Manager enables you to manage cloud storage and use cloud services in Amazon Web Services, Microsoft Azure, and Google Cloud.

Cost

Cloud Manager software is free of charge from NetApp.

For most tasks, Cloud Manager prompts you to deploy a Connector in your cloud network, which results in charges from your cloud provider for the compute instance and associated storage. You do have the option to run the Connector software on your premises.

How Cloud Manager works

Cloud Manager includes a SaaS-based interface that is integrated with NetApp Cloud Central, and Connectors that manage Cloud Volumes ONTAP and other cloud services.

Software-as-a-service

Cloud Manager is accessible through a [SaaS-based user interface](#) and APIs. This SaaS experience enables you to automatically access the latest features as they're released and to easily switch between your Cloud Central accounts and Connectors.

NetApp Cloud Central

[NetApp Cloud Central](#) provides a centralized location to access and manage [NetApp cloud services](#). With centralized user authentication, you can use the same set of credentials to access Cloud Manager and other cloud services like Cloud Insights.

When you log in to Cloud Manager for the first time, you're prompted to create a *Cloud Central account*. This account provides multi-tenancy and enables you to organize users and resources in isolated *workspaces*.

Connectors

In most cases, an Account Admin will need to deploy a *Connector* in your cloud or on-premises network. The Connector enables Cloud Manager to manage resources and processes within your public cloud environment.

A Connector should remain running at all times. It's important for the continued health and operation of the services that you enable.

For example, a Connector is a key component in the health and operation of Cloud Volumes ONTAP PAYGO systems. If a Connector is powered down, Cloud Volumes ONTAP PAYGO systems will shut down after losing communication with a Connector for longer than 14 days.

[Learn more about when Connectors are required and how they work.](#)

Getting started checklist

This checklist describes what's needed to get up and running with Cloud Manager.

A NetApp Cloud Central login

You need to sign up to [NetApp Cloud Central](#) so you can access Cloud Manager and other cloud services.

Network access from a web browser to several endpoints

Users need to access Cloud Manager from a web browser. The machine running the web browser must have connections to the following endpoints.

Endpoints	Purpose
https://cloudmanager.cloud.netapp.com	To connect you to the Cloud Manager SaaS interface.
https://api.services.cloud.netapp.com	To contact Cloud Central APIs.
https://auth0.com https://cdn.auth0.com https://netapp-cloud-account.auth0.com https://services.cloud.netapp.com	Your web browser connects to these endpoints for centralized user authentication through NetApp Cloud Central.
https://widget.intercom.io	For in-product chat that enables you to talk to NetApp cloud experts.

A network location for a Connector

After logging in to Cloud Manager, an Account Admin will need to deploy a *Connector* in a cloud provider or in your on-premises network. The Connector enables Cloud Manager to manage resources and processes within your public cloud environment. A Connector isn't required for Azure NetApp Files, Cloud Volumes Service, or Cloud Sync, but it is required for all other services and features in Cloud Manager. [Learn more about Connectors and how they work.](#)

- The network location must have an outbound internet connection.

The Connector requires outbound internet access to manage resources and processes within your public cloud environment. Outbound internet access is also required if you want to manually install the Connector on a Linux host or access the local UI running on the Connector.

[View the specific endpoints that the Connector contacts.](#)

- There's no incoming traffic to the Connector, unless you initiate it.

HTTP (80) and HTTPS (443) provide access to the local UI, which you'll use in rare circumstances. SSH (22) is only needed if you need to connect to the host for troubleshooting.

Cloud provider permissions

You need an account that has permissions to deploy the Connector in your cloud provider directly from Cloud Manager.



There are alternative ways to create a Connector. You also have the option to create a Connector from the [AWS Marketplace](#), [Azure Marketplace](#), or to [manually install the software](#).

Location	High-level steps	Detailed steps
AWS	<ol style="list-style-type: none">1. Use a JSON file that includes the required permissions to create an IAM policy in AWS.2. Attach the policy to the IAM user who will create the Connector from Cloud Manager.3. When you create the Connector, provide Cloud Manager with the AWS access key and secret key for the IAM user that has the required permissions.	Click here to view detailed steps.
Azure	<ol style="list-style-type: none">1. Use a JSON file that includes the required permissions to create a custom role in Azure.2. Assign the role to the user who will create the Connector from Cloud Manager.3. When you create the Connector, log in with the Microsoft account that has the required permissions (the login prompt that is owned and hosted by Microsoft).	Click here to view detailed steps.
Google Cloud	<ol style="list-style-type: none">1. Use a YAML file that includes the required permissions to create a custom role in Google Cloud.2. Attach that role to the user who will create the Connector from Cloud Manager.3. If you plan to use Cloud Volumes ONTAP, set up a service account that has the required permissions.4. Enable Google Cloud APIs.5. When you create the Connector, log in with the Google account that has the required permissions (the login prompt is owned and hosted by Google).	Click here to view detailed steps.

Networking for individual services

Now that your setup is complete, you're ready to start using the services available from Cloud Manager. Note that each service has its own networking requirements. Refer to the following pages for more details.

- Cloud Volumes ONTAP for AWS
- Cloud Volumes ONTAP for Azure
- Cloud Volumes ONTAP for GCP
- Data replication between ONTAP systems
- Deploying Cloud Data Sense
 - Cloud Data Sense for Cloud Volumes ONTAP and on-premises ONTAP
 - Cloud Data Sense for Azure NetApp Files
 - Cloud Data Sense for Amazon FSx for ONTAP
 - Cloud Data Sense for Amazon S3
 - Cloud Data Sense for non-NetApp NFS or CIFS file shares
- On-prem ONTAP clusters
- Cloud Tiering
 - Data tiering from ONTAP clusters to Amazon S3
 - Data tiering from ONTAP clusters to Azure Blob storage
 - Data tiering from ONTAP clusters to Google Cloud Storage
 - Data tiering from ONTAP clusters to StorageGRID
 - Data tiering from ONTAP clusters to generic S3 object storage
- Cloud Backup
 - Data backup from ONTAP clusters to Amazon S3
 - Data backup from ONTAP clusters to Azure Blob storage
 - Data backup from ONTAP clusters to Google Cloud Storage
 - Data backup from ONTAP clusters to StorageGRID

Signing up to NetApp Cloud Central

Sign up to NetApp Cloud Central so you can access NetApp's cloud services.



You can use single sign-on to log in using credentials from your corporate directory (federated identity). To learn more, go to the [Cloud Central Help Center](#) and then click **Cloud Central sign-in options**.

Steps

1. Open a web browser and go to [NetApp Cloud Central](#).
2. Click **Sign Up**.
3. Fill out the form and click **Sign Up**.

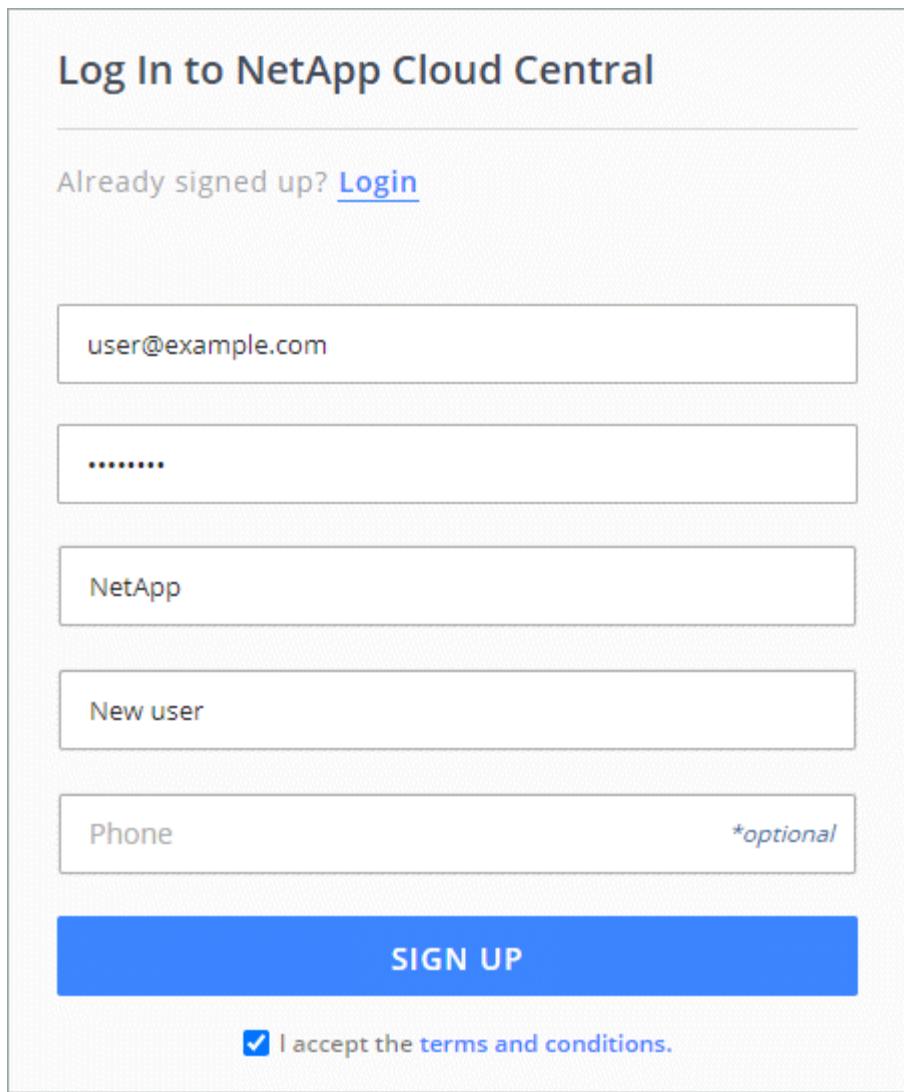
Log In to NetApp Cloud Central

Already signed up? [Login](#)

*optional

SIGN UP

I accept the [terms and conditions](#).



4. Wait for an email from NetApp Cloud Central.
5. Click the link in the email to verify your email address.

Result

You now have an active Cloud Central user login.

Logging in to Cloud Manager

The Cloud Manager interface is accessible through a SaaS-based user interface by going to <https://cloudmanager.netapp.com>.



You can use single sign-on to log in using credentials from your corporate directory (federated identity). To learn more, go to the [Cloud Central Help Center](#) and then click **Cloud Central sign-in options**.

Steps

1. Open a web browser and go to <https://cloudmanager.netapp.com>.
2. Log in using your NetApp Cloud Central credentials.



[Continue to Cloud Manager](#)

Log In to NetApp Cloud Central

Don't have an account yet? [Sign Up](#)

Email

Password

[Forgot your password?](#)

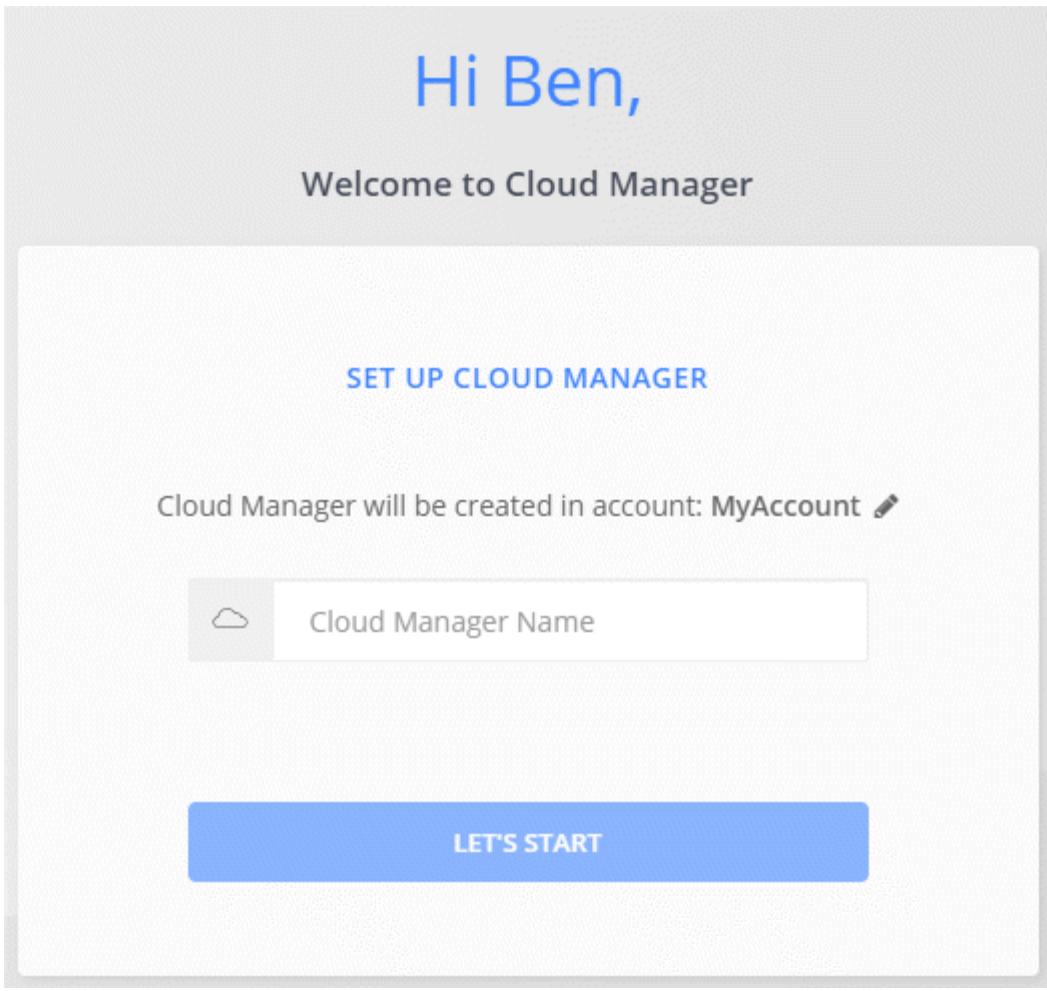
Set up a Cloud Central account

Account settings: users, workspaces, Connectors, and subscriptions

A *Cloud Central account* provides multi-tenancy and enables you to organize users and resources in isolated workspaces from within Cloud Manager.

For example, multiple users can deploy and manage Cloud Volumes ONTAP systems in isolated environments called *workspaces*. These workspaces are invisible to other users, unless they are shared.

When you first access Cloud Manager, you're prompted to select or create a Cloud Central account:



Account Admins can then modify the settings for this account by managing users (members), workspaces, Connectors, and subscriptions:

A screenshot of the "Manage Account" interface for the account "activeIQ". The top navigation bar includes tabs for Overview, Members, Workspaces (which is underlined in blue), Service Connector, Subscriptions, and a close button (X). The main content area is titled "Manage the Account's Workspaces" and features a button "+ Add New Workspace". A single workspace entry named "Workspace-1" is listed with a delete (trash) and edit icon next to it.

For step-by-step instructions, see [Setting up the Cloud Central account](#).

Account Settings

The Manage Account widget in Cloud Manager enables Account Admins to manage a Cloud Central account. If you just created your account, then you'll start from scratch. But if you've already set up an account, then you'll see *all* the users, workspaces, Connectors, and subscriptions that are associated with the account.

Overview

The Overview page shows the Account Name and the Account ID. You may need to provide your Account ID when registering some services. This page also includes some Cloud Manager configuration options.

Members

The members are NetApp Cloud Central users that you associate with your Cloud Central account. Associating a user with an account and one or more workspaces in that account enables those users to create and manage working environments in Cloud Manager.

When you associate a user, you assign them a role:

- *Account Admin*: Can perform any action in Cloud Manager.
- *Workspace Admin*: Can create and manage resources in the assigned workspace.
- *Compliance Viewer*: Can only view Cloud Data Sense compliance information and generate reports for systems that they have permission to access.
- *SnapCenter Admin*: Can use the SnapCenter Service to create application consistent backups and restore data using those backups. *This service is currently in Beta.*

Workspaces

In Cloud Manager, a workspace isolates any number of *working environments* from other working environments. Workspace Admins can't access the working environments in a workspace unless the Account Admin associates the admin with that workspace.

A working environment represents a storage system:

- A single-node Cloud Volumes ONTAP system or an HA pair
- An on-premises ONTAP cluster in your network
- An ONTAP cluster in a NetApp Private Storage configuration

Connectors

A Connector enables Cloud Manager to manage resources and processes within your public cloud environment. The Connector runs on a virtual machine instance that you deploy in your cloud provider, or on an on-prem host that you configured.

You can use a Connector with more than one NetApp cloud data service. For example, if you already have a Connector for Cloud Manager, you can select it when you set up the Cloud Tiering service.

Subscriptions

These are the NetApp subscriptions associated with the selected account.

When you subscribe to Cloud Manager from a cloud provider's marketplace, you're redirected to Cloud Central where you need to save your subscription and associate it with specific accounts.

After you've subscribed, each subscription is available from the Manage Account widget. You'll only see the subscriptions that are associated with the account that you're currently viewing.

You have the option to rename a subscription and to disassociate the subscription from one or more accounts.

For example, let's say that you have two accounts and each is billed through separate subscriptions. You might disassociate a subscription from one of the accounts so the users in that account don't accidentally choose the wrong subscription when creating a Cloud Volume ONTAP working environment.

Examples

The following examples depict how you might set up your accounts.

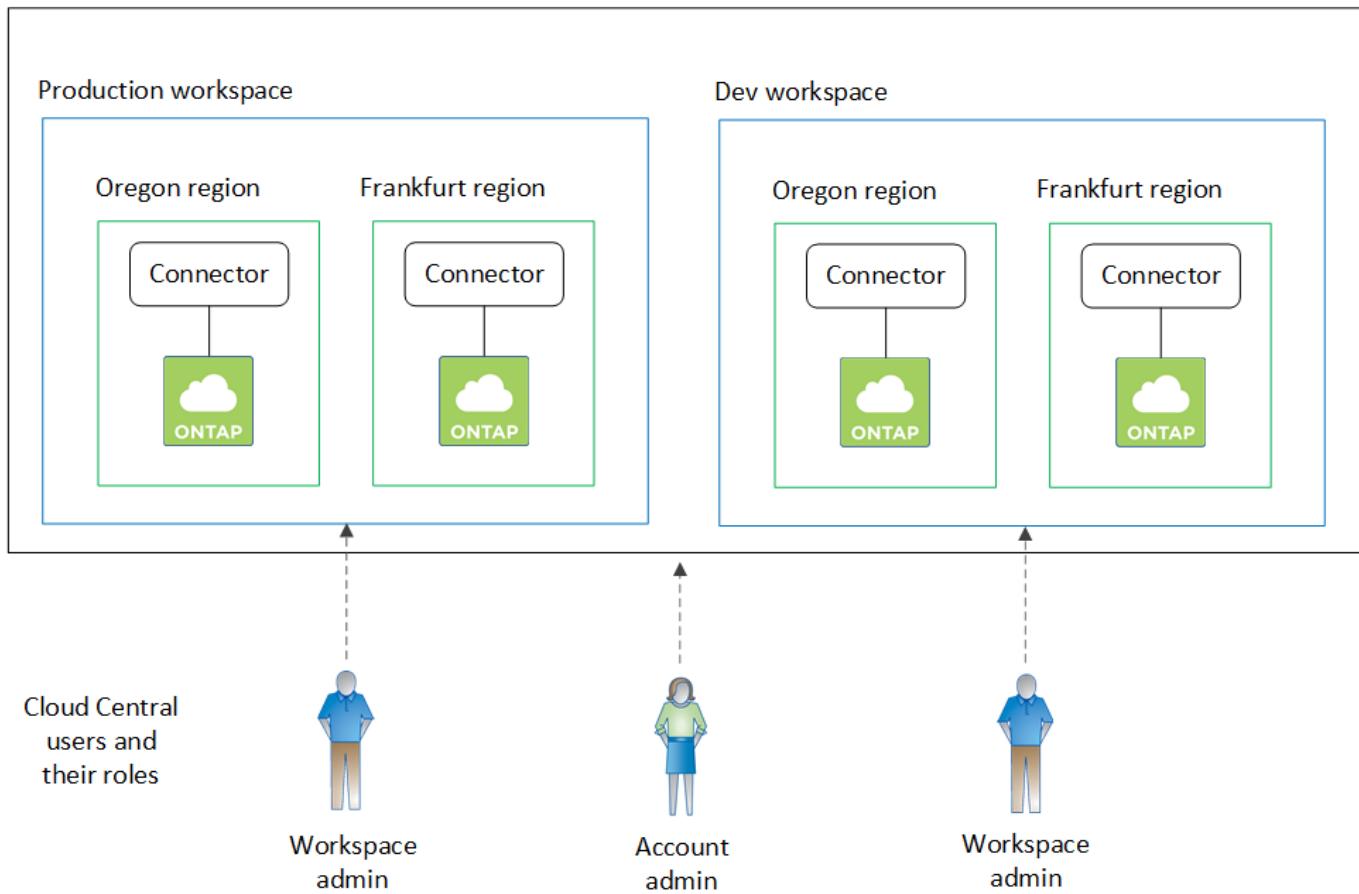


In both example images that follow, the Connector and the Cloud Volumes ONTAP systems don't actually reside *in* the NetApp Cloud Central account—they're running in a cloud provider. This is a conceptual representation of the relationship between each component.

Example 1

The following example shows an account that uses two workspaces to create isolated environments. The first workspace is for a production environment and the second is for a dev environment.

Account

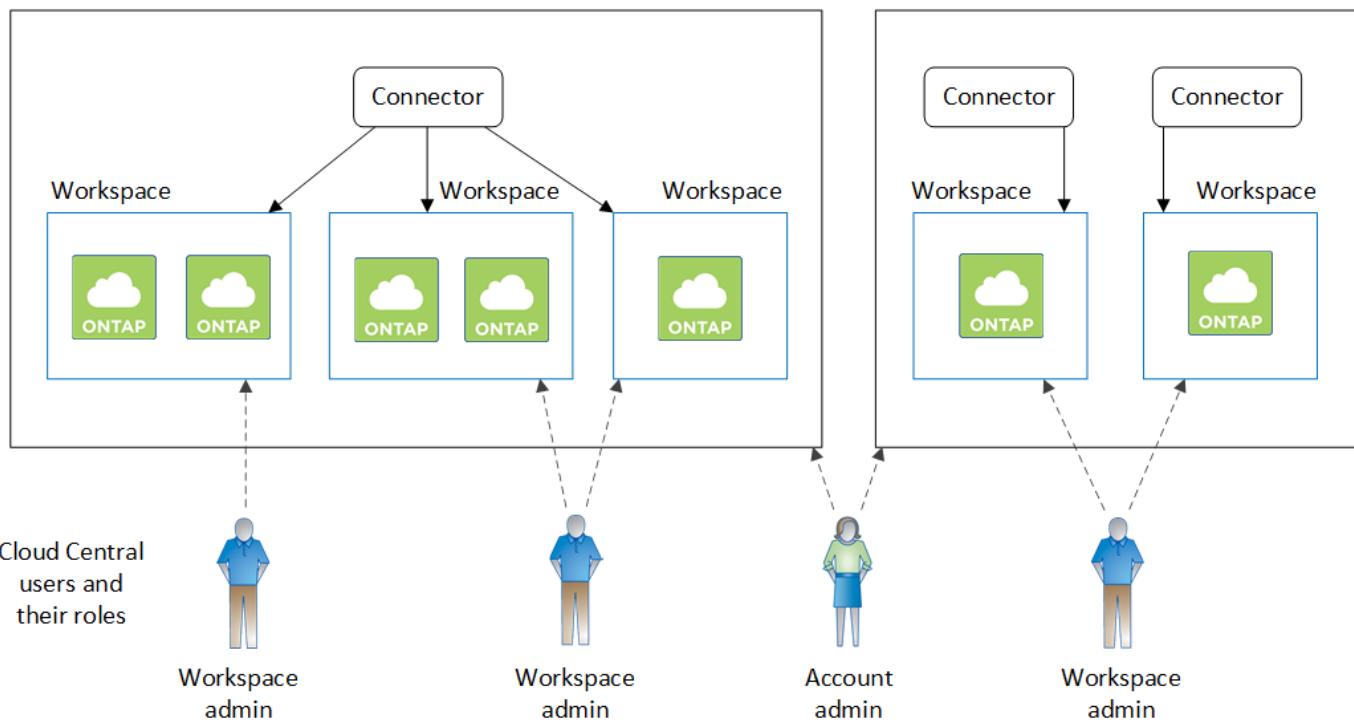


Example 2

Here's another example that shows the highest level of multi-tenancy by using two separate Cloud Central accounts. For example, a service provider might use Cloud Manager in one account to provide services for their customers, while using another account to provide disaster recovery for one of their business units.

Note that account 2 includes two separate Connectors. This might happen if you have systems in separate regions or in separate cloud providers.

Account #1



Setting up workspaces and users in the Cloud Central account

When you log in to Cloud Manager for the first time, you're prompted to create a *NetApp Cloud Central account*. This account provides multi-tenancy and enables you to organize users and resources in isolated *workspaces*.

[Learn more about how Cloud Central accounts work.](#)

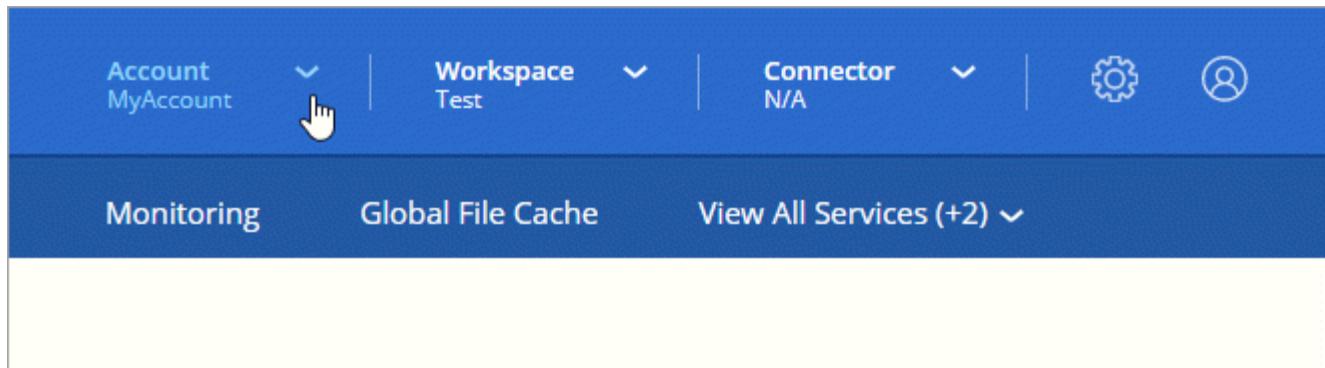
Set up your Cloud Central account so users can access Cloud Manager and access the working environments in a workspace. Just add a single user or add multiple users and workspaces.

Adding workspaces

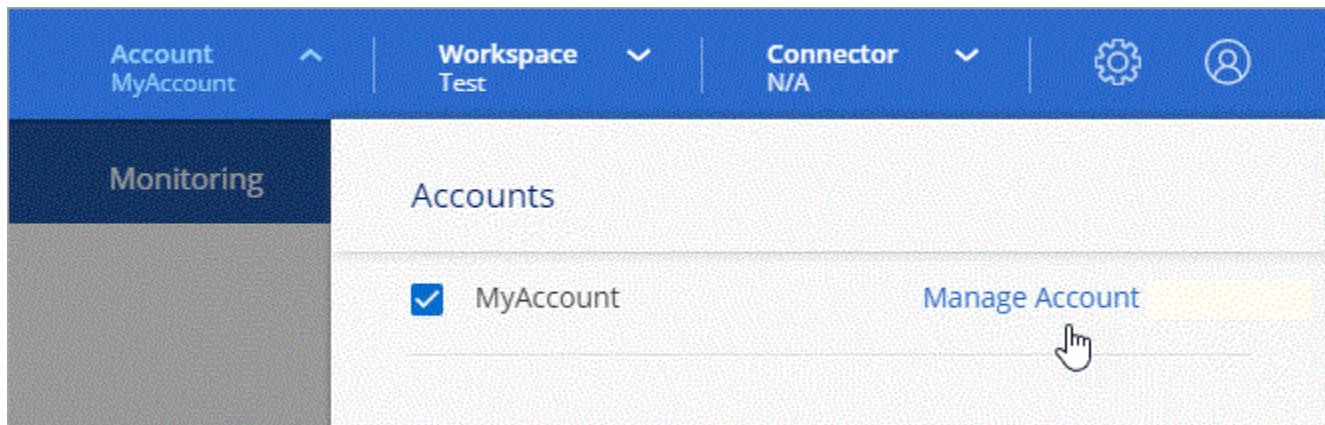
In Cloud Manager, workspaces enable you to isolate a set of working environments from other working environments and from other users. For example, you can create two workspaces and associate separate users with each workspace.

Steps

1. From the top of [Cloud Manager](#), click the **Account** drop-down.



2. Click **Manage Account** next to the currently selected account.



3. Click **Workspaces**.
4. Click **Add New Workspace**.
5. Enter a name for the workspace and click **Add**.

After you finish

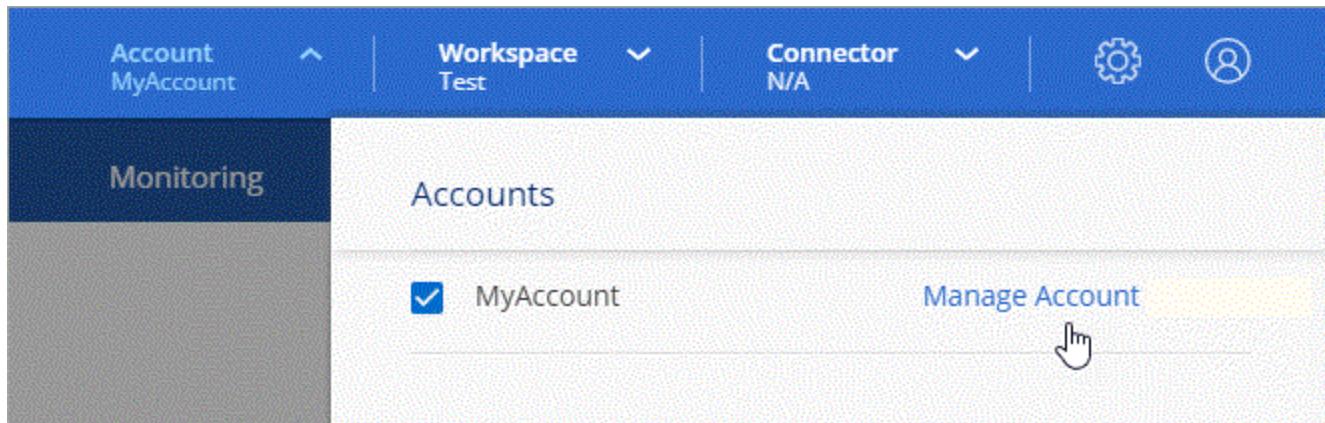
If a Workspace Admin needs access to this workspace, then you'll need to associate the user. You'll also need to associate Connectors with the workspace so Workspace Admins can use those Connectors.

Adding users

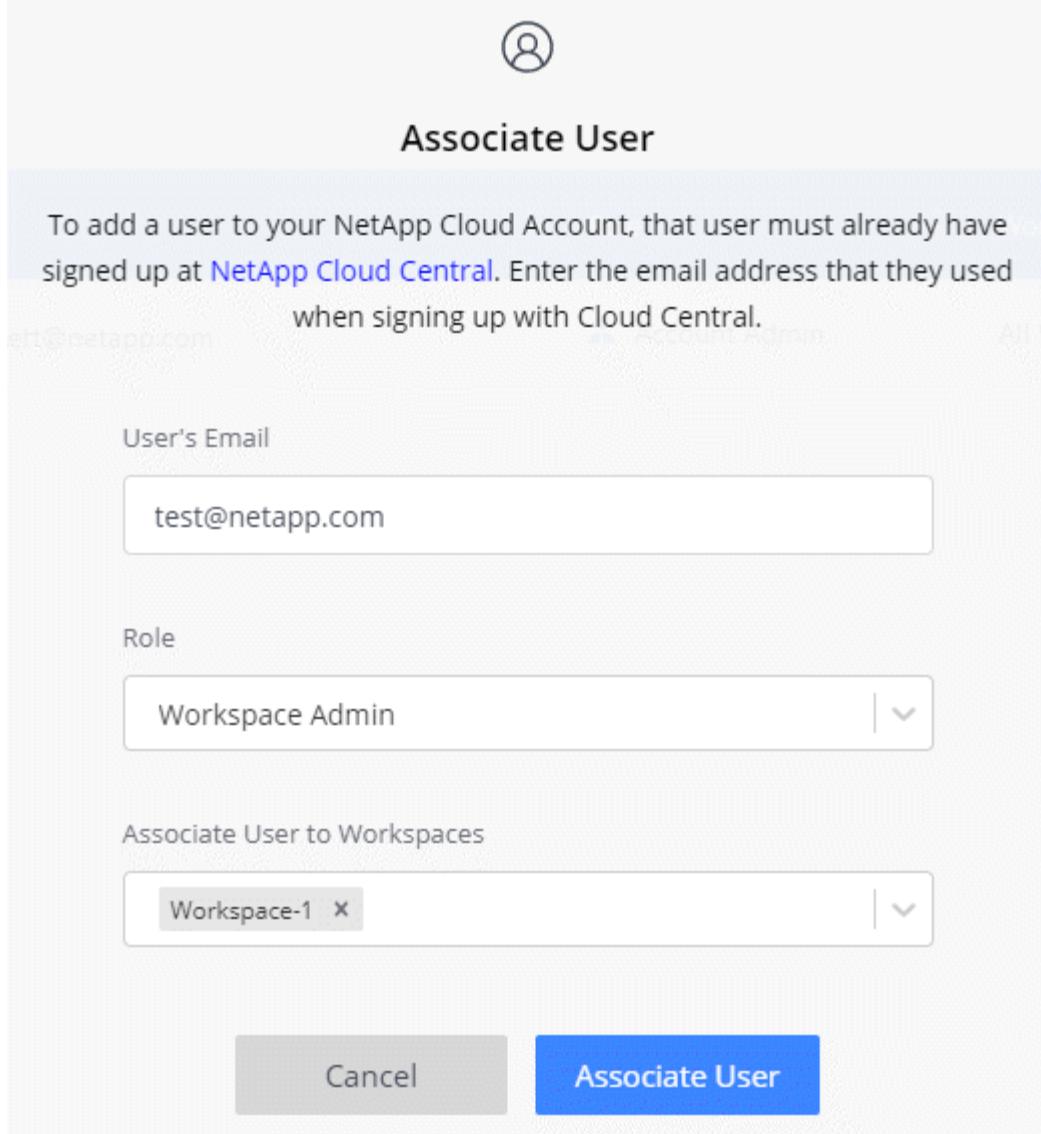
Associate Cloud Central users with the Cloud Central account so those users can create and manage working environments in Cloud Manager.

Steps

1. If the user hasn't already done so, ask the user to go to [NetApp Cloud Central](#) and sign up.
2. From the top of [Cloud Manager](#), click the **Account** drop-down and click **Manage Account**.



3. From the Members tab, click **Associate User**.
4. Enter the user's email address and select a role for the user:
 - **Account Admin:** Can perform any action in Cloud Manager.
 - **Workspace Admin:** Can create and manage resources in assigned workspaces.
 - **Compliance Viewer:** Can only view Cloud Data Sense governance and compliance information and generate reports for workspaces that they have permission to access.
 - **SnapCenter Admin:** Can use the SnapCenter Service to create application consistent backups and restore data using those backups. This service is currently in Beta.
5. If you selected an account other than Account Admin, select one or more workspaces to associate with that user.



The screenshot shows a 'Associate User' dialog box. At the top is a user icon. Below it is the title 'Associate User'. A descriptive text explains that the user must already be signed up at NetApp Cloud Central and provides a placeholder for their email address. The 'User's Email' field contains 'test@netapp.com'. The 'Role' dropdown is set to 'Workspace Admin'. Under 'Associate User to Workspaces', 'Workspace-1' is listed. At the bottom are 'Cancel' and 'Associate User' buttons.

To add a user to your NetApp Cloud Account, that user must already have signed up at [NetApp Cloud Central](#). Enter the email address that they used when signing up with Cloud Central.

User's Email

test@netapp.com

Role

Workspace Admin

Associate User to Workspaces

Workspace-1 X

Cancel Associate User

6. Click **Associate**.

Result

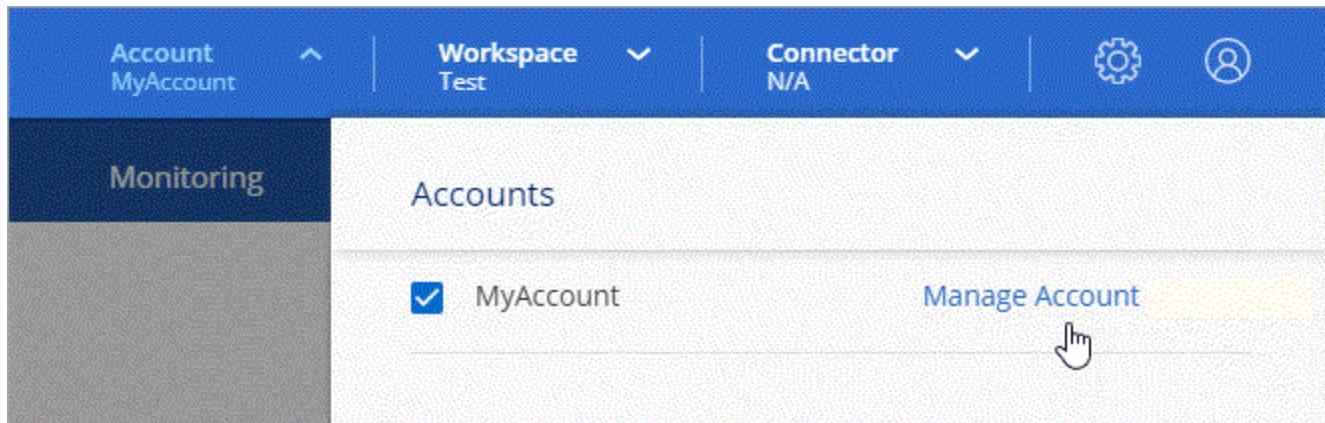
The user should receive an email from NetApp Cloud Central titled "Account Association." The email includes the information needed to access Cloud Manager.

Associating Workspace Admins with workspaces

You can associate Workspace Admins with additional workspaces at any time. Associating the user enables them to create and view the working environments in that workspace.

Steps

1. From the top of Cloud Manager, click the **Account** drop-down and click **Manage Account**.



2. From the Members tab, click the action menu in the row that corresponds to the user.

5 Members						Q
Type	Name	Email	Role	Workspace	...	
👤	Ben	[REDACTED]	⭐ Account Admin	All Workspaces	...	
👤	Tom	[REDACTED]	⭐ Account Admin	All Workspaces	...	
👤	Ben	[REDACTED]	Workspace Admin	Newone		

3. Click **Manage Workspaces**.

4. Select one or more workspaces and click **Apply**.

Result

The user can now access those workspaces from Cloud Manager, as long as the Connector was also associated with the workspaces.

Associating Connectors with workspaces

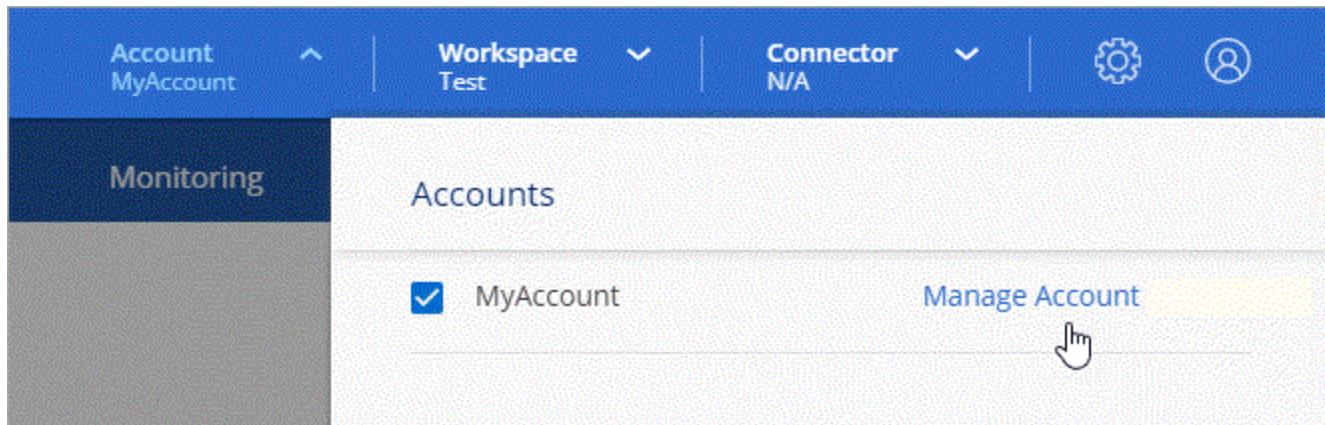
You need to associate a Connector with workspaces so Workspace Admins can use those Connectors to create Cloud Volumes ONTAP systems.

If you only have Account Admins, then associating the Connector with workspaces isn't required. Account Admins have the ability to access all workspaces in Cloud Manager by default.

[Learn more about users, workspaces, and Connectors.](#)

Steps

1. From the top of Cloud Manager, click the **Account** drop-down and click **Manage Account**.



2. Click **Connector**.
3. Click **Manage Workspaces** for the Connector that you want to associate.
4. Select one or more workspaces and click **Apply**.

Result

Workspace Admins can now use those Connectors to create Cloud Volumes ONTAP systems.

What's next?

Now that you've set up your account, you can manage it any time by removing users, managing workspaces, Connectors, and subscriptions. [Learn more](#).

Set up a Connector

Learn about Connectors

In most cases, an Account Admin will need to deploy a *Connector* in your cloud or on-premises network. The Connector enables Cloud Manager to manage resources and processes within your public cloud environment.

When a Connector is required

A Connector is required to use any of the following features within Cloud Manager:

- Cloud Volumes ONTAP
- On-premises ONTAP clusters
- Cloud Data Sense
- Kubernetes
- Cloud Backup
- Monitoring
- Tiering on-prem data
- Global File Cache
- Amazon S3 bucket discovery
- Amazon FSx for ONTAP advanced management features

A Connector is **not** required for Azure NetApp Files, Cloud Volumes Service, or Cloud Sync.



While a Connector isn't required to set up and manage Azure NetApp Files, a Connector is required if you want to use Cloud Data Sense to scan Azure NetApp Files data.

Supported locations

A Connector is supported in the following locations:

- Amazon Web Services
- Microsoft Azure
- Google Cloud
- On your premises



If you want to create a Cloud Volumes ONTAP system in Google Cloud, then you must have a Connector running in Google Cloud, as well. You can't use a Connector that's running in another location.

Connectors should remain running

A Connector should remain running at all times. It's important for the continued health and operation of the services that you enable.

For example, a Connector is a key component in the health and operation of Cloud Volumes ONTAP PAYGO systems. If a Connector is powered down, Cloud Volumes ONTAP PAYGO systems will shut down after losing communication with a Connector for longer than 14 days.

How to create a Connector

An Account Admin needs to create a Connector before a Workspace Admin can create a Cloud Volumes ONTAP working environment and use any of the other features listed above.

An Account Admin can create a Connector in a number of ways:

- Directly from Cloud Manager (recommended)
 - [Create in AWS](#)
 - [Create in Azure](#)
 - [Create in GCP](#)
- [From the AWS Marketplace](#)
- [From the Azure Marketplace](#)
- [By downloading and installing the software on an existing Linux host](#)

When you create your first Cloud Volumes ONTAP working environment, Cloud Manager will prompt you to create a Connector if you don't have one yet.

Permissions

Specific permissions are needed to create the Connector and another set of permissions are needed for the Connector instance itself.

Permissions to create a Connector

The user who creates a Connector from Cloud Manager needs specific permissions to deploy the instance in your cloud provider of choice. Cloud Manager will remind you of the permissions requirements when you create a Connector.

[View policies for each cloud provider.](#)

Permissions for the Connector instance

The Connector needs specific cloud provider permissions to perform operations on your behalf. For example, to deploy and manage Cloud Volumes ONTAP.

When you create a Connector directly from Cloud Manager, Cloud Manager creates the Connector with the permissions that it needs. There's nothing that you need to do.

If you create the Connector yourself from the AWS Marketplace, the Azure Marketplace, or by manually installing the software, then you'll need to make sure that the right permissions are in place.

[View policies for each cloud provider.](#)

When to use multiple Connectors

In some cases, you might only need one Connector, but you might find yourself needing two or more Connectors.

Here are a few examples:

- You're using a multi-cloud environment (AWS and Azure), so you have one Connector in AWS and another in Azure. Each manages the Cloud Volumes ONTAP systems running in those environments.
- A service provider might use one Cloud Central account to provide services for their customers, while using another account to provide disaster recovery for one of their business units. Each account would have separate Connectors.

Using multiple Connectors with the same working environment

You can manage a working environment with multiple Connectors at the same time for disaster recovery purposes. If one Connector goes down, you can switch to the other Connector to immediately manage the working environment.

To set up this configuration:

1. [Switch to another Connector](#)
2. Discover the existing working environment.
 - [Adding existing Cloud Volumes ONTAP systems to Cloud Manager](#)
 - [Discovering ONTAP clusters](#)
3. Set the [Capacity Management Mode](#) to **Manual** on any additional Connectors.

Only the main Connector should be set to **Automatic Mode**. If you switch to another Connector for DR purposes, then you can change the Capacity Management Mode as needed.

When to switch between Connectors

When you create your first Connector, Cloud Manager automatically uses that Connector for each additional working environment that you create. Once you create an additional Connector, you'll need to switch between them to see the working environments that are specific to each Connector.

[Learn how to switch between Connectors.](#)

The local user interface

While you should perform almost all tasks from the [SaaS user interface](#), a local user interface is still available on the Connector. This interface is needed for a few tasks that need to be performed from the Connector itself:

- [Setting a proxy server](#)
- Installing a patch (you'll typically work with NetApp personnel to install a patch)
- Downloading AutoSupport messages (usually directed by NetApp personnel when you have issues)

[Learn how to access the local UI.](#)

Connector upgrades

The Connector automatically updates its software to the latest version, as long as it has [outbound internet access](#) to obtain the software update.

Networking requirements for the Connector

Set up your networking so the Connector can manage resources and processes within your public cloud environment. The most important step is ensuring outbound internet access to various endpoints.



If your network uses a proxy server for all communication to the internet, you can specify the proxy server from the Settings page. Refer to [Configuring the Connector to use a proxy server](#).

Connection to target networks

A Connector requires a network connection to the type of working environment that you're creating and the services that you're planning to enable.

For example, if you install a Connector in your corporate network, then you must set up a VPN connection to the VPC or VNet in which you launch Cloud Volumes ONTAP.

Possible conflict with IP addresses in the 172 range

If your network has a subnet configured in the 172 range, then you might experience connectivity failures from Cloud Manager. [Learn more about this known issue.](#)

Outbound internet access

The Connector requires outbound internet access to manage resources and processes within your public cloud environment. Outbound internet access is also required if you want to manually install the Connector on a Linux host or access the local UI running on the Connector.

The following sections identify the specific endpoints.

Endpoints to manage resources in AWS

A Connector contacts the following endpoints when managing resources in AWS:



If your VPC uses a network access control list (ACL) to filter traffic, then make sure that you enable these endpoints for both outbound and inbound traffic.

Endpoints	Purpose
AWS services (amazonaws.com): <ul style="list-style-type: none">• CloudFormation• Elastic Compute Cloud (EC2)• Key Management Service (KMS)• Security Token Service (STS)• Simple Storage Service (S3) <p>The exact endpoint depends on the region in which you deploy Cloud Volumes ONTAP. Refer to AWS documentation for details.</p>	Enables the Connector to deploy and manage Cloud Volumes ONTAP in AWS.
https://api.services.cloud.netapp.com:443	API requests to NetApp Cloud Central.
https://cloud.support.netapp.com.s3.us-west-1.amazonaws.com	Provides access to software images, manifests, and templates.
https://cognito-idp.us-east-1.amazonaws.com https://cognito-identity.us-east-1.amazonaws.com https://sts.amazonaws.com https://cloud-support-netapp-com-accelerated.s3.amazonaws.com	Enables the Connector to access and download manifests, templates, and Cloud Volumes ONTAP upgrade images.
https://cloudmanagerinfraprod.azurecr.io *.blob.core.windows.net	Access to software images of container components for an infrastructure that's running Docker and provides a solution for service integrations with Cloud Manager.
https://kinesis.us-east-1.amazonaws.com	Enables NetApp to stream data from audit records.
https://cloudmanager.cloud.netapp.com	Communication with the Cloud Manager service, which includes Cloud Central accounts.
https://netapp-cloud-account.auth0.com	Communication with NetApp Cloud Central for centralized user authentication.
support.netapp.com:443 https://mysupport.netapp.com	Communication with NetApp AutoSupport. Note that the Connector communicates with support.netapp.com:443, which redirects to https://mysupport.netapp.com.

Endpoints	Purpose
https://support.netapp.com/svcgw https://support.netapp.com/ServiceGW/entitlement https://eval.lic.netapp.com.s3.us-west-1.amazonaws.com https://cloud-support-netapp-com.s3.us-west-1.amazonaws.com	Communication with NetApp for system licensing and support registration.
https://client.infra.support.netapp.com.s3.us-west-1.amazonaws.com https://cloud-support-netapp-com-accelerated.s3.us-west-1.amazonaws.com https://trigger.asup.netapp.com.s3.us-west-1.amazonaws.com	Enables NetApp to collect information needed to troubleshoot support issues.
https://ipa-signer.cloudmanager.netapp.com	Enables Cloud Manager to generate licenses (for example, a FlexCache license for Cloud Volumes ONTAP)
Various third-party locations, for example: <ul style="list-style-type: none"> • https://repo1.maven.org/maven2 • https://oss.sonatype.org/content/repositories • https://repo.typesafe.com <p>Third-party locations are subject to change.</p>	During upgrades, Cloud Manager downloads the latest packages for third-party dependencies.

Endpoints to manage resources in Azure

A Connector contacts the following endpoints when managing resources in Azure:

Endpoints	Purpose
https://management.azure.com https://login.microsoftonline.com	Enables Cloud Manager to deploy and manage Cloud Volumes ONTAP in most Azure regions.
https://management.microsoftazure.de https://login.microsoftonline.de	Enables Cloud Manager to deploy and manage Cloud Volumes ONTAP in the Azure Germany regions.
https://management.usgovcloudapi.net https://login.microsoftonline.com	Enables Cloud Manager to deploy and manage Cloud Volumes ONTAP in the Azure US Gov regions.
https://api.services.cloud.netapp.com:443	API requests to NetApp Cloud Central.
https://cloud.support.netapp.com.s3.us-west-1.amazonaws.com	Provides access to software images, manifests, and templates.

Endpoints	Purpose
https://cognito-idp.us-east-1.amazonaws.com https://cognito-identity.us-east-1.amazonaws.com https://sts.amazonaws.com https://cloud-support-netapp-com-accelerated.s3.amazonaws.com	Enables the Connector to access and download manifests, templates, and Cloud Volumes ONTAP upgrade images.
https://cloudmanagerinfraprod.azurecr.io <code>*.blob.core.windows.net</code>	Access to software images of container components for an infrastructure that's running Docker and provides a solution for service integrations with Cloud Manager.
https://kinesis.us-east-1.amazonaws.com	Enables NetApp to stream data from audit records.
https://cloudmanager.cloud.netapp.com	Communication with the Cloud Manager service, which includes Cloud Central accounts.
https://netapp-cloud-account.auth0.com	Communication with NetApp Cloud Central for centralized user authentication.
support.netapp.com:443 https://mysupport.netapp.com	Communication with NetApp AutoSupport. Note that the Connector communicates with support.netapp.com:443, which redirects to https://mysupport.netapp.com.
https://support.netapp.com/svcgw https://support.netapp.com/ServiceGW/entitlement https://eval.lic.netapp.com.s3.us-west-1.amazonaws.com https://cloud-support-netapp-com.s3.us-west-1.amazonaws.com	Communication with NetApp for system licensing and support registration.
https://client.infra.support.netapp.com.s3.us-west-1.amazonaws.com https://cloud-support-netapp-com-accelerated.s3.us-west-1.amazonaws.com https://trigger.asup.netapp.com.s3.us-west-1.amazonaws.com	Enables NetApp to collect information needed to troubleshoot support issues.
https://ipa-signer.cloudmanager.netapp.com	Enables Cloud Manager to generate licenses (for example, a FlexCache license for Cloud Volumes ONTAP)
<code>*.blob.core.windows.net</code>	Required for HA pairs when using a proxy.
Various third-party locations, for example: <ul style="list-style-type: none"> • https://repo1.maven.org/maven2 • https://oss.sonatype.org/content/repositories • https://repo.typesafe.com 	During upgrades, Cloud Manager downloads the latest packages for third-party dependencies.
Third-party locations are subject to change.	

Endpoints to manage resources in GCP

A Connector contacts the following endpoints when managing resources in GCP:

Endpoints	Purpose
https://www.googleapis.com	Enables the Connector to contact Google APIs for deploying and managing Cloud Volumes ONTAP in GCP.
https://api.services.cloud.netapp.com:443	API requests to NetApp Cloud Central.
https://cloud.support.netapp.com.s3.us-west-1.amazonaws.com	Provides access to software images, manifests, and templates.
https://cognito-idp.us-east-1.amazonaws.com https://cognito-identity.us-east-1.amazonaws.com https://sts.amazonaws.com https://cloud-support-netapp-com-accelerated.s3.amazonaws.com	Enables the Connector to access and download manifests, templates, and Cloud Volumes ONTAP upgrade images.
https://cloudmanagerinfraprod.azurecr.io *.blob.core.windows.net	Access to software images of container components for an infrastructure that's running Docker and provides a solution for service integrations with Cloud Manager.
https://kinesis.us-east-1.amazonaws.com	Enables NetApp to stream data from audit records.
https://cloudmanager.cloud.netapp.com	Communication with the Cloud Manager service, which includes Cloud Central accounts.
https://netapp-cloud-account.auth0.com	Communication with NetApp Cloud Central for centralized user authentication.
support.netapp.com:443 https://mysupport.netapp.com	Communication with NetApp AutoSupport. Note that the Connector communicates with support.netapp.com:443, which redirects to https://mysupport.netapp.com.
https://support.netapp.com/svcgw https://support.netapp.com/ServiceGW/entitlement https://eval.lic.netapp.com.s3.us-west-1.amazonaws.com https://cloud-support-netapp-com.s3.us-west-1.amazonaws.com	Communication with NetApp for system licensing and support registration.
https://client.infra.support.netapp.com.s3.us-west-1.amazonaws.com https://cloud-support-netapp-com-accelerated.s3.us-west-1.amazonaws.com https://trigger.asup.netapp.com.s3.us-west-1.amazonaws.com	Enables NetApp to collect information needed to troubleshoot support issues.
https://ipa-signer.cloudmanager.netapp.com	Enables Cloud Manager to generate licenses (for example, a FlexCache license for Cloud Volumes ONTAP)

Endpoints	Purpose
<p>Various third-party locations, for example:</p> <ul style="list-style-type: none"> • https://repo1.maven.org/maven2 • https://oss.sonatype.org/content/repositories • https://repo.typesafe.com <p>Third-party locations are subject to change.</p>	<p>During upgrades, Cloud Manager downloads the latest packages for third-party dependencies.</p>

Endpoints to install the Connector on a Linux host

You have the option to manually install the Connector software on your own Linux host. If you do, the installer for the Connector must access the following URLs during the installation process:

- <http://dev.mysql.com/get/mysql-community-release-el7-5.noarch.rpm>
- <https://dl.fedoraproject.org/pub/epel/epel-release-latest-7.noarch.rpm>
- <https://s3.amazonaws.com/aws-cli/awscli-bundle.zip>
- *.blob.core.windows.net

The host might try to update operating system packages during installation. The host can contact different mirroring sites for these OS packages.

Endpoints accessed from your web browser when using the local UI

While you should perform almost all tasks from the SaaS user interface, a local user interface is still available on the Connector. The machine running the web browser must have connections to the following endpoints:

Endpoints	Purpose
<p>The Connector host</p>	<p>You must enter the host's IP address from a web browser to load the Cloud Manager console.</p> <p>Depending on your connectivity to your cloud provider, you can use the private IP or a public IP assigned to the host:</p> <ul style="list-style-type: none"> • A private IP works if you have a VPN and direct connect access to your virtual network • A public IP works in any networking scenario <p>In any case, you should secure network access by ensuring that security group rules allow access from only authorized IPs or subnets.</p>
<p>https://auth0.com https://cdn.auth0.com https://netapp-cloud-account.auth0.com https://services.cloud.netapp.com</p>	<p>Your web browser connects to these endpoints for centralized user authentication through NetApp Cloud Central.</p>

Endpoints	Purpose
https://widget.intercom.io	For in-product chat that enables you to talk to NetApp cloud experts.

Ports and security groups

There's no incoming traffic to the Connector, unless you initiate it. HTTP and HTTPS provide access to the [local UI](#), which you'll use in rare circumstances. SSH is only needed if you need to connect to the host for troubleshooting.

Rules for the Connector in AWS

The security group for the Connector requires both inbound and outbound rules.

Inbound rules

Protocol	Port	Purpose
SSH	22	Provides SSH access to the Connector host
HTTP	80	Provides HTTP access from client web browsers to the local user interface and connections from Cloud Data Sense
HTTPS	443	Provides HTTPS access from client web browsers to the local user interface
TCP	3128	Provides the Cloud Data Sense instance with internet access, if your AWS network doesn't use a NAT or proxy

Outbound rules

The predefined security group for the Connector opens all outbound traffic. If that is acceptable, follow the basic outbound rules. If you need more rigid rules, use the advanced outbound rules.

Basic outbound rules

The predefined security group for the Connector includes the following outbound rules.

Protocol	Port	Purpose
All TCP	All	All outbound traffic
All UDP	All	All outbound traffic

Advanced outbound rules

If you need rigid rules for outbound traffic, you can use the following information to open only those ports that are required for outbound communication by the Connector.



The source IP address is the Connector host.

Service	Protocol	Port	Destination	Purpose
Active Directory	TCP	88	Active Directory forest	Kerberos V authentication
	TCP	139	Active Directory forest	NetBIOS service session
	TCP	389	Active Directory forest	LDAP
	TCP	445	Active Directory forest	Microsoft SMB/CIFS over TCP with NetBIOS framing
	TCP	464	Active Directory forest	Kerberos V change & set password (SET_CHANGE)
	TCP	749	Active Directory forest	Active Directory Kerberos V change & set password (RPCSEC_GSS)
	UDP	137	Active Directory forest	NetBIOS name service
	UDP	138	Active Directory forest	NetBIOS datagram service
	UDP	464	Active Directory forest	Kerberos key administration
API calls and AutoSupport	HTTP	443	Outbound internet and ONTAP cluster management LIF	API calls to AWS and ONTAP, and sending AutoSupport messages to NetApp
API calls	TCP	3000	ONTAP HA mediator	Communication with the ONTAP HA mediator
	TCP	8088	Backup to S3	API calls to Backup to S3
DNS	UDP	53	DNS	Used for DNS resolve by Cloud Manager
Cloud Data Sense	HTTP	80	Cloud Data Sense instance	Cloud Data Sense for Cloud Volumes ONTAP

Rules for the Connector in Azure

The security group for the Connector requires both inbound and outbound rules.

Inbound rules

Port	Protocol	Purpose
22	SSH	Provides SSH access to the Connector host
80	HTTP	Provides HTTP access from client web browsers to the local user interface
443	HTTPS	Provides HTTPS access from client web browsers to the local user interface

Outbound rules

The predefined security group for the Connector opens all outbound traffic. If that is acceptable, follow the basic outbound rules. If you need more rigid rules, use the advanced outbound rules.

Basic outbound rules

The predefined security group for the Connector includes the following outbound rules.

Port	Protocol	Purpose
All	All TCP	All outbound traffic
All	All UDP	All outbound traffic

Advanced outbound rules

If you need rigid rules for outbound traffic, you can use the following information to open only those ports that are required for outbound communication by the Connector.



The source IP address is the Connector host.

Service	Port	Protocol	Destination	Purpose
Active Directory	88	TCP	Active Directory forest	Kerberos V authentication
	139	TCP	Active Directory forest	NetBIOS service session
	389	TCP	Active Directory forest	LDAP
	445	TCP	Active Directory forest	Microsoft SMB/CIFS over TCP with NetBIOS framing
	464	TCP	Active Directory forest	Kerberos V change & set password (SET_CHANGE)
	749	TCP	Active Directory forest	Active Directory Kerberos V change & set password (RPCSEC_GSS)
	137	UDP	Active Directory forest	NetBIOS name service
	138	UDP	Active Directory forest	NetBIOS datagram service
	464	UDP	Active Directory forest	Kerberos key administration
API calls and AutoSupport	443	HTTP	Outbound internet and ONTAP cluster management LIF	API calls to AWS and ONTAP, and sending AutoSupport messages to NetApp
DNS	53	UDP	DNS	Used for DNS resolve by Cloud Manager

Rules for the Connector in GCP

The firewall rules for the Connector requires both inbound and outbound rules.

Inbound rules

Protocol	Port	Purpose
SSH	22	Provides SSH access to the Connector host
HTTP	80	Provides HTTP access from client web browsers to the local user interface
HTTPS	443	Provides HTTPS access from client web browsers to the local user interface

Outbound rules

The predefined firewall rules for the Connector opens all outbound traffic. If that is acceptable, follow the basic outbound rules. If you need more rigid rules, use the advanced outbound rules.

Basic outbound rules

The predefined firewall rules for the Connector includes the following outbound rules.

Protocol	Port	Purpose
All TCP	All	All outbound traffic
All UDP	All	All outbound traffic

Advanced outbound rules

If you need rigid rules for outbound traffic, you can use the following information to open only those ports that are required for outbound communication by the Connector.



The source IP address is the Connector host.

Service	Protocol	Port	Destination	Purpose
Active Directory	TCP	88	Active Directory forest	Kerberos V authentication
	TCP	139	Active Directory forest	NetBIOS service session
	TCP	389	Active Directory forest	LDAP
	TCP	445	Active Directory forest	Microsoft SMB/CIFS over TCP with NetBIOS framing
	TCP	464	Active Directory forest	Kerberos V change & set password (SET_CHANGE)
	TCP	749	Active Directory forest	Active Directory Kerberos V change & set password (RPCSEC_GSS)
	UDP	137	Active Directory forest	NetBIOS name service
	UDP	138	Active Directory forest	NetBIOS datagram service
	UDP	464	Active Directory forest	Kerberos key administration
API calls and AutoSupport	HTTP	443	Outbound internet and ONTAP cluster management LIF	API calls to GCP and ONTAP, and sending AutoSupport messages to NetApp
DNS	UDP	53	DNS	Used for DNS resolve by Cloud Manager

Creating a Connector in AWS from Cloud Manager

An Account Admin needs to deploy a *Connector* before you can use most Cloud Manager features. [Learn when a Connector is required](#). The Connector enables Cloud Manager to manage resources and processes within your public cloud environment.

This page describes how to create a Connector in AWS directly from Cloud Manager. You also have the option to [create the Connector from the AWS Marketplace](#), or to [download the software and install it on your own host](#).

These steps must be completed by a user who has the Account Admin role. A Workspace Admin can't create a Connector.



When you create your first Cloud Volumes ONTAP working environment, Cloud Manager will prompt you to create a Connector if you don't have one yet.

Setting up AWS permissions to create a Connector

Before you can deploy a Connector from Cloud Manager, you need to ensure that your AWS account has the correct permissions.

Steps

1. Download the Connector IAM policy from the following location:

[NetApp Cloud Manager: AWS, Azure, and GCP Policies](#)



For IAM user permissions for Amazon FSx for ONTAP, see [Create an FSx for ONTAP working environment](#).

2. From the AWS IAM console, create your own policy by copying and pasting the text from the Connector IAM policy.
3. Attach the policy that you created in the previous step to the IAM user who will create the Connector from Cloud Manager.

Result

The AWS user now has the permissions required to create the Connector from Cloud Manager. You'll need to specify AWS access keys for this user when you're prompted by Cloud Manager.

Creating a Connector in AWS

Cloud Manager enables you to create a Connector in AWS directly from its user interface.

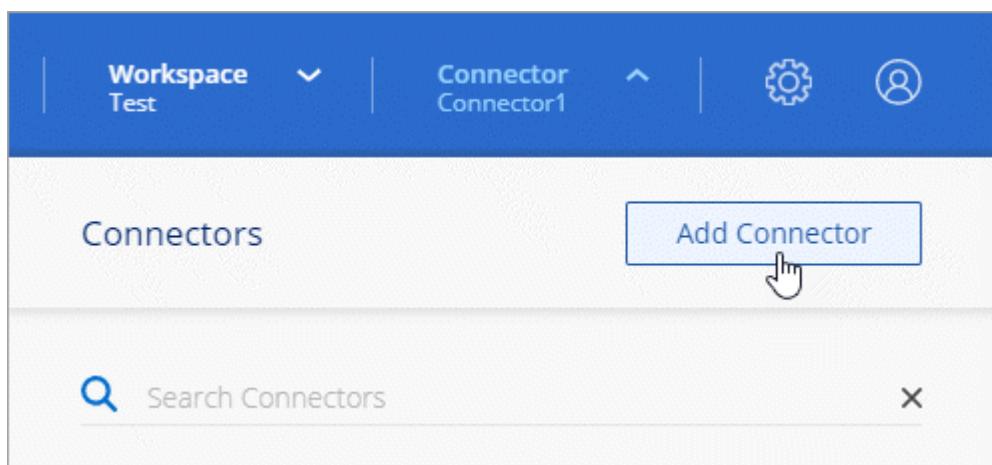
What you'll need

- An AWS access key and secret key for an IAM user who has the [required permissions](#) to create a Connector.
- A VPC, subnet, and keypair in your AWS region of choice.
- If you don't want Cloud Manager to automatically create an IAM role for the Connector, then you'll need to create your own [using this policy](#).

These permissions are for the Connector instance. It's a different set of permissions than what's provided in the first bullet above.

Steps

1. If you're creating your first Working Environment, click **Add Working Environment** and follow the prompts. Otherwise, click the **Connector** drop-down and select **Add Connector**.



2. Choose **Amazon Web Services** as your cloud provider and click **Continue**.

Remember that the Connector must have a network connection to the type of working environment that

you're creating and the services that you're planning to enable.

[Learn more about networking requirements for the Connector.](#)

3. Follow the steps in the wizard to create the Connector:

- **Get Ready:** Review what you'll need.
- **AWS Credentials:** Specify the AWS access key and secret key that meet permissions requirements and then select your region.
- **Details:** Enter a name for the instance, specify tags, and choose whether you want Cloud Manager to create a new role that has the required permissions, or if you want to select an existing role that you set up with [the required permissions](#).
- **Network:** Specify a VPC, subnet, and key pair for the instance, choose whether to enable a public IP address, and optionally specify a proxy configuration.
- **Security Group:** Choose whether to create a new security group or whether to select an existing security group that allows inbound HTTP, HTTPS, and SSH access.



There's no incoming traffic to the Connector, unless you initiate it. HTTP and HTTPS provide access to the [local UI](#), which you'll use in rare circumstances. SSH is only needed if you need to connect to the host for troubleshooting.

- **Review:** Review your selections to verify that your set up is correct.

4. Click **Add**.

The instance should be ready in about 7 minutes. You should stay on the page until the process is complete.

After you finish

You need to associate a Connector with workspaces so Workspace Admins can use those Connectors to create Cloud Volumes ONTAP systems. If you only have Account Admins, then associating the Connector with workspaces isn't required. Account Admins have the ability to access all workspaces in Cloud Manager by default. [Learn more](#).

Creating a Connector in Azure from Cloud Manager

An Account Admin needs to deploy a *Connector* before you can use most Cloud Manager features. [Learn when a Connector is required](#). The Connector enables Cloud Manager to manage resources and processes within your public cloud environment.

This page describes how to create a Connector in Azure directly from Cloud Manager. You also have the option to [create the Connector from the Azure Marketplace](#), or to [download the software and install it on your own host](#).

These steps must be completed by a user who has the Account Admin role. A Workspace Admin can't create a Connector.



When you create your first Cloud Volumes ONTAP working environment, Cloud Manager will prompt you to create a Connector if you don't have one yet.

Setting up Azure permissions to create a Connector

Before you can deploy a Connector from Cloud Manager, you need to ensure that your Azure account has the correct permissions.

Steps

1. Create a custom role using the Azure policy for the Connector:

- a. Download the [Azure policy for the Connector](#).



Right-click the link and click **Save link as...** to download the file.

- b. Modify the JSON file by adding your Azure subscription ID to the assignable scope.

Example

```
"AssignableScopes": [  
    "/subscriptions/d333af45-0d07-4154-943d-c25fbzzzzz"  
],
```

- c. Use the JSON file to create a custom role in Azure.

The following example shows how to create a custom role using the Azure CLI 2.0:

```
az role definition create --role-definition  
C:\Policy_for_Setup_As_Service_Azure.json
```

You should now have a custom role called *Azure SetupAsService*.

2. Assign the role to the user who will deploy the Connector from Cloud Manager:

- a. Open the **Subscriptions** service and select the user's subscription.
- b. Click **Access control (IAM)**.
- c. Click **Add > Add role assignment** and then add the permissions:

- Select the **Azure SetupAsService** role.



Azure SetupAsService is the default name provided in the [Connector deployment policy for Azure](#). If you chose a different name for the role, then select that name instead.

- Assign access to an **Azure AD user, group, or application**.
- Select the user account.
- Click **Save**.

Result

The Azure user now has the permissions required to deploy the Connector from Cloud Manager.

Creating a Connector in Azure

Cloud Manager enables you to create a Connector in Azure directly from its user interface.

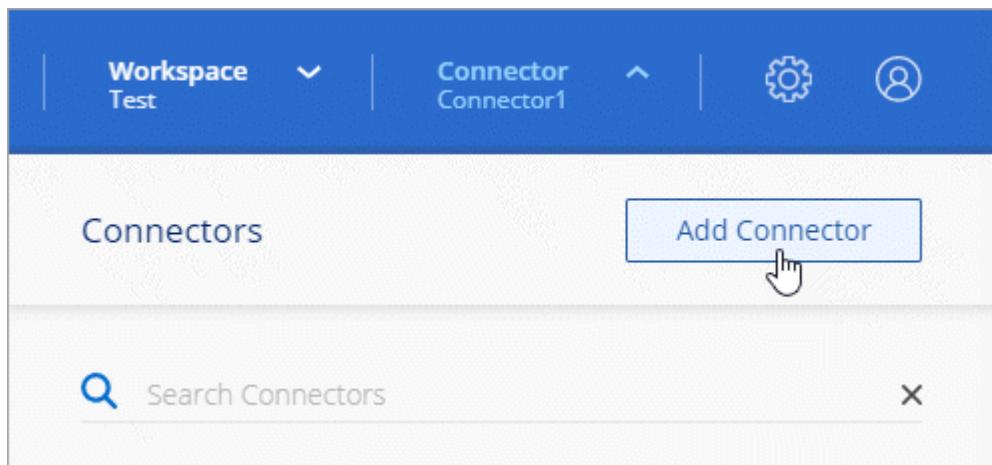
What you'll need

- The [required permissions](#) for your Azure account.
- An Azure subscription.
- A VNet and subnet in your Azure region of choice.
- If you don't want Cloud Manager to automatically create an Azure role for the Connector, then you'll need to create your own [using this policy](#).

These permissions are for the Connector instance. It's a different set of permissions than what's provided in the first bullet above.

Steps

1. If you're creating your first Working Environment, click **Add Working Environment** and follow the prompts. Otherwise, click the **Connector** drop-down and select **Add Connector**.



2. Choose **Microsoft Azure** as your cloud provider.

Remember that the Connector must have a network connection to the type of working environment that you're creating and the services that you're planning to enable.

[Learn more about networking requirements for the Connector.](#)

3. Follow the steps in the wizard to create the Connector:

- **Get Ready:** Review what you'll need.
- If you're prompted, log in to your Microsoft account, which should have the required permissions to create the virtual machine.

The form is owned and hosted by Microsoft. Your credentials are not provided to NetApp.



If you're already logged in to an Azure account, then Cloud Manager will automatically use that account. If you have multiple accounts, then you might need to log out first to ensure that you're using the right account.

- **VM Authentication:** Choose an Azure subscription, a location, a new resource group or an existing resource group, and then choose an authentication method.
- **Details:** Enter a name for the instance, specify tags, and choose whether you want Cloud Manager to create a new role that has the required permissions, or if you want to select an existing role that you set

up with [the required permissions](#).

Note that you can choose the subscriptions associated with this role. Each subscription that you choose provides the Connector with permissions to deploy Cloud Volumes ONTAP in those subscriptions.

- **Network:** Choose a VNet and subnet, whether to enable a public IP address, and optionally specify a proxy configuration.
- **Security Group:** Choose whether to create a new security group or whether to select an existing security group that allows inbound HTTP, HTTPS, and SSH access.



There's no incoming traffic to the Connector, unless you initiate it. HTTP and HTTPS provide access to the [local UI](#), which you'll use in rare circumstances. SSH is only needed if you need to connect to the host for troubleshooting.

- **Review:** Review your selections to verify that your set up is correct.

4. Click **Add**.

The virtual machine should be ready in about 7 minutes. You should stay on the page until the process is complete.

After you finish

You need to associate a Connector with workspaces so Workspace Admins can use those Connectors to create Cloud Volumes ONTAP systems. If you only have Account Admins, then associating the Connector with workspaces isn't required. Account Admins have the ability to access all workspaces in Cloud Manager by default. [Learn more](#).

Creating a Connector in GCP from Cloud Manager

An Account Admin needs to deploy a *Connector* before you can use most Cloud Manager features. [Learn when a Connector is required](#). The Connector enables Cloud Manager to manage resources and processes within your public cloud environment.

This page describes how to create a Connector in GCP directly from Cloud Manager. You also have the option to [download the software and install it on your own host](#).

These steps must be completed by a user who has the Account Admin role. A Workspace Admin can't create a Connector.



When you create your first Cloud Volumes ONTAP working environment, Cloud Manager will prompt you to create a Connector if you don't have one yet.

Setting up GCP permissions to create a Connector

Before you can deploy a Connector from Cloud Manager, you need to ensure that your GCP account has the correct permissions and that a service account is set up for the Connector VM.

Steps

1. Ensure that the GCP user who deploys Cloud Manager from NetApp Cloud Central has the permissions in the [Connector deployment policy for GCP](#).

You can create a custom role using the YAML file and then attach it to the user. You'll need to use the gcloud command line to create the role.

- Set up a service account that has the permissions that Cloud Manager needs to create and manage Cloud Volumes ONTAP systems in projects.

You'll associate this service account with the Connector VM when you create it from Cloud Manager.

- Create a role in GCP that includes the permissions defined in the [Cloud Manager policy for GCP](#). Again, you'll need to use the gcloud command line.

The permissions contained in this YAML file are different than the permissions in step 1.

- Create a GCP service account and apply the custom role that you just created.
- If you want to deploy Cloud Volumes ONTAP in other projects, [grant access by adding the service account with the Cloud Manager role to that project](#). You'll need to repeat this step for each project.

Result

The GCP user now has the permissions required to create the Connector from Cloud Manager and the service account for the Connector VM is set up.

Shared VPC Permissions

If you are using a shared VPC to deploy resources into a service project, then the following permissions are required. This table is for reference and your environment should reflect the permissions table when IAM configuration is complete.

Service Account	Creator	Hosted in	Service project permissions	Host project permissions	Purpose
Cloud Manager service account	Custom	Service project	<ul style="list-style-type: none">The permissions found in this .yaml file	<ul style="list-style-type: none">compute.networkUserdeploymentmanager.editor	Deploying and maintaining Cloud Volumes ONTAP and services in the service project
Cloud Volumes ONTAP service account	Custom	Service project	<ul style="list-style-type: none">storage.adminmember: Cloud Manager service account as serviceAccount.user	N/A	(Optional) For data tiering and Cloud Backup
Google APIs service agent	GCP	Service project	<ul style="list-style-type: none">(Default) Editor	<ul style="list-style-type: none">compute.networkUser	Interacts with GCP APIs on behalf of deployment. Allows Cloud Manager to use the shared network.

Service Account	Creator	Hosted in	Service project permissions	Host project permissions	Purpose
Google Compute Engine default service account	GCP	Service project	• (Default) Editor	• compute.networkUser	Deploys GCP instances and compute infrastructure on behalf of deployment. Allows Cloud Manager to use the shared network.

Notes:

1. `deploymentmanager.editor` is only required at the host project if you are not passing firewall rules to the deployment and are choosing to let Cloud Manager create them for you.
2. `firewall.create` and `firewall.delete` are only required if you are not passing firewall rules to the deployment and are choosing to let Cloud Manager create them for you.
3. For data tiering, the tiering service account must have the `serviceAccount.user` role on the service account, not just at the project level. Currently if you assign `serviceAccount.user` at the project level, the permissions don't show when you query the service account with `getIAMPolicy`.

Enabling Google Cloud APIs

Several APIs are required to deploy the Connector and Cloud Volumes ONTAP.

Step

1. [Enable the following Google Cloud APIs in your project.](#)
 - Cloud Deployment Manager V2 API
 - Cloud Logging API
 - Cloud Resource Manager API
 - Compute Engine API
 - Identity and Access Management (IAM) API

Creating a Connector in GCP

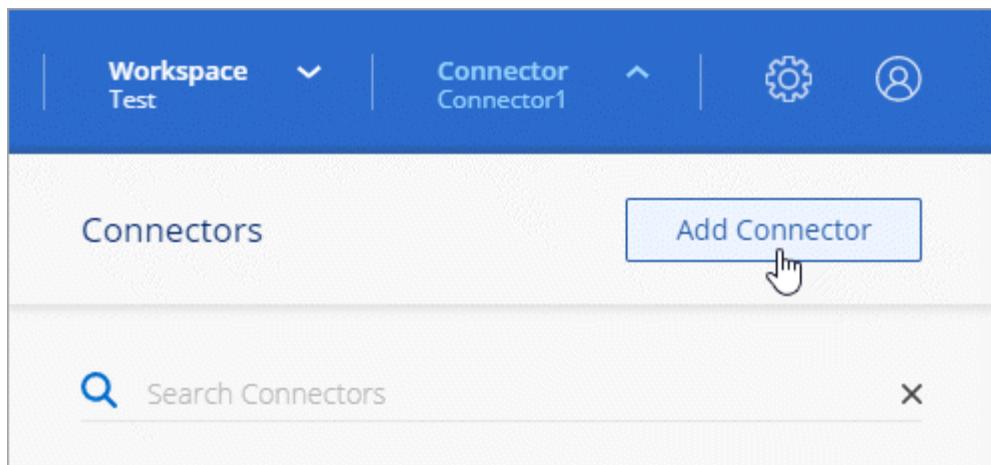
Cloud Manager enables you to create a Connector in GCP directly from its user interface.

What you'll need

- The [required permissions](#) for your Google Cloud account, as described in the first section of this page.
- A Google Cloud project.
- A service account that has the required permissions to create and manage Cloud Volumes ONTAP, as described in the first section of this page.
- A VPC and subnet in your Google Cloud region of choice.

Steps

1. If you're creating your first Working Environment, click **Add Working Environment** and follow the prompts. Otherwise, click the **Connector** drop-down and select **Add Connector**.



2. Choose **Google Cloud Platform** as your cloud provider.

Remember that the Connector must have a network connection to the type of working environment that you're creating and the services that you're planning to enable.

[Learn more about networking requirements for the Connector.](#)

3. Follow the steps in the wizard to create the Connector:

- **Get Ready:** Review what you'll need.
- If you're prompted, log in to your Google account, which should have the required permissions to create the virtual machine instance.

The form is owned and hosted by Google. Your credentials are not provided to NetApp.

- **Basic Settings:** Enter a name for the virtual machine instance, specify tags, select a project, and then select the service account that has the required permissions (refer to the section above for details).
- **Location:** Specify a region, zone, VPC, and subnet for the instance.
- **Network:** Choose whether to enable a public IP address and optionally specify a proxy configuration.
- **Firewall Policy:** Choose whether to create a new firewall policy or whether to select an existing firewall policy that allows inbound HTTP, HTTPS, and SSH access.



There's no incoming traffic to the Connector, unless you initiate it. HTTP and HTTPS provide access to the [local UI](#), which you'll use in rare circumstances. SSH is only needed if you need to connect to the host for troubleshooting.

- **Review:** Review your selections to verify that your set up is correct.

4. Click **Add**.

The instance should be ready in about 7 minutes. You should stay on the page until the process is complete.

After you finish

You need to associate a Connector with workspaces so Workspace Admins can use those Connectors to create Cloud Volumes ONTAP systems. If you only have Account Admins, then associating the Connector with workspaces isn't required. Account Admins have the ability to access all workspaces in Cloud Manager by default. [Learn more](#).

Where to go next

Now that you've logged in and set up Cloud Manager, users can start creating and discovering working environments.

- [Get started with Cloud Volumes ONTAP for AWS](#)
- [Get started with Cloud Volumes ONTAP for Azure](#)
- [Get started with Cloud Volumes ONTAP for Google Cloud](#)
- [Set up Azure NetApp Files](#)
- [Set up Amazon FSx for ONTAP](#)
- [Set up Cloud Volumes Service for AWS](#)
- [Discover an on-premises ONTAP cluster](#)
- [Discover your Amazon S3 buckets](#)

If you're an administrator, you can manage Cloud Manager settings after you create your first Connector.

- [Learn about Connectors](#)
- [Manage an HTTPS certificate for secure access](#)
- [Configure proxy settings](#)

Manage Cloud Volumes ONTAP

Learn

Learn about Cloud Volumes ONTAP

Cloud Volumes ONTAP enables you to optimize your cloud storage costs and performance while enhancing data protection, security, and compliance.

Cloud Volumes ONTAP is a software-only storage appliance that runs ONTAP data management software in the cloud. It provides enterprise-grade storage with the following key features:

- Storage efficiencies

Leverage built-in data deduplication, data compression, thin provisioning, and cloning to minimize storage costs.

- High availability

Ensure enterprise reliability and continuous operations in case of failures in your cloud environment.

- Data protection

Cloud Volumes ONTAP leverages SnapMirror, NetApp's industry-leading replication technology, to replicate on-premises data to the cloud so it's easy to have secondary copies available for multiple use cases.

Cloud Volumes ONTAP also integrates with Cloud Backup service to deliver backup and restore capabilities for protection, and long-term archive of your cloud data.

- Data tiering

Switch between high and low-performance storage pools on-demand without taking applications offline.

- Application consistency

Ensure consistency of NetApp Snapshot copies using NetApp SnapCenter.

- Data security

Cloud Volumes ONTAP supports data encryption and provides protection against viruses and ransomware.

- Privacy compliance controls

Integration with Cloud Data Sense helps you understand data context and identify sensitive data.



Licenses for ONTAP features are included with Cloud Volumes ONTAP.

[View supported Cloud Volumes ONTAP configurations](#)

[Learn more about Cloud Volumes ONTAP](#)

Storage

Disks and aggregates

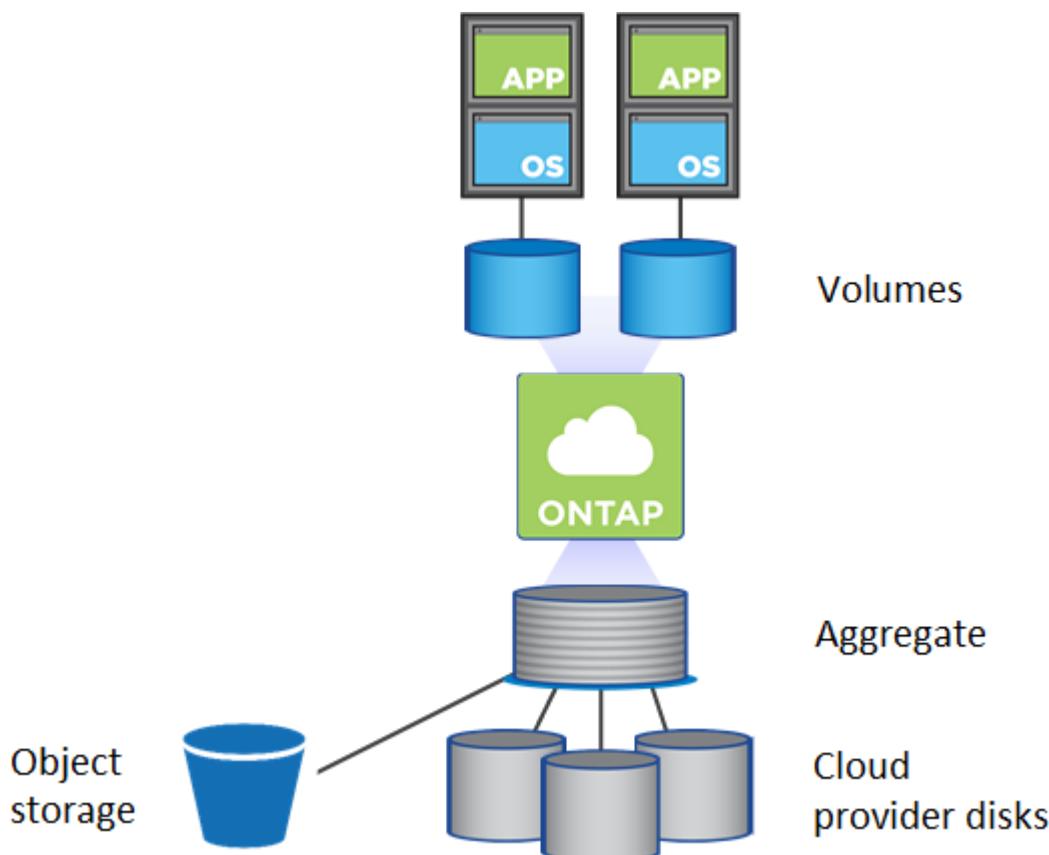
Understanding how Cloud Volumes ONTAP uses cloud storage can help you understand your storage costs.



All disks and aggregates must be created and deleted directly from Cloud Manager. You should not perform these actions from another management tool. Doing so can impact system stability, hamper the ability to add disks in the future, and potentially generate redundant cloud provider fees.

Overview

Cloud Volumes ONTAP uses cloud provider storage as disks and groups them into one or more aggregates. Aggregates provide storage to one or more volumes.



Several types of cloud disks are supported. You choose the disk type when you create a volume and the default disk size when you deploy Cloud Volumes ONTAP.



The total amount of storage purchased from a cloud provider is the *raw capacity*. The *usable capacity* is less because approximately 12 to 14 percent is overhead that is reserved for Cloud Volumes ONTAP use. For example, if Cloud Manager creates a 500 GiB aggregate, the usable capacity is 442.94 GiB.

AWS storage

In AWS, Cloud Volumes ONTAP uses EBS storage for user data and local NVMe storage as Flash Cache on some EC2 instance types.

EBS storage

In AWS, an aggregate can contain up to 6 disks that are all the same size. The maximum disk size is 16 TiB.

The underlying EBS disk type can be either General Purpose SSDs (gp3 or gp2), Provisioned IOPS SSD (io1), or Throughput Optimized HDD (st1). You can pair an EBS disk with Amazon S3 to [tier inactive data to low-cost object storage](#).



Tiering data to object storage is not recommended when using Throughput Optimized HDDs (st1).

Local NVMe storage

Some EC2 instance types include local NVMe storage, which Cloud Volumes ONTAP uses as [Flash Cache](#).

Related links

- [AWS documentation: EBS Volume Types](#)
- [Learn how to choose disk types and disk sizes for your systems in AWS](#)
- [Review storage limits for Cloud Volumes ONTAP in AWS](#)
- [Review supported configurations for Cloud Volumes ONTAP in AWS](#)

Azure storage

In Azure, an aggregate can contain up to 12 disks that are all the same size. The disk type and maximum disk size depends on whether you use a single node system or an HA pair:

Single node systems

Single node systems can use three types of Azure Managed Disks:

- *Premium SSD Managed Disks* provide high performance for I/O-intensive workloads at a higher cost.
- *Standard SSD Managed Disks* provide consistent performance for workloads that require low IOPS.
- *Standard HDD Managed Disks* are a good choice if you don't need high IOPS and want to reduce your costs.

Each managed disk type has a maximum disk size of 32 TiB.

You can pair a managed disk with Azure Blob storage to [tier inactive data to low-cost object storage](#).

HA pairs

HA pairs use Premium page blobs, which have a maximum disk size of 8 TiB.

Related links

- [Microsoft Azure documentation: Introduction to Microsoft Azure Storage](#)
- [Learn how to choose disk types and disk sizes for your systems in Azure](#)

- [Review storage limits for Cloud Volumes ONTAP in Azure](#)

GCP storage

In GCP, an aggregate can contain up to 6 disks that are all the same size. The maximum disk size is 16 TiB.

The disk type can be either *Zonal SSD persistent disks*, *Zonal Balanced persistent disks*, or *Zonal standard persistent disks*. You can pair persistent disks with a Google Storage bucket to [tier inactive data to low-cost object storage](#).

Related links

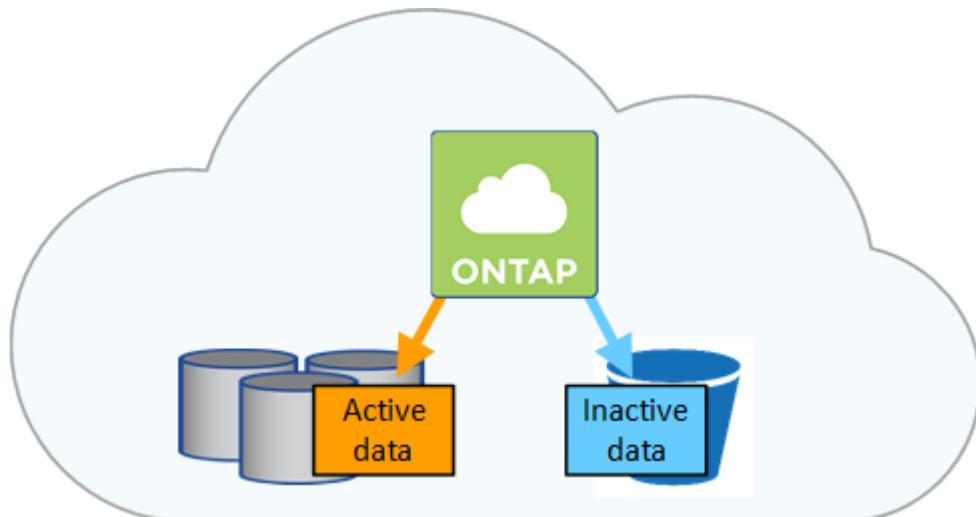
- [Google Cloud Platform documentation: Storage Options](#)
- [Review storage limits for Cloud Volumes ONTAP in GCP](#)

RAID type

The RAID type for each Cloud Volumes ONTAP aggregate is RAID0 (striping). No other RAID types are supported. Cloud Volumes ONTAP relies on the cloud provider for disk availability and durability.

Data tiering overview

Reduce your storage costs by enabling automated tiering of inactive data to low-cost object storage. Active data remains in high-performance SSDs or HDDs, while inactive data is tiered to low-cost object storage. This enables you to reclaim space on your primary storage and shrink secondary storage.



Cloud Volumes ONTAP supports data tiering in AWS, Azure, and Google Cloud Platform. Data tiering is powered by FabricPool technology.



You don't need to install a feature license to enable data tiering (FabricPool).

Data tiering in AWS

When you enable data tiering in AWS, Cloud Volumes ONTAP uses EBS as a performance tier for hot data and AWS S3 as a capacity tier for inactive data.

Performance tier

The performance tier can be General Purpose SSDs (gp3 or gp2) or Provisioned IOPS SSDs (io1).



Tiering data to object storage is not recommended when using Throughput Optimized HDDs (st1).

Capacity tier

A Cloud Volumes ONTAP system tiers inactive data to a single S3 bucket using the *Standard* storage class. Standard is ideal for frequently accessed data stored across multiple Availability Zones.



Cloud Manager creates a single S3 bucket for each working environment and names it *fabric-pool-cluster unique identifier*. A different S3 bucket is not created for each volume.

Storage classes

The default storage class for tiered data in AWS is *Standard*. If you don't plan to access the inactive data, you can reduce your storage costs by changing the storage class to one of the following: *Intelligent Tiering*, *One-Zone Infrequent Access*, or *Standard-Infrequent Access*. When you change the storage class, inactive data starts in the Standard storage class and transitions to the storage class that you selected, if the data is not accessed after 30 days.

The access costs are higher if you do access the data, so take that into consideration before you change the storage class. [Learn more about Amazon S3 storage classes](#).

You can select a storage class when you create the working environment and you can change it any time after. For details about changing the storage class, see [Tiering inactive data to low-cost object storage](#).

The storage class for data tiering is system wide—it's not per volume.

Data tiering in Azure

When you enable data tiering in Azure, Cloud Volumes ONTAP uses Azure managed disks as a performance tier for hot data and Azure Blob storage as a capacity tier for inactive data.

Performance tier

The performance tier can be either SSDs or HDDs.

Capacity tier

A Cloud Volumes ONTAP system tiers inactive data to a single Blob container using the Azure *hot* storage tier. The hot tier is ideal for frequently accessed data.



Cloud Manager creates a new storage account with a single container for each Cloud Volumes ONTAP working environment. The name of the storage account is random. A different container is not created for each volume.

Storage access tiers

The default storage access tier for tiered data in Azure is the *hot* tier. If you don't plan to access the inactive data, you can reduce your storage costs by changing to the *cool* storage tier. When you change the storage tier, inactive data starts in the hot storage tier and transitions to the cool storage tier, if the data is not accessed after 30 days.

The access costs are higher if you do access the data, so take that into consideration before you change the storage tier. [Learn more about Azure Blob storage access tiers](#).

You can select a storage tier when you create the working environment and you can change it any time after. For details about changing the storage tier, see [Tiering inactive data to low-cost object storage](#).

The storage access tier for data tiering is system wide—it's not per volume.

Data tiering in GCP

When you enable data tiering in GCP, Cloud Volumes ONTAP uses persistent disks as a performance tier for hot data and a Google Cloud Storage bucket as a capacity tier for inactive data.

Performance tier

The performance tier can be either SSD persistent disks, balanced persistent disks, or standard persistent disks.

Capacity tier

A Cloud Volumes ONTAP system tiers inactive data to a single Google Cloud Storage bucket using the *Regional* storage class.



Cloud Manager creates a single bucket for each working environment and names it fabric-pool-cluster *unique identifier*. A different bucket is not created for each volume.

Storage classes

The default storage class for tiered data is the *Standard Storage* class. If the data is infrequently accessed, you can reduce your storage costs by changing to *Nearline Storage* or *Coldline Storage*. When you change the storage class, inactive data starts in the Standard Storage class and transitions to the storage class that you selected, if the data is not accessed after 30 days.

The access costs are higher if you do access the data, so take that into consideration before you change the storage class. [Learn more about storage classes for Google Cloud Storage](#).

You can select a storage tier when you create the working environment and you can change it any time after. For details about changing the storage class, see [Tiering inactive data to low-cost object storage](#).

The storage class for data tiering is system wide—it's not per volume.

Data tiering and capacity limits

If you enable data tiering, a system's capacity limit stays the same. The limit is spread across the performance tier and the capacity tier.

Volume tiering policies

To enable data tiering, you must select a volume tiering policy when you create, modify, or replicate a volume. You can select a different policy for each volume.

Some tiering policies have an associated minimum cooling period, which sets the time that user data in a volume must remain inactive for the data to be considered "cold" and moved to the capacity tier. The cooling period starts when data is written to the aggregate.



You can change the minimum cooling period and default aggregate threshold of 50% (more on that below). [Learn how to change the cooling period](#) and [learn how to change the threshold](#).

Cloud Manager enables you to choose from the following volume tiering policies when you create or modify a volume:

Snapshot Only

After an aggregate has reached 50% capacity, Cloud Volumes ONTAP tiers cold user data of Snapshot copies that are not associated with the active file system to the capacity tier. The cooling period is approximately 2 days.

If read, cold data blocks on the capacity tier become hot and are moved to the performance tier.

All

All data (not including metadata) is immediately marked as cold and tiered to object storage as soon as possible. There is no need to wait 48 hours for new blocks in a volume to become cold. Note that blocks located in the volume prior to the All policy being set require 48 hours to become cold.

If read, cold data blocks on the cloud tier stay cold and are not written back to the performance tier. This policy is available starting with ONTAP 9.6.

Auto

After an aggregate has reached 50% capacity, Cloud Volumes ONTAP tiers cold data blocks in a volume to a capacity tier. The cold data includes not just Snapshot copies but also cold user data from the active file system. The cooling period is approximately 31 days.

This policy is supported starting with Cloud Volumes ONTAP 9.4.

If read by random reads, the cold data blocks in the capacity tier become hot and move to the performance tier. If read by sequential reads, such as those associated with index and antivirus scans, the cold data blocks stay cold and do not move to the performance tier.

None

Keeps data of a volume in the performance tier, preventing it from being moved to the capacity tier.

When you replicate a volume, you can choose whether to tier the data to object storage. If you do, Cloud Manager applies the **Backup** policy to the data protection volume. Starting with Cloud Volumes ONTAP 9.6, the **All** tiering policy replaces the backup policy.

Turning off Cloud Volumes ONTAP impacts the cooling period

Data blocks are cooled by cooling scans. During this process, blocks that haven't been used have their block temperature moved (cooled) to the next lower value. The default cooling time depends on the volume tiering policy:

- Auto: 31 days
- Snapshot Only: 2 days

Cloud Volumes ONTAP must be running for the cooling scan to work. If Cloud Volumes ONTAP is turned off, cooling will stop, as well. As a result, you can experience longer cooling times.



When Cloud Volumes ONTAP is turned off, the temperature of each block is preserved until you restart the system. For example, if the temperature of a block is 5 when you turn the system off, the temp is still 5 when you turn the system back on.

Setting up data tiering

For instructions and a list of supported configurations, see [Tiering inactive data to low-cost object storage](#).

Storage management

Cloud Manager provides simplified and advanced management of Cloud Volumes ONTAP storage.



All disks and aggregates must be created and deleted directly from Cloud Manager. You should not perform these actions from another management tool. Doing so can impact system stability, hamper the ability to add disks in the future, and potentially generate redundant cloud provider fees.

Storage provisioning

Cloud Manager makes storage provisioning for Cloud Volumes ONTAP easy by purchasing disks and managing aggregates for you. You simply need to create volumes. You can use an advanced allocation option to provision aggregates yourself, if desired.

Simplified provisioning

Aggregates provide cloud storage to volumes. Cloud Manager creates aggregates for you when you launch an instance, and when you provision additional volumes.

When you create a volume, Cloud Manager does one of three things:

- It places the volume on an existing aggregate that has sufficient free space.
- It places the volume on an existing aggregate by purchasing more disks for that aggregate.
- It purchases disks for a new aggregate and places the volume on that aggregate.

Cloud Manager determines where to place a new volume by looking at several factors: an aggregate's maximum size, whether thin provisioning is enabled, and free space thresholds for aggregates.



The Account Admin can modify free space thresholds from the **Settings** page.

Disk size selection for aggregates in AWS

When Cloud Manager creates new aggregates for Cloud Volumes ONTAP in AWS, it gradually increases the disk size in an aggregate, as the number of aggregates in the system increases. Cloud Manager does this to ensure that you can utilize the system's maximum capacity before it reaches the maximum number of data disks allowed by AWS.

For example, Cloud Manager might choose the following disk sizes for aggregates in a Cloud Volumes ONTAP Premium or BYOL system:

Aggregate number	Disk size	Max aggregate capacity
1	500 GiB	3 TiB
4	1 TiB	6 TiB
6	2 TiB	12 TiB

You can choose the disk size yourself by using the advanced allocation option.

Advanced allocation

Rather than let Cloud Manager manage aggregates for you, you can do it yourself. [From the Advanced allocation page](#), you can create new aggregates that include a specific number of disks, add disks to an existing aggregate, and create volumes in specific aggregates.

Capacity management

The Account Admin can choose whether Cloud Manager notifies you of storage capacity decisions or whether Cloud Manager automatically manages capacity requirements for you. It might help for you to understand how these modes work.

Automatic capacity management

The Capacity Management Mode is set to automatic by default. In this mode, Cloud Manager automatically purchases new disks for Cloud Volumes ONTAP instances when more capacity is needed, deletes unused collections of disks (aggregates), moves volumes between aggregates when needed, and attempts to unfail disks.

The following examples illustrate how this mode works:

- If an aggregate with 5 or fewer EBS disks reaches the capacity threshold, Cloud Manager automatically purchases new disks for that aggregate so volumes can continue to grow.

Cloud Manager checks the free space ratio every 15 minutes to determine if it needs to purchase additional disks.

- If an aggregate with 12 Azure disks reaches the capacity threshold, Cloud Manager automatically moves a volume from that aggregate to an aggregate with available capacity or to a new aggregate.

If Cloud Manager creates a new aggregate for the volume, it chooses a disk size that accommodates the size of that volume.

Note that free space is now available on the original aggregate. Existing volumes or new volumes can use that space. The space can't be returned to AWS, Azure, or GCP in this scenario.

- If an aggregate contains no volumes for more than 12 hours, Cloud Manager deletes it.

Management of LUNs with automatic capacity management

Cloud Manager's automatic capacity management doesn't apply to LUNs. When Cloud Manager creates a LUN, it disables the autogrow feature.

Manual capacity management

If the Account Admin set the Capacity Management Mode to manual, Cloud Manager displays Action Required messages when capacity decisions must be made. The same examples described in the automatic mode apply to the manual mode, but it is up to you to accept the actions.

Write speed

Cloud Manager enables you to choose normal or high write speed for Cloud Volumes ONTAP. Before you choose a write speed, you should understand the differences between the normal and high settings and risks and recommendations when using high

write speed.

High write speed is supported with all types of single node systems. It's also supported with HA pairs in AWS and Azure when using a specific instance or VM type (refer to the sections below for the list of supported instances and VM types). High write speed is not supported with HA pairs in GCP.

Normal write speed

When you choose normal write speed, data is written directly to disk. When data is written directly to disk, reduces the likelihood of data loss in the event of an unplanned system outage, or a cascading failure involving an unplanned system outage (HA pairs only).

Normal write speed is the default option.

High write speed

When you choose high write speed, data is buffered in memory before it is written to disk, which provides faster write performance. Due to this caching, there is the potential for data loss if an unplanned system outage occurs.

The amount of data that can be lost in the event of an unplanned system outage is the span of the last two consistency points. A consistency point is the act of writing buffered data to disk. A consistency point occurs when the write log is full or after 10 seconds (whichever comes first). However, the performance of the storage provided by your cloud provider can affect consistency point processing time.

When to use high write speed

High write speed is a good choice if fast write performance is required for your workload and you can withstand the risk of data loss in the event of an unplanned system outage, or a cascading failure involving an unplanned system outage (HA pairs only).

Recommendations when using high write speed

If you enable high write speed, you should ensure write protection at the application layer, or that the applications can tolerate data loss, if it occurs.

Configurations that support high write speed

Not all Cloud Volumes ONTAP configurations support high write speed. Those configurations use normal write speed by default.

AWS

If you use a single node system, Cloud Volumes ONTAP supports high write speed with all instance types.

Starting with the 9.8 release, Cloud Volumes ONTAP supports high write speed with HA pairs when using almost all supported EC2 instance types, except for m5.xlarge and r5.xlarge.

[Learn more about the Amazon EC2 instances that Cloud Volumes ONTAP supports.](#)

Azure

If you use a single node system, Cloud Volumes ONTAP supports high write speed with all VM types.

If you use an HA pair, Cloud Volumes ONTAP supports high write speed with the following VM types, starting

with the 9.8 release:

- DS5_v2
- DS14_v2
- DS15_v2
- E48s_v3

[Learn more about the Azure VM types that Cloud Volumes ONTAP supports.](#)

Google Cloud

If you use a single node system, Cloud Volumes ONTAP supports high write speed with all machine types.

Cloud Volumes ONTAP doesn't support high write speed with HA pairs in Google Cloud.

[Learn more about the Google Cloud machine types that Cloud Volumes ONTAP supports.](#)

How to select a write speed

You can choose a write speed when you create a new working environment and you can [change the write speed for an existing system](#).

What to expect if data loss occurs

If you choose high write speed and data loss occurs, the system should be able to boot up and continue to serve data without user intervention. Two EMS messages will be reported when a node runs into data loss. One is wafl.root.content.changed with the ERROR severity level event, the other is nv.check.failed with the DEBUG severity level event. Both messages must be present as an indication of data loss.

How to stop data access if data loss occurs

If you are concerned about data loss, want the applications to stop running upon data loss, and the data access to be resumed after the data loss issue is properly addressed, you can use the NVFAIL option from the CLI to achieve that goal.

To enable the NVFAIL option

```
vol modify -volume <vol-name> -nvfail on
```

To check NVFAIL settings

```
vol show -volume <vol-name> -fields nvfail
```

To disable the NVFAIL option

```
vol modify -volume <vol-name> -nvfail off
```

When data loss occurs, an NFS or iSCSI volume with NVFAIL enabled should stop serving data (there's no impact to CIFS which is a stateless protocol). For more details, refer to [How NVFAIL impacts access to NFS volumes or LUNs](#).

To check the NVFAIL state

```
vol show -fields in-nvfailed-state
```

After the data loss issue is properly addressed, you can clear the NVFAIL state and the volume will be available for data access.

To clear the NVFAIL state

```
vol modify -volume <vol-name> -in-nvfailed-state false
```

Flash Cache

Some Cloud Volumes ONTAP configurations in AWS and Azure include local NVMe storage, which Cloud Volumes ONTAP uses as *Flash Cache* for better performance.

What's Flash Cache?

Flash Cache speeds access to data through real-time intelligent caching of recently read user data and NetApp metadata. It's effective for random read-intensive workloads, including databases, email, and file services.

Supported instances in AWS

Select one of the following EC2 instance types with a new or existing Cloud Volumes ONTAP Premium or BYOL system:

- c5d.4xlarge
- c5d.9xlarge
- c5d.18xlarge
- m5d.8xlarge
- m5d.12xlarge
- r5d.2xlarge

Supported VM type in Azure

Select the Standard_L8s_v2 VM type with a single node Cloud Volumes ONTAP BYOL system in Azure.

Limitations

- Compression must be disabled on all volumes to take advantage of the Flash Cache performance improvements.

Choose no storage efficiency when creating a volume from Cloud Manager, or create a volume and then [disable data compression by using the CLI](#).

- Cache rewarming after a reboot is not supported with Cloud Volumes ONTAP.

WORM storage

You can activate write once, read many (WORM) storage on a Cloud Volumes ONTAP system to retain files in unmodified form for a specified retention period. WORM storage is powered by SnapLock technology in Enterprise mode, which means WORM files are protected at the file level.

Once a file has been committed to WORM storage, it can't be modified, even after the retention period has expired. A tamper-proof clock determines when the retention period for a WORM file has elapsed.

After the retention period has elapsed, you are responsible for deleting any files that you no longer need.

Activating WORM storage

You can activate WORM storage on a Cloud Volumes ONTAP system when you create a new working environment. This includes setting the default retention period for files.



You can't activate WORM storage on individual volumes—WORM must be activated at the system level.

The following image shows how to activate WORM storage when creating a working environment:

The screenshot shows the Cloud Manager interface for creating a new working environment. The 'WORM (write once, read many)' option is selected. Under 'Write Speed', 'Normal' is chosen, with a note explaining it reduces data loss in case of outage. 'High' is also available, buffering data in memory. An 'Activate WORM' checkbox is checked. A retention period of 15 years is set. The bottom right corner features a blue circular icon with a white envelope symbol.

Committing files to WORM

You can use an application to commit files to WORM over NFS or CIFS, or use the ONTAP CLI to autocommit files to WORM automatically. You can also use a WORM appendable file to retain data that is written incrementally, like log information.

After you activate WORM storage on a Cloud Volumes ONTAP system, you must use the ONTAP CLI for all management of WORM storage. For instructions, refer to [ONTAP documentation](#).



Cloud Volumes ONTAP support for WORM storage is equivalent to SnapLock Enterprise mode.

Limitations

- WORM storage in Cloud Volumes ONTAP operates under a "trusted storage administrator" model. While WORM files are protected from alteration or modification, volumes can be deleted by a cluster administrator even if those volumes contain unexpired WORM data.
- In addition to the trusted storage administrator model, WORM storage in Cloud Volumes ONTAP also implicitly operates under a "trusted cloud administrator" model. A cloud administrator could delete WORM data before its expiry date by removing or editing cloud storage directly from the cloud provider.
- When WORM storage is activated, data tiering to object storage can't be enabled.
- Cloud Backup must be disabled in order to enable WORM storage.

High-availability pairs

High-availability pairs in AWS

A Cloud Volumes ONTAP high availability (HA) configuration provides nondisruptive operations and fault tolerance. In AWS, data is synchronously mirrored between the two nodes.

Overview

In AWS, Cloud Volumes ONTAP HA configurations include the following components:

- Two Cloud Volumes ONTAP nodes whose data is synchronously mirrored between each other.
- A mediator instance that provides a communication channel between the nodes to assist in storage takeover and giveback processes.



The mediator instance runs the Linux operating system on a t2.micro instance and uses one EBS magnetic disk that is approximately 8 GiB.

Storage takeover and giveback

If a node goes down, the other node can serve data for its partner to provide continued data service. Clients can access the same data from the partner node because the data was synchronously mirrored to the partner.

After the node reboots, the partner must resync data before it can return the storage. The time that it takes to resync data depends on how much data was changed while the node was down.

Storage takeover, resync, and giveback are all automatic by default. No user action is required.

RPO and RTO

An HA configuration maintains high availability of your data as follows:

- The recovery point objective (RPO) is 0 seconds.
Your data is transactionally consistent with no data loss.
- The recovery time objective (RTO) is 60 seconds.
In the event of an outage, data should be available in 60 seconds or less.

HA deployment models

You can ensure the high availability of your data by deploying an HA configuration across multiple Availability Zones (AZs) or in a single AZ. You should review more details about each configuration to choose which best fits your needs.

Multiple Availability Zones

Deploying an HA configuration in multiple Availability Zones (AZs) ensures high availability of your data if a failure occurs with an AZ or an instance that runs a Cloud Volumes ONTAP node. You should understand how NAS IP addresses impact data access and storage failover.

NFS and CIFS data access

When an HA configuration is spread across multiple Availability Zones, *floating IP addresses* enable NAS client access. The floating IP addresses, which must be outside of the CIDR blocks for all VPCs in the region, can migrate between nodes when failures occur. They aren't natively accessible to clients that are outside of the VPC, unless you [set up an AWS transit gateway](#).

If you can't set up a transit gateway, private IP addresses are available for NAS clients that are outside the VPC. However, these IP addresses are static—they can't failover between nodes.

You should review requirements for floating IP addresses and route tables before you deploy an HA configuration across multiple Availability Zones. You must specify the floating IP addresses when you deploy the configuration. The private IP addresses are automatically created by Cloud Manager.

For details, see [AWS networking requirements for Cloud Volumes ONTAP HA in multiple AZs](#).

iSCSI data access

Cross-VPC data communication is not an issue since iSCSI does not use floating IP addresses.

Takeover and giveback for iSCSI

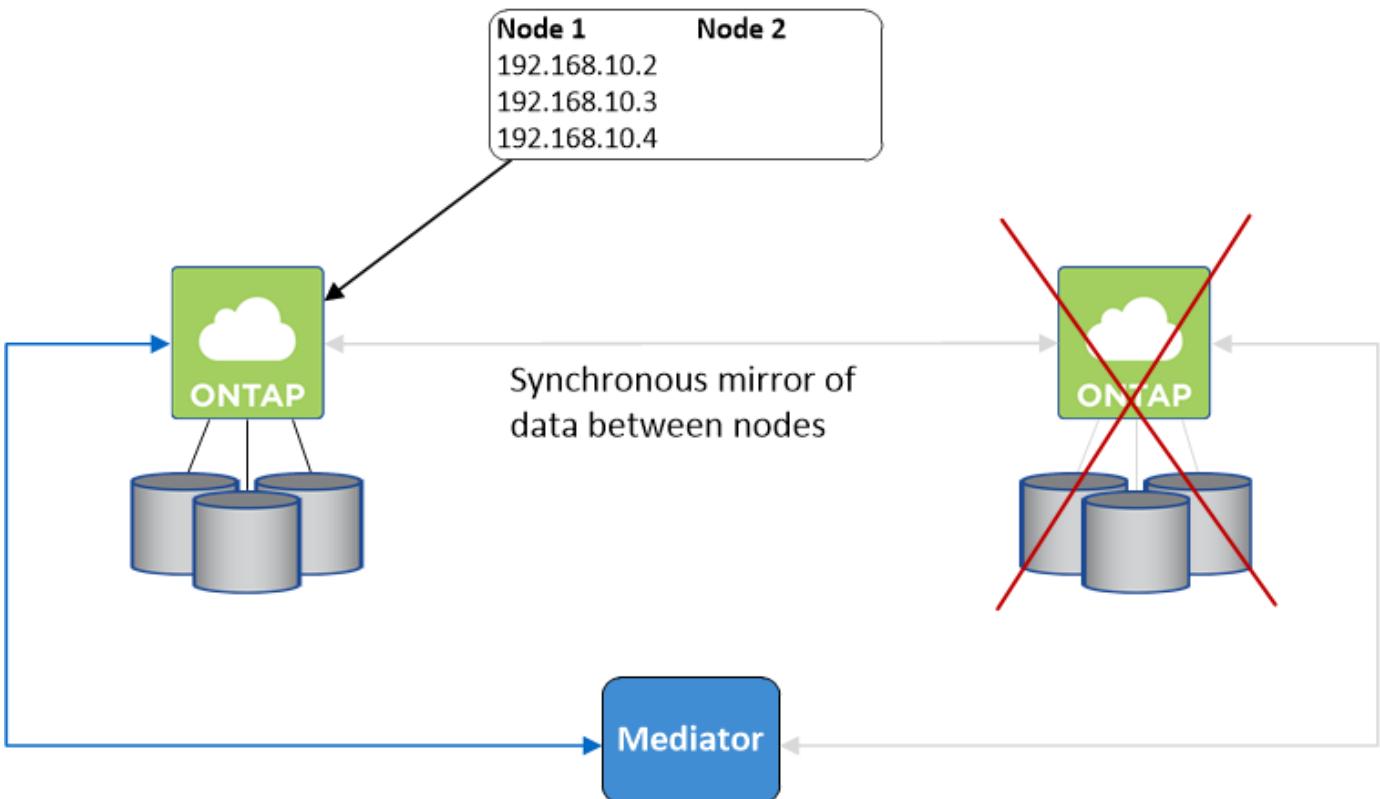
For iSCSI, Cloud Volumes ONTAP uses multipath I/O (MPIO) and Asymmetric Logical Unit Access (ALUA) to manage path failover between the active-optimized and non-optimized paths.



For information about which specific host configurations support ALUA, see the [NetApp Interoperability Matrix Tool](#) and the Host Utilities Installation and Setup Guide for your host operating system.

Takeover and giveback for NAS

When takeover occurs in a NAS configuration using floating IPs, the node's floating IP address that clients use to access data moves to the other node. The following image depicts storage takeover in a NAS configuration using floating IPs. If node 2 goes down, the floating IP address for node 2 moves to node 1.



NAS data IPs used for external VPC access cannot migrate between nodes if failures occur. If a node goes offline, you must manually remount volumes to clients outside the VPC by using the IP address on the other node.

After the failed node comes back online, remount clients to volumes using the original IP address. This step is needed to avoid transferring unnecessary data between two HA nodes, which can cause significant performance and stability impact.

You can easily identify the correct IP address from Cloud Manager by selecting the volume and clicking **Mount Command**.

Cloud Volumes ONTAP HA in a single Availability Zone

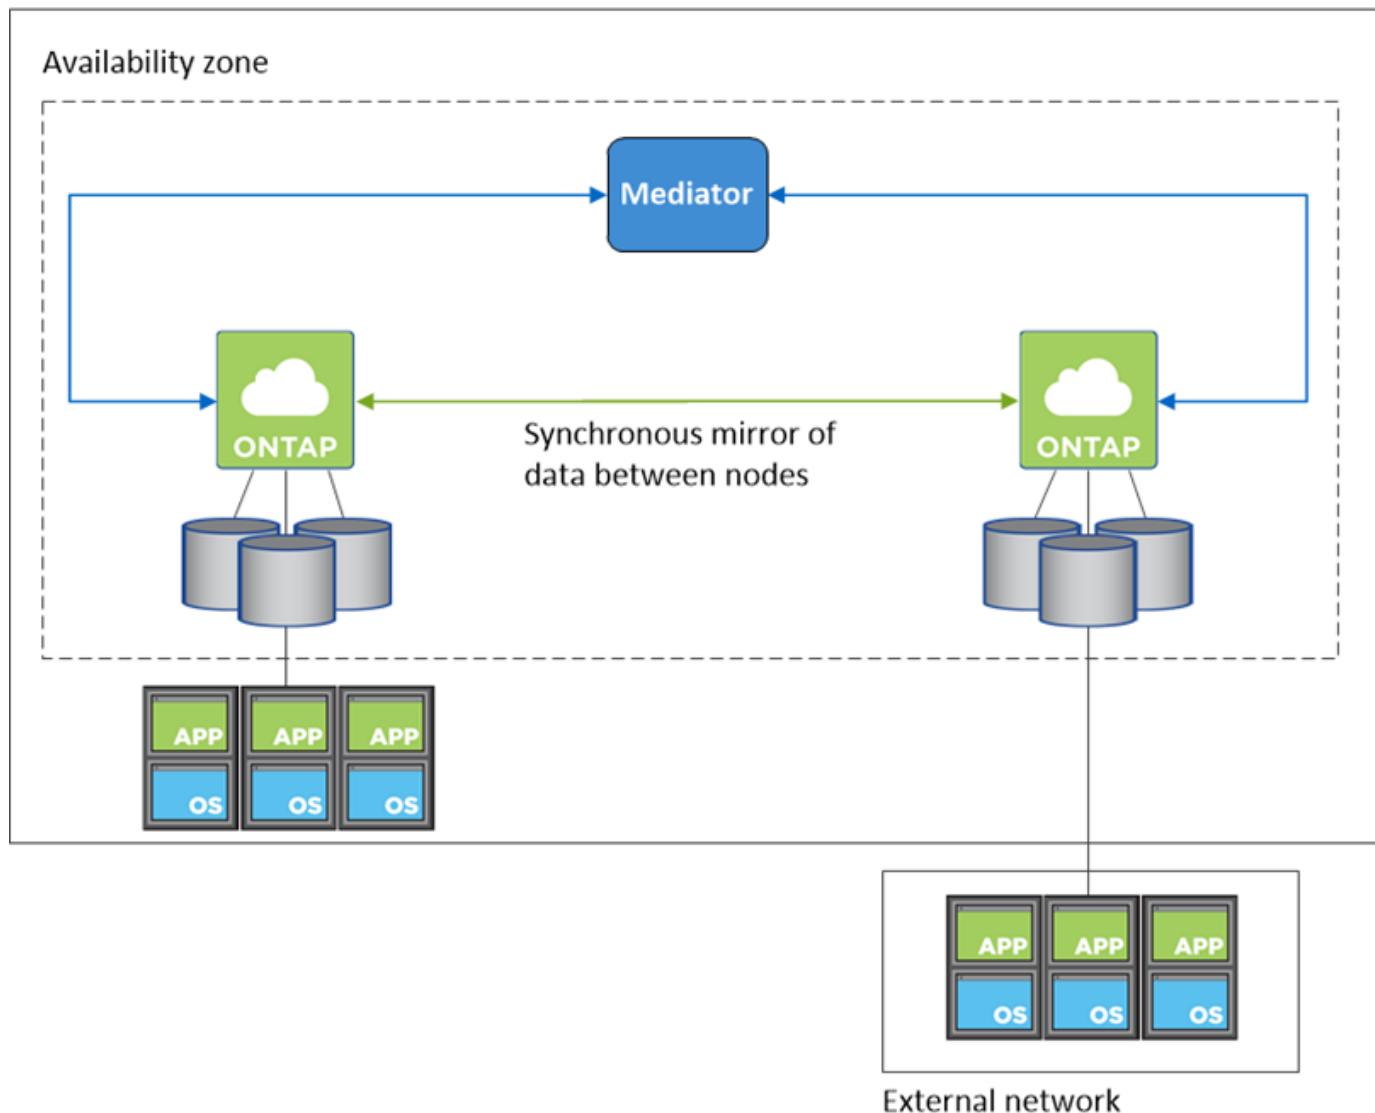
Deploying an HA configuration in a single Availability Zone (AZ) can ensure high availability of your data if an instance that runs a Cloud Volumes ONTAP node fails. All data is natively accessible from outside of the VPC.

i Cloud Manager creates an [AWS spread placement group](#) and launches the two HA nodes in that placement group. The placement group reduces the risk of simultaneous failures by spreading the instances across distinct underlying hardware. This feature improves redundancy from a compute perspective and not from disk failure perspective.

Data access

Because this configuration is in a single AZ, it does not require floating IP addresses. You can use the same IP address for data access from within the VPC and from outside the VPC.

The following image shows an HA configuration in a single AZ. Data is accessible from within the VPC and from outside the VPC.



Takeover and giveback

For iSCSI, Cloud Volumes ONTAP uses multipath I/O (MPIO) and Asymmetric Logical Unit Access (ALUA) to manage path failover between the active-optimized and non-optimized paths.



For information about which specific host configurations support ALUA, see the [NetApp Interoperability Matrix Tool](#) and the Host Utilities Installation and Setup Guide for your host operating system.

For NAS configurations, the data IP addresses can migrate between HA nodes if failures occur. This ensures client access to storage.

How storage works in an HA pair

Unlike an ONTAP cluster, storage in a Cloud Volumes ONTAP HA pair is not shared between nodes. Instead, data is synchronously mirrored between the nodes so that the data is available in the event of failure.

Storage allocation

When you create a new volume and additional disks are required, Cloud Manager allocates the same number of disks to both nodes, creates a mirrored aggregate, and then creates the new volume. For example, if two disks are required for the volume, Cloud Manager allocates two disks per node for a total of four disks.

Storage configurations

You can use an HA pair as an active-active configuration, in which both nodes serve data to clients, or as an active-passive configuration, in which the passive node responds to data requests only if it has taken over storage for the active node.



You can set up an active-active configuration only when using Cloud Manager in the Storage System View.

Performance expectations

A Cloud Volumes ONTAP HA configuration synchronously replicates data between nodes, which consumes network bandwidth. As a result, you can expect the following performance in comparison to a single-node Cloud Volumes ONTAP configuration:

- For HA configurations that serve data from only one node, read performance is comparable to the read performance of a single-node configuration, whereas write performance is lower.
- For HA configurations that serve data from both nodes, read performance is higher than the read performance of a single-node configuration, and write performance is the same or higher.

For more details about Cloud Volumes ONTAP performance, see [Performance](#).

Client access to storage

Clients should access NFS and CIFS volumes by using the data IP address of the node on which the volume resides. If NAS clients access a volume by using the IP address of the partner node, traffic goes between both nodes, which reduces performance.

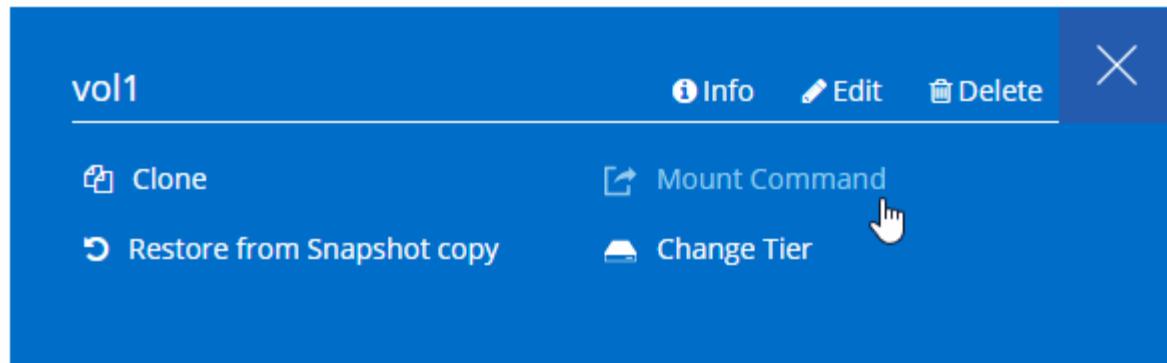


If you move a volume between nodes in an HA pair, you should remount the volume by using the IP address of the other node. Otherwise, you can experience reduced performance. If clients support NFSv4 referrals or folder redirection for CIFS, you can enable those features on the Cloud Volumes ONTAP systems to avoid remounting the volume. For details, see ONTAP documentation.

You can easily identify the correct IP address from Cloud Manager:

Volumes

2 Volumes | 0.22 TB Allocated | < 0.01 TB Used (0 TB in S3)

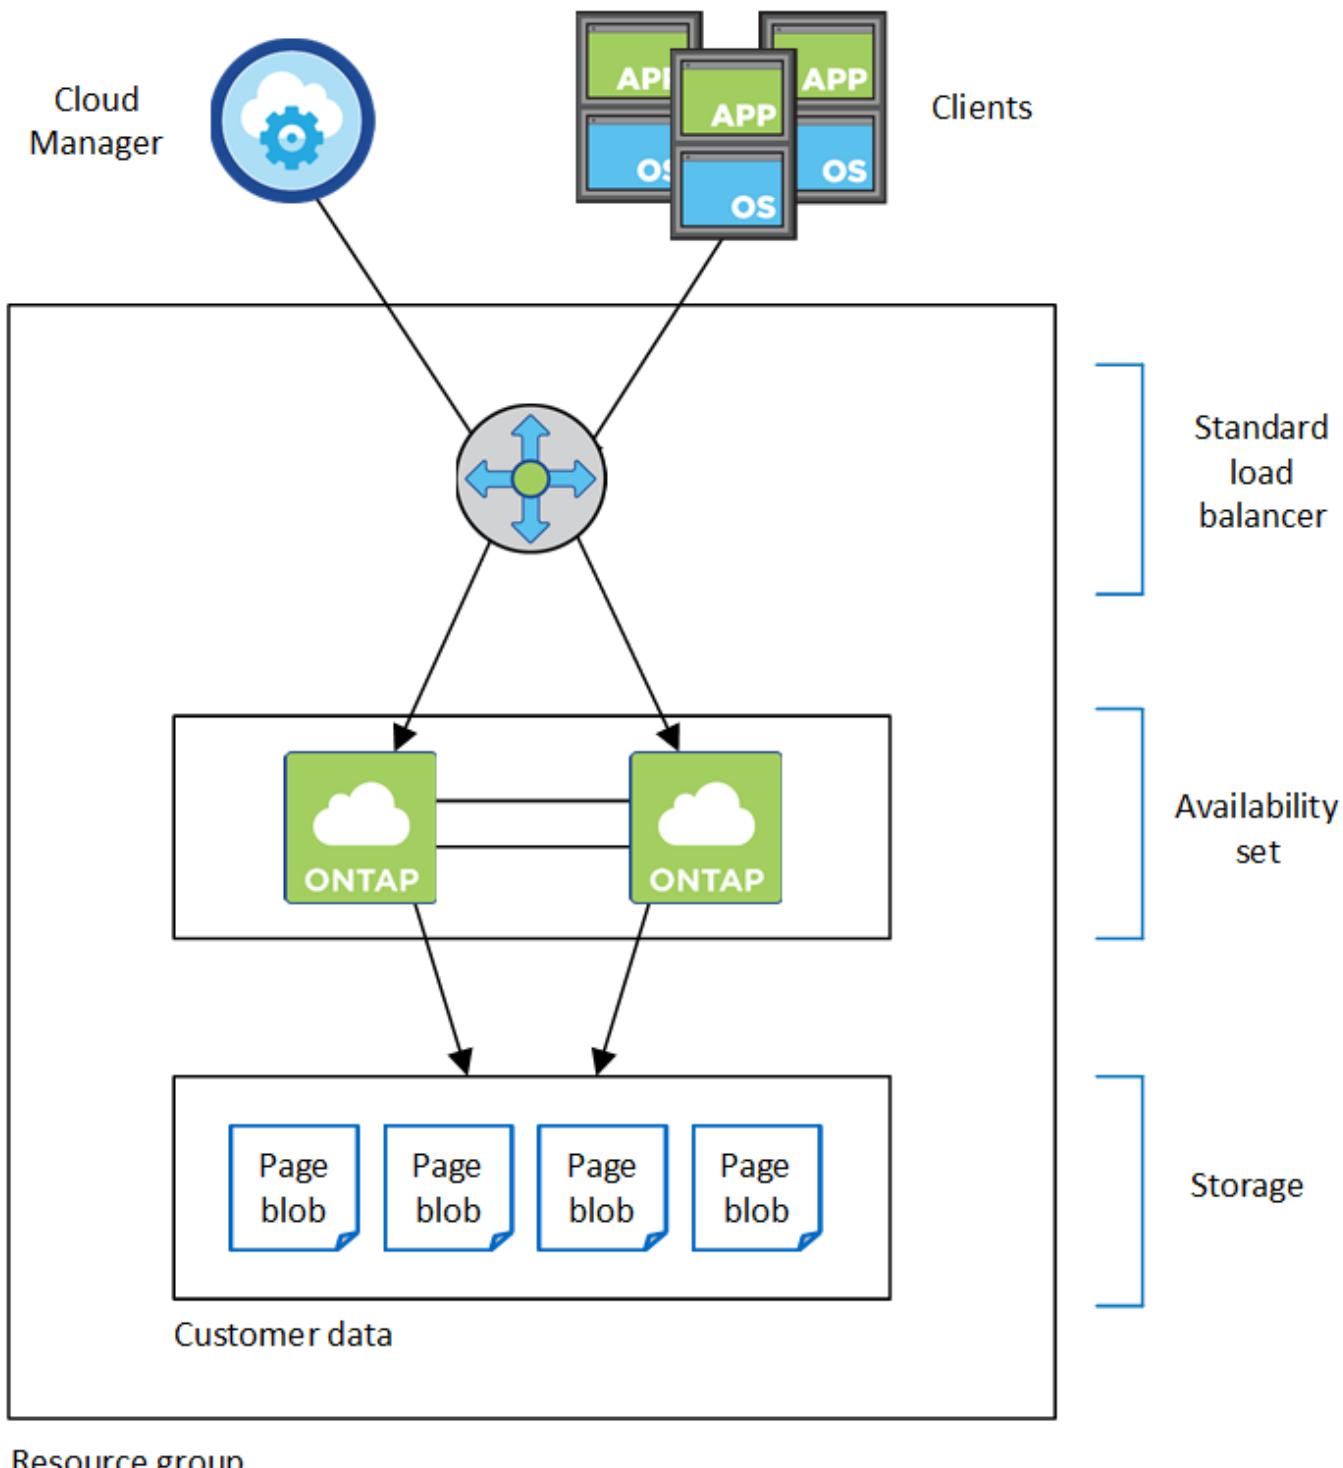


High-availability pairs in Azure

A Cloud Volumes ONTAP high availability (HA) pair provides enterprise reliability and continuous operations in case of failures in your cloud environment. In Azure, storage is shared between the two nodes.

HA components

A Cloud Volumes ONTAP HA configuration in Azure includes the following components:



Resource group

Note the following about the Azure components that Cloud Manager deploys for you:

Azure Standard Load Balancer

The load balancer manages incoming traffic to the Cloud Volumes ONTAP HA pair.

Availability Set

The Azure Availability Set is a logical grouping of the Cloud Volumes ONTAP nodes. The Availability Set ensures that the nodes are in different fault and update domains to provide redundancy and availability. [Learn more about Availability Sets in the Azure docs.](#)

Disks

Customer data resides on Premium Storage page blobs. Each node has access to the other node's storage. Additional storage is also required for [boot, root, and core data](#).

Storage accounts

- One storage account is required for managed disks.
- One or more storage accounts are required for the Premium Storage page blobs, as the disk capacity limit per storage account is reached.

[Azure documentation: Azure Storage scalability and performance targets for storage accounts](#).

- One storage account is required for data tiering to Azure Blob storage.
- Starting with Cloud Volumes ONTAP 9.7, the storage accounts that Cloud Manager creates for HA pairs are general-purpose v2 storage accounts.
- You can enable an HTTPS connection from a Cloud Volumes ONTAP 9.7 HA pair to Azure storage accounts when creating a working environment. Note that enabling this option can impact write performance. You can't change the setting after you create the working environment.

RPO and RTO

An HA configuration maintains high availability of your data as follows:

- The recovery point objective (RPO) is 0 seconds.
Your data is transactionally consistent with no data loss.
- The recovery time objective (RTO) is 60 seconds.
In the event of an outage, data should be available in 60 seconds or less.

Storage takeover and giveback

Similar to a physical ONTAP cluster, storage in an Azure HA pair is shared between nodes. Connections to the partner's storage allows each node to access the other's storage in the event of a *takeover*. Network path failover mechanisms ensure that clients and hosts continue to communicate with the surviving node. The partner *gives back* storage when the node is brought back on line.

For NAS configurations, data IP addresses automatically migrate between HA nodes if failures occur.

For iSCSI, Cloud Volumes ONTAP uses multipath I/O (MPIO) and Asymmetric Logical Unit Access (ALUA) to manage path failover between the active-optimized and non-optimized paths.



For information about which specific host configurations support ALUA, see the [NetApp Interoperability Matrix Tool](#) and the Host Utilities Installation and Setup Guide for your host operating system.

Storage takeover, resync, and giveback are all automatic by default. No user action is required.

Storage configurations

You can use an HA pair as an active-active configuration, in which both nodes serve data to clients, or as an active-passive configuration, in which the passive node responds to data requests only if it has taken over storage for the active node.

HA limitations

The following limitations affect Cloud Volumes ONTAP HA pairs in Azure:

- HA pairs are supported with Cloud Volumes ONTAP Standard, Premium, and BYOL. Explore is not supported.
- NFSv4 is not supported. NFSv3 is supported.
- HA pairs are not supported in some regions.

[See the list of supported Azure regions.](#)

[Learn how to deploy an HA system in Azure.](#)

High-availability pairs in Google Cloud Platform

A Cloud Volumes ONTAP high availability (HA) configuration provides nondisruptive operations and fault tolerance. In Google Cloud Platform, data is synchronously mirrored between the two nodes.

HA components

Cloud Volumes ONTAP HA configurations in GCP include the following components:

- Two Cloud Volumes ONTAP nodes whose data is synchronously mirrored between each other.
- A mediator instance that provides a communication channel between the nodes to assist in storage takeover and giveback processes.

The mediator runs the Linux operating system on a f1-micro instance and uses two standard persistent disks that are 10 GB each.

- One zone or three zones (recommended).

If you choose three zones, the two nodes and mediator are in separate Google Cloud zones.

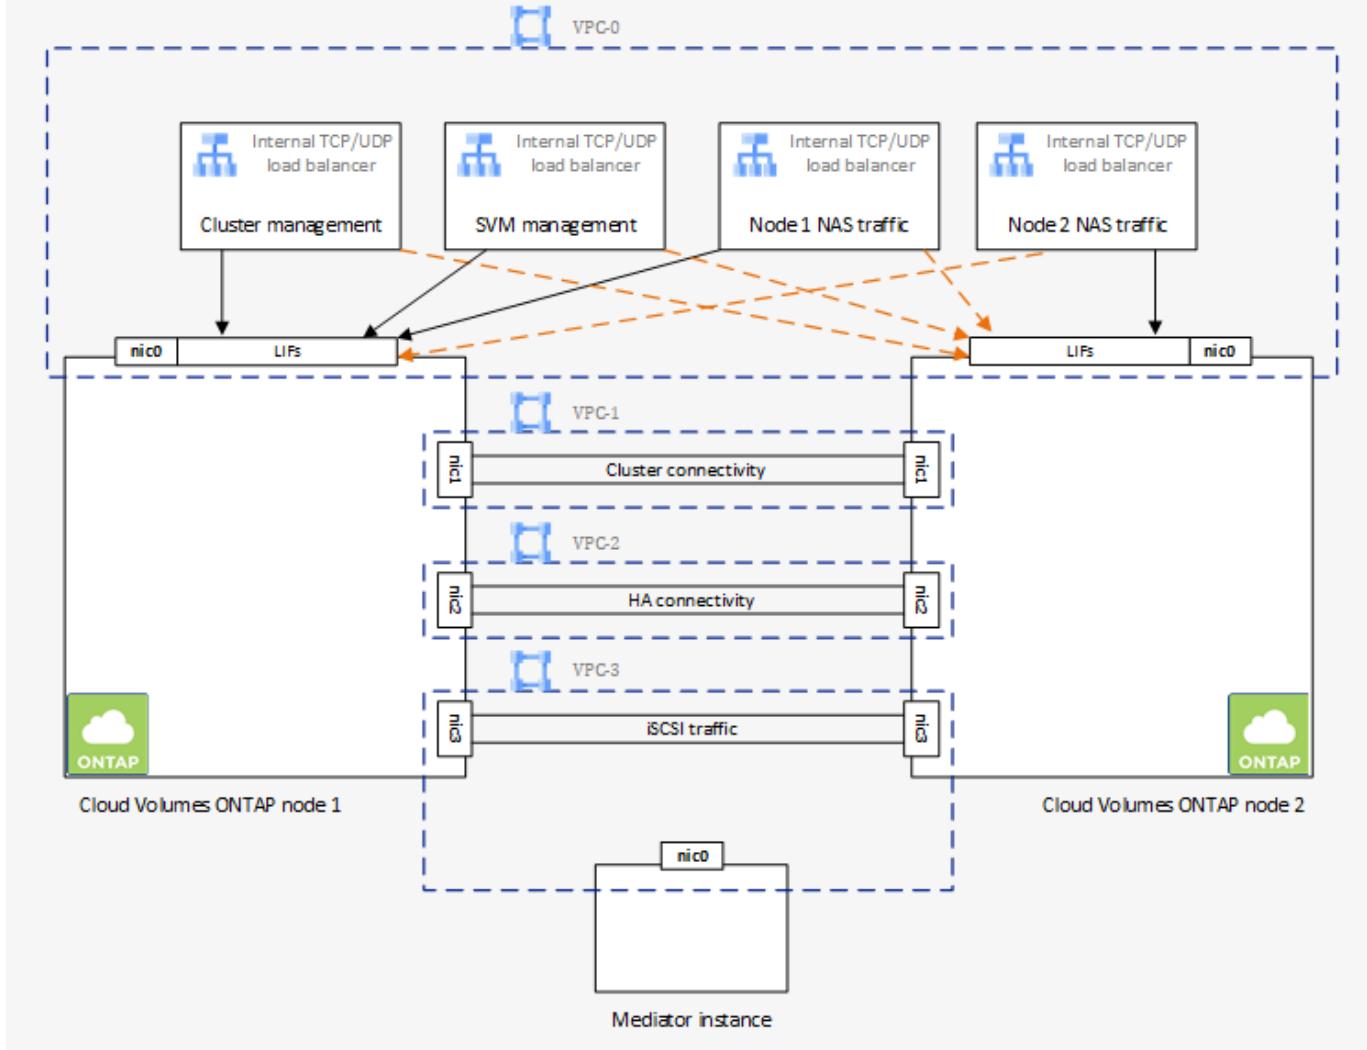
- Four Virtual Private Clouds (VPCs).

The configuration uses four VPCs because GCP requires that each network interface resides in a separate VPC network.

- Four Google Cloud internal load balancers (TCP/UDP) that manage incoming traffic to the Cloud Volumes ONTAP HA pair.

[Learn about networking requirements](#), including more details about load balancers, VPCs, internal IP addresses, subnets, and more.

The following conceptual image shows a Cloud Volumes ONTAP HA pair and its components:



Storage takeover and giveback

If a node goes down, the other node can serve data for its partner to provide continued data service. Clients can access the same data from the partner node because the data was synchronously mirrored to the partner.

After the node reboots, the partner must resync data before it can return the storage. The time that it takes to resync data depends on how much data was changed while the node was down.

Storage takeover, resync, and giveback are all automatic by default. No user action is required.

RPO and RTO

An HA configuration maintains high availability of your data as follows:

- The recovery point objective (RPO) is 0 seconds.

Your data is transactionally consistent with no data loss.

- The recovery time objective (RTO) is 60 seconds.

In the event of an outage, data should be available in 60 seconds or less.

HA deployment models

You can ensure the high availability of your data by deploying an HA configuration in multiple zones or in a single zone.

Multiple zones (recommended)

Deploying an HA configuration across three zones ensures continuous data availability if a failure occurs within a zone. Note that write performance is slightly lower compared to using a single zone, but it's minimal.

Single zone

When deployed in a single zone, a Cloud Volumes ONTAP HA configuration uses a spread placement policy. This policy ensures that an HA configuration is protected from a single point of failure within the zone, without having to use separate zones to achieve fault isolation.

This deployment model does lower your costs because there are no data egress charges between zones.

How storage works in an HA pair

Unlike an ONTAP cluster, storage in a Cloud Volumes ONTAP HA pair in GCP is not shared between nodes. Instead, data is synchronously mirrored between the nodes so that the data is available in the event of failure.

Storage allocation

When you create a new volume and additional disks are required, Cloud Manager allocates the same number of disks to both nodes, creates a mirrored aggregate, and then creates the new volume. For example, if two disks are required for the volume, Cloud Manager allocates two disks per node for a total of four disks.

Storage configurations

You can use an HA pair as an active-active configuration, in which both nodes serve data to clients, or as an active-passive configuration, in which the passive node responds to data requests only if it has taken over storage for the active node.

Performance expectations for an HA configuration

A Cloud Volumes ONTAP HA configuration synchronously replicates data between nodes, which consumes network bandwidth. As a result, you can expect the following performance in comparison to a single-node Cloud Volumes ONTAP configuration:

- For HA configurations that serve data from only one node, read performance is comparable to the read performance of a single-node configuration, whereas write performance is lower.
- For HA configurations that serve data from both nodes, read performance is higher than the read performance of a single-node configuration, and write performance is the same or higher.

For more details about Cloud Volumes ONTAP performance, see [Performance](#).

Client access to storage

Clients should access NFS and CIFS volumes by using the data IP address of the node on which the volume resides. If NAS clients access a volume by using the IP address of the partner node, traffic goes between both nodes, which reduces performance.



If you move a volume between nodes in an HA pair, you should remount the volume by using the IP address of the other node. Otherwise, you can experience reduced performance. If clients support NFSv4 referrals or folder redirection for CIFS, you can enable those features on the Cloud Volumes ONTAP systems to avoid remounting the volume. For details, see ONTAP documentation.

You can easily identify the correct IP address from Cloud Manager:

Volumes

2 Volumes | 0.22 TB Allocated | < 0.01 TB Used (0 TB in S3)

The screenshot shows a Cloud Manager interface for managing volumes. At the top, it displays '2 Volumes | 0.22 TB Allocated | < 0.01 TB Used (0 TB in S3)'. Below this, a volume named 'vol1' is listed. To the right of 'vol1', there are three buttons: 'Info' (with a blue info icon), 'Edit' (with a blue edit icon), and 'Delete' (with a blue trash icon). Below these buttons is a close button (an 'X'). Underneath the volume name, there are several actions: 'Clone' (with a blue clone icon), 'Mount Command' (with a blue icon of a drive and a right-pointing arrow), 'Restore from Snapshot copy' (with a blue circular arrow icon), and 'Change Tier' (with a blue icon of a hard drive). A hand cursor icon is positioned over the 'Mount Command' button.

Related links

- [Learn about networking requirements](#)
- [Learn how to get started in GCP](#)

Cloud Volumes ONTAP licensing

Several licensing options are available for Cloud Volumes ONTAP. Each option enables you to choose a configuration that meets your needs.

Licensing overview

The following table provides an overview of the licensing options for Cloud Volumes ONTAP.

Beyond these licensing options, you can also choose the *Freemium* offering to get started with Cloud Volumes ONTAP without purchasing a license or contract.

Charging method	Highlights	Support	Max system capacity
Capacity-based license: Essentials package	<ul style="list-style-type: none"> • Pay per TiB of capacity for one or more Cloud Volumes ONTAP systems • Provides a la carte licensing for Cloud Volumes ONTAP: <ul style="list-style-type: none"> ◦ A single node or HA system ◦ File and block storage or secondary data (DR) • Available by bringing your own license (BYOL) purchased from NetApp 	Included	2 PB
Capacity-based license: Professional package	<ul style="list-style-type: none"> • Pay per TiB of capacity for one or more Cloud Volumes ONTAP systems • Provides licensing for any Cloud Volumes ONTAP configuration (single node or HA with any storage type) • Includes volume backups using the Cloud Backup Service (only for volumes charged against this license) • Available through an AWS Marketplace annual contract or by bringing your own license (BYOL) purchased from NetApp 	Included	2 PB
PAYGO by node	<ul style="list-style-type: none"> • Pay-as-you-go by the hour through a marketplace subscription from your cloud provider • Charging is per Cloud Volumes ONTAP node • Available in three licensing options: Explore, Standard, and Premium 	Included, but you must activate support	<ul style="list-style-type: none"> • Explore: 2 TiB • Standard: 10 TiB • Premium: 368 TiB
Node-based license	<ul style="list-style-type: none"> • The previous generation BYOL for Cloud Volumes ONTAP • A node-based license is available for license renewals only 	Included	368 TiB per license

The following sections provide more details about each of these options.

Freemium offering

- A new offering that provides all Cloud Volumes ONTAP features free of charge from NetApp (cloud provider charges still apply).
- No license or contract is needed.
- Support is not included.

- You're limited to 500 GiB of provisioned capacity per Cloud Volumes ONTAP system.
- You can use up to 10 Cloud Volumes ONTAP systems with the Freemium offering per NetApp account.
- If the provisioned capacity for a Cloud Volumes ONTAP system exceeds 500 GiB, Cloud Manager converts the system to the Essentials package (which is a capacity-based license) and charging starts.

Any other systems that have less than 500 GiB of provisioned capacity stay on the Freemium offering (as long as they were deployed using the Freemium offering).

To get started with the Freemium offering, create a new Cloud Volumes ONTAP working environment and select **Freemium** when prompted to choose a charging method.

Capacity-based licenses

Capacity-based licensing enables you to pay for Cloud Volumes ONTAP per TiB of capacity. The license is associated with your NetApp account and enables you to charge multiple systems against the license, as long as enough capacity is available through the license.

For example, you could purchase a single 20 TiB license, deploy four Cloud Volumes ONTAP systems, and then allocate a 5 TiB volume to each system, for a total of 20 TiB.

Unlike the by-node charging method where a license is purchased per Cloud Volumes ONTAP system, a capacity-based license is issued to a NetApp account. The capacity is then available to the volumes on each Cloud Volumes ONTAP system deployed in that account.

Capacity-based licensing is available in the form of a *package*. When you deploy a Cloud Volumes ONTAP system, you can choose from the following packages: Essentials or Professional.

This licensing method is available for Cloud Volumes ONTAP 9.7 and later.



For each package, there is a minimum 4 TiB capacity charge. Any Cloud Volumes ONTAP instance that has less than 4 TiB of capacity will be charged at a rate of 4 TiB.

Essentials package

- Provides a la carte licensing for Cloud Volumes ONTAP:
 - A single node or HA system
 - File and block storage or secondary data for disaster recovery (DR)
- This package is available as a license (BYOL) purchased from NetApp.
- Support is included for the length of the subscription term.
- Conversions to another licensing option isn't supported.
- Each individual Cloud Volumes ONTAP system supports up to 2 PB of capacity through disks and tiering to object storage.

Professional package

- Provides licensing for any Cloud Volumes ONTAP configuration (single node or HA with any storage type).
- Includes volume backups using the Cloud Backup Service (only for volumes charged against this license).
- This package is available as an annual contract from the AWS Marketplace or as a license (BYOL) purchased from NetApp.

If you have an AWS Marketplace contract, *all* Cloud Volumes ONTAP systems that you deploy are charged against that contract. You can't mix and match a Marketplace contract with BYOL.

- Support is included for the length of the subscription term.
- Conversions to another licensing option isn't supported.
- Each individual Cloud Volumes ONTAP system supports up to 2 PB of capacity through disks and tiering to object storage.

To get started with a capacity-based license, [Contact NetApp Sales](#) and then [add your license to Cloud Manager](#).

PAYGO by node

- Requires a subscription from a cloud provider's marketplace for pay-as-you-go pricing at an hourly rate.
- Charging is per Cloud Volumes ONTAP node.
- Offers Cloud Volumes ONTAP in three different licensing options: Explore, Standard, and Premium. Each license provides support for different amounts of storage and compute.
- A 30-day free trial is available for the first Cloud Volumes ONTAP system that you deploy in a cloud provider. [Learn more about 30-day free trials](#).
 - There are no hourly software charges, but cloud provider infrastructure charges still apply (compute, storage, and networking).
 - When the free trial ends, you'll be charged hourly according to the selected license, as long as you subscribed. If you haven't subscribed, the system shuts down.

Cloud Manager prompts you to subscribe to your cloud provider's marketplace when you create a Cloud Volumes ONTAP system.

- Conversions to another licensing option isn't supported.
- Basic technical support is offered, but you must [register and activate the NetApp serial number associated with your system](#).

You can view pricing details from your cloud provider's marketplace:

- [AWS Marketplace](#)
- [Azure Marketplace](#)
- [Google Cloud Platform Marketplace](#)

To get started with PAYGO, create a Cloud Volumes ONTAP working environment and subscribe to your cloud provider's marketplace when prompted.

Node-based licenses

- The previous generation BYOL for Cloud Volumes ONTAP.
- A node-based license is available for license renewals only.
- Each Cloud Volumes ONTAP system supports up to 368 TiB of capacity per license.
- Conversions to another licensing option isn't supported.

If you want to transition to capacity-based licensing, you can purchase a license, deploy a new Cloud Volumes ONTAP system, and then replicate the data to that new system.

License management for node-based BYOL

Each Cloud Volumes ONTAP system that has a node-based BYOL must have a system license installed with an active subscription. Cloud Manager simplifies the process by managing licenses for you and by displaying a warning before they expire.

[Learn more about Cloud Volumes ONTAP licensing options.](#)

BYOL system licenses

A node-based license provides up to 368 TiB of capacity for a single node or HA pair.

You can purchase multiple licenses for a Cloud Volumes ONTAP BYOL system to allocate more than 368 TiB of capacity. For example, you might purchase two licenses to allocate up to 736 TiB of capacity to Cloud Volumes ONTAP. Or you could purchase four licenses to get up to 1.4 PB.

The number of licenses that you can purchase for a single node system or HA pair is unlimited.



Some on-premises ONTAP storage systems that you purchased may have included a free Cloud Volumes ONTAP license. You can use the license to create a new Cloud Volumes ONTAP system, or you can apply the license to an existing Cloud Volumes ONTAP system to expand the capacity. [See if you have any available licenses to use.](#)

Be aware that disk limits can prevent you from reaching the capacity limit by using disks alone. You can go beyond the disk limit by [tiering inactive data to object storage](#). For information about disk limits, refer to [storage limits in the Cloud Volumes ONTAP Release Notes](#).

License management for a new system

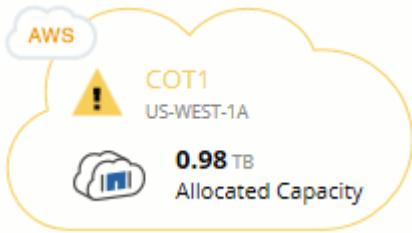
When you create a node-based BYOL system, Cloud Manager prompts you for the serial number of your license and your NetApp Support Site account. Cloud Manager uses the account to download the license file from NetApp and to install it on the Cloud Volumes ONTAP system.

[Learn how to add NetApp Support Site accounts to Cloud Manager.](#)

If Cloud Manager can't access the license file over the secure internet connection, you can [obtain the file yourself and then manually upload the file to Cloud Manager](#).

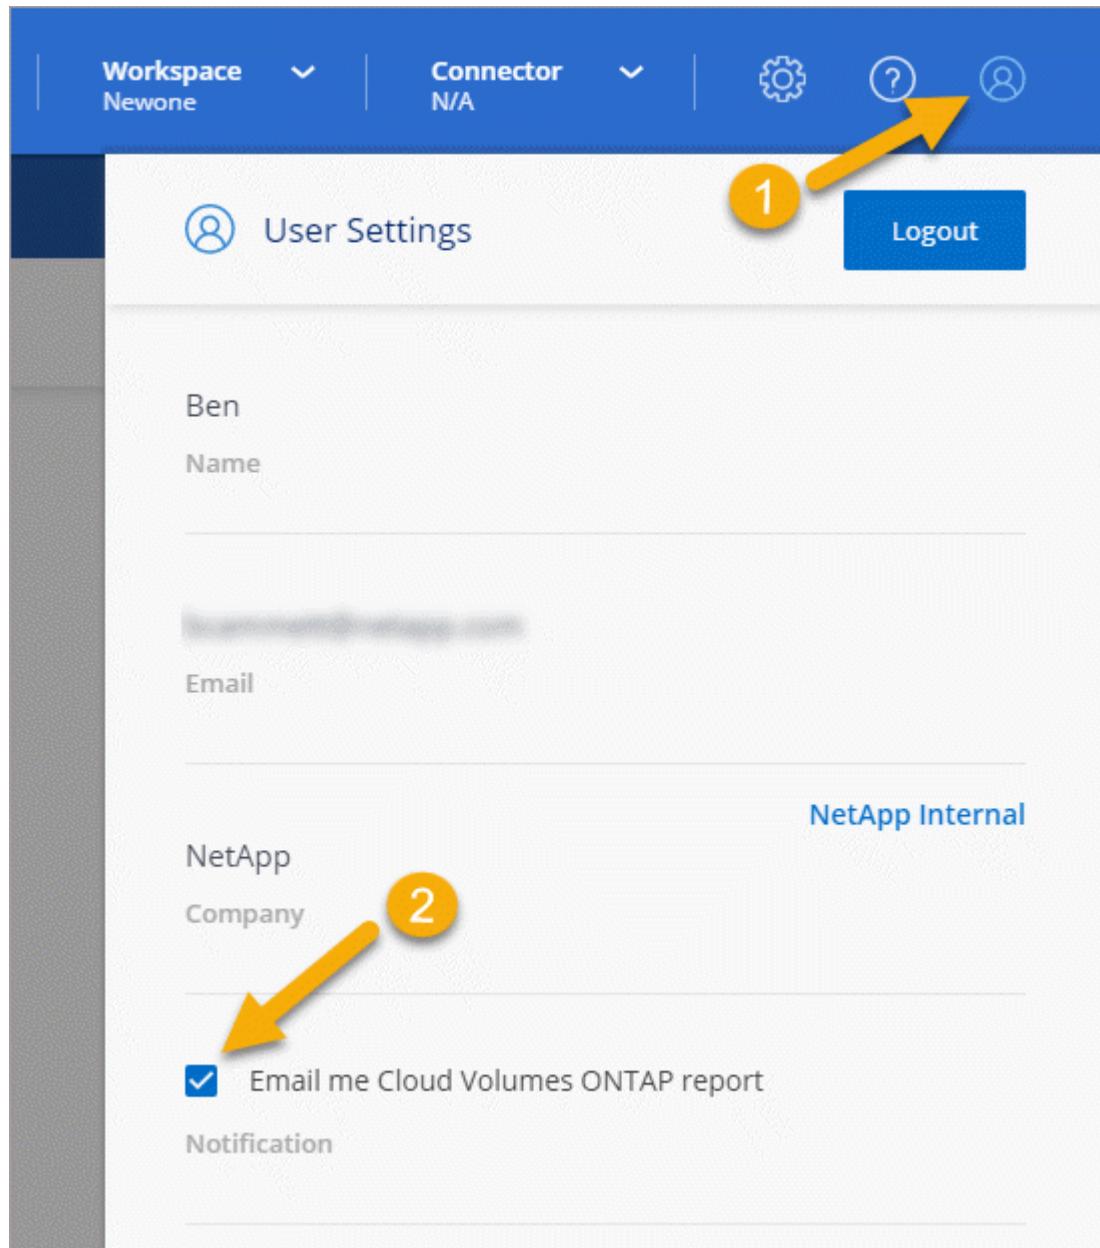
License expiration

Cloud Manager displays a warning 30 days before a node-based license is due to expire and again when the license expires. The following image shows a 30-day expiration warning that appears in the user interface:



You can select the working environment to review the message.

Cloud Manager includes a license expiration warning in the Cloud Volumes ONTAP report that's emailed to you, if you are an Account Admin and you enabled the option:



The emailed report includes the license expiration warning every 2 weeks.

If you don't renew the license in time, the Cloud Volumes ONTAP system shuts itself down. If you restart it, it shuts itself down again.

License renewal

When you renew a node-based BYOL subscription by contacting a NetApp representative, Cloud Manager automatically obtains the new license from NetApp and installs it on the Cloud Volumes ONTAP system.

If Cloud Manager can't access the license file over the secure internet connection, you can [obtain the file yourself and then manually upload the file to Cloud Manager](#).

License transfer to a new system

A node-based BYOL license is transferable between Cloud Volumes ONTAP systems when you delete an existing system and then create a new one using the same license.

For example, you might want to delete an existing licensed system and then use the license with a new BYOL system in a different VPC/VNet or cloud provider. Note that only *cloud-agnostic* serial numbers work in any cloud provider. Cloud-agnostic serial numbers start with the 908xxxx prefix.

It's important to note that your BYOL license is tied to your company and a specific set of NetApp Support Site credentials.

Evaluating

You can evaluate Cloud Volumes ONTAP before you pay for the software. Three options are available: the Freemium offering, a 30-day free trial, and an evaluation license.

If you need assistance with your proof of concept, contact [the Sales team](#) or reach out through the chat option available from [NetApp Cloud Central](#) and from within Cloud Manager.

Freemium offering

The Freemium offering enables you to use all Cloud Volumes ONTAP features free of charge from NetApp (cloud provider charges still apply). You're limited to 500 GiB of provisioned capacity per node and there's no support contract. You can have up to 10 Freemium systems.

[Learn more about the Freemium offering.](#)

30-day free trials for PAYGO

A 30-day free trial is available if you plan to pay for Cloud Volumes ONTAP as you go. You can start a 30-day free trial of Cloud Volumes ONTAP from Cloud Manager by creating your first Cloud Volumes ONTAP system in a payer's account.

There are no hourly software license charges for the instance, but infrastructure charges from your cloud provider still apply.

A free trial automatically converts to a paid hourly subscription when it expires. If you terminate the instance within the time limit, the next instance that you deploy is not part of the free trial (even if it's deployed within those 30 days).

Pay-as-you-go trials are awarded through a cloud provider and are not extendable by any means.

Evaluation license for node-based licensing

An evaluation BYOL license is an option for customers who expect to pay for Cloud Volumes ONTAP by purchasing a termed license from NetApp. You can obtain an evaluation license from your account team, your Sales Engineer, or your partner.

The evaluation key is good for 30 days, and can be used multiple times, each for 30 days (regardless of the creation day).

At the end of 30 days, daily shutdowns will occur, so it's best to plan ahead. You can apply a new BYOL license on top of the evaluation license for an in-place upgrade (this requires a restart of single node systems). Your hosted data is **not** deleted at the end of the trial period.



You can't upgrade Cloud Volumes ONTAP software when using an evaluation license.

Security

Cloud Volumes ONTAP supports data encryption and provides protection against viruses and ransomware.

Encryption of data at rest

Cloud Volumes ONTAP supports the following encryption technologies:

- NetApp encryption solutions (NVE and NAE)
- AWS Key Management Service
- Azure Storage Service Encryption
- Google Cloud Platform default encryption

You can use NetApp encryption solutions with native encryption from AWS, Azure, or GCP, which encrypt data at the hypervisor level. Doing so would provide double encryption, which might be desired for very sensitive data. When the encrypted data is accessed, it's unencrypted twice—once at the hypervisor-level (using keys from the cloud provider) and then again using NetApp encryption solutions (using keys from an external key manager).

NetApp encryption solutions (NVE and NAE)

Cloud Volumes ONTAP supports both NetApp Volume Encryption (NVE) and NetApp Aggregate Encryption (NAE) with an external key manager. NVE and NAE are software-based solutions that enable (FIPS) 140-2-compliant data-at-rest encryption of volumes.

- NVE encrypts data at rest one volume at a time. Each data volume has its own unique encryption key.
- NAE is an extension of NVE—it encrypts data for each volume, and the volumes share a key across the aggregate. NAE also allows common blocks across all volumes in the aggregate to be deduplicated.

Both NVE and NAE use AES 256-bit encryption.

[Learn more about NetApp Volume Encryption and NetApp Aggregate Encryption.](#)

Starting with Cloud Volumes ONTAP 9.7, new aggregates will have NetApp Aggregate Encryption (NAE) enabled by default after you set up an external key manager. New volumes that aren't part of an NAE aggregate will have NetApp Volume Encryption (NVE) enabled by default (for example, if you have existing aggregates that were created before setting up an external key manager).

Setting up a supported key manager is the only required step. For set up instructions, see [Encrypting volumes with NetApp encryption solutions](#).

AWS Key Management Service

When you launch a Cloud Volumes ONTAP system in AWS, you can enable data encryption using the [AWS Key Management Service \(KMS\)](#). Cloud Manager requests data keys using a customer master key (CMK).



You can't change the AWS data encryption method after you create a Cloud Volumes ONTAP system.

If you want to use this encryption option, then you must ensure that the AWS KMS is set up appropriately. For details, see [Setting up the AWS KMS](#).

Azure Storage Service Encryption

[Azure Storage Service Encryption](#) for data at rest is enabled by default for Cloud Volumes ONTAP data in Azure. No setup is required.

You can encrypt Azure managed disks on single node Cloud Volumes ONTAP systems using external keys from another account. This feature is supported using Cloud Manager APIs.

You just need to add the following to the API request when creating the single node system:

```
"azureEncryptionParameters": {  
    "key": <azure id of encryptionset>  
}
```



Customer-managed keys are not supported with Cloud Volumes ONTAP HA pairs.

Google Cloud Platform default encryption

[Google Cloud Platform data-at-rest encryption](#) is enabled by default for Cloud Volumes ONTAP. No setup is required.

While Google Cloud Storage always encrypts your data before it's written to disk, you can use Cloud Manager APIs to create a Cloud Volumes ONTAP system that uses *customer-managed encryption* keys. These are keys that you generate and manage in GCP using the Cloud Key Management Service. [Learn more](#).

ONTAP virus scanning

You can use integrated antivirus functionality on ONTAP systems to protect data from being compromised by viruses or other malicious code.

ONTAP virus scanning, called *Vscan*, combines best-in-class third-party antivirus software with ONTAP features that give you the flexibility you need to control which files get scanned and when.

For information about the vendors, software, and versions supported by Vscan, see the [NetApp Interoperability Matrix](#).

For information about how to configure and manage the antivirus functionality on ONTAP systems, see the [ONTAP 9 Antivirus Configuration Guide](#).

Ransomware protection

Ransomware attacks can cost a business time, resources, and reputation. Cloud Manager enables you to implement the NetApp solution for ransomware, which provides effective tools for visibility, detection, and remediation.

- Cloud Manager identifies volumes that are not protected by a Snapshot policy and enables you to activate the default Snapshot policy on those volumes.

Snapshot copies are read-only, which prevents ransomware corruption. They can also provide the

granularity to create images of a single file copy or a complete disaster recovery solution.

- Cloud Manager also enables you to block common ransomware file extensions by enabling ONTAP's FPolicy solution.

Ransomware Protection

Ransomware attacks can cost a business time, resources, and reputation. The NetApp solution for ransomware provides effective tools for visibility, detection, and remediation. [Learn More](#)

1 Enable Snapshot Copy Protection [i](#)

50 % Protection

1 Volumes without a Snapshot Policy

To protect your data, activate the default Snapshot policy for these volumes [i](#)

Activate Snapshot Policy

2 Block Ransomware File Extensions [i](#)

ONTAP's native FPolicy configuration monitors and blocks file operations based on a file's extension.

[View Denied File Names](#) [i](#)

Activate FPolicy

[Learn how to implement the NetApp solution for ransomware.](#)

Performance

You can review performance results to help you decide which workloads are appropriate for Cloud Volumes ONTAP.

- Cloud Volumes ONTAP for AWS

[NetApp Technical Report 4383: Performance Characterization of Cloud Volumes ONTAP in Amazon Web Services with Application Workloads.](#)

- Cloud Volumes ONTAP for Microsoft Azure

[NetApp Technical Report 4671: Performance Characterization of Cloud Volumes ONTAP in Azure with Application Workloads.](#)

- Cloud Volumes ONTAP for Google Cloud

[NetApp Technical Report 4816: Performance Characterization of Cloud Volumes ONTAP for Google Cloud.](#)

AutoSupport and Active IQ Digital Advisor

The AutoSupport component of ONTAP collects telemetry and sends it for analysis. Active IQ Digital Advisor analyzes the data from AutoSupport and provides proactive care and optimization. Using artificial intelligence, Active IQ can identify potential problems and help you resolve them before they impact your business.

Active IQ enables you to optimize your data infrastructure across your global hybrid cloud by delivering actionable predictive analytics and proactive support through a cloud-based portal and mobile app. Data-driven insights and recommendations from Active IQ are available to all NetApp customers with an active

SupportEdge contract (features vary by product and support tier).

Here are some things you can do with Active IQ:

- Plan upgrades.

Active IQ identifies issues in your environment that can be resolved by upgrading to a newer version of ONTAP and the Upgrade Advisor component helps you plan for a successful upgrade.

- View system wellness.

Your Active IQ dashboard reports any issues with wellness and helps you correct those issues. Monitor system capacity to make sure you never run out of storage space. View support cases for your system.

- Manage performance.

Active IQ shows system performance over a longer period than you can see in ONTAP System Manager. Identify configuration and system issues that are impacting your performance. Maximize efficiency. View storage efficiency metrics and identify ways to store more data in less space.

- View inventory and configuration.

Active IQ displays complete inventory and software and hardware configuration information. See when service contracts are expiring and renew them to ensure you remain supported.

Related information

- [NetApp Documentation: Active IQ Digital Advisor](#)
- [Launch Active IQ](#)
- [SupportEdge Services](#)

Default configuration for Cloud Volumes ONTAP

Understanding how Cloud Volumes ONTAP is configured by default can help you set up and administer your systems, especially if you are familiar with ONTAP because the default setup for Cloud Volumes ONTAP is different than ONTAP.

Defaults

- Cloud Volumes ONTAP is available as a single-node system and as an HA pair in AWS, Azure, and GCP.
- Cloud Manager creates one data-serving storage VM when it deploys Cloud Volumes ONTAP. Some configurations support additional storage VMs. [Learn more about managing storage VMs](#).

Starting with the Cloud Manager 3.9.5 release, logical space reporting is enabled on the initial storage VM. When space is reported logically, ONTAP reports the volume space such that all the physical space saved by the storage efficiency features are also reported as used.

- Cloud Manager automatically installs the following ONTAP feature licenses on Cloud Volumes ONTAP:
 - CIFS
 - FlexCache
 - FlexClone

- iSCSI
- NetApp Volume Encryption (only for BYOL or registered PAYGO systems)
- NFS
- SnapMirror
- SnapRestore
- SnapVault
- Several network interfaces are created by default:
 - A cluster management LIF
 - An intercluster LIF
 - An SVM management LIF on HA systems in Azure and in GCP, on single node systems in AWS, and optionally on HA systems in multiple AWS Availability Zones
 - A node management LIF (in GCP, this LIF is combined with the intercluster LIF)
 - An iSCSI data LIF
 - A CIFS and NFS data LIF



LIF failover is disabled by default for Cloud Volumes ONTAP due to EC2 requirements. Migrating a LIF to a different port breaks the external mapping between IP addresses and network interfaces on the instance, making the LIF inaccessible.

- Cloud Volumes ONTAP sends configuration backups to the Connector using HTTPS.

The backups are accessible from <https://ipaddress/occm/offboxconfig> where *ipaddress* is the IP address of the Connector host.

- Cloud Manager sets a few volume attributes differently than other management tools (System Manager or the CLI, for example).

The following table lists the volume attributes that Cloud Manager sets differently from the defaults:

Attribute	Value set by Cloud Manager
Autosize mode	grow
Maximum autosize	1,000 percent <div style="display: flex; align-items: center;"> i The Account Admin can modify this value from the Settings page. </div>
Security style	NTFS for CIFS volumes UNIX for NFS volumes
Space guarantee style	none
UNIX permissions (NFS only)	777

See the *volume create* man page for information about these attributes.

Internal disks for system data

In addition to the storage for user data, Cloud Manager also purchases cloud storage for system data.

AWS

- Two disks per node for boot and root data:
 - 9.7: 160 GiB io1 disk for boot data and a 220 GiB gp2 disk for root data
 - 9.6: 93 GiB io1 disk for boot data and a 140 GiB gp2 disk for root data
 - 9.5: 45 GiB io1 disk for boot data and a 140 GiB gp2 disk for root data
- Starting with version 9.8, a 540 GiB General Purpose SSD (gp2) for a core disk when using a C5, M5, or R5 instance type
- One EBS snapshot for each boot disk and root disk
- For HA pairs, one EBS volume for the Mediator instance, which is approximately 8 GiB

Azure (single node)

- Three Premium SSD disks:
 - One 10 GiB disk for boot data
 - One 140 GiB disk for root data
 - One 128 GiB disk for NVRAM

If the virtual machine that you chose for Cloud Volumes ONTAP supports Ultra SSDs, then the system uses an Ultra SSD for NVRAM, rather than a Premium SSD.

- One 1024 GiB Standard HDD disk for saving cores
- One Azure snapshot for each boot disk and root disk

Azure (HA pairs)

- Two 10 GiB Premium SSD disks for the boot volume (one per node)
- Two 140 GiB Premium Storage page blobs for the root volume (one per node)
- Two 1024 GiB Standard HDD disks for saving cores (one per node)
- Two 128 GiB Premium SSD disks for NVRAM (one per node)
- One Azure snapshot for each boot disk and root disk

GCP

- One 10 GiB Standard persistent disk for boot data
- One 64 GiB Standard persistent disk for root data
- One 500 GiB Standard persistent disk for NVRAM
- One 315 GiB Standard persistent disk for saving cores
- One GCP snapshot each for the boot disk and root disk

For an HA pair, there are two disks per node for root data.

Where the disks reside

Cloud Manager lays out the storage as follows:

- Boot data resides on a disk attached to the instance or virtual machine.

This disk, which contains the boot image, is not available to Cloud Volumes ONTAP.

- Root data, which contains the system configuration and logs, resides in aggr0.
- The storage virtual machine (SVM) root volume resides in aggr1.
- Data volumes also reside in aggr1.

Encryption

Boot and root disks are always encrypted in Azure and Google Cloud Platform because encryption is enabled by default in those cloud providers.

When you enable data encryption in AWS using the Key Management Service (KMS), the boot and root disks for Cloud Volumes ONTAP are encrypted, as well. This includes the boot disk for the mediator instance in an HA pair. The disks are encrypted using the CMK that you select when you create the working environment.

Get started in AWS

Getting started with Cloud Volumes ONTAP for AWS

Get started with Cloud Volumes ONTAP for AWS in a few steps.



Create a Connector

If you don't have a [Connector](#) yet, an Account Admin needs to create one. [Learn how to create a Connector in AWS.](#)

When you create your first Cloud Volumes ONTAP working environment, Cloud Manager prompts you to deploy a Connector if you don't have one yet.



Plan your configuration

Cloud Manager offers preconfigured packages that match your workload requirements, or you can create your own configuration. If you choose your own configuration, you should understand the options available to you. [Learn more.](#)



Set up your networking

- a. Ensure that your VPC and subnets will support connectivity between the Connector and Cloud Volumes ONTAP.
- b. Enable outbound internet access from the target VPC so the Connector and Cloud Volumes ONTAP can contact several endpoints.

This step is important because the Connector can't manage Cloud Volumes ONTAP without outbound internet access. If you need to limit outbound connectivity, refer to the list of endpoints for [the Connector and Cloud Volumes ONTAP](#).

- c. Set up a VPC endpoint to the S3 service.

A VPC endpoint is required if you want to tier cold data from Cloud Volumes ONTAP to low-cost object storage.

[Learn more about networking requirements.](#)



Set up the AWS KMS

If you want to use Amazon encryption with Cloud Volumes ONTAP, then you need to ensure that an active Customer Master Key (CMK) exists. You also need to modify the key policy for each CMK by adding the IAM role that provides permissions to the Connector as a *key user*. [Learn more](#).



Launch Cloud Volumes ONTAP using Cloud Manager

Click **Add Working Environment**, select the type of system that you would like to deploy, and complete the steps in the wizard. [Read step-by-step instructions](#).

Related links

- [Evaluating](#)
- [Creating a Connector from Cloud Manager](#)
- [Launching a Connector from the AWS Marketplace](#)
- [Installing the Connector software on a Linux host](#)
- [What Cloud Manager does with AWS permissions](#)

Planning your Cloud Volumes ONTAP configuration in AWS

When you deploy Cloud Volumes ONTAP in AWS, you can choose a preconfigured system that matches your workload requirements, or you can create your own configuration. If you choose your own configuration, you should understand the options available to you.

Viewing supported regions

Cloud Volumes ONTAP is supported in most AWS regions. [View the full list of supported regions](#).

Newer AWS regions must be enabled before you can create and manage resources in those regions. [Learn how to enable a region](#).

Choosing a license type

A few licensing options are available for Cloud Volumes ONTAP. Each of these licensing options enables you to choose a configuration that meets your needs.

Supported configurations for Cloud Volumes ONTAP in AWS

Choosing a supported instance

Cloud Volumes ONTAP supports several instance types, depending on the license type that you choose.

Supported configurations for Cloud Volumes ONTAP in AWS

Choosing a configuration that supports Flash Cache

Some Cloud Volumes ONTAP configurations in AWS include local NVMe storage, which Cloud Volumes ONTAP uses as *Flash Cache* for better performance. [Learn more about Flash Cache](#).

Understanding storage limits

The raw capacity limit for a Cloud Volumes ONTAP system is tied to the license. Additional limits impact the size of aggregates and volumes. You should be aware of these limits as you plan your configuration.

Storage limits for Cloud Volumes ONTAP in AWS

Sizing your system in AWS

Sizing your Cloud Volumes ONTAP system can help you meet requirements for performance and capacity. You should be aware of a few key points when choosing an instance type, disk type, and disk size:

Instance type

- Match your workload requirements to the maximum throughput and IOPS for each EC2 instance type.
- If several users write to the system at the same time, choose an instance type that has enough CPUs to manage the requests.
- If you have an application that is mostly reads, then choose a system with enough RAM.
 - [AWS Documentation: Amazon EC2 Instance Types](#)
 - [AWS Documentation: Amazon EBS–Optimized Instances](#)

EBS disk type

At a high level, the differences between EBS disk types are as follows. To learn more about the use cases for EBS disks, refer to [AWS Documentation: EBS Volume Types](#).

- *General Purpose SSD (gp3)* disks are the lowest-cost SSDs that balance cost and performance for a broad range of workloads. Performance is defined in terms of IOPS and throughput. gp3 disks are supported with Cloud Volumes ONTAP 9.7 and later.

When you select a gp3 disk, Cloud Manager fills in default IOPS and throughput values that provide performance that is equivalent to a gp2 disk based on the selected disk size. You can increase the values to get better performance at a higher cost, but we do not support lower values because it can result in inferior performance. In short, stick with the default values or increase them. Don't lower them. [Learn more about gp3 disks and their performance](#).

- *General Purpose SSD (gp2)* disks balance cost and performance for a broad range of workloads. Performance is defined in terms of IOPS.
- *Provisioned IOPS SSD (io1)* disks are for critical applications that require the highest performance at a higher cost.

- *Throughput Optimized HDD (st1)* disks are for frequently accessed workloads that require fast and consistent throughput at a lower price.



Tiering data to object storage is not recommended when using Throughput Optimized HDDs (st1).

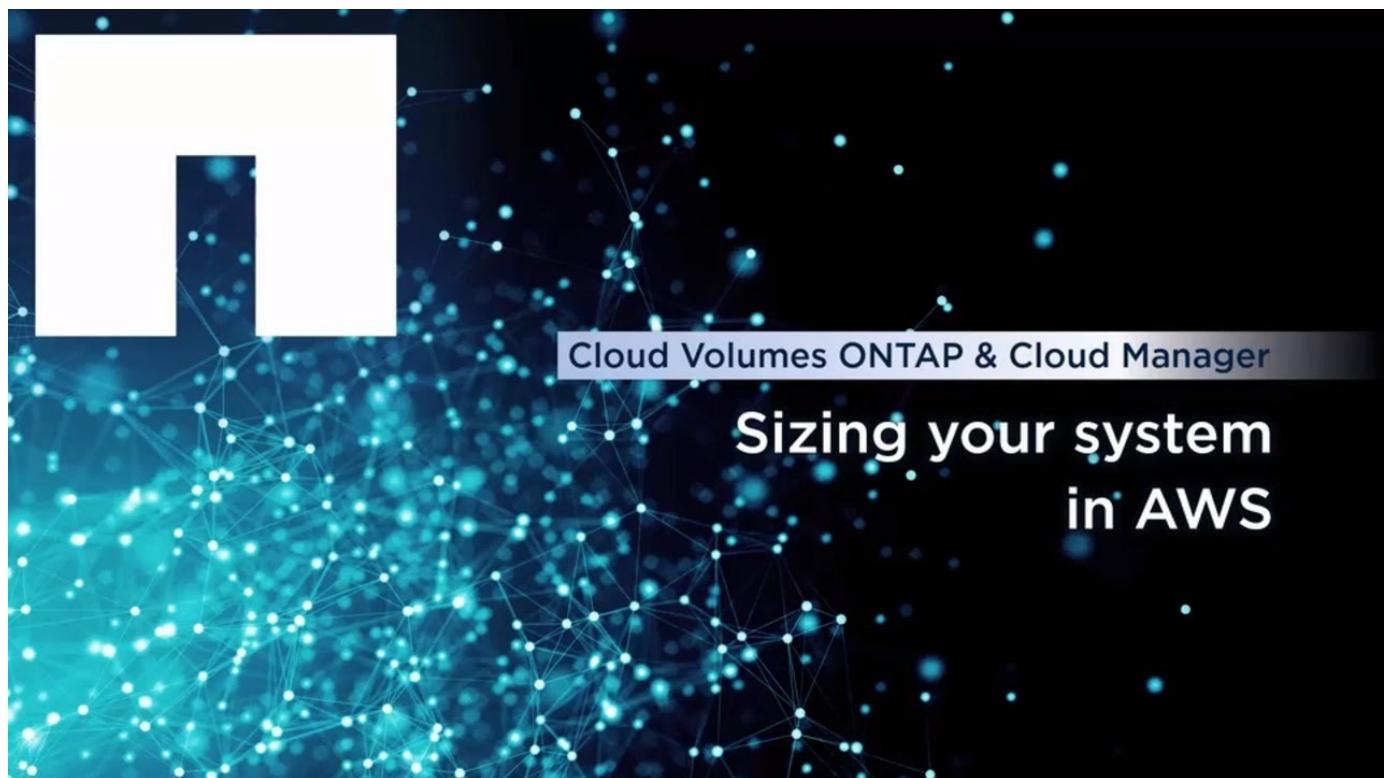
EBS disk size

You need to choose an initial disk size when you launch a Cloud Volumes ONTAP system. After that, you can [let Cloud Manager manage a system's capacity for you](#), but if you want to [build aggregates yourself](#), be aware of the following:

- All disks in an aggregate must be the same size.
- The performance of EBS disks is tied to disk size. The size determines the baseline IOPS and maximum burst duration for SSD disks and the baseline and burst throughput for HDD disks.
- Ultimately, you should choose the disk size that gives you the *sustained performance* that you need.
- Even if you do choose larger disks (for example, six 4 TiB disks), you might not get all of the IOPS because the EC2 instance can reach its bandwidth limit.

For more details about EBS disk performance, refer to [AWS Documentation: EBS Volume Types](#).

Watch the following video for more details about sizing your Cloud Volumes ONTAP system in AWS:



Preparing to deploy Cloud Volumes ONTAP in an AWS Outpost

If you have an AWS Outpost, you can deploy Cloud Volumes ONTAP in that Outpost by selecting the Outpost VPC in the Working Environment wizard. The experience is the same as any other VPC that resides in AWS. Note that you will need to first deploy a Connector in your AWS Outpost.

There are a few limitations to point out:

- Only single node Cloud Volumes ONTAP systems are supported at this time
- The EC2 instances that you can use with Cloud Volumes ONTAP are limited to what's available in your Outpost
- Only General Purpose SSDs (gp2) are supported at this time

AWS network information worksheet

When you launch Cloud Volumes ONTAP in AWS, you need to specify details about your VPC network. You can use a worksheet to collect the information from your administrator.

Network information for Cloud Volumes ONTAP

AWS information	Your value
Region	
VPC	
Subnet	
Security group (if using your own)	

Network information for an HA pair in multiple AZs

AWS information	Your value
Region	
VPC	
Security group (if using your own)	
Node 1 availability zone	
Node 1 subnet	
Node 2 availability zone	
Node 2 subnet	
Mediator availability zone	
Mediator subnet	
Key pair for the mediator	
Floating IP address for cluster management port	
Floating IP address for data on node 1	
Floating IP address for data on node 2	

AWS information	Your value
Route tables for floating IP addresses	

Choosing a write speed

Cloud Manager enables you to choose a write speed setting for Cloud Volumes ONTAP. Before you choose a write speed, you should understand the differences between the normal and high settings and risks and recommendations when using high write speed. [Learn more about write speed](#).

Choosing a volume usage profile

ONTAP includes several storage efficiency features that can reduce the total amount of storage that you need. When you create a volume in Cloud Manager, you can choose a profile that enables these features or a profile that disables them. You should learn more about these features to help you decide which profile to use.

NetApp storage efficiency features provide the following benefits:

Thin provisioning

Presents more logical storage to hosts or users than you actually have in your physical storage pool. Instead of preallocating storage space, storage space is allocated dynamically to each volume as data is written.

Deduplication

Improves efficiency by locating identical blocks of data and replacing them with references to a single shared block. This technique reduces storage capacity requirements by eliminating redundant blocks of data that reside in the same volume.

Compression

Reduces the physical capacity required to store data by compressing data within a volume on primary, secondary, and archive storage.

Set up your networking

Networking requirements for Cloud Volumes ONTAP in AWS

Cloud Manager handles the set up of networking components for Cloud Volumes ONTAP, such as IP addresses, netmasks, and routes. You need to make sure that outbound internet access is available, that enough private IP addresses are available, that the right connections are in place, and more.

General requirements

The following requirements must be met in AWS.

Outbound internet access for Cloud Volumes ONTAP nodes

Cloud Volumes ONTAP nodes require outbound internet access to send messages to NetApp AutoSupport, which proactively monitors the health of your storage.

Routing and firewall policies must allow AWS HTTP/HTTPS traffic to the following endpoints so Cloud Volumes ONTAP can send AutoSupport messages:

- <https://support.netapp.com/aods/asupmessage>
- <https://support.netapp.com/asupprod/post/1.0/postAsup>

If you have a NAT instance, you must define an inbound security group rule that allows HTTPS traffic from the private subnet to the internet.

[Learn how to configure AutoSupport.](#)

Outbound internet access for the HA mediator

The HA mediator instance must have an outbound connection to the AWS EC2 service so it can assist with storage failover. To provide the connection, you can add a public IP address, specify a proxy server, or use a manual option.

The manual option can be a NAT gateway or an interface VPC endpoint from the target subnet to the AWS EC2 service. For details about VPC endpoints, refer to [AWS Documentation: Interface VPC Endpoints \(AWS PrivateLink\)](#).

Private IP addresses

Cloud Manager automatically allocates the required number of private IP addresses to Cloud Volumes ONTAP. You need to ensure that your networking has enough private IP addresses available.

The number of LIFs that Cloud Manager allocates for Cloud Volumes ONTAP depends on whether you deploy a single node system or an HA pair. A LIF is an IP address associated with a physical port.

IP addresses for a single node system

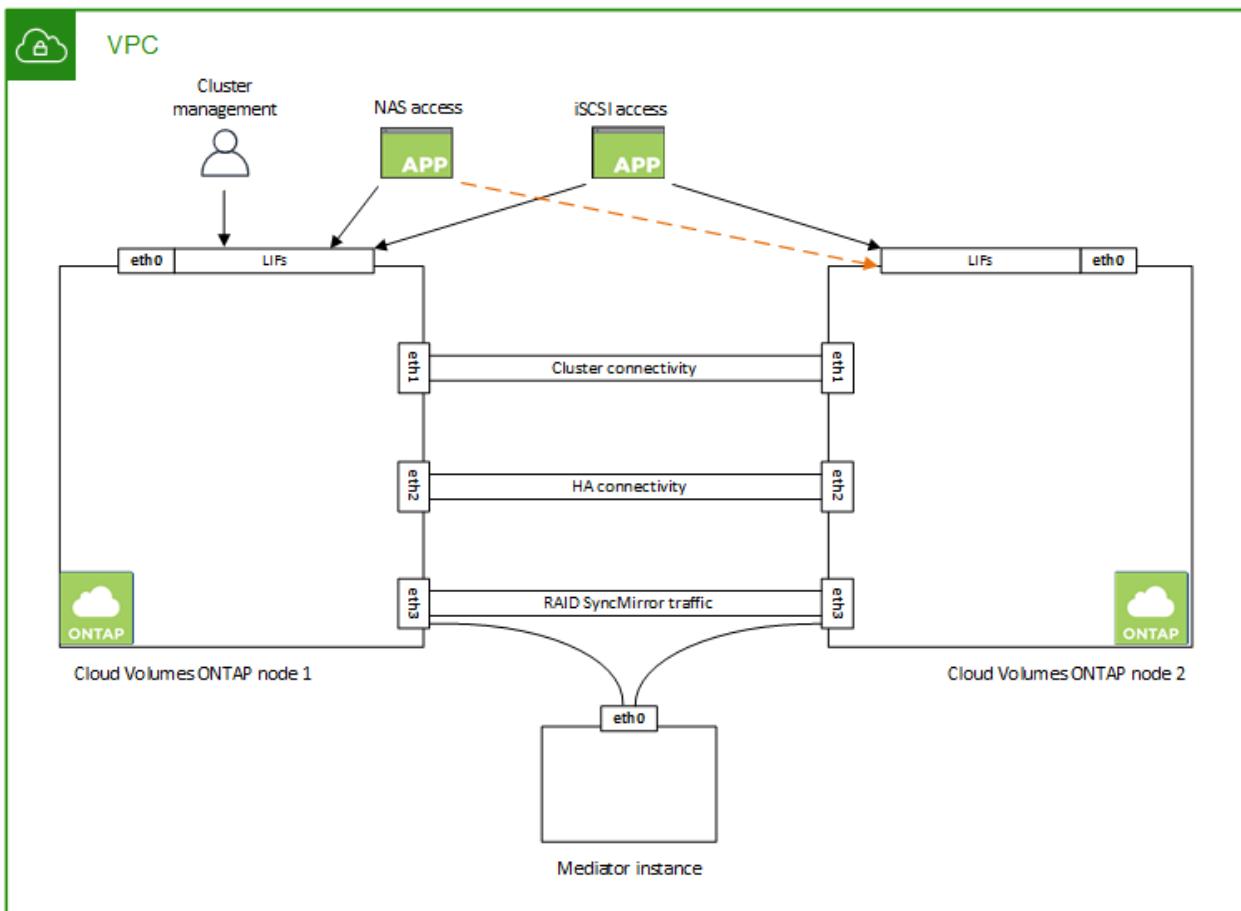
Cloud Manager allocates 6 IP addresses to a single node system:

- Cluster management LIF
- Node management LIF
- Intercluster LIF
- NAS data LIF
- iSCSI data LIF
- Storage VM management LIF

A storage VM management LIF is used with management tools like SnapCenter.

IP addresses for HA pairs

HA pairs require more IP addresses than a single node system does. These IP addresses are spread across different ethernet interfaces, as shown in the following image:



The number of private IP addresses required for an HA pair depends on which deployment model you choose. An HA pair deployed in a *single* AWS Availability Zone (AZ) requires 15 private IP addresses, while an HA pair deployed in *multiple* AZs requires 13 private IP addresses.

The following tables provide details about the LIFs that are associated with each private IP address.

LIFs for HA pairs in a single AZ

LIF	Interface	Node	Purpose
Cluster management	eth0	node 1	Administrative management of the entire cluster (HA pair).
Node management	eth0	node 1 and node 2	Administrative management of a node.
Intercluster	eth0	node 1 and node 2	Cross-cluster communication, backup, and replication.
NAS data	eth0	node 1	Client access over NAS protocols.
iSCSI data	eth0	node 1 and node 2	Client access over the iSCSI protocol.

LIF	Interface	Node	Purpose
Cluster connectivity	eth1	node 1 and node 2	Enables the nodes to communicate with each other and to move data within the cluster.
HA connectivity	eth2	node 1 and node 2	Communication between the two nodes in case of failover.
RSM iSCSI traffic	eth3	node 1 and node 2	RAID SyncMirror iSCSI traffic, as well as communication between the two Cloud Volumes ONTAP nodes and the mediator.
Mediator	eth0	Mediator	A communication channel between the nodes and the mediator to assist in storage takeover and giveback processes.

LIFs for HA pairs in multiple AZs

LIF	Interface	Node	Purpose
Node management	eth0	node 1 and node 2	Administrative management of a node.
Intercluster	eth0	node 1 and node 2	Cross-cluster communication, backup, and replication.
iSCSI data	eth0	node 1 and node 2	Client access over the iSCSI protocol. This LIF also manages the migration of floating IP addresses between nodes.
Cluster connectivity	eth1	node 1 and node 2	Enables the nodes to communicate with each other and to move data within the cluster.
HA connectivity	eth2	node 1 and node 2	Communication between the two nodes in case of failover.
RSM iSCSI traffic	eth3	node 1 and node 2	RAID SyncMirror iSCSI traffic, as well as communication between the two Cloud Volumes ONTAP nodes and the mediator.
Mediator	eth0	Mediator	A communication channel between the nodes and the mediator to assist in storage takeover and giveback processes.



When deployed in multiple Availability Zones, several LIFs are associated with [floating IP addresses](#), which don't count against the AWS private IP limit.

Security groups

You do not need to create security groups because Cloud Manager does that for you. If you need to use your own, refer to [Security group rules](#).

Connection for data tiering

If you want to use EBS as a performance tier and AWS S3 as a capacity tier, you must ensure that Cloud Volumes ONTAP has a connection to S3. The best way to provide that connection is by creating a VPC

Endpoint to the S3 service. For instructions, see [AWS Documentation: Creating a Gateway Endpoint](#).

When you create the VPC Endpoint, be sure to select the region, VPC, and route table that corresponds to the Cloud Volumes ONTAP instance. You must also modify the security group to add an outbound HTTPS rule that enables traffic to the S3 endpoint. Otherwise, Cloud Volumes ONTAP cannot connect to the S3 service.

If you experience any issues, see [AWS Support Knowledge Center: Why can't I connect to an S3 bucket using a gateway VPC endpoint?](#)

Connections to ONTAP systems

To replicate data between a Cloud Volumes ONTAP system in AWS and ONTAP systems in other networks, you must have a VPN connection between the AWS VPC and the other network—for example, an Azure VNet or your corporate network. For instructions, see [AWS Documentation: Setting Up an AWS VPN Connection](#).

DNS and Active Directory for CIFS

If you want to provision CIFS storage, you must set up DNS and Active Directory in AWS or extend your on-premises setup to AWS.

The DNS server must provide name resolution services for the Active Directory environment. You can configure DHCP option sets to use the default EC2 DNS server, which must not be the DNS server used by the Active Directory environment.

For instructions, refer to [AWS Documentation: Active Directory Domain Services on the AWS Cloud: Quick Start Reference Deployment](#).

Requirements for HA pairs in multiple AZs

Additional AWS networking requirements apply to Cloud Volumes ONTAP HA configurations that use multiple Availability Zones (AZs). You should review these requirements before you launch an HA pair because you must enter the networking details in Cloud Manager when you create the working environment.

To understand how HA pairs work, see [High-availability pairs](#).

Availability Zones

This HA deployment model uses multiple AZs to ensure high availability of your data. You should use a dedicated AZ for each Cloud Volumes ONTAP instance and the mediator instance, which provides a communication channel between the HA pair.

A subnet should be available in each Availability Zone.

Floating IP addresses for NAS data and cluster/SVM management

HA configurations in multiple AZs use floating IP addresses that migrate between nodes if failures occur. They are not natively accessible from outside the VPC, unless you [set up an AWS transit gateway](#).

One floating IP address is for cluster management, one is for NFS/CIFS data on node 1, and one is for NFS/CIFS data on node 2. A fourth floating IP address for SVM management is optional.



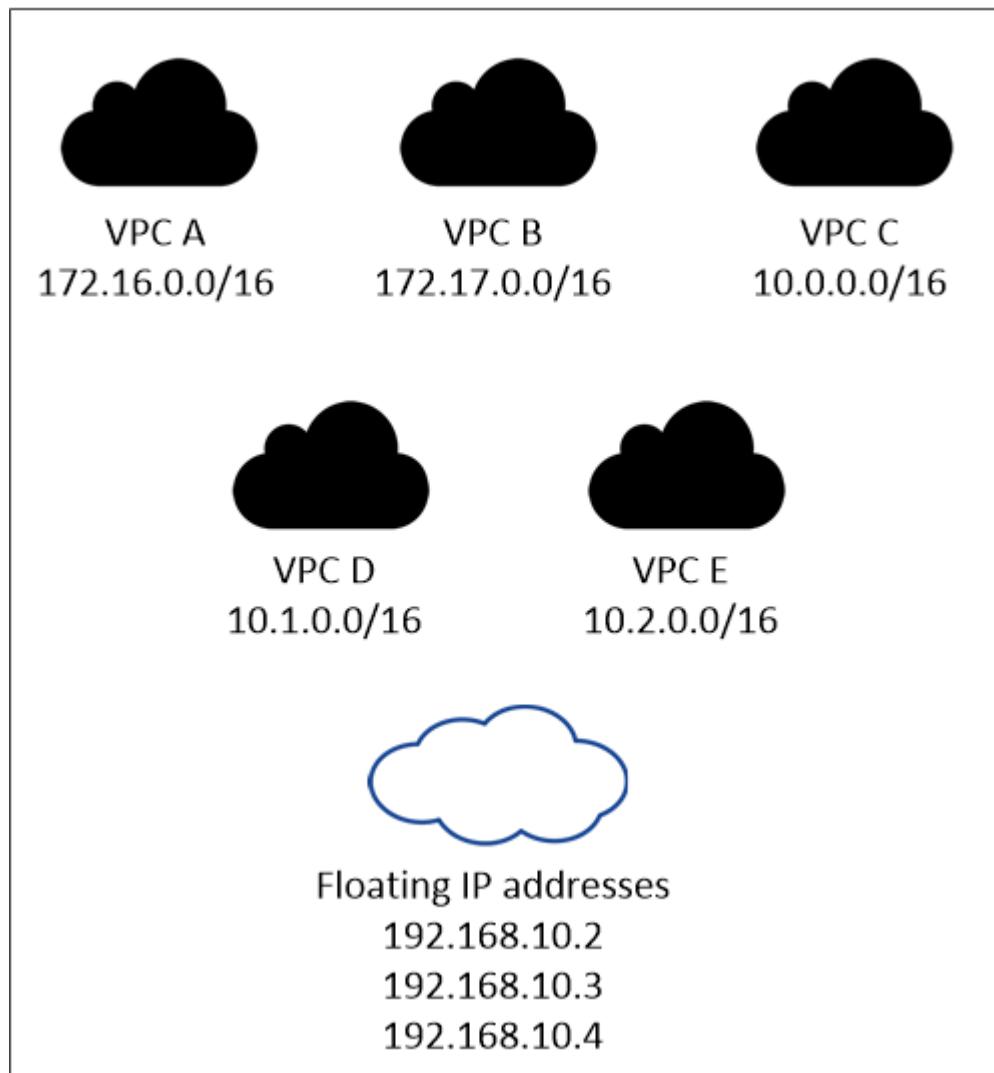
A floating IP address is required for the SVM management LIF if you use SnapDrive for Windows or SnapCenter with the HA pair. If you don't specify the IP address when you deploy the system, you can create the LIF later. For details, see [Setting up Cloud Volumes ONTAP](#).

You need to enter the floating IP addresses in Cloud Manager when you create a Cloud Volumes ONTAP HA working environment. Cloud Manager allocates the IP addresses to the HA pair when it launches the system.

The floating IP addresses must be outside of the CIDR blocks for all VPCs in the AWS region in which you deploy the HA configuration. Think of the floating IP addresses as a logical subnet that's outside of the VPCs in your region.

The following example shows the relationship between floating IP addresses and the VPCs in an AWS region. While the floating IP addresses are outside the CIDR blocks for all VPCs, they're routable to subnets through route tables.

AWS region



Cloud Manager automatically creates static IP addresses for iSCSI access and for NAS access from clients outside the VPC. You don't need to meet any requirements for these types of IP addresses.

Transit gateway to enable floating IP access from outside the VPC

If needed, [set up an AWS transit gateway](#) to enable access to an HA pair's floating IP addresses from outside the VPC where the HA pair resides.

Route tables

After you specify the floating IP addresses in Cloud Manager, you are then prompted to select the route tables that should include routes to the floating IP addresses. This enables client access to the HA pair.

If you have just one route table for the subnets in your VPC (the main route table), then Cloud Manager automatically adds the floating IP addresses to that route table. If you have more than one route table, it's very important to select the correct route tables when launching the HA pair. Otherwise, some clients might not have access to Cloud Volumes ONTAP.

For example, you might have two subnets that are associated with different route tables. If you select route table A, but not route table B, then clients in the subnet associated with route table A can access the HA pair, but clients in the subnet associated with route table B can't.

For more information about route tables, refer to [AWS Documentation: Route Tables](#).

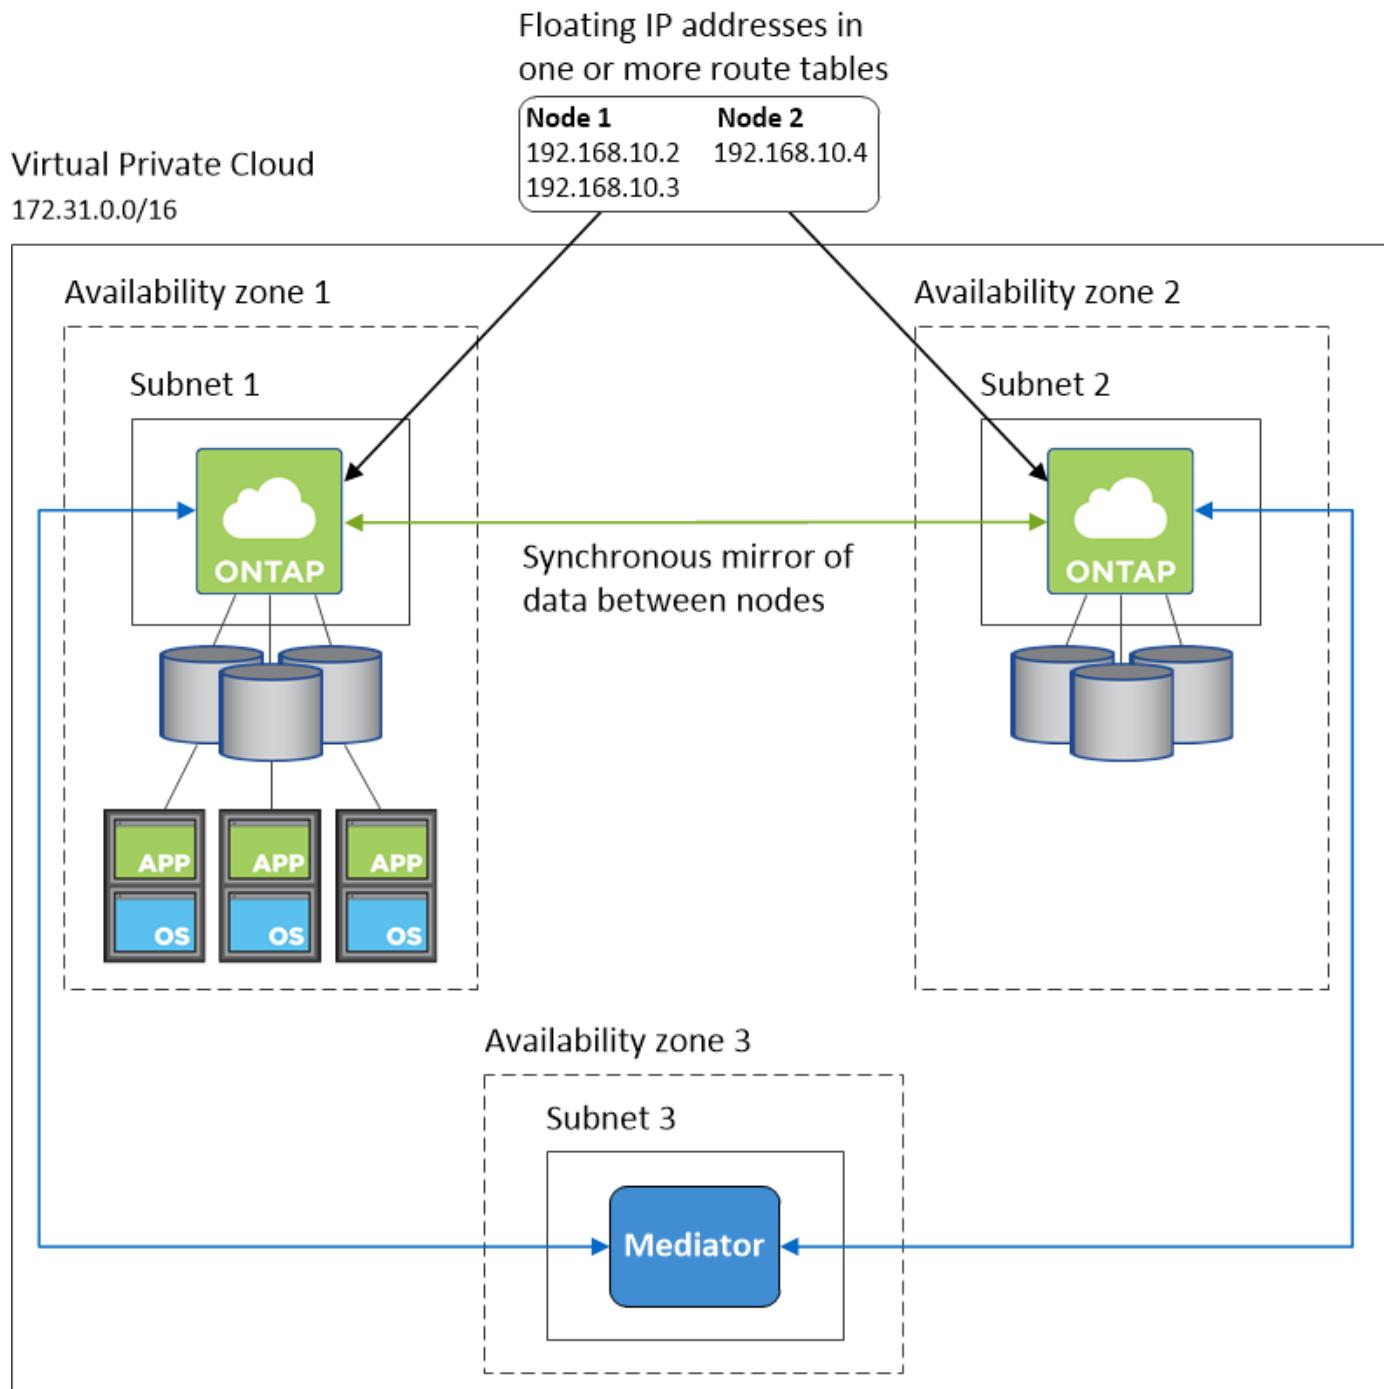
Connection to NetApp management tools

To use NetApp management tools with HA configurations that are in multiple AZs, you have two connection options:

1. Deploy the NetApp management tools in a different VPC and [set up an AWS transit gateway](#). The gateway enables access to the floating IP address for the cluster management interface from outside the VPC.
2. Deploy the NetApp management tools in the same VPC with a similar routing configuration as NAS clients.

Example HA configuration

The following image illustrates the networking components specific to an HA pair in multiple AZs: three Availability Zones, three subnets, floating IP addresses, and a route table.



Requirements for the Connector

Set up your networking so that the Connector can manage resources and processes within your public cloud environment. The most important step is ensuring outbound internet access to various endpoints.



If your network uses a proxy server for all communication to the internet, you can specify the proxy server from the Settings page. Refer to [Configuring the Connector to use a proxy server](#).

Connection to target networks

A Connector requires a network connection to the VPCs and VNets in which you want to deploy Cloud Volumes ONTAP.

For example, if you install a Connector in your corporate network, then you must set up a VPN connection to the VPC or VNet in which you launch Cloud Volumes ONTAP.

Outbound internet access

The Connector requires outbound internet access to manage resources and processes within your public cloud environment. A Connector contacts the following endpoints when managing resources in AWS:



If your VPC uses a network access control list (ACL) to filter traffic, then make sure that you enable these endpoints for both outbound and inbound traffic.

Endpoints	Purpose
AWS services (amazonaws.com): <ul style="list-style-type: none">• CloudFormation• Elastic Compute Cloud (EC2)• Key Management Service (KMS)• Security Token Service (STS)• Simple Storage Service (S3) <p>The exact endpoint depends on the region in which you deploy Cloud Volumes ONTAP. Refer to AWS documentation for details.</p>	Enables the Connector to deploy and manage Cloud Volumes ONTAP in AWS.
https://api.services.cloud.netapp.com:443	API requests to NetApp Cloud Central.
https://cloud.support.netapp.com.s3.us-west-1.amazonaws.com	Provides access to software images, manifests, and templates.
https://cognito-idp.us-east-1.amazonaws.com https://cognito-identity.us-east-1.amazonaws.com https://sts.amazonaws.com https://cloud-support-netapp-com-accelerated.s3.amazonaws.com	Enables the Connector to access and download manifests, templates, and Cloud Volumes ONTAP upgrade images.
https://cloudmanagerinfraprod.azurecr.io*.blob.core.windows.net	Access to software images of container components for an infrastructure that's running Docker and provides a solution for service integrations with Cloud Manager.
https://kinesis.us-east-1.amazonaws.com	Enables NetApp to stream data from audit records.
https://cloudmanager.cloud.netapp.com	Communication with the Cloud Manager service, which includes Cloud Central accounts.
https://netapp-cloud-account.auth0.com	Communication with NetApp Cloud Central for centralized user authentication.
support.netapp.com:443 https://mysupport.netapp.com	Communication with NetApp AutoSupport. Note that the Connector communicates with support.netapp.com:443, which redirects to https://mysupport.netapp.com.

Endpoints	Purpose
https://support.netapp.com/svcgw https://support.netapp.com/ServiceGW/entitlement https://eval.lic.netapp.com.s3.us-west-1.amazonaws.com https://cloud-support-netapp-com.s3.us-west-1.amazonaws.com	Communication with NetApp for system licensing and support registration.
https://client.infra.support.netapp.com.s3.us-west-1.amazonaws.com https://cloud-support-netapp-com-accelerated.s3.us-west-1.amazonaws.com https://trigger.asup.netapp.com.s3.us-west-1.amazonaws.com	Enables NetApp to collect information needed to troubleshoot support issues.
https://ipa-signer.cloudmanager.netapp.com	Enables Cloud Manager to generate licenses (for example, a FlexCache license for Cloud Volumes ONTAP)
Various third-party locations, for example: <ul style="list-style-type: none"> • https://repo1.maven.org/maven2 • https://oss.sonatype.org/content/repositories • https://repo.typesafe.com <p>Third-party locations are subject to change.</p>	During upgrades, Cloud Manager downloads the latest packages for third-party dependencies.

While you should perform almost all tasks from the SaaS user interface, a local user interface is still available on the Connector. The machine running the web browser must have connections to the following endpoints:

Endpoints	Purpose
The Connector host	<p>You must enter the host's IP address from a web browser to load the Cloud Manager console.</p> <p>Depending on your connectivity to your cloud provider, you can use the private IP or a public IP assigned to the host:</p> <ul style="list-style-type: none"> • A private IP works if you have a VPN and direct connect access to your virtual network • A public IP works in any networking scenario <p>In any case, you should secure network access by ensuring that security group rules allow access from only authorized IPs or subnets.</p>
https://auth0.com https://cdn.auth0.com https://netapp-cloud-account.auth0.com https://services.cloud.netapp.com	Your web browser connects to these endpoints for centralized user authentication through NetApp Cloud Central.

Endpoints	Purpose
https://widget.intercom.io	For in-product chat that enables you to talk to NetApp cloud experts.

Setting up an AWS transit gateway for HA pairs in multiple AZs

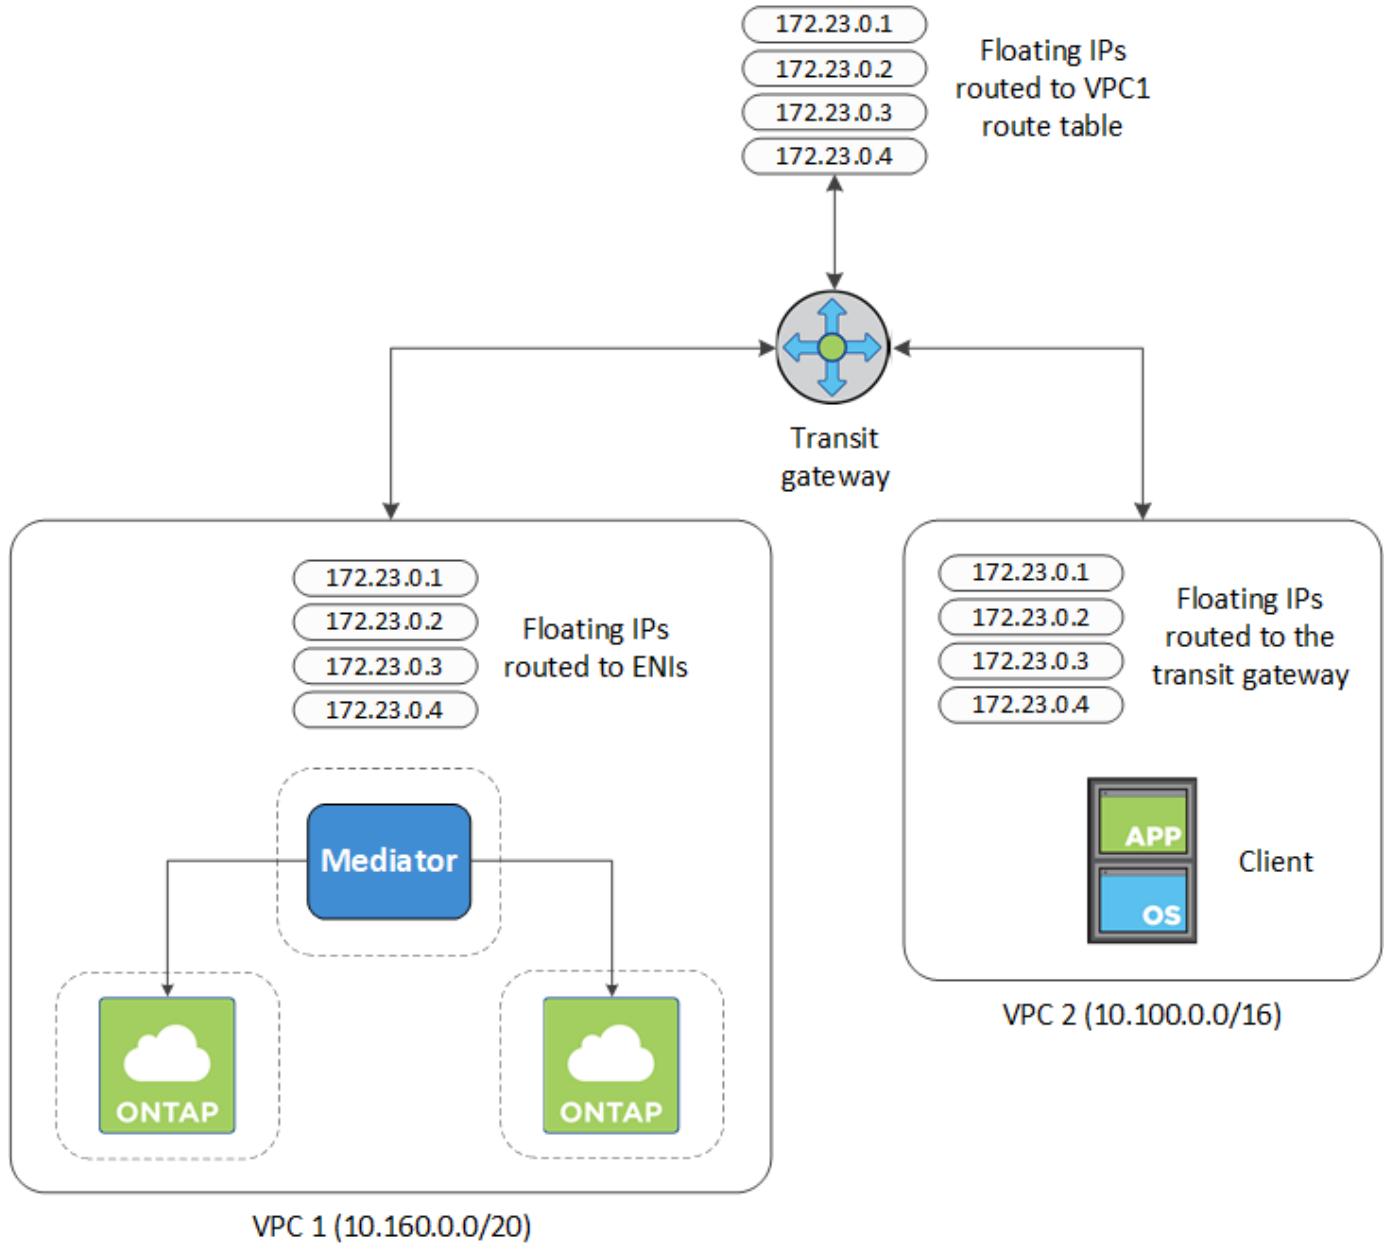
Set up an AWS transit gateway to enable access to an HA pair's [floating IP addresses](#) from outside the VPC where the HA pair resides.

When a Cloud Volumes ONTAP HA configuration is spread across multiple AWS Availability Zones, floating IP addresses are required for NAS data access from within the VPC. These floating IP addresses can migrate between nodes when failures occur, but they are not natively accessible from outside the VPC. Separate private IP addresses provide data access from outside the VPC, but they don't provide automatic failover.

Floating IP addresses are also required for the cluster management interface and the optional SVM management LIF.

If you set up an AWS transit gateway, you enable access to the floating IP addresses from outside the VPC where the HA pair resides. That means NAS clients and NetApp management tools outside the VPC can access the floating IPs.

Here's an example that shows two VPCs connected by a transit gateway. An HA system resides in one VPC, while a client resides in the other. You could then mount a NAS volume on the client using the floating IP address.



The following steps illustrate how to set up a similar configuration.

Steps

1. Create a transit gateway and attach the VPCs to the gateway.
2. Associate the VPCs with the transit gateway route table.
 - a. In the **VPC** service, click **Transit Gateway Route Tables**.
 - b. Select the route table.
 - c. Click **Associations** and then select **Create association**.
 - d. Choose the attachments (the VPCs) to associate and then click **Create association**.
3. Create routes in the transit gateway's route table by specifying the HA pair's floating IP addresses.

You can find the floating IP addresses on the Working Environment Information page in Cloud Manager. Here's an example:

NFS & CIFS access from within the VPC using Floating IP

Auto failover

Cluster Management : 172.23.0.1

Data (nfs,cifs) : Node 1: 172.23.0.2 | Node 2: 172.23.0.3

Access

SVM Management : 172.23.0.4

The following sample image shows the route table for the transit gateway. It includes routes to the CIDR blocks of the two VPCs and four floating IP addresses used by Cloud Volumes ONTAP.

Transit Gateway Route Table: tgw-rtb-0ea8ee291c7aedd3					
Details	Associations	Propagations	Routes	Tags	...
The table below will return a maximum of 1000 routes. Narrow the filter or use export routes to view more routes.					
Create route	Replace route	Delete route			
<input type="text"/> Filter by attributes or search by keyword					
CIDR	Attachment		Resource type	Route type	Route state
10.100.0.0/16	tgw-attach-05e77bd34e2ff91f8 vpc-0b2bc30e0dc8e0db1		VPC2	propagated	active
10.160.0.0/20	tgw-attach-00eba3eac3250d7db vpc-673ae603		VPC1	propagated	active
172.23.0.1/32	tgw-attach-00eba3eac3250d7db vpc-673ae603	VPC	static	active	
172.23.0.2/32	tgw-attach-00eba3eac3250d7db vpc-673ae603	VPC	static	active	
172.23.0.3/32	tgw-attach-00eba3eac3250d7db vpc-673ae603	VPC	static	active	
172.23.0.4/32	tgw-attach-00eba3eac3250d7db vpc-673ae603	VPC	static	active	

4. Modify the route table of VPCs that need to access the floating IP addresses.

- Add route entries to the floating IP addresses.
- Add a route entry to the CIDR block of the VPC where the HA pair resides.

The following sample image shows the route table for VPC 2, which includes routes to VPC 1 and the floating IP addresses.

Route Table: rtb-0569a1bd740ed033f				
Summary	Routes	Subnet Associations	Route Propagation	Tags
Edit routes				
View All routes				
Destination	Target	Status	Propagated	
10.100.0.0/16	local	active	No	
0.0.0.0/0	igw-07250bd01781e67df	active	No	
10.160.0.0/20	taw-015b7c249661ac279	active	No	VPC1
172.23.0.1/32	tgw-015b7c249661ac279	active	No	
172.23.0.2/32	tgw-015b7c249661ac279	active	No	
172.23.0.3/32	tgw-015b7c249661ac279	active	No	
172.23.0.4/32	tgw-015b7c249661ac279	active	No	Floating IP Addresses

5. Modify the route table for the HA pair's VPC by adding a route to the VPC that needs access to the floating IP addresses.

This step is important because it completes the routing between the VPCs.

The following sample image shows the route table for VPC 1. It includes a route to the floating IP addresses and to VPC 2, which is where a client resides. Cloud Manager automatically added the floating IPs to the route table when it deployed the HA pair.

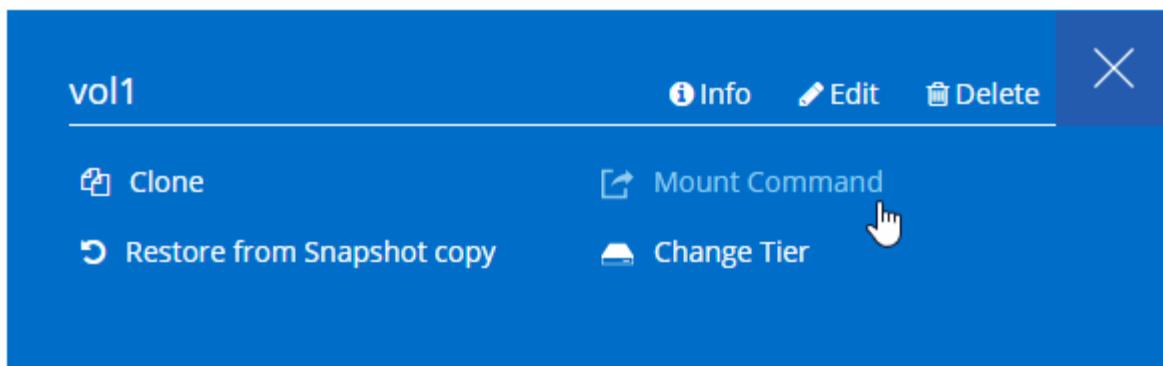
Route Table: rtb-0569a1bd740ed033f				
Summary	Routes	Subnet Associations	Route Propagation	Tags
Edit routes				
View All routes				
Destination	Target	Status		
10.160.0.0/20	local	active		
pl-68a54001 (com.amazonaws.us-west-2.s3, 54.231.160.0/19, 52.218.128.0/17, 52.92.32.0/22)	vpce-cb51a0a2	active		
0.0.0.0/0	igw-b2182dd7	active		
10.60.29.0/25	pcx-589c3331	active		
10.100.0.0/16	tgw-015b7c249661ac279	active		VPC2
10.129.0.0/20	pcx-ff7e1396	active		
172.23.0.1/32	eni-0854d4715559c3cdb	active		
172.23.0.2/32	eni-0854d4715559c3cdb	active		
172.23.0.3/32	eni-0f76681216c3108ed	active		
172.23.0.4/32	eni-0854d4715559c3cdb	active		Floating IP Addresses

6. Mount volumes to clients using the floating IP address.

You can find the correct IP address in Cloud Manager by selecting a volume and clicking **Mount Command**.

Volumes

2 Volumes | 0.22 TB Allocated | < 0.01 TB Used (0 TB in S3)



7. If you're mounting an NFS volume, configure the export policy to match the subnet of the client VPC.

[Learn how to edit a volume.](#)

Related links

- [High-availability pairs in AWS](#)
- [Networking requirements for Cloud Volumes ONTAP in AWS](#)

Security group rules for AWS

Cloud Manager creates AWS security groups that include the inbound and outbound rules that the Connector and Cloud Volumes ONTAP need to operate successfully. You might want to refer to the ports for testing purposes or if you prefer your to use own security groups.

Rules for Cloud Volumes ONTAP

The security group for Cloud Volumes ONTAP requires both inbound and outbound rules.

Inbound rules

The source for inbound rules in the predefined security group is 0.0.0.0/0.

Protocol	Port	Purpose
All ICMP	All	Pinging the instance
HTTP	80	HTTP access to the System Manager web console using the IP address of the cluster management LIF
HTTPS	443	HTTPS access to the System Manager web console using the IP address of the cluster management LIF
SSH	22	SSH access to the IP address of the cluster management LIF or a node management LIF

Protocol	Port	Purpose
TCP	111	Remote procedure call for NFS
TCP	139	NetBIOS service session for CIFS
TCP	161-162	Simple network management protocol
TCP	445	Microsoft SMB/CIFS over TCP with NetBIOS framing
TCP	635	NFS mount
TCP	749	Kerberos
TCP	2049	NFS server daemon
TCP	3260	iSCSI access through the iSCSI data LIF
TCP	4045	NFS lock daemon
TCP	4046	Network status monitor for NFS
TCP	10000	Backup using NDMP
TCP	11104	Management of intercluster communication sessions for SnapMirror
TCP	11105	SnapMirror data transfer using intercluster LIFs
UDP	111	Remote procedure call for NFS
UDP	161-162	Simple network management protocol
UDP	635	NFS mount
UDP	2049	NFS server daemon
UDP	4045	NFS lock daemon
UDP	4046	Network status monitor for NFS
UDP	4049	NFS rquotad protocol

Outbound rules

The predefined security group for Cloud Volumes ONTAP opens all outbound traffic. If that is acceptable, follow the basic outbound rules. If you need more rigid rules, use the advanced outbound rules.

Basic outbound rules

The predefined security group for Cloud Volumes ONTAP includes the following outbound rules.

Protocol	Port	Purpose
All ICMP	All	All outbound traffic
All TCP	All	All outbound traffic
All UDP	All	All outbound traffic

Advanced outbound rules

If you need rigid rules for outbound traffic, you can use the following information to open only those ports that are required for outbound communication by Cloud Volumes ONTAP.



The source is the interface (IP address) on the Cloud Volumes ONTAP system.

Service	Protocol	Port	Source	Destination	Purpose
Active Directory	TCP	88	Node management LIF	Active Directory forest	Kerberos V authentication
	UDP	137	Node management LIF	Active Directory forest	NetBIOS name service
	UDP	138	Node management LIF	Active Directory forest	NetBIOS datagram service
	TCP	139	Node management LIF	Active Directory forest	NetBIOS service session
	TCP & UDP	389	Node management LIF	Active Directory forest	LDAP
	TCP	445	Node management LIF	Active Directory forest	Microsoft SMB/CIFS over TCP with NetBIOS framing
	TCP	464	Node management LIF	Active Directory forest	Kerberos V change & set password (SET_CHANGE)
	UDP	464	Node management LIF	Active Directory forest	Kerberos key administration
	TCP	749	Node management LIF	Active Directory forest	Kerberos V change & set Password (RPCSEC_GSS)
	TCP	88	Data LIF (NFS, CIFS, iSCSI)	Active Directory forest	Kerberos V authentication
	UDP	137	Data LIF (NFS, CIFS)	Active Directory forest	NetBIOS name service
	UDP	138	Data LIF (NFS, CIFS)	Active Directory forest	NetBIOS datagram service
	TCP	139	Data LIF (NFS, CIFS)	Active Directory forest	NetBIOS service session
	TCP & UDP	389	Data LIF (NFS, CIFS)	Active Directory forest	LDAP
	TCP	445	Data LIF (NFS, CIFS)	Active Directory forest	Microsoft SMB/CIFS over TCP with NetBIOS framing
	TCP	464	Data LIF (NFS, CIFS)	Active Directory forest	Kerberos V change & set password (SET_CHANGE)
	UDP	464	Data LIF (NFS, CIFS)	Active Directory forest	Kerberos key administration
	TCP	749	Data LIF (NFS, CIFS)	Active Directory forest	Kerberos V change & set password (RPCSEC_GSS)
Backup to S3	TCP	5010	Intercluster LIF	Backup endpoint or restore endpoint	Back up and restore operations for the Backup to S3 feature

Service	Protocol	Port	Source	Destination	Purpose
Cluster	All traffic	All traffic	All LIFs on one node	All LIFs on the other node	Intercluster communications (Cloud Volumes ONTAP HA only)
	TCP	3000	Node management LIF	HA mediator	ZAPI calls (Cloud Volumes ONTAP HA only)
	ICMP	1	Node management LIF	HA mediator	Keep alive (Cloud Volumes ONTAP HA only)
DHCP	UDP	68	Node management LIF	DHCP	DHCP client for first-time setup
DHCPS	UDP	67	Node management LIF	DHCP	DHCP server
DNS	UDP	53	Node management LIF and data LIF (NFS, CIFS)	DNS	DNS
NDMP	TCP	1860 0–18 699	Node management LIF	Destination servers	NDMP copy
SMTP	TCP	25	Node management LIF	Mail server	SMTP alerts, can be used for AutoSupport
SNMP	TCP	161	Node management LIF	Monitor server	Monitoring by SNMP traps
	UDP	161	Node management LIF	Monitor server	Monitoring by SNMP traps
	TCP	162	Node management LIF	Monitor server	Monitoring by SNMP traps
	UDP	162	Node management LIF	Monitor server	Monitoring by SNMP traps
SnapMirr or	TCP	1110 4	Intercluster LIF	ONTAP intercluster LIFs	Management of intercluster communication sessions for SnapMirror
	TCP	1110 5	Intercluster LIF	ONTAP intercluster LIFs	SnapMirror data transfer
Syslog	UDP	514	Node management LIF	Syslog server	Syslog forward messages

Rules for the HA mediator external security group

The predefined external security group for the Cloud Volumes ONTAP HA mediator includes the following inbound and outbound rules.

Inbound rules

The source for inbound rules is 0.0.0.0/0.

Protocol	Port	Purpose
SSH	22	SSH connections to the HA mediator
TCP	3000	RESTful API access from the Connector

Outbound rules

The predefined security group for the HA mediator opens all outbound traffic. If that is acceptable, follow the basic outbound rules. If you need more rigid rules, use the advanced outbound rules.

Basic outbound rules

The predefined security group for the HA mediator includes the following outbound rules.

Protocol	Port	Purpose
All TCP	All	All outbound traffic
All UDP	All	All outbound traffic

Advanced outbound rules

If you need rigid rules for outbound traffic, you can use the following information to open only those ports that are required for outbound communication by the HA mediator.

Protocol	Port	Destination	Purpose
HTTP	80	Connector IP address	Download upgrades for the mediator
HTTPS	443	AWS API services	Assist with storage failover
UDP	53	AWS API services	Assist with storage failover



Rather than open ports 443 and 53, you can create an interface VPC endpoint from the target subnet to the AWS EC2 service.

Rules for the HA mediator internal security group

The predefined internal security group for the Cloud Volumes ONTAP HA mediator includes the following rules. Cloud Manager always creates this security group. You do not have the option to use your own.

Inbound rules

The predefined security group includes the following inbound rules.

Protocol	Port	Purpose
All traffic	All	Communication between the HA mediator and HA nodes

Outbound rules

The predefined security group includes the following outbound rules.

Protocol	Port	Purpose
All traffic	All	Communication between the HA mediator and HA nodes

Rules for the Connector

The security group for the Connector requires both inbound and outbound rules.

Inbound rules

Protocol	Port	Purpose
SSH	22	Provides SSH access to the Connector host
HTTP	80	Provides HTTP access from client web browsers to the local user interface and connections from Cloud Data Sense
HTTPS	443	Provides HTTPS access from client web browsers to the local user interface
TCP	3128	Provides the Cloud Data Sense instance with internet access, if your AWS network doesn't use a NAT or proxy

Outbound rules

The predefined security group for the Connector opens all outbound traffic. If that is acceptable, follow the basic outbound rules. If you need more rigid rules, use the advanced outbound rules.

Basic outbound rules

The predefined security group for the Connector includes the following outbound rules.

Protocol	Port	Purpose
All TCP	All	All outbound traffic
All UDP	All	All outbound traffic

Advanced outbound rules

If you need rigid rules for outbound traffic, you can use the following information to open only those ports that are required for outbound communication by the Connector.



The source IP address is the Connector host.

Service	Protocol	Port	Destination	Purpose
Active Directory	TCP	88	Active Directory forest	Kerberos V authentication
	TCP	139	Active Directory forest	NetBIOS service session
	TCP	389	Active Directory forest	LDAP
	TCP	445	Active Directory forest	Microsoft SMB/CIFS over TCP with NetBIOS framing
	TCP	464	Active Directory forest	Kerberos V change & set password (SET_CHANGE)
	TCP	749	Active Directory forest	Active Directory Kerberos V change & set password (RPCSEC_GSS)
	UDP	137	Active Directory forest	NetBIOS name service
	UDP	138	Active Directory forest	NetBIOS datagram service
	UDP	464	Active Directory forest	Kerberos key administration
API calls and AutoSupport	HTTP	443	Outbound internet and ONTAP cluster management LIF	API calls to AWS and ONTAP, and sending AutoSupport messages to NetApp
API calls	TCP	3000	ONTAP HA mediator	Communication with the ONTAP HA mediator
	TCP	8088	Backup to S3	API calls to Backup to S3
DNS	UDP	53	DNS	Used for DNS resolve by Cloud Manager
Cloud Data Sense	HTTP	80	Cloud Data Sense instance	Cloud Data Sense for Cloud Volumes ONTAP

Setting up the AWS KMS

If you want to use Amazon encryption with Cloud Volumes ONTAP, then you need to set up the AWS Key Management Service (KMS).

Steps

1. Ensure that an active Customer Master Key (CMK) exists.

The CMK can be an AWS-managed CMK or a customer-managed CMK. It can be in the same AWS account as Cloud Manager and Cloud Volumes ONTAP or in a different AWS account.

[AWS Documentation: Customer Master Keys \(CMKs\)](#)

2. Modify the key policy for each CMK by adding the IAM role that provides permissions to Cloud Manager as a *key user*.

Adding the IAM role as a key user gives Cloud Manager permissions to use the CMK with Cloud Volumes ONTAP.

[AWS Documentation: Editing Keys](#)

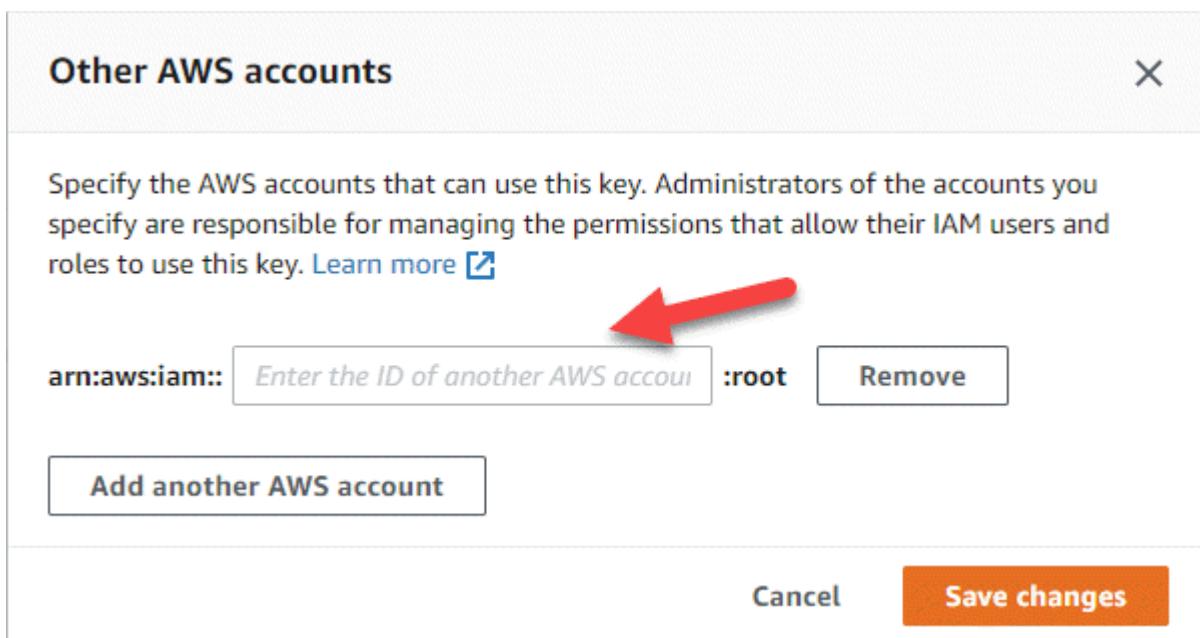
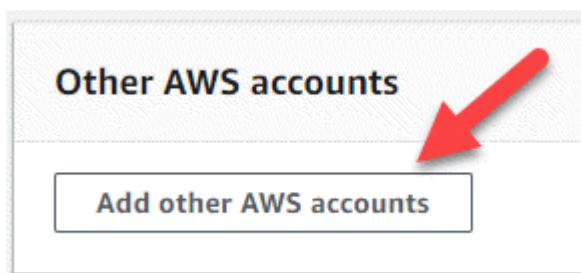
3. If the CMK is in a different AWS account, complete the following steps:

- Go to the KMS console from the account where the CMK resides.
- Select the key.
- In the **General configuration** pane, copy the ARN of the key.

You'll need to provide the ARN to Cloud Manager when you create the Cloud Volumes ONTAP system.

- In the **Other AWS accounts** pane, add the AWS account that provides Cloud Manager with permissions.

In most cases, this is the account where Cloud Manager resides. If Cloud Manager wasn't installed in AWS, it would be the account for which you provided AWS access keys to Cloud Manager.



- Now switch to the AWS account that provides Cloud Manager with permissions and open the IAM console.
- Create an IAM policy that includes the permissions listed below.
- Attach the policy to the IAM role or IAM user that provides permissions to Cloud Manager.

The following policy provides the permissions that Cloud Manager needs to use the CMK from the external AWS account. Be sure to modify the region and account ID in the "Resource" sections.

```

{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "AllowUseOfTheKey",
            "Effect": "Allow",
            "Action": [
                "kms:Encrypt",
                "kms:Decrypt",
                "kms:ReEncrypt*",
                "kms:GenerateDataKey*",
                "kms:DescribeKey"
            ],
            "Resource": [
                "arn:aws:kms:us-east-
1:externalaccountid:key/externalkeyid"
            ]
        },
        {
            "Sid": "AllowAttachmentOfPersistentResources",
            "Effect": "Allow",
            "Action": [
                "kms>CreateGrant",
                "kms>ListGrants",
                "kms:RevokeGrant"
            ],
            "Resource": [
                "arn:aws:kms:us-east-
1:externalaccountid:key/externalaccountid"
            ],
            "Condition": {
                "Bool": {
                    "kms:GrantIsForAWSResource": true
                }
            }
        }
    ]
}

```

For additional details about this process, see [AWS Documentation: Allowing External AWS Accounts to Access a CMK](#).

4. If you are using a customer-managed CMK, modify the key policy for the CMK by adding the Cloud Volumes ONTAP IAM role as a *key user*.

This step is required if you enabled data tiering on Cloud Volumes ONTAP and want to encrypt the data

stored in the S3 bucket.

You'll need to perform this step *after* you deploy Cloud Volumes ONTAP because the IAM role is created when you create a working environment. (Of course, you do have the option to use an existing Cloud Volumes ONTAP IAM role, so it's possible to perform this step before.)

[AWS Documentation: Editing Keys](#)

Launching Cloud Volumes ONTAP in AWS

You can launch Cloud Volumes ONTAP in a single-system configuration or as an HA pair in AWS.

Before you get started

You need the following to create a working environment.

- A Connector that's up and running.
 - You should have a [Connector that is associated with your workspace](#).
 - [You should be prepared to leave the Connector running at all times](#).
- An understanding of the configuration that you want to use.

You should have prepared by choosing a configuration and by obtaining AWS networking information from your administrator. For details, see [Planning your Cloud Volumes ONTAP configuration](#).

- DNS and Active Directory for CIFS configurations.

For details, see [Networking requirements for Cloud Volumes ONTAP in AWS](#).

- An understanding of what's required to choose a specific licensing option in the Add Working Environment wizard.

Licensing option	Requirement	How to meet the requirement
PAYGO free trial	A Marketplace subscription is required.	You'll have the option to subscribe to your cloud provider's marketplace from the Details & Credentials page.
Freemium	A Marketplace subscription or NetApp Support Site (NSS) account is required.	<p>You'll have the option to subscribe to your cloud provider's marketplace from the Details & Credentials page</p> <p>You can enter your NSS account on the Charging Methods and NSS Account page.</p>

Licensing option	Requirement	How to meet the requirement
Capacity-based BYOL	A Marketplace subscription or NetApp Support Site (NSS) account is required. A Marketplace subscription is recommended for capacity-based charging in the event that your account doesn't have a valid capacity-based license, or in the event that your provisioned capacity exceeds the licensed capacity.	You'll have the option to subscribe to your cloud provider's marketplace from the Details & Credentials page You can enter your NSS account on the Charging Methods and NSS Account page.
Node-based BYOL	The 20-digit serial number (license key) is required.	You'll enter the serial number on the Charging Methods and NSS Account page.

Launching a single-node Cloud Volumes ONTAP system in AWS

If you want to launch Cloud Volumes ONTAP in AWS, you need to create a new working environment in Cloud Manager.

About this task

Immediately after you create the working environment, Cloud Manager launches a test instance in the specified VPC to verify connectivity. If successful, Cloud Manager immediately terminates the instance and then starts deploying the Cloud Volumes ONTAP system. If Cloud Manager cannot verify connectivity, creation of the working environment fails. The test instance is either a t2.nano (for default VPC tenancy) or m3.medium (for dedicated VPC tenancy).

Steps

1. On the Canvas page, click **Add Working Environment** and follow the prompts.
2. **Choose a Location:** Select **Amazon Web Services** and **Cloud Volumes ONTAP Single Node**.
3. If you're prompted, [create a Connector](#).
4. **Details and Credentials:** Optionally change the AWS credentials and subscription, enter a working environment name, add tags if needed, and then enter a password.

Some of the fields in this page are self-explanatory. The following table describes fields for which you might need guidance:

Field	Description
Working Environment Name	Cloud Manager uses the working environment name to name both the Cloud Volumes ONTAP system and the Amazon EC2 instance. It also uses the name as the prefix for the predefined security group, if you select that option.

Field	Description
Add tags	<p>AWS tags are metadata for your AWS resources. Cloud Manager adds the tags to the Cloud Volumes ONTAP instance and each AWS resource associated with the instance.</p> <p>You can add up to four tags from the user interface when creating a working environment, and then you can add more after its created. Note that the API does not limit you to four tags when creating a working environment.</p> <p>For information about tags, refer to AWS Documentation: Tagging your Amazon EC2 Resources.</p>
User name and password	<p>These are the credentials for the Cloud Volumes ONTAP cluster admin account. You can use these credentials to connect to Cloud Volumes ONTAP through OnCommand System Manager or its CLI.</p>
Edit Credentials	<p>Choose the AWS credentials associated with the account where you want to deploy this system. You can also associate the AWS Marketplace subscription to use with this Cloud Volumes ONTAP system.</p> <p>Click Add Subscription to associate the selected credentials with a new AWS Marketplace subscription. The subscription can be for an annual contract or to pay for Cloud Volumes ONTAP at an hourly rate.</p> <p>Learn how to add additional AWS credentials to Cloud Manager.</p>

The following video shows how to associate a pay-as-you-go Marketplace subscription to your AWS credentials:

- ▶ https://docs.netapp.com/us-en/occm//media/video_subscribing_aws.mp4 (video)

If multiple IAM users work in the same AWS account, then each user needs to subscribe. After the first user subscribes, the AWS Marketplace informs subsequent users that they're already subscribed, as shown in the image below. While a subscription is in place for the AWS account, each IAM user needs to associate themselves with that subscription. If you see the message shown below, click the **click here** link to go to Cloud Central and complete the process.



Cloud Manager (for Cloud Volumes ONTAP)

You are currently subscribed to this product and will be charged for your accumulated usage at the end of your next billing cycle, based on the costs listed in Pricing information on the right.



Having issues signing up for your product?

If you were unable to complete the set-up process for this software, please [click here](#) to be taken to the product's registration area.

[Subscribe](#)

You are already subscribed to this product

Pricing Details

Software Fees

5. **Services:** Keep the services enabled or disable the individual services that you don't want to use with Cloud Volumes ONTAP.

- [Learn more about Cloud Data Sense](#).
- [Learn more about Cloud Backup](#).
- [Learn more about Monitoring](#).

6. Location & Connectivity: Enter the network information that you recorded in the [AWS worksheet](#).

If you have an AWS Outpost, you can deploy a single node Cloud Volumes ONTAP system in that Outpost by selecting the Outpost VPC. The experience is the same as any other VPC that resides in AWS.

The following image shows the page filled out:

Location	Connectivity
AWS Region	Security Group
US West Oregon	<input checked="" type="radio"/> Generated security group <input type="radio"/> Use existing security group
VPC	SSH Authentication Method
vpc-3a01e05f - 172.31.0.0/16	<input checked="" type="radio"/> Password <input type="radio"/> Key Pair
Subnet	
172.31.5.0/24 (OCCM subnet)	

7. Data Encryption: Choose no data encryption or AWS-managed encryption.

For AWS-managed encryption, you can choose a different Customer Master Key (CMK) from your account or another AWS account.



You can't change the AWS data encryption method after you create a Cloud Volumes ONTAP system.

[Learn how to set up the AWS KMS for Cloud Volumes ONTAP.](#)

[Learn more about supported encryption technologies.](#)

8. Charging Methods and NSS Account: Specify which charging option would you like to use with this system, and then specify a NetApp Support Site account.

- [Learn about these charging methods.](#)
- [Learn what's required in the wizard for the licensing method that you want to use.](#)

9. Cloud Volumes ONTAP Configuration (annual AWS Marketplace contract only): Review the default configuration and click **Continue** or click **Change Configuration** to select your own configuration.

If you keep the default configuration, then you only need to specify a volume and then review and approve the configuration.

10. Preconfigured Packages (hourly or BYOL only): Select one of the packages to quickly launch Cloud Volumes ONTAP, or click **Change Configuration** to select your own configuration.

If you choose one of the packages, then you only need to specify a volume and then review and approve the configuration.

11. IAM Role: It's best to keep the default option to let Cloud Manager create the role for you.

If you prefer to use your own policy, it must meet [policy requirements for Cloud Volumes ONTAP nodes](#).

12. **Licensing:** Change the Cloud Volumes ONTAP version as needed, select a license, an instance type, and the instance tenancy.

If your needs change after you launch the instance, you can modify the license or instance type later.



If a newer Release Candidate, General Availability, or patch release is available for the selected version, then Cloud Manager updates the system to that version when creating the working environment. For example, the update occurs if you select Cloud Volumes ONTAP 9.6 RC1 and 9.6 GA is available. The update does not occur from one release to another—for example, from 9.6 to 9.7.

13. **Underlying Storage Resources:** Choose settings for the initial aggregate: a disk type, a size for each disk, and whether data tiering should be enabled.

Note the following:

- The disk type is for the initial volume. You can choose a different disk type for subsequent volumes.
- The disk size is for all disks in the initial aggregate and for any additional aggregates that Cloud Manager creates when you use the simple provisioning option. You can create aggregates that use a different disk size by using the advanced allocation option.

For help choosing a disk type and size, see [Sizing your system in AWS](#).

- You can choose a specific volume tiering policy when you create or edit a volume.
- If you disable data tiering, you can enable it on subsequent aggregates.

[Learn how data tiering works.](#)

14. **Write Speed & WORM:** Choose **Normal** or **High** write speed, and activate write once, read many (WORM) storage, if desired.

[Learn more about write speed.](#)

WORM can't be enabled if Cloud Backup was enabled or if data tiering was enabled.

[Learn more about WORM storage.](#)

15. **Create Volume:** Enter details for the new volume or click **Skip**.

Some of the fields in this page are self-explanatory. The following table describes fields for which you might need guidance:

Field	Description
Size	The maximum size that you can enter largely depends on whether you enable thin provisioning, which enables you to create a volume that is bigger than the physical storage currently available to it.
Access control (for NFS only)	An export policy defines the clients in the subnet that can access the volume. By default, Cloud Manager enters a value that provides access to all instances in the subnet.

Field	Description
Permissions and Users / Groups (for CIFS only)	These fields enable you to control the level of access to a share for users and groups (also called access control lists or ACLs). You can specify local or domain Windows users or groups, or UNIX users or groups. If you specify a domain Windows user name, you must include the user's domain using the format domain\username.
Snapshot Policy	A Snapshot copy policy specifies the frequency and number of automatically created NetApp Snapshot copies. A NetApp Snapshot copy is a point-in-time file system image that has no performance impact and requires minimal storage. You can choose the default policy or none. You might choose none for transient data: for example, tempdb for Microsoft SQL Server.
Advanced options (for NFS only)	Select an NFS version for the volume: either NFSv3 or NFSv4.
Initiator group and IQN (for iSCSI only)	<p>iSCSI storage targets are called LUNs (logical units) and are presented to hosts as standard block devices.</p> <p>Initiator groups are tables of iSCSI host node names and control which initiators have access to which LUNs.</p> <p>iSCSI targets connect to the network through standard Ethernet network adapters (NICs), TCP offload engine (TOE) cards with software initiators, converged network adapters (CNAs) or dedicated host bus adapters (HBAs) and are identified by iSCSI qualified names (IQNs).</p> <p>When you create an iSCSI volume, Cloud Manager automatically creates a LUN for you. We've made it simple by creating just one LUN per volume, so there's no management involved. After you create the volume, use the IQN to connect to the LUN from your hosts.</p>

The following image shows the Volume page filled out for the CIFS protocol:

Volume Details, Protection & Protocol

<h4>Details & Protection</h4> <div style="margin-top: 10px;"> Volume Name: <input type="text" value="vol"/> Size (GB): <input type="text" value="250"/> </div> <div style="margin-top: 10px;"> Snapshot Policy: <input type="text" value="default"/> <div style="float: right;"><small>Default Policy</small></div> </div>	<h4>Protocol</h4> <div style="margin-top: 10px;"> <input checked="" type="radio"/> CIFS <input type="radio"/> NFS <input type="radio"/> iSCSI </div> <div style="margin-top: 10px;"> Share name: <input type="text" value="vol_share"/> Permissions: <input type="text" value="Full Control"/> </div> <div style="margin-top: 10px;"> Users / Groups: <input type="text" value="engineering"/> <div style="font-size: small; margin-top: -10px;">Valid users and groups separated by a semicolon</div> </div>
--	--

16. **CIFS Setup:** If you chose the CIFS protocol, set up a CIFS server.

Field	Description
DNS Primary and Secondary IP Address	The IP addresses of the DNS servers that provide name resolution for the CIFS server. The listed DNS servers must contain the service location records (SRV) needed to locate the Active Directory LDAP servers and domain controllers for the domain that the CIFS server will join.
Active Directory Domain to join	The FQDN of the Active Directory (AD) domain that you want the CIFS server to join.
Credentials authorized to join the domain	The name and password of a Windows account with sufficient privileges to add computers to the specified Organizational Unit (OU) within the AD domain.
CIFS server NetBIOS name	A CIFS server name that is unique in the AD domain.
Organizational Unit	The organizational unit within the AD domain to associate with the CIFS server. The default is CN=Computers. If you configure AWS Managed Microsoft AD as the AD server for Cloud Volumes ONTAP, you should enter OU=Computers,OU=corp in this field.
DNS Domain	The DNS domain for the Cloud Volumes ONTAP storage virtual machine (SVM). In most cases, the domain is the same as the AD domain.
NTP Server	Select Use Active Directory Domain to configure an NTP server using the Active Directory DNS. If you need to configure an NTP server using a different address, then you should use the API. See the Cloud Manager automation docs for details.

17. **Usage Profile, Disk Type, and Tiering Policy:** Choose whether you want to enable storage efficiency features and edit the volume tiering policy, if needed.

For more information, see [Understanding volume usage profiles](#) and [Data tiering overview](#).

18. **Review & Approve:** Review and confirm your selections.

- Review details about the configuration.
- Click **More information** to review details about support and the AWS resources that Cloud Manager will purchase.
- Select the **I understand...** check boxes.
- Click **Go**.

Result

Cloud Manager launches the Cloud Volumes ONTAP instance. You can track the progress in the timeline.

If you experience any issues launching the Cloud Volumes ONTAP instance, review the failure message. You can also select the working environment and click Re-create environment.

For additional help, go to [NetApp Cloud Volumes ONTAP Support](#).

After you finish

- If you provisioned a CIFS share, give users or groups permissions to the files and folders and verify that those users can access the share and create a file.
- If you want to apply quotas to volumes, use System Manager or the CLI.

Quotas enable you to restrict or track the disk space and number of files used by a user, group, or qtree.

Launching a Cloud Volumes ONTAP HA pair in AWS

If you want to launch a Cloud Volumes ONTAP HA pair in AWS, you need to create an HA working environment in Cloud Manager.

Limitation

At this time, HA pairs are not supported with AWS Outposts.

About this task

Immediately after you create the working environment, Cloud Manager launches a test instance in the specified VPC to verify connectivity. If successful, Cloud Manager immediately terminates the instance and then starts deploying the Cloud Volumes ONTAP system. If Cloud Manager cannot verify connectivity, creation of the working environment fails. The test instance is either a t2.nano (for default VPC tenancy) or m3.medium (for dedicated VPC tenancy).

Steps

1. On the Canvas page, click **Add Working Environment** and follow the prompts.
2. **Choose a Location:** Select **Amazon Web Services** and **Cloud Volumes ONTAP Single Node**.
3. **Details and Credentials:** Optionally change the AWS credentials and subscription, enter a working environment name, add tags if needed, and then enter a password.

Some of the fields in this page are self-explanatory. The following table describes fields for which you might need guidance:

Field	Description
Working Environment Name	Cloud Manager uses the working environment name to name both the Cloud Volumes ONTAP system and the Amazon EC2 instance. It also uses the name as the prefix for the predefined security group, if you select that option.
Add tags	AWS tags are metadata for your AWS resources. Cloud Manager adds the tags to the Cloud Volumes ONTAP instance and each AWS resource associated with the instance. You can add up to four tags from the user interface when creating a working environment, and then you can add more after its created. Note that the API does not limit you to four tags when creating a working environment. For information about tags, refer to AWS Documentation: Tagging your Amazon EC2 Resources .
User name and password	These are the credentials for the Cloud Volumes ONTAP cluster admin account. You can use these credentials to connect to Cloud Volumes ONTAP through OnCommand System Manager or its CLI.

Field	Description
Edit Credentials	<p>Choose the AWS credentials and marketplace subscription to use with this Cloud Volumes ONTAP system.</p> <p>Click Add Subscription to associate the selected credentials with a new AWS Marketplace subscription. The subscription can be for an annual contract or to pay for Cloud Volumes ONTAP at an hourly rate.</p> <p>If purchased a license directly from NetApp (BYOL), then an AWS subscription isn't required.</p> <p>Learn how to add additional AWS credentials to Cloud Manager.</p>

The following video shows how to associate a pay-as-you-go Marketplace subscription to your AWS credentials:

- ▶ https://docs.netapp.com/us-en/occm//media/video_subscribing_aws.mp4 (video)



If multiple IAM users work in the same AWS account, then each user needs to subscribe. After the first user subscribes, the AWS Marketplace informs subsequent users that they're already subscribed, as shown in the image below. While a subscription is in place for the AWS account, each IAM user needs to associate themselves with that subscription. If you see the message shown below, click the **click here** link to go to Cloud Central and complete the process.

4. **Services:** Keep the services enabled or disable the individual services that you don't want to use with this Cloud Volumes ONTAP system.
 - [Learn more about Cloud Data Sense.](#)
 - [Learn more about Cloud Backup.](#)
 - [Learn more about Monitoring.](#)

5. **HA Deployment Models:** Choose an HA configuration.

For an overview of the deployment models, see [Cloud Volumes ONTAP HA for AWS](#).

6. **Region & VPC:** Enter the network information that you recorded in the AWS worksheet.

The following image shows the page filled out for a multiple AZ configuration:

Region & VPC

AWS Region	VPC	Security group
US East N. Virginia	vpc-a76d91c2 - 172.31.0.0/16	Use a generated security group
<div style="display: flex; justify-content: space-around;"> <div style="width: 33%;"> Node 1: Availability Zone us-east-1a </div> <div style="width: 33%;"> Node 2: Availability Zone us-east-1b </div> <div style="width: 33%;"> Mediator: Availability Zone us-east-1c </div> </div>		
Subnet	Subnet	Subnet
172.31.8.0/24	172.31.9.0/24	172.31.2.0/24

7. **Connectivity and SSH Authentication:** Choose connection methods for the HA pair and the mediator.
8. **Floating IPs:** If you chose multiple AZs, specify the floating IP addresses.

The IP addresses must be outside of the CIDR block for all VPCs in the region. For additional details, see [AWS networking requirements for Cloud Volumes ONTAP HA in multiple AZs](#).

9. **Route Tables:** If you chose multiple AZs, select the route tables that should include routes to the floating IP addresses.

If you have more than one route table, it is very important to select the correct route tables. Otherwise, some clients might not have access to the Cloud Volumes ONTAP HA pair. For more information about route tables, refer to [AWS Documentation: Route Tables](#).

10. **Data Encryption:** Choose no data encryption or AWS-managed encryption.

For AWS-managed encryption, you can choose a different Customer Master Key (CMK) from your account or another AWS account.



You can't change the AWS data encryption method after you create a Cloud Volumes ONTAP system.

[Learn how to set up the AWS KMS for Cloud Volumes ONTAP.](#)

[Learn more about supported encryption technologies.](#)

11. **Charging Methods and NSS Account:** Specify which charging option would you like to use with this system, and then specify a NetApp Support Site account.
 - [Learn about these charging methods.](#)
 - [Learn what's required in the wizard for the licensing method that you want to use.](#)
12. **Cloud Volumes ONTAP Configuration** (annual AWS Marketplace contract only): Review the default configuration and click **Continue** or click **Change Configuration** to select your own configuration.

If you keep the default configuration, then you only need to specify a volume and then review and approve the configuration.

13. **Preconfigured Packages** (hourly or BYOL only): Select one of the packages to quickly launch Cloud Volumes ONTAP, or click **Change Configuration** to select your own configuration.

If you choose one of the packages, then you only need to specify a volume and then review and approve the configuration.

14. **IAM Role**: It's best to keep the default option to let Cloud Manager create the role for you.

If you prefer to use your own policy, it must meet [policy requirements for Cloud Volumes ONTAP nodes and the HA mediator](#).

15. **Licensing**: Change the Cloud Volumes ONTAP version as needed, select a license, an instance type, and the instance tenancy.

If your needs change after you launch the instances, you can modify the license or instance type later.



If a newer Release Candidate, General Availability, or patch release is available for the selected version, then Cloud Manager updates the system to that version when creating the working environment. For example, the update occurs if you select Cloud Volumes ONTAP 9.6 RC1 and 9.6 GA is available. The update does not occur from one release to another—for example, from 9.6 to 9.7.

16. **Underlying Storage Resources**: Choose settings for the initial aggregate: a disk type, a size for each disk, and whether data tiering should be enabled.

Note the following:

- The disk type is for the initial volume. You can choose a different disk type for subsequent volumes.
- The disk size is for all disks in the initial aggregate and for any additional aggregates that Cloud Manager creates when you use the simple provisioning option. You can create aggregates that use a different disk size by using the advanced allocation option.

For help choosing a disk type and size, see [Sizing your system in AWS](#).

- You can choose a specific volume tiering policy when you create or edit a volume.
- If you disable data tiering, you can enable it on subsequent aggregates.

[Learn how data tiering works](#).

17. **Write Speed & WORM**: Choose **Normal** or **High** write speed, and activate write once, read many (WORM) storage, if desired.

[Learn more about write speed](#).

WORM can't be enabled if Cloud Backup was enabled or if data tiering was enabled.

[Learn more about WORM storage](#).

18. **Create Volume**: Enter details for the new volume or click **Skip**.

Some of the fields in this page are self-explanatory. The following table describes fields for which you might need guidance:

Field	Description
Size	The maximum size that you can enter largely depends on whether you enable thin provisioning, which enables you to create a volume that is bigger than the physical storage currently available to it.
Access control (for NFS only)	An export policy defines the clients in the subnet that can access the volume. By default, Cloud Manager enters a value that provides access to all instances in the subnet.
Permissions and Users / Groups (for CIFS only)	These fields enable you to control the level of access to a share for users and groups (also called access control lists or ACLs). You can specify local or domain Windows users or groups, or UNIX users or groups. If you specify a domain Windows user name, you must include the user's domain using the format domain\username.
Snapshot Policy	A Snapshot copy policy specifies the frequency and number of automatically created NetApp Snapshot copies. A NetApp Snapshot copy is a point-in-time file system image that has no performance impact and requires minimal storage. You can choose the default policy or none. You might choose none for transient data: for example, tempdb for Microsoft SQL Server.
Advanced options (for NFS only)	Select an NFS version for the volume: either NFSv3 or NFSv4.
Initiator group and IQN (for iSCSI only)	<p>iSCSI storage targets are called LUNs (logical units) and are presented to hosts as standard block devices.</p> <p>Initiator groups are tables of iSCSI host node names and control which initiators have access to which LUNs.</p> <p>iSCSI targets connect to the network through standard Ethernet network adapters (NICs), TCP offload engine (TOE) cards with software initiators, converged network adapters (CNAs) or dedicated host bus adapters (HBAs) and are identified by iSCSI qualified names (IQNs).</p> <p>When you create an iSCSI volume, Cloud Manager automatically creates a LUN for you. We've made it simple by creating just one LUN per volume, so there's no management involved. After you create the volume, use the IQN to connect to the LUN from your hosts.</p>

The following image shows the Volume page filled out for the CIFS protocol:

Volume Details, Protection & Protocol

Details & Protection <div style="margin-top: 10px;"> Volume Name: <input type="text" value="vol"/> Size (GB): <input type="text" value="250"/> </div> <div style="margin-top: 10px;"> Snapshot Policy: <input type="text" value="default"/> <div style="float: right;">(i) Default Policy</div> </div>	Protocol <div style="margin-top: 10px;"> NFS CIFS iSCSI </div> <div style="margin-top: 10px;"> Share name: <input type="text" value="vol_share"/> Permissions: <input type="text" value="Full Control"/> </div> <div style="margin-top: 10px;"> Users / Groups: <input type="text" value="engineering"/> <div style="font-size: small; margin-top: -10px;">Valid users and groups separated by a semicolon</div> </div>
---	---

19. **CIFS Setup:** If you selected the CIFS protocol, set up a CIFS server.

Field	Description
DNS Primary and Secondary IP Address	The IP addresses of the DNS servers that provide name resolution for the CIFS server. The listed DNS servers must contain the service location records (SRV) needed to locate the Active Directory LDAP servers and domain controllers for the domain that the CIFS server will join.
Active Directory Domain to join	The FQDN of the Active Directory (AD) domain that you want the CIFS server to join.
Credentials authorized to join the domain	The name and password of a Windows account with sufficient privileges to add computers to the specified Organizational Unit (OU) within the AD domain.
CIFS server NetBIOS name	A CIFS server name that is unique in the AD domain.
Organizational Unit	The organizational unit within the AD domain to associate with the CIFS server. The default is CN=Computers. If you configure AWS Managed Microsoft AD as the AD server for Cloud Volumes ONTAP, you should enter OU=Computers,OU=corp in this field.
DNS Domain	The DNS domain for the Cloud Volumes ONTAP storage virtual machine (SVM). In most cases, the domain is the same as the AD domain.
NTP Server	Select Use Active Directory Domain to configure an NTP server using the Active Directory DNS. If you need to configure an NTP server using a different address, then you should use the API. See the Cloud Manager automation docs for details.

20. **Usage Profile, Disk Type, and Tiering Policy:** Choose whether you want to enable storage efficiency features and edit the volume tiering policy, if needed.

For more information, see [Understanding volume usage profiles](#) and [Data tiering overview](#).

21. **Review & Approve:** Review and confirm your selections.

- Review details about the configuration.

- b. Click **More information** to review details about support and the AWS resources that Cloud Manager will purchase.
- c. Select the **I understand...** check boxes.
- d. Click **Go**.

Result

Cloud Manager launches the Cloud Volumes ONTAP HA pair. You can track the progress in the timeline.

If you experience any issues launching the HA pair, review the failure message. You can also select the working environment and click Re-create environment.

For additional help, go to [NetApp Cloud Volumes ONTAP Support](#).

After you finish

- If you provisioned a CIFS share, give users or groups permissions to the files and folders and verify that those users can access the share and create a file.
- If you want to apply quotas to volumes, use System Manager or the CLI.

Quotas enable you to restrict or track the disk space and number of files used by a user, group, or qtree.

Get started with Cloud Volumes ONTAP in the AWS C2S environment

Similar to a standard AWS region, you can use Cloud Manager in the [AWS Commercial Cloud Services \(C2S\)](#) environment to deploy Cloud Volumes ONTAP, which provides enterprise-class features for your cloud storage. AWS C2S is a closed region specific to the U.S. Intelligence Community; the instructions on this page only apply to AWS C2S region users.

Supported features in C2S

The following features are available from Cloud Manager in the C2S environment:

- Cloud Volumes ONTAP
- Data replication
- A timeline for auditing

For Cloud Volumes ONTAP, you can create a single node system or an HA pair. Both licensing options are available: pay-as-you-go and bring your own license (BYOL).

Data tiering to S3 is also supported with Cloud Volumes ONTAP in C2S.

Limitations

None of NetApp's cloud services are available from Cloud Manager.

Because there's no internet access in the C2S environment, the following features aren't available either:

- Integration with NetApp Cloud Central
- Automated software upgrades from Cloud Manager

- NetApp AutoSupport
- AWS cost information for Cloud Volumes ONTAP resources

Deployment overview

Getting started with Cloud Volumes ONTAP in C2S includes a few steps.

1. Preparing your AWS environment.

This includes setting up networking, subscribing to Cloud Volumes ONTAP, setting up permissions, and optionally setting up the AWS KMS.

2. Installing the Connector and setting up Cloud Manager.

Before you can start using Cloud Manager to deploy Cloud Volumes ONTAP, you'll need to create a *Connector*. The Connector enables Cloud Manager to manage resources and processes within your public cloud environment (this includes Cloud Volumes ONTAP).

You'll log in to Cloud Manager from the software that gets installed on the Connector instance.

3. Launching Cloud Volumes ONTAP from Cloud Manager.

Each of these steps are described below.

Prepare your AWS environment

Your AWS environment must meet a few requirements.

Set up your networking

Set up your AWS networking so Cloud Volumes ONTAP can operate properly.

Steps

1. Choose the VPC and subnets in which you want to launch the Connector instance and Cloud Volumes ONTAP instances.
2. Ensure that your VPC and subnets will support connectivity between the Connector and Cloud Volumes ONTAP.
3. Set up a VPC endpoint to the S3 service.

A VPC endpoint is required if you want to tier cold data from Cloud Volumes ONTAP to low-cost object storage.

Subscribe to Cloud Volumes ONTAP

A Marketplace subscription is required to deploy Cloud Volumes ONTAP from Cloud Manager.

Steps

1. Go to the AWS Intelligence Community Marketplace and search for Cloud Volumes ONTAP.
2. Select the offering that you plan to deploy.
3. Review the terms and click **Accept**.
4. Repeat these steps for the other offerings, if you plan to deploy them.

You must use Cloud Manager to launch Cloud Volumes ONTAP instances. You must not launch Cloud Volumes ONTAP instances from the EC2 console.

Set up permissions

Set up IAM policies and roles that provide Cloud Manager and Cloud Volumes ONTAP with the permissions that they need to perform actions in the AWS Commercial Cloud Services environment.

You need an IAM policy and IAM role for each of the following:

- The Connector instance
- Cloud Volumes ONTAP instances
- The Cloud Volumes ONTAP HA mediator instance (if you want to deploy HA pairs)

Steps

1. Go to the AWS IAM console and click **Policies**.
2. Create a policy for the Connector instance.

```
{  
    "Version": "2012-10-17",  
    "Statement": [{  
        "Effect": "Allow",  
        "Action": [  
            "ec2:DescribeInstances",  
            "ec2:DescribeInstanceStatus",  
            "ec2:RunInstances",  
            "ec2:ModifyInstanceAttribute",  
            "ec2:DescribeRouteTables",  
            "ec2:DescribeImages",  
            "ec2:CreateTags",  
            "ec2:CreateVolume",  
            "ec2:DescribeVolumes",  
            "ec2:ModifyVolumeAttribute",  
            "ec2:DeleteVolume",  
            "ec2:CreateSecurityGroup",  
            "ec2:DeleteSecurityGroup",  
            "ec2:DescribeSecurityGroups",  
            "ec2:RevokeSecurityGroupEgress",  
            "ec2:RevokeSecurityGroupIngress",  
            "ec2:AuthorizeSecurityGroupEgress",  
            "ec2:AuthorizeSecurityGroupIngress",  
            "ec2>CreateNetworkInterface",  
            "ec2:DescribeNetworkInterfaces",  
            "ec2:DeleteNetworkInterface",  
            "ec2:ModifyNetworkInterfaceAttribute",  
            "ec2:DescribeSubnets",  
            "ec2:DescribeVpcs",  
        ]  
    }]  
}
```

```

    "ec2:DescribeDhcpOptions",
    "ec2>CreateSnapshot",
    "ec2>DeleteSnapshot",
    "ec2:DescribeSnapshots",
    "ec2:GetConsoleOutput",
    "ec2:DescribeKeyPairs",
    "ec2:DescribeRegions",
    "ec2>DeleteTags",
    "ec2:DescribeTags",
    "cloudformation>CreateStack",
    "cloudformation>DeleteStack",
    "cloudformation:DescribeStacks",
    "cloudformation:DescribeStackEvents",
    "cloudformation:ValidateTemplate",
    "iam:PassRole",
    "iam>CreateRole",
    "iam>DeleteRole",
    "iam:PutRolePolicy",
    "iam>CreateInstanceProfile",
    "iam>DeleteRolePolicy",
    "iam>AddRoleToInstanceProfile",
    "iam:RemoveRoleFromInstanceProfile",
    "iam>DeleteInstanceProfile",
    "s3:GetObject",
    "s3>ListBucket",
    "s3:GetBucketTagging",
    "s3:GetBucketLocation",
    "s3>ListAllMyBuckets",
    "kms>List*",
    "kms:Describe*",
    "ec2:AssociateIamInstanceProfile",
    "ec2:DescribeIamInstanceProfileAssociations",
    "ec2:DisassociateIamInstanceProfile",
    "ec2:DescribeInstanceAttribute",
    "ec2>CreatePlacementGroup",
    "ec2>DeletePlacementGroup"
],
"Resource": "*"
},
{
  "Sid": "fabricPoolPolicy",
  "Effect": "Allow",
  "Action": [
    "s3>DeleteBucket",
    "s3:GetLifecycleConfiguration",
    "s3:PutLifecycleConfiguration",

```

```

        "s3:PutBucketTagging",
        "s3>ListBucketVersions"
    ],
    "Resource": [
        "arn:aws-iso:s3:::fabric-pool*"
    ]
},
{
    "Effect": "Allow",
    "Action": [
        "ec2:StartInstances",
        "ec2:StopInstances",
        "ec2:TerminateInstances",
        "ec2:AttachVolume",
        "ec2:DetachVolume"
    ],
    "Condition": {
        "StringLike": {
            "ec2:ResourceTag/WorkingEnvironment": "*"
        }
    },
    "Resource": [
        "arn:aws-iso:ec2:*:*:instance/*"
    ]
},
{
    "Effect": "Allow",
    "Action": [
        "ec2:AttachVolume",
        "ec2:DetachVolume"
    ],
    "Resource": [
        "arn:aws-iso:ec2:*:*:volume/*"
    ]
}
]
}

```

3. Create a policy for Cloud Volumes ONTAP.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Action": "s3>ListAllMyBuckets",
            "Resource": "arn:aws-iso:s3:::*",
            "Effect": "Allow"
        },
        {
            "Action": [
                "s3>ListBucket",
                "s3>GetBucketLocation"
            ],
            "Resource": "arn:aws-iso:s3:::fabric-pool-*",
            "Effect": "Allow"
        },
        {
            "Action": [
                "s3GetObject",
                "s3PutObject",
                "s3DeleteObject"
            ],
            "Resource": "arn:aws-iso:s3:::fabric-pool-*",
            "Effect": "Allow"
        }
    ]
}
```

4. If you plan to deploy a Cloud Volumes ONTAP HA pair, create a policy for the HA mediator.

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "ec2>AssignPrivateIpAddresses",
                "ec2>CreateRoute",
                "ec2>DeleteRoute",
                "ec2>DescribeNetworkInterfaces",
                "ec2>DescribeRouteTables",
                "ec2>DescribeVpcs",
                "ec2>ReplaceRoute",
                "ec2>UnassignPrivateIpAddresses"
            ],
            "Resource": "*"
        }
    ]
}
```

5. Create IAM roles with the role type Amazon EC2 and attach the policies that you created in the previous steps.

Similar to the policies, you should have one IAM role for the Connector, one for the Cloud Volumes ONTAP nodes, and one for the HA mediator (if you want to deploy HA pairs).

You must select the Connector IAM role when you launch the Connector instance.

You can select the IAM roles for Cloud Volumes ONTAP and the HA mediator when you create a Cloud Volumes ONTAP working environment from Cloud Manager.

Set up the AWS KMS

If you want to use Amazon encryption with Cloud Volumes ONTAP, ensure that requirements are met for the AWS Key Management Service.

Steps

1. Ensure that an active Customer Master Key (CMK) exists in your account or in another AWS account.

The CMK can be an AWS-managed CMK or a customer-managed CMK.

2. If the CMK is in an AWS account separate from the account where you plan to deploy Cloud Volumes ONTAP, then you need to obtain the ARN of that key.

You'll need to provide the ARN to Cloud Manager when you create the Cloud Volumes ONTAP system.

3. Add the IAM role for the Cloud Manager instance to the list of key users for a CMK.

This gives Cloud Manager permissions to use the CMK with Cloud Volumes ONTAP.

Install and set up Cloud Manager

Before you can launch Cloud Volumes ONTAP systems in AWS, you must first launch the Connector instance from the AWS Marketplace and then log in and set up Cloud Manager.

Steps

1. Obtain a root certificate signed by a certificate authority (CA) in the Privacy Enhanced Mail (PEM) Base-64 encoded X.509 format. Consult your organization's policies and procedures for obtaining the certificate.

You'll need to upload the certificate during the setup process. Cloud Manager uses the trusted certificate when sending requests to AWS over HTTPS.

2. Launch the Connector instance:

- a. Go to the AWS Intelligence Community Marketplace page for Cloud Manager.
- b. On the Custom Launch tab, choose the option to launch the instance from the EC2 console.
- c. Follow the prompts to configure the instance.

Note the following as you configure the instance:

- We recommend t3.xlarge.
- You must choose the IAM role that you created when preparing your AWS environment.

- You should keep the default storage options.
 - The required connection methods for the Connector are as follows: SSH, HTTP, and HTTPS.
3. Set up Cloud Manager from a host that has a connection to the Connector instance:
- a. Open a web browser and enter the following URL: <http://ipaddress:80>
 - b. Specify a proxy server for connectivity to AWS services.
 - c. Upload the certificate that you obtained in step 1.
 - d. Complete the steps in the Setup wizard to set up Cloud Manager.
 - **System Details:** Enter a name for this instance of Cloud Manager and provide your company name.
 - **Create User:** Create the Admin user that you'll use to administer Cloud Manager.
 - **Review:** Review the details and approve the end user license agreement.
 - e. To complete installation of the CA-signed certificate, restart the Connector instance from the EC2 console.
4. After the Connector restarts, log in using the administrator user account that you created in the Setup wizard.

Launch Cloud Volumes ONTAP

You can launch Cloud Volumes ONTAP instances in the AWS Commercial Cloud Services environment by creating new working environments in Cloud Manager.

What you'll need

- If you purchased a license, you must have the license file that you received from NetApp. The license file is a .NLF file in JSON format.
- A key pair is required to enable key-based SSH authentication to the HA mediator.

Steps

1. On the Working Environments page, click **Add Working Environment**.
2. Under Create, select Cloud Volumes ONTAP or Cloud Volumes ONTAP HA.
3. Complete the steps in the wizard to launch the Cloud Volumes ONTAP system.

Note the following as you complete the wizard:

- If you want to deploy Cloud Volumes ONTAP HA in multiple Availability Zones, deploy the configuration as follows because only two AZs were available in the AWS Commercial Cloud Services environment at the time of publication:
 - Node 1: Availability Zone A
 - Node 2: Availability Zone B
 - Mediator: Availability Zone A or B
- You should leave the default option to use a generated security group.

The predefined security group includes the rules that Cloud Volumes ONTAP needs to operate successfully. If you have a requirement to use your own, you can refer to the security group section below.

- You must choose the IAM role that you created when preparing your AWS environment.

- The underlying AWS disk type is for the initial Cloud Volumes ONTAP volume.
You can choose a different disk type for subsequent volumes.
 - The performance of AWS disks is tied to disk size.
You should choose the disk size that gives you the sustained performance that you need. Refer to AWS documentation for more details about EBS performance.
 - The disk size is the default size for all disks on the system.
-  If you need a different size later, you can use the Advanced allocation option to create an aggregate that uses disks of a specific size.
- Storage efficiency features can improve storage utilization and reduce the total amount of storage that you need.

Result

Cloud Manager launches the Cloud Volumes ONTAP instance. You can track the progress in the timeline.

Security group rules

Cloud Manager creates security groups that include the inbound and outbound rules that Cloud Manager and Cloud Volumes ONTAP need to operate successfully in the cloud. You might want to refer to the ports for testing purposes or if you prefer to use your own security groups.

Security group for the Connector

The security group for the Connector requires both inbound and outbound rules.

Inbound rules

Protocol	Port	Purpose
SSH	22	Provides SSH access to the Connector host
HTTP	80	Provides HTTP access from client web browsers to the local user interface
HTTPS	443	Provides HTTPS access from client web browsers to the local user interface

Outbound rules

The predefined security group for the Connector includes the following outbound rules.

Protocol	Port	Purpose
All TCP	All	All outbound traffic
All UDP	All	All outbound traffic

Security group for Cloud Volumes ONTAP

The security group for Cloud Volumes ONTAP nodes requires both inbound and outbound rules.

Inbound rules

The source for inbound rules in the predefined security group is 0.0.0.0/0.

Protocol	Port	Purpose
All ICMP	All	Pinging the instance
HTTP	80	HTTP access to the System Manager web console using the IP address of the cluster management LIF
HTTPS	443	HTTPS access to the System Manager web console using the IP address of the cluster management LIF
SSH	22	SSH access to the IP address of the cluster management LIF or a node management LIF
TCP	111	Remote procedure call for NFS
TCP	139	NetBIOS service session for CIFS
TCP	161-162	Simple network management protocol
TCP	445	Microsoft SMB/CIFS over TCP with NetBIOS framing
TCP	635	NFS mount
TCP	749	Kerberos
TCP	2049	NFS server daemon
TCP	3260	iSCSI access through the iSCSI data LIF
TCP	4045	NFS lock daemon
TCP	4046	Network status monitor for NFS
TCP	10000	Backup using NDMP
TCP	11104	Management of intercluster communication sessions for SnapMirror
TCP	11105	SnapMirror data transfer using intercluster LIFs
UDP	111	Remote procedure call for NFS
UDP	161-162	Simple network management protocol
UDP	635	NFS mount
UDP	2049	NFS server daemon
UDP	4045	NFS lock daemon
UDP	4046	Network status monitor for NFS
UDP	4049	NFS rquotad protocol

Outbound rules

The predefined security group for Cloud Volumes ONTAP includes the following outbound rules.

Protocol	Port	Purpose
All ICMP	All	All outbound traffic

Protocol	Port	Purpose
All TCP	All	All outbound traffic
All UDP	All	All outbound traffic

External security group for the HA mediator

The predefined external security group for the Cloud Volumes ONTAP HA mediator includes the following inbound and outbound rules.

Inbound rules

The source for inbound rules is traffic from the VPC where the Connector resides.

Protocol	Port	Purpose
SSH	22	SSH connections to the HA mediator
TCP	3000	RESTful API access from the Connector

Outbound rules

The predefined security group for the HA mediator includes the following outbound rules.

Protocol	Port	Purpose
All TCP	All	All outbound traffic
All UDP	All	All outbound traffic

Internal security group for the HA mediator

The predefined internal security group for the Cloud Volumes ONTAP HA mediator includes the following rules. Cloud Manager always creates this security group. You don't have the option to use your own.

Inbound rules

The predefined security group includes the following inbound rules.

Protocol	Port	Purpose
All traffic	All	Communication between the HA mediator and HA nodes

Outbound rules

The predefined security group includes the following outbound rules.

Protocol	Port	Purpose
All traffic	All	Communication between the HA mediator and HA nodes

Get started in Azure

Getting started with Cloud Volumes ONTAP for Azure

Get started with Cloud Volumes ONTAP for Azure in a few steps.



Create a Connector

If you don't have a [Connector](#) yet, an Account Admin needs to create one. [Learn how to create a Connector in Azure](#).

When you create your first Cloud Volumes ONTAP working environment, Cloud Manager prompts you to deploy a Connector if you don't have one yet.



Plan your configuration

Cloud Manager offers preconfigured packages that match your workload requirements, or you can create your own configuration. If you choose your own configuration, you should understand the options available to you. [Learn more](#).



Set up your networking

- a. Ensure that your VNet and subnets will support connectivity between the Connector and Cloud Volumes ONTAP.
- b. Enable outbound internet access from the target VNet so the Connector and Cloud Volumes ONTAP can contact several endpoints.

This step is important because the Connector can't manage Cloud Volumes ONTAP without outbound internet access. If you need to limit outbound connectivity, refer to the list of endpoints for [the Connector and Cloud Volumes ONTAP](#).

[Learn more about networking requirements](#).



Launch Cloud Volumes ONTAP using Cloud Manager

Click **Add Working Environment**, select the type of system that you would like to deploy, and complete the steps in the wizard. [Read step-by-step instructions](#).

Related links

- [Evaluating](#)
- [Creating a Connector from Cloud Manager](#)
- [Creating a Connector from the Azure Marketplace](#)
- [Installing the Connector software on a Linux host](#)
- [What Cloud Manager does with Azure permissions](#)

Planning your Cloud Volumes ONTAP configuration in Azure

When you deploy Cloud Volumes ONTAP in Azure, you can choose a preconfigured system that matches your workload requirements, or you can create your own configuration. If you choose your own configuration, you should understand the options available to you.

Choosing a license type

Cloud Volumes ONTAP is available in two pricing options: pay-as-you-go and Bring Your Own License (BYOL). For pay-as-you-go, you can choose from three licenses: Explore, Standard, or Premium. Each license provides different capacity and compute options.

[Supported configurations for Cloud Volumes ONTAP in Azure](#)

Supported VM types

Cloud Volumes ONTAP supports several VM types, depending on the license type that you choose.

[Supported configurations for Cloud Volumes ONTAP in Azure](#)

Understanding storage limits

The raw capacity limit for a Cloud Volumes ONTAP system is tied to the license. Additional limits impact the size of aggregates and volumes. You should be aware of these limits as you plan your configuration.

[Storage limits for Cloud Volumes ONTAP in Azure](#)

Sizing your system in Azure

Sizing your Cloud Volumes ONTAP system can help you meet requirements for performance and capacity. You should be aware of a few key points when choosing a VM type, disk type, and disk size:

Virtual machine type

Look at the supported virtual machine types in the [Cloud Volumes ONTAP Release Notes](#) and then review details about each supported VM type. Be aware that each VM type supports a specific number of data disks.

- [Azure documentation: General purpose virtual machine sizes](#)
- [Azure documentation: Memory optimized virtual machine sizes](#)

Azure disk type

When you create volumes for Cloud Volumes ONTAP, you need to choose the underlying cloud storage that Cloud Volumes ONTAP uses as a disk.

HA systems use Premium page blobs. Meanwhile, single node systems can use two types of Azure Managed Disks:

- *Premium SSD Managed Disks* provide high performance for I/O-intensive workloads at a higher cost.
- *Standard SSD Managed Disks* provide consistent performance for workloads that require low IOPS.
- *Standard HDD Managed Disks* are a good choice if you don't need high IOPS and want to reduce your costs.

For additional details about the use cases for these disks, see [Microsoft Azure Documentation: What disk types are available in Azure?](#).

Azure disk size

When you launch Cloud Volumes ONTAP instances, you must choose the default disk size for aggregates. Cloud Manager uses this disk size for the initial aggregate, and for any additional aggregates that it creates when you use the simple provisioning option. You can create aggregates that use a disk size different from the default by [using the advanced allocation option](#).



All disks in an aggregate must be the same size.

When choosing a disk size, you should take several factors into consideration. The disk size impacts how much you pay for storage, the size of volumes that you can create in an aggregate, the total capacity available to Cloud Volumes ONTAP, and storage performance.

The performance of Azure Premium Storage is tied to the disk size. Larger disks provide higher IOPS and throughput. For example, choosing 1 TiB disks can provide better performance than 500 GiB disks, at a higher cost.

There are no performance differences between disk sizes for Standard Storage. You should choose disk size based on the capacity that you need.

Refer to Azure for IOPS and throughput by disk size:

- [Microsoft Azure: Managed Disks pricing](#)
- [Microsoft Azure: Page Blobs pricing](#)

Choosing a configuration that supports Flash Cache

A Cloud Volumes ONTAP configuration in Azure includes local NVMe storage, which Cloud Volumes ONTAP uses as *Flash Cache* for better performance. [Learn more about Flash Cache](#).

Azure network information worksheet

When you deploy Cloud Volumes ONTAP in Azure, you need to specify details about your virtual network. You can use a worksheet to collect the information from your administrator.

Azure information	Your value
Region	
Virtual network (VNet)	
Subnet	
Network security group (if using your own)	

Choosing a write speed

Cloud Manager enables you to choose a write speed setting for Cloud Volumes ONTAP. Before you choose a write speed, you should understand the differences between the normal and high settings and risks and recommendations when using high write speed. [Learn more about write speed](#).

Choosing a volume usage profile

ONTAP includes several storage efficiency features that can reduce the total amount of storage that you need. When you create a volume in Cloud Manager, you can choose a profile that enables these features or a profile that disables them. You should learn more about these features to help you decide which profile to use.

NetApp storage efficiency features provide the following benefits:

Thin provisioning

Presents more logical storage to hosts or users than you actually have in your physical storage pool. Instead of preallocating storage space, storage space is allocated dynamically to each volume as data is written.

Deduplication

Improves efficiency by locating identical blocks of data and replacing them with references to a single shared block. This technique reduces storage capacity requirements by eliminating redundant blocks of data that reside in the same volume.

Compression

Reduces the physical capacity required to store data by compressing data within a volume on primary, secondary, and archive storage.

Networking requirements for Cloud Volumes ONTAP in Azure

Set up your Azure networking so Cloud Volumes ONTAP systems can operate properly. This includes networking for the Connector and Cloud Volumes ONTAP.

Requirements for Cloud Volumes ONTAP

The following networking requirements must be met in Azure.

Outbound internet access for Cloud Volumes ONTAP

Cloud Volumes ONTAP requires outbound internet access to send messages to NetApp AutoSupport, which proactively monitors the health of your storage.

Routing and firewall policies must allow HTTP/HTTPS traffic to the following endpoints so Cloud Volumes ONTAP can send AutoSupport messages:

- <https://support.netapp.com/aods/asupmessage>
- <https://support.netapp.com/asupprod/post/1.0/postAsup>

[Learn how to configure AutoSupport.](#)

Security groups

You do not need to create security groups because Cloud Manager does that for you. If you need to use your own, refer to the security group rules listed below.

Number of IP addresses

Cloud Manager allocates the following number of IP addresses to Cloud Volumes ONTAP in Azure:

- Single node: 5 IP addresses
- HA pair: 16 IP addresses

Note that Cloud Manager creates an SVM management LIF on HA pairs, but not on single node systems in Azure.



A LIF is an IP address associated with a physical port. An SVM management LIF is required for management tools like SnapCenter.

Connection from Cloud Volumes ONTAP to Azure Blob storage for data tiering

If you want to tier cold data to Azure Blob storage, you don't need to set up a connection between the performance tier and the capacity tier as long as Cloud Manager has the required permissions. Cloud Manager enables a VNet service endpoint for you if the Cloud Manager policy has these permissions:

```
"Microsoft.Network/virtualNetworks/subnets/write",  
"Microsoft.Network/routeTables/join/action",
```

These permissions are included in the latest [Cloud Manager policy](#).

For details about setting up data tiering, see [Tiering cold data to low-cost object storage](#).

Connections to ONTAP systems in other networks

To replicate data between a Cloud Volumes ONTAP system in Azure and ONTAP systems in other networks, you must have a VPN connection between the Azure VNet and the other network—for example, an AWS VPC or your corporate network.

For instructions, refer to [Microsoft Azure Documentation: Create a Site-to-Site connection in the Azure portal](#).

Requirements for the Connector

Set up your networking so that the Connector can manage resources and processes within your public cloud environment. The most important step is ensuring outbound internet access to various endpoints.



If your network uses a proxy server for all communication to the internet, you can specify the proxy server from the Settings page. Refer to [Configuring the Connector to use a proxy server](#).

Connections to target networks

A Connector requires a network connection to the VPCs and VNets in which you want to deploy Cloud Volumes ONTAP.

For example, if you install a Connector in your corporate network, then you must set up a VPN connection to the VPC or VNet in which you launch Cloud Volumes ONTAP.

Outbound internet access

The Connector requires outbound internet access to manage resources and processes within your public cloud environment. A Connector contacts the following endpoints when managing resources in Azure:

Endpoints	Purpose
https://management.azure.com	Enables Cloud Manager to deploy and manage Cloud Volumes ONTAP in most Azure regions.
https://login.microsoftonline.com	

Endpoints	Purpose
https://management.microsoftazure.de https://login.microsoftonline.de	Enables Cloud Manager to deploy and manage Cloud Volumes ONTAP in the Azure Germany regions.
https://management.usgovcloudapi.net https://login.microsoftonline.com	Enables Cloud Manager to deploy and manage Cloud Volumes ONTAP in the Azure US Gov regions.
https://api.services.cloud.netapp.com:443	API requests to NetApp Cloud Central.
https://cloud.support.netapp.com.s3.us-west-1.amazonaws.com	Provides access to software images, manifests, and templates.
https://cognito-idp.us-east-1.amazonaws.com https://cognito-identity.us-east-1.amazonaws.com https://sts.amazonaws.com https://cloud-support-netapp-com-accelerated.s3.amazonaws.com	Enables the Connector to access and download manifests, templates, and Cloud Volumes ONTAP upgrade images.
https://cloudmanagerinfraprod.azurecr.io *.blob.core.windows.net	Access to software images of container components for an infrastructure that's running Docker and provides a solution for service integrations with Cloud Manager.
https://kinesis.us-east-1.amazonaws.com	Enables NetApp to stream data from audit records.
https://cloudmanager.cloud.netapp.com	Communication with the Cloud Manager service, which includes Cloud Central accounts.
https://netapp-cloud-account.auth0.com	Communication with NetApp Cloud Central for centralized user authentication.
support.netapp.com:443 https://mysupport.netapp.com	Communication with NetApp AutoSupport. Note that the Connector communicates with support.netapp.com:443, which redirects to https://mysupport.netapp.com.
https://support.netapp.com/svcgw https://support.netapp.com/ServiceGW/entitlement https://eval.lic.netapp.com.s3.us-west-1.amazonaws.com https://cloud-support-netapp-com.s3.us-west-1.amazonaws.com	Communication with NetApp for system licensing and support registration.
https://client.infra.support.netapp.com.s3.us-west-1.amazonaws.com https://cloud-support-netapp-com-accelerated.s3.us-west-1.amazonaws.com https://trigger.asup.netapp.com.s3.us-west-1.amazonaws.com	Enables NetApp to collect information needed to troubleshoot support issues.
https://ipa-signer.cloudmanager.netapp.com	Enables Cloud Manager to generate licenses (for example, a FlexCache license for Cloud Volumes ONTAP)
*.blob.core.windows.net	Required for HA pairs when using a proxy.

Endpoints	Purpose
<p>Various third-party locations, for example:</p> <ul style="list-style-type: none"> • https://repo1.maven.org/maven2 • https://oss.sonatype.org/content/repositories • https://repo.typesafe.com <p>Third-party locations are subject to change.</p>	<p>During upgrades, Cloud Manager downloads the latest packages for third-party dependencies.</p>

While you should perform almost all tasks from the SaaS user interface, a local user interface is still available on the Connector. The machine running the web browser must have connections to the following endpoints:

Endpoints	Purpose
The Connector host	<p>You must enter the host's IP address from a web browser to load the Cloud Manager console.</p> <p>Depending on your connectivity to your cloud provider, you can use the private IP or a public IP assigned to the host:</p> <ul style="list-style-type: none"> • A private IP works if you have a VPN and direct connect access to your virtual network • A public IP works in any networking scenario <p>In any case, you should secure network access by ensuring that security group rules allow access from only authorized IPs or subnets.</p>
https://auth0.com https://cdn.auth0.com https://netapp-cloud-account.auth0.com https://services.cloud.netapp.com	<p>Your web browser connects to these endpoints for centralized user authentication through NetApp Cloud Central.</p>
https://widget.intercom.io	<p>For in-product chat that enables you to talk to NetApp cloud experts.</p>

Security group rules for Cloud Volumes ONTAP

Cloud Manager creates Azure security groups that include the inbound and outbound rules that Cloud Volumes ONTAP needs to operate successfully. You might want to refer to the ports for testing purposes or if you prefer your to use own security groups.

The security group for Cloud Volumes ONTAP requires both inbound and outbound rules.

Inbound rules for single node systems

The rules listed below allow traffic, unless the description notes that it blocks specific inbound traffic.

Priority and name	Port and protocol	Source and destination	Description
1000 inbound_ssh	22 TCP	Any to Any	SSH access to the IP address of the cluster management LIF or a node management LIF
1001 inbound_http	80 TCP	Any to Any	HTTP access to the System Manager web console using the IP address of the cluster management LIF
1002 inbound_111_tcp	111 TCP	Any to Any	Remote procedure call for NFS
1003 inbound_111_udp	111 UDP	Any to Any	Remote procedure call for NFS
1004 inbound_139	139 TCP	Any to Any	NetBIOS service session for CIFS
1005 inbound_161-162 _tcp	161-162 TCP	Any to Any	Simple network management protocol
1006 inbound_161-162 _udp	161-162 UDP	Any to Any	Simple network management protocol
1007 inbound_443	443 TCP	Any to Any	HTTPS access to the System Manager web console using the IP address of the cluster management LIF
1008 inbound_445	445 TCP	Any to Any	Microsoft SMB/CIFS over TCP with NetBIOS framing
1009 inbound_635_tcp	635 TCP	Any to Any	NFS mount
1010 inbound_635_udp	635 UDP	Any to Any	NFS mount
1011 inbound_749	749 TCP	Any to Any	Kerberos
1012 inbound_2049_tcp	2049 TCP	Any to Any	NFS server daemon
1013 inbound_2049_udp	2049 UDP	Any to Any	NFS server daemon
1014 inbound_3260	3260 TCP	Any to Any	iSCSI access through the iSCSI data LIF
1015 inbound_4045- 4046_tcp	4045-4046 TCP	Any to Any	NFS lock daemon and network status monitor
1016 inbound_4045- 4046_udp	4045-4046 UDP	Any to Any	NFS lock daemon and network status monitor

Priority and name	Port and protocol	Source and destination	Description
1017 inbound_10000	10000 TCP	Any to Any	Backup using NDMP
1018 inbound_11104-11105	11104- 11105 TCP	Any to Any	SnapMirror data transfer
3000 inbound_deny _all_tcp	Any port TCP	Any to Any	Block all other TCP inbound traffic
3001 inbound_deny _all_udp	Any port UDP	Any to Any	Block all other UDP inbound traffic
65000 AllowVnetInBound	Any port Any protocol	VirtualNetwork to VirtualNetwork	Inbound traffic from within the VNet
65001 AllowAzureLoad BalancerInBound	Any port Any protocol	AzureLoadBalanc er to Any	Data traffic from the Azure Standard Load Balancer
65500 DenyAllInBound	Any port Any protocol	Any to Any	Block all other inbound traffic

Inbound rules for HA systems

The rules listed below allow traffic, unless the description notes that it blocks specific inbound traffic.



HA systems have less inbound rules than single node systems because inbound data traffic goes through the Azure Standard Load Balancer. Because of this, traffic from the Load Balancer should be open, as shown in the "AllowAzureLoadBalancerInBound" rule.

Priority and name	Port and protocol	Source and destination	Description
100 inbound_443	443 Any protocol	Any to Any	HTTPS access to the System Manager web console using the IP address of the cluster management LIF
101 inbound_111_tcp	111 Any protocol	Any to Any	Remote procedure call for NFS
102 inbound_2049_tcp	2049 Any protocol	Any to Any	NFS server daemon
111 inbound_ssh	22 Any protocol	Any to Any	SSH access to the IP address of the cluster management LIF or a node management LIF

Priority and name	Port and protocol	Source and destination	Description
121 inbound_53	53 Any protocol	Any to Any	DNS and CIFS
65000 AllowVnetInBound	Any port Any protocol	VirtualNetwork to VirtualNetwork	Inbound traffic from within the VNet
65001 AllowAzureLoadBalancerInBound	Any port Any protocol	AzureLoadBalancer to Any	Data traffic from the Azure Standard Load Balancer
65500 DenyAllInBound	Any port Any protocol	Any to Any	Block all other inbound traffic

Outbound rules

The predefined security group for Cloud Volumes ONTAP opens all outbound traffic. If that is acceptable, follow the basic outbound rules. If you need more rigid rules, use the advanced outbound rules.

Basic outbound rules

The predefined security group for Cloud Volumes ONTAP includes the following outbound rules.

Port	Protocol	Purpose
All	All TCP	All outbound traffic
All	All UDP	All outbound traffic

Advanced outbound rules

If you need rigid rules for outbound traffic, you can use the following information to open only those ports that are required for outbound communication by Cloud Volumes ONTAP.



The source is the interface (IP address) on the Cloud Volumes ONTAP system.

Service	Port	Protocol	Source	Destination	Purpose
Active Directory	88	TCP	Node management LIF	Active Directory forest	Kerberos V authentication
	137	UDP	Node management LIF	Active Directory forest	NetBIOS name service
	138	UDP	Node management LIF	Active Directory forest	NetBIOS datagram service
	139	TCP	Node management LIF	Active Directory forest	NetBIOS service session
	389	TCP & UDP	Node management LIF	Active Directory forest	LDAP
	445	TCP	Node management LIF	Active Directory forest	Microsoft SMB/CIFS over TCP with NetBIOS framing
	464	TCP	Node management LIF	Active Directory forest	Kerberos V change & set password (SET_CHANGE)
	464	UDP	Node management LIF	Active Directory forest	Kerberos key administration
	749	TCP	Node management LIF	Active Directory forest	Kerberos V change & set Password (RPCSEC_GSS)
	88	TCP	Data LIF (NFS, CIFS, iSCSI)	Active Directory forest	Kerberos V authentication
	137	UDP	Data LIF (NFS, CIFS)	Active Directory forest	NetBIOS name service
	138	UDP	Data LIF (NFS, CIFS)	Active Directory forest	NetBIOS datagram service
	139	TCP	Data LIF (NFS, CIFS)	Active Directory forest	NetBIOS service session
	389	TCP & UDP	Data LIF (NFS, CIFS)	Active Directory forest	LDAP
DHCP	445	TCP	Data LIF (NFS, CIFS)	Active Directory forest	Microsoft SMB/CIFS over TCP with NetBIOS framing
	464	TCP	Data LIF (NFS, CIFS)	Active Directory forest	Kerberos V change & set password (SET_CHANGE)
	464	UDP	Data LIF (NFS, CIFS)	Active Directory forest	Kerberos key administration
	749	TCP	Data LIF (NFS, CIFS)	Active Directory forest	Kerberos V change & set password (RPCSEC_GSS)
	68	UDP	Node management LIF	DHCP	DHCP client for first-time setup

Service	Port	Protocol	Source	Destination	Purpose
DHCP	67	UDP	Node management LIF	DHCP	DHCP server
DNS	53	UDP	Node management LIF and data LIF (NFS, CIFS)	DNS	DNS
NDMP	18600–18699	TCP	Node management LIF	Destination servers	NDMP copy
SMTP	25	TCP	Node management LIF	Mail server	SMTP alerts, can be used for AutoSupport
SNMP	161	TCP	Node management LIF	Monitor server	Monitoring by SNMP traps
	161	UDP	Node management LIF	Monitor server	Monitoring by SNMP traps
	162	TCP	Node management LIF	Monitor server	Monitoring by SNMP traps
	162	UDP	Node management LIF	Monitor server	Monitoring by SNMP traps
SnapMirror	11104	TCP	Intercluster LIF	ONTAP intercluster LIFs	Management of intercluster communication sessions for SnapMirror
	11105	TCP	Intercluster LIF	ONTAP intercluster LIFs	SnapMirror data transfer
Syslog	514	UDP	Node management LIF	Syslog server	Syslog forward messages

Security group rules for the Connector

The security group for the Connector requires both inbound and outbound rules.

Inbound rules

Port	Protocol	Purpose
22	SSH	Provides SSH access to the Connector host
80	HTTP	Provides HTTP access from client web browsers to the local user interface
443	HTTPS	Provides HTTPS access from client web browsers to the local user interface

Outbound rules

The predefined security group for the Connector opens all outbound traffic. If that is acceptable, follow the basic outbound rules. If you need more rigid rules, use the advanced outbound rules.

Basic outbound rules

The predefined security group for the Connector includes the following outbound rules.

Port	Protocol	Purpose
All	All TCP	All outbound traffic
All	All UDP	All outbound traffic

Advanced outbound rules

If you need rigid rules for outbound traffic, you can use the following information to open only those ports that are required for outbound communication by the Connector.



The source IP address is the Connector host.

Service	Port	Protocol	Destination	Purpose
Active Directory	88	TCP	Active Directory forest	Kerberos V authentication
	13 9	TCP	Active Directory forest	NetBIOS service session
	38 9	TCP	Active Directory forest	LDAP
	44 5	TCP	Active Directory forest	Microsoft SMB/CIFS over TCP with NetBIOS framing
	46 4	TCP	Active Directory forest	Kerberos V change & set password (SET_CHANGE)
	74 9	TCP	Active Directory forest	Active Directory Kerberos V change & set password (RPCSEC_GSS)
	13 7	UDP	Active Directory forest	NetBIOS name service
	13 8	UDP	Active Directory forest	NetBIOS datagram service
	46 4	UDP	Active Directory forest	Kerberos key administration
API calls and AutoSupport	44 3	HTTP	Outbound internet and ONTAP cluster management LIF	API calls to AWS and ONTAP, and sending AutoSupport messages to NetApp
DNS	53	UDP	DNS	Used for DNS resolve by Cloud Manager

Set up Cloud Volumes ONTAP to use a customer-managed key in Azure

Data is automatically encrypted on Cloud Volumes ONTAP in Azure using [Azure Storage Service Encryption](#) with a Microsoft-managed key. But you can use your own encryption key instead by following the steps on this page.

Data encryption overview

Cloud Volumes ONTAP data is automatically encrypted in Azure using [Azure Storage Service Encryption](#). The default implementation uses a Microsoft-managed key. No setup is required.

If you want to use a customer-managed key with Cloud Volumes ONTAP, then you need to complete the following steps:

1. From Azure, create a key vault and then generate a key in that vault
2. From Cloud Manager, use the API to create a Cloud Volumes ONTAP working environment that uses the key

Key rotation

If you create a new version of your key, Cloud Volumes ONTAP automatically uses the latest key version.

How data is encrypted

After you create a Cloud Volumes ONTAP working environment that is configured to use a customer-managed key, Cloud Volumes ONTAP data is encrypted as follows.

HA pairs

- All Azure storage accounts for Cloud Volumes ONTAP are encrypted using a customer-managed key.
- Any new storage accounts (for example, when you add disks or aggregates) also use the same key.

Single node

- All Azure storage accounts for Cloud Volumes ONTAP are encrypted using a customer-managed key.
- For root, boot, and data disks, Cloud Manager uses a [disk encryption set](#), which enables management of encryption keys with managed disks.
- Any new data disks also use the same disk encryption set.
- NVRAM and the core disk are encrypted using a Microsoft-managed key, instead of the customer-managed key.

Create a key vault and generate a key

The key vault must reside in the same Azure subscription and region in which you plan to create the Cloud Volumes ONTAP system.

Steps

1. [Create a key vault in your Azure subscription](#).

Note the following requirements for the key vault:

- The key vault must reside in the same region as the Cloud Volumes ONTAP system.
- The following options should be enabled:
 - **Soft-delete** (this option is enabled by default, but must *not* be disabled)
 - **Purge protection**
 - **Azure Disk Encryption for volume encryption** (for single node Cloud Volumes ONTAP systems only)

2. [Generate a key in the key vault](#).

Note the following requirements for the key:

- The key type must be **RSA**.
- The recommended RSA key size is **2048**, but other sizes are supported.

Create a working environment that uses the encryption key

After you create the key vault and generate an encryption key, you can create a new Cloud Volumes ONTAP system that is configured to use the key. These steps are supported by using the Cloud Manager API.

Required permissions

If you want to use a customer-managed key with a single node Cloud Volumes ONTAP system, ensure that the Cloud Manager Connector has the following permissions:

```
"Microsoft.Compute/diskEncryptionSets/read"  
"Microsoft.Compute/diskEncryptionSets/write",  
"Microsoft.Compute/diskEncryptionSets/delete"  
"Microsoft.KeyVault/vaults/deploy/action",  
"Microsoft.KeyVault/vaults/read",  
"Microsoft.KeyVault/vaults/accessPolicies/write"
```

You can find the latest list of permissions on the [Cloud Manager policies page](#).

These permissions aren't required for HA pairs.

Steps

1. Obtain the list of key vaults in your Azure subscription by using the following Cloud Manager API call.

For an HA pair: GET /azure/ha/metadata/vaults

For single node: GET /azure/vsa/metadata/vaults

Make note of the **name** and **resourceGroup**. You'll need to specify those values in the next step.

[Learn more about this API call](#).

2. Obtain the list of keys within the vault by using the following Cloud Manager API call.

For an HA pair: GET /azure/ha/metadata/keys-vault

For single node: GET /azure/vsa/metadata/keys-vault

Make note of the **keyName**. You'll need to specify that value (along with the vault name) in the next step.

[Learn more about this API call](#).

3. Create a Cloud Volumes ONTAP system by using the following Cloud Manager API call.

- a. For an HA pair:

```
POST /azure/ha/working-environments
```

The request body must include the following fields:

```
"azureEncryptionParameters": {  
    "key": "keyName",  
    "vaultName": "vaultName"  
}
```

[Learn more about this API call.](#)

b. For a single node system:

```
POST /azure/vsa/working-environments
```

The request body must include the following fields:

```
"azureEncryptionParameters": {  
    "key": "keyName",  
    "vaultName": "vaultName"  
}
```

[Learn more about this API call.](#)

Result

You have a new Cloud Volumes ONTAP system that is configured to use your customer-managed key for data encryption.

Launching Cloud Volumes ONTAP in Azure

You can launch a single node system or an HA pair in Azure by creating a Cloud Volumes ONTAP working environment in Cloud Manager.

What you'll need

You need the following to create a working environment.

- A Connector that's up and running.
 - You should have a [Connector that is associated with your workspace](#).
 - [You should be prepared to leave the Connector running at all times](#).
- An understanding of the configuration that you want to use.

You should have chose a configuration and obtained Azure networking information from your administrator. For details, see [Planning your Cloud Volumes ONTAP configuration](#).

- An understanding of what's required in the Add Working Environment wizard for the licensing method that you want to use.

Licensing option	Requirement	How to meet the requirement
PAYGO free trial	A Marketplace subscription is required.	You'll have the option to subscribe to your cloud provider's marketplace from the Details & Credentials page.
Freemium	A Marketplace subscription or NetApp Support Site (NSS) account is required.	You'll have the option to subscribe to your cloud provider's marketplace from the Details & Credentials page You can enter your NSS account on the Charging Methods and NSS Account page.
Capacity-based BYOL	A Marketplace subscription or NetApp Support Site (NSS) account is required. A Marketplace subscription is recommended for capacity-based charging in the event that your account doesn't have a valid capacity-based license, or in the event that your provisioned capacity exceeds the licensed capacity.	You'll have the option to subscribe to your cloud provider's marketplace from the Details & Credentials page You can enter your NSS account on the Charging Methods and NSS Account page.
Node-based BYOL	The 20-digit serial number (license key) is required.	You'll enter the serial number on the Charging Methods and NSS Account page.

About this task

When Cloud Manager creates a Cloud Volumes ONTAP system in Azure, it creates several Azure objects, such as a resource group, network interfaces, and storage accounts. You can review a summary of the resources at the end of the wizard.

Potential for Data Loss

The best practice is to use a new, dedicated resource group for each Cloud Volumes ONTAP system.

 Deploying Cloud Volumes ONTAP in an existing, shared resource group is not recommended due to the risk of data loss. While Cloud Manager can remove Cloud Volumes ONTAP resources from a shared resource group in case of deployment failure or deletion, an Azure user might accidentally delete Cloud Volumes ONTAP resources from a shared resource group.

Steps

1. On the Canvas page, click **Add Working Environment** and follow the prompts.
2. **Choose a Location:** Select **Microsoft Azure** and **Cloud Volumes ONTAP Single Node or Cloud Volumes ONTAP High Availability**.
3. If you're prompted, [create a Connector](#).
4. **Details and Credentials:** Optionally change the Azure credentials and subscription, specify a cluster name, add tags if needed, and then specify credentials.

The following table describes fields for which you might need guidance:

Field	Description
Working Environment Name	Cloud Manager uses the working environment name to name both the Cloud Volumes ONTAP system and the Azure virtual machine. It also uses the name as the prefix for the predefined security group, if you select that option.
Resource Group Tags	<p>Tags are metadata for your Azure resources. When you enter tags in this field, Cloud Manager adds them to the resource group associated with the Cloud Volumes ONTAP system.</p> <p>You can add up to four tags from the user interface when creating a working environment, and then you can add more after its created. Note that the API does not limit you to four tags when creating a working environment.</p> <p>For information about tags, refer to Microsoft Azure Documentation: Using tags to organize your Azure resources.</p>
User name and password	These are the credentials for the Cloud Volumes ONTAP cluster admin account. You can use these credentials to connect to Cloud Volumes ONTAP through System Manager or its CLI.
Edit Credentials	You can choose different Azure credentials and a different Azure subscription to use with this Cloud Volumes ONTAP system. You need to associate an Azure Marketplace subscription with the selected Azure subscription in order to deploy a pay-as-you-go Cloud Volumes ONTAP system. Learn how to add credentials .

The following video shows how to associate a Marketplace subscription to an Azure subscription:

- ▶ https://docs.netapp.com/us-en/occm//media/video_subscribing_azure.mp4 (video)

5. **Services:** Keep the services enabled or disable the individual services that you don't want to use with Cloud Volumes ONTAP.
 - [Learn more about Cloud Data Sense](#).
 - [Learn more about Cloud Backup](#).
 - [Learn more about the Monitoring service](#).
6. **Location & Connectivity:** Select a location, a resource group, a security group, and then select the checkbox to confirm network connectivity between the Connector and the target location.

The following table describes fields for which you might need guidance:

Field	Description
Location	For single node systems, you can choose the Availability Zone in which you'd like to deploy Cloud Volumes ONTAP. If you don't select an AZ, Cloud Manager will select one for you.

Field	Description
Resource Group	<p>Create a new resource group for Cloud Volumes ONTAP or use an existing resource group. The best practice is to use a new, dedicated resource group for Cloud Volumes ONTAP. While it is possible to deploy Cloud Volumes ONTAP in an existing, shared resource group, it's not recommended due to the risk of data loss. See the warning above for more details.</p> <p>You must use a dedicated resource group for each Cloud Volumes ONTAP HA pair that you deploy in Azure. Only one HA pair is supported in a resource group. Cloud Manager experiences connection issues if you try to deploy a second Cloud Volumes ONTAP HA pair in an Azure resource group.</p> <p> If the Azure account that you're using has the required permissions, Cloud Manager removes Cloud Volumes ONTAP resources from a resource group, in case of deployment failure or deletion.</p>
Security group	<p>If you choose an existing security group, then it must meet Cloud Volumes ONTAP requirements. View the default security group.</p>

7. **Charging Methods and NSS Account:** Specify which charging option would you like to use with this system, and then specify a NetApp Support Site account.
 - [Learn about these charging methods](#).
 - [Learn what's required in the wizard for the licensing method that you want to use](#).
8. **Preconfigured Packages:** Select one of the packages to quickly deploy a Cloud Volumes ONTAP system, or click **Create my own configuration**.

If you choose one of the packages, you only need to specify a volume and then review and approve the configuration.
9. **Licensing:** Change the Cloud Volumes ONTAP version as needed, select a license, and select a virtual machine type.

If your needs change after you launch the system, you can modify the license or virtual machine type later.



If a newer Release Candidate, General Availability, or patch release is available for the selected version, then Cloud Manager updates the system to that version when creating the working environment. For example, the update occurs if you select Cloud Volumes ONTAP 9.6 RC1 and 9.6 GA is available. The update does not occur from one release to another—for example, from 9.6 to 9.7.

10. **Subscribe from the Azure Marketplace:** Follow the steps if Cloud Manager could not enable programmatic deployments of Cloud Volumes ONTAP.
11. **Underlying Storage Resources:** Choose settings for the initial aggregate: a disk type, a size for each disk, and whether data tiering to Blob storage should be enabled.

Note the following:

- The disk type is for the initial volume. You can choose a different disk type for subsequent volumes.
- The disk size is for all disks in the initial aggregate and for any additional aggregates that Cloud Manager creates when you use the simple provisioning option. You can create aggregates that use a different disk size by using the advanced allocation option.

For help choosing a disk type and size, see [Sizing your system in Azure](#).

- You can choose a specific volume tiering policy when you create or edit a volume.
- If you disable data tiering, you can enable it on subsequent aggregates.

[Learn more about data tiering](#).

12. Write Speed & WORM (single node systems only): Choose **Normal** or **High** write speed, and activate write once, read many (WORM) storage, if desired.

[Learn more about write speed](#).

WORM can't be enabled if Cloud Backup was enabled or if data tiering was enabled.

[Learn more about WORM storage](#).

13. Secure Communication to Storage & WORM (HA only): Choose whether to enable an HTTPS connection to Azure storage accounts, and activate write once, read many (WORM) storage, if desired.

The HTTPS connection is from a Cloud Volumes ONTAP 9.7 HA pair to Azure storage accounts. Note that enabling this option can impact write performance. You can't change the setting after you create the working environment.

[Learn more about WORM storage](#).

14. Create Volume: Enter details for the new volume or click **Skip**.

Some of the fields in this page are self-explanatory. The following table describes fields for which you might need guidance:

Field	Description
Size	The maximum size that you can enter largely depends on whether you enable thin provisioning, which enables you to create a volume that is bigger than the physical storage currently available to it.
Access control (for NFS only)	An export policy defines the clients in the subnet that can access the volume. By default, Cloud Manager enters a value that provides access to all instances in the subnet.
Permissions and Users / Groups (for CIFS only)	These fields enable you to control the level of access to a share for users and groups (also called access control lists or ACLs). You can specify local or domain Windows users or groups, or UNIX users or groups. If you specify a domain Windows user name, you must include the user's domain using the format domain\username.

Field	Description
Snapshot Policy	A Snapshot copy policy specifies the frequency and number of automatically created NetApp Snapshot copies. A NetApp Snapshot copy is a point-in-time file system image that has no performance impact and requires minimal storage. You can choose the default policy or none. You might choose none for transient data: for example, tempdb for Microsoft SQL Server.
Advanced options (for NFS only)	Select an NFS version for the volume: either NFSv3 or NFSv4.
Initiator group and IQN (for iSCSI only)	<p>iSCSI storage targets are called LUNs (logical units) and are presented to hosts as standard block devices.</p> <p>Initiator groups are tables of iSCSI host node names and control which initiators have access to which LUNs.</p> <p>iSCSI targets connect to the network through standard Ethernet network adapters (NICs), TCP offload engine (TOE) cards with software initiators, converged network adapters (CNAs) or dedicated host bus adapters (HBAs) and are identified by iSCSI qualified names (IQNs).</p> <p>When you create an iSCSI volume, Cloud Manager automatically creates a LUN for you. We've made it simple by creating just one LUN per volume, so there's no management involved. After you create the volume, use the IQN to connect to the LUN from your hosts.</p>

The following image shows the Volume page filled out for the CIFS protocol:

The screenshot shows the 'Volume Details, Protection & Protocol' page. The 'Protocol' tab is selected, showing 'CIFS' highlighted. In the 'Details & Protection' section, 'Volume Name' is 'vol', 'Size (GB)' is '250', 'Snapshot Policy' is 'default', and a note says 'Default Policy'. In the 'Protocol' section, 'Share name' is 'vol_share', 'Permissions' is 'Full Control', and 'Users / Groups' is 'engineering'. A note at the bottom says 'Valid users and groups separated by a semicolon'.

15. CIFS Setup: If you chose the CIFS protocol, set up a CIFS server.

Field	Description
DNS Primary and Secondary IP Address	<p>The IP addresses of the DNS servers that provide name resolution for the CIFS server.</p> <p>The listed DNS servers must contain the service location records (SRV) needed to locate the Active Directory LDAP servers and domain controllers for the domain that the CIFS server will join.</p>

Field	Description
Active Directory Domain to join	The FQDN of the Active Directory (AD) domain that you want the CIFS server to join.
Credentials authorized to join the domain	The name and password of a Windows account with sufficient privileges to add computers to the specified Organizational Unit (OU) within the AD domain.
CIFS server NetBIOS name	A CIFS server name that is unique in the AD domain.
Organizational Unit	<p>The organizational unit within the AD domain to associate with the CIFS server. The default is CN=Computers.</p> <p>To configure Azure AD Domain Services as the AD server for Cloud Volumes ONTAP, you should enter OU=AADDC Computers or OU=AADDC Users in this field.</p> <p>Azure Documentation: Create an Organizational Unit (OU) in an Azure AD Domain Services managed domain</p>
DNS Domain	The DNS domain for the Cloud Volumes ONTAP storage virtual machine (SVM). In most cases, the domain is the same as the AD domain.
NTP Server	Select Use Active Directory Domain to configure an NTP server using the Active Directory DNS. If you need to configure an NTP server using a different address, then you should use the API. See the Cloud Manager automation docs for details.

16. **Usage Profile, Disk Type, and Tiering Policy:** Choose whether you want to enable storage efficiency features and change the volume tiering policy, if needed.

For more information, see [Understanding volume usage profiles](#) and [Data tiering overview](#).

17. **Review & Approve:** Review and confirm your selections.

- Review details about the configuration.
- Click **More information** to review details about support and the Azure resources that Cloud Manager will purchase.
- Select the **I understand...** check boxes.
- Click **Go**.

Result

Cloud Manager deploys the Cloud Volumes ONTAP system. You can track the progress in the timeline.

If you experience any issues deploying the Cloud Volumes ONTAP system, review the failure message. You can also select the working environment and click **Re-create environment**.

For additional help, go to [NetApp Cloud Volumes ONTAP Support](#).

After you finish

- If you provisioned a CIFS share, give users or groups permissions to the files and folders and verify that those users can access the share and create a file.
- If you want to apply quotas to volumes, use System Manager or the CLI.

Quotas enable you to restrict or track the disk space and number of files used by a user, group, or qtree.

Get started in GCP

Getting started with Cloud Volumes ONTAP for Google Cloud

Get started with Cloud Volumes ONTAP for GCP in a few steps.



Create a Connector

If you don't have a [Connector](#) yet, an Account Admin needs to create one. [Learn how to create a Connector in GCP.](#)

When you create your first Cloud Volumes ONTAP working environment, Cloud Manager prompts you to deploy a Connector if you don't have one yet.



Plan your configuration

Cloud Manager offers preconfigured packages that match your workload requirements, or you can create your own configuration. If you choose your own configuration, you should understand the options available to you.

[Learn more about planning your configuration.](#)



Set up your networking

- a. Ensure that your VPC and subnets will support connectivity between the Connector and Cloud Volumes ONTAP.
- b. If you plan to enable data tiering, [configure the Cloud Volumes ONTAP subnet for Private Google Access](#).
- c. If you're deploying an HA pair, ensure that you have four VPCs, each with their own subnet.
- d. If you're using a shared VPC, provide the *Compute Network User* role to the Connector service account.
- e. Enable outbound internet access from the target VPC so the Connector and Cloud Volumes ONTAP can contact several endpoints.

This step is important because the Connector can't manage Cloud Volumes ONTAP without outbound internet access. If you need to limit outbound connectivity, refer to the list of endpoints for [the Connector and Cloud Volumes ONTAP](#).

[Learn more about networking requirements.](#)



Set up a service account

Cloud Volumes ONTAP requires a Google Cloud service account for two purposes. The first is when you enable [data tiering](#) to tier cold data to low-cost object storage in Google Cloud. The second is when you enable the [Cloud Backup Service](#) to back up volumes to low-cost object storage.

You can set up one service account and use it for both purposes. The service account must have the **Storage Admin** role.

[Read step-by-step instructions.](#)

5

Enable Google Cloud APIs

Enable the following Google Cloud APIs in your project. These APIs are required to deploy the Connector and Cloud Volumes ONTAP.

- Cloud Deployment Manager V2 API
- Cloud Logging API
- Cloud Resource Manager API
- Compute Engine API
- Identity and Access Management (IAM) API

6

Launch Cloud Volumes ONTAP using Cloud Manager

Click **Add Working Environment**, select the type of system that you would like to deploy, and complete the steps in the wizard. [Read step-by-step instructions](#).

Related links

- [Evaluating](#)
- [Creating a Connector from Cloud Manager](#)
- [Installing the Connector software on a Linux host](#)
- [What Cloud Manager does with GCP permissions](#)

Planning your Cloud Volumes ONTAP configuration in Google Cloud

When you deploy Cloud Volumes ONTAP in Google Cloud, you can choose a preconfigured system that matches your workload requirements, or you can create your own configuration. If you choose your own configuration, you should understand the options available to you.

Choosing a license type

Cloud Volumes ONTAP is available in two pricing options: pay-as-you-go and Bring Your Own License (BYOL). For pay-as-you-go, you can choose from three licenses: Explore, Standard, or Premium. Each license provides different capacity and compute options.

[Supported configurations for Cloud Volumes ONTAP in GCP](#)

Supported machine types

Cloud Volumes ONTAP supports several machine types, depending on the license type that you choose.

[Supported configurations for Cloud Volumes ONTAP in GCP](#)

Understanding storage limits

The raw capacity limit for a Cloud Volumes ONTAP system is tied to the license. Additional limits impact the size of aggregates and volumes. You should be aware of these limits as you plan your configuration.

[Storage limits for Cloud Volumes ONTAP in GCP](#)

Sizing your system in GCP

Sizing your Cloud Volumes ONTAP system can help you meet requirements for performance and capacity. You should be aware of a few key points when choosing a machine type, disk type, and disk size:

Machine type

Look at the supported machine types in the [Cloud Volumes ONTAP Release Notes](#) and then review details from Google about each supported machine type. Match your workload requirements to the number of vCPUs and memory for the machine type. Note that each CPU core increases networking performance.

Refer to the following for more details:

- [Google Cloud documentation: N1 standard machine types](#)
- [Google Cloud documentation: Performance](#)

GCP disk type

When you create volumes for Cloud Volumes ONTAP, you need to choose the underlying cloud storage that Cloud Volumes ONTAP uses for a disk. The disk type can be any of the following:

- *Zonal SSD persistent disks*: SSD persistent disks are best for workloads that require high rates of random IOPS.
- *Zonal Balanced persistent disks*: These SSDs balance performance and cost by providing lower IOPS per GB.
- *Zonal Standard persistent disks* : Standard persistent disks are economical and can handle sequential read/write operations.

For more details, see [Google Cloud documentation: Zonal Persistent disks \(Standard and SSD\)](#).

GCP disk size

You need to choose an initial disk size when you deploy a Cloud Volumes ONTAP system. After that you can let Cloud Manager manage a system's capacity for you, but if you want to build aggregates yourself, be aware of the following:

- All disks in an aggregate must be the same size.
- Determine the space that you need, while taking performance into consideration.
- The performance of persistent disks scales automatically with disk size and the number of vCPUs available to the system.

Refer to the following for more details:

- [Google Cloud documentation: Zonal Persistent disks \(Standard and SSD\)](#)
- [Google Cloud documentation: Optimizing Persistent Disk and Local SSD Performance](#)

GCP network information worksheet

When you deploy Cloud Volumes ONTAP in GCP, you need to specify details about your virtual network. You can use a worksheet to collect the information from your administrator.

Network information for a single-node system

GCP information	Your value
Region	
Zone	
VPC network	
Subnet	
Firewall policy (if using your own)	

Network information for an HA pair in multiple zones

GCP information	Your value
Region	
Zone for Node 1	
Zone for Node 2	
Zone for the mediator	
VPC-0 and subnet	
VPC-1 and subnet	
VPC-2 and subnet	
VPC-3 and subnet	
Firewall policy (if using your own)	

Network information for an HA pair in a single zone

GCP information	Your value
Region	
Zone	
VPC-0 and subnet	
VPC-1 and subnet	
VPC-2 and subnet	
VPC-3 and subnet	
Firewall policy (if using your own)	

Choosing a write speed

Cloud Manager enables you to choose a write speed setting for Cloud Volumes ONTAP, except for high availability (HA) pairs in Google Cloud. Before you choose a write speed, you should understand the differences between the normal and high settings and risks and recommendations when using high write speed. [Learn more about write speed](#).

Choosing a volume usage profile

ONTAP includes several storage efficiency features that can reduce the total amount of storage that you need. When you create a volume in Cloud Manager, you can choose a profile that enables these features or a profile that disables them. You should learn more about these features to help you decide which profile to use.

NetApp storage efficiency features provide the following benefits:

Thin provisioning

Presents more logical storage to hosts or users than you actually have in your physical storage pool. Instead of preallocating storage space, storage space is allocated dynamically to each volume as data is written.

Deduplication

Improves efficiency by locating identical blocks of data and replacing them with references to a single shared block. This technique reduces storage capacity requirements by eliminating redundant blocks of data that reside in the same volume.

Compression

Reduces the physical capacity required to store data by compressing data within a volume on primary, secondary, and archive storage.

Networking requirements for Cloud Volumes ONTAP in GCP

Set up your Google Cloud Platform networking so Cloud Volumes ONTAP systems can operate properly. This includes networking for the Connector and Cloud Volumes ONTAP.

If you want to deploy an HA pair, you should [learn how HA pairs work in GCP](#).

Requirements for Cloud Volumes ONTAP

The following requirements must be met in GCP.

Internal load balancers

Cloud Manager automatically creates four Google Cloud internal load balancers (TCP/UDP) that manage incoming traffic to the Cloud Volumes ONTAP HA pair. No setup is required from your end. We've listed this as a requirement simply to inform you of the network traffic and to mitigate any security concerns.

One load balancer is for cluster management, one is for storage VM (SVM) management, one is for NAS traffic to node 1, and the last is for NAS traffic to node 2.

The setup for each load balancer is as follows:

- One shared private IP address
- One global health check

By default, the ports used by the health check are 63001, 63002, and 63003.

- One regional TCP backend service
- One regional UDP backend service
- One TCP forwarding rule
- One UDP forwarding rule
- Global access is disabled

Even though global access is disabled by default, enabling it post deployment is supported. We disabled it because cross region traffic will have significantly higher latencies. We wanted to ensure that you didn't have a negative experience due to accidental cross region mounts. Enabling this option is specific to your business needs.

One or multiple zones for HA pairs

You can ensure the high availability of your data by deploying an HA configuration across multiple or in a single zone. Cloud Manager will prompt you to choose multiple zones or a single zone when you create the HA pair.

- Multiple zones (recommended)

Deploying an HA configuration across three zones ensures continuous data availability if a failure occurs within a zone. Note that write performance is slightly lower compared to using a single zone, but it's minimal.

- Single zone

When deployed in a single zone, a Cloud Volumes ONTAP HA configuration uses a spread placement policy. This policy ensures that an HA configuration is protected from a single point of failure within the zone, without having to use separate zones to achieve fault isolation.

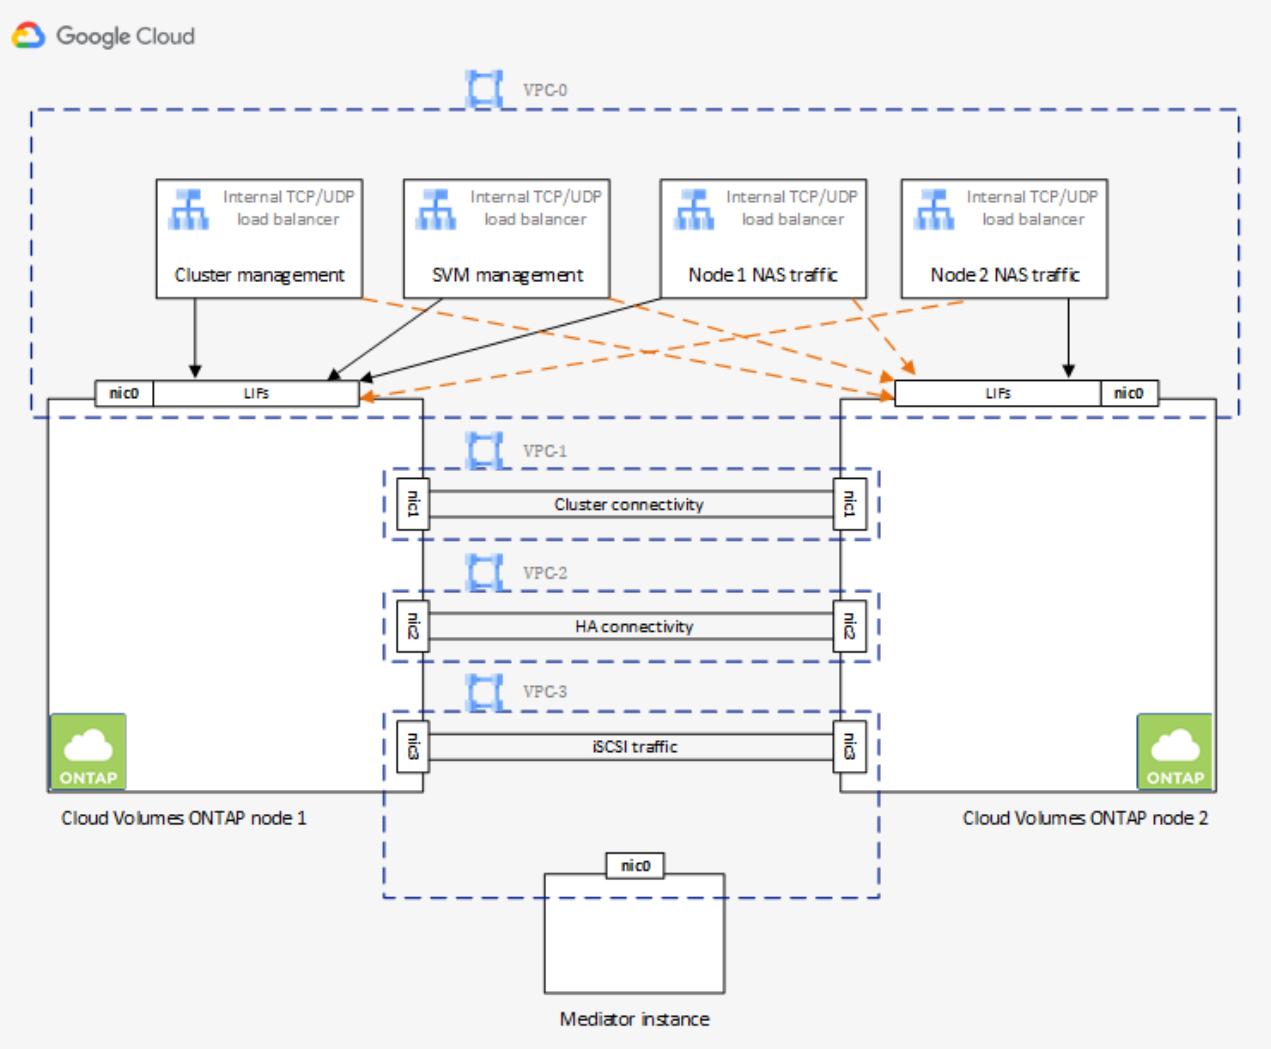
This deployment model does lower your costs because there are no data egress charges between zones.

Four Virtual Private Clouds for HA pairs

Four Virtual Private Clouds (VPCs) are required for an HA configuration. Four VPCs are required because GCP requires that each network interface resides in a separate VPC network.

Cloud Manager will prompt you to choose four VPCs when you create the HA pair:

- VPC-0 for inbound connections to the data and nodes
- VPC-1, VPC-2, and VPC-3 for internal communication between the nodes and the HA mediator



Subnets for HA pairs

A private subnet is required for each VPC.

If you place the Connector in VPC-0, then you will need to enable Private Google Access on the subnet to access the APIs and to enable data tiering.

The subnets in these VPCs must have distinct CIDR ranges. They can't have overlapping CIDR ranges.

One Virtual Private Cloud for single node systems

One VPC is required for a single node system.

Shared VPCs

Cloud Volumes ONTAP and the Connector are supported in a Google Cloud shared VPC and also in standalone VPCs.

For a single node system, the VPC can be either a shared VPC or a standalone VPC.

For an HA pair, four VPCs are required. Each of those VPCs can be either shared or standalone. For example, VPC-0 could be a shared VPC, while VPC-1, VPC-2, and VPC-3 could be standalone VPCs.

A shared VPC enables you to configure and centrally manage virtual networks across multiple projects. You can set up shared VPC networks in the *host project* and deploy the Connector and Cloud Volumes ONTAP

virtual machine instances in a *service project*. [Google Cloud documentation: Shared VPC overview](#).

The only requirement when using a shared VPC is to provide the [Compute Network User role](#) to the Connector service account. Cloud Manager needs these permissions to query the firewalls, VPC, and subnets in the host project.

Outbound internet access for Cloud Volumes ONTAP

Cloud Volumes ONTAP requires outbound internet access to send messages to NetApp AutoSupport, which proactively monitors the health of your storage.

Routing and firewall policies must allow HTTP/HTTPS traffic to the following endpoints so Cloud Volumes ONTAP can send AutoSupport messages:

- <https://support.netapp.com/aods/asupmessage>
- <https://support.netapp.com/asupprod/post/1.0/postAsup>

[Learn how to configure AutoSupport](#).



If you're using an HA pair, the HA mediator doesn't require outbound internet access.

Private IP addresses

Cloud Manager allocates the following number of private IP addresses to Cloud Volumes ONTAP in GCP:

- **Single node:** 3 or 4 private IP addresses

You can skip creation of the storage VM (SVM) management LIF if you deploy Cloud Volumes ONTAP using the API and specify the following flag:

```
skipSvmManagementLif: true
```

A LIF is an IP address associated with a physical port. A storage VM (SVM) management LIF is required for management tools like SnapCenter.

- **HA pair:** 15 or 16 private IP addresses
 - 7 or 8 private IP addresses for VPC-0

You can skip creation of the storage VM (SVM) management LIF if you deploy Cloud Volumes ONTAP using the API and specify the following flag:

```
skipSvmManagementLif: true
```

- Two private IP addresses for VPC-1
- Two private IP addresses for VPC-2
- Three private IP addresses for VPC-3

Firewall rules

You don't need to create firewall rules because Cloud Manager does that for you. If you need to use your own, refer to the firewall rules listed below.

Note that two sets of firewall rules are required for an HA configuration:

- One set of rules for HA components in VPC-0. These rules enable data access to Cloud Volumes

ONTAP. [Learn more](#).

- Another set of rules for HA components in VPC-1, VPC-2, and VPC-3. These rules are open for inbound & outbound communication between the HA components. [Learn more](#).

Connection from Cloud Volumes ONTAP to Google Cloud Storage for data tiering

If you want to tier cold data to a Google Cloud Storage bucket, the subnet in which Cloud Volumes ONTAP resides must be configured for Private Google Access (if you're using an HA pair, this is the subnet in VPC-0). For instructions, refer to [Google Cloud documentation: Configuring Private Google Access](#).

For additional steps required to set up data tiering in Cloud Manager, see [Tiering cold data to low-cost object storage](#).

Connections to ONTAP systems in other networks

To replicate data between a Cloud Volumes ONTAP system in GCP and ONTAP systems in other networks, you must have a VPN connection between the VPC and the other network—for example, your corporate network.

For instructions, refer to [Google Cloud documentation: Cloud VPN overview](#).

Requirements for the Connector

Set up your networking so that the Connector can manage resources and processes within your public cloud environment. The most important step is ensuring outbound internet access to various endpoints.



If your network uses a proxy server for all communication to the internet, you can specify the proxy server from the Settings page. Refer to [Configuring the Connector to use a proxy server](#).

Connection to target networks

A Connector requires a network connection to the VPCs in which you want to deploy Cloud Volumes ONTAP. If you're deploying an HA pair, then the Connector needs a connection to VPC-0 only.

Outbound internet access

The Connector requires outbound internet access to manage resources and processes within your public cloud environment. A Connector contacts the following endpoints when managing resources in GCP:

Endpoints	Purpose
https://www.googleapis.com	Enables the Connector to contact Google APIs for deploying and managing Cloud Volumes ONTAP in GCP.
https://api.services.cloud.netapp.com:443	API requests to NetApp Cloud Central.
https://cloud.support.netapp.com.s3.us-west-1.amazonaws.com	Provides access to software images, manifests, and templates.
https://cognito-idp.us-east-1.amazonaws.com https://cognito-identity.us-east-1.amazonaws.com https://sts.amazonaws.com https://cloud-support-netapp-com-accelerated.s3.amazonaws.com	Enables the Connector to access and download manifests, templates, and Cloud Volumes ONTAP upgrade images.

Endpoints	Purpose
https://cloudmanagerinfraprod.azurecr.io *.blob.core.windows.net	Access to software images of container components for an infrastructure that's running Docker and provides a solution for service integrations with Cloud Manager.
https://kinesis.us-east-1.amazonaws.com	Enables NetApp to stream data from audit records.
https://cloudmanager.cloud.netapp.com	Communication with the Cloud Manager service, which includes Cloud Central accounts.
https://netapp-cloud-account.auth0.com	Communication with NetApp Cloud Central for centralized user authentication.
support.netapp.com:443 https://mysupport.netapp.com	Communication with NetApp AutoSupport. Note that the Connector communicates with support.netapp.com:443, which redirects to https://mysupport.netapp.com.
https://support.netapp.com/svcgw https://support.netapp.com/ServiceGW/entitlement https://eval.lic.netapp.com.s3.us-west-1.amazonaws.com https://cloud-support-netapp-com.s3.us-west-1.amazonaws.com	Communication with NetApp for system licensing and support registration.
https://client.infra.support.netapp.com.s3.us-west-1.amazonaws.com https://cloud-support-netapp-com-accelerated.s3.us-west-1.amazonaws.com https://trigger.asup.netapp.com.s3.us-west-1.amazonaws.com	Enables NetApp to collect information needed to troubleshoot support issues.
https://ipa-signer.cloudmanager.netapp.com	Enables Cloud Manager to generate licenses (for example, a FlexCache license for Cloud Volumes ONTAP)
Various third-party locations, for example: <ul style="list-style-type: none"> • https://repo1.maven.org/maven2 • https://oss.sonatype.org/content/repositories • https://repo.typesafe.com Third-party locations are subject to change.	During upgrades, Cloud Manager downloads the latest packages for third-party dependencies.

While you should perform almost all tasks from the SaaS user interface, a local user interface is still available on the Connector. The machine running the web browser must have connections to the following endpoints:

Endpoints	Purpose
The Connector host	<p>You must enter the host's IP address from a web browser to load the Cloud Manager console.</p> <p>Depending on your connectivity to your cloud provider, you can use the private IP or a public IP assigned to the host:</p> <ul style="list-style-type: none"> • A private IP works if you have a VPN and direct connect access to your virtual network • A public IP works in any networking scenario <p>In any case, you should secure network access by ensuring that security group rules allow access from only authorized IPs or subnets.</p>
https://auth0.com https://cdn.auth0.com https://netapp-cloud-account.auth0.com https://services.cloud.netapp.com	Your web browser connects to these endpoints for centralized user authentication through NetApp Cloud Central.
https://widget.intercom.io	For in-product chat that enables you to talk to NetApp cloud experts.

Firewall rules for Cloud Volumes ONTAP

Cloud Manager creates GCP firewall rules that include the inbound and outbound rules that Cloud Volumes ONTAP needs to operate successfully. You might want to refer to the ports for testing purposes or if you prefer your to use own firewall rules.

The firewall rules for Cloud Volumes ONTAP requires both inbound and outbound rules.

If you're deploying an HA configuration, these are the firewall rules for Cloud Volumes ONTAP in VPC-0.

Inbound rules

The source for inbound rules in the predefined firewall is 0.0.0.0/0.

To create your own firewall, ensure that you add all networks that need to communicate with Cloud Volumes ONTAP, but also ensure to add both address ranges to allow the internal Google Load Balancer to function correctly. These addresses are 130.211.0.0/22 and 35.191.0.0/16. For more information, refer to [Google Cloud documentation: Load Balancer Firewall Rules](#).

Protocol	Port	Purpose
All ICMP	All	Pinging the instance
HTTP	80	HTTP access to the System Manager web console using the IP address of the cluster management LIF
HTTPS	443	HTTPS access to the System Manager web console using the IP address of the cluster management LIF
SSH	22	SSH access to the IP address of the cluster management LIF or a node management LIF

Protocol	Port	Purpose
TCP	111	Remote procedure call for NFS
TCP	139	NetBIOS service session for CIFS
TCP	161-162	Simple network management protocol
TCP	445	Microsoft SMB/CIFS over TCP with NetBIOS framing
TCP	635	NFS mount
TCP	749	Kerberos
TCP	2049	NFS server daemon
TCP	3260	iSCSI access through the iSCSI data LIF
TCP	4045	NFS lock daemon
TCP	4046	Network status monitor for NFS
TCP	10000	Backup using NDMP
TCP	11104	Management of intercluster communication sessions for SnapMirror
TCP	11105	SnapMirror data transfer using intercluster LIFs
TCP	63001-63050	Load balance probe ports to determine which node is healthy (required for HA pairs only)
UDP	111	Remote procedure call for NFS
UDP	161-162	Simple network management protocol
UDP	635	NFS mount
UDP	2049	NFS server daemon
UDP	4045	NFS lock daemon
UDP	4046	Network status monitor for NFS
UDP	4049	NFS rquotad protocol

Outbound rules

The predefined security group for Cloud Volumes ONTAP opens all outbound traffic. If that is acceptable, follow the basic outbound rules. If you need more rigid rules, use the advanced outbound rules.

Basic outbound rules

The predefined security group for Cloud Volumes ONTAP includes the following outbound rules.

Protocol	Port	Purpose
All ICMP	All	All outbound traffic
All TCP	All	All outbound traffic
All UDP	All	All outbound traffic

Advanced outbound rules

If you need rigid rules for outbound traffic, you can use the following information to open only those ports that are required for outbound communication by Cloud Volumes ONTAP.



The source is the interface (IP address) on the Cloud Volumes ONTAP system.

Service	Protocol	Port	Source	Destination	Purpose
Active Directory	TCP	88	Node management LIF	Active Directory forest	Kerberos V authentication
	UDP	137	Node management LIF	Active Directory forest	NetBIOS name service
	UDP	138	Node management LIF	Active Directory forest	NetBIOS datagram service
	TCP	139	Node management LIF	Active Directory forest	NetBIOS service session
	TCP & UDP	389	Node management LIF	Active Directory forest	LDAP
	TCP	445	Node management LIF	Active Directory forest	Microsoft SMB/CIFS over TCP with NetBIOS framing
	TCP	464	Node management LIF	Active Directory forest	Kerberos V change & set password (SET_CHANGE)
	UDP	464	Node management LIF	Active Directory forest	Kerberos key administration
	TCP	749	Node management LIF	Active Directory forest	Kerberos V change & set Password (RPCSEC_GSS)
	TCP	88	Data LIF (NFS, CIFS, iSCSI)	Active Directory forest	Kerberos V authentication
	UDP	137	Data LIF (NFS, CIFS)	Active Directory forest	NetBIOS name service
	UDP	138	Data LIF (NFS, CIFS)	Active Directory forest	NetBIOS datagram service
	TCP	139	Data LIF (NFS, CIFS)	Active Directory forest	NetBIOS service session
	TCP & UDP	389	Data LIF (NFS, CIFS)	Active Directory forest	LDAP
	TCP	445	Data LIF (NFS, CIFS)	Active Directory forest	Microsoft SMB/CIFS over TCP with NetBIOS framing
	TCP	464	Data LIF (NFS, CIFS)	Active Directory forest	Kerberos V change & set password (SET_CHANGE)
	UDP	464	Data LIF (NFS, CIFS)	Active Directory forest	Kerberos key administration
	TCP	749	Data LIF (NFS, CIFS)	Active Directory forest	Kerberos V change & set password (RPCSEC_GSS)

Service	Protocol	Port	Source	Destination	Purpose
Cluster	All traffic	All traffic	All LIFs on one node	All LIFs on the other node	Intercluster communications (Cloud Volumes ONTAP HA only)
	TCP	3000	Node management LIF	HA mediator	ZAPI calls (Cloud Volumes ONTAP HA only)
	ICMP	1	Node management LIF	HA mediator	Keep alive (Cloud Volumes ONTAP HA only)
DHCP	UDP	68	Node management LIF	DHCP	DHCP client for first-time setup
DHCPS	UDP	67	Node management LIF	DHCP	DHCP server
DNS	UDP	53	Node management LIF and data LIF (NFS, CIFS)	DNS	DNS
NDMP	TCP	1860 0–18 699	Node management LIF	Destination servers	NDMP copy
SMTP	TCP	25	Node management LIF	Mail server	SMTP alerts, can be used for AutoSupport
SNMP	TCP	161	Node management LIF	Monitor server	Monitoring by SNMP traps
	UDP	161	Node management LIF	Monitor server	Monitoring by SNMP traps
	TCP	162	Node management LIF	Monitor server	Monitoring by SNMP traps
	UDP	162	Node management LIF	Monitor server	Monitoring by SNMP traps
SnapMirr or	TCP	1110 4	Intercluster LIF	ONTAP intercluster LIFs	Management of intercluster communication sessions for SnapMirror
	TCP	1110 5	Intercluster LIF	ONTAP intercluster LIFs	SnapMirror data transfer
Syslog	UDP	514	Node management LIF	Syslog server	Syslog forward messages

Firewall rules for VPC-1, VPC-2, and VPC-3

In GCP, an HA configuration is deployed across four VPCs. The firewall rules needed for the HA configuration in VPC-0 are [listed above for Cloud Volumes ONTAP](#).

Meanwhile, the predefined firewall policy that Cloud Manager creates for instances in VPC-1, VPC-2, and VPC-3 enables ingress communication over *all* protocols and ports. These rules enable communication between HA nodes.

Communication from the HA nodes to the HA mediator takes place over port 3260 (iSCSI).

Firewall rules for the Connector

The firewall rules for the Connector requires both inbound and outbound rules.

Inbound rules

Protocol	Port	Purpose
SSH	22	Provides SSH access to the Connector host
HTTP	80	Provides HTTP access from client web browsers to the local user interface
HTTPS	443	Provides HTTPS access from client web browsers to the local user interface

Outbound rules

The predefined firewall rules for the Connector opens all outbound traffic. If that is acceptable, follow the basic outbound rules. If you need more rigid rules, use the advanced outbound rules.

Basic outbound rules

The predefined firewall rules for the Connector includes the following outbound rules.

Protocol	Port	Purpose
All TCP	All	All outbound traffic
All UDP	All	All outbound traffic

Advanced outbound rules

If you need rigid rules for outbound traffic, you can use the following information to open only those ports that are required for outbound communication by the Connector.



The source IP address is the Connector host.

Service	Protocol	Port	Destination	Purpose
Active Directory	TCP	88	Active Directory forest	Kerberos V authentication
	TCP	139	Active Directory forest	NetBIOS service session
	TCP	389	Active Directory forest	LDAP
	TCP	445	Active Directory forest	Microsoft SMB/CIFS over TCP with NetBIOS framing
	TCP	464	Active Directory forest	Kerberos V change & set password (SET_CHANGE)
	TCP	749	Active Directory forest	Active Directory Kerberos V change & set password (RPCSEC_GSS)
	UDP	137	Active Directory forest	NetBIOS name service
	UDP	138	Active Directory forest	NetBIOS datagram service
	UDP	464	Active Directory forest	Kerberos key administration
API calls and AutoSupport	HTTP	443	Outbound internet and ONTAP cluster management LIF	API calls to GCP and ONTAP, and sending AutoSupport messages to NetApp
DNS	UDP	53	DNS	Used for DNS resolve by Cloud Manager

Create a service account for data tiering and backups

Cloud Volumes ONTAP requires a Google Cloud service account for two purposes. The first is when you enable [data tiering](#) to tier cold data to low-cost object storage in Google Cloud. The second is when you enable the [Cloud Backup Service](#) to back up volumes to low-cost object storage.

Cloud Volumes ONTAP uses the service account to access and manage one bucket for tiered data and another bucket for backups.

You can set up one service account and use it for both purposes. The service account must have the **Storage Admin** role.

Steps

1. In the Google Cloud console, [go to the Service accounts page](#).
2. Select your project.
3. Click **Create service account** and provide the required information.
 - a. **Service account details:** Enter a name and description.
 - b. **Grant this service account access to project:** Select the **Storage Admin** role.

Create service account

Service account details

2 Grant this service account access to project (optional)

Grant this service account access to OCCM-Dev so that it has permission to complete specific actions on the resources in your project. [Learn more](#)

Select a role Condition

Role	Description	Condition
storage admin	Cloud ML Service Agent Cloud ML service agent can act as log writer, Cloud Storage admin, Artifact Registry Reader, BigQuery writer, and service account access token creator.	
Storage Admin	Storage Admin Full control of GCS resources.	 optional)
Storage HMAC Key Admin	Storage HMAC Key Admin Full control of GCS HMAC Keys.	
Storage Object Admin	Storage Object Admin Full control of GCS objects.	

DONE

- c. **Grant users access to this service account:** Add the Connector service account as a *Service Account User* to this new service account.

This step is required for data tiering only. It's not required for the Cloud Backup Service.

Create service account

Service account details

Grant this service account access to project (optional)

3 Grant users access to this service account (optional)

Grant access to users or groups that need to perform actions as this service account. [Learn more](#)

Service account users role

netapp-cloud-manager@iam.gserviceaccount.com 



Grant users the permissions to deploy jobs and VMs with this service account

Service account admins role



Grant users the permission to administer this service account

DONE

CANCEL

What's next?

You'll need to select the service account later when you create a Cloud Volumes ONTAP working environment.

Details and Credentials

default-project Google Cloud Project	gcp-sub2 Marketplace Subscription	Edit Project
<div style="display: flex; justify-content: space-between;"> <div style="width: 45%;"> <p>Details</p> <p>Working Environment Name (Cluster Name)</p> <input type="text" value="cloudvolumesontap"/> </div> <div style="width: 45%;"> <p>Credentials</p> <p>User Name</p> <input type="text" value="admin"/> </div> </div> <hr/> <div style="margin-top: 10px;"> <div style="display: flex; align-items: center;"> Service Account <small>i</small> <input checked="" type="checkbox"/> </div> <div style="margin-top: 10px;"> <p>Service Account Name</p> <input type="text" value="account1"/> </div> </div> <div style="display: flex; justify-content: space-between; margin-top: 10px;"> + Add Labels Optional Field Up to four labels </div>		

Using customer-managed encryption keys with Cloud Volumes ONTAP

While Google Cloud Storage always encrypts your data before it's written to disk, you can use the Cloud Manager API to create a Cloud Volumes ONTAP system that uses *customer-managed encryption keys*. These are keys that you generate and manage in GCP using the Cloud Key Management Service.

Steps

1. Ensure that the Cloud Manager Connector service account has the correct permissions at the project level, in the project where the key is stored.

The permissions are provided by the [Cloud Manager YAML file](#) by default, but may not be applied if you use an alternate project for the Cloud Key Management Service.

The permissions are as follows:

- `cloudkms.cryptoKeyVersions.list`
- `cloudkms.cryptoKeys.get`
- `cloudkms.cryptoKeys.list`
- `cloudkms.keyRings.list`

2. Ensure that the service account for the [Google Compute Engine Service Agent](#) has Cloud KMS Encrypter/Decrypter permissions on the key.

The name of the service account uses the following format: "service-[service_project_number]@compute-system.iam.gserviceaccount.com".

[Google Cloud Documentation: Using IAM with Cloud KMS - Granting roles on a resource](#)

3. Obtain the "id" of the key by invoking the get command for the /gcp/vsa/metadata/gcp-encryption-keys API call or by choosing "Copy Resource Name" on the key in the GCP console.
4. If using customer-managed encryption keys and tiering data to object storage, Cloud Manager attempts to utilize the same keys that are used to encrypt the persistent disks. But you'll first need to enable Google Cloud Storage buckets to use the keys:
 - a. Find the Google Cloud Storage service agent by following the [Google Cloud Documentation: Getting the Cloud Storage service agent](#).
 - b. Navigate to the encryption key and assign the Google Cloud Storage service agent with Cloud KMS Encrypter/Decrypter permissions.

For more information, refer to [Google Cloud Documentation: Using customer-managed encryption keys](#)

5. Use the "GcpEncryption" parameter with your API request when creating a working environment.

Example

```
"gcpEncryptionParameters": {  
    "key": "projects/project-1/locations/us-east4/keyRings/keyring-1/cryptoKeys/generatedkey1"  
}
```

Refer to the [Cloud Manager automation docs](#) for more details about using the "GcpEncryption" parameter.

Launching Cloud Volumes ONTAP in GCP

You can launch Cloud Volumes ONTAP in a single-node configuration or as an HA pair in Google Cloud Platform.

Before you get started

You need the following to create a working environment.

- A Connector that's up and running.
 - You should have a [Connector that is associated with your workspace](#).
 - [You should be prepared to leave the Connector running at all times](#).
 - The service account associated with the Connector [should have the latest permissions](#).
- An understanding of the configuration that you want to use.

You should have prepared by choosing a configuration and by obtaining GCP networking information from your administrator. For details, see [Planning your Cloud Volumes ONTAP configuration](#).

- An understanding of what's required to choose a specific licensing option in the Add Working Environment wizard.

Licensing option	Requirement	How to meet the requirement
PAYGO free trial	A Marketplace subscription is required.	You'll have the option to subscribe to your cloud provider's marketplace from the Details & Credentials page.
Freemium	A Marketplace subscription or NetApp Support Site (NSS) account is required.	You'll have the option to subscribe to your cloud provider's marketplace from the Details & Credentials page You can enter your NSS account on the Charging Methods and NSS Account page.
Capacity-based BYOL	A Marketplace subscription or NetApp Support Site (NSS) account is required. A Marketplace subscription is recommended for capacity-based charging in the event that your account doesn't have a valid capacity-based license, or in the event that your provisioned capacity exceeds the licensed capacity.	You'll have the option to subscribe to your cloud provider's marketplace from the Details & Credentials page You can enter your NSS account on the Charging Methods and NSS Account page.
Node-based BYOL	The 20-digit serial number (license key) is required.	You'll enter the serial number on the Charging Methods and NSS Account page.

- Google Cloud APIs should be [enabled in your project](#):
 - Cloud Deployment Manager V2 API
 - Cloud Logging API
 - Cloud Resource Manager API
 - Compute Engine API
 - Identity and Access Management (IAM) API

Launching a single-node system in GCP

Create a working environment in Cloud Manager to launch Cloud Volumes ONTAP in GCP.

Steps

1. On the Canvas page, click **Add Working Environment** and follow the prompts.
2. **Choose a Location:** Select **Google Cloud** and **Cloud Volumes ONTAP**.
3. If you're prompted, [create a Connector](#).
4. **Details & Credentials:** Select a project, specify a cluster name, optionally select a Service Account, optionally add labels, and then specify credentials.

The following table describes fields for which you might need guidance:

Field	Description
Working Environment Name	Cloud Manager uses the working environment name to name both the Cloud Volumes ONTAP system and the GCP VM instance. It also uses the name as the prefix for the predefined security group, if you select that option.
Service Account Name	If you plan to use data tiering or Cloud Backup with Cloud Volumes ONTAP, then you need to enable Service Account and select a service account that has the predefined Storage Admin role. Learn how to create a service account .
Add Labels	<p>Labels are metadata for your GCP resources. Cloud Manager adds the labels to the Cloud Volumes ONTAP system and GCP resources associated with the system.</p> <p>You can add up to four labels from the user interface when creating a working environment, and then you can add more after its created. Note that the API does not limit you to four labels when creating a working environment.</p> <p>For information about labels, refer to Google Cloud Documentation: Labeling Resources.</p>
User name and password	These are the credentials for the Cloud Volumes ONTAP cluster admin account. You can use these credentials to connect to Cloud Volumes ONTAP through System Manager or its CLI.
Edit Project	<p>Select the project where you want Cloud Volumes ONTAP to reside. The default project is the project where Cloud Manager resides.</p> <p>If you don't see any additional projects in the drop-down list, then you haven't yet associated the Cloud Manager service account with other projects. Go to the Google Cloud console, open the IAM service, and select the project. Add the service account with the Cloud Manager role to that project. You'll need to repeat this step for each project.</p> <p> This is the service account that you set up for Cloud Manager, as described on this page.</p> <p>Click Add Subscription to associate the selected credentials with a subscription.</p> <p>To create a pay-as-you-go Cloud Volumes ONTAP system, you need to select a GCP project that's associated with a subscription to Cloud Volumes ONTAP from the GCP Marketplace.</p>

The following video shows how to associate a pay-as-you-go Marketplace subscription to your GCP project:

► https://docs.netapp.com/us-en/occm//media/video_subscribing_gcp.mp4 (video)

5. **Services:** Select the services that you want to use on this system. In order to select Cloud Backup, or to use Tiering, you must have specified the Service Account in step 3.
6. **Location & Connectivity:** Select a location, choose a firewall policy, and select the checkbox to confirm network connectivity to Google Cloud storage for data tiering.

If you want to tier cold data to a Google Cloud Storage bucket, the subnet in which Cloud Volumes ONTAP

resides must be configured for Private Google Access. For instructions, refer to [Google Cloud Documentation: Configuring Private Google Access](#).

7. **Charging Methods and NSS Account:** Specify which charging option would you like to use with this system, and then specify a NetApp Support Site account.
 - [Learn about these charging methods.](#)
 - [Learn what's required in the wizard for the licensing method that you want to use.](#)
8. **Preconfigured Packages:** Select one of the packages to quickly deploy a Cloud Volumes ONTAP system, or click **Create my own configuration**.

If you choose one of the packages, you only need to specify a volume and then review and approve the configuration.
9. **Licensing:** Change the Cloud Volumes ONTAP version as needed, select a license, and select a virtual machine type.

If your needs change after you launch the system, you can modify the license or virtual machine type later.



If a newer Release Candidate, General Availability, or patch release is available for the selected version, then Cloud Manager updates the system to that version when creating the working environment. For example, the update occurs if you select Cloud Volumes ONTAP 9.6 RC1 and 9.6 GA is available. The update does not occur from one release to another—for example, from 9.6 to 9.7.

10. **Underlying Storage Resources:** Choose settings for the initial aggregate: a disk type and the size for each disk.

The disk type is for the initial volume. You can choose a different disk type for subsequent volumes.

The disk size is for all disks in the initial aggregate and for any additional aggregates that Cloud Manager creates when you use the simple provisioning option. You can create aggregates that use a different disk size by using the advanced allocation option.

For help choosing a disk type and size, see [Sizing your system in GCP](#).

11. **Write Speed & WORM:** Choose **Normal** or **High** write speed, and activate write once, read many (WORM) storage, if desired.

Choosing a write speed is supported with single node systems only.

[Learn more about write speed](#).

WORM can't be enabled if Cloud Backup was enabled or if data tiering was enabled.

[Learn more about WORM storage](#).

12. **Data Tiering in Google Cloud Platform:** Choose whether to enable data tiering on the initial aggregate, choose a storage class for the tiered data, and then either select a service account that has the predefined Storage Admin role (required for Cloud Volumes ONTAP 9.7 or later), or select a GCP account (required for Cloud Volumes ONTAP 9.6).

Note the following:

- Cloud Manager sets the service account on the Cloud Volumes ONTAP instance. This service account provides permissions for data tiering to a Google Cloud Storage bucket. Be sure to add the Connector service account as a user of the tiering service account, otherwise, you can't select it from Cloud Manager.
- For help with adding a GCP account, see [Setting up and adding GCP accounts for data tiering with 9.6](#).
- You can choose a specific volume tiering policy when you create or edit a volume.
- If you disable data tiering, you can enable it on subsequent aggregates, but you'll need to turn off the system and add a service account from the GCP console.

[Learn more about data tiering](#).

13. **Create Volume:** Enter details for the new volume or click **Skip**.

Some of the fields in this page are self-explanatory. The following table describes fields for which you might need guidance:

Field	Description
Size	The maximum size that you can enter largely depends on whether you enable thin provisioning, which enables you to create a volume that is bigger than the physical storage currently available to it.
Access control (for NFS only)	An export policy defines the clients in the subnet that can access the volume. By default, Cloud Manager enters a value that provides access to all instances in the subnet.
Permissions and Users / Groups (for CIFS only)	These fields enable you to control the level of access to a share for users and groups (also called access control lists or ACLs). You can specify local or domain Windows users or groups, or UNIX users or groups. If you specify a domain Windows user name, you must include the user's domain using the format domain\username.
Snapshot Policy	A Snapshot copy policy specifies the frequency and number of automatically created NetApp Snapshot copies. A NetApp Snapshot copy is a point-in-time file system image that has no performance impact and requires minimal storage. You can choose the default policy or none. You might choose none for transient data: for example, tempdb for Microsoft SQL Server.
Advanced options (for NFS only)	Select an NFS version for the volume: either NFSv3 or NFSv4.

Field	Description
Initiator group and IQN (for iSCSI only)	<p>iSCSI storage targets are called LUNs (logical units) and are presented to hosts as standard block devices.</p> <p>Initiator groups are tables of iSCSI host node names and control which initiators have access to which LUNs.</p> <p>iSCSI targets connect to the network through standard Ethernet network adapters (NICs), TCP offload engine (TOE) cards with software initiators, converged network adapters (CNAs) or dedicated host bus adapters (HBAs) and are identified by iSCSI qualified names (IQNs).</p> <p>When you create an iSCSI volume, Cloud Manager automatically creates a LUN for you. We've made it simple by creating just one LUN per volume, so there's no management involved. After you create the volume, use the IQN to connect to the LUN from your hosts.</p>

The following image shows the Volume page filled out for the CIFS protocol:

The screenshot shows the 'Volume Details, Protection & Protocol' page. On the left, under 'Details & Protection', there are fields for 'Volume Name' (vol) and 'Size (GB)' (250). Below these are 'Snapshot Policy' (default) and a note about 'Default Policy'. On the right, under 'Protocol', the 'CIFS' tab is selected, showing 'Share name' (vol_share), 'Permissions' (Full Control), and 'Users / Groups' (engineering). A note at the bottom states 'Valid users and groups separated by a semicolon'.

14. CIFS Setup: If you chose the CIFS protocol, set up a CIFS server.

Field	Description
DNS Primary and Secondary IP Address	The IP addresses of the DNS servers that provide name resolution for the CIFS server. The listed DNS servers must contain the service location records (SRV) needed to locate the Active Directory LDAP servers and domain controllers for the domain that the CIFS server will join.
Active Directory Domain to join	The FQDN of the Active Directory (AD) domain that you want the CIFS server to join.
Credentials authorized to join the domain	The name and password of a Windows account with sufficient privileges to add computers to the specified Organizational Unit (OU) within the AD domain.
CIFS server NetBIOS name	A CIFS server name that is unique in the AD domain.

Field	Description
Organizational Unit	The organizational unit within the AD domain to associate with the CIFS server. The default is CN=Computers.
DNS Domain	The DNS domain for the Cloud Volumes ONTAP storage virtual machine (SVM). In most cases, the domain is the same as the AD domain.
NTP Server	Select Use Active Directory Domain to configure an NTP server using the Active Directory DNS. If you need to configure an NTP server using a different address, then you should use the API. See the Cloud Manager automation docs for details.

15. **Usage Profile, Disk Type, and Tiering Policy:** Choose whether you want to enable storage efficiency features and change the volume tiering policy, if needed.

For more information, see [Understanding volume usage profiles](#) and [Data tiering overview](#).

16. **Review & Approve:** Review and confirm your selections.

- Review details about the configuration.
- Click **More information** to review details about support and the GCP resources that Cloud Manager will purchase.
- Select the **I understand...** check boxes.
- Click **Go**.

Result

Cloud Manager deploys the Cloud Volumes ONTAP system. You can track the progress in the timeline.

If you experience any issues deploying the Cloud Volumes ONTAP system, review the failure message. You can also select the working environment and click **Re-create environment**.

For additional help, go to [NetApp Cloud Volumes ONTAP Support](#).

After you finish

- If you provisioned a CIFS share, give users or groups permissions to the files and folders and verify that those users can access the share and create a file.
- If you want to apply quotas to volumes, use System Manager or the CLI.

Quotas enable you to restrict or track the disk space and number of files used by a user, group, or qtree.

Launching an HA pair in GCP

Create a working environment in Cloud Manager to launch Cloud Volumes ONTAP in GCP.

Steps

- On the Canvas page, click **Add Working Environment** and follow the prompts.
- Choose a Location:** Select **Google Cloud** and **Cloud Volumes ONTAP HA**.
- Details & Credentials:** Select a project, specify a cluster name, optionally select a Service Account, optionally add labels, and then specify credentials.

The following table describes fields for which you might need guidance:

Field	Description
Working Environment Name	Cloud Manager uses the working environment name to name both the Cloud Volumes ONTAP system and the GCP VM instance. It also uses the name as the prefix for the predefined security group, if you select that option.
Service Account Name	If you plan to use the Tiering or Cloud Backup services, you need to enable the Service Account switch and then select the Service Account that has the predefined Storage Admin role.
Add Labels	<p>Labels are metadata for your GCP resources. Cloud Manager adds the labels to the Cloud Volumes ONTAP system and GCP resources associated with the system.</p> <p>You can add up to four labels from the user interface when creating a working environment, and then you can add more after its created. Note that the API does not limit you to four labels when creating a working environment.</p> <p>For information about labels, refer to Google Cloud Documentation: Labeling Resources.</p>
User name and password	These are the credentials for the Cloud Volumes ONTAP cluster admin account. You can use these credentials to connect to Cloud Volumes ONTAP through System Manager or its CLI.
Edit Project	<p>Select the project where you want Cloud Volumes ONTAP to reside. The default project is the project where Cloud Manager resides.</p> <p>If you don't see any additional projects in the drop-down list, then you haven't yet associated the Cloud Manager service account with other projects. Go to the Google Cloud console, open the IAM service, and select the project. Add the service account with the Cloud Manager role to that project. You'll need to repeat this step for each project.</p> <p> This is the service account that you set up for Cloud Manager, as described in step 4b on this page.</p> <p>Click Add Subscription to associate the selected credentials with a subscription.</p> <p>To create a pay-as-you-go Cloud Volumes ONTAP system, you need to select a GCP project that's associated with a subscription to Cloud Volumes ONTAP from the GCP Marketplace.</p>

The following video shows how to associate a pay-as-you-go Marketplace subscription to your GCP project:

► https://docs.netapp.com/us-en/occm//media/video_subscribing_gcp.mp4 (video)

4. **Services:** Select the services that you want to use on this system. In order to select Cloud Backup, or to use Tiering, you must have specified the Service Account in step 3.
5. **HA Deployment Models:** Choose multiple zones (recommended) or a single zone for the HA configuration. Then select a region and zones.

[Learn more about HA deployment models.](#)

6. **Connectivity:** Select four different VPCs for the HA configuration, a subnet in each VPC, and then choose a firewall policy.

[Learn more about networking requirements.](#)

7. **Charging Methods and NSS Account:** Specify which charging option would you like to use with this system, and then specify a NetApp Support Site account.

- [Learn about these charging methods.](#)
- [Learn what's required in the wizard for the licensing method that you want to use.](#)

8. **Preconfigured Packages:** Select one of the packages to quickly deploy a Cloud Volumes ONTAP system, or click **Create my own configuration**.

If you choose one of the packages, you only need to specify a volume and then review and approve the configuration.

9. **Licensing:** Change the Cloud Volumes ONTAP version as needed, select a license, and select a virtual machine type.

If your needs change after you launch the system, you can modify the license or virtual machine type later.



If a newer Release Candidate, General Availability, or patch release is available for the selected version, then Cloud Manager updates the system to that version when creating the working environment. For example, the update occurs if you select Cloud Volumes ONTAP 9.8 RC1 and 9.8 GA is available. The update does not occur from one release to another—for example, from 9.7 to 9.8.

10. **Underlying Storage Resources:** Choose settings for the initial aggregate: a disk type and the size for each disk.

The disk type is for the initial volume. You can choose a different disk type for subsequent volumes.

The disk size is for all disks in the initial aggregate and for any additional aggregates that Cloud Manager creates when you use the simple provisioning option. You can create aggregates that use a different disk size by using the advanced allocation option.

For help choosing a disk type and size, see [Sizing your system in GCP](#).

11. **WORM:** Activate write once, read many (WORM) storage, if desired.

WORM can't be enabled if data tiering was enabled. [Learn more about WORM storage.](#)

12. **Data Tiering in Google Cloud Platform:** Choose whether to enable data tiering on the initial aggregate, choose a storage class for the tiered data, and then select a service account that has the predefined Storage Admin role.

Note the following:

- Cloud Manager sets the service account on the Cloud Volumes ONTAP instance. This service account provides permissions for data tiering to a Google Cloud Storage bucket. Be sure to add the Connector service account as a user of the tiering service account, otherwise, you can't select it from Cloud Manager.

- You can choose a specific volume tiering policy when you create or edit a volume.
- If you disable data tiering, you can enable it on subsequent aggregates, but you'll need to turn off the system and add a service account from the GCP console.

[Learn more about data tiering.](#)

13. **Create Volume:** Enter details for the new volume or click **Skip**.

Some of the fields in this page are self-explanatory. The following table describes fields for which you might need guidance:

Field	Description
Size	The maximum size that you can enter largely depends on whether you enable thin provisioning, which enables you to create a volume that is bigger than the physical storage currently available to it.
Access control (for NFS only)	An export policy defines the clients in the subnet that can access the volume. By default, Cloud Manager enters a value that provides access to all instances in the subnet.
Permissions and Users / Groups (for CIFS only)	These fields enable you to control the level of access to a share for users and groups (also called access control lists or ACLs). You can specify local or domain Windows users or groups, or UNIX users or groups. If you specify a domain Windows user name, you must include the user's domain using the format domain\username.
Snapshot Policy	A Snapshot copy policy specifies the frequency and number of automatically created NetApp Snapshot copies. A NetApp Snapshot copy is a point-in-time file system image that has no performance impact and requires minimal storage. You can choose the default policy or none. You might choose none for transient data: for example, tempdb for Microsoft SQL Server.
Advanced options (for NFS only)	Select an NFS version for the volume: either NFSv3 or NFSv4.
Initiator group and IQN (for iSCSI only)	iSCSI storage targets are called LUNs (logical units) and are presented to hosts as standard block devices. Initiator groups are tables of iSCSI host node names and control which initiators have access to which LUNs. iSCSI targets connect to the network through standard Ethernet network adapters (NICs), TCP offload engine (TOE) cards with software initiators, converged network adapters (CNAs) or dedicated host bus adapters (HBAs) and are identified by iSCSI qualified names (IQNs). When you create an iSCSI volume, Cloud Manager automatically creates a LUN for you. We've made it simple by creating just one LUN per volume, so there's no management involved. After you create the volume, use the IQN to connect to the LUN from your hosts .

The following image shows the Volume page filled out for the CIFS protocol:

Volume Details, Protection & Protocol

Details & Protection <div style="margin-top: 10px;"> Volume Name: <input type="text" value="vol"/> Size (GB): <input type="text" value="250"/> </div> <div style="margin-top: 10px;"> Snapshot Policy: <input type="text" value="default"/> <div style="float: right;">(i) Default Policy</div> </div>	Protocol <div style="margin-top: 10px;"> NFS CIFS iSCSI </div> <div style="margin-top: 10px;"> Share name: <input type="text" value="vol_share"/> Permissions: <input type="text" value="Full Control"/> </div> <div style="margin-top: 10px;"> Users / Groups: <input type="text" value="engineering"/> <div style="font-size: small; margin-top: -10px;">Valid users and groups separated by a semicolon</div> </div>
---	---

14. **CIFS Setup:** If you chose the CIFS protocol, set up a CIFS server.

Field	Description
DNS Primary and Secondary IP Address	The IP addresses of the DNS servers that provide name resolution for the CIFS server. The listed DNS servers must contain the service location records (SRV) needed to locate the Active Directory LDAP servers and domain controllers for the domain that the CIFS server will join.
Active Directory Domain to join	The FQDN of the Active Directory (AD) domain that you want the CIFS server to join.
Credentials authorized to join the domain	The name and password of a Windows account with sufficient privileges to add computers to the specified Organizational Unit (OU) within the AD domain.
CIFS server NetBIOS name	A CIFS server name that is unique in the AD domain.
Organizational Unit	The organizational unit within the AD domain to associate with the CIFS server. The default is CN=Computers.
DNS Domain	The DNS domain for the Cloud Volumes ONTAP storage virtual machine (SVM). In most cases, the domain is the same as the AD domain.
NTP Server	Select Use Active Directory Domain to configure an NTP server using the Active Directory DNS. If you need to configure an NTP server using a different address, then you should use the API. See the Cloud Manager automation docs for details.

15. **Usage Profile, Disk Type, and Tiering Policy:** Choose whether you want to enable storage efficiency features and change the volume tiering policy, if needed.

For more information, see [Understanding volume usage profiles](#) and [Data tiering overview](#).

16. **Review & Approve:** Review and confirm your selections.

- a. Review details about the configuration.
- b. Click **More information** to review details about support and the GCP resources that Cloud Manager will purchase.

c. Select the **I understand...** check boxes.

d. Click **Go**.

Result

Cloud Manager deploys the Cloud Volumes ONTAP system. You can track the progress in the timeline.

If you experience any issues deploying the Cloud Volumes ONTAP system, review the failure message. You can also select the working environment and click **Re-create environment**.

For additional help, go to [NetApp Cloud Volumes ONTAP Support](#).

After you finish

- If you provisioned a CIFS share, give users or groups permissions to the files and folders and verify that those users can access the share and create a file.
- If you want to apply quotas to volumes, use System Manager or the CLI.

Quotas enable you to restrict or track the disk space and number of files used by a user, group, or qtree.

Provision and manage storage

Provisioning storage

You can provision additional storage for your Cloud Volumes ONTAP systems from Cloud Manager by managing volumes and aggregates.

 All disks and aggregates must be created and deleted directly from Cloud Manager. You should not perform these actions from another management tool. Doing so can impact system stability, hamper the ability to add disks in the future, and potentially generate redundant cloud provider fees.

Creating FlexVol volumes

If you need more storage after you launch your initial Cloud Volumes ONTAP system, you can create new FlexVol volumes for NFS, CIFS, or iSCSI from Cloud Manager.

A Cloud Manager feature called "templates" enables you to create volumes that are optimized for the workload requirements for certain applications; such as databases or streaming services. If your organization has created volume templates that you should use, follow [these steps](#).

About this task

When you create an iSCSI volume, Cloud Manager automatically creates a LUN for you. We've made it simple by creating just one LUN per volume, so there's no management involved. After you create the volume, [use the IQN to connect to the LUN from your hosts](#).

 You can create additional LUNs from System Manager or the CLI.

Before you begin

If you want to use CIFS in AWS, you must have set up DNS and Active Directory. For details, see [Networking requirements for Cloud Volumes ONTAP for AWS](#).

Steps

1. On the Canvas page, double-click the name of the Cloud Volumes ONTAP system on which you want to provision FlexVol volumes.
2. Create a new volume on any aggregate or on a specific aggregate:

Action	Steps
Create a new volume and let Cloud Manager choose the containing aggregate	Click Add Volume > New volume .
Create a new volume on a specific aggregate	<ol style="list-style-type: none">a. Click the menu icon, and then click Advanced > Advanced allocation.b. Click the menu for an aggregate.c. Click Create volume.

3. Enter details for the new volume, and then click **Continue**.

Some of the fields in this page are self-explanatory. The following table describes fields for which you might need guidance:

Field	Description
Size	The maximum size that you can enter largely depends on whether you enable thin provisioning, which enables you to create a volume that is bigger than the physical storage currently available to it.
Access control (for NFS only)	An export policy defines the clients in the subnet that can access the volume. By default, Cloud Manager enters a value that provides access to all instances in the subnet.
Permissions and Users / Groups (for CIFS only)	These fields enable you to control the level of access to a share for users and groups (also called access control lists or ACLs). You can specify local or domain Windows users or groups, or UNIX users or groups. If you specify a domain Windows user name, you must include the user's domain using the format domain\username.
Snapshot Policy	A Snapshot copy policy specifies the frequency and number of automatically created NetApp Snapshot copies. A NetApp Snapshot copy is a point-in-time file system image that has no performance impact and requires minimal storage. You can choose the default policy or none. You might choose none for transient data: for example, tempdb for Microsoft SQL Server.
Advanced options (for NFS only)	Select an NFS version for the volume: either NFSv3 or NFSv4.

Field	Description
Initiator group and IQN (for iSCSI only)	<p>iSCSI storage targets are called LUNs (logical units) and are presented to hosts as standard block devices.</p> <p>Initiator groups are tables of iSCSI host node names and control which initiators have access to which LUNs.</p> <p>iSCSI targets connect to the network through standard Ethernet network adapters (NICs), TCP offload engine (TOE) cards with software initiators, converged network adapters (CNAs) or dedicated host bus adapters (HBAs) and are identified by iSCSI qualified names (IQNs).</p> <p>When you create an iSCSI volume, Cloud Manager automatically creates a LUN for you. We've made it simple by creating just one LUN per volume, so there's no management involved. After you create the volume, use the IQN to connect to the LUN from your hosts.</p>

4. If you chose the CIFS protocol and the CIFS server has not been set up, specify details for the server in the Create a CIFS Server dialog box, and then click **Save and continue**:

Field	Description
DNS Primary and Secondary IP Address	<p>The IP addresses of the DNS servers that provide name resolution for the CIFS server.</p> <p>The listed DNS servers must contain the service location records (SRV) needed to locate the Active Directory LDAP servers and domain controllers for the domain that the CIFS server will join.</p>
Active Directory Domain to join	The FQDN of the Active Directory (AD) domain that you want the CIFS server to join.
Credentials authorized to join the domain	The name and password of a Windows account with sufficient privileges to add computers to the specified Organizational Unit (OU) within the AD domain.
CIFS server NetBIOS name	A CIFS server name that is unique in the AD domain.
Organizational Unit	<p>The organizational unit within the AD domain to associate with the CIFS server. The default is CN=Computers.</p> <ul style="list-style-type: none"> • To configure AWS Managed Microsoft AD as the AD server for Cloud Volumes ONTAP, you should enter OU=Computers,OU=corp in this field. • To configure Azure AD Domain Services as the AD server for Cloud Volumes ONTAP, you should enter OU=AADDC Computers or OU=AADDC Users in this field. Azure Documentation: Create an Organizational Unit (OU) in an Azure AD Domain Services managed domain
DNS Domain	The DNS domain for the Cloud Volumes ONTAP storage virtual machine (SVM). In most cases, the domain is the same as the AD domain.

Field	Description
NTP Server	Select Use Active Directory Domain to configure an NTP server using the Active Directory DNS. If you need to configure an NTP server using a different address, then you should use the API. See the Cloud Manager automation docs for details.

5. On the Usage Profile, Disk Type, and Tiering Policy page, choose whether you want to enable storage efficiency features, choose a disk type, and edit the tiering policy, if needed.

For help, refer to the following:

- [Understanding volume usage profiles](#)
- [Sizing your system in AWS](#)
- [Sizing your system in Azure](#)
- [Sizing your system in GCP](#)
- [Data tiering overview](#)

6. Click **Go**.

Result

Cloud Volumes ONTAP provisions the volume.

After you finish

If you provisioned a CIFS share, give users or groups permissions to the files and folders and verify that those users can access the share and create a file.

If you want to apply quotas to volumes, you must use System Manager or the CLI. Quotas enable you to restrict or track the disk space and number of files used by a user, group, or qtree.

Creating volumes from templates

If your organization has created Cloud Volumes ONTAP volume templates so you can deploy volumes that are optimized for the workload requirements for certain applications, follow the steps in this section.

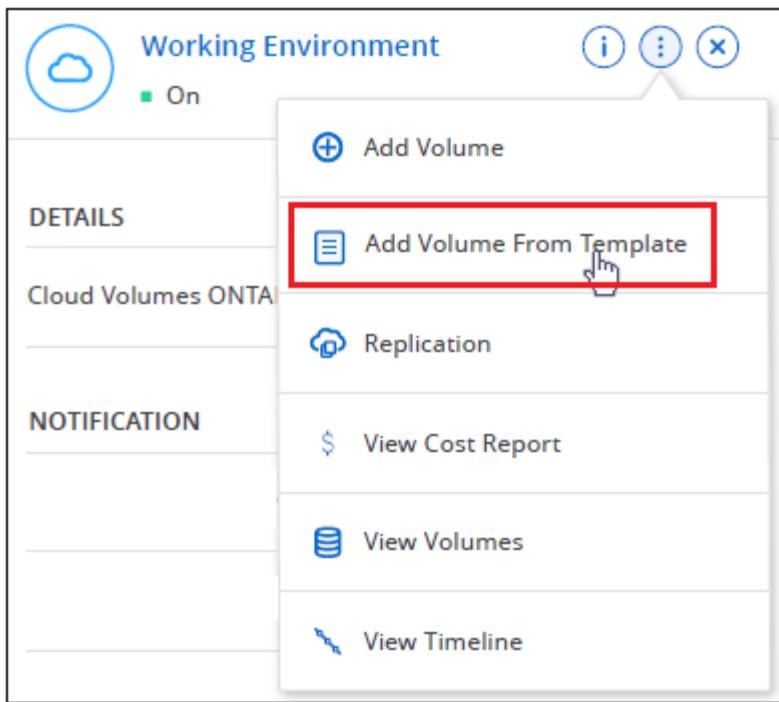
The template should make your job easier because certain volume parameters will already be defined in the template, such as disk type, size, protocol, snapshot policy, cloud provider, and more. When a parameter is already predefined, you can just skip to the next volume parameter.



You can only create NFS or CIFS volumes when using templates.

Steps

1. On the Canvas page, click the name of the Cloud Volumes ONTAP system on which you want to provision a volume.
2. Click > **Add Volume From Template**.



3. In the *Select Template* page, select the template that you want to use to create the volume and click **Next**.

The screenshot shows the 'Select Template' page. At the top, there are two numbered steps: '1 Select Template' and '2 Define Parameters'. Below that is a table titled 'Select Template' with the heading '2 Templates'. The table has columns: 'Template Name', 'Template Description', 'Created by', 'Last Modified', and 'Parameters'. There are two rows of data:

- The first row contains 'AWS CVO volume for production' (with a checked checkbox and highlighted by a red box), 'High efficiency with daily backup', 'Rabin', 'Apr 05 2021, 1:39:30 pm', and a 'View' link.
- The second row contains 'AWS CVO volume for staging', 'High efficiency with weekly backup', 'Rabin', 'Apr 05 2021, 1:34:20 pm', and a 'View' link.

Template Name	Template Description	Created by	Last Modified	Parameters
AWS CVO volume for production	High efficiency with daily backup	Rabin	Apr 05 2021, 1:39:30 pm	View
AWS CVO volume for staging	High efficiency with weekly backup	Rabin	Apr 05 2021, 1:34:20 pm	View

The *Define Parameters* page is displayed.

Define Parameters

Enter your values for the actions. Parameters that are locked by the template are not editable.

Actions

```

graph TD
    A[Create Volume in Cloud Volumes  
ONTAP (1)] --> B[Enable Cloud Backup (1)]
    
```

Show read-only parameters

Details ⓘ

Volume Name
Volume Name should start with "staging"

Volume Size (GB)
Minimum value is 160, Maximum value is 185

Protection ⓘ

Snapshot Policy
Default

Usage Profile ⓘ

Storage Efficiency No Storage Efficiency

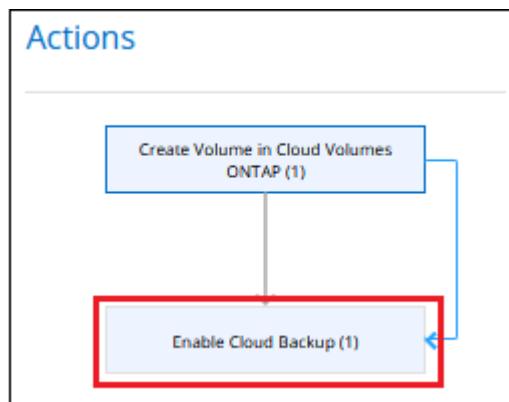
Disk Type

Disk Type
GP2 - General Purpose SSD

Note: You can click the checkbox **Show read-only parameters** to show all the fields that have been locked by the template if you want to see the values for those parameters. By default these predefined fields are hidden and only the fields you need to complete are shown.

4. In the *Context* area, the Working Environment is filled in with the name of the working environment you started with. You need to select the **Storage VM** where the volume will be created.
5. Add values for all of the parameters that are not hard-coded from the template. See [creating volumes](#) for details about all the parameters you need to complete to deploy a Cloud Volumes ONTAP volume.
6. If there are no other Actions that you need to define (for example, configuring Cloud Backup), click **Run Template**.

If there are other actions, click the action in the left pane to display the parameters you need to complete.

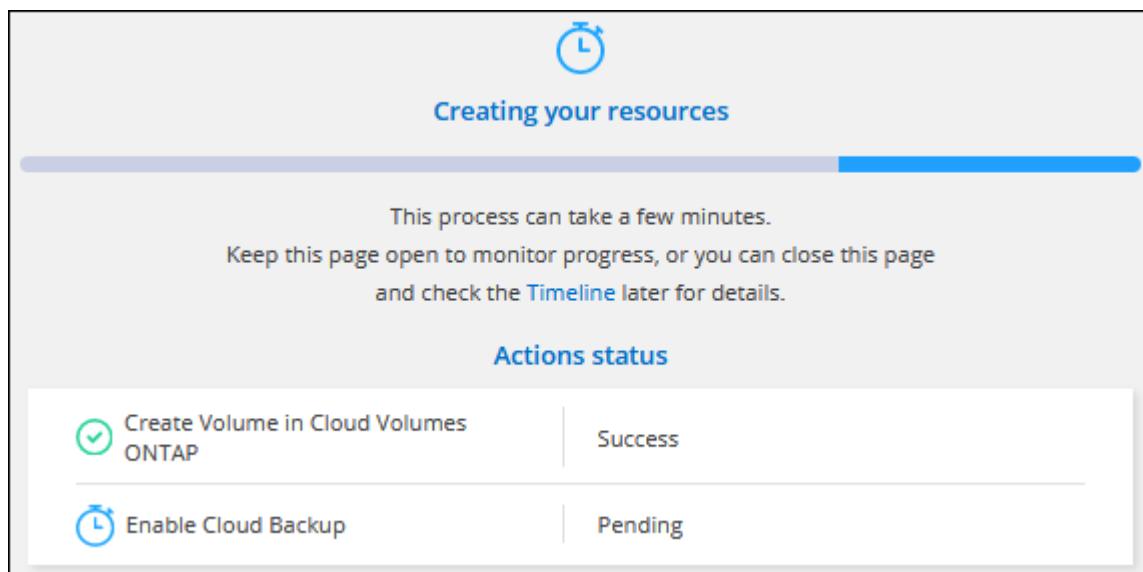


For example, if the Enable Cloud Backup action requires that you select a backup policy, you can do that now.

7. Click **Run Template**.

Result

Cloud Volumes ONTAP provisions the volume and displays a page so that you can see the progress.



Additionally, if any secondary action is implemented in the template, for example, enabling Cloud Backup on the volume, that action is also performed.

After you finish

If you provisioned a CIFS share, give users or groups permissions to the files and folders and verify that those users can access the share and create a file.

If you want to apply quotas to volumes, you must use System Manager or the CLI. Quotas enable you to restrict or track the disk space and number of files used by a user, group, or qtree.

Creating FlexVol volumes on the second node in an HA configuration

By default, Cloud Manager creates volumes on the first node in an HA configuration. If you need an active-active configuration, in which both nodes serve data to clients, you must create aggregates and volumes on the second node.

Steps

1. On the Canvas page, double-click the name of the Cloud Volumes ONTAP working environment on which you want to manage aggregates.
2. Click the menu icon and then click **Advanced > Advanced allocation**.
3. Click **Add Aggregate** and then create the aggregate.
4. For Home Node, choose the second node in the HA pair.
5. After Cloud Manager creates the aggregate, select it and then click **Create volume**.
6. Enter details for the new volume, and then click **Create**.

After you finish

You can create additional volumes on this aggregate if required.



For HA pairs deployed in multiple AWS Availability Zones, you must mount the volume to clients by using the floating IP address of the node on which the volume resides.

Creating aggregates

You can create aggregates yourself or let Cloud Manager do it for you when it creates volumes. The benefit of creating aggregates yourself is that you can choose the underlying disk size, which enables you to size your aggregate for the capacity or the performance that you need.

Steps

1. On the Canvas page, double-click the name of the Cloud Volumes ONTAP instance on which you want to manage aggregates.
2. Click the menu icon, and then click **Advanced > Advanced allocation**.
3. Click **Add Aggregate** and then specify details for the aggregate.

For help with disk type and disk size, see [Planning your configuration](#).

4. Click **Go**, and then click **Approve and Purchase**.

Connecting a LUN to a host

When you create an iSCSI volume, Cloud Manager automatically creates a LUN for you. We've made it simple by creating just one LUN per volume, so there's no management involved. After you create the volume, use the IQN to connect to the LUN from your hosts.

Note the following:

1. Cloud Manager's automatic capacity management doesn't apply to LUNs. When Cloud Manager creates a LUN, it disables the autogrow feature.
2. You can create additional LUNs from System Manager or the CLI.

Steps

1. On the Canvas page, double-click the Cloud Volumes ONTAP working environment on which you want to manage volumes.
2. Select a volume, and then click **Target IQN**.
3. Click **Copy** to copy the IQN name.
4. Set up an iSCSI connection from the host to the LUN.
 - [ONTAP 9 iSCSI express configuration for Red Hat Enterprise Linux: Starting the iSCSI sessions with the target](#)
 - [ONTAP 9 iSCSI express configuration for Windows: Starting iSCSI sessions with the target](#)

Using FlexCache volumes to accelerate data access

A FlexCache volume is a storage volume that caches NFS read data from an origin (or source) volume. Subsequent reads to the cached data result in faster access to that data.

You can use FlexCache volumes to speed up access to data or to offload traffic from heavily accessed volumes. FlexCache volumes help improve performance, especially when clients need to access the same data repeatedly, because the data can be served directly without having to access the origin volume. FlexCache volumes work well for system workloads that are read-intensive.

Cloud Manager does not provide management of FlexCache volumes at this time, but you can use the ONTAP CLI or ONTAP System Manager to create and manage FlexCache volumes:

- [FlexCache Volumes for Faster Data Access Power Guide](#)
- [Creating FlexCache volumes in System Manager](#)

Starting with the 3.7.2 release, Cloud Manager generates a FlexCache license for all new Cloud Volumes ONTAP systems. The license includes a 500 GiB usage limit.



To generate the license, Cloud Manager needs to access <https://ipa-signer.cloudmanager.netapp.com>. Make sure that this URL is accessible from your firewall.



Managing existing storage

Cloud Manager enables you to manage volumes, aggregates, and CIFS servers. It also prompts you to move volumes to avoid capacity issues.

Managing existing volumes

You can manage existing volumes as your storage needs change. You can view, edit, clone, restore, and delete volumes.

Steps

1. On the Canvas page, double-click the Cloud Volumes ONTAP working environment on which you want to manage volumes.
2. Manage your volumes:

Task	Action
View information about a volume	Select a volume, and then click Info .

Task	Action
Edit a volume (read-write volumes only)	<p>a. Select a volume, and then click Edit.</p> <p>b. Modify the volume's Snapshot policy, NFS protocol version, NFS access control list (export policy), or share permissions, and then click Update.</p> <p> If you need custom Snapshot policies, you can create them by using System Manager.</p>
Clone a volume	<p>a. Select a volume, and then click Clone.</p> <p>b. Modify the clone name as needed, and then click Clone.</p> <p>This process creates a FlexClone volume. A FlexClone volume is a writable, point-in-time copy that is space-efficient because it uses a small amount of space for metadata, and then only consumes additional space as data is changed or added.</p> <p>To learn more about FlexClone volumes, see the ONTAP 9 Logical Storage Management Guide.</p>
Restore data from a Snapshot copy to a new volume	<p>a. Select a volume, and then click Restore from Snapshot copy.</p> <p>b. Select a Snapshot copy, enter a name for the new volume, and then click Restore.</p>
Create a Snapshot copy on demand	<p>a. Select a volume, and then click Create a Snapshot copy.</p> <p>b. Change the name, if needed, and then click Create.</p>
Get the NFS mount command	<p>a. Select a volume, and then click Mount Command.</p> <p>b. Click Copy.</p>
View the target iQN for an iSCSI volume	<p>a. Select a volume, and then click Target iQN.</p> <p>b. Click Copy.</p> <p>c. Use the IQN to connect to the LUN from your hosts.</p>
Change the underlying disk type	<p>a. Select a volume, and then click Change Disk Type & Tiering Policy.</p> <p>b. Select the disk type, and then click Change.</p> <p> Cloud Manager moves the volume to an existing aggregate that uses the selected disk type or it creates a new aggregate for the volume.</p>

Task	Action
Change the tiering policy	<p>a. Select a volume, and then click Change Disk Type & Tiering Policy. b. Click Edit Policy. c. Select a different policy and click Change.</p> <p> Cloud Manager moves the volume to an existing aggregate that uses the selected disk type with tiering, or it creates a new aggregate for the volume.</p>
Delete a volume	<p>a. Select a volume, and then click Delete. b. Click Delete again to confirm.</p>

Resizing a volume

By default, a volume automatically grows to a maximum size when it's out of space. The default value is 1,000, which means the volume can grow to 11 times its size. This value is configurable in a Connector's settings.

If you need to resize your volume, you can do it through [ONTAP System Manager](#). Be sure to take your system's capacity limits into consideration as you resize volumes. Go to the [Cloud Volumes ONTAP Release Notes](#) for more details.

Managing existing aggregates

Manage aggregates yourself by adding disks, viewing information about the aggregates, and by deleting them.

Before you begin

If you want to delete an aggregate, you must have first deleted the volumes in the aggregate.

About this task

If an aggregate is running out of space, you can move volumes to another aggregate by using OnCommand System Manager.

Steps

1. On the Canvas page, double-click the Cloud Volumes ONTAP working environment on which you want to manage aggregates.
2. Click the menu icon and then click **Advanced > Advanced allocation**.
3. Manage your aggregates:

Task	Action
View information about an aggregate	Select an aggregate and click Info .
Create a volume on a specific aggregate	Select an aggregate and click Create volume .

Task	Action
Add disks to an aggregate	<p>a. Select an aggregate and click Add AWS disks or Add Azure disks.</p> <p>b. Select the number of disks that you want to add and click Add.</p>  All disks in an aggregate must be the same size.
Delete an aggregate	<p>a. Select an aggregate that does not contain any volumes and click Delete.</p> <p>b. Click Delete again to confirm.</p>

Modifying the CIFS server

If you change your DNS servers or Active Directory domain, you need to modify the CIFS server in Cloud Volumes ONTAP so that it can continue to serve storage to clients.

Steps

1. From the working environment, click the menu icon and then click **Advanced > CIFS setup**.
2. Specify settings for the CIFS server:

Task	Action
DNS Primary and Secondary IP Address	<p>The IP addresses of the DNS servers that provide name resolution for the CIFS server.</p> <p>The listed DNS servers must contain the service location records (SRV) needed to locate the Active Directory LDAP servers and domain controllers for the domain that the CIFS server will join.</p>
Active Directory Domain to join	The FQDN of the Active Directory (AD) domain that you want the CIFS server to join.
Credentials authorized to join the domain	The name and password of a Windows account with sufficient privileges to add computers to the specified Organizational Unit (OU) within the AD domain.
CIFS server NetBIOS name	A CIFS server name that is unique in the AD domain.
Organizational Unit	<p>The organizational unit within the AD domain to associate with the CIFS server. The default is CN=Computers.</p> <p>If you configure AWS Managed Microsoft AD as the AD server for Cloud Volumes ONTAP, you should enter OU=Computers,OU=corp in this field.</p>
DNS Domain	The DNS domain for the Cloud Volumes ONTAP storage virtual machine (SVM). In most cases, the domain is the same as the AD domain.
NTP Server	Select Use Active Directory Domain to configure an NTP server using the Active Directory DNS. If you need to configure an NTP server using a different address, then you should use the API. See the Cloud Manager automation docs for details.

3. Click **Save**.

Result

Cloud Volumes ONTAP updates the CIFS server with the changes.

Moving a volume

Move volumes for capacity utilization, improved performance, and to satisfy service-level agreements.

You can move a volume in System Manager by selecting a volume and the destination aggregate, starting the volume move operation, and optionally monitoring the volume move job. When using System Manager, a volume move operation finishes automatically.

Steps

1. Use System Manager or the CLI to move the volumes to the aggregate.

In most situations, you can use System Manager to move volumes.

For instructions, see the [ONTAP 9 Volume Move Express Guide](#).

Moving a volume when Cloud Manager displays an Action Required message

Cloud Manager might display an Action Required message that says moving a volume is necessary to avoid capacity issues, but that it cannot provide recommendations to correct the issue. If this happens, you need to identify how to correct the issue and then move one or more volumes.

Steps

1. [Identify how to correct the issue](#).
2. Based on your analysis, move volumes to avoid capacity issues:
 - [Move volumes to another system](#).
 - [Move volumes to another aggregate on the same system](#).

Identifying how to correct capacity issues

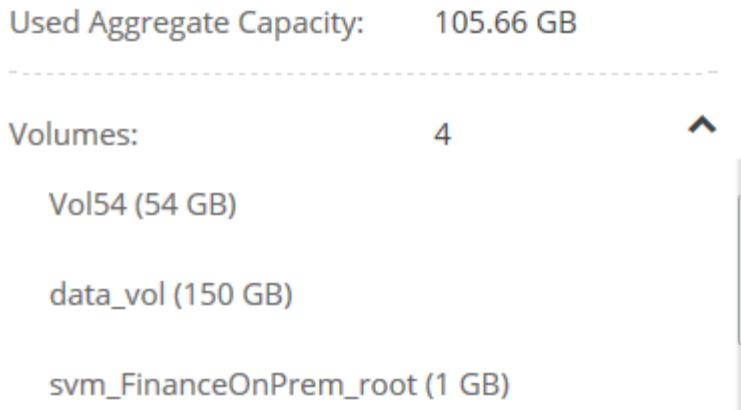
If Cloud Manager cannot provide recommendations for moving a volume to avoid capacity issues, you must identify the volumes that you need to move and whether you should move them to another aggregate on the same system or to another system.

Steps

1. View the advanced information in the Action Required message to identify the aggregate that has reached its capacity limit.

For example, the advanced information should say something similar to the following: Aggregate aggr1 has reached its capacity limit.

2. Identify one or more volumes to move out of the aggregate:
 - a. In the working environment, click the menu icon, and then click **Advanced > Advanced allocation**.
 - b. Select the aggregate, and then click **Info**.
 - c. Expand the list of volumes.



- Review the size of each volume and choose one or more volumes to move out of the aggregate.

You should choose volumes that are large enough to free space in the aggregate so that you avoid additional capacity issues in the future.

- If the system has not reached the disk limit, you should move the volumes to an existing aggregate or a new aggregate on the same system.

For details, see [Moving volumes to another aggregate to avoid capacity issues](#).

- If the system has reached the disk limit, do any of the following:

- Delete any unused volumes.
- Rearrange volumes to free space on an aggregate.

For details, see [Moving volumes to another aggregate to avoid capacity issues](#).

- Move two or more volumes to another system that has space.

For details, see [Moving volumes to another system to avoid capacity issues](#).

Moving volumes to another system to avoid capacity issues

You can move one or more volumes to another Cloud Volumes ONTAP system to avoid capacity issues. You might need to do this if the system reached its disk limit.

About this task

You can follow the steps in this task to correct the following Action Required message:

Moving a volume is necessary to avoid capacity issues; however, Cloud Manager cannot perform this action for you because the system has reached the disk limit.

Steps

- Identify a Cloud Volumes ONTAP system that has available capacity, or deploy a new system.
- Drag and drop the source working environment on the target working environment to perform a one-time data replication of the volume.

For details, see [Replicating data between systems](#).

3. Go to the Replication Status page, and then break the SnapMirror relationship to convert the replicated volume from a data protection volume to a read/write volume.

For details, see [Managing data replication schedules and relationships](#).

4. Configure the volume for data access.

For information about configuring a destination volume for data access, see the [ONTAP 9 Volume Disaster Recovery Express Guide](#).

5. Delete the original volume.

For details, see [Managing existing volumes](#).

Moving volumes to another aggregate to avoid capacity issues

You can move one or more volumes to another aggregate to avoid capacity issues.

About this task

You can follow the steps in this task to correct the following Action Required message:

Moving two or more volumes is necessary to avoid capacity issues; however, Cloud Manager cannot perform this action for you.

Steps

1. Verify whether an existing aggregate has available capacity for the volumes that you need to move:
 - a. In the working environment, click the menu icon, and then click **Advanced > Advanced allocation**.
 - b. Select each aggregate, click **Info**, and then view the available capacity (aggregate capacity minus used aggregate capacity).

aggr1

Aggregate Capacity: 442.94 GB

Used Aggregate Capacity: 105.66 GB

2. If needed, add disks to an existing aggregate:
 - a. Select the aggregate, and then click **Add disks**.
 - b. Select the number of disks to add, and then click **Add**.
3. If no aggregates have available capacity, create a new aggregate.

For details, see [Creating aggregates](#).

4. Use System Manager or the CLI to move the volumes to the aggregate.

5. In most situations, you can use System Manager to move volumes.

For instructions, see the [ONTAP 9 Volume Move Express Guide](#).

Reasons why a volume move might perform slowly

Moving a volume might take longer than you expect if any of the following conditions are true for Cloud Volumes ONTAP:

- The volume is a clone.
- The volume is a parent of a clone.
- The source or destination aggregate has a single Throughput Optimized HDD (st1) disk.
- The Cloud Volumes ONTAP system is in AWS and one aggregate uses an older naming scheme for objects. Both aggregates have to use the same name format.

An older naming scheme is used if data tiering was enabled on an aggregate in the 9.4 release or earlier.

- The encryption settings don't match on the source and destination aggregates, or a rekey is in progress.
- The *-tiering-policy* option was specified on the volume move to change the tiering policy.
- The *-generate-destination-key* option was specified on the volume move.

Tiering inactive data to low-cost object storage

You can reduce storage costs for Cloud Volumes ONTAP by combining an SSD or HDD performance tier for hot data with an object storage capacity tier for inactive data. Data tiering is powered by FabricPool technology. For a high-level overview, see [Data tiering overview](#).

To set up data tiering, you need to do the following:



Choose a supported configuration

Most configurations are supported. If you have a Cloud Volumes ONTAP Standard, Premium, or BYOL system running the most recent version, then you should be good to go. [Learn more](#).



Ensure connectivity between Cloud Volumes ONTAP and object storage

- For AWS, you'll need a VPC Endpoint to S3. [Learn more](#).
- For Azure, you won't need to do anything as long as Cloud Manager has the required permissions. [Learn more](#).
- For GCP, you need to configure the subnet for Private Google Access and set up a service account. [Learn more](#).

3

Ensure that you have an aggregate with tiering enabled

Data tiering must be enabled on an aggregate in order to enable data tiering on a volume. You should be aware of the requirements for new volumes and for existing volumes. [Learn more](#).

4

Choose a tiering policy when creating, modifying, or replicating a volume

Cloud Manager prompts you to choose a tiering policy when you create, modify, or replicate a volume.

- [Tiering data on read-write volumes](#)
- [Tiering data on data protection volumes](#)

What's not required for data tiering?

- You don't need to install a feature license to enable data tiering.
- You don't need to create the capacity tier (an S3 bucket, Azure Blob container, or GCP bucket). Cloud Manager does that for you.
- You don't need to enable data tiering at the system level.

Cloud Manager creates an object store for cold data when the system is created, [as long as there are no connectivity or permissions issues](#). After that, you just need to enable data tiering on volumes (and in some cases, [on aggregates](#)).

Configurations that support data tiering

You can enable data tiering when using specific configurations and features:

- Data tiering is supported with Cloud Volumes ONTAP Standard, Premium, and BYOL, starting with the following versions:
 - Version 9.2 in AWS
 - Version 9.4 in Azure with single node systems
 - Version 9.6 in Azure with HA pairs
 - Version 9.6 in GCP



Data tiering is not supported in Azure with the DS3_v2 virtual machine type.

- In AWS, the performance tier can be General Purpose SSDs (gp3 or gp2) or Provisioned IOPS SSDs (io1).



Tiering data to object storage is not recommended when using Throughput Optimized HDDs (st1).

- In Azure, the performance tier can be Premium SSD managed disks, Standard SSD managed disks, or Standard HDD managed disks.
- In GCP, the performance tier can be either SSD persistent disks, balanced persistent disks, or standard persistent disks.
- Data tiering is supported with encryption technologies.

- Thin provisioning must be enabled on volumes.

Requirements

Depending on your cloud provider, certain connections and permissions must be set up so that Cloud Volumes ONTAP can tier cold data to object storage.

Requirements to tier cold data to AWS S3

Ensure that Cloud Volumes ONTAP has a connection to S3. The best way to provide that connection is by creating a VPC Endpoint to the S3 service. For instructions, see [AWS Documentation: Creating a Gateway Endpoint](#).

When you create the VPC Endpoint, be sure to select the region, VPC, and route table that corresponds to the Cloud Volumes ONTAP instance. You must also modify the security group to add an outbound HTTPS rule that enables traffic to the S3 endpoint. Otherwise, Cloud Volumes ONTAP cannot connect to the S3 service.

If you experience any issues, see [AWS Support Knowledge Center: Why can't I connect to an S3 bucket using a gateway VPC endpoint?](#)

Requirements to tier cold data to Azure Blob storage

You don't need to set up a connection between the performance tier and the capacity tier as long as Cloud Manager has the required permissions. Cloud Manager enables a VNet service endpoint for you if the Cloud Manager policy has these permissions:

```
"Microsoft.Network/virtualNetworks/subnets/write",
"Microsoft.Network/routeTables/join/action",
```

The permissions are included in the latest [Cloud Manager policy](#).

Requirements to tier cold data to a Google Cloud Storage bucket

- The subnet in which Cloud Volumes ONTAP resides must be configured for Private Google Access. For instructions, refer to [Google Cloud Documentation: Configuring Private Google Access](#).
- You need a service account that meets the following requirements:
 - It must have the predefined Storage Admin role.
 - The Connector service account must be a *Service Account User* of this tiering service account.

[Learn how to set up a service account.](#)

- To encrypt the bucket with customer-managed encryption keys, enable the Google Cloud storage bucket to use the key.

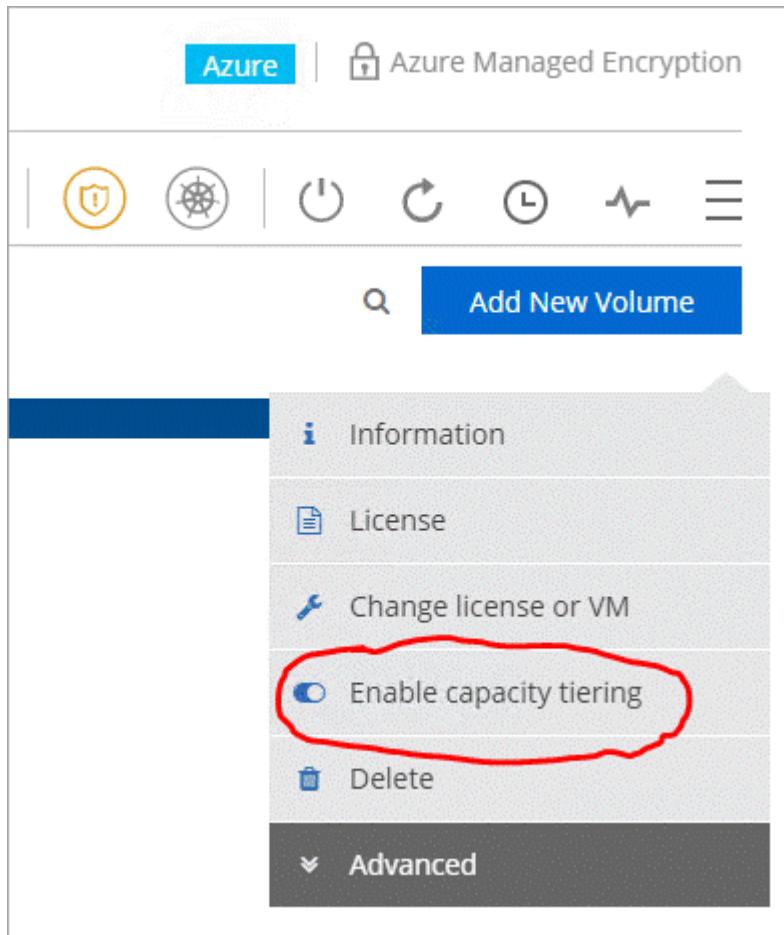
[Learn how to use customer-managed encryption keys with Cloud Volumes ONTAP.](#)

Enabling data tiering after implementing the requirements

Cloud Manager creates an object store for cold data when the system is created, as long as there are no connectivity or permissions issues. If you didn't implement the requirements listed above until after you created the system, then you'll need to manually enable tiering, which creates the object store.

Steps

1. Ensure that you've met all requirements.
2. On the Canvas page, double-click the name of the Cloud Volumes ONTAP instance.
3. Click the menu icon and select **Enable capacity tiering**.



You'll only see this option if data tiering couldn't be enabled when Cloud Manager created the system.

4. Click **Enable** so Cloud Manager can create the object store that this Cloud Volumes ONTAP system will use for tiered data.

Ensuring that tiering is enabled on aggregates

Data tiering must be enabled on an aggregate in order to enable data tiering on a volume. You should be aware of the requirements for new volumes and for existing volumes.

- **New volumes**

If you're enabling data tiering on a new volume, then you don't need to worry about enabling data tiering on an aggregate. Cloud Manager creates the volume on an existing aggregate that has tiering enabled, or it creates a new aggregate for the volume if a data tiering-enabled aggregate doesn't already exist.

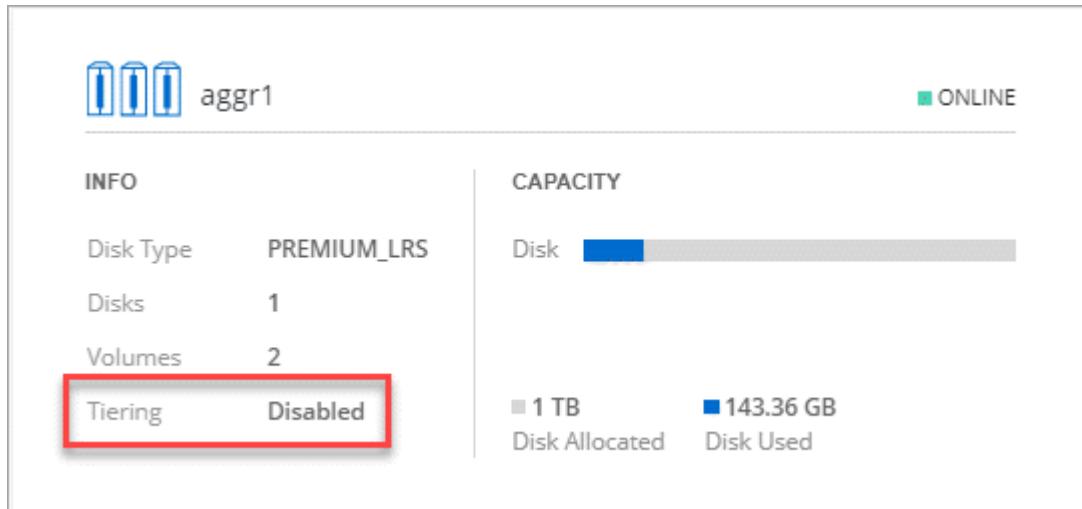
- **Existing volumes**

If you want to enable data tiering on an existing volume, then you'll need to ensure that data tiering is

enabled on the underlying aggregate. If data tiering isn't enabled on the existing aggregate, then you'll need to use System Manager to attach an existing aggregate to the object store.

Steps to confirm whether tiering is enabled on an aggregate

1. Open the working environment in Cloud Manager.
2. Click the menu icon, click **Advanced**, and then click **Advanced allocation**.
3. Verify whether tiering is enabled or disabled on the aggregate.



Steps to enable tiering on an aggregate

1. In System Manager, click **Storage > Tiers**.
2. Click the action menu for the aggregate and select **Attach Cloud Tiers**.
3. Select the cloud tier to attach and click **Save**.

What's next?

You can now enable data tiering on new and existing volumes, as explained in the next section.

Tiering data from read-write volumes

Cloud Volumes ONTAP can tier inactive data on read-write volumes to cost-effective object storage, freeing up the performance tier for hot data.

Steps

1. In the working environment, create a new volume or change the tier of an existing volume:

Task	Action
Create a new volume	Click Add New Volume .
Modify an existing volume	Select the volume and click Change Disk Type & Tiering Policy .

2. Select a tiering policy.

For a description of these policies, see [Data tiering overview](#).

Example

 **Tiering data to object storage**

i **Volume Tiering Policy**

- All - Immediately tiers all data (not including metadata) to object storage.
- Auto - Tiers cold Snapshot copies and cold user data from the active file system to object storage.
- Snapshot Only - Tiers cold Snapshot copies to object storage
- None - Data tiering is disabled.

i **Working Environment S3 Storage classes: Standard**

Cloud Manager creates a new aggregate for the volume if a data tiering-enabled aggregate does not already exist.

Tiering data from data protection volumes

Cloud Volumes ONTAP can tier data from a data protection volume to a capacity tier. If you activate the destination volume, the data gradually moves to the performance tier as it is read.

Steps

1. On the Canvas page, select the working environment that contains the source volume, and then drag it to the working environment to which you want to replicate the volume.
2. Follow the prompts until you reach the tiering page and enable data tiering to object storage.

Example

 **S3 Tiering**

i What are storage tiers?

Enabled Disabled

Note: If you enable S3 tiering, thin provisioning must be enabled on volumes created in this aggregate.

For help with replicating data, see [Replicating data to and from the cloud](#).

Changing the storage class for tiered data

After you deploy Cloud Volumes ONTAP, you can reduce your storage costs by changing the storage class for inactive data that hasn't been accessed for 30 days. The access costs are higher if you do access the data, so you must take that into consideration before you change the storage class.

The storage class for tiered data is system wide—it's not per volume.

For information about supported storage classes, see [Data tiering overview](#).

Steps

1. From the working environment, click the menu icon and then click **Storage Classes** or **Blob Storage Tiering**.
2. Choose a storage class and then click **Save**.

Changing the free space ratio for data tiering

The free space ratio for data tiering defines how much free space is required on Cloud Volumes ONTAP SSDs/HDDs when tiering data to object storage. The default setting is 10% free space, but you can tweak the setting based on your requirements.

For example, you might choose less than 10% free space to ensure that you are utilizing the purchased capacity. Cloud Manager can then purchase additional disks for you when additional capacity is required (up until you reach the disk limit for the aggregate).

The ratio is important for disaster recovery scenarios because as data is read from the object store, Cloud Volumes ONTAP moves the data to SSDs/HDDs to provide better performance. If there isn't sufficient space, then Cloud Volumes ONTAP can't move the data. Take this into consideration when changing the ratio so that you can meet your business requirements.



This setting applies to Cloud Volumes ONTAP only.

Steps

1. In the upper right of the Cloud Manager console, click the **Settings** icon, and select **Connector Settings**.



2. Under **Capacity**, click **Aggregate Capacity Thresholds - Free Space Ratio for Data Tiering**.
3. Change the free space ratio based on your requirements and click **Save**.

Changing the cooling period for the auto tiering policy

If you enabled data tiering on a Cloud Volumes ONTAP volume using the *auto* tiering policy, you can adjust the default cooling period based on your business needs. This action is supported using the API only.

The cooling period is the number of days that user data in a volume must remain inactive before it is considered "cold" and moved to object storage.

The default cooling period for the auto tiering policy is 31 days. You can change the cooling period as follows:

- 9.8 or later: 2 days to 183 days
- 9.7 or earlier: 2 days to 63 days

Step

1. Use the *minimumCoolingDays* parameter with your API request when creating a volume or modifying an existing volume.

Using Cloud Volumes ONTAP as persistent storage for Kubernetes

Cloud Manager can automate the deployment of NetApp Trident on Kubernetes clusters so you can use Cloud Volumes ONTAP as persistent storage for containers.

Trident is a fully-supported open source project maintained by NetApp. Trident integrates natively with Kubernetes and its Persistent Volume framework to seamlessly provision and manage volumes from systems running any combination of NetApp's storage platforms. [Learn more about Trident](#).



The Kubernetes feature isn't supported with on-prem ONTAP clusters. It's supported with Cloud Volumes ONTAP only.

Quick start

Get started quickly by following these steps or scroll down to the remaining sections for full details.



Review prerequisites

Ensure that your environment can meet the prerequisites, which includes connectivity between Kubernetes clusters and Cloud Volumes ONTAP, connectivity between Kubernetes clusters and a Connector, a minimum Kubernetes version of 1.14, at least one worker node in a cluster, and more. [See the complete list](#).



Add your Kubernetes clusters to Cloud Manager

In Cloud Manager, click **K8s** and discover clusters directly from your cloud provider's managed service or import a cluster by providing a kubeconfig file.



Connect your clusters to Cloud Volumes ONTAP

After you add a Kubernetes cluster, click **Connect to Working Environment** to connect the cluster to one or more Cloud Volumes ONTAP systems.



Start provisioning Persistent Volumes

Request and manage Persistent Volumes using native Kubernetes interfaces and constructs. Cloud Manager creates NFS and iSCSI storage classes that you can use when provisioning Persistent Volumes.

[Learn more about provisioning your first volume with Trident for Kubernetes](#).

Reviewing prerequisites

Before you get started, ensure that your Kubernetes clusters and Connector meet specific requirements.

Kubernetes cluster requirements

- Network connectivity is required between a Kubernetes cluster and the Connector and between a Kubernetes cluster and Cloud Volumes ONTAP.

Both the Connector and Cloud Volumes ONTAP need a connection to the Kubernetes API endpoint:

- For managed clusters, set up a route between a cluster's VPC and the VPC where the Connector and Cloud Volumes ONTAP reside.
- For other clusters, the IP address of the primary node or load balancer (as listed in the kubeconfig file) must be reachable by the Connector and Cloud Volumes ONTAP, and it must present a valid TLS certificate.
- A Kubernetes cluster can be in any location that has the network connectivity listed above.
- A Kubernetes cluster must be running version 1.14 at a minimum.

The maximum supported version is defined by Trident. [Click here to see the maximum supported Kubernetes version.](#)

- A Kubernetes cluster must have at least one worker node.
- For clusters running in Amazon Elastic Kubernetes Service (Amazon EKS), each cluster needs an IAM role added in order to resolve a permissions error. After you add the cluster, Cloud Manager will prompt you with the exact eksctl command that resolves the error.

[Learn about IAM permissions boundaries.](#)

- For clusters running in Azure Kubernetes Service (AKS), those clusters must be assigned the *Azure Kubernetes Service RBAC Cluster Admin* role. This is required so Cloud Manager can install Trident and configure storage classes on the cluster.
- For clusters running in Google Kubernetes Engine (GKE), those clusters must not use the default Container Optimized OS. You should switch them to use Ubuntu.

GKE defaults to using the Google [container-optimized image](#), which doesn't have the utilities that Trident needs to mount volumes.

Connector requirements

Ensure that the following permissions are in place for the Connector.

Required permissions to discover and manage EKS clusters

The Connector needs Admin permissions to discover and manage Kubernetes clusters running in Amazon Elastic Kubernetes Service (EKS):

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "VisualEditor0",
            "Effect": "Allow",
            "Action": "eks:*",
            "Resource": "*"
        }
    ]
}
```

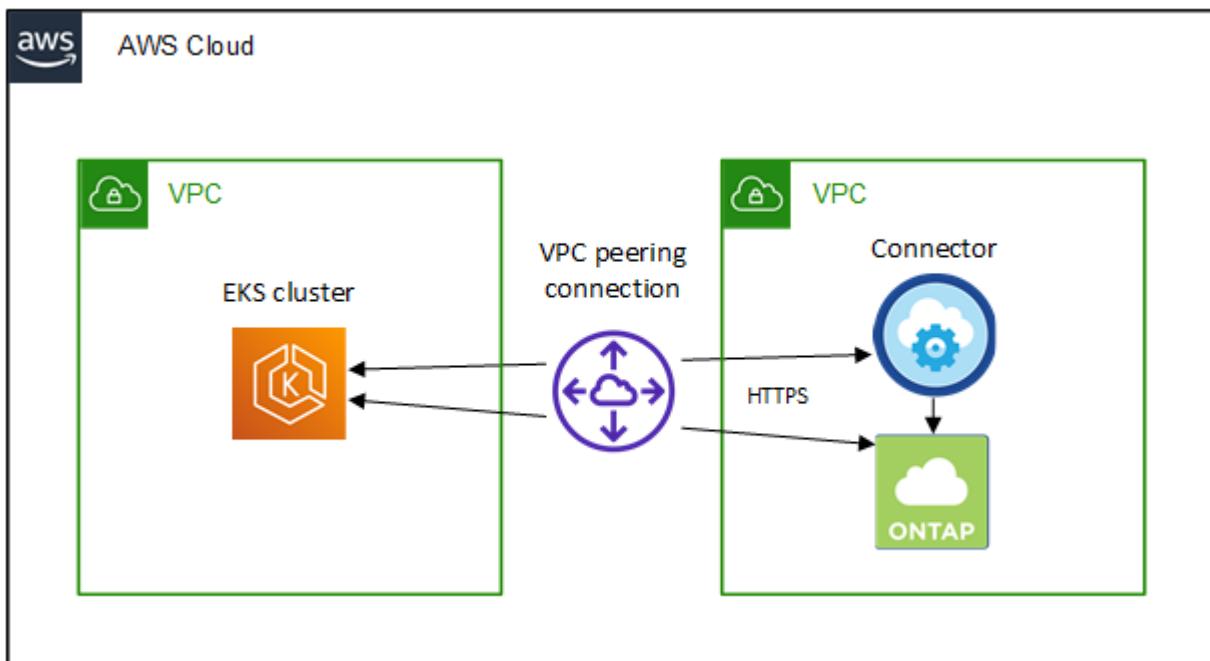
Required permissions to discover and manage GKE clusters

The Connector needs the following permissions to discover and manage Kubernetes clusters running in Google Kubernetes Engine (GKE):

```
container.*
```

Example setup

The following image shows an example of a Kubernetes cluster running in Amazon Elastic Kubernetes Service (Amazon EKS) and its connections to the Connector and Cloud Volumes ONTAP. In this example, VPC peering provides a connection between the VPC for the EKS cluster and the VPC for the Connector and Cloud Volumes ONTAP.



Adding Kubernetes clusters

Add Kubernetes clusters to Cloud Manager by discovering the clusters running in your cloud provider's managed Kubernetes service or by importing a cluster's kubeconfig file.

Steps

1. At the top of Cloud Manager, click **K8s**.
2. Click **Add Cluster**.
3. Choose one of the available options:
 - Click **Discover Clusters** to discover the managed clusters that Cloud Manager has access to based on permissions that you provided to the Connector.

For example, if your Connector is running in Google Cloud, Cloud Manager uses the permissions from the Connector's service account to discover clusters running in Google Kubernetes Engine (GKE).

- Click **Import Cluster** to import a cluster using a kubeconfig file.

After you upload the file, Cloud Manager verifies connectivity to the cluster and saves an encrypted copy of the kubeconfig file.

Result

Cloud Manager adds the Kubernetes cluster. You can now connect the cluster to Cloud Volumes ONTAP.

Connecting a cluster to Cloud Volumes ONTAP

Connect a Kubernetes cluster to Cloud Volumes ONTAP so you can use Cloud Volumes ONTAP as persistent storage for containers.

Steps

1. At the top of Cloud Manager, click **K8s**.
2. Click **Connect to Working Environment** for the cluster that you just added.

Status	Cluster Name	Added By	Version	Working Environments	Date Added
✓	kubernetes	Import	v1.18.0	Connect to Working Environment	August 30, 2020

3. Select a working environment and click **Continue**.
4. Choose the NetApp storage class to use as the default storage class for the Kubernetes cluster and click **Continue**.

When a user creates a persistent volume, the Kubernetes cluster can use this storage class as the backend storage by default.

5. Choose whether to use default auto export policies or whether to add a custom CIDR block.

6. Click **Add Working Environment**.

Result

Cloud Manager connects the working environment to the cluster, which can take up to 15 minutes.

Managing your clusters

Cloud Manager enables you to manage your Kubernetes clusters by changing the default storage class, upgrading Trident, and more.

Changing the default storage class

Make sure that you've set a Cloud Volumes ONTAP storage class as the default storage class so clusters use Cloud Volumes ONTAP as the backend storage.

Steps

1. At the top of Cloud Manager, click **K8s**.
2. Click the name of the Kubernetes cluster.
3. In the **Storage Classes** table, click the actions menu on the far right for the storage class that you'd like to set as the default.

Storage Class ID	Provisioner	Volumes	Labels	
Gp2	aws	0		...
NFS Single Node	NetApp	0		...
NFS High Availability [Default]	NetApp	0		...
iSCSI High Availability	NetApp	0		Set as Default
iSCSI Single Node	NetApp	0		

4. Click **Set as Default**.

Upgrading Trident

You can upgrade Trident from Cloud Manager when a new version is available.

Steps

1. At the top of Cloud Manager, click **K8s**.
2. Click the name of the Kubernetes cluster.
3. If a new version is available, click **Upgrade** next to the Trident version.

The screenshot shows the 'Cluster Details' page for the 'kubernetes-baldwin' cluster. At the top, there are buttons for 'Connect to Working Environment' and 'Remove Cluster'. Below the cluster name, there's a summary table with columns: Status (Running), Cluster Version (1.15), Added by (Discovery), Volumes (2), VPC (vpc-0485a0b201c3a1f2d), Date Added (September 3, 2020), and Trident Version (v20.04). The 'v20.04' entry has a green 'Upgrade' button to its right. A red arrow points to this 'Upgrade' button.

Updating the kubeconfig file

If you added your cluster to Cloud Manager by importing the kubeconfig file, you can upload the latest kubeconfig file to Cloud Manager at any time. You might do this if you've updated the credentials, if you've changed users or roles, or if something changed that affects the cluster, user, namespaces, or authentication.

Steps

1. At the top of Cloud Manager, click **K8s**.
2. Click the name of the Kubernetes cluster.
3. Click **Update Kubeconfig**.
4. When prompted through your web browser, select the updated kubeconfig file and click **Open**.

Result

Cloud Manager updates information about the Kubernetes cluster based on the latest kubeconfig file.

Disconnecting a cluster

When you disconnect a cluster from Cloud Volumes ONTAP, you can no longer use that Cloud Volumes ONTAP system as persistent storage for containers. Existing Persistent Volumes are not deleted.

Steps

1. At the top of Cloud Manager, click **K8s**.
2. Click the name of the Kubernetes cluster.
3. In the **Working Environments** table, click the actions menu on the far right for the working environment that you want to disconnect.

kubernetes

Status: Running | Cluster Version: v1.18.0 | Added by: Import | Volumes: 0 | VPC: - | Date Added: August 30, 2020 | Trident Version: Unknown | Provider: -

1 Working Environments

Name	Provider	Region	Zone	Subnet	Capacity
ishai0ntap4k8	Google Cloud	asia-east1	asia-east1-a	10.140.0.0/20	0.00 used of 10 TB available

Disconnect

4. Click **Disconnect**.

Result

Cloud Manager disconnects the cluster from the Cloud Volumes ONTAP system.

Removing a cluster

Remove decommissioned clusters from Cloud Manager after you disconnect all working environments from the cluster.

Steps

1. At the top of Cloud Manager, click **K8s**.
2. Click the name of the Kubernetes cluster.
3. Click **Remove Cluster**.

k8sdemo

Status: Running | Cluster Version: v1.16.11-gke.5 | Added by: Import | Volumes: 0 | VPC: - | Date Added: September 1, 2020 | Trident is not yet installed. Connect this cluster to a Working Environment. | Provider: Google Cloud

Remove Cluster

Encrypting volumes with NetApp encryption solutions

Cloud Volumes ONTAP supports both NetApp Volume Encryption (NVE) and NetApp Aggregate Encryption (NAE) with an external key manager. NVE and NAE are software-based solutions that enable (FIPS) 140-2-compliant data-at-rest encryption of volumes. [Learn more about these encryption solutions.](#)

Starting with Cloud Volumes ONTAP 9.7, new aggregates will have NAE enabled by default after you set up an

external key manager. New volumes that aren't part of an NAE aggregate will have NVE enabled by default (for example, if you have existing aggregates that were created before setting up an external key manager).

Cloud Volumes ONTAP doesn't support onboard key management.

What you'll need

Your Cloud Volumes ONTAP system should be registered with NetApp support. Starting with Cloud Manager 3.7.1, a NetApp Volume Encryption license is automatically installed on each Cloud Volumes ONTAP system that is registered with NetApp Support.

- [Adding NetApp Support Site accounts to Cloud Manager](#)
- [Registering pay-as-you-go systems](#)



Cloud Manager doesn't install the NVE license on systems that reside in the China region.

Steps

1. Review the list of supported key managers in the [NetApp Interoperability Matrix Tool](#).



Search for the **Key Managers** solution.

2. [Connect to the Cloud Volumes ONTAP CLI](#).

3. Install SSL certificates and connect to the external key management servers.

[ONTAP 9 NetApp Encryption Power Guide: Configuring external key management](#)

Manage storage VMs

Manage storage VMs in Cloud Manager

A storage VM is a virtual machine running within ONTAP that provides storage and data services to your clients. You might know this as an *SVM* or a *vserver*. Cloud Volumes ONTAP is configured with one storage VM by default, but some configurations support additional storage VMs.

Supported number of storage VMs

Multiple storage VMs are supported with Cloud Volumes ONTAP in AWS and in Azure with certain configurations. Go to the [Cloud Volumes ONTAP Release Notes](#) to verify the supported number of storage VMs for your version of Cloud Volumes ONTAP.

Work with multiple storage VMs

Cloud Manager supports any additional storage VMs that you create from System Manager or the CLI.

For example, the following image shows how you can choose a storage VM when you create a volume.

Details & Protection

Storage VM Name

i

Volume Name

i

Size (GiB)

i

Snapshot Policy

v

i Default Policy

And the following image shows how you can choose a storage VM when replicating a volume to another system.

Destination Volume Name

Destination Storage VM Name

v

Destination Aggregate

v

Modify the name of the default storage VM

Cloud Manager automatically names the single storage VM that it creates for Cloud Volumes ONTAP. You can modify the name of the storage VM if you have strict naming standards. For example, you might want the name to match how you name the storage VMs for your ONTAP clusters.

If you created any additional storage VMs for Cloud Volumes ONTAP, then you can't rename the storage VMs from Cloud Manager. You'll need to do so directly from Cloud Volumes ONTAP by using System Manager or

the CLI.

Steps

1. From the working environment, click the menu icon, and then click **Information**.
2. Click the edit icon to the right of the storage VM name.

The screenshot shows a 'Working Environment Information' page for an ONTAP system. It displays the following details:

ONTAP	
Serial Number:	[REDACTED]
System ID:	system-id-capacitytest
Cluster Name:	capacitytest
ONTAP Version:	9.7RC1
Date Created:	Jul 6, 2020 07:42:02 am
Storage VM Name:	svm_capacitytest 

3. In the Modify SVM Name dialog box, change the name, and then click **Save**.

Manage storage VMs for disaster recovery

Cloud Manager doesn't provide any setup or orchestration support for storage VM disaster recovery. You must use System Manager or the CLI.

- [SVM Disaster Recovery Preparation Express Guide](#)
- [SVM Disaster Recovery Express Guide](#)

Create data-serving storage VMs for Cloud Volumes ONTAP in AWS

A storage VM is a virtual machine running within ONTAP that provides storage and data services to your clients. You might know this as an *SVM* or a *vserver*. Cloud Volumes ONTAP is configured with one storage VM by default, but some configurations support additional storage VMs.

To create additional data-serving storage VMs, you need to allocate IP addresses in AWS and then run ONTAP commands based on your Cloud Volumes ONTAP configuration.

Supported number of storage VMs

Multiple storage VMs are supported with Cloud Volumes ONTAP BYOL in AWS with an add-on license, starting with the 9.7 release. Go to the [Cloud Volumes ONTAP Release Notes](#) to verify the supported number of storage VMs for your version of Cloud Volumes ONTAP.

All other Cloud Volumes ONTAP configurations support one data-serving storage VM and one destination storage VM used for disaster recovery. You can activate the destination storage VM for data access if there's an outage on the source storage VM.

Verify limits for your configuration

Each EC2 instance supports a maximum number of private IPv4 addresses per network interface. You need to verify the limit before you allocate IP addresses in AWS for the new storage VM.

Steps

1. Go the [Storage limits section in the Cloud Volumes ONTAP Release Notes](#).
2. Identify the maximum number of IP addresses per interface for your instance type.
3. Make note of this number because you'll need it in the next section when you allocate IP addresses in AWS.

Allocate IP addresses in AWS

Private IPv4 addresses must be assigned to port e0a in AWS before you create LIFs for the new storage VM.

Note that an optional management LIF for a storage VM requires a private IP address on a single node system and on an HA pair in a single AZ. This management LIF provides a connection to management tools like SnapCenter.

Steps

1. Log in to AWS and open the EC2 service.
2. Select the Cloud Volumes ONTAP instance and click **Networking**.
If you're creating a storage VM on an HA pair, select node 1.
3. Scroll down to **Network interfaces** and click the **Interface ID** for port e0a.

Name	Insta...	Instance state	Instance type	Status check
danielleAws	i-070...	Running	m5.2xlarge	2/2 check
occmTiering0702	i-0a7...	Stopped	m5.2xlarge	-
<input checked="" type="checkbox"/> cvoTiering1	i-02a...	Stopped	m5.2xlarge	-

Interface ID	Description
eni-07c301...	Interface for Node & Cluster Management, Inter-Cluster Communication, and Data - e0a

4. Select the network interface and click **Actions > Manage IP addresses**.
5. Expand the list of IP addresses for e0a.
6. Verify the IP addresses:
 - a. Count the number of allocated IP addresses to confirm that the port has room for additional IPs.

You should have identified the maximum number of supported IP addresses per interface in the previous section of this page.
 - b. Optional: Go to the CLI for Cloud Volumes ONTAP and run **network interface show** to confirm that each of these IP addresses are in use.

If an IP address isn't in use, then you can use it with the new storage VM.
7. Back in the AWS Console, click **Assign new IP address** to assign additional IP addresses based on the amount that you need for the new storage VM.
 - Single node system: One unused secondary private IP is required.

An optional secondary private IP is required if you want to create a management LIF on the storage VM.
 - HA pair in a single AZ: One unused secondary private IP is required on node 1.

An optional secondary private IP is required if you want to create a management LIF on the storage VM.
 - HA pair in multiple AZs: One unused secondary private IP is required on each node.
8. If you're allocating the IP address on an HA pair in a single AZ, enable **Allow secondary private IPv4 addresses to be reassigned**.
9. Click **Save**.
10. If you have an HA pair in multiple AZs, then you'll need to repeat these steps for node 2.

Create a storage VM on a single node system

These steps create a new storage VM on a single node system. One private IP address is required to create a NAS LIF and another optional private IP address is needed if you want to create a management LIF.

Steps

1. Create the storage VM and a route to the storage VM.

```
vserver create -rootvolume-security-style unix -rootvolume root_svm_2  
-snapshot-policy default -vserver svm_2 -aggregate aggr1
```

```
network route create -destination 0.0.0.0/0 -vserver svm_2 -gateway  
subnet_gateway
```

2. Create a NAS LIF.

```
network interface create -auto-revert true -vserver svm_2 -service  
-policy default-data-files -home-port e0a -address private_ip_x -netmask  
node1Mask -lif ip_nas_2 -home-node cvo-node
```

Where *private_ip_x* is an unused secondary private IP on e0a.

3. Optional: Create a storage VM management LIF.

```
network interface create -auto-revert true -vserver svm_2 -service  
-policy default-management -home-port e0a -address private_ip_y -netmask  
node1Mask -lif ip_svm_mgmt_2 -home-node cvo-node
```

Where *private_ip_y* is another unused secondary private IP on e0a.

Create a storage VM on an HA pair in a single AZ

These steps create a new storage VM on an HA pair in a single AZ. One private IP address is required to create a NAS LIF and another optional private IP address is needed if you want to create a management LIF.

Both of these LIFs get allocated on node 1. The private IP addresses can move between nodes if failures occur.

Steps

1. Create the storage VM and a route to the storage VM.

```
vserver create -rootvolume-security-style unix -rootvolume root_svm_2  
-snapshot-policy default -vserver svm_2 -aggregate aggr1
```

```
network route create -destination 0.0.0.0/0 -vserver svm_2 -gateway  
subnet_gateway
```

2. Create a NAS LIF on node 1.

```
network interface create -auto-revert true -vserver svm_2 -service  
-policy default-data-files -home-port e0a -address private_ip_x -netmask  
node1Mask -lif ip_nas_2 -home-node cvo-node1
```

Where *private_ip_x* is an unused secondary private IP on e0a of cvo-node1. This IP address can be relocated to the e0a of cvo-node2 in case of takeover because the service policy default-data-files indicates that IPs can migrate to the partner node.

3. Optional: Create a storage VM management LIF on node 1.

```
network interface create -auto-revert true -vserver svm_2 -service  
-policy default-management -home-port e0a -address private_ip_y -netmask  
node1Mask -lif ip_svm_mgmt_2 -home-node cvo-node1
```

Where *private_ip_y* is another unused secondary private IP on e0a.

Create a storage VM on an HA pair in multiple AZs

These steps create a new storage VM on an HA pair in multiple AZs.

A *floating* IP address is required for a NAS LIF and is optional for a management LIF. These floating IP addresses don't require you to allocate private IPs in AWS. Instead, the floating IPs are automatically configured in the AWS route table to point to a specific node's ENI in the same VPC.

In order for floating IPs to work with ONTAP, a private IP address must be configured on every storage VM on each node. This is reflected in the steps below where an iSCSI LIF is created on node 1 and on node 2.

Steps

1. Create the storage VM and a route to the storage VM.

```
vserver create -rootvolume-security-style unix -rootvolume root_svm_2  
-snapshot-policy default -vserver svm_2 -aggregate aggr1
```

```
network route create -destination 0.0.0.0/0 -vserver svm_2 -gateway  
subnet_gateway
```

2. Create a NAS LIF on node 1.

```
network interface create -auto-revert true -vserver svm_2 -service  
-policy default-data-files -home-port e0a -address floating_ip -netmask  
node1Mask -lif ip_nas_floating_2 -home-node cvo-node1
```

- The floating IP address must be outside of the CIDR blocks for all VPCs in the AWS region in which you deploy the HA configuration. 192.168.209.27 is an example floating IP address. [Learn more about choosing a floating IP address](#).
- `-service-policy default-data-files` indicates that IPs can migrate to the partner node.

3. Optional: Create a storage VM management LIF on node 1.

```
network interface create -auto-revert true -vserver svm_2 -service  
-policy default-management -home-port e0a -address floating_ip -netmask  
node1Mask -lif ip_svm_mgmt_2 -home-node cvo-node1
```

4. Create an iSCSI LIF on node 1.

```
network interface create -vserver svm_2 -service-policy default-data-blocks -home-port e0a -address private_ip -netmask node1Mask -lif ip_node1_iscsi_2 -home-node cvo-node1
```

- This iSCSI LIF is required to support LIF migration of the floating IPs in the storage VM. It doesn't have to be an iSCSI LIF, but it can't be configured to migrate between nodes.
- `-service-policy default-data-block` indicates that an IP address does not migrate between nodes.
- `private_ip` is an unused secondary private IP address on eth0 (e0a) of cvo_node1.

5. Create an iSCSI LIF on node 2.

```
network interface create -vserver svm_2 -service-policy default-data-blocks -home-port e0a -address private_ip -netmaskNode2Mask -lif ip_node2_iscsi_2 -home-node cvo-node2
```

- This iSCSI LIF is required to support LIF migration of the floating IPs in the storage VM. It doesn't have to be an iSCSI LIF, but it can't be configured to migrate between nodes.
- `-service-policy default-data-block` indicates that an IP address does not migrate between nodes.
- `private_ip` is an unused secondary private IP address on eth0 (e0a) of cvo_node2.

Create data-serving storage VMs for Cloud Volumes ONTAP in Azure

A storage VM is a virtual machine running within ONTAP that provides storage and data services to your clients. You might know this as an *SVM* or a *vserver*. Cloud Volumes ONTAP is configured with one storage VM by default, but additional storage VMs are supported when running Cloud Volumes ONTAP in Azure.

To create additional data-serving storage VMs, you need to allocate IP addresses in Azure and then run ONTAP commands to create the storage VM and data LIFs.

Supported number of storage VMs

Multiple storage VMs are supported with Cloud Volumes ONTAP BYOL in Azure with an add-on license starting with the 9.9.0 release. Go to the [Cloud Volumes ONTAP Release Notes](#) to verify the supported number of storage VMs for your version of Cloud Volumes ONTAP.

All other Cloud Volumes ONTAP configurations support one data-serving storage VM and one destination storage VM used for disaster recovery. You can activate the destination storage VM for data access if there's an outage on the source storage VM.

Allocate IP addresses in Azure

Follow the steps below for your configuration: either a single node system, an HA pair using iSCSI, or an HA pair using NFS/SMB.

Single node

IP addresses must be assigned to nic0 in Azure before you create a storage VM and allocate LIFs.

You'll need to create an IP address for data LIF access and another optional IP address for a storage VM (SVM) management LIF. This management LIF provides a connection to management tools like SnapCenter.

Steps

1. Log in to the Azure portal and open the **Virtual machine** service.
2. Click the name of the Cloud Volumes ONTAP VM.
3. Click **Networking**.
4. Click the name of the network interface for nic0.
5. Under **Settings**, click **IP configurations**.
6. Click **Add**.
7. Enter a name for the IP configuration, select **Dynamic**, and then click **OK**.
8. Click the name of the IP configuration that you just created, change the **Assignment** to **Static**, and click **Save**.
9. If you want to create an SVM management LIF, repeat these steps on node 1.

After you finish

Copy the private IP addresses that you just created. You'll need to specify those IP addresses when you create LIFs for the new storage VM.

HA pair using iSCSI

iSCSI IP addresses must be assigned to nic0 in Azure before you create a storage VM and allocate LIFs. IPs for iSCSI are assigned to nic0 and not the load balancer because iSCSI uses ALUA for failover.

You'll need to create an IP address for data LIF access from node 1, another IP address for data LIF access from node 2, and another optional IP address for a storage VM (SVM) management LIF. This management LIF provides a connection to management tools like SnapCenter.

Steps

1. Log in to the Azure portal and open the **Virtual machine** service.
2. Click the name of the Cloud Volumes ONTAP VM for node 1.
3. Click **Networking**.
4. Click the name of the network interface for nic0.
5. Under **Settings**, click **IP configurations**.
6. Click **Add**.
7. Enter a name for the IP configuration, select **Dynamic**, and then click **OK**.
8. Click the name of the IP configuration that you just created, change the **Assignment** to **Static**, and click **Save**.
9. Repeat these steps on node 2.
10. If you want to create an SVM management LIF, repeat these steps on node 1.

After you finish

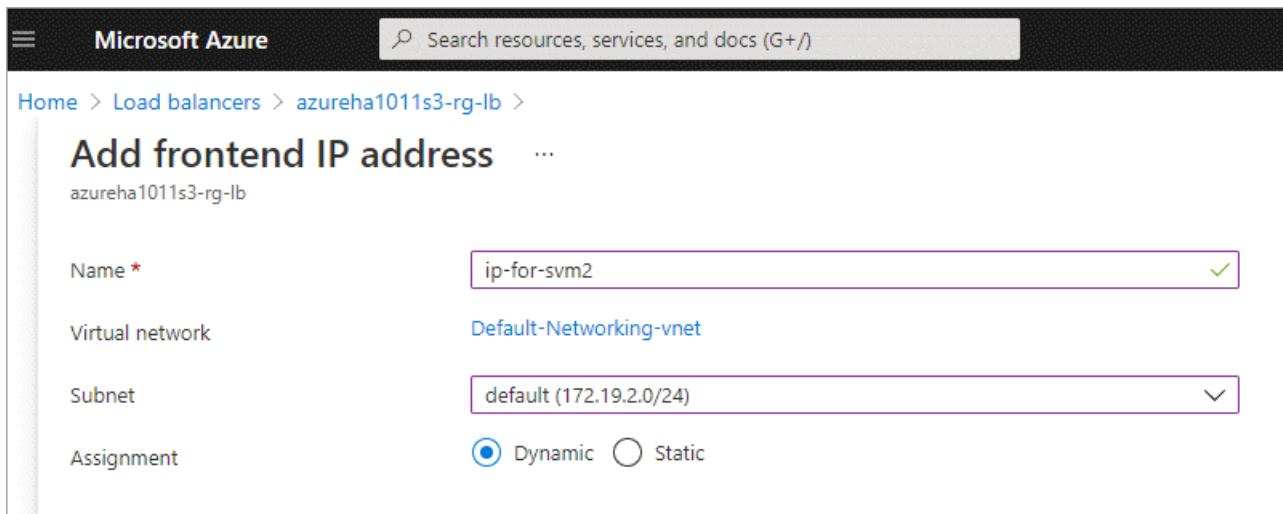
Copy the private IP addresses that you just created. You'll need to specify those IP addresses when you create LIFs for the new storage VM.

HA pair using NFS/SMB

IP addresses that you use for NFS and SMB data are allocated in the load balancer so that the IP addresses can migrate to the other node in case failover events occur.

Steps

1. In the Azure portal, open the **Load balancers** service.
2. Click the name of the load balancer for the HA pair.
3. Create one frontend IP configuration for data LIF access from node 1, another for data LIF access from node 2 (HA pairs only), and another optional frontend IP for a storage VM (SVM) management LIF.
 - a. Under **Settings**, click **Frontend IP configuration**.
 - b. Click **Add**.
 - c. Enter a name for the frontend IP, select the subnet for the Cloud Volumes ONTAP HA pair, and leave **Dynamic** selected.



The screenshot shows the Azure portal interface for adding a frontend IP address. The page title is "Add frontend IP address". The "Name" field contains "ip-for-svm2". The "Virtual network" is set to "Default-Networking-vnet". The "Subnet" dropdown shows "default (172.19.2.0/24)". The "Assignment" section has two radio buttons: "Dynamic" (selected) and "Static". A green checkmark is visible next to the "Dynamic" button, indicating it is the active selection.

- d. Click the name of the frontend IP configuration that you just created, change the **Assignment** to **Static**, and click **Save**.
4. Add a health probe for each frontend IP that you just created.
 - a. Under the load balancer's **Settings**, click **Health probes**.
 - b. Click **Add**.
 - c. Enter a name for the health probe and enter a port number that's between 63005 and 65000. Keep the default values for the other fields.

It's important that the port number is between 63005 and 65000. For example, if you are creating three health probes, you could enter probes that use the port numbers 63005, 63006, and 63007.

Microsoft Azure

Search resources, services, and

Home > Load balancers > azureha1011s3-rg-lb >

Add health probe

azureha1011s3-rg-lb

Name *	svm2-health-probe1	✓
Protocol *	TCP	▼
Port *	63005	✓
Interval *	5	seconds
Unhealthy threshold *	2	consecutive failures
Used by	Not used	

5. Create new load balancing rules for each frontend IP.

- Under the load balancer's **Settings**, click **Load balancing rules**.
- Click **Add** and enter the required information:
 - Name:** Enter a name for the rule.
 - IP Version:** Select **IPv4**.
 - Frontend IP address:** Select one of the frontend IP addresses that you just created.
 - HA Ports:** Enable this option.
 - Backend pool:** Keep the default Backend pool that was already selected.
 - Health probe:** Select the health probe that you created for the selected frontend IP.
 - Session persistence:** Select **None**.
 - Floating IP:** Select **Enabled**.

Add load balancing rule

chandanaTcpRst3-rg-lb

ⓘ A load balancing rule distributes incoming traffic that is sent to a selected IP address and port combination across a group of backend pool instances. Only backend instances that the health probe considers healthy receive new traffic.

Name *

jimmy_new_rule



IP Version *

IPv4 IPv6

Frontend IP address *

10.1.0.156 (dataAFIP)



HA Ports



Backend pool

backendPool (2 virtual machines)



Health probe

dataAProbe (TCP:63002)



Session persistence

None



Floating IP

After you finish

Ensure that the network security group rules for Cloud Volumes ONTAP allows the load balancer to send TCP probes for the health probes that were created in step 4 above. Note that this is allowed by default.

Create a storage VM and LIFs

These steps create a new storage VM on a single node system or on an HA pair. One IP address is required for data LIF access from node 1, another IP address for data LIF access from node 2 (HA pairs only), and another optional IP address for a storage VM (SVM) management LIF. This management LIF provides a connection to management tools like SnapCenter.

Use the commands below that match the data access protocol for the storage VM, which is either NAS or iSCSI.

Steps

1. Create the storage VM and a route to the storage VM.

```
vserver create -vserver <svm-name> -subtype default -rootvolume <root-volume-name> -rootvolume-security-style unix
```

```
network route create -destination 0.0.0.0/0 -vserver <svm-name> -gateway <ip-of-gateway-server>
```

2. Create data LIFs:

- Use the following command to create a NAS LIF on node 1.

```
network interface create -vserver <svm-name> -lif <lif-name> -role data -data-protocol cifs,nfs -address <nfs-ip-address> -netmask <length> -length <length> -home-node <name-of-node1> -status-admin up -failover-policy system-defined -firewall-policy data -home-port e0a -auto-revert true -failover-group Default -probe-port <port-number-for-azure-health-probe1>
```

If this is a single node system, then you should change the value of the `-failover-policy` parameter to *disabled*.

- Use the following command to create a NAS LIF on node 2 (for HA pairs only).

```
network interface create -vserver <svm-name> -lif <lif-name> -role data -data-protocol cifs,nfs -address <nfs-cifs-ip-address> -netmask <length> -length <length> -home-node <name-of-node2> -status-admin up -failover-policy system-defined -firewall-policy data -home-port e0a -auto-revert true -failover-group Default -probe-port <port-number-for-azure-health-probe2>
```

- Use the following command to create an iSCSI LIF on node 1.

```
network interface create -vserver <svm-name> -home-port e0a -address <iscsi-ip-address> -lif <lif-name> -home-node <name-of-node1> -data-protocol iscsi
```

- Use the following command to create an iSCSI LIF on node 2 (for HA pairs only).

```
network interface create -vserver <svm-name> -home-port e0a -address <iscsi-ip-address> -lif <lif-name> -home-node <name-of-node2> -data-protocol iscsi
```

3. Optional: Create a storage VM management LIF on node 1.

```
network interface create -vserver <svm-name> -lif <lif-name> -role data  
-data-protocol none -address <svm-mgmt-ip-address> -netmask-length  
<length> -home-node node1 -status-admin up -failover-policy system-  
defined -firewall-policy mgmt -home-port e0a -auto-revert false  
-failover-group Default -probe-port <port-number-for-azure-health-  
probe3>
```

Monitor performance

Learn about the Monitoring service

The Monitoring service gives you insights into the health and performance of your Cloud Volumes ONTAP instances and helps you troubleshoot and optimize the performance of your cloud storage environment.

Features

- Automatically monitor all volumes
- View volume performance data in terms of IOPS, throughput, and latency
- Identify performance issues to minimize impact on your users and apps

Supported cloud providers

The Monitoring service is supported with Cloud Volumes ONTAP for AWS and Cloud Volumes ONTAP for Azure.

Cost

NetApp doesn't charge you for using the Monitoring service, but Cloud Manager launches a virtual machine in your VPC to facilitate monitoring. This VM results in charges from your cloud provider.

How the Monitoring service works

Cloud Manager leverages [NetApp's Cloud Insights service](#) to provide monitoring.

At a high-level, Cloud Insights integration with Cloud Manager works like this:

1. You enable the Monitoring service on Cloud Volumes ONTAP.
2. Cloud Manager configures your environment. It does the following:
 - a. Creates a Cloud Insights tenant (also called *environment*) and associates all users in your Cloud Central account to the tenant.

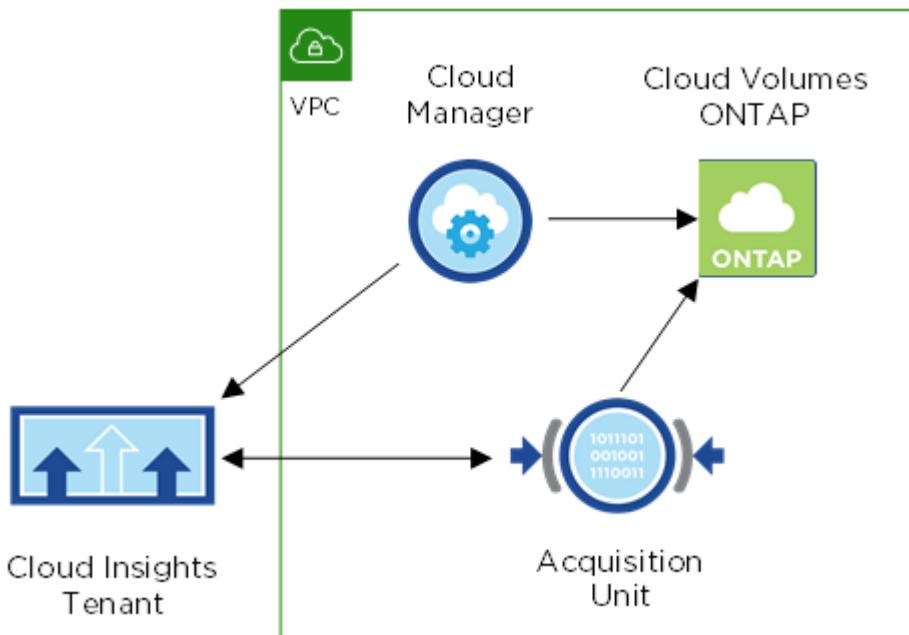
If you already have a tenant, then Cloud Manager uses that existing tenant.

- b. Enables a free trial of Cloud Insights.

On the 29th day, your plan automatically transitions from the Trial Version to the [Basic Edition](#).

- c. Deploys a virtual machine in your VPC/VNet called an Acquisition Unit. The Acquisition Unit facilitates monitoring of volumes (this is the VM mentioned in the Cost section above).
 - d. Connects the Acquisition Unit to Cloud Volumes ONTAP and to the Cloud Insights tenant.
3. In Cloud Manager, you click Monitoring and use the performance data to troubleshoot and optimize performance.

The following image shows the relationship between these components in an AWS VPC:



The Acquisition Unit

When you enable Monitoring, Cloud Manager deploys an Acquisition Unit in the same subnet as the Connector.

An *Acquisition Unit* collects performance data from Cloud Volumes ONTAP and sends it to the Cloud Insights tenant. Cloud Manager then queries that data and presents it to you.

Note the following about the Acquisition Unit instance:

- In AWS, the Acquisition Unit runs on a t3.xlarge instance with a 100 GiB GP2 volume.
- In Azure, the Acquisition Unit runs on a D4_v3 virtual machine with a 30 GiB Standard SSD.
- The instance is named *CloudInsights-AU* with a generated hash (UUID) concatenated to it. For example: *CloudInsights-AU-FAN7FqeH*
- Only one Acquisition Unit is deployed per Connector.
- The instance must be running to access performance information in the Monitoring tab.

Cloud Insights tenant

If you don't already have a *tenant*, Cloud Manager sets one up for you when you enable Monitoring. A Cloud Insights tenant enables you to access the performance data that the Acquisition Unit collects. The tenant is a secure data partition within the NetApp Cloud Insights service.

Cloud Insights web interface

The Monitoring tab in Cloud Manager provides basic performance data for your volumes. You can go to the Cloud Insights web interface from your browser to perform more in-depth monitoring and to configure alerts for your Cloud Volumes ONTAP systems.

Free trial and subscription

Cloud Manager enables a free trial of Cloud Insights to provide performance data within Cloud Manager and for you to explore the features that Cloud Insights Standard Edition has to offer.

On the 29th day, your plan automatically transitions from the Trial Version to the [Basic Edition](#).

You have the option to subscribe to the Standard or Premium editions to gain access to additional features from the Cloud Insights web interface.

[Learn how to subscribe to Cloud Insights.](#)

Monitoring Cloud Volumes ONTAP

Complete a few steps to start monitoring Cloud Volumes ONTAP performance.

Quick start

Get started quickly by following these steps or scroll down to the remaining sections for full details.



Verify support for your configuration

- You need a Cloud Volumes ONTAP system running in AWS or Azure.
- For AWS, you need a Connector running version 3.8.4 or later.
- For Azure, you need a Connector running version 3.9.3 or later.



Enable Monitoring on your new or existing system

- New working environments: Be sure to keep Monitoring enabled when you create the working environment (it's enabled by default).
- Existing working environments: Select a working environment and click **Start Monitoring**.



View performance data

Click **Monitoring** and view performance data for your volumes.

Requirements

Read the following requirements to make sure that you have a supported configuration.

Supported Cloud Volumes ONTAP versions

Any version of Cloud Volumes ONTAP in AWS or in Azure.

Supported Connector

- For AWS, you need a Connector running version 3.8.4 or later.
- For Azure, you need a Connector running version 3.9.3 or later.



You can view a Connector's version by clicking the icon and then **Support > Connector**.

Email address for Cloud Central

The email address for your Cloud Central user account should be your business email address. Free email domains like gmail and hotmail aren't supported when creating a Cloud Insights tenant.

Networking for the Acquisition Unit

The Acquisition Unit uses 2-way/mutual authentication to connect to the Cloud Insights server. The client certificate must be passed to the Cloud Insights server to be authenticated. To accomplish this, the proxy must be set up to forward the http request to the Cloud Insights server without decrypting the data.

The Acquisition Unit uses the following two endpoints to communicate with Cloud Insights. If you have a firewall between the Acquisition Unit server and Cloud Insights, you need these endpoints when configuring firewall rules:

```
https://aulogin.<Cloud Insights Domain>
https://<your-tenant-ID>.<Cloud Insights Domain>
```

For example:

```
https://aulogin.c01.cloudinsights.netapp.com
https://cg0c586a-ee05-45rb-a5ac-
333b5ae7718d7.c01.cloudinsights.netapp.com
```

Contact us through the in-product chat if you need help identifying your Cloud Insights domain and tenant ID.

Networking for the Connector

Similar to the Acquisition Unit, the Connector must have outbound connectivity to the Cloud Insights tenant. But the endpoint that the Connector contacts is slightly different. It contacts the tenant host URL using the shortened tenant ID:

```
https://<your-short-tenant-ID>.<Cloud Insights Domain>
```

For example:

```
https://abcd12345.c01.cloudinsights.netapp.com
```

Again, you can contact us through the in-product chat if you need help identifying the tenant host URL.

Enabling monitoring on a new system

The Monitoring service is enabled by default in the working environment wizard. Be sure to keep the option enabled.

Steps

1. Click **Create Cloud Volumes ONTAP**.
2. Select Amazon Web Services or Microsoft Azure as the cloud provider and then choose a single node or HA system.
3. Fill out the Details & Credentials page.
4. On the Services page, leave the Monitoring service enabled and click **Continue**.

The screenshot shows a configuration interface for the NetApp Monitoring service. At the top, there's a header with a monitoring icon and the word "Monitoring". To the right is a blue toggle switch that is turned on. Below the header, a descriptive text block explains what NetApp Monitoring does. Underneath this, there are two sections: "ADVANTAGES" and "CLARIFICATIONS", each containing a bulleted list of points.

ADVANTAGES	CLARIFICATIONS
<ul style="list-style-type: none">➤ Automatically monitor all volumes - no configuration is required➤ Prevent performance issues from impacting your users and apps	<ul style="list-style-type: none">➤ Activation is free, but requires deploying a small-size cloud instance which will incur charges by your cloud provider➤ Monitoring can be disabled at any time

Enabling monitoring on an existing system

Enable monitoring at any time from the working environment.

Steps

1. At the top of Cloud Manager, click **Canvas**.
2. Select a working environment.
3. In the pane on the right, click **Start Monitoring**.

The screenshot shows the Cloud Manager interface for environment CVO2, which is currently set to AWS. The dashboard displays four service sections: Cloud Compliance (Off), Backup to Cloud (On, 1 volume backed up), Kubernetes (Off), and Monitoring (Off). The Monitoring section includes a prominent 'Start Monitoring' button with a hand cursor icon.

Service	Status	Details
Cloud Compliance	Off	Enable Compliance
Backup to Cloud	On	1 Volume Backed Up
Kubernetes	Off	Activate Kubernetes
Monitoring	Off	Start Monitoring

Monitoring your volumes

Monitor performance by viewing IOPS, throughput, and latency for each of your volumes.

Steps

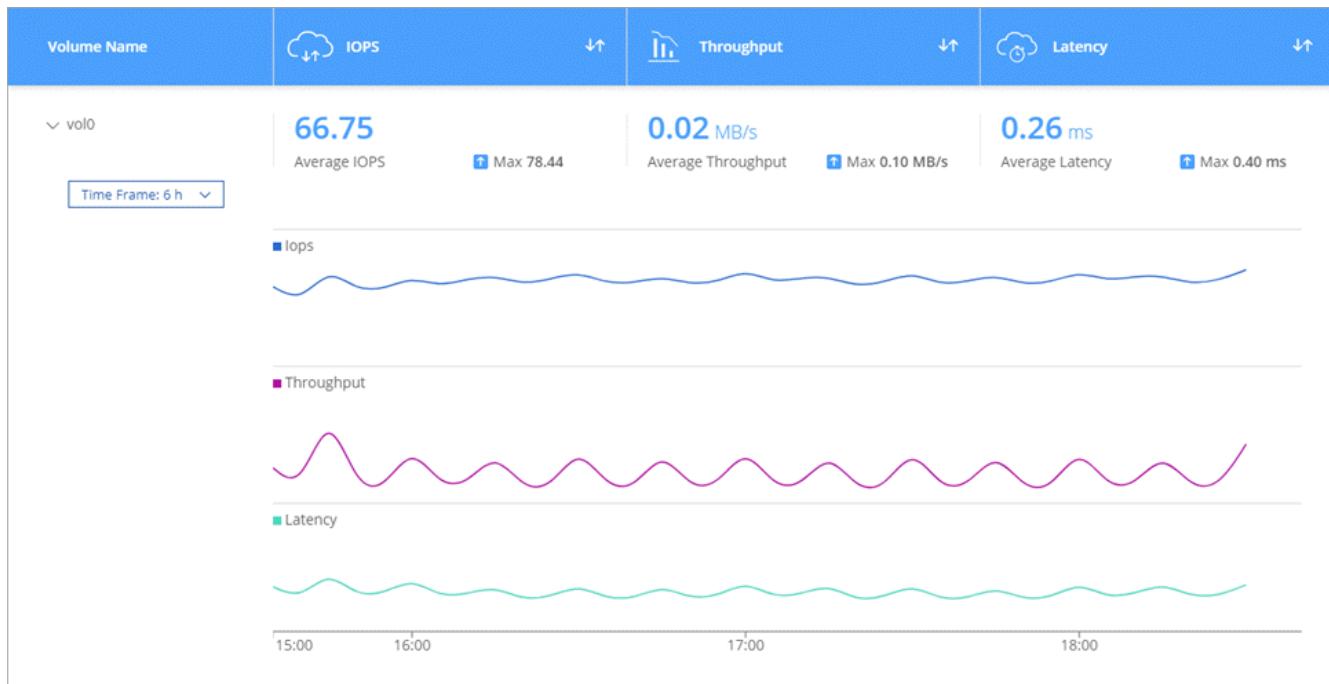
1. At the top of Cloud Manager, click **Monitoring**.
2. Filter the contents of the dashboard to get the information that you need.
 - Select a specific working environment.
 - Select a different timeframe.
 - Select a specific SVM.
 - Search for a specific volume.

The following image highlights each of these options:

The screenshot shows the Cloud Manager interface with the 'Monitoring' tab selected. At the top, there are filters for 'testaws2' (selected), 'Time Frame: 6 h', and 'SVM: All'. Below this, two volume rows are displayed:

Volume Name	IOPS	Throughput	Latency
> vol0	66.75 Average IOPS	0.02 MB/s Average Throughput	0.26 ms Average Latency
> volaws2	50.02 Average IOPS	12.60 MB/s Average Throughput	0.19 ms Average Latency

- Click a volume in the table to expand the row and view a timeline for IOPS, throughput, and latency.



- Use the data to identify performance issues to minimize impact on your users and apps.

Getting more information from Cloud Insights

The Monitoring tab in Cloud Manager provides basic performance data for your volumes. You can go to the Cloud Insights web interface from your browser to perform more in-depth monitoring and to configure alerts for your Cloud Volumes ONTAP systems.

Steps

- At the top of Cloud Manager, click **Monitoring**.
- Click the **Cloud Insights** link.

The screenshot shows the Cloud Manager interface. At the top, there are two tabs: 'Monitoring' and 'Tiering'. In the top right corner, there is a blue link labeled 'Cloud Insights' with a small arrow icon. Below the tabs, there are two dropdown menus: 'Time Frame: 12 h' and 'SVM: All'. The main content area is titled 'Node 2' and displays two metrics: '0.04 MB/s' (Average Network Throughput) and '8.02 %' (Average CPU Utilization), each accompanied by a small icon.

Result

Cloud Insights open in a new browser tab. If you need help, refer to the [Cloud Insights documentation](#).

Disabling monitoring

If you no longer want to monitor Cloud Volumes ONTAP, you can disable the service at any time.



If you disable monitoring from each of your working environments, you'll need to delete the virtual machine instance yourself. The instance is named *AcquisitionUnit* with a generated hash (UUID) concatenated to it. For example: *AcquisitionUnit-FAN7FqeH*

Steps

1. At the top of Cloud Manager, click **Canvas**.
2. Select a working environment.
3. In the pane on the right, click the icon and select **Deactivate Scan**.

Improving protection against ransomware

Ransomware attacks can cost a business time, resources, and reputation. Cloud Manager enables you to implement the NetApp solution for ransomware, which provides effective tools for visibility, detection, and remediation.

Steps

1. From the working environment, click the **Ransomware** icon.

2. Implement the NetApp solution for ransomware:

- Click **Activate Snapshot Policy**, if you have volumes that do not have a Snapshot policy enabled.

NetApp Snapshot technology provides the industry's best solution for ransomware remediation. The key to a successful recovery is restoring from uninfected backups. Snapshot copies are read-only, which prevents ransomware corruption. They can also provide the granularity to create images of a single file copy or a complete disaster recovery solution.

- Click **Activate FPolicy** to enable ONTAP's FPolicy solution, which can block file operations based on a file's extension.

This preventative solution improves protection from ransomware attacks by blocking common ransomware file types.

The default FPolicy scope blocks files that have the following extensions:

micro, encrypted, locked, crypto, crypt, crinf, r5a, XRNT, XTBL, R16M01D05, pzdc, good, LOL!, OMG!, RDM, RRK, encryptedRS, crjoker, EnCiPhErEd, LeChiffre



Cloud Manager creates this scope when you activate FPolicy on Cloud Volumes ONTAP. The list is based on common ransomware file types. You can customize the blocked file extensions by using the `vserver fpolicy policy scope` commands from the Cloud Volumes ONTAP CLI.

Ransomware Protection

Ransomware attacks can cost a business time, resources, and reputation. The NetApp solution for ransomware provides effective tools for visibility, detection, and remediation. [Learn More](#)

1 Enable Snapshot Copy Protection ⓘ



1 Volumes without a Snapshot Policy

To protect your data, activate the default Snapshot policy for these volumes ⓘ

[Activate Snapshot Policy](#)

2 Block Ransomware File Extensions ⓘ



ONTAP's native FPolicy configuration monitors and blocks file operations based on a file's extension.

[View Denied File Names](#) ⓘ

[Activate FPolicy](#)

Administer

Manage licenses for Cloud Volumes ONTAP

The Digital Wallet page enables you to manage licenses for Cloud Volumes ONTAP from a single location. You can add new licenses and update existing licenses.

You can manage the following types of licenses for Cloud Volumes ONTAP:

- *Capacity-based licenses*, which enable you to pay for Cloud Volumes ONTAP per TiB of capacity.

[Contact NetApp Sales](#) to purchase a capacity-based license.

- *Node-based licenses*, which are any of the following:

- BYOL licenses purchased from NetApp.

This includes Eval licenses that you can later convert to BYOL licenses.

To get started with a node-based BYOL, [contact NetApp Sales](#).

- Hourly pay-as-you-go (PAYGO) subscriptions from your cloud provider's marketplace.

This includes 30-day free trials.

To get started with PAYGO, create a Cloud Volumes ONTAP working environment and subscribe to your cloud provider's marketplace when prompted.

[Learn more about Cloud Volumes ONTAP licenses.](#)

Manage capacity-based licenses

Manage your capacity-based licenses to ensure that your NetApp account has enough capacity for your Cloud Volumes ONTAP systems.

View your account's capacity

View the licensed capacity and provisioned capacity by package to ensure that you have enough room for your data volumes.

Steps

1. Click **All Services > Digital Wallet > Capacity-Based Licenses**.
2. View the licensed capacity and provisioned capacity for each package.

Serial Number	Package Type	Package Flavor	License Capacity	License Expiry	Actions
90120130000000000021	Professional	--	100 TB	January 1, 2022	...

3. If needed, purchase additional licensed capacity and then add the license to your account.

Add purchased licenses to your account

After you purchase licenses for your NetApp account, you need to add the licenses to Cloud Manager so that the capacity is available for Cloud Volumes ONTAP.

What you'll need

- You need to provide Cloud Manager the serial number of the license or the license file.
- If you want to enter the serial number, you first need to [add your NetApp Support Site account to Cloud Manager](#). This is the NetApp Support Site account that's authorized to access the serial number.

Steps

1. Click **All Services > Digital Wallet > Capacity-Based Licenses**.
2. Click **Add License**.
3. Enter the serial number for your capacity-based license or upload the license file.

If you entered a serial number, you also need to select the NetApp Support Site account that's authorized to access the serial number.

4. Click **Add License**.

Update a capacity-based license

If you purchased additional capacity or extended the term of your license, then you'll need to update the license in Cloud Manager.

What you'll need

The license file (or *files* if you have an HA pair).

Steps

1. Click **All Services > Digital Wallet > Capacity-Based Licenses**.
2. Click the action menu next to the license and select **Update License**.
3. Upload the license file.
4. Click **Upload License**.

Remove a capacity-based license

If a capacity-based license expired and is no longer in use, then you can remove it at any time.

Steps

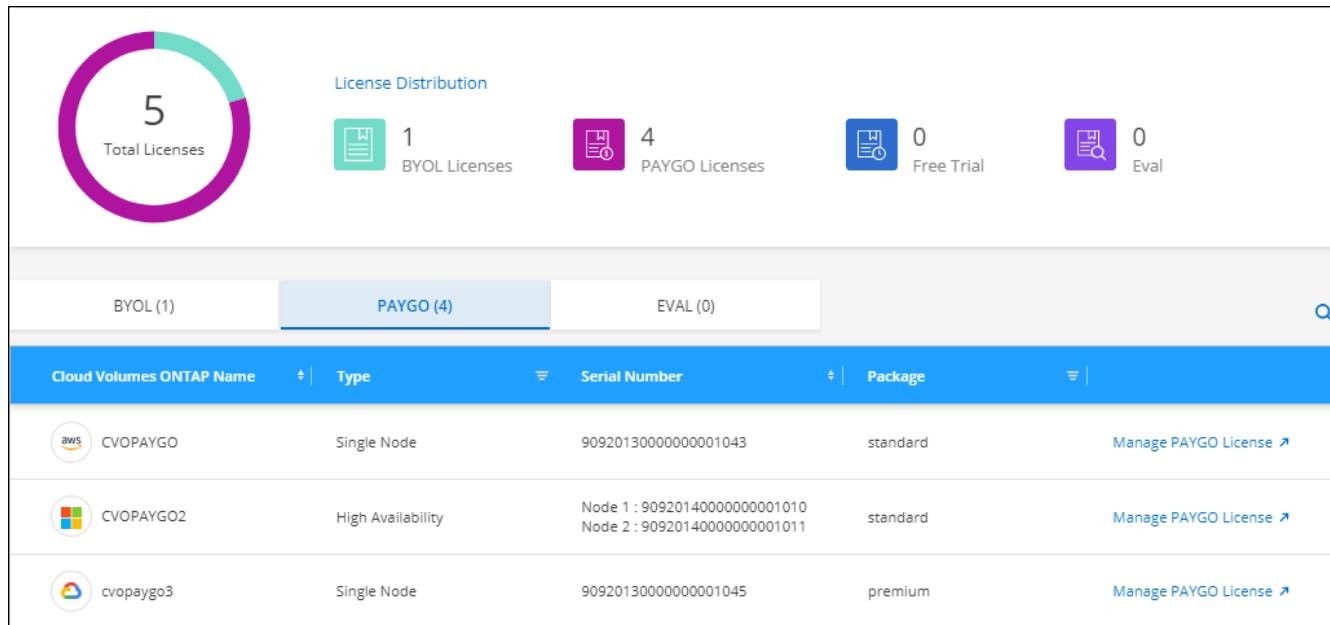
1. Click **All Services > Digital Wallet > Capacity-Based Licenses**.
2. Click the action menu next to the license and select **Remove License**.
3. Click **Remove** to confirm.

Manage node-based PAYGO licenses

The Digital Wallet page enables you to view details about each of your PAYGO Cloud Volumes ONTAP systems, including the serial number and PAYGO license type.

Steps

1. Click **All Services > Digital Wallet > Node-Based Licenses**.
2. Click **PAYGO**.
3. View details in the table about each of your PAYGO licenses.



4. If needed, click **Manage PAYGO License** to change the PAYGO license or to change the instance type.

Manage node-based BYOL licenses

Manage your node-based licenses to ensure that each Cloud Volumes ONTAP system has a valid license with the required capacity.

Obtain a system license file

In most cases, Cloud Manager can automatically obtain your license file using your NetApp Support Site account. But if it can't, then you'll need to manually upload the license file. If you don't have the license file, you can obtain it from netapp.com.

Steps

1. Go to the [NetApp License File Generator](#) and log in using your NetApp Support Site credentials.
2. Enter your password, choose your product, enter the serial number, confirm that you have read and accepted the privacy policy, and then click **Submit**.

Example

3. Choose whether you want to receive the serialnumber.NLF JSON file through email or direct download.

Update a system license

When you renew a BYOL subscription by contacting a NetApp representative, Cloud Manager automatically obtains the new license from NetApp and installs it on the Cloud Volumes ONTAP system.

If Cloud Manager can't access the license file over the secure internet connection, you can obtain the file yourself and then manually upload the file to Cloud Manager.

Steps

1. Click **All Services > Digital Wallet > Node-Based Licenses**.
2. In the **BYOL** tab, expand the details for a Cloud Volumes ONTAP system.
3. Click the action menu next to the system license and select **Update License**.
4. Upload the license file (or files if you have an HA pair).
5. Click **Update License**.

Result

Cloud Manager updates the license on the Cloud Volumes ONTAP system.

Manage extra capacity licenses

You can purchase extra capacity licenses for a Cloud Volumes ONTAP BYOL system to allocate more than the 368 TiB of capacity that's provided with a BYOL system license. For example, you might purchase one extra license capacity to allocate up to 736 TiB of capacity to Cloud Volumes ONTAP. Or you could purchase three extra capacity licenses to get up to 1.4 PB.

The number of licenses that you can purchase for a single node system or HA pair is unlimited.

Add capacity licenses

Purchase an extra capacity license by contacting us through the chat icon in the lower-right of Cloud Manager. After you purchase the license, you can apply it to a Cloud Volumes ONTAP system.

Steps

1. Click **All Services > Digital Wallet > Node-Based Licenses**.
2. In the **BYOL** tab, expand the details for a Cloud Volumes ONTAP system.
3. Click **Add Capacity License**.
4. Enter the serial number or upload the license file (or files if you have an HA pair).
5. Click **Add Capacity License**.

Update capacity licenses

If you extended the term of an extra capacity license, you'll need to update the license in Cloud Manager.

Steps

1. Click **All Services > Digital Wallet > Node-Based Licenses**.
2. In the **BYOL** tab, expand the details for a Cloud Volumes ONTAP system.
3. Click the action menu next to the capacity license and select **Update License**.
4. Upload the license file (or files if you have an HA pair).
5. Click **Update License**.

Remove capacity licenses

If an extra capacity license expired and is no longer in use, then you can remove it at any time.

Steps

1. Click **All Services > Digital Wallet > Node-Based Licenses**.
2. In the **BYOL** tab, expand the details for a Cloud Volumes ONTAP system.
3. Click the action menu next to the capacity license and select **Remove License**.
4. Click **Remove**.

Convert an Eval license to a BYOL

An evaluation license is good for 30 days. You can apply a new BYOL license on top of the evaluation license for an in-place upgrade.

When you convert an Eval license to a BYOL, Cloud Manager restarts the Cloud Volumes ONTAP system.

- For a single-node system, the restart results in I/O interruption during the reboot process.
- For an HA pair, the restart initiates takeover and giveback to continue serving I/O to clients.

[Learn more about Evaluation licenses.](#)

Steps

1. Click **All Services > Digital Wallet > Node-Based Licenses**.
2. Click **Eval**.
3. In the table, click **Convert to BYOL License** for a Cloud Volumes ONTAP system.
4. Enter the serial number or upload the license file.
5. Click **Convert License**.

Result

Cloud Manager starts the conversion process. Cloud Volumes ONTAP automatically restarts as part of this process. When it's back up, the licensing information will reflect the new license.

Registering pay-as-you-go systems

Support from NetApp is included with Cloud Volumes ONTAP Explore, Standard, and Premium systems, but you must first activate support by registering the systems with

NetApp.

Registering a PAYGO system with NetApp is required to upgrade ONTAP software using any of the methods described on this page.



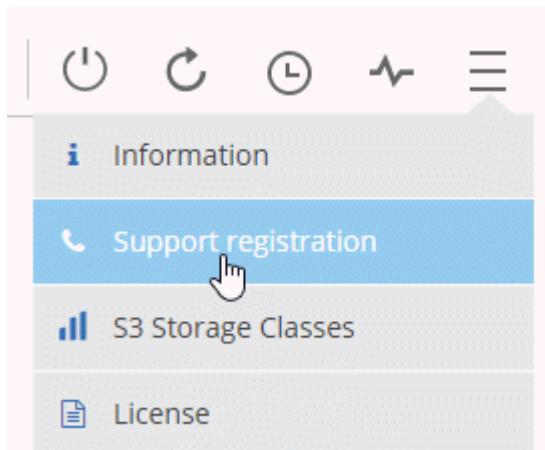
A system that isn't registered for support will still receive the software update notifications that appear in Cloud Manager when a new version is available. But you will need to register the system before you can upgrade the software.

Steps

1. If you have not yet added your NetApp Support Site account to Cloud Manager, go to **Account Settings** and add it now.

[Learn how to add NetApp Support Site accounts.](#)

2. On the Canvas page, double-click the name of the system that you want to register.
3. Click the menu icon and then click **Support registration**:



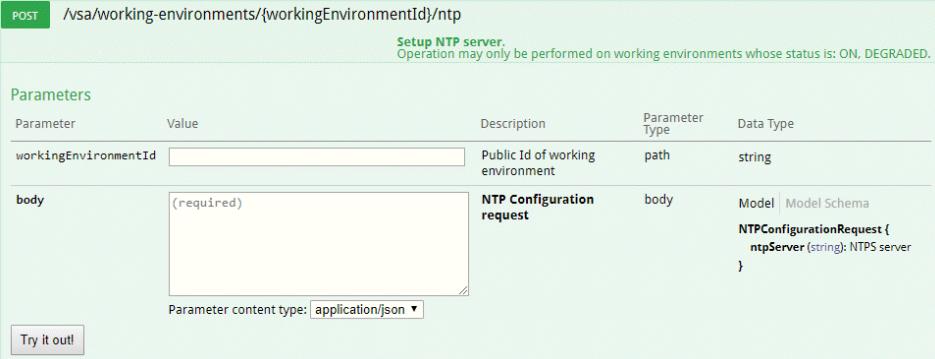
4. Select a NetApp Support Site account and click **Register**.

Result

Cloud Manager registers the system with NetApp.

Setting up Cloud Volumes ONTAP

After you deploy Cloud Volumes ONTAP, you can set it up by synchronizing the system time using NTP and by performing a few optional tasks from either System Manager or the CLI.

Task	Description
<p>Synchronize the system time using NTP</p>	<p>Specifying an NTP server synchronizes the time between the systems in your network, which can help prevent issues due to time differences.</p> <p>Specify an NTP server using the Cloud Manager API or from the user interface when you set up a CIFS server.</p> <ul style="list-style-type: none"> • Modifying the CIFS server • Cloud Manager automation docs <p>For example, here's the API for a single-node system in AWS:</p>  <pre> POST /vsa/working-environments/{workingEnvironmentId}/ntp Setup NTP server. Operation may only be performed on working environments whose status is: ON, DEGRADED. Parameters Parameter Value Description Parameter Type Data Type workingEnvironmentId [] Public Id of working environment path string body (required) NTP Configuration request body Model<NTPConfigurationRequest> NTPConfigurationRequest { ntpServer(string): NTPs server } Parameter content type: application/json ▾ Try it out! </pre>
<p>Optional: Configure AutoSupport</p>	<p>AutoSupport proactively monitors the health of your system and automatically sends messages to NetApp technical support by default.</p> <p>If the Account Admin added a proxy server to Cloud Manager before you launched your instance, Cloud Volumes ONTAP is configured to use that proxy server for AutoSupport messages.</p> <p>You should test AutoSupport to ensure that it can send messages. For instructions, see the System Manager Help or the ONTAP 9 System Administration Reference.</p>
<p>Optional: Configure EMS</p>	<p>The Event Management System (EMS) collects and displays information about events that occur on Cloud Volumes ONTAP systems. To receive event notifications, you can set event destinations (email addresses, SNMP trap hosts, or syslog servers) and event routes for a particular event severity.</p> <p>You can configure EMS using the CLI. For instructions, see the ONTAP 9 EMS Configuration Express Guide.</p>

Task	Description
Optional: Change the backup location of configuration files	<p>Cloud Volumes ONTAP automatically creates configuration backup files that contain information about the configurable options that it needs to operate properly.</p> <p>By default, Cloud Volumes ONTAP backs up the files to the Connector host every eight hours. If you want to send the backups to an alternate location, you can change the location to an FTP or HTTP server in your data center or in AWS. For example, you might already have a backup location for your FAS storage systems.</p> <p>You can change the backup location using the CLI. See the ONTAP 9 System Administration Reference.</p>

Upgrading Cloud Volumes ONTAP software

Cloud Manager includes several options that you can use to upgrade to the current Cloud Volumes ONTAP release. You should prepare Cloud Volumes ONTAP systems before you upgrade the software.

Requirements

You should be aware of the following requirements before you start the Cloud Volumes ONTAP upgrade process.

Software upgrades must be completed by Cloud Manager

Upgrades of Cloud Volumes ONTAP must be completed from Cloud Manager. You should not upgrade Cloud Volumes ONTAP by using System Manager or the CLI. Doing so can impact system stability.

Cloud Volumes ONTAP must be registered with NetApp Support

Cloud Volumes ONTAP must be registered with NetApp support in order to upgrade the software using any of the methods described on this page. This applies to both PAYGO and BYOL. You'll need to [manually register PAYGO systems](#), while BYOL systems are registered by default.

 A system that isn't registered for support will still receive the software update notifications that appear in Cloud Manager when a new version is available. But you will need to register the system before you can upgrade the software.

A note about downgrades

Cloud Manager doesn't support downgrading Cloud Volumes ONTAP to a previous version. Contact NetApp technical support for help with downgrades.

Ways to upgrade Cloud Volumes ONTAP

Cloud Manager displays a notification in Cloud Volumes ONTAP working environments when a new version of Cloud Volumes ONTAP is available:

The screenshot shows the Cloud Manager interface. At the top, there's a cloud icon and the text "Visual View". Below that, a system card for "cloudvolumesontap1" is shown, indicating it is "On | AWS". The card has three icons: a blue circle with an exclamation mark, a vertical ellipsis, and a close button. Underneath the card, a red box highlights the "NOTIFICATIONS" section. Inside this section, a star icon indicates a "New version available". To the right of the notification is a small arrow icon. Below the notifications, there's a "SERVICES" section with two items: "Cloud Compliance" (status "On") and "No Personal Files Found".

You can start the upgrade process from this notification, which automates the process by obtaining the software image from an S3 bucket, installing the image, and then restarting the system. For details, see [Upgrading Cloud Volumes ONTAP from Cloud Manager notifications](#).



For HA systems in AWS, Cloud Manager might upgrade the HA mediator as part of the upgrade process.

Cloud Manager also provides the following advanced options for upgrading Cloud Volumes ONTAP software with both PAYGO and BYOL:

- Software upgrades using an image on an external URL

This option is helpful if Cloud Manager can't access the S3 bucket to upgrade the software or if you were provided with a patch.

For details, see [Upgrading Cloud Volumes ONTAP by using an HTTP or FTP server](#).

- Software upgrades using the alternate image on the system

You can use this option to upgrade by making the alternate software image the default image. This option is not available for HA pairs.

For details, see [Upgrading Cloud Volumes ONTAP by using a local image](#).

Preparing to upgrade Cloud Volumes ONTAP software

Before performing an upgrade, you must verify that your systems are ready and make any required configuration changes.

- Understanding supported upgrade paths
- Planning for downtime
- Verifying that automatic giveback is still enabled
- Suspending SnapMirror transfers
- Verifying that aggregates are online

Understanding supported upgrade paths

Refer to the "Upgrade notes" page in the Cloud Volumes ONTAP Release Notes for the target version.

Planning for downtime

When you upgrade a single-node system, the upgrade process takes the system offline for up to 25 minutes, during which I/O is interrupted.

Upgrading an HA pair is nondisruptive and I/O is uninterrupted. During this nondisruptive upgrade process, each node is upgraded in tandem to continue serving I/O to clients.

Verifying that automatic giveback is still enabled

Automatic giveback must be enabled on a Cloud Volumes ONTAP HA pair (this is the default setting). If it isn't, then the operation will fail.

[ONTAP 9 Documentation: Commands for configuring automatic giveback](#)

Suspending SnapMirror transfers

If a Cloud Volumes ONTAP system has active SnapMirror relationships, it is best to suspend transfers before you update the Cloud Volumes ONTAP software. Suspending the transfers prevents SnapMirror failures. You must suspend the transfers from the destination system.



Even though Cloud Backup uses an implementation of SnapMirror to create backup files (called SnapMirror Cloud), backups do not need to be suspended when a system is upgraded.

About this task

These steps describe how to use System Manager for version 9.3 and later.

Steps

1. [Log in to System Manager](#) from the destination system.
2. Click **Protection > Relationships**.
3. Select the relationship and click **Operations > Quiesce**.

Verifying that aggregates are online

Aggregates for Cloud Volumes ONTAP must be online before you update the software. Aggregates should be online in most configurations, but if they are not, then you should bring them online.

About this task

These steps describe how to use System Manager for version 9.3 and later.

Steps

1. In the working environment, click the menu icon, and then click **Advanced > Advanced allocation**.
2. Select an aggregate, click **Info**, and then verify that the state is online.

The screenshot shows a detailed view of an aggregate named 'aggr1'. The following information is displayed:

- Aggregate Capacity: 88.57 GB
- Used Aggregate Capacity: 1.07 GB
- Volumes: 2
- AWS Disks: 1
- State: online

The 'State: online' field is highlighted with a red rectangular border.

3. If the aggregate is offline, use System Manager to bring the aggregate online:
 - a. [Log in to System Manager](#).
 - b. Click **Storage > Aggregates & Disks > Aggregates**.
 - c. Select the aggregate, and then click **More Actions > Status > Online**.

Upgrading Cloud Volumes ONTAP from Cloud Manager notifications

Cloud Manager notifies you when a new version of Cloud Volumes ONTAP is available. Click the notification to start the upgrade process.

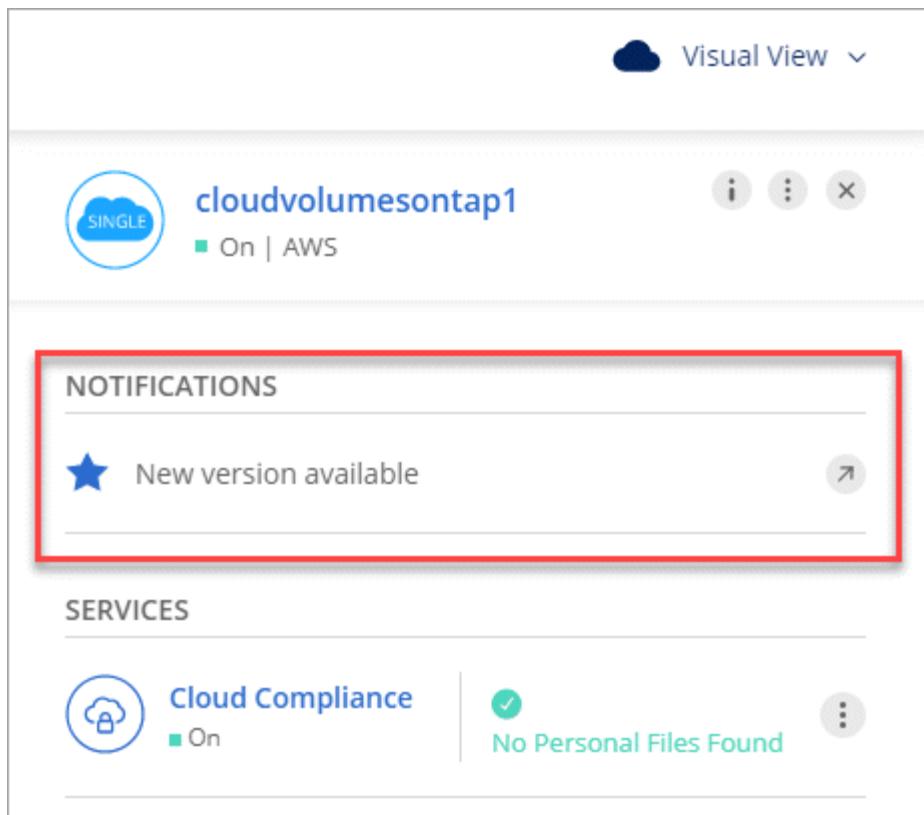
Before you begin

Cloud Manager operations such as volume or aggregate creation must not be in progress for the Cloud Volumes ONTAP system.

Steps

1. Click **Canvas**.
2. Select a working environment.

A notification appears in the right pane if a new version is available:



3. If a new version is available, click **Upgrade**.
4. In the Release Information page, click the link to read the Release Notes for the specified version, and then select the **I have read...** check box.
5. In the End User License Agreement (EULA) page, read the EULA, and then select **I read and approve the EULA**.
6. In the Review and Approve page, read the important notes, select **I understand...**, and then click **Go**.

Result

Cloud Manager starts the software upgrade. You can perform actions on the working environment once the software update is complete.

After you finish

If you suspended SnapMirror transfers, use System Manager to resume the transfers.

Upgrading Cloud Volumes ONTAP by using an HTTP or FTP server

You can place the Cloud Volumes ONTAP software image on an HTTP or FTP server and then initiate the software upgrade from Cloud Manager. You might use this option if Cloud Manager can't access the S3 bucket to upgrade the software.

Steps

1. Set up an HTTP server or FTP server that can host the Cloud Volumes ONTAP software image.
2. If you have a VPN connection to the virtual network, you can place the Cloud Volumes ONTAP software image on an HTTP server or FTP server in your own network. Otherwise, you must place the file on an HTTP server or FTP server in the cloud.
3. If you use your own security group for Cloud Volumes ONTAP, ensure that the outbound rules allow HTTP or FTP connections so Cloud Volumes ONTAP can access the software image.



The predefined Cloud Volumes ONTAP security group allows outbound HTTP and FTP connections by default.

4. Obtain the software image from [the NetApp Support Site](#).
5. Copy the software image to the directory on the HTTP or FTP server from which the file will be served.
6. From the working environment in Cloud Manager, click the menu icon, and then click **Advanced > Update Cloud Volumes ONTAP**.
7. On the update software page, choose **Select an image available from a URL**, enter the URL, and then click **Change Image**.
8. Click **Proceed** to confirm.

Result

Cloud Manager starts the software update. You can perform actions on the working environment once the software update is complete.

After you finish

If you suspended SnapMirror transfers, use System Manager to resume the transfers.

Upgrading Cloud Volumes ONTAP by using a local image

Each Cloud Volumes ONTAP system can hold two software images: the current image that is running, and an alternate image that you can boot. Cloud Manager can change the alternate image to be the default image.

Steps

1. From the working environment, click the menu icon, and then click **Advanced > Update Cloud Volumes ONTAP**.
2. On the update software page, select the alternate image, and then click **Change Image**.
3. Click **Proceed** to confirm.

Result

Cloud Manager starts the software update. You can perform actions on the working environment once the software update is complete.

After you finish

If you suspended SnapMirror transfers, use System Manager to resume the transfers.

Modifying Cloud Volumes ONTAP systems

You might need to change the configuration of Cloud Volumes ONTAP systems as your storage needs change. For example, you can change between pay-as-you-go configurations, change the instance or VM type, and more.

Changing the instance or machine type for Cloud Volumes ONTAP

You can choose from several instance or machine types when you launch Cloud Volumes ONTAP in AWS, Azure, or GCP. You can change the instance or machine type at any time if you determine that it is undersized or oversized for your needs.

About this task

- Automatic giveback must be enabled on a Cloud Volumes ONTAP HA pair (this is the default setting). If it isn't, then the operation will fail.

[ONTAP 9 Documentation: Commands for configuring automatic giveback](#)

- Changing the instance or machine type affects cloud provider service charges.
- The operation restarts Cloud Volumes ONTAP.

For single node systems, I/O is interrupted.

For HA pairs, the change is nondisruptive. HA pairs continue to serve data.



Cloud Manager gracefully changes one node at a time by initiating takeover and waiting for give back. NetApp's QA team tested both writing and reading files during this process and didn't see any issues on the client side. As connections changed, we did see retries on the I/O level, but the application layer overcame these short "re-wire" of NFS/CIFS connections.

Steps

- From the working environment, click the menu icon, and then click **Change license or instance** for AWS, **Change license or VM** for Azure, or **Change license or machine** for GCP.
- If you are using a pay-as-you-go configuration, you can optionally choose a different license.
- Select an instance or machine type, select the check box to confirm that you understand the implications of the change, and then click **OK**.

Result

Cloud Volumes ONTAP reboots with the new configuration.

Changing between hourly (pay-as-you-go) configurations

After you launch an hourly, pay-as-you-go Cloud Volumes ONTAP system, you can change between the Explore, Standard, and Premium configurations at any time by modifying the license. Changing the license increases or decreases the raw capacity limit and enables you to choose from different instance or VM types.

About this task

Note the following about changing between pay-as-you-go licenses:

- The operation restarts Cloud Volumes ONTAP.

For single node systems, I/O is interrupted.

For HA pairs, the change is nondisruptive. HA pairs continue to serve data.

- Changing the instance or machine type affects cloud provider service charges.

Steps

- From the working environment, click the menu icon, and then click **Change license or instance** for AWS, **Change license or VM** for Azure, or **Change license or machine** for GCP.
- Select a license type and an instance type or machine type, select the check box to confirm that you understand the implications of the change, and then click **OK**.

Result

Cloud Volumes ONTAP reboots with the new license, instance type or machine type, or both.

Moving to an alternate Cloud Volumes ONTAP configuration

If you want to switch between a pay-by-the-hour (PAYGO) subscription, an annual contract (AWS only), or a BYOL subscription, then you need to create a new Cloud Volumes ONTAP working environment and then replicate data from the existing working environment to the new working environment. The same is true if you want to switch between a single Cloud Volumes ONTAP system and an HA pair.

Steps

1. Create a new Cloud Volumes ONTAP working environment.

[Launching Cloud Volumes ONTAP in AWS](#)

[Launching Cloud Volumes ONTAP in Azure](#)

[Launching Cloud Volumes ONTAP in GCP](#)

2. Set up one-time data replication between the systems for each volume that you must replicate.
3. Terminate the Cloud Volumes ONTAP system that you no longer need by [deleting the original working environment](#).

Changing write speed to normal or high

Cloud Manager enables you to choose a normal or high write speed for Cloud Volumes ONTAP. The default write speed is normal. You can change to high write speed if fast write performance is required for your workload.

High write speed is supported with all types of single node systems. It's also supported with HA pairs in AWS and Azure when using a specific instance or VM type ([click here to see the list of supported instances and VM types](#)). High write speed is not supported with HA pairs in GCP.

Before you change the write speed, you should [understand the differences between the normal and high settings](#).

About this task

- Ensure that operations such as volume or aggregate creation are not in progress.
- Be aware that this change restarts Cloud Volumes ONTAP.

Steps

1. From the working environment, click the menu icon, and then click **Advanced > Writing Speed**.
2. Select **Normal** or **High**.

If you choose High, then you'll need to read the "I understand..." statement and confirm by checking the box.

3. Click **Save**, review the confirmation message, and then click **Proceed**.

Modifying the storage VM name

Cloud Manager automatically names the single storage VM (SVM) that it creates for Cloud Volumes ONTAP. You can modify the name of the SVM if you have strict naming standards. For example, you might want the name to match how you name the SVMs for your ONTAP clusters.

But if you created any additional SVMs for Cloud Volumes ONTAP, then you can't rename the SVMs from

Cloud Manager. You'll need to do so directly from Cloud Volumes ONTAP by using System Manager or the CLI.

Steps

1. From the working environment, click the menu icon, and then click **Information**.
2. Click the edit icon to the right of the storage VM name.

Working Environment Information

ONTAP

Serial Number: [REDACTED]

System ID: system-id-capacitytest

Cluster Name: capacitytest

ONTAP Version: 9.7RC1

Date Created: Jul 6, 2020 07:42:02 am

Storage VM Name: **svm_capacitytest**  

3. In the Modify SVM Name dialog box, change the name, and then click **Save**.

Changing the password for Cloud Volumes ONTAP

Cloud Volumes ONTAP includes a cluster admin account. You can change the password for this account from Cloud Manager, if needed.



You should not change the password for the admin account through System Manager or the CLI. The password will not be reflected in Cloud Manager. As a result, Cloud Manager cannot monitor the instance properly.

Steps

1. From the working environment, click the menu icon, and then click **Advanced > Set password**.
2. Enter the new password twice and then click **Save**.

The new password must be different than one of the last six passwords that you used.

Changing route tables associated with HA pairs in multiple AWS AZs

You can modify the AWS route tables that include routes to the floating IP addresses for an HA pair. You might

do this if new NFS or CIFS clients need to access an HA pair in AWS.

Steps

1. From the working environment, click the menu icon and then click **Information**.
2. Click **Route Tables**.
3. Modify the list of selected route tables and then click **Save**.

Result

Cloud Manager sends an AWS request to modify the route tables.

Managing the state of Cloud Volumes ONTAP

You can stop and start Cloud Volumes ONTAP from Cloud Manager to manage your cloud compute costs.

Scheduling automatic shutdowns of Cloud Volumes ONTAP

You might want to shut down Cloud Volumes ONTAP during specific time intervals to lower your compute costs. Rather than do this manually, you can configure Cloud Manager to automatically shut down and then restart systems at specific times.

About this task

- When you schedule an automatic shutdown of your Cloud Volumes ONTAP system, Cloud Manager postpones the shutdown if an active data transfer is in progress.

Cloud Manager shuts down the system after the transfer is complete.

- This task schedules automatic shutdowns of both nodes in an HA pair.
- Snapshots of boot and root disks are not created when turning off Cloud Volumes ONTAP through scheduled shutdowns.

Snapshots are automatically created only when performing a manual shutdown, as described in the next section.

Steps

1. From the working environment, click the clock icon:



2. Specify the shutdown schedule:

- a. Choose whether you want to shut down the system every day, every weekday, every weekend, or any combination of the three options.
- b. Specify when you want to turn off the system and for how long you want it turned off.

Example

The following image shows a schedule that instructs Cloud Manager to shut down the system every Saturday at 12:00 a.m. for 48 hours. Cloud Manager restarts the system every Monday at 12:00 a.m.

<input type="checkbox"/> Turn off every weekday	turn off at	08	:	00	PM	for	12	Hours (1-24)
Mon, Tue, Wed, Thu, Fri								
<hr/>								
<input checked="" type="checkbox"/> Turn off every weekend	turn off at	12	:	00	AM	for	48	Hours (1-48)
Sat								

3. Click **Save**.

Result

Cloud Manager saves the schedule. The clock icon changes to indicate that a schedule is set:



Stopping Cloud Volumes ONTAP

Stopping Cloud Volumes ONTAP saves you from accruing compute costs and creates snapshots of the root and boot disks, which can be helpful for troubleshooting.



To reduce costs, Cloud Manager periodically deletes older snapshots of root and boot disks. Only the two most recent snapshots are retained for both the root and boot disks.

About this task

When you stop an HA pair, Cloud Manager shuts down both nodes.

Steps

1. From the working environment, click the **Turn off** icon.



2. Keep the option to create snapshots enabled because the snapshots can enable system recovery.

3. Click **Turn Off**.

It can take up to a few minutes to stop the system. You can restart systems at a later time from the working environment page.

Monitoring AWS resource costs

Cloud Manager enables you to view the resource costs associated with running Cloud Volumes ONTAP in AWS. You can also see how much money you saved by using NetApp features that can reduce storage costs.

About this task

Cloud Manager updates the costs when you refresh the page. You should refer to AWS for final cost details.

Step

1. Verify that Cloud Manager can obtain cost information from AWS:

- a. Ensure that the IAM policy that provides Cloud Manager with permissions includes the following actions:

```
"ce:GetReservationUtilization",
"ce:GetDimensionValues",
"ce:GetCostAndUsage",
"ce:GetTags"
```

These actions are included in the latest [Cloud Manager policy](#). New systems deployed from NetApp Cloud Central automatically include these permissions.

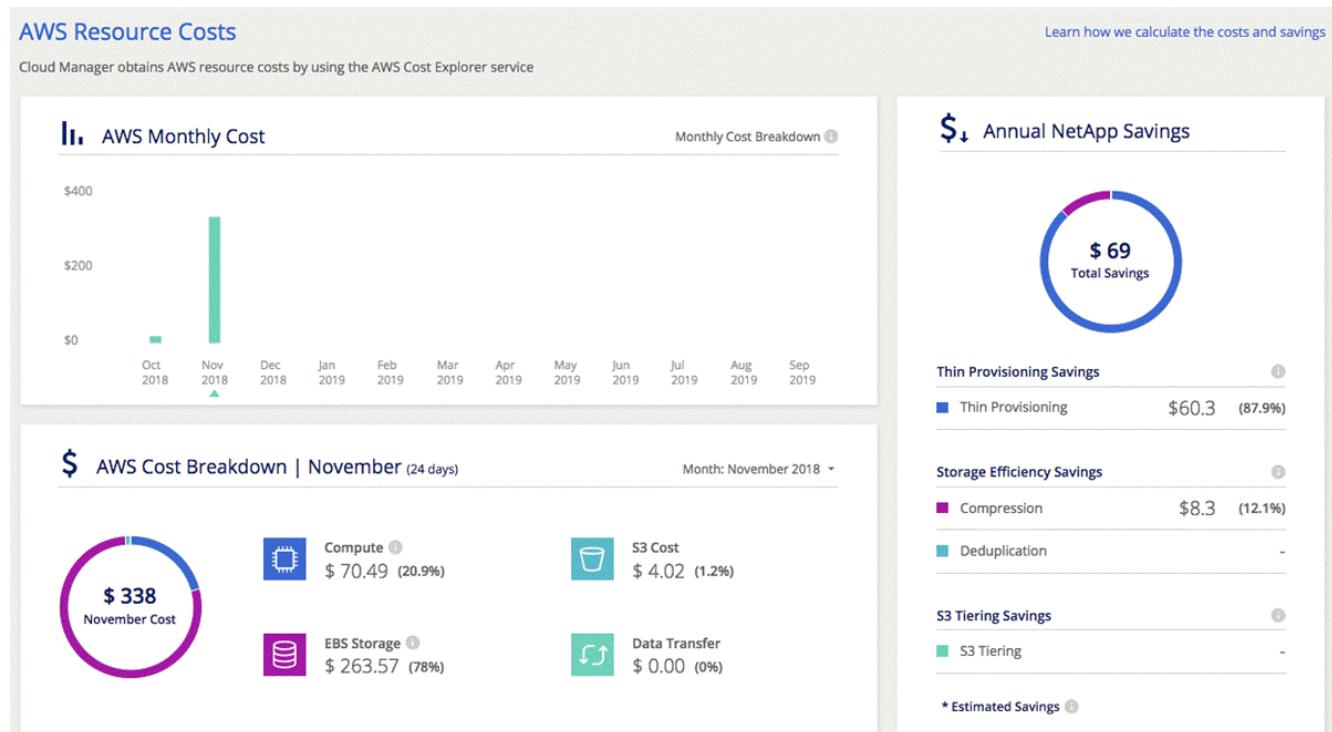
- b. [Activate the WorkingEnvironmentId tag](#).

To track your AWS costs, Cloud Manager assigns a cost allocation tag to Cloud Volumes ONTAP instances. After you create your first working environment, activate the **WorkingEnvironmentId** tag. User-defined tags don't appear on AWS billing reports until you activate them in the Billing and Cost Management console.

2. On the Canvas page, select a Cloud Volumes ONTAP working environment and then click **Cost**.

The Cost page displays costs for the current and previous months and shows your annual NetApp savings, if you enabled NetApp's cost-saving features on volumes.

The following image shows a sample Cost page:



Connecting to Cloud Volumes ONTAP

If you need to perform advanced management of Cloud Volumes ONTAP, you can do so using OnCommand System Manager or the command line interface.

Connecting to System Manager

You might need to perform some Cloud Volumes ONTAP tasks from System Manager, which is a browser-based management tool that runs on the Cloud Volumes ONTAP system. For example, you need to use System Manager if you want to create LUNs.

Before you begin

The computer from which you are accessing Cloud Manager must have a network connection to Cloud Volumes ONTAP. For example, you might need to log in to Cloud Manager from a jump host that's in your cloud provider network.



When deployed in multiple AWS Availability Zones, Cloud Volumes ONTAP HA configurations use a floating IP address for the cluster management interface, which means external routing is not available. You must connect from a host that is part of the same routing domain.

Steps

1. From the Canvas page, double-click the Cloud Volumes ONTAP system that you want to manage with System Manager.
2. Click the menu icon, and then click **Advanced > System Manager**.
3. Click **Launch**.

System Manager loads in a new browser tab.

4. At the login screen, enter **admin** in the User Name field, enter the password that you specified when you created the working environment, and then click **Sign In**.

Result

The System Manager console loads. You can now use it to manage Cloud Volumes ONTAP.

Connecting to the Cloud Volumes ONTAP CLI

The Cloud Volumes ONTAP CLI enables you to run all administrative commands and is a good choice for advanced tasks or if you are more comfortable using the CLI. You can connect to the CLI using Secure Shell (SSH).

Before you begin

The host from which you use SSH to connect to Cloud Volumes ONTAP must have a network connection to Cloud Volumes ONTAP. For example, you might need to use SSH from a jump host in AWS or Azure.



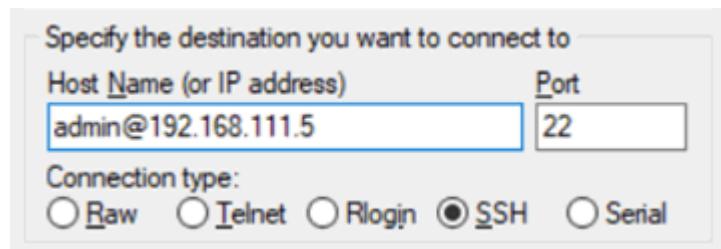
When deployed in multiple AZs, Cloud Volumes ONTAP HA configurations use a floating IP address for the cluster management interface, which means external routing is not available. You must connect from a host that is part of the same routing domain.

Steps

1. In Cloud Manager, identify the IP address of the cluster management interface:
 - a. On the Canvas page, select the Cloud Volumes ONTAP system.
 - b. Copy the cluster management IP address that appears in the right pane.
2. Use SSH to connect to the cluster management interface IP address using the admin account.

Example

The following image shows an example using PuTTY:



3. At the login prompt, enter the password for the admin account.

Example

```
 Password: *****
 COT2::>
```

Adding existing Cloud Volumes ONTAP systems to Cloud Manager

You can discover and add existing Cloud Volumes ONTAP systems to Cloud Manager. You might do this if you deployed a new Cloud Manager system.

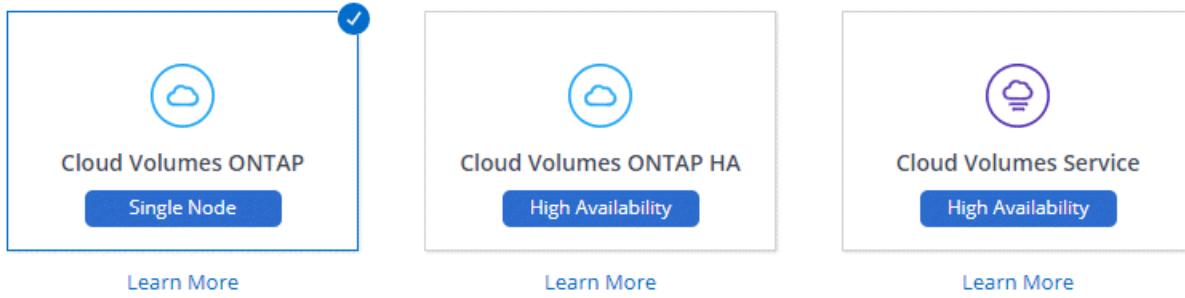
Before you begin

You must know the password for the Cloud Volumes ONTAP admin user account.

Steps

1. On the Canvas page, click **Add Working Environment**.
2. Select the cloud provider in which the system resides.
3. Choose the type of Cloud Volumes ONTAP system.
4. Click the link to discover an existing system.

Define Your Working Environment



You're about to create a new **Cloud Volumes ONTAP** system in AWS.

If you want to discover an existing Cloud Volumes ONTAP in AWS, [click here](#).

5. On the Region page, choose the region where the instances are running, and then select the instances.
6. On the Credentials page, enter the password for the Cloud Volumes ONTAP admin user, and then click **Go**.

Result

Cloud Manager adds the Cloud Volumes ONTAP instances to the workspace.

Overriding CIFS locks for Cloud Volumes ONTAP HA in Azure

The Account Admin can enable a setting in Cloud Manager that prevents issues with Cloud Volumes ONTAP storage failover during Azure maintenance events. When you enable this setting, Cloud Volumes ONTAP vetoes CIFS locks and resets active CIFS sessions.

About this task

Microsoft Azure schedules periodic maintenance events on its virtual machines. When a maintenance event occurs on a node in a Cloud Volumes ONTAP HA pair, the HA pair initiates storage takeover. If there are active CIFS sessions during this maintenance event, the locks on CIFS files can prevent storage failover.

If you enable this setting, Cloud Volumes ONTAP will veto the locks and reset the active CIFS sessions. As a result, the HA pair can complete storage failover during these maintenance events.



This process might be disruptive to CIFS clients. Data that is not committed from CIFS clients could be lost.

What you'll need

You need to create a Connector before you can change Cloud Manager settings. [Learn how](#).

Steps

1. In the upper right of the Cloud Manager console, click the Settings icon, and select **Connector Settings**.



2. Under **Azure**, click **Azure CIFS locks for Azure HA working environments**.

3. Click the checkbox to enable the feature and then click **Save**.

Using an Azure Private Link with Cloud Volumes ONTAP

By default, Cloud Manager enables an Azure Private Link connection between Cloud Volumes ONTAP and its associated storage accounts. A Private Link secures connections between endpoints in Azure and provides performance benefits. [Learn more](#).

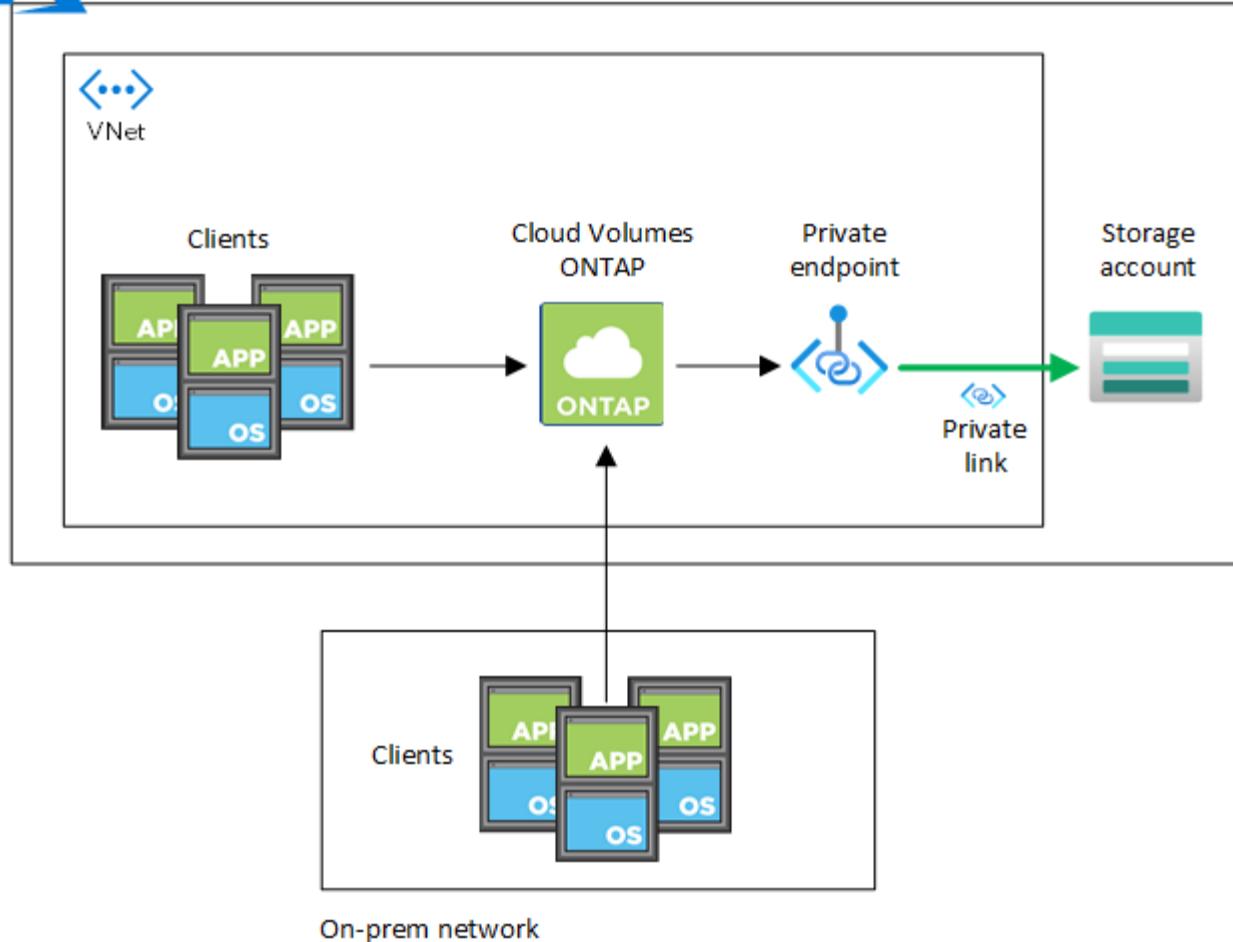
In most cases, there's nothing that you need to do—Cloud Manager manages the Azure Private Link for you. But if you use Azure Private DNS, then you'll need to edit a configuration file. You can also disable the Private Link connection, if desired.

How Private Link connections work with Cloud Volumes ONTAP

When Cloud Manager deploys Cloud Volumes ONTAP in Azure, it creates a private endpoint in the resource group. The private endpoint is associated with the storage account for Cloud Volumes ONTAP. As a result, access to Cloud Volumes ONTAP storage travels through the Microsoft backbone network.

Client access goes through the private link when clients are within the same VNet as Cloud Volumes ONTAP, within peered VNets, or in your on-premises network when using a private VPN or ExpressRoute connection to the VNet.

Here's an example that shows client access over a private link from within the same VNet and from an on-prem network that has either a private VPN or ExpressRoute connection.



Providing Cloud Manager with details about your Azure Private DNS

If you use [Azure Private DNS](#), then you need to modify a configuration file on each Connector. Otherwise, Cloud Manager can't enable the Azure Private Link connection between Cloud Volumes ONTAP and its associated storage accounts.

Note that the DNS name must match Azure DNS naming requirements as shown in [Azure documentation](#).

Steps

1. SSH to the Connector host and log in.
2. Navigate to the following directory: `/opt/application/netapp/cloudmanager/docker_occm/data`
3. Edit `app.conf` by modifying the following parameters as shown:

```
"user-private-dns-zone-settings": {
    "use-existing": true,
    "resource-group": "<resource group name of the DNS zone>",
    "subscription": "<subscription ID>"
}
```

The subscription parameter is required only if the Private DNS Zone exists in a different subscription than

the Connector.

4. Save the file and log off the Connector.

A reboot isn't required.

Disabling Azure Private Link connections

If required for your Azure configuration, you can disable the Azure Private Link connection between Cloud Volumes ONTAP and storage accounts.

Steps

1. In the upper right of the Cloud Manager console, click the Settings icon, and select **Connector Settings**.
2. Under **Azure**, click **Use Azure Private Link**.
3. Deselect **Private Link connection between Cloud Volumes ONTAP and storage accounts**.
4. Click **Save**.

Deleting a Cloud Volumes ONTAP working environment

You should always delete Cloud Volumes ONTAP systems from Cloud Manager, rather than from your cloud provider's console. For example, if you terminate a licensed Cloud Volumes ONTAP instance from your cloud provider, then you can't use the license key for another instance. You must delete the working environment from Cloud Manager to release the license.

When you delete a working environment, Cloud Manager terminates Cloud Volumes ONTAP instances and deletes disks and snapshots.

Resources managed by other services like backups for Cloud Backup and instances for Cloud Data Sense and Monitoring are not deleted when you delete a working environment. You'll need to manually delete them yourself. If you don't, then you'll continue to receive charges for these resources.



When Cloud Manager deploys Cloud Volumes ONTAP in your cloud provider, it enables termination protection on the instances. This option helps prevent accidental termination.

Steps

1. If you enabled Cloud Backup on the working environment, determine whether the backed up data is still required and then [delete the backups, if necessary](#).

Cloud Backup is independent from Cloud Volumes ONTAP by design. Cloud Backup doesn't automatically delete backups when you delete a Cloud Volumes ONTAP system, and there is no current support in the UI to delete the backups after the system has been deleted.

2. If you enabled Cloud Data Sense or Monitoring on this working environment and no other working environments use those services, then you'll need to delete the instances for those services.

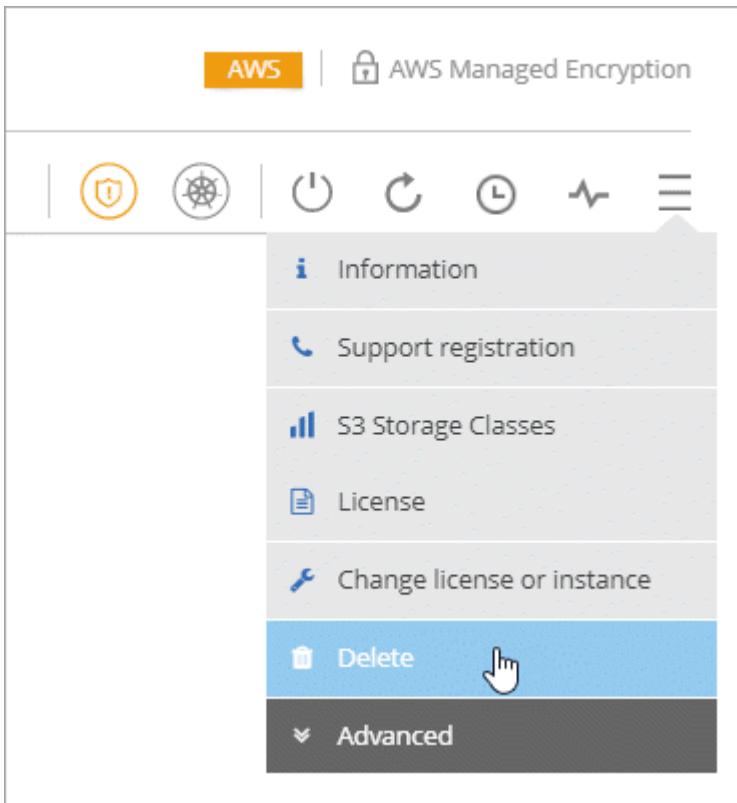
- [Learn more about the Cloud Data Sense instance](#).
- [Learn more about the Monitoring Acquisition Unit](#).

3. Delete the Cloud Volumes ONTAP working environment.

- a. On the Canvas page, double-click the name of the Cloud Volumes ONTAP working environment that

you want to delete.

- b. Click menu icon and then click **Delete**.



- c. Type the name of the working environment and then click **Delete**.

It can take up to 5 minutes to delete the working environment.

Provision volumes using a file service

Azure NetApp Files

Learn about Azure NetApp Files

Azure NetApp Files enables enterprises to migrate and run their performance-intensive and latency-sensitive core, business-critical applications in Azure with no need to refactor for the cloud.

Features

- Support for multiple protocols enables "lift & shift" of both Linux & Windows applications to run seamlessly in Azure.
- Multiple performance tiers allow for close alignment with workload performance requirements.
- Leading certifications including SAP HANA, GDPR, and HIPAA enables migration of the most demanding workloads to Azure.

Additional features in Cloud Manager

- Migrate NFS or SMB data to Azure NetApp Files directly from Cloud Manager. Data migrations are powered by NetApp's Cloud Sync service. [Learn more](#).
- Using Artificial Intelligence (AI) driven technology, Cloud Data Sense can help you understand data context and identify sensitive data that resides in your Azure NetApp Files accounts. [Learn more](#).

Cost

[View Azure NetApp Files pricing](#).

Note that your subscription and charging are maintained by the Azure NetApp Files service and not by Cloud Manager.

Supported regions

[View supported Azure regions](#).

Requesting access

You need to be granted access to Azure NetApp Files by [submitting an online request](#). You'll need to wait for approval from the Azure NetApp Files team before you can proceed.

Getting help

For technical support issues associated with Azure NetApp Files, use the Azure portal to log a support request to Microsoft. Select your associated Microsoft subscription and select the **Azure NetApp Files** service name under **Storage**. Provide the remaining information required to create your Microsoft support request.

For issues related to Cloud Sync and Azure NetApp Files, you can start with NetApp using your Cloud Sync serial number directly from the Cloud Sync service. You will need to access the Cloud Sync service through the link in Cloud Manager. [View the process to enable Cloud Sync support](#).

Related links

- [NetApp Cloud Central: Azure NetApp Files](#)
- [Azure NetApp Files documentation](#)
- [Cloud Sync documentation](#)

Setting up and discovering Azure NetApp Files

Create an Azure NetApp Files working environment in Cloud Manager to create and manage NetApp accounts, capacity pools, volumes, and snapshots.

If you haven't set up Azure NetApp Files yet, you'll need to complete all of the steps on this page.

If you already set up Azure NetApp Files from outside of Cloud Manager, then you simply need to set up an Azure AD application and then create the Azure NetApp Files working environment.

Quick start

Get started quickly by following these steps or scroll down to the remaining sections for full details.



Request access

Submit an [online request](#) to be granted access to Azure NetApp Files.



Set up an Azure AD application

From Azure, grant permissions to an Azure AD application and copy the application (client) ID, the directory (tenant) ID, and the value of a client secret.



Create an Azure NetApp Files working environment

In Cloud Manager, click **Add Working Environment > Microsoft Azure > Azure NetApp Files** and then provide details about the AD application.

Requesting access

You need to be granted access to Azure NetApp Files by [submitting an online request](#). You'll need to wait for approval from the Azure NetApp Files team before you can proceed.

Setting up an Azure AD application

Cloud Manager needs permissions to set up and manage Azure NetApp Files. You can grant the required permissions to an Azure account by creating and setting up an Azure AD application and by obtaining the Azure credentials that Cloud Manager needs.

Creating the AD application

Create an Azure Active Directory (AD) application and service principal that Cloud Manager can use for role-

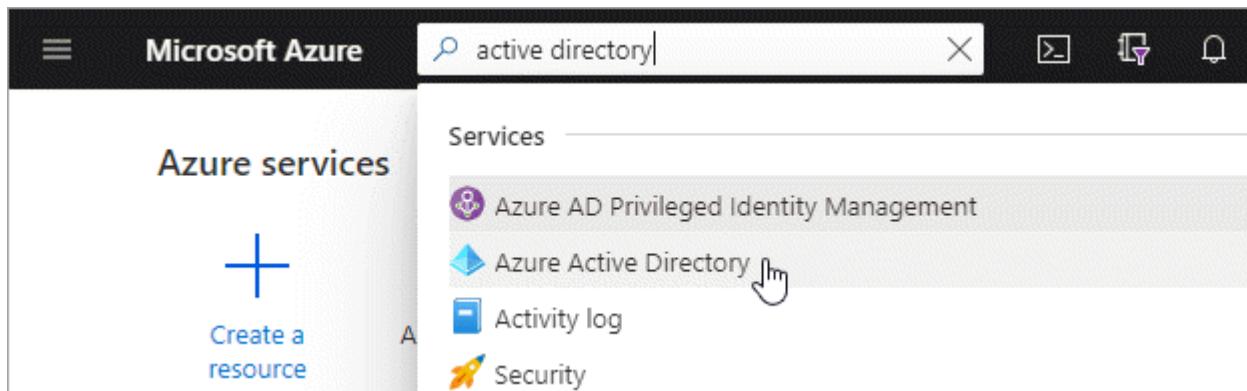
based access control.

Before you begin

You must have the right permissions in Azure to create an Active Directory application and to assign the application to a role. For details, refer to [Microsoft Azure Documentation: Required permissions](#).

Steps

1. From the Azure portal, open the **Azure Active Directory** service.



The screenshot shows the Microsoft Azure portal interface. At the top, there is a search bar with the text "active directory". Below the search bar, the "Azure services" sidebar is visible, featuring a "Create a resource" button and icons for "Azure AD Privileged Identity Management", "Azure Active Directory" (which has a hand cursor icon over it), "Activity log", and "Security". The "Azure Active Directory" item is highlighted.

2. In the menu, click **App registrations**.

3. Create the application:

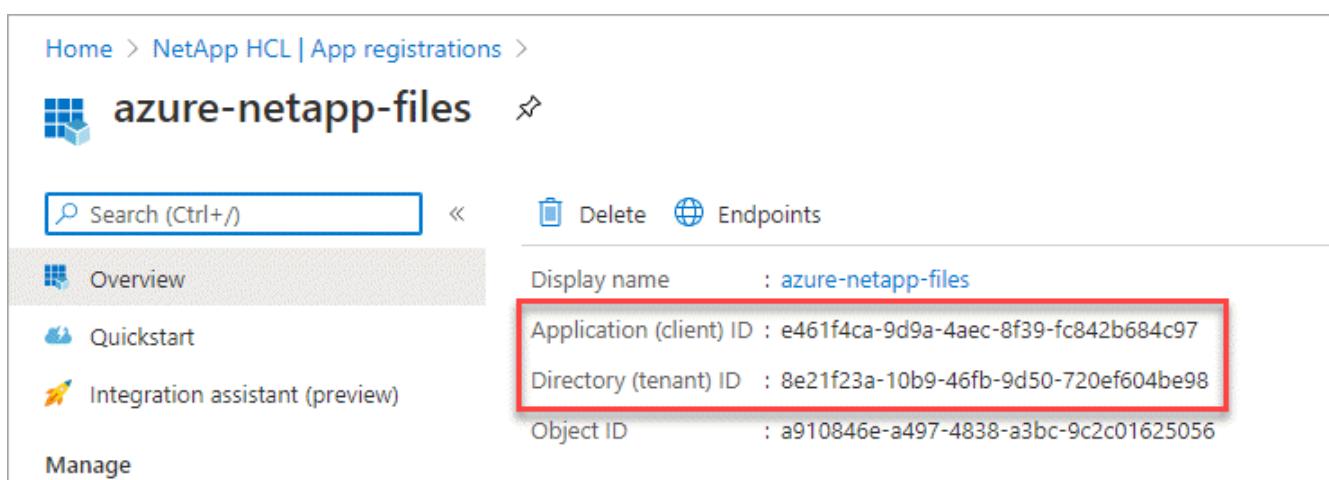
- a. Click **New registration**.

- b. Specify details about the application:

- **Name:** Enter a name for the application.
- **Account type:** Select an account type (any will work with Cloud Manager).
- **Redirect URI:** You can leave this blank.

- c. Click **Register**.

4. Copy the **Application (client) ID** and the **Directory (tenant) ID**.



The screenshot shows the "App registrations" page in the Azure portal. The application "azure-netapp-files" is selected. On the left, there is a navigation menu with "Overview" (selected), "Quickstart", and "Integration assistant (preview)". On the right, there are buttons for "Delete" and "Endpoints". Below the buttons, the application details are listed:
Display name : azure-netapp-files
Application (client) ID : e461f4ca-9d9a-4aec-8f39-fc842b684c97
Directory (tenant) ID : 8e21f23a-10b9-46fb-9d50-720ef604be98
Object ID : a910846e-a497-4838-a3bc-9c2c01625056

When you create the Azure NetApp Files working environment in Cloud Manager, you need to provide the application (client) ID and the directory (tenant) ID for the application. Cloud Manager uses the IDs to programmatically sign in.

5. Create a client secret for the application so Cloud Manager can use it to authenticate with Azure AD:

- a. Click **Certificates & secrets > New client secret**.
- b. Provide a description of the secret and a duration.
- c. Click **Add**.
- d. Copy the value of the client secret.

Description	Expires	Value
Azure NetApp Files	7/30/2022	3gywMgvF1rxtle8jU1po6~... 

Result

Your AD application is now setup and you should have copied the application (client) ID, the directory (tenant) ID, and the value of the client secret. You need to enter this information in Cloud Manager when you add an Azure NetApp Files working environment.

Assigning the app to a role

You must bind the service principal to your Azure subscription and assign it a custom role that has the required permissions.

Steps

1. [Create a custom role in Azure](#).

The following steps describe how to create the role from the Azure portal.

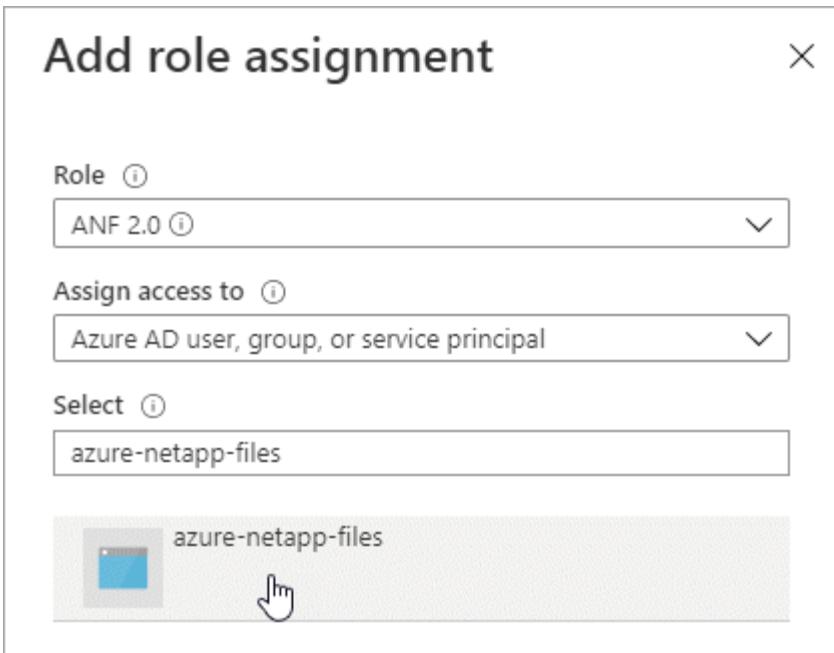
- a. Open the subscription and click **Access control (IAM)**.
- b. Click **Add > Add custom role**.

The screenshot shows the Microsoft Azure portal interface. At the top, there's a navigation bar with 'Microsoft Azure' and a search bar. Below that, the path 'Home > OCCM Dev' is shown. The main title is 'OCCM Dev | Access control (IAM)'. On the left, there's a sidebar with links: 'Overview', 'Activity log', 'Access control (IAM)' (which is highlighted with a blue circle 'a'), 'Tags', 'Diagnose and solve problems', and 'Security'. To the right of the sidebar is a content area with a search bar, a 'Download role assignments' button, and a context menu. The context menu includes 'Add role assignment', 'Add co-administrator (disabled)', 'Add custom role' (with a mouse cursor icon over it), 'View my access', and 'Check access'. A blue circle 'b' is placed over the 'Add custom role' menu item.

- c. In the **Basics** tab, enter a name and description for the role.
- d. Click **JSON** and click **Edit** which appears at the top right of the JSON format.
- e. Add the following permissions under *actions*:

```
"actions": [  
    "Microsoft.NetApp/*",  
    "Microsoft.Resources/resources/read",  
    "Microsoft.Resources/subscriptions/resourceGroups/read",  
  
    "Microsoft.Resources/subscriptions/resourcegroups/resources/read",  
    "Microsoft.Resources/subscriptions/resourceGroups/write",  
    "Microsoft.Network/virtualNetworks/read",  
    "Microsoft.Insights/Metrics/Read"  
],
```

- f. Click **Save**, click **Next**, and then click **Create**.
2. Now assign the application to the role that you just created:
 - a. From the Azure portal, open the subscription and click **Access control (IAM)** > **Add** > **Add role assignment**.
 - b. Select the custom role that you created.
 - c. Keep **Azure AD user, group, or service principal** selected.
 - d. Search for the name of the application (you can't find it in the list by scrolling).



- e. Select the application and click **Save**.

The service principal for Cloud Manager now has the required Azure permissions for that subscription.

Creating an Azure NetApp Files working environment

Set up an Azure NetApp Files working environment in Cloud Manager so you can start creating volumes.

1. From the Canvas page, click **Add Working Environment**.
2. Select **Microsoft Azure** and then **Azure NetApp Files**.
3. Provide details about the AD application that you previously set up.

Azure NetApp Files Credentials

Working Environment Name

ANF

Application (client) ID

e461f4ca-9d9a-4aec-8f39-fc842b684c97

Client Secret

.....

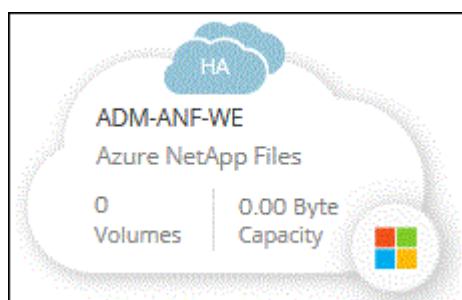
Directory (tenant) ID

8e21f23a-10b9-46fb-9d50-720ef604be98

4. Click **Add**.

Result

You should now have an Azure NetApp Files working environment.



What's next?

[Start creating and managing volumes.](#)

Creating and managing volumes for Azure NetApp Files

After you set up your working environment, you can create and manage Azure NetApp Files accounts, capacity pools, volumes, and snapshots.

Creating volumes

You can create NFS or SMB volumes in a new or existing Azure NetApp Files account.

A Cloud Manager feature called "templates" enables you to create volumes that are optimized for the workload requirements for certain applications; such as databases or streaming services. If your organization has created volume templates that you should use, follow [these steps](#).

Before you begin

- If you want to use SMB, you must have set up DNS and Active Directory.
- When planning to create an SMB volume, you must have a Windows Active Directory server available to which you can connect. You will enter this information when creating the volume.

Steps

1. Open the Azure NetApp Files working environment.

2. Click **Add New Volume**.

3. Provide the required information on each page:

- **Azure NetApp Files Account:** Choose an existing Azure NetApp Files account or create a new account. When creating a new account you can also choose the Resource Group that you want to use.

- **Capacity Pool:** Select an existing capacity pool or create a new capacity pool.

If you create a new capacity pool, you need to specify a size and select a [service level](#).

The minimum size for the capacity pool is 4 TB. You can specify a size in multiples of 4 TB.

- **Details & Tags:** Enter a volume name and size, the VNet and subnet where the volume should reside, and optionally specify tags for the volume.

- **Protocol:** Choose the NFS or SMB protocol and enter the required information.

Here's an example of details for NFS.

Protocol

Select the volume's protocol: **NFS Protocol** SMB Protocol

Protocol	Export Policy
Volume Path vol1	Allowed Client & Access 192.168.1.22/24 <input checked="" type="radio"/> Read & Write <input type="radio"/> Read Only X
Select NFS Version: <input checked="" type="checkbox"/> NFSv3 <input type="checkbox"/> NFSv4.1	192.168.1.22/24 <input checked="" type="radio"/> Read & Write <input type="radio"/> Read Only X
+ Add Export Policy Rule (Up to 5)	

Here's an example of details for SMB. You'll need to provide Active Directory information on the next page when you set up your first SMB volume.

The screenshot shows a 'Protocol' configuration screen. At the top, there is a heading 'Protocol'. Below it, a section titled 'Select the volume's protocol:' contains two radio buttons: 'NFS Protocol' (unchecked) and 'SMB Protocol' (checked). A 'Share Name' field below has the value 'vol1' entered into it.

4. If you want this volume to be created based on a snapshot of an existing volume, select the snapshot from the Snapshot Name drop-down list.
5. Click **Add Volume**.

Result

The new volume is added to the working environment.

Continue with [mounting the cloud volume](#).

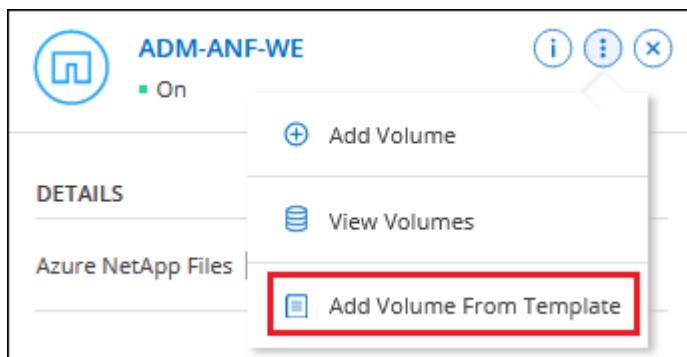
Creating volumes from templates

If your organization has created ANF volume templates so you can deploy volumes that are optimized for the workload requirements for certain applications, follow the steps in this section.

The template should make your job easier because certain volume parameters will already be defined in the template, such as capacity pool, size, protocol, VNet and subnet where the volume should reside, and more. When a parameter is already predefined, you can just skip to the next volume parameter.

Steps

1. On the Canvas page, click the Azure NetApp Files working environment on which you want to provision a volume.
2. Click > **Add Volume From Template**.



3. In the *Select Template* page, select the template that you want to use to create the volume and click **Next**.

The screenshot shows the 'Select Template' page with a header indicating the user is on step 1 of 2. The main area displays a table of templates under the heading '1 Templates'. The first row in the table is highlighted with a blue background, representing the selected template: 'Azure ANF volume for staging'. The table columns include 'Template Name', 'Template Description', 'Created by', 'Last Modified', and 'Parameters'. A search bar and a 'View' link are also visible at the top right.

The *Define Parameters* page is displayed.

The screenshot shows the 'Define Parameters' page. At the top, a message says 'Enter your values for the actions. Parameters that are locked by the template are not editable.' Below this, there are sections for 'Actions', 'Volume Details', 'Tags', and 'Protocol'.

- Actions:** A button labeled 'Create Volume in Azure NetApp Files (1)' is shown.
- Volume Details:**
 - Volume Name:** A field with placeholder text 'Volume Name should start with "staging"'.
 - Volume Size (GB):** A field with placeholder text 'Minimum value is 160, Maximum value is 185'.
- Tags:** A section with a '+ Add Tags' button.
- Protocol:** A section where 'NFSv3' is selected as the protocol.
- Volume Path:** A field containing the path 'volPatsdscwwwq'.
- Export Policy Rules (up to 5):** A section for defining export policy rules.

A checkbox 'Show read-only parameters' is located at the top right of the page.

Note: You can click the checkbox **Show read-only parameters** to show all the fields that have been locked by the template if you want to see the values for those parameters. By default these predefined fields are hidden and only the fields you need to complete are shown.

- Add values for all of the parameters that are not hard-coded from the template. See [creating volumes](#) for details about all the parameters you need to complete to deploy an ANF volume.
- Click **Run Template** after you have defined all the parameters needed for this volume.

Result

Cloud Manager provisions the volume and displays a page so that you can see the progress.

This process can take a few minutes.
Keep this page open to monitor progress, or you can close this page
and check the [Timeline](#) later for details.

Actions status

Create Volume in Azure NetApp Files	Pending
-------------------------------------	---------

Then the new volume is added to the working environment.

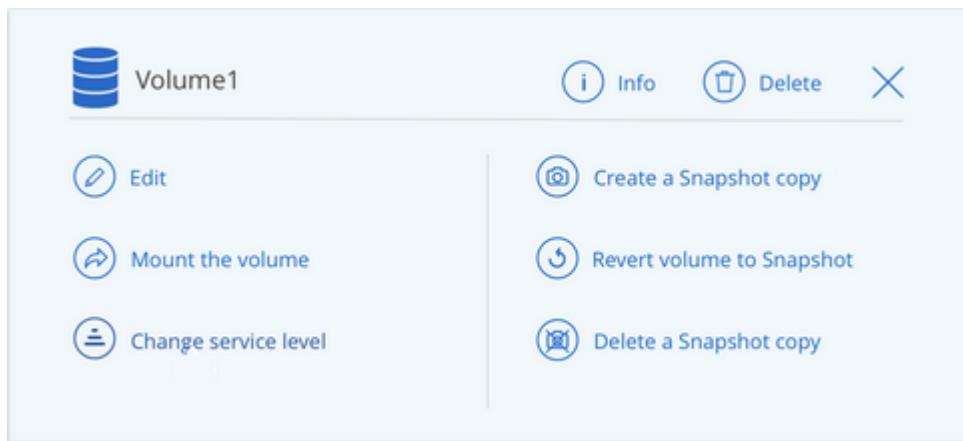
Continue with mounting the cloud volume.

Mounting volumes

Access mounting instructions from within Cloud Manager so you can mount the volume to a host.

Steps

1. Open the working environment.
2. Hover over the volume and select **Mount the volume**.



3. Follow the instructions to mount the volume.

Editing a volume's size and tags

After you create a volume, you can modify its size and tags at any time.

Steps

1. Open the working environment.
2. Hover over the volume and select **Edit**.
3. Modify the size and tags as needed.
4. Click **Apply**.

Changing the volume's service level

After you create a volume, you can change the service level at any time as long as the destination capacity pool already exists.

Steps

1. Open the working environment.
2. Hover over the volume and select **Change service level**.
3. Select the capacity pool that provides the service level that you want.
4. Click **Change**.

Result

The volume is moved to the other capacity pool with no impact to the volume.

Managing Snapshot copies

Snapshot copies provide a point-in-time copy of your volume. Create Snapshot copies, restore the data to a new volume, and delete Snapshot copies.

Steps

1. Open the working environment.
2. Hover over the volume and choose one of the available options to manage Snapshot copies:
 - **Create a Snapshot copy**
 - **Revert volume to Snapshot**
 - **Delete a Snapshot copy**
3. Follow the prompts to complete the selected action.

Deleting volumes

Delete the volumes that you no longer need.

Steps

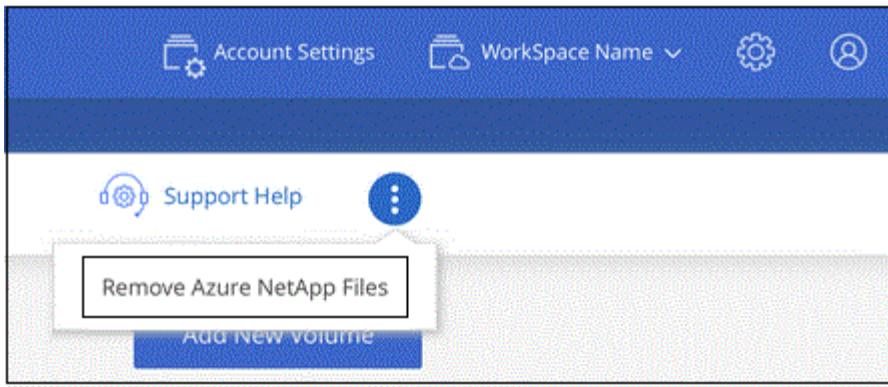
1. Open the working environment.
2. Hover over the volume and click **Delete**.
3. Confirm that you want to delete the volume.

Removing Azure NetApp Files

This action removes Azure NetApp Files from Cloud Manager. It doesn't delete your Azure NetApp Files account or volumes. You can add Azure NetApp Files back to Cloud Manager at any time.

Steps

1. Open the Azure NetApp Files working environment.
2. At the top right of the page, select the actions menu and click **Remove Azure NetApp Files**.



3. Click **Remove** to confirm.

Amazon FSx for ONTAP

Learn about Amazon FSx for ONTAP

Amazon FSx for ONTAP is a fully managed service allowing customers to launch and run file systems powered by NetApp's ONTAP storage operating system. FSx for ONTAP provides the same features, performance, and administrative capabilities NetApp customers use on premises, with the simplicity, agility, security, and scalability of a native AWS service.

Features

- No need to configure or manage storage devices, software, or backups.
- Support for CIFS, NFSv3, NFSv4.0, NFSv4.1, and SMB v2.0 - v3.1.1 protocols.
- Low cost, virtually unlimited data storage capacity using available Infrequently Accessed (IA) storage tier.
- Certified to run on latency-sensitive applications including Oracle RAC.
- Choice of bundled and pay-as-you-go pricing.

Additional features in Cloud Manager

- Using a Connector in AWS and Cloud Manager, you can create and manage volumes, replicate data, and integrate FSx for ONTAP with NetApp cloud services, such as Data Sense and Cloud Sync.
- Using Artificial Intelligence (AI) driven technology, Cloud Data Sense can help you understand data context and identify sensitive data that resides in your FSx for ONTAP accounts. [Learn more](#).
- Using NetApp Cloud Sync, you can automate data migration to any target in the cloud or on premises. [Learn more](#)

Cost

Your FSx for ONTAP account is maintained by AWS and not by Cloud Manager. [Amazon FSx for ONTAP getting started guide](#)

There is an additional cost associated with using the Connector in AWS and the optional data services such as Cloud Sync and Data Sense.

Supported regions

[View supported Amazon regions.](#)

Getting help

Amazon FSx for ONTAP is an AWS first-party solution. For questions or technical support issues associated with your AWS FSx file system, infrastructure or any AWS solution using this service, use the Support Center in your AWS console to open a support case to AWS. Select the “FSx for ONTAP” service and appropriate category. Provide the remaining information required to create your AWS support case.

For general questions specific to Cloud Manager or Cloud Manager micro-services, you can start with the in-line Cloud Manager chat.

For technical support issues specific to Cloud Manager or micro-services within, you can open a NetApp support ticket using your Cloud Manager account level serial number. You will need to register your Cloud Manager serial number to activate support.

Limitations

- Cloud Manager can replicate data only from on-premises or Cloud Volumes ONTAP to FSx for ONTAP.
- At this time iSCSI volumes can be created using the ONTAP CLI or ONTAP API.

Get started with Amazon FSx for ONTAP

Get started with Amazon FSx for ONTAP in a few steps.

You can get started with FSx for ONTAP in just a few steps.



Create an FSx for ONTAP working environment

You must create an Amazon FSx for ONTAP working environment before adding volumes. You will need an AWS access key and secret key for an [IAM user with FSx for ONTAP permissions](#).



Create a Connector

You must have a [Connector for AWS](#) to open the FSx for ONTAP working environment, create volumes, or perform other actions. When a Connector is required, Cloud Manager will prompt you if one is not already added.



Add volumes

You can create FSx for ONTAP volumes using Cloud Manager.



Manage your volumes

Use Cloud Manager to manage your volumes and configure additional services such as replication, Cloud

Sync, and Data Sense.

Related links

- [Creating a Connector from Cloud Manager](#)
- [Launching a Connector from the AWS Marketplace](#)
- [Installing the Connector software on a Linux host](#)

Set up permissions for FSx for ONTAP

To create or manage your Amazon FSx for ONTAP working environment, you need an AWS access key and secret key for an IAM user role with FSx for ONTAP permissions. These permissions are different from the permissions required to create a Connector in AWS.

 You can create a new IAM user role with FSx for ONTAP permissions or edit an existing IAM user role to include the additional FSx for ONTAP permissions. We recommend the latter to avoid having to use multiple keys for your Connector and for FSx for ONTAP access.

Steps

To grant FSx for ONTAP permissions to an IAM user role:

1. From the AWS IAM console, create a new policy or edit an existing policy to include the following actions for FSx for ONTAP:

```
"ec2:Describe*"  
"kms:Describe*"  
"kms>List*"  
"fsx:*"  
"iam>CreateServiceLinkedRole"  
"ec2>CreateTags"
```

[AWS Documentation: Creating IAM Policies](#)

2. Attach the policy you created in the previous step to the IAM user role.

[AWS Documentation: Creating IAM Roles](#)

[AWS Documentation: Adding and Removing IAM Policies](#)

Result

The AWS user now has permissions required for FSx for ONTAP in Cloud Manager.

Related links

- [AWS credentials and permissions](#)
- [Create a Connector in AWS](#)
- [Managing AWS credentials for Cloud Manager](#)

- What Cloud Manager does with AWS permissions

Create an Amazon FSx for ONTAP working environment

Using Cloud Manager you can create an FSx for ONTAP working environment to add volumes and manage additional data services.

Create an Amazon FSx for ONTAP working environment

The first step is to create an FSx for ONTAP working environment. If you already created an FSx for ONTAP file system in the AWS Management Console, you can [discover it using Cloud Manager](#).

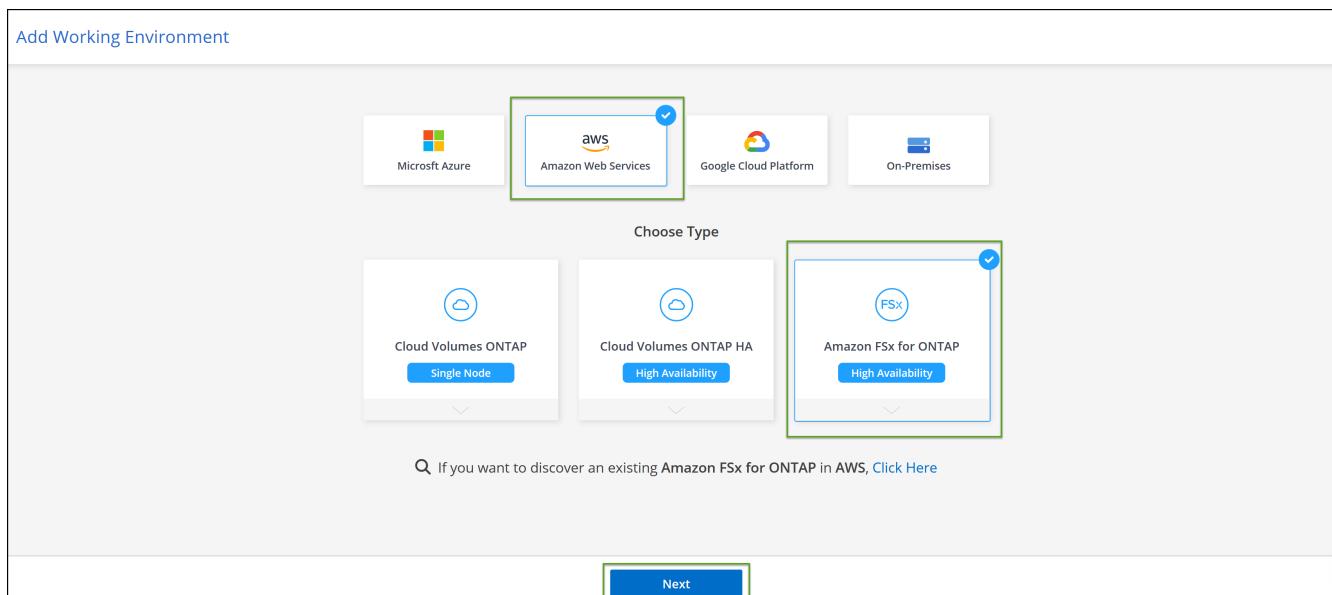
Before you begin

Before creating your FSx for ONTAP working environment in Cloud Manager, you will need:

- An AWS access key and secret key for an IAM user with the [required FSx for ONTAP permissions](#).
- The region and VPN information for where you will create the FSx for ONTAP instance.

Steps

1. In Cloud Manager, add a new Working Environment, select the location **Amazon Web Services**, and click **Next**.
2. Select **Amazon FSx for ONTAP** and click **Next**.



3. You can select existing FSx for ONTAP credentials or create new credentials using your AWS access key and secret key. Click to verify your IAM user policy adheres to [FSx for ONTAP requirements](#).
4. Provide information about your FSx for ONTAP instance:
 - a. Enter the working environment name you want to use.
 - b. Optionally, you can create tags by clicking the plus sign and entering a tag name and value.
 - c. Enter and confirm the ONTAP Cluster password you want to use.
 - d. Select the option to use the same password for your SVM user or set a different password.
 - e. Click **Next**.

Add FSx for ONTAP

Details and Credentials

Details	Credentials
Working Environment Name <input type="text" value="myfsxenvironment"/>	User Name <input type="text" value="fsxadmin"/>
Tags + Add Tags	<i>Optional</i>
	ONTAP Cluster Password <input type="password" value="*****"/>
	Confirm ONTAP Cluster Password <input type="password" value="*****"/>
	<input checked="" type="checkbox"/> Use the same password for SVM user (vsadmin)

[Previous](#) [Next](#)

5. Provide region and VPC information:

- a. Select a region and VPC with subnets in at least two Availability Zones so each node is in a dedicated Availability Zone.
- b. Accept the default security group or select a different one. [AWS security groups](#) control inbound and outbound traffic. These are configured by your AWS admin and are associated with your [AWS elastic network interface \(ENI\)](#).
- c. Select an Availability Zone and subnet for each node.
- d. Click **Next**.

Add FSx for ONTAP

Region and VPC

Region <input type="text" value="us-east-2 US East (Ohio)"/>	VPC <input type="text" value="VPC4QA - 10.0.0.0/16"/>	Security Group <input type="text" value="Default security group"/>
Node 1	Node 2	
Availability Zone <input type="text" value="us-east-2b"/>	Availability Zone <input type="text" value="us-east-2c"/>	
Subnet <input type="text" value="10.0.4.0/24"/>	Subnet <input type="text" value="10.0.3.0/24"/>	

[Previous](#) [Next](#)

6. Leave *CIDR Range* empty and click **Next** to automatically set an available range. Optionally, you can use [AWS Transit Gateway](#) to manually configure a range.

Add FSx for ONTAP

Floating IP

Floating IP addresses are required for cluster and SVM access and for NFS and CIFS data access.

Floating IPs can migrate between HA nodes if failures occur. To access the data from outside the VPC, you can set up an [AWS transit gateway](#).

CIDR Range	<i>Optional</i>
Example: 10.10.10.10/24	

Notice: You must specify a CIDR block that is outside of the CIDR blocks for all VPCs in the selected AWS region.

[Previous](#) [Next](#)

- Select route tables that include routes to the floating IP addresses. If you have just one route table for the subnets in your VPC (the main route table), Cloud Manager automatically adds the floating IP addresses to that route table. Click **Next** to continue.

Add FSx for ONTAP

Route Tables

Select the route tables that should include routes to the floating IP addresses. This enables client access to volumes. Clients associated with unselected route tables won't have access to volumes.

[Learn More](#)

2 Route table

<input type="checkbox"/>	Name	Main	ID	Associate with Subnets	Tags	
<input checked="" type="checkbox"/>	VPC4QA	Yes	rtb-0880ec9d aeb55d630	2 Subnets	2	▼
<input type="checkbox"/>	No tag name	No	rtb-0e0c7d9e a4cf05d66	1 Subnet	1	▼

Notice: The main route table is the default for the VPC

[Previous](#) [Next](#)

- Accept the default AWS master key or click **Change Key** to select a different AWS Customer Master Key (CMK). For more information on CMK, see [Setting up the AWS KMS](#). Click **Next** to continue.

Add FSx for ONTAP

Data Encryption

AWS Managed Encryption

AWS is responsible for data encryption and decryption operations. Key management is handled by AWS key management services.

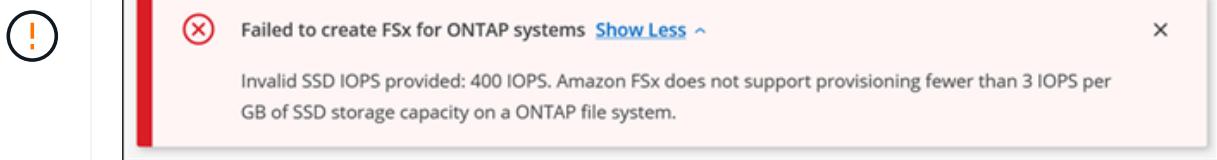
Default Master Key: aws/fsx

[Previous](#) [Next](#)

9. Configure your storage:

- Select the throughput, capacity, and unit.
- You can optionally specify an IOPS value. If you don't specify an IOPS value, Cloud Manager will set a default value based on 3 IOPS per GiB of the total capacity entered. For example, if you enter 2000 GiB for the total capacity and no value for the IOPS, the effective IOPS value will be set to 6000.

If you specify an IOPS value that does not meet the minimum requirements, you'll receive an error when adding the working environment.



- Click **Next**.

Add FSx for ONTAP

Storage Configuration

SSD Disk Properties

Throughput	Optional	Capacity	Unit
<input type="text" value="512 MBps"/>		<input type="text" value="3"/>	<input type="text" value="TiB"/>
IOPS Value	<input type="text" value="400"/>	Notice: The current version of FSx does not allow changing the capacity after creation. Also, note that the capacity drives the cost of the service.	

[Previous](#) [Next](#)

10. Review your configuration:
- Click the tabs to review your ONTAP properties, provider properties, and networking configuration.
 - Click **Previous** to make changes to any settings.
 - Click **Add** to accept the settings and create your Working Environment.

Review

 myfsxenvironment
FSx for ONTAP | HA | Multiple AZs

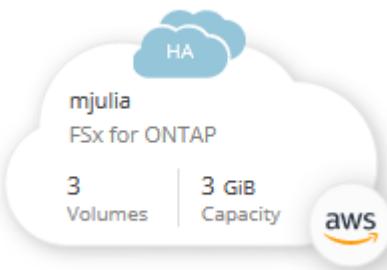
Overview

ONTAP Properties		Provider Properties	Networking
HA Deployment Model	Multiple Availability Zone		
Capacity	3 TiB		
Throughput	512 MBps		

[Previous](#) [Add](#)

Result

Cloud Manager displays your FSx for ONTAP configuration on the Canvas page.



You can now add volumes to your FSx for ONTAP working environment using Cloud Manager.

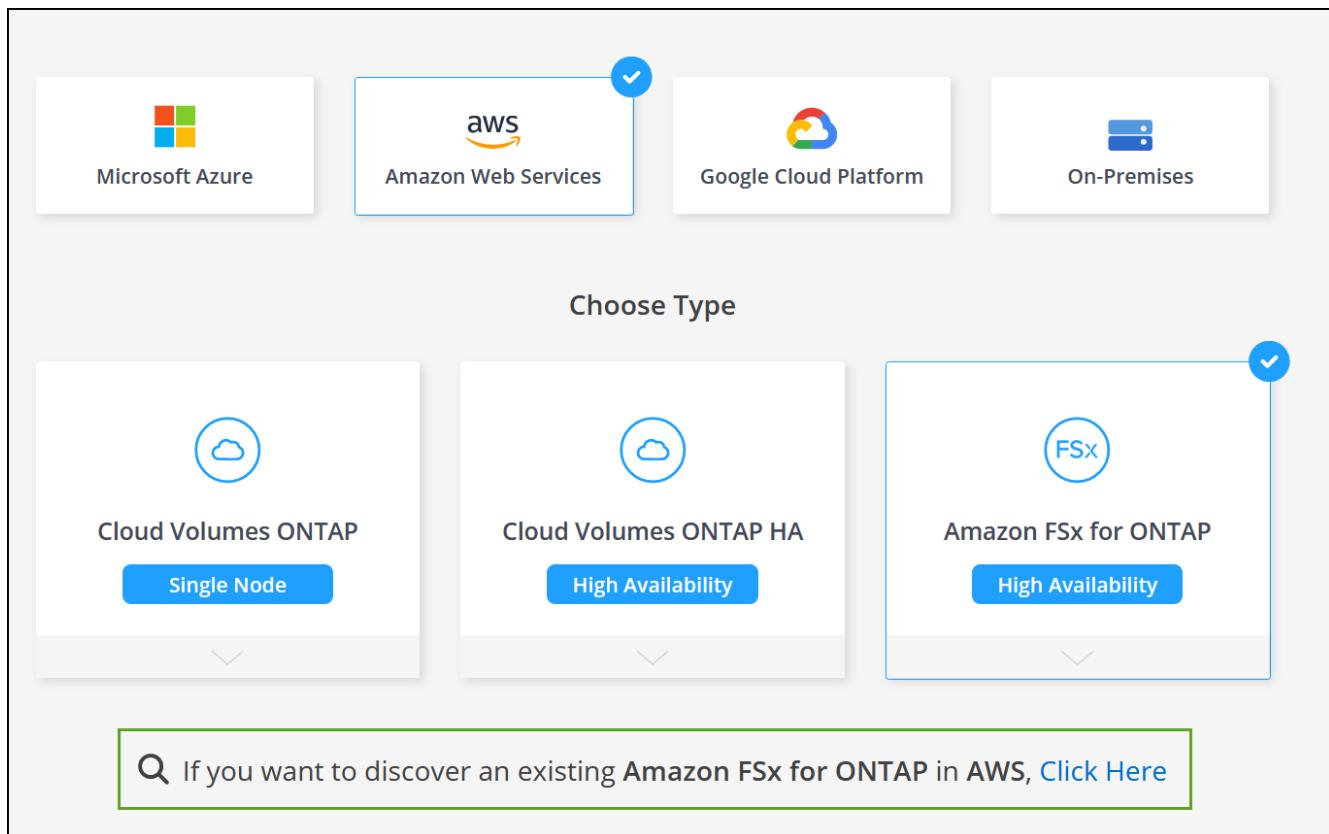
Discover an existing FSx for ONTAP file system

If you created an FSx for ONTAP file system using the AWS Management Console or if you want to restore a working environment you previously removed, you can discover it using Cloud Manager.

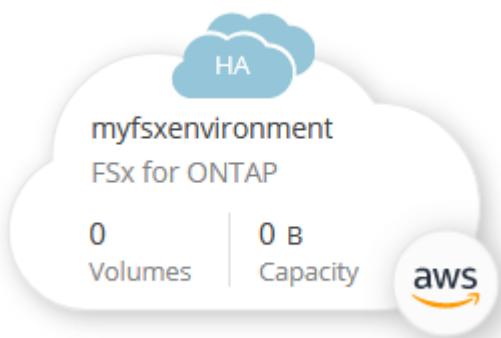
Steps

- In Cloud Manager, click **Add Working Environment**, select **Amazon Web Services**.

2. Select **Amazon FSx for ONTAP** and click **Click Here**.



3. Select existing credentials or create new credentials. Click **Next**.
4. Select the AWS region and the working environment you want to add.



5. Click **Add**.

Result

Cloud Manager displays your discovered FSx for ONTAP file system.

Create and manage volumes for Amazon FSx for ONTAP

After you set up your working environment, you can create and manage FSx for ONTAP volumes, clones, and snapshots, change tiering policies, and remove or delete FSx for ONTAP.

Creating volumes

You can create and edit NFS and CIFS volumes in a new or existing FSx for ONTAP working environment. Volumes created using ONTAP CLI will also be visible in your FSx for ONTAP working environment.

Before you begin

You need:

- An active [Connector in AWS](#).



You do not need a Connector in AWS to remove or delete a working environment.

- If you want to use SMB, you must have set up DNS and Active Directory.

Steps

1. Open the FSx for ONTAP working environment.
2. If you don't have a Connector enabled, you'll be prompted to add one.

A Connector is required to manage your FSx for ONTAP instance

A Connector enables Cloud Manager to manage resources and processes within your public cloud environment. For FSx for ONTAP, this includes creating volumes, replicating data, managing snapshots, integrating with NetApp's cloud services, and more.

[Add Connector](#)

[Back to Canvas](#)

3. Click the **Volumes** tab
4. Click **Add Volume**.

There are no volumes

To start using FSx for ONTAP, add your first volume.

[Add Volume](#)

5. **Volume Details and Protection:**

- Enter a name for your new volume.
- Enter the volume size. Note that the volume size will grow with usage.
- Select a snapshot policy. By default, a snapshot is taken every hour (keeping the last six copies), every day (keeping the last two copies), and every week (keeping the last two copies).
- Click **Next**.

Volume Details & Protection

Volume Name: myfsxvol

Size (GiB): 3

Snapshot Policy: default

[default policy](#)

6. Protocol: Select the an NFS or CIFS volume protocol.

a. For NFS:

- Select an Access Control policy.
- Select the NFS versions.
- Select a Custom Export Policy. Click the information icon for valid value criteria.

1 Details & Protection 2 Protocol 3 Usage Profile & Tiering Policy 4 Review

Volume Protocol

Select the volume's protocol: NFS Protocol CIFS Protocol

Access Control: Custom_export_policy

Select NFS Version: NFSv3 NFSv4

Custom Export Policy: 10.20.0.0/16

b. For CIFS:

- Enter a Share Name.
- Enter users or groups separated by a semicolon.
- Select the permission level for the volume.



If this is the first CIFS volume for this working environment, you will be prompted to configure CIFS connectivity using an *Active Directory* or *Workgroup* setup.

- If you select an Active Directory setup, you'll need to provide the following configuration information.

Field	Description
DNS Primary IP Address	The IP addresses of the DNS servers that provides name resolution for the CIFS server. The listed DNS server must contain the service location records (SRV) needed to locate the Active Directory LDAP servers and domain controllers for the domain the CIFS server will join.
Active Directory Domain to join	The FQDN of the Active Directory (AD) domain you want the CIFS server to join.
Credentials authorized to join the domain	The name and password of a Windows account with sufficient privileges to add computers to the specified Organizational Unit (OU) within the AD domain.
CIFS server NetBIOS name	A CIFS server name that is unique in the AD domain.
Organizational Unit	The organizational unit within the AD domain to associate with the CIFS server. The default is CN=Computers.
DNS Domain	The DNS domain for the storage virtual machine (SVM). In most cases, the domain is the same as the AD domain.
NTP Server	Select Enable NTP Server Configuration to configure an NTP server using the Active Directory DNS. If you need to configure an NTP server using a different address, then you should use the API. See the Cloud Manager automation docs for details.

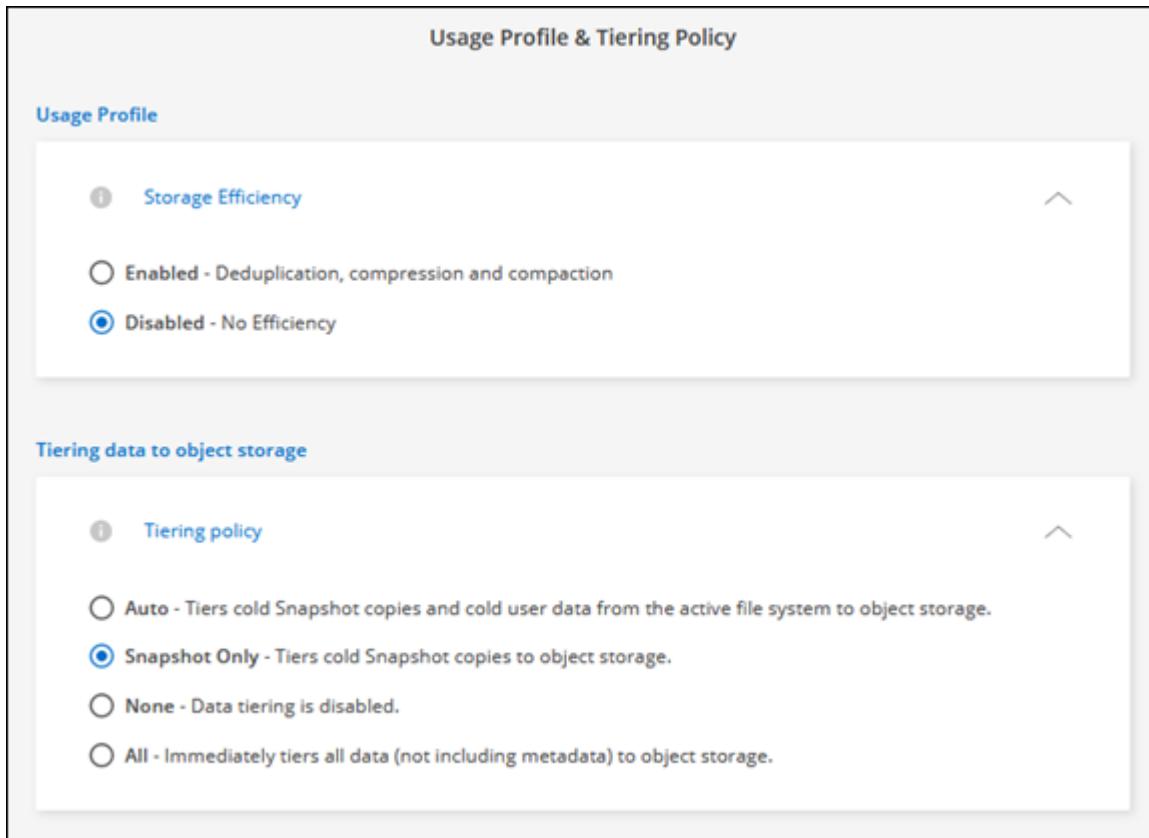
- If you select a Workgroup setup, enter the server and workgroup name for a workgroup configured

for CIFS.

- c. Click **Next**.

7. Usage Profile and Tiering:

- a. By default, **Storage Efficiency** is disabled. You can change this setting to enable deduplication and compression.
- b. By default, **Tiering Policy** is set to **Snapshot Only**. You can select a different tiering policy based on your needs.
- c. Click **Next**.



8. **Review:** Review your volume configuration. Click **Previous** to change settings or click **Add** to create the volume.

Result

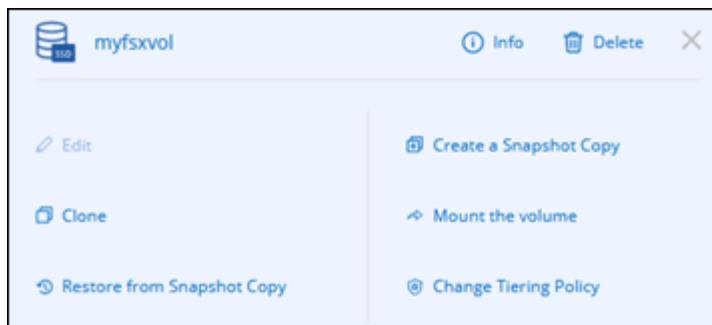
The new volume is added to the working environment.

Mounting volumes

Access mounting instructions from within Cloud Manager so you can mount the volume to a host.

Steps

1. Open the working environment.
2. Open the volume menu and select **Mount the volume**.



3. Follow the instructions to mount the volume.

Editing volumes

After you create a volume, you can modify it at any time.

Steps

1. Open the working environment.
2. Open the volume menu and select **Edit**.
 - a. For NFS, you can modify the size and tags.
 - b. For CIFS, you can modify the share name, users, permissions, and Snapshot policy as needed.
3. Click **Apply**.

Cloning the volume

After you create a volume, you can create a new read-write volume from a new Snapshot.

Steps

1. Open the working environment.
2. Open the volume menu and select **Clone**.
3. Enter a name for the cloned volume.
4. Click **Clone**.

Managing Snapshot copies

Snapshot copies provide a point-in-time copy of your volume. Create Snapshot copies and restore the data to a new volume.

Steps

1. Open the working environment.
2. Open the volume menu and choose one of the available options to manage Snapshot copies:
 - **Create a Snapshot copy**
 - **Restore from a Snapshot copy**
3. Follow the prompts to complete the selected action.

Changing the tiering policy

Change the tiering policy for the volume.

Steps

1. Open the working environment.
2. Open the volume menu and select **Change Tiering policy**.
3. Select a new volume tiering policy and click **Change**.

Replicating and syncing data

You can replicate data between storage environments using Cloud Manager. To configure FSx for ONTAP replication, see [replicating data between systems](#).

You can create sync relationships using Cloud Sync in Cloud Manager. To configure sync relationships, see [create sync relationships](#).

Deleting volumes

Delete the volumes that you no longer need.

Before you begin

You cannot delete a volume that was previously part of a SnapMirror relationship using Cloud Manager. SnapMirror volumes must be deleted using the AWS Management Console or CLI.

Steps

1. Open the working environment.
2. Open the volume menu and select **Delete**.
3. Enter the working environment name and confirm that you want to delete the volume. It can take up to an hour before the volume is completely removed from Cloud Manager.



If you try to delete a cloned volume, you will receive an error.

Removing FSx for ONTAP from the workspace

You can remove FSx for ONTAP from Cloud Manager without deleting your FSx for ONTAP account or volumes. You can add the FSx for ONTAP working environment back to Cloud Manager at any time.

Steps

1. Open the working environment. If you don't have a Connector in AWS, you will see the prompt screen. You can ignore this and proceed with removing the working environment.
2. At the top right of the page, select the actions menu and click **Remove from workspace**.

The screenshot shows the 'Overview' tab of the FSx for ONTAP working environment 'My0xWorkingEnv'. The main area displays 'Capacity Distribution' with values: 80.25 TiB Provisioned, 60.35 TiB EBS Used, and 40.71 TiB S3 Used. Below this, there are two boxes: one showing 57 Volumes and another showing 4 Data Replication Relationships. On the right, the 'INFORMATION' section lists details like ARN, AWS Filesystem ID, Cluster Management IP, AWS Account, and Encryption status. Action buttons at the top right include 'Timeline', 'Edit', 'Remove from workspace', and 'Delete'.

3. Click **Remove** to remove FSx for ONTAP from Cloud Manager.

Deleting the FSx for ONTAP working environment

You can delete the FSx for ONTAP from Cloud Manager.

Before you begin

- You must delete all volumes associated with the file system.
- You cannot delete a working environment that contains failed volumes. Failed volumes must be deleted using the AWS Management Console or CLI prior to deleting FSx for ONTAP files system.



This action will delete all resources associated with the working environment. This action cannot be undone.

Steps

1. Open the working environment. If you don't have a Connector in AWS, you will see the prompt screen. You can ignore this and proceed to deleting the working environment.
2. At the top right of the page, select the actions menu and click **Delete**.

The screenshot shows the same 'Overview' tab for 'My0xWorkingEnv'. The 'INFORMATION' section is identical to the first screenshot. In the top right corner, the 'Delete' button is highlighted with a green border, indicating it is the target of the next step.

3. Enter the name of the working environment and click **Delete**.

Cloud Volumes Service for AWS

Learn about Cloud Volumes Service for AWS

NetApp Cloud Volumes Service for AWS is a cloud native file service that provides NAS volumes over NFS and SMB with all-flash performance. This service enables any workload, including legacy applications, to run in the AWS cloud.



With the launch of [Amazon FSx for ONTAP](#), you can no longer create new CVS for AWS working environments in Cloud Manager. However, if you had previously added CVS for AWS working environments to Cloud Manager, you can continue to create and manage volumes.

Benefits of using Cloud Volumes Service for AWS

Cloud Volumes Service for AWS provides the following benefits:

- Fully managed service, therefore no need to configure or manage storage devices
- Support for NFSv3 and NFSv4.1, and SMB 3.0 and 3.1.1 NAS protocols
- Secure access to Linux and Windows Elastic Container Service (ECS) instances, with support including the following:
 - Amazon Linux 2, Red Hat Enterprise Linux 7.5, SLES 12 SP3, and Ubuntu 16.04 LTS
 - Windows Server 2008 R2, Windows Server 2012 R2, and Windows Server 2016
- Choice of bundled and pay-as-you-go pricing

Cost

Volumes created by the Cloud Volumes Service for AWS are charged based on your subscription to the service, not through Cloud Manager.

There are no charges to discover a Cloud Volumes Service for AWS region or volume from Cloud Manager.

Quick start

Get started quickly by following these steps, or go to the next sections for full details.

1

Verify support for your configuration

You must have set up AWS for Cloud Volumes Service and subscribed to one of the [NetApp Cloud Volumes Service offerings on the AWS Marketplace](#) and have an existing CVS for AWS working environment configured in Cloud Manager to create and manage volumes.

2

Create, mount, and manage cloud volumes

Using an existing CVS for AWS working environment, you can create, mount, and manage cloud volumes for existing CVS for AWS subscriptions using Cloud Manager.

Getting help

Use the Cloud Manager chat for general service questions.

For technical support issues associated with your cloud volumes, use your 20 digit "930" serial number located in the "Support" tab of the Cloud Volumes Service user interface. Use this support ID when opening a web ticket or calling for support. Be sure to activate your Cloud Volumes Service serial number for support from the Cloud Volumes Service user interface. [Those steps are explained here.](#)

Limitations

- Cloud Manager doesn't support data replication between working environments when using Cloud Volumes Service volumes.
- Removing your Cloud Volumes Service for AWS subscription from Cloud Manager isn't supported. You can do this only through the Cloud Volumes Service for AWS interface.

Related links

- [NetApp Cloud Central: Cloud Volumes Service for AWS](#)
- [NetApp Cloud Volumes Service for AWS documentation](#)

Managing Cloud Volumes Service for AWS

Cloud Manager enables you to create cloud volumes based on your [Cloud Volumes Service for AWS](#) subscription. You can also discover cloud volumes that you have already created from the Cloud Volumes Service interface and add them to a working environment.



With the launch of [Amazon FSx for ONTAP](#), you can no longer create new CVS for AWS working environments in Cloud Manager. However, if you had previously added CVS for AWS working environments to Cloud Manager, you can continue to create and manage volumes.

Create cloud volumes

For configurations where volumes already exist in the Cloud Volumes Service working environment you can use these steps to add new volumes.

For configurations where no volumes exist, you can create your first volume directly from Cloud Manager after you have set up your Cloud Volumes Service for AWS subscription. In the past, the first volume had to be created directly in the Cloud Volumes Service user interface.

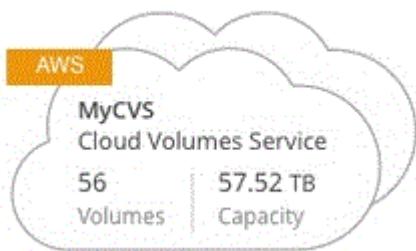
Before you begin

- If you want to use SMB in AWS, you must have set up DNS and Active Directory.
- When planning to create an SMB volume, you must have a Windows Active Directory server available to which you can connect. You will enter this information when creating the volume. Also, make sure that the Admin user is able to create a machine account in the Organizational unit (OU) path specified.
- You will need this information when creating the first volume in a new region/working environment:
 - AWS account ID: A 12-digit Amazon account identifier with no dashes. To find your account ID, refer to this [AWS topic](#).
 - Classless Inter-Domain Routing (CIDR) Block: An unused IPv4 CIDR block. The network prefix must

range between /16 and /28, and it must also fall within the ranges reserved for private networks (RFC 1918). Do not choose a network that overlaps your VPC CIDR allocations.

Steps

1. Select a CVS for AWS working environment and click **Add New Volume**.



2. If you are adding the first volume to the working environment in the region, you have to add AWS networking information.
 - a. Enter the IPv4 range (CIDR) for the region.
 - b. Enter the 12-digit AWS account ID (with no dashes) to connect your Cloud Volumes account to your AWS account.
 - c. Click **Continue**.

Network Setup

! Your Cloud Volumes Service account isn't connected to your AWS account yet. Enter information about your AWS networking to connect the accounts. For details, see the [Cloud Volumes Service for AWS Account Setup document](#).

CIDR (IPv4)	AWS Account ID
192.168.0.0/28	123456789012345

3. The Accepting Virtual Interfaces page describes some steps you will need to perform after you add the volume so that you are prepared to complete that step. Just click **Continue** again.
4. In the Details & Tags page, enter details about the volume:
 - a. Enter a name for the volume.
 - b. Specify a size within the range of 100 GiB to 90,000 GiB (equivalent to 88 TiBs).
[Learn more about allocated capacity](#).
 - c. Specify a service level: Standard, Premium, or Extreme.
[Learn more about service levels](#).
 - d. Enter one or more tag names to categorize the volume if you want.
 - e. Click **Continue**.

5. In the Protocol page, select NFS, SMB, or Dual Protocol and then define the details. Required entries for NFS and SMB are shown in separate sections below.
6. In the Volume Path field, specify the name of the volume export you will see when you mount the volume.
7. If you select Dual-protocol you can select the security style by selecting NTFS or UNIX. Security styles affect the file permission type used and how permissions can be modified.
 - UNIX uses NFSv3 mode bits, and only NFS clients can modify permissions.
 - NTFS uses NTFS ACLs, and only SMB clients can modify permissions.
8. For NFS:
 - a. In the NFS Version field, select NFSv3, NFSv4.1, or both depending on your requirements.
 - b. Optionally, you can create an export policy to identify the clients that can access the volume. Specify the:
 - Allowed clients by using an IP address or Classless Inter-Domain Routing (CIDR).
 - Access rights as Read & Write or Read Only.
 - Access protocol (or protocols if the volume allows both NFSv3 and NFSv4.1 access) used for users.
 - Click **+ Add Export Policy Rule** if you want to define additional export policy rules.

The following image shows the Volume page filled out for the NFS protocol:

Protocol

Select the volume's protocol: NFS Protocol SMB Protocol Dual Protocol

Export Policy

Volume Path	Allowed Client & Access
vol1	192.168.1.2/24 <input checked="" type="radio"/> Read & Write <input type="radio"/> Read Only
Select NFS Version:	Select NFS Version: <input checked="" type="checkbox"/> NFSv3 <input type="checkbox"/> NFSv4.1
	192.168.1.22/24 <input checked="" type="radio"/> Read & Write <input type="radio"/> Read Only
	Select NFS Version: <input type="checkbox"/> NFSv3 <input checked="" type="checkbox"/> NFSv4.1

9. For SMB:
 - a. You can enable SMB session encryption by checking the box for SMB Protocol Encryption.
 - b. You can integrate the volume with an existing Windows Active Directory server by completing the fields in the Active directory section:

Field	Description
DNS Primary IP Address	The IP addresses of the DNS servers that provide name resolution for the SMB server. Use a comma to separate the IP addresses when referencing multiple servers, for example, 172.31.25.223, 172.31.2.74..
Active Directory Domain to join	The FQDN of the Active Directory (AD) domain that you want the SMB server to join. When using AWS Managed Microsoft AD, use the value from the "Directory DNS name" field.
SMB Server NetBIOS name	A NetBIOS name for the SMB server that will be created.
Credentials authorized to join the domain	The name and password of a Windows account with sufficient privileges to add computers to the specified Organizational Unit (OU) within the AD domain.
Organizational Unit	The organizational unit within the AD domain to associate with the SMB server. The default is CN=Computers for connections to your own Windows Active Directory server. If you configure AWS Managed Microsoft AD as the AD server for the Cloud Volumes Service, you should enter OU=Computers,OU=corp in this field.

The following image shows the Volume page filled out for the SMB protocol:

SMB Connectivity Setup

DNS Primary IP Address 127.0.0.1	User Name administrator
Active Directory Domain to Join yourdomain.com up to 107 characters	Password
SMB Server NetBIOS Name WENName	Organizational Unit CN=Computers



You should follow the guidance on AWS security group settings to enable cloud volumes to integrate with Windows Active Directory servers correctly. See [AWS security group settings for Windows AD servers](#) for more information.

10. In the Volume from Snapshot page, if you want this volume to be created based on a snapshot of an existing volume, select the snapshot from the Snapshot Name drop-down list.
11. In the Snapshot Policy page, you can enable Cloud Volumes Service to create snapshot copies of your volumes based on a schedule. You can do this now or edit the volume later to define the snapshot policy.

See [Creating a snapshot policy](#) for more information about snapshot functionality.

12. Click **Add Volume**.

The new volume is added to the working environment.

After you finish

If this is the first volume created in this AWS subscription, you need to launch the AWS Management Console to accept the two virtual interface that will be used in this AWS region to connect all your cloud volumes. See the [NetApp Cloud Volumes Service for AWS Account Setup Guide](#) for details.

You must accept the interfaces within 10 minutes after clicking the **Add Volume** button or the system may time out. If this happens, email cvs-support@netapp.com with your AWS Customer ID and NetApp Serial Number. Support will fix the issue and you can restart the onboarding process.

Then continue with [Mounting the cloud volume](#).

Mount the cloud volume

You can mount a cloud volume to your AWS instance. Cloud volumes currently support NFSv3 and NFSv4.1 for Linux and UNIX clients, and SMB 3.0 and 3.1.1 for Windows clients.

Note: Please use the highlighted protocol/dialect supported by your client.

Steps

1. Open the working environment.
2. Hover over the volume and click **Mount the volume**.

NFS and SMB volumes display mount instructions for that protocol. Dual-protocol volumes provide both sets of instructions.

3. Hover over the commands and copy them to your clipboard to make this process easier. Just add the destination directory/mount point at the end of the command.

NFS example:

Mount the volume - testk

Setting up your instance

1. Open an SSH client and connect to your instance.

2. Install the nfs client on your instance.

On Red Hat Enterprise Linux or SuSE Linux instance:

```
$ sudo yum install -y nfs-utils
```



On an Ubuntu or Debian instance:

```
$ sudo apt-get install nfs-common
```



Mounting your volume

1. Create a new directory on your instance:

```
$ sudo mkdir /dir
```



2. Mount your NFSv3 volume using the command below:

```
sudo mount -t nfs -o rw,hard,rsize=65536,wsize=65536,vers=3,tc...
```



3. Mount your NFSv4.1 volume using the command below:

```
sudo mount -t nfs -o rw,hard,rsize=65536,wsize=65536,vers=4,1,t...
```



The maximum I/O size defined by the `rsize` and `wsize` options is 1048576, however 65536 is the recommended default for most use cases.

Note that Linux clients will default to NFSv4.1 unless the version is specified with the `vers=<nfs_version>` option.

SMB example:

Mount the volume - <Volume Name>

Mapping your network drive

1. Click the Start button and then click on Computer.
2. Click Map Network Drive.
3. In the Drive list, click any available drive letter.
4. In the Folder box, type this:

```
\test.cv-pm.local\silly-condescending-mcnulty
```



To connect every time you log on to your computer, check the **Reconnect at logon** option.

5. Click Finish.

4. Connect to your Amazon Elastic Compute Cloud (EC2) instance by using an SSH or RDP client, and then follow the mount instructions for your instance.

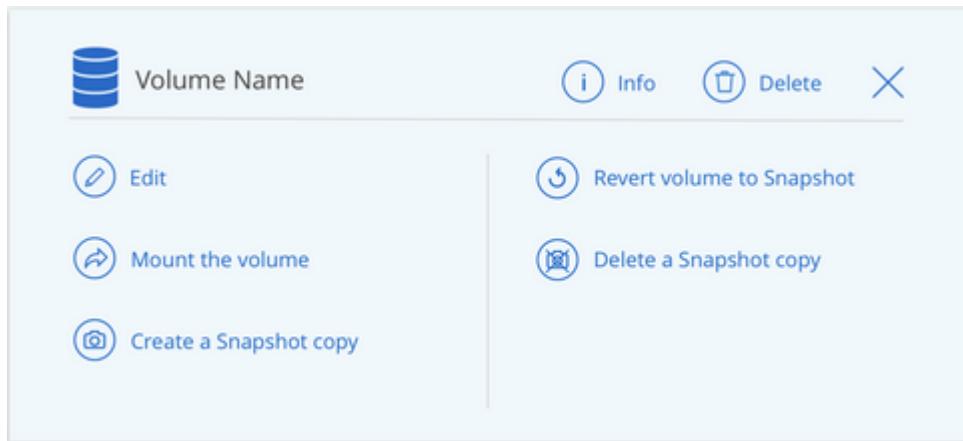
After completing the steps in the mount instructions, you have successfully mounted the cloud volume to your AWS instance.

Managing existing volumes

You can manage existing volumes as your storage needs change. You can view, edit, restore, and delete volumes.

Steps

1. Open the working environment.
2. Hover over the volume.



3. Manage your volumes:

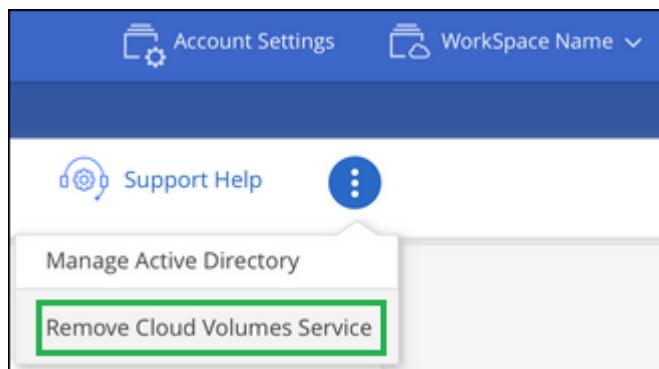
Task	Action
View information about a volume	Select a volume, and then click Info .
Edit a volume (including snapshot policy)	a. Select a volume, and then click Edit . b. Modify the volume's properties and then click Update .
Get the NFS or SMB mount command	a. Select a volume, and then click Mount the volume . b. Click Copy to copy the command(s).
Create a Snapshot copy on demand	a. Select a volume, and then click Create a Snapshot copy . b. Change the snapshot name, if needed, and then click Create .
Replace the volume with the contents of a Snapshot copy	a. Select a volume, and then click Revert volume to Snapshot . b. Select a Snapshot copy and click Revert .
Delete a Snapshot copy	a. Select a volume, and then click Delete a Snapshot copy . b. Select the Snapshot copy you want to delete and click Delete . c. Click Delete again to confirm.
Delete a volume	a. Unmount the volume from all clients: ◦ On Linux clients, use the <code>umount</code> command. ◦ On Windows clients, click Disconnect network drive . b. Select a volume, and then click Delete . c. Click Delete again to confirm.

Remove Cloud Volumes Service from Cloud Manager

You can remove a Cloud Volumes Service for AWS subscription and all existing volumes from Cloud Manager. The volumes are not deleted, they are just removed from the Cloud Manager interface.

Steps

1. Open the working environment.



2. Click the  button at the top of the page and click **Remove Cloud Volumes Service**.
3. In the confirmation dialog box, click **Remove**.

Manage Active Directory configuration

If you change your DNS servers or Active Directory domain, you need to modify the SMB server in Cloud Volumes Services so that it can continue to serve storage to clients.

You can also delete the link to an Active Directory if you no longer need it.

Steps

1. Open the working environment.
2. Click the  button at the top of the page and click **Manage Active Directory**.
3. If no Active Directory is configured, you can add one now. If one is configured, you can modify the settings or delete it using the  button.
4. Specify the settings for the Active Directory that you want to join:

Field	Description
DNS Primary IP Address	The IP addresses of the DNS servers that provide name resolution for the SMB server. Use a comma to separate the IP addresses when referencing multiple servers, for example, 172.31.25.223, 172.31.2.74.
Active Directory Domain to join	The FQDN of the Active Directory (AD) domain that you want the SMB server to join. When using AWS Managed Microsoft AD, use the value from the "Directory DNS name" field.
SMB Server NetBIOS name	A NetBIOS name for the SMB server that will be created.
Credentials authorized to join the domain	The name and password of a Windows account with sufficient privileges to add computers to the specified Organizational Unit (OU) within the AD domain.
Organizational Unit	The organizational unit within the AD domain to associate with the SMB server. The default is CN=Computers for connections to your own Windows Active Directory server. If you configure AWS Managed Microsoft AD as the AD server for the Cloud Volumes Service, you should enter OU=Computers,OU=corp in this field.

5. Click **Save** to save your settings.

Manage cloud volumes snapshots

You can create a snapshot policy for each volume so that you can recover or restore the entire contents of a volume from an earlier time. You can also create an on-demand snapshot of a cloud volume when needed.

Create an on-demand snapshot

You can create an on-demand snapshot of a cloud volume if you want to create a snapshot with the current

volume state.

Steps

1. Open the working environment.
2. Hover over the volume and click **Create a snapshot copy**.
3. Enter a name for the snapshot, or use the automatically generated name, and click **Create**.

Create a Snapshot Copy - <Volume Name>

A NetApp Snapshot copy is a read-only, point-in-time image of a volume. The image protects your data with no performance impact and requires minimal storage.

Snapshot Copy Name

manually.2020-05-04_1722

Create

Create or modify a snapshot policy

You can create or modify a snapshot policy as necessary for a cloud volume. You define the snapshot policy from the *Snapshot Policy* tab either when creating a volume or when editing a volume.

Steps

1. Open the working environment.
2. Hover over the volume and click **Edit**.
3. From the *Snapshot Policy* tab, move the enable snapshots slider to the right.
4. Define the schedule for snapshots:
 - a. Select the frequency: **Hourly**, **Daily**, **Weekly**, or **Monthly**
 - b. Select the number of snapshots you want to keep.
 - c. Select the day, hour, and minute when the snapshot should be taken.

Schedule Snapshot Policies:

<input checked="" type="checkbox"/> Hourly	Number of Snapshot to Keep	Minute	
	<input type="text" value="12"/>	<input type="text" value="30"/>	
<input type="checkbox"/> Daily	Number of Snapshot to Keep	Hour Minute	
	<input type="text" value="0"/>	<input type="text" value="0"/> <input type="text" value="0"/>	
<input checked="" type="checkbox"/> Weekly	Number of Snapshot to Keep	Days	Hour Minute
	<input type="text" value="3"/>	<input type="button" value="Sunday X"/>	<input type="text" value="0"/> <input type="text" value="0"/>
<input type="checkbox"/> Monthly	Number of Snapshot to Keep	Days	Hour Minute
	<input type="text" value="0"/>	<input type="checkbox"/> Sunday <input type="checkbox"/> Monday <input type="checkbox"/> Tuesday	<input type="text" value="0"/> <input type="text" value="0"/>

5. Click **Add volume** or **Update volume** to save your policy settings.

Disable a snapshot policy

You can disable a snapshot policy to stop snapshots from being created for a short period of time while retaining your snapshot policy settings.

Steps

1. Open the working environment.
2. Hover over the volume and click **Edit**.
3. From the *Snapshot Policy* tab, move the enable snapshots slider to the left.

Enable automatic Snapshot copies



When disabled, Cloud Volumes Service does not create Snapshot copies of your volumes.

4. Click **Update volume**.

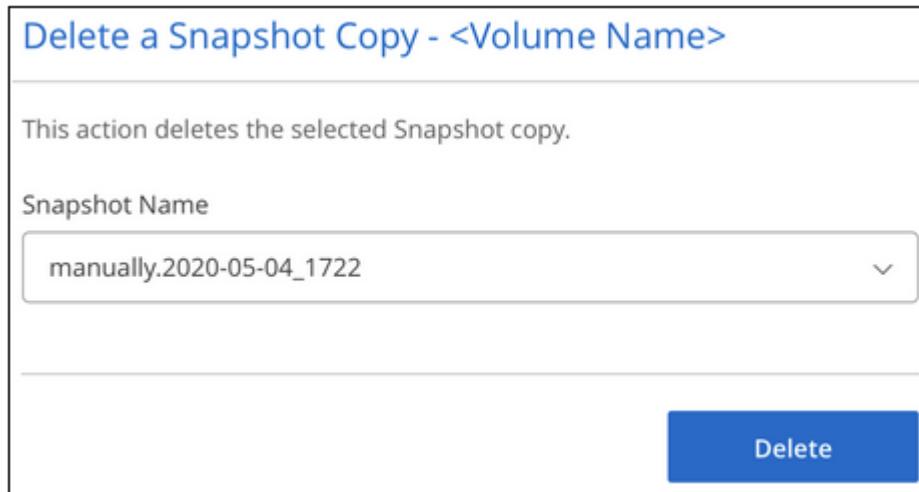
When you want to re-enable the snapshot policy, move the enable snapshots slider to the right and click **Update volume**.

Delete a snapshot

You can delete a snapshot from the Volumes page.

Steps

1. Open the working environment.
2. Hover over the volume and click **Delete a Snapshot copy**.
3. Select the snapshot from the drop-down list and click **Delete**.



4. In the confirmation dialog box, click **Delete**.

Revert a volume from a snapshot

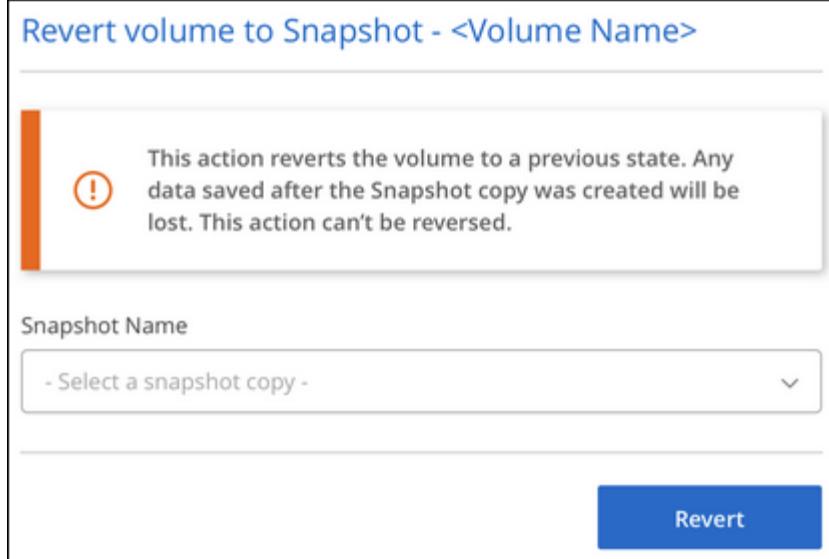
You can revert a volume to an earlier point in time from an existing snapshot.

When you revert a volume, the content of the snapshot overwrites the existing volume configuration. Any changes that were made to the data in the volume after the snapshot was created are lost.

Note that clients do not need to remount the volume after the revert operation.

Steps

1. Open the working environment.
2. Hover over the volume and click **Revert volume to Snapshot**.
3. Select the snapshot that you want to use to restore the existing volume from the drop-down list and click **Revert**.



Reference

Service levels and allocated capacity

The cost for Cloud Volumes Service for AWS is based on the *service level* and the *allocated capacity* that you select. Selecting the appropriate service level and capacity helps you meet your storage needs at the lowest cost.

Considerations

Storage needs include two fundamental aspects:

- The storage *capacity* for holding data
- The storage *bandwidth* for interacting with data

If you consume more storage space than the capacity you selected for the volume, the following considerations apply:

- You will be billed for the additional storage capacity that you consume at the price defined by your service level.
- The amount of storage bandwidth available to the volume does not increase until you increase the allocated capacity size or change the service level.

Service levels

Cloud Volumes Service for AWS supports three service levels. You specify your service level when you create or modify the volume.

The service levels are catered to different storage capacity and storage bandwidth needs:

- **Standard** (capacity)

If you want capacity at the lowest cost, and your bandwidth needs are limited, then the Standard service level might be most appropriate for you. An example is using the volume as a backup target.

- Bandwidth: 16 KB of bandwidth per GB provisioned capacity
- **Premium** (a balance of capacity and performance)

If your application has a balanced need for storage capacity and bandwidth, then the Premium service level might be most appropriate for you. This level is less expensive per MB/s than the Standard service level, and it is also less expensive per GB of storage capacity than the Extreme service level.

- Bandwidth: 64 KB of bandwidth per GB provisioned capacity
- **Extreme** (performance)

The Extreme service level is least expensive in terms of storage bandwidth. If your application demands storage bandwidth without the associated demand for lots of storage capacity, then the Extreme service level might be most appropriate for you.

- Bandwidth: 128 KB of bandwidth per GB provisioned capacity

Allocated capacity

You specify your allocated capacity for the volume when you create or modify the volume.

While you would select your service level based on your general, high-level business needs, you should select your allocated capacity size based on the specific needs of applications, for example:

- How much storage space the applications need
- How much storage bandwidth per second the applications or the users require

Allocated capacity is specified in GBs. A volume's allocated capacity can be set within the range of 100 GB to 100,000 GB (equivalent to 100 TBs).

Number of inodes

Volumes less than or equal to 1 TB can use up to 20 million inodes. The number of inodes increase by 20 million for each TB you allocate, up to a maximum of 100 million inodes.

- <= 1TB = 20 million inodes
- >1 TB to 2 TB = 40 million inodes
- >2 TB to 3 TB = 60 million inodes
- >3 TB to 4 TB = 80 million inodes
- >4 TB to 100 TB = 100 million inodes

Bandwidth

The combination of both the service level and the allocated capacity you select determines the maximum bandwidth for the volume.

If your applications or users need more bandwidth than your selections, you can change the service level or increase the allocated capacity. The changes do not disrupt data access.

Selecting the service level and the allocated capacity

To select the most appropriate service level and allocated capacity for your needs, you need to know how much capacity and bandwidth you require at the peak or the edge.

List of service levels and allocated capacity

The leftmost column indicates the capacity, and the other columns define the MB/s available at each capacity point based on service level.

See [Contract subscription pricing](#) and [Metered subscription pricing](#) for complete details on pricing.

Capacity (TB)	Standard (MB/s)	Premium (MB/s)	Extreme (MB/s)
0.1 (100 GB)	1.6	6.4	12.8
1	16	64	128
2	32	128	256
3	48	192	384
4	64	256	512
5	80	320	640
6	96	384	768
7	112	448	896
8	128	512	1,024
9	144	576	1,152
10	160	640	1,280
11	176	704	1,408
12	192	768	1,536
13	208	832	1,664
14	224	896	1,792
15	240	960	1,920
16	256	1,024	2,048
17	272	1,088	2,176
18	288	1,152	2,304
19	304	1,216	2,432
20	320	1,280	2,560
21	336	1,344	2,688
22	352	1,408	2,816
23	368	1,472	2,944
24	384	1,536	3,072
25	400	1,600	3,200
26	416	1,664	3,328
27	432	1,728	3,456
28	448	1,792	3,584

Capacity (TB)	Standard (MB/s)	Premium (MB/s)	Extreme (MB/s)
29	464	1,856	3,712
30	480	1,920	3,840
31	496	1,984	3,968
32	512	2,048	4,096
33	528	2,112	4,224
34	544	2,176	4,352
35	560	2,240	4,480
36	576	2,304	4,500
37	592	2,368	4,500
38	608	2,432	4,500
39	624	2,496	4,500
40	640	2,560	4,500
41	656	2,624	4,500
42	672	2,688	4,500
43	688	2,752	4,500
44	704	2,816	4,500
45	720	2,880	4,500
46	736	2,944	4,500
47	752	3,008	4,500
48	768	3,072	4,500
49	784	3,136	4,500
50	800	3,200	4,500
51	816	3,264	4,500
52	832	3,328	4,500
53	848	3,392	4,500
54	864	3,456	4,500
55	880	3,520	4,500
56	896	3,584	4,500
57	912	3,648	4,500
58	928	3,712	4,500
59	944	3,776	4,500
60	960	3,840	4,500
61	976	3,904	4,500

Capacity (TB)	Standard (MB/s)	Premium (MB/s)	Extreme (MB/s)
62	992	3,968	4,500
63	1,008	4,032	4,500
64	1,024	4,096	4,500
65	1,040	4,160	4,500
66	1,056	4,224	4,500
67	1,072	4,288	4,500
68	1,088	4,352	4,500
69	1,104	4,416	4,500
70	1,120	4,480	4,500
71	1,136	4,500	4,500
72	1,152	4,500	4,500
73	1,168	4,500	4,500
74	1,184	4,500	4,500
75	1,200	4,500	4,500
76	1,216	4,500	4,500
77	1,232	4,500	4,500
78	1,248	4,500	4,500
79	1,264	4,500	4,500
80	1,280	4,500	4,500
81	1,296	4,500	4,500
82	1,312	4,500	4,500
83	1,328	4,500	4,500
84	1,344	4,500	4,500
85	1,360	4,500	4,500
86	1,376	4,500	4,500
87	1,392	4,500	4,500
88	1,408	4,500	4,500
89	1,424	4,500	4,500
90	1,440	4,500	4,500
91	1,456	4,500	4,500
92	1,472	4,500	4,500
93	1,488	4,500	4,500
94	1,504	4,500	4,500

Capacity (TB)	Standard (MB/s)	Premium (MB/s)	Extreme (MB/s)
95	1,520	4,500	4,500
96	1,536	4,500	4,500
97	1,552	4,500	4,500
98	1,568	4,500	4,500
99	1,584	4,500	4,500
100	1,600	4,500	4,500

Example 1

For example, your application requires 25 TB of capacity and 100 MB/s of bandwidth. At 25 TB of capacity, the Standard service level would provide 400 MB/s of bandwidth at a cost of \$2,500 (estimate: see current pricing), making Standard the most suitable service level in this case.

Example 2

For example, your application requires 12 TB of capacity and 800 MB/s of peak bandwidth. Although the Extreme service level can meet the demands of the application at the 12 TB mark, it is more cost-effective (estimate: see current pricing) to select 13 TB at the Premium service level.

AWS security group settings for Windows AD servers

If you use Windows Active Directory (AD) servers with cloud volumes, you should familiarize yourself with the guidance on AWS security group settings. The settings enable cloud volumes to integrate with AD correctly.

By default, the AWS security group applied to an EC2 Windows instance does not contain inbound rules for any protocol except RDP. You must add rules to the security groups that are attached to each Windows AD instance to enable inbound communication from Cloud Volumes Service. The required ports are as follows:

Service	Port	Protocol
AD Web Services	9389	TCP
DNS	53	TCP
DNS	53	UDP
ICMPv4	N/A	Echo Reply
Kerberos	464	TCP
Kerberos	464	UDP
Kerberos	88	TCP
Kerberos	88	UDP
LDAP	389	TCP

Service	Port	Protocol
LDAP	389	UDP
LDAP	3268	TCP
NetBIOS name	138	UDP
SAM/LSA	445	TCP
SAM/LSA	445	UDP
Secure LDAP	636	TCP
Secure LDAP	3269	TCP
w32time	123	UDP

If you are deploying and managing your AD installation domain controllers and member servers on an AWS EC2 instance, you will require several security group rules to allow traffic for the Cloud Volumes Service. Below is an example of how to implement these rules for AD applications as part of the AWS CloudFormation template.

```
{
    "AWSTemplateFormatVersion" : "2010-09-09",
    "Description" : "Security Group for AD",
    "Parameters" :
    {
        "VPC" :
        {
            "Type" : "AWS::EC2::VPC::Id",
            "Description" : "VPC where the Security Group will belong:"
        },
        "Name" :
        {
            "Type" : "String",
            "Description" : "Name Tag of the Security Group:"
        },
        "Description" :
        {
            "Type" : "String",
            "Description" : "Description Tag of the Security Group:",
            "Default" : "Security Group for Active Directory for CVS"
        },
        "CIDRrangeforTCPandUDP" :
        {
            "Type" : "String",
            "Description" : "CIDR Range for the UDP ports
445,138,464,389,53,123 and for the TCP ports
464,339,3389,3268,88,636,9389,445 and 0-65535: *CIDR range format:
10.0.0.0/24"
        }
    }
}
```

```

        },
    },
    "Resources" :
    {
        "ADSGWest" :
        {
            "Type" : "AWS::EC2::SecurityGroup",
            "Properties" :
            {
                "GroupDescription" : { "Ref" : "Description" },
                "VpcId" : { "Ref" : "VPC" },
                "SecurityGroupIngress" : [
                    {
                        "IpProtocol" : "udp",
                        "CidrIp" : { "Ref" : "CIDRrangeforTCPandUDP" },
                        "FromPort" : "445",
                        "ToPort" : "445"
                    },
                    {
                        "IpProtocol" : "udp",
                        "CidrIp" : { "Ref" : "CIDRrangeforTCPandUDP" },
                        "FromPort" : "138",
                        "ToPort" : "138"
                    },
                    {
                        "IpProtocol" : "udp",
                        "CidrIp" : { "Ref" : "CIDRrangeforTCPandUDP" },
                        "FromPort" : "464",
                        "ToPort" : "464"
                    },
                    {
                        "IpProtocol" : "tcp",
                        "CidrIp" : { "Ref" : "CIDRrangeforTCPandUDP" },
                        "FromPort" : "464",
                        "ToPort" : "464"
                    },
                    {
                        "IpProtocol" : "udp",
                        "CidrIp" : { "Ref" : "CIDRrangeforTCPandUDP" },
                        "FromPort" : "389",
                        "ToPort" : "389"
                    },
                    {
                        "IpProtocol" : "udp",
                        "CidrIp" : { "Ref" : "CIDRrangeforTCPandUDP" },
                        "FromPort" : "53",
                        "ToPort" : "53"
                    }
                ]
            }
        }
    }
}

```

```

        "ToPort" : "53"
    },
{
    "IpProtocol" : "tcp",
    "CidrIp" : {"Ref" : "CIDRrangeforTCPandUDP"},
    "FromPort" : "339",
    "ToPort" : "339"
},
{
    "IpProtocol" : "udp",
    "CidrIp" : {"Ref" : "CIDRrangeforTCPandUDP"},
    "FromPort" : "123",
    "ToPort" : "123"
},
{
    "IpProtocol" : "tcp",
    "CidrIp" : {"Ref" : "CIDRrangeforTCPandUDP"},
    "FromPort" : "3389",
    "ToPort" : "3389"
},
{
    "IpProtocol" : "tcp",
    "CidrIp" : {"Ref" : "CIDRrangeforTCPandUDP"},
    "FromPort" : "3268",
    "ToPort" : "3268"
},
{
    "IpProtocol" : "tcp",
    "CidrIp" : {"Ref" : "CIDRrangeforTCPandUDP"},
    "FromPort" : "88",
    "ToPort" : "88"
},
{
    "IpProtocol" : "tcp",
    "CidrIp" : {"Ref" : "CIDRrangeforTCPandUDP"},
    "FromPort" : "636",
    "ToPort" : "636"
},
{
    "IpProtocol" : "tcp",
    "CidrIp" : {"Ref" : "CIDRrangeforTCPandUDP"},
    "FromPort" : "3269",
    "ToPort" : "3269"
},
{
    "IpProtocol" : "tcp",

```

```

        "CidrIp" : {"Ref" : "CIDRrangeforTCPandUDP"},  

        "FromPort" : "53",  

        "ToPort" : "53"  

    },  

    {  

        "IpProtocol" : "tcp",  

        "CidrIp" : {"Ref" : "CIDRrangeforTCPandUDP"},  

        "FromPort" : "0",  

        "ToPort" : "65535"  

    },  

    {  

        "IpProtocol" : "tcp",  

        "CidrIp" : {"Ref" : "CIDRrangeforTCPandUDP"},  

        "FromPort" : "9389",  

        "ToPort" : "9389"  

    },  

    {  

        "IpProtocol" : "tcp",  

        "CidrIp" : {"Ref" : "CIDRrangeforTCPandUDP"},  

        "FromPort" : "445",  

        "ToPort" : "445"  

    }  

}  

]  

}  

}  

},  

"Outputs" :  

{  

    "SecurityGroupID" :  

    {  

        "Description" : "Security Group ID",  

        "Value" : { "Ref" : "ADSGWest" }  

    }  

}
}

```

Cloud Volumes Service for GCP

Learn about Cloud Volumes Service for Google Cloud

NetApp Cloud Volumes Service for Google Cloud enables you to quickly add multi-protocol workloads as well as build and deploy both Windows-based and UNIX-based apps.

Key features:

- Migrate data between on-premises and Google Cloud.
- Provision volumes from 1 to 100 TiB in seconds.
- Multiprotocol support (you can create an NFS or SMB volume).
- Protect data with automated, efficient snapshots.
- Accelerate app development with rapid cloning.

Cost

Volumes created by the Cloud Volumes Service for Google Cloud are charged to your subscription to the service, not through Cloud Manager.

[View pricing](#)

There are no charges to discover a Cloud Volumes Service for Google Cloud region or volume from Cloud Manager.

Supported regions

[View supported Google Cloud regions.](#)

Before you get started

Cloud Manager can discover existing Cloud Volumes Service for GCP subscriptions and volumes. See the [NetApp Cloud Volumes Service for Google Cloud documentation](#) if you haven't set up your subscription yet.

Getting help

Use the Cloud Manager chat for general questions about Cloud Volumes Service operation in Cloud Manager.

For general questions about Cloud Volumes Service for Google Cloud, email NetApp's Google Cloud Team at gcinfo@netapp.com.

For technical issues associated with your cloud volumes, you can create a technical support case from the Google Cloud Console. See [obtaining support](#) for details.

Limitations

- Cloud Manager doesn't support data replication between working environments when using Cloud Volumes Service volumes.
- Deleting your Cloud Volumes Service for Google Cloud subscription from Cloud Manager isn't supported. You can do this only through the Google Cloud Console.

Related links

- [NetApp Cloud Central: Cloud Volumes Service for Google Cloud](#)
- [NetApp Cloud Volumes Service for Google Cloud documentation](#)

Set up Cloud Volumes Service for Google Cloud

Create a Cloud Volumes Service for Google Cloud working environment in Cloud Manager to create and manage volumes and snapshots.

Quick start

Get started quickly by following these steps, or go to the next section for full details.

1

Enable the Cloud Volumes Service API

From Google, enable the Cloud Volumes Service for GCP API so that Cloud Manager can manage the subscription and cloud volumes.

2

Create a GCP service account and download credentials

From Google, create a GCP service account and role so that Cloud Manager can access your Cloud Volumes Service for GCP account.

3

Create a Cloud Volumes Service for GCP working environment

In Cloud Manager, click **Add Working Environment > Google Cloud > Cloud Volumes Service** and then provide details about the service account and Google Cloud project.

Enable the Cloud Volumes Service API

In Google Cloud Shell, run the following command to enable the Cloud Volumes Service API:

```
gcloud --project=<my-cvs-project> services enable cloudvolumesgcp-api.netapp.com
```

Give Cloud Manager access to the Cloud Volumes Service for GCP account

You must complete the following tasks so that Cloud Manager can access your Google Cloud project:

- Create a new service account
- Add the new service account member to your project and assign it specific roles (permissions)
- Create and download a key pair for the service account that is used to authenticate to Google

Steps

1. In the Google Cloud console, [go to the Service accounts page](#).
2. Click **Select a project**, choose your project, and click **Open**.
3. Click **Create service account**.
4. Enter the service account name (friendly display name) and description.

The Cloud Console generates a service account ID based on this name. Edit the ID if necessary - you cannot change the ID later.

5. To set access controls now, click **Create** and then **DONE** from the bottom of the page, and continue to the next step.
6. From the *IAM page* click **Add** and fill out the fields in the *Add Members* page:
 - a. In the New Members field, enter the full service account ID, for example, user1-service-account-cvs@project1.iam.gserviceaccount.com.
 - b. Add these roles:
 - *NetApp Cloud Volumes Admin*
 - *Compute Network Viewer*
 - c. Click **Save**.
7. Click the Service Account name, and then from the *Service account details* page, click **Add key > Create new key**.
8. Select **JSON** as the key type and click **Create**.

By clicking **Create** your new public/private key pair is generated and downloaded to your system. It serves as the only copy of the private key. Store this file securely because it can be used to authenticate as your service account.

For detailed steps, see the Google Cloud topics [Creating and managing service accounts](#), [Granting, changing, and revoking access to resources](#), and [Creating and managing service account keys](#).

Create a Cloud Volumes Service for GCP working environment

Set up a Cloud Volumes Service for GCP working environment in Cloud Manager so you can start creating volumes.

Regardless of whether you have already created volumes from the Google Cloud Console, or if you just signed up for Cloud Volumes Service for GCP and have no volumes yet, the first step is to create a working environment for the volumes based on your GCP subscription.

If cloud volumes already exist for this subscription, then the volumes will appear in the new working environment. If you haven't added any cloud volumes yet for the GCP subscription, then you do that after you create the new working environment.



If you have subscriptions and volumes in multiple GCP projects, you need to perform this task for each project.

Before you begin

You must have the following information available when adding a subscription for each project:

- Service account credentials (JSON private key you downloaded)
- Project name

Steps

1. In Cloud Manager, add a new Working Environment, select the location **Google Cloud**, and click **Continue**.
2. Select **Cloud Volumes Service** and click **Continue**.

Add Working Environment Wizard Define Your Working Environment

↑ Previous Step

Microsoft Azure

aws

Google Cloud

On-Premises ONTAP

Cloud Volumes ONTAP
Single Node

Cloud Volumes ONTAP HA
High Availability

Cloud Volumes Service
High Availability

You're about to set up Cloud Volumes Service for GCP.

Continue

3. Provide information about your Cloud Volumes Service subscription:
 - a. Enter the Working Environment Name you want to use.
 - b. Copy/paste the JSON private key you downloaded in the previous steps.
 - c. Select the name of your Google Cloud project.
 - d. Click **Continue**.

Cloud Volumes Service Credentials

Working Environment Name

Service Account Credentials

Paste the contents of the JSON file here

Apply

Project

- Select project -

Result

Cloud Manager displays your Cloud Volumes Service for Google Cloud working environment.



If cloud volumes already exist for this subscription, then the volumes appear in the new working environment. You can add additional cloud volumes from Cloud Manager.

If no cloud volumes exist for this subscription, create them now.

What's next?

[Start creating and managing volumes](#).

Create and manage volumes for Cloud Volumes Service for Google Cloud

Cloud Manager enables you to create cloud volumes based on your [Cloud Volumes Service for Google Cloud](#) subscription. You can also edit certain attributes of a volume, get the relevant mount commands, create snapshot copies, and delete cloud volumes.

Create cloud volumes

You can create NFS or SMB volumes in a new or existing Cloud Volumes Service for Google Cloud account. Cloud volumes currently support NFSv3 and NFSv4.1 for Linux and UNIX clients, and SMB 3.x for Windows clients.

Before you begin

- If you want to use SMB in GCP, you must have set up DNS and Active Directory.
- When planning to create an SMB volume, you must have a Windows Active Directory server available to which you can connect. You will enter this information when creating the volume. Also, make sure that the Admin user is able to create a machine account in the Organizational unit (OU) path specified.

Steps

1. Select the working environment and click **Add New Volume**.
2. In the Details & Location page, enter details about the volume:
 - a. Enter a name for the volume.
 - b. Specify a size within the range of 1 TiB (1024 GiB) to 100 TiB.
[Learn more about allocated capacity](#).
 - c. Specify a service level: Standard, Premium, or Extreme.
[Learn more about service levels](#).
 - d. Select the Google Cloud region.
 - e. Select the VPC Network from which the volume will be accessible. Note that the VPC cannot be changed or edited after the volume is created.
 - f. Click **Continue**.
3. In the Protocol page, select NFS or SMB and then define the details. Required entries for NFS and SMB are shown in separate sections below.
4. For NFS:
 - a. In the Volume Path field, specify the name of the volume export you will see when you mount the volume.

- b. Select NFSv3, NFSv4.1, or both depending on your requirements.
- c. Optionally, you can create an export policy to identify the clients that can access the volume. Specify the:
 - Allowed clients by using an IP address or Classless Inter-Domain Routing (CIDR).
 - Access rights as Read & Write or Read Only.
 - Access protocol (or protocols if the volume allows both NFSv3 and NFSv4.1 access) used for users.
 - Click **+ Add Export Policy Rule** if you want to define additional export policy rules.

The following image shows the Volume page filled out for the NFS protocol:

Protocol

Select the volume's protocol: **NFS Protocol** SMB Protocol

Protocol	Export Policy
Volume Path <input type="text" value="vol1"/>	Allowed Client & Access <input type="text" value="0.0.0.0/24"/> <input checked="" type="radio"/> Read & Write <input type="radio"/> Read Only
Select NFS Version: <input checked="" type="checkbox"/> NFSv3 <input type="checkbox"/> NFSv4.1	Select NFS Version: <input checked="" type="checkbox"/> NFSv3 <input type="checkbox"/> NFSv4.1

+ Add Export Policy Rule (Up to 5)

5. For SMB:

- a. In the Volume Path field, specify the name of the volume export you will see when you mount the volume and click **Continue**.
- b. If Active Directory has been set up, you see the configuration. If it is the first volume being set up and no Active Directory has been set up, you can enable SMB session encryption in the SMB Connectivity Setup page:

Field	Description
DNS Primary IP Address	The IP addresses of the DNS servers that provide name resolution for the SMB server. Use a comma to separate the IP addresses when referencing multiple servers, for example, 172.31.25.223, 172.31.2.74..
Active Directory Domain to join	The FQDN of the Active Directory (AD) domain that you want the SMB server to join.
SMB Server NetBIOS name	A NetBIOS name for the SMB server that will be created.
Credentials authorized to join the domain	The name and password of a Windows account with sufficient privileges to add computers to the specified Organizational Unit (OU) within the AD domain.

Field	Description
Organizational Unit	The organizational unit within the AD domain to associate with the SMB server. The default is CN=Computers for connections to your own Windows Active Directory server.

The following image shows the Volume page filled out for the SMB protocol:

The screenshot shows the 'SMB Connectivity Setup' page. It contains the following fields:

- DNS Primary IP Address:** 127.0.0.1
- User Name:** administrator
- Active Directory Domain to Join:** yourdomain.com up to 107 characters
- Password:** (empty field)
- SMB Server NetBIOS Name:** WENName
- Organizational Unit:** CN=Computers

6. Click **Continue**.

7. If you want to create the volume based on a snapshot of an existing volume, select the snapshot from the Snapshot Name drop-down list. Otherwise just click **Continue**.
8. In the Snapshot Policy page, you can enable Cloud Volumes Service to create snapshot copies of your volumes based on a schedule. You can do this now by moving the selector to the right, or you can edit the volume later to define the snapshot policy.

See [Creating a snapshot policy](#) for more information about snapshot functionality.

9. Click **Add Volume**.

The new volume is added to the working environment.

Continue with [Mounting the cloud volume](#).

Mount cloud volumes

Access mounting instructions from within Cloud Manager so you can mount the volume to a host.

Note: Please use the highlighted protocol/dialect supported by your client.

Steps

1. Open the working environment.
2. Hover over the volume and click **Mount the volume**.

NFS and SMB volumes display mount instructions for that protocol.

3. Hover over the commands and copy them to your clipboard to make this process easier. Just add the destination directory/mount point at the end of the command.

NFS example:

Mount the volume - testk

Setting up your instance

1. Open an SSH client and connect to your instance.
2. Install the nfs client on your instance.

On Red Hat Enterprise Linux or SuSE Linux instance:

```
$ sudo yum install -y nfs-utils
```

On an Ubuntu or Debian instance:

```
$ sudo apt-get install nfs-common
```

Mounting your volume

1. Create a new directory on your instance:

```
$ sudo mkdir /dir
```
2. Mount your NFSv3 volume using the command below:

```
sudo mount -t nfs -o rw,hard,rsize=65536,wsize=65536,vers=3,tc...
```
3. Mount your NFSv4.1 volume using the command below:

```
sudo mount -t nfs -o rw,hard,rsize=65536,wsize=65536,vers=4,1,t...
```

The maximum I/O size defined by the `rsize` and `wszie` options is 1048576, however 65536 is the recommended default for most use cases.

Note that Linux clients will default to NFSv4.1 unless the version is specified with the `vers=<nfs_version>` option.

SMB example:

Mount the volume - <Volume Name>

Mapping your network drive

1. Click the Start button and then click on Computer.
2. Click Map Network Drive.
3. In the Drive list, click any available drive letter.
4. In the Folder box, type this:

```
\test.cv-pm.local\silly-condescending-mcnulty
```



To connect every time you log on to your computer, check the **Reconnect at logon** option.

5. Click Finish.

4. Map your network drive by following the mount instructions for your instance.

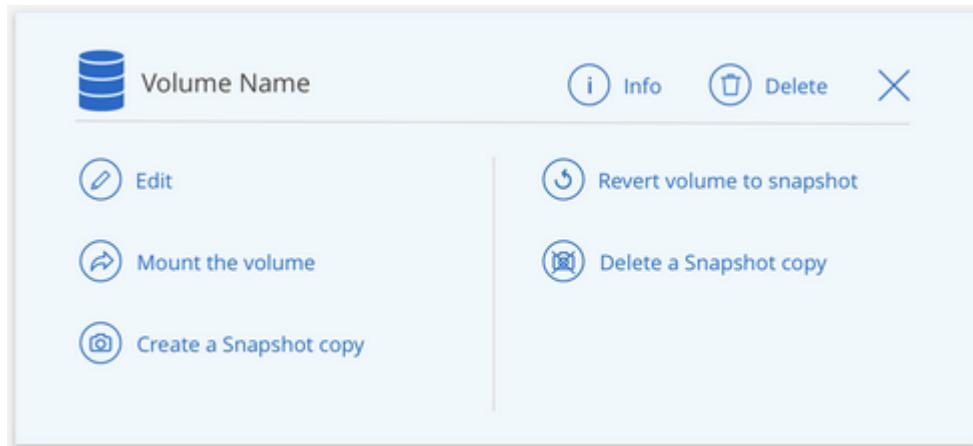
After completing the steps in the mount instructions, you have successfully mounted the cloud volume to your GCP instance.

Manage existing volumes

You can manage existing volumes as your storage needs change. You can view, edit, restore, and delete volumes.

Steps

1. Open the working environment.
2. Hover over the volume.



3. Manage your volumes:

Task	Action
View information about a volume	Click Info .
Edit a volume (including snapshot policy)	a. Click Edit . b. Modify the volume's properties and then click Update .
Get the NFS or SMB mount command	a. Click Mount the volume . b. Click Copy to copy the command(s).
Create a Snapshot copy on demand	a. Click Create a Snapshot copy . b. Change the name, if needed, and then click Create .
Replace the volume with the contents of a Snapshot copy	a. Click Revert volume to snapshot . b. Select a Snapshot copy and click Restore .
Delete a Snapshot copy	a. Click Delete a Snapshot copy . b. Select the snapshot and click Delete . c. Click Delete again when prompted to confirm.
Delete a volume	a. Unmount the volume from all clients: ◦ On Linux clients, use the <code>umount</code> command. ◦ On Windows clients, click Disconnect network drive . b. Select a volume, and then click Delete . c. Click Delete again to confirm.

Remove Cloud Volumes Service from Cloud Manager

You can remove a Cloud Volumes Service for Google Cloud subscription and all existing volumes from Cloud Manager. The volumes are not deleted, they are just removed from the Cloud Manager interface.

Steps

1. Open the working environment.
2. Click the  button at the top of the page and click **Remove Cloud Volumes Service**.
3. In the confirmation dialog box, click **Remove**.

Manage Active Directory configuration

If you change your DNS servers or Active Directory domain, you need to modify the SMB server in Cloud Volumes Services so that it can continue to serve storage to clients.

Steps

1. Open the working environment.

2. Click the  button at the top of the page and click **Manage Active Directory**. If no Active Directory is configured, you can add one now. If one is configured, you can modify or delete the settings using the  button.
3. Specify the settings for the SMB server:

Field	Description
DNS Primary IP Address	The IP addresses of the DNS servers that provide name resolution for the SMB server. Use a comma to separate the IP addresses when referencing multiple servers, for example, 172.31.25.223, 172.31.2.74.
Active Directory Domain to join	The FQDN of the Active Directory (AD) domain that you want the SMB server to join.
SMB Server NetBIOS name	A NetBIOS name for the SMB server that will be created.
Credentials authorized to join the domain	The name and password of a Windows account with sufficient privileges to add computers to the specified Organizational Unit (OU) within the AD domain.
Organizational Unit	The organizational unit within the AD domain to associate with the SMB server. The default is CN=Computers for connections to your own Windows Active Directory server.

4. Click **Save** to save your settings.

Manage cloud volumes snapshots

You can create a snapshot policy for each volume so that you can recover or restore the entire contents of a volume from an earlier time. You can also create an on-demand snapshot of a cloud volume when needed.

Create an on-demand snapshot

You can create an on-demand snapshot of a cloud volume if you want to create a snapshot with the current volume state.

Steps

1. Open the working environment.
2. Hover over the volume and click **Create a snapshot copy**.
3. Enter a name for the snapshot, or use the automatically generated name, and click **Create**.

Create a Snapshot Copy - <Volume Name>

A NetApp Snapshot copy is a read-only, point-in-time image of a volume. The image protects your data with no performance impact and requires minimal storage.

Snapshot Copy Name

manually.2020-05-04_1722

Create

The snapshot is created.

Create or modify a snapshot policy

You can create or modify a snapshot policy as necessary for a cloud volume. You define the snapshot policy from the *Snapshot Policy* tab either when creating a volume or when editing a volume.

Steps

1. Open the working environment.
2. Hover over the volume and click **Edit**.
3. From the *Snapshot Policy* tab, move the enable snapshots slider to the right.
4. Define the schedule for snapshots:
 - a. Select the frequency: **Hourly**, **Daily**, **Weekly**, or **Monthly**
 - b. Select the number of snapshots you want to keep.
 - c. Select the day, hour, and minute when the snapshot should be taken.

Schedule Snapshot Policies:

<input checked="" type="checkbox"/> Hourly	Number of Snapshot to Keep	Minute	
	<input type="text" value="12"/>	<input type="text" value="30"/>	
<input type="checkbox"/> Daily	Number of Snapshot to Keep	Hour Minute	
	<input type="text" value="0"/>	<input type="text" value="0"/> <input type="text" value="0"/>	
<input checked="" type="checkbox"/> Weekly	Number of Snapshot to Keep	Days	Hour Minute
	<input type="text" value="3"/>	<input type="button" value="Sunday X"/>	<input type="text" value="0"/> <input type="text" value="0"/>
<input type="checkbox"/> Monthly	Number of Snapshot to Keep	Days	Hour Minute
	<input type="text" value="0"/>	<input type="checkbox"/> Sunday <input type="checkbox"/> Monday <input type="checkbox"/> Tuesday	<input type="text" value="0"/> <input type="text" value="0"/>

5. Click **Add volume** or **Update volume** to save your policy settings.

Disable a snapshot policy

You can disable a snapshot policy to stop snapshots from being created for a short period of time while retaining your snapshot policy settings.

Steps

1. Open the working environment.
2. Hover over the volume and click **Edit**.
3. From the *Snapshot Policy* tab, move the enable snapshots slider to the left.

Enable automatic Snapshot copies



When disabled, Cloud Volumes Service does not create Snapshot copies of your volumes.

4. Click **Update volume**.

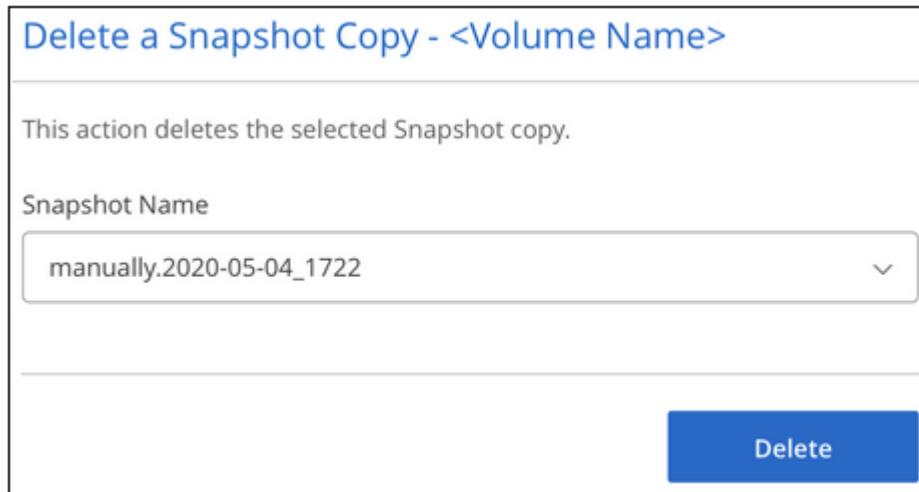
When you want to re-enable the snapshot policy, move the enable snapshots slider to the right and click **Update volume**.

Delete a snapshot

You can delete a snapshot if it is no longer needed.

Steps

1. Open the working environment.
2. Hover over the volume and click **Delete a Snapshot copy**.
3. Select the snapshot from the drop-down list and click **Delete**.



4. In the confirmation dialog box, click **Delete**.

Restore a snapshot to a new volume

You can restore a snapshot to a new volume as necessary.

Steps

1. Open the working environment.
2. Hover over the volume and click **Restore to a new volume**.
3. Select the snapshot that you want to use to create the new volume from the drop-down list.
4. Enter a name for the new volume and click **Restore**.

Restore to a new volume - <Volume Name>

This operation restores data from a Snapshot copy to a new volume.

Snapshot Name

manually.2020-05-04_1722



Restored Volume Name:

vol_restore

Restore

The volume is created in the working environment.

5. If you need to change any of the volume attributes, such as volume path or service level:
 - a. Hover over the volume and click **Edit**.
 - b. Make your changes and click **Update volume**.

After you finish

Continue with [Mounting the cloud volume](#).

Manage on-prem ONTAP clusters

Discovering ONTAP clusters

Cloud Manager can discover the ONTAP clusters in your on-premises environment, in a NetApp Private Storage configuration, and in the IBM Cloud. Adding on-prem clusters to the Cloud Manager Canvas enables you to manage these clusters using the same tools as your Cloud Volumes ONTAP and other cloud storage solutions.

In addition to being able to provision storage on those systems, adding these systems to Cloud Manager also makes it easy to configure critical cloud-based services for those clusters. This includes replicating data to the cloud, backing up data to the cloud, tiering cold data to the cloud, and running compliance scans on that data.

What you'll need

- A Connector installed in a cloud provider or on your premises.

If you want to tier cold data to the cloud, then you should review requirements for the Connector based on where you plan to tier cold data.

- [Learn about Connectors](#)
- [Switching between Connectors](#)
- [Learn about Cloud Tiering](#)

- The cluster management IP address and the password for the admin user account to add the cluster to Cloud Manager.

Cloud Manager discovers ONTAP clusters using HTTPS. If you use custom firewall policies, they must meet the following requirements:

- The Connector host must allow outbound HTTPS access through port 443.

If the Connector is in the cloud, all outbound communication is allowed by the predefined security group.

- The ONTAP cluster must allow inbound HTTPS access through port 443.

The default "mgmt" firewall policy allows inbound HTTPS access from all IP addresses. If you modified this default policy, or if you created your own firewall policy, you must associate the HTTPS protocol with that policy and enable access from the Connector host.

- A valid set of NetApp Support Site credentials for accessing the Active IQ page. See how to [add NetApp Support Site accounts to Cloud Manager](#).

Checking for on-premises clusters that have not been added to Cloud Manager

You can use the Active IQ service in Cloud Manager to discover, view, and manage your on-prem ONTAP clusters that are under a support contract.

 If your support contract expires, the systems are removed from the Active IQ page. However, you can continue to manage these systems in their working environment. See how to [renew your support contract from Active IQ Digital Advisor](#).

Steps

1. From Cloud Manager, click the **Active IQ** tab.
2. Select the Active IQ login associated with your NSS account if necessary.

The screenshot shows the Active IQ dashboard with the following sections:

- On-Premise Cluster Inventory:** A donut chart showing 1 Discovered and 41 Undiscovered clusters.
- Licenses:** A donut chart showing 12 Used and 18 Unused licenses.
- Firmware Updates for On-Premise Clusters:** 7 Updates Available.
- Cloud Ready Workloads:** 1375 Workloads.

Below the dashboard is a table with the following data:

Cluster Name	Cluster Status	OS Version	IP Address
durlabu01	Discovered	9.8RC1	10. [REDACTED]
durbclu99	Undiscovered	9.8	10. [REDACTED]

Buttons at the bottom right of the table include "Discover" and "Download CSV".

Your ONTAP clusters that have a valid support contract are displayed with a status of whether they have been discovered in Cloud Manager.

3. If you want to manage the clusters that are listed as "Undiscovered" using Cloud Manager, click **Discover** and follow the prompts to add them.

Discovering clusters from the Active IQ page

You can discover your ONTAP clusters and add them to a working environment from the Active IQ page.

Steps

1. From the **Active IQ** page, click the **Cluster Inventory** tab.

The screenshot shows the Cluster Inventory tab with the following data:

Cluster Name	Cluster Status	OS Version	IP Address
durlabu01	Discovered	9.8RC1	10. [REDACTED]
durbclu99	Undiscovered	9.8	10. [REDACTED]
b1rclu02	Undiscovered	9.7P7	10. [REDACTED]

Buttons at the bottom right of the table include "Discover" for each undiscovered cluster.

2. Click **Discover** for the cluster that you want to manage through Cloud Manager.
3. On the *Choose a Location* page **On-Premises ONTAP** is pre-selected, so just click **Continue**.

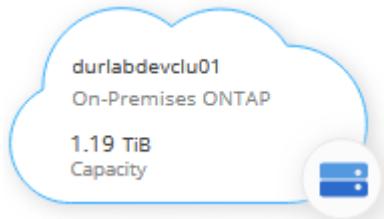
4. On the *ONTAP Cluster Details* page, enter the password for the admin user account and click **Add**.

Note that the cluster management IP address is populated based on information from Active IQ.

5. On the *Details & Credentials* page the cluster name is added as the Working Environment Name, so just click **Go**.

Result

Cloud Manager discovers the cluster and adds it to a working environment in the Canvas using the cluster name as the working environment name.



You can enable services for this cluster in the right panel to replicate data to and from the cluster, set up data tiering to the cloud, back up volumes to the cloud, or run compliance scans on the volumes. You can also create new volumes or launch System Manager to perform advanced tasks.

Discovering clusters from the Canvas page

You can discover your ONTAP clusters and add them to a working environment from the Canvas page. These steps can be used in cases where the cluster is not listed in the Active IQ page because it currently has no support contract.

Steps

1. On the Canvas page, click **Add Working Environment** and select **On-Premises ONTAP**.
2. If you're prompted, create a Connector.

Refer to the links above for more details.

3. On the *ONTAP Cluster Details* page, enter the cluster management IP address, the password for the admin user account, and click **Add**.
4. On the *Details & Credentials* page, enter a name and description for the working environment, and then click **Go**.

Result

Cloud Manager discovers the cluster and adds it to a working environment in the Canvas.

You can enable services for this cluster in the right panel to replicate data to and from the cluster, set up data tiering to the cloud, back up volumes to the cloud, or run compliance scans on the volumes. You can also create new volumes or launch System Manager to perform advanced tasks.

Using Active IQ data to manage ONTAP clusters

The Active IQ page in Cloud Manager shows you any undiscovered ONTAP clusters in your on-premises environments, whether any clusters require updated disk or shelf firmware, and if you are using all the Cloud Volumes ONTAP licenses that you were granted when you purchased the on-prem systems. This information is provided to Cloud Manager from [Active IQ Digital Advisor](#).

Viewing unused Cloud Volumes ONTAP licenses

Many on-premises ONTAP storage system packages that you purchased included a free Cloud Volumes ONTAP license so you can try the NetApp cloud storage offerings in Cloud Manager. You can use the license to create a new Cloud Volumes ONTAP instance, or you can apply the license to an existing Cloud Volumes ONTAP instance to expand the capacity by 368 TiB.

You can see whether you have any unused Cloud Volumes ONTAP licenses based on your NetApp Support Site credentials.

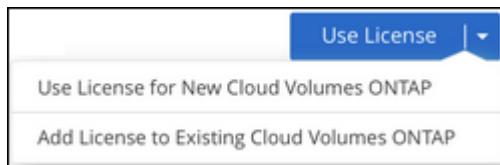
Steps

1. In Cloud Manager, click the **Active IQ** tab.
2. Click the **Licenses** tab in the lower portion of the page.

Cluster Inventory (42)		Licenses (30)		Firmware Updates (7)		Cloud Ready Workloads (1375)	
Serial Number		License Type	Hyperscaler	Model Type	Expires	AutoSupport	
903201300000000001514	BYOL	AWS	Single	December 31, 2022	No	Use License	▼
908201300000000001141	BYOL	AWS	Single	N/A	Yes		
908201300000000001142	BYOL	AWS	Single	December 31, 2022	Yes	Use License	▼
908201300000000001143	BYOL	AWS	HA	December 31, 2022	Yes	Use License	▼

A **Use License** button appears for each unused license.

3. If you want to activate and start using the license, click **Use License**.



See the tasks below to learn about the options for using the available licenses.

Using unused Cloud Volumes ONTAP licenses

You can use unused licenses to create a new Cloud Volumes ONTAP instance or to extend the capacity of the license on an existing Cloud Volumes ONTAP instance. The capacity of the license is 368 TiB.

The *Expires* column indicates the last day the license is active. When creating a new Cloud Volumes ONTAP system this is the date the license expires. When updating an existing Cloud Volumes ONTAP system this

indicates the length of time the existing license is extended.

The *License Type*, *Hyperscaler*, and *Model Type* columns describe the type of Cloud Volumes ONTAP license it is. For example, **BYOL | Single | Azure** means the license is a "bring-your-own" license for a "single node" Cloud Volumes ONTAP system deployed in "Microsoft Azure". The values that can appear in this column are shown in the table.

Column	Values
License Type	PAYGO BYOL
Hyperscaler	Azure AWS GCP All Providers
Model Type	Single HA

When creating a *new* Cloud Volumes ONTAP system, this is the type of system you are deploying. For example, using the sample license (**BYOL | Single | Azure**), you can create a single-node Cloud Volumes ONTAP system in Azure with entitlement for up to 368 TiB. This license can't be used to create an HA system or to deploy an instance in AWS.

When updating an *existing* Cloud Volumes ONTAP system, this indicates the type of system that can have the capacity for its existing license extended. Using the sample license again, you can extend the license for any single-node Cloud Volumes ONTAP system in Azure. This license can't be used to extend the license for an HA system or for an instance deployed in AWS.

Creating a new Cloud Volumes ONTAP system with the unused license

Follow these steps to create a new Cloud Volumes ONTAP instance with the unused license.

Steps

1. Click **Use License** and select **Use License for a new Cloud Volumes ONTAP**.
2. In the "Use License..." page, verify the license information and click **Use License**.

In most cases you will be directed to the **Details & Credentials** page for creating the working environment for the Cloud Volumes ONTAP system because both the cloud provider and number of nodes are defined by the license.

If you are using a license defined as "All Providers", then you are directed to the **Choose a Location** page so you can pick the cloud provider first, before completing the **Details & Credentials** page.

3. Follow the steps to create the working environment and your first volume.

See the following sections depending on the cloud provider on which you are deploying the Cloud Volumes ONTAP system.

- [Launching Cloud Volumes ONTAP in Azure](#)
- [Launching Cloud Volumes ONTAP in AWS](#)
- [Launching Cloud Volumes ONTAP in GCP](#)

Extending the license capacity for an existing Cloud Volumes ONTAP system

If you have a currently deployed Cloud Volumes ONTAP system that matches the license requirements of one of the free licenses (meaning the same cloud provider, number of nodes, etc.), you can follow these steps to extend the capacity of the license by 368 TiB.

Steps

1. Click **Use License** and select **Add License to existing Cloud Volumes ONTAP**.

The screenshot shows the 'Add License to Existing Cloud Volumes ONTAP' dialog box. It has a header 'Add License to Existing Cloud Volumes ONTAP'. Below it is a 'License Information' section with four fields: Serial Number (90419737477578510576), License Type (BYOL), Provider (AWS), and License Type (Single). The next section is 'Select Cloud Volumes ONTAP', which contains a dropdown menu showing 'Cloud_Volumes_ONTAP_001 | AWS | Single'. A note below says 'Note: Only Cloud Volumes ONTAP systems that match the license parameters are displayed.' At the bottom are two buttons: 'Add License' and 'Close'.

2. In the "Add License..." page, select the Cloud Volumes ONTAP system where you want to extend the license and click **Add License**.

A confirmation dialog is displayed.

The screenshot shows a confirmation dialog box with a green checkmark icon and the text 'License added Successfully'. Below it, a message states 'License: 90419737477578510576 (BYOL | AWS | Single) added successfully to Cloud Volumes ONTAP "Cloud_Volumes_ONTAP_Name"' and a red-bordered link 'Go to "Cloud_Volumes_ONTAP_Name" License Page'. At the bottom is a 'Close' button.

3. You can click **Close** to close the window and return to the Active IQ page, or you can click the link to go to the Cloud Volumes ONTAP Licensing page to view more details about licensing for that system.

Downloading new disk and shelf firmware

You can see whether any of your discovered ONTAP clusters need to have their shelf or disk firmware updated. And you can download the Ansible playbook to upgrade the firmware.

Note: The ability to view and download new firmware is available only when you have subscribed to certain support plans.

Steps

1. From the Active IQ page, click the **Firmware Updates** tab.

Cluster Inventory (42)	Licenses (30)	Firmware Updates (7)	Cloud Ready Workloads (1375)	Download All
Cluster Name	Cluster Status	Disk Firmware	Shelf Firmware	
durbkpclu99	● Undiscovered	✗ Update Available	No Updates Available	
durdevnasclu01	● Undiscovered	✗ Update Available	No Updates Available	
durlabdevclu01	● Discovered	No Updates Available	No Updates Available	
blrprdclu02	● Undiscovered	No Updates Available	No Updates Available	

If any cluster require new firmware, a **Download All** button appears.

2. Click **Download All** and save the zip file.
3. Unzip the zip file and see the following instructions to [update your storage system firmware](#).

Result

Your firmware is updated. The next time your ONTAP system sends an AutoSupport message to Active IQ, the status in the *Firmware Updates* page will be updated to show that updates are no longer needed.

Viewing on-prem workloads that are candidates for the cloud

Certain workloads or volumes are ideal to move to a Cloud Volumes ONTAP system from your on-prem ONTAP clusters. Some of the advantages include reduced costs and improved performance and resiliency. The *Cloud Ready Workloads* tab provides a list of these workloads from your discovered ONTAP clusters.

Cluster Inventory (42)	Licenses (30)	Firmware Updates (7)	Cloud Ready Workloads (1375)	
Cluster Name	Cluster Status	SVM Name	Volume Name	Workload Type
hioprclu02	● Undiscovered	vsvhiopax01prd	volpaxprd_hanabackup01	SAP HANA
hioprclu02	● Undiscovered	svmhiocdb02prd	voicdbprd_sqluserdata01	MSSQL
durdevclu02	● Discovered	vsvdurmrx01spd	volpaxdev_hana_data	SAP HANA
durdevclu02	● Discovered	vsvdurmrx01spd	volpaxstg_hana_backup	SAP HANA
durdevclu02	● Discovered	vsvdurerp01spd	xdperpspd_oradata02	ORACLE

The supported workloads that are called out on this page include: SAP, SAP HANA, Oracle, File share, and SharePoint.

Lift and shift is an approach for migrating your apps to the cloud. It means moving an application and its associated data to a cloud platform without redesigning the app. See more information about [lift and shift](#).

Managing storage for ONTAP clusters

After you discover your ONTAP cluster from Cloud Manager, you can open the working environment to provision and manage storage.

Creating volumes for ONTAP clusters

Cloud Manager enables you to provision NFS, CIFS, and iSCSI volumes on ONTAP clusters.

A Cloud Manager feature called "templates" enables you to create volumes that are optimized for the workload requirements for certain applications; such as databases or streaming services. If your organization has created volume templates that you should use, follow [these steps](#).

Before you begin

The data protocols must be set up on the cluster using System Manager or the CLI.

About this task

You can create volumes on existing aggregates. You can't create new aggregates from Cloud Manager.

Steps

1. On the Canvas page, double-click the name of the ONTAP cluster on which you want to provision volumes.
2. Click **Add New Volume**.
3. On the Create New Volume page, enter details for the volume, and then click **Create**.

Some of the fields in this page are self-explanatory. The following table describes fields for which you might need guidance:

Field	Description
Size	The maximum size that you can enter largely depends on whether you enable thin provisioning, which enables you to create a volume that is bigger than the physical storage currently available to it.
Snapshot Policy	A Snapshot copy policy specifies the frequency and number of automatically created NetApp Snapshot copies. A NetApp Snapshot copy is a point-in-time file system image that has no performance impact and requires minimal storage. You can choose the default policy or none. You might choose none for transient data: for example, tempdb for Microsoft SQL Server.
Access control (for NFS only)	An export policy defines the clients in the subnet that can access the volume. By default, Cloud Manager enters a value that provides access to all instances in the subnet.

Field	Description
Permissions and Users / Groups (for CIFS only)	These fields enable you to control the level of access to a share for users and groups (also called access control lists or ACLs). You can specify local or domain Windows users or groups, or UNIX users or groups. If you specify a domain Windows user name, you must include the user's domain using the format domain\username.
Initiator group and IQN (for iSCSI only)	<p>iSCSI storage targets are called LUNs (logical units) and are presented to hosts as standard block devices.</p> <p>Initiator groups are tables of iSCSI host node names and control which initiators have access to which LUNs.</p> <p>iSCSI targets connect to the network through standard Ethernet network adapters (NICs), TCP offload engine (TOE) cards with software initiators, converged network adapters (CNAs) or dedicated host bus adapters (HBAs) and are identified by iSCSI qualified names (IQNs).</p> <p>When you create an iSCSI volume, Cloud Manager automatically creates a LUN for you. We've made it simple by creating just one LUN per volume, so there's no management involved. After you create the volume, select it, click Target IQN, and then use the IQN to connect to the LUN from your hosts.</p>
Usage Profile	Usage profiles define the NetApp storage efficiency features that are enabled for a volume.

Creating volumes from templates

If your organization has created on-premises ONTAP volume templates so you can deploy volumes that are optimized for the workload requirements for certain applications, follow the steps in this section.

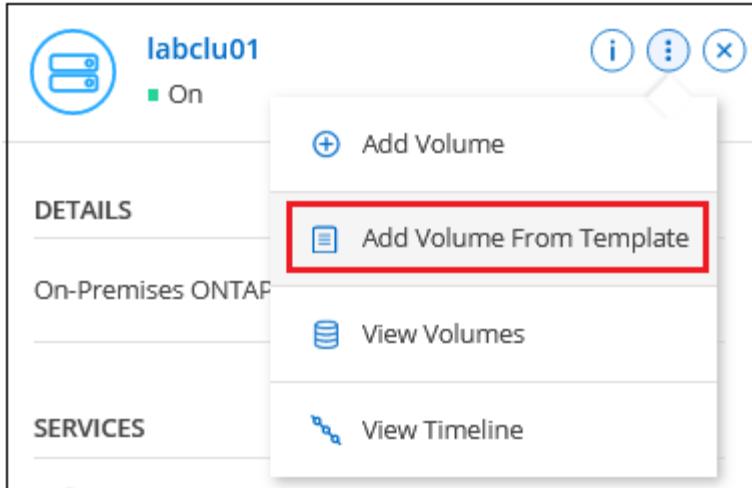
The template should make your job easier because certain volume parameters will already be defined in the template, such as disk type, size, protocol, snapshot policy, and more. When a parameter is already predefined, you can just skip to the next volume parameter.



You can only create NFS or CIFS volumes when using templates.

Steps

1. On the Canvas page, click the name of the on-premises ONTAP system on which you want to provision a volume.
2. Click > Add Volume From Template.



3. In the *Select Template* page, select the template that you want to use to create the volume and click **Next**.

Template Name	Template Description	Created by	Last Modified	Parameters
Onprem ONTAP volume	Volumes for application testing	Tacki	Jun 04 2021, 2:39:51 pm	View

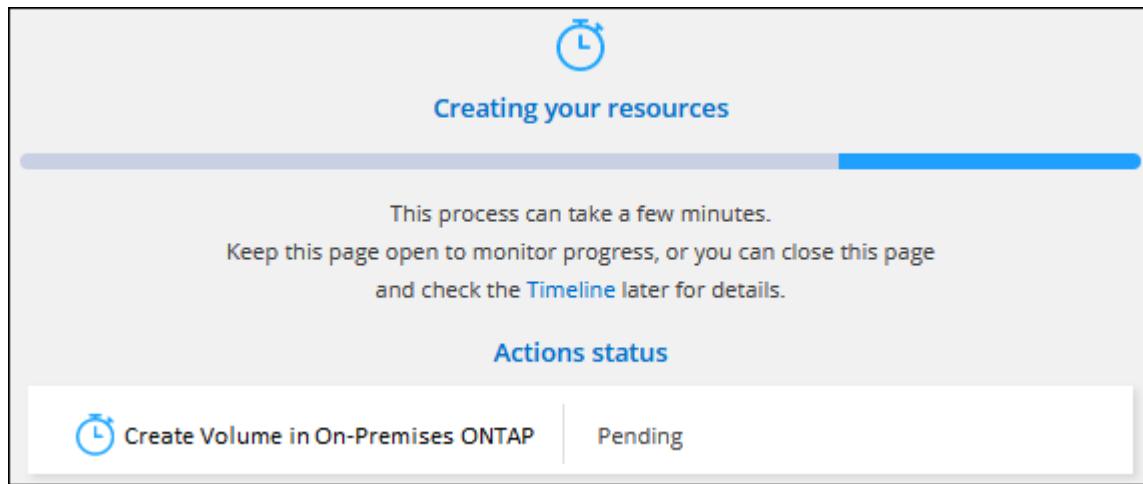
The *Define Parameters* page is displayed.

Note: You can click the checkbox **Show read-only parameters** to show all the fields that have been locked by the template if you want to see the values for those parameters. By default these predefined fields are hidden and only the fields you need to complete are shown.

4. In the *Context* area, the Working Environment is filled in with the name of the working environment you started with. You need to select the **Storage VM** and **Aggregate** where the volume will be created.
5. Add values for all of the parameters that are not hard-coded from the template. See [creating volumes](#) for details about all the parameters you need to complete to deploy an on-prem ONTAP volume.
6. Click **Run Template** after you have defined all the parameters needed for this volume.

Result

Cloud Manager provisions the volume and displays a page so that you can see the progress.



Then the new volume is added to the working environment.

Additionally, if any secondary action is implemented in the template, for example, enabling Cloud Backup on the volume, that action is also performed.

After you finish

If you provisioned a CIFS share, give users or groups permissions to the files and folders and verify that those users can access the share and create a file.

Replicating data

You can replicate data between Cloud Volumes ONTAP systems and ONTAP clusters by choosing a one-time data replication, which can help you move data to and from the cloud, or a recurring schedule, which can help with disaster recovery or long-term retention.

[Click here for more details.](#)

Backing up data

You can back up data from your on-premises ONTAP system to low-cost object storage in the cloud by using the Cloud Manager Cloud Backup service. This service provides backup and restore capabilities for protection and long-term archive of your cloud data.

[Click here for more details.](#)

Tiering data to the cloud

Extend your data center to the cloud by automatically tiering inactive data from ONTAP clusters to object

storage.

[Click here for more details.](#)

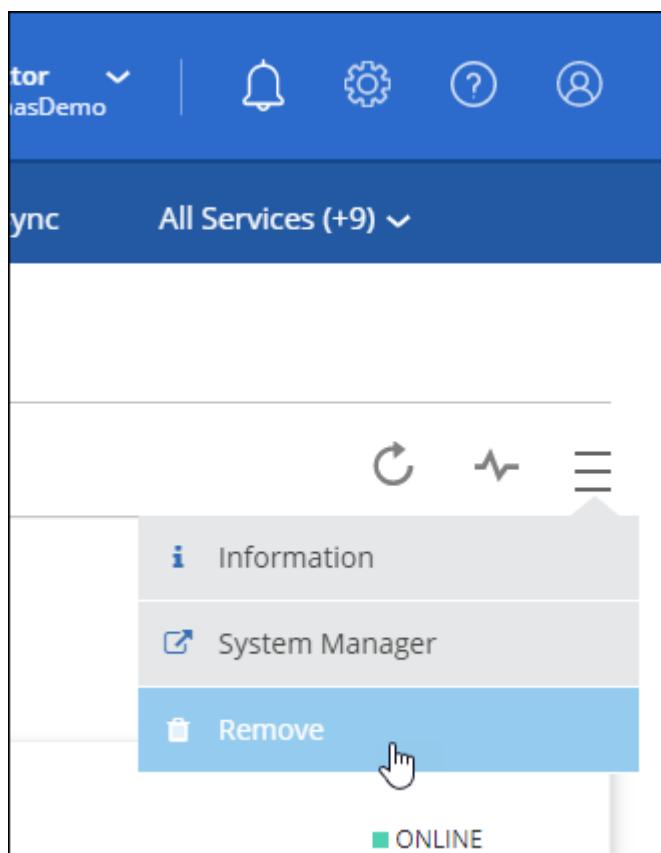
Remove an on-prem ONTAP working environment

Remove an on-premises ONTAP working environment if you no longer want to manage it from Cloud Manager.

Removing the working environment doesn't affect the ONTAP cluster. You can rediscover it from Cloud Manager at any time.

Steps

1. On the Canvas page, double-click the name of the on-premises ONTAP working environment.
2. Click the menu icon and select **Remove**.



3. Click **Remove** to confirm.

Replicate ONTAP data

Learn about the Replication service

NetApp SnapMirror replicates data at high speeds over LAN or WAN, so you get high data availability and fast data replication in both virtual and traditional environments. When you replicate data to NetApp storage systems and continually update the secondary data, your data is kept current and remains available whenever you need it. No external replication servers are required.

Features

- Replicate data between ONTAP storage systems to support backup and disaster recovery to the cloud or between clouds.
- Ensure the reliability of your DR environment with high availability.
- Efficient block-level replication between ONTAP storage is fast and efficient, with granular recovery points for both DR and backup.

Cost

NetApp doesn't charge you for using the Replication service, but you'll need to check your cloud provider for applicable data ingress and egress charges.

Supported working environments

Cloud Manager enables data replication between the following types of working environments.

Source working environment	Supported target working environments
Cloud Volumes ONTAP	<ul style="list-style-type: none">• Amazon FSx for ONTAP• Cloud Volumes ONTAP• On-prem ONTAP cluster
On-prem ONTAP cluster	<ul style="list-style-type: none">• Amazon FSx for ONTAP• Cloud Volumes ONTAP• On-prem ONTAP cluster

How data replication works

Cloud Manager simplifies data replication between volumes on separate ONTAP systems using SnapMirror and SnapVault technologies. You simply need to identify the source volume and the destination volume, and then choose a replication policy and schedule.

For Cloud Volumes ONTAP, Cloud Manager purchases the required disks, configures relationships, applies the replication policy, and then initiates the baseline transfer between volumes.



The baseline transfer includes a full copy of the source data. Subsequent transfers contain differential copies of the source data.

Supported data protection configurations

Cloud Manager supports simple, fanout, and cascade data protection configurations:

- In a simple configuration, replication occurs from volume A to volume B.
- In a fanout configuration, replication occurs from volume A to multiple destinations.
- In a cascade configuration, replication occurs from volume A to volume B and from volume B to volume C.

Replicating data between systems

You can replicate data between ONTAP working environments by choosing a one-time data replication for data transfer, or a recurring schedule for disaster recovery or long-term retention. For example, you can set up data replication from an on-prem ONTAP system to Cloud Volumes ONTAP for disaster recovery.

Data replication requirements

Before you can replicate data, you should confirm that specific requirements are met for Cloud Volumes ONTAP, on-prem ONTAP clusters, or Amazon FSx for ONTAP.

Working environments

If you haven't done so already, you need to create the working environments for the source and target in the data replication relationship.

- [Create an Amazon FSx for ONTAP working environment](#)
- [Launch Cloud Volumes ONTAP in AWS](#)
- [Launch Cloud Volumes ONTAP in Azure](#)
- [Launch Cloud Volumes ONTAP in GCP](#)
- [Add existing Cloud Volumes ONTAP systems](#)
- [Discover ONTAP clusters](#)

Version requirements

You should verify that the source and destination volumes are running compatible ONTAP versions before replicating data. For details, see the [Data Protection Power Guide](#).

Requirements specific to Cloud Volumes ONTAP

- The instance's security group must include the required inbound and outbound rules: specifically, rules for ICMP and ports 11104 and 11105.

These rules are included in the predefined security group.

- To replicate data between two Cloud Volumes ONTAP systems in different subnets, the subnets must be routed together (this is the default setting).
- To replicate data between two Cloud Volumes ONTAP systems in different cloud providers, you must have a VPN connection between the virtual networks.

Requirements specific to ONTAP clusters

- An active SnapMirror license must be installed.
- If the cluster is on your premises, you should have a connection from your corporate network to your virtual network in AWS, Azure, or GCP. This is typically a VPN connection.
- ONTAP clusters must meet additional subnet, port, firewall, and cluster requirements.

For details, see the [Cluster and SVM Peering Express Guide](#).

Requirements specific to Amazon FSx for ONTAP

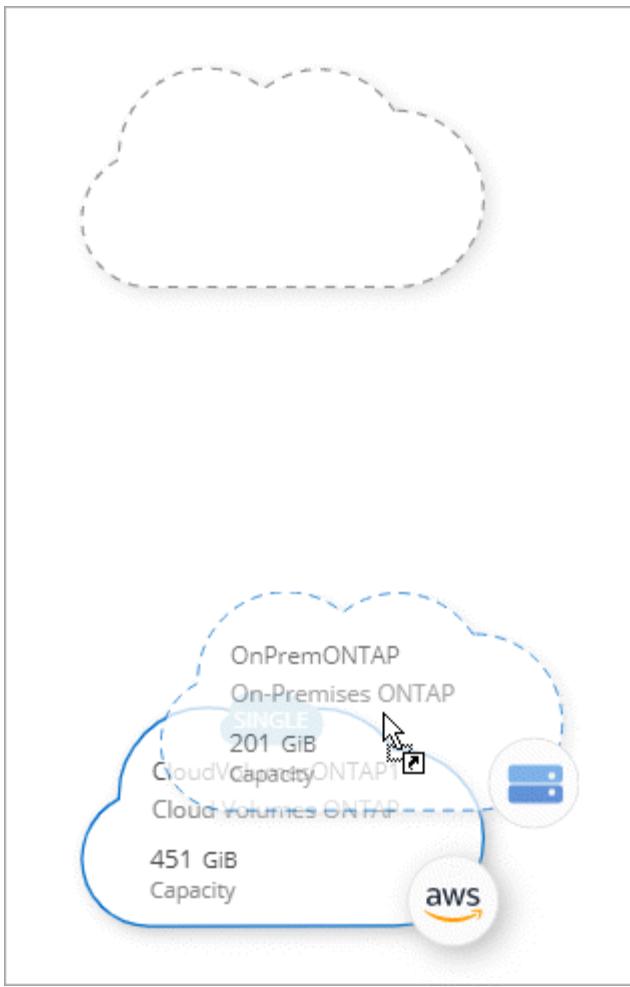
- An Amazon FSx for ONTAP working environment must be the target in the data replication relationship.
The source can be Cloud Volumes ONTAP or an on-prem ONTAP cluster.
- If Cloud Volumes ONTAP is the source, ensure connectivity between VPCs by enabling VPC peering or by using a Transit Gateway.
- If an on-prem ONTAP cluster is the source, ensure connectivity between your on-premises network and the AWS VPC by using a Direct Connect or VPN connection.

Setting up data replication between systems

You can replicate data by choosing a one-time data replication, which can help you move data to and from the cloud, or a recurring schedule, which can help with disaster recovery or long-term retention.

Steps

1. On the Canvas page, select the working environment that contains the source volume, and then drag it to the working environment to which you want to replicate the volume.



2. **Source and Destination Peering Setup:** If this page appears, select all of the intercluster LIFs for the cluster peer relationship.

The intercluster network should be configured so that cluster peers have *pair-wise full-mesh connectivity*, which means that each pair of clusters in a cluster peer relationship has connectivity among all of their intercluster LIFs.

These pages appear if an ONTAP cluster that has multiple LIFs is the source or destination.

3. **Source Volume Selection:** Select the volume that you want to replicate.
4. **Destination Disk Type and Tiering:** If the target is a Cloud Volumes ONTAP system, select the destination disk type and choose whether you want to enable data tiering.
5. **Destination Volume Name:** Specify the destination volume name and choose the destination aggregate.

If the destination is an ONTAP cluster, you must also specify the destination storage VM.

6. **Max Transfer Rate:** Specify the maximum rate (in megabytes per second) at which data can be transferred.

You should limit the transfer rate. An unlimited rate might negatively impact the performance of other applications and it might impact your internet performance.

7. **Replication Policy:** Choose a default policy or click **Additional Policies**, and then select one of the advanced policies.

For help, [learn about replication policies](#).

If you choose a custom backup (SnapVault) policy, the labels associated with the policy must match the labels of the Snapshot copies on the source volume. For more information, [learn how backup policies work](#).

8. **Schedule:** Choose a one-time copy or a recurring schedule.

Several default schedules are available. If you want a different schedule, you must create a new schedule on the *destination* cluster using System Manager.

9. **Review:** Review your selections and click **Go**.

Result

Cloud Manager starts the data replication process. You can view details about the volume relationship in the Replication service.

Managing data replication schedules and relationships

After you set up data replication between two systems, you can manage the data replication schedule and relationship from Cloud Manager.

Steps

1. Click **Replication**.
2. Review the status of the data replication relationships to verify that they are healthy.



If the Status of a relationship is idle and the Mirror State is uninitialized, you must initialize the relationship from the destination system for the data replication to occur according to the defined schedule. You can initialize the relationship by using System Manager or the command-line interface (CLI). These states can appear when the destination system fails and then comes back online.

3. Click the action menu for a volume relationship and choose one of the available actions.

The screenshot shows a table titled "1 Volume Relationship". The columns are: Health Status, Source Volume, Target Volume, Transfer Time, Status, Mirror State, and Last Successful Transfer. The first row contains: checked, vol1 OnPremONTAP, vol1_copy CloudVolumesONTAP1, 1 second, idle, snapmirrored, Apr 26, 2021, 1:53:54 PM, 0 B. To the right of the table is a vertical action menu with the following options: Information, Break, Reverse Resync, Edit Schedule, Edit Max Transfer Rate, Update, and Delete. The "Information" option is highlighted.

1 Volume Relationship						
Health Status	Source Volume	Target Volume	Transfer Time	Status	Mirror State	Last Successful Transfer
checked	vol1 OnPremONTAP	vol1_copy CloudVolumesONTAP1	1 second	idle	snapmirrored	Apr 26, 2021, 1:53:54 PM 0 B

- Information
- Break
- Reverse Resync
- Edit Schedule
- Edit Max Transfer Rate
- Update
- Delete

The following table describes the available actions:

Action	Description
Information	Shows you details about the volume relationship: transfer information, last transfer information, details about the volume, and information about the protection policy assigned to the relationship.
Break	<p>Breaks the relationship between the source and destination volumes, and activates the destination volume for data access.</p> <p>This option is typically used when the source volume cannot serve data due to events such as data corruption, accidental deletion, or an offline state.</p> <p>For information about configuring a destination volume for data access and reactivating a source volume, see the ONTAP 9 Volume Disaster Recovery Express Guide.</p>
Resync	<p>Reestablishes a broken relationship between volumes and resumes data replication according to the defined schedule.</p> <p> When you resynchronize the volumes, the contents on the destination volume are overwritten by the contents on the source volume.</p> <p>To perform a reverse resync, which resynchronizes the data from the destination volume to the source volume, see the ONTAP 9 Volume Disaster Recovery Express Guide.</p>
Reverse Resync	<p>Reverses the roles of the source and destination volumes. Contents from the original source volume are overwritten by contents of the destination volume. This is helpful when you want to reactivate a source volume that went offline.</p> <p>Any data written to the original source volume between the last data replication and the time that the source volume was disabled is not preserved.</p>
Edit Schedule	Enables you to choose a different schedule for data replication.
Edit Max Transfer Rate	Enables you to edit the maximum rate (in kilobytes per second) at which data can be transferred.
Update	Starts an incremental transfer to update the destination volume.
Delete	Deletes the data protection relationship between the source and destination volumes, which means that data replication no longer occurs between the volumes. This action does not activate the destination volume for data access. This action also deletes the cluster peer relationship and the storage VM (SVM) peer relationship, if there are no other data protection relationships between the systems.

Result

After you select an action, Cloud Manager updates the relationship or schedule.

Learn about replication policies

You might need help choosing a replication policy when you set up data replication in Cloud Manager. A replication policy defines how the storage system replicates data from a source volume to a destination volume.

What replication policies do

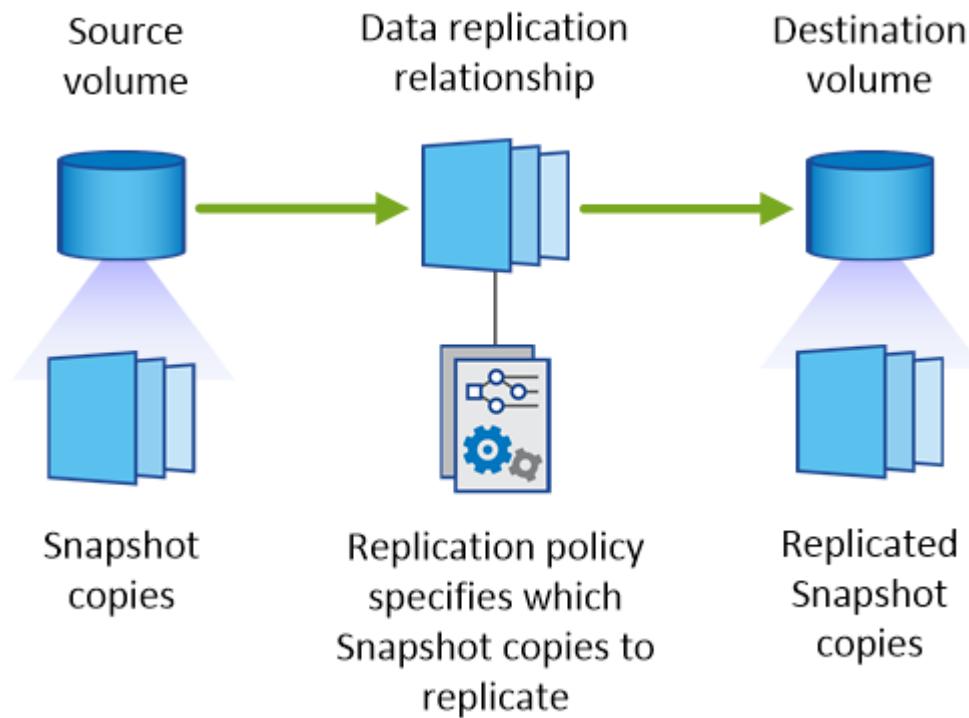
The ONTAP operating system automatically creates backups called Snapshot copies. A Snapshot copy is a read-only image of a volume that captures the state of the file system at a point in time.

When you replicate data between systems, you replicate Snapshot copies from a source volume to a destination volume. A replication policy specifies which Snapshot copies to replicate from the source volume to the destination volume.



Replication policies are also referred to as *protection* policies because they are powered by SnapMirror and SnapVault technologies, which provide disaster recovery protection and disk-to-disk backup and recovery.

The following image shows the relationship between Snapshot copies and replication policies:



Types of replication policies

There are three types of replication policies:

- A *Mirror* policy replicates newly created Snapshot copies to a destination volume.

You can use these Snapshot copies to protect the source volume in preparation for disaster recovery or for one-time data replication. You can activate the destination volume for data access at any time.

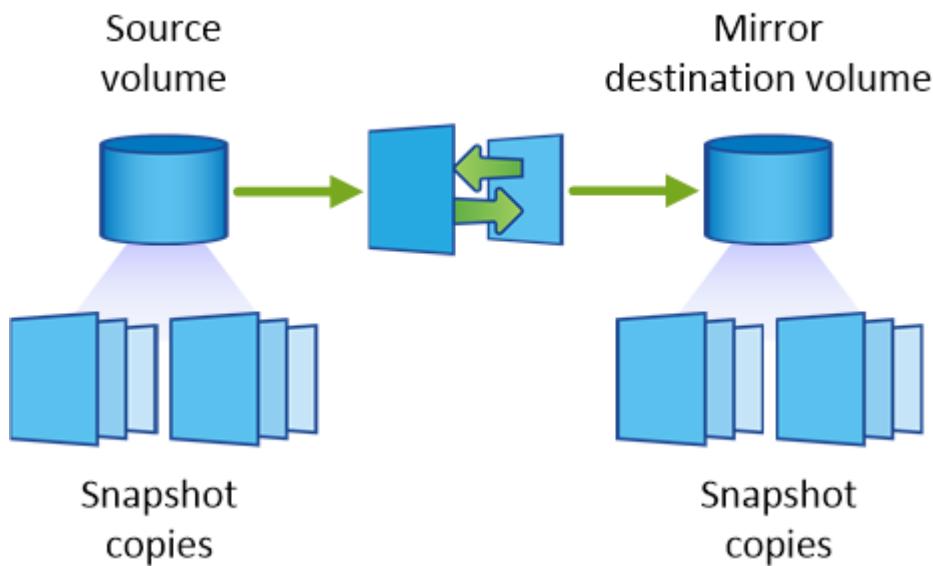
- A *Backup* policy replicates specific Snapshot copies to a destination volume and typically retains them for a longer period of time than you would on the source volume.

You can restore data from these Snapshot copies when data is corrupted or lost, and retain them for standards compliance and other governance-related purposes.

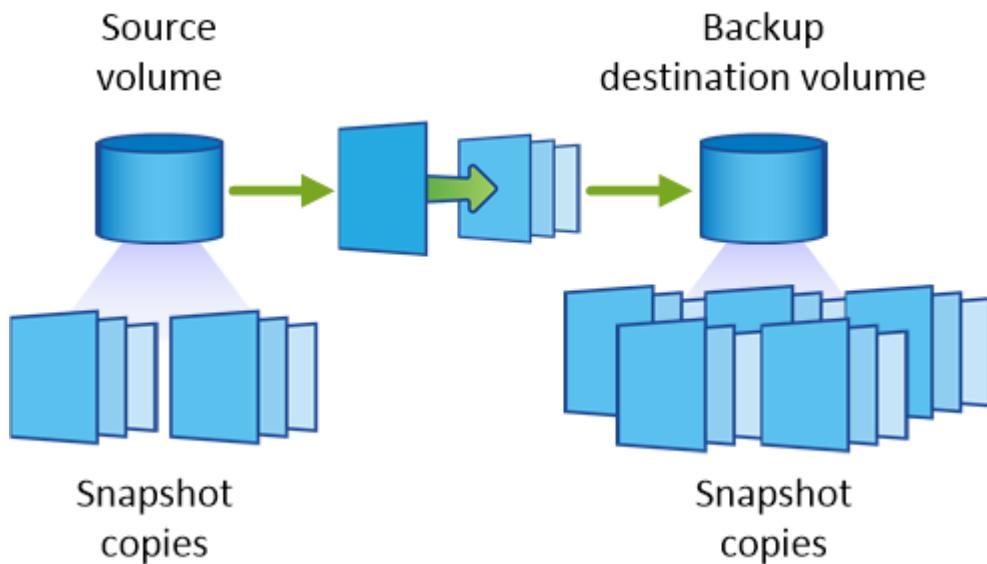
- A *Mirror and Backup* policy provides both disaster recovery and long-term retention.

Each system includes a default Mirror and Backup policy, which works well for many situations. If you find that you need custom policies, you can create your own using System Manager.

The following images show the difference between the Mirror and Backup policies. A Mirror policy mirrors the Snapshot copies available on the source volume.



A Backup policy typically retains Snapshot copies longer than they are retained on the source volume:



How Backup policies work

Unlike Mirror policies, Backup (SnapVault) policies replicate specific Snapshot copies to a destination volume. It is important to understand how Backup policies work if you want to use your own policies instead of the default policies.

Understanding the relationship between Snapshot copy labels and Backup policies

A Snapshot policy defines how the system creates Snapshot copies of volumes. The policy specifies when to create the Snapshot copies, how many copies to retain, and how to label them. For example, a system might

create one Snapshot copy every day at 12:10 a.m., retain the two most recent copies, and label them "daily".

A Backup policy includes rules that specify which labeled Snapshot copies to replicate to a destination volume and how many copies to retain. The labels defined in a Backup policy must match one or more labels defined in a Snapshot policy. Otherwise, the system cannot replicate any Snapshot copies.

For example, a Backup policy that includes the labels "daily" and "weekly" results in replication of Snapshot copies that include only those labels. No other Snapshot copies are replicated, as shown in the following image:

Default policies and custom policies

The default Snapshot policy creates hourly, daily, and weekly Snapshot copies, retaining six hourly, two daily, and two weekly Snapshot copies.

You can easily use a default Backup policy with the default Snapshot policy. The default Backup policies replicate daily and weekly Snapshot copies, retaining seven daily and 52 weekly Snapshot copies.

If you create custom policies, the labels defined by those policies must match. You can create custom policies using System Manager.

Use templates to standardize resource creation

Learn about Application Templates

The Application Templates service enables you to standardize resource creation in your working environments. For example, you can hard-code required parameters in a "volume template" that are later applied when a storage admin creates a volume. This can include required disk type, size, protocol, snapshot policy, cloud provider, and more. You can also turn on certain services, like Cloud Backup, for every created volume.

Templates make it easy for your storage admins to create volumes that are optimized for the workload requirements for each deployed application; such as databases, email, or streaming services. And it makes life easier for your storage architects knowing that each volume is created optimally for each application.

Features

Application Templates offer the following features and benefits:

- Automates and improves the design and development of your infrastructure
- Provides a single location to activate different NetApp Cloud services; like Cloud Backup and Cloud Data Sense
- Identifies resources that have been changed and are no longer compliant with the template (using the "drift" feature)

What is "drift"?

"Drift" allows Cloud Manager to monitor the parameter values used when a resource is created with the template. At this time, "drift" can identify when a resource has been changed so you can manually make adjustments to bring it back into compliance with the template. In the future we'll be able to send you notifications when a resource is out of compliance, or even reverse a user's change so that all resources created from a template are brought back into compliance automatically.

[Learn more about drift.](#)

Pricing and licenses

The Application Templates feature requires no licensing and is free to use by all Cloud Manager users.



Templates enable you to apply a cloud service onto a created resources, for example, enable Cloud Backup on every volume. In this case there is a cost for using the Backup service and for the object storage space used by the backup files.

Available template actions

A template is a chain of "actions" that have some pre-defined values. You can build templates that include the following actions:

Resource actions:

- Create a Cloud Volumes ONTAP volume

- Create an Azure NetApp Files volume
- Create an on-premises ONTAP volume
- Create a Cloud Volumes ONTAP working environment (single node on AWS)

Services actions:

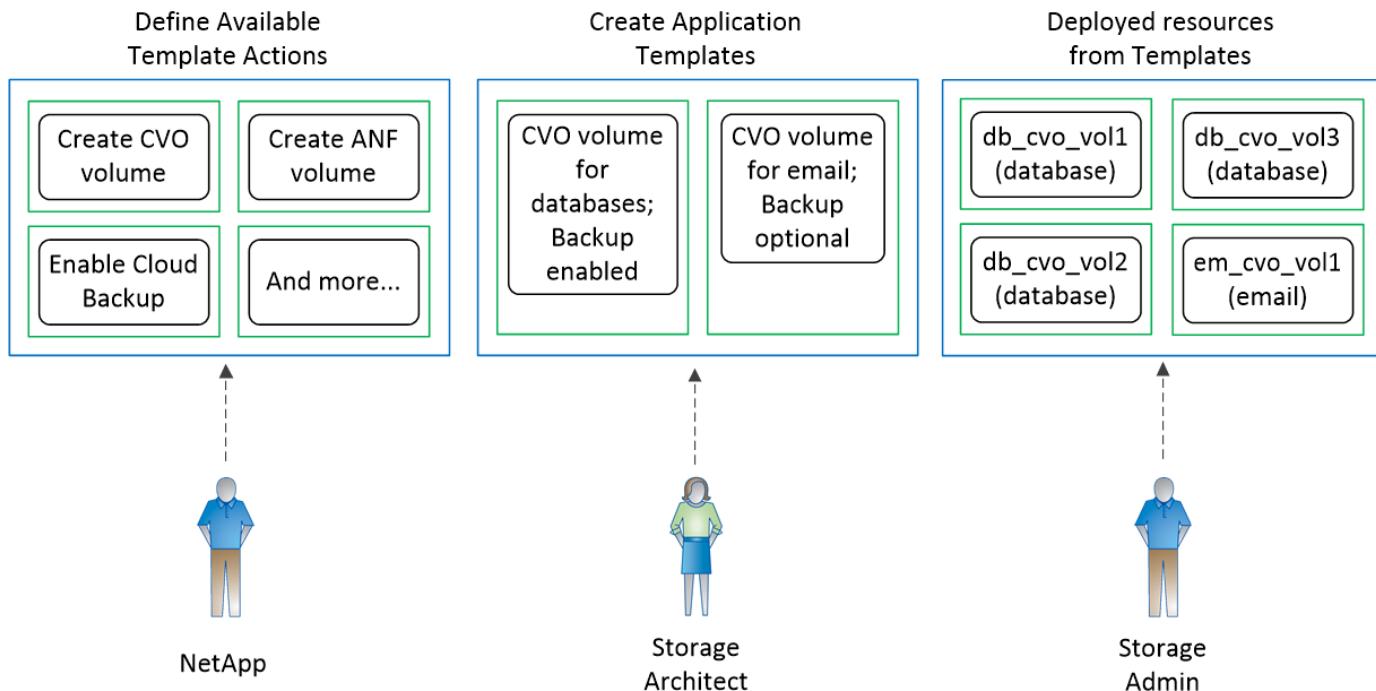
- Activate [Cloud Backup](#) on volumes (not applicable for Azure NetApp Files)
- Activate [Cloud Data Sense](#) on volumes
- Activate [Replication](#) on volumes (not applicable for Azure NetApp Files)

For example, you can create a template that creates a Cloud Volumes ONTAP volume. Or that that creates a Cloud Volumes ONTAP volume and then enables Cloud Backup on that volume. Or that that creates a Cloud Volumes ONTAP volume, and then enables Cloud Backup *and* Cloud Data Sense on that volume.

More actions will be added over time by NetApp.

How Application Templates work

The Application Templates service is made up of 3 parts. The available template "actions", the customized application template, and the deployed resource as a result of running the template. The following image shows the relationship between each component:



At a high level, Templates work like this:

1. NetApp defines the available template "actions".

For example, an "action" to create a Cloud Volumes ONTAP volume or an Azure NetApp Files volume.

2. Your storage architect selects the "actions" they want to use to create an Application Template, and then they hard-code certain values for the listed parameters.

For example, they select high speed disks and a large amount of RAM for Cloud Volumes ONTAP volumes that will be used to carry the workloads for Oracle databases. And they require that backups are made for each volume.

3. Your storage admins use the templates to create resources that are optimized for the application they will be used for.

For example, they create a volume that will be used for an Oracle database by using the volume template created for databases.

4. The service tracks certain resource settings defined in the template using the "drift" feature as determined by your storage architect.

Limitations

- The Application Templates service is not supported in any of the Gov Cloud regions or in "dark" sites.
- You can't use a template to create a Cloud Volumes ONTAP volume on an existing aggregate. New volumes are created in a new aggregate.

Template building blocks

There are certain features you can use when building a template that enable you pass values between actions (like a volume name), conditionally branch to connect actions together (enable backups on a new volume), and that help your users customize resources when using the template.

Special template controls

Before you start creating your template, you should understand some special options that you can set when pre-populating a value for a parameter in a template.

Enable Storage Efficiency	<input type="checkbox"/> Editable	<input type="checkbox"/> Drift
<input checked="" type="radio"/> Storage Efficiency	<input type="radio"/> No Storage Efficiency	

Editable checkbox

Check this box to let the storage admin override the pre-populated value you have entered in the template. This gives the storage admin a suggestion for what the value should be, but it allows them to customize the value when creating the resource.

When unchecked, the user can't change the value and the hard-coded value in the template is always used when the admin deploys a resource.

Drift checkbox

Check this box so that Cloud Manager monitors the hard-coded value you entered for a parameter when a resource is created with the template. Later, you can run a Drift Report to see which fields configured with Drift are no longer compliant with the template settings.

When unchecked, the user can change the value to any value after the resource has been created.



For the drift feature to work, after you have defined drift for some parameters in the template, you must enable the drift feature for the template. This is the last step when creating a template. Drift doesn't work if it is enabled for a parameter but hasn't been enabled on the template.

Using a regular expression (regex) in fields

There are a few fields within templates that allow you to enter a regex to define the value that your admin can enter in the field; for example "Volume Name" and "Share Name".

The screenshot shows a user interface for defining a template field. The field is labeled "Volume Name". Below the label is a text input box containing the regular expression "Regex: ^[a-zA-Z][0-9a-zA-Z]{0,149}\$". A red box highlights this input field. To the right of the input box are standard UI controls: a close button ("X") and a dropdown arrow.

As an example, if you enter "`^[a-zA-Z][0-9a-zA-Z]{0,149}$`" as the regex for the volume name, it means that *"the name should start with an alphabetic character, it can contain only numbers, letters, or the underscore, and it should be 150 or fewer characters in length"*.

Pass values between template actions

Templates have the ability to use information from a previous action to populate a field in a future action. For example, when defining the name of the volume that will have Cloud Backup functionality enabled, you can instruct the Backup action to use the value the storage admin entered as the name of the volume from the Create Cloud Volumes ONTAP action.

There are three types of information that the AppTemplate service can use:

- Input value - This is the actual value the storage admin entered into a field in a previous template action.
- Output value - This is the value Cloud Manager generates after creating a resource from a previous template action.
- Enter your own value - This is a value that you enter; it is not accessed from a previous action in the template.

For example, to enable compliance scanning on a volume, the Cloud Data Sense service needs both the "volume name" that the storage admin enters (the Input value), and the "volume uuid" that Cloud Manager generates when it creates the volume (the Output value).

The following illustration shows how to enter this information in the Cloud Data Sense action section of the template.

Activate Cloud Data Sense on Volume

Action Definition

Details

Volume Name 	<input type="checkbox"/> Editable <input type="checkbox"/> Drift
Get input value from action	
Create Volume in On-Premises ONTAP	
Volume Name 	
Volume UUID 	<input type="checkbox"/> Editable <input type="checkbox"/> Drift
Get output value from action	
Create Volume in On-Premises ONTAP	
uuid	

Get the volume name:

From the value the storage admin entered...

In the earlier action used to create the volume...

in the "Volume Name" field.

Get the volume UUID:

From the value Cloud Manager generates...

In the earlier action used to create the volume...

in the "uuid" field.

Use conditions to perform different actions based on logic statements

Conditions tell the template to execute different actions depending on whether the condition is true or false when the storage admin runs the template. You add a condition by selecting the **Logical If** action:

Add New Action

X

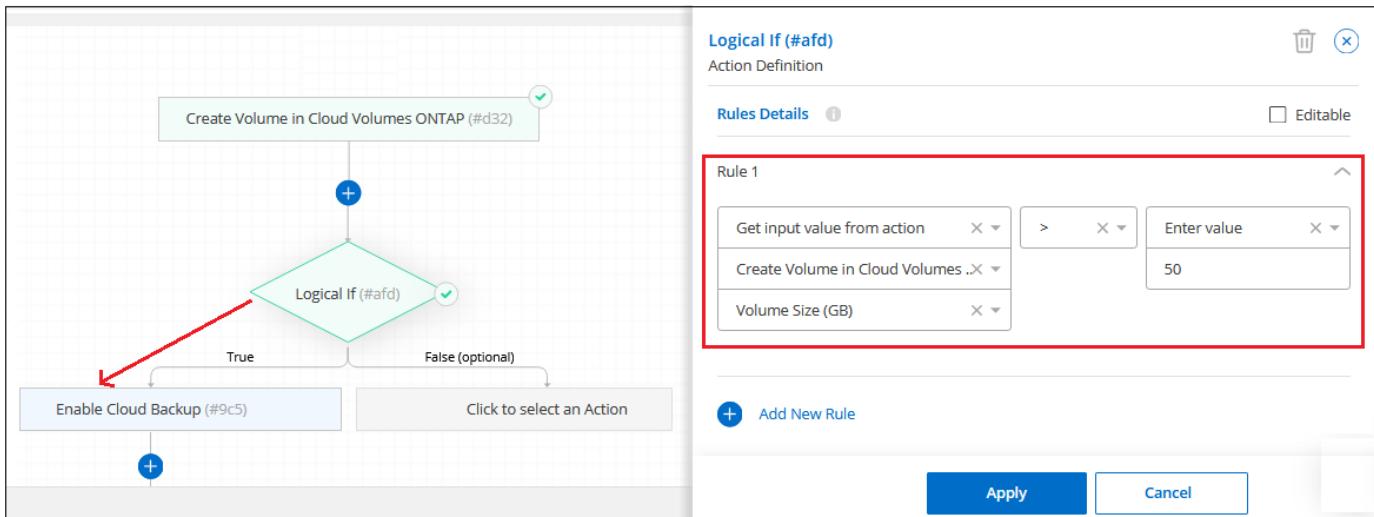
ACTIONS - SERVICES

- Activate Cloud Data Sense on Volume
- Enable Cloud Backup
- Enable Volume Replication to Cloud Volumes ONTAP
- Enable Volume Replication to On-Premises ONTAP

FLOW CONTROLS

- Logical If

For example, you may have a guideline that if a volume has a capacity larger than 50 GB then Cloud Backup is required to be enabled on that volume. If the volume has a smaller capacity, then Cloud Backup is not enabled. You can implement this in your template as shown below.



Conditions consist of two parts:

- Rules - The item you are checking for being either true or false.
- And/Or statement - Allows you to use multiple rules to further refine whether additional actions are added.

A Rule is made up of three parts:

Source field - The location from which you will get the value to compare.

- Get input value from action - The actual value the storage admin entered into a field in a previous template action.
- Get output value from action - The value Cloud Manager generated after creating a resource from a previous template action.
- Enter value - This is a value that you enter; it is not accessed from a previous action in the template. This can be a value from a resource that already exists; for example an existing volume.

Operator - The operator used for the comparison. The options are **Equal**, **Not Equal**, **Greater Than**, **Less Than**, **Greater Than or Equal**, **Less Than or Equal**.

Field value - The actual value you are comparing. The options are the same as those for the *Source field*.

An And/Or statement enables you to conditionally add more actions for users when they run the template based on whether multiple rules are evaluated as True or False. **And** requires all rules to be true or false, and **Or** requires just one of the rules needs to be true or false.

When using both an And and Or statement with your rules, the evaluation process follows standard mathematical order where “AND” precedes “OR”. For example:

- <Rule1> OR <Rule2> AND <Rule3>

This statement is evaluated in the following order:
<Rule1> OR (<Rule2> AND <Rule3>)

Build application templates for your organization

Select one or more of the NetApp-provided "actions" and quickly build an application

template that your organization can use to start optimizing the creation of resources.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.



Verify required prerequisites

- Before users can create a volume for a Cloud Volumes ONTAP, on-premises ONTAP, or Azure NetApp Files system using a template, make sure they have access to an appropriate working environment where the volume will be deployed.
- If you plan to add a Cloud service "action" to your template, such as [Cloud Backup](#) or [Cloud Data Sense](#), ensure that the service is active and licensed in your environment.



Launch the Application Templates service

Select the **AppTemplate** service, click the **Editor** tab, and select the template.



Build the template by selecting "actions" and defining parameters

Follow the creation steps and define the actions that will be performed by the template.

Requirements

Read the following requirements to make sure that you have a supported configuration.

- If you don't already have a Connector, [see how to create Connectors](#) for AWS, Azure, and GCP.
- When creating a Cloud Volumes ONTAP volume template, make sure you have a Cloud Volumes ONTAP working environment available for your users. See how to launch a Cloud Volumes ONTAP system in [AWS](#), [Azure](#), or in [GCP](#).
- When creating an on-premises ONTAP volume template, make sure you have an on-premises ONTAP working environment available for your users. See how to [discover an on-premises ONTAP system](#) in Cloud Manager.
- When creating an Azure NetApp Files volume template, make sure you have an Azure NetApp Files working environment available for your users. See how to [create an Azure NetApp Files working environment](#) in Cloud Manager.
- If you plan to enable Cloud Backup in the template, ensure that your environment has an active and licensed Cloud Backup service.
- If you plan to enable Cloud Data Sense in the template, ensure that your environment has an active and licensed Cloud Data Sense service.
- If you plan to enable Replication in the template, and the template is for an on-premises ONTAP volume, the ONTAP cluster must have an active SnapMirror license.

Examples of creating resources using templates

Resource templates enable you to create new volumes or a new Cloud Volumes ONTAP working environment.

Create a template for a Cloud Volumes ONTAP volume

See [how to provision Cloud Volumes ONTAP volumes](#) for details about all the parameters you need to complete in the Cloud Volumes ONTAP volume template.

For this example we'll create a template named "CVO volume for databases" and include the following 2 actions:

- Create Cloud Volumes ONTAP Volume

Make the volume for the AWS environment, configure it with 100 GB of storage, set the Snapshot Policy to "default", and enable Storage Efficiency.

- Enable Cloud Backup

Create daily backups with a retention value of 30 copies.

Steps

1. Select the **AppTemplate** service, click the **Templates** tab, and click **Add New Template**.

The *Select_a_Template* page is displayed.

Select a Template to Get Started

Choose a template for the type of resource you want to create and then customize the parameters for what's required in your environment.

For even greater flexibility, choose the "blank" template to mix and match configurations and to apply additional Cloud Manager services.

Start with a template of:

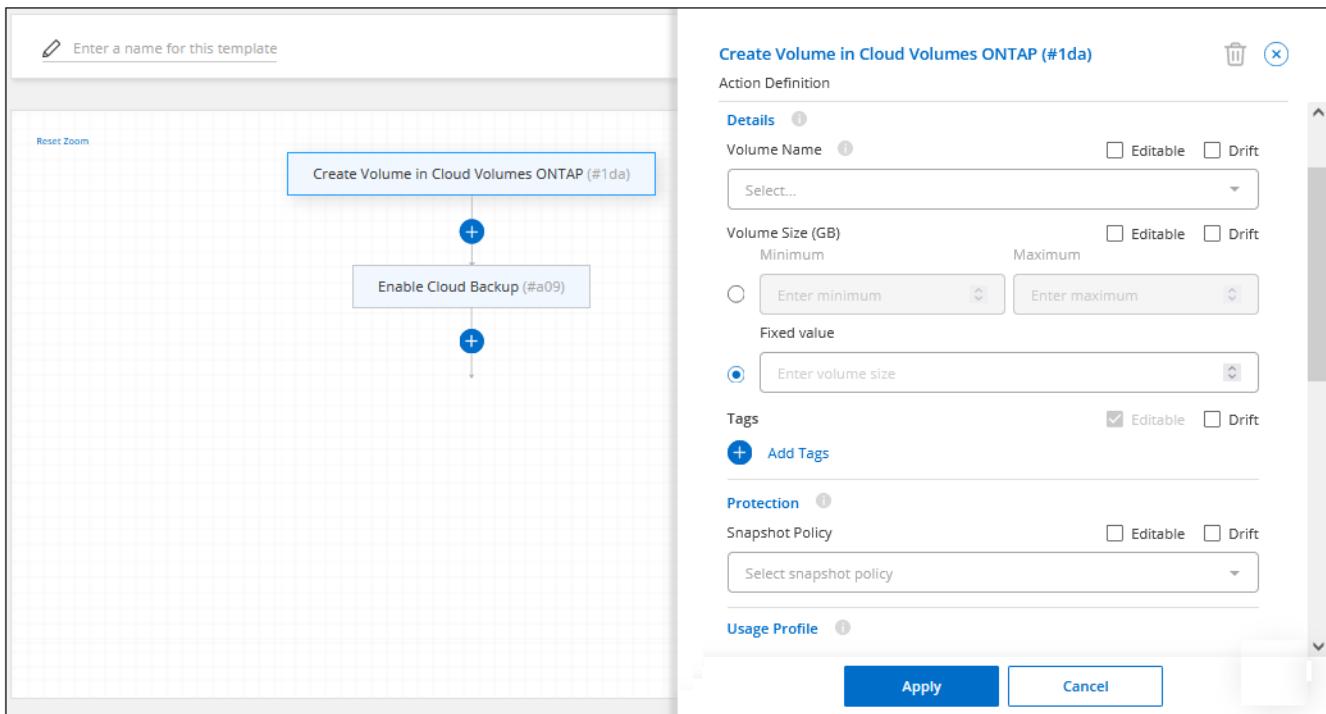
Blank template **Volume for Cloud Volumes ONTAP** **Volume for Cloud Volumes ONTAP + Backup** **Volume for Azure NetApp files**

Get Started **Cancel**

The screenshot shows a user interface for selecting a template. At the top, a title says 'Select a Template to Get Started'. Below it, there's a descriptive paragraph about choosing a template for the type of resource. Then, another paragraph suggests using a 'blank' template for flexibility. Under the heading 'Start with a template of:', there are four options arranged in a grid. The first two are in one row, and the last two are in the next. The third option, 'Volume for Cloud Volumes ONTAP + Backup', is highlighted with a blue border and a small checkmark icon in the top right corner of its box. The other three options are plain boxes: 'Blank template', 'Volume for Cloud Volumes ONTAP', and 'Volume for Azure NetApp files'. At the bottom, there are two large buttons: a blue 'Get Started' button on the left and a white 'Cancel' button on the right.

2. Select **Volume for Cloud Volumes ONTAP + Backup** as the type of resource you want to create, and click **Get Started**.

The *Create Volume in Cloud Volumes ONTAP Action Definition* page is displayed.



3. **Action Name:** Optionally, enter a customized action name instead of the default value.

4. **Details:** Enter the volume name and size.

Field	Description
Volume Name	<p>Click in the field and select one of the 5 options. You can let the admin enter any name by selecting Free Text, or you can specify that the volume name must have a certain prefix or suffix, that it <i>contains</i> certain characters, or that it follows rules from a regular expression (regex) you enter.</p> <p>For example, you could specify that "db" be a required prefix, suffix, or contains; requiring the user to add volume names like "db_vol1", "vol1_db", or "vol_db_1".</p>
Volume Size	You can specify a range of allowable values, or you can specify a fixed size. This value is in GB. For our example we can add a fixed value 100 .
Tags	Enter a name and value pair for a tag that you want to associate with this volume. For example, you could add "Cost Center" as the tag name and the cost center code "6655829" as the value. You can associate more than one tag with a volume by adding more tag name and value pairs.

5. **Protection:** Choose whether this volume will have Snapshot copies created by selecting "Default" or some other policy, or choose "None" if you do not want to create Snapshot copies.
6. **Usage Profile:** Choose whether or not NetApp storage efficiency features are applied to the volume. This includes Thin Provisioning, Deduplication, and Compression. For our example, keep storage efficiency enabled.
7. **Disk Type:** Choose the cloud storage provider and the type of disk. For some disk selections you can also select a minimum and maximum IOPS or Throughput (MB/s) value; basically defining a certain Quality of Service (QoS).
8. **Protocol Options:** Select **NFS** or **SMB** to set the protocol of the volume. And then provide the protocol details.

NFS Fields	Description
Access Control	Choose whether access controls are needed to access the volume.
Export Policy	Create an export policy to define the clients in the subnet that can access the volume.
NFS Version	Select the NFS version for the volume: either <i>NFSv3</i> or <i>NFSv4</i> , or you can select both.

SMB Fields	Description
Share Name	Click in the field and select one of the 5 options. You can let the admin enter any name (Free Text) or you can specify that the share name must have a certain prefix or suffix, that it <i>contains</i> certain characters, or that it follows rules from a regular expression (regex) you enter.
Permissions	Select the level of access to a share for users and groups (also called access control lists, or ACLs).
Users / Groups	Specify local or domain Windows users or groups, or UNIX users or groups. If you specify a domain Windows user name, you must include the user's domain using the format domain\username.

9. **Tiering:** Choose the tiering policy that you would like applied to the volume, or set this to "None" if you do not want to tier cold data from this volume to object storage.

See [volume tiering policies](#) for an overview, and see [Tiering inactive data to object storage](#) to make sure your environment is set up for tiering.

10. **Context:** Enter the Cloud Volumes ONTAP working environment context; if required.

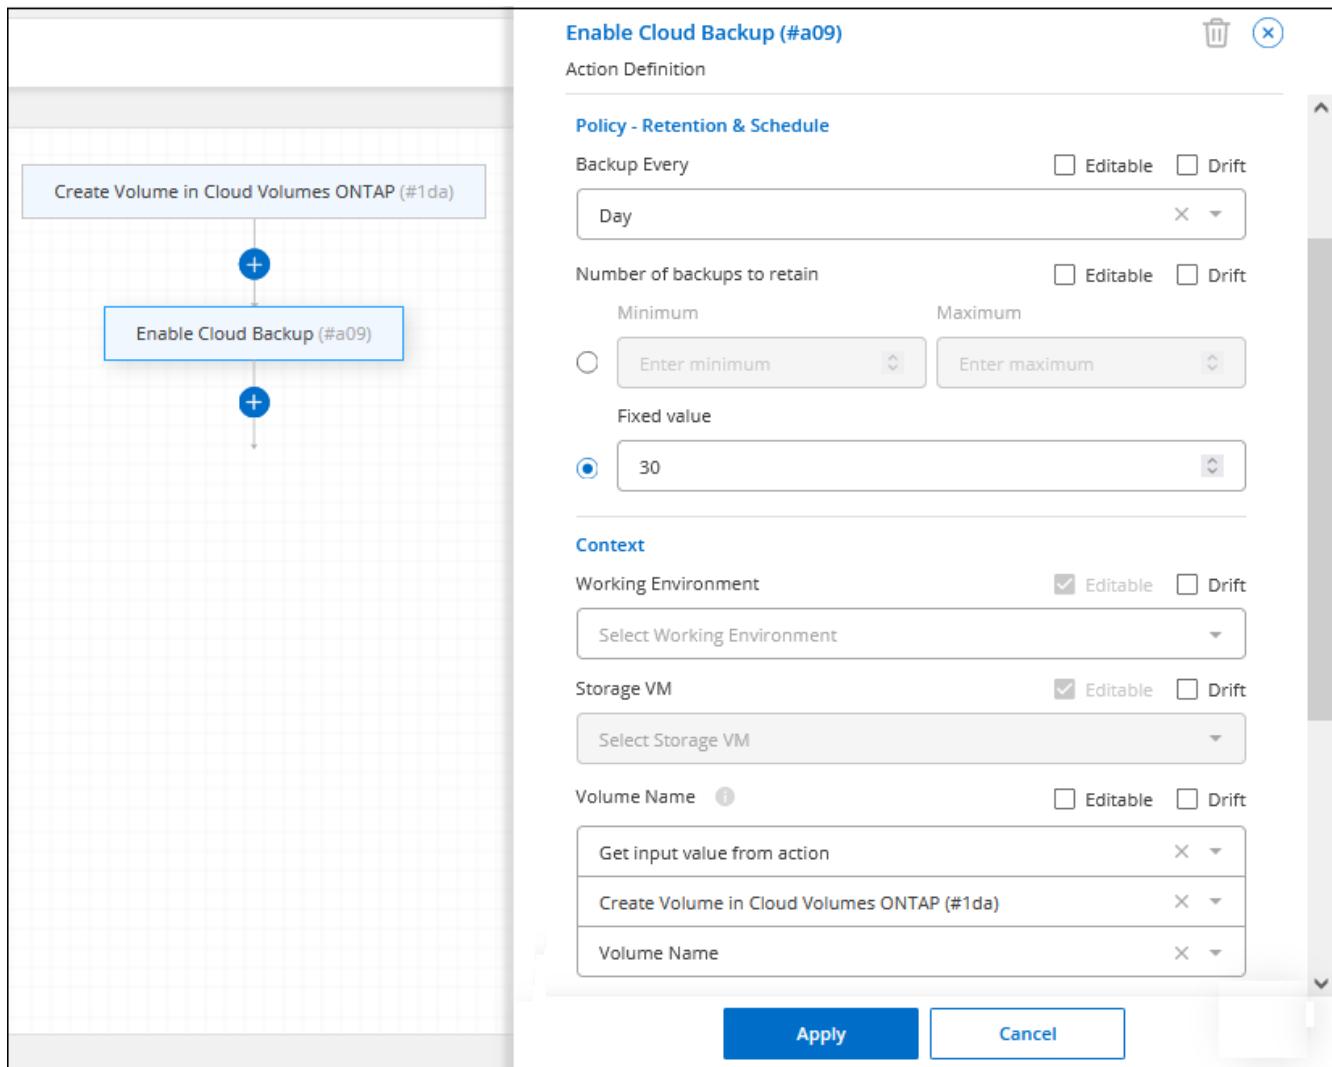
When users launch the template from an existing working environment, this information gets filled in automatically.

When users launch the template from the Templates Dashboard (not in a working environment context), then they need to select the working environment and the SVM where the volume will be created. That's why these fields are marked as "Editable".

11. Click **Apply** after you have defined the parameters needed for this action.

If the template values are correctly completed, a green checkmark is added to the "Create Volume in Cloud Volumes ONTAP" box.

12. Click the **Enable Cloud Backup** box and the *Enable Cloud Backup Action Definition* dialog is displayed so you can fill in the Cloud Backup details.



13. Define the backup policy to create daily backups with a 30-day retention value.
 14. Below the Volume Name field there are three fields you use to indicate which volume will have backup enabled. See [how to complete these fields](#).
 15. Click **Apply** and the Cloud Backup dialog is saved.
 16. Enter the template name **CVO volume for databases** (for this example) in the top left.
 17. Click **Settings & Drift** to provide a more detailed description so that this template can be distinguished from other similar templates, and so you can enable Drift for the overall template, and then click **Apply**.
- Drift allows Cloud Manager to monitor the hard-coded values you entered for parameters when creating this template.
18. Click **Save Template**.

Result

The template is created and you are returned to the Templates Dashboard where your new template appears.

See [what you should tell your users about templates](#).

Create a template for an Azure NetApp Files volume

Creating a template for an Azure NetApp Files volume is done in the same manner as creating a template for a Cloud Volumes ONTAP volume.

See [how to provision Azure NetApp Files volumes](#) for details about all the parameters you need to complete in the ANF volume template.

Steps

1. Select the **AppTemplate** service, click the **Templates** tab, and click **Add New Template**.

The *Select_a_Template* page is displayed.

Select a Template to Get Started

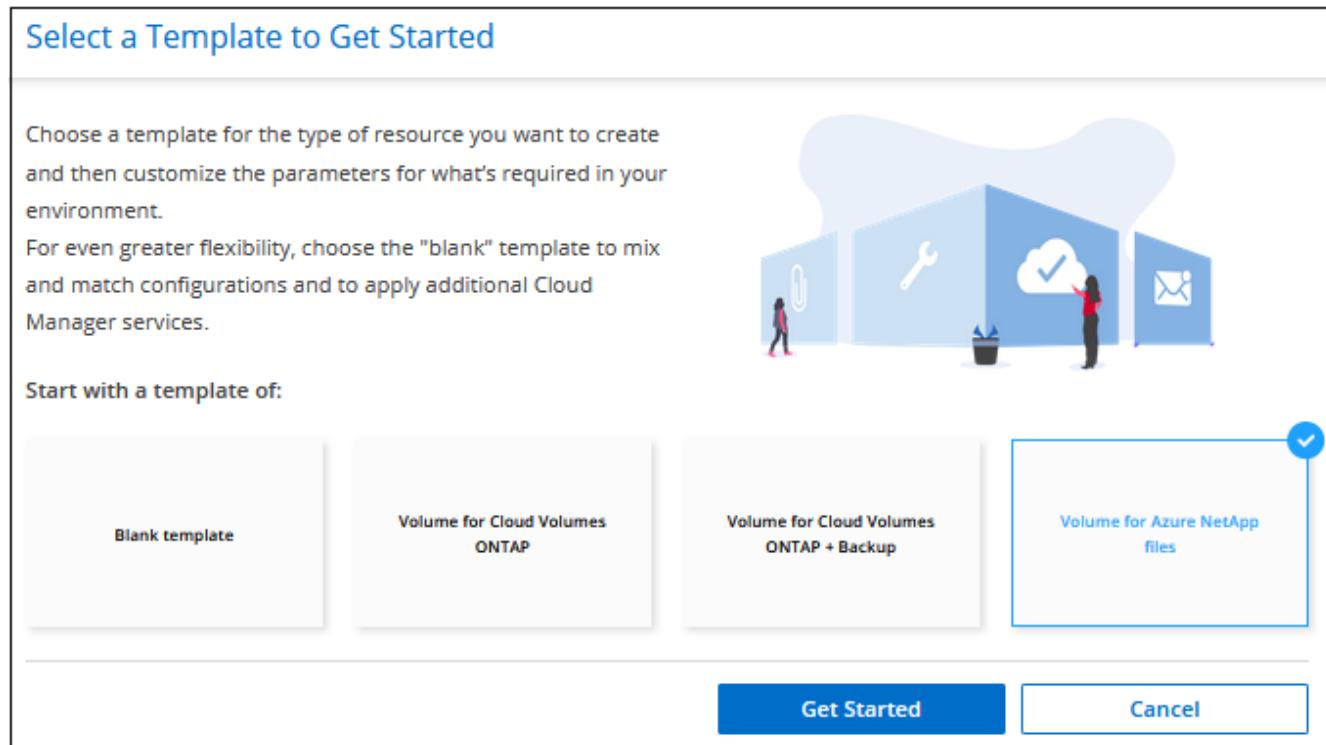
Choose a template for the type of resource you want to create and then customize the parameters for what's required in your environment.

For even greater flexibility, choose the "blank" template to mix and match configurations and to apply additional Cloud Manager services.

Start with a template of:

Blank template **Volume for Cloud Volumes
ONTAP** **Volume for Cloud Volumes
ONTAP + Backup** **Volume for Azure NetApp
files**

Get Started **Cancel**



2. Select **Volume for Azure NetApp Files** as the type of resource you want to create, and click **Get Started**.

The *Create Volume in Azure NetApp Files Action Definition* page is displayed.

The screenshot shows a configuration dialog for creating a volume in Azure NetApp Files. The left side of the dialog displays a template editor interface with a 'Create Volume in Azure NetApp Files (#23c)' action definition selected. The right side is the configuration form for this action.

Action Definition

Volume Details

- Volume Name:** A dropdown menu with options: Select..., Editable (unchecked), Drift (unchecked).
- Volume Size (GB):** A section with 'Minimum' and 'Maximum' fields (both Editable unchecked, Drift unchecked) and a 'Fixed value' field (selected, Editable checked, Drift unchecked). The 'Enter volume size in GB' input field contains 'Enter volume size in GB'.
- Tags:** A section with 'Add Tags' (button with plus sign) and 'Tags' (checkbox checked, Editable checked, Drift unchecked).
- Protocol:** A section with 'Protocol' (radio buttons: NFSv3 selected, NFSv4.1, SMB), 'Editable' (unchecked), and 'Drift' (unchecked).
- Volume Path:** A dropdown menu with 'Select...' (Editable unchecked, Drift unchecked).

Buttons: Apply (blue button) and Cancel (white button).

3. **Action Name:** Optionally, enter a customized action name instead of the default value.
4. **Volume Details:** Enter a volume name and size, and optionally specify tags for the volume.

Field	Description
Volume Name	<p>Click in the field and select one of the 5 options. You can let the admin enter any name by selecting Free Text, or you can specify that the volume name must have a certain prefix or suffix, that it <i>contains</i> certain characters, or that it follows rules from a regular expression (regex) you enter.</p> <p>For example, you could specify that "db" be a required prefix, suffix, or contains; requiring the user to add volume names like "db_vol1", "vol1_db", or "vol_db_1".</p>
Volume Size	You can specify a range of allowable values, or you can specify a fixed size. This value is in GB.
Tags	Enter a name and value pair for a tag that you want to associate with this volume. For example, you could add "Cost Center" as the tag name and the cost center code "6655829" as the value. You can associate more than one tag with a volume by adding more tag name and value pairs.

5. **Protocol:** Select **NFSv3**, **NFSv4.1**, or **SMB** to set the protocol of the volume. And then provide the protocol details.

NFS Fields	Description
Volume Path	Select one of the 5 options. You can let the admin enter any path by selecting Free Text , or you can specify that the path name must have a certain prefix or suffix, that it <i>contains</i> certain characters, or that it follows rules from a regular expression (regex) you enter.

NFS Fields	Description
Export Policy Rules	Create an export policy to define the clients in the subnet that can access the volume.

SMB Fields	Description
Volume Path	Select one of the 5 options. You can let the admin enter any path by selecting Free Text , or you can specify that the path name must have a certain prefix or suffix, that it <i>contains</i> certain characters, or that it follows rules from a regular expression (regex) you enter.

6. **Context:** Enter the Azure NetApp Files working environment, details for a new or an existing Azure NetApp Files account, and other details.

Field	Description
Working Environment	When storage admin users launch the template from an existing working environment, this information gets filled in automatically. When users launch the template from the Templates Dashboard (not in a working environment context), then they need to select the working environment where the volume will be created.
NetApp Account Name	Enter the name you want to use for the account.
Azure Subscription ID	Enter the Azure Subscription ID. This is the full ID in a format similar to "2b04f26-7de6-42eb-9234-e2903d7s327".
Region	Enter the region using the internal region name .
Resource Group Name	Enter the name of the Resource Group you want to use.
Capacity Pool Name	Enter the name of an existing capacity pool.
Subnet	Enter the VNet and subnet. This value includes the full path, in a format similar to "/subscriptions/<subscription_id>/resourceGroups/<resource_group>/providers/Microsoft.Network/virtualNetworks/<vpc_name>/subnets/<subhet_na me>".

7. **Snapshot Copy:** Enter the Snapshot ID for an existing volume Snapshot if you want this new volume to be created using characteristics from an existing volume.
8. Click **Apply** after you have defined the parameters needed for this action.
9. Enter the name you want to use for the template in the top left.
10. Click **Settings & Drift** to provide a more detailed description so that this template can be distinguished from other similar templates, and so you can enable Drift for the overall template, and then click **Apply**.

Drift allows Cloud Manager to monitor the hard-coded values you entered for parameters when creating this template.
11. Click **Save Template**.

Result

The template is created and you are returned to the Templates Dashboard where your new template appears.

See [what you should tell your users about templates](#).

Create a template for an on-premises ONTAP volume

See [how to provision on-premises ONTAP volumes](#) for details about all the parameters you need to complete in the on-premises ONTAP volume template.

Steps

1. Select the **AppTemplate** service, click the **Templates** tab, and click **Add New Template**.

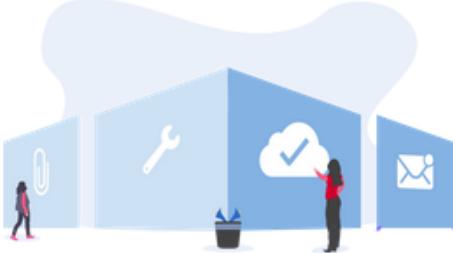
The *Select_a_Template* page is displayed.

Select a Template to Get Started

Choose a template for the type of resource you want to create and then customize the parameters for what's required in your environment.

For even greater flexibility, choose the "blank" template to mix and match configurations and to apply additional Cloud Manager services.

Start with a template of:

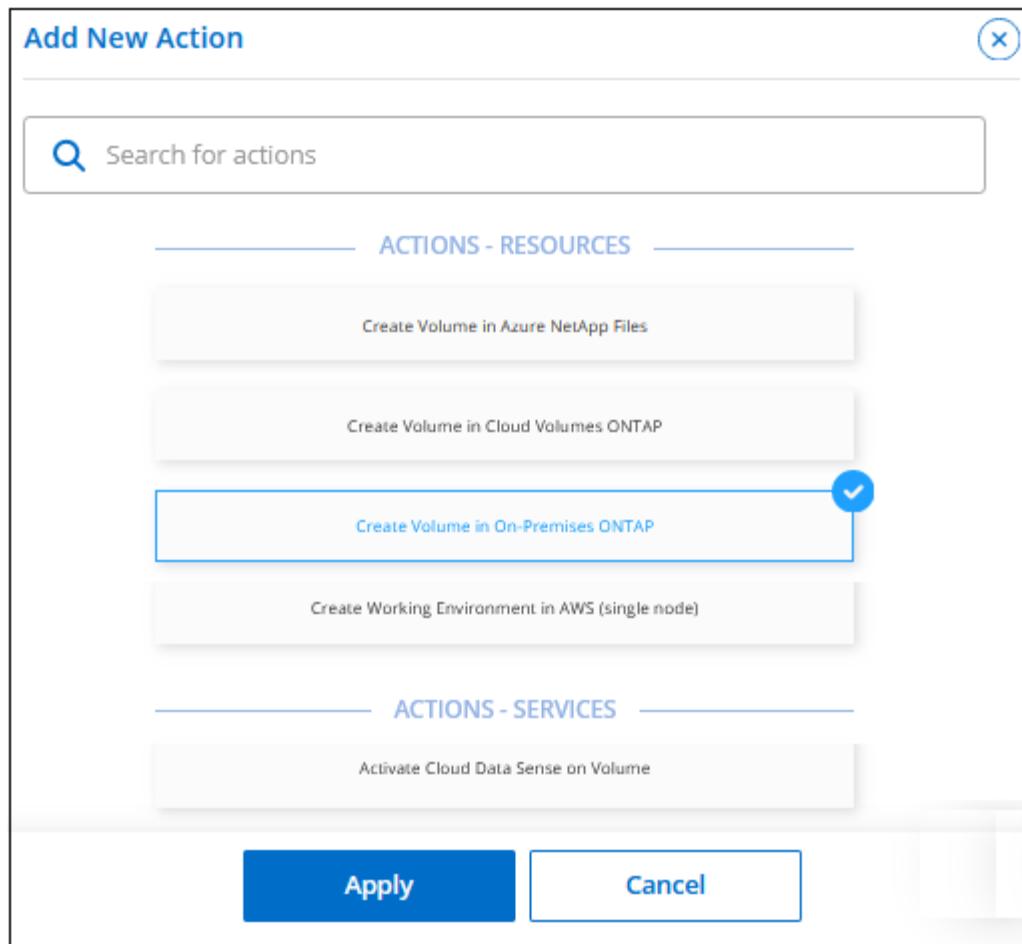


- Blank template
- Volume for Cloud Volumes
ONTAP + Backup
- Volume for Cloud Volumes
ONTAP + Backup + Compliance
- Volume for Azure NetApp files +
Compliance

Get Started **Cancel**

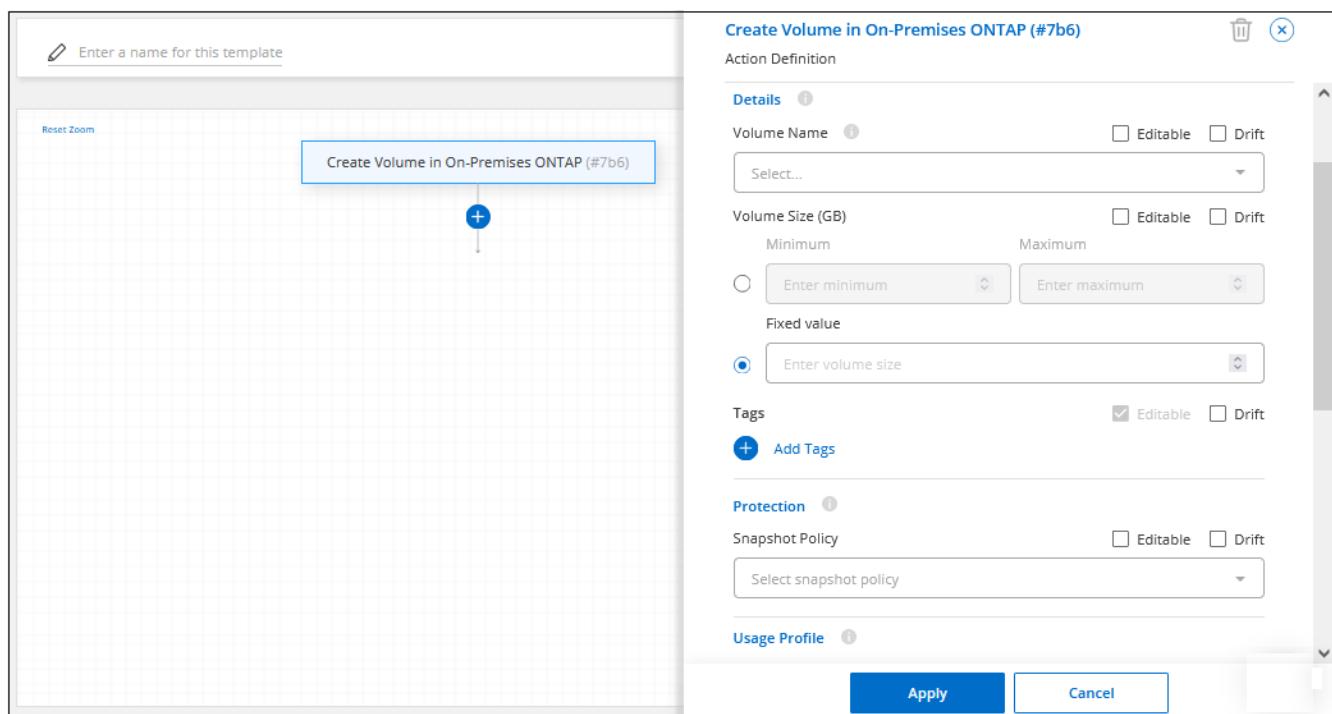
2. Select **Blank template** and click **Get Started**.

The *Add New Action* page is displayed.



3. Select **Create Volume in On-Premises ONTAP** as the type of resource you want to create, and click **Apply**.

The *Create Volume in On-Premises ONTAP Action Definition* page is displayed.



4. **Action Name:** Optionally, enter a customized action name instead of the default value.

5. **Details:** Enter the volume name and size.

Field	Description
Volume Name	<p>Click in the field and select one of the 5 options. You can let the admin enter any name by selecting Free Text, or you can specify that the volume name must have a certain prefix or suffix, that it <i>contains</i> certain characters, or that it follows rules from a regular expression (regex) you enter.</p> <p>For example, you could specify that "db" be a required prefix, suffix, or contains; requiring the user to add volume names like "db_vol1", "vol1_db", or "vol_db_1".</p>
Volume Size	You can specify a range of allowable values, or you can specify a fixed size. This value is in GB. For our example we can add a fixed value 100 .
Tags	Enter a name and value pair for a tag that you want to associate with this volume. For example, you could add "Cost Center" as the tag name and the cost center code "6655829" as the value. You can associate more than one tag with a volume by adding more tag name and value pairs.

6. **Protection:** Choose whether this volume will have Snapshot copies created by selecting "Default" or some other policy, or choose "None" if you do not want to create Snapshot copies.

7. **Usage Profile:** Choose whether or not NetApp storage efficiency features are applied to the volume. This includes Thin Provisioning, Deduplication, and Compression.

8. **Protocol Options:** Select **NFS** or **SMB** to set the protocol of the volume. And then provide the protocol details.

NFS Fields	Description
Access Control	Choose whether access controls are needed to access the volume.
Export Policy	Create an export policy to define the clients in the subnet that can access the volume.
NFS Version	Select the NFS version for the volume: either <i>NFSv3</i> or <i>NFSv4</i> , or you can select both.

SMB Fields	Description
Share Name	Click in the field and select one of the 5 options. You can let the admin enter any name (Free Text) or you can specify that the share name must have a certain prefix or suffix, that it <i>contains</i> certain characters, or that it follows rules from a regular expression (regex) you enter.
Permissions	Select the level of access to a share for users and groups (also called access control lists, or ACLs).
Users / Groups	Specify local or domain Windows users or groups, or UNIX users or groups. If you specify a domain Windows user name, you must include the user's domain using the format domain\username.

9. **Context:** Enter the on-premises ONTAP working environment context; if required.

When users launch the template from an existing working environment, this information gets filled in automatically.

When users launch the template from the Templates Dashboard (not in a working environment context), then they need to select the working environment, the SVM, and the aggregate where the volume will be created.

10. Click **Apply** after you have defined the parameters needed for this action.

If the template values are correctly completed, a green checkmark is added to the "Create Volume in On-Premises ONTAP" box.

11. Enter the template name in the top left.

12. Click **Settings & Drift** to provide a more detailed description so that this template can be distinguished from other similar templates, and so you can enable Drift for the overall template, and then click **Apply**.

Drift allows Cloud Manager to monitor the hard-coded values you entered for parameters when creating this template.

13. Click **Save Template**.

Result

The template is created and you are returned to the Template Dashboard where your new template appears.

See [what you should tell your users about templates](#).

Create a template for a Cloud Volumes ONTAP working environment

You can create a Cloud Volumes ONTAP working environment using templates.

-  • This support is provided only for AWS environments at this time, and only for single-node clusters.
• This template doesn't create the first volume in the working environment. You must add a "Create Volume in Cloud Volumes ONTAP" action in the template to create the volume.

See [how to launch a single-node Cloud Volumes ONTAP system in AWS](#) for the prerequisites that must be in place, and for details about all the parameters you need to complete in this template.

Steps

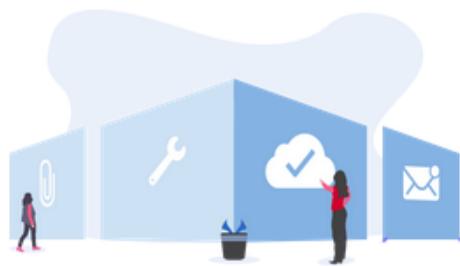
1. Select the **AppTemplate** service, click the **Templates** tab, and click **Add New Template**.

The *Select_a_Template* page is displayed.

Select a Template to Get Started

Choose a template for the type of resource you want to create and then customize the parameters for what's required in your environment.

For even greater flexibility, choose the "blank" template to mix and match configurations and to apply additional Cloud Manager services.



Start with a template of:

Blank template

Volume for Cloud Volumes
ONTAP + Backup

Volume for Cloud Volumes
ONTAP + Backup + Compliance

Volume for Azure NetApp files +
Compliance

Get Started Cancel

2. Select **Blank template** and click **Get Started**.

The *Add New Action* page is displayed.

Add New Action

🔍

ACTIONS - RESOURCES

Create Volume in Azure NetApp Files

Create Volume in Cloud Volumes ONTAP

Create Volume in On-Premises ONTAP

Create Working Environment in AWS (single node)

ACTIONS - SERVICES

Activate Cloud Data Sense on Volume

Apply Cancel

3. Select **Create Working Environment in AWS (single node)** as the type of resource you want to create, and click **Apply**.

The *Create Working Environment in AWS (single node)* page is displayed.

The screenshot shows the 'Create Working Environment in AWS (single node)' configuration page. It includes sections for Action Definition, Details and Credentials, and Tags. The Action Name is set to 'Create Working Environment in AWS (single node) (#a22)'. Under Details and Credentials, there is a dropdown menu for credentials. The Working Environment Name is set to 'Select...'. Under Tags, the 'Editable' checkbox is checked. A 'Add Tags' button is also present.

4. **Action Name:** Optionally, enter a customized action name instead of the default value.
5. **Details and Credentials:** Select the AWS credentials to use, enter a working environment name, and add tags, if needed.

Some of the fields in this page are self-explanatory. The following table describes fields for which you might need guidance:

Field	Description
Credentials	These are the credentials for the Cloud Volumes ONTAP cluster admin account. You can use these credentials to connect to Cloud Volumes ONTAP through ONTAP System Manager or its CLI.
Working Environment Name	Cloud Manager uses the working environment name to name both the Cloud Volumes ONTAP system and the Amazon EC2 instance. It also uses the name as the prefix for the predefined security group, if you select that option. Click in the field and select one of the 5 options. You can let the admin enter any name by selecting Free Text , or you can specify that the working environment name must have a certain prefix or suffix, that it <i>contains</i> certain characters, or that it follows rules from a regular expression (regex) you enter.

Field	Description
Tags	<p>AWS tags are metadata for your AWS resources. Cloud Manager adds the tags to the Cloud Volumes ONTAP instance and each AWS resource associated with the instance.</p> <p>For information about tags, refer to AWS Documentation: Tagging your Amazon EC2 Resources.</p>

6. **Location & Connectivity:** Enter the network information that you recorded in the [AWS worksheet](#). This includes the AWS Region, VPC, Subnet, and Security Group.

If you have an AWS Outpost, you can deploy a single node Cloud Volumes ONTAP system in that Outpost by selecting the Outpost VPC. The experience is the same as any other VPC that resides in AWS.

7. **Authentication Method:** Select the SSH authentication method you want to use; either a password or a key pair.
8. **Data Encryption:** Choose no data encryption or AWS-managed encryption.

For AWS-managed encryption, you can choose a different Customer Master Key (CMK) from your account or another AWS account.

[Learn how to set up the AWS KMS for Cloud Volumes ONTAP](#).

9. **Charging Method:** Specify which charging option would you like to use with this system.

[Learn about these charging methods](#).

10. **NSS Account:** Select a NetApp Support Site account.

11. **Preconfigured Packages:** Select one of the four preconfigured packages that will determine several factors for volumes created in the working environment.

12. **SMB Configuration:** If you plan to deploy volumes using SMB on this working environment, you can set up a CIFS server and related configuration elements.

13. Click **Apply** after you have defined the parameters needed for this action.

If the template values are correctly completed, a green checkmark is added to the "Create Working Environment in AWS (single node)" box.

14. You may want to add another action in this template to create a volume for this working environment. If so, click  and add that action. See how to [Create a template for a Cloud Volumes ONTAP volume](#) for details.

15. Enter the template name in the top left.

16. Click **Settings & Drift** to provide a more detailed description so that this template can be distinguished from other similar templates, and so you can enable Drift for the overall template, and then click **Apply**.

Drift allows Cloud Manager to monitor the hard-coded values you entered for parameters when creating this template.

17. Click **Save Template**.

Result

The template is created and you are returned to the Template Dashboard where your new template appears.

See [what you should tell your users about templates](#).

Examples of enabling services using templates

Service templates enable you to activate Cloud Backup, Cloud Data Sense, or Replication (SnapMirror) services on a newly created volume.

Add Backup functionality to a volume

When creating a volume template, you can add in the template that you want to create backups of the volume periodically using the [Cloud Backup](#) service. This action is not applicable for Azure NetApp Files volumes.

Enable Cloud Backup

Action Definition

Context

Working Environment Editable Drift
Select Working Environment

Storage VM Editable Drift
Select Storage VM

Policy - Retention & Schedule

Backup Every Editable Drift
Week

Number of backups to retain Editable Drift

Minimum Maximum

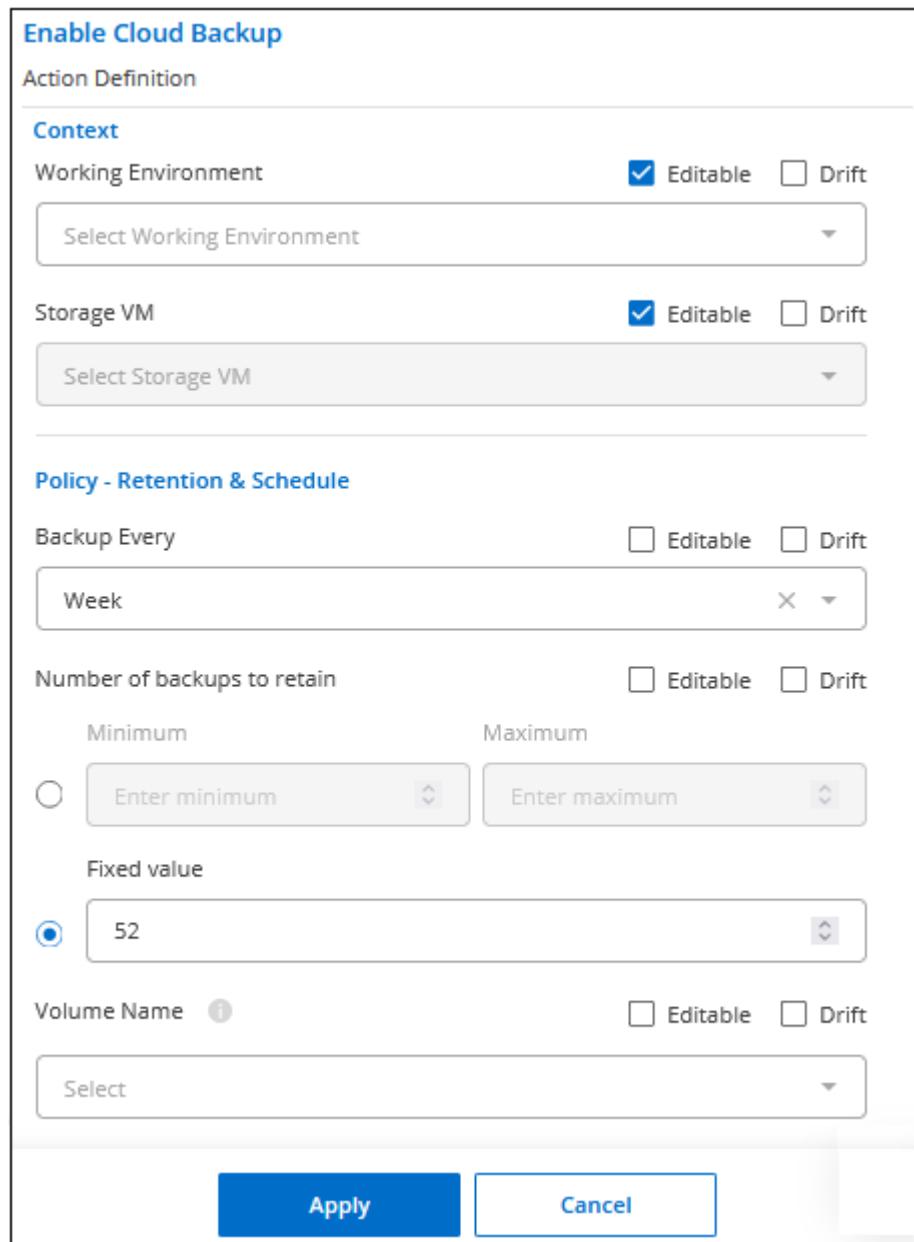
Enter minimum Enter maximum

Fixed value
52

Volume Name Editable Drift
Select

Buttons

Apply **Cancel**



- Context:** You can enter a working environment Name and storage VM name if you are using this action in a template without first creating a volume. Otherwise, leave these fields as "Editable."
- Policy:** Define the backup policy to create daily, weekly, or monthly backups with a specific number of backup copies to retain.

3. **Volume Name:** Typically the volume is the one created prior to the backup action in the same template. In this case, see how to [complete the fields](#) within the volume name to indicate that volume.
4. Click **Apply** to save your changes.

Add Data Sense functionality to a volume

When creating a volume template, you can add in the template that you want to scan the volume for compliance and classification using the [Cloud Data Sense](#) service.

Activate Cloud Data Sense on Volume

Action Definition

Context

Working Environment	<input checked="" type="checkbox"/> Editable <input type="checkbox"/> Drift
Select Working Environment	
Volume Name	<input type="checkbox"/> Editable <input type="checkbox"/> Drift
Select	
Volume UUID	<input type="checkbox"/> Editable <input type="checkbox"/> Drift
Select	
Volume Path	<input type="checkbox"/> Editable <input type="checkbox"/> Drift
Select	
Protocol	<input type="checkbox"/> Editable <input type="checkbox"/> Drift
Select	

Buttons

Apply **Cancel**

1. **Working Environment:** You can enter a working environment Name if you are using this action in a template without first creating a volume. Otherwise, leave this field as "Editable."
2. **Volume Name:** Typically the volume is the one created prior to the Data Sense action in the same template. In this case, see how to [complete the fields](#) within the volume name to indicate that volume.
3. **Volume UUID:** Data Sense needs the UUID of the volume before it can scan the volume. See how to [complete the three fields](#) below the volume name to indicate that volume.
4. **Volume Path:** Typically this is the mount point from the volume you are creating in the template. So you can get the "mountpoint" value from the output of that volume.
5. **Protocol:** Data Sense needs to know the protocol so it can scan the data. So you can get the "Protocol" value from the output of that volume.
6. Click **Apply** to save your changes.

Add Replication functionality to a volume

When creating a volume template, you can add in the template that you want to replicate the data in the volume to another volume using the [Replication](#) service. You can replicate data to a Cloud Volumes ONTAP cluster or to an on-prem ONTAP cluster.



This action is not applicable for Azure NetApp Files volumes.

Replication functionality consists of three parts: selecting the source volume, selecting the destination volume, and defining the replication settings. Each section is described below.

1. **Source Details:** Enter the details about the source volume you want to replicate:

The screenshot shows the 'Source Details' configuration pane. It includes fields for 'Source Working Environment' (with 'Editable' and 'Drift' checkboxes), 'Source Storage VM' (with 'Editable' and 'Drift' checkboxes), 'Source Volume Name' (with 'Editable' and 'Drift' checkboxes), and 'Source intercluster LIF IPs' (with 'Editable' checked and 'Drift' checkbox). A button at the bottom left says '+ Add Source intercluster LIF IPs'.

Field	Editable	Drift
Source Working Environment	<input type="checkbox"/>	<input type="checkbox"/>
Source Storage VM	<input type="checkbox"/>	<input type="checkbox"/>
Source Volume Name	<input type="checkbox"/>	<input type="checkbox"/>
Source intercluster LIF IPs	<input checked="" type="checkbox"/>	<input type="checkbox"/>

- Select the working environment where the volume resides.
- Select the storage VM on which the volume resides.
- Typically the volume is the one created prior to the replication action in the same template. In this case, see how to [complete the fields](#) within the Source Volume Name field to indicate that volume.
- Replication requires that the source and destination working environments are connected through their intercluster LIFs. Enter the intercluster LIF IP address for the source working environment.

To get this information: double-click the working environment, click the menu icon, and click Information.

2. **Destination Details:** Enter the details about the destination volume that will be created by the replication operation:

Destination Details

Destination Working Environment	<input type="checkbox"/> Editable <input type="checkbox"/> Drift
Select destination Working Environment	<input type="button" value="▼"/>
Destination Storage VM	<input type="checkbox"/> Editable <input type="checkbox"/> Drift
Select destination Storage VM	<input type="button" value="▼"/>
Destination Aggregate Name	<input type="checkbox"/> Editable <input type="checkbox"/> Drift
Select destination Aggregate	<input type="button" value="▼"/>
Destination Volume name	<input type="checkbox"/> Editable <input type="checkbox"/> Drift
Select...	<input type="button" value="▼"/>
Destination intercluster LIF IPs	<input checked="" type="checkbox"/> Editable <input type="checkbox"/> Drift
Intercluster LIF IP (1)	<input type="button" value="x"/>
+ Add Destination intercluster LIF IPs	
Destination Provider	<input type="checkbox"/> Editable <input type="checkbox"/> Drift
AWS	<input type="button" value="x"/> <input type="button" value="▼"/>
Destination Disk Type	<input type="checkbox"/> Editable <input type="checkbox"/> Drift
GP2 - General Purpose SSD	<input type="button" value="x"/> <input type="button" value="▼"/>

- a. Select the working environment where the volume will be created.
- b. Select the storage VM on which the volume will reside.
- c. Select the aggregate on which the volume will reside.
- d. For the destination volume, click in the field and select one of the 5 options. You can let the admin enter any name by selecting **Free Text**, or you can specify that the volume name must have a certain prefix or suffix, that it *contains* certain characters, or that it follows rules from a regular expression (regex) you enter.
- e. Replication requires that the source and destination working environments are connected through their intercluster LIFs. Enter the intercluster LIF IP address for the destination working environment.
- f. When replicating a volume to a Cloud Volumes ONTAP cluster (not to an on-prem ONTAP cluster), you need to specify the Destination Provider (AWS, Azure, or GCP) and the type of disk that will be used for the new volume.

3. **Replication Details:** Enter the details about the type and frequency of the replication operation:

Replication Details ⓘ

Replication Policy ⓘ	<input type="checkbox"/> Editable	<input type="checkbox"/> Drift
<input type="button" value="▼"/>		
Schedule	<input type="checkbox"/> Editable	<input type="checkbox"/> Drift
<input type="button" value="▼"/>		
<input type="checkbox"/> Enable replication health monitoring	<input type="checkbox"/> Editable	<input type="checkbox"/> Drift
Enable Transfer Rate Limit	<input type="checkbox"/> Editable	<input type="checkbox"/> Drift
<input checked="" type="radio"/> Limit transfer rate	<input type="radio"/> Unlimited (recommended for DR only machines)	
Transfer Rate Limit (KB/s) ⓘ	<input type="checkbox"/> Editable	<input type="checkbox"/> Drift
Minimum	Maximum	
<input type="radio"/> Enter minimum	<input type="radio"/> Enter maximum	
Fixed value		
<input type="radio"/> Enter a value for transfer rate limit		

- a. Select the [replication policy](#) that you want to use.
 - b. Choose a one-time copy or a recurring replication schedule.
 - c. Enable replication health monitoring if you want the drift report to include the replication health of the SnapMirror relationship along with the lag time, status, and last transfer time. [See what this looks like in the drift report.](#)
 - d. Select whether you want to set a transfer rate limit, and then enter the maximum rate (in kilobytes per second) at which data can be transferred. You can enter a fixed value, or you can provide a minimum and maximum and let the storage admin select a value in that range.
4. Click **Apply** to save your changes.

What to do after you have created the template

After you have created a template, you should inform your storage administrators to use the template when creating new working environments and volumes.

You can point them to [Creating resources using templates](#) for details.

Edit or delete a template

You can modify a template if you need to change any of the parameters. After you save your changes, all future resources created from the template will use the new parameter values.

You can also delete a template if you no longer need it. Deleting a template does not affect any of the resources that were created with the template. However, no Drift compliance checking can be done after the

template is deleted.

The screenshot shows a list of templates. The first template, 'CVO volume for databases', has a context menu open. The menu items are: Run Drift, Run Template, Edit Template, and Delete Template. The 'Edit Template' option is highlighted with a red box.

Check resources for template compliance

Cloud Manager can monitor the parameter values used when a resource was created with a template using the "drift" feature. Drift identifies resources that have been changed and that are no longer compliant with the template settings.

At this time, drift identifies the changed parameters in a resource — you must manually make adjustments to the resource to bring it back into compliance with the template.

How does drift work

Drift identifies non-compliant parameters like this:

1. When creating a template, you turn drift on for certain parameters that you do not want users to change. For example, you may require that Snapshot copies are created using the "Default" policy for all volumes created from a template.

The screenshot shows the 'Protection' section of a template configuration. It includes a 'Snapshot Policy' dropdown set to 'Default' and a 'Drift' checkbox which is checked and highlighted with a red box.

2. You turn drift on for the template, and then save the template.

Settings & Drift

Settings

Template Description

Volumes for application testing

Drift

Drift Off
We do not monitor any future changes to resources created with this template.

Drift On
We will monitor future changes to resources created with this template. You can run a Drift Report to see which fields configured with Drift are no longer compliant with the template settings.

3. Storage admins run the template to create volumes.
4. Later, a storage admin edits a volume and disables Snapshot copies.
5. You run drift checking on all templates, and the AppTemplates service compares the Snapshot copies template setting to the current setting in the volume. Any non-compliant values are flagged so you can fix the incorrect setting.

The Drift Dashboard

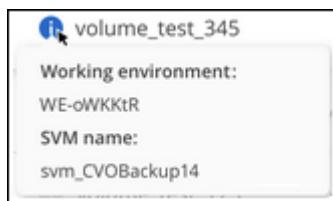
The Drift Dashboard shows the total number of resources (for example, volumes) that have been created using your templates, the number that are still compliant with the template, the number that are not compliant (drifted), and the number that were created with Drift disabled.

ID	Template Name	Action Type	Drift Policy & Status	Create Timestamp	Drift Time
volume_test_234	AWS CVO volume for staging	Create Volume in Cloud Volumes...	ON Drifted ↗	May 04 2021, 11:28:21 am	May 04 2021
volume_test_456	AWS CVO volume for staging	Create Volume in Cloud Volumes...	ON Compliant	May 04 2021, 11:28:21 am	May 04 2021
volume_test_123	AWS CVO volume for staging	Create Volume in Cloud Volumes...	ON Not Checked	May 04 2021, 11:28:21 am	--

- The controls at the top of each column allow you to sort the results in numerical or alphabetical order.
- The enables you to filter the results by Template Name, Drift Policy & Status, and Action Type. For example:

ID	Status	Description
volume_test_345	ON Drifted	S CVO volume for staging
	ON Compliant	
	OFF Not Checked	

- The search bar enables you to search for a specific volume name or template name.
- To find out more details about the actual resource (or volume), such as the working environment and storage VM, you can click the



Populate the Drift Dashboard

You must run drift checking on a template before it populates the values in the Drift Dashboard.

You can run drift checking for all templates from the Templates Dashboard:

You can run drift checking on a single template from the Templates Dashboard:

Template name	Resource type	Created by	Last Modified	Parameters	Resources created	Actions
Cloud Volume ONTAP	Cloud Volume ONTAP Volume	Jhon Doe	Dec 16, 2020 02:13:59 pm UTC	10	View	
Cloud Volume ONTAP	Cloud Volume ONTAP Volume	Jhon Doe	Dec 16, 2020 02:13:59 pm UTC	10	View	
Cloud Volume ONTAP	Cloud Volume ONTAP Volume	Jhon Doe	Dec 16, 2020 02:13:59 pm UTC	10	View	
Cloud Volume ONTAP	Cloud Volume ONTAP Volume	Jhon Doe	Dec 16, 2020 02:13:59 pm UTC	10	View	
Cloud Volume ONTAP	Cloud Volume ONTAP Volume	Jhon Doe	Dec 16, 2020 02:13:59 pm UTC	10	View	

Create a drift report for non-compliant resources

You can view a drift report for a single resource, or run a report to download a report for all resources. Using this report you can assign actions to system admins to make changes that bring the resources back into compliance with the template.

You can click the Drift icon for a resource in the Drift Dashboard to view a list of the parameters in each resource that is non-compliant.

The screenshot shows a table titled "7/150 Resources". The columns are: ID, Template Name, Action Type, Drift Policy & Status, Create Timestamp, and Drift Time. One row is highlighted, showing "volume_test_234" with "Drifted" status. A red box highlights the "Drifted" status. An arrow points from this row to a modal window titled "Resource Differences" for "volume_test_345". The modal lists four properties: Size (Not Equal), Storage efficiency (Not Equal), Tiering policy (Removed), and Provider volume type (Not Equal). The "Current Value" column shows values 100, False, None, and GP3 respectively. A "Close" button is at the bottom right of the modal.

ID	Template Name	Action Type	Drift Policy & Status	Create Timestamp	Drift Time
volume_test_234	AWS CVO volume for staging	Create Volume in Cloud Volumes...	ON Drifted ↗	May 04 2021, 11:28:21 am	May 04 2021
volume_test_234	AWS CVO volume for staging	Create Volume in Cloud Volumes...	ON Drifted ↗	May 04 2021, 11:28:21 am	May 04 2021
volume_test_345	AWS CVO volum				

Resource Differences

volume_test_345

Property	Change	Expected Value	Current Value
Size	Not Equal	50	100
Storage efficiency	Not Equal	True	False
Tiering policy	Removed	None	--
Provider volume type	Not Equal	GP2	GP3

Close

To view a drift report for resources that have been created from your templates, click to download a .CSV file. The drift report reflects what is currently filtered on the page - it does not show all resources unless you haven't applied any filters on the page.

Replication health details in the drift report

When [enabling Replication on a volume using templates](#), you can choose to show more detailed replication information in the drift report by enabling drift in the "Enable replication health monitoring" field. When enabled, the drift report shows whether the health of the Replication relationship is healthy or unhealthy (drifted), along with the SnapMirror lag time, status, and last transfer time.

This screenshot shows the replication details for a SnapMirror relationship that is unhealthy in the drift report.

Throughput	ADD	--	575								
Health	Not Equal	True	False								
Monitor <table border="1"> <thead> <tr> <th>Property</th> <th>Value</th> </tr> </thead> <tbody> <tr> <td>Lagtime</td> <td>5 Days</td> </tr> <tr> <td>Status</td> <td>Idle</td> </tr> <tr> <td>Last transfer end time</td> <td>May 04 2021, 11:28:21 am</td> </tr> </tbody> </table>				Property	Value	Lagtime	5 Days	Status	Idle	Last transfer end time	May 04 2021, 11:28:21 am
Property	Value										
Lagtime	5 Days										
Status	Idle										
Last transfer end time	May 04 2021, 11:28:21 am										

Note: As replication is initially being applied to the volume the health will be returned as "False", meaning it is unhealthy. After a few minutes the real replication status will be displayed.

Create resources using templates

Select one of the application templates that your organization has built to create working environments or volumes that are optimized for specific workloads and applications.

Templates also enable you to activate [Cloud Backup](#), [Cloud Data Sense](#), and [Replication \(SnapMirror\)](#) on the created volumes.

Templates enable you to create volumes for Cloud Volumes ONTAP, Azure NetApp Files, and on-premises ONTAP systems.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.

1

Verify required prerequisites

- Before you can create a volume for a Cloud Volumes ONTAP, on-premises ONTAP, or Azure NetApp Files system using a template, make sure you have access to an appropriate working environment where the volume will be deployed.
- If the template activates a cloud service on the volume, such as [Cloud Backup](#) or [Cloud Data Sense](#), ensure that the service is active and licensed in your environment.

2

Launch the Application Templates service

Select the **AppTemplate** service and click the **Templates** tab.



Build the resource by running the template and defining parameters

Select the template, click **Run Template**, and enter values in the editable fields to create the resource.

Requirements

Read the following requirements to make sure that you have a supported configuration.

- If you don't already have a Connector, [see how to create Connectors](#) for AWS, Azure, and GCP.
- When creating a Cloud Volumes ONTAP volume, make sure you have a Cloud Volumes ONTAP working environment available.
- When creating an on-premises ONTAP volume, make sure you have an on-premises ONTAP working environment available.
- When creating an Azure NetApp Files volume, make sure you have an Azure NetApp Files working environment available.
- If the template activates a cloud service on the volume, such as [Cloud Backup](#), [Cloud Data Sense](#), or [Replication \(SnapMirror\)](#), ensure that the service is active and licensed in your environment.

Select and run the volume template

There are multiple ways to select and run a template to create new volumes:

- Run the volume template from the working environment
- Run the volume template from the Templates dashboard

Regardless of the method you choose, the details about the required volume parameters that you must define are available in these sections:

- [How to provision Cloud Volumes ONTAP volumes](#)
- [How to provision Azure NetApp Files volumes](#)
- [How to provision on-premises ONTAP volumes](#)

Run a volume template from the working environment

You can add a volume to an existing working environment from the *Working Environment* page and from the *Volume Details* page.

Steps

1. From the *Working Environment* page or from the *Volume Details* page, click **Add Volume From Template**.

The *Templates Dashboard* is displayed and it lists only those templates that are applicable to the selected working environment—for example, it only shows Cloud Volumes ONTAP templates.

2. Click **...** and Run Template.

The *Add Volume from Template* page appears.

3. Enter values in the editable fields to create the volume and click **Run Template**.

Run a volume template from the Templates Dashboard

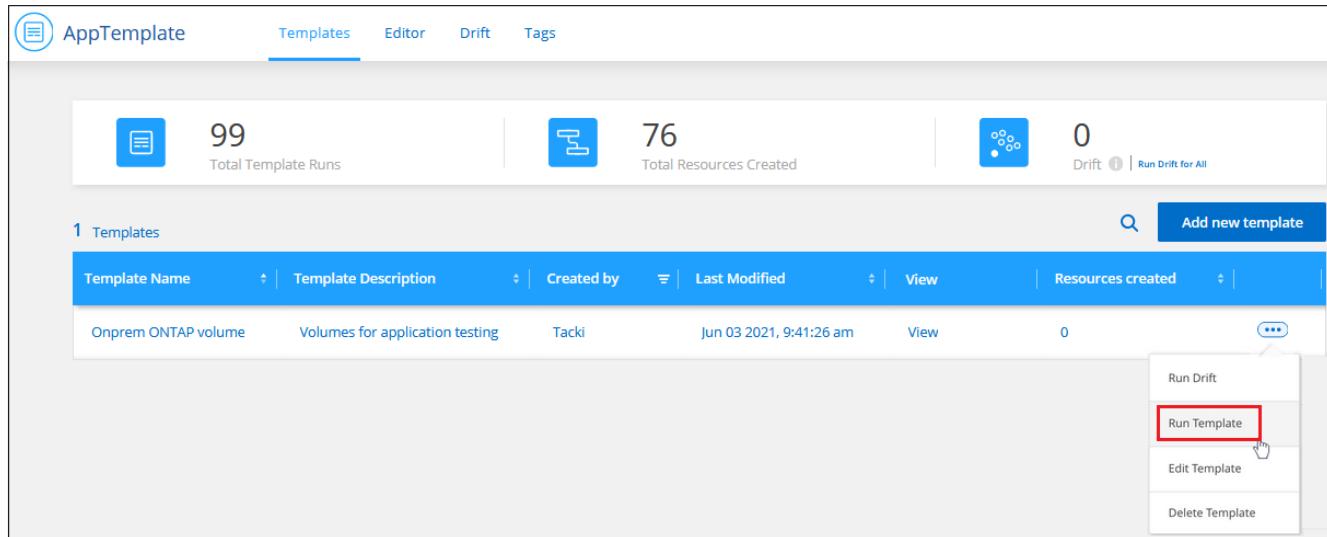
You can add a volume to an existing working environment from the *Templates Dashboard*.

Steps

1. Select the **AppTemplate** service and click the **Templates** tab.

The *Templates Dashboard* is displayed.

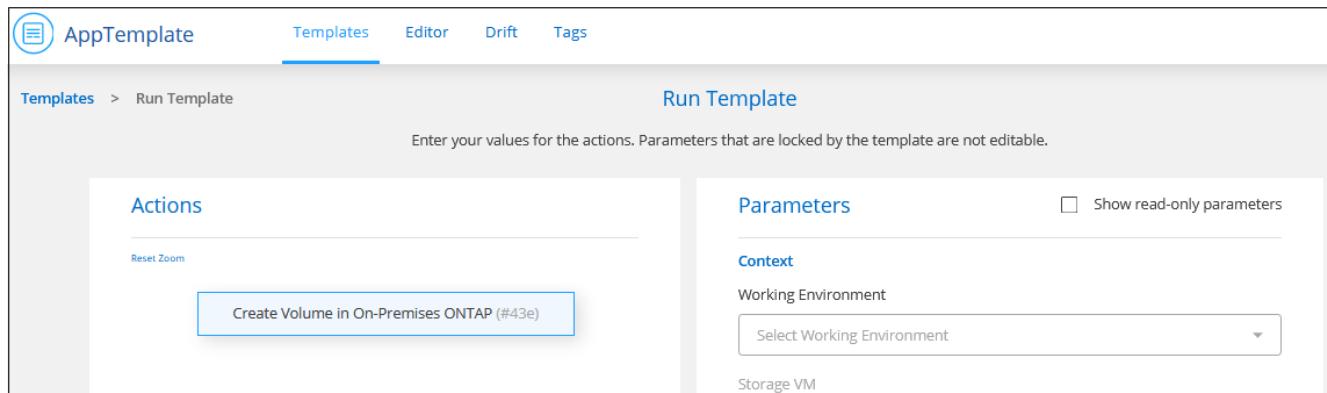
2. For the template that you want to use, click **...** and **Run Template**.



The screenshot shows the AppTemplate service's Templates dashboard. At the top, there are summary statistics: 99 Total Template Runs, 76 Total Resources Created, and 0 Drift. Below this is a table header for 'Templates' with columns: Template Name, Template Description, Created by, Last Modified, View, and Resources created. A single row is visible: 'Onprem ONTAP volume' with 'Volumes for application testing' as the description, 'Tacki' as the creator, 'Jun 03 2021, 9:41:26 am' as the last modified date, a 'View' link, and 0 resources created. To the right of this row is a context menu with options: Run Drift, Run Template (which is highlighted with a red box), Edit Template, and Delete Template. A cursor arrow points to the 'Run Template' option.

The *Run Template* page appears.

3. Enter values in editable fields to create the volume and click **Run Template**.



The screenshot shows the 'Run Template' page for the 'Onprem ONTAP volume' template. The top navigation bar shows 'Templates > Run Template'. The main area has two sections: 'Actions' and 'Parameters'. The 'Actions' section contains a button labeled 'Create Volume in On-Premises ONTAP (#43e)'. The 'Parameters' section has a checkbox 'Show read-only parameters' followed by a 'Context' section which is currently set to 'Working Environment' and has a dropdown button labeled 'Select Working Environment'.

Note that when you run the template from the dashboard that you need to select the working environment and other variables (for example, the storage VM and/or aggregate). When you run the template from the working environment, the working environment gets filled in automatically.

Select and run a working environment template

You can create a new working environment from the *Templates Dashboard* if your company has created a template for this functionality.

If you have any questions about the details required to create the working environment, see [Launching Cloud](#)

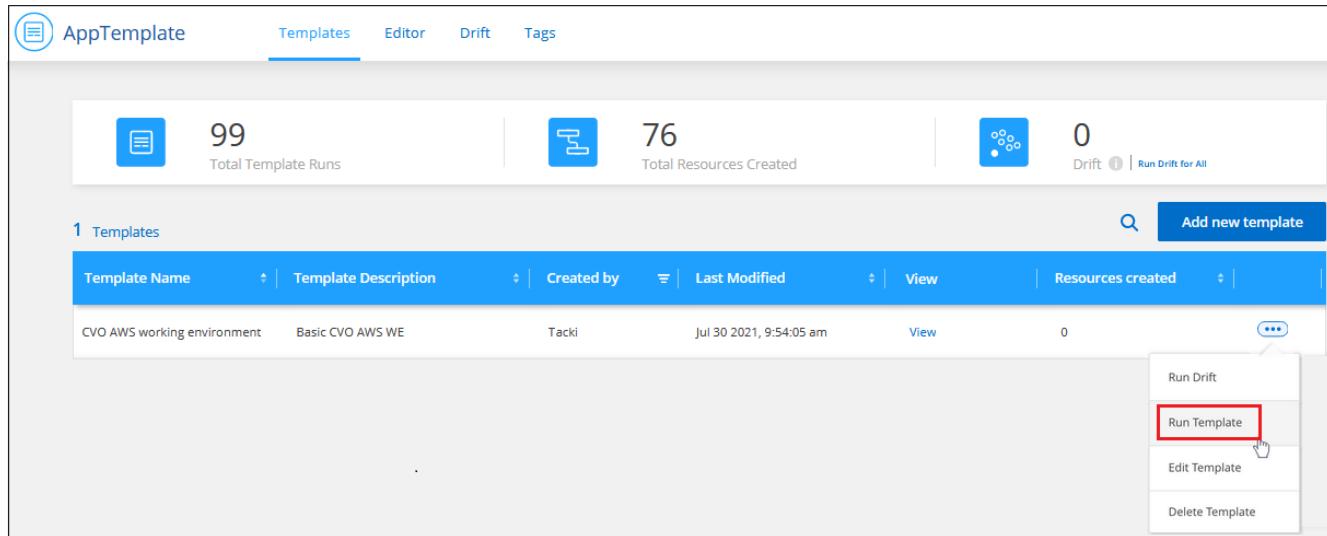
Volumes ONTAP in AWS.

Steps

1. Select the **AppTemplate** service and click the **Templates** tab.

The *Templates Dashboard* is displayed.

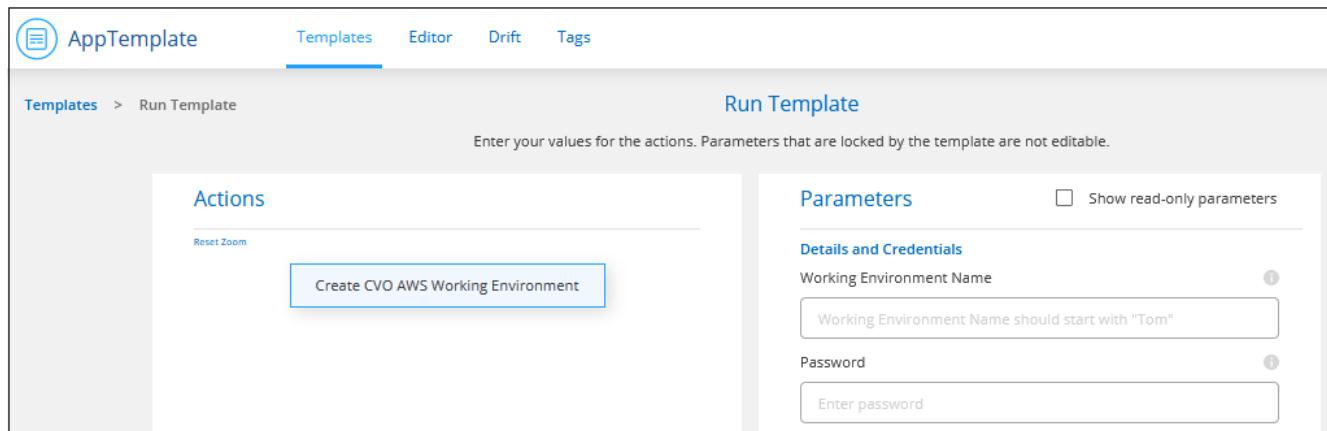
2. For the template that you want to use, click **...** and **Run Template**.



The screenshot shows the AppTemplate service's Templates dashboard. At the top, there are summary statistics: 99 Total Template Runs, 76 Total Resources Created, and 0 Drift items. Below this is a table listing one template: "CVO AWS working environment". The table columns include Template Name, Template Description, Created by, Last Modified, View, and Resources created. To the right of the table is a context menu with options: Run Drift, Run Template (which is highlighted with a red box), Edit Template, and Delete Template. A mouse cursor is hovering over the "Run Template" option.

The *Run Template* page appears.

3. Enter values in editable fields to create the volume and click **Run Template**.



The screenshot shows the "Run Template" page for the "CVO AWS working environment" template. The top navigation bar includes "Templates" (which is selected), "Editor", "Drift", and "Tags". The main area has a breadcrumb "Templates > Run Template" and a title "Run Template" with the sub-instruction "Enter your values for the actions. Parameters that are locked by the template are not editable." On the left, under "Actions", there is a button labeled "Create CVO AWS Working Environment". On the right, under "Parameters", there are two fields: "Working Environment Name" (with a note "Working Environment Name should start with 'Tom'") and "Password" (with a note "Enter password"). There is also a checkbox "Show read-only parameters" which is unchecked.

Back up data to the cloud

Learn about Cloud Backup

Cloud Backup is a service for Cloud Volumes ONTAP and on-premises ONTAP clusters that delivers backup and restore capabilities for protection, and long-term archive of your data. Backups are automatically generated and stored in an object store in your public or private cloud account, independent of volume Snapshot copies used for near-term recovery or cloning.

When necessary, you can restore an entire *volume*, or one or more *files*, from a backup to the same or different Cloud Volumes ONTAP or on-premises ONTAP cluster.

[Learn more about Cloud Backup.](#)

Features

- Back up independent copies of your data volumes to low-cost object storage.
- Back up from cloud to cloud, and from on-premises ONTAP systems to public or private cloud.
- Backups can reside on a different subscription/account or different region than your Cloud Volumes ONTAP system.
- Backup data is secured with AES-256 bit encryption at-rest and TLS 1.2 HTTPS connections in-flight.
- Support for up to 4,000 backups of a single volume.
- Restore data from a specific point in time.
- Restore a volume, or individual files, to the source system or to a different system.
- Restore data to a working environment using a different subscription/account or that is in a different region.
- Restores data on a block level, placing the data directly in the location you specify, all while preserving the original ACLs.
- Browsable file catalog for selecting individual files for single file restore.

Supported working environments and object storage providers

Cloud Backup enables you to back up volumes from the following working environments to object storage in the following cloud providers:

Source Working Environment	Backup File Destination
Cloud Volumes ONTAP in AWS	Amazon S3
Cloud Volumes ONTAP in Azure	Azure Blob
Cloud Volumes ONTAP in Google	Google Cloud Storage
On-premises ONTAP system	Amazon S3 Azure Blob Google Cloud Storage NetApp StorageGRID

You can restore a volume, or individual files, from a backup file to the following working environments:

Backup File Location	Destination Working Environment	
	Volume Restore	File Restore
Amazon S3	Cloud Volumes ONTAP in AWS On-premises ONTAP system	Cloud Volumes ONTAP in AWS On-premises ONTAP system
Azure Blob	Cloud Volumes ONTAP in Azure On-premises ONTAP system	Cloud Volumes ONTAP in Azure On-premises ONTAP system
Google Cloud Storage	Cloud Volumes ONTAP in Google On-premises ONTAP system	
NetApp StorageGRID	On-premises ONTAP system	

Cost

There are two types of costs associated with using Cloud Backup: resource charges and service charges.

Resource charges

Resource charges are paid to the cloud provider for storage and for running a virtual machine-instance in the cloud.

- For Backup, you pay your cloud provider for object storage costs. (There are no storage costs when creating backups on your StorageGRID systems.)

Since Cloud Backup preserves the storage efficiencies of the source volume, you pay the cloud provider object storage costs for the data *after* ONTAP efficiencies (for the smaller amount of data after deduplication and compression have been applied).

- For File Restore, you pay your cloud provider for compute costs only when the Restore instance is running.

The instance resides in the same subnet as the Connector, and it runs only when browsing a backup file to locate the individual files you want to restore. The instance is turned off when not in use to save costs.

- In AWS, the Restore instance runs on an [m5n.xlarge instance](#) with 4 CPUs, 16 GiB memory, and EBS Only instance storage. The operating system image is Amazon Linux 2.

In regions where m5n.xlarge instance isn't available, Restore runs on an m5.xlarge instance instead.

- In Azure, the Restore virtual machine runs on a [Standard_D4s_v3 VM](#) with 4 CPUs, 16 GiB memory, and a 32 GB disk. The operating system image is CentOS 7.5).

The instance is named *Cloud-Restore-Instance* with your Account ID concatenated to it. For example: *Cloud-Restore-Instance-MyAccount*.

- For Volume Restore there is no cost because no separate instance or virtual machine is required.

Service charges

Backup service charges are paid to NetApp and cover both the cost to *create* backups and to *restore* volumes, or files, from those backups. You pay only for the data that you protect, calculated by the target backup capacity *before* ONTAP efficiencies.

There are three ways to pay for the Backup service. The first option is to subscribe from your cloud provider, which enables you to pay per month based on the amount of backed up data. The second option is to get an annual contract - this is only available through AWS. The third option is to purchase licenses directly from NetApp. Read the [Licensing](#) section for details.

Licensing

Cloud Backup is available in three licensing options: Pay As You Go (PAYGO), an annual contract from the AWS Marketplace, and Bring Your Own License (BYOL). A 30-day free trial is available if you don't have a license.

Free trial

When using the 30-day free trial, you are notified about the number of free trial days that remain. At the end of your free trial, backups stop being created. You must subscribe to the service or purchase a license to continue using the service.

Backup are not deleted when the service is disabled. You'll continue to be charged by your cloud provider for object storage costs for the capacity that your backups use unless you delete the backups.

Pay-as-you-go subscription

Cloud Backup offers consumption-based licensing in a pay-as-you-go model. The licensing costs are based on target backup capacity (before ONTAP storage efficiencies). After subscribing through your cloud provider's marketplace, you pay per GB for data that's backed up—there's no up-front payment. You are billed by your cloud provider through your monthly bill.

You should subscribe even if you have a free trial or if you bring your own license (BYOL):

- Subscribing ensures that there's no disruption of service after your free trial ends.

When the trial ends, you'll be charged hourly according to the amount of data that you back up.

- If you back up more data than allowed by your BYOL license, then data backup continues through your pay-as-you-go subscription.

For example, if you have a 10 TB BYOL license, all capacity beyond the 10 TB is charged through the pay-as-you-go subscription.

You won't be charged from your pay-as-you-go subscription during your free trial or if you haven't exceeded your BYOL license.

[Learn how to set up a pay-as-you-go subscription.](#)

Annual contract (AWS only)

Two annual contracts are available from the AWS Marketplace:

- An annual contract that enables you to back up Cloud Volumes ONTAP data and on-premises ONTAP data.

You'll also need to pay for Cloud Volumes ONTAP using this annual contract subscription, since you can assign only one active subscription to your AWS credentials in Cloud Manager.

- A Professional Package that enables you to bundle Cloud Volumes ONTAP and Cloud Backup Service by using an annual contract for 12 months, 24 months, or 36 months. This option doesn't enable you to back up on-prem data.

You can set up the annual contract when you create a Cloud Volumes ONTAP working environment and Cloud Manager will prompt you to subscribe to the AWS Marketplace.

[Learn how to set up yearly AWS contracts.](#)

Bring your own license

BYOL is term-based (1YR/2YR/3YR) and capacity-based in 1 TB increments, based on the logical (before ONTAP storage efficiencies) backed up capacity. You pay NetApp to use the service for a period of time, say 1 year, and for a maximum amount backup capacity, say 10 TB, and you'll need to pay your cloud provider for object storage costs (as described earlier).



Backup to StorageGRID does require a BYOL license, but there is no cost for cloud provider storage space in this case.

You'll receive a serial number that you enter in the Cloud Manager Digital Wallet page to enable the service. When either limit is reached you'll need to renew the license. The Backup BYOL license applies to all Cloud Volumes ONTAP and on-premises ONTAP systems associated with your [Cloud Central account](#).

[Learn how to manage your BYOL licenses.](#)

BYOL license considerations

When using a Cloud Backup BYOL license, Cloud Manager displays a warning in the user interface when backups are nearing the capacity limit or nearing the license expiration date. You receive these warnings:

- When backups have reached 80% of licensed capacity, and again when you have reached the limit
- 30 days before a license is due to expire, and again when the license expires

Use the chat icon in the lower right of the Cloud Manager interface to renew your license when you see these warnings.

Two things can happen when your license expires:

- If the account you are using for your ONTAP systems has a marketplace account, the backup service continues to run, but you are shifted over to a PAYGO licensing model. You are charged by your cloud provider for object storage costs, and by NetApp for backup licensing costs, for the capacity that your backups are using.
- If the account you are using for your ONTAP systems does not have a marketplace account, the backup service continues to run, but you will continue to see the warnings.

Once you renew your BYOL subscription, Cloud Manager automatically obtains the new license from NetApp and installs it. If Cloud Manager can't access the license file over the secure internet connection, you can obtain the file yourself and manually upload it to Cloud Manager. For instructions, see [how to update a Cloud Backup license](#).

Systems that were shifted over to a PAYGO license are returned to the BYOL license automatically. And systems that were running without a license will stop seeing the warnings and will be charged for backup activity that occurred while the license was expired.

How Cloud Backup works

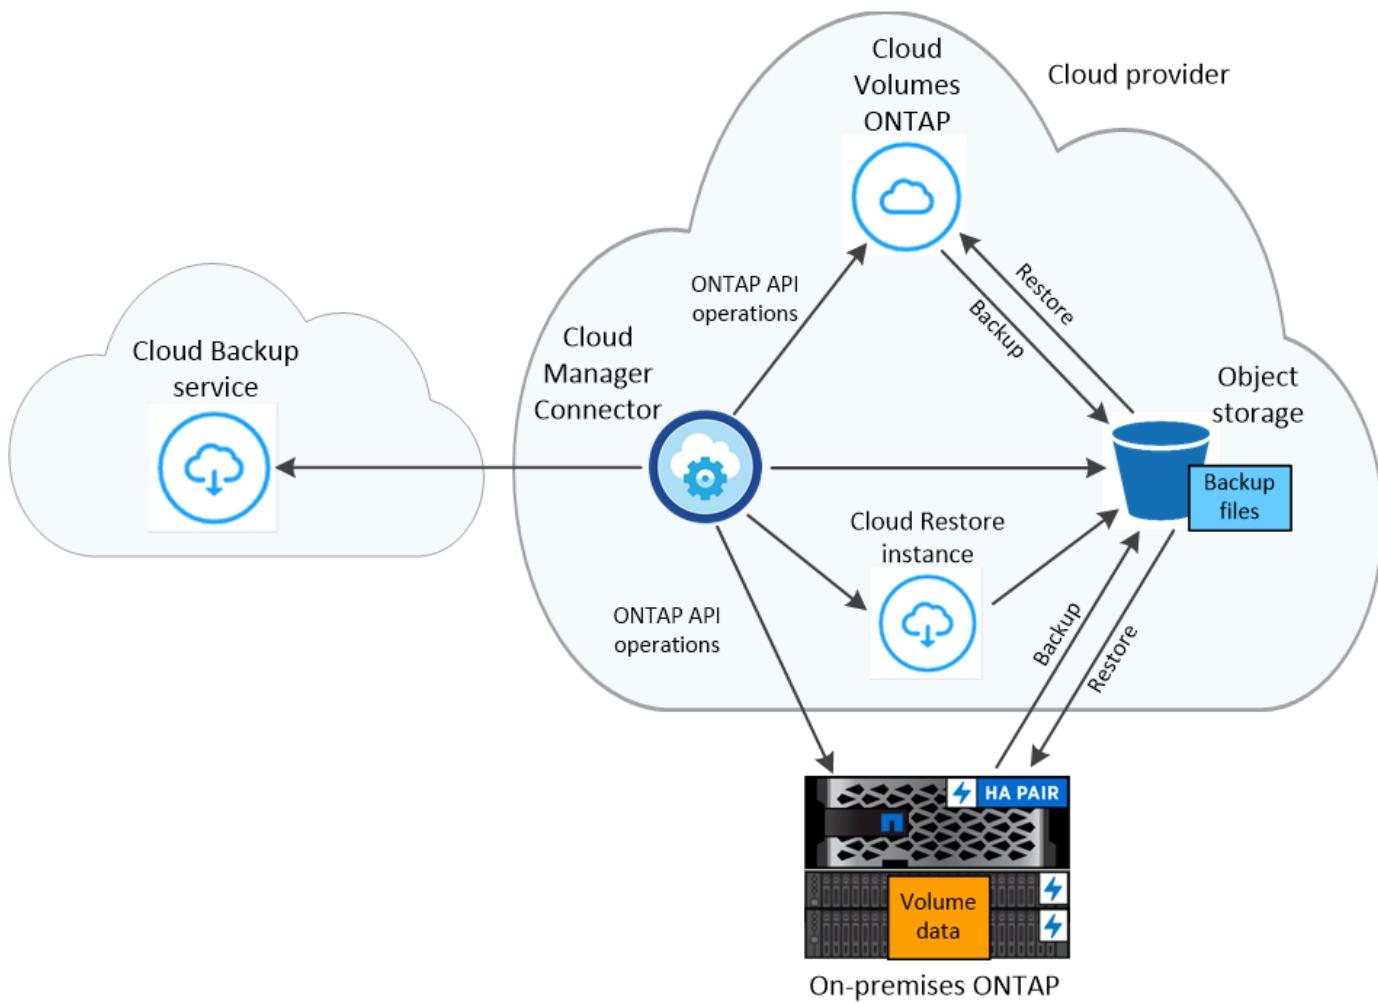
When you enable Cloud Backup on a Cloud Volumes ONTAP or on-premises ONTAP system, the service performs a full backup of your data. Volume snapshots are not included in the backup image. After the initial backup, all additional backups are incremental, which means that only changed blocks and new blocks are backed up. This keeps network traffic to a minimum.

In most cases you will use Cloud Manager for all backup and restore operations. However, starting with ONTAP 9.9.1 you can initiate volume backup operations of your on-premises ONTAP clusters using ONTAP System Manager. [See how to use System Manager to back up your volumes to the cloud using Cloud Backup.](#)



Any actions taken directly from your cloud provider environment to manage or change backup files may corrupt the files and will result in an unsupported configuration.

The following image shows the relationship between each component:



Where backups reside

Backup copies are stored in an object store that Cloud Manager creates in your cloud account. You identify the region when you enable the service.

There's one object store per Cloud Volumes ONTAP or on-premises ONTAP system. Cloud Manager names the object store as follows: "netapp-backup-clusteruuid". Be sure not to delete this object store.

- In AWS, Cloud Manager enables the [Amazon S3 Block Public Access feature](#) on the S3 bucket.
- In Azure, Cloud Manager uses a new or existing resource group with a storage account for the Blob container. Cloud Manager [blocks public access to your blob data](#) by default.
- In GCP, Cloud Manager uses a new or existing project with a storage account for the Google Cloud Storage bucket.
- In StorageGRID, Cloud Manager uses an existing storage account for the object store bucket.

Supported storage classes or access tiers

- In AWS, backups start in the *Standard* storage class and transition to the *Standard-Infrequent Access* storage class after 30 days.
- In Azure, backups are associated with the *Cool* access tier.
- In GCP, backups are associated with the *Standard* storage class by default.

You can also use the lower cost *Nearline* storage class, or the *Coldline* or *Archive* storage classes. See the Google topic [Storage classes](#) for information about changing the storage class.

- In StorageGRID, backups are associated with the *Standard* storage class.

Backup settings are system wide

When you enable Cloud Backup, all the volumes you identify on the system are backed up to the cloud.

The schedule and number of backups to retain are defined at the system level. The backup settings affect all volumes on the system.

The schedule is hourly, daily, weekly, monthly, or a combination

You can choose a combination of hourly, daily, weekly, and monthly backups of all volumes. You can also select one of the system-defined policies that provide backups and retention for 3 months, 1 year, and 7 years. These policies are:

Backup Policy Name	Backups per interval...			Max. Backups
	Daily	Weekly	Monthly	
Netapp3MonthsRetention	30	13	3	46
Netapp1YearRetention	30	13	12	55
Netapp7YearsRetention	30	53	84	167

Backup protection policies that you have created on the system using ONTAP System Manager or the ONTAP CLI are also available as selections.

You can also [create an on-demand backup of a volume](#) from the Backup Dashboard in addition to those backup files created from the scheduled backups.

Once you have reached the maximum number of backups for a category, or interval, older backups are removed so you always have the most current backups.

Note that the retention period for backups of data protection volumes is the same as defined in the source SnapMirror relationship. You can change this if you want by using the API.

Backups are taken at midnight

- Hourly backups start 5 minutes past the hour, every hour.
- Daily backups start just after midnight each day.
- Weekly backups start just after midnight on Sunday mornings.
- Monthly backups start just after midnight on the first day of each month.

The start time is based on the time zone set on each source ONTAP system. You can't schedule backup operations at a user-specified time from the UI. For more information, contact your System Engineer.

Backup copies are associated with your Cloud Central account

Backup copies are associated with the [Cloud Central account](#) in which Cloud Manager resides.

If you have multiple Cloud Manager systems in the same Cloud Central account, each Cloud Manager system will display the same list of backups. That includes the backups associated with Cloud Volumes ONTAP and on-premises ONTAP instances from other Cloud Manager systems.

Supported volumes

Cloud Backup supports FlexVol read-write volumes and data protection (DP) volumes.

FlexGroup volumes and SnapLock volumes aren't currently supported.

FabricPool tiering policy considerations

There are certain things you need to be aware of when the volume you are backing up resides on a FabricPool aggregate and it has an assigned policy other than `none`:

- The first backup of a FabricPool-tiered volume requires retrieval of all local and all tiered data (from the object store). This operation could cause a one-time increase in cost to read the data from your cloud provider.
 - Subsequent backups are incremental and do not have this effect.
 - If the tiering policy is assigned to the volume when it is initially created you will not see this issue.
- Consider the impact of backups before assigning the `all` tiering policy to volumes. Because data is tiered immediately, Cloud Backup will read data from the cloud tier rather than from the local tier. Because concurrent backup operations share the network link to the cloud object store, performance degradation might occur if network resources become saturated. In this case, you may want to proactively configure multiple network interfaces (LIFs) to decrease this type of network saturation.
- A backup operation does not "reheat" the cold data tiered in object storage.

Limitations

- When making backups from on-premises ONTAP systems to public cloud storage, the Connector must be deployed in the cloud.
- When making backups from on-premises ONTAP systems to StorageGRID (private cloud), the Connector must be deployed on premises.

- When backing up data protection (DP) volumes, relationships with the following SnapMirror labels will not be backed up to cloud:
 - app_consistent
 - all_source_snapshot
- In Azure, if you enable Cloud Backup when Cloud Volumes ONTAP is deployed, Cloud Manager creates the resource group for you and you cannot change it. If you want to pick your own resource group when enabling Cloud Backup, **disable** Cloud Backup when deploying Cloud Volumes ONTAP and then enable Cloud Backup and choose the resource group from the Cloud Backup Settings page.
- When backing up volumes from Cloud Volumes ONTAP systems, volumes that you create outside of Cloud Manager aren't automatically backed up. For example, if you create a volume from the ONTAP CLI, ONTAP API, or System Manager, then the volume won't be automatically backed up. If you want to back up these volumes, you would need to disable Cloud Backup and then enable it again.
- ILM (tiering) from the object storage, or direct write to AWS Glacier or similar lower tier object storage, is not supported.
- SVM-DR and SM-BC configurations are not supported.
- MetroCluster (MCC) backup is supported from ONTAP secondary only: MCC > SnapMirror > ONTAP > Cloud Backup Service > object storage.
- WORM/Compliance mode on an object store is not supported.

Single File Restore limitations

- Single file restore can restore up to 100 individual files at a time. There is currently no support for restoring folders/directories.
- The file being restored must be using the same language as the language on the destination volume. You will receive an error message if the languages are not the same.
- Single file restore is not supported when using the same account with different Cloud Managers in different subnets.

Get started

Backing up Cloud Volumes ONTAP data to Amazon S3

Complete a few steps to get started backing up data from Cloud Volumes ONTAP to Amazon S3.

Quick start

Get started quickly by following these steps or scroll down to the remaining sections for full details.



Verify support for your configuration

- You're running Cloud Volumes ONTAP 9.6 or later in AWS.
- You have a valid cloud provider subscription for the storage space where your backups will be located.
- You have subscribed to the [Cloud Manager Marketplace Backup offering](#), an [AWS annual contract](#), or you have purchased [and activated](#) a Cloud Backup BYOL license from NetApp.

- The IAM role that provides the Cloud Manager Connector with permissions includes S3 permissions from the latest [Cloud Manager policy](#).

2

Enable Cloud Backup on your new or existing system

- New systems: Cloud Backup is enabled by default in the working environment wizard. Be sure to keep the option enabled.
- Existing systems: Select the working environment and click **Enable** next to the Backup & Restore service in the right-panel, and then follow the setup wizard.



3

Enter the provider details

Select the AWS Account and the region where you want to create the backups. You can also choose your own customer-managed key for data encryption instead of using the default Amazon S3 encryption key.

Provider Settings

Provider Information <div style="margin-top: 10px;"> AWS Account <div style="border: 1px solid #ccc; padding: 2px; width: 100%;">AWS_Account_1</div> </div> <div style="margin-top: 10px;"> AWS Access Key <div style="border: 1px solid #ccc; padding: 2px; width: 100%;">Enter AWS Access Key</div> </div> <div style="margin-top: 10px;"> AWS Secret Key <div style="border: 1px solid #ccc; padding: 2px; width: 100%;">Enter AWS Secret Key</div> </div>	Location & Connectivity <div style="margin-top: 10px;"> Region <div style="border: 1px solid #ccc; padding: 2px; width: 100%;">us-east-2</div> </div> <div style="margin-top: 10px;"> Encryption <div style="display: flex; justify-content: space-between;"> Encryption Key Type: AWS SSE-S3 Change Key </div> </div>
--	--

4

Define the backup policy

The default policy backs up volumes every day and retains the most recent 30 backup copies of each volume. Change to hourly, daily, weekly, or monthly backups, or select one of the system-defined policies that provide more options. You can also change the number of backup copies to retain.

Define Policy

Policy - Retention & Schedule	<input checked="" type="radio"/> Create a New Policy <input type="radio"/> Select an Existing Policy <table style="width: 100%; border-collapse: collapse;"> <tr> <td style="width: 15%;"><input type="checkbox"/> Hourly</td> <td style="width: 15%;">Number of backups to retain</td> <td style="width: 15%; text-align: right;">24</td> <td style="width: 15%; text-align: right;"><input type="button" value="▼"/></td> <td style="width: 15%; text-align: right;"><input type="button" value="▲"/></td> </tr> <tr> <td><input checked="" type="checkbox"/> Daily</td> <td>Number of backups to retain</td> <td style="text-align: right;">30</td> <td style="text-align: right;"><input type="button" value="▼"/></td> <td style="text-align: right;"><input type="button" value="▲"/></td> </tr> <tr> <td><input type="checkbox"/> Weekly</td> <td>Number of backups to retain</td> <td style="text-align: right;">52</td> <td style="text-align: right;"><input type="button" value="▼"/></td> <td style="text-align: right;"><input type="button" value="▲"/></td> </tr> <tr> <td><input type="checkbox"/> Monthly</td> <td>Number of backups to retain</td> <td style="text-align: right;">12</td> <td style="text-align: right;"><input type="button" value="▼"/></td> <td style="text-align: right;"><input type="button" value="▲"/></td> </tr> </table>	<input type="checkbox"/> Hourly	Number of backups to retain	24	<input type="button" value="▼"/>	<input type="button" value="▲"/>	<input checked="" type="checkbox"/> Daily	Number of backups to retain	30	<input type="button" value="▼"/>	<input type="button" value="▲"/>	<input type="checkbox"/> Weekly	Number of backups to retain	52	<input type="button" value="▼"/>	<input type="button" value="▲"/>	<input type="checkbox"/> Monthly	Number of backups to retain	12	<input type="button" value="▼"/>	<input type="button" value="▲"/>
<input type="checkbox"/> Hourly	Number of backups to retain	24	<input type="button" value="▼"/>	<input type="button" value="▲"/>																	
<input checked="" type="checkbox"/> Daily	Number of backups to retain	30	<input type="button" value="▼"/>	<input type="button" value="▲"/>																	
<input type="checkbox"/> Weekly	Number of backups to retain	52	<input type="button" value="▼"/>	<input type="button" value="▲"/>																	
<input type="checkbox"/> Monthly	Number of backups to retain	12	<input type="button" value="▼"/>	<input type="button" value="▲"/>																	
DP Volumes Data protection volume backups use the same retention period as defined in the source SnapMirror relationship by default. Use the API if you want to change this value																					
S3 Bucket Cloud Manager will create the S3 bucket after you complete the wizard																					

5

Select the volumes that you want to back up

Identify which volumes you want to back up in the Select Volumes page.

6

Restore your data, as needed

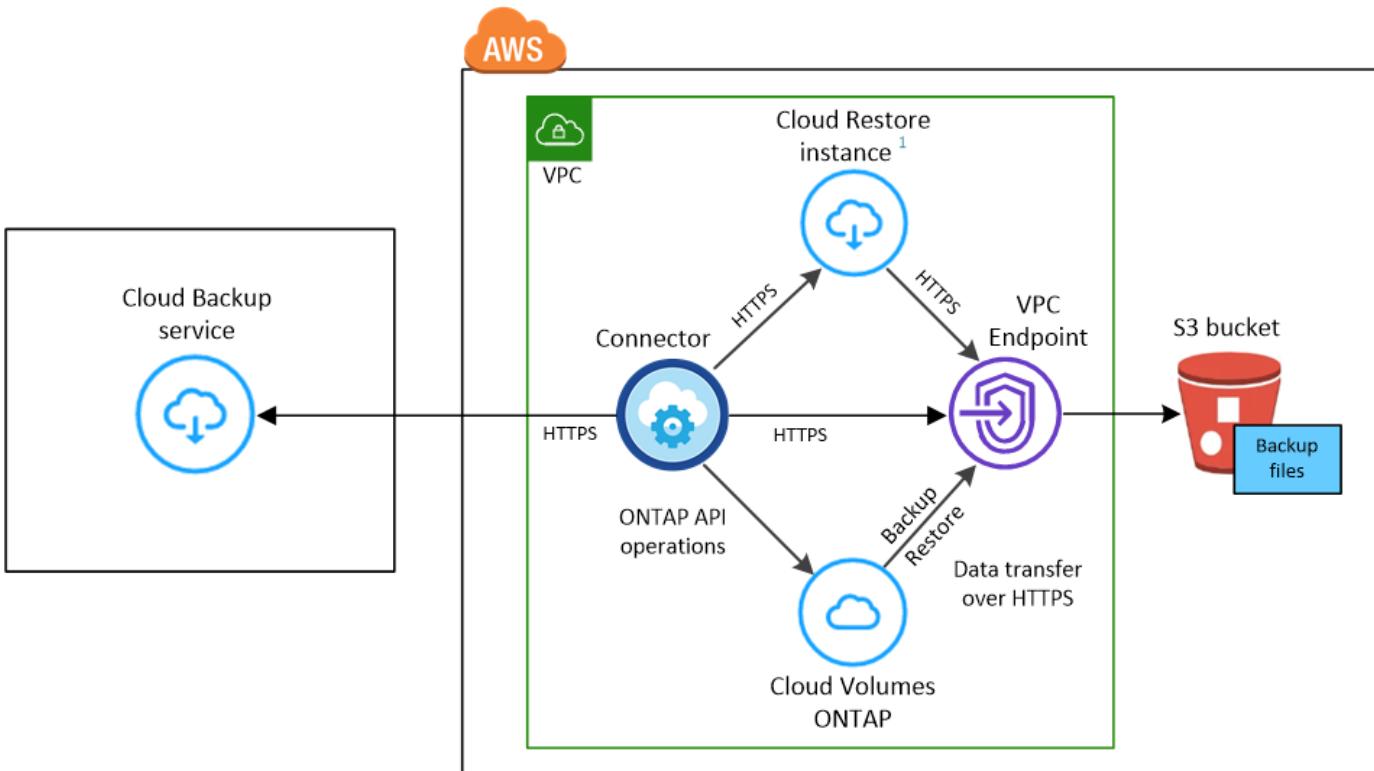
Choose to restore an entire backup to a new volume, or to restore individual files from the backup to an existing volume. You can restore data to a Cloud Volumes ONTAP system in AWS, or to an on-premises ONTAP system.

See [Restoring volume data from backup files](#) for details.

Requirements

Read the following requirements to make sure that you have a supported configuration before you start backing up volumes to S3.

The following image shows each component and the connections that you need to prepare between them:



¹ Cloud Restore instance is active only during single-file restore operations.

When the Cloud Restore instance is deployed in the cloud, it is located in the same subnet as the Connector.

Supported ONTAP versions

Cloud Volumes ONTAP 9.6 and later.

License requirements

For Cloud Backup PAYGO licensing, a Cloud Manager subscription is available in the AWS Marketplace that enables deployments of Cloud Volumes ONTAP and Cloud Backup. You need to [subscribe to this Cloud Manager subscription](#) before you enable Cloud Backup. Billing for Cloud Backup is done through this subscription.

For an annual contract that enables you to back up both Cloud Volumes ONTAP data and on-premises ONTAP data, you need to subscribe from the [AWS Marketplace page](#) and then [associate the subscription with your AWS credentials](#).

For an annual contract that enables you to bundle Cloud Volumes ONTAP and Cloud Backup Service, you must set up the annual contract when you create a Cloud Volumes ONTAP working environment. This option doesn't enable you to back up on-prem data.

For Cloud Backup BYOL licensing, you need the serial number from NetApp that enables you to use the service for the duration and capacity of the license. [Learn how to manage your BYOL licenses](#).

And you need to have an AWS account for the storage space where your backups will be located.

Supported AWS regions

Cloud Backup is supported in all AWS regions [where Cloud Volumes ONTAP is supported](#).

Required setup for creating backups in a different AWS account

By default, backups are created using the same account as the one used for your Cloud Volumes ONTAP system. If you want to use a different AWS account for your backups, you must [log in to the AWS portal and link the two accounts](#).

Required information for using customer-managed keys for data encryption

You can choose your own customer-managed keys for data encryption in the activation wizard instead of using the default Amazon S3 encryption keys. In this case you'll need to have the encryption managed keys already set up. [See how to use your own keys](#).

AWS Backup permissions required

The IAM role that provides Cloud Manager with permissions must include S3 permissions from the latest [Cloud Manager policy](#).

Here are the specific permissions from the policy:

```
{  
    "Sid": "backupPolicy",  
    "Effect": "Allow",  
    "Action": [  
        "s3:DeleteBucket",  
        "s3:GetLifecycleConfiguration",  
        "s3:PutLifecycleConfiguration",  
        "s3:PutBucketTagging",  
        "s3>ListBucketVersions",  
        "s3:GetObject",  
        "s3>DeleteObject",  
        "s3>ListBucket",  
        "s3>ListAllMyBuckets",  
        "s3:GetBucketTagging",  
        "s3:GetBucketLocation",  
        "s3:GetBucketPolicyStatus",  
        "s3:GetBucketPublicAccessBlock",  
        "s3:GetBucketAcl",  
        "s3:GetBucketPolicy",  
        "s3:PutBucketPublicAccessBlock"  
    ],  
    "Resource": [  
        "arn:aws:s3:::netapp-backup-*"  
    ]  
},
```

AWS Restore permissions required

The following EC2 permissions are needed for the IAM role that provides Cloud Manager with permissions so that it can start, stop, and terminate the Cloud Restore instance:

```

    "Action": [
        "ec2:DescribeInstanceTypeOfferings",
        "ec2:StartInstances",
        "ec2:StopInstances",
        "ec2:TerminateInstances"
    ],

```

Required outbound internet access for AWS deployments

The Cloud Restore instance requires outbound internet access. If your virtual or physical network uses a proxy server for internet access, ensure that the instance has outbound internet access to contact the following endpoints.

Endpoints	Purpose
http://amazonlinux.us-east-1.amazonaws.com/2/extras/docker/stable/x86_64/4bf88ee77c395ffe1e0c3ca68530dfb3a683ec65a4a1ce9c0ff394be50e922b2/	CentOS package for the Cloud Restore Instance AMI.
http://cloudmanagerinfraprod.azurecr.io https://cloudmanagerinfraprod.azurecr.io	Cloud Restore Instance image repository.

Enabling Cloud Backup on a new system

Cloud Backup is enabled by default in the working environment wizard. Be sure to keep the option enabled.

See [Launching Cloud Volumes ONTAP in AWS](#) for requirements and details for creating your Cloud Volumes ONTAP system.

Steps

1. Click **Create Cloud Volumes ONTAP**.
2. Select Amazon Web Services as the cloud provider and then choose a single node or HA system.
3. Fill out the Details & Credentials page.
4. On the Services page, leave the service enabled and click **Continue**.


Backup to Cloud
(On)

Integrated backup for Cloud Volumes ONTAP based on SnapMirror and Snapshot technologies. Backup copies are maintained in S3 buckets. Backups stored in S3 are charged separately from Cloud Volumes ONTAP.

ADVANTAGES <ul style="list-style-type: none"> ✓ Automatically back up all volumes. ✓ Creates new backup copy every day. ✓ Retains backups for 30 days. 	CLARIFICATIONS <ul style="list-style-type: none"> ➤ Backup settings are editable after working environment creation.
--	--

5. Complete the pages in the wizard to deploy the system.

Result

Cloud Backup is enabled on the system and backs up volumes every day and retains the most recent 30 backup copies.

What's next?

You can [start and stop backups for volumes](#) or [change the backup schedule](#) and you can [restore entire volumes](#) or individual files from a backup file.

Enabling Cloud Backup on an existing system

Enable Cloud Backup at any time directly from the working environment.

Steps

1. Select the working environment and click **Enable** next to the Backup & Restore service in the right-panel.



2. Select the provider details and click **Next**.

- a. The AWS Account used to store the backups. This can be a different account than where the Cloud Volumes ONTAP system resides.

If you want to use a different AWS account for your backups, you must [log in to the AWS portal and link the two accounts](#).

- b. The region where the backups will be stored. This can be a different region than where the Cloud Volumes ONTAP system resides.

- c. Whether you'll use the default Amazon S3 encryption keys or choose your own customer-managed keys from your AWS account to manage encryption of your data. ([See how to use your own keys](#)).

A screenshot of a configuration window titled 'Provider Settings'. The window is divided into several sections: 'Provider Information' containing fields for 'AWS Account' (set to 'AWS_Account_1'), 'AWS Access Key' (with a placeholder 'Enter AWS Access Key'), and 'AWS Secret Key' (with a placeholder 'Enter AWS Secret Key'); 'Location & Connectivity' containing a 'Region' dropdown set to 'us-east-2'; and 'Encryption' containing a note 'Encryption Key Type: AWS SSE-S3' and a 'Change Key' button with a gear icon.

3. Define the backup schedule and retention value and click **Next**.

Define Policy

Policy - Retention & Schedule

Create a New Policy Select an Existing Policy

<input type="checkbox"/> Hourly	Number of backups to retain	24
<input checked="" type="checkbox"/> Daily	Number of backups to retain	30
<input type="checkbox"/> Weekly	Number of backups to retain	52
<input type="checkbox"/> Monthly	Number of backups to retain	12

DP Volumes Data protection volume backups use the same retention period as defined in the source SnapMirror relationship by default. Use the API if you want to change this value

S3 Bucket Cloud Manager will create the S3 bucket after you complete the wizard

See the list of existing policies.

4. Select the volumes that you want to back up and click **Activate Backup**.

Select Volumes						
57 Volumes						
<input checked="" type="checkbox"/>	Volume Name	Volume Type	SVM Name	Used Capacity	Allocated Capacity	Backup Status
<input checked="" type="checkbox"/>	Volume_Name_1	RW	SVM_Name_1	0.25 TB	10 TB	<input type="radio"/> Not Active
<input checked="" type="checkbox"/>	Volume_Name_2	RW	SVM_Name_2	0.25 TB	10 TB	<input type="radio"/> Not Active
<input checked="" type="checkbox"/>	Volume_Name_3	RW	SVM_Name_3	0.25 TB	10 TB	<input type="radio"/> Not Active
<input checked="" type="checkbox"/>	Volume_Name_4	DP	SVM_Name_4	0.25 TB	10 TB	<input type="radio"/> Not Active
<input checked="" type="checkbox"/>	Volume_Name_5	RW	SVM_Name_5	0.25 TB	10 TB	<input type="radio"/> Not Active

- To back up all volumes, check the box in the title row (Volume Name).
- To back up individual volumes, check the box for each volume (Volume_1).

Result

Cloud Backup starts taking the initial backups of each selected volume and the Backup Dashboard is displayed so you can monitor the state of the backups.

What's next?

You can [start and stop backups for volumes or change the backup schedule](#) and you can [restore entire volumes or individual files from a backup file](#).

Backing up Cloud Volumes ONTAP data to Azure Blob storage

Complete a few steps to get started backing up data from Cloud Volumes ONTAP to

Azure Blob storage.

Quick start

Get started quickly by following these steps or scroll down to the remaining sections for full details.

1

Verify support for your configuration

- You're running Cloud Volumes ONTAP 9.7 or later in Azure.
- You have a valid cloud provider subscription for the storage space where your backups will be located.
- You have subscribed to the [Cloud Manager Marketplace Backup offering](#), or you have purchased [and activated](#) a Cloud Backup BYOL license from NetApp.

2

Enable Cloud Backup on your new or existing system

- New systems: Cloud Backup is enabled by default in the working environment wizard. Be sure to keep the option enabled.
- Existing systems: Select the working environment and click **Enable** next to the Backup & Restore service in the right-panel, and then follow the setup wizard.



3

Enter the provider details

Select the provider subscription and region, and choose whether you want to create a new resource group or use an already existing resource group. You can also choose your own customer-managed keys for data encryption instead of using the default Microsoft-managed encryption key.

Provider Settings

Azure Subscription	Region
Azure_Subscription_1	Default_CM_Region
Resource Group	Encryption Managed Keys
<input checked="" type="radio"/> Create a new	<input checked="" type="radio"/> Microsoft-managed
<input type="radio"/> Use an existing	<input type="radio"/> Customer-managed
Resource Group Name	

4

Define the backup policy

The default policy backs up volumes every day and retains the most recent 30 backup copies of each volume. Change to hourly, daily, weekly, or monthly backups, or select one of the system-defined policies that provide more options. You can also change the number of backup copies to retain.

Define Policy

Policy - Retention & Schedule Create a New Policy Select an Existing Policy

<input type="checkbox"/> Hourly	Number of backups to retain	24
<input checked="" type="checkbox"/> Daily	Number of backups to retain	30
<input type="checkbox"/> Weekly	Number of backups to retain	52
<input type="checkbox"/> Monthly	Number of backups to retain	12

DP Volumes	Data protection volume backups use the same retention period as defined in the source SnapMirror relationship by default. Use the API if you want to change this value
Storage Account	Cloud Manager will create the storage account after you complete the wizard

5

Select the volumes that you want to back up

Identify which volumes you want to back up in the Select Volumes page.

6

Restore your data, as needed

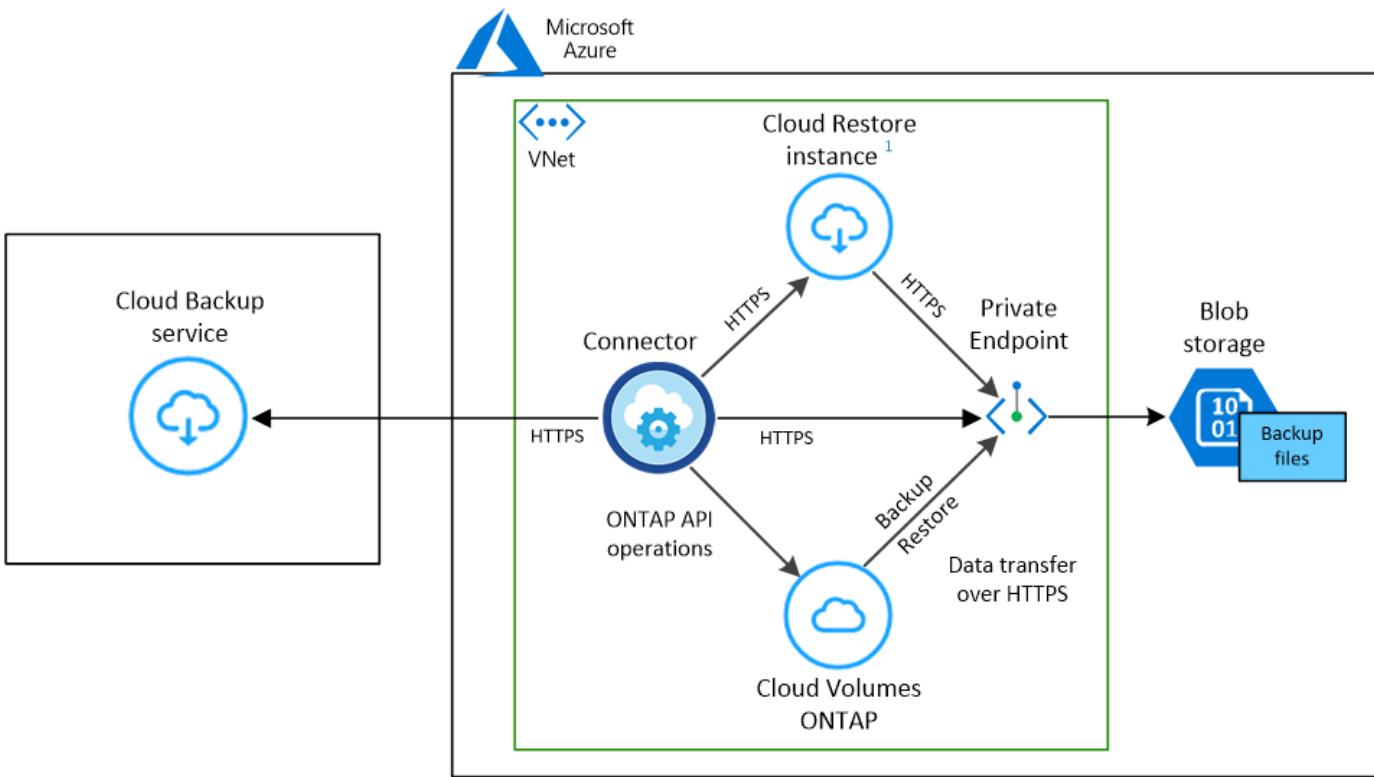
Choose to restore an entire backup to a new volume, or to restore individual files from the backup to an existing volume. You can restore data to a Cloud Volumes ONTAP system in Azure, or to an on-premises ONTAP system.

See [Restoring volume data from backup files](#) for details.

Requirements

Read the following requirements to make sure that you have a supported configuration before you start backing up volumes to Azure Blob storage.

The following image shows each component and the connections that you need to prepare between them:



¹ Cloud Restore instance is active only during single-file restore operations.

When the Cloud Restore virtual machine is deployed in the cloud, it is located in the same subnet as the Connector.

Supported ONTAP versions

Cloud Volumes ONTAP 9.7 and later.

License requirements

For Cloud Backup PAYGO licensing, a subscription through the Azure Marketplace is required before you enable Cloud Backup. Billing for Cloud Backup is done through this subscription. [You can subscribe from the Details & Credentials page of the working environment wizard.](#)

For Cloud Backup BYOL licensing, you need the serial number from NetApp that enables you to use the service for the duration and capacity of the license. [Learn how to manage your BYOL licenses.](#)

And you need to have a Microsoft Azure subscription for the storage space where your backups will be located.

Supported Azure regions

Cloud Backup is supported in all Azure regions [where Cloud Volumes ONTAP is supported](#).

Required setup for creating backups in a different Azure subscription

By default, backups are created using the same subscription as the one used for your Cloud Volumes ONTAP system. If you want to use a different Azure subscription for your backups, you must [log in to the Azure portal and link the two subscriptions](#).

Required information for using customer-managed keys for data encryption

You can use your own customer-managed keys for data encryption in the activation wizard instead of using the default Microsoft-managed encryption keys. In this case you will need to have the Azure Subscription, Key Vault name, and the Key. [See how to use your own keys](#).

Required outbound internet access for Azure deployments

The Cloud Restore virtual machine requires outbound internet access. If your virtual or physical network uses a proxy server for internet access, ensure that the instance has outbound internet access to contact the following endpoints.

Endpoints	Purpose
http://olcentgbl.trafficmanager.net https://olcentgbl.trafficmanager.net	Provides CentOS packages for the Cloud Restore virtual machine.
http://cloudmanagerinfraprod.azurecr.io https://cloudmanagerinfraprod.azurecr.io	Cloud Restore virtual machine image repository.

Enabling Cloud Backup on a new system

Cloud Backup is enabled by default in the working environment wizard. Be sure to keep the option enabled.

See [Launching Cloud Volumes ONTAP in Azure](#) for requirements and details for creating your Cloud Volumes ONTAP system.



If you want to pick the name of the resource group, **disable** Cloud Backup when deploying Cloud Volumes ONTAP. Follow the steps for [enabling Cloud Backup on an existing system](#) to enable Cloud Backup and choose the resource group.

Steps

1. Click **Create Cloud Volumes ONTAP**.
2. Select Microsoft Azure as the cloud provider and then choose a single node or HA system.
3. In the Define Azure Credentials page, enter the credentials name, client ID, client secret, and directory ID, and click **Continue**.
4. Fill out the Details & Credentials page and be sure that an Azure Marketplace subscription is in place, and click **Continue**.
5. On the Services page, leave the service enabled and click **Continue**.

The screenshot shows the 'Backup to Cloud' service configuration page. At the top, there's a toggle switch set to 'On' and a help icon. Below the title, a descriptive text explains that integrated backup uses SnapMirror and Snapshot technologies, stored in Storage Accounts. The 'ADVANTAGES' section lists three benefits: automatic backup of all volumes, daily creation of new backup copies, and retention for 30 days. The 'CLARIFICATIONS' section notes that backup settings are editable after the environment is created.

6. Complete the pages in the wizard to deploy the system.

Result

Cloud Backup is enabled on the system and backs up volumes every day and retains the most recent 30 backup copies.

What's next?

You can start and stop backups for volumes or change the backup schedule and you can restore entire volumes or individual files from a backup file.

Enabling Cloud Backup on an existing system

Enable Cloud Backup at any time directly from the working environment.

Steps

1. Select the working environment and click **Enable** next to the Backup & Restore service in the right-panel.



2. Select the provider details and click **Next**:

- a. The Azure subscription used to store the backups. This can be a different subscription than where the Cloud Volumes ONTAP system resides.

If you want to use a different Azure subscription for your backups, you must [log in to the Azure portal and link the two subscriptions](#).

- b. The region where the backups will be stored. This can be a different region than where the Cloud Volumes ONTAP system resides.
- c. The resource group that manages the Blob container - you can create a new resource group or select an existing resource group.
- d. Whether you'll use the default Microsoft-managed encryption key or choose your own customer-managed keys to manage encryption of your data. ([See how to use your own keys](#)).

A screenshot of a 'Provider Settings' configuration page. At the top center is the title 'Provider Settings'. Below it are several input fields:

- 'Azure Subscription': A dropdown menu showing 'Azure_Subscription_1'.
- 'Region': A dropdown menu showing 'Default_CM_Region'.
- 'Resource Group': A section with a radio button for 'Create a new' (selected) and another for 'Use an existing'.
- 'Encryption Managed Keys': A section with a radio button for 'Microsoft-managed' (selected) and another for 'Customer-managed'.

3. In the *Define Policy* page, select the backup schedule and retention value and click **Next**.

Define Policy

Policy - Retention & Schedule	<input checked="" type="radio"/> Create a New Policy <input type="radio"/> Select an Existing Policy												
<table><tr><td><input type="checkbox"/> Hourly</td><td>Number of backups to retain</td><td>24</td></tr><tr><td><input checked="" type="checkbox"/> Daily</td><td>Number of backups to retain</td><td>30</td></tr><tr><td><input type="checkbox"/> Weekly</td><td>Number of backups to retain</td><td>52</td></tr><tr><td><input type="checkbox"/> Monthly</td><td>Number of backups to retain</td><td>12</td></tr></table>		<input type="checkbox"/> Hourly	Number of backups to retain	24	<input checked="" type="checkbox"/> Daily	Number of backups to retain	30	<input type="checkbox"/> Weekly	Number of backups to retain	52	<input type="checkbox"/> Monthly	Number of backups to retain	12
<input type="checkbox"/> Hourly	Number of backups to retain	24											
<input checked="" type="checkbox"/> Daily	Number of backups to retain	30											
<input type="checkbox"/> Weekly	Number of backups to retain	52											
<input type="checkbox"/> Monthly	Number of backups to retain	12											
DP Volumes	Data protection volume backups use the same retention period as defined in the source SnapMirror relationship by default. Use the API if you want to change this value												
Storage Account	Cloud Manager will create the storage account after you complete the wizard												

See the list of existing policies.

4. Select the volumes that you want to back up and click **Activate Backup**.

Select Volumes						
57 Volumes						
<input checked="" type="checkbox"/>	Volume Name	Volume Type	SVM Name	Used Capacity	Allocated Capacity	Backup Status
<input checked="" type="checkbox"/>	Volume_Name_1	RW	SVM_Name_1	0.25 TB	10 TB	<input type="radio"/> Not Active
<input checked="" type="checkbox"/>	Volume_Name_2	RW	SVM_Name_2	0.25 TB	10 TB	<input type="radio"/> Not Active
<input checked="" type="checkbox"/>	Volume_Name_3	RW	SVM_Name_3	0.25 TB	10 TB	<input type="radio"/> Not Active
<input checked="" type="checkbox"/>	Volume_Name_4	DP	SVM_Name_4	0.25 TB	10 TB	<input type="radio"/> Not Active
<input checked="" type="checkbox"/>	Volume_Name_5	RW	SVM_Name_5	0.25 TB	10 TB	<input type="radio"/> Not Active

- To back up all volumes, check the box in the title row (Volume Name).
- To back up individual volumes, check the box for each volume (Volume_1).

Result

Cloud Backup starts taking the initial backups of each selected volume and the Backup Dashboard is displayed so you can monitor the state of the backups.

What's next?

You can [start and stop backups for volumes or change the backup schedule](#) and you can [restore entire volumes or individual files from a backup file](#).

Backing up Cloud Volumes ONTAP data to Google Cloud Storage

Complete a few steps to get started backing up data from Cloud Volumes ONTAP to

Google Cloud Storage.

Quick start

Get started quickly by following these steps or scroll down to the remaining sections for full details.

1

Verify support for your configuration

- You're running Cloud Volumes ONTAP 9.7P5 or later in GCP.
- You have a valid GCP subscription for the storage space where your backups will be located.
- You have a service account in your Google Cloud Project that has the predefined Storage Admin role.
- You have subscribed to the [Cloud Manager Marketplace Backup offering](#), or you have purchased [and activated](#) a Cloud Backup BYOL license from NetApp.

2

Enable Cloud Backup on your new or existing system

- New systems: Cloud Backup can be enabled when you complete the new working environment wizard.
- Existing systems: Select the working environment and click **Enable** next to the Backup & Restore service in the right-panel, and then follow the setup wizard.



3

Enter the provider details

Select the Google Cloud Project where you want the Google Cloud Storage bucket to be created for backups.

A screenshot of a 'Provider Settings' dialog box. At the top, the title 'Provider Settings' is in blue. Below it, the section 'Google Cloud Project' has a dropdown menu set to 'Default Project'. The section 'Region' also has a dropdown menu set to 'us-east-2'.

4

Define the backup policy

The default policy backs up volumes every day and retains the most recent 30 backup copies of each volume. Change to hourly, daily, weekly, or monthly backups, or select one of the system-defined policies that provide more options.

Define Policy

Policy - Retention & Schedule	<input checked="" type="radio"/> Create a New Policy <input type="radio"/> Select an Existing Policy
	<input type="checkbox"/> Hourly Number of backups to retain <input type="text" value="24"/>
	<input checked="" type="checkbox"/> Daily Number of backups to retain <input type="text" value="30"/>
	<input type="checkbox"/> Weekly Number of backups to retain <input type="text" value="52"/>
	<input type="checkbox"/> Monthly Number of backups to retain <input type="text" value="12"/>
DP Volumes Data protection volume backups use the same retention period as defined in the source SnapMirror relationship by default. Use the API if you want to change this value	
Google Cloud Storage Bucket Cloud Manager will create the Google Cloud Storage Bucket after you complete the wizard	

5

Select the volumes that you want to back up

Identify which volumes you want to back up in the Select Volumes page.

6

Restore your data, as needed

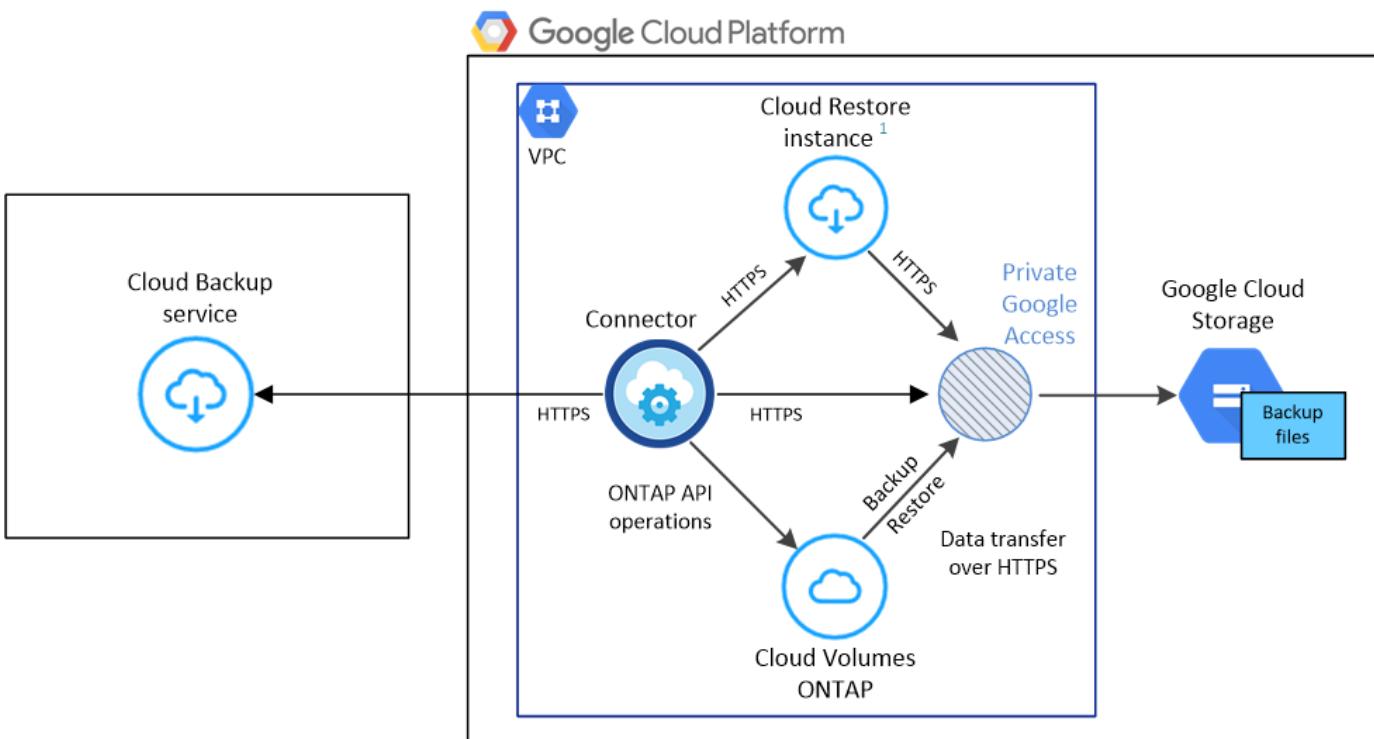
Restore a backup to a new volume. You can restore data to a Cloud Volumes ONTAP system in Google. A Service Account is required on the Cloud Volumes ONTAP system where you are performing the restore.

See [Restoring volume data from backup files](#) for details.

Requirements

Read the following requirements to make sure that you have a supported configuration before you start backing up volumes to Google Cloud storage.

The following image shows each component and the connections that you need to prepare between them:



¹ Cloud Restore instance is active only during single-file restore operations.

Supported ONTAP versions

Cloud Volumes ONTAP 9.7P5 and later.

Supported GCP regions

Cloud Backup is supported in all GCP regions [where Cloud Volumes ONTAP is supported](#).

License requirements

For Cloud Backup PAYGO licensing, a subscription through the [GCP Marketplace](#) is required before you enable Cloud Backup. Billing for Cloud Backup is done through this subscription. [You can subscribe from the Details & Credentials page of the working environment wizard](#).

For Cloud Backup BYOL licensing, you need the serial number from NetApp that enables you to use the service for the duration and capacity of the license. [Learn how to manage your BYOL licenses](#).

And you need to have a Google subscription for the storage space where your backups will be located.

GCP Service Account

You need to have a service account in your Google Cloud Project that has the predefined Storage Admin role. [Learn how to create a service account](#).

Enabling Cloud Backup on a new system

Cloud Backup can be enabled when you complete the working environment wizard to create a new Cloud Volumes ONTAP system.

You must have a Service Account already configured. If you don't select a service account when you create the Cloud Volumes ONTAP system, then you'll need to turn off the system and add the service account to Cloud Volumes ONTAP from the GCP console.

See [Launching Cloud Volumes ONTAP in GCP](#) for requirements and details for creating your Cloud Volumes

ONTAP system.

Steps

1. On the Working Environments page, click **Add Working Environment** and follow the prompts.
2. **Choose a Location:** Select **Google Cloud Platform**.
3. **Choose Type:** Select **Cloud Volumes ONTAP** (either single-node or high-availability).
4. **Details & Credentials:** Enter the following information:
 - a. Click **Edit Project** and select a new project if the one you want to use is different than the default Project (where Cloud Manager resides).
 - b. Specify the cluster name.
 - c. Enable the **Service Account** switch and select the Service Account that has the predefined Storage Admin role. This is required to enable backups and tiering.
 - d. Specify the credentials.

Make sure that a GCP Marketplace subscription is in place.

Details & Credentilas

Project1 Google Cloud Project	MPAWSSubscription1222 Marketplace Subscription	Edit Project
Details		Credentials
Working Environment Name (Cluster Name) <input type="text" value="TamiVSA"/>		User Name <input type="text" value="admin"/>
Service Account <small>i</small> <input checked="" type="checkbox"/>		Password <input type="password" value="*****"/>
Service Account Name <input type="text" value="ServiceAccount1"/>		Confirm Password <input type="password" value="*****"/>
Add Labels <small>Optional Field Up to four labels</small>		

5. **Services:** Leave the Cloud Backup service enabled and click **Continue**.

Services

<input checked="" type="checkbox"/> Backup to Cloud	<input checked="" type="checkbox"/> <small>Cloud</small>
---	--

6. Complete the pages in the wizard to deploy the system as described in [Launching Cloud Volumes ONTAP in GCP](#).

Result

Cloud Backup is enabled on the system and backs up the volume you created every day and retains the most recent 30 backup copies.

You can start and stop backups for volumes or change the backup schedule and you can restore entire volumes or individual files from a backup file.

Enabling Cloud Backup on an existing system

You can enable Cloud Backup at any time directly from the working environment.

Steps

1. Select the working environment and click **Enable** next to the Backup & Restore service in the right-panel.



2. Select the Google Cloud Project and region where you want the Google Cloud Storage bucket to be created for backups, and click **Next**.

A screenshot of a 'Provider Settings' configuration window. It contains two dropdown menus: 'Google Cloud Project' set to 'Default Project' and 'Region' set to 'us-east-2'.

Note that the Project must have a Service Account that has the predefined Storage Admin role.

3. In the *Define Policy* page, select the backup schedule and retention value and click **Next**.

A screenshot of the 'Define Policy' configuration window. It features several sections:

- Policy - Retention & Schedule:** Contains radio buttons for 'Create a New Policy' (selected) and 'Select an Existing Policy'. Below are four schedule options: 'Hourly' (unchecked), 'Daily' (checked, with a retention of 24), 'Weekly' (unchecked), and 'Monthly' (unchecked).
- DP Volumes:** A note stating that data protection volume backups use the same retention period as defined in the source SnapMirror relationship by default. It advises using the API if changes are needed.
- Google Cloud Storage Bucket:** A note indicating that the Cloud Manager will create the Google Cloud Storage Bucket after completing the wizard.

[See the list of existing policies.](#)

4. Select the volumes that you want to back up and click **Activate Backup**.

Select Volumes						
57 Volumes						
<input checked="" type="checkbox"/>	Volume Name	Volume Type	SVM Name	Used Capacity	Allocated Capacity	Backup Status
<input checked="" type="checkbox"/>	Volume_Name_1	RW	SVM_Name_1	0.25 TB	10 TB	Not Active
<input checked="" type="checkbox"/>	Volume_Name_2	RW	SVM_Name_2	0.25 TB	10 TB	Not Active
<input checked="" type="checkbox"/>	Volume_Name_3	RW	SVM_Name_3	0.25 TB	10 TB	Not Active
<input checked="" type="checkbox"/>	Volume_Name_4	DP	SVM_Name_4	0.25 TB	10 TB	Not Active
<input checked="" type="checkbox"/>	Volume_Name_5	RW	SVM_Name_5	0.25 TB	10 TB	Not Active

- To back up all volumes, check the box in the title row (**Volume Name**).
- To back up individual volumes, check the box for each volume (**Volume_1**).

Result

Cloud Backup starts taking the initial backups of each selected volume and the Backup Dashboard is displayed so you can monitor the state of the backups.

What's next?

You can [start and stop backups for volumes or change the backup schedule](#) and you can [restore entire volumes or individual files from a backup file](#).

Backing up on-premises ONTAP data to Amazon S3

Complete a few steps to get started backing up data from your on-premises ONTAP systems to Amazon S3 storage.

TIP

In most cases you'll use Cloud Manager for all backup and restore operations. However, starting with ONTAP 9.9.1 you can initiate volume backup operations of your on-premises ONTAP clusters using ONTAP System Manager. [See how to use System Manager to back up your volumes to the cloud using Cloud Backup](#).

A Beta feature released in January 2021 allows you to run compliance scans on the backed up volumes from your on-premises systems. Typically, compliance scans are free up to 1 TB of data, and then a cost for the service is applied for data over 1 TB. When combining Backup and Data Sense for your on-premises volumes, the cost for scans on those on-prem volumes is free. Learn more about how [Cloud Data Sense](#) can get your business applications and cloud environments privacy ready.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.

1

Verify support for your configuration

- You have discovered the on-premises cluster and added it to a working environment in Cloud Manager. See [Discovering ONTAP clusters](#) for details.
 - The cluster is running ONTAP 9.7P5 or later.
 - The cluster has a SnapMirror license — it is included as part of the Premium Bundle or Data Protection Bundle.
 - The cluster must have the required network connections to S3 storage and to the Connector.
- The Connector must have the required network connections to S3 storage and to the cluster, and the required permissions.
- You have a valid AWS subscription for the object storage space where your backups will be located.
- You have an AWS Account with an access key and secret key, and the [required permissions](#) so the ONTAP cluster can back up and restore data.

2

Enable Cloud Backup on the system

Select the working environment and click **Enable** next to the Backup & Compliance service in the right-panel, and then follow the setup wizard.



3

Select the cloud provider and enter the provider details

Select Amazon Web Services as your provider and then enter the provider details. You'll need to select the AWS Account and the region where you want to create the backups. You can also choose your own customer-managed key for data encryption instead of using the default Amazon S3 encryption key.

A screenshot of a 'Provider Settings' dialog box. The title 'Provider Settings' is at the top center. The dialog is divided into two main sections: 'Provider Information' on the left and 'Location & Connectivity' on the right. In the 'Provider Information' section, there are three dropdown menus: 'AWS Account' (set to 'AWS_Account_1'), 'AWS Access Key' (with a placeholder 'Enter AWS Access Key'), and 'AWS Secret Key' (with a placeholder 'Enter AWS Secret Key'). In the 'Location & Connectivity' section, there are two dropdown menus: 'Region' (set to 'us-east-2') and 'Encryption' (with a placeholder 'Enter Encryption Key'). Below the 'Encryption' field, it says 'Encryption Key Type: AWS SSE-S3' and has a 'Change Key' link.

4

Select the cluster IPspace and optionally select an AWS PrivateLink connection

Select the IPspace in the ONTAP cluster where the volumes reside. You can also choose to use an existing AWS PrivateLink configuration for a more secure connection to the VPC from your on-prem data center.

Networking

IPspace

IP_Space_1

Private Link Configuration

Select Private Link

Name	VPC	Endpoint ID
Private_Link_Name_001	vpce0-012345678901234567890 (Default)	vpce0-012345678901234567890
Private_Link_Name_002	vpce0-012345678901234567890 (k8s)	vpce0-012345678901234567890

5

Define the backup policy

The default policy backs up volumes every day and retains the most recent 30 backup copies of each volume. Change to hourly, daily, weekly, or monthly backups, or select one of the system-defined policies that provide more options. You can also change the number of backup copies to retain.

Define Policy

Policy - Retention & Schedule

Create a New Policy Select an Existing Policy

Select Policy

Default Policy (30 Daily)

DP Volumes

Data protection volume backups use the same retention period as defined in the source SnapMirror relationship by default. Use the API if you want to change this value

6

Select the volumes that you want to back up

Identify which volumes you want to back up from the cluster.

7

Activate Compliance scans on the backed up volumes (optional)

Choose whether you want to have Cloud Data Sense scan the volumes that are backed up in the cloud.

8

Restore your data, as needed

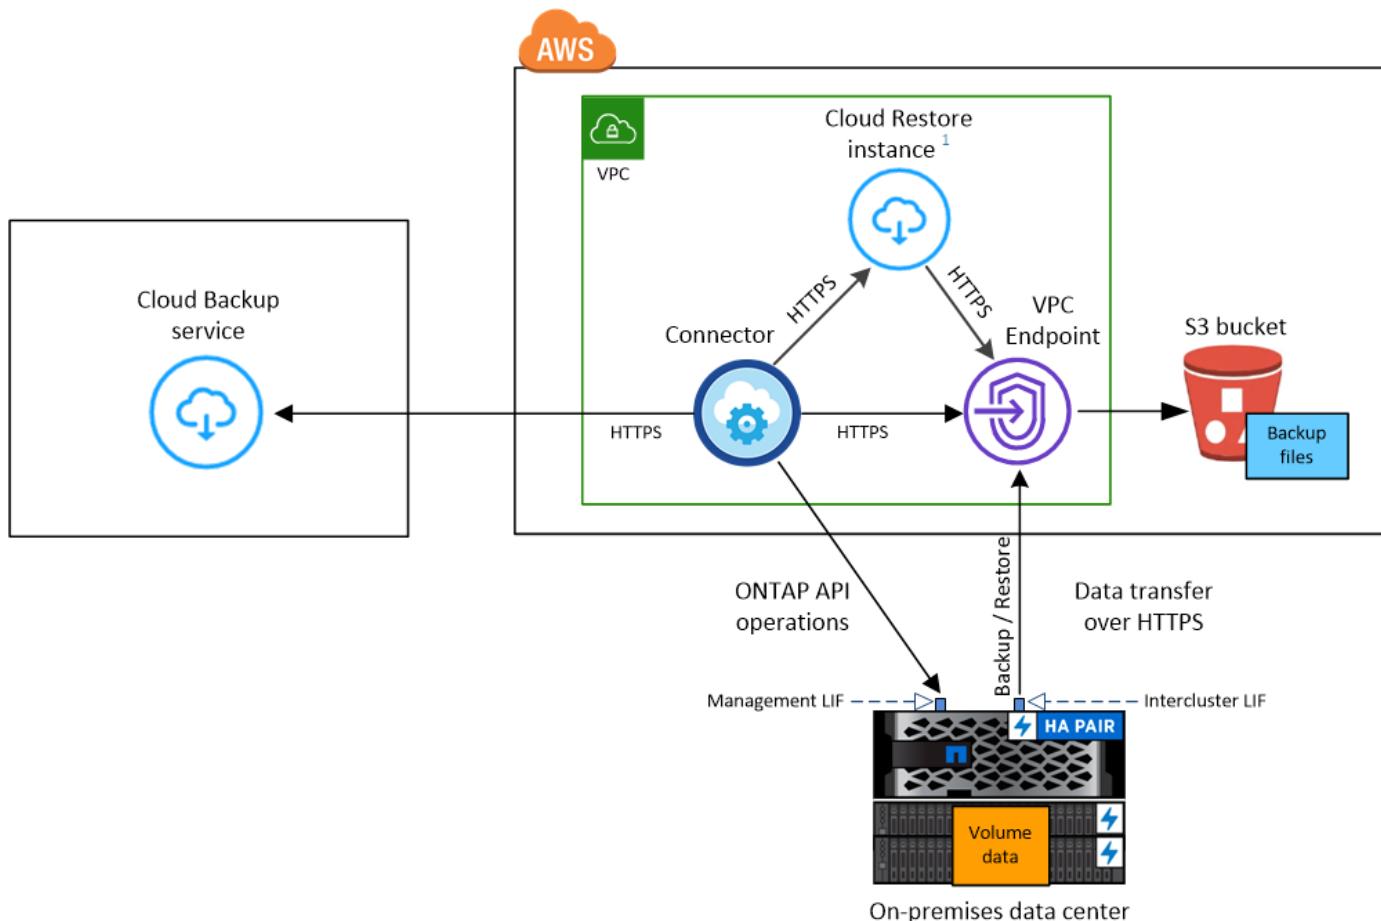
Choose to restore an entire backup to a new volume, or to restore individual files from the backup to an existing volume. You can restore data to a Cloud Volumes ONTAP system in AWS, or to an on-premises ONTAP system.

See [Restoring volume data from backup files](#) for details.

Requirements

Read the following requirements to make sure you have a supported configuration before you start backing up on-premises volumes to S3 storage.

The following image shows each component and the connections that you need to prepare between them:



Note that when the Cloud Restore instance is deployed in the cloud, it is located in the same subnet as the Connector.

Preparing your ONTAP clusters

You need to discover your on-premises ONTAP clusters in Cloud Manager before you can start backing up volume data.

[Learn how to discover a cluster.](#)

ONTAP requirements

- ONTAP 9.7P5 and later.
- A SnapMirror license (included as part of the Premium Bundle or Data Protection Bundle).

Note: The "Hybrid Cloud Bundle" is not required when using the Cloud Backup service.

See how to [manage your cluster licenses](#).

- Time and time zone are set correctly.

See how to [configure your cluster time](#).

Cluster networking requirements

- The ONTAP cluster initiates an HTTPS connection over port 443 from the intercluster LIF to Amazon S3 storage for backup and restore operations.

ONTAP reads and writes data to and from object storage. The object storage never initiates, it just responds.

- ONTAP requires an inbound connection from the Connector to the cluster management LIF. The Connector can reside in an AWS VPC.
- An intercluster LIF is required on each ONTAP node that hosts the volumes you want to back up. The LIF must be associated with the *IPspace* that ONTAP should use to connect to object storage. [Learn more about IPspaces](#).

When you set up Cloud Backup, you are prompted for the IPspace to use. You should choose the IPspace that each LIF is associated with. That might be the "Default" IPspace or a custom IPspace that you created.

- The nodes' intercluster LIFs are able to access the internet.
- DNS servers have been configured for the storage VM where the volumes are located. See how to [configure DNS services for the SVM](#).
- Note that if you use are using a different IPspace than the Default, then you might need to create a static route to get access to the object storage.
- Update firewall rules, if necessary, to allow Cloud Backup service connections from ONTAP to object storage through port 443 and name resolution traffic from the storage VM to the DNS server over port 53 (TCP/UDP).

Creating or switching Connectors

A Connector is required to back up data to the cloud, and the Connector must be in an AWS VPC when backing up data to AWS S3 storage. You can't use a Connector that's deployed on-premises. You'll either need to create a new Connector or make sure that the currently selected Connector resides in the correct provider.

- [Learn about Connectors](#)

- [Creating a Connector in AWS](#)
- [Switching between Connectors](#)

Preparing networking for the Connector

Ensure that the Connector has the required networking connections.

Steps

1. Ensure that the network where the Connector is installed enables the following connections:
 - An outbound internet connection to the Cloud Backup service over port 443 (HTTPS)
 - An HTTPS connection over port 443 to your S3 object storage
 - An HTTPS connection over port 443 to your ONTAP clusters
2. Enable a VPC Endpoint to S3. This is needed if you have a Direct Connect or VPN connection from your ONTAP cluster to the VPC and you want communication between the Connector and S3 to stay in your AWS internal network.

Supported regions

You can create backups from on-premises systems to Amazon S3 in all regions [where Cloud Volumes ONTAP is supported](#). You specify the region where the backups will be stored when you set up the service.

License requirements

Before your 30-day free trial of the Cloud Backup service expires, you need to subscribe to a pay-as-you-go (PAYGO) Cloud Manager Marketplace offering from AWS, or purchase and activate a Cloud Backup BYOL license from NetApp. These licenses are for the account and can be used across multiple systems.

- For Cloud Backup PAYGO licensing, you'll need a subscription to the [AWS Cloud Manager Marketplace](#) offering to continue using Cloud Backup. Billing for Cloud Backup is done through this subscription.
- For Cloud Backup BYOL licensing, you don't need a subscription. You need the serial number from NetApp that enables you to use the service for the duration and capacity of the license. [Learn how to manage your BYOL licenses](#).

You need to have an AWS subscription for the object storage space where your backups will be located.

A SnapMirror license is required on the cluster. Note that the "Hybrid Cloud Bundle" is not required when using Cloud Backup.

Preparing Amazon S3 for backups

When you are using Amazon S3, you must configure permissions for the Connector to create and manage the S3 bucket, and you must configure permissions so the on-premises ONTAP cluster can read and write to the S3 bucket.

Steps

1. Confirm that the following S3 permissions (from the latest [Cloud Manager policy](#)) are part of the IAM role that provides the Connector with permissions:

```
{
    "Sid": "backupPolicy",
    "Effect": "Allow",
    "Action": [
        "s3:DeleteBucket",
        "s3:GetLifecycleConfiguration",
        "s3:PutLifecycleConfiguration",
        "s3:PutBucketTagging",
        "s3>ListBucketVersions",
        "s3:GetObject",
        "s3>ListBucket",
        "s3>ListAllMyBuckets",
        "s3:GetBucketTagging",
        "s3:GetBucketLocation",
        "s3:GetBucketPolicyStatus",
        "s3:GetBucketPublicAccessBlock",
        "s3:GetBucketAcl",
        "s3:GetBucketPolicy",
        "s3:PutBucketPublicAccessBlock"
    ],
    "Resource": [
        "arn:aws:s3:::netapp-backup-*"
    ]
},
}
```

2. Add the following EC2 permissions to the IAM role that provides the Connector with permissions so that it can start, stop, and terminate the Cloud Restore instance:

```
"Action": [
    "ec2:DescribeInstanceTypeOfferings",
    "ec2:StartInstances",
    "ec2:StopInstances",
    "ec2:TerminateInstances"
],
```

3. During the Backup wizard you will be prompted to enter an access key and secret key. For that, you will need to create an IAM user with the following permissions. Cloud Backup passes these credentials on to the ONTAP cluster so that ONTAP can backup and restore data to the S3 bucket.

```

"s3>ListAllMyBuckets",
"s3>ListBucket",
"s3:GetBucketLocation",
"s3GetObject",
"s3PutObject",
"s3DeleteObject"

```

See the [AWS Documentation: Creating a Role to Delegate Permissions to an IAM User](#) for details.

4. If your virtual or physical network uses a proxy server for internet access, ensure that the Cloud Restore instance has outbound internet access to contact the following endpoints.

Endpoints	Purpose
http://amazonlinux.us-east-1.amazonaws.com/2/extras/docker/stable/x86_64/4bf88ee77c395ffe1e0c3ca68530dfb3a683ec65a4a1ce9c0ff394be50e922b2/	CentOS package for the Cloud Restore Instance AMI.
http://cloudmanagerinfraprod.azurecr.io https://cloudmanagerinfraprod.azurecr.io	Cloud Restore Instance image repository.

5. You can choose your own custom-managed keys for data encryption in the activation wizard instead of using the default Amazon S3 encryption keys. In this case you'll need to have the encryption managed keys already set up. [See how to use your own keys](#).
6. If you want to have a more secure connection over the public internet from your on-prem data center to the VPC, there is an option to select an AWS PrivateLink connection in the activation wizard. It is required if you are connecting your on-premises system via VPN/DirectConnect. In this case you'll need to have created an Interface endpoint configuration using the Amazon VPC console or the command line. [See details about using AWS PrivateLink](#).

Note that you'll also need to modify the security group configuration that is associated with the Cloud Manager Connector. You must change the policy to "Custom" (from "Full Access"), and you must add the permissions from the backup policy as shown earlier (above).

Policy*

- Full Access - Allow access by any user or service within the VPC using credentials from any AWS accounts to any resources in this AWS service. All policies — IAM user policies, VPC endpoint policies, and AWS service-specific policies (e.g. Amazon S3 bucket policies, any S3 ACL policies) — must grant the necessary permissions for access to succeed.
- Custom

Use the [policy creation tool](#) to generate a policy, then paste the generated policy below.

```
{
  "Sid": "backupPolicy",
  "Effect": "Allow",
  "Action": [
    "s3:DeleteBucket",
    "s3:GetLifecycleConfiguration",
    "s3:PutLifecycleConfiguration",
    "s3:PutBucketTagging",
    "s3>ListBucketVersions",
    "s3GetObject",
    "s3:DeleteObject",
    "s3>ListBucket",
    "s3>ListAllMyBuckets",
    "s3GetBucketTagging",
    "s3GetBucketLocation",
    ...
  ],
  "Resource": "*"
}
```

Enabling Cloud Backup

Enable Cloud Backup at any time directly from the on-premises working environment.

Steps

- From the Canvas, select the working environment and click **Enable** next to the Backup & Compliance service in the right-panel.



- Select Amazon Web Services as your provider and click **Next**.
- Enter the provider details. Note that you can't change this information after the service has started.
 - The AWS Account, the AWS Access Key, and the Secret Key used to store the backups.

The access key and secret key are for the user you created to give the ONTAP cluster access to the S3 bucket.

- The AWS region where the backups will be stored.
- Whether you'll use the default Amazon S3 encryption keys or choose your own customer-managed keys from your AWS account to manage encryption of your data. ([See how to use your own keys](#)).

Provider Settings

Provider Information <div style="margin-top: 10px;"> AWS Account <input type="text" value="AWS_Account_1"/> </div> <div style="margin-top: 10px;"> AWS Access Key <input type="text" value="Enter AWS Access Key"/> </div> <div style="margin-top: 10px;"> AWS Secret Key <input type="text" value="Enter AWS Secret Key"/> </div>	Location & Connectivity <div style="margin-top: 10px;"> Region <input type="text" value="us-east-2"/> </div> <div style="margin-top: 10px;"> Encryption <small>Encryption Key Type: AWS SSE-S3</small> </div> <div style="margin-top: 10px;"> Change Key </div>
--	--

4. Click **Next** after you've entered the provider details.
5. Enter the networking details and click **Next**.
 - a. The IPspace in the ONTAP cluster where the volumes you want to back up reside. The intercluster LIFs for this IPspace must have outbound internet access.
 - b. Optionally, choose whether you'll use an AWS PrivateLink that you have previously configured. [See details about using an AWS PrivateLink](#).

Networking

IPspace <input type="text" value="IP_Space_1"/>	<div style="border: 1px solid #ccc; padding: 5px; margin-bottom: 10px;"> <input checked="" type="checkbox"/> Private Link Configuration </div> <div style="border: 1px solid #ccc; padding: 5px; background-color: #f9f9f9;"> <p>Select Private Link</p> <table border="1" style="width: 100%; border-collapse: collapse;"> <thead> <tr style="background-color: #e0e0e0;"> <th>Name</th> <th>VPC</th> <th>Endpoint ID</th> </tr> </thead> <tbody> <tr> <td><input type="radio"/> Private_Link_Name_001</td> <td>vpce0-012345678901234567890 (Default)</td> <td>vpce0-012345678901234567890</td> </tr> <tr> <td><input type="radio"/> Private_Link_Name_002</td> <td>vpce0-012345678901234567890 (k8s)</td> <td>vpce0-012345678901234567890</td> </tr> </tbody> </table> </div>	Name	VPC	Endpoint ID	<input type="radio"/> Private_Link_Name_001	vpce0-012345678901234567890 (Default)	vpce0-012345678901234567890	<input type="radio"/> Private_Link_Name_002	vpce0-012345678901234567890 (k8s)	vpce0-012345678901234567890
Name	VPC	Endpoint ID								
<input type="radio"/> Private_Link_Name_001	vpce0-012345678901234567890 (Default)	vpce0-012345678901234567890								
<input type="radio"/> Private_Link_Name_002	vpce0-012345678901234567890 (k8s)	vpce0-012345678901234567890								

6. Select an existing backup schedule and retention value, or define a new backup policy, and click **Next**.

Define Policy

Policy - Retention & Schedule Create a New Policy Select an Existing Policy

Select Policy

Default Policy (30 Daily)

DP Volumes Data protection volume backups use the same retention period as defined in the source SnapMirror relationship by default. Use the API if you want to change this value

See [the list of existing policies](#).

7. Select the volumes that you want to back up.

- To back up all volumes, check the box in the title row (Volume Name).
- To back up individual volumes, check the box for each volume (Volume_1).

Select Volumes						
57 Volumes						
<input checked="" type="checkbox"/>	Volume Name	Volume Type	SVM Name	Used Capacity	Allocated Capacity	Backup Status
<input checked="" type="checkbox"/>	Volume_Name_1	RW	SVM_Name_1	0.25 TB	10 TB	<input type="radio"/> Not Active
<input checked="" type="checkbox"/>	Volume_Name_2	RW	SVM_Name_2	0.25 TB	10 TB	<input type="radio"/> Not Active
<input checked="" type="checkbox"/>	Volume_Name_3	RW	SVM_Name_3	0.25 TB	10 TB	<input type="radio"/> Not Active
<input checked="" type="checkbox"/>	Volume_Name_4	DP	SVM_Name_4	0.25 TB	10 TB	<input type="radio"/> Not Active
<input checked="" type="checkbox"/>	Volume_Name_5	RW	SVM_Name_5	0.25 TB	10 TB	<input type="radio"/> Not Active

8. Click **Activate Backup** and Cloud Backup starts taking the initial backups of your volumes.

You are prompted whether you want to run compliance scans on the backed up volumes. Cloud Data Sense scans are free when you run them on the backed up volumes (except for the [cost of the deployed Cloud Data Sense instance](#)).

Activate Compliance on your Backed Up Volumes

You have successfully activated Backup to Cloud on 12 Volumes in your working environment "Name 1".



Data Sense

- › Cloud Compliance offer automated controls for data privacy regulations such as the GDPR, CCPA and more.
- › Driven by powerful artificial intelligence algorithms, Cloud Compliance gets your business application data and cloud environments privacy ready.

[Go to Compliance](#)

[Close](#)

9. Click **Go to Compliance** to activate compliance scans on the volumes. (If you choose **Close** and not to scan these backed up volumes, you can always [enable this functionality](#) later from Cloud Data Sense.)
 - If an instance of Cloud Data Sense is already deployed in your environment, you are directed to the Configuration page to select the volumes you want to scan in each on-premises working environment that has backups. See [how to choose the volumes](#).

The screenshot shows the Cloud Data Sense Configuration page. At the top, there's a navigation bar with tabs: Data Sense, Governance, Compliance, Investigation, Policies, and Configuration (which is underlined). Below the tabs, it says "(2/20) Working Environments". There are several filter buttons: CVO, ANF, S3, DB, ONEDR, and BACKUP (which is highlighted with a red box). To the right of these filters is a 'Clear filters' link. Below the filters, there's a list of working environments. One entry is shown: "Working Environment 1 (back up) Cloud Backup of ONTAP BETA". At the bottom of the list area, there are two buttons: "Activate Compliance for all Backed Up Volumes" (with a red arrow pointing to it) and "or select Volumes".

- If Cloud Data Sense has not been deployed, you are directed to the Compliance page where you can choose to deploy Compliance in the cloud or in your premises. We strongly recommend deploying it in the cloud. Go [here](#) for installation requirements and instructions.

The screenshot shows the Data Sense interface. At the top left is a logo with a padlock and the text "Data Sense". Below it is a link "How does it work?". A section titled "Always-on Privacy & Compliance Controls" contains text about automated controls for data privacy regulations like GDPR, CCPA, and HIPAA. It also states that Data Sense is driven by powerful AI algorithms to keep business application data and cloud environments privacy ready. Two buttons are present: "Deploy Data Sense in the Cloud" and "Deploy Data Sense On-Premises", with the former being highlighted by a red border. To the right is a "Compliance Status" dashboard featuring a circular progress bar, a "Data Distribution" chart, and two tables showing file counts for Personal Files and Sensitive Personal Files across categories like Email Address and Credit Card.

After you have deployed Compliance you can choose the volumes you want to scan as described above.

Result

Cloud Backup backs up your volumes from the on-premises ONTAP system, and optionally, Cloud Data Sense runs compliance scans on the backed up volumes.

What's next?

You can [start and stop backups for volumes or change the backup schedule](#) and you can [restore entire volumes or individual files from a backup file](#).

You can also [view the results of the compliance scans](#) and review other features of Cloud Data Sense that can help you understand data context and identify sensitive data in your organization.



The scan results are not available immediately because Cloud Backup has to finish creating the backups before Cloud Data Sense can start compliance scans.

Backing up on-premises ONTAP data to Azure Blob storage

Complete a few steps to get started backing up data from your on-premises ONTAP systems to Azure Blob storage.

TIP

In most cases you'll use Cloud Manager for all backup and restore operations. However, starting with ONTAP 9.9.1 you can initiate volume backup operations of your on-premises ONTAP clusters using ONTAP System Manager. [See how to use System Manager to back up your volumes to the cloud using Cloud Backup.](#)

A Beta feature released in January 2021 allows you to run compliance scans on the backed up volumes from your on-premises systems. Typically, compliance scans are free up to 1 TB of data, and then a cost for the service is applied for data over 1 TB. When combining Backup and Data Sense for your on-premises volumes, the cost for scans on those on-prem volumes is free. Learn more about how [Cloud Data Sense](#) can get your business applications and cloud environments privacy ready.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.

1

Verify support for your configuration

- You have discovered the on-premises cluster and added it to a working environment in Cloud Manager. See [Discovering ONTAP clusters](#) for details.
 - The cluster is running ONTAP 9.7P5 or later.
 - The cluster has a SnapMirror license — it is included as part of the Premium Bundle or Data Protection Bundle.
 - The cluster must have the required network connections to Blob storage and to the Connector.
- The Connector must have the required network connections to Blob storage and to the cluster, and the required permissions.
- You have a valid Azure subscription for the object storage space where your backups will be located.

2

Enable Cloud Backup on the system

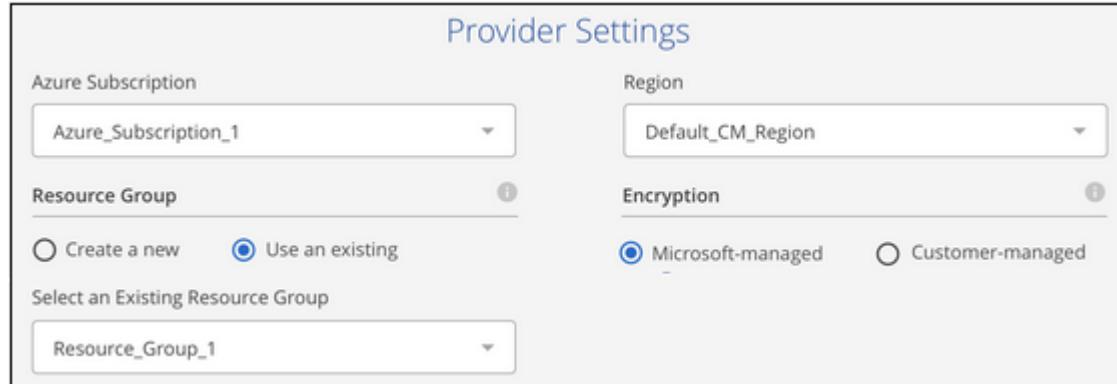
Select the working environment and click **Enable** next to the Backup & Compliance service in the right-panel, and then follow the setup wizard.



3

Select the cloud provider and enter the provider details

Select Microsoft Azure as your provider and then enter the provider details. You'll need to select the Azure Subscription and the region where you want to create the backups. You can also choose your own customer-managed key for data encryption instead of using the default Microsoft-managed encryption key.



Provider Settings	
Azure Subscription	Region
Azure_Subscription_1	Default_CM_Region
Resource Group	Encryption
<input type="radio"/> Create a new <input checked="" type="radio"/> Use an existing	<input checked="" type="radio"/> Microsoft-managed <input type="radio"/> Customer-managed
Select an Existing Resource Group	
Resource_Group_1	

4

Select the cluster IPspace and optional use of a private VNet endpoint

Select the IPspace in the ONTAP cluster where the volumes reside. You can also choose to use an existing Azure Private Endpoint for a more secure connection to the VNet from your on-prem data center.

Networking

IPspace
IP_Space_1

Private Endpoint Configuration

VNet
Select VNet

Subnet
Select Subnet

5

Define the backup policy

The default policy backs up volumes every day and retains the most recent 30 backup copies of each volume. Change to hourly, daily, weekly, or monthly backups, or select one of the system-defined policies that provide more options. You can also change the number of backup copies to retain.

Define Policy

Policy - Retention & Schedule Create a New Policy Select an Existing Policy

Select Policy
Default Policy (30 Daily)

DP Volumes Data protection volume backups use the same retention period as defined in the source SnapMirror relationship by default. Use the API if you want to change this value

6

Select the volumes that you want to back up

Identify which volumes you want to back up from the cluster.

7

Activate Compliance scans on the backed up volumes (optional)

Choose whether you want to have Cloud Data Sense scan the volumes that are backed up in the cloud.

8

Restore your data, as needed

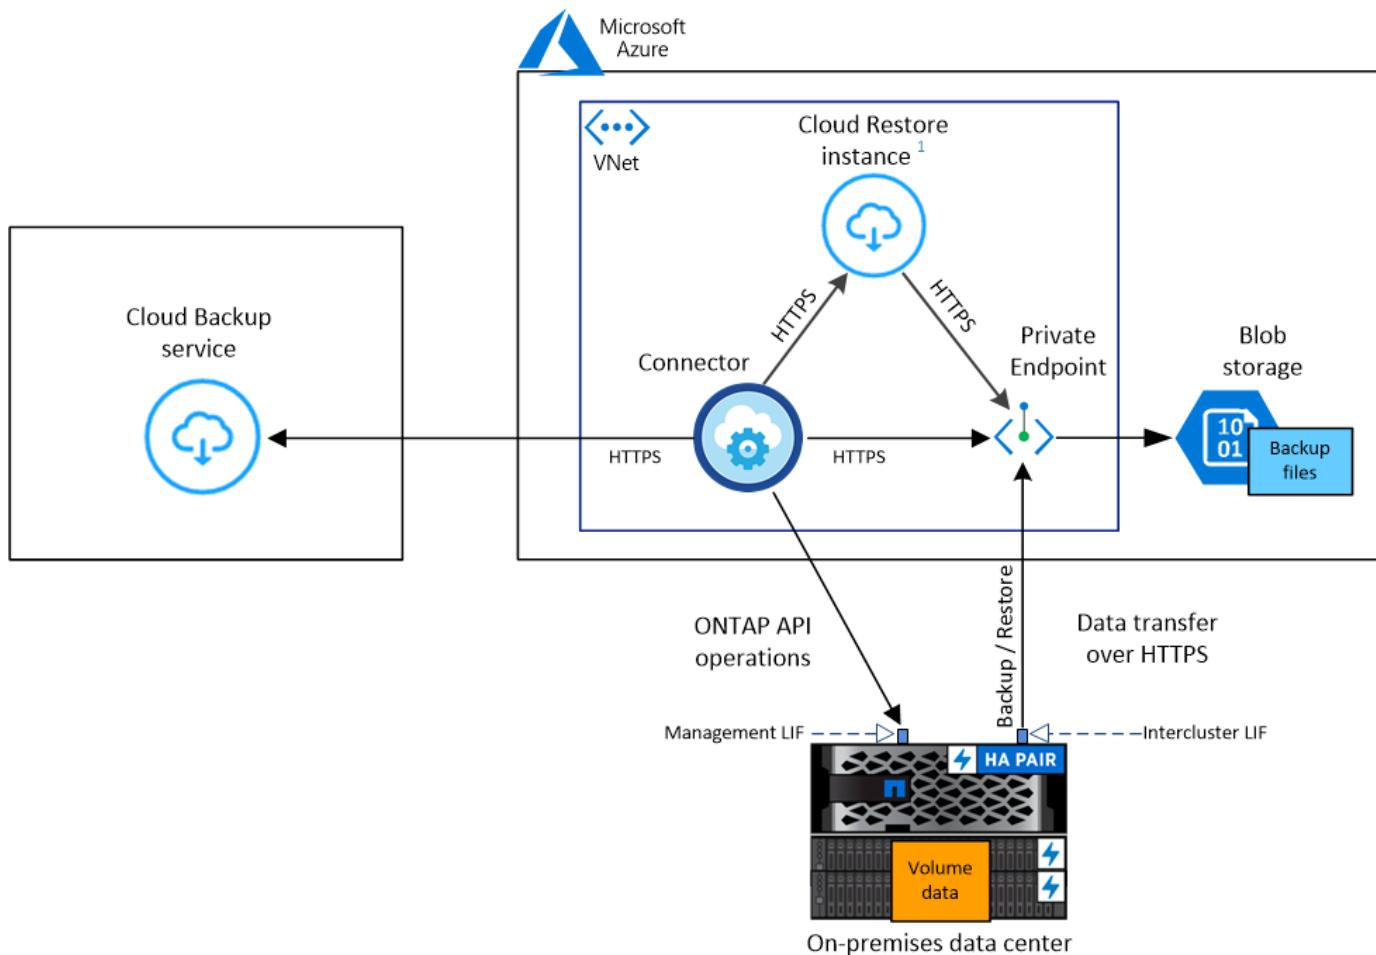
Choose to restore an entire backup to a new volume, or to restore individual files from the backup to an existing volume. You can restore data to a Cloud Volumes ONTAP system in Azure, or to an on-premises ONTAP system.

See [Restoring volume data from backup files](#) for details.

Requirements

Read the following requirements to make sure you have a supported configuration before you start backing up on-premises volumes to Azure Blob storage.

The following image shows each component and the connections that you need to prepare between them:



Note that when the Cloud Restore instance is deployed in the cloud, it is located in the same subnet as the Connector.

Preparing your ONTAP clusters

You need to discover your on-premises ONTAP clusters in Cloud Manager before you can start backing up volume data.

[Learn how to discover a cluster.](#)

ONTAP requirements

- ONTAP 9.7P5 and later.
- A SnapMirror license (included as part of the Premium Bundle or Data Protection Bundle).

Note: The "Hybrid Cloud Bundle" is not required when using the Cloud Backup service.

See how to [manage your cluster licenses](#).

- Time and time zone are set correctly.

See how to [configure your cluster time](#).

Cluster networking requirements

- The ONTAP cluster initiates an HTTPS connection over port 443 from the intercluster LIF to Azure Blob storage for backup and restore operations.

ONTAP reads and writes data to and from object storage. The object storage never initiates, it just responds.

- ONTAP requires an inbound connection from the Connector to the cluster management LIF. The Connector can reside in an Azure VNet.
- An intercluster LIF is required on each ONTAP node that hosts the volumes you want to back up. The LIF must be associated with the *IPspace* that ONTAP should use to connect to object storage. [Learn more about IPspaces](#).

When you set up Cloud Backup, you are prompted for the IPspace to use. You should choose the IPspace that each LIF is associated with. That might be the "Default" IPspace or a custom IPspace that you created.

- The nodes' and intercluster LIFs are able to access the internet.
- DNS servers have been configured for the storage VM where the volumes are located. See how to [configure DNS services for the SVM](#).
- Note that if you use are using a different IPspace than the Default, then you might need to create a static route to get access to the object storage.
- Update firewall rules, if necessary, to allow Cloud Backup service connections from ONTAP to object storage through port 443 and name resolution traffic from the storage VM to the DNS server over port 53 (TCP/UDP).

Creating or switching Connectors

A Connector is required to back up data to the cloud, and the Connector must be in an Azure VNet when backing up data to Azure Blob storage. You can't use a Connector that's deployed on-premises. You'll either need to create a new Connector or make sure that the currently selected Connector resides in the correct provider.

- [Learn about Connectors](#)
- [Creating a Connector in Azure](#)
- [Switching between Connectors](#)

Preparing networking for the Connector

Ensure that the Connector has the required networking connections.

Steps

1. Ensure that the network where the Connector is installed enables the following connections:
 - An outbound internet connection to the Cloud Backup service over port 443 (HTTPS)

- An HTTPS connection over port 443 to your Blob object storage
 - An HTTPS connection over port 443 to your ONTAP clusters
2. Enable a VNet Private Endpoint to Azure storage. This is needed if you have an ExpressRoute or VPN connection from your ONTAP cluster to the VNet and you want communication between the Connector and Blob storage to stay in your virtual private network.

Supported regions

You can create backups from on-premises systems to Azure Blob in all regions [where Cloud Volumes ONTAP is supported](#). You specify the region where the backups will be stored when you set up the service.

License requirements

Before your 30-day free trial of the Cloud Backup service expires, you need to subscribe to a pay-as-you-go (PAYGO) Cloud Manager Marketplace offering from Azure, or purchase and activate a Cloud Backup BYOL license from NetApp. These licenses are for the account and can be used across multiple systems.

- For Cloud Backup PAYGO licensing, you'll need a subscription to the [Azure Cloud Manager Marketplace](#) offering to continue using Cloud Backup. Billing for Cloud Backup is done through this subscription.
- For Cloud Backup BYOL licensing, you don't need a subscription. You need the serial number from NetApp that enables you to use the service for the duration and capacity of the license. [Learn how to manage your BYOL licenses](#).

You need to have an Azure subscription for the object storage space where your backups will be located.

A SnapMirror license is required on the cluster. Note that the "Hybrid Cloud Bundle" is not required when using Cloud Backup.

Preparing Azure Blob storage for backups

1. If your virtual or physical network uses a proxy server for internet access, ensure that the Cloud Restore virtual machine has outbound internet access to contact the following endpoints.

Endpoints	Purpose
http://olcentgbl.trafficmanager.net https://olcentgbl.trafficmanager.net	Provides CentOS packages for the Cloud Restore virtual machine.
http://cloudmanagerinfraprod.azurecr.io https://cloudmanagerinfraprod.azurecr.io	Cloud Restore virtual machine image repository.

2. You can choose your own custom-managed keys for data encryption in the activation wizard instead of using the default Microsoft-managed encryption keys. In this case you will need to have the Azure Subscription, Key Vault name, and the Key. [See how to use your own keys](#).
3. If you want to have a more secure connection over the public internet from your on-prem data center to the VNet, there is an option to configure an Azure Private Endpoint in the activation wizard. In this case you will need to know the VNet and Subnet for this connection. [See details about using a Private Endpoint](#).

Enabling Cloud Backup

Enable Cloud Backup at any time directly from the on-premises working environment.

Steps

- From the Canvas, select the working environment and click **Enable** next to the Backup & Compliance service in the right-panel.



- Select Microsoft Azure as your provider and click **Next**.
- Enter the provider details. Note that you can't change this information after the service has started.
 - The Azure subscription used for backups and the Azure region where the backups will be stored.
 - The resource group that manages the Blob container - you can create a new resource group or select an existing resource group.
 - Whether you will use the default Microsoft-managed encryption key or choose your own customer-managed keys to manage encryption of your data. ([See how to use your own keys](#)).

A screenshot of a "Provider Settings" dialog box. It contains several input fields:

- "Azure Subscription": A dropdown menu showing "Azure_Subscription_1".
- "Region": A dropdown menu showing "Default_CM_Region".
- "Resource Group":
 - A dropdown menu showing "Resource_Group_1".
 - Below it are two radio buttons: "Create a new" (unchecked) and "Use an existing" (checked).
 - A link "Select an Existing Resource Group" is visible.
- "Encryption":
 - A dropdown menu showing "Microsoft-managed".
 - Below it are two radio buttons: "Microsoft-managed" (checked) and "Customer-managed" (unchecked).

- Click **Next** after you've entered the provider details.
- Enter the networking details and click **Next**.
 - The IPspace in the ONTAP cluster where the volumes you want to back up reside. The intercluster LIFs for this IPspace must have outbound internet access.
 - Optionally, choose whether you will configure an Azure Private Endpoint. [See details about using a Private Endpoint](#).

A screenshot of a "Networking" dialog box. It contains the following fields:

- "IPspace": A dropdown menu showing "IP_Space_1".
- "Private Endpoint Configuration": A toggle switch that is currently off (grayed out).
- "VNet": A dropdown menu labeled "Select VNet".
- "Subnet": A dropdown menu labeled "Select Subnet".

- In the *Define Policy* page, select an existing backup schedule and retention value, or define a new backup

policy, and click **Next**.

Define Policy

Policy - Retention & Schedule Create a New Policy Select an Existing Policy

Select Policy

Default Policy (30 Daily)

DP Volumes Data protection volume backups use the same retention period as defined in the source SnapMirror relationship by default. Use the API if you want to change this value

See the list of existing policies.

7. Select the volumes that you want to back up.

- To back up all volumes, check the box in the title row (**Volume Name**).
- To back up individual volumes, check the box for each volume (**Volume_1**).

Select Volumes						
57 Volumes						
<input checked="" type="checkbox"/>	Volume Name	Volume Type	SVM Name	Used Capacity	Allocated Capacity	Backup Status
<input checked="" type="checkbox"/>	Volume_Name_1	RW	SVM_Name_1	0.25 TB	10 TB	<input type="radio"/> Not Active
<input checked="" type="checkbox"/>	Volume_Name_2	RW	SVM_Name_2	0.25 TB	10 TB	<input type="radio"/> Not Active
<input checked="" type="checkbox"/>	Volume_Name_3	RW	SVM_Name_3	0.25 TB	10 TB	<input type="radio"/> Not Active
<input checked="" type="checkbox"/>	Volume_Name_4	DP	SVM_Name_4	0.25 TB	10 TB	<input type="radio"/> Not Active
<input checked="" type="checkbox"/>	Volume_Name_5	RW	SVM_Name_5	0.25 TB	10 TB	<input type="radio"/> Not Active

8. Click **Activate Backup** and Cloud Backup starts taking the initial backups of your volumes.

You are prompted whether you want to run compliance scans on the backed up volumes. Cloud Data Sense scans are free when you run them on the backed up volumes (except for the [cost of the deployed Cloud Data Sense instance](#)).

Activate Compliance on your Backed Up Volumes

You have successfully activated Backup to Cloud on 12 Volumes in your working environment "Name 1".



Data Sense

- > Cloud Compliance offer automated controls for data privacy regulations such as the GDPR, CCPA and more.
- > Driven by powerful artificial intelligence algorithms, Cloud Compliance gets your business application data and cloud environments privacy ready.

[Go to Compliance](#)

[Close](#)

9. Click **Go to Compliance** to activate compliance scans on the volumes. (If you choose **Close** and not to scan these backed up volumes, you can always [enable this functionality](#) later from Cloud Data Sense.)

- If an instance of Cloud Data Sense is already deployed in your environment, you are directed to the Configuration page to select the volumes you want to scan in each on-premises working environment that has backups. See [how to choose the volumes](#).

The screenshot shows the Cloud Data Sense Configuration page. At the top, there's a navigation bar with tabs: Data Sense, Governance, Compliance, Investigation, Policies, and Configuration (which is underlined). Below the navigation bar, there's a section titled "(2/20) Working Environments". Underneath this, there's a filter bar with several buttons: CVO, ANF, S3, DB, ONEDR, and BACKUP (which is highlighted with a red box). To the right of the filter bar is a 'Clear filters' link. Below the filter bar, there's a list of working environments. The first item is "Working Environment 1 (back up)", described as "Cloud Backup of ONTAP" and marked as "BETA". At the bottom of the page, there are two prominent buttons: "Activate Compliance for all Backed Up Volumes" (with a red arrow pointing to it) and "or select Volumes".

- If Cloud Data Sense has not been deployed, you are directed to the Compliance page where you can choose to deploy Compliance in the cloud or in your premises. We strongly recommend deploying it in the cloud. Go [here](#) for installation requirements and instructions.

The screenshot shows the Data Sense interface. At the top left is a lock icon and the text "Data Sense". Below it is a link "How does it work?". A section titled "Always-on Privacy & Compliance Controls" contains text about automated controls for data privacy regulations like GDPR, CCPA, and HIPAA. It also states that Data Sense is driven by powerful AI algorithms to keep business application data and cloud environments privacy ready. Two buttons are present: "Deploy Data Sense in the Cloud" and "Deploy Data Sense On-Premises", with the former being highlighted with a red border. To the right is a "Compliance Status" dashboard featuring a circular progress bar, a "Data Distribution" chart, and several data tables. One table shows "Personal Files" with 28,000 entries, another shows "Sensitive Personal Files" with 7,000 entries, and others show "Email Address", "Credit Card", "Health", and "Ethnicity" data.

After you have deployed Compliance you can choose the volumes you want to scan as described above.

Result

Cloud Backup backs up your volumes from the on-premises ONTAP system, and optionally, Cloud Data Sense runs compliance scans on the backed up volumes.

What's next?

You can [start and stop backups for volumes or change the backup schedule](#) and you can [restore entire volumes or individual files](#) from a backup file.

You can also [view the results of the compliance scans](#) and review other features of Cloud Data Sense that can help you understand data context and identify sensitive data in your organization.



The scan results are not available immediately because Cloud Backup has to finish creating the backups before Cloud Data Sense can start compliance scans.

Backing up on-premises ONTAP data to Google Cloud Storage

Complete a few steps to get started backing up data from your on-premises ONTAP systems to Google Cloud Storage.

TIP

In most cases you'll use Cloud Manager for all backup and restore operations. However, starting with ONTAP 9.9.1 you can initiate volume backup operations of your on-premises ONTAP clusters using ONTAP System Manager. [See how to use System Manager to back up your volumes to the cloud using Cloud Backup.](#)

A Beta feature released in January 2021 allows you to run compliance scans on the backed up volumes from your on-premises systems. Typically, compliance scans are free up to 1 TB of data, and then a cost for the service is applied for data over 1 TB. When combining Backup and Data Sense for your on-premises volumes, the cost for scans on those on-prem volumes is free. Learn more about how [Cloud Data Sense](#) can get your business applications and cloud environments privacy ready.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.

1

Verify support for your configuration

- You have discovered the on-premises cluster and added it to a working environment in Cloud Manager. See [Discovering ONTAP clusters](#) for details.
 - The cluster is running ONTAP 9.7P5 or later.
 - The cluster has a SnapMirror license — it is included as part of the Premium Bundle or Data Protection Bundle.
 - The cluster must have the required network connections to Google storage and to the Connector.
- The Connector must have the required network connections to Google storage and to the cluster.
- You have a valid Google subscription for the object storage space where your backups will be located.
- You have a Google account with an access key and secret key so the ONTAP cluster can back up and restore data.

2

Enable Cloud Backup on the system

Select the working environment and click **Enable** next to the Backup & Compliance service in the right-panel, and then follow the setup wizard.

**3**

Select the cloud provider and enter the provider details

Select Google Cloud as your provider and then enter the provider details. You also need to specify the IPspace in the ONTAP cluster where the volumes reside.

4

Define the backup policy

The default policy backs up volumes every day and retains the most recent 30 backup copies of each volume. Change to hourly, daily, weekly, or monthly backups, or select one of the system-defined policies that provide more options. You can also change the number of backup copies to retain.

Define Policy

Policy - Retention & Schedule

Create a New Policy Select an Existing Policy

Select Policy

Default Policy (30 Daily)

DP Volumes

Data protection volume backups use the same retention period as defined in the source SnapMirror relationship by default. Use the API if you want to change this value

5

Select the volumes that you want to back up

Identify which volumes you want to back up from the cluster.

6

Activate Compliance scans on the backed up volumes (optional)

Choose whether you want to have Cloud Data Sense scan the volumes that are backed up in the cloud.

7

Restore your data, as needed

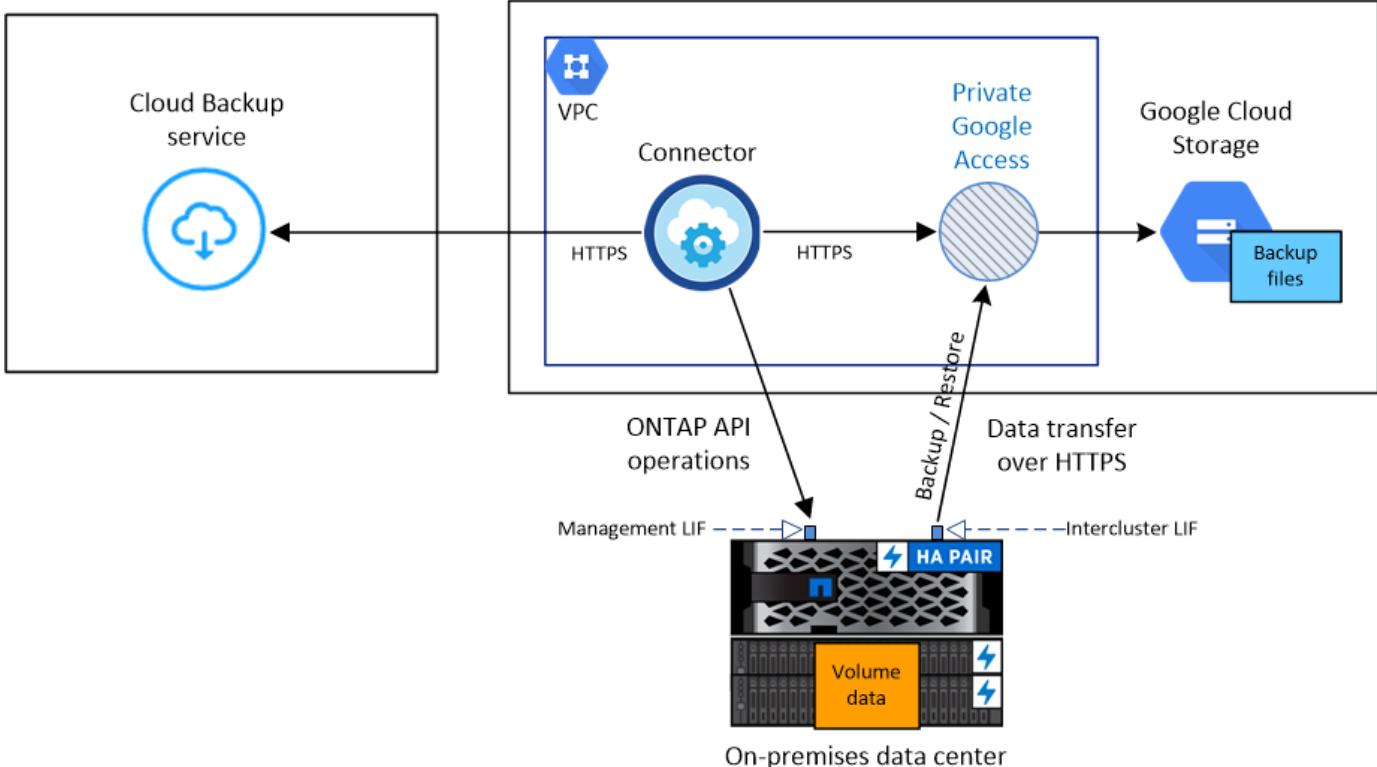
Restore a backup to a new volume. You can restore data to a Cloud Volumes ONTAP system in Google, or to an on-premises ONTAP system.

See [Restoring volume data from backup files](#) for details.

Requirements

Read the following requirements to make sure you have a supported configuration before you start backing up on-premises volumes to Google Cloud storage.

The following image shows each component and the connections that you need to prepare between them:



Note that the Cloud Restore instance is not shown in this diagram because single-file restore is not currently supported in GCP.

Preparing your ONTAP clusters

You need to discover your on-premises ONTAP clusters in Cloud Manager before you can start backing up volume data.

[Learn how to discover a cluster.](#)

ONTAP requirements

- ONTAP 9.7P5 and later.
- A SnapMirror license (included as part of the Premium Bundle or Data Protection Bundle).

Note: The "Hybrid Cloud Bundle" is not required when using the Cloud Backup service.

See how to [manage your cluster licenses](#).

- Time and time zone are set correctly.

See how to [configure your cluster time](#).

Cluster networking requirements

- The ONTAP cluster initiates an HTTPS connection over port 443 from the intercluster LIF to Google Cloud storage for backup and restore operations.

ONTAP reads and writes data to and from object storage. The object storage never initiates, it just responds.

- ONTAP requires an inbound connection from the Connector to the cluster management LIF. The Connector can reside in a Google Cloud Platform VPC.
- An intercluster LIF is required on each ONTAP node that hosts the volumes you want to back up. The LIF must be associated with the IPspace that ONTAP should use to connect to object storage. [Learn more about IPspaces](#).

When you set up Cloud Backup, you are prompted for the IPspace to use. You should choose the IPspace that each LIF is associated with. That might be the "Default" IPspace or a custom IPspace that you created.

- The nodes' intercluster LIFs are able to access the internet.
- DNS servers have been configured for the storage VM where the volumes are located. See how to [configure DNS services for the SVM](#).
- Note that if you use are using a different IPspace than the Default, then you might need to create a static route to get access to the object storage.
- Update firewall rules, if necessary, to allow Cloud Backup service connections from ONTAP to object storage through port 443 and name resolution traffic from the storage VM to the DNS server over port 53 (TCP/UDP).

Creating or switching Connectors

A Connector is required to back up data to the cloud, and the Connector must be in a Google Cloud Platform VPC when backing up data to Google Cloud storage. You can't use a Connector that's deployed on-premises. You'll either need to create a new Connector or make sure that the currently selected Connector resides in the correct provider.

- [Learn about Connectors](#)
- [Creating a Connector in GCP](#)
- [Switching between Connectors](#)

Preparing networking for the Connector

Ensure that the Connector has the required networking connections.

Steps

1. Ensure that the network where the Connector is installed enables the following connections:
 - An outbound internet connection to the Cloud Backup service over port 443 (HTTPS)
 - An HTTPS connection over port 443 to your Google Cloud storage
 - An HTTPS connection over port 443 to your ONTAP clusters
2. Enable Private Google Access on the subnet where you plan to deploy the Connector. [Private Google Access](#) is needed if you have a direct connection from your ONTAP cluster to the VPC and you want communication between the Connector and Google Cloud Storage to stay in your virtual private network.

Note that Private Google Access works with VM instances that have only internal (private) IP addresses (no external IP addresses).

Supported regions

You can create backups from on-premises systems to Google Cloud storage in all regions [where Cloud Volumes ONTAP is supported](#). You specify the region where the backups will be stored when you set up the

service.

License requirements

Before your 30-day free trial of the Cloud Backup service expires, you need to subscribe to a pay-as-you-go (PAYGO) Cloud Manager Marketplace offering from Google, or purchase and activate a Cloud Backup BYOL license from NetApp. These licenses are for the account and can be used across multiple systems.

- For Cloud Backup PAYGO licensing, you'll need a subscription to the [Google Cloud Manager Marketplace](#) offering to continue using Cloud Backup. Billing for Cloud Backup is done through this subscription.
- For Cloud Backup BYOL licensing, you don't need a subscription. You need the serial number from NetApp that enables you to use the service for the duration and capacity of the license. [Learn how to manage your BYOL licenses.](#)

You need to have a Google subscription for the object storage space where your backups will be located.

A SnapMirror license is required on the cluster. Note that the "Hybrid Cloud Bundle" is not required when using Cloud Backup.

Preparing Google Cloud Storage for backups

When you set up backup, you need to provide storage access keys for a service account that has Storage Admin permissions. A service account enables Cloud Backup to authenticate and access Cloud Storage buckets used to store backups. The keys are required so that Google Cloud Storage knows who is making the request.

Steps

1. [Create a service account that has the predefined Storage Admin role.](#)
2. Go to [GCP Storage Settings](#) and create access keys for the service account:
 - a. Select a project, and click **Interoperability**. If you haven't already done so, click **Enable interoperability access**.
 - b. Under **Access keys for service accounts**, click **Create a key for a service account**, select the service account that you just created, and click **Create Key**.

You'll need to enter the keys in Cloud Backup later when you configure the backup service.

Enabling Cloud Backup

Enable Cloud Backup at any time directly from the on-premises working environment.

Steps

1. From the Canvas, select the working environment and click **Enable** next to the Backup & Compliance service in the right-panel.



2. Select Google Cloud as your provider and click **Next**.

3. Enter the provider details. Note that you can't change this information after the service has started.
 - a. The Google Cloud Project where you want the Google Cloud Storage bucket to be created for backups.
(The Project must have a Service Account that has the predefined Storage Admin role.)
 - b. The Google Access Key and Secret Key used to store the backups.
 - c. The Google region where the backups will be stored.
 - d. The IPspace in the ONTAP cluster where the volumes you want to back up reside. The intercluster LIFs for this IPspace must have outbound internet access.

Provider Settings

Provider Information	Location & Connectivity
Google Cloud Project <input type="text" value="Cloud Manager Default Project"/>	Region <input type="text" value="Cloud Manager Default Region"/>
Google Cloud Access Key <input type="text" value="Enter Google Cloud Access Key"/>	IPspace <input type="text" value="IP_Space_1"/>
Google Cloud Secret Key <input type="text" value="Enter Google Cloud Secret Key"/>	

4. Click **Next** after you've entered the provider details.
5. In the *Define Policy* page, select an existing backup schedule and retention value, or define a new backup policy, and click **Next**.

Define Policy

Policy - Retention & Schedule	<input type="radio"/> Create a New Policy <input checked="" type="radio"/> Select an Existing Policy
Select Policy <input type="text" value="Default Policy (30 Daily)"/>	
DP Volumes	Data protection volume backups use the same retention period as defined in the source SnapMirror relationship by default. Use the API if you want to change this value

See [the list of existing policies](#).

6. Select the volumes that you want to back up.
 - To back up all volumes, check the box in the title row (**Volume Name**).
 - To back up individual volumes, check the box for each volume (**Volume_1**).

Select Volumes						
57 Volumes						
	Volume Name	Volume Type	SVM Name	Used Capacity	Allocated Capacity	Backup Status
<input checked="" type="checkbox"/>	Volume_Name_1	RW	SVM_Name_1	0.25 TB	10 TB	Not Active
<input checked="" type="checkbox"/>	Volume_Name_2	RW	SVM_Name_2	0.25 TB	10 TB	Not Active
<input checked="" type="checkbox"/>	Volume_Name_3	RW	SVM_Name_3	0.25 TB	10 TB	Not Active
<input checked="" type="checkbox"/>	Volume_Name_4	DP	SVM_Name_4	0.25 TB	10 TB	Not Active
<input checked="" type="checkbox"/>	Volume_Name_5	RW	SVM_Name_5	0.25 TB	10 TB	Not Active

7. Click **Activate Backup** and Cloud Backup starts taking the initial backups of your volumes.

Result

Cloud Backup starts taking the initial backups of each selected volume and the Backup Dashboard is displayed so you can monitor the state of the backups.

What's next?

You can [start and stop backups for volumes or change the backup schedule](#) and you can [restore entire volumes from a backup file](#).

Backing up on-premises ONTAP data to StorageGRID

Complete a few steps to get started backing up data from your on-premises ONTAP systems to object storage in your NetApp StorageGRID systems.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.



Verify support for your configuration

- You have discovered the on-premises cluster and added it to a working environment in Cloud Manager. See [Discovering ONTAP clusters](#) for details.
 - The cluster is running ONTAP 9.7P5 or later.
 - The cluster has a SnapMirror license — it is included as part of the Premium Bundle or Data Protection Bundle.
 - The cluster must have the required network connections to StorageGRID and to the Connector.
- You have a Connector installed on your premises.
 - Networking for the Connector enables an outbound HTTPS connection to the ONTAP cluster and to StorageGRID.
- You have purchased [and activated](#) a Cloud Backup BYOL license from NetApp.
- Your StorageGRID has version 10.3 or later with access keys that have S3 permissions.

2

Enable Cloud Backup on the system

Select the working environment and click **Enable** next to the Backup & Compliance service in the right-panel, and then follow the setup wizard.

**3**

Enter the StorageGRID details

Select StorageGRID as the provider, and then enter the StorageGRID details. You also need to specify the IPspace in the ONTAP cluster where the volumes reside.

A screenshot of a 'Provider Settings' dialog box. The title 'Provider Settings' is at the top center. On the left, under 'Provider Information', there are three input fields: 'Storage Server' (placeholder 'Enter Storage Server'), 'Access Key' (placeholder 'Access Key'), and 'Secret Key' (placeholder 'Secret Key'). On the right, under 'Connectivity', there is a section labeled 'IPspace' with a dropdown menu showing 'IP_Space_1'.**4**

Define the backup policy

The default policy backs up volumes every day and retains the most recent 30 backup copies of each volume. Change to hourly, daily, weekly, or monthly backups, or select one of the system-defined policies that provide more options.

Define Policy

Policy - Retention & Schedule

Create a New Policy Select an Existing Policy

Select Policy

Default Policy (30 Daily) ▾

DP Volumes

Data protection volume backups use the same retention period as defined in the source SnapMirror relationship by default. Use the API if you want to change this value

5

Select the volumes that you want to back up

Identify which volumes you want to back up from the cluster.

6

Restore your data, as needed

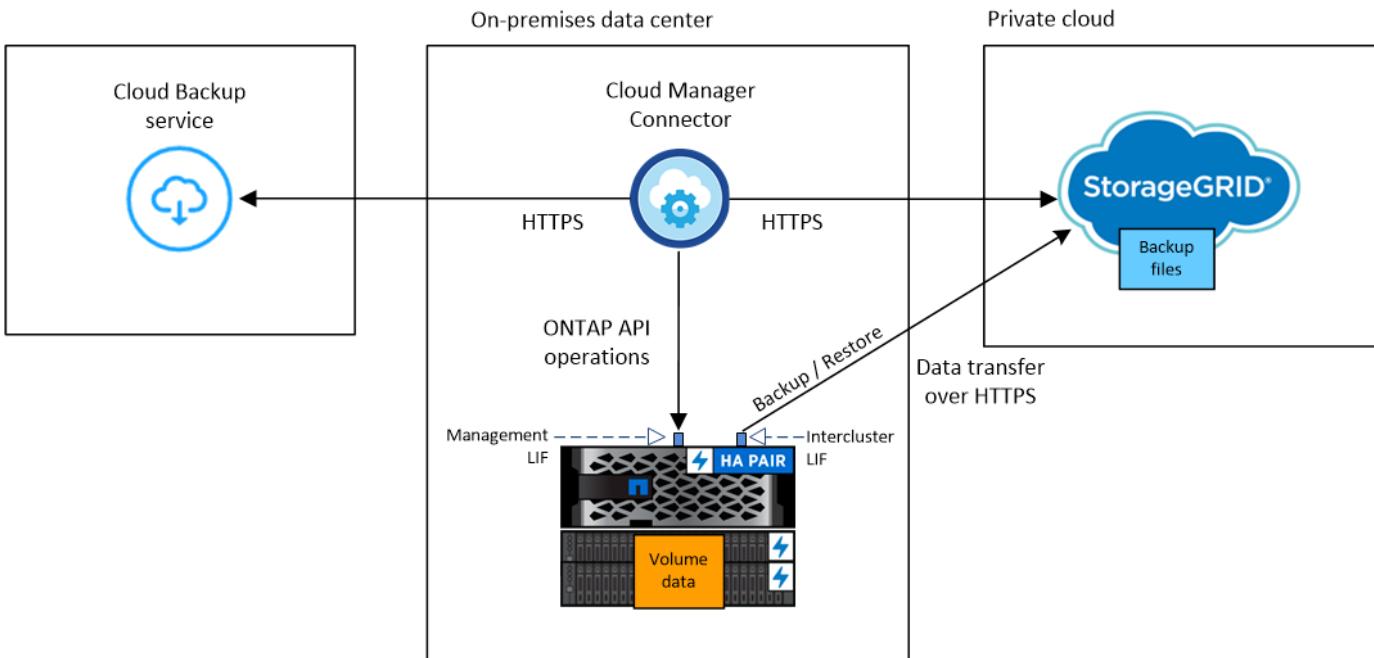
If necessary, choose the backup file to restore an entire backup to a new volume on an on-premises ONTAP system.

See [Restoring volume data from backup files](#) for details.

Requirements

Read the following requirements to make sure you have a supported configuration before you start backing up on-premises volumes to StorageGRID.

The following image shows each component when backing up an on-prem ONTAP system to StorageGRID and the connections that you need to prepare between them:



Note that the Cloud Restore instance is not shown in this diagram because single-file restore is not currently supported when using StorageGRID.

Preparing your ONTAP clusters

You need to discover your on-premises ONTAP clusters in Cloud Manager before you can start backing up volume data.

[Learn how to discover a cluster.](#)

ONTAP requirements

- ONTAP 9.7P5 and later.
- A SnapMirror license (included as part of the Premium Bundle or Data Protection Bundle).

Note: The "Hybrid Cloud Bundle" is not required when using the Cloud Backup service.

[See how to manage your cluster licenses.](#)

- Time and time zone are set correctly.

[See how to configure your cluster time.](#)

Cluster networking requirements

- The ONTAP cluster initiates an HTTPS connection over a user-specified port from the intercluster LIF to StorageGRID for backup and restore operations. The port is configurable during backup setup.

ONTAP reads and writes data to and from object storage. The object storage never initiates, it just responds.

- ONTAP requires an inbound connection from the Connector to the cluster management LIF. The Connector must reside on your premises.
- An intercluster LIF is required on each ONTAP node that hosts the volumes you want to back up. The LIF must be associated with the *IPspace* that ONTAP should use to connect to object storage. [Learn](#)

[more about IPspaces.](#)

When you set up Cloud Backup, you are prompted for the IPspace to use. You should choose the IPspace that each LIF is associated with. That might be the "Default" IPspace or a custom IPspace that you created.

- The nodes' intercluster LIFs are able to access the internet.
- DNS servers have been configured for the storage VM where the volumes are located. See how to [configure DNS services for the SVM](#).
- Note that if you use are using a different IPspace than the Default, then you might need to create a static route to get access to the object storage.
- Update firewall rules, if necessary, to allow Cloud Backup service connections from ONTAP to object storage through the port you specified (typically port 443) and name resolution traffic from the storage VM to the DNS server over port 53 (TCP/UDP).

Preparing StorageGRID

StorageGRID must meet the following requirements. See the [StorageGRID documentation](#) for more information.

Supported StorageGRID versions

StorageGRID 10.3 and later is supported.

S3 credentials

When you set up backup to StorageGRID, the backup wizard prompts you for an S3 access key and secret key for a service account. A service account enables Cloud Backup to authenticate and access the StorageGRID buckets used to store backups. The keys are required so that StorageGRID knows who is making the request.

These access keys must be associated with a user who has the following permissions:

```
"s3>ListAllMyBuckets",
"s3>ListBucket",
"s3GetObject",
"s3PutObject",
"s3DeleteObject",
"s3CreateBucket"
```

Object versioning

You must not enable StorageGRID object versioning on the object store bucket.

Creating or switching Connectors

When backing up data to StorageGRID, a Connector must be available on your premises. You'll either need to install a new Connector or make sure that the currently selected Connector resides on-prem.

- [Learn about Connectors](#)
- [Connector host requirements](#)
- [Installing the Connector on an existing Linux host](#)

- [Switching between Connectors](#)

Preparing networking for the Connector

Ensure that the Connector has the required networking connections.

Steps

1. Ensure that the network where the Connector is installed enables the following connections:
 - An outbound internet connection to the Cloud Backup service over port 443 (HTTPS)
 - An HTTPS connection over port 443 to StorageGRID
 - An HTTPS connection over port 443 to your ONTAP clusters

License requirements

Before your 30-day free trial of the Cloud Backup service expires, you need to purchase and activate a Cloud Backup BYOL license from NetApp. This license is for the account and can be used across multiple systems.

You'll need the serial number from NetApp that enables you to use the service for the duration and capacity of the license. [Learn how to manage your BYOL licenses](#).

TIP

PAYGO licensing is not currently supported when backing up files to StorageGRID.

A SnapMirror license is required on the cluster. Note that the "Hybrid Cloud Bundle" is not required when using Cloud Backup.

Enabling Cloud Backup to StorageGRID

Enable Cloud Backup at any time directly from the on-premises working environment.

Steps

1. From the Canvas, select the on-premises working environment and click **Enable** next to the Backup & Compliance service in the right-panel.



2. Select **StorageGRID** as the provider, click **Next**, and then enter the provider details:
 - a. The FQDN of the StorageGRID server and the port that ONTAP should use for HTTPS communication with StorageGRID; for example: s3.eng.company.com:8082
 - b. The Access Key and the Secret Key used to access the bucket to store backups.
 - c. The IPspace in the ONTAP cluster where the volumes you want to back up reside. The intercluster LIFs for this IPspace must have outbound internet access.

Selecting the correct IPspace ensures that Cloud Backup can set up a connection from ONTAP to your StorageGRID object storage.

Provider Settings

<p>Provider Information</p> <p>Storage Server Enter Storage Server</p> <p>Access Key Access Key</p> <p>Secret Key Secret Key</p>	<p>Connectivity</p> <p>IPspace IP_Space_1</p>
--	---

Note that you cannot change this information after the service has started.

3. In the *Define Policy* page, select the backup schedule and retention value and click **Next**.

Define Policy

<p>Policy - Retention & Schedule</p> <p><input type="radio"/> Create a New Policy <input checked="" type="radio"/> Select an Existing Policy</p> <p>Select Policy Default Policy (30 Daily)</p>	<p>DP Volumes Data protection volume backups use the same retention period as defined in the source SnapMirror relationship by default. Use the API if you want to change this value</p>
---	---

See [the list of existing policies](#).

4. Select the volumes that you want to back up.

- To back up all volumes, check the box in the title row (**Volume Name**).
- To back up individual volumes, check the box for each volume (Volume_1).

Select Volumes						
57 Volumes						
	Volume Name	Volume Type	SVM Name	Used Capacity	Allocated Capacity	Backup Status
<input checked="" type="checkbox"/>	Volume_Name_1	RW	SVM_Name_1	0.25 TB	10 TB	Not Active
<input checked="" type="checkbox"/>	Volume_Name_2	RW	SVM_Name_2	0.25 TB	10 TB	Not Active
<input checked="" type="checkbox"/>	Volume_Name_3	RW	SVM_Name_3	0.25 TB	10 TB	Not Active
<input checked="" type="checkbox"/>	Volume_Name_4	DP	SVM_Name_4	0.25 TB	10 TB	Not Active
<input checked="" type="checkbox"/>	Volume_Name_5	RW	SVM_Name_5	0.25 TB	10 TB	Not Active

5. Click **Activate Backup** and Cloud Backup starts taking the initial backups of each selected volume and the Backup Dashboard is displayed so you can monitor the state of the backups.

Result

Cloud Backup backs up your volumes from the on-premises ONTAP system.

What's next?

You can [start and stop backups for volumes or change the backup schedule](#) and you can [restore entire volumes from a backup file](#).

Set up licensing for Cloud Backup

As mentioned in the [Licensing overview](#), you can pay for Cloud Backup using a pay-as-you-go (PAYGO) subscription through your cloud provider, an annual contract through AWS, or a bring-your-own license (BYOL) from NetApp. If you want to pay as you go or use an annual contract, then you need to subscribe from the marketplace for the cloud provider to which you want to back up data. There's no need to subscribe from every marketplace.

Set up a PAYGO subscription

For pay-as-you-go you'll need to pay your cloud provider for object storage costs and NetApp for backup licensing costs. The licensing costs are based on target backup capacity (*before* ONTAP storage efficiencies). Use these links to subscribe to Cloud Backup from your cloud provider marketplace:

- AWS: [Go to the Cloud Manager Marketplace offering for pricing details](#).
- Azure: [Go to the Cloud Manager Marketplace offering for pricing details](#).
- GCP: [Go to the Cloud Manager Marketplace offering for pricing details](#).

Subscribe to yearly contracts through AWS

There are two annual contracts available from the AWS Marketplace:

- An annual contract that enables you to back up Cloud Volumes ONTAP data and on-premises ONTAP data.

Go to the [AWS Marketplace page](#) to view pricing details.

If you want to use this option, set up your subscription from the Marketplace page and then [associate the subscription with your AWS credentials](#). Note that you'll also need to pay for Cloud Volumes ONTAP using this annual contract subscription since you can assign only one active subscription to your AWS credentials in Cloud Manager.

- A Professional Package that enables you to bundle Cloud Volumes ONTAP and Cloud Backup service by using an annual contract for 1, 2, or 3 years. Payment is per TiB. This option doesn't enable you to back up on-premises ONTAP data.

Go to the [AWS Marketplace page](#) to view pricing details and go to the [Cloud Volumes ONTAP Release Notes](#) to learn more about this licensing option.

If you want to use this option, you can set up the annual contract when you create a Cloud Volumes ONTAP working environment and Cloud Manager prompts you to subscribe to the AWS Marketplace.

Use a Cloud Backup BYOL license

Bring-your-own licenses from NetApp provide 1-, 2-, or 3-year terms. You use the Digital Wallet page in Cloud Manager to manage BYOL licenses for the Cloud Backup service. You can add new licenses and update existing licenses.

Obtain your Cloud Backup license file

After you have purchased your Cloud Backup license, you activate the license in Cloud Manager by entering the Cloud Backup serial number and NSS account, or by uploading the NLF license file. The steps below show how to get the NLF license file if you plan to use that method.

Steps

1. Sign in to the [NetApp Support Site](#) and click **Systems > Software Licenses**.
2. Enter your Cloud Backup license serial number.

The screenshot shows the "Software Licenses" page from the NetApp Support Site. At the top, there is a search bar labeled "Serial Number" containing "481*". Below the search bar is a table header with columns: Serial #, Cluster SN, License Name, License Key, Host ID, Value, and End Date. In the table, there is one row with data: "4810", "CLOUD_BKP_SERVICE", and a red box highlights the "Get NetApp License File" button. The rest of the table columns are empty.

3. Under **License Key**, click **Get NetApp License File**.
4. Enter your Cloud Manager Account ID (this is called a Tenant ID on the support site) and click **Submit** to download the license file.

Get License

SERIAL NUMBER:	4810 [REDACTED]
LICENSE:	CLOUD_BKP_SERVICE
SALES ORDER:	3005 [REDACTED]
TENANT ID:	<input type="text" value="Enter Tenant ID"/>

Example: account-xxxxxxxx

[Cancel](#) Submit

You can find your Cloud Manager Account ID by selecting the **Account** drop-down from the top of Cloud Manager, and then clicking **Manage Account** next to your account. Your Account ID is in the Overview tab.

Add Cloud Backup BYOL licenses to your account

After you purchase a Cloud Backup license for your NetApp account, you need to add the license to Cloud Manager to use the Cloud Backup service.

Steps

1. Click **All Services > Digital Wallet > Cloud Backup Licenses**.
2. Click **Add Backup License**.
3. In the *Add Cloud Backup License* dialog, enter the license information and click **Add Backup License**:
 - If you have the backup license serial number and know your NSS account, select the **Enter Serial Number** option and enter that information.

If your NetApp Support Site account isn't available from the drop-down list, [add the NSS account to Cloud Manager](#).

- If you have the backup license file, select the **Upload License File** option and follow the prompts to attach the file.

Add Cloud Backup License

A Backup License must be installed with an active subscription. A Backup license enables you to use Cloud Backup for a certain period of time and for a maximum amount of backup space.

<input checked="" type="radio"/> Enter Serial Number	<input type="radio"/> Upload License File
Serial Number	
<input type="text" value="Enter Serial Number"/>	
NetApp Support Site Account	
<input type="text" value="Select Support Site Account"/>	
Add Backup License Cancel	

Enter Serial Number Upload License File

To install a license, follow these instructions:

- 1 Obtain the license file from the "System > Software Licenses" tab at [NetApp Support Site](#). You will need to provide your cloud service serial number and Cloud Manager Account ID.
- 2 Click Upload File and then select the file.

Upload License File

Upload

Add Backup License [Cancel](#)

Result

Cloud Manager adds the license so that your Cloud Backup service is active.

Update a Cloud Backup BYOL license

If your licensed term is nearing the expiration date, or if your licensed capacity is reaching the limit, you'll be notified. This status also appears in the Digital Wallet page.

You can update your Cloud Backup license before it expires so that there is no interruption in your ability to back up and restore your data.

Steps

1. Click the chat icon in the lower-right of Cloud Manager to request an extension or capacity add-on to your Cloud Backup license for the particular serial number.

After you pay for the license and it is registered with the NetApp Support Site, in most cases, Cloud Manager can automatically obtain your updated license file and the Cloud Backup Licenses page will reflect the change in 5 to 10 minutes.

2. If Cloud Manager can't automatically update the license, then you'll need to manually upload the license file.
 - a. You can [obtain the license file from the NetApp Support Site](#).
 - b. On the *Cloud Backup Licenses* page, click **Update Backup License**.
 - c. In the *Update Cloud Backup License* dialog, enter the license information and click **Update Backup License**.

Result

Cloud Manager updates the license so that your Cloud Backup service continues to be active.

Managing backups for Cloud Volumes ONTAP and on-premises ONTAP systems

You can manage backups for Cloud Volumes ONTAP and on-premises ONTAP systems by changing the backup schedule, enabling/disabling volume backups, creating an on-demand backup, deleting backups, and more.



Do not manage or change backup files directly from your cloud provider environment. This may corrupt the files and will result in an unsupported configuration.

Viewing the volumes that are being backed up

You can view a list of all the volumes that are currently being backed up in the Backup Dashboard.

Steps

1. Click the **Backup & Restore** tab.
2. Click the **Backup** tab and the Backup Dashboard is displayed.

All Backup Working Environments

Backup Settings

57 Backups

Source Working Environment	Source Volume	Source SVM	Last Backup	Backups	Backup Status
CVO_AWS	Source Volume Name	Source SVM Name	May 22 2019, 00:00:00	2,050 Backups	Active
CVO_AWS	Source Volume Name	Source SVM Name	May 22 2019, 00:00:00	2,050 Backups	Active
CVO_AWS	Source Volume Name	Source SVM Name	May 22 2019, 00:00:00	2,050 Backups	Active

If you are looking for specific volumes in certain working environments, you can refine the list by working environment and volume, or you can use the search filter.

Changing the schedule and backup retention

The default policy backs up volumes every day and retains the most recent 30 backup copies of each volume. You can change to a combination of hourly, daily, weekly, or monthly backups and you can change the number of backup copies to retain. You can also select one of the system-defined policies that provide scheduled backups for 3 months, 1 year, and 7 years.

Changing the backup policy affects both new volumes created after you change the schedule, and any existing volumes that were using the original policy.

Steps

1. From the Backup Dashboard, select **Backup Settings**.

Backup Settings

All Backup Working Environments

Protected Volumes Status

5 Working Environments

57 Protected Volumes

15.1 TB Total Backup Capacity

57 Healthy Backup Volumes

0 Failed Backup Volumes

2. From the *Backup Settings* page, click **...** for the working environment where you want to change the settings and select **Modify Policy**.

CVO_AWS

Cloud Volumes ONTAP | On

Active

All (57) Protected Volumes

30 Daily | 13 Weekly | 3 Monthly

Modify Policy

Manage Volumes

Delete All Backups

Deactivate Backup

Unregister

- From the *Modify Policy* page, change the schedule and backup retention and then click **Save**.

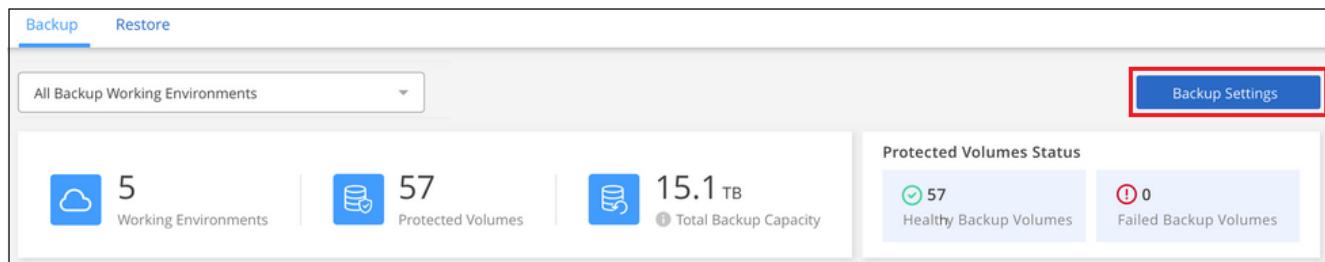
The screenshot shows the 'Modify Policy' configuration interface. Under 'Backup Policy', the 'Create a New Policy' radio button is selected. It includes options for Hourly, Daily, Weekly, and Monthly backups, each with a dropdown to set the number of backups to retain. The 'Daily' option is checked, with 'Number of backups to retain' set to 30. Below this, there's a note about DP Volumes and an S3 Bucket creation note.

Starting and stopping backups of volumes

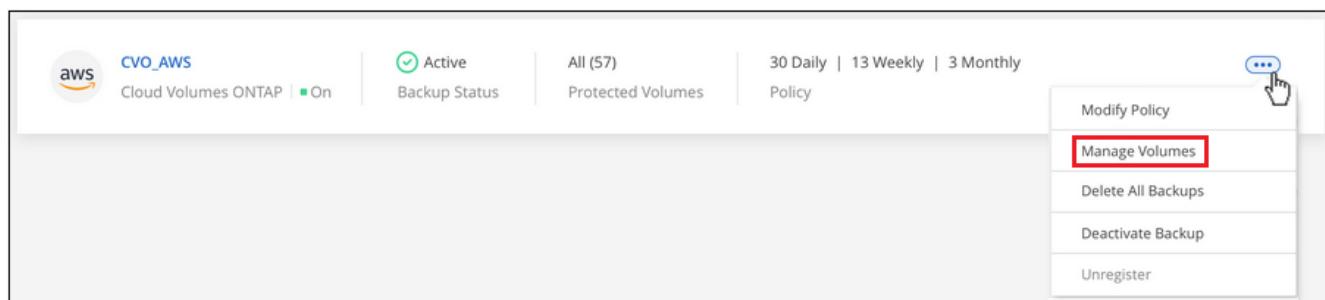
You can stop backing up a volume if you do not need backup copies of that volume and you do not want to pay for the cost to store the backups. You can also add a new volume to the backup list if it is not currently being backed up.

Steps

- From the Backup Dashboard, select **Backup Settings**.



- From the *Backup Settings* page, click **...** for the working environment and select **Manage Volumes**.



- Select the checkbox for volumes that you want to start backing up, and deselect the checkbox for volumes that you want to stop backing up.

Manage Volumes						
57 Volumes 25 Selected Volumes						
	Volume Name	Volume Type	SVM Name	Used Capacity	Allocated Capacity	Backup Status
<input type="checkbox"/>	Volume_1	RW	SVM_1	2.25 TB	10 TB	Active
<input type="checkbox"/>	Volume_2	RW	SVM_1	2.25 TB	10 TB	Not Active
<input checked="" type="checkbox"/>	Volume_3	RW	SVM_1	2.25 TB	10 TB	Not Active
<input type="checkbox"/>	Volume_4	DP i	SVM_2	2.25 TB	10 TB	Active

- Click **Save** to commit your changes.

Note: When stopping a volume from being backed up you'll continue to be charged by your cloud provider for object storage costs for the capacity that the backups use unless you [delete the backups](#).

Creating a manual volume backup at any time

You can create an on-demand backup at any time to capture the current state of the volume. This can be useful if very important changes have been made to a volume and you don't want to wait for the next scheduled backup to protect that data, or if the volume is not currently being backed up and you want to capture its current state.

The backup name includes the timestamp so you can identify your on-demand backup from other scheduled backups.

Steps

- From the Backup Dashboard, click **...** for the volume and select **Backup Now**.

Backup		Restore		Protected Volumes Status			
				57 Working Environments		Backup Settings	
Cloud 1 Working Environments		Cloud 57 Protected Volumes		Cloud 15.1 TB Total Backup Capacity		Protected Volumes Status	
				Green 57 Healthy Backup Volumes		Red 0 Failed Backup Volumes	
57 Backups						Details & Backup List	
Source Working Environment		Source Volume		Source SVM		Last Backup	Backups
AWS CVO_AWS On		Volume_1 On		SVM_1		May 22 2019, 00:00:00	2,050 Backups
AWS CVO_AWS On		Volume_2 On		SVM_1		May 22 2019, 00:00:00	2,050 Backups
AWS CVO_AWS On		Volume_3 On		SVM_1		May 22 2019, 00:00:00	2,050 Backups

The Backup Status column for that volume displays "In Progress" until the backup is created.

Viewing the list of backups for each volume

You can view the list of all backup files that exist for each volume. This page displays details about the source volume, destination location, and backup details such as last backup taken, the current backup policy, backup file size, and more.

This page also enables you to perform the following tasks:

- Delete all backup files for the volume
- Delete individual backup files for the volume
- Download a backup report for the volume

Steps

1. From the Backup Dashboard, click **...** for the source volume and select **Details & Backup List**.

The screenshot shows the Backup Dashboard interface. At the top, there are tabs for 'Backup' (selected) and 'Restore'. Below the tabs, there's a dropdown menu set to 'All Backup Working Environments'. On the right, there's a 'Backup Settings' button. In the center, there are summary statistics: 1 Working Environment, 57 Protected Volumes, and 15.1 TB Total Backup Capacity. To the right, there's a 'Protected Volumes Status' section showing 57 Healthy Backup Volumes and 0 Failed Backup Volumes. Below these stats, a table titled '57 Backups' lists three entries. Each entry includes columns for 'Source Working Environment' (CVO_AWS), 'Source Volume' (Volume_1, Volume_2, Volume_3), 'Source SVM' (SVM_1), 'Last Backup' (May 22 2019, 00:00:00), 'Backups' (2,050 Backups), and 'Backup Status' (Active). The third row has a red box around the 'Details & Backup List' button in the 'Backup Status' column.

Source Working Environment	Source Volume	Source SVM	Last Backup	Backups	Backup Status
CVO_AWS On	Volume_1 On	SVM_1	May 22 2019, 00:00:00	2,050 Backups	Active Details & Backup List
CVO_AWS On	Volume_2 On	SVM_1	May 22 2019, 00:00:00	2,050 Backups	Backup Now
CVO_AWS On	Volume_3 On	SVM_1	May 22 2019, 00:00:00	2,050 Backups	Pause Backups

The list of all backup files is displayed along with details about the source volume, destination location, and backup details.

The screenshot shows the NetApp Cloud Backup interface. At the top, there are three main sections: 'Source' (with details like Working Environment, Type, Provider, Volume, and SVM), 'Destination' (with details like Cloud Provider, Bucket, Region, and Account ID), and 'Backup Information' (with Relationship Status, Last Backup, Lag Duration, Backups, and Backup Policy). Below these is a summary section for '2,050 Backups' with a search bar and filter options ('Select Timeframe', 'Actions'). A table lists three backup entries: 'Backup_2020_Jan' (May 22 2019, 00:00:00), 'Backup_2020_Mar' (May 22 2019, 00:00:00), and 'Backup_2020_Apr' (May 22 2019, 00:00:00), each with a '...' button.

Deleting backups

Cloud Backup enables you to delete a single backup file, delete all backups for a volume, or delete all backups of all volumes in a working environment. You might want to delete all backups if you no longer need the backups or if you deleted the source volume and want to remove all backups.

Note that deleting all backups does not disable further backups of this volume or the working environment. If you want to stop creating backups of a volume, you can disable backups [as described here](#). If you want to stop creating backups of all volumes in a working environment, you can deactivate backups [as described here](#).

! If you plan to delete a Cloud Volumes ONTAP or on-premises ONTAP system that has backups, you must delete the backups **before** deleting the system. Cloud Backup doesn't automatically delete backups when you delete a system, and there is no current support in the UI to delete the backups after the system has been deleted. You'll continue to be charged for object storage costs for any remaining backups.

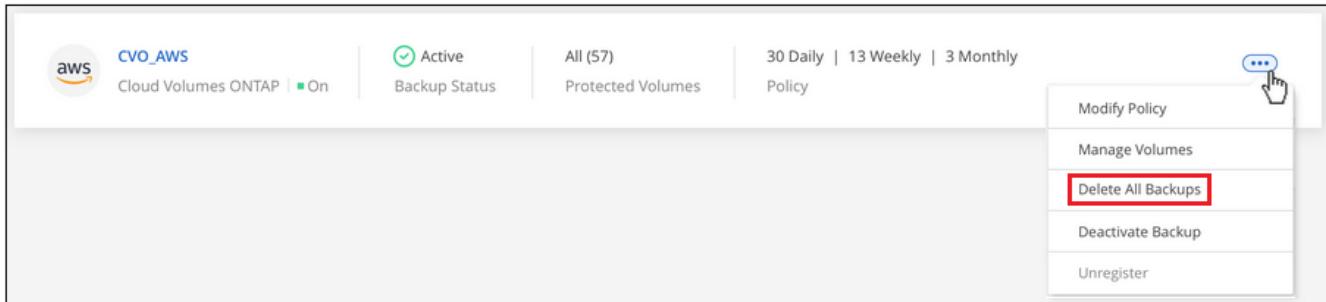
Deleting all backup files for a working environment

Steps

- From the Backup Dashboard, select **Backup Settings**.

The screenshot shows the 'Backup' tab of the Backup Dashboard. It includes a dropdown menu for 'All Backup Working Environments', a 'Backup Settings' button (which is highlighted with a red box), and a summary section for 'Protected Volumes Status' showing 57 healthy volumes and 0 failed volumes.

- Click **...** for the working environment where you want to delete all backups and select **Delete All Backups**.



3. In the confirmation dialog box, enter the name of the working environment and click **Delete**.

Deleting all backup files for a volume

Steps

1. From the Backup Dashboard, click **...** for the source volume and select **Details & Backup List**.

A screenshot of the Backup Dashboard showing the 'Details & Backup List' option for a specific volume. The dashboard displays 1 Working Environment, 57 Protected Volumes, and 15.1 TB Total Backup Capacity. The 'Protected Volumes Status' section shows 57 Healthy Backup Volumes and 0 Failed Backup Volumes. Below, a table lists 57 backups for three volumes: Volume_1, Volume_2, and Volume_3. The 'Backups' column header has a dropdown arrow pointing to a menu where 'Details & Backup List' is highlighted with a red box.

The list of all backup files is displayed.

The screenshot shows the NetApp Backup Dashboard interface. It consists of three main sections: **Source**, **Destination**, and **Backup Information**.

- Source:** Details include Working Environment (Working Environment Name), Type (Cloud Volumes ONTAP (HA)), Provider (AWS), Volume (Volume Name), and SVM (SVM Name).
- Destination:** Details include Cloud Provider (AWS), Bucket (Backup Bucket Name), Region (US East (N. Virginia)), and Account ID (012345678901234567890).
- Backup Information:** Relationship Status is Failed. Last Backup was on May 22 2019, 00:00:00. Lag Duration is 28 days ago. There have been 2,050 Backups. The Backup Policy is Netapp7YearsRetention.

Below these sections is a table titled "2,050 Backups" with columns for **Backup Name** and **Date**. The table lists three entries:

Backup Name	Date	Actions
Backup_2020_Jan	May 22 2019, 00:00:00	...
Backup_2020_Mar	May 22 2019, 00:00:00	...
Backup_2020_Apr	May 22 2019, 00:00:00	...

2. Click Actions > Delete all Backups.

The screenshot shows the same Backup Dashboard interface as above, but with the **Actions** button open. A dropdown menu appears with the following options:

- Delete All Backups** (highlighted with a red box)
- Download Backup Report

3. In the confirmation dialog box, enter the volume name and click **Delete**.

Deleting a single backup file for a volume

You can delete a single backup file. This feature is available only if the volume backup was created from a system with ONTAP 9.8 or greater.

Steps

- From the Backup Dashboard, click **...** for the source volume and select **Details & Backup List**.

The screenshot shows the NetApp Backup interface. At the top, there are tabs for 'Backup' (which is selected) and 'Restore'. Below the tabs, a dropdown menu shows 'All Backup Working Environments'. On the right, there's a 'Backup Settings' button. The main area displays summary statistics: 1 Working Environment, 57 Protected Volumes, and 15.1 TB Total Backup Capacity. To the right, a 'Protected Volumes Status' section shows 57 Healthy Backup Volumes and 0 Failed Backup Volumes. Below this, a table lists 57 Backups. The columns include Source Working Environment, Source Volume, Source SVM, Last Backup, Backups, and Backup Status. The first three rows show entries for 'CVO_AWS' with volumes 'Volume_1', 'Volume_2', and 'Volume_3' respectively, all in 'SVM_1'. The 'Backup Status' column indicates they are 'Active'. A red box highlights the 'Details & Backup List' link next to the first entry.

The list of all backup files is displayed.

This screenshot shows the detailed configuration for a backup relationship. It is divided into three main sections: 'Source', 'Destination', and 'Backup Information'. The 'Source' section includes fields for Working Environment, Type (Cloud Volumes ONTAP (HA)), Provider (AWS), Volume (Volume Name), and SVM (SVM Name). The 'Destination' section includes fields for Cloud Provider (AWS), Bucket (Backup Bucket Name), Region (US East (N. Virginia)), and Account ID (012345678901234567890). The 'Backup Information' section shows Relationship Status (Failed), Last Backup (May 22 2019, 00:00:00), Lag Duration (28 days ago), Backups (2,050), and Backup Policy (Netapp7YearsRetention). Below this, a table lists 2,050 Backups. The columns are 'Backup Name' and 'Date'. The first three entries are 'Backup_2020_Jan' (May 22 2019, 00:00:00), 'Backup_2020_Mar' (May 22 2019, 00:00:00), and 'Backup_2020_Apr' (May 22 2019, 00:00:00). A red box highlights the 'Actions' button at the top right of the backup list table.

2. Click **...** for the volume backup file you want to delete and click **Delete**.

This screenshot shows the same backup list as the previous one, but with a focus on the 'Delete' action. A red box highlights the 'Delete' button in the context menu that appears when hovering over the three-dot menu for the 'Backup_2020_Feb' entry. The table lists four backup entries: 'Backup_2020_Feb' (May 22 2019, 00:00:00), 'Backup_2020_Jan' (May 22 2019, 00:00:00), and 'Backup_2020_Mar' (May 22 2019, 00:00:00).

- In the confirmation dialog box, click **Delete**.

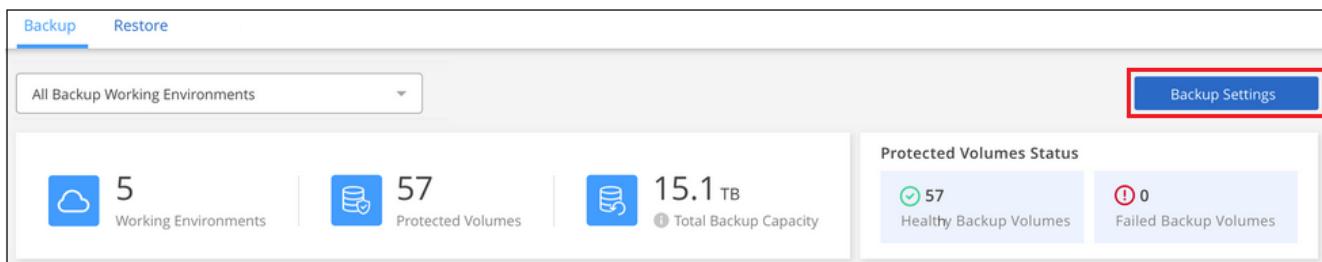
Disabling Cloud Backup for a working environment

Disabling Cloud Backup for a working environment disables backups of each volume on the system, and it also disables the ability to restore a volume. Any existing backups will not be deleted. This does not unregister the backup service from this working environment - it basically allows you to pause all backup and restore activity for a period of time.

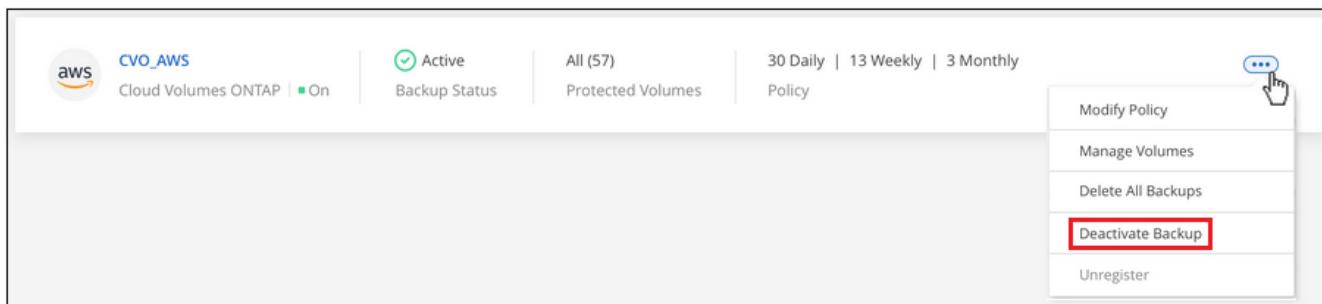
Note that you'll continue to be charged by your cloud provider for object storage costs for the capacity that your backups use unless you [delete the backups](#).

Steps

- From the Backup Dashboard, select **Backup Settings**.



- From the *Backup Settings* page, click **...** for the working environment where you want to disable backups and select **Deactivate Backup**.



- In the confirmation dialog box, click **Deactivate**.



An **Activate Backup** button appears for that working environment while backup is disabled. You can click this button when you want to re-enable backup functionality for that working environment.

Unregistering Cloud Backup for a working environment

You can unregister Cloud Backup for a working environment if you no longer want to use backup functionality and you want to stop being charged for backups in that working environment. Typically this feature is used when you're planning to delete a working environment and you want to cancel the backup service.

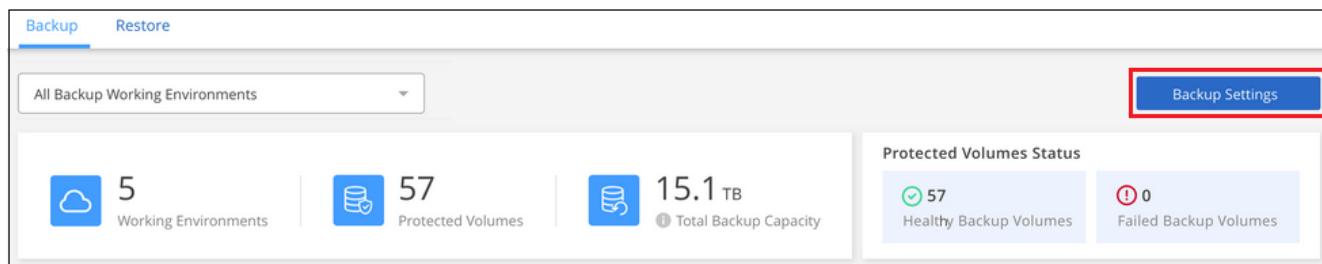
Before you can unregister Cloud Backup, you must perform the following steps, in this order:

- Deactivate Cloud Backup for the working environment
- Delete all backups for that working environment

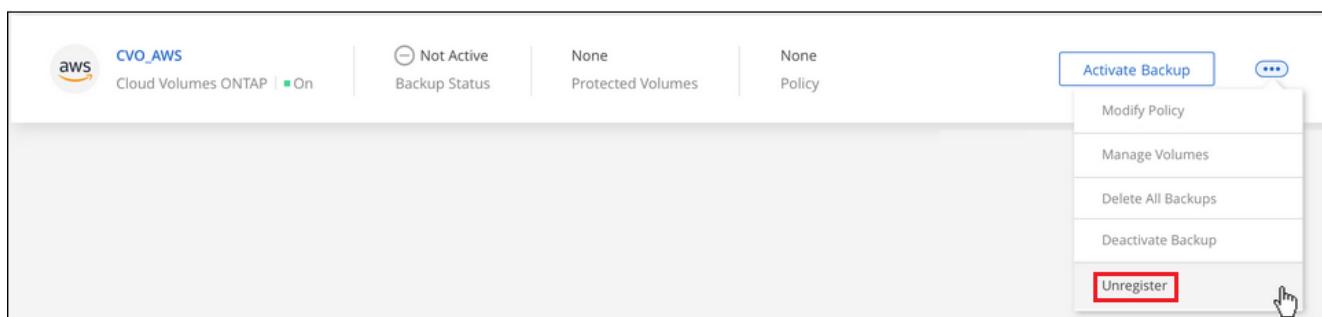
The unregister option is not available until these two actions are complete.

Steps

1. From the Backup Dashboard, select **Backup Settings**.



2. From the *Backup Settings* page, click **...** for the working environment where you want to unregister the backup service and select **Unregister**.



3. In the confirmation dialog box, click **Unregister**.

Restoring data from backup files

Backups are stored in an object store in your cloud account so that you can restore data from a specific point in time. You can restore an entire volume from a saved backup file, or if you only need to restore a few files, you can restore individual files from a saved backup file.

You can restore an entire volume to the same working environment, to a different working environment that's using the same cloud account, or to an on-premises ONTAP system. See [Restoring a volume from a backup](#).

You can restore files to a volume in the same working environment, to a volume in a different working environment that's using the same cloud account, or to a volume on an on-premises ONTAP system. See [Restoring files from a backup](#).

Supported working environments and object storage providers

You can restore a volume, or individual files, from a backup file to the following working environments:

Backup File Location	Destination Working Environment	
	Volume Restore	File Restore
Amazon S3	Cloud Volumes ONTAP in AWS On-premises ONTAP system	Cloud Volumes ONTAP in AWS On-premises ONTAP system

Backup File Location	Destination Working Environment	
Azure Blob	Cloud Volumes ONTAP in Azure On-premises ONTAP system	Cloud Volumes ONTAP in Azure On-premises ONTAP system
Google Cloud Storage	Cloud Volumes ONTAP in Google On-premises ONTAP system	
NetApp StorageGRID	On-premises ONTAP system	

The Restore Dashboard

You access the Restore Dashboard by clicking the **Backup & Restore** tab from the top of Cloud Manager, or you can click  > **View Restore Dashboard** from the Backup & Restore service from the Services panel.



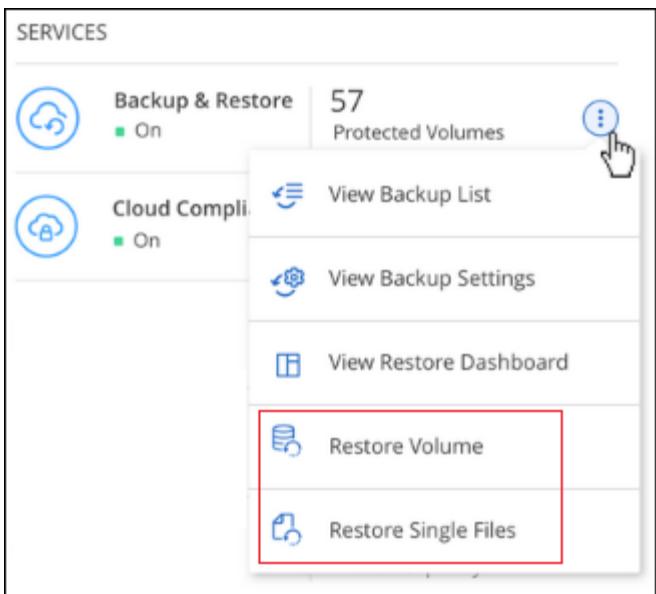
The Cloud Backup service must already be activated for at least one working environment.

Volume Name	Restore Date	Status
VolumeVeryLongName1	Nov 19 2020, 4:15:05 pm	Completed Successfully
Source	WorkingEnvironmentLongName	Backup2020_4_2
Destination	WorkingEnvironmentLongName	Nov 19 2020, 4:15:05 pm
VolumeVeryLongName2	Nov 19 2020, 4:15:05 pm	In Progress
VolumeVeryLongName3	Nov 19 2020, 4:15:05 pm	Completed Successfully

The Restore Dashboard provides buttons for you to restore volumes and files. Clicking the *Restore Volumes* or *Restore Files* buttons starts a wizard that walks you through the steps to restore that data.

The dashboard also provides a list of all the volumes and all the files you have restored in case you need a history of previous restore actions. You can expand the row for each restored volume or file to view the details about the source and destination locations for the volume or file.

Note that you can also initiate a volume or file restore operation from a working environment in the Services panel. When started from this location the source working environment selection is automatically filled with the name of the current working environment.



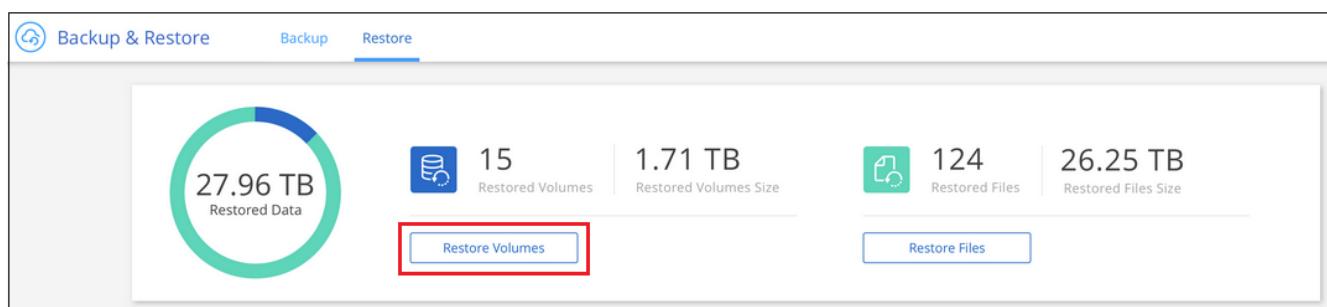
Restoring a volume from a backup file

When you restore a volume from a backup file, Cloud Manager creates a *new* volume using the data from the backup. You can restore the data to a volume in the same working environment or to a different working environment that's located in the same cloud account as the source working environment. You can also restore files to an on-premises ONTAP system.

You should know the name of the volume you want to restore and the date of the backup file you want to use to create the newly restored volume.

Steps

1. Select the **Backup & Restore** tab.
2. Click the **Restore** tab and the Restore Dashboard is displayed.
3. Click **Restore Volumes**.



4. In the *Select Source* page, navigate to the backup file for the volume you want to restore. Select the **Working Environment**, the **Volume**, and the **Backup** that has the date/time stamp that you want to restore.

1 Select Source 2 Select Destination

Select Source

Selected Working Environment
Working Environment Name 3

Selected Volume
Volume Very Long Name

Selected Backup >
Backup_2020_Feb

Backup Name	Date
Backup_2020_Jan	Nov 19 2020, 4:15:05 pm
<input checked="" type="checkbox"/> Backup_2020_Feb	Nov 19 2020, 4:15:05 pm
Backup_2020_Mar	Nov 19 2020, 4:15:05 pm
Backup_2020_Apr	Nov 19 2020, 4:15:05 pm

5. Click **Continue**.
6. In the *Select Destination* page, select the **Working Environment** where you want to restore the volume.

1 Select Source 2 Select Destination

Select Destination

Select Working Environment >

Destination Volume

Working Environment Name Type Provider

Working Environment Name	Type	Provider
Working Environment 3 On <input type="button" value="Source Working Environment"/>	Cloud Volumes ONTAP	Azure
<input checked="" type="checkbox"/> Working Environment 2 On	Cloud Volumes ONTAP	Azure

7. If you select an on-premises ONTAP system and you haven't already configured the cluster connection to the object storage, you are prompted for additional information:
 - When restoring from Amazon S3, select the AWS Account and the Access Key and Secret Key to access the object storage, the region where the backups are stored, and the IPspace in the ONTAP cluster where the destination volumes reside.
 - When restoring from Azure Blob, select the Azure Subscription to access the object storage, the region where the backups are stored, and the IPspace in the ONTAP cluster where the destination volumes reside.
 - When restoring from Google Cloud Storage, select the Google Cloud Project and the Access Key and Secret Key to access the object storage, the region where the backups are stored, and the IPspace in the ONTAP cluster where the destination volumes reside.
 - When restoring from StorageGRID, select the Access Key and Secret Key needed to access the object storage, and the IPspace in the ONTAP cluster where the destination volumes reside.
8. Select the Storage VM where the volume will reside and enter the name you want to use for the restored volume. By default, **<source_volume_name>_Restore** is used as the volume name.

Select Destination

Selected Working Environment
Working Environment Name 2

Destination Volume >
Source_Volume_Name_Restore

i A new volume will be created in the working environment based on the backup snapshot you selected

Storage VM	Storage VM 1	Volume Information	Volume Size: 100 GB
Aggregate	Aggregate 1	Snapshot Policy:	Default
Volume Name	Source_Volume_Name_Restore	NFS Protocol:	Custom export policy, 10.20.0.0/16
		Storage Efficiency:	ON
		Disk Type:	GP2
		Tiering:	all

You can select the Aggregate that the volume will use for its' capacity only when restoring a volume to an on-premises ONTAP system.

- Click **Restore** and you are returned to the Restore Dashboard so you can review the progress of the restore operation.

Result

Cloud Manager creates a new volume based on the backup you selected. You can [manage this new volume](#) as required.

Restoring files from a backup

If you only need to restore a few files from a volume, you can choose to restore individual files instead of restoring the entire volume. You can restore files to a volume in the same working environment, or to a different working environment that's using the same cloud account. You can also restore files to an on-premises ONTAP system.

All the files are restored to the same destination volume that you choose. If you want to restore files to different volumes, you need run the restore process a second time.

Prerequisites

- The ONTAP version must be 9.6 or greater in your Cloud Volumes ONTAP or on-premises ONTAP systems to perform file restore operations.
- Restoring individual files from a backup file uses a separate Restore instance/virtual machine. See the [AWS Requirements](#) or [Azure Requirements](#) to make sure your environment is ready.
- Restoring files also requires that specific EC2 permissions are added to the user role that provides Cloud Manager with permissions. [Make sure all the permissions are configured correctly](#).
- AWS cross-account restore requires manual action in the cloud provider console. See the AWS topic [granting cross-account bucket permissions](#) for details.

File Restore process

The process goes like this:

- When you want to restore one or more files from a volume, click the **Restore** tab, click **Restore Files**, and select the backup file in which the file (or files) reside.

2. The Restore instance starts up and displays the folders and files that exist within the backup file.

Note: The Restore instance is deployed in your cloud providers' environment the first time you restore a file.

3. Choose the file (or files) that you want to restore from that backup.
4. Select the location where you want the file(s) to be restored (the working environment, volume, and folder), and click **Restore**.
5. The file(s) are restored, and then the Restore instance is shut down to save costs after a period of inactivity.

Restoring files from a backup file

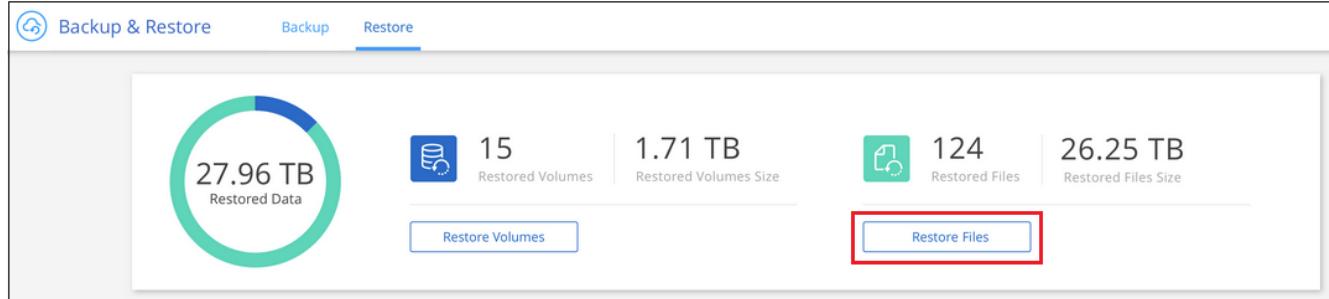
Follow these steps to restore files from a volume backup to a volume. You should know the name of the volume and the date of the backup file that you want to use to restore the file, or files. This functionality uses Live Browsing so that you can view the list of directories and files within the backup file.

The following video shows a quick walkthrough of restoring a single file:

[| <https://img.youtube.com/vi/ROAY6gPL9N0/maxresdefault.jpg>

Steps

1. Select the **Backup & Restore** tab.
2. Click the **Restore** tab and the Restore Dashboard is displayed.
3. Click the **Restore Files** button.



4. In the *Select Source* page, navigate to the backup file for the volume that contains the files you want to restore. Select the **Working Environment**, the **Volume**, and the **Backup** that has the date/time stamp from which you want to restore files.

1 Select Source 2 Select Files 3 Select Destination

Select Source

56 Backups

Backup Name	Date	Time Zone
Backup_2020_Jan	September 30 2020 00:00:00	NY, USA (GMT-4)
Backup_2020_Feb	September 30 2020 00:00:00	NY, USA (GMT-4)
Backup_2020_Mar	September 30 2020 00:00:00	NY, USA (GMT-4)
Backup_2020_Apr	September 30 2020 00:00:00	NY, USA (GMT-4)

- Click **Continue** and the Restore instance is started. After a few minutes the Restore instance displays the list of folders and files from the volume backup.

Note: The Restore instance is deployed in your cloud providers' environment the first time you restore a file, so this step could take a few minutes longer the first time.

Select Files

Select Files

Folders & Files

All Folders & Files > Folder A Very Long Name

Name	Last Modified	Size
File D Very Long Name	September 30 2020 00:00:00	1.25 MB
File E Very Long Name	September 30 2020 00:00:00	1.25 MB

- In the **Select Files** page, select the file or files that you want to restore and click **Continue**.
 - You can click the search icon and enter the name of the file to navigate directly to the file.
 - You can click the file name if you see it.
 - You can navigate down levels in folders using the **>** button at the end of the row to find the file.

As you select files they are added to the left side of the page so you can see the files that you have already chosen. You can remove a file from this list if needed by clicking the **x** next to the file name.

- In the **Select Destination** page, select the **Working Environment** where you want to restore the files.

1 Select Source 2 Select Files 3 Select Destination

Select Destination

5 Working Environments

Working Environment Name	Type	Provider
Working Environment 3 * On Source Working Environment	Cloud Volumes ONTAP	Azure
Working Environment 1 * On	Cloud Volumes ONTAP	Azure
Working Environment 2 * On	On-Premises	--

If you select an on-premises cluster and you haven't already configured the cluster connection to the object storage, you are prompted for additional information:

- When restoring from Amazon S3, enter the IPspace in the ONTAP cluster where the destination volumes reside, and the AWS Access Key and Secret Key needed to access the object storage.
- When restoring from Azure Blob, enter the IPspace in the ONTAP cluster where the destination volumes reside.

8. Then select the **Volume** and the **Folder** where you want to restore the files.

Name	Last Modified	Size
Folder A Very Long Name	September 30 2020 00:00:00	---
Folder B Very Long Name	September 30 2020 00:00:00	---
Folder C Very Long Name	September 30 2020 00:00:00	---
Folder D Very Long Name	September 30 2020 00:00:00	---

You have a few options for the location when restoring files.

- When you have chosen **Select Target Folder**, as shown above:
 - You can select any folder.
 - You can hover over a folder and click **>** at the end of the row to drill down into subfolders, and then select a folder.
 - If you have selected the same destination Working Environment and Volume as where the source file was located, you can select **Maintain Source Folder Path** to restore the file, or all files, to the same folder where they existed in the source structure. All the same folders and sub-folders must already exist; folders are not created.
9. Click **Restore** and you are returned to the Restore Dashboard so you can review the progress of the restore operation.

The Restore instance is shut down after a certain period of inactivity to save you money so that you incur costs only when it is active.

Cross-account and cross-region configurations

These topics describe how to configure Cloud Backup for cross account configurations when using different cloud providers.

- [Configure Cloud Backup for multi-account access in AWS](#)
- [Configure Cloud Backup for multi-account access in Azure](#)

Configure backup for multi-account access in AWS

Cloud Backup enables you to create backup files in an AWS account that is different than where your source volumes reside. And both of those accounts can be different than the account where the Cloud Manager Connector resides.

Just follow the steps below to set up your configuration in this manner.

Set up VPC peering between accounts

1. Log in to second account and Create Peering Connection:
 - a. Select a local VPC: Select the VPC of the second account.
 - b. Select another VPC: Enter the account ID of the first account.
 - c. Select the Region where the Cloud Manager Connector is running. In this test setup both accounts are running in same region.
 - d. VPC ID: Log into first account and enter the acceptor VPC ID. This is the VPC ID of the Cloud Manager Connector.

The screenshot shows the 'Create Peering Connection' page in the AWS Management Console. At the top, it says 'Peering Connections > Create Peering Connection'. The main section is titled 'Create Peering Connection'. It has a 'Peering connection name tag' input field containing 'cbs-multi-account'. Below that, under 'Select a local VPC to peer with', there's a dropdown menu showing 'vpc-82f55afa'. A table titled 'CIDRs' lists a single entry: '10.0.0.0/16' with a status of 'associated'. Under 'Select another VPC to peer with', there are two sections: 'Account' (radio buttons for 'My account' and 'Another account', with 'Another account' selected) and 'Region' (radio buttons for 'This region (us-east-1)' and 'Another Region', with 'This region (us-east-1)' selected). There is also an input field for 'VPC ID (Acceptor)*' containing 'vpc-116d9174'.

A Success dialog displays.

✓ Success

A VPC peering connection (pcx-049758069d9b7c140) has been requested.

The owner of **vpc-116d9174** must accept the peering connection.

Requester VPC owner	733004784675 (This account)	Acceptor VPC owner	464262061435
Requester VPC ID	vpc-82f55afa	Acceptor VPC ID	vpc-116d9174
Requester VPC Region	us-east-1	Acceptor VPC Region	us-east-1
Requester VPC CIDRs	10.0.0.0/16	Acceptor VPC CIDRs	-

The status of the peering connection shows as Pending Acceptance.

Name	Peering Connectivity	Status	Requester VPC	Acceptor VPC	Requester CIDRs	Acceptor CIDRs	Requester Owner	Acceptor Owner
cbs-multi-ac...	pcx-049758069d9...	Pending Acceptance	vpc-82f55afa VP...	vpc-116d9174	10.0.0.0/16	-	733004784675	464262061435
cbs-multi-peer	pcx-05f2d310cb7f...	Deleted	vpc-82f55afa VP...	vpc-116d9174	-	-	733004784675	464262061435
New_Peering	pcx-6d55ca04	Active	vpc-b16c90d4 V...	vpc-fc2aa39a De...	172.31.0.0/16	192.168.0.0/16	733004784675	733004784675

2. Log into the first account and accept the peering request:

The screenshot shows the AWS VPC Peering Connections list. A context menu is open over a row where the status is 'Pending Acceptance'. The menu options are 'Accept Request' (highlighted in yellow), 'Reject Request', 'Delete VPC Peering Connection', 'Edit ClassicLink Settings', and 'Add/Edit Tags'. The main table lists four peering connections, with the first one being the pending request.

Name	Status	Requester VPC	Acceptor VPC	Requester CIDRs	Acceptor CIDRs	Requester Owner	Acceptor Owner
cbs-multi-ac...	Pending Acceptance	vpc-82f55afa VP...	vpc-116d9174	10.0.0.0/16	-	733004784675	464262061435
cbs-multi-peer	Deleted	vpc-82f55afa VP...	vpc-116d9174	-	-	733004784675	464262061435
New_Peering	Active	vpc-b16c90d4 V...	vpc-fc2aa39a De...	172.31.0.0/16	192.168.0.0/16	733004784675	733004784675

The dialog box is titled 'Accept VPC Peering Connection Request'. It asks, 'Are you sure you want to accept this VPC peering connection request (pcx-049758069d9b7c140)?'. Below are the details of the request:

Requester Account ID	733004784675	Acceptor Account ID	464262061435 (This account)
Requester VPC ID	vpc-82f55afa	Acceptor VPC ID	vpc-116d9174
Requester VPC Region	us-east-1	Acceptor VPC Region	us-east-1
Requester VPC CIDR	10.0.0.0/16	Acceptor VPC CIDR	-

At the bottom are 'Cancel' and 'Yes, Accept' buttons.

a. Click Yes.

The dialog box is titled 'Accept VPC Peering Connection Request'. It displays the message: 'Your VPC Peering Connection has been established.' It also states: 'To send and receive traffic across this VPC peering connection, you must add a route to the peered VPC in one or more of your VPC route tables. [Learn more](#)' and 'Modify my route tables now'. At the bottom is a 'Close' button.

The connection now shows as Active. We have also added a Name tag to identify the peering connection called **cbs-multi-account**.

	Name	Peering Connection	Status	Requester VPC	Acceptor VPC	Requester CIDRs	Acceptor CIDRs	Requester Owner	Acceptor Owner
	pcx-004715531514cb0d8	estycvoconnect	Active	vpc-0647747d M...	vpc-116d9174	10.2.0.0/24	172.31.0.0/16	464262061435	464262061435
	cbs-multi-account	pcx-0305041f9cc2dfbd1	Active	vpc-116d9174	vpc-445d4f21	172.31.0.0/16	10.129.0.0/20	464262061435	759995470648
	hill-vpc-peer-chen	pcx-0d0e5c7fc4360254d	Active	vpc-0d12df59528f...	vpc-824dc0e4 nf...	10.0.0.0/24	10.20.30.0/24	464262061435	464262061435

b. Refresh the peering connection in the second account and notice that the status changes to Active.

	Name	Peering Connection	Status	Requester VPC	Acceptor VPC	Requester CIDRs	Acceptor CIDRs	Requester Owner	Acceptor Owner
	cbs-multi-account	pcx-049758069d9b7c140	Active	vpc-82f55afa VP...	vpc-116d9174	10.0.0.0/16	172.31.0.0/16	733004784675	464262061435
	New_Peering	pcx-6d55ca04	Active	vpc-b16c90d4 V...	vpc-fc2aa39a De...	172.31.0.0/16	192.168.0.0/16	733004784675	733004784675

Add a route to the route tables in both accounts

1. Go to VPC > Subnet > Route table.

Details			
Subnet ID subnet-4d315328	State Available	VPC vpc-116d9174	IPv4 CIDR 172.31.64.0/20
Available IPv4 addresses 3587	IPv6 CIDR -	Availability Zone us-east-1a	Availability Zone ID use1-az1
Network border group us-east-1	Route table rtb-4da55528	Network ACL acl-c3738a6	Default subnet Yes
Auto-assign public IPv4 address Yes	Auto-assign IPv6 address No	Auto-assign customer-owned IPv4 address No	Customer-owned IPv4 pool -
Outpost ID -	Owner 464262061435	Subnet ARN arn:aws:ec2:us-east-1:464262061435:subnet/subnet-4d315328	

Flow logs **Route table** Network ACL Sharing Tags

2. Click on the Routes tab.

Route Table ID : rtb-4da55528 Add filter							
	Name	Route Table ID	Explicit subnet association	Edge associations	Main	VPC ID	Owner
		rtb-4da55528	subnet-4d315328	-	Yes	vpc-116d9174	464262061435

Route Table: rtb-4da55528

Summary **Routes** Subnet Associations Edge Associations Route Propagation Tags

Edit routes

View All routes

Destination	Target	Status	Propagated
172.31.0.0/16	local	active	No
pl-63a5400a	vpc-e-098587ed33c36408c	active	No

3. Click Edit routes.

Edit routes

Destination	Target	Status	Propagated	
172.31.0.0/16	local	active	No	
10.20.30.0/24	pcx-0791b47f6f9a27d65	active	No	×
10.129.0.0/20	pcx-0305041f9cc2dfbdb	active	No	×
Add route				
* Required				Cancel Save routes

4. Click **Add route**, and from the Target drop-down list select **Peering Connection**, and then select the peering connection that you created.

- a. In the Destination, enter the other account's subnet CIDR.

Destination	Target	Status	Propagated	
172.31.0.0/16	local	active	No	
10.20.30.0/24	pcx-0791b47f6f9a27d65	active	No	×
10.129.0.0/20	pcx-0305041f9cc2dfbdb	active	No	×
10.0.0.0/24	pcx-		No	×
Add route				
* Required				Cancel Save routes

- b. Click **Save routes** and a Success dialog displays.

Route Tables > Edit routes
Edit routes
Routes successfully edited
Close

Add the second AWS account credentials in Cloud Manager

1. Add the second AWS account, for example, *Saran-XCP-Dev*.

Credentials		+ Add Credentials	
3 Credentials			
 Instance Profile 464262061435 AWS Account ID aws-sub-a2 Subscription	Credential Type: AWS Keys CBS-SR-OCCMOCCM1620912870830... IAM Role 2 Working Environments	 Saran-XCP-Dev 733004784675 AWS Account ID aws-sub-a2 Subscription	Credential Type: AWS Keys AKIA2VKT5MQRZRAWW3HI AWS Access Key 0 Working Environments

2. In the Discover Cloud Volumes ONTAP page, select the newly added credentials.

Choose an AWS region and then select the working environment that you want to discover.

AWS Region
US East | N. Virginia

aws AWS Credentials

Credential Name

- Saran-XCP-Dev | Account ID: 733004784675
- Instance Profile | Account ID: 464262061435

To add new AWS credentials, go to the [Credentials settings](#).

Apply **Cancel**

3. Select the Cloud Volumes ONTAP system you want to discover from second account. You can also deploy a new Cloud Volumes ONTAP system in the second account.

Add an Existing Cloud Volumes ONTAP		Region				
		This working environment will be created in Cloud Provider Account: Saran-XCP-Dev Account ID: 733004784675 Switch Account				
↑ Previous Step		Choose an AWS region and then select the working environment that you want to discover.				
		AWS Region US East N. Virginia				
Cloud Volumes ONTAP instances found						
Name	VPC Name	Availability Zone	Subnet Id	Cloud Formation Name	Cluster Address	Type
cbscvo01	VPC-NAT	us-east-1f	subnet-68e8d464	cbscvo01	10.0.0.80	Cloud Volumes ONTAP
testbyolliraz	VPC for VSA	us-east-1a	subnet-c1d99699	testbyolliraz	172.31.5.142	Cloud Volumes ONTAP
idanAwsHa991001	VPC for VSA	us-east-1a	subnet-c1d99699	idanAwsHa991001	172.31.5.234,172.31.5.110	HA Cloud Volumes ONTAP
Continue						

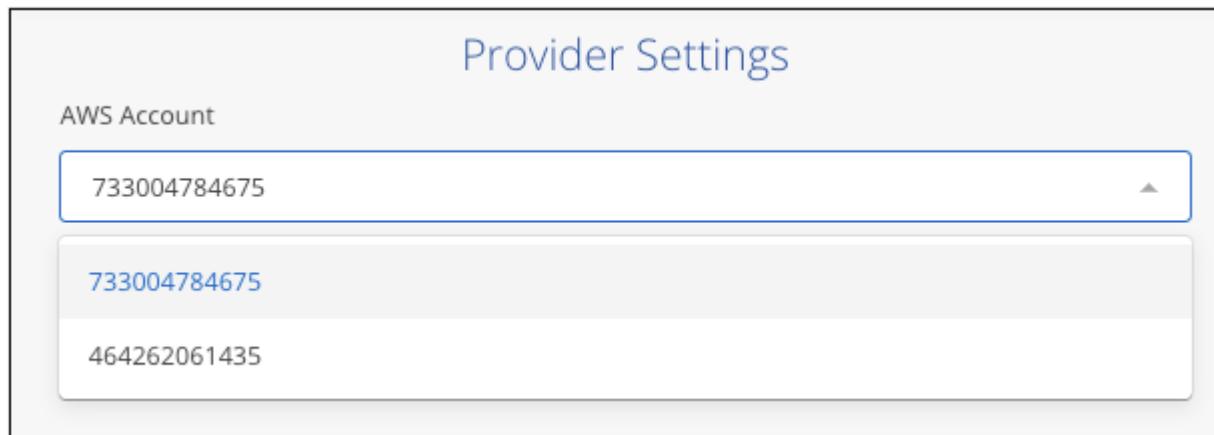
The Cloud Volumes ONTAP system from the second account is now added to Cloud Manager which is running in a different account.

The screenshot shows the Cloud Manager interface with the following components:

- Top Bar:** Account (cbs), Workspace (Workspace-1), Connector (LBS-SR-ULCM), and various navigation icons.
- Canvas:** Shows three Cloud Volumes ONTAP systems:
 - A central system labeled "cbssrnonprem" with "SINGLE" storage, "Cloud Volumes ONTAP", and a status of "Off".
 - A left system labeled "cbscvo01" with "SINGLE" storage, "Cloud Volumes ONTAP", "2 GiB Capacity", and an "aws" icon.
 - A bottom system labeled "CbsSrCV099Aws" with "SINGLE" storage, "Cloud Volumes ONTAP", "11 GiB Capacity", and an "aws" icon.
- Amazon S3 Bucket:** A cloud icon labeled "Amazon S3" with "134 Buckets" and "8 Regions".
- Details Panel (cbscvo01):**
 - DETAILS:** Cloud Volumes ONTAP | AWS | Single
 - NOTIFICATIONS:** New Version Available
 - SERVICES:**
 - Replication: Off, Enable
 - Backup & Restore: Off, Enable
 - K8s: Off, Connect a Cluster
 - Data Sense & Compliance: Off, Enable
 - Monitoring: Off, Enable

Enable backup in the other AWS account

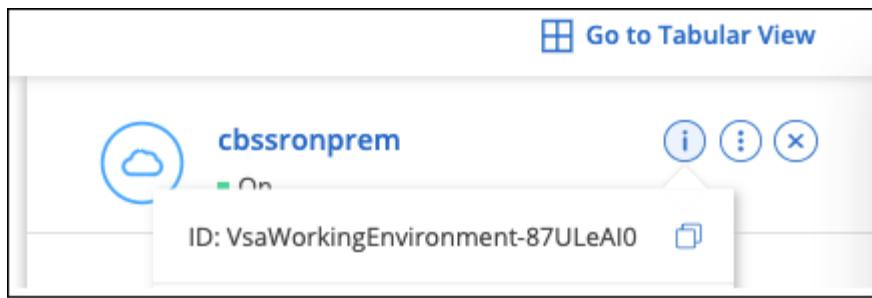
1. In Cloud Manager, enable backup for the Cloud Volumes ONTAP system running in the first account, but select the second account as the location for creating the backup files.



2. Then select a backup policy and the volumes you want to back up, and Cloud Backup attempts to create a new bucket in the selected account.

However, adding the bucket to the Cloud Volumes ONTAP system will fail because Cloud Backup uses the instance profile to add the bucket and the Cloud Manager instance profile doesn't have access to the resources in the second account.

3. Get the working environment ID for the Cloud Volumes ONTAP system.



Cloud Backup creates every bucket with the prefix Netapp-backup- and will include the working environment ID; for example: 87ULEA10

4. In the EC2 portal, go to S3 and search for the bucket with name ending with 87ULEA10 and you'll see the bucket name displayed as Netapp-backup-vsa87ulea10.

Name	AWS Region	Access	Creation date
netapp-backup-vsa87uleai0	US East (N. Virginia) us-east-1	Bucket and objects not public	May 25, 2021, 17:50:08 (UTC+05:30)

5. Click on the bucket, then click the Permissions tab, and then click **Edit** in the Bucket policy section.

6. Add a bucket policy for the newly created bucket to provide access to the Cloud Manager's AWS account, and then Save the changes.

```

{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "PublicRead",
      "Effect": "Allow",
      "Principal": {
        "AWS": "arn:aws:iam::464262061435:root"
      },
      "Action": [
        "s3>ListBucket",
        "s3:GetBucketLocation",
        "s3GetObject",
        "s3PutObject",
        "s3DeleteObject"
      ],
      "Resource": [
        "arn:aws:s3::::netapp-backup-vsa87uleai0",
        "arn:aws:s3::::netapp-backup-vsa87uleai0/*"
      ]
    }
  ]
}

```

Note that "AWS": "arn:aws:iam::464262061435:root" gives complete access to this bucket for all resources in account 464262061435. If you want to reduce it to specific role, level, you can update the policy with specific role(s). If you are adding individual roles, ensure that occm role also added, otherwise backups will not get updated in the Cloud Backup UI.

For example: "AWS": "arn:aws:iam::464262061435:role/cvo-instance-profile-version10-d8e-iamInstanceRole-IKJPJ1HC2E7R"

7. Retry enabling Cloud Backup on the Cloud Volumes ONTAP system and this time it should be successful.

Configure backup for multi-account access in Azure

Cloud Backup enables you to create backup files in an Azure account that is different than where your source volumes reside. And both of those accounts can be different than the account where the Cloud Manager Connector resides.

Just follow the steps below to set up your configuration in this manner.

Set up VNet peering between accounts

Note that if you want Cloud Manager to manage your Cloud Volumes ONTAP system in a different account/region, then you need to setup VNet peering. VNet peering is not required for storage account connectivity.

1. Log in to the Azure portal and from home, select Virtual Networks.
2. Select the subscription you are using as subscription 1 and click on the VNet where you want to set up peering.

The screenshot shows the Azure Virtual networks page. At the top, there are navigation links: Home > Virtual networks. Below the header, there are several filter options: Subscription == OCCM Dev, Resource group == all, Location == all, and a 'Filter for any field...' search bar. The main table lists 60 records, showing columns for Name, Resource group, and Location. The first row, 'cbsnetwork', is highlighted with a red box.

Name	Resource group	Location
cbsnetwork	occm_group_eastasia	East Asia
Vnet1	occm_group_australiaeast	Australia East
Vnet1	occm_group_australiasoutheast	Australia Southeast

3. Select **cbsnetwork** and from the left panel, click on **Peerings**, and then click **Add**.

The screenshot shows the 'Add Peering' configuration page. It includes fields for Subscription (OCCM Automation), Virtual network (cbse2evnet), and various traffic settings for both directions. The 'Virtual network gateway or Route Server' section also has options for using local or remote gateways, or none. A prominent blue 'Add' button is at the bottom.

4. Enter the following information on the Peering page and then click **Add**.

- Peering link name for this network: you can give any name to identify the peering connection.
- Remote virtual network peering link name: enter a name to identify the remote VNet.
- Keep all the selections as default values.

- Under subscription, select the subscription 2.
- Virtual network, select the virtual network in subscription 2 to which you want to set up the peering.

The screenshot shows the 'Peerings' blade for a virtual network named 'cbsnetwork'. The left sidebar contains navigation links such as Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, Settings, Address space, Connected devices, Subnets, DDoS protection, Firewall, Security, DNS servers, and Peerings. The main area features a search bar and buttons for 'Add' and 'Refresh'. A table lists a single peering entry:

Name	Peering status	Peer
cbsnetwork	Connected	cbse2evnet

5. Perform the same steps in subscription 2 VNet and specify the subscription and remote VNet details of subscription 1.

Subscription * ⓘ

OCCM Dev

Virtual network *

cbsnetwork

Traffic to remote virtual network ⓘ

Allow (default)

Block all traffic to the remote virtual network

Traffic forwarded from remote virtual network ⓘ

Allow (default)

Block traffic that originates from outside this virtual network

Virtual network gateway or Route Server ⓘ

Use this virtual network's gateway or Route Server

Use the remote virtual network's gateway or Route Server

None (default)

Add

The peering settings are added.

cbse2evnet | Peerings

Virtual network

Search (Cmd+/) < Add Refresh

Filter by name...

Name	Peering status	Peer
cbsnetworkpeer	Connected	cbsnetwork

- Overview
- Activity log
- Access control (IAM)
- Tags
- Diagnose and solve problems
- Settings
 - Address space
 - Connected devices
 - Subnets
 - DDoS protection
 - Firewall
 - Security
 - DNS servers
 - Peerings

Create a private endpoint for the storage account

Now you need to create a private endpoint for the storage account. In this example, the storage account is created in subscription 1 and the Cloud Volumes ONTAP system is running in subscription 2.



You need network contributor permission to perform the following action.

```
{  
  "id": "/subscriptions/d333af45-0d07-4154-  
943dc25fbbce1b18/providers/Microsoft.Authorization/roleDefinitions/4d97b98  
b-1d4f-4787-a291-c67834d212e7",  
  "properties": {  
    "roleName": "Network Contributor",  
    "description": "Lets you manage networks, but not access to them.",  
    "assignableScopes": [  
      "/"  
    ],  
    "permissions": [  
      {  
        "actions": [  
          "Microsoft.Authorization/*/read",  
          "Microsoft.Insights/alertRules/*",  
          "Microsoft.Network/*",  
          "Microsoft.ResourceHealth/availabilityStatuses/read",  
          "Microsoft.Resources/deployments/*",  
          "Microsoft.Resources/subscriptions/resourceGroups/read",  
          "Microsoft.Support/*"  
        ],  
        "notActions": [],  
        "dataActions": [],  
        "notDataActions": []  
      }  
    ]  
  }  
}
```

1. Go to the storage account > Networking > Private endpoint connections and click **+ Private endpoint**.

The screenshot shows the Azure Storage account 'netappbackupvsatxdkmfym' under the 'Networking' section. The 'Private endpoint connections' tab is selected. On the left, there's a sidebar with links like Activity log, Tags, Diagnose and solve problems, Access Control (IAM), Data migration, Events, and Storage Explorer (preview). Below that are sections for Data storage (Containers, File shares, Queues, Tables) and Security + networking (Networking, Firewall rules). The main area shows a table for private endpoint connections with columns for Connection name, Connection state, Private endpoint, and Description. A search bar and filter dropdown are at the top of the table.

2. In the Private Endpoint *Basics* page:

- Select subscription 2 (where the Cloud Manager Connector and Cloud Volumes ONTAP system are deployed) and the resource group.
- Enter an endpoint name.
- Select the region.

The screenshot shows the 'Create a private endpoint' Basics page. The steps are numbered 1 to 5: Basics (selected), Resource, Configuration, Tags, and Review + create. The Basics step asks to use private endpoints to privately connect to a service or resource. It includes fields for Project details (Subscription: OCCM Dev, Resource group: cbsoccmdevcvo-rg) and Instance details (Name: cbse2e, Region: (Asia Pacific) East Asia).

3. In the *Resource* page, select Target sub-resource as **blob**.

Create a private endpoint

✓ Basics 2 Resource 3 Configuration 4 Tags 5 Review + create

Private Link offers options to create private endpoints for different Azure resources, like your private link service, a SQL server, or an Azure storage account. Select which resource you would like to connect to using this private endpoint. [Learn more](#)

Subscription	OCCM Dev (d333af45-0d07-4154-943d-c25fbbe1b18)
Resource type	Microsoft.Storage/storageAccounts
Resource	test150521
Target sub-resource *	<input type="text" value="blob"/>

4. In the Configuration page:

- Select the virtual network and subnet.
- Click the **Yes** radio button to "Integrate with private DNS zone".

Create a private endpoint

✓ Basics ✓ Resource 3 Configuration 4 Tags 5 Review + create

Networking

To deploy the private endpoint, select a virtual network subnet. [Learn more](#)

Virtual network *	<input type="text" value="cbsnetwork"/>
Subnet *	<input type="text" value="default (10.2.0.0/24)"/>

Tip: If you have a network security group (NSG) enabled for the subnet above, it will be disabled for private endpoints on this subnet only. Other resources on the subnet will still have NSG enforcement.

Private DNS integration

To connect privately with your private endpoint, you need a DNS record. We recommend that you integrate your private endpoint with a private DNS zone. You can also utilize your own DNS servers or create DNS records using the host files on your virtual machines. [Learn more](#)

Integrate with private DNS zone Yes No

Configuration name	Subscription	Private DNS zone
privatelink-blob-core-...	OCCM Dev	privatelink.blob.core.windows.net

[Review + create](#) < Previous Next : Tags >

5. In the Private DNS zone list, ensure that the Private Zone is selected from the correct Region, and click **Review + Create**.

Configuration name	Subscription	Private DNS zone
privatelink-blob-core-...	OCCM Dev	privatelink.blob.core.windows.net
		Filter private DNS zones
		occm_group_centralus
		privatelink.blob.core.windows.net
		occm_group_eastus
		privatelink.blob.core.windows.net
		occm_group_eastus2
		privatelink.blob.core.windows.net

Now the storage account (in subscription 1) has access to the Cloud Volumes ONTAP system which is running in subscription 2.

6. Retry enabling Cloud Backup on the Cloud Volumes ONTAP system and this time it should be successful.

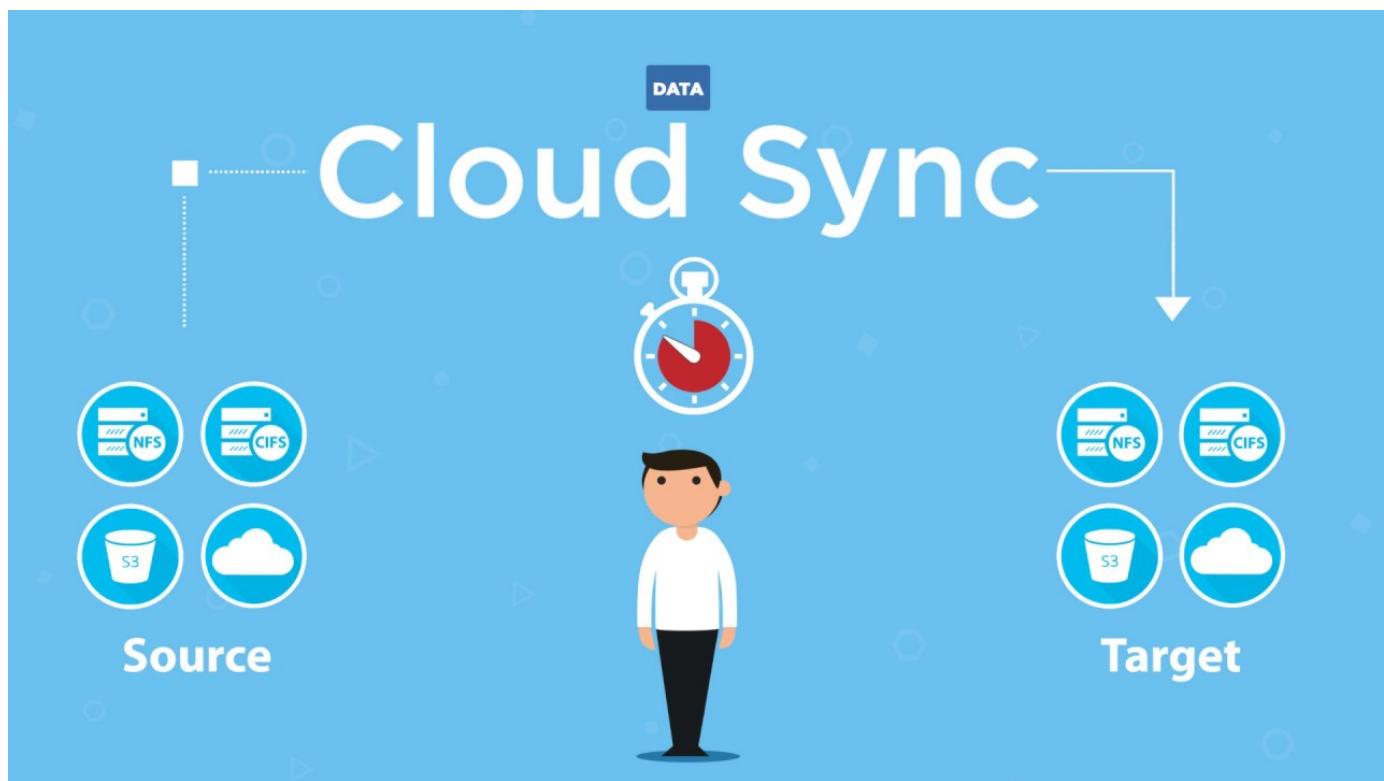
Copy and synchronize data

Cloud Sync overview

The NetApp Cloud Sync service offers a simple, secure, and automated way to migrate your data to any target, in the cloud or on your premises. Whether it's a file-based NAS dataset (NFS or SMB), Amazon Simple Storage Service (S3) object format, a NetApp StorageGRID® appliance, or any other cloud provider object store, Cloud Sync can convert and move it for you.

Features

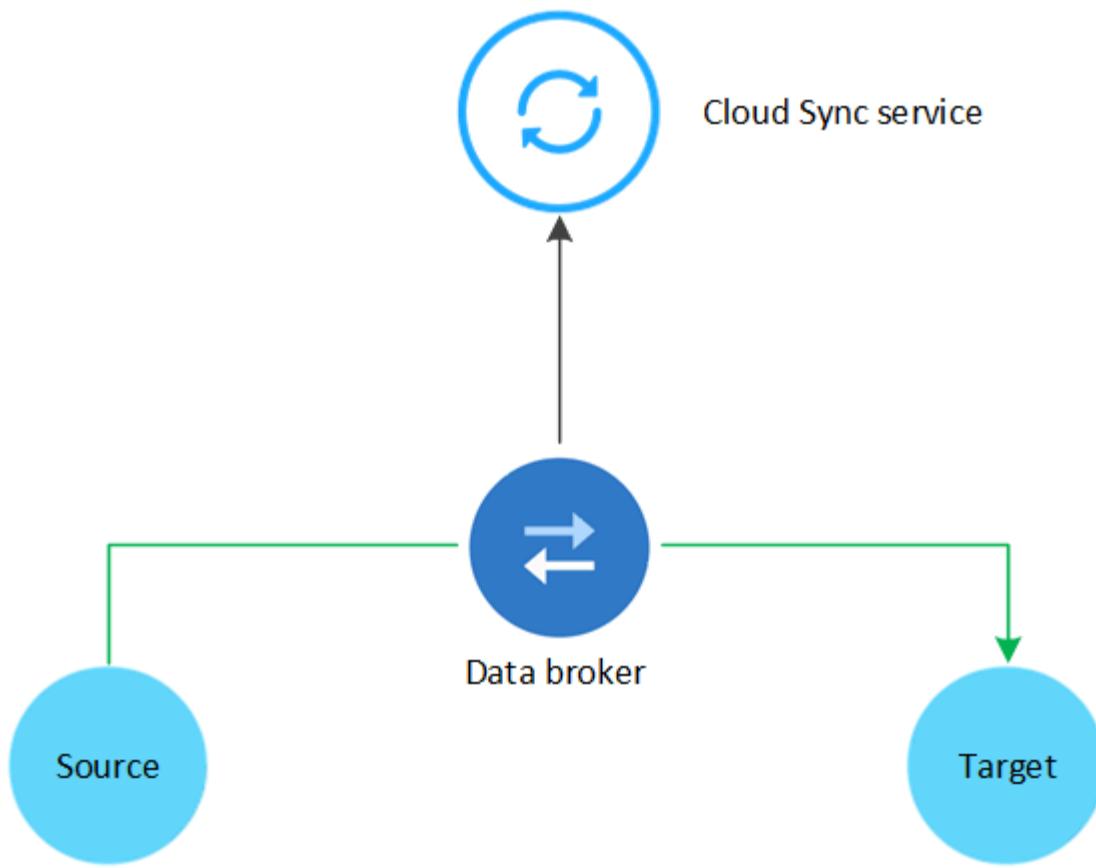
Watch the following video for an overview of Cloud Sync:



How Cloud Sync works

Cloud Sync is a software-as-a-service (SaaS) platform that consists of a data broker, a cloud-based interface available through Cloud Manager, and a source and target.

The following image shows the relationship between Cloud Sync components:



The NetApp data broker software syncs data from a source to a target (this is called a *sync relationship*). You can run the data broker in AWS, Azure, Google Cloud Platform, or on your premises. The data broker needs an outbound internet connection over port 443 so it can communicate with the Cloud Sync service and contact a few other services and repositories. [View the list of endpoints](#).

After the initial copy, the service syncs any changed data based on the schedule that you set.

Supported storage types

Cloud Sync supports the following storage types:

- Any NFS server
- Any SMB server
- Amazon EFS
- Amazon FSx for ONTAP
- Amazon S3
- Azure Blob
- Azure NetApp Files
- Cloud Volumes Service
- Cloud Volumes ONTAP
- Google Cloud Storage
- IBM Cloud Object Storage
- On-premises ONTAP cluster

- ONTAP S3 Storage
- StorageGRID

[Review the supported sync relationships.](#)

Cost

There are two types of costs associated with using Cloud Sync: resource charges and service charges.

Resource charges

Resource charges are related to the compute and storage costs for running the data broker in the cloud.

Service charges

There are two ways to pay for sync relationships after your 14-day free trial ends. The first option is to subscribe from AWS or Azure, which enables you to pay hourly or annually. The second option is to purchase licenses directly from NetApp. Read the following sections for more details.

Marketplace subscription

Subscribing to the Cloud Sync service from AWS or Azure enables you to pay at an hourly rate, or to pay annually. [You can subscribe through either AWS or Azure](#), depending on where you want to be billed.

Hourly subscriptions

With an hourly pay-as-you-go subscription, the Cloud Sync service charges hourly based on the number of sync relationships that you create.

- [View pricing in Azure](#)
- [View pay-as-you-go pricing in AWS](#)

Annual subscriptions

An annual subscription provides a license for 20 sync relationships that you pay for up front. If you go above 20 sync relationships and you've subscribed through Azure, you pay for the additional relationships by the hour.

[View annual pricing in AWS](#)

Licenses from NetApp

Another way to pay for sync relationships up front is by purchasing licenses directly from NetApp. Each license enables you to create up to 20 sync relationships.

You can use these licenses with an AWS or Azure subscription. For example, if you have 25 sync relationships, you can pay for the first 20 sync relationships using a license and then pay-as-you-go from AWS or Azure with the remaining 5 sync relationships.

[Learn how to purchase licenses and add them to Cloud Sync.](#)

License terms

Customers who purchase a Bring Your Own License (BYOL) to the Cloud Sync service should be aware of limitations associated with the license entitlement.

- Customers are entitled to leverage the BYOL license for a term not to exceed one year from the date of

delivery.

- Customers are entitled to leverage the BYOL license to establish and not to exceed a total of 20 individual connections between a source and a target (each a “sync relationship”).
- A customer’s entitlement expires at the conclusion of the one-year license term, irrespective as to whether Customer has reached the 20 sync relationship limitation.
- In the event the Customer chooses to renew its license, unused sync relationships associated from the previous license grant DO NOT roll over to the license renewal.

Data privacy

NetApp doesn’t have access to any credentials that you provide while using the Cloud Sync service. The credentials are stored directly on the data broker machine, which resides in your network.

Depending on the configuration that you choose, Cloud Sync might prompt you for credentials when you create a new relationship. For example, when setting up a relationship that includes an SMB server, or when deploying the data broker in AWS.

These credentials are always saved directly to the data broker itself. The data broker resides on a machine in your network, whether it’s on premises or in your cloud account. The credentials are never made available to NetApp.

The credentials are locally encrypted on the data broker machine using HashiCorp Vault.

Limitations

- Cloud Sync is not supported in China.
- In addition to China, the Cloud Sync data broker is not supported in the following regions:
 - AWS GovCloud (US)
 - Azure US Gov
 - Azure US DoD

Get started

Quick start for Cloud Sync

Getting started with the Cloud Sync service includes a few steps.



1 Prepare your source and target

Verify that your source and target are supported and set up. The most important requirement is to verify connectivity between the data broker and the source and target locations. [Learn more](#).



2 Prepare a location for the NetApp data broker

The NetApp data broker software syncs data from a source to a target (this is called a *sync relationship*). You can run the data broker in AWS, Azure, Google Cloud Platform, or on your premises. The data broker needs an outbound internet connection over port 443 so it can communicate with the Cloud Sync service and contact

a few other services and repositories. [View the list of endpoints.](#)

Cloud Sync guides you through the installation process when you create a sync relationship, at which point you can deploy the data broker in the cloud or download an install script for your own Linux host.

- [Review AWS installation](#)
- [Review Azure installation](#)
- [Review GCP installation](#)
- [Review Linux host installation](#)



Create your first sync relationship

Log in to [Cloud Manager](#), click **Sync**, and then drag and drop your selections for the source and target. Follow the prompts to complete the setup. [Learn more.](#)



Pay for your sync relationships after your free trial ends

Subscribe from AWS or Azure to pay-as-you-go or to pay annually. Or purchase licenses directly from NetApp. Just go to the License Settings page in Cloud Sync to set it up. [Learn more.](#)

Preparing the source and target

Prepare to sync data by verifying that your source and target are supported and setup.

Supported sync relationships

Cloud Sync enables you to sync data from a source to a target (this is called a *sync relationship*). You should understand the supported relationships before you get started.

Source location	Supported target locations
Amazon EFS	<ul style="list-style-type: none"> • Amazon EFS • Amazon FSx for ONTAP • Amazon S3 • Azure Blob • Azure NetApp Files • Cloud Volumes ONTAP • Cloud Volumes Service • Google Cloud Storage • IBM Cloud Object Storage • NFS server • On-premises ONTAP cluster • SMB server • StorageGRID
Amazon FSx for ONTAP	<ul style="list-style-type: none"> • Amazon EFS • Amazon FSx for ONTAP • Amazon S3 • Azure Blob • Azure NetApp Files • Cloud Volumes ONTAP • Cloud Volumes Service • Google Cloud Storage • IBM Cloud Object Storage • NFS server • On-premises ONTAP cluster • SMB Server • StorageGRID

Source location	Supported target locations
Amazon S3	<ul style="list-style-type: none"> • Amazon EFS • Amazon FSx for ONTAP • Amazon S3 • Azure Blob • Azure NetApp Files • Box ^{1,2} • Cloud Volumes ONTAP • Cloud Volumes Service • Google Cloud Storage • IBM Cloud Object Storage • NFS server • On-premises ONTAP cluster • SMB Server • StorageGRID
Azure Blob	<ul style="list-style-type: none"> • Amazon EFS • Amazon FSx for ONTAP • Amazon S3 • Azure Blob • Azure NetApp Files • Cloud Volumes ONTAP • Cloud Volumes Service • Google Cloud Storage • IBM Cloud Object Storage • NFS server • On-premises ONTAP cluster • SMB Server • StorageGRID

Source location	Supported target locations
Azure NetApp Files	<ul style="list-style-type: none"> • Amazon EFS • Amazon FSx for ONTAP • Amazon S3 • Azure Blob • Azure NetApp Files • Cloud Volumes ONTAP • Cloud Volumes Service • Google Cloud Storage • IBM Cloud Object Storage • NFS server • On-premises ONTAP cluster • SMB server • StorageGRID
Box ¹	<ul style="list-style-type: none"> • Amazon S3 • IBM Cloud Object Storage • NFS server • SMB Server • StorageGRID
Cloud Volumes ONTAP	<ul style="list-style-type: none"> • Amazon EFS • Amazon FSx for ONTAP • Amazon S3 • Azure Blob • Azure NetApp Files • Cloud Volumes ONTAP • Cloud Volumes Service • Google Cloud Storage • IBM Cloud Object Storage • NFS server • On-premises ONTAP cluster • SMB Server • StorageGRID

Source location	Supported target locations
Cloud Volumes Service	<ul style="list-style-type: none"> • Amazon EFS • Amazon FSx for ONTAP • Amazon S3 • Azure Blob • Azure NetApp Files • Cloud Volumes ONTAP • Cloud Volumes Service • Google Cloud Storage • IBM Cloud Object Storage • NFS server • On-premises ONTAP cluster • SMB Server • StorageGRID
Google Cloud Storage	<ul style="list-style-type: none"> • Amazon EFS • Amazon FSx for ONTAP • Amazon S3 • Azure Blob • Azure NetApp Files • Cloud Volumes ONTAP • Cloud Volumes Service • Google Cloud Storage • IBM Cloud Object Storage • NFS server • On-premises ONTAP cluster • ONTAP S3 Storage • SMB Server • StorageGRID

Source location	Supported target locations
IBM Cloud Object Storage	<ul style="list-style-type: none"> • Amazon EFS • Amazon FSx for ONTAP • Amazon S3 • Azure Blob • Azure NetApp Files • Box ^{1,2} • Cloud Volumes ONTAP • Cloud Volumes Service • Google Cloud Storage • IBM Cloud Object Storage • NFS server • On-premises ONTAP cluster • SMB Server • StorageGRID
NFS server	<ul style="list-style-type: none"> • Amazon EFS • Amazon FSx for ONTAP • Amazon S3 • Azure Blob • Azure NetApp Files • Cloud Volumes ONTAP • Cloud Volumes Service • Google Cloud Storage • IBM Cloud Object Storage • NFS server • On-premises ONTAP cluster • SMB Server • StorageGRID

Source location	Supported target locations
On-prem ONTAP cluster	<ul style="list-style-type: none"> • Amazon EFS • Amazon FSx for ONTAP • Amazon S3 • Azure Blob • Azure NetApp Files • Cloud Volumes ONTAP • Cloud Volumes Service • Google Cloud Storage • IBM Cloud Object Storage • NFS server • On-premises ONTAP cluster • SMB Server • StorageGRID
ONTAP S3 Storage	<ul style="list-style-type: none"> • Google Cloud Storage • SMB server • StorageGRID • ONTAP S3 Storage
SFTP ¹	S3
SMB server	<ul style="list-style-type: none"> • Amazon EFS • Amazon FSx for ONTAP • Amazon S3 • Azure Blob • Azure NetApp Files • Cloud Volumes ONTAP • Cloud Volumes Service • Google Cloud Storage • IBM Cloud Object Storage • NFS server • On-premises ONTAP cluster • ONTAP S3 Storage • SMB Server • StorageGRID

Source location	Supported target locations
StorageGRID	<ul style="list-style-type: none"> • Amazon EFS • Amazon FSx for ONTAP • Amazon S3 • Azure Blob • Azure NetApp Files • Box ^{1,2} • Cloud Volumes ONTAP • Cloud Volumes Service • Google Cloud Storage • IBM Cloud Object Storage • NFS server • On-premises ONTAP cluster • ONTAP S3 Storage • SMB Server • StorageGRID

Notes:

1. Sync relationships with this source/target are supported by using the Cloud Sync API only.
2. When you create a sync relationship from Amazon S3 to Box, you must use a data broker group that has a unified configuration where the following settings are set to 1:

- Scanner Concurrency
- Scanner Processes Limit
- Transferrer Concurrency
- Transferrer Processes Limit

[Learn how to define a unified configuration for a data broker group.](#)

3. You can choose a specific Azure Blob storage tier when a Blob container is the target:
 - Hot storage
 - Cool storage
4. You can choose a specific S3 storage class when Amazon S3 is the target:
 - Standard (this is the default class)
 - Intelligent-Tiering
 - Standard-Infrequent Access
 - One Zone-Infrequent Access
 - Glacier
 - Glacier Deep Archive

5. You can choose a specific storage class when a Google Cloud Storage bucket is the target:

- Standard
- Nearline
- Coldline
- Archive

Source and target requirements

Verify that your source and targets meet the following requirements.

Networking

- The source and target must have a network connection to the data broker.

For example, if an NFS server is in your data center and the data broker is in AWS, then you need a network connection (VPN or Direct Connect) from your network to the VPC.

- NetApp recommends configuring the source, target, and data broker to use a Network Time Protocol (NTP) service. The time difference between the three components should not exceed 5 minutes.

Target directory

When you create a sync relationship, Cloud Sync enables you to select an existing target directory and then optionally create a new folder inside that directory. So be sure that your preferred target directory already exists.

Permissions to read directories

In order to show every directory or folder in a source or target, Cloud Sync needs read permissions on the directory or folder.

NFS

Permissions must be defined on the source/target with uid/gid on files and directories.

Object storage

- For AWS and Google Cloud, the data broker must have list object permissions (these permissions are provided by default if you follow the data broker installation steps).
- For Azure, StorageGRID, and IBM, the credentials that you enter when setting up a sync relationship must have list object permissions.

SMB

The SMB credentials that you enter when setting up a sync relationship must have list folder permissions.



The data broker ignores the following directories by default: .snapshot, ~snapshot, .copy-offload

Amazon S3 bucket requirements

Make sure that your Amazon S3 bucket meets the following requirements.

Supported data broker locations for Amazon S3

Sync relationships that include S3 storage require a data broker deployed in AWS or on your premises. In either case, Cloud Sync prompts you to associate the data broker with an AWS account during installation.

- [Learn how to deploy the AWS data broker](#)
- [Learn how to install the data broker on a Linux host](#)

Supported AWS regions

All regions are supported except for the China and GovCloud (US) regions.

Permissions required for S3 buckets in other AWS accounts

When setting up a sync relationship, you can specify an S3 bucket that resides in an AWS account that isn't associated with the data broker.

The [permissions included in this JSON file](#) must be applied to that S3 bucket so the data broker can access it. These permissions enable the data broker to copy data to and from the bucket and to list the objects in the bucket.

Note the following about the permissions included in the JSON file:

1. <BucketName> is the name of the bucket that resides in the AWS account that isn't associated with the data broker.
2. <RoleARN> should be replaced with one of the following:
 - If the data broker was manually installed on a Linux host, *RoleARN* should be the ARN of the AWS user for which you provided AWS credentials when deploying the data broker.
 - If the data broker was deployed in AWS using the CloudFormation template, *RoleARN* should be the ARN of the IAM role created by the template.

You can find the Role ARN by going to the EC2 console, selecting the data broker instance, and clicking the IAM role from the Description tab. You should then see the Summary page in the IAM console that contains the Role ARN.

Summary Delete role

Role ARN	arn:aws:iam::123456789012:role/tanyaBroker0304-DataBrokeriamRole-1VMHWXMW3AQ05	
Role description	Edit	

Azure Blob storage requirements

Make sure that your Azure Blob storage meets the following requirements.

Supported data broker locations for Azure Blob

The data broker can reside in any location when a sync relationship includes Azure Blob storage.

Supported Azure regions

All regions are supported except for the China, US Gov, and US DoD regions.

Connection string required for relationships that include Azure Blob and NFS/SMB

When creating a sync relationship between an Azure Blob container and an NFS or SMB server, you need to provide Cloud Sync with the storage account connection string:

The screenshot shows the 'Access keys' page for an Azure Storage account. On the left, there's a navigation menu with links like Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, and Storage Explorer (preview). Below that is a 'Settings' section with 'Access keys' highlighted. The main content area displays the 'key1' key and its value, which is partially obscured. A red box highlights the 'Connection string' input field, which also contains partially obscured text. The 'Storage account name' field above it contains 'a63cde60b553020'. There are download icons next to each of these fields.

If you want to sync data between two Azure Blob containers, then the connection string must include a [shared access signature \(SAS\)](#). You also have the option to use a SAS when syncing between a Blob container and an NFS or SMB server.

The SAS must allow access to the Blob service and all resource types (Service, Container, and Object). The SAS must also include the following permissions:

- For the source Blob container: Read and List
- For the target Blob container: Read, Write, List, Add, and Create

a63cde60b553020 - Shared access signature

Storage account

Search (Ctrl+ /)

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

Storage Explorer (preview)

Settings

Access keys

CORS

Configuration

Encryption

Shared access signature

Firewalls and virtual networks

Advanced Threat Protection (pr...)

Properties

Locks

Allowed services

Blob File Queue Table

Allowed resource types

Service Container Object

Allowed permissions

Read Write Delete List Add Create Update Process

Start and expiry date/time

Start: 2018-10-23 10:07:32 AM

End: 2019-10-23 6:07:32 PM

(UTC-04:00) --- Current Time Zone ---

Allowed IP addresses

for example, 168.1.5.65 or 168.1.5.65-168.1.5.70

Allowed protocols

HTTPS only HTTPS and HTTP

Signing key

key1

Generate SAS and connection string

Azure NetApp Files requirement

Use the Premium or Ultra service level when you sync data to or from Azure NetApp Files. You might experience failures and performance issues if the disk service level is Standard.



Consult a solutions architect if you need help determining the right service level. The volume size and volume tier determines the throughput that you can get.

[Learn more about Azure NetApp Files service levels and throughput.](#)

Box requirements

Cloud Sync supports [Box](#) in a sync relationship when using the Cloud Sync API.

In order to copy from Box, you'll need to provide the following credentials:

- clientId
- clientSecret
- publicKeyId
- privateKey

- passphrase
- enterpriseId

[Learn more about using the API to create a sync relationship from Box.](#)

Google Cloud Storage bucket requirements

Make sure that your Google Cloud Storage bucket meets the following requirements.

Supported data broker locations for Google Cloud Storage

Sync relationships that include Google Cloud Storage require a data broker deployed in GCP or on your premises. Cloud Sync guides you through the data broker installation process when you create a sync relationship.

- [Learn how to deploy the GCP data broker](#)
- [Learn how to install the data broker on a Linux host](#)

Supported GCP regions

All regions are supported.

Permissions required for buckets in other Google Cloud projects

When setting up a sync relationship, you can choose from Google Cloud buckets in different projects, if you provide the required permissions to the data broker's service account. [Learn how to set up the service account.](#)

ONTAP requirements

If the sync relationship includes Cloud Volumes ONTAP or an on-prem ONTAP cluster and you selected NFSv4 or later, then you'll need to enable NFSv4 ACLs on the ONTAP system. This is required to copy the ACLs.

Permissions for a SnapMirror destination

If the source for a sync relationship is a SnapMirror destination (which is read-only), "read/list" permissions are sufficient to sync data from the source to a target.

NFS server requirements

- The NFS server can be a NetApp system or a non-NetApp system.
- The file server must allow the data broker host to access the exports.
- NFS versions 3, 4.0, 4.1, and 4.2 are supported.

The desired version must be enabled on the server.

- If you want to sync NFS data from an ONTAP system, ensure that access to the NFS export list for an SVM is enabled (`vserver nfs modify -vserver svm_name -showmount enabled`).



The default setting for showmount is *enabled* starting with ONTAP 9.2.

ONTAP S3 Storage requirements

When you set up a sync relationship that includes [ONTAP S3 Storage](#), you'll need to provide the following:

- The IP address of the LIF that's connected to ONTAP S3
- The access key and secret key that ONTAP is configured to use

SMB server requirements

- The SMB server can be a NetApp system or a non-NetApp system.
- You need to provide Cloud Sync with credentials that have permissions on the SMB server.
 - For a source SMB server, the following permissions are required: list and read.

Members of the Backup Operators group are supported with a source SMB server.

- For a target SMB server, the following permissions are required: list, read, and write.

- The file server must allow the data broker host to access the exports.
- SMB versions 1.0, 2.0, 2.1, 3.0 and 3.11 are supported.
- Grant the "Administrators" group with "Full Control" permissions to the source and target folders.

If you don't grant this permission, then the data broker might not have sufficient permissions to get the ACLs on a file or directory. If this occurs, you'll receive the following error: "getxattr error 95"

SMB limitation for hidden directories and files

An SMB limitation affects hidden directories and files when syncing data between SMB servers. If any of the directories or files on the source SMB server were hidden through Windows, the hidden attribute isn't copied to the target SMB server.

SMB sync behavior due to case-insensitivity limitation

The SMB protocol is case-insensitive, which means uppercase and lowercase letters are treated as being the same. This behavior can result in overwritten files and directory copy errors, if a sync relationship includes an SMB server and data already exists on the target.

For example, let's say that there's a file named "a" on the source and a file named "A" on the target. When Cloud Sync copies the file named "a" to the target, file "A" is overwritten by file "a" from the source.

In the case of directories, let's say that there's a directory named "b" on the source and a directory named "B" on the target. When Cloud Sync tries to copy the directory named "b" to the target, Cloud Sync receives an error that says the directory already exists. As a result, Cloud Sync always fails to copy the directory named "b."

The best way to avoid this limitation is to ensure that you sync data to an empty directory.

Networking overview for Cloud Sync

Networking for Cloud Sync includes connectivity between the data broker and the source and target locations, and an outbound internet connection from the data broker over port 443.

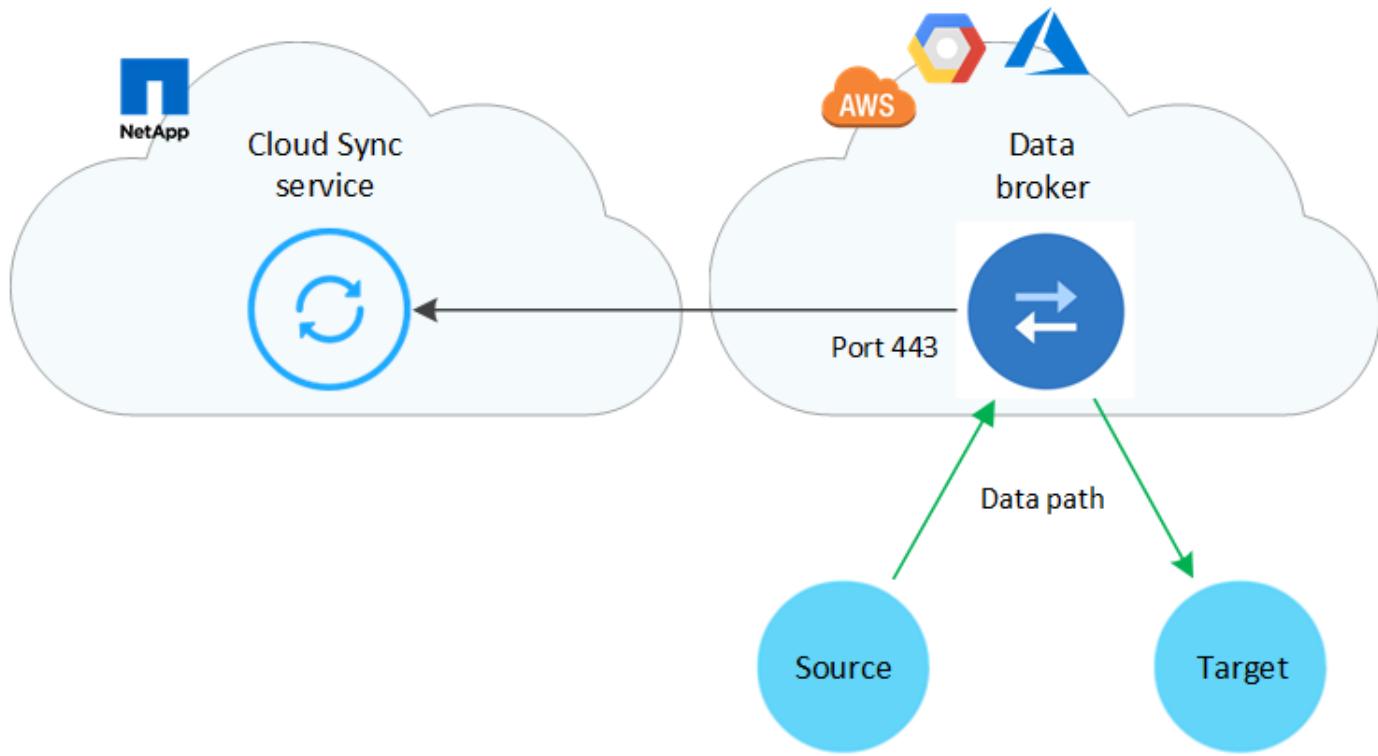
Data broker location

You can install the data broker in the cloud or on your premises.

Data broker in the cloud

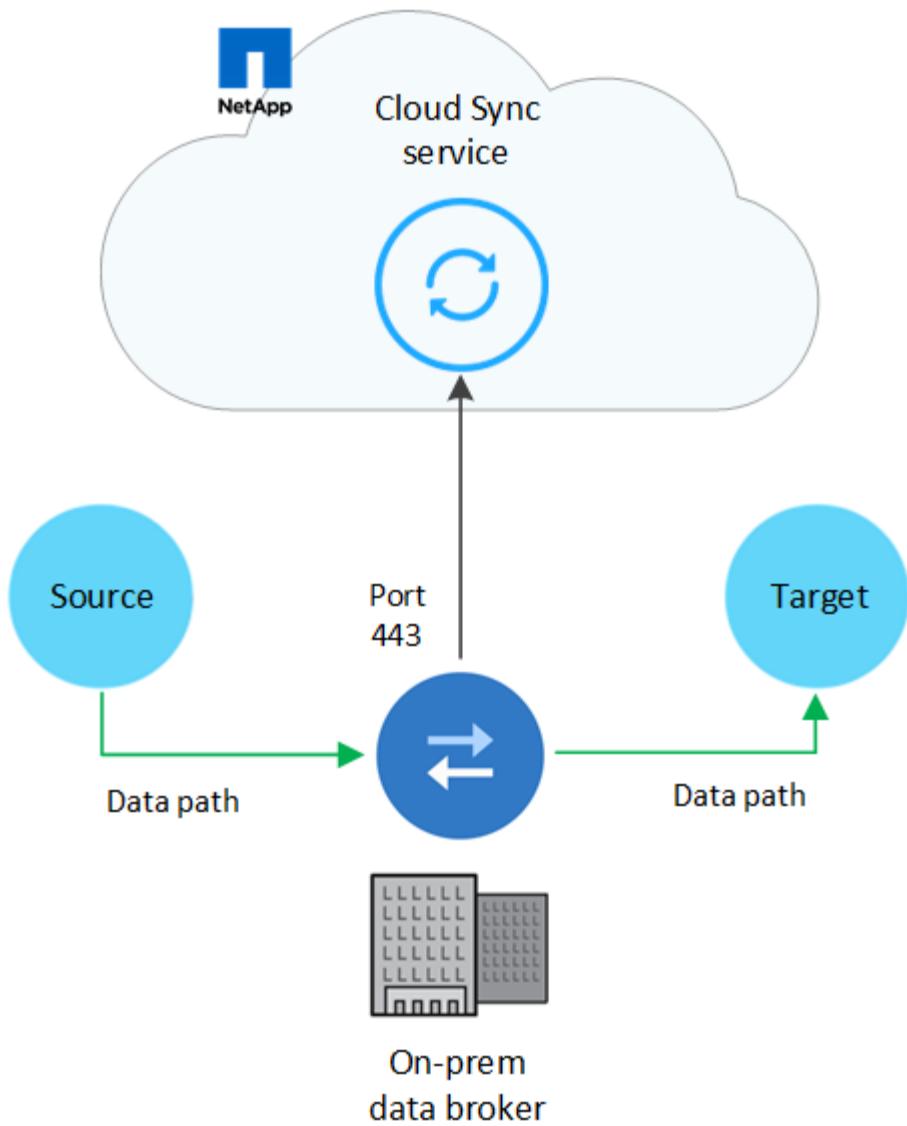
The following image shows the data broker running in the cloud, in either AWS, GCP, or Azure. The source and target can be in any location, as long as there's a connection to the data broker. For example, you might have a VPN connection from your data center to your cloud provider.

- When Cloud Sync deploys the data broker in AWS, Azure, or GCP, it creates a security group that enables the required outbound communication.



Data broker on your premises

The following image shows the data broker running on-prem, in a data center. Again, the source and target can be in any location, as long as there's a connection to the data broker.



Networking requirements

- The source and target must have a network connection to the data broker.

For example, if an NFS server is in your data center and the data broker is in AWS, then you need a network connection (VPN or Direct Connect) from your network to the VPC.

- The data broker needs an outbound internet connection so it can poll the Cloud Sync service for tasks over port 443.
- NetApp recommends configuring the source, target, and data broker to use a Network Time Protocol (NTP) service. The time difference between the three components should not exceed 5 minutes.

Networking endpoints

The NetApp data broker requires outbound internet access over port 443 to communicate with the Cloud Sync service and to contact a few other services and repositories. Your local web browser also requires access to endpoints for certain actions. If you need to limit outbound connectivity, refer to the following list of endpoints when configuring your firewall for outbound traffic.

Data broker endpoints

The data broker contacts the following endpoints:

Endpoints	Purpose
olcentgbl.trafficmanager.net:443	To contact a repository for updating CentOS packages for the data broker host. This endpoint is contacted only if you manually install the data broker on a CentOS host.
rpm.nodesource.com:443 registry.npmjs.org:443 nodejs.org:443	To contact repositories for updating Node.js, npm, and other 3rd party packages used in development.
tgz.pm2.io:443	To access a repository for updating PM2, which is a 3rd party package used to monitor Cloud Sync.
sqs.us-east-1.amazonaws.com:443 kinesis.us-east-1.amazonaws.com:443	To contact the AWS services that Cloud Sync uses for operations (queuing files, registering actions, and delivering updates to the data broker).
s3.region.amazonaws.com:443 For example: s3.us-east-2.amazonaws.com:443 See AWS documentation for a list of S3 endpoints	To contact Amazon S3 when a sync relationship includes an S3 bucket.
s3.us-east-1.amazonaws.com:443	When you download data broker logs from Cloud Sync, the data broker zips its logs directory and uploads the logs to a predefined S3 bucket in the us-east-1 region.
cf.cloudsync.netapp.com:443 repo.cloudsync.netapp.com:443	To contact the Cloud Sync service.
support.netapp.com:443	To contact NetApp support when using a BYOL license for sync relationships.
fedoraproject.org:443	To install 7z on the data broker virtual machine during installation and updates. 7z is needed to send AutoSupport messages to NetApp technical support.
sts.amazonaws.com:443	To verify AWS credentials when the data broker is deployed in AWS or when it's deployed on your premises and AWS credentials are provided. The data broker contacts this endpoint during deployment, when it's updated, and when it's restarted.

Web browser endpoints

Your web browser needs access to the following endpoint to download logs for troubleshooting purposes:

logs.cloudsync.netapp.com:443

How to install a data broker

Installing the data broker in AWS

When you create a new data broker, choose the AWS Data Broker option to deploy the data broker software on a new EC2 instance in a VPC. Cloud Sync guides you through the installation process, but the requirements and steps are repeated on this page to help you prepare for installation.

You also have the option to install the data broker on an existing Linux host in the cloud or on your premises.

[Learn more.](#)

Supported AWS regions

All regions are supported except for the China and GovCloud (US) regions.

Networking requirements

- The data broker needs an outbound internet connection so it can poll the Cloud Sync service for tasks over port 443.

When Cloud Sync deploys the data broker in AWS, it creates a security group that enables the required outbound communication. Note that you can configure the data broker to use a proxy server during the installation process.

If you need to limit outbound connectivity, see [the list of endpoints that the data broker contacts](#).

- NetApp recommends configuring the source, target, and data broker to use a Network Time Protocol (NTP) service. The time difference between the three components should not exceed 5 minutes.

Permissions required to deploy the data broker in AWS

The AWS user account that you use to deploy the data broker must have the permissions included in [this NetApp-provided policy](#).

Requirements to use your own IAM role with the AWS data broker

When Cloud Sync deploys the data broker, it creates an IAM role for the data broker instance. You can deploy the data broker using your own IAM role, if you prefer. You might use this option if your organization has strict security policies.

The IAM role must meet the following requirements:

- The EC2 service must be allowed to assume the IAM role as a trusted entity.
- [The permissions defined in this JSON file](#) must be attached to the IAM role so the data broker can function properly.

Follow the steps below to specify the IAM role when deploying the data broker.

Installing the data broker

You can install a data broker in AWS when you create a sync relationship.

Steps

- Click **Create New Sync**.

2. On the **Define Sync Relationship** page, choose a source and target and click **Continue**.

Complete the steps until you reach the **Data Broker** page.

3. On the **Data Broker** page, click **Create Data Broker** and then select **Amazon Web Services**.

If you already have a data broker, you'll need to click the  icon first.

4. Enter a name for the data broker and click **Continue**.

5. Enter an AWS access key so Cloud Sync can create the data broker in AWS on your behalf.

The keys aren't saved or used for any other purposes.

If you'd rather not provide access keys, click the link at the bottom of the page to use a CloudFormation template instead. When you use this option, you don't need to provide credentials because you are logging in directly to AWS.

The following video shows how to launch the data broker instance using a CloudFormation template:

► https://docs.netapp.com/us-en/occm//media/video_cloud_sync.mp4 (video)

6. If you entered an AWS access key, select a location for the instance, select a key pair, choose whether to enable a public IP address, and then select an existing IAM role, or leave the field blank so Cloud Sync creates the role for you.

If you choose your own IAM role, [you'll need to provide the required permissions](#).

Basic Settings

Location Region <input type="text" value="US West Oregon"/> VPC <input type="text" value="vpc-3c46c059 - 10.60.21.0/25"/> Subnet <input type="text" value="10.60.21.0/25"/>	Connectivity Key Pair <input type="text" value="newKey"/> Enable Public IP? <input checked="" type="radio"/> Enable <input type="radio"/> Disable IAM Role (optional) 
--	---

7. Specify a proxy configuration, if a proxy is required for internet access in the VPC.
8. After the data broker is available, click **Continue** in Cloud Sync.

The following image shows a successfully deployed instance in AWS:

The screenshot shows a search results page titled 'Select a NetApp Data Broker'. At the top left, there is a blue circular icon with a white checkmark and the number '1'. To its right, it says 'NetApp Data Brokers'. On the far right is a magnifying glass icon. Below the header, there is a table with one row. The first column contains the AWS logo and the word 'aws' followed by a placeholder 'name'. The second column lists 'US West (Oregon)' under 'Region' and 'us-west-2c' under 'Availability Zone'. The third column lists '10.60.21.0/25 | vpc-3c46c059' under 'VPC' and '10.60.21.0/25 | subnet-e7f526be' under 'Subnet'. The fourth column lists '10.60.21.5' under 'Private IP' and 'i-0fc5c97e2f5f22c20' under 'Instance ID'. The fifth column shows a green circular icon with a checkmark and the word 'Active'. The table has a thin blue border.

9. Complete the pages in the wizard to create the new sync relationship.

Result

You have deployed a data broker in AWS and created a new sync relationship. You can use this data broker with additional sync relationships.

Details about the data broker instance

Cloud Sync creates a data broker in AWS using the following configuration.

Instance type

m5n.xlarge when available in the region, otherwise m5.xlarge

vCPUs

4

RAM

16 GB

Operating system

Amazon Linux 2

Disk size and type

10 GB GP2 SSD

Installing the data broker in Azure

When you create a new data broker, choose the Azure Data Broker option to deploy the data broker software on a new virtual machine in a VNet. Cloud Sync guides you through the installation process, but the requirements and steps are repeated on this page to help you prepare for installation.

You also have the option to install the data broker on an existing Linux host in the cloud or on your premises. [Learn more.](#)

Supported Azure regions

All regions are supported except for the China, US Gov, and US DoD regions.

Networking requirements

- The data broker needs an outbound internet connection so it can poll the Cloud Sync service for tasks over port 443.

When Cloud Sync deploys the data broker in Azure, it creates a security group that enables the required outbound communication.

If you need to limit outbound connectivity, see [the list of endpoints that the data broker contacts](#).

- NetApp recommends configuring the source, target, and data broker to use a Network Time Protocol (NTP) service. The time difference between the three components should not exceed 5 minutes.

Authentication method

When you deploy the data broker, you'll need to choose an authentication method: a password or an SSH public-private key pair.

For help with creating a key pair, refer to [Azure Documentation: Create and use an SSH public-private key pair for Linux VMs in Azure](#).

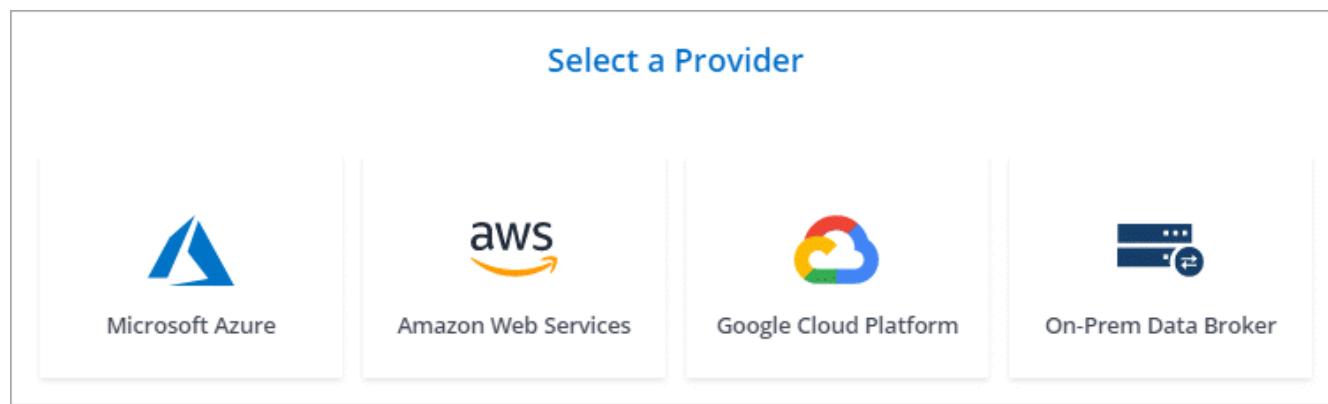
Installing the data broker

You can install a data broker in Azure when you create a sync relationship.

Steps

- Click **Create New Sync**.
- On the **Define Sync Relationship** page, choose a source and target and click **Continue**.
Complete the pages until you reach the **Data Broker** page.
- On the **Data Broker** page, click **Create Data Broker** and then select **Microsoft Azure**.

If you already have a data broker, you'll need to click the  icon first.



- Enter a name for the data broker and click **Continue**.
- If you're prompted, log in to your Microsoft account. If you're not prompted, click **Log in to Azure**.

The form is owned and hosted by Microsoft. Your credentials are not provided to NetApp.

- Choose a location for the data broker and enter basic details about the virtual machine.

Location	Virtual Machine
Subscription <div style="border: 1px solid #ccc; padding: 5px; width: fit-content;">OCCM Dev</div>	VM Name <div style="border: 1px solid #ccc; padding: 5px; width: fit-content;">netappdatabroker</div>
Azure Region <div style="border: 1px solid #ccc; padding: 5px; width: fit-content;">West US 2</div>	User Name <div style="border: 1px solid #ccc; padding: 5px; width: fit-content;">databroker</div>
VNet <div style="border: 1px solid #ccc; padding: 5px; width: fit-content;">Vnet1</div>	Authentication Method: <input checked="" type="radio"/> Password <input type="radio"/> Public Key
Subnet <div style="border: 1px solid #ccc; padding: 5px; width: fit-content;">Subnet1</div>	Enter Password <div style="border: 1px solid #ccc; padding: 5px; width: fit-content;">*****</div>
	Resource Group: <input checked="" type="radio"/> Generate a new group <input type="radio"/> Use an existing group

7. Specify a proxy configuration, if a proxy is required for internet access in the VNet.
8. Click **Continue** and keep the page open until the deployment is complete.

The process can take up to 7 minutes.

9. In Cloud Sync, click **Continue** once the data broker is available.
10. Complete the pages in the wizard to create the new sync relationship.

Result

You have deployed a data broker in Azure and created a new sync relationship. You can use this data broker with additional sync relationships.

Getting a message about needing admin consent?

If Microsoft notifies you that admin approval is required because Cloud Sync needs permission to access resources in your organization on your behalf, then you have two options:

1. Ask your AD admin to provide you with the following permission:

In Azure, go to **Admin Centers > Azure AD > Users and Groups > User Settings** and enable **Users can consent to apps accessing company data on their behalf**.

2. Ask your AD admin to consent on your behalf to **CloudSync-AzureDataBrokerCreator** using the following URL (this is the admin consent endpoint):

```
https://login.microsoftonline.com/{FILL HERE YOUR TENANT ID}/v2.0/adminconsent?client_id=8ee4ca3a-bafa-4831-97cc-5a38923cab85&redirect_uri=https://cloudsync.netapp.com&scope=https://management.azure.com/user_impersonationhttps://graph.microsoft.com/User.Read
```

As shown in the URL, our app URL is <https://cloudsync.netapp.com> and the application client ID is 8ee4ca3a-bafa-4831-97cc-5a38923cab85.

Details about the data broker VM

Cloud Sync creates a data broker in Azure using the following configuration.

VM type

Standard DS4 v2

vCPUs

8

RAM

28 GB

Operating system

CentOS 7.7

Disk size and type

64 GB Premium SSD

Installing the data broker in Google Cloud Platform

When you create a new data broker, choose the GCP Data Broker option to deploy the data broker software on a new virtual machine instance in a VPC. Cloud Sync guides you through the installation process, but the requirements and steps are repeated on this page to help you prepare for installation.

You also have the option to install the data broker on an existing Linux host in the cloud or on your premises. [Learn more](#).

Supported GCP regions

All regions are supported.

Networking requirements

- The data broker needs an outbound internet connection so it can poll the Cloud Sync service for tasks over port 443.

When Cloud Sync deploys the data broker in GCP, it creates a security group that enables the required outbound communication.

If you need to limit outbound connectivity, see [the list of endpoints that the data broker contacts](#).

- NetApp recommends configuring the source, target, and data broker to use a Network Time Protocol (NTP) service. The time difference between the three components should not exceed 5 minutes.

Permissions required to deploy the data broker in GCP

Ensure that the GCP user who deploys the data broker has the following permissions:

- `compute.networks.list`
- `compute.regions.list`
- `deploymentmanager.deployments.create`
- `deploymentmanager.deployments.delete`
- `deploymentmanager.operations.get`
- `iam.serviceAccounts.list`

Permissions required for the service account

When you deploy the data broker, you need to select a service account that has the following permissions:

- `logging.logEntries.create`
- `resourcemanager.projects.get`
- `storage.buckets.get`
- `storage.buckets.list`
- `storage.objects.*`
- `iam.serviceAccounts.signJwt`



The "iam.serviceAccounts.signJwt" permission is required only if you're planning to set up the data broker to use an external HashiCorp vault.

Installing the data broker

You can install a data broker in GCP when you create a sync relationship.

Steps

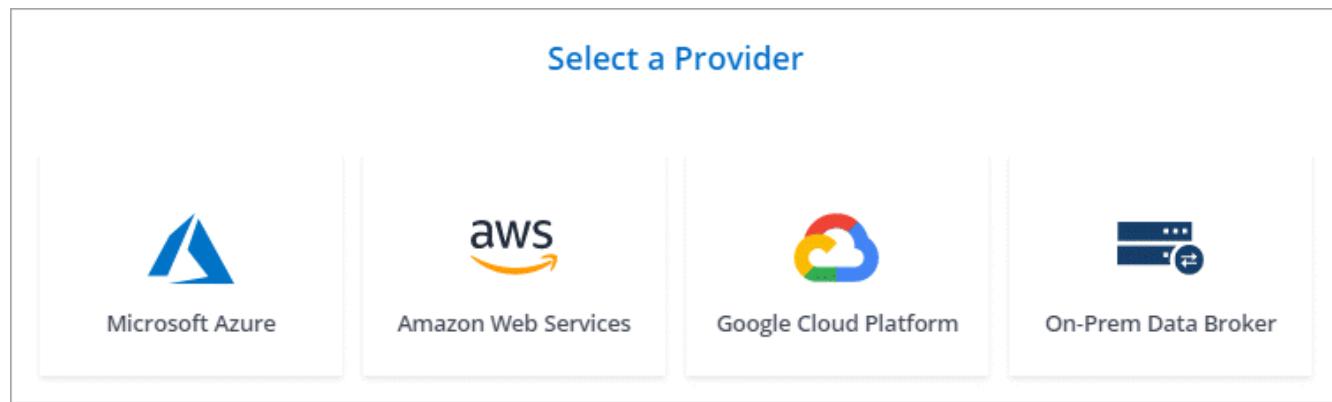
1. Click **Create New Sync**.

2. On the **Define Sync Relationship** page, choose a source and target and click **Continue**.

Complete the steps until you reach the **Data Broker** page.

3. On the **Select a Data Broker** page, click the **Create Data Broker** and then select **Google Cloud Platform**.

If you already have a data broker, you'll need to click the  icon first.



4. Enter a name for the data broker and click **Continue**.

5. If you're prompted, log in with your Google account.

The form is owned and hosted by Google. Your credentials are not provided to NetApp.

6. Select a project and service account and then choose a location for the data broker, including whether you want to enable or disable a public IP address.

If you don't enable a public IP address, then you'll need to define a proxy server in the next step.

Basic Settings

<p>Project</p> <p>Project</p> <div style="border: 1px solid #ccc; padding: 5px; width: 100%;">OCCM-Dev</div>	<p>Location</p> <p>Region</p> <div style="border: 1px solid #ccc; padding: 5px; width: 100%;">us-west1</div>
<p>Service Account</p> <div style="border: 1px solid #ccc; padding: 5px; width: 100%;">test</div>	<p>Zone</p> <div style="border: 1px solid #ccc; padding: 5px; width: 100%;">us-west1-a</div>
<p>Select a Service Account that includes these permissions</p>	
<p>VPC</p> <div style="border: 1px solid #ccc; padding: 5px; width: 100%;">default</div>	
<p>Subnet</p> <div style="border: 1px solid #ccc; padding: 5px; width: 100%;">default</div>	
<p>Public IP</p> <div style="border: 1px solid #ccc; padding: 5px; width: 100%;">Enable</div>	

7. Specify a proxy configuration, if a proxy is required for internet access in the VPC.

If a proxy is required for internet access, then the proxy must be in Google Cloud and use the same service account as the data broker.

8. Once the data broker is available, click **Continue** in Cloud Sync.

The instance takes approximately 5 to 10 minutes to deploy. You can monitor the progress from the Cloud Sync service, which automatically refreshes when the instance is available.

9. Complete the pages in the wizard to create the new sync relationship.

Result

You've deployed a data broker in GCP and created a new sync relationship. You can use this data broker with additional sync relationships.

Providing permissions to use buckets in other Google Cloud projects

When you create a sync relationship and choose Google Cloud Storage as the source or target, Cloud Sync enables you to choose from the buckets that the data broker's service account has permissions to use. By default, this includes the buckets that are in the *same* project as the data broker service account. But you can choose buckets from *other* projects if you provide the required permissions.

Steps

1. Open the Google Cloud Platform console and load the Cloud Storage service.

2. Click the name of the bucket that you'd like to use as a source or target in a sync relationship.
3. Click **Permissions**.
4. Click **Add**.
5. Enter the name of the data broker's service account.
6. Select a role that provides [the same permissions as shown above](#).
7. Click **Save**.

Result

When you set up a sync relationship, you can now choose that bucket as the source or target in the sync relationship.

Details about the data broker VM instance

Cloud Sync creates a data broker in Google Cloud using the following configuration.

Machine type

n1-standard-4

vCPUs

4

RAM

15 GB

Operating system

Red Hat Enterprise Linux 7.7

Disk size and type

20 GB HDD pd-standard

Installing the data broker on a Linux host

When you create a new data broker, choose the On-Prem Data Broker option to install the data broker software on an on-premises Linux host, or on an existing Linux host in the cloud. Cloud Sync guides you through the installation process, but the requirements and steps are repeated on this page to help you prepare for installation.

Linux host requirements

- **Operating system:**

- CentOS 7.0, 7.7, and 8.0

- CentOS Stream is not supported.

- Red Hat Enterprise Linux 7.7 and 8.0

- Ubuntu Server 20.04 LTS

- SUSE Linux Enterprise Server 15 SP1

The command `yum update all` must be run on the host before you install the data broker.

A Red Hat Enterprise Linux system must be registered with Red Hat Subscription Management. If it is not registered, the system cannot access repositories to update required 3rd party software during installation.

- **RAM:** 16 GB
- **CPU:** 4 cores
- **Free disk space:** 10 GB
- **SELinux:** We recommend that you disable [SELinux](#) on the host.

SELinux enforces a policy that blocks data broker software updates and can block the data broker from contacting endpoints required for normal operation.

- **OpenSSL:** OpenSSL must be installed on the Linux host.

Networking requirements

- The Linux host must have a connection to the source and target.
- The file server must allow the Linux host to access the exports.
- Port 443 must be open on the Linux host for outbound traffic to AWS (the data broker constantly communicates with the Amazon SQS service).
- NetApp recommends configuring the source, target, and data broker to use a Network Time Protocol (NTP) service. The time difference between the three components should not exceed 5 minutes.

Enabling access to AWS

If you plan to use the data broker with a sync relationship that includes an S3 bucket, then you should prepare the Linux host for AWS access. When you install the data broker, you'll need to provide AWS keys for an AWS user that has programmatic access and specific permissions.

Steps

1. Create an IAM policy using [this NetApp-provided policy](#). [View AWS instructions](#).
2. Create an IAM user that has programmatic access. [View AWS instructions](#).

Be sure to copy the AWS keys because you need to specify them when you install the data broker software.

Enabling access to Google Cloud

If you plan to use the data broker with a sync relationship that includes a Google Cloud Storage bucket, then you should prepare the Linux host for GCP access. When you install the data broker, you'll need to provide a key for a service account that has specific permissions.

Steps

1. Create a GCP service account that has Storage Admin permissions, if you don't already have one.
2. Create a service account key saved in JSON format. [View GCP instructions](#).

The file should contain at least the following properties: "project_id", "private_key", and "client_email"



When you create a key, the file gets generated and downloaded to your machine.

3. Save the JSON file to the Linux host.

Enabling access to Microsoft Azure

Access to Azure is defined per relationship by providing a storage account and a connection string in the Sync Relationship wizard.

Installing the data broker

You can install a data broker on a Linux host when you create a sync relationship.

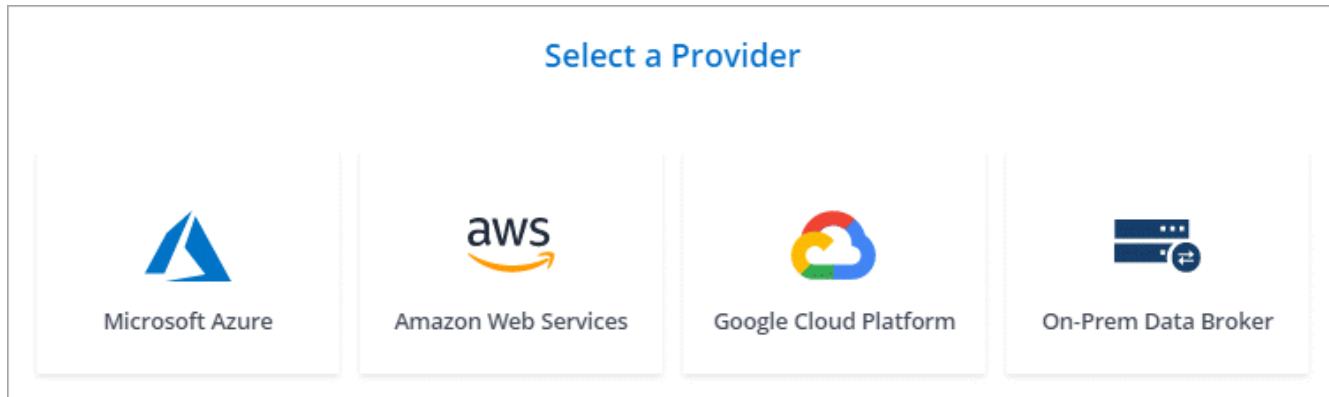
Steps

1. Click **Create New Sync**.
2. On the **Define Sync Relationship** page, choose a source and target and click **Continue**.

Complete the steps until you reach the **Data Broker** page.

3. On the **Data Broker** page, click **Create Data Broker** and then select **On-Prem Data Broker**.

If you already have a data broker, you'll need to click the  icon first.



Even though the option is labeled **On-Prem Data Broker**, it applies to a Linux host on your premises or in the cloud.

4. Enter a name for the data broker and click **Continue**.

The instructions page loads shortly. You'll need to follow these instructions—they include a unique link to download the installer.

5. On the instructions page:
 - Select whether to enable access to **AWS**, **Google Cloud**, or both.
 - Select an installation option: **No proxy**, **Use proxy server**, or **Use proxy server with authentication**.
 - Use the commands to download and install the data broker.

The following steps provide details about each possible installation option. Follow the instructions page to get the exact command based on your installation option.

- d. Download the installer:
 - No proxy:

```
curl <URI> -o data_broker_installer.sh
```

- Use proxy server:

```
curl <URI> -o data_broker_installer.sh -x <proxy_host>:<proxy_port>
```

- Use proxy server with authentication:

```
curl <URI> -o data_broker_installer.sh -x  
<proxy_username>:<proxy_password>@<proxy_host>:<proxy_port>
```

URI

Cloud Sync displays the URI of the installation file on the instructions page, which loads when you follow the prompts to deploy the On-Prem Data Broker. That URI isn't repeated here because the link is generated dynamically and can be used only once. [Follow these steps to obtain the URI from Cloud Sync.](#)

- Switch to superuser, make the installer executable and install the software:



Each command listed below includes parameters for AWS access and GCP access. Follow the instructions page to get the exact command based on your installation option.

- No proxy configuration:

```
sudo -s  
chmod +x data_broker_installer.sh  
./data_broker_installer.sh -a <aws_access_key> -s <aws_secret_key> -g  
<absolute_path_to_the_json_file>
```

- Proxy configuration:

```
sudo -s  
chmod +x data_broker_installer.sh  
./data_broker_installer.sh -a <aws_access_key> -s <aws_secret_key> -g  
<absolute_path_to_the_json_file> -h <proxy_host> -p <proxy_port>
```

- Proxy configuration with authentication:

```
sudo -s  
chmod +x data_broker_installer.sh  
./data_broker_installer.sh -a <aws_access_key> -s <aws_secret_key> -g  
<absolute_path_to_the_json_file> -h <proxy_host> -p <proxy_port> -u  
<proxy_username> -w <proxy_password>
```

AWS keys

These are the keys for the user that you should have prepared [following these steps](#). The AWS keys are stored on the data broker, which runs in your on-premises or cloud network. NetApp doesn't use the keys outside of the data broker.

JSON file

This is the JSON file that contains a service account key that you should have prepared [following these steps](#).

6. Once the data broker is available, click **Continue** in Cloud Sync.
7. Complete the pages in the wizard to create the new sync relationship.

Create sync relationships

When you create a sync relationship, the Cloud Sync service copies files from the source to the target. After the initial copy, the service syncs any changed data every 24 hours.

Before you can create some types of sync relationships, you'll first need to create a working environment in Cloud Manager.

Create sync relationships for specific types of working environment

If you want to create sync relationships for any of the following, then you first need to create or discover the working environment:

- Amazon FSx for ONTAP
- Azure NetApp Files
- Cloud Volumes ONTAP
- On-prem ONTAP clusters

Steps

1. Create or discover the working environment.
 - [Create an Amazon FSx for ONTAP working environment](#)
 - [Setting up and discovering Azure NetApp Files](#)
 - [Launching Cloud Volumes ONTAP in AWS](#)
 - [Launching Cloud Volumes ONTAP in Azure](#)
 - [Launching Cloud Volumes ONTAP in GCP](#)
 - [Adding existing Cloud Volumes ONTAP systems](#)
 - [Discovering ONTAP clusters](#)
2. Click **Canvas**.
3. Select a working environment that matches any of the types listed above.
4. Select the action menu next to Sync.

The screenshot shows the Cloud Manager interface. On the left, under 'Canvas', there is a 'CloudVolumesONTAP' entry with a 'SINGLE' icon, showing 'Cloud Volumes ONTAP' and '51 GiB Capacity'. To its right is an 'aws' icon with an orange arrow pointing towards it. On the right side, there is a detailed view of the 'CloudVolumesONTAP' location. It shows the status 'On' with a green checkmark. Below this are four configuration items: 'Compliance' (Off, with an 'Enable' button), 'Monitoring' (On), 'File Cache' (Off), and 'Sync' (On). The 'Sync' section also displays '559.16TiB Data Synced'. A hand cursor is hovering over the three-dot menu icon next to the 'Sync' section.

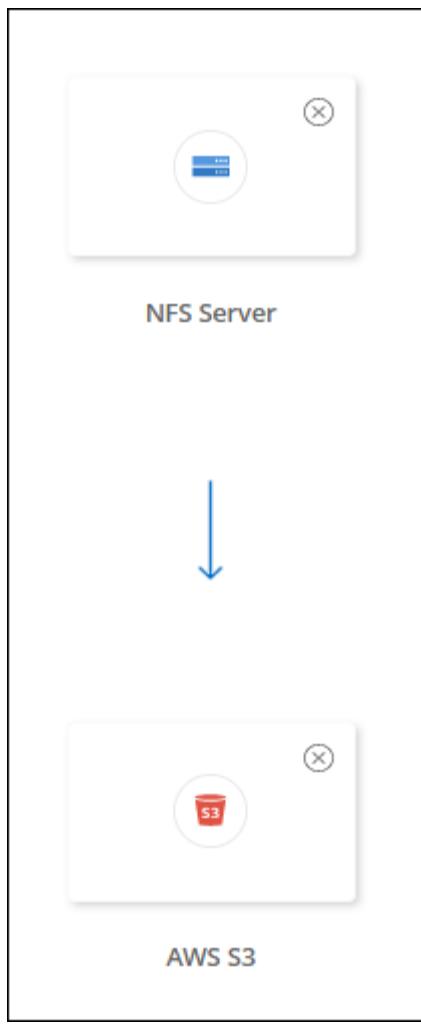
5. Select **Sync data from this location** or **Sync data to this location** and follow the prompts to set up the sync relationship.

Create other types of sync relationships

Use these steps to sync data to or from a supported storage type other than Amazon FSx for ONTAP, Azure NetApp Files, Cloud Volumes ONTAP, or on-prem ONTAP clusters. The steps below provide an example that shows how to set up a sync relationship from an NFS server to an S3 bucket.

1. In Cloud Manager, click **Sync**.
2. On the **Define Sync Relationship** page, choose a source and target.

The following steps provide an example of how to create a sync relationship from an NFS server to an S3 bucket.



3. On the **NFS Server** page, enter the IP address or fully qualified domain name of the NFS server that you want to sync to AWS.
4. On the **Data Broker** page, follow the prompts to create a data broker virtual machine in AWS, Azure, or Google Cloud Platform, or to install the data broker software on an existing Linux host.

For more details, refer to the following pages:

- [Installing the data broker in AWS](#)
- [Installing the data broker in Azure](#)
- [Installing the data broker in GCP](#)
- [Installing the data broker on a Linux host](#)

5. After you install the data broker, click **Continue**.

The following image shows a successfully deployed data broker in AWS:

- On the **Directories** page, select a top-level directory or subdirectory.

If Cloud Sync is unable to retrieve the exports, click **Add Export Manually** and enter the name of an NFS export.



If you want to sync more than one directory on the NFS server, then you must create additional sync relationships after you are done.

- On the **AWS S3 Bucket** page, select a bucket:

- Drill down to select an existing folder within the bucket or to select a new folder that you create inside the bucket.
- Click **Add to the list** to select an S3 bucket that is not associated with your AWS account. [Specific permissions must be applied to the S3 bucket.](#)

- On the **Bucket Setup** page, set up the bucket:

- Choose whether to enable S3 bucket encryption and then select an AWS KMS key, enter the ARN of a KMS key, or select AES-256 encryption.
- Select an S3 storage class. [View the supported storage classes.](#)

- On the **Settings** page, define how source files and folders are synced and maintained in the target

location:

Schedule

Choose a recurring schedule for future syncs or turn off the sync schedule. You can schedule a relationship to sync data as often as every 1 minute.

Retries

Define the number of times that Cloud Sync should retry to sync a file before skipping it.

Compare By

Choose whether Cloud Sync should compare certain attributes when determining whether a file or directory has changed and should be synced again.

Even if you uncheck these attributes, Cloud Sync still compares the source to the target by checking the paths, file sizes, and file names. If there are any changes, then it syncs those files and directories.

You can choose to enable or disable Cloud Sync from comparing the following attributes:

- **mtime**: The last modified time for a file. This attribute isn't valid for directories.
- **uid**, **gid**, and **mode**: Permission flags for Linux.

Copy for Objects

Enable this option to copy object storage metadata and tags. If a user changes the metadata on the source, Cloud Sync copies this object in the next sync, but if a user changes the tags on the source (and not the data itself), Cloud Sync doesn't copy the object in the next sync.

You can't edit this option after you create the relationship.

Copying tags is supported with sync relationships that include an S3-compatible endpoint (S3, StorageGRID, or IBM Cloud Object Storage).

Copying metadata is supported with "cloud-to-cloud" relationships between any of the following endpoints:

- AWS S3
- Azure Blob
- Google Cloud Storage
- IBM Cloud Object Storage
- StorageGRID

Recently Modified Files

Choose to exclude files that were recently modified prior to the scheduled sync.

Delete Files on Source

Choose to delete files from the source location after Cloud Sync copies the files to the target location. This option includes the risk of data loss because the source files are deleted after they're copied.

If you enable this option, you also need to change a parameter in the local.json file on the data broker. Open the file and change the parameter named `workers.transferrer.delete-on-source` to `true`.

Delete Files on Target

Choose to delete files from the target location, if they were deleted from the source. The default is to never deletes files from the target location.

File Types

Define the file types to include in each sync: files, directories, and symbolic links.

Exclude File Extensions

Specify file extensions to exclude from the sync by typing the file extension and pressing **Enter**. For example, type *log* or *.log* to exclude *.log files. A separator isn't required for multiple extensions. The following video provides a short demo:

► [https://docs.netapp.com/us-en/occm//media/video_file_extensions.mp4 \(video\)](https://docs.netapp.com/us-en/occm//media/video_file_extensions.mp4)

File Size

Choose to sync all files regardless of their size or just files that are in a specific size range.

Date Modified

Choose all files regardless of their last modified date, files modified after a specific date, before a specific date, or between a time range.

10. On the **Tags/Metadata** page, choose whether to save a key-value pair as a tag on all files transferred to the S3 bucket or to assign a metadata key-value pair on all files.

The screenshot shows the AWS Cloud Sync configuration interface. At the top, there are tabs: AWS S3 Bucket, Settings, Tags/Metadata (which is selected), and Review. Below the tabs, the section title is "Relationship Tags". It says "Cloud Sync assigns the relationship tags to all of the files transferred to the S3 bucket. This enables you to search for the transferred files by using the tag values." There are two radio button options: "Save on Object's Tags" (selected) and "Save On Object's Metadata". Below these are fields for "Tag Key" (Up to 128 characters) and "Tag Value" (Up to 256 characters). At the bottom left is a "Add Relationship Tag" button, and at the bottom right is an "Optional Field | [Up to 5]" note.



This same feature is available when syncing data to StorageGRID and IBM Cloud Object Storage. For Azure and Google Cloud Storage, only the metadata option is available.

11. Review the details of the sync relationship and then click **Create Relationship**.

Result

Cloud Sync starts syncing data between the source and target.

Paying for sync relationships after your free trial ends

There are two ways to pay for sync relationships after your 14-day free trial ends. The first option is to

subscribe from AWS or Azure to pay-as-you-go or to pay annually. The second option is to purchase licenses directly from NetApp.

You can subscribe from either the AWS Marketplace or the Azure Marketplace. You can't subscribe from both.

You have the option to use licenses from NetApp with a Marketplace subscription. For example, if you have 25 sync relationships, you can pay for the first 20 sync relationships using a license and then pay-as-you-go from AWS or Azure with the remaining 5 sync relationships.

[Learn more about how licenses work.](#)

What if I don't immediately pay after my free trial ends?

You won't be able to create any additional relationships. Existing relationships are not deleted, but you cannot make any changes to them until you subscribe or enter a license.

Subscribing from AWS

AWS enables you to pay-as-you-go or to pay annually.

Steps to pay-as-you-go

1. Click **Sync > Licensing**.
2. Select **AWS**
3. Click **Subscribe** and then click **Continue**.
4. Subscribe from the AWS Marketplace, and then log back in to the Cloud Sync service to complete the registration.

The following video shows the process:

► https://docs.netapp.com/us-en/occm//media/video_cloud_sync_registering.mp4 (video)

Steps to pay annually

1. [Go to the AWS Marketplace page](#).
2. Click **Continue to Subscribe**.
3. Select your contract options and click **Create contract**.

Subscribing from Azure

Azure enables you to pay-as-you-go or to pay annually.

What you'll need

An Azure user account that has Contributor or Owner permissions in the relevant subscription.

Steps

1. Click **Sync > Licensing**.
2. Select **Azure**.
3. Click **Subscribe** and then click **Continue**.
4. In the Azure portal, click **Create**, select your options, and click **Subscribe**.

Select **Monthly** to pay by the hour, or **Yearly** to pay for a year up front.

5. When deployment is complete, click the name of the SaaS resource in the notification pop-up.
6. Click **Configure Account** to return to Cloud Sync.

The following video shows the process:

► https://docs.netapp.com/us-en/occm//media/video_cloud_sync_registering_azure.mp4 (video)

Purchasing licenses from NetApp and adding them to Cloud Sync

To pay for your sync relationships up front, you must purchase one or more licenses and add them to the Cloud Sync service.

What you'll need

You'll need the serial number for your license and the user name and password for the NetApp Support Site account that the license is associated with.

Steps

1. Purchase a license by [contacting NetApp](#).
2. In Cloud Manager, click **Sync > Licensing**.
3. Click **Add License** and add the required information:
 - a. Enter the serial number.
 - b. Select the NetApp Support Site account that is associated with the license that you're adding:
 - If your account was already added to Cloud Manager, select it from the drop-down list.
 - If your account wasn't added yet, click **Add NSS Credentials**, enter the user name and password, click **Register**, and then select it from the drop-down list.
 - c. Click **Add**.

Tutorials

Copying ACLs between SMB shares

Cloud Sync can copy access control lists (ACLs) between a source SMB share and a target SMB share. If needed, you can manually preserve the ACLs yourself by using robocopy.

Choices

- [Set up Cloud Sync to automatically copy ACLs](#)
- [Manually copy the ACLs yourself](#)

Setting up Cloud Sync to copy ACLs between SMB servers

Copy ACLs between SMB servers by enabling a setting when you create a relationship or after you create a relationship.

Note that this feature is available for new sync relationships created after the 23 Feb 2020 release. If you'd like to use this feature with existing relationships created prior to that date, then you'll need to recreate the

relationship.

What you'll need

- A new sync relationship or an existing sync relationship created after the 23 Feb 2020 release.
- Any type of data broker.

This feature works with *any* type of data broker: the AWS, Azure, Google Cloud Platform, or on-prem data broker. The on-prem data broker can run [any supported operating system](#).

- For NFS, you'll need to use NFS version 4 or later.

The source and target versions must be the same. For example, 4.0 on the source and 4.0 on the target is supported. But 4.0 on the source and 4.1 on the target isn't supported because the versions are different.

Steps for a new relationship

1. From Cloud Sync, click **Create New Sync**.
2. Drag and drop **SMB Server** to the source and target and click **Continue**.
3. On the **SMB Server** page:
 - a. Enter a new SMB server or select an existing server and click **Continue**.
 - b. Enter credentials for the SMB server.
 - c. Select **Copy Access Control Lists to the target** and click **Continue**.

The screenshot shows a dialog box titled "Select an SMB Source". At the top right is a close button (X). Below the title, it says "SMB Version : 2.1 ▾". Underneath, there's a section for "Selected SMB Server:" showing a circular icon with "SMB" and the IP address "10.20.30.152". To the right is a delete button (X). Below this, there's a "Define SMB Credentials:" section with fields for "User Name" (containing "user1"), "Password" (containing "*****"), and "Domain (Optional)". Further down is an "ACL - Access Control List" section with a checked checkbox for "Copy Access Control Lists to the target". At the bottom, there's a note: "Notice: Copying ACLs can affect sync performance. You can change this setting after you create the relationship."

4. Follow the remaining prompts to create the sync relationship.

Steps for an existing relationship

1. Hover over the sync relationship and click the action menu.
2. Click **Settings**.

3. Select **Copy Access Control Lists to the target**.

4. Click **Save Settings**.

Result

When syncing data, Cloud Sync preserves the ACLs between the source and target SMB shares.

Manually copying ACLs

You can manually preserve ACLs between SMB shares by using the Windows robocopy command.

Steps

1. Identify a Windows host that has full access to both SMB shares.
2. If either of the endpoints require authentication, use the **net use** command to connect to the endpoints from the Windows host.

You must perform this step before you use robocopy.

3. From Cloud Sync, create a new relationship between the source and target SMB shares or sync an existing relationship.
4. After the data sync is complete, run the following command from the Windows host to sync the ACLs and ownership:

```
robocopy /E /COPY:SOU /secfix [source] [target] /w:0 /r:0 /XD ~snapshots  
/UNILOG:"[logfilepath]"
```

Both *source* and *target* should be specified using the UNC format. For example: \\<server>\<share>\<path>

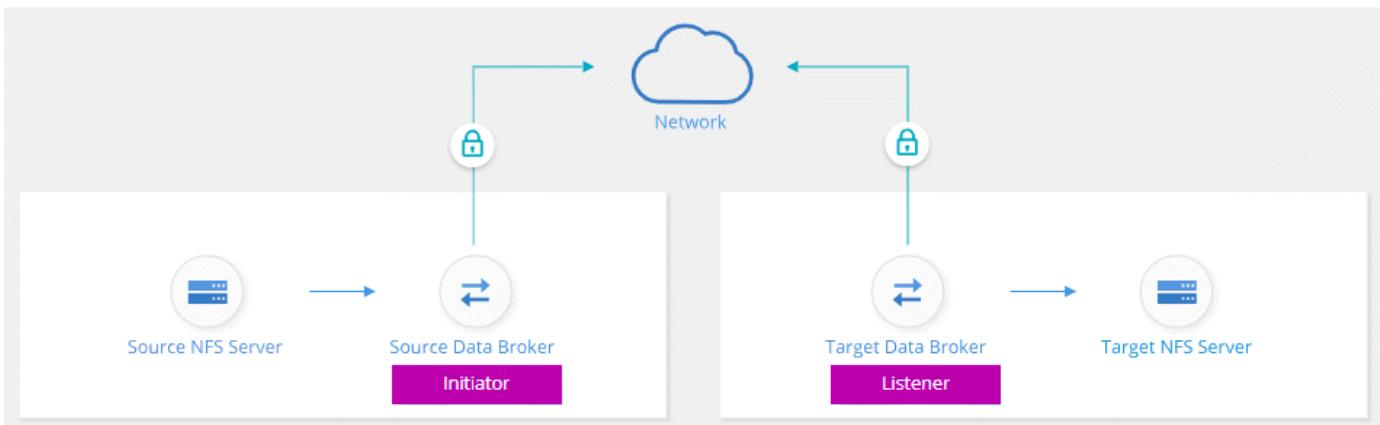
Syncing NFS data using data-in-flight encryption

If your business has strict security policies, you can sync NFS data using data-in-flight encryption. This feature is supported from an NFS server to another NFS server and from Azure NetApp Files to Azure NetApp Files.

For example, you might want to sync data between two NFS servers that are in different networks. Or you might need to securely transfer data on Azure NetApp Files across subnets or regions.

How data-in-flight encryption works

Data-in-flight encryption encrypts NFS data when it's sent over the network between two data brokers. The following image shows a relationship between two NFS servers and two data brokers:



One data broker functions as the *initiator*. When it's time to sync data, it sends a connection request to the other data broker, which is the *listener*. That data broker listens for requests on port 443. You can use a different port, if needed, but be sure to check that the port is not in use by another service.

For example, if you sync data from an on-premises NFS server to a cloud-based NFS server, you can choose which data broker listens for the connection requests and which sends them.

Here's how in-flight encryption works:

1. After you create the sync relationship, the initiator starts an encrypted connection with the other data broker.
2. The source data broker encrypts data from the source using TLS 1.3.
3. It then sends the data over the network to the target data broker.
4. The target data broker decrypts the data before sending it to the target.
5. After the initial copy, the service syncs any changed data every 24 hours. If there is data to sync, the process starts with the initiator opening an encrypted connection with the other data broker.

If you prefer to sync data more frequently, [you can change the schedule after you create the relationship](#).

Supported NFS versions

- For NFS servers, data-in-flight encryption is supported with NFS versions 3, 4.0, 4.1, and 4.2.
- For Azure NetApp Files, data-in-flight encryption is supported with NFS versions 3 and 4.1.

Proxy server limitation

If you create an encrypted sync relationship, the encrypted data is sent over HTTPS and isn't routable through a proxy server.

What you'll need to get started

Be sure to have the following:

- Two NFS servers that meet [source and target requirements](#) or Azure NetApp Files in two subnets or regions.
- The IP addresses or fully qualified domain names of the servers.
- Network locations for two data brokers.

You can select an existing data broker but it must function as the initiator. The listener data broker must be a new data broker.

If you have not yet deployed a data broker, review the data broker requirements. Because you have strict security policies, be sure to review the networking requirements, which includes outbound traffic from port 443 and the [internet endpoints](#) that the data broker contacts.

- Review AWS installation
- Review Azure installation
- Review GCP installation
- Review Linux host installation

Syncing NFS data using data-in-flight encryption

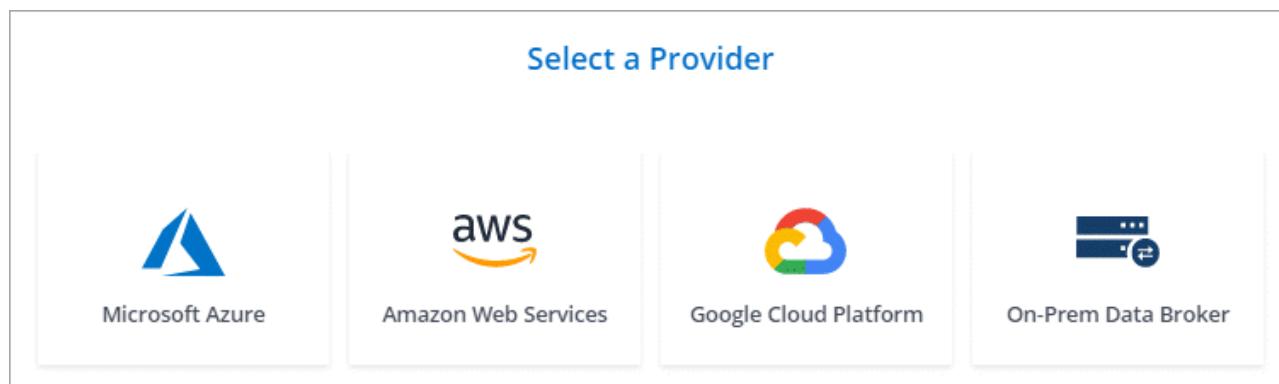
Create a new sync relationship between two NFS servers or between Azure NetApp Files, enable the in-flight encryption option, and follow the prompts.

Steps

1. Click **Create New Sync**.
2. Drag and drop **NFS Server** to the source and target locations or **Azure NetApp Files** to the source and target locations and select **Yes** to enable data-in-flight encryption.
3. Follow the prompts to create the relationship:
 - a. **NFS Server/Azure NetApp Files:** Choose the NFS version and then specify a new NFS source or select an existing server.
 - b. **Define Data Broker Functionality:** Define which data broker *listens* for connection requests on a port and which one *initiates* the connection. Make your choice based on your networking requirements.
 - c. **Data Broker:** Follow the prompts to add a new source data broker or select an existing data broker.

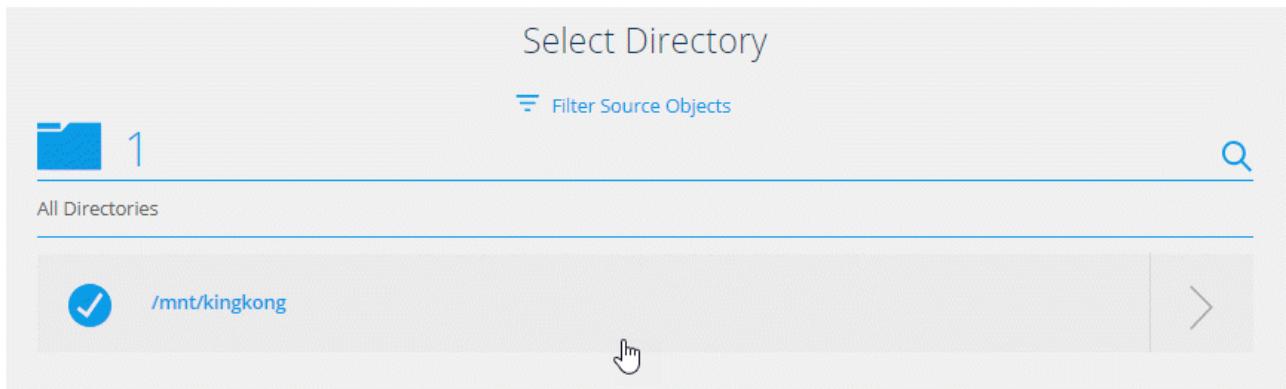
If the source data broker acts as the listener, then it must be a new data broker.

If you need a new data broker, Cloud Sync prompts you with the installation instructions. You can deploy the data broker in the cloud or download an installation script for your own Linux host.



- d. **Directories:** Choose the directories that you want to sync by selecting all directories, or by drilling down and selecting a subdirectory.

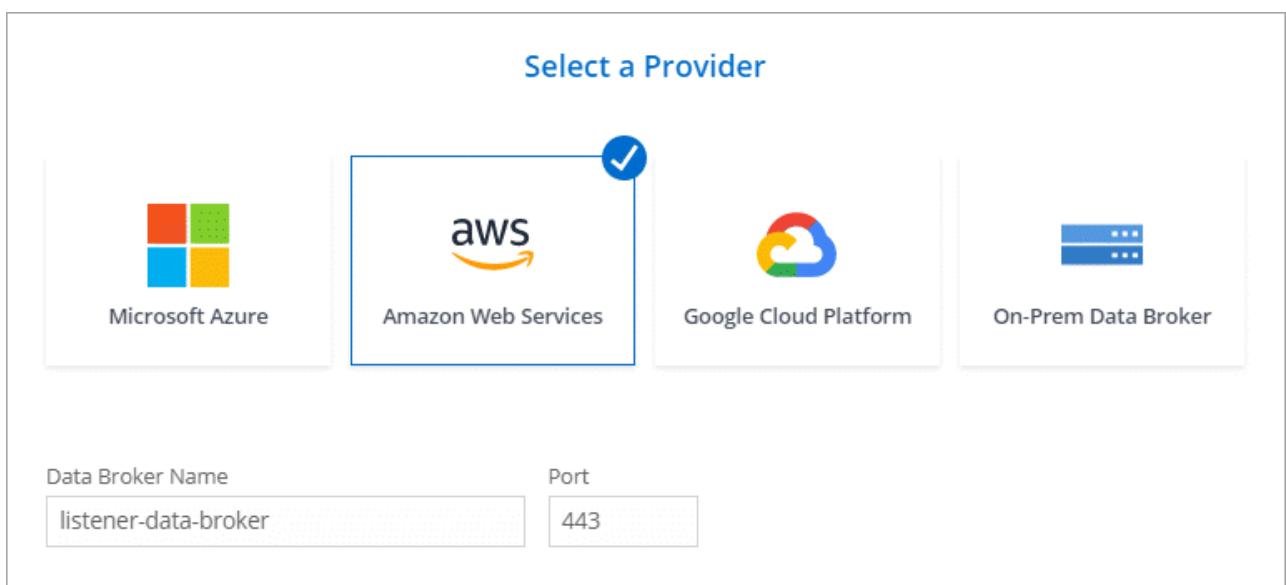
Click **Filter Source Objects** to modify settings that define how source files and folders are synced and maintained in the target location.



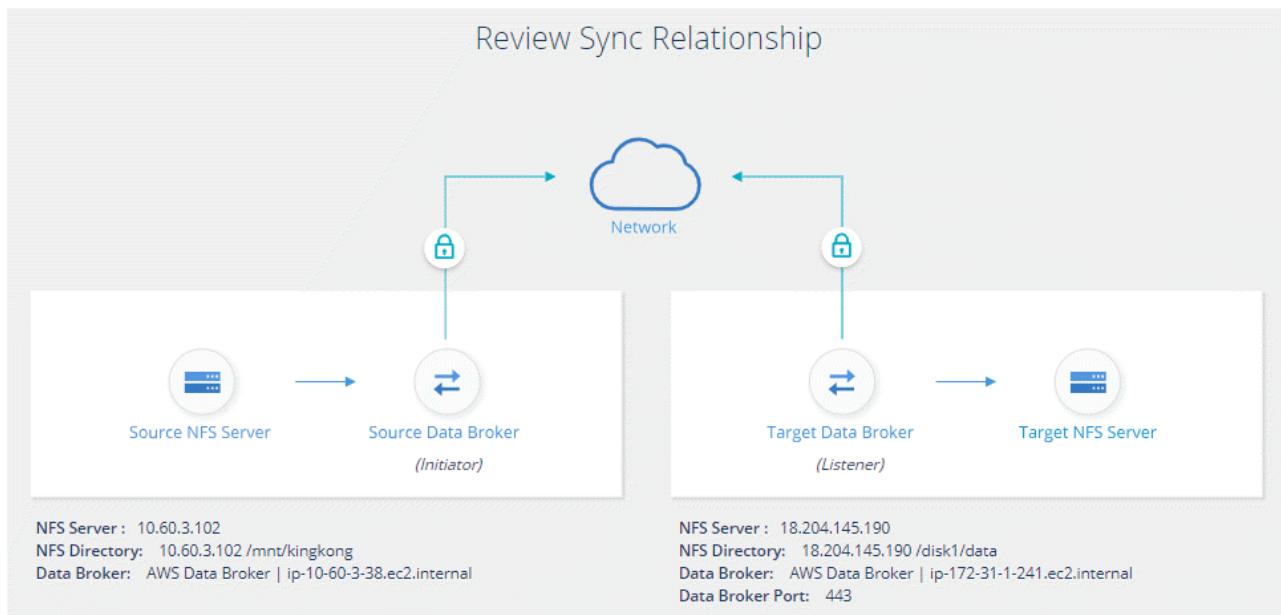
- e. **Target NFS Server/Target Azure NetApp Files:** Choose the NFS version and then enter a new NFS target or select an existing server.
- f. **Target Data Broker:** Follow the prompts to add a new source data broker or select an existing data broker.

If the target data broker acts as the listener, then it must be a new data broker.

Here's an example of the prompt when the target data broker functions as the listener. Notice the option to specify the port.



- g. **Target Directories:** Select a top-level directory, or drill down to select an existing subdirectory or to create a new folder inside an export.
- h. **Settings:** Define how source files and folders are synced and maintained in the target location.
- i. **Review:** Review the details of the sync relationship and then click **Create Relationship**.



Result

Cloud Sync starts creating the new sync relationship. When it's done, click **View in Dashboard** to view details about the new relationship.

Setting up the data broker to use an external HashiCorp Vault

When you create a sync relationship that requires Amazon S3, Azure, or Google Cloud credentials, you need to specify those credentials through the Cloud Sync user interface or API. An alternative is to set up the data broker to access the credentials (or *secrets*) directly from an external HashiCorp Vault.

This feature is supported through the Cloud Sync API with sync relationships that require Amazon S3, Azure, or Google Cloud credentials.

1 Prepare the vault

Prepare the vault to supply credentials to the data broker by setting up the URLs. The URLs to the secrets in the vault must end with *Creds*.

2 Prepare the data broker

Prepare the data broker to fetch credentials from the external vault by modifying the local config file for the data broker.

3 Create a sync relationship using the API

Now that everything is set up, you can send an API call to create a sync relationship that uses your vault to get the secrets.

Preparing the vault

You'll need to provide Cloud Sync with the URL to the secrets in your vault. Prepare the vault by setting up those URLs. You need to set up URLs to the credentials for each source and target in the sync relationships that you plan to create.

The URL must be set up as follows:

```
/<path>/<requestid>/<endpoint-protocol>Creds
```

Path

The prefix path to the secret. This can be any value that's unique to you.

Request ID

A request ID that you need to generate. You'll need to provide the ID in one of the headers in the API POST request when you create the sync relationship.

Endpoint protocol

One of the following protocols, as defined [in the post relationship v2 documentation](#): S3, AZURE, or GCP (each must be in uppercase).

Creds

The URL must end with *Creds*.

Examples

The following examples show URLs to secrets.

Example for the full URL and path for source credentials

```
http://example.vault.com:8200/my-path/all-secrets/hb312vdasr2/S3Creds
```

As you can see in the example, the prefix path is */my-path/all-secrets/*, the request ID is *hb312vdasr2* and the source endpoint is S3.

Example for the full URL and path for target credentials

```
http://example.vault.com:8200/my-path/all-secrets/n32hcbnejk2/AZURECreds
```

The prefix path is */my-path/all-secrets/*, the request ID is *n32hcbnejk2*, and the target endpoint is Azure.

Preparing the data broker

Prepare the data broker to fetch credentials from the external vault by modifying the local config file for the data broker.

Steps

1. SSH to the data broker.
2. Edit the `local.json` file that resides in `/opt/netapp/databroker/config`.
3. Set `enable` to **true** and set the config parameter fields under `external-integrations.hashicorp` as follows:

enabled

- Valid values: true/false
- Type: Boolean

- Default value: false
- True: The data broker gets secrets from your own external HashiCorp Vault
- False: The data broker stores credentials in its local vault

url

- Type: string
- Value: The URL to your external vault

path

- Type: string
- Value: Prefix path to the secret with your credentials

Reject-unauthorized

- Determines if you want the data broker to reject unauthorized external vault
- Type: Boolean
- Default: false

auth-method

- The authentication method that the data broker should use to access credentials from the external vault
- Type: string
- Valid values: "aws-iam" / "role-app" / "gcp-iam"

role-name

- Type: string
- Your role name (in case you use aws-iam or gcp-iam)

Secretid & rootid

- Type: string (in case you use app-role)

Namespace

- Type: string
- Your namespace (X-Vault-Namespace header if needed)

Example for aws-role authentication

```
{
    "external-integrations": {
        "hashicorp": {
            "enabled": true,
            "url": "https://example.vault.com:8200",
            "path": "/my-path/all-secrets",
            "reject-unauthorized": false,
            "auth-method": "aws-role",
            "aws-role": {
                "role-name": "my-role"
            }
        }
    }
}
```

Example for gcp-iam authentication

```
{
    "external-integrations": {
        "hashicorp": {
            "enabled": true,
            "url": "http://ip-10-20-30-55.ec2.internal:8200",
            "path": "v1/secret",
            "namespace": "",
            "reject-unauthorized": true,
            "auth-method": "gcp-iam",
            "aws-iam": {
                "role-name": ""
            },
            "app-role": {
                "root_id": "",
                "secret_id": ""
            }
        },
        "gcp-iam": {
            "role-name": "my-iam-role"
        }
    }
}
```

Setting up permissions when using gcp-iam authentication

If you're using the *gcp-iam* authentication method, then the data broker must have the following GCP permission:

- `iam.serviceAccounts.signJwt`

Learn more about GCP permission requirements for the data broker.

Creating a new sync relationship using secrets from the vault

Now that everything is set up, you can send an API call to create a sync relationship that uses your vault to get the secrets.

Post the relationship using the Cloud Sync REST API.

Headers:

```
Authorization: Bearer <user-token>
Content-Type: application/json
x-account-id: <accountid>
x-netapp-external-request-id-src: request ID as part of path for source
credentials
x-netapp-external-request-id-trg: request ID as part of path for target
credentials
Body: post relationship v2 body
```

- To obtain a user token and your Cloud Central account ID, refer to this page in the documentation.
- To build a body for your post relationship, refer to the relationships-v2 API call.

Example

Example for the POST request:

```

url: https://api.cloudsync.netapp.com/api/relationships-v2
headers:
"x-account-id": "CS-SasdW"
"x-netapp-external-request-id-src": "hb312vdasr2"
"Content-Type": "application/json"
"Authorization": "Bearer eyJhbGciOiJSUzI1NiIsInR5cCI6IkpxVCIsImtpZCI6Ik..."
Body:
{
  "dataBrokerId": "5e6e111d578dtyuu1555sa60",
  "source": {
    "protocol": "s3",
    "s3": {
      "provider": "sgws",
      "host": "1.1.1.1",
      "port": "443",
      "bucket": "my-source"
    },
    "target": {
      "protocol": "s3",
      "s3": {
        "bucket": "my-target-bucket"
      }
    }
  }
}

```

Managing sync relationships

You can manage sync relationships at any time by immediately syncing data, changing schedules, and more.

Performing an immediate data sync

Rather than wait for the next scheduled sync, you can press a button to immediately sync data between the source and target.

Steps

1. From the **Sync Dashboard**, hover over the sync relationship and click the action menu.

Source	Target	NetApp Data Broker	Schedule	Sync Status
nfs://172.31.91.49/disk1/data/d...	nfs://172.31.91.49/disk2/target/...	vadimBroker1	ON	Synced Successfully
nfs://172.31.91.49/disk1/data/e...	nfs://172.31.91.49/disk2/target/...	vadimBroker1	ON	Synced Successfully

2. Click **Sync Now** and then click **Sync** to confirm.

The screenshot shows the 'Syncs' section of the Cloud Sync interface. There are two sync relationships listed:

- Source:** nfs://172.31.91.49/disk1/data/... **Target:** s3://vadim-service-2test-e... **Status:** 3 da **Action:** Sync Now (highlighted with a cursor)
- Source:** nfs://172.31.91.49/disk1/data/... **Target:** nfs://172.31.91.49/disk2/t... **Status:** 3 data brokers OFF **Action:** Sync Completed

Result

Cloud Sync starts the data sync process for the relationship.

Accelerating sync performance

Accelerate the performance of a sync relationship by adding an additional data broker to the relationship. The additional data broker must be a *new* data broker.

How this works

If the existing data brokers in the relationship are used in other sync relationships, then Cloud Sync automatically adds the new data broker to those relationships, as well.

For example, let's say you have three relationships:

- Relationship 1 uses data broker A
- Relationship 2 uses data broker B
- Relationship 3 uses data broker A

You want to accelerate the performance of relationship 1 so you add a new data broker to that relationship (data broker C). Because data broker A is also used in relationship 3, the new data broker is automatically added to relationship 3, as well.

Steps

1. Ensure that at least one of the existing data brokers in the relationship are online.
2. Hover over the sync relationship and click the action menu.
3. Click **Accelerate**.

The screenshot shows the 'Syncs' section of the Cloud Sync interface. The same two sync relationships are listed as in the previous screenshot. The 'Accelerate' button for the top relationship is highlighted with a cursor.

4. Follow the prompts to create a new data broker.

Result

Cloud Sync adds the new data broker to the sync relationships. The performance of the next data sync should be accelerated.

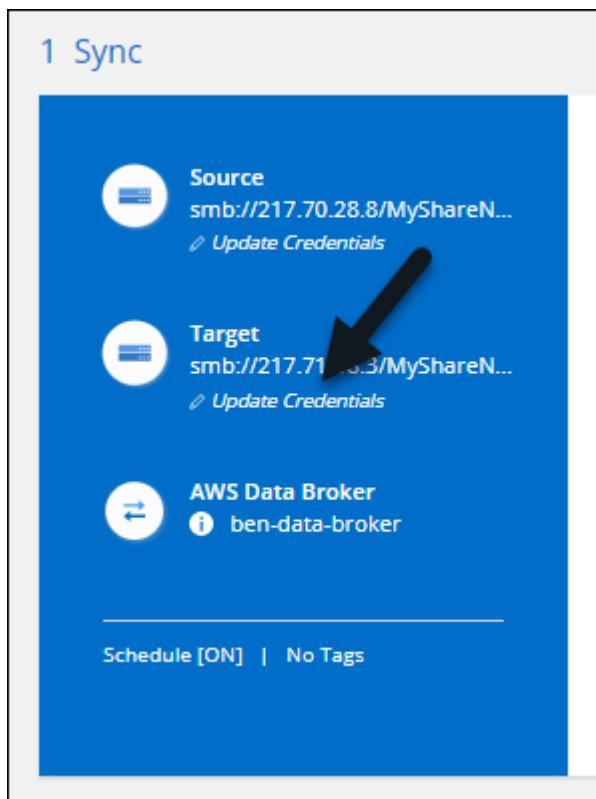
Updating credentials

You can update the data broker with the latest credentials of the source or target in an existing sync relationship. Updating the credentials can help if your security policies require you to update credentials on a periodic basis.

Updating credentials is supported with any source or target that Cloud Sync requires credentials for: Azure Blob, IBM Cloud Object Storage, StorageGRID, ONTAP S3 Storage, SFTP, and SMB servers.

Steps

1. From the **Sync Dashboard**, go to a sync relationship that requires credentials and then click **Update Credentials**.



2. Enter the credentials and click **Update**.

A note about SMB servers: if the domain is new, then you'll need to specify it when you update the credentials. If the domain hasn't changed, then you don't need to enter it again.

If you entered a domain when you created the sync relationship, but you don't enter a new domain when you update the credentials, then Cloud Sync will keep using the original domain that you provided.

Result

Cloud Sync updates the credentials on the data broker. It can take up 10 minutes until the data broker starts using the updated credentials for data syncs.

Changing the settings for a sync relationship

Modify settings that define how source files and folders are synced and maintained in the target location.

1. Hover over the sync relationship and click the action menu.
2. Click **Settings**.
3. Modify any of the settings.

The screenshot shows the Cloud Sync Settings page. At the top, there's a 'General' tab and a 'Files and Directories' tab. Under 'General', there are two sections: 'Schedule' (set to 'ON | Every 1 Day') and 'Retries' (set to 'Retry 3 times before skipping file'). Under 'Files and Directories', there are several settings: 'Compare By' (set to 'The following attributes (and size): uid, gid, mtime, mode'), 'Recently Modified Files' (set to 'Exclude files that are modified up to 30 Seconds before a scheduled sync'), 'Delete Files On Source' (set to 'Never delete files from the source location'), 'Delete Files On Target' (set to 'Never delete files from the target location'), 'File Types' (set to 'Include All: Files, Directories, Symbolic Links'), 'Exclude File Extensions' (set to 'None'), 'File Size' (set to 'All'), and 'Date Modified' (set to 'All'). At the bottom of this section is a 'Reset to defaults' button.

Here's a brief description of each setting:

Schedule

Choose a recurring schedule for future syncs or turn off the sync schedule. You can schedule a relationship to sync data as often as every 1 minute.

Retries

Define the number of times that Cloud Sync should retry to sync a file before skipping it.

Compare By

Choose whether Cloud Sync should compare certain attributes when determining whether a file or directory has changed and should be synced again.

Even if you uncheck these attributes, Cloud Sync still compares the source to the target by checking the

paths, file sizes, and file names. If there are any changes, then it syncs those files and directories.

You can choose to enable or disable Cloud Sync from comparing the following attributes:

- **mtime**: The last modified time for a file. This attribute isn't valid for directories.
- **uid, gid, and mode**: Permission flags for Linux.

Copy for Objects

You can't edit this option after you create the relationship.

Recently Modified Files

Choose to exclude files that were recently modified prior to the scheduled sync.

Delete Files on Source

Choose to delete files from the source location after Cloud Sync copies the files to the target location. This option includes the risk of data loss because the source files are deleted after they're copied.

If you enable this option, you also need to change a parameter in the local.json file on the data broker. Open the file and change the parameter named `workers.transferrer.delete-on-source` to **true**.

Delete Files on Target

Choose to delete files from the target location, if they were deleted from the source. The default is to never deletes files from the target location.

File Types

Define the file types to include in each sync: files, directories, and symbolic links.

Exclude File Extensions

Specify file extensions to exclude from the sync by typing the file extension and pressing **Enter**. For example, type `log` or `.log` to exclude `*.log` files. A separator isn't required for multiple extensions. The following video provides a short demo:

► https://docs.netapp.com/us-en/occm//media/video_file_extensions.mp4 (video)

File Size

Choose to sync all files regardless of their size or just files that are in a specific size range.

Date Modified

Choose all files regardless of their last modified date, files modified after a specific date, before a specific date, or between a time range.

Copy Access Control Lists to the target

Choose to copy access control lists (ACLs) between source SMB shares and target SMB shares. Note that this option is only available for sync relationships created after the 23 Feb 2020 release.

4. Click **Save Settings**.

Result

Cloud Sync modifies the sync relationship with the new settings.

Deleting relationships

You can delete a sync relationship, if you no longer need to sync data between the source and target. This action does not delete the data broker instance and it does not delete data from the target.

Steps

1. Hover over the sync relationship and click the action menu.
2. Click **Delete** and then click **Delete** again to confirm.

Result

Cloud Sync deletes the sync relationship.

Manage data brokers

A data broker syncs data from a source location to a target location. A data broker is required for each sync relationship that you create. Manage data brokers by adding a new data broker to a group, by viewing information about data brokers, and more.

Data broker groups

Grouping data brokers together can help improve the performance of sync relationships.

Determining the number of data brokers

In many cases, a single data broker can meet the performance requirements for a sync relationship. If it doesn't, you can accelerate sync performance by adding additional data brokers to the group. But you should first check other factors that can impact sync performance. [Learn more about how to determine when multiple data brokers are required.](#)

Groups can manage several relationships

A data broker group can manage one or more sync relationships at a time.

For example, let's say you have three relationships:

- Relationship 1 uses data broker A
- Relationship 2 uses data broker B
- Relationship 3 uses data broker A

You want to accelerate the performance of relationship 1 so you add a new data broker (data broker C) to the group. Because data broker A is also used to manage relationship 3, having two data brokers in the group also accelerates the performance of this relationship.

New data brokers only

You can only add new data brokers to a group. You can't add existing data brokers to a group.

Add a new data broker

There are several ways to create a new data broker:

- When creating a new sync relationship

[Learn how to create a new data broker when creating a sync relationship.](#)

- From the **Manage Data Brokers** page by clicking **Add New Data Broker** which creates the data broker in a new group
- From the **Manage Data Brokers** page by creating a new data broker in an existing group

Things you should know

- You can't add data brokers to a group that manages an encrypted sync relationship.
- If you want to create a data broker in an existing group, the data broker must be an on-prem data broker or the same type of data broker.

For example, if a group includes an AWS data broker, then you can create an AWS data broker or on-prem data broker in that group. You can't create an Azure data broker or GCP data broker because they aren't the same data broker type.

Steps to create a data broker in a new group

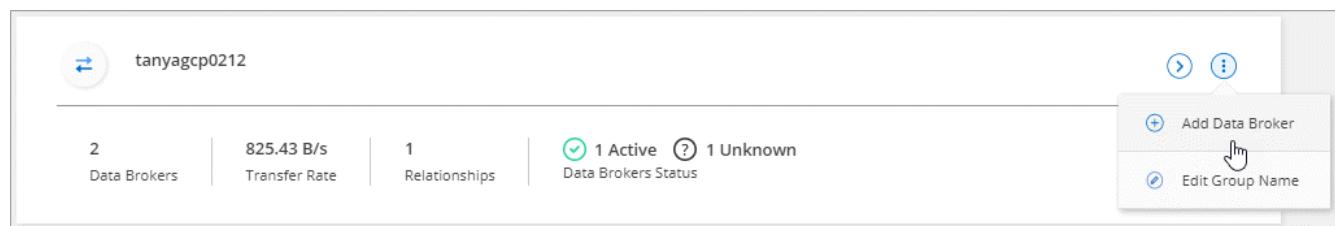
1. Click **Sync > Manage Data Brokers**.
2. Click **Add New Data Broker**.
3. Follow the prompts to create the data broker.

For help, refer to the following pages:

- [Installing the data broker in AWS](#)
- [Installing the data broker in Azure](#)
- [Installing the data broker in GCP](#)
- [Installing the data broker on a Linux host](#)

Steps to create a data broker in an existing group

1. Click **Sync > Manage Data Brokers**.
2. Click the action menu and select **Add Data Broker**.



3. Follow the prompts to create the data broker.

For help, refer to the following pages:

- [Installing the data broker in AWS](#)
- [Installing the data broker in Azure](#)
- [Installing the data broker in GCP](#)
- [Installing the data broker on a Linux host](#)

Security recommendations

To ensure the security of your data broker machine, NetApp recommends the following:

- SSH should not permit X11 Forwarding
- SSH should not permit TCP connection forwarding
- SSH should not permit tunnels
- SSH should not accept client environment variables

These security recommendations can help prevent unauthorized connections to the data broker machine.

View a data broker's configuration

You might want to view details about a data broker to identify things like its host name, IP address, available CPU and RAM, and more.

Cloud Sync provides the following details about a data broker:

- Basic information: Instance ID, host name, etc.
- Network: Region, network, subnet, private IP, etc.
- Software: Linux distribution, data broker version, etc.
- Hardware: CPU and RAM
- Configuration: Details about the data broker's two kinds of main processes—scanner and transcoder



The scanner scans the source and target and decides what should be copied. The transcoder does the actual copying. NetApp personnel might use these configuration details to suggest actions that can optimize performance.

Steps

1. Click **Sync > Manage Data Brokers**.
2. Click to expand the list of data brokers in a group.
3. Click to view details about a data broker.

tanyagcp0212

2 Data Brokers | 968.5 B/s Transfer Rate | 1 Relationships | 1 Active 1 Unknown Data Brokers Status

Information	5fc766b3d3e3664b9e116... Broker ID	288871247573080556 Instance ID	tanyagcp0212-mnx-data-... Host Name	cloudsync-dev-214020 Project Id
Network	us-east1-b Region	default Network	255.255.240.0 Subnet	10.142.0.37 Private IP
Software	linux Linux Distribution & Version	1.5.4 Vault Version	14.15.1 Node Version	1.3.0.18650-73f960d-integ Data Broker Version
Hardware	4 Available CPUs	62.22 MB Available RAM		
Configuration	50 Scanner Concurrency	4 Scanner CPUs	50 Transferrer Concurrency	4 Transferrer CPUs

Remove a data broker from a group

You might remove a data broker from a group if it's no longer needed or if the initial deployment failed. This action only deletes the data broker from Cloud Sync's records. You'll need to manually delete the data broker and any additional cloud resources yourself.

Things you should know

- Cloud Sync deletes a group when you remove the last data broker from the group.
- You can't remove the last data broker from a group if there is a relationship using that group.

Steps

- Click **Sync > Manage Data Brokers**.
- Click to expand the list of data brokers in a group.
- Click the action menu for a data broker and select **Remove Data Broker**.

tanyagcp0212

2 Data Brokers | 968.5 B/s Transfer Rate | 1 Relationships | 1 Active 1 Unknown Data Brokers Status

tanyagcp0212	GCP	Transfer Rate: 968.5 B/s	Active
tanya1	ONPREM	Transfer Rate: N/A	Unknown

⋮ Remove Data Broker

4. Click Remove Data Broker.

Result

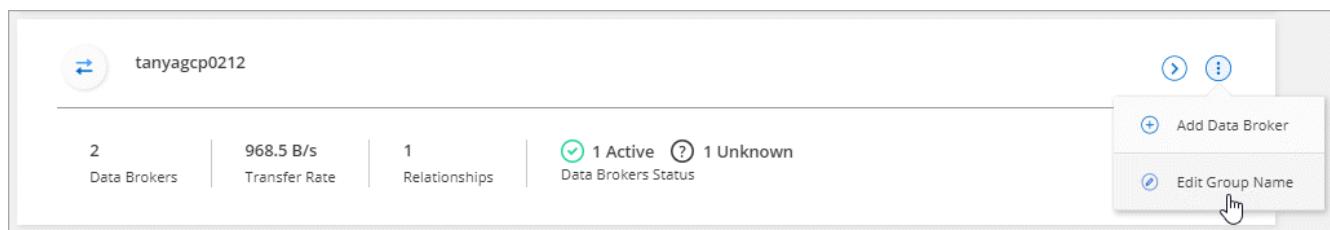
Cloud Sync removes the data broker from the group.

Edit a group's name

Change the name of a data broker group at any time.

Steps

1. Click **Sync > Manage Data Brokers**.
2. Click the action menu and select **Edit Group Name**.



3. Enter a new name and click **Save**.

Result

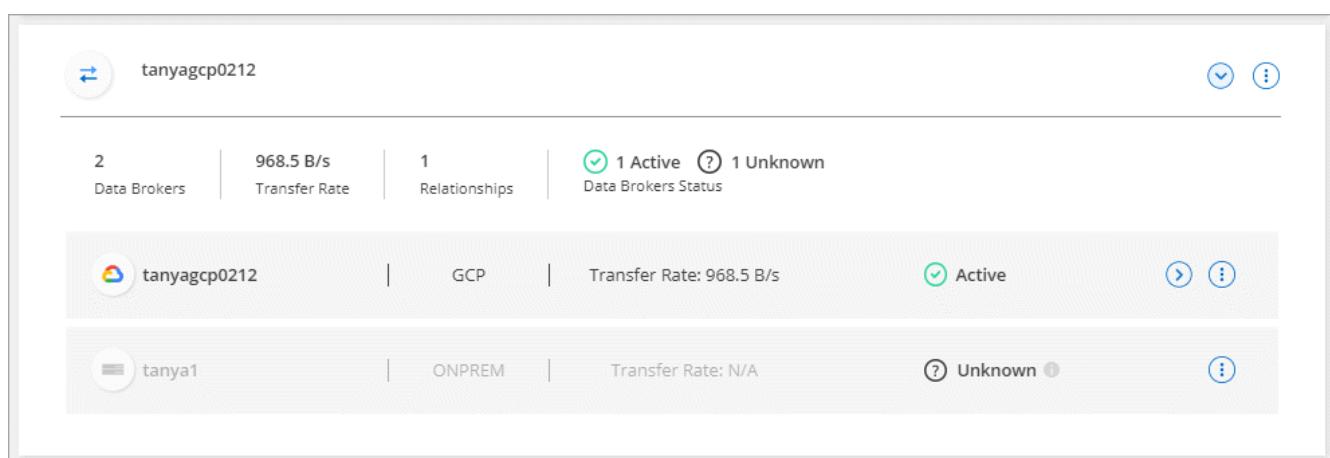
Cloud Sync updates the name of the data broker group.

Address issues with a data broker

Cloud Sync displays a status for each data broker that can help you troubleshoot issues.

Steps

1. Identify any data brokers that have a status of "Unknown" or "Failed."



2. Hover over the icon to see the failure reason.

3. Correct the issue.

For example, you might need to simply restart the data broker if it's offline, or you might need to remove data broker if the initial deployment failed.

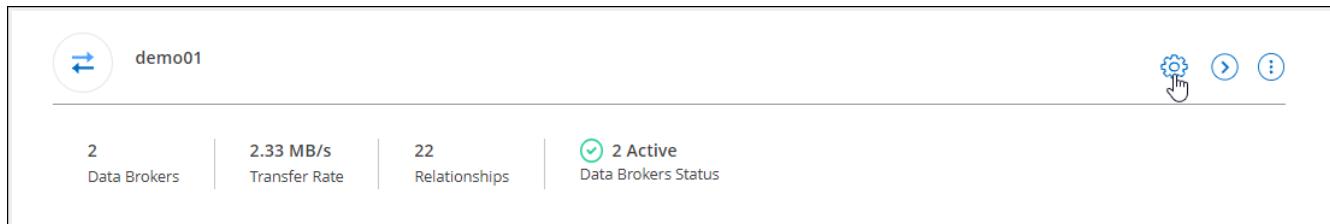
Define a unified configuration for a data broker group

If a sync relationship encounters errors during the sync process, unifying the concurrency of the data broker group can help to decrease the number of sync errors. Be aware that changes to the group's configuration can affect performance by slowing down the transfer.

We don't recommend changing the configuration on your own. You should consult with NetApp to understand when to change the configuration and how to change it.

Steps

1. Click **Manage Data Brokers**.
2. Click the Settings icon for a data broker group.



3. Change the settings as needed and then click **Unify Configuration**.

Note the following:

- You can pick and choose which settings to change—you don't need to change all four at once.
- After a new configuration is sent to a data broker, the data broker automatically restarts and uses the new configuration.
- It can take up to a minute until this change takes place and is visible in the Cloud Sync interface.
- If a data broker isn't running, its configuration won't change because Cloud Sync can't communicate with it. The configuration will change after the data broker restarts.
- After you set a unified configuration, any new data brokers will automatically use the new configuration.

Creating and viewing reports to tune your configuration

Create and view reports to get information that you can use with the help of NetApp personnel to tune a data broker's configuration and improve performance.

Each report provides in-depth details about a path in a sync relationship. For example, the report for a file system shows how many directories and files there are, the distribution of file size, how deep and wide the directories are, and more.

Creating reports

Each time that you create a report, Cloud Sync scans the path and then compiles the details into a report.

Steps

1. Click **Sync > Reports**.

The paths (source or target) in each of your sync relationships display in a table.

2. In the **Reports Actions** column, go to a specific path and click **Create**, or click the action menu and select **Create New**.
3. When the report is ready, click action menu and select **View**.

Here's a sample report for a file system path.

And here's a sample report for object storage.

Viewing report errors

The Paths table identifies whether any errors are present in the most recent report. An error identifies an issue that Cloud Sync faced when scanning the path.

For example, a report might contain permission denied errors. This type of error can affect Cloud Sync's ability to scan the entire set of files and directories.

After you view the list of errors, you can then address the issues and run the report again.

Steps

1. Click **Sync > Reports**.
2. In the **Errors** column, identify whether any errors are present in a report.
3. If errors are present, click the arrow next to the number of errors.

20 Paths								Q
Path	Reports #	Last Report Date	Last Scan Duration	Last Report Status	Errors	Reports Actions		
nfs://1.1.1.1/data1/success	2 Reports	Apr 12, 2021 11:34 AM	Less than a minute	Completed	None	...		
nfs://4.4.4.4/data1/success	2 Reports	Apr 19, 2021 3:21 PM	Less than a minute	Completed	None	...		
s3://phoebe/failed	1 Report	Apr 20, 2021 3:30 PM	Less than a minute	Failed	None	...		
s3://phoebe	2 Reports	Apr 19, 2021 1:23 PM	Less than a minute	Completed	1 ↗	...		
s3://large	1 Report	Apr 20, 2021 3:30 PM	Less than a minute	Completed	1 ↗	...		

4. Use the information in the error to correct the issue.

After you resolve the issue, the error shouldn't appear the next time that you run the report.

Deleting reports

You might delete a report if it contained an error that you fixed, or if the report is related to a sync relationship that you removed.

Steps

1. Click **Sync > Reports**.

2. In the **Reports Actions** column, click the action menu for a path and select **Delete last report** or **Delete all reports**.
3. Confirm that you want to delete the report or reports.

Uninstalling the data broker

If needed, run an uninstall script to remove the data broker and the packages and directories that were created when the data broker was installed.

Steps

1. Log in to the data broker host.
2. Change to the data broker directory: /opt/netapp/databroker
3. Run the following commands:

```
chmod +x uninstaller-DataBroker.sh  
./uninstaller-DataBroker.sh
```

4. Press 'y' to confirm the uninstallation.

Cloud Sync APIs

The Cloud Sync capabilities that are available through the web UI are also available through the RESTful API.

Getting started

To get started with the Cloud Sync API, you need to obtain a user token and your Cloud Central account ID. You'll need to add the token and account ID to the Authorization header when making API calls.

Steps

1. Obtain a user token from NetApp Cloud Central.

```
POST https://netapp-cloud-account.auth0.com/oauth/token  
Header: Content-Type: application/json  
Body:  
{  
    "username": "<user_email>",  
    "scope": "profile",  
    "audience": "https://api.cloud.netapp.com",  
    "client_id": "UaVhOIXMWQs5i1WdDxauXe5Mqkb34NJQ",  
    "grant_type": "password",  
    "password": "<user_password>"  
}
```

2. Obtain your Cloud Central account ID.

```
GET https://api.cloudsync.netapp.com/api/accounts  
Headers: Authorization: Bearer <user_token>  
Content-Type: application/json
```

This API will return a response like the following:

```
[  
 {  
   "accountId": "account-JeL97Ry3",  
   "name": "Test"  
 }  
]
```

3. Add the user token and account ID in the Authorization header of each API call.

Example

The following example shows an API call to create a data broker in Microsoft Azure. You would simply replace <user_token> and <accountId> with the token and ID that you obtained in the previous steps.

```
POST https://api.cloudsync.netapp.com/api/data-brokers  
Headers: Authorization: Bearer <user_token>  
Content-Type: application/json  
x-account-id: <accountId>  
Body: { "name": "databroker1", "type": "AZURE" }
```

What should I do when the token expires?

The user token from NetApp Cloud Central has an expiration date. To refresh the token, you need to call the API from step 1 again.

The API response includes an "expires_in" field that states when the token expires.

API reference

Documentation for each Cloud Sync API is available from <https://api.cloudsync.netapp.com/docs>.

Using list APIs

List APIs are asynchronous APIs, so the result does not return immediately (for example: GET /data-brokers/{id}/list-nfs-export-folders and GET /data-brokers/{id}/list-s3-buckets). The only response from the server is HTTP status 202. To get the actual result, you must use the GET /messages/client API.

Steps

1. Call the list API that you want to use.
2. Use the GET /messages/client API to view the result of the operation.

3. Use the same API by appending it with the ID that you just received: GET

```
http://api.cloudsync.netapp.com/api/messages/client?last=<id_from_step_2>
```

Note that the ID changes each time that you call the GET /messages/client API.

Example

When you call the list-s3-buckets API, a result is not immediately returned:

```
GET http://api.cloudsync.netapp.com/api/data-brokers/<data-broker-id>/list-s3-buckets
Headers: Authorization: Bearer <user_token>
Content-Type: application/json
x-account-id: <accountId>
```

The result is HTTP status code 202, which means the message was accepted, but was not processed yet.

To get the result of the operation, you need to use the following API:

```
GET http://api.cloudsync.netapp.com/api/messages/client
Headers: Authorization: Bearer <user_token>
Content-Type: application/json
x-account-id: <accountId>
```

The result is an array with one object that includes an ID field. The ID field represents the last message that the server sent. For example:

```
[{"header": {"requestId": "init", "clientId": "init", "agentId": "init"}, "payload": {"init": {}}, "id": "5801"}]
```

You would now make the following API call using the ID that you just received:

```
GET  
http://api.cloudsync.netapp.com/api/messages/client?last=<id_from_step_2>  
Headers: Authorization: Bearer <user_token>  
Content-Type: application/json  
x-account-id: <accountId>
```

The result is an array of messages. Inside each message is a payload object, which consists of the name of the operation (as key) and its result (as value). For example:

```
[  
  {  
    "payload": {  
      "list-s3-buckets": [  
        {  
          "tags": [  
            {  
              "Value": "100$"  
              "Key": "price"  
            }  
          ],  
          "region": {  
            "displayName": "US West (Oregon)",  
            "name": "us-west-2"  
          },  
          "name": "small"  
        }  
      ]  
    },  
    "header": {  
      "requestId": "f687ac55-2f0c-40e3-9fa6-57fb8c4094a3",  
      "clientId": "5beb032f548e6e35f4ed1ba9",  
      "agentId": "5bed61f4489fb04e34a9aac6"  
    },  
    "id": "5802"  
  }  
]
```

Using the API to create a sync relationship from Box

Cloud Sync supports [Box](#) as the source in a sync relationship when using the Cloud Sync API. [View supported sync relationships for Box.](#)

In order to copy from Box, you'll need to provide the following credentials:

- clientId
- clientSecret
- publicKeyId
- privateKey
- passphrase
- enterpriseId

Steps

1. Create a private key configuration file on the Box dev console.

The json file should look like this:

```
{
  "boxAppSettings": {
    "clientID": "<clientId>",
    "clientSecret": "<clientSecret>",
    "appAuth": {
      "publicKeyID": "<publicKeyID>",
      "privateKey": "<privatkey>",
      "passphrase": "<passphrase>"
    }
  },
  "enterpriseID": "<enterpriseId>"
}
```

2. Use the Cloud Sync API to create a sync relationship from Box to a supported target.

Here's an example for creating a sync relationship from Box to Amazons S3.

```
{
  "dataBrokerId": "60ede6f862e3cf7d31325037",
  "source": {
    "protocol": "box",
    "box": {
      "folderId": "0",
      "credentials": {
        "clientId": "<clientId>",
        "clientSecret": "<clientSecret>",
        "publicKeyId": "<publicKeyId>",
        "privateKey": "<privateKey>",
        "passphrase": "<passphrase>",
        "enterpriseId": "<enterpriseId>"
      }
    }
  },
}
```

```

"target": {
    "protocol": "s3",
    "s3": {
        "bucket": "chen02",
        "prefix": "fromBox20",
        "region": "us-east-1",
    }
},
"settings": {
    "gracePeriod": 0,
    "deleteOnSource": false,
    "deleteOnTarget": false,
    "objectTagging": false,
    "copyAcl": false,
    "retries": 3,
    "fileTypes": {
        "files": true,
        "directories": true,
        "symlinks": true
    },
    "copyProperties": {
        "metadata": false,
        "tags": false
    },
    "schedule": {
        "syncInDays": 1,
        "syncInHours": 0,
        "syncInMinutes": 0,
        "isEnabled": true,
        "syncWhenCreated": true
    }
}
}

```

Cloud Sync technical FAQ

This FAQ can help if you're just looking for a quick answer to a question.

Getting started

The following questions relate to getting started with Cloud Sync.

How does Cloud Sync work?

Cloud Sync uses the NetApp data broker software to sync data from a source to a target (this is called a *sync*

relationship).

The data broker controls the sync relationships between your sources and targets. After you set up a sync relationship, Cloud Sync analyzes your source system and breaks it up into multiple replication streams to push to your selected target data.

After the initial copy, the service syncs any changed data based on the schedule that you set.

How does the 14-day free trial work?

The 14-day free trial starts when you sign up for the Cloud Sync service. You're not subject to NetApp charges for Cloud Sync relationships you create for 14 days. However, all resource charges for any data broker that you deploy still applies.

How much does Cloud Sync cost?

There are two types of costs associated with using Cloud Sync: service charges and resource charges.

Service charges

For pay-as-you-go pricing, Cloud Sync service charges are hourly, based on the number of sync relationships that you create.

- [View pay-as-you-go pricing in AWS](#)
- [View annual pricing in AWS](#)
- [View pricing in Azure](#)

Cloud Sync licenses are also available through your NetApp representative. Each license enables 20 sync relationships for 12 months.

[Learn more about licenses.](#)



Cloud Sync relationships are free for Cloud Volumes Service and Azure NetApp Files.

Resource charges

The resource charges are related to the compute and storage costs for running the data broker in the cloud.

How is Cloud Sync billed?

There are two ways to pay for sync relationships after your 14-day free trial ends. The first option is to subscribe from AWS or Azure, which enables you to pay-as-you-go or to pay annually. The second option is to purchase licenses directly from NetApp.

Can I use Cloud Sync outside the cloud?

Yes, you can use Cloud Sync in a non-cloud architecture. The source and target can reside on-premises and so can the data broker.

Note the following key points about using Cloud Sync outside of the cloud:

- For on-premises synchronization, a private Amazon S3 bucket is available through NetApp StorageGRID.
- The data broker does need an internet connection to communicate with the Cloud Sync service.

- If you don't purchase a license directly from NetApp, you will need an AWS or Azure account for the PAYGO Cloud Sync service billing.

How do I access Cloud Sync?

Cloud Sync is available from Cloud Manager in the **Sync** tab.

Supported sources and targets

The following questions related to the source and targets that are supported in a sync relationship.

Which sources and targets does Cloud Sync support?

Cloud Sync supports many different types of sync relationships. [View the entire list.](#)

What versions of NFS and SMB does Cloud Sync support?

Cloud Sync supports NFS version 3 and later, and SMB version 1 and later.

[Learn more about sync requirements.](#)

When Amazon S3 is the target, can the data be tiered to a specific S3 storage class?

Yes, you can choose a specific S3 storage class when AWS S3 is the target:

- Standard (this is the default class)
- Intelligent-Tiering
- Standard-Infrequent Access
- One Zone-Infrequent Access
- Glacier
- Glacier Deep Archive

What about storage tiers for Azure Blob storage?

You can choose a specific Azure Blob storage tier when a Blob container is the target:

- Hot storage
- Cool storage

Do you support Google Cloud storage tiers?

Yes, you can choose a specific storage class when a Google Cloud Storage bucket is the target:

- Standard
- Nearline
- Coldline
- Archive

Networking

The following questions relate to networking requirements for Cloud Sync.

What are the networking requirements for Cloud Sync?

The Cloud Sync environment requires that the data broker is connected with the source and the target through the selected protocol (NFS, SMB, EFS) or object storage API (Amazon S3, Azure Blob, IBM Cloud Object Storage).

In addition, the data broker needs an outbound internet connection over port 443 so it can communicate with the Cloud Sync service and contact a few other services and repositories.

For more details, [review networking requirements](#).

Can I use a proxy server with the data broker?

Yes.

Cloud Sync supports proxy servers with or without basic authentication. If you specify a proxy server when you deploy a data broker, all HTTP and HTTPS traffic from the data broker is routed through the proxy. Note that non-HTTP traffic such as NFS or SMB can't be routed through a proxy server.

The only proxy server limitation is when using data-in-flight encryption with an NFS or Azure NetApp Files sync relationship. The encrypted data is sent over HTTPS and isn't routable through a proxy server.

Data synchronization

The following questions relate to how data synchronization works.

How often does synchronization occur?

The default schedule is set for daily synchronization. After the initial synchronization, you can:

- Modify the sync schedule to your desired number of days, hours, or minutes
- Disable the sync schedule
- Delete the sync schedule (no data will be lost; only the sync relationship will be removed)

What is the minimum sync schedule?

You can schedule a relationship to sync data as often as every 1 minute.

Does the data broker retry when a file fails to sync? Or does it timeout?

The data broker doesn't timeout when a single file fails to transfer. Instead, the data broker retries 3 times before skipping the file. The retry value is configurable in the settings for a sync relationship.

[Learn how to change the settings for a sync relationship.](#)

What if I have a very large dataset?

If a single directory contains 600,000 files or more, [contact us](#) so that we can help you configure the data broker to handle the payload. We might need to add additional memory to the data broker machine.

Note that there's no limit to the total number of files in the mount point. The extra memory is required for large directories with 600,000 files or more, regardless of their level in the hierarchy (top directory or subdirectory).

Security

The following questions related to security.

Is Cloud Sync secure?

Yes. All Cloud Sync service networking connectivity is done using [Amazon Simple Queue Service \(SQS\)](#).

All communication between the data broker and Amazon S3, Azure Blob, Google Cloud Storage, and IBM Cloud Object Storage is done through the HTTPS protocol.

If you're using Cloud Sync with on-premises (source or destination) systems, here's a few recommended connectivity options:

- An AWS Direct Connect, Azure ExpressRoute, or Google Cloud Interconnect connection, which is non-internet routed (and can only communicate with the cloud networks that you specify)
- A VPN connection between your on-premises gateway device and your cloud networks
- For extra secure data transfer with S3 buckets, Azure Blob storage, or Google Cloud Storage, an Amazon Private S3 Endpoint, Azure Virtual Network service endpoints, or Private Google Access may be established.

Any of these methods establishes a secure connection between your on-premises NAS servers and a Cloud Sync data broker.

Is data encrypted by Cloud Sync?

- Cloud Sync supports data-in-flight encryption between source and target NFS servers. [Learn more](#).
- Encryption is not supported with SMB.
- When an Amazon S3 bucket is the target in a sync relationship, you can choose whether to enable data encryption using AWS KMS encryption or AES-256 encryption.

Permissions

The following questions relate to data permissions.

Are SMB data permissions synced to the target location?

You can set up Cloud Sync to preserve access control lists (ACLs) between a source SMB share and a target SMB share. Or you can manually copy the ACLs yourself. [Learn how to copy ACLs between SMB shares](#).

Are NFS data permissions synced to the target location?

Cloud Sync automatically copies NFS permissions between NFS servers as follows:

- NFS version 3: Cloud Sync copies the permissions and the user group owner.
- NFS version 4: Cloud Sync copies the ACLs.

Object storage metadata

Cloud Sync copies object storage metadata from the source to the target for the following types of sync relationships:

- Amazon S3 → Amazon S3 ¹
- Amazon S3 → StorageGRID
- StorageGRID → Amazon S3
- StorageGRID → StorageGRID
- StorageGRID → Google Cloud Storage
- Google Cloud Storage → StorageGRID ¹
- Google Cloud Storage → IBM Cloud Object Storage ¹
- Google Cloud Storage → Amazon S3 ¹
- Amazon S3 → Google Cloud Storage
- IBM Cloud Object Storage → Google Cloud Storage
- StorageGRID → IBM Cloud Object Storage
- IBM Cloud Object Storage → StorageGRID
- IBM Cloud Object Storage → IBM Cloud Object Storage

¹ For these sync relationships, you need to [enable a setting when you create the sync relationship](#).

Performance

The following questions relate to Cloud Sync performance.

What does the progress indicator for a sync relationship represent?

The sync relationship shows the throughput of the data broker's network adapter. If you accelerated sync performance by using multiple data brokers, then the throughput is the sum of all traffic. This throughput refreshes every 20 seconds.

I'm experiencing performance issues. Can we limit the number of concurrent transfers?

The data broker can sync 4 files at a time. If you have very large files (multiple TiBs each), it can take a long time to complete the transfer process and performance might be impacted.

Limiting the number of concurrent transfers can help. [Contact us for help](#).

Why am I experiencing low performance with Azure NetApp Files?

When you sync data to or from Azure NetApp Files, you might experience failures and performance issues if the disk service level is Standard.

Change the service level to Premium or Ultra to enhance the sync performance.

[Learn more about Azure NetApp Files service levels and throughput](#).

Why am I experiencing low performance with Cloud Volumes Service for AWS?

When you sync data to or from a cloud volume, you might experience failures and performance issues if the level of performance for the cloud volume is Standard.

Change the Service level to Premium or Extreme to enhance the sync performance.

How many data brokers are required?

When you create a new relationship, you start with a single data broker (unless you selected an existing data broker that belongs to an accelerated sync relationship). In many cases, a single data broker can meet the performance requirements for a sync relationship. If it doesn't, you can accelerate sync performance by adding additional data brokers. But you should first check other factors that can impact sync performance.

Multiple factors can impact data transfer performance. The overall sync performance might be impacted due to network bandwidth, latency, and network topology, as well as the data broker VM specs and storage system performance. For example, a single data broker in a sync relationship can reach 100 MB/s, while disk throughput on the target might only allow 64 MB/s. As a result, the data broker keeps trying to copy the data, but the target can't meet the performance of the data broker.

So be sure to check the performance of your networking and the disk throughput on the target.

Then you can consider accelerating sync performance by adding an additional data broker to share the load of that relationship. [Learn how to accelerate sync performance](#).

Deleting things

The following questions relate to deleting sync relationships and data from sources and targets.

What happens if I delete my Cloud Sync relationship?

Deleting a relationship stops all future data syncs and terminates payment. Any data that was synced to the target remains as-is.

What happens if I delete something from my source server? Is it removed from the target too?

By default, if you have an active sync relationship, the item deleted on the source server is not deleted from the target during the next synchronization. But there is an option in the sync settings for each relationship, where you can define that Cloud Sync will delete files in the target location if they were deleted from the source.

[Learn how to change the settings for a sync relationship.](#)

What happens if I delete something from my target? Is it removed from my source too?

If an item is deleted from the target, it will not be removed from the source. The relationship is one-way—from source to target. On the next sync cycle, Cloud Sync compares the source to the target, identifies that the item is missing, and Cloud Sync copies it again from the source to the target.

Troubleshooting

[NetApp Knowledgebase: Cloud Sync FAQ: Support and Troubleshooting](#)

Data broker deep dive

The following question relates to the data broker.

Can you explain the architecture of the data broker?

Sure. Here are the most important points:

- The data broker is a node.js application running on a Linux host.
- Cloud Sync deploys the data broker as follows:
 - AWS: From an AWS CloudFormation template
 - Azure: From Azure Resource Manager
 - Google: From Google Cloud Deployment Manager
 - If you use your own Linux host, you need to manually install the software
- The data broker software automatically upgrades itself to the latest version.
- The data broker uses AWS SQS as a reliable and secure communication channel and for control and monitoring. SQS also provides a persistency layer.
- You can add additional data brokers to a relationship to increase transfer speed and add high availability. There is service resiliency if one data broker fails.

Gain insight into data privacy

Learn about Cloud Data Sense

Cloud Data Sense is a data governance service for Cloud Manager that scans your corporate on-premises and cloud data sources and working environments to map and classify data, and to identify private information. This can help reduce your security and compliance risk, decrease storage costs, and assist with your data migration projects.



Cloud Compliance was renamed to **Cloud Data Sense** in June 2021.

[Learn about the use cases for Cloud Data Sense.](#)

Features

Cloud Data Sense provides several tools that can help you with your compliance efforts. You can use Data Sense to:

- Identify Personal Identifiable Information (PII)
- Identify a wide scope of sensitive information as required by GDPR, CCPA, PCI, and HIPAA privacy regulations
- Respond to Data Subject Access Requests (DSAR)
- Notify Cloud Manager users through email when files contain certain PII (you define this criteria using [Policies](#))
- View and modify [Azure Information Protection \(AIP\) labels](#) in your files
- Add a custom tags to files (for example, "needs to be moved") and assign a Cloud Manager user so that person can own the change to the files
- Move and delete files

Cloud Data Sense also provides tools that can help with your governance efforts. You can use Cloud Data Sense to:

- Identify the stale data, non-business data, duplicate files, and very large files in your systems.
You can use this information to decide whether you want to move, delete, or tier some files to less expensive object storage.
- View the size of data and whether any of the data contains sensitive information prior to moving it.

This is useful if you are planning to migrate data from on-premises locations to the cloud.

Supported working environments and data sources

Cloud Data Sense can scan data from the following types of working environments and data sources:

- Cloud Volumes ONTAP in AWS
- Cloud Volumes ONTAP in Azure

- On-premises ONTAP clusters
- Azure NetApp Files
- Amazon FSx for ONTAP
- Amazon S3
- Non-NetApp file shares
- Object storage (that uses S3 protocol)
- Databases
- OneDrive accounts

 A Beta feature released in January 2021 allows you to run compliance scans *for free* on the backup files created from your on-prem ONTAP volumes (created using [Cloud Backup](#)). This gives you a choice whether you want to have Cloud Data Sense scan your on-prem ONTAP volumes directly, or scan the backup files made from those volumes.

Cost

- The cost to use Cloud Data Sense depends on the amount of data that you're scanning. The first 1 TB of data that Data Sense scans in a Cloud Manager workspace is free. This includes all data from all working environments and data sources. A subscription to the AWS or Azure Marketplace is required to continue scanning data after that point. See [pricing](#) for details.

[Learn how to subscribe.](#)

Note: This subscription is not needed to scan backup files created from your on-prem ONTAP systems.

- Installing Cloud Data Sense in the cloud requires deploying a cloud instance, which results in charges from the cloud provider where it is deployed. See [the type of instance that is deployed for each cloud provider](#). There is no cost if you install Data Sense on an on-premises system.
- Cloud Data Sense requires that you have deployed a Connector. In many cases you already have a Connector because of other storage and services you are using in Cloud Manager. The Connector instance results in charges from the cloud provider where it is deployed. See the [type of instance that is deployed for each cloud provider](#).

Data transfer costs

Data transfer costs depend on your setup. If the Cloud Data Sense instance and data source are in the same Availability Zone and region, then there are no data transfer costs. But if the data source, such as a Cloud Volumes ONTAP cluster or S3 Bucket, is in a *different* Availability Zone or region, then you'll be charged by your cloud provider for data transfer costs. See these links for more details:

- [AWS: Amazon EC2 Pricing](#)
- [Microsoft Azure: Bandwidth Pricing Details](#)

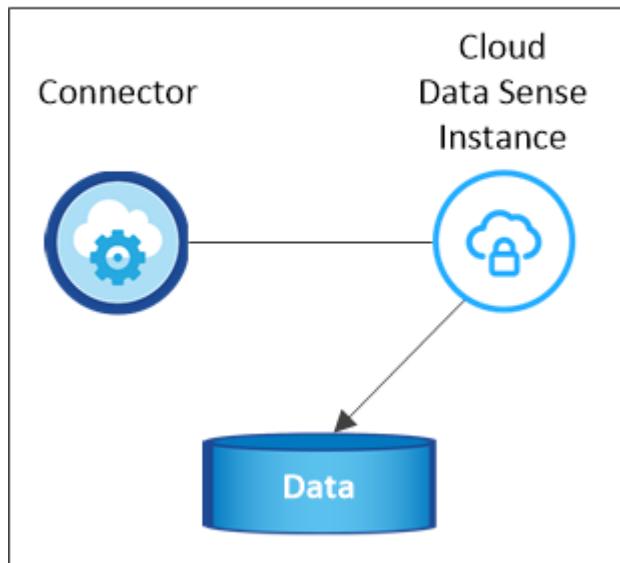
The Cloud Data Sense instance

When you deploy Data Sense in the cloud, Cloud Manager deploys the instance in the same subnet as the Connector. [Learn more about Connectors.](#)



If the Connector is installed on-prem, it deploys the Cloud Data Sense instance in same VPC or VNet as the first Cloud Volumes ONTAP system in the request. You can install Data Sense on-prem as well.

VPC or VNet



Note the following about the default instance:

- In AWS, Cloud Data Sense runs on an [m5.4xlarge instance](#) with a 500 GB GP2 disk. The operating system image is Amazon Linux 2 (Red Hat 7.3.1).

In regions where m5.4xlarge isn't available, Data Sense runs on an m4.4xlarge instance instead.

- In Azure, Cloud Data Sense runs on a [Standard_D16s_v3 VM](#) with a 512 GB disk. The operating system image is CentOS 7.8.
- The instance is named *CloudCompliance* with a generated hash (UUID) concatenated to it. For example: *CloudCompliance-16bb6564-38ad-4080-9a92-36f5fd2f71c7*
- Only one Data Sense instance is deployed per Connector.
- Upgrades of Data Sense software is automated—you don't need to worry about it.



The instance should remain running at all times because Cloud Data Sense continuously scans the data.

Using a smaller instance type

You can deploy Data Sense on a system with fewer CPUs and less RAM, but there are some limitations when using these less powerful systems.

System size	Specs	Limitations
Extra Large (default)	16 CPUs, 64 GB RAM	None
Medium	8 CPUs, 32 GB RAM	Slower scanning, and can only scan up to 1 million files.

System size	Specs	Limitations
Small	8 CPUs, 16 GB RAM	Same limitations as "Medium", plus the ability to identify data subject names inside files is disabled.

When deploying Data Sense in the cloud, email ng-contact-data-sense@netapp.com for assistance if you want to use one of these smaller systems.

When deploying Data Sense on-premises, just use a Linux host with these specifications.

How Cloud Data Sense works

At a high-level, Cloud Data Sense works like this:

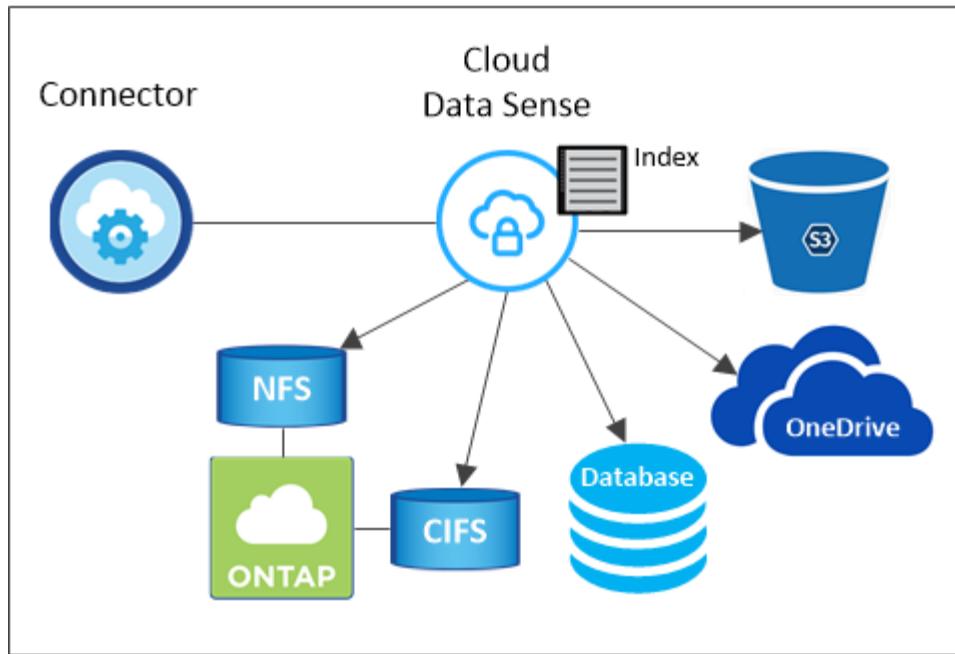
1. You deploy an instance of Data Sense in Cloud Manager.
2. You enable high-level mapping or deep-level scanning on one or more working environments or data sources.
3. Data Sense scans the data using an AI learning process.
4. You click **Data Sense** and use the provided dashboards and reporting tools to help in your compliance efforts.

How scans work

After you enable Cloud Data Sense and select the volumes, buckets, database schemas, or OneDrive users you want to scan, it immediately starts scanning the data to identify personal and sensitive data. It maps your organizational data, categorizes each file, and identifies and extracts entities and predefined patterns in the data. The result of the scan is an index of personal information, sensitive personal information, data categories, and file types.

Data Sense connects to the data like any other client by mounting NFS and CIFS volumes. NFS volumes are automatically accessed as read-only, while you need to provide Active Directory credentials to scan CIFS volumes.

VPC or VNet



After the initial scan, Data Sense continuously scans your data to detect incremental changes (this is why it's important to keep the instance running).

You can enable and disable scans at the volume level, at the bucket level, at the database schema level, and at the OneDrive user level.

What's the difference between Mapping and Classification scans

Cloud Data Sense enables you to run a general "mapping" scan on selected working environments and data sources. Mapping provides only a high-level overview of your data, whereas Classification provides deep-level scanning of your data. Mapping can be done on your data sources very quickly because it does not access files to see the data inside.

Many users like this functionality because they want to quickly scan their data to identify the data sources that require more research - and then they can enable classification scans only on those required data sources.

The table below shows some of the differences:

Feature	Classification	Mapping
Scan speed	Slow	Fast
List of file types and used capacity	Yes	Yes
Number of files and used capacity	Yes	Yes
Age and size of files	Yes	Yes
Ability to run a Data Mapping Report	Yes	Yes
Data Investigation page to view file details	Yes	No
Search for names within files	Yes	No
Create policies that provide custom search results	Yes	No

Feature	Classification	Mapping
Categorize data using AIP labels and Status tags	Yes	No
Delete and move source files	Yes	No
Ability to run other reports	Yes	No

Information that Cloud Data Sense indexes

Data Sense collects, indexes, and assigns categories to your data (files). The data that Data Sense indexes includes the following:

Standard metadata

Cloud Data Sense collects standard metadata about files: the file type, its size, creation and modification dates, and so on.

Personal data

Personally identifiable information such as email addresses, identification numbers, or credit card numbers. [Learn more about personal data](#).

Sensitive personal data

Special types of sensitive information, such as health data, ethnic origin, or political opinions, as defined by GDPR and other privacy regulations. [Learn more about sensitive personal data](#).

Categories

Cloud Data Sense takes the data that it scanned and divides it into different types of categories. Categories are topics based on AI analysis of the content and metadata of each file. [Learn more about categories](#).

Types

Cloud Data Sense takes the data that it scanned and breaks it down by file type. [Learn more about types](#).

Name entity recognition

Cloud Data Sense uses AI to extract natural persons' names from documents. [Learn about responding to Data Subject Access Requests](#).

Networking overview

Cloud Manager deploys the Cloud Data Sense instance with a security group that enables inbound HTTP connections from the Connector instance.

When using Cloud Manager in SaaS mode, the connection to Cloud Manager is served over HTTPS, and the private data sent between your browser and the Data Sense instance are secured with end-to-end encryption, which means NetApp and third parties can't read it.

If you need to use the local user interface instead of the SaaS user interface for any reason, you can still [access the local UI](#).

Outbound rules are completely open. Internet access is needed to install and upgrade the Data Sense software and to send usage metrics.

If you have strict networking requirements, [learn about the endpoints that Cloud Data Sense contacts](#).

User access to compliance information

The role each user has been assigned provides different capabilities within Cloud Manager and within Cloud Data Sense:

- An **Account Admin** can manage compliance settings and view compliance information for all working environments.
- A **Workspace Admin** can manage compliance settings and view compliance information only for systems that they have permissions to access. If a Workspace Admin can't access a working environment in Cloud Manager, then they can't see any compliance information for the working environment in the Data Sense tab.
- Users with the **Compliance Viewer** role can only view compliance information and generate reports for systems that they have permission to access. These users cannot enable/disable scanning of volumes, buckets, or database schemas.

[Learn more about Cloud Manager roles](#) and how to [add users with specific roles](#).

Get started

Deploy Cloud Data Sense

Complete a few steps to deploy Cloud Data Sense in your Cloud Manager workspace. You can deploy Data Sense in the cloud or on an on-premises system.

The on-prem installation may be a good option if you prefer to scan on-premises ONTAP systems using a Data Sense instance that's also located on premises — but this is not a requirement. The software functions exactly the same way regardless of which installation method you choose.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.



Create a Connector

If you don't already have a Connector, create a Connector in AWS or Azure. See [creating a Connector in AWS](#) or [creating a Connector in Azure](#).

You can also [deploy the Connector on-premises](#) on an existing Linux host in your network or in the cloud.



Review prerequisites

Ensure that your environment can meet the prerequisites. This includes outbound internet access for the instance, connectivity between the Connector and Cloud Data Sense over port 80, and more. [See the complete list](#).

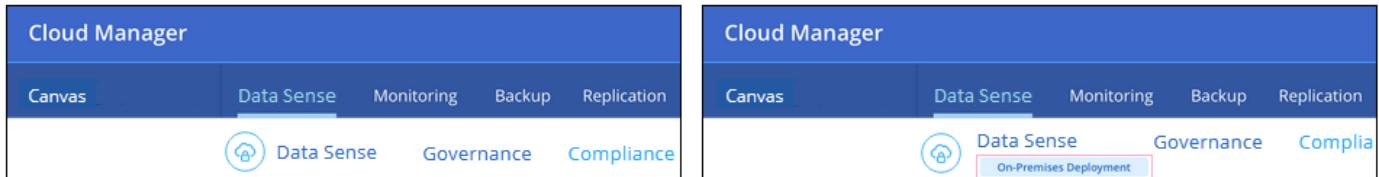
- When installed in the cloud, the default configuration requires 16 vCPUs for the Cloud Data Sense instance. See [more details about the instance type](#).
- When installed on premises, you need a Linux system that meets the [following requirements](#).

3

Deploy Cloud Data Sense

Launch the installation wizard to deploy the Cloud Data Sense instance.

You can deploy Cloud Data Sense in the cloud or in an on-premises location. The only difference you'll notice in the UI is the words "On-Premises Deployment".



4

Subscribe to the Cloud Data Sense service

The first 1 TB of data that Cloud Data Sense scans in Cloud Manager is free. A subscription to the AWS or Azure Marketplace is required to continue scanning data after that point.

Creating a Connector

If you don't already have a Connector, create a Connector in AWS or Azure. See [creating a Connector in AWS](#) or [creating a Connector in Azure](#). In most cases you will probably have a Connector set up before you attempt to activate Cloud Data Sense because most [Cloud Manager features require a Connector](#), but there are cases where you'll need to set one up now.

There are some scenarios where you have to use a Connector that's deployed in a specific cloud provider:

- When scanning data in Cloud Volumes ONTAP in AWS, Amazon FSx for ONTAP, or in AWS S3 buckets, you use a connector in AWS.
- When scanning data in Cloud Volumes ONTAP in Azure or in Azure NetApp Files, you use a connector in Azure.

On-prem ONTAP systems, non-NetApp file shares, generic S3 Object storage, databases, and OneDrive folders can be scanned using any of these cloud Connectors.

Note that you can also [deploy the Connector on-premises](#) on an existing Linux host in your network or in the cloud. Some users planning to install Data Sense on-prem may also choose to install the Connector on-prem.

As you can see, there may be some situations where you need to use [multiple Connectors](#).



If you're planning on scanning Azure NetApp Files volumes, you need to make sure you're deploying in the same region as the volumes you wish to scan.

Reviewing prerequisites

Review the following prerequisites to make sure that you have a supported configuration before you deploy Cloud Data Sense.

Enable outbound internet access from Cloud Data Sense

Cloud Data Sense requires outbound internet access. If your virtual or physical network uses a proxy server for internet access, ensure that the Data Sense instance has outbound internet access to contact the following endpoints. When you deploy Data Sense in the cloud, it's located in the same subnet as the Connector.

Review the appropriate table below depending on whether you are deploying Cloud Data Sense in AWS, Azure, or on-premises.

Required endpoints for AWS deployments:

Endpoints	Purpose
https://cloudmanager.cloud.netapp.com	Communication with the Cloud Manager service, which includes Cloud Central accounts.
https://netapp-cloud-account.auth0.com https://auth0.com	Communication with NetApp Cloud Central for centralized user authentication.
https://cloud-compliance-support-netapp.s3.us-west-2.amazonaws.com https://hub.docker.com https://auth.docker.io https://registry-1.docker.io https://index.docker.io/ https://dseasb33srnrrn.cloudfront.net/ https://production.cloudflare.docker.com/	Provides access to software images, manifests, and templates.
https://kinesis.us-east-1.amazonaws.com	Enables NetApp to stream data from audit records.
https://cognito-idp.us-east-1.amazonaws.com https://cognito-identity.us-east-1.amazonaws.com https://user-feedback-store-prod.s3.us-west-2.amazonaws.com https://customer-data-production.s3.us-west-2.amazonaws.com	Enables Cloud Data Sense to access and download manifests and templates, and to send logs and metrics.

Required endpoints for Azure and On-Prem deployments:

Endpoints	Purpose
https://cloudmanager.cloud.netapp.com	Communication with the Cloud Manager service, which includes Cloud Central accounts.
https://netapp-cloud-account.auth0.com https://auth0.com	Communication with NetApp Cloud Central for centralized user authentication.

Endpoints	Purpose
https://support.compliance.cloudmanager.cloud.netapp.com/ https://hub.docker.com https://auth.docker.io https://registry-1.docker.io https://index.docker.io/ https://dseasb33srnrrn.cloudfront.net/ https://production.cloudflare.docker.com/	Provides access to software images, manifests, templates, and to send logs and metrics.
https://support.compliance.cloudmanager.cloud.netapp.com/	Enables NetApp to stream data from audit records.
On-premises installs only: https://github.com/docker https://download.docker.com https://rhui3.us-west-2.aws.ce.redhat.com https://github-production-release-asset-2e65be.s3.amazonaws.com https://pypi.org https://pypi.python.org https://files.pythonhosted.org http://mirror.centos.org http://mirrorlist.centos.org http://mirror.centos.org/centos/7/extras/x86_64/Packages/container-selinux-2.107-3.el7.noarch.rpm	Provides prerequisite packages for installation.

Ensure that Cloud Manager has the required permissions

Ensure that Cloud Manager has permissions to deploy resources and create security groups for the Cloud Data Sense instance. You can find the latest Cloud Manager permissions in [the policies provided by NetApp](#).

Check your vCPU limits

When installed in the cloud, ensure that your cloud provider's vCPU limit allows for the deployment of an instance with 16 cores. You'll need to verify the vCPU limit for the relevant instance family in the region where Cloud Manager is running. [See the required instance types](#).

See the following links for more details on vCPU limits:

- [AWS documentation: Amazon EC2 service quotas](#)
- [Azure documentation: Virtual machine vCPU quotas](#)

Note that you can deploy Data Sense on a system with fewer CPUs and less RAM, but there are limitations when using these systems. See [Using a smaller instance type](#) for details.

Ensure that Cloud Manager can access Cloud Data Sense

Ensure connectivity between the Connector and the Cloud Data Sense instance. The security group for the Connector must allow inbound and outbound traffic over port 80 to and from the Data Sense instance.

This connection enables deployment of the Data Sense instance and enables you to view information in the Compliance and Governance tabs.

Ensure that you can keep Cloud Data Sense running

The Cloud Data Sense instance needs to stay on to continuously scan your data.

Ensure web browser connectivity to Cloud Data Sense

After Cloud Data Sense is enabled, ensure that users access the Cloud Manager interface from a host that has a connection to the Data Sense instance.

The Data Sense instance uses a private IP address to ensure that the indexed data isn't accessible to the internet. As a result, the web browser that you use to access Cloud Manager must have a connection to that private IP address. That connection can come from a direct connection to AWS or Azure (for example, a VPN), or from a host that's inside the same network as the Data Sense instance.

Deploying the Cloud Data Sense instance in the cloud

Deploying an instance of Cloud Data Sense in the cloud is the most common deployment model. But you have the option to [deploy the Compliance software on a Linux host](#) in your network or in the cloud.

The Data Sense software functions exactly the same way regardless of which installation method you choose.

Steps

1. In Cloud Manager, click **Data Sense**.
2. Click **Activate Cloud Data Sense**.

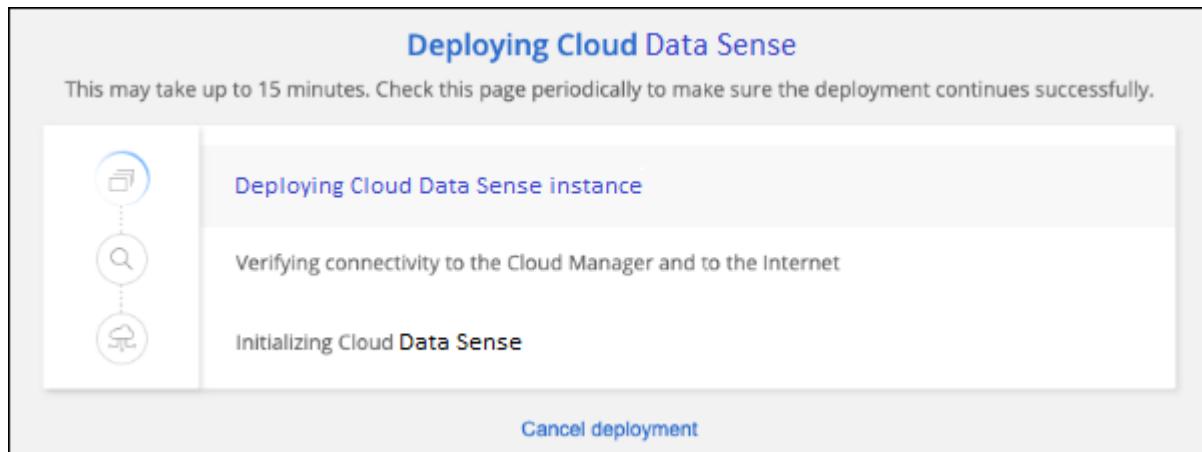
The screenshot shows the Cloud Manager interface with the 'Data Sense' section selected. On the left, there's a 'How does it work?' link and a large heading 'Always-on Privacy & Compliance Controls'. Below this, there's a brief description of automated controls for data privacy regulations like GDPR, CCPA, and HIPAA. A prominent blue button labeled 'Activate Cloud Data Sense' is highlighted with a red box. To the right, there's a 'Compliance Status' section with a circular progress bar, a 'Data Distribution' chart, and several data statistics tables.

Category	Value
Personal Files	28,000
Email Address	2,700 Files
Credit Card	2,700 Files
Sensitive Personal Files	7,000
Health	2,700 Files
Ethnicity	2,700 Files

3. Click **Activate Data Sense** to start the cloud deployment wizard.

The screenshot shows the 'Select where to deploy Data Sense' wizard. It has two main options: 'Deploy Data Sense in the Cloud' (which is 'Recommended') and 'Deploy Data Sense On-Premises'. Each option has an associated icon and an 'Activate Data Sense' button. A note below the first option states: 'We recommend deploying Data Sense in the Cloud. Selecting this option will deploy the instance in the same location as the Cloud Manager Connector instance.' There are also collapse/expand arrows on the right side of the wizard.

4. The wizard displays progress as it goes through the deployment steps. It will stop and ask for input if it runs into any issues.



5. When the instance is deployed, click **Continue to configuration** to go to the *Configuration* page.

Result

Cloud Manager deploys the Cloud Data Sense instance in your cloud provider.

What's Next

From the Configuration page you can select the data sources that you want to scan.

You can also [subscribe to the Cloud Data Sense service](#) at this time. You will not be charged until the amount of data exceeds 1 TB.

Deploying the Cloud Data Sense instance on premises

You can download and install the Data Sense software on a Linux host in your network if you do not want to [deploy it in the cloud](#).

The Data Sense software functions exactly the same way regardless of which installation method you choose.

For typical configurations you'll install the software on a single host system. For very large configurations where you'll be scanning petabytes of data, you can include additional hosts as *scanner nodes* to provide additional processing power.



Cloud Data Sense is currently unable to scan S3 buckets and Azure NetApp Files when the software is installed on premises. In these cases you'll need to deploy a separate Connector and instance of Data Sense in the cloud and [switch between Connectors](#) for your different data sources.

Host requirements

- Operating system: Red Hat Enterprise Linux or CentOS version 8.0 or 8.1
 - Version 7.8 can be used, but the Linux kernel version must be 4.14 or greater
 - The OS must be capable of installing the docker engine (for example, disable the *firewalld* service if needed)
- RAM: 64 GB (swap memory must be disabled on the host)

- CPU: 16 cores
- Disk: 500 GB SSD

Note that you can deploy Data Sense on a system with fewer CPUs and less RAM, but there are limitations when using these systems. See [Using a smaller instance type](#) for details.

- A Red Hat Enterprise Linux system must be registered with Red Hat Subscription Management. If it's not registered, the system can't access repositories to update required 3rd party software during installation.
- Make sure port 8080 is open so you can see the installation progress in Cloud Manager.
- Root privileges are required to install Cloud Data Sense.

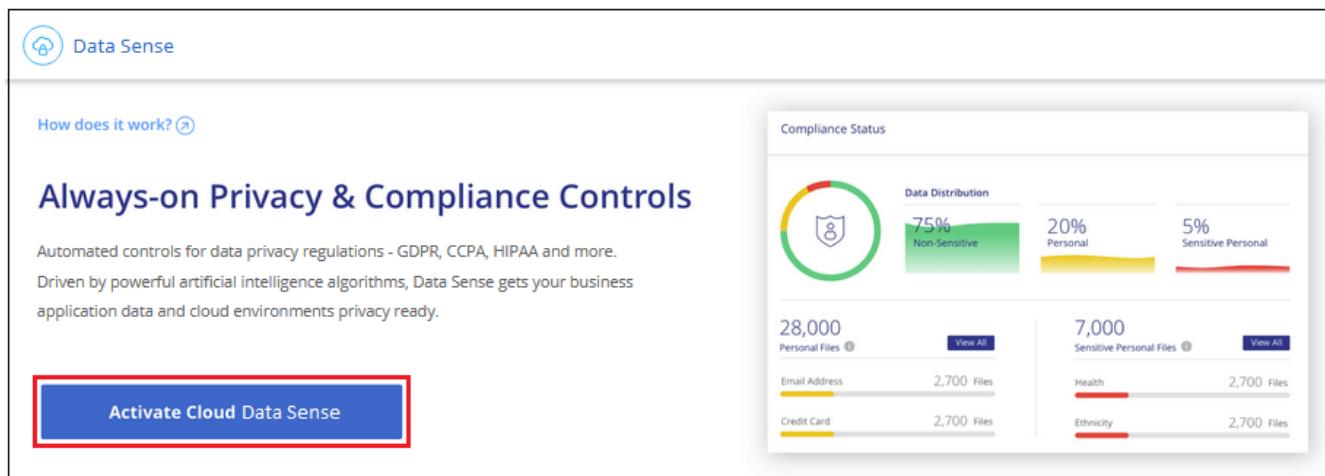
See [Reviewing prerequisites](#) for the full list of requirements and endpoints that Cloud Data Sense must be able to reach over the internet.

Single-host installation for typical configurations

Follow these steps when installing Data Sense software on a single on-premises host.

Steps

1. Download the Cloud Data Sense software from the [NetApp Support Site](#).
2. Copy the installer file to the Linux host you plan to use (using `scp` or some other method).
3. In Cloud Manager, click **Data Sense**.
4. Click **Activate Cloud Data Sense**.

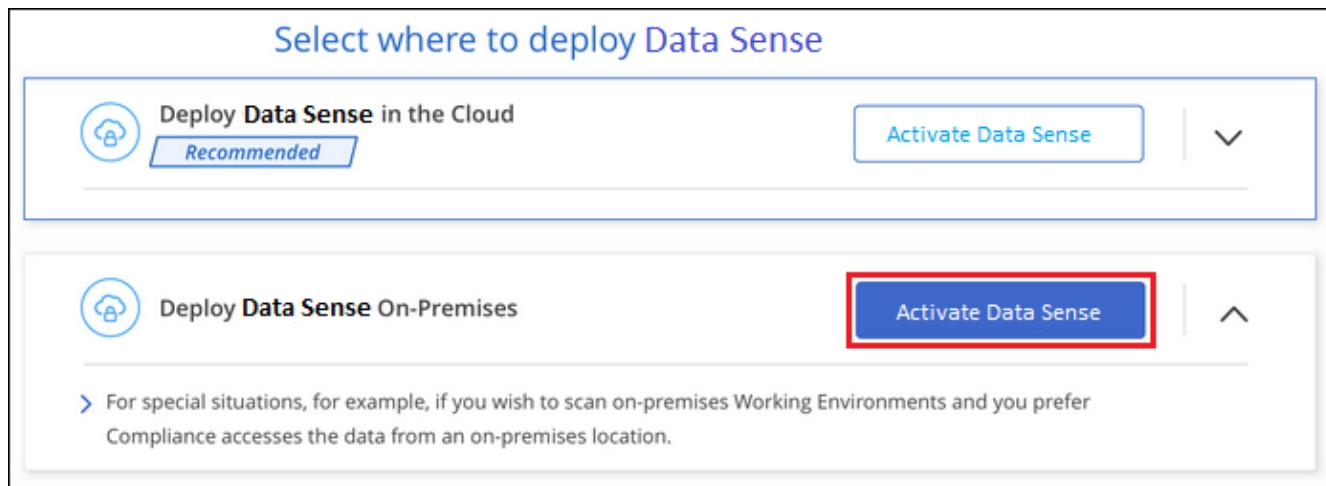


The screenshot shows the Cloud Manager interface for Data Sense. At the top, there's a navigation bar with a cloud icon and the text "Data Sense". Below it, a link "How does it work? ⓘ" is visible. The main content area has a heading "Always-on Privacy & Compliance Controls" with the subtext "Automated controls for data privacy regulations - GDPR, CCPA, HIPAA and more. Driven by powerful artificial intelligence algorithms, Data Sense gets your business application data and cloud environments privacy ready." To the right, there's a "Compliance Status" section featuring a circular progress bar with a shield icon, a "Data Distribution" chart, and several data summary cards:

Category	Value	Details
Personal Files	28,000	View All
Email Address	2,700 Files	
Credit Card	2,700 Files	
Sensitive Personal Files	7,000	View All
Health	2,700 Files	
Ethnicity	2,700 Files	

A large blue button at the bottom left is labeled "Activate Cloud Data Sense", which is highlighted with a red rectangular box.

5. Click **Activate Data Sense** to start the on-prem deployment wizard.



6. In the *Deploy Cloud Data Sense On Premises* dialog, copy the provided command and paste it in a text file so you can use it later. For example:

```
sudo ./install.sh -a 12345 -c 27AG75 -t 2198qq
```

7. Unzip the installer file on the host machine:

```
tar -xzf cc_onprem_installer.tar.gz
```

8. When prompted by the installer, you can enter the required values in a series of prompts, or you can enter the complete command in the first prompt:

Enter parameters as prompted:	Enter the full command:
<ol style="list-style-type: none"> Paste the information you copied from step 6: <code>sudo ./install.sh -a <account_id> -c <agent_id> -t <token></code> Enter the IP address or host name of the Data Sense host machine so it can be accessed by the Connector instance. Enter the IP address or host name of the Cloud Manager Connector host machine so it can be accessed by the Data Sense instance. Enter proxy details as prompted. If your Cloud Manager already uses a proxy, there is no need to enter this information again here since Data Sense will automatically use the proxy used by Cloud Manager. 	<p>Alternatively, you can create the whole command in advance and enter it in the first prompt:</p> <pre>sudo ./install.sh -a <account_id> -c <agent_id> -t <token> --host <ds_host> --cm-host <cm_host> --proxy-host <proxy_host> --proxy-port <proxy_port> --proxy-scheme <proxy_scheme> --proxy-user <proxy_user> --proxy-password <proxy_password></pre>

Variable values:

- *account_id* = NetApp Account ID
- *agent_id* = Connector ID
- *token* = jwt user token
- *ds_host* = IP address or host name of the Data Sense Linux system.

- *cm_host* = IP address or host name of the Cloud Manager Connector system.
- *proxy_host* = IP or host name of the proxy server if the host is behind a proxy server.
- *proxy_port* = Port to connect to the proxy server (default 80).
- *proxy_scheme* = Connection scheme: https or http (default http).
- *proxy_user* = Authenticated user to connect to the proxy server, if basic authentication is required.
- *proxy_password* = Password for the user name that you specified.

Result

The Cloud Data Sense installer installs packages, installs docker, registers the installation, and installs Data Sense. Installation can take 10 to 20 minutes.

If there is connectivity over port 8080 between the host machine and the Connector instance, you will see the installation progress in the Data Sense tab in Cloud Manager.

What's Next

From the Configuration page you can select the data sources that you want to scan.

You can also [subscribe to the Cloud Data Sense service](#) at this time. You will not be charged until the amount of data exceeds 1 TB. A subscription to the AWS or Azure Marketplace can be used when you have deployed Data Sense on an on-premises system.

Multi-host installation for large configurations

Follow these steps when installing Data Sense software on multiple on-premises hosts.

When using multiple host systems, the primary system is called the *Manager node* and the additional systems that provide extra processing power are call *Scanner nodes*.

Requirements

- See [Reviewing prerequisites](#) for the full list of requirements and endpoints that Cloud Data Sense must be able to reach over the internet.
- The host requirements are the same for Scanner nodes as they are for Manager nodes. See [Host requirements](#) for details.
- You must have the IP addresses of the scanner node hosts that you plan to use.
- The following ports and protocols must be enabled on all hosts:

Port	Protocols	Description
2377	TCP	Cluster management communications
7946	TCP, UDP	Inter-node communication
4789	UDP	Overlay network traffic
50	ESP	Encrypted IPsec overlay network (ESP) traffic
111	TCP, UDP	NFS Server for sharing files between the hosts (needed from each scanner node to manager node)
2049	TCP, UDP	NFS Server for sharing files between the hosts (needed from each scanner node to manager node)

Steps

1. Follow steps 1 through 7 from the [Single-host installation](#) on the manager node.
2. As shown in step 8, when prompted by the installer, you can enter the required values in a series of prompts, or you can enter the complete command in the first prompt.

In addition to the variables available for a single-host installation, a new option **-n <node_ip>** is used to specify the IP addresses of the scanner nodes. Multiple node IPs are separated by a comma.

For example, this command adds 3 scanner nodes:

```
sudo ./install.sh -a <account_id> -c <agent_id> -t <token> --host <ds_host>
--cm-host <cm_host> -n <node_ip1>,<node_ip2>,<node_ip3> --proxy-host
<proxy_host> --proxy-port <proxy_port> --proxy-scheme <proxy_scheme> --proxy
--user <proxy_user> --proxy-password <proxy_password>
```

3. Before the manager node installation completes, a dialog displays the installation command needed for the scanner nodes. Copy the command and save it in a text file. For example:

```
sudo ./node_install.sh -m 10.11.12.13 -t ABCDEF-1-3u69m1-1s35212
```

4. On **each** scanner node host:

- a. Copy the Data Sense installer file (*cc_onprem_installer.tar.gz*) to the host machine (using `scp` or some other method).
- b. Unzip the installer file.
- c. Paste and execute the command that you copied in step 3.

When the installation finishes on all scanner nodes and they have been joined to the manager node, the manager node installation finishes as well.

Result

The Cloud Data Sense installer finishes installing packages, docker, and registers the installation. Installation can take 10 to 20 minutes.

What's Next

From the Configuration page you can select the data sources that you want to scan.

You can also [subscribe to the Cloud Data Sense service](#) at this time. You will not be charged until the amount of data exceeds 1 TB. A subscription to the AWS or Azure Marketplace can be used when you have deployed Data Sense on an on-premises system.

Subscribing to the Cloud Data Sense service

The first 1 TB of data that Cloud Data Sense scans in a Cloud Manager workspace is free. A subscription to the AWS or Azure Marketplace is required to continue scanning data after that point.

You can subscribe at any time and you will not be charged until the amount of data exceeds 1 TB. You can always see the total amount of data that is being scanned from the Data Sense Dashboard. And the *Subscribe Now* button makes it easy to subscribe when you are ready.

Note: If you are prompted by Cloud Data Sense to subscribe, but you already have an Azure subscription, you're probably using the old **Cloud Manager** subscription and you need to change to the new **NetApp Cloud Manager** subscription. See [Changing to the new NetApp Cloud Manager plan in Azure](#) for details.

Steps

These steps must be completed by a user who has the *Account Admin* role.

1. In the upper right of the Cloud Manager console, click the Settings icon, and select **Credentials**.

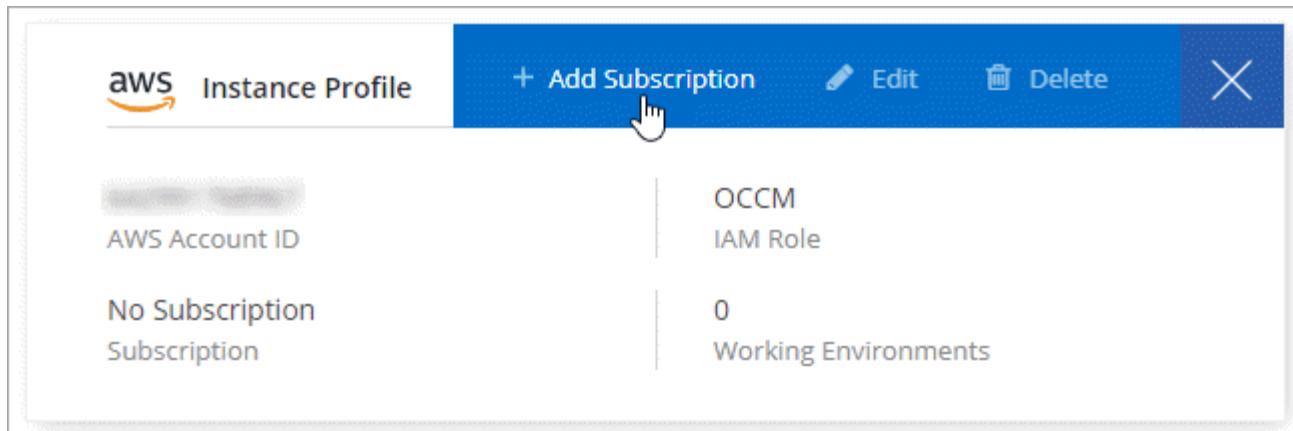


2. Find the credentials for the AWS Instance Profile or Azure Managed Service Identity.

The subscription must be added to the Instance Profile or Managed Service Identity. Charging won't work otherwise.

If you already have a subscription (shown below), then you're all set—there's nothing else that you need to do.

3. If you don't have a subscription yet, hover over the credentials and click the action menu.
4. Click **Add Subscription**.



5. Click **Add Subscription**, click **Continue**, and follow the steps.

The following video shows how to associate a Marketplace subscription to an AWS subscription:

- ▶ https://docs.netapp.com/us-en/occm//media/video_subscribing_aws.mp4 (video)

The following video shows how to associate a Marketplace subscription to an Azure subscription:

- ▶ https://docs.netapp.com/us-en/occm//media/video_subscribing_azure.mp4 (video)

Changing to the new Cloud Manager plan in Azure

Cloud Data Sense (Cloud Compliance) was added to the Azure Marketplace subscription named **NetApp Cloud Manager** as of October 2020. If you already have the original Azure **Cloud Manager** subscription it will not allow you to use Cloud Data Sense.

You need to follow these steps to change to the new **NetApp Cloud Manager** subscription before you can start using Cloud Data Sense.



If your existing Subscription was issued with a special private offer, you need to contact NetApp so that we can issue a new special private offer with Data Sense included.

Steps

1. In the upper right of the Cloud Manager console, click the Settings icon, and select **Credentials**.
2. Find the credentials for the Azure Managed Service Identity that you want to change the subscription for and hover over the credentials and click **Associate Subscription**.

The details for your current Marketplace Subscription are displayed.

3. Log in to the [Azure portal](#) and select **Software as a Service (SaaS)**.
4. Select the subscription for which you want to change the plan and click **Change Plan**.

The screenshot shows the Microsoft Azure portal interface. On the left, there's a list of SaaS subscriptions under 'Software as a Service (SaaS)'. One subscription, 'shiranSub3008', is selected and highlighted with a red box. The main pane displays 'Offer and plan details' for this subscription. It shows the 'Cloud Manager - Monthly' plan from NetApp, which is currently 'Subscribed'. Below this, there's a link to 'View this offer in Marketplace'. To the right, there's a 'Billing term & price' section showing 'Monthly' billing at '\$0.00 per month'. Underneath, there's a 'Plus:' section listing various cloud storage and management services with their respective prices.

5. In the Change Plan page, select the **NetApp Cloud Manager** plan and click the **Change Plan** button.

The screenshot shows the 'Change plan' page. It lists two software plans: 'NetApp Cloud Manager' (selected and highlighted with a red box) and 'Cloud Manager'. The 'NetApp Cloud Manager' plan is described as 'PLAN - INCLUDES DATA SENSE' and costs '\$0.00 per month'. The 'Cloud Manager' plan is described as 'OLD PLAN - DOES NOT INCLUDE DATA SENSE' and also costs '\$0.00 per month'. Both plans include a 'Current plan' note. At the bottom of the page, there are 'Change plan' and 'Cancel' buttons, with the 'Change plan' button highlighted with a red box.

6. Return to Cloud Manager, select the subscription, and hover over the "i" above subscription in the Credentials card to verify your subscription has changed.

Activate scanning on your data sources

Getting started with Cloud Data Sense for Cloud Volumes ONTAP and on-premises ONTAP

Complete a few steps to get started with Cloud Data Sense for Cloud Volumes ONTAP and on-premises ONTAP systems.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.



Discover the data sources that contain the data you want to scan

Before you can scan volumes, you must add the systems as working environments in Cloud Manager:

- For Cloud Volumes ONTAP systems, these working environments should already be available in Cloud Manager
- For on-premises ONTAP systems, [Cloud Manager must discover the ONTAP clusters](#)



Deploy the Cloud Data Sense instance

[Deploy Cloud Data Sense in Cloud Manager](#) if there isn't already an instance deployed.



Enable Cloud Data Sense and select the volumes to scan

Click **Data Sense**, select the **Configuration** tab, and activate compliance scans for volumes in specific working environments.



Ensure access to volumes

Now that Cloud Data Sense is enabled, ensure that it can access all volumes.

- The Cloud Data Sense instance needs a network connection to each Cloud Volumes ONTAP subnet or on-prem ONTAP system.
- Security groups for Cloud Volumes ONTAP must allow inbound connections from the Data Sense instance.
- Make sure these ports are open to the Data Sense instance:
 - For NFS – ports 111 and 2049.
 - For CIFS – ports 139 and 445.
- NFS volume export policies must allow access from the Data Sense instance.
- Data Sense needs Active Directory credentials to scan CIFS volumes.

[Click Compliance > Configuration > Edit CIFS Credentials](#) and provide the credentials.

5

Manage the volumes you want to scan

Select or deselect the volumes that you want to scan and Cloud Data Sense will start or stop scanning them.

Discovering the data sources that you want to scan

If the data sources you want to scan are not already in your Cloud Manager environment, you can add them to the canvas at this time.

Your Cloud Volumes ONTAP systems should already be available in the Canvas in Cloud Manager. For on-premises ONTAP systems you need to have [Cloud Manager discover these clusters](#).

Deploying the Cloud Data Sense instance

[Deploy Cloud Data Sense](#) if there isn't already an instance deployed.

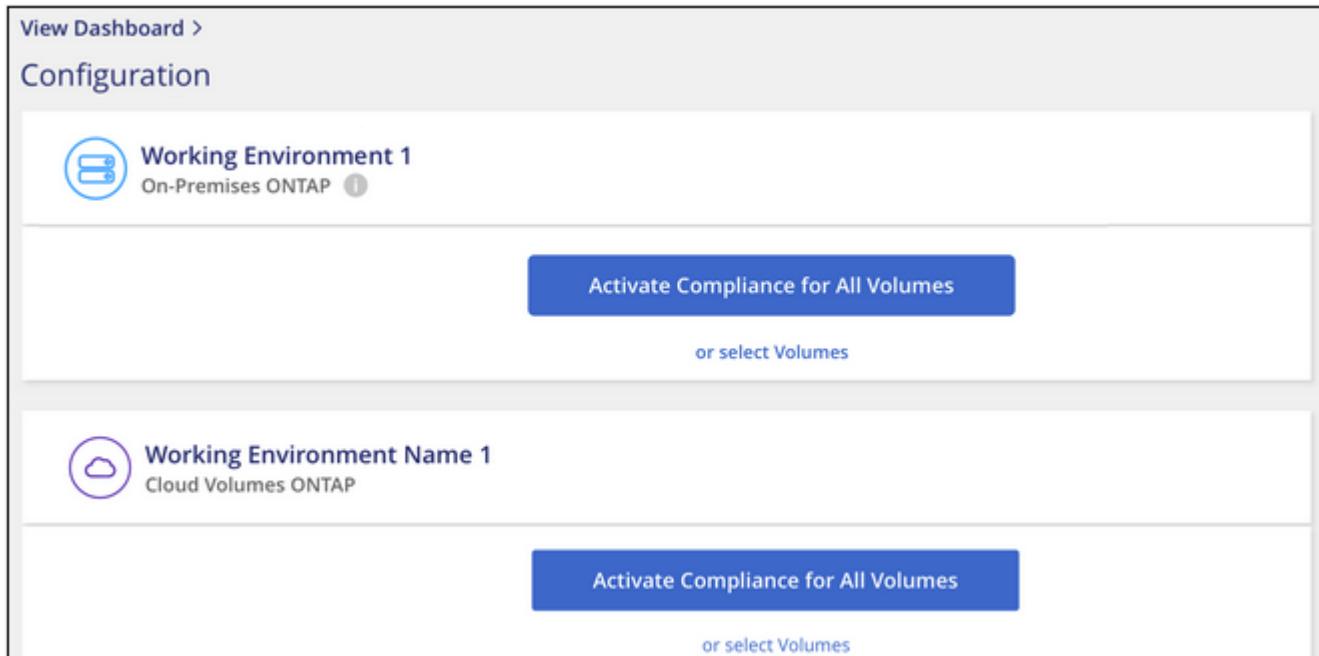
Cloud Data Sense can be deployed in the cloud or in an on-premises location when scanning Cloud Volumes ONTAP or on-premises ONTAP systems.

Enabling Cloud Data Sense in your working environments

You can enable Cloud Data Sense on Cloud Volumes ONTAP systems (in AWS or Azure) and on on-premises ONTAP clusters.

 Following these steps for on-prem ONTAP systems scans the volumes directly on the on-prem ONTAP system. If you are already creating backup files from those on-prem systems using [Cloud Backup](#), you can run compliance scans on the backup files in the cloud instead. Go to [Scanning backup files from on-premises ONTAP systems](#) to scan the volumes by scanning the backup files.

1. At the top of Cloud Manager, click **Data Sense** and then select the **Configuration** tab.



The screenshot shows the Cloud Manager Configuration tab. It displays two working environments:

- Working Environment 1**: On-Premises ONTAP. It features a blue icon with three horizontal bars. Below the icon, the text "Working Environment 1" and "On-Premises ONTAP" is displayed. A blue button labeled "Activate Compliance for All Volumes" is present, along with a link "or select Volumes".
- Working Environment Name 1**: Cloud Volumes ONTAP. It features a purple icon with a cloud shape. Below the icon, the text "Working Environment Name 1" and "Cloud Volumes ONTAP" is displayed. A blue button labeled "Activate Compliance for All Volumes" is present, along with a link "or select Volumes".

2. To scan all volumes in a working environment, click **Activate Scanning for All Volumes**.

When enabled in this manner, full "mapping and classification" scanning is performed on all volumes.

If you want to enable scanning only for certain volumes, or if you only want to perform "mapping-only" scanning, click **or select Volumes** and then choose the volumes you want to scan.

See [Enabling and disabling compliance scans on volumes](#) for details.

Result

Cloud Data Sense starts scanning the volumes you selected in the working environment. Results will be available in the Compliance dashboard as soon as Cloud Data Sense finishes the initial scans. The time that it takes depends on the amount of data—it could be a few minutes or hours.

Verifying that Cloud Data Sense has access to volumes

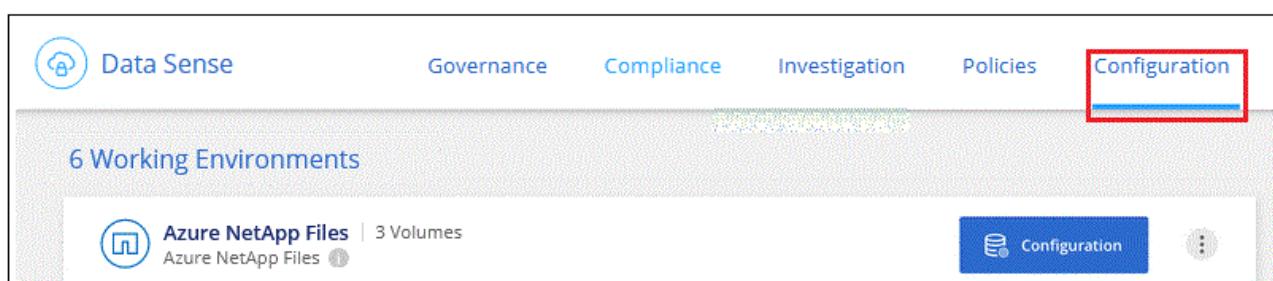
Make sure that Cloud Data Sense can access volumes by checking your networking, security groups, and export policies. You'll need to provide Data Sense with CIFS credentials so it can access CIFS volumes.

Steps

1. Make sure that there's a network connection between the Cloud Data Sense instance and each network that includes volumes for Cloud Volumes ONTAP or on-prem ONTAP clusters.
2. Ensure that the security group for Cloud Volumes ONTAP allows inbound traffic from the Data Sense instance.

You can either open the security group for traffic from the IP address of the Data Sense instance, or you can open the security group for all traffic from inside the virtual network.

3. Ensure the following ports are open to the Data Sense instance:
 - For NFS – ports 111 and 2049.
 - For CIFS – ports 139 and 445.
4. Ensure that NFS volume export policies include the IP address of the Data Sense instance so it can access the data on each volume.
5. If you use CIFS, provide Data Sense with Active Directory credentials so it can scan CIFS volumes.
 - a. At the top of Cloud Manager, click **Data Sense**.
 - b. Click the **Configuration** tab.

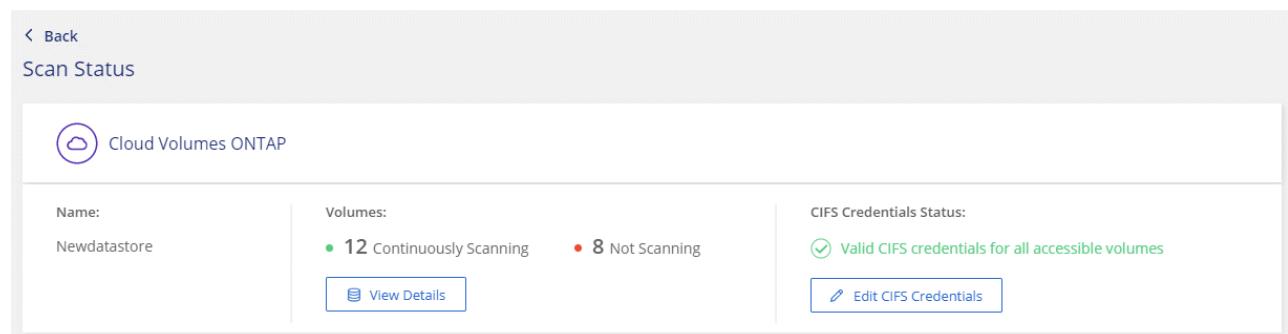


- c. For each working environment, click **Edit CIFS Credentials** and enter the user name and password that Data Sense needs to access CIFS volumes on the system.

The credentials can be read-only, but providing admin credentials ensures that Data Sense can read any data that requires elevated permissions. The credentials are stored on the Cloud Data Sense

instance.

After you enter the credentials, you should see a message that all CIFS volumes were authenticated successfully.



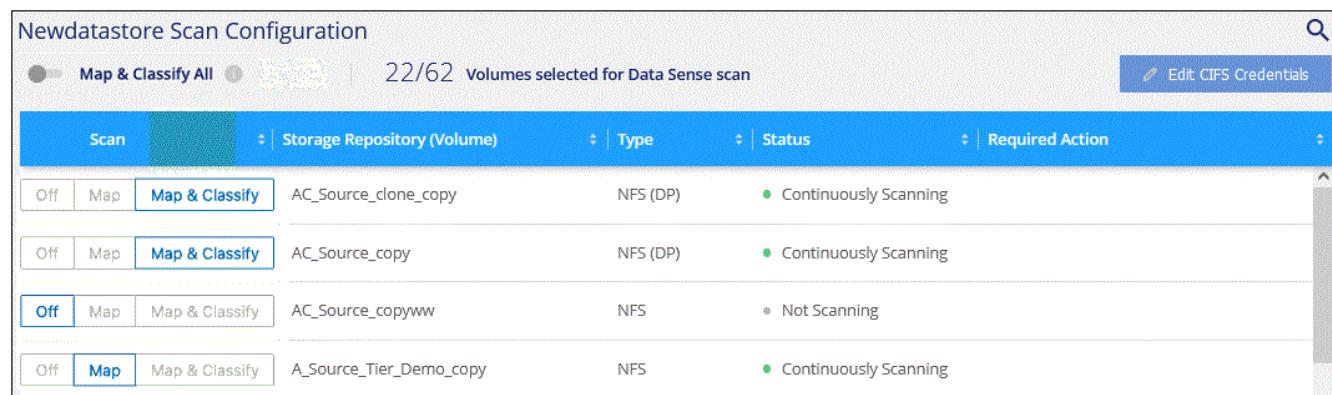
Name: Newdatastore

Volumes: 12 Continuously Scanning, 8 Not Scanning

CIFS Credentials Status: Valid CIFS credentials for all accessible volumes

- On the **Configuration** page, click **View Details** to review the status for each CIFS and NFS volume and correct any errors.

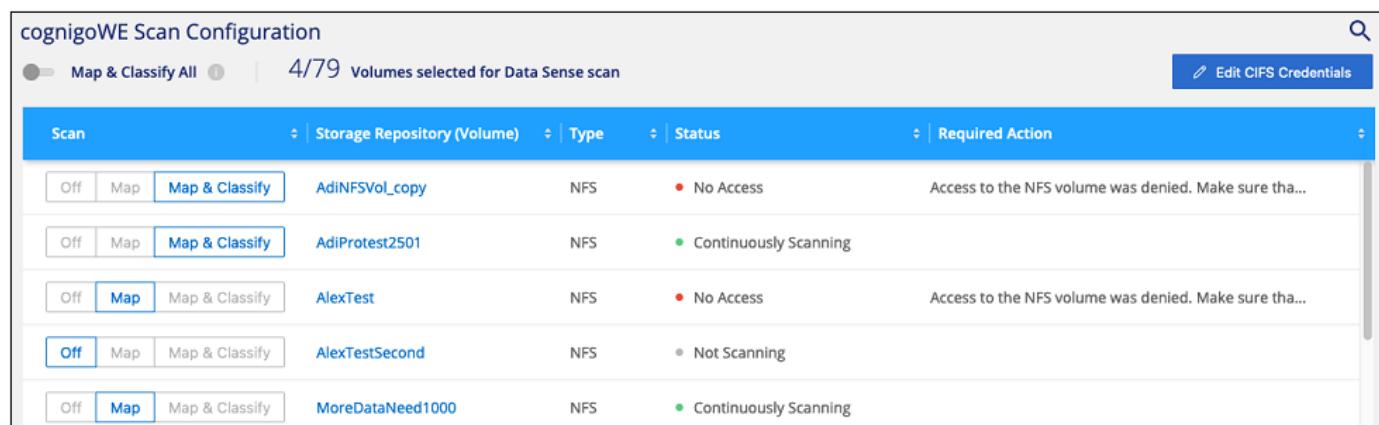
For example, the following image shows three volumes; one of which Cloud Data Sense can't scan due to network connectivity issues between the Data Sense instance and the volume.



Scan	Storage Repository (Volume)	Type	Status	Required Action
Off	AC_Source_clone_copy	NFS (DP)	Continuously Scanning	
Off	AC_Source_copy	NFS (DP)	Continuously Scanning	
Off	AC_Source_copyww	NFS	Not Scanning	
Off	A_Source_Tier_Demo_copy	NFS	Continuously Scanning	

Enabling and disabling compliance scans on volumes

You can stop or start mapping-only scans, or mapping and classification scans, in a working environment at any time from the Configuration page. We recommend that you scan all volumes.



Scan	Storage Repository (Volume)	Type	Status	Required Action
Off	AdINFSVol_copy	NFS	No Access	Access to the NFS volume was denied. Make sure tha...
Off	AdIProtect2501	NFS	Continuously Scanning	
Off	AlexTest	NFS	No Access	Access to the NFS volume was denied. Make sure tha...
Off	AlexTestSecond	NFS	Not Scanning	
Off	MoreDataNeed1000	NFS	Continuously Scanning	

To:	Do this:
Enable mapping-only scans on a volume	Click Map
Enable full scanning on a volume	Click Map & Classify
Enable full scanning on all volumes	Move the Map & Classify All slider to the right
Disable scanning on a volume	Click Off
Disable scanning on all volumes	Move the Map & Classify All slider to the left



New volumes added to the working environment are automatically scanned only when the **Map & Classify All** setting is enabled. When this setting is disabled, you'll need to activate mapping and/or full scanning on each new volume you create in the working environment.

Scanning backup files from on-premises ONTAP systems

If you don't want Cloud Data Sense to scan volumes directly on your on-prem ONTAP systems, a Beta feature released in January 2021 allows you to run compliance scans on backup files created from your on-prem ONTAP volumes. So if you're already creating backup files using [Cloud Backup](#), you can use this feature to run compliance scans on those backup files.

The Compliance scans you run on backup files are **free** - no Cloud Data Sense subscription or license is needed.

Note: When Data Sense scans backup files it uses permissions granted through the Cloud Restore instance to access the backup files. Typically the Restore instance powers down when not actively restoring files, but it remains **On** when scanning backup files. See [more information about the Restore instance](#).

Steps

If you want to scan the backup files from on-prem ONTAP systems:

1. At the top of Cloud Manager, click **Data Sense** and then select the **Configuration** tab.
2. From the list of working environments, click the **BACKUP** button from the list of filters.

All the on-premises ONTAP working environments that have backup files are listed. If you don't have any backup files from an on-prem system, then the working environment is not shown.

The screenshot shows the Cloud Manager interface with the 'Data Sense' tab selected. The 'Configuration' tab is highlighted. Below it, the 'Working Environments' section displays '(2/20) Working Environments'. A filter bar at the top includes buttons for 'CVO', 'ANF', 'S3', 'DB', 'ONEDR', a red-bordered 'BACKUP' button, and 'Clear filters'. A 'Working Environment 1 (back up)' entry is listed, showing it's a 'Cloud Backup of ONTAP' in 'BETA' mode. At the bottom, there's a blue button labeled 'Activate Compliance for all Backed Up Volumes' with a red arrow pointing to it, and another button below it labeled 'or select Volumes' with a red arrow pointing to it.

3. To scan all backed up volumes in a working environment, click **Activate Compliance for all backed up Volumes**.

To scan only certain backed up volumes in a working environment, click **or select Volumes** and then choose the backup files (volumes) that you want to scan.

See [Enabling and disabling compliance scans on volumes](#) for details.

Scanning on-prem volumes versus backups of those volumes

When you view the entire list of working environments you will see two listings for each on-prem cluster if they have backed up files.

The screenshot shows the 'Configuration' section of a software interface. It lists two working environments:

- Working Environment 1** (On-Premises ONTAP):
 - Icon: Hard drive icon.
 - Name: Working Environment 1
 - Status: On-Premises ONTAP
 - Action Buttons:
 - Red box labeled 1: Activate Compliance for All Volumes
 - Blue box: or select Volumes
- Working Environment 1 (back up)** (Cloud Backup of ONTAP):
 - Icon: Cloud icon.
 - Name: Working Environment 1 (back up)
 - Status: Cloud Backup of ONTAP
 - Action Buttons:
 - Red box labeled 2: Activate Compliance for all backed up Volumes
 - Blue box: or select Volumes

The first item is the on-prem cluster and the actual volumes.

The second item is the backup files of those volumes from that same on-prem cluster.

Choose the first option to scan the volumes on the on-prem system. Choose the second option to scan the backup files from those volumes. Do not scan both on-prem volumes and backup files of the same cluster.

Scanning data protection volumes

By default, data protection (DP) volumes are not scanned because they are not exposed externally and Cloud Data Sense cannot access them. These are the destination volumes for SnapMirror operations from an on-premises ONTAP system or from a Cloud Volumes ONTAP system.

Initially, the volume list identifies these volumes as **Type DP** with the **Status Not Scanning** and the **Required Action Enable Access to DP volumes**.

'Working Environment Name' Configuration

Map & Classify All | 22/28 Volumes selected for compliance scan | Enable Access to DP Volumes | Edit CIFS Credentials

Scan	Storage Repository (Volume)	Type	Status	Required Action
Off Map Map & Classify	VolumeName1	DP	● Not Scanning	Enable access to DP Volumes
Off Map Map & Classify	VolumeName2	NFS	● Continuously Scanning	
Off Map Map & Classify	VolumeName3	CIFS	● Not Scanning	

Steps

If you want to scan these data protection volumes:

1. Click **Enable Access to DP volumes** at the top of the page.
2. Review the confirmation message and click **Enable Access to DP volumes** again.
 - Volumes that were initially created as NFS volumes in the source ONTAP system are enabled.
 - Volumes that were initially created as CIFS volumes in the source ONTAP system require that you enter CIFS credentials to scan those DP volumes. If you already entered Active Directory credentials so that Cloud Data Sense can scan CIFS volumes you can use those credentials, or you can specify a different set of Admin credentials.

Provide Active Directory Credentials

Use existing CIFS Scanning Credentials (user1@domain2) Use Custom Credentials

Active Directory Domain DNS IP Address

DP Volumes, created from a SnapMirror relationship, do not allow external access by default. Continuing will create NFS shares from DP Volumes which have been activated for **Data Sense**. The shares' export policies will allow access only from the Cloud **Data Sense** instance. [Learn More](#)

Enable Access to DP Volumes **Cancel**

Provide Active Directory Credentials

Use existing CIFS Scanning Credentials (user1@domain2) Use Custom Credentials

Username Password

Active Directory Domain DNS IP Address

DP Volumes, created from a SnapMirror relationship, do not allow external access by default. Continuing will create NFS shares from DP Volumes which have been activated for **Data Sense**. The shares' export policies will allow access only from the Cloud **Data Sense** instance. [Learn More](#)

Enable Access to DP Volumes **Cancel**

3. Activate each DP volume that you want to scan [the same way you enabled other volumes](#), or use the **Activate Compliance for all Volumes** control to enable all volumes, including all DP volumes.

Result

Once enabled, Cloud Data Sense creates an NFS share from each DP volume that was activated for scanning. The share export policies only allow access from the Data Sense instance.

Note: If you had no CIFS data protection volumes when you initially enabled access to DP volumes, and later add some, the button **Enable Access to CIFS DP** appears at the top of the Configuration page. Click this button and add CIFS credentials to enable access to these CIFS DP volumes.

 Active Directory credentials are only registered in the storage VM of the first CIFS DP volume, so all DP volumes on that SVM will be scanned. Any volumes that reside on other SVMs will not have the Active Directory credentials registered, so those DP volumes won't be scanned.

Getting started with Cloud Data Sense for Azure NetApp Files

Complete a few steps to get started with Cloud Data Sense for Azure NetApp Files.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.



1 Discover the data sources that contain the data you want to scan

Before you can scan Azure NetApp Files volumes, [Cloud Manager must be set up to discover the configuration](#).



2 Deploy the Cloud Data Sense instance

[Deploy Cloud Data Sense in Cloud Manager](#) if there isn't already an instance deployed.



3 Enable Cloud Data Sense and select the volumes to scan

Click **Compliance**, select the **Configuration** tab, and activate compliance scans for volumes in specific working environments.



4 Ensure access to volumes

Now that Cloud Data Sense is enabled, ensure that it can access all volumes.

- The Cloud Data Sense instance needs a network connection to each Azure NetApp Files subnet.
- Make sure these ports are open to the Data Sense instance:
 - For NFS – ports 111 and 2049.
 - For CIFS – ports 139 and 445.
- NFS volume export policies must allow access from the Data Sense instance.
- Data Sense needs Active Directory credentials to scan CIFS volumes.

Click **Compliance > Configuration > Edit CIFS Credentials** and provide the credentials.



5 Manage the volumes you want to scan

Select or deselect the volumes that you want to scan and Cloud Data Sense will start or stop scanning them.

Discovering the data sources that you want to scan

If the data sources you want to scan are not already in your Cloud Manager environment, you can add them to the canvas at this time.

For Azure NetApp Files, [Cloud Manager must be set up to discover the configuration](#).

Deploying the Cloud Data Sense instance

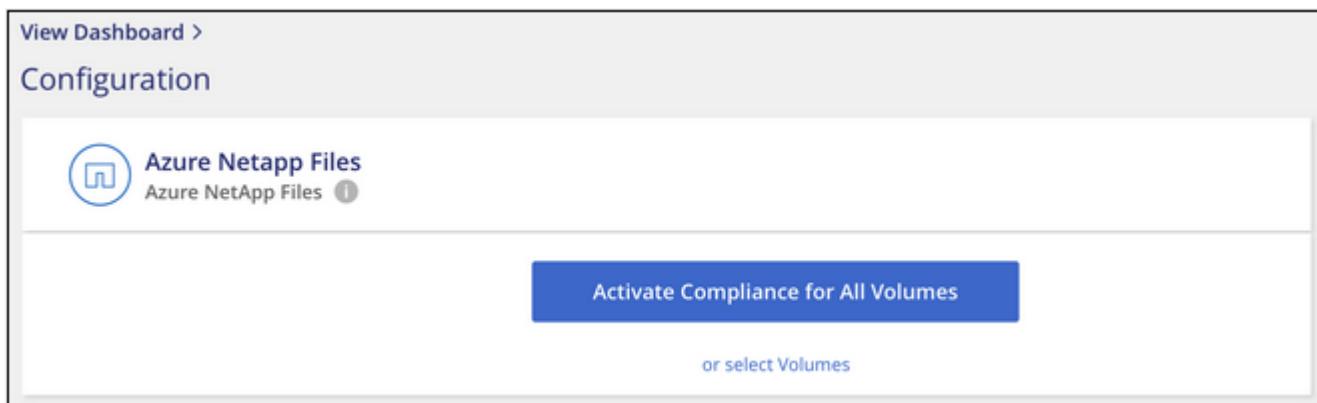
[Deploy Cloud Data Sense](#) if there isn't already an instance deployed.

Data Sense must be deployed in the cloud when scanning Azure NetApp Files volumes, and it must be deployed in the same region as the volumes you wish to scan.

Enabling Cloud Data Sense in your working environments

You can enable Cloud Data Sense on your Azure NetApp Files volumes.

1. At the top of Cloud Manager, click **Data Sense** and then select the **Configuration** tab.



2. To scan all volumes in a working environment, click **Activate Scanning for All Volumes**.

When enabled in this manner, full "mapping and classification" scanning is performed on all volumes.

If you want to enable scanning only for certain volumes, or if you only want to perform "mapping-only" scanning, click **or select Volumes** and then choose the volumes you want to scan.

See [Enabling and disabling compliance scans on volumes](#) for details.

Result

Cloud Data Sense starts scanning the volumes you selected in the working environment. Results will be available in the Compliance dashboard as soon as Cloud Data Sense finishes the initial scans. The time that it takes depends on the amount of data—it could be a few minutes or hours.

Verifying that Cloud Data Sense has access to volumes

Make sure that Cloud Data Sense can access volumes by checking your networking, security groups, and export policies. You'll need to provide Data Sense with CIFS credentials so it can access CIFS volumes.

Steps

1. Make sure that there's a network connection between the Cloud Data Sense instance and each network that includes volumes for Azure NetApp Files.



For Azure NetApp Files, Cloud Data Sense can only scan volumes that are in the same region as Cloud Manager.

2. Ensure the following ports are open to the Data Sense instance:
 - For NFS – ports 111 and 2049.
 - For CIFS – ports 139 and 445.
3. Ensure that NFS volume export policies include the IP address of the Data Sense instance so it can access the data on each volume.
4. If you use CIFS, provide Data Sense with Active Directory credentials so it can scan CIFS volumes.
 - a. At the top of Cloud Manager, click **Data Sense**.
 - b. Click the **Configuration** tab.

The screenshot shows the Data Sense configuration interface. At the top, there is a navigation bar with tabs: Data Sense, Governance, Compliance, Investigation, Policies, and Configuration. The Configuration tab is highlighted with a red box. Below the navigation bar, the page title is "6 Working Environments". Under this title, there is a card for "Azure NetApp Files" which shows "3 Volumes" and "Azure NetApp Files". To the right of this card is a blue button labeled "Configuration" with a gear icon. At the bottom right of the page is a three-dot menu icon.

- c. For each working environment, click **Edit CIFS Credentials** and enter the user name and password that Data Sense needs to access CIFS volumes on the system.

The credentials can be read-only, but providing admin credentials ensures that Data Sense can read any data that requires elevated permissions. The credentials are stored on the Cloud Data Sense instance.

After you enter the credentials, you should see a message that all CIFS volumes were authenticated successfully.

The screenshot shows the Scan Status page for "Cloud Volumes ONTAP". At the top left, there is a back arrow and the page title "Scan Status". Below the title, there is a section for "Cloud Volumes ONTAP" with a cloud icon. The page displays the following information:

- Name:** Newdatastore
- Volumes:** 12 Continuously Scanning (green dot) and 8 Not Scanning (red dot).
- CIFS Credentials Status:** Valid CIFS credentials for all accessible volumes (green checkmark).

 There are two buttons at the bottom: "View Details" and "Edit CIFS Credentials".

5. On the **Configuration** page, click **View Details** to review the status for each CIFS and NFS volume and correct any errors.

For example, the following image shows three volumes; one of which Cloud Data Sense can't scan due to network connectivity issues between the Data Sense instance and the volume.

Newdatastore Scan Configuration

Map & Classify All | 22/62 Volumes selected for Data Sense scan | Edit CIFS Credentials

Scan	Storage Repository (Volume)	Type	Status	Required Action
Off Map Map & Classify	AC_Source_clone_copy	NFS (DP)	● Continuously Scanning	
Off Map Map & Classify	AC_Source_copy	NFS (DP)	● Continuously Scanning	
Off Map Map & Classify	AC_Source_copyww	NFS	● Not Scanning	
Off Map Map & Classify	A_Source_Tier_Demo_copy	NFS	● Continuously Scanning	

Enabling and disabling compliance scans on volumes

You can stop or start mapping scans, or mapping and classification scans, in a working environment at any time from the Configuration page. We recommend that you scan all volumes.

cognigoWE Scan Configuration

Map & Classify All | 4/79 Volumes selected for Data Sense scan | Edit CIFS Credentials

Scan	Storage Repository (Volume)	Type	Status	Required Action
Off Map Map & Classify	AdiNFSVol_copy	NFS	● No Access	Access to the NFS volume was denied. Make sure tha...
Off Map Map & Classify	AdiProtest2501	NFS	● Continuously Scanning	
Off Map Map & Classify	AlexTest	NFS	● No Access	Access to the NFS volume was denied. Make sure tha...
Off Map Map & Classify	AlexTestSecond	NFS	● Not Scanning	
Off Map Map & Classify	MoreDataNeed1000	NFS	● Continuously Scanning	

To:	Do this:
Enable mapping-only scans on a volume	Click Map
Enable full scanning on a volume	Click Map & Classify
Enable full scanning on all volumes	Move the Map & Classify All slider to the right
Disable scanning on a volume	Click Off
Disable scanning on all volumes	Move the Map & Classify All slider to the left



New volumes added to the working environment are automatically scanned only when the **Activate Compliance for all Volumes** setting is enabled. When this setting is disabled, you'll need to activate scanning on each new volume you create in the working environment.

Get started with Cloud Data Sense for Amazon FSx for ONTAP

Complete a few steps to get started with Cloud Data Sense for FSx for ONTAP.

Before you begin

- You need an active Connector in AWS to deploy and manage Data Sense.

- The security group you selected when creating the working environment must allow traffic from the Cloud Data Sense instance. You can find the associated security group using the ENI connected to the FSx for ONTAP file system and edit it using the AWS Management Console.

[AWS security groups for Linux instances](#)

[AWS security groups for Windows instances](#)

[AWS elastic network interfaces \(ENI\)](#)

Quick start

Get started quickly by following these steps or scroll down for full details.



Discover the data sources that contain the data you want to scan

Before you can scan FSx for ONTAP volumes, [you must have an FSx working environment with volumes configured](#).



Deploy the Cloud Data Sense instance

[Deploy Cloud Data Sense in Cloud Manager](#) if there isn't already an instance deployed.



Enable Cloud Data Sense and select the volumes to scan

Click **Data Sense**, select the **Configuration** tab, and activate compliance scans for volumes in specific working environments.



Ensure access to volumes

Now that Cloud Data Sense is enabled, ensure that it can access all volumes.

- The Cloud Data Sense instance needs a network connection to each FSx for ONTAP subnet.
- Make sure the following ports are open to the Data Sense instance.
 - For NFS – ports 111 and 2049.
 - For CIFS – ports 139 and 445.
- NFS volume export policies must allow access from the Data Sense instance.
- Data Sense needs Active Directory credentials to scan CIFS volumes.

Click **Compliance > Configuration > Edit CIFS Credentials** and provide the credentials.



Manage the volumes you want to scan

Select or deselect the volumes you want to scan and Cloud Data Sense will start or stop scanning them.

Discovering the data sources that you want to scan

If the data sources you want to scan are not already in your Cloud Manager environment, you can add them to the canvas at this time.

For FSx for ONTAP, [Cloud Manager must be set up to discover the configuration](#).

Deploying the Cloud Data Sense instance

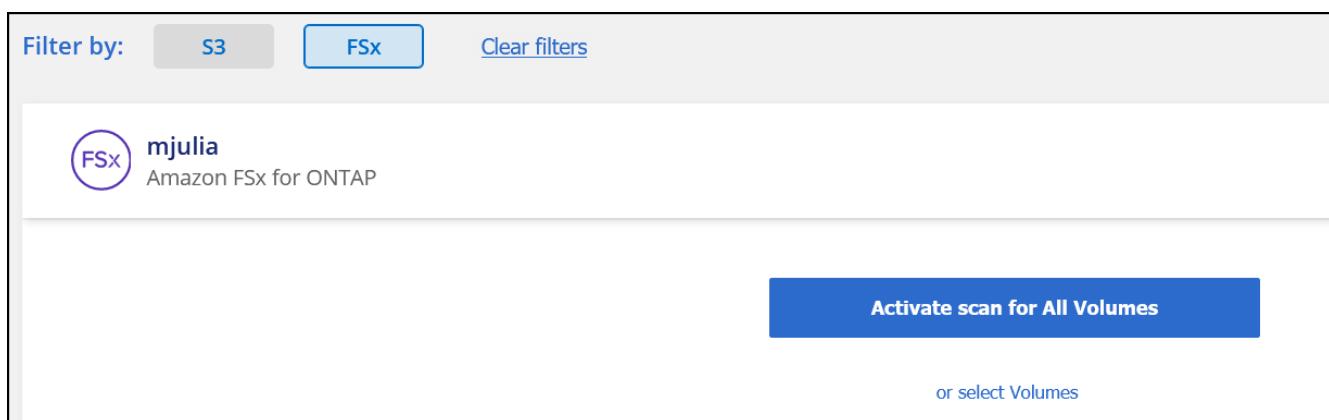
[Deploy Cloud Data Sense](#) if there isn't already an instance deployed.

Data Sense should be deployed in the same AWS network as the Connector for AWS and the FSx volumes you wish to scan.

Enabling Cloud Data Sense in your working environments

You can enable Cloud Data Sense for FSx for ONTAP volumes.

1. At the top of Cloud Manager, click **Data Sense** and then select the **Configuration** tab.



2. To scan all volumes in a working environment, click **Activate Scanning for All Volumes**.

When enabled in this manner, full "mapping and classification" scanning is performed on all volumes.

If you want to enable scanning only for certain volumes, or if you only want to perform "mapping-only" scanning, click **or select Volumes** and then choose the volumes you want to scan.

See [Enabling and disabling compliance scans on volumes](#) for details.

Result

Cloud Data Sense starts scanning the volumes you selected in the working environment. Results will be available in the Compliance dashboard as soon as Cloud Data Sense finishes the initial scans. The time that it takes depends on the amount of data—it could be a few minutes or hours.

Verifying that Cloud Data Sense has access to volumes

Make sure Cloud Data Sense can access volumes by checking your networking, security groups, and export policies.

You'll need to provide Data Sense with CIFS credentials so it can access CIFS volumes.

Steps

1. On the **Configuration** page, click **View Details** to review the status and correct any errors.

For example, the following image shows a volume Cloud Data Sense can't scan due to network connectivity issues between the Data Sense instance and the volume.

Scan	Storage Repository (Volume)	Type	Status	Required Action
Off	Map & Classify	Jrmclone	NFS	No Access Check network connectivity between the Data Sense ...

2. Make sure there's a network connection between the Cloud Data Sense instance and each network that includes volumes for FSx for ONTAP.



For FSx for ONTAP, Cloud Data Sense can scan volumes only in the same region as Cloud Manager.

3. Ensure the following ports are open to the Data Sense instance.

- For NFS – ports 111 and 2049.
- For CIFS – ports 139 and 445.

4. Ensure NFS volume export policies include the IP address of the Data Sense instance so it can access the data on each volume.
5. If you use CIFS, provide Data Sense with Active Directory credentials so it can scan CIFS volumes.

- a. At the top of Cloud Manager, click **Data Sense**.
- b. Click the **Configuration** tab.
- c. For each working environment, click **Edit CIFS Credentials** and enter the user name and password that Data Sense needs to access CIFS volumes on the system.

The credentials can be read-only, but providing admin credentials ensures that Data Sense can read any data that requires elevated permissions. The credentials are stored on the Cloud Data Sense instance.

After you enter the credentials, you should see a message that all CIFS volumes were authenticated successfully.

Enabling and disabling compliance scans on volumes

You can stop or start mapping scans, or mapping and classification scans, in a working environment at any time from the Configuration page. We recommend that you scan all volumes.

cognigoWE Scan Configuration					
Map & Classify All		4/79 Volumes selected for Data Sense scan	Edit CIFS Credentials		
Scan	Storage Repository (Volume)	Type	Status	Required Action	
Off	Map & Classify	AdINFSVol_copy	NFS	No Access	Access to the NFS volume was denied. Make sure tha...
Off	Map & Classify	AdIProtect2501	NFS	Continuously Scanning	
Off	Map	Map & Classify	AlexTest	No Access	Access to the NFS volume was denied. Make sure tha...
Off	Map	Map & Classify	AlexTestSecond	Not Scanning	
Off	Map	Map & Classify	MoreDataNeed1000	Continuously Scanning	

To:	Do this:
Enable mapping-only scans on a volume	Click Map
Enable full scanning on a volume	Click Map & Classify
Enable full scanning on all volumes	Move the Map & Classify All slider to the right
Disable scanning on a volume	Click Off
Disable scanning on all volumes	Move the Map & Classify All slider to the left



New volumes added to the working environment are automatically scanned only when the **Activate Compliance for all Volumes** setting is enabled. When this setting is disabled, you'll need to activate scanning on each new volume you create in the working environment.

Getting started with Cloud Data Sense for Amazon S3

Cloud Data Sense can scan your Amazon S3 buckets to identify the personal and sensitive data that resides in S3 object storage. Cloud Data Sense can scan any bucket in the account, regardless if it was created for a NetApp solution.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.



Set up the S3 requirements in your cloud environment

Ensure that your cloud environment can meet the requirements for Cloud Data Sense, including preparing an IAM role and setting up connectivity from Data Sense to S3. [See the complete list.](#)



Deploy the Cloud Data Sense instance

Deploy [Cloud Data Sense](#) if there isn't already an instance deployed.



Activate Data Sense on your S3 working environment

Select the Amazon S3 working environment, click **Enable**, and select an IAM role that includes the required permissions.



Select the buckets to scan

Select the buckets that you'd like to scan and Cloud Data Sense will start scanning them.

Reviewing S3 prerequisites

The following requirements are specific to scanning S3 buckets.

Set up an IAM role for the Cloud Data Sense instance

Cloud Data Sense needs permissions to connect to the S3 buckets in your account and to scan them. Set up an IAM role that includes the permissions listed below. Cloud Manager prompts you to select an IAM role when you enable Data Sense on the Amazon S3 working environment.

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": [  
                "s3:Get*",  
                "s3>List*",  
                "s3:PutObject",  
            ],  
            "Resource": "*"  
        },  
        {  
            "Effect": "Allow",  
            "Action": [  
                "iam:GetPolicyVersion",  
                "iam:GetPolicy",  
                "iam>ListAttachedRolePolicies"  
            ],  
            "Resource": [  
                "arn:aws:iam::*:policy/*",  
                "arn:aws:iam::*:role/*"  
            ]  
        }  
    ]  
}
```

Provide connectivity from Cloud Data Sense to Amazon S3

Cloud Data Sense needs a connection to Amazon S3. The best way to provide that connection is through a VPC Endpoint to the S3 service. For instructions, see [AWS Documentation: Creating a Gateway Endpoint](#).

When you create the VPC Endpoint, be sure to select the region, VPC, and route table that corresponds to the Cloud Data Sense instance. You must also modify the security group to add an outbound HTTPS rule that enables traffic to the S3 endpoint. Otherwise, Data Sense can't connect to the S3 service.

If you experience any issues, see [AWS Support Knowledge Center: Why can't I connect to an S3 bucket using a gateway VPC endpoint?](#)

An alternative is to provide the connection by using a NAT Gateway.



You can't use a proxy to get to S3 over the internet.

Deploying the Cloud Data Sense instance

Deploy Cloud Data Sense in Cloud Manager if there isn't already an instance deployed.

You need to deploy the instance in an AWS Connector so that Cloud Manager automatically discovers the S3 buckets in this AWS account and displays them in an Amazon S3 working environment.

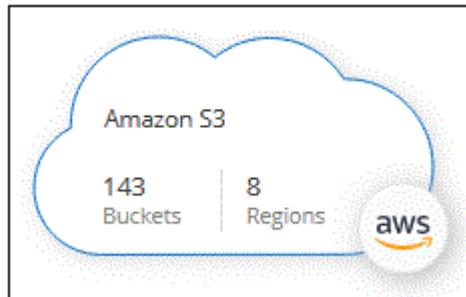
Note: Deploying Cloud Data Sense in an on-premises location is not currently supported when scanning S3 buckets.

Activating Data Sense on your S3 working environment

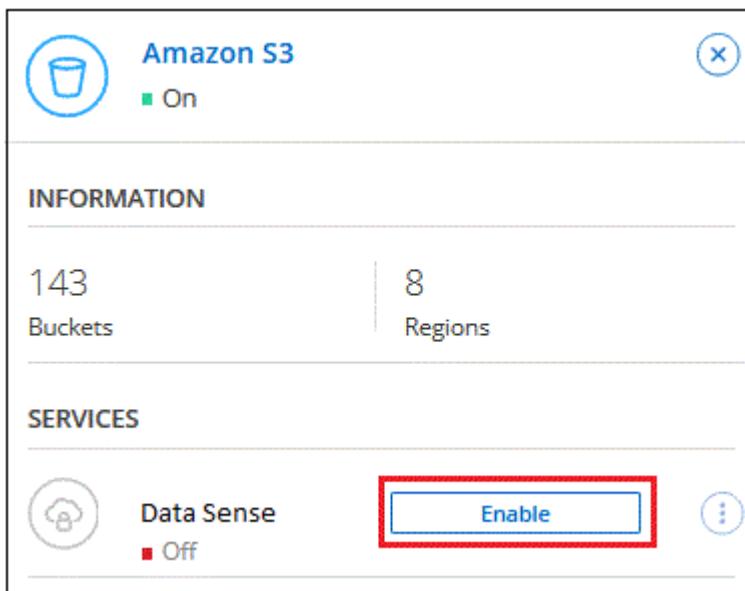
Enable Cloud Data Sense on Amazon S3 after you verify the prerequisites.

Steps

1. At the top of Cloud Manager, click **Canvas**.
2. Select the Amazon S3 working environment.



3. In the Data Sense pane on the right, click **Enable**.



4. When prompted, assign an IAM role to the Cloud Data Sense instance that has [the required permissions](#).

Assign an AWS IAM Role for Cloud Data Sense

To enable Cloud Data Sense on Amazon S3 buckets, select an existing IAM Role.

Make sure that your AWS IAM Role has the permission defined in the [Policy Requirements](#).

Select IAM Role

occm

VPC Endpoint for Amazon S3 Required

A VPC endpoint to the Amazon S3 service is required so [Data Sense](#) can securely scan the data.

Alternatively, ensure that the [Data Sense](#) instance has direct access to the internet via a NAT Gateway or Internet Gateway.

Free for the 1st TB

Over 1 TB you pay only for what you use. [Learn more about pricing.](#)

Enable

Cancel

5. Click **Enable**.

 You can also enable compliance scans for a working environment from the Configuration page by clicking the  button and selecting **Activate Data Sense**.

Result

Cloud Manager assigns the IAM role to the instance.

Enabling and disabling compliance scans on S3 buckets

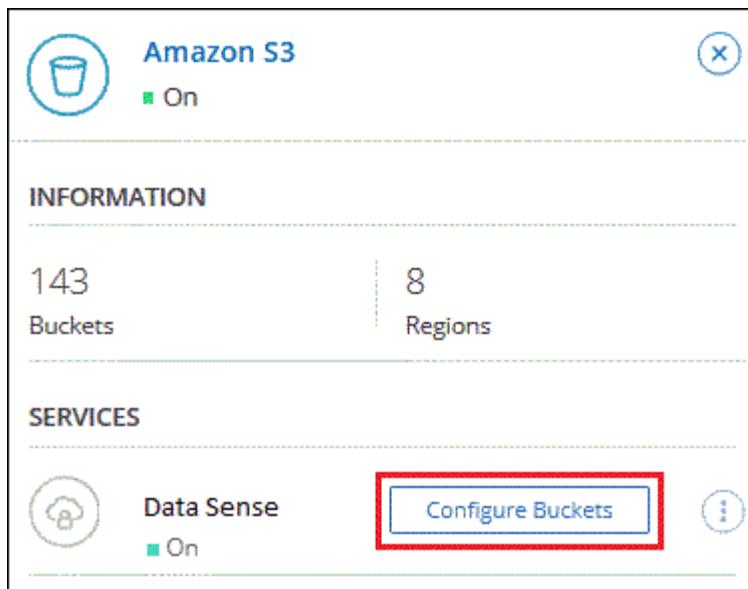
After Cloud Manager enables Cloud Data Sense on Amazon S3, the next step is to configure the buckets that you want to scan.

When Cloud Manager is running in the AWS account that has the S3 buckets you want to scan, it discovers those buckets and displays them in an Amazon S3 working environment.

Cloud Data Sense can also [scan S3 buckets that are in different AWS accounts](#).

Steps

1. Select the Amazon S3 working environment.
2. In the pane on the right, click **Configure Buckets**.



3. Enable mapping-only scans, or mapping and classification scans, on your buckets.

Amazon S3 Configuration				
15/28 Buckets in Scan Scope.				
Scan	Bucket Name	Status	Required Action	
<input type="button" value="Off"/> <input type="button" value="Map"/> <input type="button" value="Map & Classify"/>	BucketName1	● Not Scanning	Add Credentials	
<input type="button" value="Off"/> <input type="button" value="Map"/> <input type="button" value="Map & Classify"/>	BucketName2	● Continuously Scanning		
<input type="button" value="Off"/> <input type="button" value="Map"/> <input type="button" value="Map & Classify"/>	BucketName3	● Not Scanning		

To:	Do this:
Enable mapping-only scans on a bucket	Click Map
Enable full scans on a bucket	Click Map & Classify
Disable scanning on a bucket	Click Off

Result

Cloud Data Sense starts scanning the S3 buckets that you enabled. If there are any errors, they'll appear in the Status column, alongside the required action to fix the error.

Scanning buckets from additional AWS accounts

You can scan S3 buckets that are under a different AWS account by assigning a role from that account to access the existing Cloud Data Sense instance.

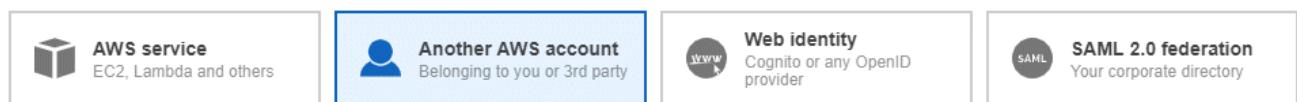
Steps

1. Go to the target AWS account where you want to scan S3 buckets and create an IAM role by selecting **Another AWS account**.

Create role

1 2 3 4

Select type of trusted entity



Allows entities in other accounts to perform actions in this account. [Learn more](#)

Specify accounts that can use this role

Account ID* ⓘ

Options Require external ID (Best practice when a third party will assume this role)
 Require MFA ⓘ

Be sure to do the following:

- Enter the ID of the account where the Cloud Data Sense instance resides.
- Change the **Maximum CLI/API session duration** from 1 hour to 12 hours and save that change.
- Attach the Cloud Data Sense IAM policy. Make sure it has the required permissions.

```
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Action": [  
                "s3:Get*",  
                "s3>List*",  
                "s3:PutObject",  
            ],  
            "Resource": "*"  
        },  
    ]  
}
```

2. Go to the source AWS account where the Data Sense instance resides and select the IAM role that is attached to the instance.
 - a. Change the **Maximum CLI/API session duration** from 1 hour to 12 hours and save that change.
 - b. Click **Attach policies** and then click **Create policy**.
 - c. Create a policy that includes the "sts:AssumeRole" action and specify the ARN of the role that you created in the target account.

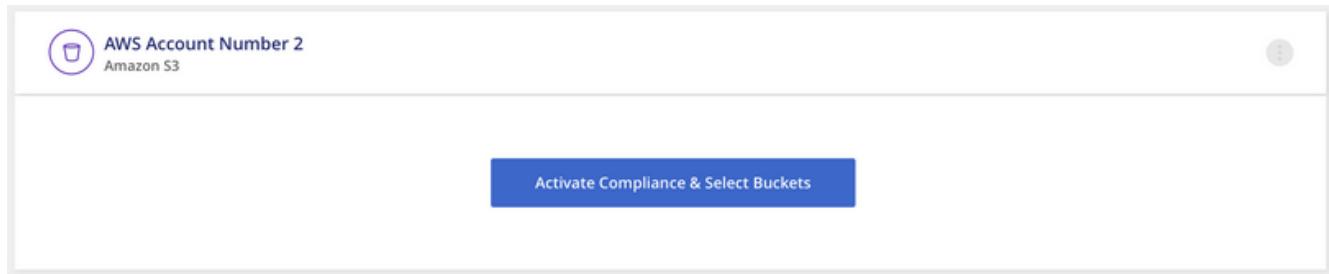
```

{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": "sts:AssumeRole",
            "Resource": "arn:aws:iam::<ADDITIONAL-ACCOUNT-ID>:role/<ADDITIONAL_ROLE_NAME>"
        },
        {
            "Effect": "Allow",
            "Action": [
                "iam:GetPolicyVersion",
                "iam:GetPolicy",
                "iam>ListAttachedRolePolicies"
            ],
            "Resource": [
                "arn:aws:iam::*:policy/*",
                "arn:aws:iam::*:role/*"
            ]
        }
    ]
}

```

The Cloud Data Sense instance profile account now has access to the additional AWS account.

3. Go to the **Amazon S3 Configuration** page and the new AWS account is displayed. Note that it can take a few minutes for Cloud Data Sense to sync the new account's working environment and show this information.



4. Click **Activate Data Sense & Select Buckets** and select the buckets you want to scan.

Result

Cloud Data Sense starts scanning the new S3 buckets that you enabled.

Scanning database schemas

Complete a few steps to start scanning your database schemas with Cloud Data Sense.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.



Review database prerequisites

Ensure that your database is supported and that you have the information necessary to connect to the database.



Deploy the Cloud Data Sense instance

[Deploy Cloud Data Sense](#) if there isn't already an instance deployed.



Add the database server

Add the database server that you want to access.



Select the schemas

Select the schemas that you want to scan.

Reviewing prerequisites

Review the following prerequisites to make sure that you have a supported configuration before you enable Cloud Data Sense.

Supported databases

Cloud Data Sense can scan schemas from the following databases:

- Amazon Relational Database Service (Amazon RDS)
- MongoDB
- MySQL
- Oracle
- PostgreSQL
- SAP HANA
- SQL Server (MSSQL)



The statistics gathering feature **must be enabled** in the database.

Database requirements

Any database with connectivity to the Cloud Data Sense instance can be scanned, regardless of where it is hosted. You just need the following information to connect to the database:

- IP Address or host name
- Port
- Service name (only for accessing Oracle databases)
- Credentials that allow read access to the schemas

When choosing a user name and password, it's important to choose one that has full read permissions to all the schemas and tables you want to scan. We recommend that you create a dedicated user for the Cloud Data Sense system with all the required permissions.

Note: For MongoDB, a read-only Admin role is required.

Adding the database server

You must have [deployed an instance of Cloud Data Sense in Cloud Manager already](#).

Add the database server where the schemas reside.

1. From the Working Environments Configuration page, click **Add Data Source > Add Database Server**.

The screenshot shows the 'Working Environments' configuration page with a title '(2/20) Working Environments'. At the top right is a blue button labeled 'Add Data Source'. Below it is a dropdown menu with three options: 'Add Database Server' (highlighted with a red box), 'Add OneDrive Account', and 'Add AWS S3 accounts'. The main area displays a table with three rows. The first row shows 'Working Environment Name 1 | 127 Volumes' under 'Cloud Volumes ONTAP'. The second row shows '87' volumes with the status 'Continuously Scanning' and a 'View Details' link. The third row shows '28' volumes with the status 'Not Scanning' and a 'View Details' link. Below the table, the text 'Continuously scanning all selected Volumes' is visible.

2. Enter the required information to identify the database server.
 - a. Select the database type.
 - b. Enter the port and the host name or IP address to connect to the database.
 - c. For Oracle databases, enter the Service name.
 - d. Enter the credentials so that Cloud Data Sense can access the server.
 - e. Click **Add DB Server**.

Add DB Server

To activate Compliance on Databases, first add a Database Server. After this step, you'll be able to select which Database Schemas you would like to activate Compliance for.

Database

Database Type

Host Name or IP Address

Port

Service Name

Credentials

Username

Password

Add DB Server

Cancel

The database is added to the list of working environments.

Enabling and disabling compliance scans on database schemas

You can stop or start full scanning of your schemas at any time.



There is no option to select mapping-only scans for database schemas.

1. From the **Configuration** page, click the **Configuration** button for the database you want to configure.

2. Select the schemas that you want to scan by moving the slider to the right.

'Working Environment Name' Configuration			
28/28 Schemas selected for compliance scan			Edit Credentials
Scan	Schema Name	Status	Required Action
<input checked="" type="checkbox"/>	DB1 - SchemaName1	● Not Scanning	Add Credentials
<input checked="" type="checkbox"/>	DB1 - SchemaName2	● Continuously Scanning	
<input checked="" type="checkbox"/>	DB1 - SchemaName3	● Continuously Scanning	
<input checked="" type="checkbox"/>	DB1 - SchemaName4	● Continuously Scanning	

Result

Cloud Data Sense starts scanning the database schemas that you enabled. If there are any errors, they'll appear in the Status column, alongside the required action to fix the error.

Scanning OneDrive accounts

Complete a few steps to start scanning files in your user's OneDrive folders with Cloud Data Sense.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.

1

Review OneDrive prerequisites

Ensure that you have the Admin credentials to log into the OneDrive account.

2

Deploy the Cloud Data Sense instance

[Deploy Cloud Data Sense](#) if there isn't already an instance deployed.

3

Add the OneDrive account

Using Admin user credentials, log into the OneDrive account that you want to access so that it is added as a new working environment.

4

Add the users and select the users to scan

Add the list of users from the OneDrive account that you want to scan and select the type of scanning. You can add up to 100 users at time.

Reviewing OneDrive requirements

Review the following prerequisites to make sure that you have a supported configuration before you enable Cloud Data Sense.

- You must have the Admin login credentials for the OneDrive for Business account that provides read access to all user files.
- You will need a line-separated list of the email addresses for all the users whose OneDrive folders you want to scan.

Adding the OneDrive account

You must have [deployed an instance of Cloud Data Sense in Cloud Manager already](#).

Add the OneDrive account where the user files reside.

Steps

1. From the Working Environments Configuration page, click **Add Data Source > Add OneDrive Account**.

The screenshot shows the 'Working Environments' configuration page. At the top right, there is a 'Add Data Source' button with a dropdown arrow. The dropdown menu lists several options: 'Add File Shares Group', 'Add Database Server', 'Add OneDrive Account' (which is highlighted with a red box), and 'Add AWS S3 accounts'. The main area displays a list of working environments, with one entry for 'Working Environment Name 1' showing 127 volumes. Below this, there are three status cards: 'Continuously Scanning' (87 volumes), 'Not Scanning' (28 volumes), and 'Continuously scanning a selected Volumes'.

2. In the Add a OneDrive account dialog, click **Sign in to OneDrive**.
3. In the Microsoft page that appears, select the OneDrive account and enter the required Admin user and password, then click **Accept** to allow Cloud Data Sense to read data from this account.

The OneDrive account is added to the list of working environments.

Adding OneDrive users to compliance scans

You can add individual OneDrive users, or all of your OneDrive users, so that their files will be scanned by Cloud Data Sense.

Steps

1. From the **Configuration** page, click the **Configuration** button for the OneDrive account.

The screenshot shows the 'Configuration' page for a OneDrive account. At the top right, there is a 'Add Data Source' button with a dropdown arrow. The main area displays a list of users for 'OneDrive Account 1', which has 41 users. To the right of the user list, there is a 'Configuration' button, which is highlighted with a red box. There is also a three-dot ellipsis button.

2. If this is the first time adding users for this OneDrive account, click **Add your first OneDrive users**.

'Working Environment Name' Configuration



No OneDrive users are being scanned

+ Add your first OneDrive users

If you are adding additional users from a OneDrive account, click **Add OneDrive users**.

Working Environment 4 Configuration

24 users are being scanned for compliance

+ Add OneDrive users

Scan	Username	Status	Required Action
Off Map Map & Classify	user2@example.com	Continuously Scanning	...
Off Map Map & Classify	user3@example.com	Continuously Scanning	...

3. Add the email addresses for the users whose files you want to scan - one email address per line (up to 100 maximum per session) - and click **Add Users**.

Add OneDrive users

Provide a list of OneDrive users for Cloud Data Sense to scan their data, line-separated. You can add up to 100 users at a time.

Type or paste below the OneDrive user accounts to add

User Accounts

```
user@example.com
user@example.com
user@example.com
user@example.com
user@example.com
user@example.com
user@example.com
```

Add Users

Cancel

A confirmation dialog displays the number of users who were added.

If the dialog lists any users who could not be added, capture this information so that you can resolve the issue. In some cases you can re-add the user with a corrected email address.

4. Enable mapping-only scans, or mapping and classification scans, on user files.

To:	Do this:
Enable mapping-only scans on user files	Click Map
Enable full scans on user files	Click Map & Classify
Disable scanning on user files	Click Off

Result

Cloud Data Sense starts scanning the files for the users you added, and the results are displayed in the Dashboard and in other locations.

Removing a OneDrive user from compliance scans

If users leave the company or if their email address changes, you can remove individual OneDrive users from having their files scanned at any time. Just click **Remove OneDrive User** from the Configuration page.

The screenshot shows a configuration interface for managing OneDrive users. At the top, it says 'Working Environment 4 Configuration' and indicates '24 users are being scanned for compliance'. Below this is a table with columns: Scan, Username, Status, and Required Action. The 'Scan' column has buttons for Off, Map, and Map & Classify. The 'Username' column lists 'user2@example.com'. The 'Status' column shows a green dot and the text 'Continuously Scanning'. The 'Required Action' column contains a red-bordered button labeled 'Remove OneDrive User' with a delete icon. A red box highlights this button.

Scanning file shares

Complete a few steps to start scanning non-NetApp NFS or CIFS file shares directly with Cloud Data Sense. These file shares can reside on-premises or in the cloud.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.



Review file share prerequisites

For CIFS (SMB) shares, ensure that you have credentials to access the shares.



Deploy the Cloud Data Sense instance

Deploy Cloud Data Sense if there isn't already an instance deployed.

**3**

Create a group to hold the file shares

The group is a container for the file shares that you want to scan, and it is used as the working environment name for those file shares.

**4**

Add the file shares and select the shares to scan

Add the list of file shares that you want to scan and select the type of scanning. You can add up to 100 file shares at a time.

Reviewing file share requirements

Review the following prerequisites to make sure that you have a supported configuration before you enable Cloud Data Sense.

- The shares can be hosted anywhere, including in the cloud or on-premises. These are file shares that reside on non-NetApp storage systems.
- There needs to be network connectivity between the Data Sense instance and the shares.
- Make sure these ports are open to the Data Sense instance:
 - For NFS – ports 111 and 2049.
 - For CIFS – ports 139 and 445.
- You will need the list of shares you want to add in the format <host_name>:<share_path>. You can enter the shares individually, or you can supply a line-separated list of the file shares you want to scan.
- For CIFS (SMB) shares, ensure that you have Active Directory credentials that provide read access to the shares. Admin credentials are preferred in case Cloud Data Sense needs to scan any data that requires elevated permissions.

Creating the group for the file shares

You must have [deployed an instance of Cloud Data Sense in Cloud Manager already](#).

You must add a files shares "group" before you can add your file shares. The group is a container for the file shares that you want to scan, and the group name is used as the working environment name for those file shares.

You can mix NFS and CIFS shares in the same group, however, all CIFS file shares in a group need to be using the same Active Directory credentials. If you plan to add CIFS shares that use different credentials, you must make a separate group for each unique set of credentials.

Steps

1. From the Working Environments Configuration page, click **Add Data Source > Add File Shares Group**.

(2/20) Working Environments

Filter by: CVO ANF S3 DB ONEDR BACKUP Clear filters

Working Environment Name 1 | 127 Volumes
Cloud Volumes ONTAP

87 <i>Continuously Scanning</i> View Details	28 <i>Not Scanning</i> View Details	Continuously scanning a selected Volumes
--	---	--

Add File Shares Group
Add Database Server
Add OneDrive Account
Add AWS S3 accounts

2. In the Add Files Shares Group dialog, enter the name for the group of shares and click **Continue**.

The new File Shares Group is added to the list of working environments.

Adding file shares to a group

You add file shares to the File Shares Group so that the files in those shares will be scanned by Cloud Data Sense. You add the shares in the format <host_name>:<share_path>.

You can add individual file shares, or you can supply a line-separated list of the file shares you want to scan. You can add up to 100 shares at a time.

When adding both NFS and CIFS shares in a single group, you'll need to run through the process twice - once adding NFS shares, and then again adding the CIFS shares.

Steps

1. From the *Working Environments* page, click the **Configuration** button for the File Shares Group.

(1/20) Working Environments

Filter by: CVO ANF S3 DB ONEDR SHARES Clear filters

Shares Group 1 | 41 Shares
File Shares Group

Configuration

2. If this is the first time adding file shares for this File Shares Group, click **Add your first Shares**.

'Working Environment Name' Configuration

No Shares are being scanned

+ Add your first Shares

If you are adding file shares to an existing group, click **Add Shares**.

Working Environment 2 Configuration

2/22 Shares selected for compliance scan

Scan	Share name	Protocol	Status	Required Action
Off	Map	Map & Classify	Sharepath2	CIFS Continuously Scanning
Off	Map	Map & Classify	Sharepath3	NFS Continuously Scanning

- Select the protocol for the file shares you are adding, add the file shares that you want to scan - one file share per line - and click **Continue**.

When adding CIFS (SMB) shares, you need to enter the Active Directory credentials that provide read access to the shares. Admin credentials are preferred.

Adding Shares

Directly add any NFS or CIFS (SMB) File Shares, located in the cloud or on-premises.

Select Protocol

You'll be able to add additional shares from the other protocol later.

NFS CIFS (SMB)

Type or paste below the Shares to add

Provide a list of shares, line-separated. You can add up to 100 at a time (able to add more later).

Hostname:/SHAREPATH
 Hostname:/SHAREPATH
 Hostname:/SHAREPATH

NFS CIFS (SMB)

Provide CIFS Credentials ⓘ
 Username ⓘ Password

Continue **Cancel**

A confirmation dialog displays the number of shares that were added.

If the dialog lists any shares that could not be added, capture this information so that you can resolve the issue. In some cases you can re-add the share with a corrected host name or share name.

- Enable mapping-only scans, or mapping and classification scans, on each file share.

To:	Do this:
Enable mapping-only scans on file shares	Click Map
Enable full scans on file shares	Click Map & Classify
Disable scanning on file shares	Click Off

Result

Cloud Data Sense starts scanning the files in the file shares you added, and the results are displayed in the Dashboard and in other locations.

Removing a file share from compliance scans

If you no longer need to scan certain file shares, you can remove individual file shares from having their files scanned at any time. Just click **Remove Share** from the Configuration page.

The screenshot shows the 'Working Environment 2 Configuration' page. At the top, there are buttons for '+ Add Shares' and 'Edit CIFS Credentials'. Below that, it says '2/22 Shares selected for compliance scan'. A table lists a single share: 'Share name' is 'Sharepath1', 'Protocol' is 'NFS', 'Status' is 'Not Scanning' (indicated by a red dot), and 'Required Action' is 'Add new credentials'. In the bottom right corner of the share row, there is a red-bordered button labeled 'Remove Share' with a delete icon.

Scanning object storage that uses S3 protocol

Complete a few steps to start scanning data within object storage directly with Cloud Data Sense. Data Sense can scan data from any Object Storage service which uses the Simple Storage Service (S3) protocol. This includes NetApp StorageGRID, IBM Cloud Object Store, Azure Blob (using MinIO), Linode, B2 Cloud Storage, and more.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.

1

Review object storage prerequisites

You need to have the endpoint URL to connect with the object storage service.

You need to have the Access Key and Secret Key from the object storage provider so that Cloud Data Sense can access the buckets.

2

Deploy the Cloud Data Sense instance

[Deploy Cloud Data Sense](#) if there isn't already an instance deployed.

3

Add the Object Storage Service

Add the object storage service to Cloud Data Sense.

4

Select the buckets to scan

Select the buckets that you'd like to scan and Cloud Data Sense will start scanning them.

Reviewing object storage requirements

Review the following prerequisites to make sure that you have a supported configuration before you enable Cloud Data Sense.

- You need to have the endpoint URL to connect with the object storage service.
- You need to have the Access Key and Secret Key from the object storage provider so that Data Sense can access the buckets.
- Support for Azure Blob requires that you use the [MinIO service](#).

Adding the object storage service to Cloud Data Sense

You must have [deployed an instance of Cloud Data Sense in Cloud Manager already](#).

Add the object storage service.

Steps

1. From the Working Environments Configuration page, click **Add Data Source > Add Object Storage Service**.

The screenshot shows the 'Working Environments' configuration page. At the top, there's a header '(2/20) Working Environments' with buttons for 'Integrate AIP Labels' and 'Add Data Source'. Below the header, there's a filter bar with buttons for 'CVO' (selected), 'ANF', 'S3', 'DB', 'ONEDR', 'BACKUP', and 'Clear filters'. The main area displays a list of working environments. One environment, 'Working Environment Name 1 | 127 Volumes' (Cloud Volumes ONTAP), is shown with 87 volumes, labeled 'Continuously Scanning' with a green checkmark icon. Another environment has 28 volumes, labeled 'Not Scanning' with a red X icon. To the right, a sidebar lists options: 'Add File Shares Group', 'Add Database Server', 'Add OneDrive Account', 'Add AWS S3 accounts' (with a link), and 'Add Object Storage Service', which is highlighted with a red box. There's also a 'View Details' link for each volume count.

2. In the Add Object Storage Service dialog, enter the details for the object storage service and click **Continue**.
 - a. Enter the name you want to use for the Working Environment. This name should reflect the name of the object storage service to which you are connecting.
 - b. Enter the Endpoint URL to access the object storage service.
 - c. Enter the Access Key and Secret Key so that Cloud Data Sense can access the buckets in the object storage.

Add Object Storage Service

Cloud Data Sense can scan data from any Object Storage service which uses the S3 protocol. This includes NetApp StorageGRID, IBM Object Store, and more.

To continue, enter the following information. In the next steps you'll need to select the buckets you want to scan.

Name the Working Environment

Endpoint URL

Access Key

Secret Key

[Continue](#)

[Cancel](#)

Result

The new Object Storage Service is added to the list of working environments.

Enabling and disabling compliance scans on object storage buckets

After you enable Cloud Data Sense on your Object Storage Service, the next step is to configure the buckets that you want to scan. Data Sense discovers those buckets and displays them in the working environment you created.

Steps

1. In the Configuration page, click **Configuration** from the Object Storage Service working environment.

(1/20) Working Environments

Filter by: CVO ANF S3 DB ONEDR BCKUP OB.STG Clear filters

Rstor Integrated | 41 Buckets
Object Storage Service

[Configuration](#)

Bucket Status	Count	Description
Continuous Scanning	23	Continuously Scanning View Details
All Buckets selected for Compliance	0	All Buckets selected for Compliance
Continuously scanning all selected Buckets	0	Continuously scanning all selected Buckets

2. Enable mapping-only scans, or mapping and classification scans, on your buckets.

Rstor Integrated Configuration

3/55 Buckets selected for Compliance scan



Scan	Storage Repository (Bucket) ↓↑	Status ↓↑	Required Action ↓↑
Off	Map	Map & Classify	logs-759995470648-us-east-1 ● Not Scanning
Off	Map	Map & Classify	logs-759995470648-us-west-2 ● Not Scanning
Off	Map	Map & Classify	carstock ● Continuously Scanning

To:	Do this:
Enable mapping-only scans on a bucket	Click Map
Enable full scans on a bucket	Click Map & Classify
Disable scanning on a bucket	Click Off

Result

Cloud Data Sense starts scanning the buckets that you enabled. If there are any errors, they'll appear in the Status column, alongside the required action to fix the error.

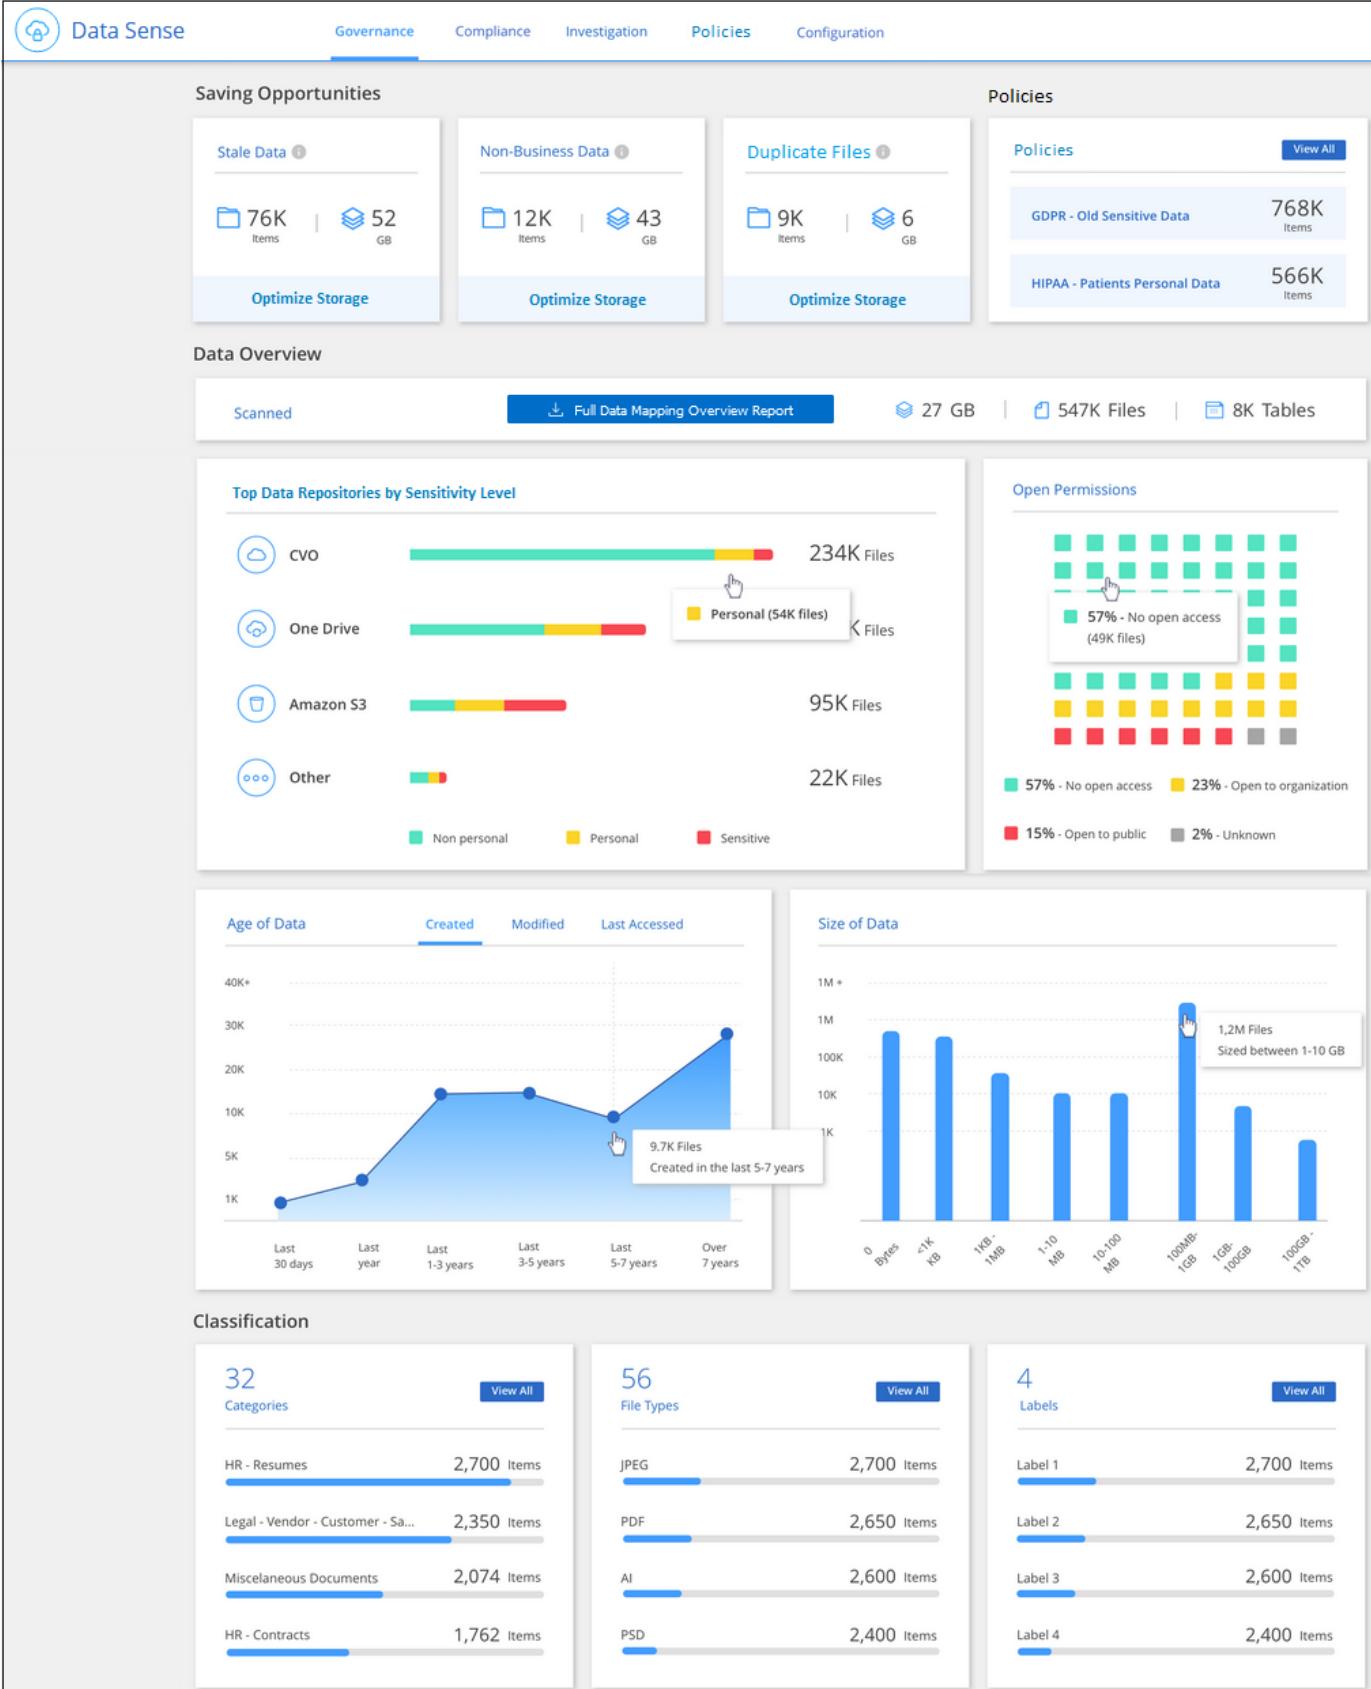
Viewing governance details about the data stored in your organization

Gain control of the costs related to the data on your organizations' storage resources. Cloud Data Sense identifies the amount of stale data, non-business data, duplicate files, and very large files in your systems so you can decide whether you want to remove or tier some files to less expensive object storage.

Additionally, if you are planning to migrate data from on-premises locations to the cloud, you can view the size of the data and whether any of the data contains sensitive information prior to moving it.

The Governance dashboard

The Governance dashboard provides information so that you can increase the efficiency and control the costs related to the data stored on your storage resources.



Saving Opportunities

You may want to investigate the items in the *Saving Opportunities* area to see if there is any data you should delete or tier to less expensive object storage. Click each item to view the filtered results in the Investigation

page.

- **Stale Data** - Data that was last modified over 3 years ago.
- **Non-Business Data** - Data considered not to be business related, based on their Category or File Type. This includes:
 - Application Data
 - Audio
 - Executables
 - Images
 - Logs
 - Videos
 - Miscellaneous (general "other" category)
- **Duplicate Files** - Files that are duplicated in other locations in the data sources you are scanning. [See what types of duplicate files are displayed.](#)

Data Overview

A quick overview of all the data that is being scanned. Click the button to download a full data mapping report that includes Usage Capacity, Age of Data, Size of Data, and File Types for all working environments and data sources. See [Data Mapping Report](#) for complete details.

Policies with the largest number of results

Click the name of a Policy in the *Policy* area to display the results in the Investigation page. Click **View All** to view the list of all available Policies.

Click [here](#) to learn more about Policies.

Top data repositories listed by data sensitivity

The *Top Data Repositories by Sensitivity Level* area lists up to the top four data repositories (working environments and data sources) that contain the most sensitive items. The bar chart for each working environment is divided into:

- Non-Personal data
- Personal data
- Sensitive Personal data

You can hover over each section to see the total number of items in each category.

Data listed by types of Open Permissions

The *Open Permissions* area shows the percentage for each type of permissions that exist for all files that are being scanned. The chart shows the following types of permissions:

- No Open Access
- Open to Organization
- Open to Public

- Unknown Access

You can hover over each section to see the total number of files in each category. Click each area to view the filtered results in the Investigation page so that you can investigate further.



Files in OneDrive accounts and in databases are not represented in this chart.

Age of Data and Size of Data graphs

You may want to investigate the items in the *Age* and *Size* graphs to see if there is any data you should delete or tier to less expensive object storage.

You can hover over a point in the charts to see details about the age or size of the data in that category. Click to view all the files filtered by that age or size range.

- **Age of Data graph** - Categorizes data based on the time it was created, the last time it was accessed, or the last time it was modified.
- **Size of Data graph** - Categorizes data based on size.

Most identified data Classifications

The *Classification* area provides a list of the most identified [Categories](#), [File types](#), and [AIP Labels](#) in your scanned data.

Categories

Categories can help you understand what's happening with your data by showing you the types of information that you have. For example, a category like "resumes" or "employee contracts" can include sensitive data. When you investigate the results, you might find that employee contracts are stored in an insecure location. You can then correct that issue.

See [Viewing files by categories](#) for more information.

File types

Reviewing your file types can help you control your sensitive data because you might find that certain file types are not stored correctly.

See [Viewing file types](#) for more information.

AIP labels

If you have subscribed to Azure Information Protection (AIP), you can classify and protect documents and files by applying labels to content. Reviewing the most used AIP labels that are assigned to files enables you to see which labels are most used in your files.

See [AIP Labels](#) for more information.

Viewing compliance details about the data stored in your organization

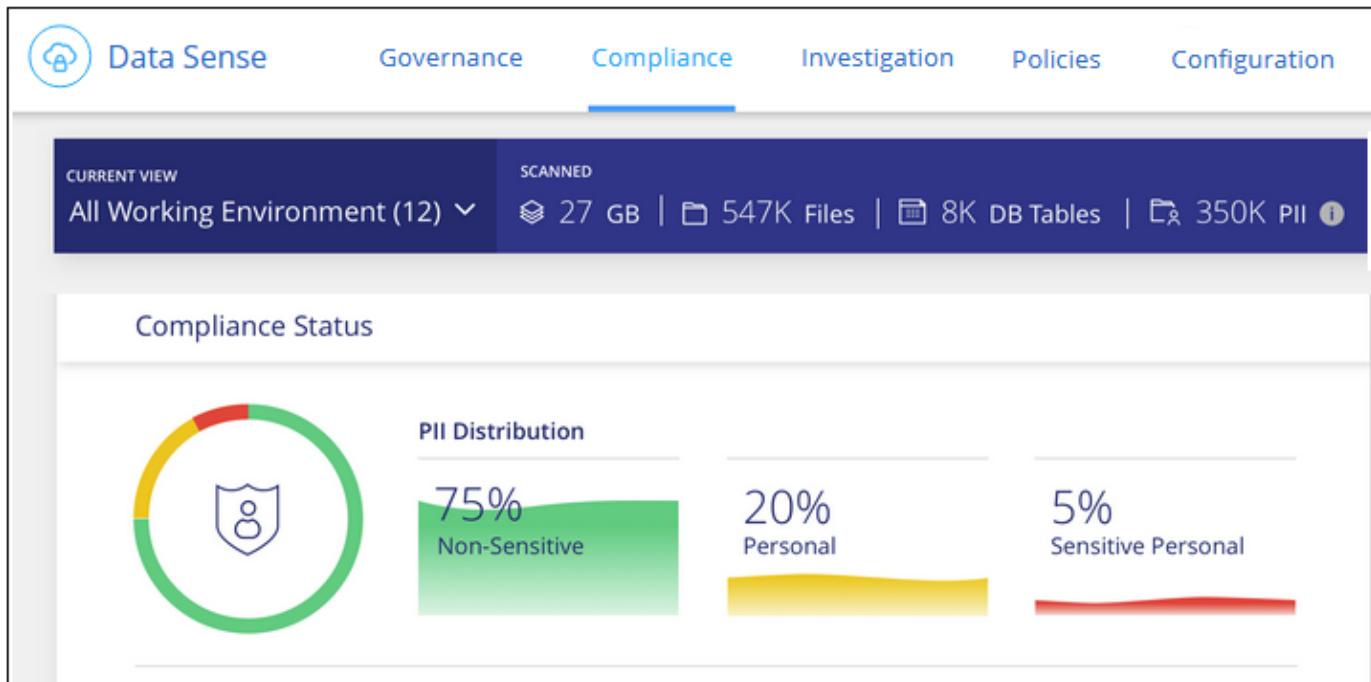
Gain control of your private data by viewing details about the personal data and sensitive

personal data in your organization. You can also gain visibility by reviewing the categories and file types that Cloud Data Sense found in your data.



The capabilities described in this section are available only if you have chosen to perform a full classification scan on your data sources. Data sources that have had a mapping-only scan do not show file-level details.

By default, the Cloud Data Sense dashboard displays compliance data for all working environments and databases.



If you want to see data for only some of the working environments, [select those working environments](#).

You can also filter the results from the Data Investigation page and download a report of the results as a CSV file. See [Filtering data in the Data Investigation page](#) for details.

Viewing files that contain personal data

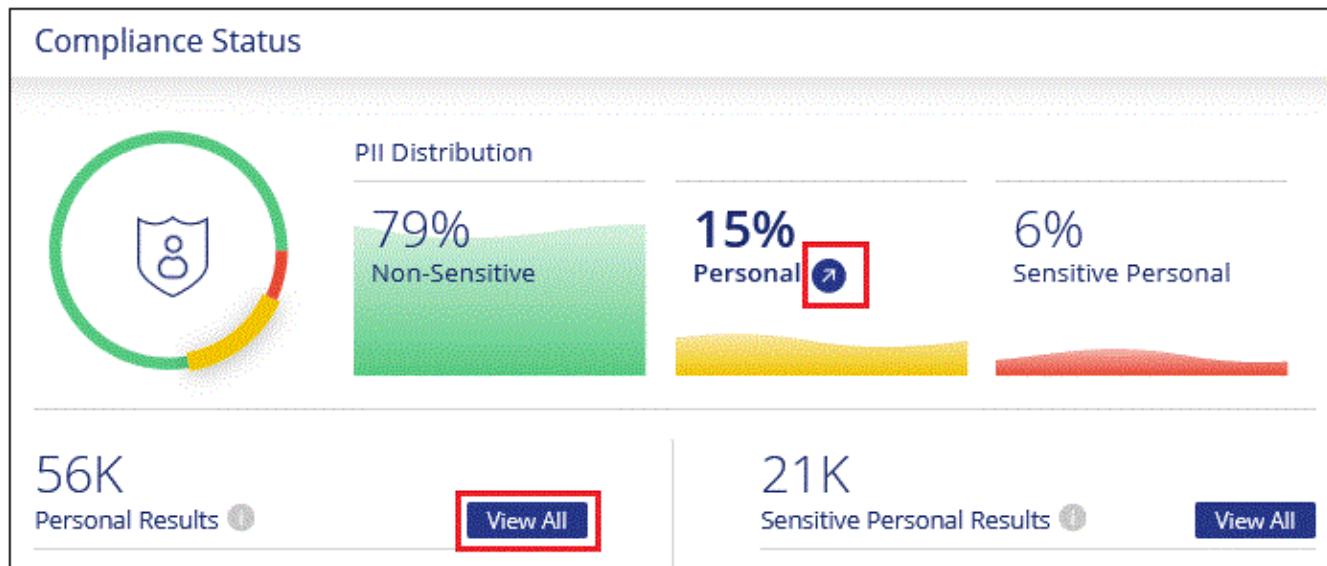
Cloud Data Sense automatically identifies specific words, strings, and patterns (Regex) inside the data. For example, Personal Identification Information (PII), credit card numbers, social security numbers, bank account numbers, and more. [See the full list](#).

Additionally, if you have added a database server to be scanned, the *Data Fusion* feature allows you to scan your files to identify whether unique identifiers from your databases are found in those files or other databases. See [Adding personal data identifiers using Data Fusion](#) for details.

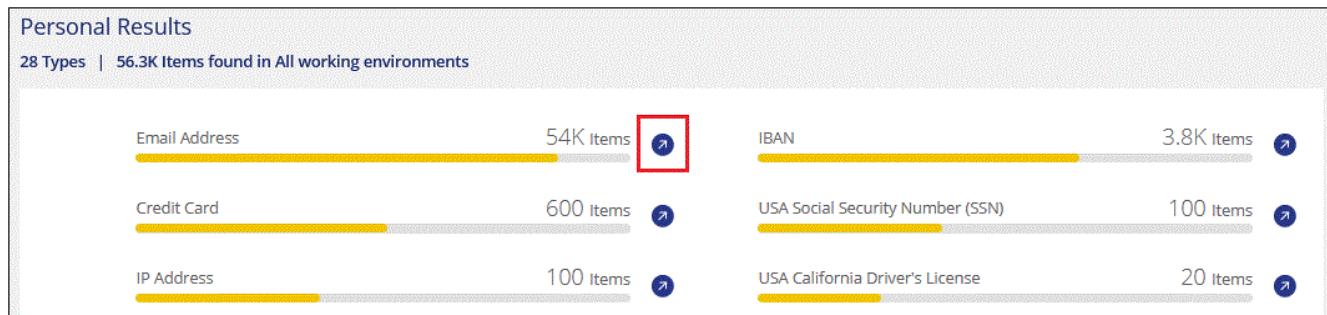
For some types of personal data, Data Sense uses *proximity validation* to validate its findings. The validation occurs by looking for one or more predefined keywords in proximity to the personal data that was found. For example, Data Sense identifies a U.S. social security number (SSN) as a SSN if it sees a proximity word next to it—for example, *SSN* or *social security*. The [table of personal data](#) shows when Data Sense uses proximity validation.

Steps

1. At the top of Cloud Manager, click **Data Sense** and click the **Compliance** tab.
2. To investigate the details for all personal data, click the icon next to the personal data percentage.



3. To investigate the details for a specific type of personal data, click **View All** and then click the **Investigate Results** icon for a specific type of personal data; for example, email addresses.



4. Investigate the data by searching, sorting, expanding details for a specific file, clicking **Investigate Results** to see masked information, or by downloading the file list.

The screenshot shows the Cloud Data Sense interface. At the top, there are tabs for 'Unstructured (54K Files)' and 'Structured (3 Tables)'. A search bar is at the top right. Below the tabs, a blue header bar has 'File Name' and dropdown menus for 'Personal', 'Sensitive Personal', 'Data Subjects', and 'File Type'. The main content area shows a file named 'customer-data.xls'. To its right are icons for S3 (blue), 838 (grey), 0 (grey), 63 (red box), and XLS. Below these are several details: Working Environment (Account): S3 - 759995470648; Storage Repository (Bucket): compliancefiles; File Path: /Patterns/NEW SSN/custo...; Category: Miscellaneous Spreadsheets; File Size: 142.35 KB; Last Modified: 2019-12-16 12:18; Open Permissions: NOT PUBLIC; and Duplicates: 2. There are also 'Assign a Label to this file' and 'Delete this file' buttons. A link 'Give feedback on this result' is at the bottom right.

Viewing files that contain sensitive personal data

Cloud Data Sense automatically identifies special types of sensitive personal information, as defined by privacy regulations such as [articles 9 and 10 of the GDPR](#). For example, information regarding a person's health, ethnic origin, or sexual orientation. See the full list.

Cloud Data Sense uses artificial intelligence (AI), natural language processing (NLP), machine learning (ML), and cognitive computing (CC) to understand the meaning of the content that it scans in order to extract entities and categorize it accordingly.

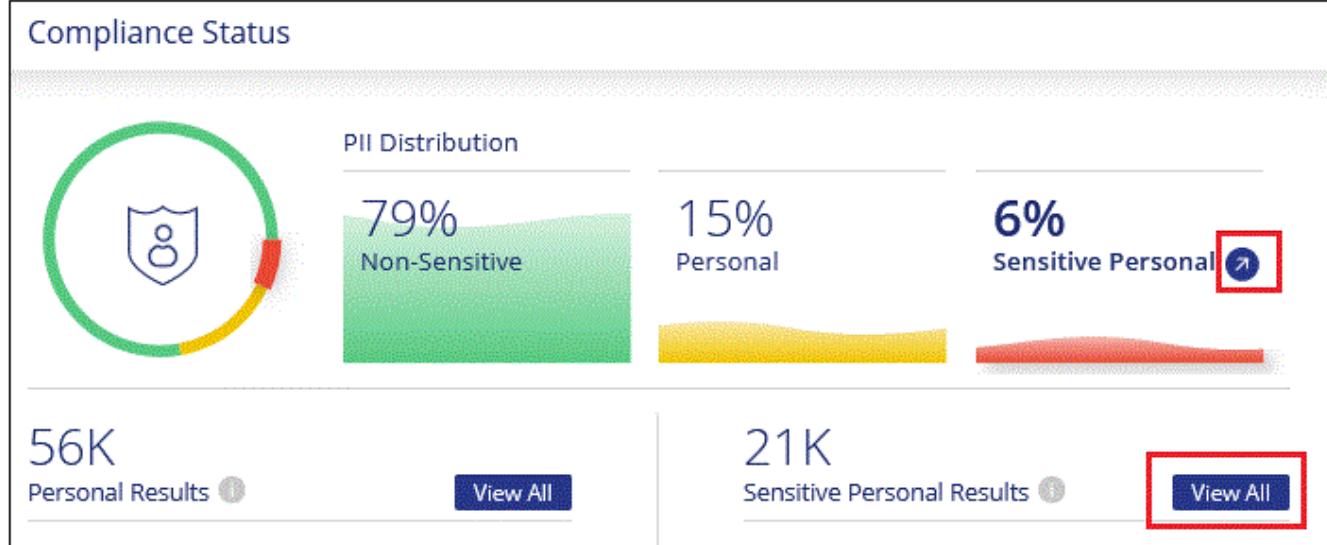
For example, one sensitive GDPR data category is ethnic origin. Because of its NLP abilities, Data Sense can distinguish the difference between a sentence that reads "George is Mexican" (indicating sensitive data as specified in article 9 of the GDPR), versus "George is eating Mexican food."



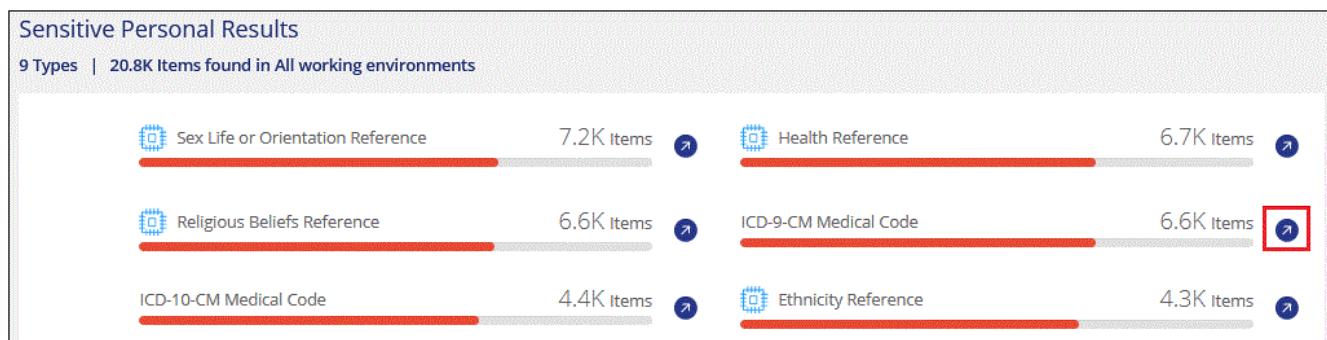
Only English is supported when scanning for sensitive personal data. Support for more languages will be added later.

Steps

1. At the top of Cloud Manager, click **Data Sense** and click the **Compliance** tab.
2. To investigate the details for all sensitive personal data, click the icon next to the sensitive personal data percentage.



- To investigate the details for a specific type of sensitive personal data, click **View All** and then click the **Investigate Results** icon for a specific type of sensitive personal data.



- Investigate the data by searching, sorting, expanding details for a specific file, clicking **Investigate Results** to see masked information, or by downloading the file list.

Viewing files by categories

Cloud Data Sense takes the data that it scanned and divides it into different types of categories. Categories are topics based on AI analysis of the content and metadata of each file. [See the list of categories](#).

Categories can help you understand what's happening with your data by showing you the types of information that you have. For example, a category like resumes or employee contracts can include sensitive data. When you investigate the results, you might find that employee contracts are stored in an insecure location. You can then correct that issue.



Only English is supported for categories. Support for more languages will be added later.

Steps

- At the top of Cloud Manager, click **Data Sense** and click the **Compliance** tab.
- Click the **Investigate Results** icon for one of the top 4 categories directly from the main screen, or click **View All** and then click the icon for any of the categories.



3. Investigate the data by searching, sorting, expanding details for a specific file, clicking **Investigate Results** to see masked information, or by downloading the file list.

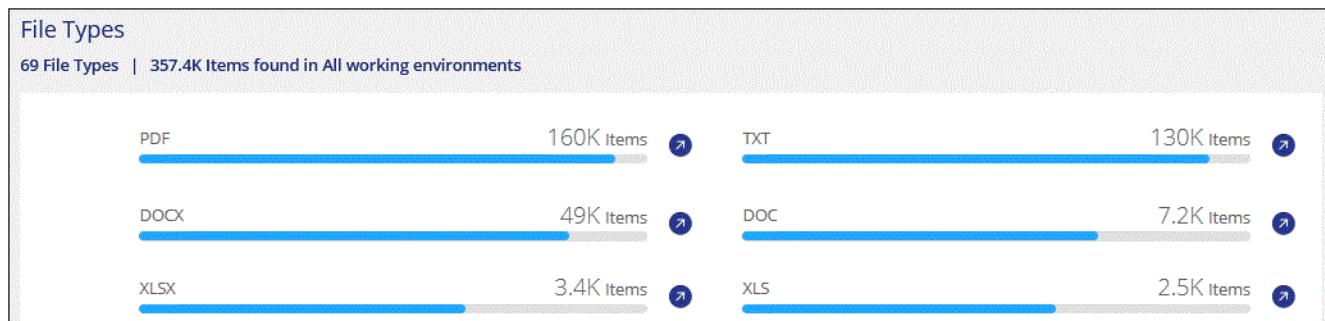
Viewing files by file types

Cloud Data Sense takes the data that it scanned and breaks it down by file type. Reviewing your file types can help you control your sensitive data because you might find that certain file types are not stored correctly. [See the list of file types](#).

For example, you might be storing CAD files that include very sensitive information about your organization. If they are unsecured, you can take control of the sensitive data by restricting permissions or moving the files to another location.

Steps

1. At the top of Cloud Manager, click **Data Sense** and click the **Compliance** tab.
2. Click the **Investigate Results** icon for one of the top 4 file types directly from the main screen, or click **View All** and then click the icon for any of the file types.



3. Investigate the data by searching, sorting, expanding details for a specific file, clicking **Investigate Results** to see masked information, or by downloading the file list.

Viewing file metadata

In the Data Investigation results pane you can click for any single file to view the file metadata.

In addition to showing you the working environment and volume where the file resides, the metadata shows much more information, including the file permissions, file owner, whether there are duplicates of this file, and assigned AIP label (if you have [integrated AIP in Cloud Data Sense](#)). This information is useful if you're planning to [create Policies](#) because you can see all the information that you can use to filter your data.

Note that not all information is available for all data sources - just what is appropriate for that data source. For example, volume name, permissions, and AIP labels are not relevant for database files.

When viewing the details for a single file there are a few actions you can take on the file:

- You can move the file to any NFS share. See [Moving source files to an NFS share](#) for details.
- You can delete the file. See [Deleting source files](#) for details.
- You can assign a certain Status to the file. See [Applying Status tags](#) for details.
- You can assign the file to a Cloud Manager user to be responsible for any follow-up actions that need to be done on the file. See [Assigning users to a file](#) for details.
- If you have integrated AIP labels with Cloud Data Sense, you can assign a label to this file, or change to a different label if one already exists. See [Assigning AIP labels manually](#) for details.

Viewing permissions for files

To view a list of all users or groups who have access to a file, and the types of permissions they have, click [View all Permissions](#).

The screenshot shows a file named "Expense Report TPO-1060.pdf" with the following details:

- Working Environment:** WorkingEnvironment1
- Repository:** Volume Name
- File Path:** /Prod/labs-base/Expense Report TPO-1060.pdf
- Category:** Legal
- File Size:** 22 MB
- Last Modified:** 2019-08-06 07:51
- Open Permissions:** NO OPEN PERMISSIONS
- File Owner:** Avy

A red box highlights the "View all Permissions" button, which is labeled "View all Permissions" with a cursor icon pointing at it.

A modal window titled "Permissions list for file Expense Report TPO-1060.pdf" displays the following table:

Group or User	Read	Write
user1@company.com	✓	✗
user2@company.com	✓	✓
dist_list_IT@company.com	✓	✗
user4@company.com	✓	✓

A "Close" button is located at the bottom right of the modal.

This button is available only for files in CIFS shares.

Checking for duplicate files in your storage systems

You can view if duplicate files are being stored in your storage systems. This is useful if you want to identify areas where you can save storage space. It can also be helpful to make sure certain files that have specific permissions or sensitive information are not unnecessarily duplicated in your storage systems.

You can download the list of duplicate files and send it to your storage admin so they can decide which files, if any, can be deleted. Or you can [delete the file](#) yourself if you are confident that a specific version of the file is not needed.

Viewing all duplicated files

If you want a list of all files that are duplicated in the working environments and data sources you are scanning, you can use the filter called **Duplicates > Has duplicates** in the Data Investigation page.

All files with duplicates from all file types (not including databases), with a minimum size of 50 MB, and/or containing personal or sensitive personal information, will show in the Results page.

Viewing if a specific file is duplicated

If you want to see if a single file has duplicates, in the Data Investigation results pane you can click **▼** for any single file to view the file metadata. If there are duplicates of a certain file, this information appears next to the **Duplicates** field.

To view the list of duplicate files and where they are located, click **View Details**. In the next page click **View Duplicates** to view the files in the Investigation page.

The screenshot shows the Cloud Data Sense interface. On the left, there's a summary bar with icons for Last Modified (2019-08-06 07:51), Open Permissions (NO OPEN PERMISSIONS), File Owner (Asaf Ley), and Duplicates (3). A red box highlights the 'View Details' button next to the duplicates count. A modal window titled 'Duplicates of File 'Name 1'' is open, showing 'Duplicates: 3', 'Total Size of all Duplicates: 1GB', and 'File Hash: xxxxxx'. It has 'View Duplicates' and 'Close' buttons, with 'View Duplicates' also highlighted by a red box. Below the modal, a table lists three items under '3 items', with columns for File Name, Personal, Sensitive Personal, Data Subjects, File Type, and a dropdown arrow. Each row contains a checkbox, the file name 'Expense Report EXP-TPO-106038887654', the label 'cvo', counts 6, 3, and 16, and PDF file type.

<input type="checkbox"/> File Name	Personal	Sensitive Personal	Data Subjects	File Type	
<input type="checkbox"/> Expense Report EXP-TPO-106038887654 cvo 6 3 16 PDF ▾					
<input type="checkbox"/> Expense Report EXP-TPO-106038887654 cvo 6 3 16 PDF ▾					
<input type="checkbox"/> Expense Report EXP-TPO-106038887654 cvo 6 3 16 PDF ▾					



You can use the "file hash" value provided in this page and enter it directly in the Investigation page to search for a specific duplicate file at any time - or to be used in a Policy.

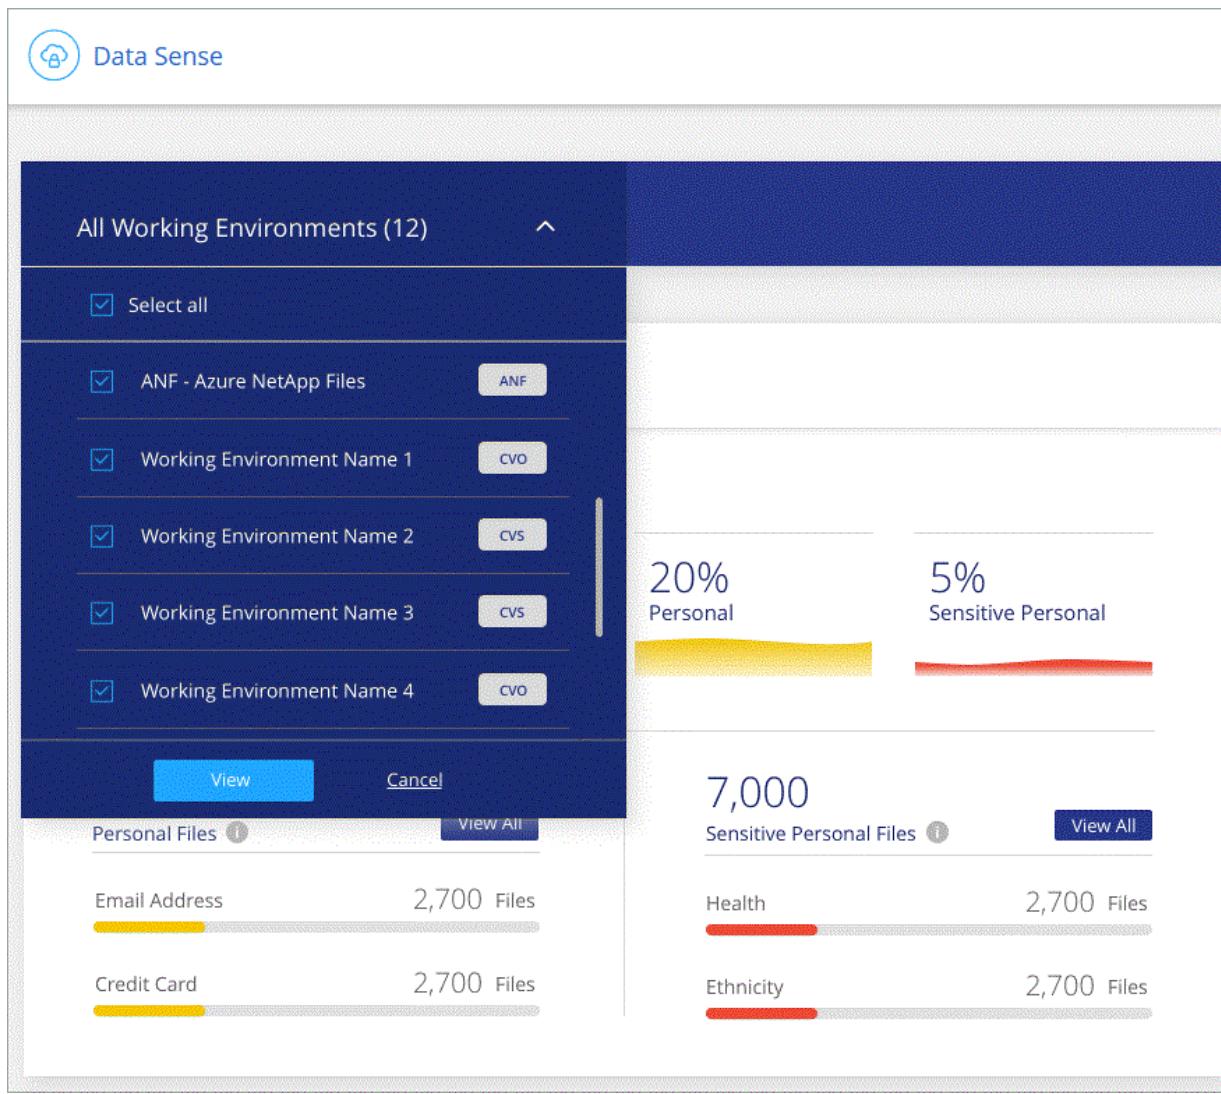
Viewing Dashboard data for specific working environments

You can filter the contents of the Cloud Data Sense dashboard to see compliance data for all working environments and databases, or for just specific working environments.

When you filter the dashboard, Data Sense scopes the compliance data and reports to just those working environments that you selected.

Steps

1. Click the filter drop-down, select the working environments that you'd like to view data for, and click **View**.



Filtering data in the Data Investigation page

You can filter the contents of the investigation page to display only the results you want to see. If you want to save a CSV version of the content as a report after you have refined it, click the button.

Data Investigation		Unstructured (32K Files)		Structured (323 DB Tables)			
FILTERS		File Name		Personal	Sensitive Personal	Data Subjects	File Type
<input type="text"/> Search filters	x	Expense Report EXP-TPO-10603888765435	cvo	6	3	16	PDF
Policies	+	Expense Report EXP-TPO-10603888765435	cvo	6	3	16	PDF
Working Environment	4	Expense Report EXP-TPO-10603888765435	cvo	6	3	16	PDF
Storage Repository	+	Expense Report EXP-TPO-10603888765435	cvo	6	3	16	PDF
Category	+	Expense Report EXP-TPO-10603888765435	cvo	6	3	16	PDF
Private Data	6	Expense Report EXP-TPO-10603888765435	cvo	6	3	16	PDF
File Type	+	Expense Report EXP-TPO-10603888765435	cvo	6	3	16	PDF

- The top-level tabs allow you to view data from files (unstructured data) or from databases (structured data).
- The controls at the top of each column allow you to sort the results in numerical or alphabetical order.
- The left-pane filters enable you to refine the results by selecting from the following attributes:
 - Policies
 - Open Permissions
 - File Owner
 - Label
 - Working Environment
 - Storage Repository
 - File Path
 - Category
 - Sensitivity Level
 - Personal Data
 - Sensitive Personal Data
 - Data Subject
 - File Type
 - File Size
 - Created date
 - Last Modified date
 - Last Accessed date (For the [types of files that we scan for content](#) this is the last time Data Sense scanned the file.)
 - Has Duplicates
 - Status
 - Assigned To
 - File Hash

- The *Policies* filter at the top of the Filters pane lists the custom filters that provide commonly requested combinations of filters; like a saved database query or Favorites list. Go [here](#) to view the list of predefined Policies and to see how you can create your own custom Policies.

What's included in each file list report (CSV file)

From each Investigation page you can click the  button to download file lists (in CSV format) that include details about the identified files. If Data Sense is scanning both Structured (database tables) and Unstructured (files) data, there are two reports contained in the downloaded ZIP file.

If there are more than 10,000 results, only the top 10,000 appear in the list.

The **Unstructured Data Report** includes the following information:

- File name
- Location type
- Working environment
- Storage repository
- Protocol type
- File path
- File type
- Created time
- Last modified
- Last accessed
- File size
- File owner
- Category
- Personal information
- Sensitive personal information
- Deletion detection date

A deletion detection date identifies the date that the file was deleted or moved. This enables you to identify when sensitive files have been moved. Deleted files aren't part of the file number count that appears in the dashboard or on the Investigation page. The files only appear in the CSV reports.

The **Structured Data Report** includes the following information:

- DB Table name
- Location type
- Working environment
- Storage repository
- Column count
- Row count
- Personal information

- Sensitive personal information

Managing your private data

Cloud Data Sense provides many ways for you to manage your private data. Some functionality just makes it easier to see the data that is most important to you, and other functionality allows you to make changes to the data.

- Using the "Policy" functionality you can create your own custom search queries so that you can easily see the results by clicking one button.
- You can send email alerts to Cloud Manager users when certain critical Policies return results.
- You can add Tags to files that you want to mark for organization or for some type of follow-up.
- If you are subscribed to [Azure Information Protection \(AIP\)](#) to classify and protect your files, you can use Cloud Data Sense to manage those AIP labels.
- You can delete files that seem insecure or too risky to leave in your storage system, or that you have identified as a duplicate.



The capabilities described in this section are available only if you have chosen to perform a full classification scan on your data sources. Data sources that have had a mapping-only scan do not show file-level details.

Controlling your data using Policies

Policies are like a favorites list of custom filters that provide search results in the Investigation page for commonly requested compliance queries. Cloud Data Sense provides a set of predefined Policies based on common customer requests. You can create custom Policies that provide results for searches specific to your organization.

Policies provide the following functionality:

- [Predefined Policies](#) from NetApp based on user requests
- Ability to create your own custom Policies
- Launch the Investigation page with the results from your Policies in one click
- Send email alerts to Cloud Manager users when certain critical Policies return results so you can get notifications to protect your data
- Assign AIP (Azure Information Protection) labels automatically to all files that match the criteria defined in a Policy

The **Policies** tab in the Compliance Dashboard lists all the predefined and custom Policies available on this instance of Cloud Data Sense.

The screenshot shows the 'Policies List' section of the Data Sense interface. It displays two predefined policies:

- GDPR - Old Sensitive Data**: A Predefined Policy. Email notifications are set to ON. The description states: "Data with European IDs for GDPR received from XDR database, sharing with Legal area in charged of Jon Doe. Also, this lines can take up to 2 lines, we need to limit character input so the description does not take more than 2 lines here in the component."
- HIPAA - Patients Personal Data**: Last modified on 17-10-20. Email notifications are set to OFF. The description states: "The Health Insurance Portability and Accountability Act of 1996 (HIPAA) is a federal law that required the creation of national standards to protect sensitive patient health information from being disclosed without the patient's consent or knowledge."

In addition, Policies appear in the list of Filters in the Investigation page.

Viewing Policy results in the Investigation page

To display the results for a Policy in the Investigation page, click the button for a specific Policy, and then select **Investigate Results**.

The screenshot shows the 'Policies List' section of the Data Sense interface. For the 'GDPR - Old Sensitive Data' policy, the 'More options' button has been clicked, revealing a dropdown menu. The 'Investigate Results' option is highlighted with a red box.

Creating custom Policies

You can create your own custom Policies that provide results for searches specific to your organization.

Steps

1. From the Data Investigation page, define your search by selecting all the filters you want to use. See [Filtering data in the Data Investigation page](#) for details.
2. Once you have all the filter characteristics just the way you want them, click **Create Policy from this search**.

The screenshot shows the 'Data Investigation' interface with the 'FILTERS' tab selected. At the top right is a 'Clear All' button. Below it is a search bar with the placeholder 'Search filters' and a clear 'X' icon. A list of filter categories is shown, each with a '+' sign to expand: 'Policies', 'Working Environment' (with a count of 4), 'Storage Repository', 'Category', 'Private Data' (with a count of 6), and 'File Type'. At the bottom of the filter list is a blue button with white text that reads 'Create Policy from this search', which is highlighted with a red rectangular border.

3. Name the Policy and select other actions that can be performed by the Policy:
 - a. Enter a unique name and description.
 - b. Optionally, check the box if you want notification emails sent to Cloud Manager users, and choose the interval at which the email is sent.
 - c. Optionally, check the box to automatically assign AIP labels to files that match the Policy parameters, and select the label. (Only if you have already integrated AIP labels. Learn more about [AIP labels](#).)
 - d. Click **Create Policy**.

Create Policy

This will create a new Policy according to the current selected filters and search term. You can view or delete this later from the "Policies" tab.

Note it may take up to 15 minutes for results to be displayed for a new Policy.

Name this Policy

New Policy to view all files that were created over 60 days ago

Give it a detailed description that explains what it searches for

See if any files greater than 60 days old should be deleted from the file system.

Send email updates about this Policy to Cloud Manager users on this account every Day ▾

Automatically label this Policy's matches with: Select a label ▾

Create Policy

Cancel

Result

The new Policy appears in the Policies tab.

Editing Policies

You can modify certain parts of a Policy depending on the type of Policy:

- Custom Policies - You can modify the *Name*, the *Description*, whether email notifications are sent, and whether AIP labels are added.
- Predefined Policies - You can modify only whether email notifications are sent and whether AIP labels are added.



If you need to change the filter parameters for a custom Policy, you'll need to create a new Policy with the parameters you want, and then delete the old Policy.

To modify a Policy, click the **Edit** button, enter your changes on the *Edit Policy* page, and click **Save Policy**.

Deleting Policies

You can delete any custom Policy that you created if you no longer need it. You can't delete any of the predefined Policies.

To delete a Policy, click the button for a specific Policy, click **Delete Policy**, and then click **Delete Policy** again in the confirmation dialog.

Applying tags to manage your scanned files

You can add a tag to files that you want to mark for some type of follow-up. For example, you may have found some duplicate files and you want to delete one of them, but you need to check to see which one should be deleted. You could add a tag of "Check to delete" to the file so you know this file requires some research and some type of future action.

Data Sense enables you to view the tags that are assigned to files, add or remove tags from files, and change the name or delete an existing tag.

Note that the tag is not added to the file in the same way as AIP Labels are part of the file metadata. The tag is just seen by Cloud Manager users using Cloud Data Sense so you can see if a file needs to be deleted or checked for some type of follow-up.



Tags assigned to files in Cloud Data Sense are not related to the tags you can add to resources, such as volumes or virtual machine instances. Data Sense tags are applied at the file level.

Viewing files that have certain tags applied

You can view all the files that have specific tags assigned.

1. Click the **Investigation** tab from Cloud Data Sense.
2. In the Data Investigation page, click **Tags** in the Filters pane and then select the required tags.

The screenshot shows the 'FILTERS' pane of the Cloud Data Sense interface. At the top, there's a 'Clear All' button. Below it is a search bar labeled 'Search filters'. The 'Working Environment' filter is set to 4 items. The 'Tags' filter is set to 6 items, with 'Action Required' checked (indicated by a blue circle with a white checkmark) and the other three options ('Not For Review', 'For Review', 'Approved Deletion') unchecked (indicated by empty circles).

The Investigation Results pane displays all the files that have those tags assigned.

Assigning tags to files

You can add tags to a single file or to a group of files.

To add a tag to a single file:

Steps

1. In the Data Investigation results pane, click **▼** for the file to expand the file metadata details.
2. Click the **Tags** field and the currently assigned tags are displayed.

3. Add the tag or tags:

- To assign an existing tag, click in the **New Tag...** field and start typing the name of the tag. When the tag you are looking for appears, select it and press **Enter**.
- To create a new tag and assign it to the file, click in the **New Tag...** field, enter the name of the new tag, and press **Enter**.

The screenshot shows the Data Investigation results pane for an expense report. The top bar includes navigation icons and a PDF download button. Below the bar, the file's title is "Expense Report EXP-TPO-106038887654". The "Tags" section shows "Approved Deletion" and a "New" button. A dropdown menu titled "Tags: 5 Tags" is open, listing "Approved Deletion", "New", "For Review", and "Another Tag". A sub-menu for "Approved" is also open, showing "New Tag", "Approved" (which is selected), "For Review", and "Another Tag". Other metadata listed include "Working Environment: WorkingEnvironment1", "Repository: Volume Name", "File Path: /Prod/labs-base/Expense Report EXP-TPO-1060388.pdf", "Category: Legal", "File Size: 22 MB", and "Last Modified: 2019-08-06 07:51". The "Open Permissions" section shows "NO OPEN PERMISSIONS" and a "View all Permissions" link.

The tag appears in the file metadata.

To add a tag to multiple files:

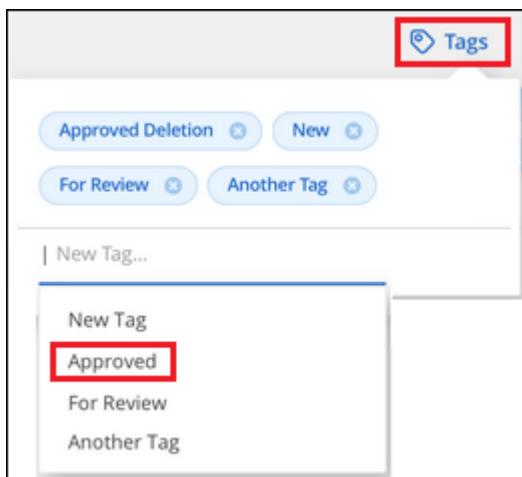
Steps

1. In the Data Investigation results pane, select the file, or files, that you want to tag.

The screenshot shows the Data Investigation results pane with 2345 items. The "Tags" button is highlighted. The results table lists four expense reports. The first two rows have the "File Name" column checked, indicating they are selected. The columns include "File Name", "cvo", "Personal", "Sensitive Personal", "Data Subjects", "File Type", and a dropdown arrow. The selected files are "Expense Report EXP-TPO-106038887654" and "Expense Report EXP-TPO-106038887654".

- To select all files on the current page, check the box in the title row (**File Name**). (You can't select files from more than one page.)
 - To select individual files, check the box for each file (**Volume_1**).
2. From the button bar, click **Tags** and the currently assigned tags are displayed.
 3. Add the tag or tags:
 - To assign an existing tag, click in the **New Tag...** field and start typing the name of the tag. When the tag you are looking for appears, select it and press **Enter**.

- To create a new tag and assign it to the file, click in the **New Tag...** field, enter the name of the new tag, and press **Enter**.



- Approve adding the tags in the confirmation dialog and the tags are added to the metadata for all selected files.

Deleting tags from files

You can delete a tag if you don't need to use it anymore.

Just click the **x** for an existing tag.



If you had selected multiple files, the tag is removed from all the files.

Assigning users to manage certain files

You can assign a Cloud Manager user to a specific file, or to multiple files, so that person can be responsible for any follow-up actions that need to be done on the file. This capability is often used with the feature to add custom Status tags to a file.

For example, you might have a file that contains certain personal data that allows too many users read and write access (open permissions). So you could assign the Status tag "Change permissions" and assign this file to user "Joan Smith" so they can decide how to fix the issue. When they have fixed the issue they could change the Status tag to "Completed".

Note that the user name is not added to the file as part of the file metadata - it is just seen by Cloud Manager users when using Cloud Data Sense.

A new Filter in the Investigation page enables you to easily view all files that have the same person in the "Assigned To" field.

To assign a user to a single file:

Steps

- In the Data Investigation results pane, click **▼** for the file to expand the file metadata details.
- Click the **Assigned to** field and select the user name.

The screenshot shows the Data Investigation results pane. On the left, a list of file metadata includes: Working Environment: WorkingEnvironment1, Repository: Volume Name, File Path: /Prod/labs-base/Expense Report EXP-TPO-1060388.pdf, Category: Legal, File Size: 22 MB, Created: 2013-01-05 08:22, Last Modified: 2019-08-06 07:51, and Last Accessed: 2019-08-06 07:51. On the right, there are two dropdown menus: 'Status: To Check' and 'Assigned to: Ariel'. Below these is a search bar labeled 'Search' and a list of users: None, Ariel (highlighted with a red box), and Nadav.

The User name appears in the file metadata.

To assign a user to multiple files:

Steps

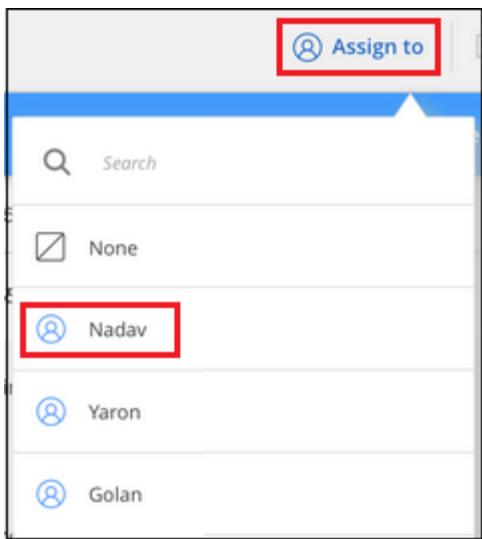
1. In the Data Investigation results pane, select the file, or files, that you want to assign to a user.

The screenshot shows the Data Investigation results pane with 2345 items. The 'Assign to' button is highlighted. A table lists four files: Expense Report EXP-TPO-10603887654 (checked), Expense Report EXP-TPO-10603887654 (checked), Expense Report EXP-TPO-10603887654 (unchecked), and Expense Report EXP-TPO-10603887654 (unchecked). The checked files are highlighted with a red box.

File Name	Tags	Personal	Sensitive Personal	Data Subjects	File Type
Expense Report EXP-TPO-10603887654	cvo	6	3	16	PDF
Expense Report EXP-TPO-10603887654	cvo	6	3	6	PDF
Expense Report EXP-TPO-10603887654	cvo	6	3	6	PDF
Expense Report EXP-TPO-10603887654	cvo	6	3	6	PDF

- To select all files on the current page, check the box in the title row (File Name). (You can't select files from more than one page.)
- To select individual files, check the box for each file (Volume_1).

2. From the button bar, click **Assign to** and select the user name:



The user is added to the metadata for all selected files.

Categorizing your data using AIP labels

You can manage AIP labels in the files that Cloud Data Sense is scanning if you have subscribed to [Azure Information Protection \(AIP\)](#). AIP enables you to classify and protect documents and files by applying labels to content. Data Sense enables you to view the labels that are already assigned to files, add labels to files, and change labels when a label already exists.

Cloud Data Sense supports AIP labels within the following file types: .DOC, .DOCX, .PDF, .PPTX, .XLS, .XLSX.

Note that you can't currently change labels in files larger than 30 MB. For OneDrive accounts the maximum file size is 4 MB.



If a file has a label which doesn't exist anymore in AIP, Cloud Data Sense considers it as a file without a label.

Integrating AIP labels in your workspace

Before you can manage AIP labels, you need to integrate the AIP label functionality into Cloud Data Sense by signing into your existing Azure account. Once enabled, you can manage AIP labels within files for all [working environments and data sources](#) in your Cloud Manager workspace.

Requirements

- You must have an account and an Azure Information Protection license.
- You must have the login credentials for the Azure account.
- If you plan to change labels in files that reside in Amazon S3 buckets, ensure that the permission s3:PutObject is included in the IAM role. See [setting up the IAM role](#).

Steps

1. From the Cloud Data Sense Configuration page, click **Integrate AIP Labels**.

(2/20) Working Environments

Filter by: CVO ANF S3 DB ONEDR Clear filters

Add Data Source | ▾

2. In the Integrate AIP Labels dialog, click **Sign in to Azure**.
3. In the Microsoft page that appears, select the account and enter the required credentials.
4. Return to the Cloud Data Sense tab and you'll see the message "*AIP Labels were integrated successfully with the account <account_name>*".
5. Click **Close** and you'll see the text *AIP Labels integrated* at the top of the page.

(2/20) Working Environments

Filter by: CVO ANF S3 DB ONEDR Clear filters

AIP Labels integrated ▾ Add Data Source | ▾

Result

You can view and assign AIP labels from the results pane of the Investigation page. You can also assign AIP labels to files using Policies.

Viewing AIP labels in your files

You can view the current AIP label that is assigned to a file.

In the Data Investigation results pane, click **▼** for the file to expand the file metadata details.

File Name	Personal	Sensitive Personal	Data Subjects	File Type
Expense Report EXP-TPO-10603888765435	cvo	6	3	16 PDF
Expense Report EXP-TPO-10603888765435	cvo	6	3	16 PDF

Working Environment: WorkingEnvironment1

Repository: Volume Name

Label: Finance

Assigning AIP labels manually

You can add, change, and remove AIP labels from your files using Cloud Data Sense.

Follow these steps to assign an AIP label to a single file.

Steps

- In the Data Investigation results pane, click for the file to expand the file metadata details.

The screenshot shows the Data Investigation results pane. At the top, there are two tabs: "Unstructured (32K Files)" (selected) and "Structured (323 DB Tables)". Below the tabs is a header row with columns: File Name, Personal, Sensitive Personal, Data Subjects, and File Type. The first row of data is for an "Expense Report EXP-TPO-10603888765435". The "File Type" column shows "PDF" with a dropdown arrow. The "Data Subjects" column shows values 6, 3, and 16, each with a circular edit icon. The "Personal" column shows "cvo". In the "Data Subjects" column, the value "16" is highlighted with a red box. To the right of the table, there is a sidebar with file details: Working Environment (WorkingEnvironment1), Repository (Volume Name), File Path (/Prod/labs-base/Expense Report EXP-TPO-1060388.pdf), Category (Legal), File Size (22 MB), Last Modified (2019-08-06 07:51), Open Permissions (NO OPEN PERMISSIONS), and File Owner (Assaf Vol). On the far right, there is a "Label" section with a button labeled "Assign a Label to this file" and a dropdown menu containing "General", "Finance" (highlighted with a red box), and "Confidential".

- Click **Assign a Label to this file** and then select the label.

The label appears in the file metadata.

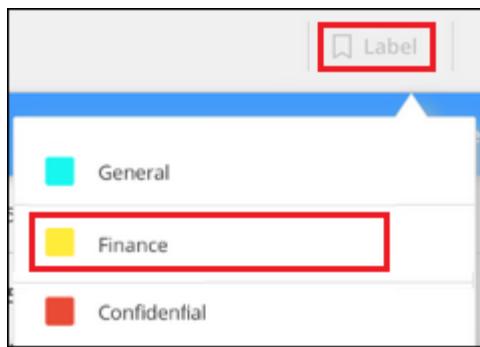
To assign an AIP label to multiple files:

Steps

- In the Data Investigation results pane, select the file, or files, that you want to label.

The screenshot shows the Data Investigation results pane with a total of 2345 items. The title bar includes buttons for Tags, Assign to, Label, Move, and Delete. Below the title bar is a table with columns: File Name, Personal, Sensitive Personal, Data Subjects, File Type, and a dropdown arrow. There are four rows of data. The first two rows have the "File Name" checkbox checked (highlighted with a red box). The third and fourth rows have the "File Name" checkbox unchecked. The "File Type" column for all rows is "PDF". The "Data Subjects" column for the first two rows shows values 6, 3, and 16, while for the last two rows it shows 6. The "Personal" column shows "cvo".

- To select all files on the current page, check the box in the title row (**File Name**). (You can't select files from more than one page.)
 - To select individual files, check the box for each file (**Volume_1**).
- From the button bar, click **Label** and select the AIP label:



The AIP label is added to the metadata for all selected files.

Assigning AIP labels automatically with Policies

You can assign an AIP label to all the files that meet the criteria of the Policy. You can specify the AIP label when creating the Policy, or you can add the label when editing any Policy.

Labels are added or updated in files continuously as Cloud Data Sense scans your files.

Depending on whether a label is already applied to a file, and the classification level of the label, the following actions are taken when changing a label:

If the file...	Then...
Has no label	The label is added
Has an existing label of a lower level of classification	The higher level label is added
Has an existing label of a higher level of classification	The higher level label is retained
Is assigned a label both manually and by a Policy	The higher level label is added
Is assigned two different labels by two Policies	The higher level label is added

Follow these steps to add an AIP label to an existing Policy.

Steps

1. From the Policies List page, click **Edit** for the Policy where you want to add (or change) the AIP label.

The screenshot shows the Policies List page in the Data Sense interface. There are two policies listed:

- GDPR - Old Sensitive Data**: Predefined Policy. Label: General. E-mail notifications: Monthly. Edit button.
- HIPAA - Patients Personal Data**: Last modified: 17-10-20. Label: OFF. E-mail notifications: OFF. Edit button (highlighted with a red box).

A green oval highlights the 'Label: OFF' field for the HIPAA policy.

2. In the Edit Policy page, check the box to enable automatic labels for files that match the Policy parameters, and select the label (for example, **General**).

The screenshot shows the 'Edit Policy' interface. It includes fields for 'Name this Policy' (set to 'HIPAA - Patient Personal Data') and 'Give it a description to quickly identify it' (set to 'Files containing patient health information that is more than 30 days old'). Below these, there are two checked checkboxes: one for sending email updates and another for automatically labeling matches. The second checkbox is followed by a dropdown menu titled 'select label' with three options: 'General' (highlighted with a red box), 'Finance', and 'Confidential'. A 'Cancel' button is also visible in the bottom right corner of the dropdown.

3. Click **Save Policy** and the label appears in the Policy description.



If a Policy was configured with a label, but the label has since been removed from AIP, the label name is turned to OFF and the label is not assigned anymore.

Removing the AIP integration

If you no longer want the ability to manage AIP labels in files, you can remove the AIP account from the Cloud Data Sense interface.

Note that no changes are made to the labels you have added using Data Sense. The labels that exist in files will stay as they currently exist.

Steps

1. From the *Configuration* page, click **AIP Labels integrated > Remove Integration**.

The screenshot shows the 'Configuration' page with the 'AIP Labels integrated' section. It features a green checkmark icon next to the text 'AIP Labels integrated' and a red-bordered button labeled 'Remove Integration' with a red circle and minus sign icon.

2. Click **Remove Integration** from the confirmation dialog.

Sending email alerts when non-compliant data is found

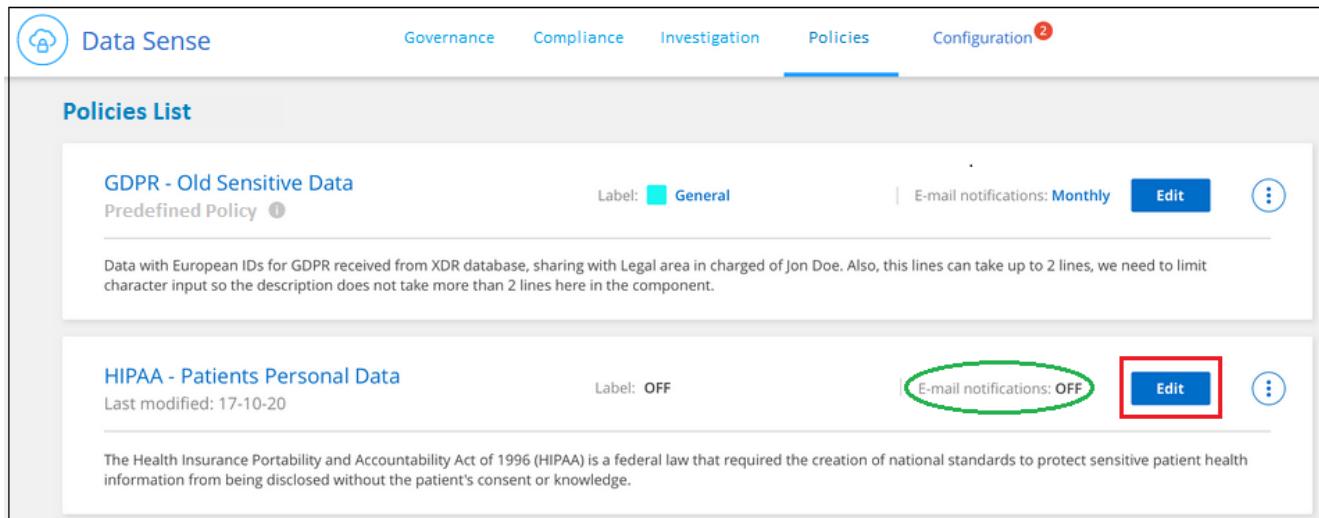
Cloud Data Sense can send email alerts to Cloud Manager users when certain critical Policies return results so you can get notifications to protect your data. You can choose to send the email notifications on a daily, weekly, or monthly basis.

You can configure this setting when creating the Policy or when editing any Policy.

Follow these steps to add email updates to an existing Policy.

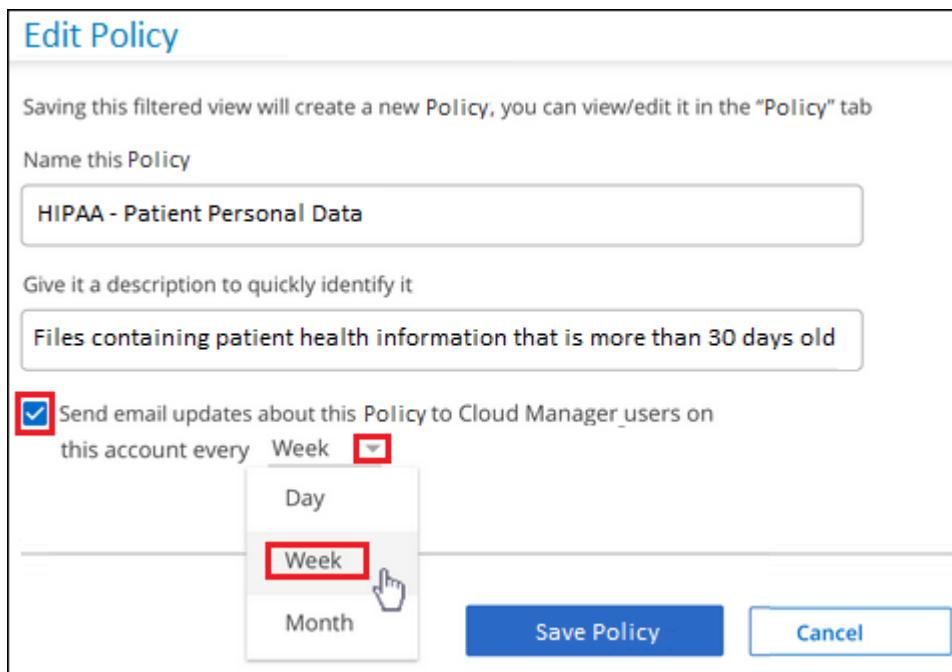
Steps

1. From the Policies List page, click **Edit** for the Policy where you want to add (or change) the email setting.



The screenshot shows the Policies List page in the Data Sense interface. It lists two policies: "GDPR - Old Sensitive Data" (Predefined Policy) and "HIPAA - Patients Personal Data". The "GDPR" policy has "General" as its label and "Monthly" as its email notification interval. The "HIPAA" policy has "OFF" as its label and "E-mail notifications: OFF" (which is circled in green). Both policies have an "Edit" button (which is highlighted with a red box) and a more options button.

2. In the Edit Policy page, check the box if you want notification emails sent to Cloud Manager users, and choose the interval at which the email is sent (for example, every **Week**).



The screenshot shows the "Edit Policy" page for the "HIPAA - Patient Personal Data" policy. It includes fields for naming the policy ("Name this Policy: HIPAA - Patient Personal Data") and giving it a description ("Give it a description to quickly identify it: Files containing patient health information that is more than 30 days old"). A checkbox labeled "Send email updates about this Policy to Cloud Manager users on this account every" is checked, and a dropdown menu shows "Week" selected (with a red box around it). At the bottom are "Save Policy" and "Cancel" buttons.

3. Click **Save Policy** and the interval at which the email is sent appears in the Policy description.

Result

The first email is sent now if there are any results from the Policy - but only if any files meet the Policy criteria. No personal information is sent in the notification emails. The email indicates that there are files that match the Policy criteria, and it provides a link to the Policy results.

Deleting source files

You can permanently remove source files that seem insecure or too risky to leave in your storage system, or that you have identified as a duplicate. This action is permanent and there is no undo.



You can't delete files that reside in databases or files that reside in volume Backups.

Requirements

You must have the Account Admin or Workspace Admin role to delete files.

Deleting files requires the following permissions:

- For NFS data – the export policy needs to be defined with write permissions.
- For CIFS data – the CIFS credentials need to have write permissions.
- For S3 data - the IAM role must include the following permission: s3:DeleteObject

Steps

1. In the Data Investigation results pane, select the file, or files, that you want to delete.

2345 items							+ Set Status	Assign to	Label	Move	Delete
<input type="checkbox"/>	File Name	Personal	Sensitive Personal	Data Subjects	File Type						
<input checked="" type="checkbox"/>	Expense Report EXP-TPO-106038887654	cvo	6	3	16	PDF					
<input checked="" type="checkbox"/>	Expense Report EXP-TPO-106038887654	cvo	6	3	6	PDF					
<input type="checkbox"/>	Expense Report EXP-TPO-106038887654	cvo	6	3	6	PDF					
<input type="checkbox"/>	Expense Report EXP-TPO-106038887654	cvo	6	3	6	PDF					

- To select all files on the current page, check the box in the title row (File Name). (You can't select files from more than one page.)
 - To select individual files, check the box for each file (Volume_1).
2. From the button bar, click **Delete**.
 3. Because the delete operation is permanent, you must type "permanently delete" in the subsequent *Delete File* dialog and click **Delete File**.

Note that you can also delete an individual file when viewing the metadata details for a file. Just click **Delete this file**.

Unstructured (32K Files)		Structured (323 DB Tables)			
File Name		Personal	Sensitive Personal	Data Subjects	File Type
<input type="checkbox"/> Expense Report EXP-TPO-10603888765435	cvo	6	3	16	PDF
<input type="checkbox"/> Expense Report EXP-TPO-10603888765435	cvo	6	3	16	PDF
Working Environment: WorkingEnvironment1 Repository: Volume Name File Path: /Prod/labs-base/Expense Report EXP-TPO-1060388.pdf					
Assign a Label to this file ▼ Delete this file					

Moving source files to an NFS share

You can move source files that Data Sense is scanning to any NFS share. The NFS share does not need to be integrated with Data Sense (see [Scanning file shares](#)).



You can't move files that reside in databases or files that reside in volume Backups.

Requirements

You must have the Account Admin or Workspace Admin role to move files.

Moving files requires that the NFS share allows access from the Data Sense instance.

Steps

1. In the Data Investigation results pane, select the file, or files, that you want to move.

2345 items		+ Set Status	Assign to	Label	Move	Delete
<input type="checkbox"/>	File Name	Personal	Sensitive Personal	Data Subjects	File Type	
<input checked="" type="checkbox"/>	Expense Report EXP-TPO-106038887654	cvo	6	3	16	PDF
<input checked="" type="checkbox"/>	Expense Report EXP-TPO-106038887654	cvo	6	3	6	PDF
<input type="checkbox"/>	Expense Report EXP-TPO-106038887654	cvo	6	3	6	PDF
<input type="checkbox"/>	Expense Report EXP-TPO-106038887654	cvo	6	3	6	PDF

- To select all files on the current page, check the box in the title row (File Name). (You can't select files from more than one page.)
 - To select individual files, check the box for each file (Volume_1).
2. From the button bar, click **Move**.

Move Multiple Files

This file will be moved to the destination folder you provide and will no longer be available at its current location.

Moving files is supported only to destination folders in NFS Shares. Any NFS Share is supported, no matter where its hosted, as long as the Share's export policy allows access from the Data Connector instance IP Address.

File Name 1

File Name 2

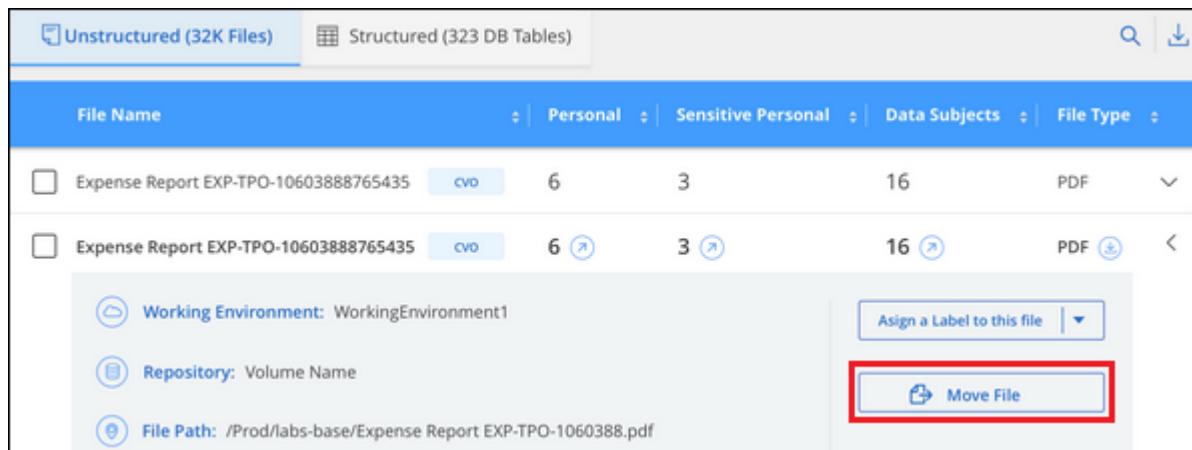
File Name 3

Enter the *NFS destination folder path* to continue

Move Files **Cancel**

3. In the *Move Files* dialog, enter the name of the NFS share where all selected files will be moved in the format <host_name>:/<share_path>, and click **Move Files**.

Note that you can also move an individual file when viewing the metadata details for a file. Just click **Move file**.



The screenshot shows a list of files under the "Unstructured (32K Files)" tab. Two files are selected, both labeled "Expense Report EXP-TPO-10603888765435". The second file's details panel is open, displaying its working environment ("WorkingEnvironment1"), repository ("Volume Name"), and file path ("/Prod/labs-base/Expense Report EXP-TPO-1060388.pdf"). A red box highlights the "Move File" button in the bottom right corner of this details panel.

List of predefined Policies

Cloud Data Sense provides the following system-defined Policies:

Name	Description	Logic
S3 publicly-exposed private data	S3 Objects containing personal or sensitive personal information, with open Public read access.	(S3 Public) AND contains personal OR sensitive personal info)

Name	Description	Logic
PCI DSS – Stale data over 30 days	Files containing Credit Card information, last modified over 30 days ago.	Contains credit card AND last modified over 30 days
HIPAA – Stale data over 30 days	Files containing Health information, last modified over 30 days ago.	Contains health data (defined same way as in HIPAA report) AND last modified over 30 days
Private data – Stale over 7 years	Files containing personal or sensitive personal information, last modified over 7 years ago.	Files containing personal or sensitive personal information, last modified over 7 years ago
GDPR – European citizens	Files containing more than 5 identifiers of an EU country's citizens or DB Tables containing identifiers of an EU country's citizens.	Files containing over 5 identifiers of an (one) EU citizens or DB Tables containing rows with over 15% of columns with one country's EU identifiers. (any one of the national identifiers of the European countries. Does not include Brazil, California, USA SSN, Israel, South Africa)
CCPA – California residents	Files containing over 10 California Driver's License identifiers or DB Tables with this identifier.	Files containing over 10 California Driver's License identifiers OR DB Tables containing California Driver's license
Data Subject names – High risk	Files with over 50 Data Subject names.	Files with over 50 Data Subject names
Email Addresses – High risk	Files with over 50 Email Addresses, or DB Columns with over 50% of their rows containing Email Addresses	Files with over 50 Email Addresses, or DB Columns with over 50% of their rows containing Email Addresses
Personal data – High risk	Files with over 20 Personal data identifiers, or DB Columns with over 50% of their rows containing Personal data identifiers.	Files with over 20 personal, or DB Columns with over 50% of their rows containing personal
Sensitive Personal data – High risk	Files with over 20 Sensitive Personal data identifiers, or DB Columns with over 50% of their rows containing Sensitive Personal data.	Files with over 20 sensitive personal, or DB Columns with over 50% of their rows containing sensitive personal

Adding personal data identifiers using Data Fusion

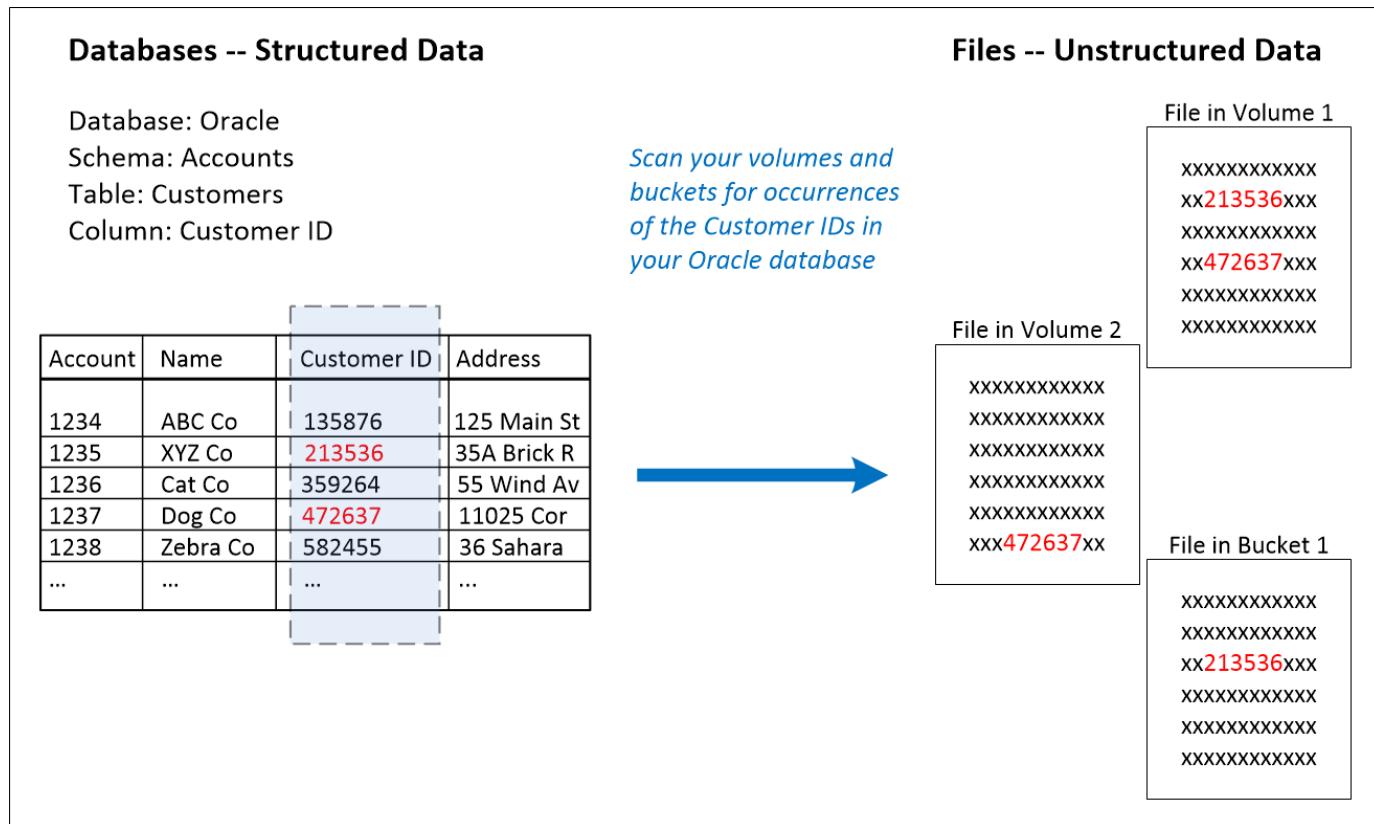
A feature we call *Data Fusion* allows you to scan your organizations' data to identify whether unique identifiers from your databases are found in files or other databases - basically making your own list of "personal data" that is identified in Cloud Data Sense scans. This gives you the full picture about where potentially sensitive data resides in *all* your files.



The capabilities described in this section are available only if you have chosen to perform a full classification scan on your data sources. Data sources that have had a mapping-only scan do not show file-level details.

Creating custom personal data identifiers from your databases

You can choose the additional identifiers that Cloud Data Sense will look for in its' scans by selecting a specific column, or columns, in a database table. For example, the diagram below shows how data fusion is used to scan your volumes, buckets, and databases for occurrences of all your Customer IDs from your Oracle database.



As you can see, two unique Customer IDs have been found in two volumes and in one S3 bucket. Any matches in database tables will also be identified.

Steps

You must have [added at least one database server](#) to Cloud Data Sense before you can add data fusion sources.

1. In the Configuration page, click **Manage Data Fusion** in the database where the source data resides.

The screenshot shows the Oracle DB 1 configuration interface. At the top, it displays "Oracle DB 1 | 41 Schemas" and "Oracle". On the right, there are buttons for "Manage Data Fusion" (highlighted with a red box), "Configuration", and three dots. Below this, there are two sections: "No Schemas selected for Compliance" and "Not Scanning" (with 7 items). The "Not Scanning" section includes a "View Details" link.

2. Click **Add Data Fusion source** on the next page.

3. In the *Add Data Fusion Source* page:

- Select the Database Schema from the drop-down menu.
- Enter the Table name in that schema.
- Enter the Column, or Columns, that contain the unique identifiers you want to use.

When adding multiple columns, enter each column name, or table view name, on a separate line.

4. Click **Add Data Fusion Source**.

The Data Fusion inventory page displays the database source columns that you have configured for Cloud Data Sense to scan.

The screenshot shows the 'DB Name 1' Data Fusion inventory page. It has a header with "+ Add Data Fusion source". Below it, a message states: "With Data Fusion, Cloud Compliance can identify occurrences of your organization's unique identifiers found in your unstructured data stores, using structured data indexes containing those unique identifiers as a source reference. [Learn More](#)". A table lists configured sources:

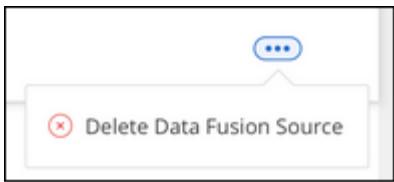
Database Schema	Table	Data Fusion Source Columns	...
SchemaName1	Table 1	Column 12, Column 14, Column 18	...
SchemaName2	Table 2	Column 12, Column 14, Column 18	...

Results

After the next scan, the results will include this new information in the Dashboard under the "Personal" results section, and in the Investigation page in the "Personal Data" filter. Each source column you added appears in the filter list as "Table.Column", for example Customers.Customer ID.

Deleting a Data Fusion source

If at some point you decide not to scan your files using a certain Data Fusion source, you can select the source row from the Data Fusion inventory page and click **Delete Data Fusion Source**.



Viewing compliance reports

Cloud Data Sense provides reports that you can use to better understand the status of your organization's data privacy program.

By default, the Cloud Data Sense dashboards display compliance and governance data for all working environments and databases. If you want to view reports that contain data for only some of the working environments, [select those working environments](#).



The reports described in this section are available only if you have chosen to perform a full classification scan on your data sources. Data sources that have had a mapping-only scan can only generate the Data Mapping Report.



NetApp can't guarantee 100% accuracy of the personal data and sensitive personal data that Cloud Data Sense identifies. You should always validate the information by reviewing the data.

Privacy Risk Assessment Report

The Privacy Risk Assessment Report provides an overview of your organization's privacy risk status, as required by privacy regulations such as GDPR and CCPA. The report includes the following information:

Compliance status

A [severity score](#) and the distribution of data, whether it's non-sensitive, personal, or sensitive personal.

Assessment overview

A breakdown of the types of personal data found, as well as the categories of data.

Data subjects in this assessment

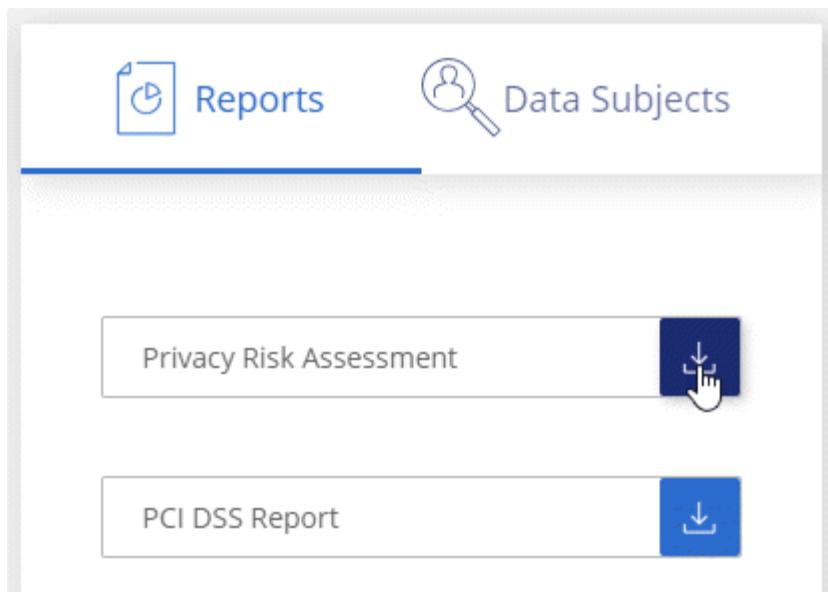
The number of people, by location, for which national identifiers were found.

Generating the Privacy Risk Assessment Report

Go to the Data Sense tab to generate the report.

Steps

1. At the top of Cloud Manager, click **Data Sense**.
2. Click **Compliance**, and then click the download icon next to **Privacy Risk Assessment** under **Reports**.



Result

Cloud Data Sense generates a PDF report that you can review and send to other groups as needed.

Severity score

Cloud Data Sense calculates the severity score for the Privacy Risk Assessment Report on the basis of three variables:

- The percentage of personal data out of all data.
- The percentage of sensitive personal data out of all data.
- The percentage of files that include data subjects, determined by national identifiers such as national IDs, Social Security numbers, and tax ID numbers.

The logic used to determine the score is as follows:

Severity score	Logic
0	All three variables are exactly 0%
1	One of the variables are larger than 0%
2	One of the variables are larger than 3%
3	Two of the variables are larger than 3%
4	Three of the variables are larger than 3%
5	One of the variables are larger than 6%
6	Two of the variables are larger than 6%
7	Three of the variables are larger than 6%
8	One of the variables are larger than 15%
9	Two of the variables are larger than 15%
10	Three of the variables are larger than 15%

PCI DSS Report

The Payment Card Industry Data Security Standard (PCI DSS) Report can help you identify the distribution of credit card information across your files. The report includes the following information:

Overview

How many files contain credit card information and in which working environments.

Encryption

The percentage of files containing credit card information that are on encrypted or unencrypted working environments. This information is specific to Cloud Volumes ONTAP.

Ransomware Protection

The percentage of files containing credit card information that are on working environments that do or don't have ransomware protection enabled. This information is specific to Cloud Volumes ONTAP.

Retention

The timeframe in which the files were last modified. This is helpful because you shouldn't keep credit card information for longer than you need to process it.

Distribution of Credit Card Information

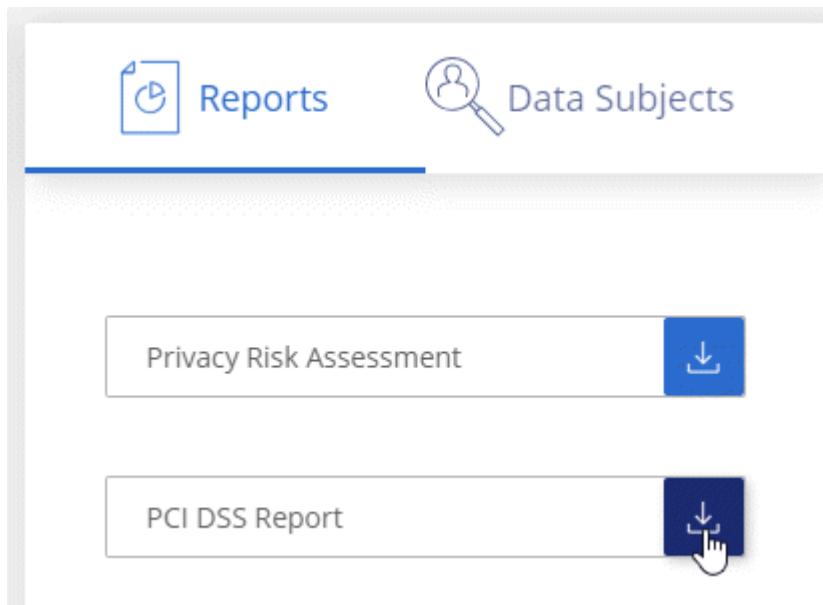
The working environments where the credit card information was found and whether encryption and ransomware protection are enabled.

Generating the PCI DSS Report

Go to the Data Sense tab to generate the report.

Steps

1. At the top of Cloud Manager, click **Data Sense**.
2. Click **Compliance**, and then click the download icon next to **PCI DSS Report** under **Reports**.



Result

Cloud Data Sense generates a PDF report that you can review and send to other groups as needed.

HIPAA Report

The Health Insurance Portability and Accountability Act (HIPAA) Report can help you identify files containing health information. It is designed to aid in your organization's requirement to comply with HIPAA data privacy laws. The information Cloud Data Sense looks for includes:

- Health reference pattern
- ICD-10-CM Medical code
- ICD-9-CM Medical code
- HR – Health category
- Health Application Data category

The report includes the following information:

Overview

How many files contain health information and in which working environments.

Encryption

The percentage of files containing health information that are on encrypted or unencrypted working environments. This information is specific to Cloud Volumes ONTAP.

Ransomware Protection

The percentage of files containing health information that are on working environments that do or don't have ransomware protection enabled. This information is specific to Cloud Volumes ONTAP.

Retention

The timeframe in which the files were last modified. This is helpful because you shouldn't keep health information for longer than you need to process it.

Distribution of Health Information

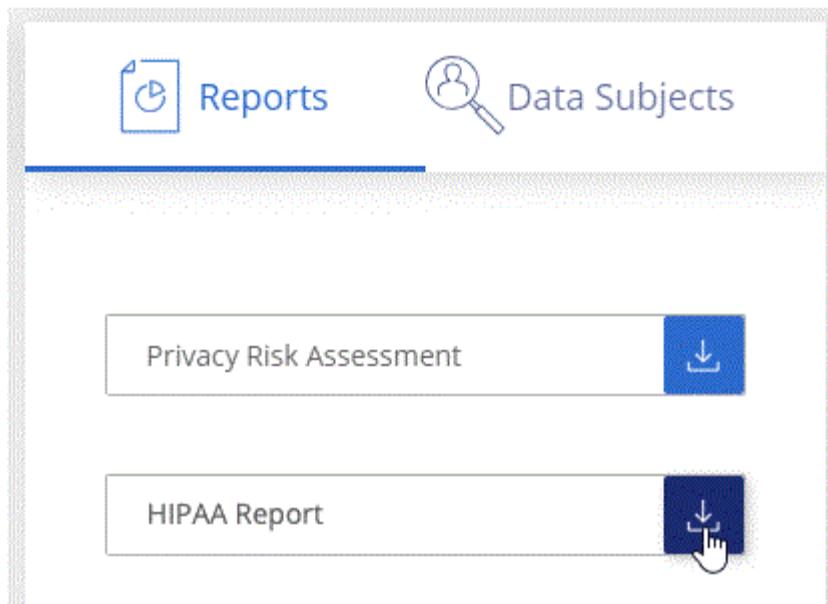
The working environments where the health information was found and whether encryption and ransomware protection are enabled.

Generating the HIPAA Report

Go to the Data Sense tab to generate the report.

Steps

1. At the top of Cloud Manager, click **Data Sense**.
2. Click **Compliance**, and then click the download icon next to **HIPAA Report** under **Reports**.



Result

Cloud Data Sense generates a PDF report that you can review and send to other groups as needed.

Data Mapping Report

The Data Mapping Report provides an overview of the data being stored in your corporate data sources to assist you with decisions of migration, back up, security, and compliance processes. The report first lists an overview report summarizing all your working environments and data sources, and then provides a breakdown for each working environment.

The report includes the following information:

Usage Capacity

For all working environments: Lists the number of files and the used capacity for each working environment.
For single working environments: Lists the files that are using the most capacity.

Age of Data

Provides three charts and graphs for when files were created, last modified, or last accessed. Lists the number of files, and their used capacity, based on certain date ranges.

Size of Data

Lists the number of files that exist within certain size ranges in your working environments.

File Types

Lists the total number of files and the used capacity for each type of file being stored in your working environments.

Generating the Data Mapping Report

Go to the Data Sense tab to generate the report.

Steps

1. At the top of Cloud Manager, click **Data Sense**.

2. Click **Governance**, and then click the **Full Data Mapping Overview Report** button from the Governance Dashboard.



Result

Cloud Data Sense generates a PDF report that you can review and send to other groups as needed.

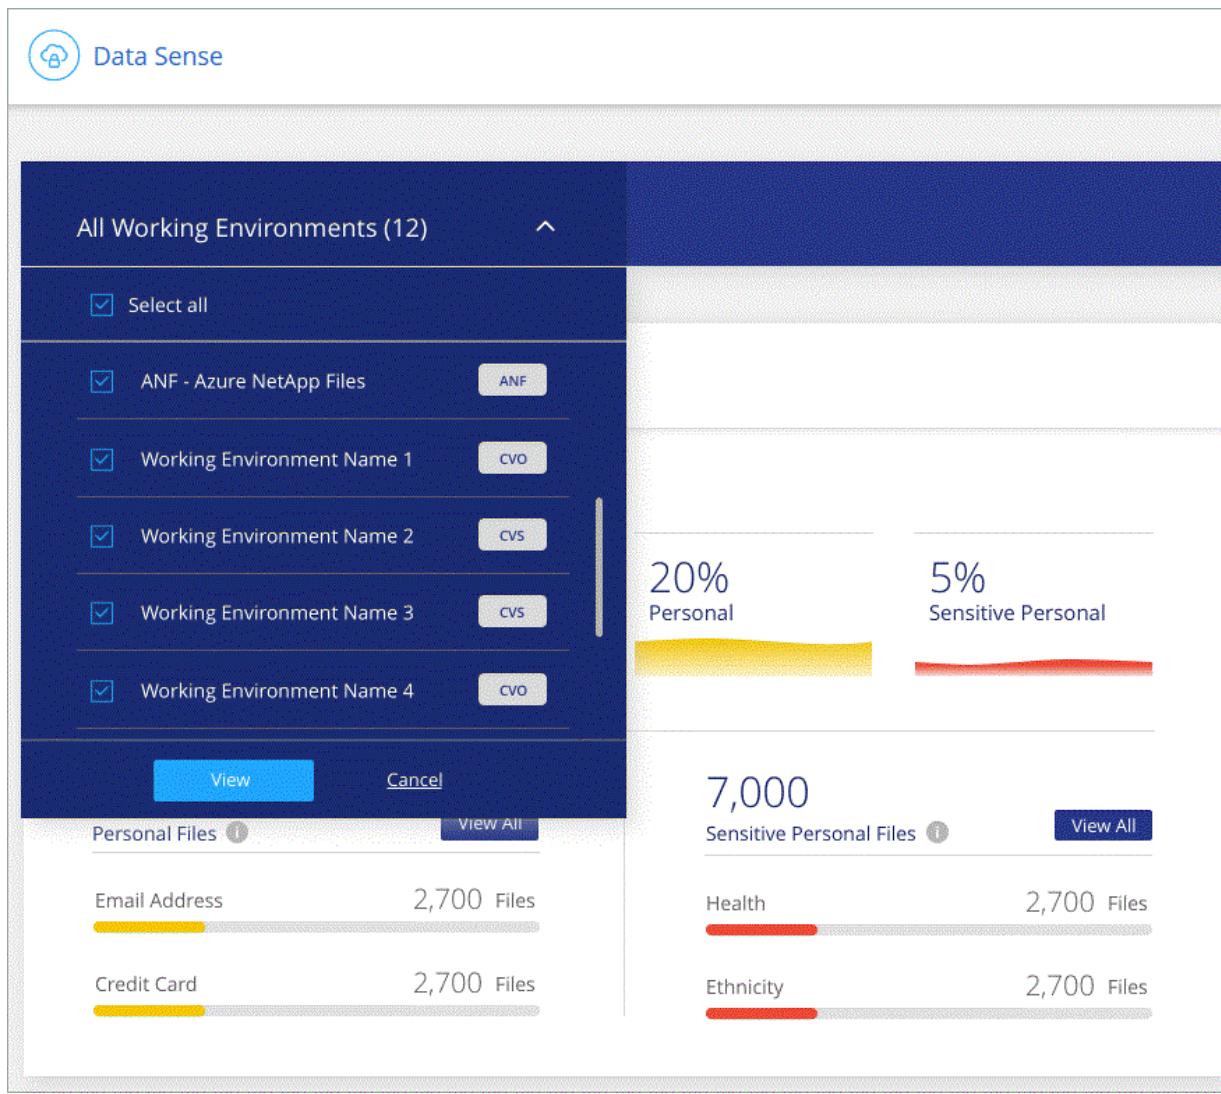
Selecting the working environments for reports

You can filter the contents of the Cloud Data Sense Compliance dashboard to see compliance data for all working environments and databases, or for just specific working environments.

When you filter the dashboard, Data Sense scopes the compliance data and reports to just those working environments that you selected.

Steps

1. Click the filter drop-down, select the working environments that you'd like to view data for, and click **View**.



Responding to a Data Subject Access Request

Respond to a Data Subject Access Request (DSAR) by searching for a subject's full name or known identifier (such as an email address) and then downloading a report. The report is designed to aid in your organization's requirement to comply with GDPR or similar data privacy laws.



The DSAR capabilities are available only if you have chosen to perform a full classification scan on your data sources. Data sources that have had a mapping-only scan do not provide file-level details.



NetApp can't guarantee 100% accuracy of the personal data and sensitive personal data that Cloud Data Sense identifies. You should always validate the information by reviewing the data.

What is a Data Subject Access Request?

Privacy regulations such as the European GDPR grant data subjects (such as customers or employees) the right to access their personal data. When a data subject requests this information, this is known as a DSAR.

(data subject access request). Organizations are required to respond to these requests "without undue delay", and at the latest within one month of receipt.

How can Cloud Data Sense help you respond to a DSAR?

When you perform a data subject search, Cloud Data Sense finds all of the files, buckets, and databases that have that person's name or identifier in it. Data Sense checks the latest pre-indexed data for the name or identifier. It doesn't initiate a new scan.

After the search is complete, you can then download the list of files for a Data Subject Access Request report. The report aggregates insights from the data and puts it into legal terms that you can send back to the person.

Searching for data subjects and downloading reports

Search for the data subject's full name or known identifier and then download a file list report or DSAR report. You can search by [any personal information type](#).

Only English is supported when searching for the names of data subjects. Support for more languages will be added later.



Data subject search is not supported within databases at this time.

Steps

1. At the top of Cloud Manager, click **Data Sense**.
2. Click **Data Subjects**.
3. Search for the data subject's full name or known identifier.

Here's an example that shows a search for the name *john doe*:

The screenshot shows the 'Data Subjects' section of the Cloud Data Sense interface. At the top, there are two navigation icons: 'Reports' (document icon) and 'Data Subjects' (person icon with a magnifying glass). Below the navigation is a back button labeled '< Back'. The main search results are displayed under the heading 'john doe Results'. It shows a summary: '203 Files Found'. Below this, there are two buttons: 'Download DSAR Report' with a download icon and 'Investigate Results' with a magnifying glass icon.

4. Choose one of the available options:

- **Download DSAR Report:** A formal response to the access request that you can send to the data subject. This report contains automatically-generated information based on data that Cloud Data Sense found on the data subject and is designed to be used as a template. You should complete the form and review it internally before sending it to the data subject.
- **Investigate Results:** A page that enables you to investigate the data by searching, sorting, expanding details for a specific file, and by downloading the file list.



If there are more than 10,000 results, only the top 10,000 appear in the file list.

Categories of private data

There are many types of private data that Cloud Data Sense can identify in your volumes, Amazon S3 buckets, databases, and OneDrive folders. See the categories below.



If you need Cloud Data Sense to identify other private data types, such as additional national ID numbers or healthcare identifiers, email ng-contact-data-sense@netapp.com with your request.

Types of personal data

The personal data found in files can be general personal data or national identifiers. The third column identifies whether Cloud Data Sense uses [proximity validation](#) to validate its findings for the identifier.

Note that you can add to the list of personal data that is found in your files if you are scanning a database server. The *Data Fusion* feature allows you to choose the additional identifiers that Cloud Data Sense will look for in its' scans by selecting columns in a database table. See [Adding personal data identifiers using Data Fusion](#) for details.

Type	Identifier	Proximity validation?
General	Email address	No
	Credit card number	No
	IBAN number (International Bank Account Number)	No
	IP address	No

Type	Identifier	Proximity validation?
National Identifiers	Austrian SSN	Yes
	Belgian ID (Numero National)	Yes
	Brazilian ID (CPF)	Yes
	Bulgarian ID (UCN)	Yes
	California Driver's License	Yes
	Croatian ID (OIB)	Yes
	Cyprus Tax Identification Number (TIC)	Yes
	Czech/Slovak ID	Yes
	Danish ID (CPR)	Yes
	Dutch ID (BSN)	Yes
	Estonian ID	Yes
	Finnish ID (HETU)	Yes
	French Tax Identification Number (SPI)	Yes
	German Tax Identification Number (Steuerliche Identifikationsnummer)	Yes
	Greek ID	Yes
	Hungarian Tax Identification Number	Yes
	Irish ID (PPS)	Yes
	Israeli ID	Yes
	Italian Tax Identification Number	Yes
	Latvian ID	Yes
	Lithuanian ID	Yes
	Luxembourg ID	Yes
	Maltese ID	Yes
	Polish ID (PESEL)	Yes
	Portuguese Tax Identification Number (NIF)	Yes
	Romanian ID (CNP)	Yes
	Slovenian ID (EMSO)	Yes
	South African ID	Yes
	Spanish Tax Identification Number	Yes
	Swedish ID	Yes
	U.K. ID (NINO)	Yes
	USA Social Security Number (SSN)	Yes

Types of sensitive personal data

The sensitive personal data that Cloud Data Sense can find in files includes the following:

Criminal Procedures Reference

Data concerning a natural person's criminal convictions and offenses.

Ethnicity Reference

Data concerning a natural person's racial or ethnic origin.

Health Reference

Data concerning a natural person's health.

ICD-9-CM Medical Codes

Codes used in the medical and health industry.

ICD-10-CM Medical Codes

Codes used in the medical and health industry.

Philosophical Beliefs Reference

Data concerning a natural person's philosophical beliefs.

Political Opinions Reference

Data concerning a natural person's political opinions.

Religious Beliefs Reference

Data concerning a natural person's religious beliefs.

Sex Life or Orientation Reference

Data concerning a natural person's sex life or sexual orientation.

Types of categories

Cloud Data Sense categorizes your data as follows:

Finance

- Balance Sheets
- Purchase Orders
- Invoices
- Quarterly Reports

HR

- Background Checks
- Compensation Plans
- Employee Contracts
- Employee Reviews
- Health
- Resumes

Legal

- NDAs
- Vendor-Customer contracts

Marketing

- Campaigns
- Conferences

Operations

- Audit Reports

Sales

- Sales Orders

Services

- RFI
- RFP
- SOW
- Training

Support

- Complaints and Tickets

Metadata categories

- Application Data
- Archive Files
- Audio
- Business Application Data
- CAD Files
- Code
- Database and index files
- Design Files
- Email Application Data
- Executables
- Financial Application Data
- Health Application Data
- Images
- Logs
- Miscellaneous Documents
- Miscellaneous Presentations
- Miscellaneous Spreadsheets
- Videos

Types of files

Cloud Data Sense scans all files for category and metadata insights and displays all file types in the file types section of the dashboard.

But when Data Sense detects Personal Identifiable Information (PII), or when it performs a DSAR search, only the following file formats are supported:

.CSV, .DCM, .DICOM, .DOC, .DOCX, .JSON, .PDF, .PPTX, .RTF, .TXT, .XLS, and .XLSX.

Accuracy of information found

NetApp can't guarantee 100% accuracy of the personal data and sensitive personal data that Cloud Data Sense identifies. You should always validate the information by reviewing the data.

Based on our testing, the table below shows the accuracy of the information that Data Sense finds. We break it down by *precision* and *recall*:

Precision

The probability that what Data Sense finds has been identified correctly. For example, a precision rate of 90% for personal data means that 9 out of 10 files identified as containing personal information, actually contain personal information. 1 out of 10 files would be a false positive.

Recall

The probability for Data Sense to find what it should. For example, a recall rate of 70% for personal data means that Data Sense can identify 7 out of 10 files that actually contain personal information in your organization. Data Sense would miss 30% of the data and it won't appear in the dashboard.

We are constantly improving the accuracy of our results. Those improvements will be automatically available in future Data Sense releases.

Type	Precision	Recall
Personal data - General	90%-95%	60%-80%
Personal data - Country identifiers	30%-60%	40%-60%
Sensitive personal data	80%-95%	20%-30%
Categories	90%-97%	60%-80%

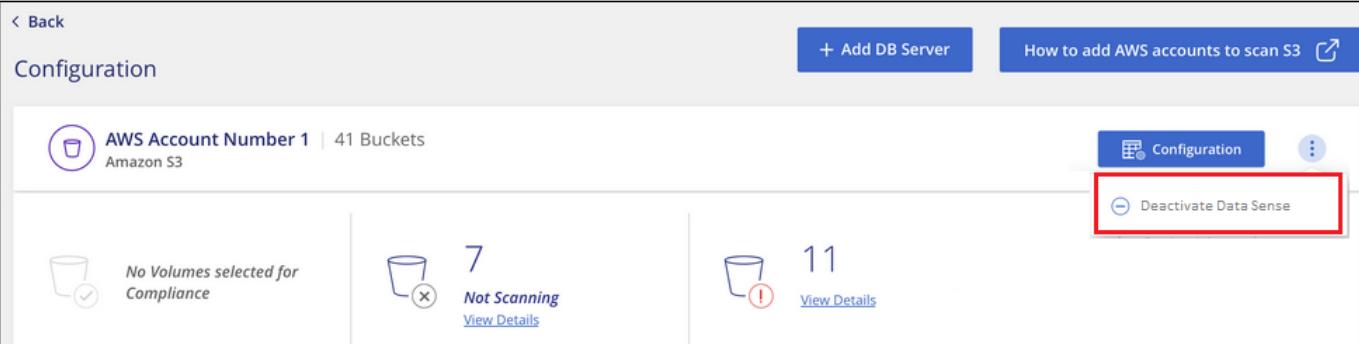
Removing data sources from Cloud Data Sense

If you need to, you can stop Cloud Data Sense from scanning one or more working environments, databases, file share groups, or OneDrive accounts. You can also delete the Cloud Data Sense instance if you no longer want to use Data Sense with your working environments.

Deactivating compliance scans for a working environment

When you deactivate scans, Cloud Data Sense no longer scans the data on the working environment and it removes the indexed compliance insights from the Data Sense instance (the data from the working environment itself isn't deleted).

- From the *Configuration* page, click the  button in the row for the working environment, and then click **Deactivate Data Sense**.



The screenshot shows the Cloud Data Sense Configuration page for an AWS account. The top navigation bar includes 'Back', '+ Add DB Server', and 'How to add AWS accounts to scan S3'. The main area displays statistics: 'AWS Account Number 1 | 41 Buckets' and 'Amazon S3'. Below this, there are three rows of status indicators: 'No Volumes selected for Compliance' (green), '7 Not Scanning' (yellow), and '11 View Details' (red). A modal window is overlaid on the 'Not Scanning' section, containing options like 'Deactivate Data Sense' and 'Remove DB Server', with 'Deactivate Data Sense' highlighted by a red box.

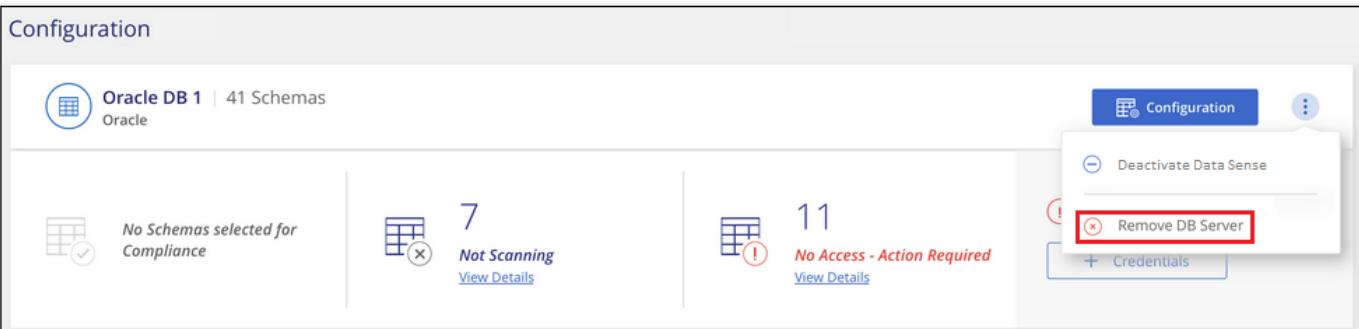


You can also disable compliance scans for a working environment from the Services panel when you select the working environment.

Removing a database from Cloud Data Sense

If you no longer want to scan a certain database, you can delete it from the Cloud Data Sense interface and stop all scans.

- From the *Configuration* page, click the  button in the row for the database, and then click **Remove DB Server**.



The screenshot shows the Cloud Data Sense Configuration page for an Oracle database. The top navigation bar includes 'Back', '+ Add DB Server', and 'How to add AWS accounts to scan S3'. The main area displays statistics: 'Oracle DB 1 | 41 Schemas' and 'Oracle'. Below this, there are three rows of status indicators: 'No Schemas selected for Compliance' (green), '7 Not Scanning' (yellow), and '11 No Access - Action Required' (red). A modal window is overlaid on the 'No Access' section, containing options like 'Deactivate Data Sense' and 'Remove DB Server', with 'Remove DB Server' highlighted by a red box.

Removing a OneDrive account from Cloud Data Sense

If you no longer want to scan user files from a certain OneDrive account, you can delete the account from the Cloud Data Sense interface and stop all scans.

Steps

- From the *Configuration* page, click the  button in the row for the OneDrive account, and then click **Remove OneDrive Account**.

Configuration

Add Data Source ▾



OneDrive Account 1 | 41 Users
OneDrive

Configuration



✖ Remove OneDrive Account

2. Click **Delete Account** from the confirmation dialog.

Removing a group of file shares from Cloud Data Sense

If you no longer want to scan user files from a file shares group, you can delete the File Shares Group from the Cloud Data Sense interface and stop all scans.

Steps

1. From the *Configuration* page, click the button in the row for the File Shares Group, and then click **Remove File Shares Group**.

Configuration

Add Data Source ▾



Shares Group 1 | 41 Shares
File Shares Group

Configuration



✖ Remove Shares Group

2. Click **Delete Group of Shares** from the confirmation dialog.

Reducing the Data Sense scan speed

Data scans have a negligible impact on your storage systems and on your data. However, if you are concerned with even a very small impact, you can configure Data Sense to perform "slow" scans. When enabled, slow scanning is used on all data sources - you can't configure slow scanning for a single working environment or data source.



The scan speed can't be reduced when scanning databases.

Steps

1. From the bottom of the *Configuration* page, move the slider to the right to activate slow scanning.

The top of the Configuration page indicates that slow scanning is enabled.

2. You can disable slow scanning by clicking **Disable** from this message.

Deleting the Cloud Data Sense instance

You can delete the Cloud Data Sense instance if you no longer want to use Data Sense. Deleting the instance also deletes the associated disks where the indexed data resides.

1. Go to your cloud provider's console and delete the Cloud Data Sense instance.

The instance is named *CloudCompliance* with a generated hash (UUID) concatenated to it. For example:
CloudCompliance-16bb6564-38ad-4080-9a92-36f5fd2f71c7

Frequently asked questions about Cloud Data Sense

This FAQ can help if you're just looking for a quick answer to a question.

What is Cloud Data Sense?

Cloud Data Sense (Cloud Compliance) is a cloud offering that uses Artificial Intelligence (AI) driven technology to help organizations understand data context and identify sensitive data across your storage systems. The systems can be Azure NetApp Files configurations, Amazon FSx for ONTAP, Cloud Volumes ONTAP systems hosted in AWS or Azure, Amazon S3 buckets, on-prem ONTAP systems, non-NetApp file shares, generic S3 object storage, databases, and OneDrive accounts.

Cloud Data Sense provides pre-defined parameters (such as sensitive information types and categories) to

address new data compliance regulations for data privacy and sensitivity, such as GDPR, CCPA, HIPAA, and more.

Why should I use Cloud Data Sense?

Cloud Data Sense can empower you with data to help you:

- Comply with data compliance and privacy regulations.
- Comply with data retention policies.
- Easily locate and report on specific data in response to data subjects, as required by GDPR, CCPA, HIPAA, and other data privacy regulations.

What are the common use cases for Cloud Data Sense?

- Identify Personal Identifiable Information (PII).
- Identify a wide scope of sensitive information as required by GDPR and CCPA privacy regulations.
- Comply with new and upcoming data privacy regulations.

[Learn more about the use cases for Cloud Data Sense.](#)

What types of data can be scanned with Cloud Data Sense?

Cloud Data Sense supports scanning of unstructured data over NFS and CIFS protocols that are managed by Cloud Volumes ONTAP, Azure NetApp Files, Amazon FSx for ONTAP, and on-prem ONTAP systems. Data Sense can also scan data stored on Amazon S3 buckets, in generic S3 object storage, and non-NetApp file shares.

Additionally, Data Sense can scan databases that are located anywhere, and user files from OneDrive accounts.

[Learn how scans work.](#)

Which cloud providers are supported?

Cloud Data Sense operates as part of Cloud Manager and currently supports AWS and Azure. This provides your organization with unified privacy visibility across different cloud providers.

How do I access Cloud Data Sense?

Cloud Data Sense is operated and managed through Cloud Manager. You can access Data Sense features from the **Data Sense** tab in Cloud Manager.

How does Cloud Data Sense work?

Cloud Data Sense deploys another layer of Artificial Intelligence alongside your Cloud Manager system and storage systems. It then scans the data on volumes, buckets, databases, and OneDrive accounts and indexes the data insights that are found.

[Learn more about how Cloud Data Sense works.](#)

How much does Cloud Data Sense cost?

The cost to use Cloud Data Sense depends on the amount of data that you're scanning. The first 1 TB of data that Data Sense scans in a Cloud Manager workspace is free. A subscription to the AWS or Azure Marketplace is required to continue scanning data after that point. See [pricing](#) for details.

What type of instance or VM is required for Cloud Data Sense?

- In AWS, Cloud Data Sense runs by default on an m5.4xlarge instance with a 500 GB GP2 disk.
- In Azure, Cloud Data Sense runs by default on a Standard_D16s_v3 VM with a 512 GB disk.

You can also download and install Data Sense software on a Linux host in your network or in the cloud. Everything works the same and you continue to manage your scan configuration and results through Cloud Manager. See [Deploying Cloud Data Sense on premises](#) for system requirements and installation details.

Note that you can deploy Data Sense on a system with fewer CPUs and less RAM, but there are limitations when using these systems. See [Using a smaller instance type](#) for details.



Cloud Data Sense is currently unable to scan S3 buckets and ANF files when it is installed on premises.

[Learn more about how Cloud Data Sense works.](#)

How often does Cloud Data Sense scan my data?

Data changes frequently, so Cloud Data Sense scans your data continuously with no impact to your data. While the initial scan of your data might take longer, subsequent scans only scan the incremental changes, which reduces system scan times.

[Learn how scans work.](#)

Data scans have a negligible impact on your storage systems and on your data. However, if you are concerned with even a very small impact, you can configure Data Sense to perform "slow" scans. [See how to reduce the scan speed.](#)

Does Cloud Data Sense offer reports?

Yes. The information offered by Cloud Data Sense can be relevant to other stakeholders in your organizations, so we enable you to generate reports to share the insights.

The following reports are available for Data Sense:

Privacy Risk Assessment report

Provides privacy insights from your data and a privacy risk score. [Learn more.](#)

Data Subject Access Request report

Enables you to extract a report of all files that contain information regarding a data subject's specific name or personal identifier. [Learn more.](#)

PCI DSS report

Helps you identify the distribution of credit card information across your files. [Learn more.](#)

HIPAA report

Helps you identify the distribution of health information across your files. [Learn more.](#)

Data Mapping report

Provides information about the size and number of files in your working environments. This includes usage capacity, age of data, size of data, and file types. [Learn more.](#)

Reports on a specific information type

Reports are available that include details about the identified files that contain personal data and sensitive personal data. You can also see files broken down by category and file type. [Learn more.](#)

Does scan performance vary?

Scan performance can vary based on the network bandwidth and the average file size in your cloud environment. It can also depend on the size characteristics of the host system (either in the cloud or on-premises). See [The Cloud Data Sense instance](#) and [Deploying Cloud Data Sense](#) for more information.

When initially adding new data sources you can also choose to only perform a "mapping" scan instead of a full "classification" scan. Mapping can be done on your data sources very quickly because it does not access files to see the data inside. [See the difference between a mapping and classification scan.](#)

Which file types are supported?

Cloud Data Sense scans all files for category and metadata insights and displays all file types in the file types section of the dashboard.

When Data Sense detects Personal Identifiable Information (PII), or when it performs a DSAR search, only the following file formats are supported:

.CSV, .DCM, .DICOM, .DOC, .DOCX, .JSON, .PDF, .PPTX, .RTF, .TXT, .XLS, and .XLSX.

How do I enable Cloud Data Sense?

First you need to deploy an instance of Cloud Data Sense in Cloud Manager. Once the instance is running, you can enable it on existing working environments and databases from the **Data Sense** tab or by selecting a specific working environment.

[Learn how to get started.](#)



Activating Cloud Data Sense results in an immediate initial scan. Scan results display shortly after.

How do I disable Cloud Data Sense?

You can disable Cloud Data Sense from scanning an individual working environment, database, file share group, or OneDrive account from the Data Sense Configuration page.

[Learn more.](#)



To completely remove the Cloud Data Sense instance, you can manually remove the Data Sense instance from your cloud provider's portal.

What happens if data tiering is enabled on Cloud Volumes ONTAP?

You might want to enable Cloud Data Sense on a Cloud Volumes ONTAP system that tiers cold data to object storage. If data tiering is enabled, Data Sense scans all of the data—data that's on disks and cold data tiered to object storage.

The compliance scan doesn't heat up the cold data—it stays cold and tiered to object storage.

Can I use Cloud Data Sense to scan on-premises ONTAP storage?

Yes. As long as you have discovered the on-prem ONTAP cluster as a working environment in Cloud Manager, you can scan any of the volume data.

Alternatively, you can run compliance scans on backup files created from your on-prem ONTAP volumes. So if you're already creating backup files from your on-prem systems using [Cloud Backup](#), you can run compliance scans on those backup files.

[Learn more.](#)

Can Cloud Data Sense send notifications to my organization?

Yes. In conjunction with the Policies feature, you can send email alerts to Cloud Manager users (daily, weekly, or monthly) when a Policy returns results so you can get notifications to protect your data. Learn more about [Policies](#).

You can also download status reports from the Investigation page in .CSV format that you can share internally in your organization.

Can I customize the service to my organization's needs?

Cloud Data Sense provides out-of-the-box insights to your data. These insights can be extracted and used for your organization's needs.

Additionally, you can use the **Data Fusion** capability to have Data Sense scan all your data based on criteria found in specific columns in databases you are scanning — essentially allowing you to make your own custom personal data types.

[Learn more.](#)

Can Cloud Data Sense work with the AIP labels I have embedded in my files?

Yes. You can manage AIP labels in the files that Cloud Data Sense is scanning if you have subscribed to [Azure Information Protection \(AIP\)](#). You can view the labels that are already assigned to files, add labels to files, and change existing labels.

[Learn more.](#)

Can I limit Cloud Data Sense information to specific users?

Yes, Cloud Data Sense is fully integrated with Cloud Manager. Cloud Manager users can only see information for the working environments they are eligible to view according to their workspace privileges.

Additionally, if you want to allow certain users to just view Data Sense scan results without having the ability to manage Data Sense settings, you can assign those users the *Cloud Compliance Viewer* role.

[Learn more.](#)

Tier on-premises data to the cloud

Learn about Cloud Tiering

NetApp's Cloud Tiering service extends your data center to the cloud by automatically tiering inactive data from on-premises ONTAP clusters to object storage. This frees valuable space on the cluster for more workloads, without making changes to the application layer. Cloud Tiering can reduce costs in your data center and enables you to switch from a CAPEX model to an OPEX model.

The Cloud Tiering service leverages the capabilities of *FabricPool*. FabricPool is a NetApp Data Fabric technology that enables automated tiering of data to low-cost object storage. Active (hot) data remains on the local tier (on-premises ONTAP aggregates), while inactive (cold) data is moved to the cloud tier—all while preserving ONTAP data efficiencies.

Note: Starting with ONTAP 9.8 you can tier data from aggregates consisting of HDDs in addition to high-performance SSDs. Prior to that release you can tier data only from AFF systems or FAS systems with all-SSD aggregates. See [Requirements for using FabricPool](#) for details.

Features

Cloud Tiering offers automation, monitoring, reports, and a common management interface:

- Automation makes it easier to set up and manage data tiering from on-prem ONTAP clusters to the cloud
- You can choose the default cloud provider storage class/access tier, or use lifecycle management to move older tiered data to a more cost-effective tier
- A single pane of glass removes the need to independently manage FabricPool across several clusters
- Reports show the amount of active and inactive data on each cluster
- A tiering health status helps you identify and correct issues as they occur
- If you have Cloud Volumes ONTAP systems, you'll find them in the Cluster Dashboard so you get a full view of data tiering in your hybrid cloud infrastructure

For more details about the value that Cloud Tiering provides, [check out the Cloud Tiering page on NetApp Cloud Central](#).



Cloud Volumes ONTAP systems are read-only from Cloud Tiering. You set up tiering for Cloud Volumes ONTAP from the working environment in Cloud Manager.

Supported object storage providers

You can tier inactive data from an on-premises ONTAP cluster to the following object storage providers:

- Amazon S3
- Microsoft Azure Blob
- Google Cloud Storage
- S3-compatible object storage
- NetApp StorageGRID

Object storage tiers

Each ONTAP cluster tiers inactive data to a single object store. When you set up data tiering, you have the choice to add a new bucket/container or to select an existing bucket/container, along with a storage class or access tier.

- [Learn about supported AWS S3 storage classes](#)
- [Learn about supported Azure Blob access tiers](#)
- [Learn about supported Google Cloud storage classes](#)

Cloud Tiering uses the cloud provider default storage class/access tier for your inactive data. However, you can apply a lifecycle rule so that the data automatically transitions from the default storage class to another storage class after a certain number of days. This can help keep your costs down by moving very cold data to less expensive storage.



You can't select lifecycle rules for data tiered to StorageGRID or S3-compatible storage.

Pricing and licenses

Pay for Cloud Tiering through a pay-as-you-go subscription, a bring-your-own Cloud Tiering license, or a combination of both. A 30-day free trial is available for your first cluster if you don't have a license.

There are no charges when tiering data to StorageGRID. Neither a BYOL license or PAYGO registration is required.

[View pricing details.](#)

30-day free trial

If you don't have a Cloud Tiering license, a 30-day free trial of Cloud Tiering starts when you set up tiering to your first cluster. After that 30-day free trial ends, you'll need to pay for Cloud Tiering through a pay-as-you-go subscription, a BYOL license, or a combination of both.

If your free trial ends and you haven't subscribed or added a license, then ONTAP no longer tiers cold data to object storage, but existing data is still available for access.

Pay-as-you-go subscription

Cloud Tiering offers consumption-based licensing in a pay-as-you-go model. After subscribing through your cloud provider's marketplace, you pay per GB for data that's tiered—there's no up-front payment. You are billed by your cloud provider through your monthly bill.

You should subscribe even if you have a free trial or if you bring your own license (BYOL):

- Subscribing ensures that there's no disruption of service after your free trial ends.

When the trial ends, you'll be charged hourly according to the amount of data that you tier.

- If you tier more data than allowed by your BYOL license, then data tiering continues through your pay-as-you-go subscription.

For example, if you have a 10 TB license, all capacity beyond the 10 TB is charged through the pay-as-you-go subscription.

You won't be charged from your pay-as-you-go subscription during your free trial or if you haven't exceeded your Cloud Tiering BYOL license.

[Learn how to set up a pay-as-you-go subscription.](#)

Bring your own license

Bring your own license by purchasing a **Cloud Tiering** license from NetApp. You can purchase 2-, 12-, 24-, or 36-month term licenses and specify any amount of tiering capacity. The BYOL Cloud Tiering license is a *floating* license that you can use across multiple on-premises ONTAP clusters. The total tiering capacity that you define in your Cloud Tiering license can be used by all of your on-prem clusters.

After you purchase a Cloud Tiering license, you'll need to use the Digital Wallet page in Cloud Manager to add the license. [See how to use a Cloud Tiering BYOL license.](#)

As noted above, we recommend that you set up a pay-as-you-go subscription, even if you have purchased a BYOL license.

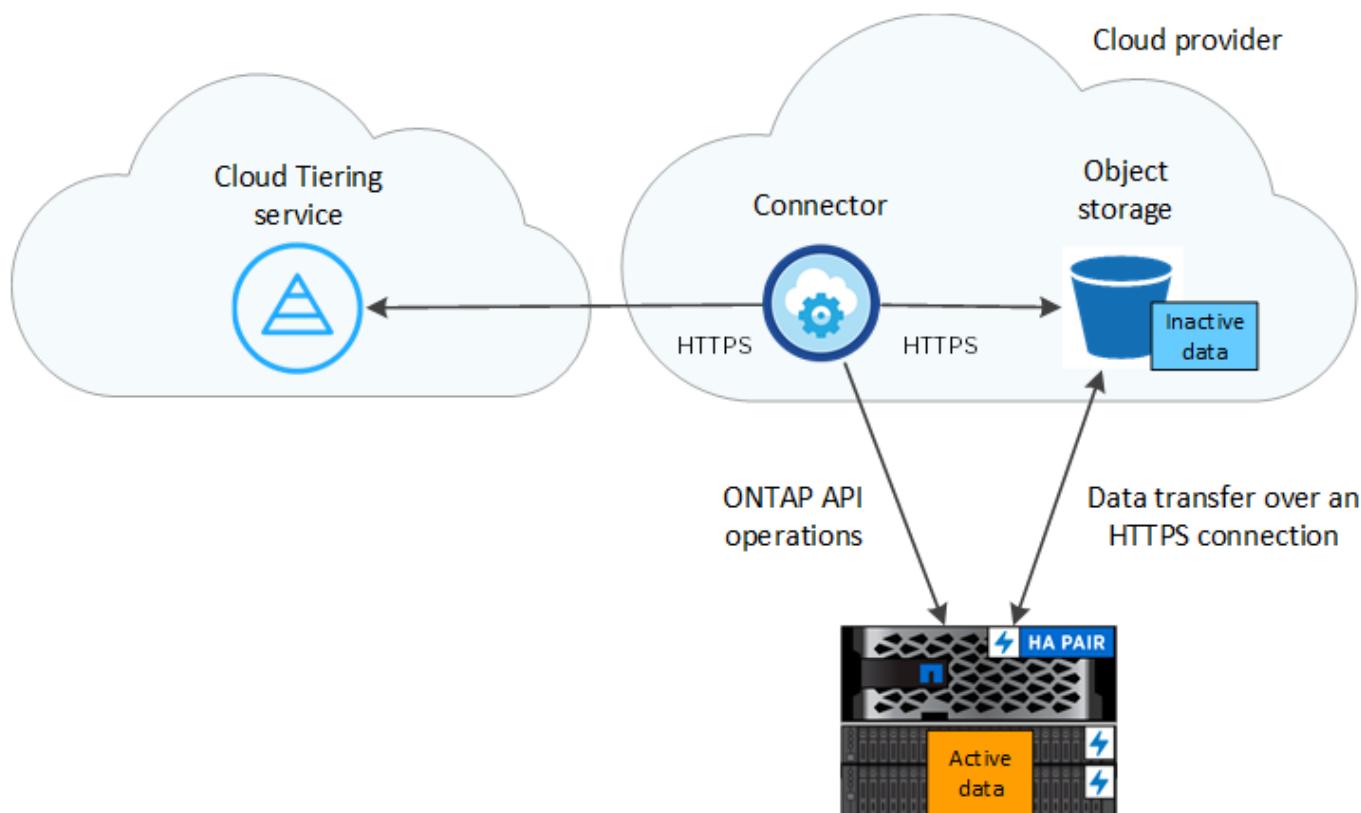


Starting August 21, 2021 the old **FabricPool** license is being replaced by the **Cloud Tiering** license. [Read more about how the Cloud Tiering license is different than the FabricPool license.](#)

How Cloud Tiering works

Cloud Tiering is a NetApp-managed service that uses FabricPool technology to automatically tier inactive (cold) data from your on-premises ONTAP clusters to object storage in your public cloud or private cloud. Connections to ONTAP take place from a Connector.

The following image shows the relationship between each component:



At a high level, Cloud Tiering works like this:

1. You discover your on-prem cluster from Cloud Manager.
2. You set up tiering by providing details about your object storage, including the bucket/container, a storage class or access tier, and lifecycle rules for the tiered data.
3. Cloud Manager configures ONTAP to use the object storage provider and discovers the amount of active and inactive data on the cluster.
4. You choose the volumes to tier and the tiering policy to apply to those volumes.
5. ONTAP starts tiering inactive data to the object store as soon as the data has reached the thresholds to be considered inactive (see [Volume tiering policies](#)).
6. If you have applied a lifecycle rule to the tiered data (only available for some providers), older tiered data is moved to a more cost-effective tier after a certain number of days.

Volume tiering policies

When you select the volumes that you want to tier, you choose a *volume tiering policy* to apply to each volume. A tiering policy determines when or whether the user data blocks of a volume are moved to the cloud.

You can also adjust the **cooling period**. This is the number of days that user data in a volume must remain inactive before it is considered "cold" and moved to object storage. For tiering policies that allow you to adjust the cooling period, the valid values are 2 to 183 days when using ONTAP 9.8 and later, and 2 to 63 days for earlier ONTAP versions; 2 to 63 is the recommended best practice.

No Policy (None)

Keeps the data on a volume in the performance tier, preventing it from being moved to the cloud tier.

Cold snapshots (Snapshot only)

ONTAP tiers cold Snapshot blocks in the volume that are not shared with the active file system to object storage. If read, cold data blocks on the cloud tier become hot and are moved to the performance tier.

Data is tiered only after an aggregate has reached 50% capacity and when the data has reached the cooling period. The default number of cooling days is 2, but you can adjust this number.



Re-heated data is written back to the performance tier only if there is space. If the performance tier capacity is more than 70% full, blocks continue to be accessed from the cloud tier.

Cold user data & snapshots (Auto)

ONTAP tiers all cold blocks in the volume (not including metadata) to object storage. The cold data includes not just Snapshot copies, but also cold user data from the active file system.

If read by random reads, cold data blocks on the cloud tier become hot and are moved to the performance tier. If read by sequential reads, such as those associated with index and antivirus scans, cold data blocks on the cloud tier stay cold and are not written to the performance tier. This policy is available starting with ONTAP 9.4.

Data is tiered only after an aggregate has reached 50% capacity and when the data has reached the cooling period. The default number of cooling days is 31, but you can adjust this number.



Re-heated data is written back to the performance tier only if there is space. If the performance tier capacity is more than 70% full, blocks continue to be accessed from the cloud tier.

All user data (All)

All data (not including metadata) is immediately marked as cold and tiered to object storage as soon as possible. There is no need to wait 48 hours for new blocks in a volume to become cold. Note that blocks located in the volume prior to the All policy being set require 48 hours to become cold.

If read, cold data blocks on the cloud tier stay cold and are not written back to the performance tier. This policy is available starting with ONTAP 9.6.

Take the following into consideration before you choose this tiering policy:

- Tiering data immediately reduces storage efficiencies (inline only).
- You should use this policy only if you are confident that cold data on the volume will not change.
- Object storage is not transactional and will result in significant fragmentation if subjected to change.
- Consider the impact of SnapMirror transfers before assigning the All tiering policy to source volumes in data protection relationships.

Because data is tiered immediately, SnapMirror will read data from the cloud tier rather than the performance tier. This will result in slower SnapMirror operations—possibly slowing other SnapMirror operations later in queue—even if they are using different tiering policies.

- Cloud Backup is similarly affected by volumes set with a tiering policy. [See tiering policy considerations with Cloud Backup](#).

All DP user data (Backup)

All data on a data protection volume (not including metadata) is immediately moved to the cloud tier. If read, cold data blocks on the cloud tier stay cold and are not written back to the performance tier (starting with ONTAP 9.4).



This policy is available for ONTAP 9.5 or earlier. It was replaced with the **All** tiering policy starting with ONTAP 9.6.

Get started

Tiering data from on-premises ONTAP clusters to Amazon S3

Free space on your on-prem ONTAP clusters by tiering inactive data to Amazon S3.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.



1 Prepare to tier data to Amazon S3

You need the following:

- An on-prem ONTAP cluster that's running ONTAP 9.2 or later and has an HTTPS connection to Amazon S3. [Learn how to discover a cluster](#).
- An AWS account that has an access key and [the required permissions](#) so the ONTAP cluster can tier inactive data in and out of S3.
- A Connector installed in an AWS VPC or on your premises.
- Networking for the Connector that enables an outbound HTTPS connection to the ONTAP cluster, to S3 storage, and to the Cloud Tiering service.

2

Set up tiering

In Cloud Manager, select an on-prem working environment, click **Enable** for the Tiering service, and follow the prompts to tier data to Amazon S3.

3

Set up licensing

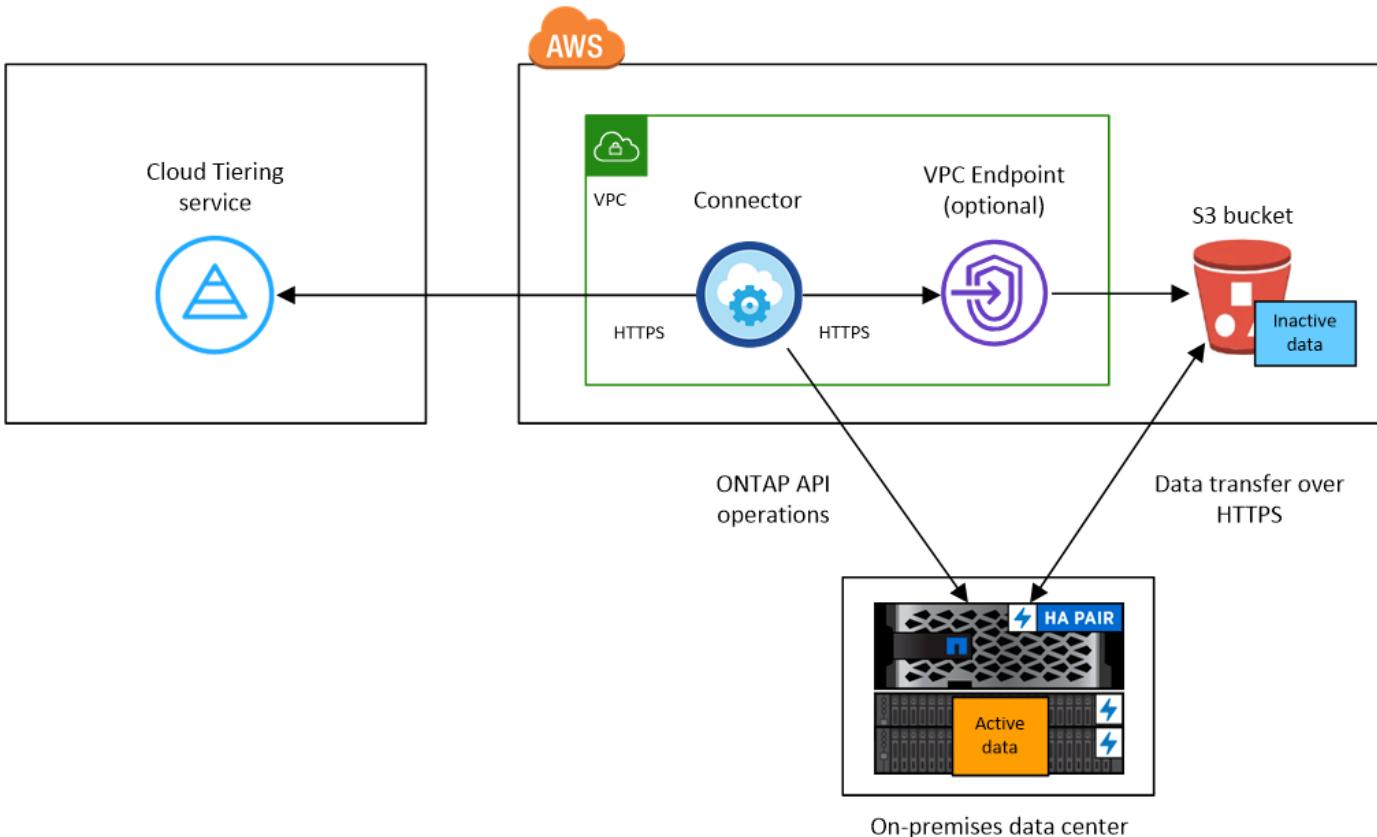
After your free trial ends, pay for Cloud Tiering through a pay-as-you-go subscription, an ONTAP Cloud Tiering BYOL license, or a combination of both:

- To subscribe from the AWS Marketplace, [go to the Cloud Manager Marketplace offering](#), click **Subscribe**, and then follow the prompts.
- To pay using a Cloud Tiering BYOL license, [contact us if you need to purchase one](#), and then [add it to your account from the Cloud Manager Digital Wallet](#).

Requirements

Verify support for your ONTAP cluster, set up your networking, and prepare your object storage.

The following image shows each component and the connections that you need to prepare between them:



Communication between a Connector and S3 is for object storage setup only. The Connector can reside on your premises, instead of in the cloud.

Preparing your ONTAP clusters

Your ONTAP clusters must meet the following requirements when tiering data to Amazon S3.

Supported ONTAP platforms

- When using ONTAP 9.8 and later: You can tier data from AFF systems, or FAS systems with all-SSD aggregates or all-HDD aggregates.
- When using ONTAP 9.7 and earlier: You can tier data from AFF systems, or FAS systems with all-SSD aggregates.

Supported ONTAP version

ONTAP 9.2 or later

Cluster networking requirements

- The ONTAP cluster initiates an HTTPS connection over port 443 to Amazon S3.

ONTAP reads and writes data to and from object storage. The object storage never initiates, it just responds.

Although AWS Direct Connect provides better performance and lower data transfer charges, it's not required between the ONTAP cluster and S3. But doing so is the recommended best practice.

- An inbound connection is required from the Connector, which can reside in an AWS VPC or on your premises.

A connection between the cluster and the Cloud Tiering service is not required.

- An intercluster LIF is required on each ONTAP node that hosts the volumes you want to tier. The LIF must be associated with the IPspace that ONTAP should use to connect to object storage.

When you set up data tiering, Cloud Tiering prompts you for the IPspace to use. You should choose the IPspace that each LIF is associated with. That might be the "Default" IPspace or a custom IPspace that you created. Learn more about [LIFs](#) and [IPspaces](#).

Supported volumes and aggregates

The total number of volumes that Cloud Tiering can tier might be less than the number of volumes on your ONTAP system. That's because volumes can't be tiered from some aggregates. Refer to ONTAP documentation for [functionality or features not supported by FabricPool](#).



Cloud Tiering supports FlexGroup volumes, starting with ONTAP 9.5. Setup works the same as any other volume.

Discovering an ONTAP cluster

You need to create an on-prem ONTAP working environment in Cloud Manager before you can start tiering cold data.

[Learn how to discover a cluster](#).

Creating or switching Connectors

A Connector is required to tier data to the cloud. When tiering data to AWS S3, you can use a Connector that's in an AWS VPC or on your premises. You'll either need to create a new Connector or make sure that the currently selected Connector resides in AWS or on-prem.

- [Learn about Connectors](#)
- [Creating a Connector in AWS](#)
- [Connector host requirements](#)
- [Installing the Connector on an existing Linux host](#)
- [Switching between Connectors](#)

Preparing networking for the Connector

Ensure that the Connector has the required networking connections. A Connector can be installed on-prem or in AWS.

Steps

1. Ensure that the network where the Connector is installed enables the following connections:
 - An outbound internet connection to the Cloud Tiering service over port 443 (HTTPS)
 - An HTTPS connection over port 443 to S3
 - An HTTPS connection over port 443 to your ONTAP clusters
2. If needed, enable a VPC Endpoint to S3.

A VPC Endpoint to S3 is recommended if you have a Direct Connect or VPN connection from your ONTAP cluster to the VPC and you want communication between the Connector and S3 to stay in your AWS

internal network.

Preparing Amazon S3

When you set up data tiering to a new cluster, you're prompted to create an S3 bucket or to select an existing S3 bucket in the AWS account where the Connector is set up. The AWS account must have permissions and an access key that you can enter in Cloud Tiering. The ONTAP cluster uses the access key to tier data in and out of S3.



If you are planning to configure Cloud Tiering to use a lower cost storage class where your tiered data will transition to after a certain number of days, you must not select any life cycle rules when setting up the bucket in your AWS account. Cloud Tiering manages the life cycle transitions.

Steps

1. Provide the following permissions to the IAM user:

```
"s3>ListAllMyBuckets",
"s3>ListBucket",
"s3:GetBucketLocation",
"s3:GetObject",
"s3:PutObject",
"s3>DeleteObject"
```

[AWS Documentation: Creating a Role to Delegate Permissions to an IAM User](#)

2. Create or locate an access key.

Cloud Tiering passes the access key on to the ONTAP cluster. The credentials are not stored in the Cloud Tiering service.

[AWS Documentation: Managing Access Keys for IAM Users](#)

Tiering inactive data from your first cluster to Amazon S3

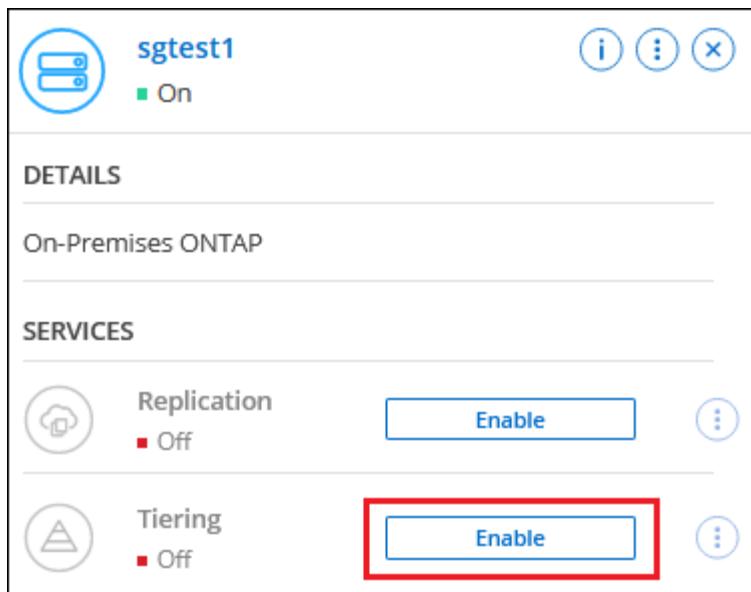
After you prepare your AWS environment, start tiering inactive data from your first cluster.

What you'll need

- An on-premises working environment.
- An AWS access key for an IAM user who has the required S3 permissions.

Steps

1. Select an on-prem cluster.
2. Click **Enable** for the Tiering service.



3. **Choose your provider:** This page appears only when using an on-prem Connector. Select **Amazon Web Services** and click **Continue**.

4. Complete the steps on the **Tiering Setup** page:

- S3 Bucket:** Add a new S3 bucket or select an existing S3 bucket that starts with the prefix *fabric-pool* and click **Continue**.

When using an on-prem Connector, you must enter the AWS Account ID that provides access to the existing S3 bucket or new S3 bucket that will be created.

The *fabric-pool* prefix is required because the IAM policy for the Connector enables the instance to perform S3 actions on buckets named with that exact prefix. For example, you could name the S3 bucket *fabric-pool-AFF1*, where AFF1 is the name of the cluster.

- Storage Class Life Cycle:** Cloud Tiering manages the life cycle transitions of your tiered data. Data starts in the *Standard* class, but you can create a rule to move the data to another class after a certain number of days.

Select the S3 storage class that you want to transition the tiered data to and the number of days before the data will be moved, and click **Continue**. For example, the screenshot below shows that tiered data is moved from the *Standard* class to the *Standard-IA* class after 45 days in object storage.

If you choose **Keep data in this storage class**, then the data remains in the *Standard* storage class and no rules are applied. [See supported storage classes](#).

Storage Class Life Cycle Management

We'll move the tiered data through the access tiers that you include in the life cycle. [Learn more about Amazon S3 storage classes.](#)

STORAGE CLASS SETUP

Standard

Move data from Standard after days

Keep data in this storage class

Standard-IA No Time Limit

Standard-IA

Intelligent-Tiering

One Zone-IA

Note that the life cycle rule is applied to all objects in the selected bucket.

- c. **Credentials:** Enter the access key ID and secret key for an IAM user who has the required S3 permissions, and click **Continue**.

The IAM user must be in the same AWS account as the bucket that you selected or created on the **S3 Bucket** page.

- d. **Cluster Network:** Select the IPspace that ONTAP should use to connect to object storage, and click **Continue**.

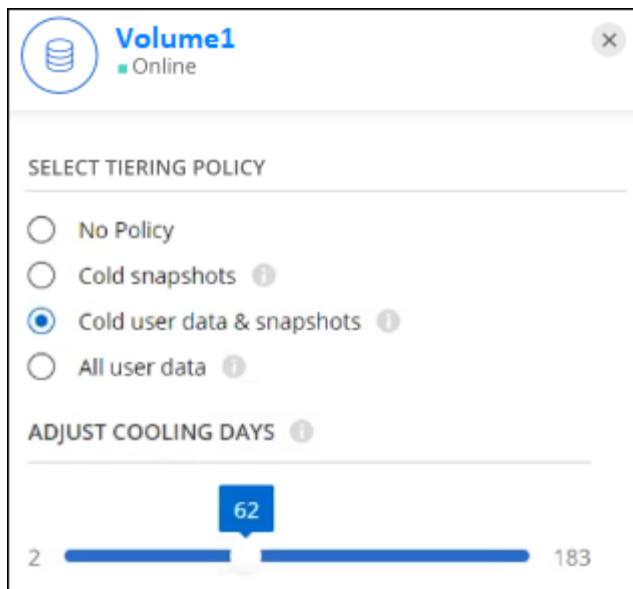
Selecting the correct IPspace ensures that Cloud Tiering can set up a connection from ONTAP to your cloud provider's object storage.

5. On the *Tier Volumes* page, select the volumes that you want to configure tiering for and launch the *Tiering Policy* page:

- To select all volumes, check the box in the title row (**Volume Name**) and click **Configure volumes**.
- To select multiple volumes, check the box for each volume (**Volume_1**) and click **Configure volumes**.
- To select a single volume, click the row (or icon) for the volume.

6. In the *Tiering Policy* dialog, select a tiering policy, optionally adjust the cooling days for the selected volumes, and click **Apply**.

[Learn more about volume tiering policies and cooling days.](#)



Result

You've successfully set up data tiering from volumes on the cluster to S3 object storage.

What's next?

[Be sure to subscribe to the Cloud Tiering service.](#)

You can also add additional clusters or review information about the active and inactive data on the cluster. For details, see [Managing data tiering from your clusters](#).

Tiering data from on-premises ONTAP clusters to Azure Blob storage

Free space on your on-prem ONTAP clusters by tiering inactive data to Azure Blob storage.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.

1

Prepare to tier data to Azure Blob storage

You need the following:

- An on-prem ONTAP cluster that's running ONTAP 9.4 or later and has an HTTPS connection to Azure Blob storage. [Learn how to discover a cluster](#).
- A Connector installed in an Azure VNet or on your premises.
- Networking for a Connector that enables an outbound HTTPS connection to the ONTAP cluster in your data center, to Azure storage, and to the Cloud Tiering service.

2

Set up tiering

In Cloud Manager, select an on-prem working environment, click **Enable** for the Tiering service, and follow the

prompts to tier data to Azure Blob storage.

3

Set up licensing

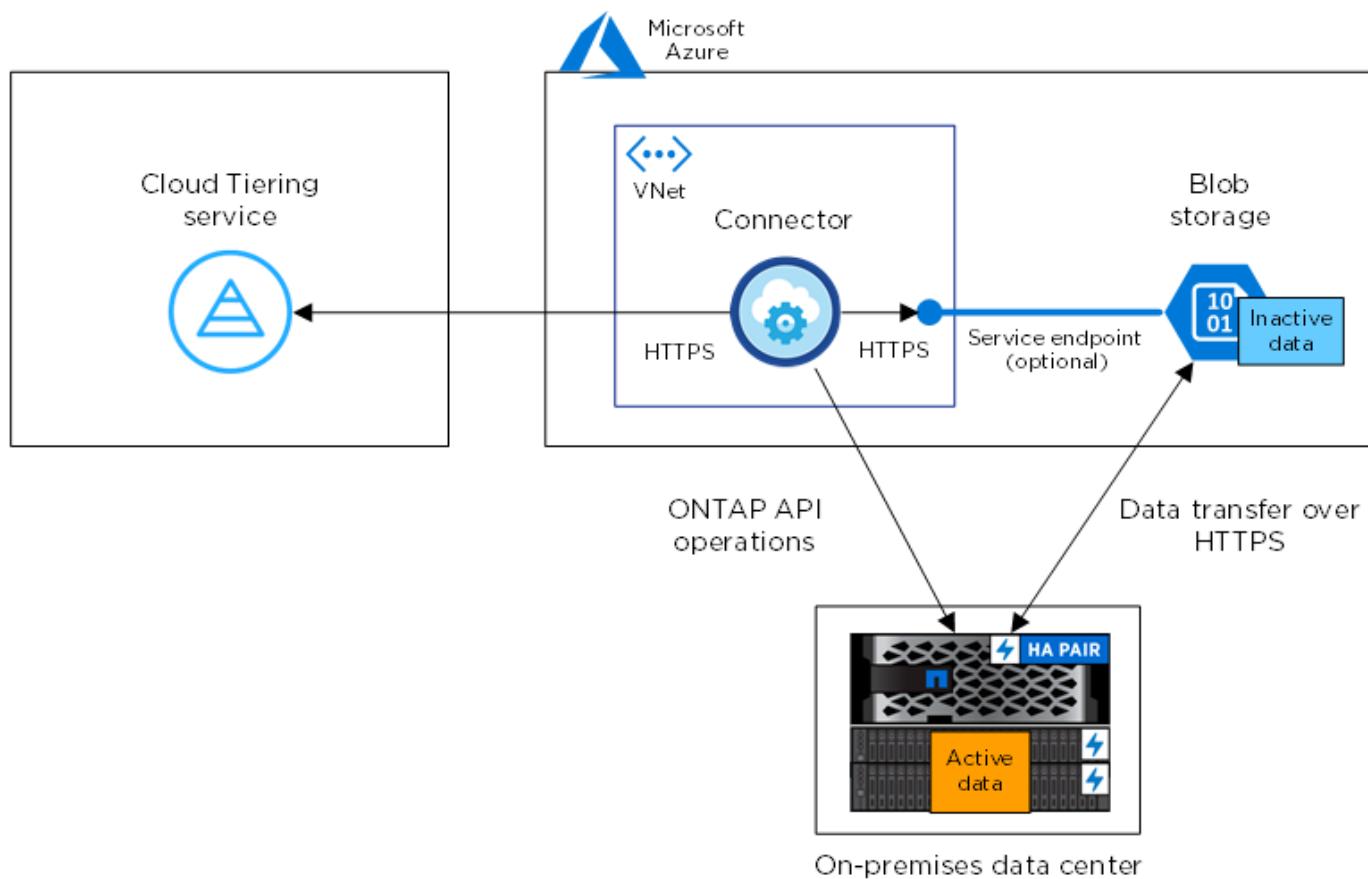
After your free trial ends, pay for Cloud Tiering through a pay-as-you-go subscription, an ONTAP Cloud Tiering BYOL license, or a combination of both:

- To subscribe from the Azure Marketplace, [go to the Cloud Manager Marketplace offering](#), click **Subscribe**, and then follow the prompts.
- To pay using a Cloud Tiering BYOL license, [contact us if you need to purchase one](#), and then [add it to your account from the Cloud Manager Digital Wallet](#).

Requirements

Verify support for your ONTAP cluster, set up your networking, and prepare your object storage.

The following image shows each component and the connections that you need to prepare between them:



Communication between the Connector and Blob storage is for object storage setup only. The Connector can reside on your premises, instead of in the cloud.

Preparing your ONTAP clusters

Your ONTAP clusters must meet the following requirements when tiering data to Azure Blob storage.

Supported ONTAP platforms

- When using ONTAP 9.8 and later: You can tier data from AFF systems, or FAS systems with all-SSD aggregates or all-HDD aggregates.
- When using ONTAP 9.7 and earlier: You can tier data from AFF systems, or FAS systems with all-SSD aggregates.

Supported ONTAP version

ONTAP 9.4 or later

Cluster networking requirements

- The ONTAP cluster initiates an HTTPS connection over port 443 to Azure Blob storage.
- ONTAP reads and writes data to and from object storage. The object storage never initiates, it just responds.
- Although ExpressRoute provides better performance and lower data transfer charges, it's not required between the ONTAP cluster and Azure Blob storage. But doing so is the recommended best practice.
- An inbound connection is required from the Connector, which can reside in an Azure VNet or on your premises.
- A connection between the cluster and the Cloud Tiering service is not required.
- An intercluster LIF is required on each ONTAP node that hosts the volumes you want to tier. The LIF must be associated with the *IPspace* that ONTAP should use to connect to object storage.

When you set up data tiering, Cloud Tiering prompts you for the *IPspace* to use. You should choose the *IPspace* that each LIF is associated with. That might be the "Default" *IPspace* or a custom *IPspace* that you created. Learn more about [LIFs](#) and [IPspaces](#).

Supported volumes and aggregates

The total number of volumes that Cloud Tiering can tier might be less than the number of volumes on your ONTAP system. That's because volumes can't be tiered from some aggregates. Refer to ONTAP documentation for [functionality or features not supported by FabricPool](#).



Cloud Tiering supports FlexGroup volumes, starting with ONTAP 9.5. Setup works the same as any other volume.

Discovering an ONTAP cluster

You need to create an on-prem ONTAP working environment in Cloud Manager before you can start tiering cold data.

[Learn how to discover a cluster.](#)

Creating or switching Connectors

A Connector is required to tier data to the cloud. When tiering data to Azure Blob storage, you can use a Connector that's in an Azure VNet or on your premises. You'll either need to create a new Connector or make sure that the currently selected Connector resides in Azure or on-prem.

- [Learn about Connectors](#)

- [Creating a Connector in Azure](#)
- [Connector host requirements](#)
- [Installing the Connector on an existing Linux host](#)
- [Switching between Connectors](#)

Verify that you have the necessary Connector permissions

If you created the Connector using Cloud Manager version 3.9.7 or greater, then you're all set.

If you created the Connector using an earlier version of Cloud Manager, then you'll need to edit the permission list to add 2 required permissions:

```
Microsoft.Storage/storageAccounts/managementPolicies/read  
Microsoft.Storage/storageAccounts/managementPolicies/write
```

Preparing networking for the Connector

Ensure that the Connector has the required networking connections. A Connector can be installed on-prem or in Azure.

Steps

1. Ensure that the network where the Connector is installed enables the following connections:
 - An outbound internet connection to the Cloud Tiering service over port 443 (HTTPS)
 - An HTTPS connection over port 443 to Azure Blob storage
 - An HTTPS connection over port 443 to your ONTAP clusters
2. If needed, enable a VNet service endpoint to Azure storage.

A VNet service endpoint to Azure storage is recommended if you have an ExpressRoute or VPN connection from your ONTAP cluster to the VNet and you want communication between the Connector and Blob storage to stay in your virtual private network.

Preparing Azure Blob storage

When you set up tiering, you need to identify the resource group you want to use, and the storage account and Azure container that belong to the resource group. A storage account enables Cloud Tiering to authenticate and access the Blob container used for data tiering.

Cloud Tiering supports only the General Purpose v2 and Premium Block Blob types of storage accounts.

 If you are planning to configure Cloud Tiering to use a lower cost access tier where your tiered data will transition to after a certain number of days, you must not select any life cycle rules when setting up the container in your Azure account. Cloud Tiering manages the life cycle transitions.

Tiering inactive data from your first cluster to Azure Blob storage

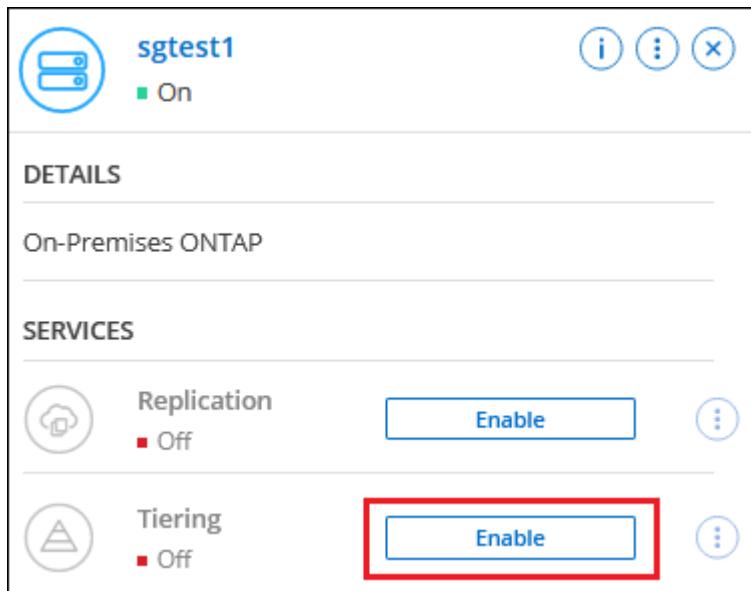
After you prepare your Azure environment, start tiering inactive data from your first cluster.

What you'll need

An on-premises working environment.

Steps

1. Select an on-prem cluster.
2. Click **Enable** for the Tiering service.



3. **Choose your provider:** This page appears only when using an on-prem Connector. Select **Microsoft Azure** and click **Continue**.
4. Complete the steps on the **Tiering Setup** page:
 - a. **Resource Group:** Select a resource group where an existing container is managed, or where you would like to create a new container for tiered data, and click **Continue**.
 - b. **Azure Container:** Add a new Blob container to a storage account, or select an existing container, and click **Continue**.

When using an on-prem Connector, you must enter the Azure Subscription that provides access to the existing container or new container that will be created.

The storage account and containers that appear in this step belong to the resource group that you selected in the previous step.

- c. **Access Tier Life Cycle:** Cloud Tiering manages the life cycle transitions of your tiered data. Data starts in the *Hot* class, but you can create a rule to move the data to the *Cool* class after a certain number of days.

Select the access tier that you want to transition the tiered data to and the number of days before the data will be moved, and click **Continue**. For example, the screenshot below shows that tiered data is moved from the *Hot* class to the *Cool* class after 45 days in object storage.

If you choose **Keep data in this access tier**, then the data remains in the *Hot* access tier and no rules are applied. [See supported access tiers](#).

Access Tier Life Cycle Management

We'll move the tiered data through the access tiers that you include in the life cycle. [Learn more about Azure Blob storage access tiers.](#)

ACCESS TIER SETUP

Hot

- Move data from Hot to Cool after days
 Keep data in this storage class



Cool

No Time Limit 

Note that the life cycle rule is applied to all blob containers in the selected storage account.

[Verify that you have the necessary Connector permissions](#) for the life cycle management feature.

- d. **Cluster Network:** Select the IPspace that ONTAP should use to connect to object storage, and click **Continue**.

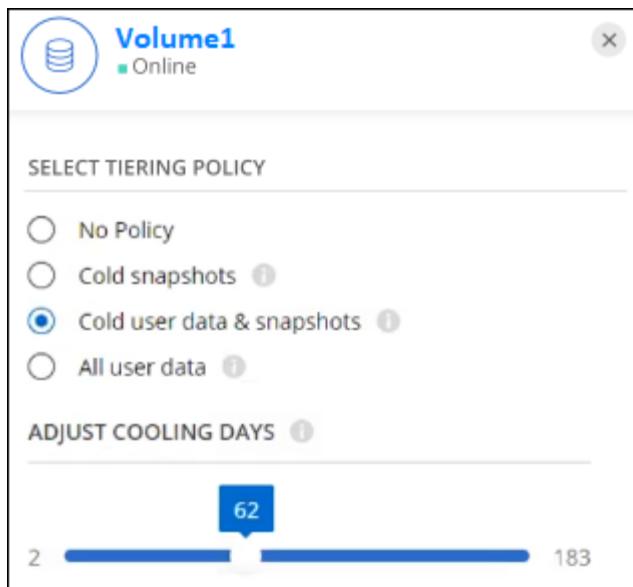
Selecting the correct IPspace ensures that Cloud Tiering can set up a connection from ONTAP to your cloud provider's object storage.

5. On the *Tier Volumes* page, select the volumes that you want to configure tiering for and launch the *Tiering Policy* page:

- To select all volumes, check the box in the title row ( **Volume Name**) and click **Configure volumes**.
- To select multiple volumes, check the box for each volume ( **Volume_1**) and click **Configure volumes**.
- To select a single volume, click the row (or  icon) for the volume.

6. In the *Tiering Policy* dialog, select a tiering policy, optionally adjust the cooling days for the selected volumes, and click **Apply**.

[Learn more about volume tiering policies and cooling days.](#)



Result

You've successfully set up data tiering from volumes on the cluster to Azure Blob object storage.

What's next?

[Be sure to subscribe to the Cloud Tiering service.](#)

You can also add additional clusters or review information about the active and inactive data on the cluster. For details, see [Managing data tiering from your clusters](#).

Tiering data from on-premises ONTAP clusters to Google Cloud Storage

Free space on your on-prem ONTAP clusters by tiering inactive data to Google Cloud Storage.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.

1

Prepare to tier data to Google Cloud Storage

You need the following:

- An on-prem ONTAP cluster that's running ONTAP 9.6 or later and has an HTTPS connection to Google Cloud Storage. [Learn how to discover a cluster](#).
- A service account that has the predefined Storage Admin role and storage access keys.
- A Connector installed in a Google Cloud Platform VPC.
- Networking for the Connector that enables an outbound HTTPS connection to the ONTAP cluster in your data center, to Google Cloud Storage, and to the Cloud Tiering service.

2

Set up tiering

In Cloud Manager, select an on-prem working environment, click **Enable** for the Tiering service, and follow the prompts to tier data to Google Cloud Storage.

3

Set up licensing

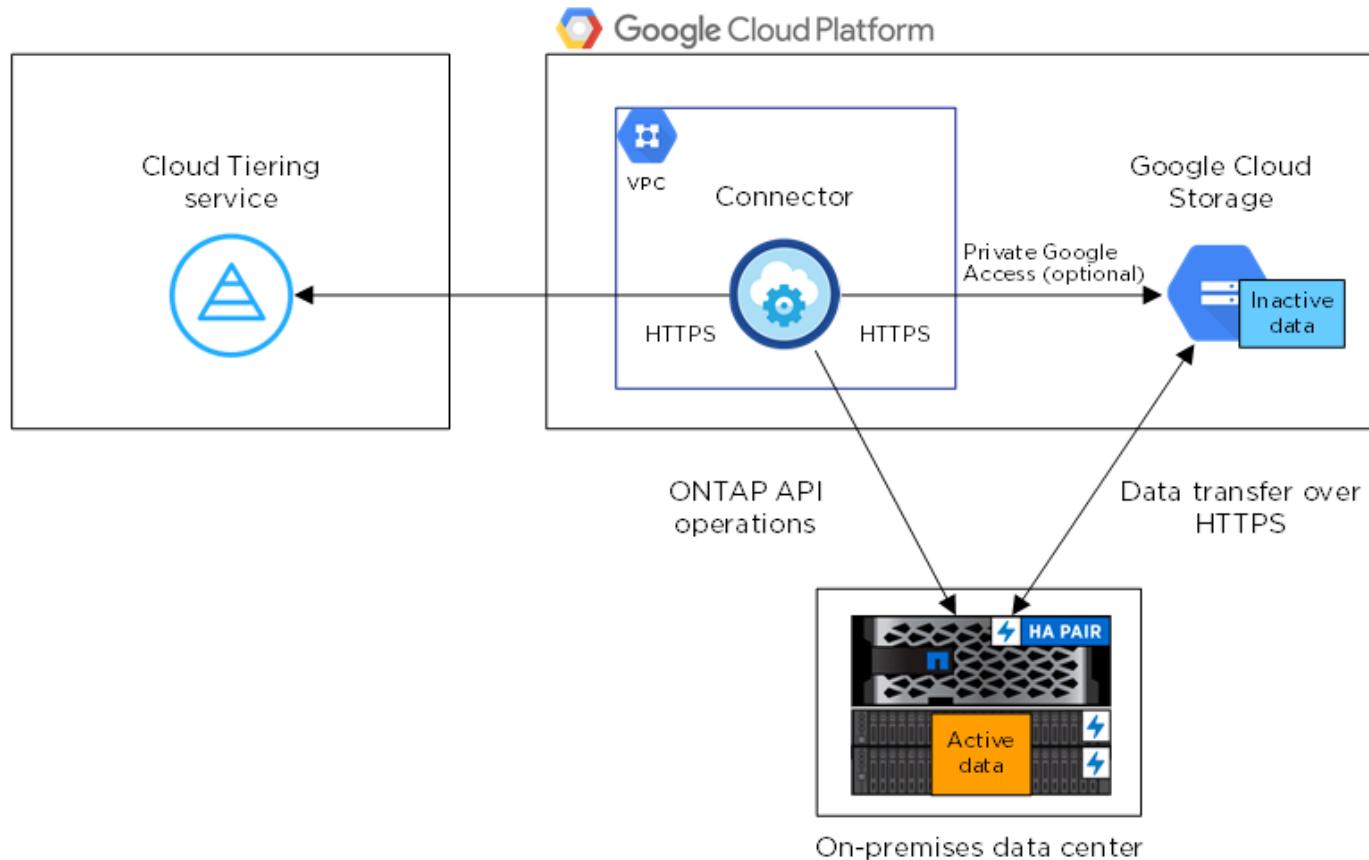
After your free trial ends, pay for Cloud Tiering through a pay-as-you-go subscription, an ONTAP Cloud Tiering BYOL license, or a combination of both:

- To subscribe from the GCP Marketplace, [go to the Cloud Manager Marketplace offering](#), click **Subscribe**, and then follow the prompts.
- To pay using a Cloud Tiering BYOL license, [contact us if you need to purchase one](#), and then [add it to your account from the Cloud Manager Digital Wallet](#).

Requirements

Verify support for your ONTAP cluster, set up your networking, and prepare your object storage.

The following image shows each component and the connections that you need to prepare between them:



Communication between the Connector and Google Cloud Storage is for object storage setup only.

Preparing your ONTAP clusters

Your ONTAP clusters must meet the following requirements when tiering data to Google Cloud Storage.

Supported ONTAP platforms

- When using ONTAP 9.8 and later: You can tier data from AFF systems, or FAS systems with all-SSD aggregates or all-HDD aggregates.
- When using ONTAP 9.7 and earlier: You can tier data from AFF systems, or FAS systems with all-SSD aggregates.

Supported ONTAP versions

ONTAP 9.6 or later

Cluster networking requirements

- The ONTAP cluster initiates an HTTPS connection over port 443 to Google Cloud Storage.

ONTAP reads and writes data to and from object storage. The object storage never initiates, it just responds.

Although a Google Cloud Interconnect provides better performance and lower data transfer charges, it's not required between the ONTAP cluster and Google Cloud Storage. But doing so is the recommended best practice.

- An inbound connection is required from the Connector, which resides in a Google Cloud Platform VPC.

A connection between the cluster and the Cloud Tiering service is not required.

- An intercluster LIF is required on each ONTAP node that hosts the volumes you want to tier. The LIF must be associated with the *IPspace* that ONTAP should use to connect to object storage.

When you set up data tiering, Cloud Tiering prompts you for the IPspace to use. You should choose the IPspace that each LIF is associated with. That might be the "Default" IPspace or a custom IPspace that you created. Learn more about [LIFs](#) and [IPspaces](#).

Supported volumes and aggregates

The total number of volumes that Cloud Tiering can tier might be less than the number of volumes on your ONTAP system. That's because volumes can't be tiered from some aggregates. Refer to ONTAP documentation for [functionality or features not supported by FabricPool](#).



Cloud Tiering supports FlexGroup volumes. Setup works the same as any other volume.

Discovering an ONTAP cluster

You need to create an on-prem ONTAP working environment in Cloud Manager before you can start tiering cold data.

[Learn how to discover a cluster.](#)

Creating or switching Connectors

A Connector is required to tier data to the cloud. When tiering data to Google Cloud Storage, a Connector must be available in a Google Cloud Platform VPC. You'll either need to create a new Connector or make sure that the currently selected Connector resides in GCP.

- [Learn about Connectors](#)
- [Creating a Connector in GCP](#)
- [Switching between Connectors](#)

Preparing networking for the Connector

Ensure that the Connector has the required networking connections.

Steps

1. Ensure that the VPC where the Connector is installed enables the following connections:
 - An outbound internet connection to the Cloud Tiering service over port 443 (HTTPS)
 - An HTTPS connection over port 443 to Google Cloud Storage
 - An HTTPS connection over port 443 to your ONTAP clusters
2. Optional: Enable Private Google Access on the subnet where you plan to deploy the Connector.

[Private Google Access](#) is recommended if you have a direct connection from your ONTAP cluster to the VPC and you want communication between the Connector and Google Cloud Storage to stay in your virtual private network. Note that Private Google Access works with VM instances that have only internal (private) IP addresses (no external IP addresses).

Preparing Google Cloud Storage

When you set up tiering, you need to provide storage access keys for a service account that has Storage Admin permissions. A service account enables Cloud Tiering to authenticate and access Cloud Storage buckets used for data tiering. The keys are required so that Google Cloud Storage knows who is making the request.

 If you are planning to configure Cloud Tiering to use lower cost storage classes where your tiered data will transition to after a certain number of days, you must not select any life cycle rules when setting up the bucket in your GCP account. Cloud Tiering manages the life cycle transitions.

Steps

1. [Create a service account that has the predefined Storage Admin role](#).
2. Go to [GCP Storage Settings](#) and create access keys for the service account:
 - a. Select a project, and click **Interoperability**. If you haven't already done so, click **Enable interoperability access**.
 - b. Under **Access keys for service accounts**, click **Create a key for a service account**, select the service account that you just created, and click **Create Key**.

You'll need to [enter the keys in Cloud Tiering](#) later when you set up tiering.

Tiering inactive data from your first cluster to Google Cloud Storage

After you prepare your Google Cloud environment, start tiering inactive data from your first cluster.

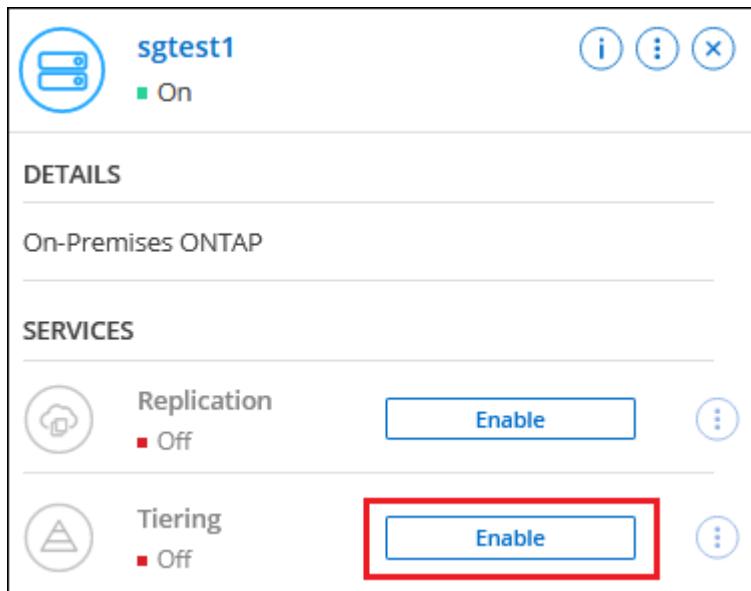
What you'll need

- [An on-premises working environment](#).

- Storage access keys for a service account that has the Storage Admin role.

Steps

1. Select an on-prem cluster.
2. Click **Enable** for the Tiering service.



3. Complete the steps on the **Tiering Setup** page:
 - a. **Bucket:** Add a new Google Cloud Storage bucket or select an existing bucket.
 - b. **Storage Class Life Cycle:** Cloud Tiering manages the life cycle transitions of your tiered data. Data starts in the *Standard* class, but you can create rules to move the data to other classes after a certain number of days.

Select the Google Cloud storage class that you want to transition the tiered data to and the number of days before the data will be moved, and click **Continue**. For example, the screenshot below shows that tiered data is moved from the *Standard* class to the *Nearline* class after 30 days in object storage, and then to the *Coldline* class after 60 days in object storage.

If you choose **Keep data in this storage class**, then the data remains in the that storage class. See [supported storage classes](#).

Storage Class Life Cycle Management

We'll move the tiered data through the storage classes that you include in the life cycle. [Learn more about Google Cloud Storage classes.](#)

STORAGE CLASS SETUP

Standard

- Move data from Standard to Nearline after days
 Keep data in this storage class



Nearline

- Move data from Nearline to Coldline after days
 Keep data in this storage class



Coldline

- Move data from Coldline to Archive after days
 Keep data in this storage class



Archive

No Time Limit

Note that the life cycle rule is applied to all objects in the selected bucket.

c. **Credentials:** Enter the storage access key and secret key for a service account that has the Storage Admin role.

d. **Cluster Network:** Select the IPspace that ONTAP should use to connect to object storage.

Selecting the correct IPspace ensures that Cloud Tiering can set up a connection from ONTAP to your cloud provider's object storage.

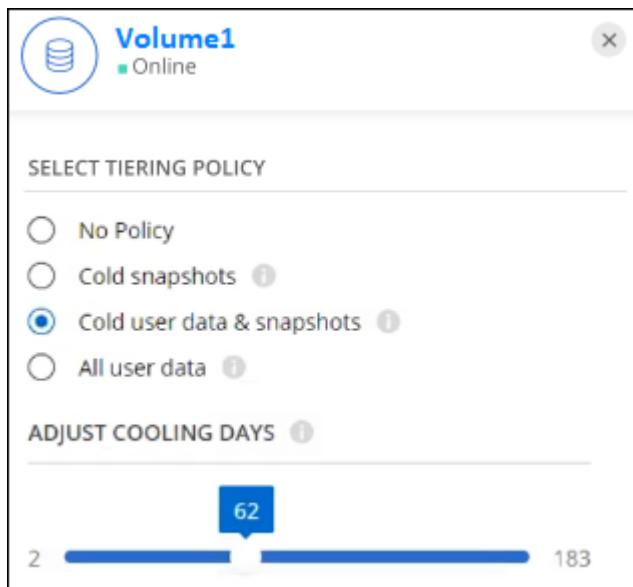
4. Click **Continue** to select the volumes that you want to tier.

5. On the *Tier Volumes* page, select the volumes that you want to configure tiering for and launch the Tiering Policy page:

- To select all volumes, check the box in the title row (**Volume Name**) and click **Configure volumes**.
- To select multiple volumes, check the box for each volume (**Volume_1**) and click **Configure volumes**.
- To select a single volume, click the row (or  icon) for the volume.

6. In the *Tiering Policy* dialog, select a tiering policy, optionally adjust the cooling days for the selected volumes, and click **Apply**.

Learn more about volume tiering policies and cooling days.



Result

You've successfully set up data tiering from volumes on the cluster to Google Cloud object storage.

What's next?

[Be sure to subscribe to the Cloud Tiering service.](#)

You can also add additional clusters or review information about the active and inactive data on the cluster. For details, see [Managing data tiering from your clusters](#).

Tiering data from on-premises ONTAP clusters to StorageGRID

Free space on your on-prem ONTAP clusters by tiering inactive data to StorageGRID.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.



Prepare to tier data to StorageGRID

You need the following:

- An on-prem ONTAP cluster that's running ONTAP 9.4 or later, and a connection over a user-specified port to StorageGRID. [Learn how to discover a cluster](#).
- StorageGRID 10.3 or later with AWS access keys that have S3 permissions.
- A Connector installed on your premises.
- Networking for the Connector that enables an outbound HTTPS connection to the ONTAP cluster, to StorageGRID, and to the Cloud Tiering service.

2

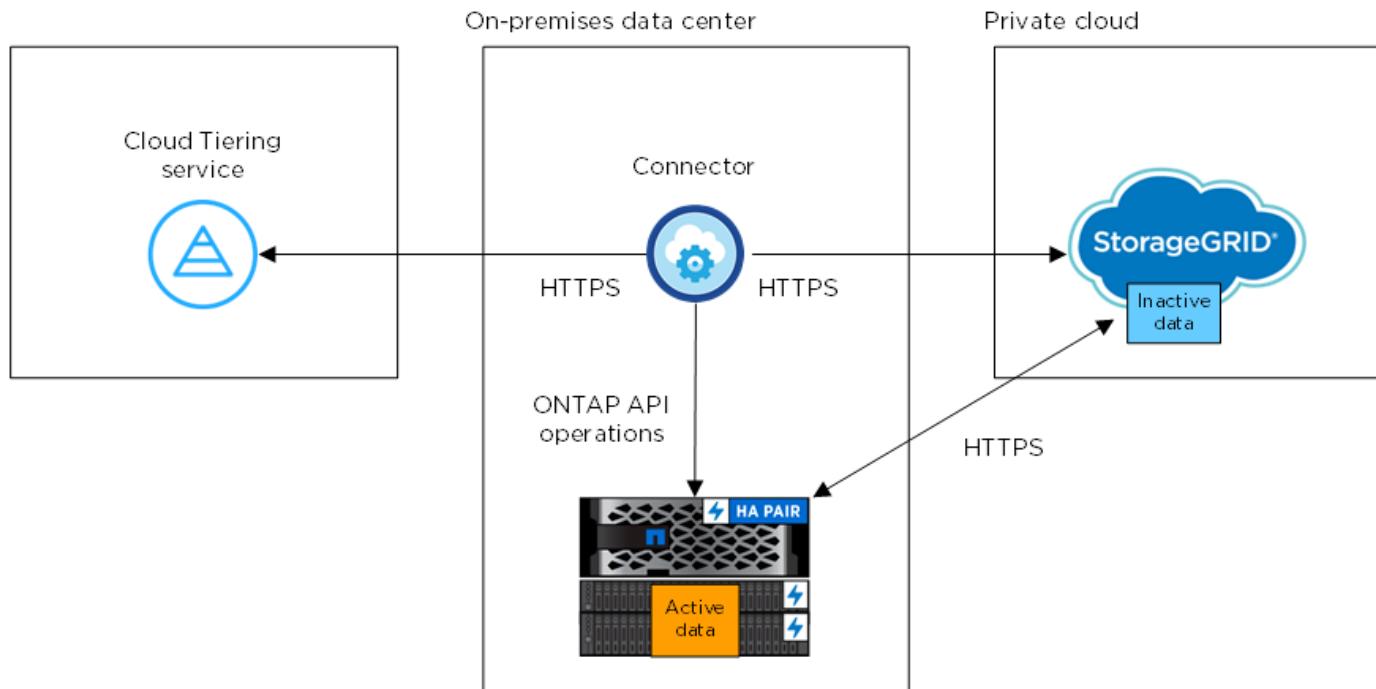
Set up tiering

In Cloud Manager, select an on-prem working environment, click **Enable** for the Tiering service, and follow the prompts to tier data to StorageGRID.

Requirements

Verify support for your ONTAP cluster, set up your networking, and prepare your object storage.

The following image shows each component and the connections that you need to prepare between them:



Communication between the Connector and StorageGRID is for object storage setup only.

Preparing your ONTAP clusters

Your ONTAP clusters must meet the following requirements when tiering data to StorageGRID.

Supported ONTAP platforms

- When using ONTAP 9.8 and later: You can tier data from AFF systems, or FAS systems with all-SSD aggregates or all-HDD aggregates.
- When using ONTAP 9.7 and earlier: You can tier data from AFF systems, or FAS systems with all-SSD aggregates.

Supported ONTAP version

ONTAP 9.4 or later

Licensing

A Cloud Tiering license isn't required in your Cloud Manager account, nor is a FabricPool license required on the ONTAP cluster, when tiering data to StorageGRID.

Cluster networking requirements

- The ONTAP cluster initiates an HTTPS connection over a user-specified port to StorageGRID (the port is configurable during tiering setup).

ONTAP reads and writes data to and from object storage. The object storage never initiates, it just responds.

- An inbound connection is required from the Connector, which must reside on your premises.

A connection between the cluster and the Cloud Tiering service is not required.

- An intercluster LIF is required on each ONTAP node that hosts the volumes you want to tier. The LIF must be associated with the *IPspace* that ONTAP should use to connect to object storage.

When you set up data tiering, Cloud Tiering prompts you for the IPspace to use. You should choose the IPspace that each LIF is associated with. That might be the "Default" IPspace or a custom IPspace that you created. Learn more about [LIFs](#) and [IPspaces](#).

Supported volumes and aggregates

The total number of volumes that Cloud Tiering can tier might be less than the number of volumes on your ONTAP system. That's because volumes can't be tiered from some aggregates. Refer to ONTAP documentation for [functionality or features not supported by FabricPool](#).



Cloud Tiering supports FlexGroup volumes, starting with ONTAP 9.5. Setup works the same as any other volume.

Discovering an ONTAP cluster

You need to create an on-prem ONTAP working environment in the Cloud Manager Canvas before you can start tiering cold data.

[Learn how to discover a cluster.](#)

Preparing StorageGRID

StorageGRID must meet the following requirements.

Supported StorageGRID versions

StorageGRID 10.3 and later is supported.

S3 credentials

When you set up tiering to StorageGRID, you need to provide Cloud Tiering with an S3 access key and secret key. Cloud Tiering uses the keys to access your buckets.

These access keys must be associated with a user who has the following permissions:

```
"s3>ListAllMyBuckets",
"s3>ListBucket",
"s3GetObject",
"s3PutObject",
"s3DeleteObject",
"s3CreateBucket"
```

Object versioning

You must not enable StorageGRID object versioning on the object store bucket.

Creating or switching Connectors

A Connector is required to tier data to the cloud. When tiering data to StorageGRID, a Connector must be available on your premises. You'll either need to install a new Connector or make sure that the currently selected Connector resides on-prem.

- [Learn about Connectors](#)
- [Connector host requirements](#)
- [Installing the Connector on an existing Linux host](#)
- [Switching between Connectors](#)

Preparing networking for the Connector

Ensure that the Connector has the required networking connections.

Steps

1. Ensure that the network where the Connector is installed enables the following connections:
 - An outbound internet connection to the Cloud Tiering service over port 443 (HTTPS)
 - An HTTPS connection over port 443 to StorageGRID
 - An HTTPS connection over port 443 to your ONTAP cluster

Tiering inactive data from your first cluster to StorageGRID

After you prepare your environment, start tiering inactive data from your first cluster.

What you'll need

- [An on-premises working environment.](#)
- The FQDN of the StorageGRID server and the port that will be used for HTTPS communications.
- An AWS access key that has the required S3 permissions.

Steps

1. Select an on-prem cluster.
2. Click **Enable** for the Tiering service.

The screenshot shows the StorageGRID configuration interface for a provider named 'sgtest1'. The 'DETAILS' section indicates 'On-Premises ONTAP'. The 'SERVICES' section contains two items: 'Replication' (status: Off, button: 'Enable') and 'Tiering' (status: Off, button: 'Enable'). The 'Tiering' button is highlighted with a red box.

3. **Choose your provider:** Select **StorageGRID** and click **Continue**.

4. Complete the steps on the **Tiering Setup** page:

- Server:** Enter the FQDN of the StorageGRID server, the port that ONTAP should use for HTTPS communication with StorageGRID, and the access key and secret key for an account that has the required S3 permissions.
- Bucket:** Add a new bucket or select an existing bucket that starts with the prefix *fabric-pool* and click **Continue**.

The *fabric-pool* prefix is required because the IAM policy for the Connector enables the instance to perform S3 actions on buckets named with that exact prefix. For example, you could name the S3 bucket *fabric-pool-AFF1*, where AFF1 is the name of the cluster.

- Cluster Network:** Select the IPspace that ONTAP should use to connect to object storage and click **Continue**.

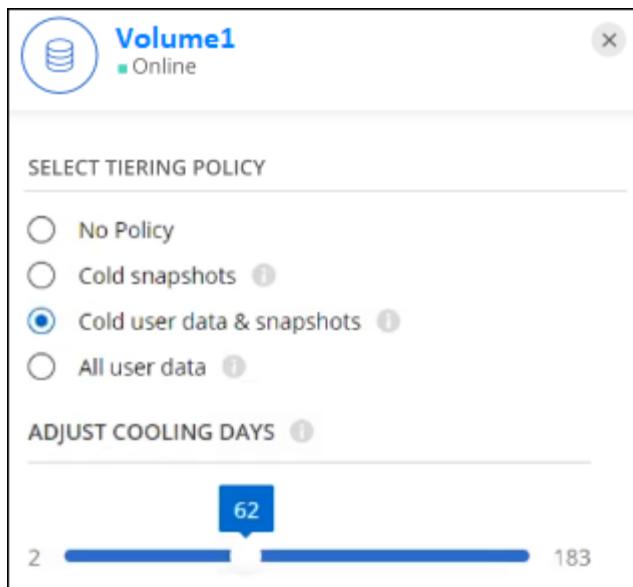
Selecting the correct IPspace ensures that Cloud Tiering can set up a connection from ONTAP to StorageGRID object storage.

5. On the *Tier Volumes* page, select the volumes that you want to configure tiering for and launch the *Tiering Policy* page:

- To select all volumes, check the box in the title row (**Volume Name**) and click **Configure volumes**.
- To select multiple volumes, check the box for each volume (**Volume_1**) and click **Configure volumes**.
- To select a single volume, click the row (or icon) for the volume.

6. In the *Tiering Policy* dialog, select a tiering policy, optionally adjust the cooling days for the selected volumes, and click **Apply**.

[Learn more about volume tiering policies and cooling days.](#)



Result

You've successfully set up data tiering from volumes on the cluster to StorageGRID.

What's next?

You can add additional clusters or review information about the active and inactive data on the cluster. For details, see [Managing data tiering from your clusters](#).

Tiering data from on-premises ONTAP clusters to S3 object storage

Free space on your on-prem ONTAP clusters by tiering inactive data to any object storage service which uses the Simple Storage Service (S3) protocol.

Customers who want to use object stores that are not officially supported as a cloud tier can do so using these instructions. Customers must test and confirm that the object store meets their requirements.



NetApp does not support nor is liable for any issues arising from any third-party Object Store Service, specifically where it does not have agreed support arrangements with the third party with whom the product originated. It is acknowledged and agreed that NetApp shall not be liable for any associated damage or otherwise be required to provide support on that third-party product.

Quick start

Get started quickly by following these steps, or scroll down to the remaining sections for full details.



1 Prepare to tier data to S3-compatible object storage

You need the following:

- An on-prem ONTAP cluster that's running ONTAP 9.8 or later, and a connection over a user-specified port to the S3-compatible object storage. [Learn how to discover a cluster](#).

- The FQDN, Access Key, and Secret Key for the object storage server so that the ONTAP cluster can access the bucket.
- A Connector installed on your premises.
- Networking for the Connector that enables an outbound HTTPS connection to the ONTAP cluster, to the S3-compatible object storage, and to the Cloud Tiering service.

2

Set up tiering

In Cloud Manager, select an on-prem working environment, click **Enable** for the Tiering service, and follow the prompts to tier data to S3-compatible object storage.

3

Set up licensing

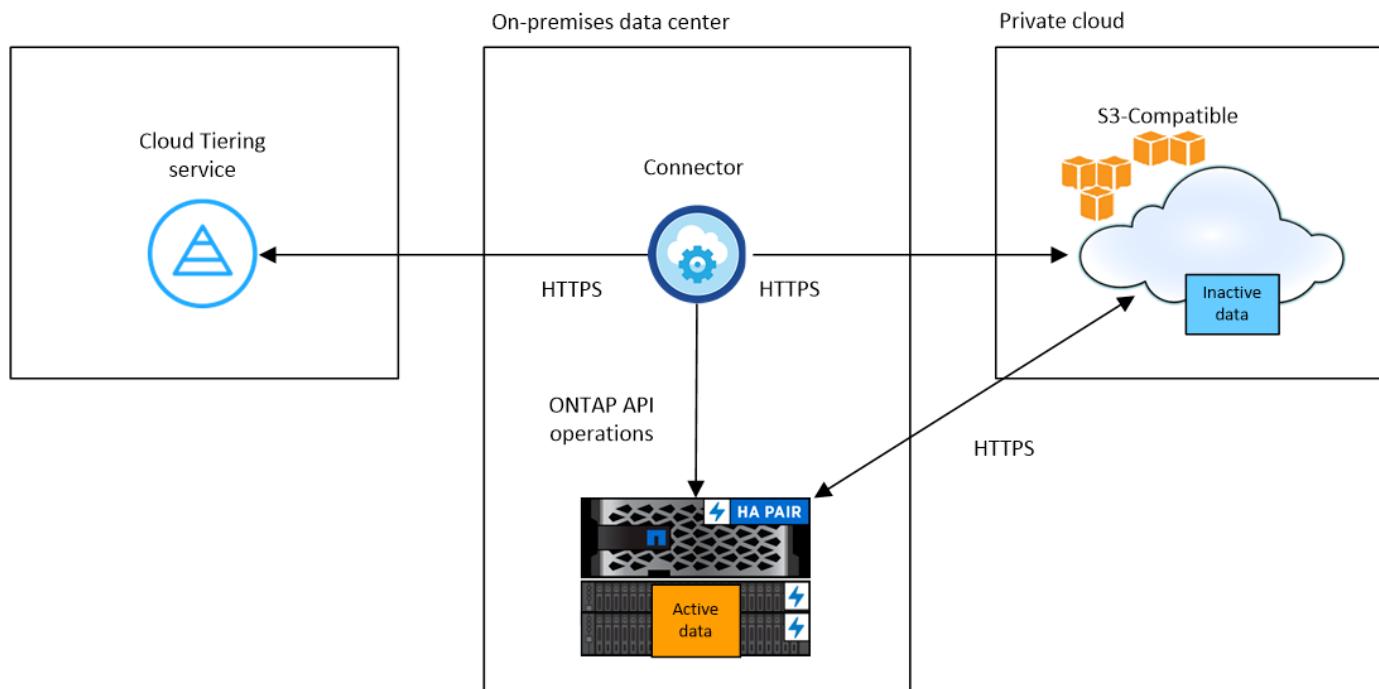
Pay for Cloud Tiering through a pay-as-you-go subscription from your cloud provider, a NetApp Cloud Tiering bring-your-own-license, or a combination of both:

- To subscribe to the Cloud Manager PAYGO offering from the [AWS Marketplace](#), [Azure Marketplace](#), or [GCP Marketplace](#), click **Subscribe** and follow the prompts.
- To pay using a Cloud Tiering BYOL license, [contact us if you need to purchase one](#), and then [add it to your account from the Cloud Manager Digital Wallet](#).

Requirements

Verify support for your ONTAP cluster, set up your networking, and prepare your object storage.

The following image shows each component and the connections that you need to prepare between them:





Communication between the Connector and the S3-compatible object storage server is for object storage setup only.

Preparing your ONTAP clusters

Your ONTAP clusters must meet the following requirements when tiering data to S3-compatible object storage.

Supported ONTAP platforms

You can tier data from AFF systems, or FAS systems with all-SSD aggregates or all-HDD aggregates.

Supported ONTAP version

ONTAP 9.8 or later

Cluster networking requirements

- The ONTAP cluster initiates an HTTPS connection over a user-specified port to S3-compatible object storage (the port is configurable during tiering setup).

ONTAP reads and writes data to and from object storage. The object storage never initiates, it just responds.

- An inbound connection is required from the Connector, which must reside on your premises.

A connection between the cluster and the Cloud Tiering service is not required.

- An intercluster LIF is required on each ONTAP node that hosts the volumes you want to tier. The LIF must be associated with the *IPspace* that ONTAP should use to connect to object storage.

When you set up data tiering, Cloud Tiering prompts you for the *IPspace* to use. You should choose the *IPspace* that each LIF is associated with. That might be the "Default" *IPspace* or a custom *IPspace* that you created. Learn more about [LIFs](#) and [IPspaces](#).

Supported volumes and aggregates

The total number of volumes that Cloud Tiering can tier might be less than the number of volumes on your ONTAP system. That's because volumes can't be tiered from some aggregates. Refer to ONTAP documentation for [functionality or features not supported by FabricPool](#).



Cloud Tiering supports both FlexVol and FlexGroup volumes.

Discovering an ONTAP cluster

You need to create an on-prem ONTAP working environment in the Cloud Manager Canvas before you can start tiering cold data.

[Learn how to discover a cluster.](#)

Preparing S3-compatible object storage

S3-compatible object storage must meet the following requirements.

S3 credentials

When you set up tiering to S3-compatible object storage, you're prompted to create an S3 bucket or to select an existing S3 bucket. You need to provide Cloud Tiering with an S3 access key and secret key. Cloud Tiering uses the keys to access your bucket.

These access keys must be associated with a user who has the following permissions:

```
"s3>ListAllMyBuckets",  
"s3>ListBucket",  
"s3.GetObject",  
"s3.PutObject",  
"s3>DeleteObject",  
"s3>CreateBucket"
```

Creating or switching Connectors

A Connector is required to tier data to the cloud. When tiering data to S3-compatible object storage, a Connector must be available on your premises. You'll either need to install a new Connector or make sure that the currently selected Connector resides on-prem.

- [Learn about Connectors](#)
- [Connector host requirements](#)
- [Installing the Connector on an existing Linux host](#)
- [Switching between Connectors](#)

Preparing networking for the Connector

Ensure that the Connector has the required networking connections.

Steps

1. Ensure that the network where the Connector is installed enables the following connections:
 - An outbound internet connection to the Cloud Tiering service over port 443 (HTTPS)
 - An HTTPS connection over port 443 to S3-compatible object storage
 - An HTTPS connection over port 443 to your ONTAP cluster

Tiering inactive data from your first cluster to S3-compatible object storage

After you prepare your environment, start tiering inactive data from your first cluster.

What you'll need

- [An on-premises working environment.](#)
- The FQDN of the S3-compatible object storage server and the port that will be used for HTTPS communications.
- An access key and secret key that has the required S3 permissions.

Steps

1. Select an on-prem cluster.
2. In the right panel, click **Enable** for the Tiering service.

The screenshot shows the 'sgtest1' volume configuration page. At the top, there's a status icon (green 'On') and three action buttons (info, more, delete). Below is a 'DETAILS' section for 'On-Premises ONTAP'. Under 'SERVICES', there are two sections: 'Replication' (status: Off, button: Enable) and 'Tiering' (status: Off, button: Enable, which is highlighted with a red box). Each service has a three-dot menu icon to its right.

3. **Choose your provider:** Select **S3 Compatible** and click **Continue**.

4. Complete the steps on the **Tiering Setup** page:

- Server:** Enter the FQDN of the S3-compatible object storage server, the port that ONTAP should use for HTTPS communication with the server, and the access key and secret key for an account that has the required S3 permissions.
- Bucket:** Add a new bucket or select an existing bucket and click **Continue**.
- Cluster Network:** Select the IPspace that ONTAP should use to connect to object storage and click **Continue**.

Selecting the correct IPspace ensures that Cloud Tiering can set up a connection from ONTAP to your S3-compatible object storage.

5. On the **Success** page click **Continue** to set up your volumes now.

6. On the **Tier Volumes** page, select the volumes that you want to configure tiering for and click **Configure volumes**:

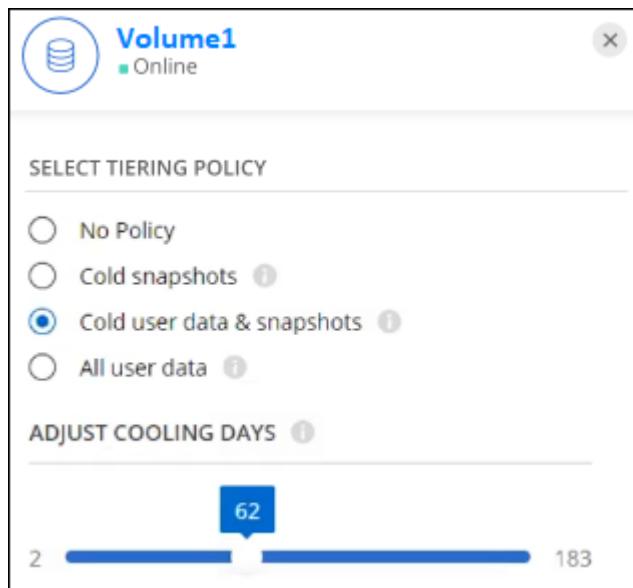
- To select all volumes, check the box in the title row (**Volume Name**) and click **Configure volumes**.
- To select multiple volumes, check the box for each volume (**Volume_1**) and click **Configure volumes**.
- To select a single volume, click the row (or icon) for the volume.

The screenshot shows the 'Tier Volumes' page with three volumes listed: vol2, vol1, and vol3. The header includes a 'Configure volumes' button (1), a search bar, and a 'Tiering Policy' dropdown. The volume table has columns for Volume Name, SVM Name, Volume Size, Used Size, Snapshot Used Size, Cold Data, Tier Status, and Tiering Policy. Row 1 (vol2) has an empty checkbox. Row 2 (vol1) has a checked checkbox (2) and an edit icon (3). Row 3 (vol3) has a checked checkbox.

	Volume Name	SVM Name	Volume Size	Used Size	Snapshot Used Size	Cold Data	Tier Status	Tiering Policy
<input type="checkbox"/>	vol2	svm_AFF1	200 GB	4.11 MB	424 KB	2.88 MB	70 %	<input checked="" type="checkbox"/> Tiered Volume Cold snapshots
<input checked="" type="checkbox"/>	vol1	svm_AFF1	200 GB	1 MB	0 B	716.8 KB	70 %	<input checked="" type="checkbox"/> Tiered Volume Cold snapshots
<input checked="" type="checkbox"/>	vol3	svm_AFF1	200 GB	1 MB	0 B	716.8 KB	70 %	<input checked="" type="checkbox"/> Tiered Volume Cold snapshots

7. In the *Tiering Policy* dialog, select a tiering policy, optionally adjust the cooling days for the selected volumes, and click **Apply**.

[Learn more about volume tiering policies and cooling days.](#)



Result

You've successfully set up data tiering from volumes on the cluster to S3-compatible object storage.

What's next?

[Be sure to subscribe to the Cloud Tiering service.](#)

You can add additional clusters or review information about the active and inactive data on the cluster. For details, see [Managing data tiering from your clusters](#).

Set up licensing for Cloud Tiering

A 30-day free trial of Cloud Tiering starts when you set up tiering from your first cluster. After the free trial ends, you'll need to pay for Cloud Tiering through a pay-as-you-go Cloud Manager subscription from your cloud provider's marketplace, a BYOL license from NetApp, or a combination of both.

A few notes before you read any further:

- If you've already subscribed to the Cloud Manager subscription (PAYGO) in your cloud provider's marketplace, then you're automatically subscribed to Cloud Tiering from on-premises ONTAP systems as well. You'll see an active subscription in the Cloud Tiering **Licensing** tab. You won't need to subscribe again.
- The BYOL Cloud Tiering license is a *floating* license that you can use across multiple on-premises ONTAP clusters in your Cloud Manager account. This is different than in the past where you purchased a *FabricPool* license for each cluster.
- There are no charges when tiering data to StorageGRID, so neither a BYOL license or PAYGO registration is required. This tiered data doesn't count against the capacity purchased in your license.

Learn more about how licensing works for Cloud Tiering.

Use a Cloud Tiering PAYGO subscription

Pay-as-you-go subscriptions from your cloud provider's marketplace enable you to license the use of Cloud Volumes ONTAP systems and many Cloud Data Services, such as Cloud Tiering.

Subscribing from the AWS Marketplace

Subscribe to Cloud Tiering from the AWS Marketplace to set up a pay-as-you-go subscription for data tiering from ONTAP clusters to AWS S3.

Steps

1. In Cloud Manager, click **Tiering > Licensing**.
2. Click **Subscribe** under AWS Marketplace and then click **Continue**.
3. Subscribe from the [AWS Marketplace](#), and then log back in to Cloud Central to complete the registration.

The following video shows the process:

► https://docs.netapp.com/us-en/occm//media/video_subscribing_aws_tiering.mp4 (video)

Subscribing from the Azure Marketplace

Subscribe to Cloud Tiering from the Azure Marketplace to set up a pay-as-you-go subscription for data tiering from ONTAP clusters to Azure Blob storage.

Steps

1. In Cloud Manager, click **Tiering > Licensing**.
2. Click **Subscribe** under Azure Marketplace and then click **Continue**.
3. Subscribe from the [Azure Marketplace](#), and then log back in to Cloud Central to complete the registration.

The following video shows the process:

► https://docs.netapp.com/us-en/occm//media/video_subscribing_azure_tiering.mp4 (video)

Subscribing from the GCP Marketplace

Subscribe to Cloud Tiering from the GCP Marketplace to set up a pay-as-you-go subscription for data tiering from ONTAP clusters to Google Cloud storage.

Steps

1. In Cloud Manager, click **Tiering > Licensing**.
2. Click **Subscribe** under GCP Marketplace and then click **Continue**.
3. Subscribe from the [GCP Marketplace](#), and then log back in to Cloud Central to complete the registration.

The following video shows the process:

► https://docs.netapp.com/us-en/occm//media/video_subscribing_gcp_tiering.mp4 (video)

Use a Cloud Tiering BYOL license

Bring-your-own licenses from NetApp provide 2-, 12-, 24-, or 36-month terms. The **BYOL Cloud Tiering** license is a *floating* license that you can use across multiple on-premises ONTAP clusters in your Cloud Manager account. The total tiering capacity defined in your Cloud Tiering license is shared among **all** of your on-prem clusters, making initial licensing and renewal easy.

If you don't have a Cloud Tiering license, contact us to purchase one:

- [Send email to purchase a license](#).
- Click the chat icon in the lower-right of Cloud Manager to request a license.

You use the Digital Wallet page in Cloud Manager to manage Cloud Tiering BYOL licenses. You can add new licenses and update existing licenses.

New Cloud Tiering BYOL licensing starting August 21, 2021

The new **Cloud Tiering** license was introduced on August 21, 2021 for tiering configurations that are supported within Cloud Manager using the Cloud Tiering service. The **FabricPool** license that you may have used in the past to tier on-premises ONTAP data to the cloud is being retained only for configurations that aren't supported within Cloud Manager.

If you have one of these configurations, you'll continue to use a FabricPool license on each cluster (managed using System Manager or the CLI):

- ONTAP installations in the GovCloud or in Dark Sites
- ONTAP clusters in a MetroCluster configuration
- ONTAP clusters using FabricPool Mirror functionality
- ONTAP clusters that are tiering data to IBM Cloud Object Storage or Alibaba Cloud Object Storage

Cloud Manager currently supports tiering to the following cloud storage: Amazon S3, Azure Blob storage, Google Cloud Storage, and S3-compatible object storage.

Note that tiering to StorageGRID does not require a FabricPool or Cloud Tiering license.

If you are currently using FabricPool licensing, you're not affected until your FabricPool license reaches its expiration date or maximum capacity. Contact NetApp when you need to update your license, or earlier to make sure there is no interruption in your ability to tier data to the cloud.

- If you're using a configuration that is supported in Cloud Manager, your FabricPool licenses will be converted to Cloud Tiering licenses and they'll appear in the Digital Wallet. When those initial licenses expire, you'll need to update the Cloud Tiering licenses.
- If you're using a configuration that is not supported in Cloud Manager, then you'll continue using a FabricPool license. [See how to license tiering using System Manager](#).

Here are some things you need to know about the two licenses:

Cloud Tiering license	FabricPool license
It is a <i>floating</i> license that you can use across multiple on-premises ONTAP clusters.	It is a per-cluster license that you purchase and license for every cluster.

Cloud Tiering license	FabricPool license
It is registered in Cloud Manager in the Digital Wallet.	It is applied to individual clusters using System Manager or the ONTAP CLI.
Tiering configuration and management is done through the Cloud Tiering service in Cloud Manager.	Tiering configuration and management is done through System Manager or the ONTAP CLI.
Once configured, you can use the tiering service without a license for 30 days using the free trial.	Once configured, you can tier the first 10 TB of data for free.

Obtain your Cloud Tiering license file

After you have purchased your Cloud Tiering license, you activate the license in Cloud Manager by entering the Cloud Tiering serial number and NSS account, or by uploading the NLF license file. The steps below show how to get the NLF license file if you plan to use that method.

Steps

1. Sign in to the [NetApp Support Site](#) and click **Systems > Software Licenses**.
2. Enter your Cloud Tiering license serial number.

Software Licenses						
Serial Number		License Key				
Serial #		Cluster SN	License Name	License Key	Host ID	Value
4810	CLOUD_TIERING_SVC	Get NetApp License File	100	12/31/9998		

3. Under **License Key**, click **Get NetApp License File**.
4. Enter your Cloud Manager Account ID (this is called a Tenant ID on the support site) and click **Submit** to download the license file.

Get License	
SERIAL NUMBER:	4810
LICENSE:	CLOUD_TIERING_SVC
SALES ORDER:	3005
TENANT ID:	<input type="text" value="Enter Tenant ID"/>
Example: account-xxxxxxxx	
Cancel	Submit

You can find your Cloud Manager Account ID by selecting the **Account** drop-down from the top of Cloud Manager, and then clicking **Manage Account** next to your account. Your Account ID is in the Overview tab.

Add Cloud Tiering BYOL licenses to your account

After you purchase a Cloud Tiering license for your Cloud Manager account, you need to add the license to Cloud Manager to use the Cloud Tiering service.

Steps

1. Click **All Services > Digital Wallet > Data Services Licenses**.
2. Click **Add License**.
3. In the *Add License* dialog, enter the license information and click **Add License**:
 - If you have the tiering license serial number and know your NSS account, select the **Enter Serial Number** option and enter that information.
If your NetApp Support Site account isn't available from the drop-down list, [add the NSS account to Cloud Manager](#).
 - If you have the tiering license file, select the **Upload License File** option and follow the prompts to attach the file.

The screenshot shows the 'Add License' dialog box. It has two main tabs: 'Enter Serial Number' (which is selected) and 'Upload License File'. Under the 'Enter Serial Number' tab, there is a 'Serial Number' input field containing 'Enter Serial Number' and a 'NetApp Support Site Account' dropdown menu containing 'Select Support Site Account'. At the bottom are 'Add License' and 'Cancel' buttons. The 'Upload License File' tab shows a 'Upload License File' input field with an 'Upload' button and a note about obtaining the file from the NetApp Support Site.

Result

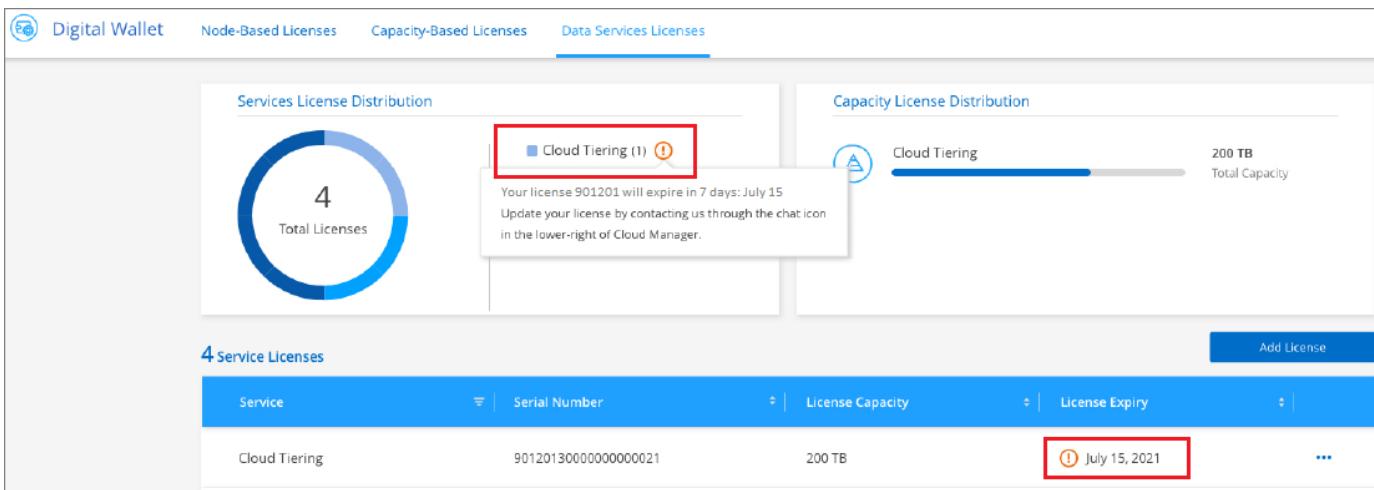
Cloud Manager adds the license so that your Cloud Tiering service is active.

Update a Cloud Tiering BYOL license

If your licensed term is nearing the expiration date, or if your licensed capacity is reaching the limit, you'll be notified in Cloud Tiering.



This status also appears in the Digital Wallet page.



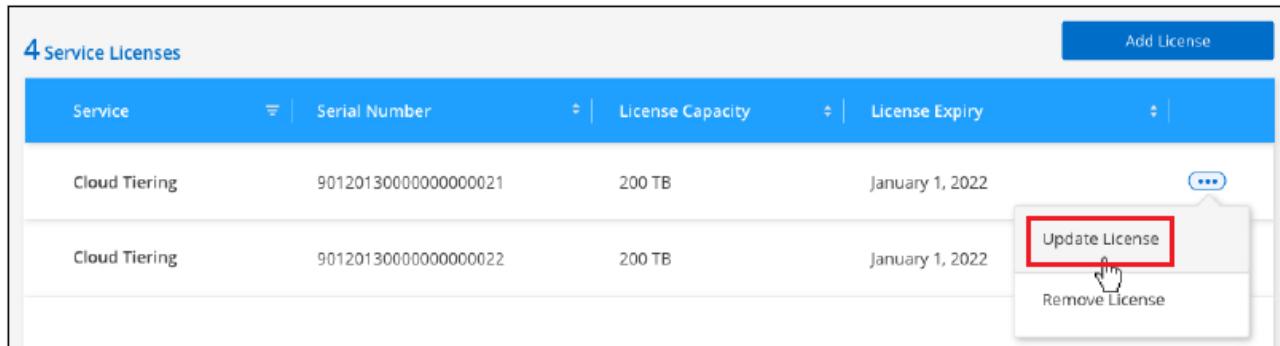
You can update your Cloud Tiering license before it expires so that there is no interruption in your ability to tier your data to the cloud.

Steps

1. Click the chat icon in the lower-right of Cloud Manager to request an extension or capacity add-on to your Cloud Tiering license for the particular serial number.

After you pay for the license and it is registered with the NetApp Support Site, in most cases, Cloud Manager can automatically obtain your updated license file and the Data Services Licenses page will reflect the change in 5 to 10 minutes.

2. If Cloud Manager can't automatically update the license, then you'll need to manually upload the license file.
 - a. You can [obtain the license file from the NetApp Support Site](#).
 - b. On the Digital Wallet page in the *Data Services Licenses* tab, click **...** for the service serial number you are updating, and click **Update License**.



- c. In the *Update License* page, upload the license file and click **Update License**.

Result

Cloud Manager updates the license so that your Cloud Tiering service continues to be active.

Measure network latency and throughput performance

Run a Cloud Performance Test to measure network latency and throughput performance from an ONTAP cluster to an object store before and after setting up data tiering. The test

also identifies any failures that occurred.

Here are sample performance results:

Your cluster performance results			
Operation	Size	Avg.Latency (ms)	Throughput
PUT	4 MB	502	408.06 MB
GET	4 KB	79	15.05 MB
GET	8 KB	197	28.35 MB
GET	32 KB	291	109.71 MB
GET	256 KB	361	714.39 MB

Before you get started

It's best to run this check when the cluster is under 50% CPU utilization.

Steps for a cluster that hasn't been set up for tiering

1. At the top of Cloud Manager, click **Tiering**.
2. From the **Cluster Dashboard**, click the menu icon for a cluster and select **Cloud Performance Test**.
3. Review the details and click **Continue**.
4. Follow the prompts to provide the required information.

The information that you need to provide is the same as if you were setting up tiering on the cluster.

5. Optionally continue to the Tier Volumes wizard to complete the setup.

Steps for a cluster that has been set up for tiering

1. At the top of Cloud Manager, click **Tiering**.
2. From the **Cluster Dashboard**, click the menu icon for a cluster and select **Cloud Performance Test**.
3. Select a node from the drop-down list.
4. View the results or recheck the performance.

Managing data tiering from your clusters

Now that you've set up data tiering from your on-prem ONTAP clusters, you can tier data from additional volumes, change a volume's tiering policy, discover additional clusters,

and more.

Tiering data from additional volumes

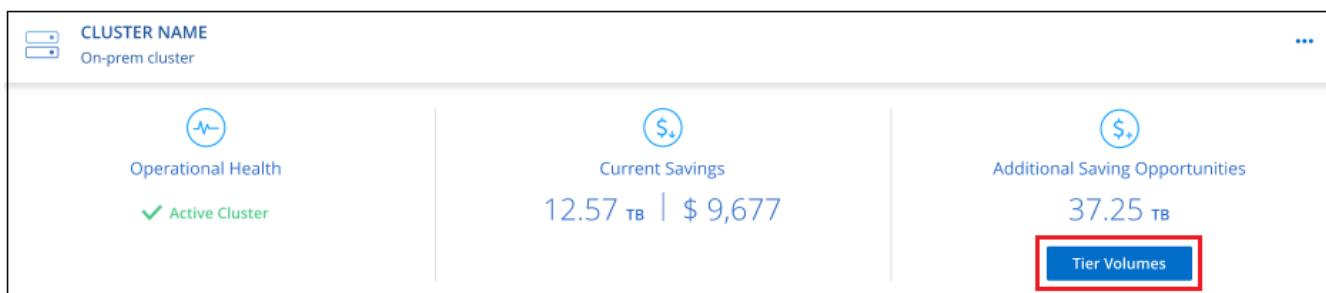
Set up data tiering for additional volumes at any time—for example, after creating a new volume.



You don't need to configure the object storage because it was already configured when you initially set up tiering for the cluster. ONTAP will tier inactive data from any additional volumes to the same object store.

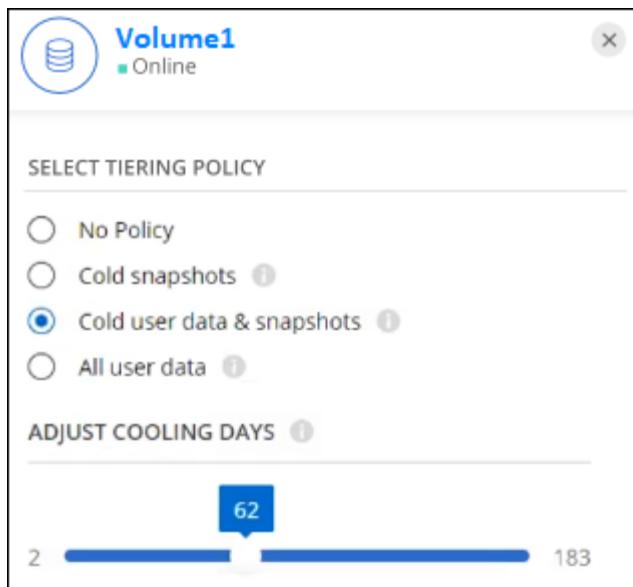
Steps

1. At the top of Cloud Manager, click **Tiering**.
2. From the **Cluster Dashboard**, click **Tier Volumes** for the cluster.



3. On the *Tier Volumes* page, select the volumes that you want to configure tiering for and launch the **Tiering Policy** page:
 - To select all volumes, check the box in the title row (**Volume Name**) and click **Configure volumes**.
 - To select multiple volumes, check the box for each volume (**Volume_1**) and click **Configure volumes**.
 - To select a single volume, click the row (or icon) for the volume.
4. In the *Tiering Policy* dialog, select a tiering policy, optionally adjust the cooling days for the selected volumes, and click **Apply**.

[Learn more about volume tiering policies and cooling days.](#)



Result

The selected volumes start to have their data tiered to the cloud.

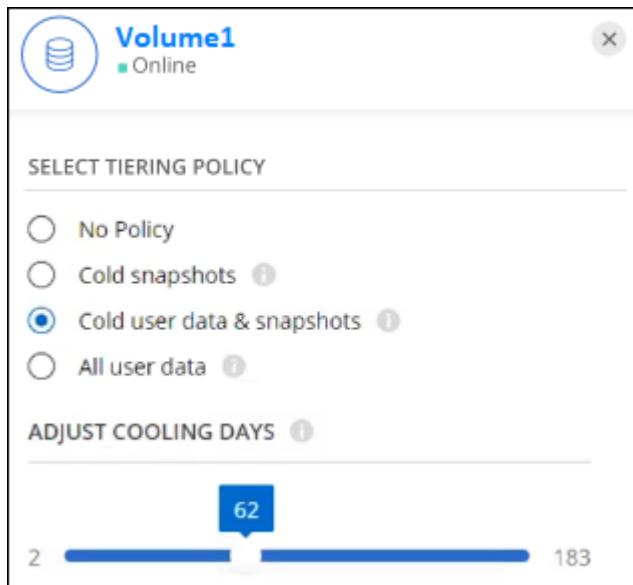
Changing a volume's tiering policy

Changing the tiering policy for a volume changes how ONTAP tiers cold data to object storage. The change starts from the moment that you change the policy. It changes only the subsequent tiering behavior for the volume—it does not retroactively move data to the cloud tier.

Steps

1. At the top of Cloud Manager, click **Tiering**.
2. From the **Cluster Dashboard**, click **Tier Volumes** for the cluster.
3. Click the row for a volume, select a tiering policy, optionally adjust the cooling days, and click **Apply**.

[Learn more about volume tiering policies and cooling days.](#)



Note: If you see options to "Retrieve Tiered Data", see [Migrating data from the cloud tier back to the performance tier](#) for details.

Result

The tiering policy is changed and data begins to be tiered based on the new policy.

Migrating data from the cloud tier back to the performance tier

Tiered data that is accessed from the cloud may be "re-heated" and moved back to the performance tier. However, if you want to proactively promote data to the performance tier from the cloud tier, you can do this in the *Tiering Policy* dialog. This capability is available when using ONTAP 9.8 and greater.

You might do this if you want to stop using tiering on a volume, or if you decide to keep all user data on the performance tier, but keep Snapshot copies on the cloud tier.

There are two options:

Option	Description	Affect on Tiering Policy
Bring back all data	Retrieves all volume data and Snapshot copies tiered in the cloud and promotes them to the performance tier.	Tiering policy is changed to "No policy".
Bring back active file system	Retrieves only active file system data tiered in the cloud and promotes it to the performance tier (Snapshot copies remain in the cloud).	Tiering policy is changed to "Cold snapshots".

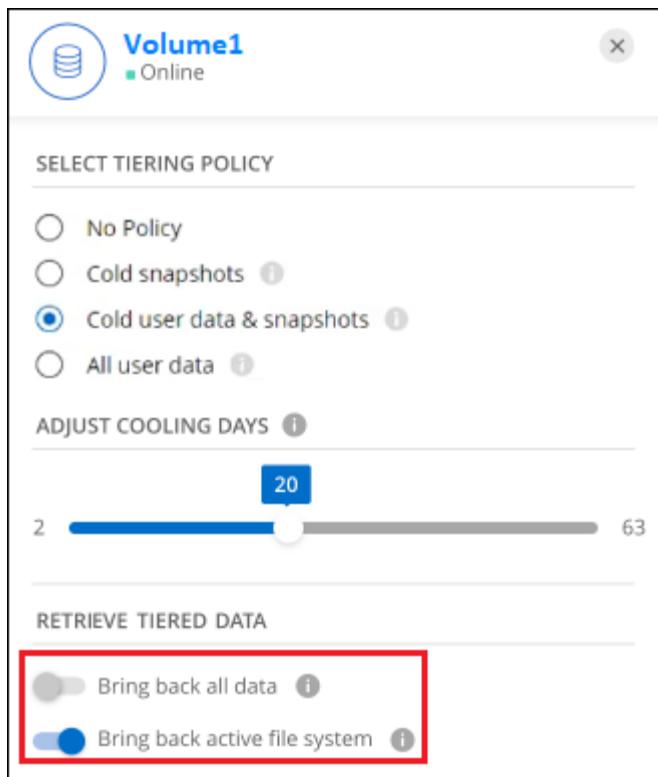


You may be charged by your cloud provider based on that amount of data transferred off the cloud.

Steps

Make sure you have enough space in the performance tier for all the data that is being moved from the cloud.

1. At the top of Cloud Manager, click **Tiering**.
2. From the **Cluster Dashboard**, click **Tier Volumes** for the cluster.
3. Click the icon for the volume, choose the retrieval option you want to use, and click **Apply**.



Result

The tiering policy is changed and the tiered data starts to be migrated back to the performance tier. Depending on the amount of data in the cloud, the transfer process could take some time.

Managing tiering settings on aggregates

Each aggregate in your on-prem ONTAP systems has two settings that you can adjust: the tiering fullness threshold and whether inactive data reporting is enabled.

Tiering fullness threshold

Setting the threshold to a lower number reduces the amount of data required to be stored on the performance tier before tiering takes place. This might be useful for large aggregates that contain little active data.

Setting the threshold to a higher number increases the amount of data required to be stored on the performance tier before tiering takes place. This might be useful for solutions designed to tier only when aggregates are near maximum capacity.

Inactive data reporting

Inactive data reporting (IDR) uses a 31-day cooling period to determine which data is considered inactive. The amount of cold data that is tiered is dependent on the tiering policies set on volumes. This amount might be different than the amount of cold data detected by IDR using a 31-day cooling period.

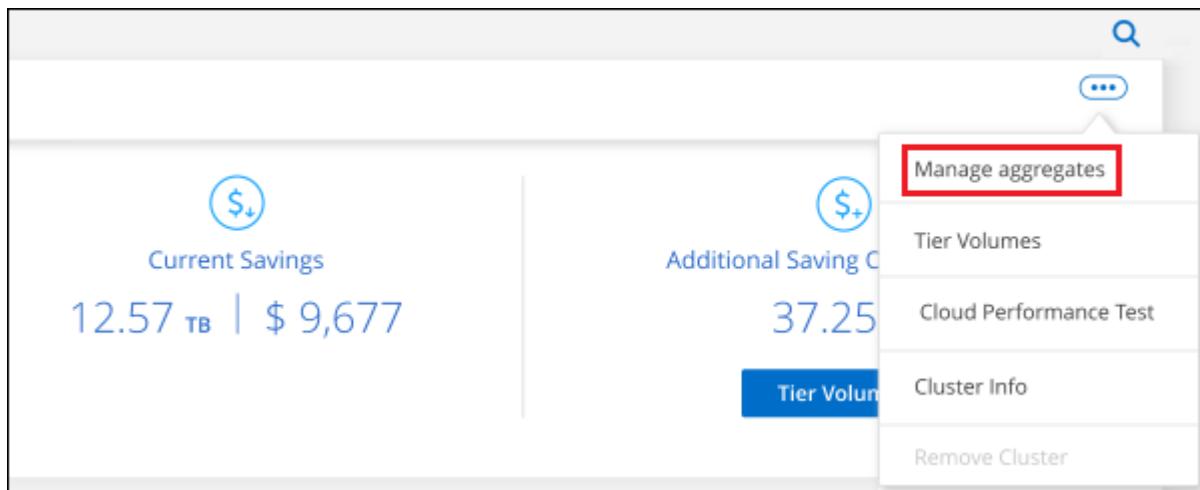


It's best to keep IDR enabled because it helps to identify your inactive data and savings opportunities. IDR must remain enabled if data tiering was enabled on an aggregate.

Steps

1. At the top of Cloud Manager, click **Tiering**.

2. From the **Cloud Tiering** page, click the menu icon for a cluster and select **Manage Aggregates**.



3. On the **Manage Aggregates** page, click the icon for an aggregate in the table.
4. Modify the fullness threshold and choose whether to enable or disable inactive data reporting.

The screenshot shows the 'Manage Aggregates' page for an aggregate named 'aggr1'. It includes a header with an arrow icon and the aggregate name. Below it, there's a section titled 'Tier data when the aggregate is this full:' with a progress bar from 0% to 100%. A cloud icon is positioned at the 50% mark. At the bottom, there's a section titled 'Activate inactive data reporting' with a blue toggle switch that is turned on.

5. Click **Apply**.

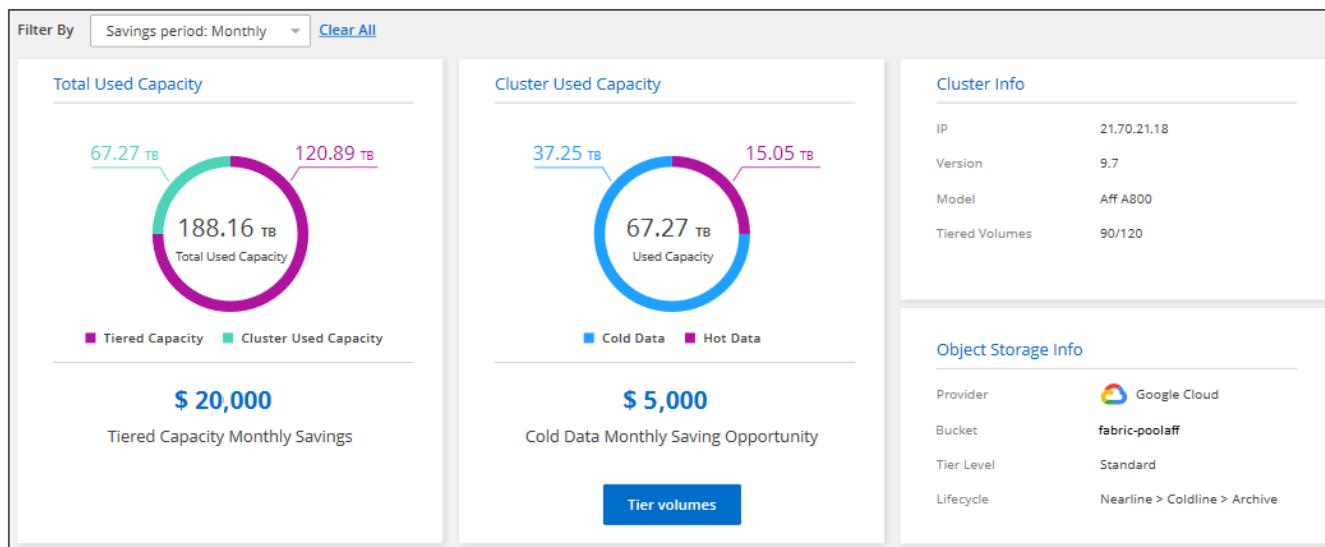
Reviewing tiering info for a cluster

You might want to see how much data is in the cloud tier and how much data is on disks. Or, you might want to see the amount of hot and cold data on the cluster's disks. Cloud Tiering provides this information for each cluster.

Steps

1. At the top of Cloud Manager, click **Tiering**.
2. From the **Cluster Dashboard**, click the menu icon for a cluster and select **Cluster info**.
3. Review details about the cluster.

Here's an example:

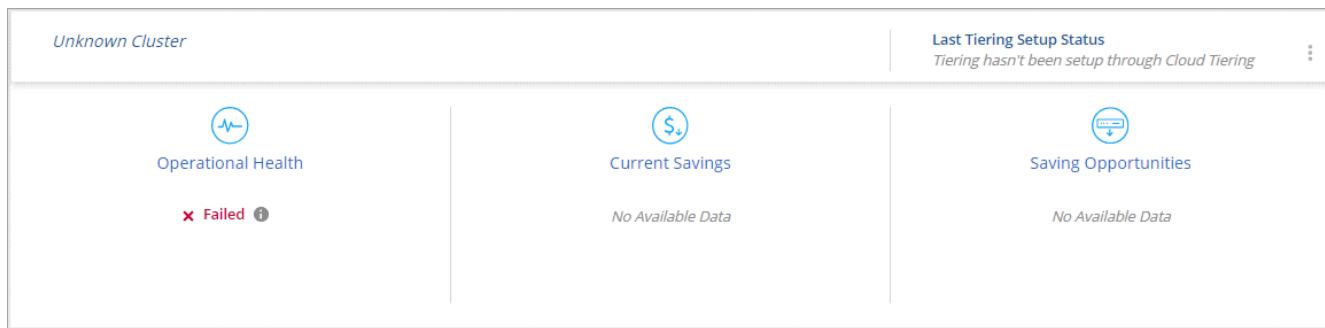


Fixing operational health

Failures can happen. When they do, Cloud Tiering displays a "Failed" operational health status on the Cluster Dashboard. The health reflects the status of the ONTAP system and Cloud Manager.

Steps

1. Identify any clusters that have an operational health of "Failed."



2. Hover over the icon to see the failure reason.

3. Correct the issue:

- a. Verify that the ONTAP cluster is operational and that it has an inbound and outbound connection to your object storage provider.
- b. Verify that Cloud Manager has outbound connections to the Cloud Tiering service, to the object store, and to the ONTAP clusters that it discovers.

Discovering additional clusters from Cloud Tiering

You can add your undiscovered on-prem ONTAP clusters to Cloud Manager from the *Tiering Cluster Dashboard* so that you can enable tiering for the cluster.

Note that buttons also appear on the *Tiering On-Prem Overview* page for you to discover additional clusters.

Steps

1. From Cloud Tiering, click the **Cluster Dashboard** tab.

- To see any undiscovered clusters, click **Show Undiscovered Clusters**.

The screenshot shows the Cloud Tiering Cluster Dashboard. At the top, there are three categories: 'All clusters' (6), 'On-prem cluster' (3), and 'Cloud Volumes ONTAP' (3). Below these, a button labeled 'Show Undiscovered Clusters' is highlighted with a red box.

If your NSS credentials are saved in Cloud Manager, the clusters in your account are displayed in the list.

If your NSS credentials are not saved in Cloud Manager, you are first prompted to add your credentials before you can see the undiscovered clusters.

- Scroll down the page to locate the clusters.

The screenshot shows the Cluster Details page. At the top, it displays cluster counts: All clusters (6), On-prem cluster (3), Cloud Volumes ONTAP (3), and Undiscovered Clusters (3). Below this, a section titled 'Clusters (6/6)' shows operational health, current savings, and additional saving opportunities. A specific cluster entry for 'On-prem cluster' with IP address '127.0.0.1' is shown. At the bottom, a callout box says 'To Optimize Data Tiering, Discover Your Cluster' with the instruction 'Please make sure your connector and cluster are on the same network'. A blue button labeled 'Discover Cluster' is highlighted with a red box.

- Click **Discover Cluster** for the cluster that you want to manage through Cloud Manager and implement data tiering.
- On the *Choose a Location* page **On-Premises ONTAP** is pre-selected, so just click **Continue**.
- On the *ONTAP Cluster Details* page, enter the password for the admin user account and click **Add**.

Note that the cluster management IP address is populated based on information from your NSS account.

- On the *Details & Credentials* page the cluster name is added as the Working Environment Name, so just click **Go**.

Result

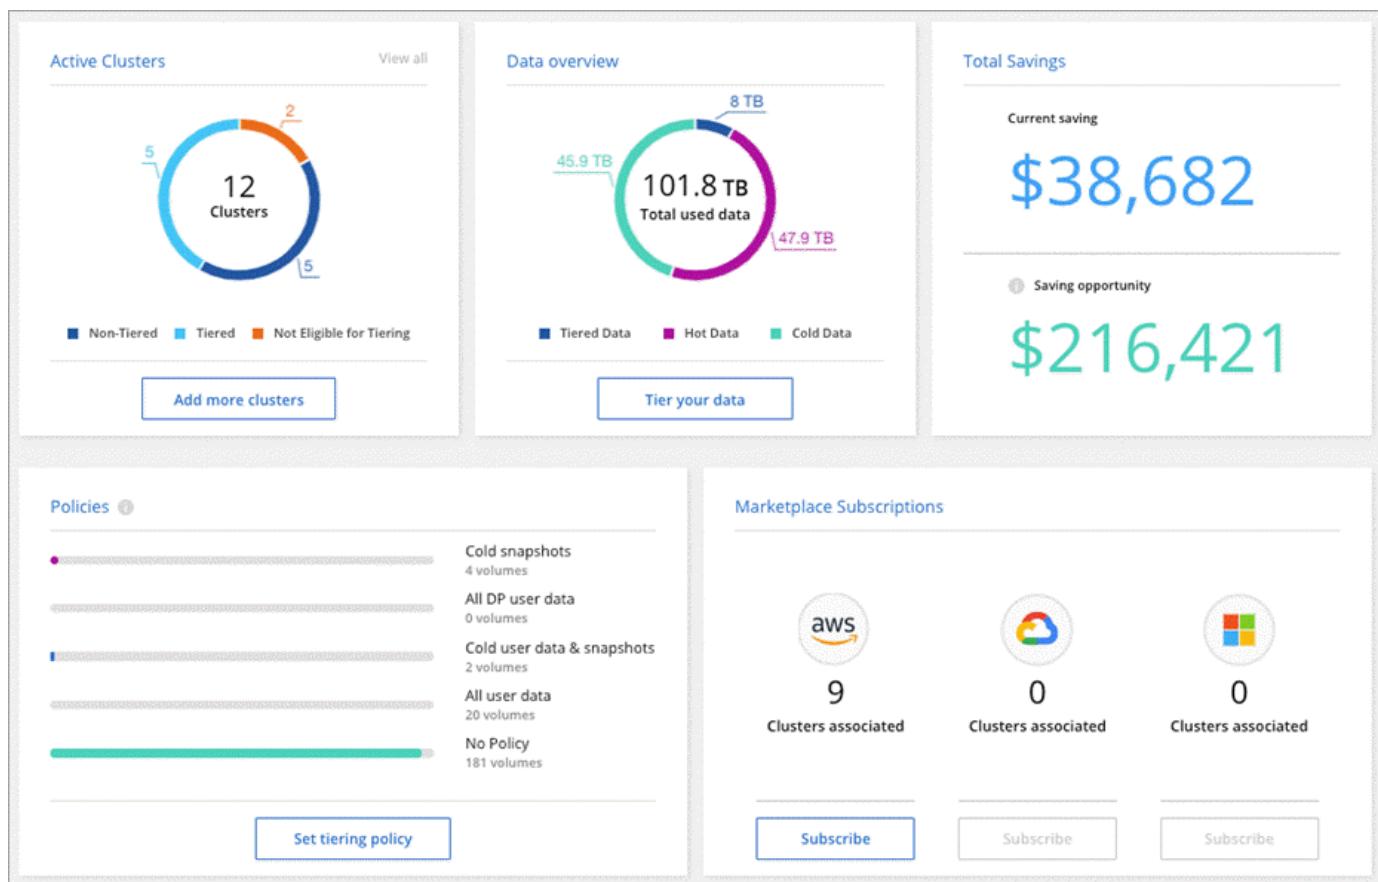
Cloud Manager discovers the cluster and adds it to a working environment in the Canvas using the cluster name as the working environment name.

You can enable the Tiering service or other services for this cluster in the right panel.

Get an overview of data tiering from your clusters

Cloud Tiering provides an aggregated view of data tiering from each of your on-premises clusters. This overview provides a clear picture of your environment and enables you to take proper actions.

Cloud Tiering provides the following details about your environment:



Active Clusters

The number of clusters that are currently tiering data to the cloud, the clusters that aren't tiering data to the cloud, and the number of clusters that don't support data tiering.

Data Overview

The amount of data that was tiered to the cloud, and the amount of hot and cold data on the cluster.

Total Savings

The amount of money that you've saved by tiering data to the cloud, as well as the amount of money that you could save by tiering more data to the cloud.

Policies

The number of times that each tiering policy has been applied to a volume.

Marketplace Subscriptions

The number of clusters associated with each type of Marketplace Subscription and an indication about your subscription status.

Steps

1. Click **Tiering > On-Prem Overview**.

Cloud Tiering technical FAQ

This FAQ can help if you're just looking for a quick answer to a question.

ONTAP

The following questions relate to ONTAP.

What are the requirements for my ONTAP cluster?

It depends on where you tier the cold data. Refer to the following:

- [Tiering data from on-premises ONTAP clusters to Amazon S3](#)
- [Tiering data from on-premises ONTAP clusters to Azure Blob storage](#)
- [Tiering data from on-premises ONTAP clusters to Google Cloud Storage](#)
- [Tiering data from on-premises ONTAP clusters to StorageGRID](#)
- [Tiering data from on-premises ONTAP clusters to S3 object storage](#)

Does Cloud Tiering enable inactive data reporting?

Yes, Cloud Tiering enables inactive data reporting on each aggregate. This setting enables us to identify the amount of inactive data that can be tiered to low-cost object storage.



Cloud Tiering enables inactive data reporting on HDD aggregates if the cluster is running ONTAP 9.6 or later.

Can I tier data from NAS volumes and SAN volumes?

You can use Cloud Tiering to tier data from NAS volumes to the public cloud and from SAN volumes to a private cloud using StorageGRID.

What about Cloud Volumes ONTAP?

If you have Cloud Volumes ONTAP systems, you'll find them in the Cluster Dashboard so you get a full view of data tiering in your hybrid cloud infrastructure.

From the Cluster Dashboard, you can view tiering information similar to an on-prem ONTAP cluster: operational health, current savings, savings opportunities, details about volumes and aggregates, and more.

Cloud Volumes ONTAP systems are read-only from Cloud Tiering. You can't set up data tiering on Cloud Volumes ONTAP from Cloud Tiering. [You set up tiering for Cloud Volumes ONTAP from the working environment in Cloud Manager](#).

Object storage

The following questions relate to object storage.

Which object storage providers are supported?

Amazon S3, Azure Blob storage, Google Cloud Storage, NetApp StorageGRID, and S3-compatible object storage providers are supported.

Can I use my own bucket/container?

Yes, you can. When you set up data tiering, you have the choice to add a new bucket/container or to select an existing bucket/container.

Which regions are supported?

- [Supported AWS regions](#)
- [Supported Azure regions](#)
- [Supported Google Cloud regions](#)

Which S3 storage classes are supported?

Cloud Tiering supports data tiering to the *Standard*, *Standard-Infrequent Access*, *One Zone-IA*, or *Intelligent* storage classes. See [Supported S3 storage classes](#) for more details.

Which Azure Blob access tiers are supported?

Cloud Tiering supports data tiering to the *Hot* or *Cool* access tiers for your inactive data. See [Supported Azure Blob access tiers](#) for more details.

Which storage classes are supported for Google Cloud Storage?

Cloud Tiering supports data tiering to the *Standard*, *Nearline*, *Coldline*, and *Archive* storage classes. See [Supported Google Cloud storage classes](#) for more details.

Does Cloud Tiering use one object store for the entire cluster or one per aggregate?

One object store for the entire cluster.

Can I apply policies to my object store to move data around independent of tiering?

Yes. You can enable life cycle management so that Cloud Tiering transitions data from the default storage class/access tier to a more cost-effective tier after a certain number of days.

The life cycle rule is applied to all objects in the selected bucket for Amazon S3 and Google Cloud storage, and to all containers in the selected storage account for Azure Blob.

Connectors

The following questions relate to Connectors.

Where does the Connector need to be installed?

- When tiering data to S3, the Connector can reside in an AWS VPC or on your premises.
- When tiering data to Blob storage, the Connector can reside in an Azure VNet or on your premises.
- When tiering data to Google Cloud Storage, the Connector must reside in a Google Cloud Platform VPC.

- When tiering data to StorageGRID or other S3-Compatible storage providers, the Connector must reside on your premises.

Networking

The following questions relate to networking.

What are the networking requirements?

- The ONTAP cluster initiates an HTTPS connection over port 443 to your object storage provider. ONTAP reads and writes data to and from object storage. The object storage never initiates, it just responds.
- For StorageGRID, the ONTAP cluster initiates an HTTPS connection over a user-specified port to StorageGRID (the port is configurable during tiering setup).
- A Connector needs an outbound HTTPS connection over port 443 to your ONTAP clusters, to the object store, and to the Cloud Tiering service.

For more details, see:

- [Tiering data from on-premises ONTAP clusters to Amazon S3](#)
- [Tiering data from on-premises ONTAP clusters to Azure Blob storage](#)
- [Tiering data from on-premises ONTAP clusters to Google Cloud Storage](#)
- [Tiering data from on-premises ONTAP clusters to StorageGRID](#)
- [Tiering data from on-premises ONTAP clusters to S3 object storage](#)

Permissions

The following questions relate to permissions.

What permissions are required in AWS?

Permissions are required to [manage the S3 bucket](#).

What permissions are required in Azure?

No extra permissions are needed outside of the permissions that you need to provide to Cloud Manager.

What permissions are required in Google Cloud Platform?

Storage Admin permissions are needed for a [service account that has storage access keys](#).

What permissions are required for StorageGRID?

[S3 permissions are needed](#).

What permissions are required for S3-compatible object storage?

[S3 permissions are needed](#).

Reference

Supported S3 storage classes and regions

Cloud Tiering supports several S3 storage classes and most regions.

Supported S3 storage classes

When you set up data tiering to AWS, Cloud Tiering automatically uses the *Standard* storage class for your inactive data. Cloud Tiering can apply a lifecycle rule so the data transitions from the *Standard* storage class to another storage class after a certain number of days. You can choose from the following storage classes:

- Standard-Infrequent Access
- One Zone-IA
- Intelligent-Tiering (where AWS automatically moves data between two tiers — Frequent Access and Infrequent Access - when access patterns change)

If you do not choose another storage class, then the data remains in the *Standard* storage class and no rules are applied.

When you configure a Cloud Tiering lifecycle rule, you must not configure any lifecycle rules when setting up the bucket in your AWS account.

[Learn about S3 storage classes.](#)

Supported AWS regions

Cloud Tiering supports the following AWS regions.

Asia Pacific

- Mumbai
- Seoul
- Singapore
- Sydney
- Tokyo

Europe

- Frankfurt
- Ireland
- London
- Paris
- Stockholm

North America

- Canada Central

- GovCloud (US-West) – starting with ONTAP 9.3
- US East (N. Virginia)
- US East (Ohio)
- US West (N. California)
- US West (Oregon)

South America

- São Paulo

Supported Azure Blob access tiers and regions

Cloud Tiering supports two access tiers and most regions.

Supported Azure Blob access tiers

When you set up data tiering to Azure, Cloud Tiering automatically uses the *Hot* access tier for your inactive data. Cloud Tiering can apply a lifecycle rule so the data transitions from the *Hot* access tier to the *Cool* access tier after a certain number of days.

If you do not choose the *Cool* access tier, then the data remains in the *Hot* access tier and no rules are applied.

When you configure a Cloud Tiering lifecycle rule, you must not configure any lifecycle rules when setting up the container in your Azure account.

[Learn about Azure Blob access tiers.](#)

Supported Azure regions

Cloud Tiering supports the following Azure regions.

Africa

- South Africa North

Asia Pacific

- Australia East
- Australia Southeast
- East Asia
- Japan East
- Japan West
- Korea Central
- Korea South
- Southeast Asia

Europe

- France Central
- Germany Central
- Germany Northeast
- North Europe
- UK South
- UK West
- West Europe

North America

- Canada Central
- Canada East
- Central US
- East US
- East US 2
- North Central US
- South Central US
- West US
- West US 2
- West Central US

South America

- Brazil South

Supported Google Cloud storage classes and regions

Cloud Tiering supports several Google Cloud storage classes and most regions.

Supported GCP storage classes

When you set up data tiering to GCP, Cloud Tiering automatically uses the *Standard* storage class for your inactive data. Cloud Tiering can apply a lifecycle rule so the data transitions from the *Standard* storage class to other storage classes after a certain number of days. You can choose from the following storage classes:

- Nearline
- Coldline
- Archive

If you do not choose another storage class, then the data remains in the *Standard* storage class and no rules are applied.

When you configure a Cloud Tiering lifecycle rule, you must not configure any lifecycle rules when setting up the bucket in your Google account.

[Learn about Google Cloud Storage classes.](#)

Supported Google Cloud regions

Cloud Tiering supports the following regions.

Americas

- Iowa
- Los Angeles
- Montreal
- N. Virginia
- Oregon
- Sao-Paulo
- South Carolina

Asia Pacific

- Hong Kong
- Mumbai
- Osaka
- Singapore
- Sydney
- Taiwan
- Tokyo

Europe

- Belgium
- Finland
- Frankfurt
- London
- Netherlands
- Zurich

Organize your resources using tags

Learn about tagging

Cloud Manager enables you to apply tags to your *existing* resources to help organize and manage those resources. Tags are metadata that you can use to group resources to identify applications, environments, regions, billing codes, cloud providers, and more.

Tags consists of a **tag key** and a **tag value**. For example, you can create a tag key called "Environment" and then add tag values of "Production" and "Test". After they are applied to your resources, you can quickly search for and view resources that match the key/value pair.

You can add tag key/value pairs to *new* resources when you create a working environment or an Azure NetApp Files volume. You can also define tag key/value pairs in [Cloud Manager templates that you build](#) for your storage admins and DevOps engineers.

You can add new tags using the Tagging service, and you can change or delete existing tags.

Features

The Tagging service offers the following features and benefits:

- Create tag keys and tag values that match the terms you use in your environment
- Organize the resources in your environment for easier monitoring and management
- Add, remove, and edit tag keys and tag values by resource type
- Tag ONTAP resources and resources in your environment from AWS and Azure.

Pricing and licenses

The ability to tag your resources requires no licensing and is free to use by all Cloud Manager users with the Account Admin or Workspace Admin role.

Resources that you can tag

You can apply tags to the following resources.

Provider	Service	Resource
ONTAP	Cloud Volumes ONTAP	Aggregate Storage VM Volume
	On-premise ONTAP	Aggregate Storage VM Volume
	Azure NetApp Files	Volume

Provider	Service	Resource
AWS	EC2	Instance Security Group Subnet Volume VPC
Azure	Compute	Snapshot Virtual Machine
	Network	Security Group Virtual Network
	Resource	Resource Group
	Storage	Storage Account
GCP	Compute	Instance
	Storage	Bucket

For information about AWS EC2 tags, refer to [AWS Documentation: Tagging your Amazon EC2 Resources](#).

For information about Azure tags, refer to [Azure Documentation: Tagging your Azure resources](#).

For information about Google labels, refer to [Google Cloud Documentation: Tagging your Google Cloud resources](#).

Prerequisites

Verify your AWS Connector permissions

If you created the Connector using Cloud Manager version 3.9.10 or greater, then you're all set. If you created the Connector using an earlier version of Cloud Manager, then you'll need to add some required permissions for the Cloud Manager IAM role to tag AWS EC2 instances:

```
{
  "Action": [
    "ec2:CreateTags",
    "ec2>DeleteTags",
    "ec2:DescribeTags",
    "tag:getResources",
    "tag:getTagKeys",
    "tag:getTagValues",
    "tag:TagResources",
    "tag:UntagResources"
  ],
  "Resource": "*",
  "Effect": "Allow",
  "Sid": "tagServicePolicy"
}
```

Verify your Azure Connector permissions

If you created the Connector using Cloud Manager version 3.9.10 or greater, then you're all set. If you created the Connector using an earlier version of Cloud Manager, then you'll need to add some required permissions for the Cloud Manager Operator IAM role to tag Azure resources:

```
{  
  "id": "<ID>",  
  "properties": {  
    "roleName": "Cloud Manager Operator-<ID>",  
    "description": "Cloud Manager Operator",  
    "assignableScopes": [  
      "/subscriptions/<SUBSCRIPTION-ID>"  
    ],  
    "permissions": [  
      {  
        "actions": [  
          "Microsoft.Resources/tags/read",  
          "Microsoft.Resources/tags/write",  
          "Microsoft.Resources/tags/delete",  
          "Microsoft.ClassicCompute/virtualMachines/read"  
        ],  
        "notActions": [],  
        "dataActions": [],  
        "notDataActions": []  
      }  
    ]  
  }  
}
```

Tag rules and restrictions

The following rules apply when creating tag keys and tag values:

- Maximum key length: 128 characters
- Maximum key value length: 256 characters
- Valid tag and tag value characters: letters, numbers, spaces, and special characters (_ , @, &, * , etc.)
- Tags are case upper/lower sensitive.
- Maximum tags per resource: 30
- Per resource, each tag key must be unique

Tag examples

Key	Values
Env	production test
Dept	finance sales eng
Owner	admin storage

Manage tags for your resources

You can view, add, modify, and delete the tags assigned to your existing resources using the Cloud Manager Tagging service. This allows you to organize and simplify the management of your resources.

Search for resources that have certain tags

If you want to see all the resources that have a certain tag, or a certain tag and tag key value, you can search for those tags. You can search across all resources, or just within certain resource categories.

Steps

1. Select the **AppTemplate** service and click the **Tags** tab.
2. In the *Resource Type* field, select the resource, for example, **ONTAP:CVO:VOLUME** to search across all Cloud Volumes ONTAP volumes.
3. In the *Tag Key* field, select the tag, for example, **Env** to restrict the search to volumes with the "Env" tag.
4. In the *Tag Value* field, select the key value, for example, **production** to restrict the search to volumes with the tag name "Env" and the tag value "production".

5. Click the **+** to add this search criteria to the Search area.

6. If you are done with your search, click **Search** and the search results appear in the Resources section.

11 Resources				Manage Tags	
ID	Tags	Type			
<input type="checkbox"/> labsod01dev	Env : production		ONTAP:CVO:VOLUME		
<input type="checkbox"/> urlabsod1	Env : production		ONTAP:CVO:VOLUME		
<input type="checkbox"/> urlabsod01dev	Env : production		ONTAP:CVO:VOLUME		

If you want to add additional search criteria, follow steps 2 through 5 again and then click **Search**.

Search rules

The following rules apply when defining your search:

- After choosing a Resource Type, you can leave the Tag Key field and Tag Value fields empty if you want to list all the resources that have any key and any key value.
- You can choose a single search, or you can define multiple searches to refine the results in the Resource section.
- When defining multiple sets of search criteria:
 - If the criteria for two searches are for different resource types, then this is treated as an "OR" operation and the result shows the resources from both searches. For example, the following search returns all Azure NetApp Files volumes that have the "Environment:demo" tag value and all Cloud Volumes ONTAP volumes that have the "Environment:demo" tag value.

ONTAP:ANF:VOLUME :	<input type="text" value="Environment : demo"/>
ONTAP:CVO:VOLUME :	<input type="text" value="Environment : demo"/>

- If the criteria for two searches are for the same resource type, then this is treated as an "AND" operation and the result shows only the resources that match **both** searches. For example, the following search returns Azure NetApp Files volumes that have both the "Environment:demo" tag value and "Group:Finance" tag value.

ONTAP:ANF:VOLUME :	<input type="text" value="Environment : demo"/>	<input type="text" value="Group : Finance"/>
--------------------	---	--

- If you have defined multiple search criteria and then decide you want to remove one, just click the to remove it from the Search area.

Add tags to existing resources

You can apply tags to a single resource or to multiple resources. The resources could have existing tags or have no current tags.

[See the list of resources that you can tag at this time.](#)

Steps

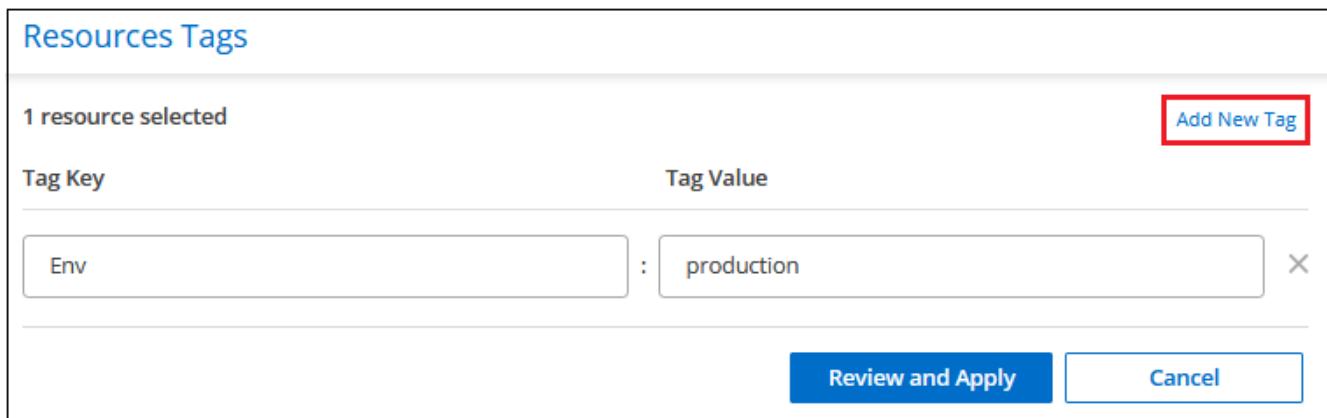
- From the **Tags** tab, create your search criteria and click **Search**.

- Select the resource, or resources, that you want to tag.
 - To select all resources on the page, check the box in the title row (**ID**).
 - To select multiple resources, check the box for each resource (**Volume_1**).
 - To select a single resource, click the  icon for the resource.



<input type="checkbox"/>	ID	Tags	Type
<input type="checkbox"/>	labsod01dev	Env : production 	ONTAP:CVO:VOLUME
<input checked="" type="checkbox"/>	urlabsod1		ONTAP:CVO:VOLUME
<input type="checkbox"/>	urlabsod01dev		ONTAP:CVO:VOLUME

- Click **Manage Tags** and the *Resource Tags* dialog appears. Any existing tags will appear in this dialog.



Resources Tags

1 resource selected

Tag Key	Tag Value
Env	: production

Add New Tag

Review and Apply Cancel

- Click **Add New Tag** and a blank Tag Key and Tag Value row appears in the dialog.
- Enter the tag key and tag value. Add more tags at this time if you want more tags on this resource, and then click **Review and Apply**.
- If your changes look correct in the *Reviewing* page, click **Save** and the new tag is added to the resource, or to all selected resources.

Change tag values for a resource

You can change the tags that are assigned to resources, and you can change the tag value that is applied to an existing tag.

Steps

- From the **Tags** tab, create your search criteria and click **Search**.
- Select the resource, or resources, on which you want to change tags.
- Click **Manage Tags** and the *Resource Tags* dialog appears.
- Enter a new value for the tag value and click **Review and Apply**.
- If your change looks correct in the *Reviewing* page, click **Save** and the tag value is changed for the resource, or for all selected resources.

Delete tags from resources

You can delete a tag key/value pair from a single resource or from multiple resources.

Steps

1. From the **Tags** tab, create your search criteria and click **Search**.
2. Select the resource, or resources, from which you want to remove tags.
3. Click **Manage Tags** and the *Resource Tags* dialog appears.

4. Click the  for the tag key/value pair you want to delete and the row is removed, and then click **Review and Apply**.
5. If your change looks correct in the *Reviewing* page, click **Save** and the tag key/value pair is removed from the resource, or from all selected resources.

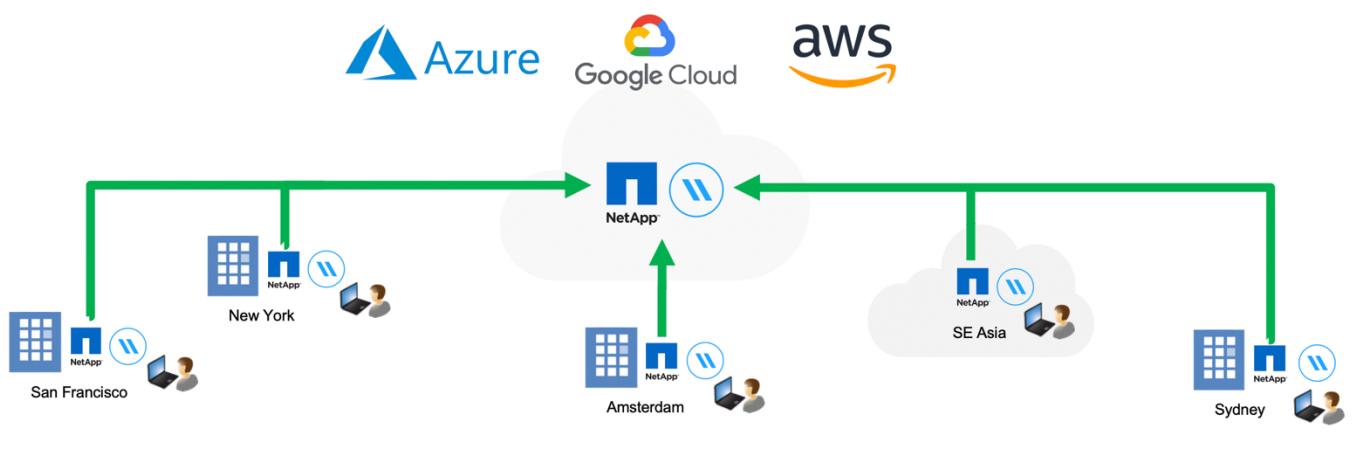
Enable real-time global file sharing

Learn about Global File Cache

NetApp Global File Cache enables you to consolidate silos of distributed file servers into one cohesive global storage footprint in the public cloud. This creates a globally accessible file system in the cloud that all remote locations can use as if they were local.

Overview

Implementing Global File Cache results in a single, centralized storage footprint, versus a distributed storage architecture that requires local data management, backup, security management, storage, and infrastructure footprint in each location.



Features

Global File Cache enables the following features:

- Consolidate and centralize your data into the public cloud and leverage the scalability and performance from enterprise-grade storage solutions
- Create a single set of data for users globally and leverage intelligent file caching to improve global data access, collaboration, and performance
- Rely on a self-sustaining, self-managing cache, and eliminate full data copies and backups. Utilize local file caching for active data and cut storage costs
- Transparent access from branch locations through a global namespace with real time central file locking

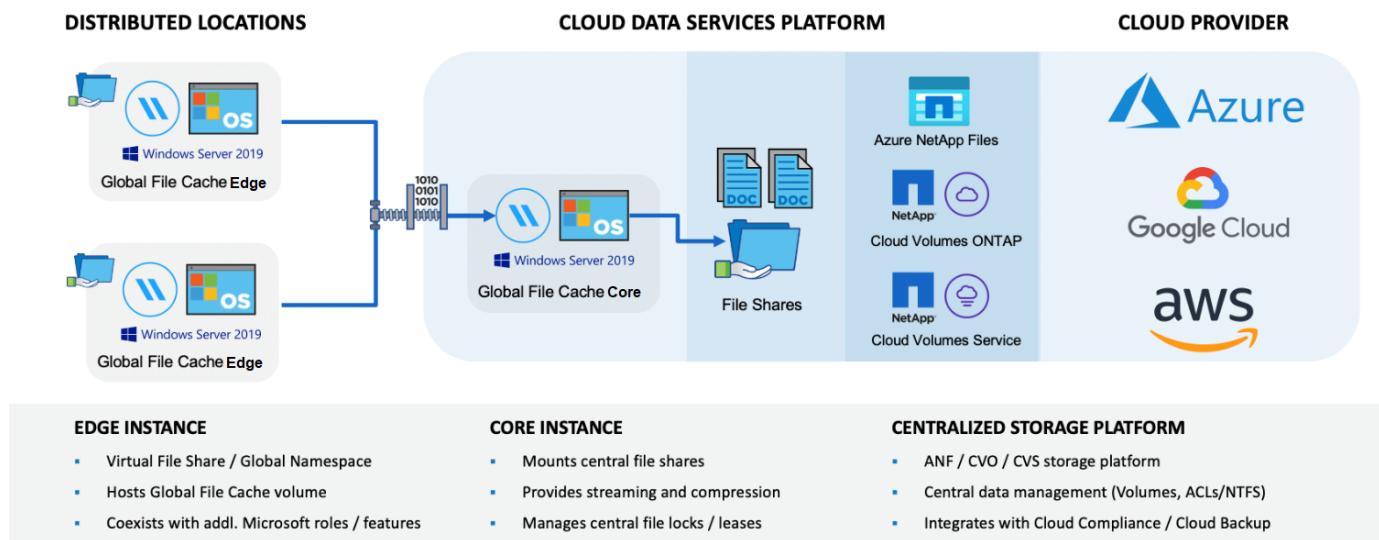
See more about Global File Cache features and use cases [here](#).

Global File Cache components

Global File Cache consists of the following components:

- Global File Cache Management Server
- Global File Cache Core
- Global File Cache Edge (deployed in your remote locations)

The Global File Cache Core instance mounts to your corporate file shares hosted on your backend storage platform of choice (such as Cloud Volumes ONTAP, Cloud Volumes Service, and Azure NetApp Files) and creates the Global File Cache Fabric that provides the ability to centralize and consolidate unstructured data into a single set of data, whether it resides on one or multiple storage platforms in the public cloud.



Supported storage platforms

The supported storage platforms for Global File Cache differ depending on the deployment option you select.

Automated deployment options

Global File Cache is supported with the following types of working environments when deployed using Cloud Manager:

- Cloud Volumes ONTAP in Azure
- Cloud Volumes ONTAP in AWS

This configuration lets you deploy and manage the entire Global File Cache server-side deployment, including Global File Cache Management Server and Global File Cache Core, from within Cloud Manager.

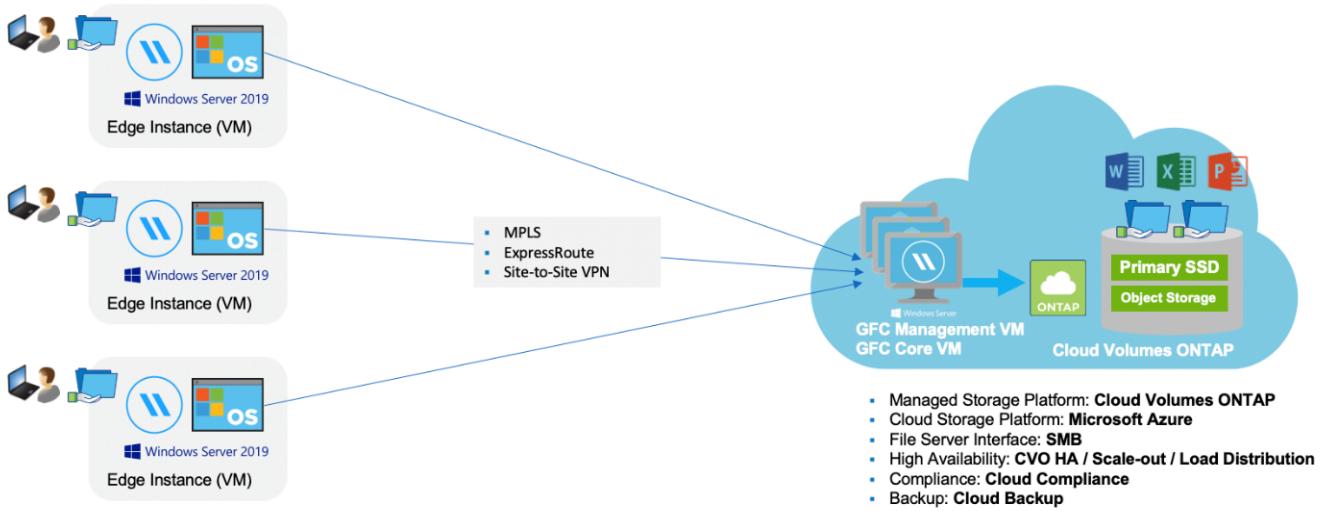
Manual deployment options

Global File Cache configurations are also supported with Cloud Volumes ONTAP, Cloud Volumes Service, Azure NetApp Files, and Amazon FSx for ONTAP systems installed on public cloud storage infrastructure. On-premises solutions are also available on NetApp AFF and FAS platforms. In these installations the Global File Cache server-side components must be configured and deployed manually, not using Cloud Manager.

See the [NetApp Global File Cache User Guide](#) for details.

How Global File Cache works

Global File Cache creates a software fabric that caches active data sets in remote offices globally. As a result, business users are guaranteed transparent data access and optimal performance on a global scale.



The topology referenced in this example is a hub and spoke model, whereby the network of remote offices/locations is accessing one common set of data in the cloud. The key points of this example are:

- Centralized data store:
 - Enterprise public cloud storage platform, such as Cloud Volumes ONTAP
- Global File Cache Fabric:
 - Extension of the central data store to the remote locations
 - Global File Cache Core instance, mounting to corporate file shares (SMB).
 - Global File Cache Edge instances running in each remote location.
 - Presents a virtual file share in each remote location that provides access to central data.
 - Hosts the Intelligent File Cache on a custom-sized NTFS volume (D:\).
- Network configuration:
 - Multiprotocol Label Switching (MPLS), ExpressRoute, or VPN connectivity
- Integration with customer's Active Directory domain services.
- DFS namespace for the use of a global namespace (recommended).

Cost

The cost to use Global File Cache depends on the type of installation you have chosen.

- All installations require that you deploy one or more volumes in the cloud (Cloud Volumes ONTAP, Cloud Volumes Service, or Azure NetApp Files). This results in charges from the selected cloud provider.
- All installations also require that you deploy two or more virtual machines (VMs) in the cloud. This results in charges from the selected cloud provider.
 - Global File Cache Management Server:

In Azure, this runs on a D2s_V3 or equivalent (2 vCPU/8 GB RAM) VM with 127 GB premium SSD

In AWS, this runs on a m4.large or equivalent (2 vCPU/8 GB RAM) instance with 127 GB general purpose SSD

- Global File Cache Core:

In Azure, this runs on a D4s_V3 or equivalent (4 vCPU/16 GB RAM) VM with 127 GB premium SSD

In AWS, this runs on a m4.xlarge or equivalent (4 vCPU/16 GB RAM) instance with 127 GB general purpose SSD

- When installed with Cloud Volumes ONTAP in Azure or AWS (the supported configurations deployed completely through Cloud Manager), there is a charge of \$3,000 per site (for each Global File Cache Edge instance), per year.
- When installed using the manual deployment options the pricing is different. To see a high-level estimate of costs, see [Calculate Your Savings Potential](#) or consult your Global File Cache Solutions Engineer to discuss the best options for your enterprise deployment.

Licensing

Global File Cache includes a software-based License Management Server (LMS), which allows you to consolidate your license management and deploy licenses to all Core and Edge instances using an automated mechanism.

When you deploy your first Core instance in the datacenter or cloud, you can choose to designate that instance as the LMS for your organization. This LMS instance is configured once, connects to the subscription service (over HTTPS) and validates your subscription using the customer ID provided by our support/operations department upon enablement of the subscription. After you have made this designation, you associate your Edge instances with the LMS by providing your customer ID and the IP address of the LMS instance.

When you purchase additional Edge licenses or renew your subscription, our support/operations department updates the license details, for example, the number of sites or subscription end date. After the LMS queries the subscription service, the license details are automatically updated on the LMS instance and will apply to your GFC Core and Edge instances.

See the [NetApp Global File Cache User Guide](#) for additional details about licensing.

Limitations

- The version of Global File Cache supported within Cloud Manager requires that the backend storage platform used as your central storage must be a working environment where you have deployed a Cloud Volumes ONTAP single node or HA pair in Azure or AWS.

Other storage platforms and other cloud providers are not supported at this time using Cloud Manager, but can be deployed using legacy deployment procedures.

These other configurations, for example, Global File Cache using Cloud Volumes ONTAP or Cloud Volumes Service on Microsoft Azure, Google Cloud, or AWS continue to be supported using the legacy procedures. See [Global File Cache overview and onboarding](#) for details.

Release Notes

What's new in version 1.1.0

Released: 9 June 2021

- The "Edge Synchronization" feature has been added. This feature keeps multiple Edges at a remote office

in sync and the data is always cached/warm. When a file is flushed/fetched at one Edge, then the same file on all Edges participating in Edge Sync is updated and cached. The Optimus UI is enhanced for Edge Sync configuration. See section 8.4 in the [NetApp Global File Cache User Guide](#) for details.

- The OpenSSL module is used for secure communication between GFC Edge and GFC Core. This module has been upgraded to latest and more secure version (ver 1.1.1k).
- The GFC License Registration Page displays the number of licenses when activated through a NetApp subscription.

Fixed issues in version 1.1.0

- The Configuration UI (Optimus) now displays the Edge hostname correctly in the Edit Pre-population page.
- An issue in parsing the json response from the NetApp License subscription server during license registration is fixed.
- Pre-population (TappN) logs are enhanced to log the file transfer status of all the files.

Fixed issues in version 1.0.4

Released: 7 April 2021

- A QoS issue has been fixed where QoS parameters were applied to *all* users instead of just the user who surpassed the limits.
- NSS credentials are now validated before being applied to the licensing system.
- Race conditions between the GFC software and certain antivirus software have been fixed.
- A crash no longer occurs when a locking hierarchy is violated during drop lease operations.

What's new in version 1.0.3

Released: 21 March 2021

- The Global File Cache software installer now adds process information to exclude all GFC processes from Windows Defender software on-demand scanning.
- A new Policy Configuration tab has been added to the Configuration Console to add pre-population jobs from the GFC Core.
- Many enhancements have been made so that the software uses less memory while increasing performance and stability.

Fixed issues in version 1.0.3

- The NetApp license registration and GFC subscription-based registration is only allowed by the cloud-based License Management Server (LMS). The on-premises LMS should use legacy licensing.
- The Configuration UI now adds jobs in pre-population when the License Manager Client Service (LMC) uses hostnames instead of IP address for the LMS during configuration.

What's new in version 1.0.2

Released: 6 Jan 2021

- Global File Cache now supports subscription ID based licensing. The Subscription Number is emailed to you after you register your system. See [Licensing Global File Cache](#) for details.
- The requirement for the GFC Core joining the domain has been relaxed. Now a username (part of Backup

Operators group on the Cloud Volumes ONTAP system) is used to configure GFC services.

- The License Manager Server now has an additional HTTPS endpoint to de-register any registration.

Fixed issues in version 1.0.2

- Duplicate GFC icons are no longer displayed on the desktop during upgrade.
- The License Manager Client Service now retains the correct licensing information when the service is restarted.

Before you begin to deploy Global File Cache

There are many requirements you need to be aware of before you begin to deploy Global File Cache in the cloud and in your remote offices.

Global File Cache Core design considerations

Depending on your requirements, you may need to deploy one or multiple Global File Cache Core instances to create the Global File Cache Fabric. The Core instance is designed to act as a traffic cop between your distributed Global File Cache Edge instances and the data center file server resources, for example, file shares, folders, and files.

When you are designing your Global File Cache deployment you need to determine what's right for your environment in terms of scale, availability of resources, and in terms of redundancy. Global File Cache Core can be deployed in the following ways:

- GFC Core stand-alone instance
- GFC Core Load Distributed design (Cold Standby)

See [Sizing guidelines](#) to understand the maximum number of Edge instances and total users that each configuration can support:

Consult your Global File Cache Solutions Engineer to discuss the best options for your enterprise deployment.

Sizing guidelines

There are a few sizing guideline ratios that you need to keep in mind when configuring the initial system. You should revisit these ratios after some usage history has accumulated to make sure you are using the system optimally. These include:

- Global File Cache Edges/Core ratio
- Distributed users/Global File Cache Edge ratio
- Distributed users/Global File Cache Core ratio

Number of Edge Instances per Core Instance

Our guidelines recommend up to 10 Edge instances per Global File Cache Core instance, with a maximum of 20 Edges per Global File Cache Core instance. This is dependent to a significant degree upon the type and mean file size of the most common workload. In some cases, with more common workloads you can add more Edge instances per Core, but in these cases you should contact NetApp Support to correctly size the number of Edge and Core instances depending on the types and sizes of the file sets.



You can leverage multiple Global File Cache Edge and Core instances simultaneously to scale out your infrastructure depending on the requirements.

Number of concurrent users per Edge instance

Global File Cache Edge handles the heavy lifting in terms of caching algorithms and file-level differencing. A single Global File Cache Edge instance can serve up to 400 users per dedicated physical Edge instance, and up to 200 users for dedicated virtual deployments. This is dependent to a significant degree upon the type and mean file size of the most common workload. For larger collaborative file types, guide towards 50% of the maximum users per Global File Cache Edge lower boundary (depending on physical or virtual deployment). For more common Office items with a mean file size <1MB, guide towards the 100% users per Global File Cache Edge upper boundary (depending on physical or virtual deployment).



Global File Cache Edge detects whether it is running on a virtual or physical instance and it will limit the number of SMB connections to the local virtual file share to the maximum of 200 or 400 concurrent connections.

Number of concurrent users per Core instance

The Global File Cache Core instance is extremely scalable, with a recommended concurrent user count of 3,000 users per Core. This is dependent to a significant degree upon the type and mean file size of the most common workload.

Consult your Global File Cache Solutions Engineer to discuss the best options for your enterprise deployment.

Prerequisites

The prerequisites described in this section are for the components installed in the cloud: the Global File Cache Management Server and the Global File Cache Core.

Global File Cache Edge prerequisites are described [here](#).

Cloud Manager instance

When using Cloud Volumes ONTAP for Azure as your storage platform, ensure that Cloud Manager has permissions as shown in the latest [Cloud Manager policy for Azure](#).

Newly created instances will have all the required permissions by default. If you deployed your instance prior to version 3.8.7 (August 3, 2020), then you will need to add these items.

```
"Microsoft.Resources/deployments/operationStatuses/read",
"Microsoft.Insights/Metrics/Read",
"Microsoft.Compute/virtualMachines/extensions/write",
"Microsoft.Compute/virtualMachines/extensions/read",
"Microsoft.Compute/virtualMachines/extensions/delete",
"Microsoft.Compute/virtualMachines/delete",
"Microsoft.Network/networkInterfaces/delete",
"Microsoft.Network/networkSecurityGroups/delete",
"Microsoft.Resources/deployments/delete",
```

Storage platform (volumes)

The back-end storage platform – in this case, your deployed Cloud Volumes ONTAP instance - should present SMB file shares. Any shares that will be exposed through Global File Cache must allow the Everyone group Full Control at the share level, while restricting permissions through NTFS permissions.

If you have not set up at least one SMB file share on the Cloud Volumes ONTAP instance, then you need to have the following information ready so you can configure this information during installation:

- Active Directory domain name, name server IP address, Active Directory admin credentials.
- The name and size of the volume you want to create, the name of the aggregate on which the volume will be created, and the share name.

We recommend that the volume is large enough to accommodate the total data set for the application along with the ability to scale accordingly as the data set grows. If you have multiple aggregates in the working environment, see [Managing existing aggregates](#) to determine which aggregate has the most available space for the new volume.

Global File Cache Management Server

This Global File Cache Management Server requires external access over HTTPS (TCP port 443) to connect to the cloud provider subscription service and to access these URLs:

- <https://talonazuremicroservices.azurewebsites.net>
- <https://talonlicensing.table.core.windows.net>

This port must be excluded from any WAN optimization devices or firewall restriction policies for the Global File Cache software to operate properly.

The Global File Cache Management Server also requires a unique (geographical) NetBIOS name for the instance (such as GFC-MS1).



One Management Server can support multiple Global File Cache Core instances deployed in different working environments. When deployed from Cloud Manager, each working environment has its own separate backend storage and would not contain the same data.

Global File Cache Core

This Global File Cache Core listens on TCP port range 6618-6630. Depending on your firewall or Network Security Group (NSG) configuration you may need to explicitly allow access to these ports through Inbound Port Rules. Also, these ports must be excluded from any WAN optimization devices or firewall restriction policies for the Global File Cache software to operate properly.

The Global File Cache Core requirements are:

- A unique (geographical) NetBIOS name for the instance (such as GFC-CORE1)
- Active Directory domain name
 - Global File Cache instances should be joined to your Active Directory domain.
 - Global File Cache instances should be managed in a Global File Cache specific Organizational Unit (OU) and excluded from inherited company GPOs.
- Service account. The services on this Global File Cache Core run as a specific domain user account. This account, also known as the Service Account, must have the following privileges on each of the SMB

servers that will be associated with the Global File Cache Core instance:

- The provisioned Service Account must be a domain user.

Depending on the level of restrictions and GPOs in the network environment, this account might require domain admin privileges.

- It must have "Run as a Service" privileges.
- The password should be set to "Never Expire".
- The account option "User Must Change Password at Next Logon" should be DISABLED (unchecked).
- It must be a member of the back-end file server Built-in Backup Operators group (this is automatically enabled when deployed through Cloud Manager).

License Management Server

- The Global File Cache License Management Server (LMS) should be configured on a Microsoft Windows Server 2016 Standard or Datacenter edition or Windows Server 2019 Standard or Datacenter edition, preferably on the Global File Cache Core instance in the datacenter or cloud.
- If you require a separate Global File Cache LMS instance, you need to install the latest Global File Cache software installation package on a pristine Microsoft Windows Server instance.
- The LMS instance needs to be able to connect to the subscription service (Azure Services / public internet) using HTTPS (TCP port 443).
- The Core and Edge instances need to connect to the LMS instance using HTTPS (TCP port 443).

Networking (External Access)

The Global File Cache LMS requires external access over HTTPS (TCP port 443) to the following URLs.

- If you are using GFC subscription-based licensing:
 - <https://rest.zuora.com/v1/subscriptions/<subscription-no>>
 - <https://rest.zuora.com/oauth/token>
- If you are using NetApp NSS-based licensing:
 - <https://login.netapp.com>
 - https://login.netapp.com/ms_oauth/oauth2/endpoints
 - https://login.netapp.com/ms_oauth/oauth2/endpoints/oauthservice/tokens
- If you are using NetApp legacy-based licensing:
 - <https://talonazuremicroservices.azurewebsites.net>
 - <https://talonlicensing.table.core.windows.net>

Networking

- Firewall: TCP ports should be allowed between Global File Cache Edge and Core instances.
- Global File Cache TCP Ports: 443 (HTTPS), 6618–6630.
- Network optimization devices (such as Riverbed Steelhead) must be configured to pass-thru Global File Cache specific ports (TCP 6618-6630).

Getting started

You use Cloud Manager to deploy the Global File Cache Management Server and Global File Cache Core software in the working environment.

Enable Global File Cache using Cloud Manager

In this configuration you will deploy the Global File Cache Management Server and Global File Cache Core in the same working environment where you created your Cloud Volumes ONTAP system using Cloud Manager.

Watch [this video](#) to see the steps from start to finish.

Quick start

Get started quickly by following these steps or scroll down to the remaining sections for full details:



Deploy Cloud Volumes ONTAP

Deploy Cloud Volumes ONTAP in Azure or AWS and configure SMB file shares. For more information, see [Launching Cloud Volumes ONTAP in Azure](#) or [Launching Cloud Volumes ONTAP in AWS](#).



Deploy the Global File Cache Management Server

Deploy an instance of the Global File Cache Management Server in the same working environment as the instance of Cloud Volumes ONTAP.



Deploy the Global File Cache Core

Deploy an instance, or multiple instances, of the Global File Cache Core in the same working environment as the instance of Cloud Volumes ONTAP and join it to your Active Directory domain.



License Global File Cache

Configure the Global File Cache License Management Server (LMS) service on a Global File Cache Core instance. You will need your NSS Credentials or a Customer ID and Subscription Number provided by NetApp to activate your subscription.



Deploy the Global File Cache Edge instances

See [Deploying Global File Cache Edge instances](#) to deploy the Global File Cache Edge instances in each remote location. This step is not done using Cloud Manager.

Deploy Cloud Volumes ONTAP as your storage platform

In the current release, Global File Cache supports Cloud Volumes ONTAP deployed in Azure or AWS. For

detailed prerequisites, requirements, and deployment instructions, see [Launching Cloud Volumes ONTAP in Azure](#) or [Launching Cloud Volumes ONTAP in AWS](#).

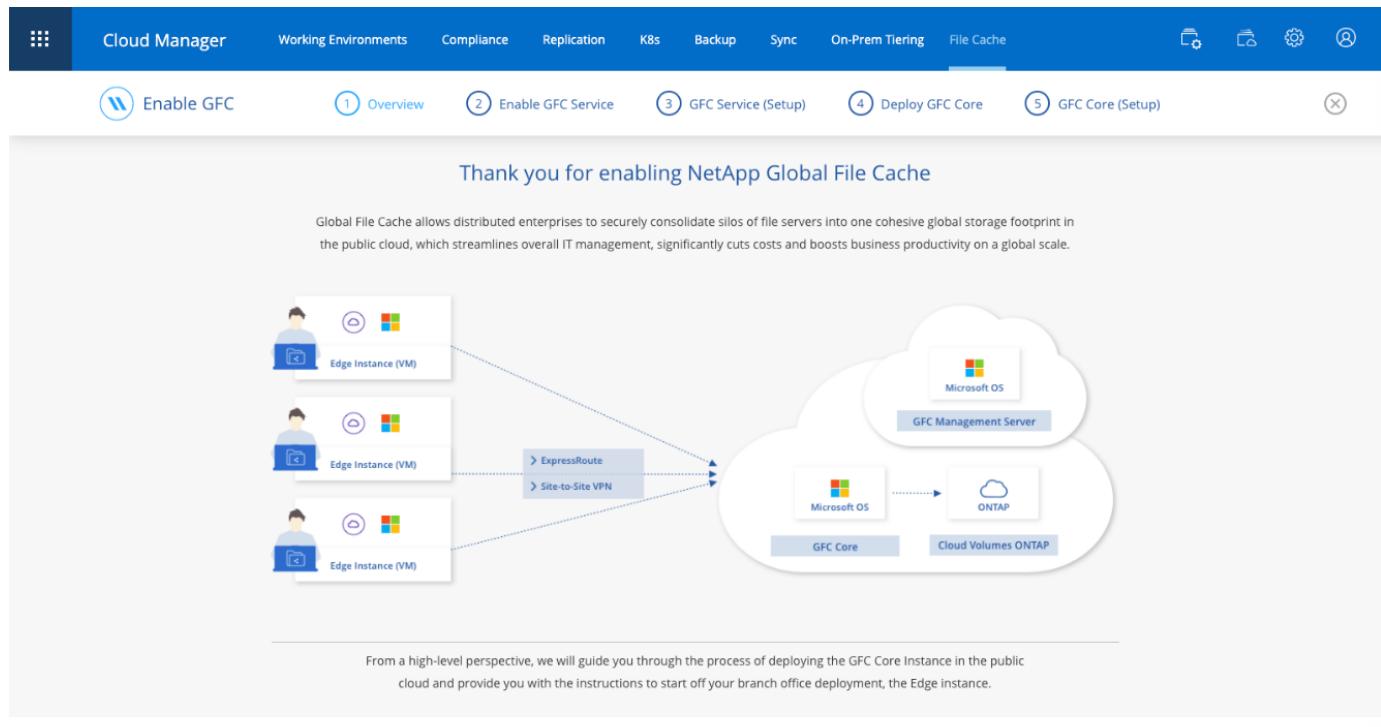
Note the following additional Global File Cache requirement:

- You should configure SMB file shares on the instance of Cloud Volumes ONTAP.

If no SMB file shares are set up on the instance, then you are prompted to configure the SMB shares during the installation of the Global File Cache components.

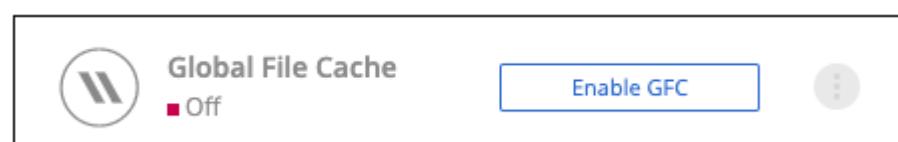
Enable Global File Cache in your working environment

The Global File Cache wizard walks you through the steps to deploy the Global File Cache Management Server instance and the Global File Cache Core instance, as highlighted below.



Steps

1. Select the working environment where you deployed Cloud Volumes ONTAP.
2. In the Services panel, click **Enable GFC**.
3. Read the Overview page and click **Continue**.
4. If no SMB shares are available on the Cloud Volumes ONTAP instance, you are prompted to enter the SMB Server and SMB Share details to create the share now. For details about the SMB configuration, see [Storage platform](#).



3. Read the Overview page and click **Continue**.
 4. If no SMB shares are available on the Cloud Volumes ONTAP instance, you are prompted to enter the SMB Server and SMB Share details to create the share now. For details about the SMB configuration, see [Storage platform](#).
- When finished, click **Continue** to create the SMB share.

SMB Setup

SMB Server	SMB Share
Active Directory Domain <input type="text" value="gfc.netapp.com"/>	Volume Name <input type="text" value="Enter Volume Name"/> Volume Size(GB) <input type="text"/>
Name Server IP Address <input type="text" value="10.0.2.4"/>	Select Aggregate <input type="text" value="Select Aggregate"/>
Active Directory Admin User <input type="text" value="cvoadmin"/>	Share Name <input type="text" value="Enter Share Name"/>
Active Directory Admin Password <input type="password" value="*****"/>	Thin provisioning Enabled 
	Deduplication Enabled 

5. On the Global File Cache Service page, enter the number of Global File Cache Edge instances you plan to deploy, and then make sure your system meets the requirements for Network Configuration and Firewall Rules, Active Directory settings, and Antivirus exclusions. See [Prerequisites](#) for more details.

Enable Global File Cache Service

Licensing Global File Cache:

Once you've completed this deployment process, you will need your NSS Credentials to activate your subscription. If you haven't purchased or received your NetApp Global File Cache licenses, which are available as an Edge-based license, they can be purchased through your NetApp Partner or NetApp Sales Representative.

How many edge instances are you planning to deploy?

10



Before you begin:

Here are the most important requirements for your environment before you can deploy the NetApp Global File Cache solution:

Configure the required Network Configuration and Firewall Rules for Global File Cache



Create a "Service Account" in your Active Directory domain: GFC.NETAPP.COM



Update Antivirus Exclusions for your Windows Server infrastructure by committing the required exclusions to your Antivirus services



For more information on all the solution requirements [Click Here](#)

Continue

6. After you have verified that the requirements have been met, or that you have the information to meet these requirements, click **Continue**.
7. Enter the admin credentials you will use to access to the Global File Cache Management Server VM and click **Enable GFC Service**. For Azure you enter the credentials as a user name and password; for AWS you select the appropriate key pair. You can change the VM/instance name if you want.

Global File Cache Service (Setup)

Information

Subscription Name OCCM Dev

Azure Region eastus

VNet Vnet1

Subnet Subnet2

Resource Group occm_group_eastus

Credentials & Virtual Machine

Local Admin Name

GFCAdmin

Local Admin Password

VM Name

GFC-MS1|

Enable GFC Service

8. After the Global File Cache Management Service is successfully deployed, click **Continue**.
9. For the Global File Cache Core, enter the admin user credentials to join the Active Directory domain, and the service account user credentials. Then click **Continue**.
 - The Global File Cache Core instance must be deployed in the same Active Directory domain as the Cloud Volumes ONTAP instance.
 - The service account is a domain user and it is part of the BUILTIN\Backup Operators group on the Cloud Volumes ONTAP instance.

Deploy Global File Cache Core

Active Directory and Admin Credentials

Provide administrative credentials to join the GFC Core instance to the Active Directory domain

Join Active Directory Domain i

Admin User i

Admin Password i

Account User Credentials

Provide Service Account credentials

Service Account User i

Service Account Password i

Continue

10. Enter the admin credentials you will use to access to the Global File Cache Core VM and click **Deploy GFC Core**. For Azure you enter the credentials as a user name and password; for AWS you select the appropriate key pair. You can change the VM/instance name if you want.

Global File Cache Core (Setup)

Information

Subscription Name	Subscription_1234567891234...
Region	East US Virginia
VNet	VNet_1234567
Subnet	10.0.0.0/24
Resource Group	Resource Group 1

Credentials & Virtual Machine

Local Admin Name

Admin@netapp.com

Local Admin Password

VM Name

GFC-CORE-1234

Local Admin Name & Password are inherited from the Global File Cache Management Service. The Virtual Machine Name is associated to your Cloud Manager Account

Deploy GFC Core

- After the Global File Cache Core is successfully deployed, click **Go to Dashboard**.

The screenshot shows the Global File Cache Management Instance dashboard. It displays a single instance with the following details:

Hostname: www.working-environment-1.com	Status: ON				
IP Address: 141.226.210.219	Region: East US	VNet: VNet1	Subnet: 10.10.10.10/24	RGName: Resource Group	CPU Utilization: 26%

1 Working Environment

Name: Working Environment_1	Type: High Availability	Status: ON	Core Instances: 2	Add Core Instance
-----------------------------	-------------------------	------------	-------------------	-------------------

Instance Core 1 | ON

Hostname: www.working-environment-1.com	IP Address: 141.226.210.219	CPU Utilization: 26%	Network Inbound: 2.5 TB	Network Outbound: 2.5 TB	Deploy GFC Edge
---	-----------------------------	----------------------	-------------------------	--------------------------	-----------------

The Dashboard shows that the Management Server instance and the Core instance are both **On** and working.

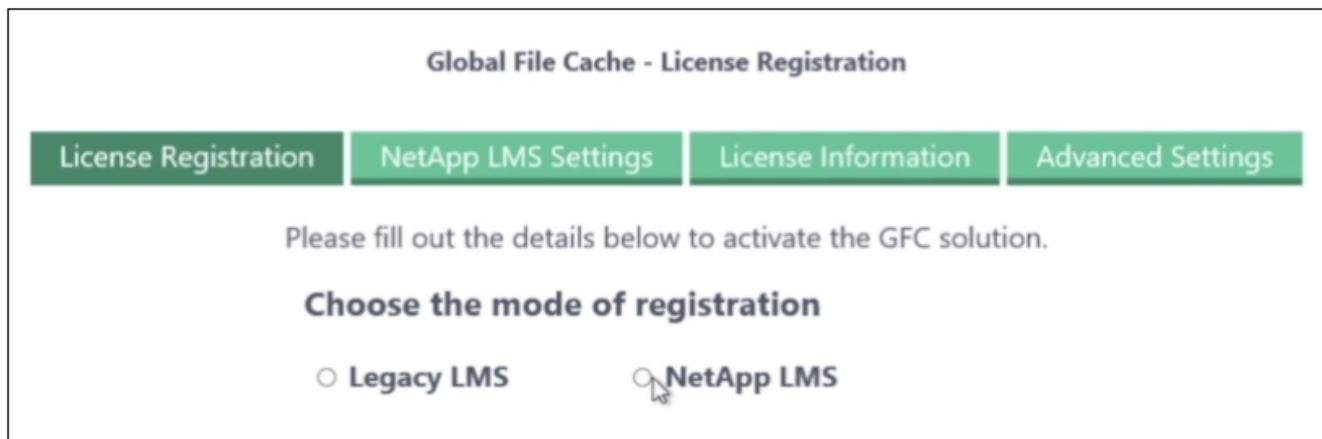
License your Global File Cache installation

Before you can use Global File Cache, you need to configure the Global File Cache License Management Server (LMS) service on a Global File Cache Core instance. You will need your NSS Credentials or a Customer ID and Subscription Number provided NetApp to activate your subscription.

In this example, we will configure the LMS service on a Core instance that you just deployed in the public cloud. This is a one-time process that sets up your LMS service.

Steps

1. Open the Global File Cache License Registration page on the Global File Cache Core (the Core you are designating as your LMS service) using the following URL. Replace <ip_address> with the IP address of the Global File Cache Core:
https://<ip_address>/lms/api/v1/config/lmsconfig.html
2. Click “**Continue to this website (not recommended)**” to continue. A page that allows you to configure the LMS, or check existing license information, is displayed.



Global File Cache - License Registration

License Registration NetApp LMS Settings License Information Advanced Settings

Please fill out the details below to activate the GFC solution.

Choose the mode of registration

Legacy LMS NetApp LMS

3. Choose the mode of registration:
 - “NetApp LMS” is used for customers who have purchased NetApp Global File Cache Edge licenses from NetApp or its certified partners. (Preferred)
 - “Legacy LMS” is used for existing or trial customers who have received a Customer ID through NetApp Support. (This option has been deprecated.)
4. For this example, click **NetApp LMS**, enter your Customer ID (preferably your email address), and click **Register LMS**.

Global File Cache - License Registration

[License Registration](#)

[NetApp LMS Settings](#)

[License Information](#)

[Advanced Settings](#)

Please fill out the details below to activate the GFC solution.

Choose the mode of registration

Legacy LMS

NetApp LMS

Customer Id:

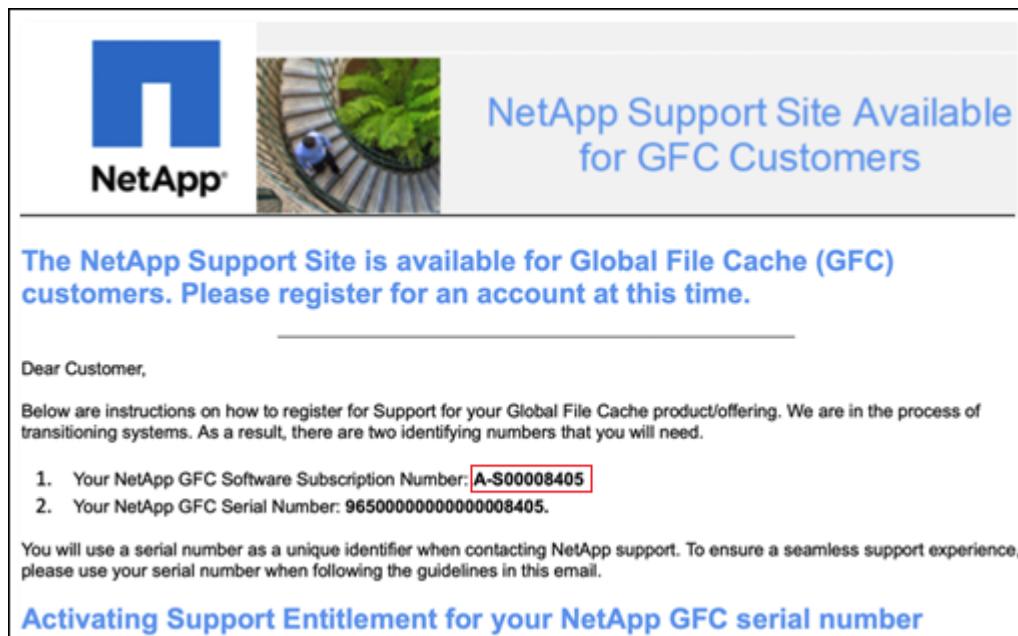
john.doe@company.com

X

* Choose a unique identifier for your GFC deployment, preferably your email address

[REGISTER LMS](#)

5. Check for a confirmation email from NetApp that includes your GFC Software Subscription Number and Serial Number.



The NetApp Support Site is available for Global File Cache (GFC) customers. Please register for an account at this time.

Dear Customer,

Below are instructions on how to register for Support for your Global File Cache product/offering. We are in the process of transitioning systems. As a result, there are two identifying numbers that you will need.

1. Your NetApp GFC Software Subscription Number: **A-S00008405**
2. Your NetApp GFC Serial Number: **965000000000000000008405**.

You will use a serial number as a unique identifier when contacting NetApp support. To ensure a seamless support experience, please use your serial number when following the guidelines in this email.

Activating Support Entitlement for your NetApp GFC serial number

6. Click the **NetApp LMS Settings** tab.
7. Select **GFC License Subscription**, enter your GFC Software Subscription Number, and click **Submit**.

Global File Cache - License Registration

License Registration	NetApp LMS Settings	License Information	Advanced Settings
----------------------	---------------------	---------------------	-------------------

NSS Credentials GFC License Subscription

GFC License Subscription: X

SUBMIT

You will see a message that your GFC License Subscription was registered successfully and activated for the LMS instance. Any subsequent purchases will automatically be added to the GFC License Subscription.

8. Optionally, you can click the **License Information** tab to view all your GFC license information.

What's Next?

If you have determined that you need to deploy multiple Global File Cache Cores to support your configuration, click **Add Core Instance** from the Dashboard and follow the deployment wizard.

After you have completed your Core deployment, you need to [deploy the Global File Cache Edge instances](#) in each of your remote offices.

Deploy additional Core instances

If your configuration requires more than one Global File Cache Core to be installed because of a large number of Edge instances, you can add another Core to the working environment.

When deploying Edge instances, you will configure some to connect to the first Core and others to the second Core. Both Core instances access the same backend storage (your Cloud Volumes ONTAP instance) in the working environment.

1. From the Global File Cache Dashboard, click **Add Core Instance**.

1 Working Environment									
Name	High Availability Type	■ ON Status	2 Core Instances	Add Core Instance					
Instance Core 1 ■ ON <table border="1" style="width: 100%; border-collapse: collapse;"> <tr> <td style="width: 25%; padding: 5px;"> www.working-environment-1.com Hostname </td> <td style="width: 25%; padding: 5px;"> 141.226.210.219 IP Address </td> <td style="width: 25%; padding: 5px;"> 26% CPU Utilization </td> <td style="width: 25%; padding: 5px;"> 2.5 TB Network Inbound </td> <td style="width: 25%; padding: 5px;"> 2.5 TB Network Outbound </td> </tr> </table> <div style="text-align: right; margin-top: 10px;"> Deploy GFC Edge </div>					www.working-environment-1.com Hostname	141.226.210.219 IP Address	26% CPU Utilization	2.5 TB Network Inbound	2.5 TB Network Outbound
www.working-environment-1.com Hostname	141.226.210.219 IP Address	26% CPU Utilization	2.5 TB Network Inbound	2.5 TB Network Outbound					

2. Enter the admin user credentials to join the Active Directory domain, and the service account user credentials. Then click **Continue**.

- The Global File Cache Core instance must be in the same Active Directory domain as the Cloud Volumes ONTAP instance.
- The service account is a domain user and it is part of the BUILTIN\Backup Operators group on the Cloud Volumes ONTAP instance.

Deploy Global File Cache Core

Active Directory and Admin Credentials

Provide administrative credentials to join the GFC Core instance to the Active Directory domain

Join Active Directory Domain

Active Directory Domain

Admin User

Enter Admin User

Admin Password

Enter Admin Password

Account User Credentials

Provide Service Account credentials

Service Account User

Enter Service Account User

Service Account Password

Enter Service Account Password

Continue

3. Enter the admin credentials you will use to access to the Global File Cache Core VM and click **Deploy GFC Core**. For Azure you enter the credentials as a user name and password; for AWS you select the appropriate key pair. You can change the VM name if you want.

Global File Cache Core (Setup)

Information

Subscription Name	Subscription_1234567891234...
Region	East US Virginia
VNet	VNet_1234567
Subnet	10.0.0.0/24
Resource Group	Resource Group 1

Credentials & Virtual Machine

Local Admin Name

Admin@netapp.com

Local Admin Password

VM Name

GFC-CORE-1234

Local Admin Name & Password are inherited from the Global File Cache Management Service. The Virtual Machine Name is associated to your Cloud Manager Account

Deploy GFC Core

- After the Global File Cache Core is successfully deployed, click **Go to Dashboard**.

The screenshot shows the 'Working Environment' section of the dashboard. It displays two entries:

Name	Type	Status	Core Instances	Add Core Instance
Working Environment_1	High Availability	ON	2	Add Core Instance

Under 'Working Environment_1', there are two instances:

Hostname	IP Address	CPU Utilization	Network Inbound	Network Outbound	Action
www.working-environment-1.com	141.226.210.219	26%	2.5 TB	2.5 TB	Deploy GFC Edge
Instance Core 1	141.226.210.219	26%	2.5 TB	2.5 TB	Deploy GFC Edge

The second instance is highlighted with a blue border.

The Dashboard reflects the second Core instance for the working environment.

Before you begin to deploy Global File Cache Edge instances

There are many requirements you need to be aware of before you begin to install Global File Cache Edge software in your remote offices.

Download required resources

Download the Global File Cache Virtual Templates you are planning to use in your branch offices, the software installation package, and additional reference documentation:

- Windows Server 2016 Virtual Template:

[Windows Server 2016 .OVA including NetApp GFC \(VMware VSphere 6.5+\)](#)
[Windows Server 2016 .VHDX including NetApp GFC \(Microsoft Hyper-v\)](#)

- Windows Server 2019 Virtual Template:

[Windows Server 2019 .OVA including NetApp GFC \(VMware VSphere 6.5+\)](#)
[Windows Server 2019 .VHDX including NetApp GFC \(Microsoft Hyper-v\)](#)

- Global File Cache Edge Software:

[NetApp GFC Software Installation Package \(.EXE\)](#)

- Global File Cache SCOM (Microsoft Systems Center Operations Manager) Package:

[NetApp GFC SCOM Management Package \(.ZIP\)](#)

- Global File Cache Documentation:

[NetApp Global File Cache User Guide \(.PDF\)](#)

Designing and deploying Global File Cache Edge

Depending on your requirements, you might need to deploy one or multiple Global File Cache Edge instances based on the concurrent user sessions in a branch office. The Edge instance presents the virtual file share to the end users within the branch office, which has been transparently extended from the associated Global File Cache Core instance. The Global File Cache Edge should contain a D:\ NTFS volume, which contains the cached files within the branch office.

 For the Global File Cache Edge, it is important to understand the [sizing guidelines](#). This will assist you in making the correct design for your Global File Cache deployment. You would also need to determine what's right for your environment in terms of scale, availability of resources, and in terms of redundancy.

Global File Cache Edge instance

When deploying a Global File Cache Edge instance, you need to provision a single VM, either by deploying Windows Server 2016 Standard or Datacenter Edition, or Windows Server 2019 Standard or Datacenter Edition, or using the Global File Cache .OVA or .VHD template, which includes the Windows Server operating system of choice and Global File Cache software.

Quick steps

1. Deploy the Global File Cache Virtual Template, or Windows Server 2016 VM, or Windows Server 2019 Standard or Datacenter edition.
2. Ensure the VM is connected to the network, joined to the domain, and accessible through RDP.

3. Install the latest Global File Cache Edge software.
4. Identify the Global File Cache Management Server and Core instance.
5. Configure the Global File Cache Edge instance.

Global File Cache Edge requirements

Global File Cache Edge is designed to function across all platforms supporting Windows Server 2016 and 2019, bringing simplified IT to corporate remote offices and beyond. Critically, Global File Cache can be deployed on your existing hardware infrastructure, virtualization, or hybrid/public cloud environments in almost every case if they meet a few base-level requirements.

Global File Cache Edge requires the following hardware and software resources to function optimally. For more information about overall sizing guidelines, see [Sizing guidelines](#).

Hardened server appliance

The Global File Cache installation package creates a hardened software appliance on any Microsoft Windows Server instance. *Do Not Uninstall* the Global File Cache Package. Uninstalling Global File Cache will impact the functionality of the server instance and might require a full rebuild of the server instance.

Physical hardware requirements

- Minimum 4 CPU cores
- Minimum 16 GB RAM
- Dedicated single or redundant 1 Gbps NIC
- 10k RPM SAS HDD or SSD (preferred)
- RAID controller with write-back caching functionality enabled

Virtual deployment requirements

Hypervisor platforms are known to be subject to performance degradation from a storage subsystem perspective (for example, latency). For optimal performance using Global File Cache, a physical server instance with SSD is recommended.

For best performance in virtual environments, in addition to the physical host requirements, the following requirements and resource reservations must be met:

Microsoft Hyper-V 2012 R2 and later:

- Processor (CPU): CPUs must be set as **Static**: Minimum: 4 vCPU cores.
- Memory (RAM): Minimum: 16 GB set as **Static**.
- Hard-disk provisioning: Hard Disks must be configured as **Fixed Disk**.

VMware vSphere 6.x and later:

- Processor (CPU): Reservation of CPU Cycles must be set. Minimum: 4 vCPU cores @ 10000 MHz.
- Memory (RAM): Minimum: Reservation of 16 GB.
- Hard-disk provisioning:
 - Disk Provisioning must be set as **Thick Provisioned Eager Zeroed**.

- Hard Disk Shares must be set to **High**.
- Devices.hotplug must be set to **False** using the vSphere Client to prevent Microsoft Windows from presenting Global File Cache drives as removable.
- Networking: Network Interface must be set to **VMXNET3** (may require VM Tools).

Global File Cache runs on Windows Server 2016 and 2019, hence the virtualization platform needs to support the operating system, as well as integration with utilities enhancing the performance of the VM's guest operating system and management of the VM, such as VM Tools.

Partition sizing requirements

- C:\ - minimum 250 GB (system/boot volume)
- D:\ - minimum 1 TB (separate data volume for Global File Cache Intelligent File Cache*)

*Minimum size is 2x the active data set. The cache volume (D:\) can be extended and is only restricted by the limitations of the Microsoft Windows NTFS file system.

Global File Cache Intelligent File Cache disk requirements

Disk Latency on the Global File Cache Intelligent File Cache disk (D:\) should deliver < 0.5ms average I/O disk latency and 1MiBps throughput per concurrent user.

For more information, see the [NetApp Global File Cache User Guide](#).

Networking

- Firewall: TCP ports should be allowed between the Global File Cache Edge and Management Server and Core instances.

Global File Cache TCP Ports: 443 (HTTPS - LMS), 6618 – 6630.

- Network optimization devices (such as Riverbed Steelhead) must be configured to pass-thru Global File Cache specific ports (TCP 6618-6630).

Client workstation and application best practices

Global File Cache transparently integrates into customer's environments, allowing users to access centralized data using their client workstations, running enterprise applications. Using Global File Cache, data is accessed through a direct drive mapping or through a DFS namespace. For more information about the Global File Cache Fabric, Intelligent File Caching, and key aspects of the software, consult the [Before you begin to Deploy Global File Cache](#) section.

To ensure an optimal experience and performance, it is important to comply with the Microsoft Windows Client requirements and best practices as outlined in the Global File Cache User Guide. This applies to all versions of Microsoft Windows.

For more information, see the [NetApp Global File Cache User Guide](#).

Firewall and Antivirus best practices

While Global File Cache makes a reasonable effort to validate that the most common antivirus application suites are compatible with Global File Cache, NetApp cannot guarantee and is not responsible for any incompatibilities or performance issues caused by these programs, or their associated updates, service packs,

or modifications.

Global File Cache does not recommend the installation nor application of monitoring or antivirus solutions on any Global File Cache enabled instance (Core or Edge). Should a solution be installed, by choice or by policy, the following best practices and recommendations must be applied. For common antivirus suites, see Appendix A in the [NetApp Global File Cache User Guide](#).

Firewall settings

- Microsoft firewall:
 - Retain firewall settings as default.
 - Recommendation: Leave Microsoft firewall settings and services at the default setting of OFF, and not started for standard Global File Cache Edge instances.
 - Recommendation: Leave Microsoft firewall settings and services at the default setting of ON, and started for Edge instances that also run the Domain Controller role.
- Corporate firewall:
 - Global File Cache Core instance listens on TCP ports 6618-6630, ensure that Global File Cache Edge instances can connect to these TCP ports.
 - Global File Cache instances require communications to the Global File Cache Management Server on TCP port 443 (HTTPS).
- Network optimization solutions/devices must be configured to pass-thru Global File Cache specific ports.

Antivirus best practices

This section helps you to understand the requirements when running antivirus software on a Windows Server instance running Global File Cache. Global File Cache has tested most commonly used antivirus products including Cylance, McAfee, Symantec, Sophos, Trend Micro, Kaspersky and Windows Defender for use in conjunction with Global File Cache.



Adding antivirus to an Edge appliance can introduce a 10–20% impact on user performance.

For more information, see the [NetApp Global File Cache User Guide](#).

Configure exclusions

Antivirus software or other third-party indexing or scanning utilities should never scan drive D:\ on the Edge instance. These scans of Edge server drive D:\ will result in numerous file open requests for the entire cache namespace. This will result in file fetches over the WAN to all file servers being optimized at the data center. WAN connection flooding and unnecessary load on the Edge instance will occur resulting in performance degradation.

In addition to the D:\ drive, the following Global File Cache directory and processes should generally be excluded from all antivirus applications:

- C:\Program Files\TalonFAST\
- C:\Program Files\TalonFAST\Bin\LMClientService.exe
- C:\Program Files\TalonFAST\Bin\LMServerService.exe
- C:\Program Files\TalonFAST\Bin\Optimus.exe

- C:\Program Files\TalonFAST\Bin\tafsexport.exe
- C:\Program Files\TalonFAST\Bin\tafsutils.exe
- C:\Program Files\TalonFAST\Bin\tapp.exe
- C:\Program Files\TalonFAST\Bin\TappN.exe
- C:\Program Files\TalonFAST\Bin\FTLSummaryGenerator.exe
- C:\Program Files\TalonFAST\Bin\RFASTSetupWizard.exe
- C:\Program Files\TalonFAST\Bin\TService.exe
- C:\Program Files\TalonFAST\Bin\tum.exe
- C:\Program Files\TalonFAST\FastDebugLogs\
- C:\Windows\System32\drivers\tfast.sys
- \\?\TafsMtPt:\ or \\?\TafsMtPt*
- \Device\TalonCacheFS\
- \\?\GLOBALROOT\Device\TalonCacheFS\
- \\?\GLOBALROOT\Device\TalonCacheFS*

NetApp Support policy

Global File Cache instances are designed specifically for Global File Cache as the primary application running on a Windows Server 2016 and 2019 platform. Global File Cache requires priority access to platform resources, for example, disk, memory, network interfaces, and can place high demands on these resources. Virtual deployments require memory/CPU reservations and high-performance disks.

- For branch office deployments of Global File Cache, supported services and applications on the server running Global File Cache are limited to:
 - DNS/DHCP
 - Active Directory domain controller (Global File Cache must be on a separate volume)
 - Print services
 - Microsoft System Center Configuration Manager (SCCM)
 - Global File Cache approved client-side system agents and anti-virus applications
- NetApp Support and maintenance applies only to Global File Cache.
- Line of business productivity software, which are typically resource intensive, for example, database servers, mail servers, and so on, are not supported.
- The customer is responsible for any non-Global File Cache software which might be installed on the server running Global File Cache:
 - If any third-party software package causes software or resource conflicts with Global File Cache or performance is compromised, Global File Cache's support organization might require the customer to disable or remove the software from the server running Global File Cache.
 - It is the customer's responsibility for all installation, integration, support, and upgrade of any software added to the server running the Global File Cache application.
- Systems management utilities/agents such as antivirus tools and licensing agents might be able to coexist. However, except for the supported services and applications listed above, these applications are not

supported by Global File Cache and the same guidelines as above must still be followed:

- It is the customer’s responsibility for all installation, integration, support, and upgrade of any software added.
- If a customer does install any third-party software package that causes, or is suspected to be causing, software or resource conflicts with Global File Cache or performance is compromised, there might be a requirement by Global File Cache’s support organization to disable/remove the software.

Deploy Global File Cache Edge instances

After you have verified that your environment meets all the requirements, you install Global File Cache Edge software in each remote office.

Before you begin

To complete Global File Cache Edge configuration tasks, you need the following information:

- Static IP addresses for each Global File Cache instance
- Subnet mask
- Gateway IP address
- The FQDN you wish to assign to each Global File Cache server
- The DNS suffix (optional)
- The user name and password of an administrative user in the domain
- The FQDN and/or IP address of the associated Core servers
- A volume to be used as the Intelligent File Cache. It is recommended this be at least 2x the size of the active dataset. This should be formatted as NTFS and assigned as D:\.

Commonly used TCP ports

There are several TCP ports used by Global File Cache services. It is mandatory that the devices can communicate on these ports and they be excluded from any WAN optimization devices or firewall restriction policies:

- Global File Cache Licensing TCP Port: 443
- Global File Cache TCP Ports: 6618-6630

Deploy the Global File Cache Virtual Template

The virtual template (.OVA and .VHD) images contain the latest release of the Global File Cache software. If you are deploying Global File Cache using the .OVA or .VHD virtual machine (VM) template, follow the steps as outlined in this section. It is assumed that you understand how to deploy the .OVA or .VHD template on the designated hypervisor platform.

Ensure that VM preferences, including resource reservations, are in line with the requirements as outlined in [Virtual deployment requirements](#).

Steps

1. Extract the package from the template you downloaded.

2. Deploy the virtual template. Refer to the following videos before you start the deployment:
 - [Deploy the Virtual Template on VMware](#)
 - [Deploy the Virtual Template on Hyper-V](#)
3. After the Virtual Template has been deployed, and you have configured the VM settings, start the VM.
4. During initial boot, when the Windows Server 2016 or 2019 operating system is preparing for first use, complete the out-of-the-box experience by installing the correct drivers and installing the necessary components for the respective hardware.
5. When the base install of the Global File Cache Edge instance has been completed, the Windows Server 2016 or 2019 operating system will guide you through an initial configuration wizard to configure operating system specifics such as localization and product key.
6. After the initial configuration wizard has completed, log in locally to the Windows Server 2016 or 2019 operating system with the following credentials:
 - User name: **FASTAdmin**
 - Password: **Tal0nFAST!**
7. Configure your Windows Server VM, join to the organization's Active Directory domain, and proceed to the Global File Cache Edge configuration section.

Configure the Global File Cache Edge instance

The Global File Cache Edge instance connects to a Global File Cache Core to provide users at the branch office access to data center file server resources.



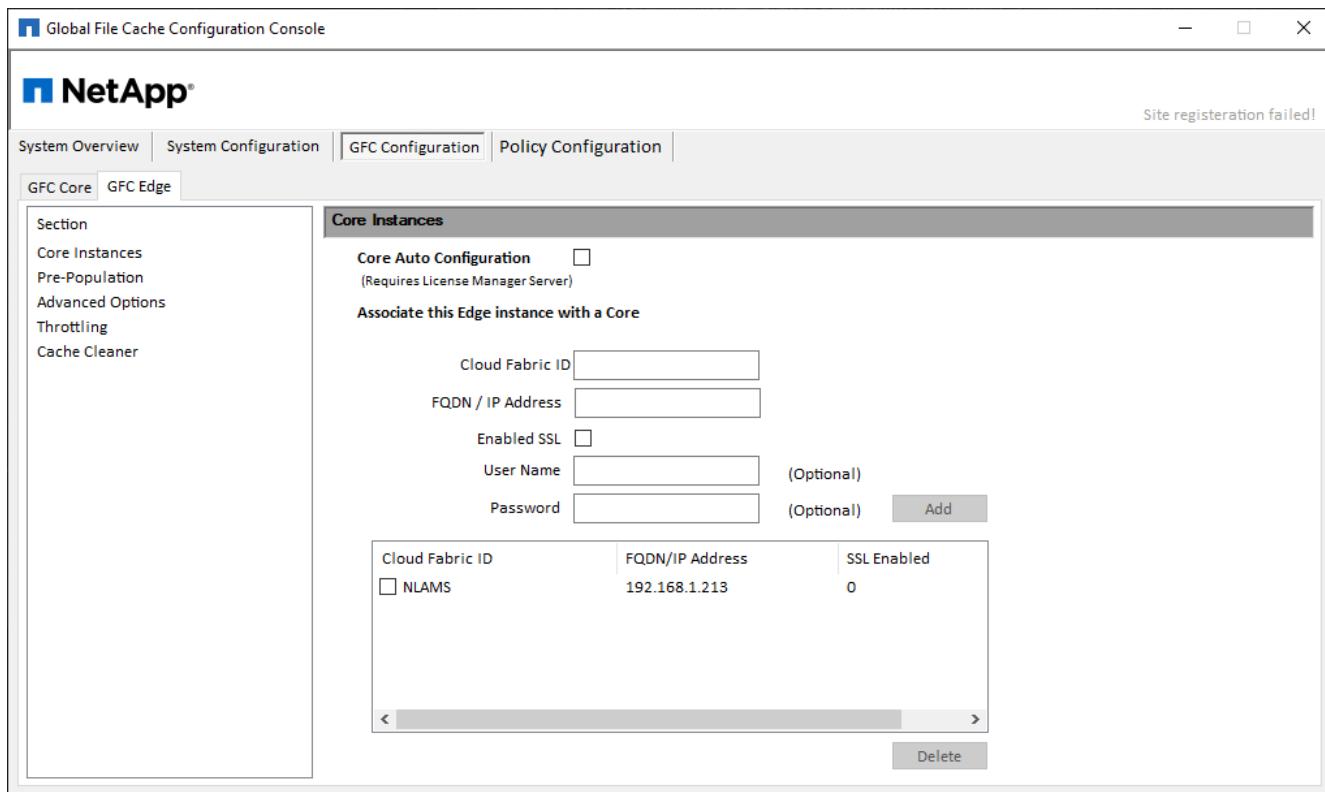
The Edge instance must be licensed as part of your Cloud Volumes ONTAP deployment prior to beginning the configuration. See [Licensing](#) for more information about licensing.

If your configuration requires more than one Global File Cache Core to be installed because of a large number of Edge instances, you will configure some Edge instances to connect to the first Core and others to connect to the second Core. Make sure you have the FQDN or IP address, and other required information, for the correct Core instance.

To configure the Edge instance, complete the following steps:

Steps

1. Click **Perform** next to the unchecked Core Configuration step listed in the "Edge Configuration Steps" section of the Initial Configuration assistant. This opens a new tab, GFC Edge, and shows the section *Core Instances*.
2. Provide the **Cloud Fabric ID** of the Global File Cache Core server. The Cloud Fabric ID is typically the NetBIOS name or the geographical location of the backend file server.
3. Provide the **FQDN/IP Address** of the Global File Cache Core server:
 - a. (Optional) Check the **SSL** box to enable SSL support for enhanced encryption from the Edge to the Core.
 - b. Enter the User Name and Password, which are the credentials of the Service Account used on the Core.
4. Click **Add** to confirm the addition of the Global File Cache Core appliance. A confirmation box will appear. Click **OK** to dismiss it.



Update Global File Cache Edge software

Global File Cache frequently releases updates to the software, either patches, enhancements, or new features/functionality. Although the virtual template (.OVA and .VHD) images contain the latest release of the Global File Cache software, it is possible that a newer version is available on the NetApp Support Download portal.

Ensure that your Global File Cache instances are up to date with the latest version.



This software package can also be used for pristine installations on Microsoft Windows Server 2016 Standard or Datacenter edition, or Windows Server 2019 Standard or Datacenter edition, or used as part of your upgrade strategy.

Below you can find the steps required to update the Global File Cache installation package:

Steps

1. After saving the latest installation package to the desired Windows Server instance, double-click it to run the installation executable.
2. Click **Next** to continue the process.
3. Click **Next** to continue.
4. Accept the Licensing Agreement and click **Next**.
5. Select the desired Installation Destination Location.

NetApp recommends that you use the default installation location.

6. Click **Next** to continue.
7. Select the Start Menu Folder.

8. Click **Next** to continue.
9. Verify your installation selections and click **Install** to begin the installation.

The installation process will start.

10. After the installation has completed, reboot the server when prompted.

What's Next?

For details about Global File Cache Edge advanced configuration, see the [NetApp Global File Cache User Guide](#).

End-user training

You will want to train your users on the best practices for accessing the shared files through Global File Cache.

This is the final phase of the Global File Cache deployment, the end-user implementation phase.

In order to prepare and streamline the end user on-boarding process, use the email template below that will help you to educate end users on what it means to work in a "central data" environment. This will help your users leverage all of the benefits of the Global File Cache solution. We have also published a video that can be shared to "train" users where needed.

Customize and forward the following resources to end users to prepare them for roll-out:

- User Training video
[End user training video](#)
- Email Template
[Mac Email Template \(.emltpl\)](#)

[Windows Email Template \(.msg\)](#)
- Onboarding Communications
[Word Document \(.docx\)](#)

See Chapter 13 in the [NetApp Global File Cache User Guide](#) for additional material.

Additional information

Use the following links to learn more about Global File Cache and other NetApp products:

- Global File Cache FAQ
 - See a list of frequently asked questions and answers [here](#)
- [NetApp Global File Cache User Guide](#)
- NetApp Product Documentation
 - See additional documentation for NetApp cloud products [here](#)
 - See additional documentation for all NetApp products [here](#)
- Customer support for Global File Cache users with Cloud Volumes ONTAP is available through these channels:

- Guided Problem Solving, Case Management, Knowledgebase, Downloads, Tools, and more go [here](#)
- Login to the NetApp Support at <https://mysupport.netapp.com> with your NSS credentials
- For immediate assistance for a P1 issue call: +1 856.481.3990 (Option 2)
- Customer support for Global File Cache users utilizing Cloud Volumes Services and Azure NetApp Files is available through standard support from your provider. Please contact Google Customer Support or Microsoft Customer Support respectively.

Optimize cloud compute costs

Learn about the Compute service

By leveraging [Spot's Cloud Analyzer service](#), Cloud Manager can provide a high-level cost analysis of your cloud compute spending and identify potential savings.

Cloud Analyzer is a cloud infrastructure management solution that uses advanced analytics to provide visibility and insights into your cloud costs. It shows you where you can optimize those costs and lets you implement that optimization using Spot's portfolio of continuous optimization products in just a few clicks.

Features

- A cost analysis that shows current cost for the month, projected monthly costs, and missed savings
- A view of spend efficiency by account, including the estimated additional savings
- A link to Spot's Cloud Analyzer for more in-depth details about the spending for all accounts

Supported cloud providers

This service is supported with AWS.

Cost

There's no cost to use this service through Cloud Manager.

How Cloud Analyzer works with Cloud Manager

At a high-level, Cloud Analyzer integration with Cloud Manager works like this:

1. You click **Compute** and connect your AWS management account.
2. NetApp configures your environment as follows:
 - a. Creates an organization in the Spot platform.
 - b. Sends an email welcoming you to Spot.

You can log in to the Spot service using the same single-sign on credentials that you use with Cloud Central and Cloud Manager.

- c. Cloud Analyzer starts processing your AWS account data.
3. In Cloud Manager, the Compute page refreshes and you use the information to gain insights on past, current, and future cloud costs.
 4. You click **Get Full Analysis** at any time to go to Spot's Cloud Analyzer, which provides a full analysis of your cloud spend and savings opportunities.

Data security

Cloud Analyzer data is encrypted at rest and no credentials are stored for any account.

Start optimizing your cloud compute costs

Connect your AWS account and then view the analysis to start optimizing your cloud compute costs.

Connect Cloud Analyzer to your AWS account

Click **Compute** and connect your AWS payer account.

Steps

1. Click **Compute**.
2. Click **Add AWS Credentials to Start**.
3. Follow the steps on the page to connect your AWS account:
 - a. Log in to your AWS management account.
 - b. Set up cost and usage reports on the AWS account.
 - c. Run the CloudFormation template.
 - d. Paste the Spot RoleARN.

[View more details about these steps.](#)

Connect your AWS Account to Optimize Costs

Connecting your billing data will allow Cloud Analyzer to access your Cost and Usage data.

Step 1

Log in to your AWS Master Payer account. [Log in](#)

Step 2

Set up your Cost and Usage Reports on your AWS account.
[\(Learn How\)](#) or skip this if the report is already enabled.)

Enter the bucket name where the report is located:

Step 3

Open CloudFormation with Spot template. [Run Template](#)

Under capabilities, mark "I acknowledge that AWS CloudFormation might create IAM resources" and click 'Create'.

Step 4

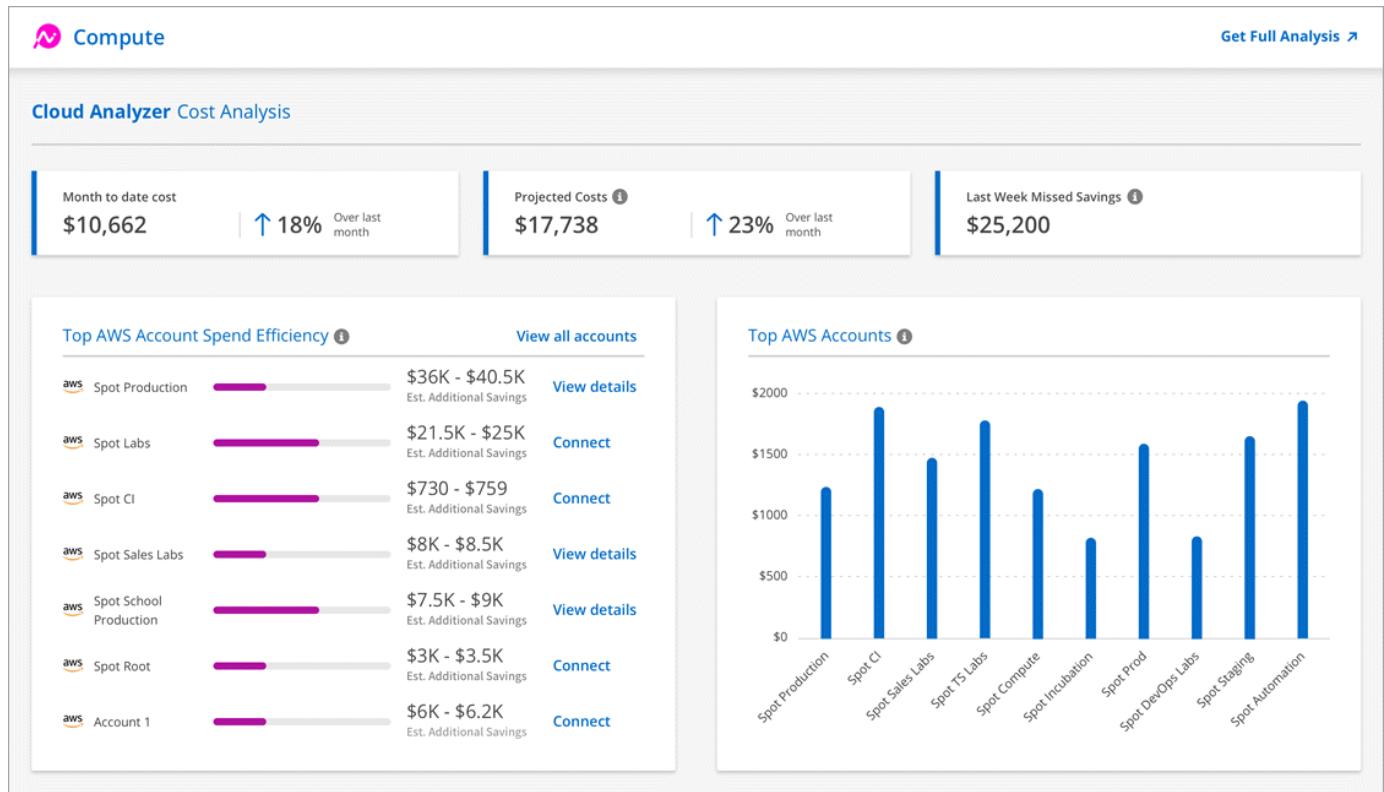
Copy the Spot RoleARN from the Output tab and paste below.

Result

Cloud Analyzer starts processing your AWS account data. If you have multiple accounts, Cloud Analyzer starts with read-only capabilities for all linked accounts under the management account. If you want to get more details about the potential savings for those accounts, then you'll need to connect them, as well. You can find more details about that process in the section below.

Analyze your compute costs

After Cloud Analyzer processes your account data, the Compute tab shows you insights on past, current, and future cloud costs.



Month to date cost

The total cost of your workloads from the beginning of the current month to present.

Projected Costs

The forecasted cost at the end of the month based on analysis of your usage pattern.

Last Week Missed Savings

Savings that could have been achieved in the previous seven days using optimization of spot instances and reservations.

Top AWS Account Spend Efficiency

The top 10 accounts according to the greatest amount of estimated additional savings.

Each account is assigned an efficiency score based on current and additional potential savings. The estimated additional savings indicates how much can be further saved by leveraging the use of spot and reserved instances.

You can take the following actions to further optimize your accounts:

- **View details:** View your cost optimization opportunities by going to Spot's Cloud Analyzer.
- **Connect:** Connect an account that is not yet managed. You will be directed to the wizard that connects the account.

Top AWS Accounts

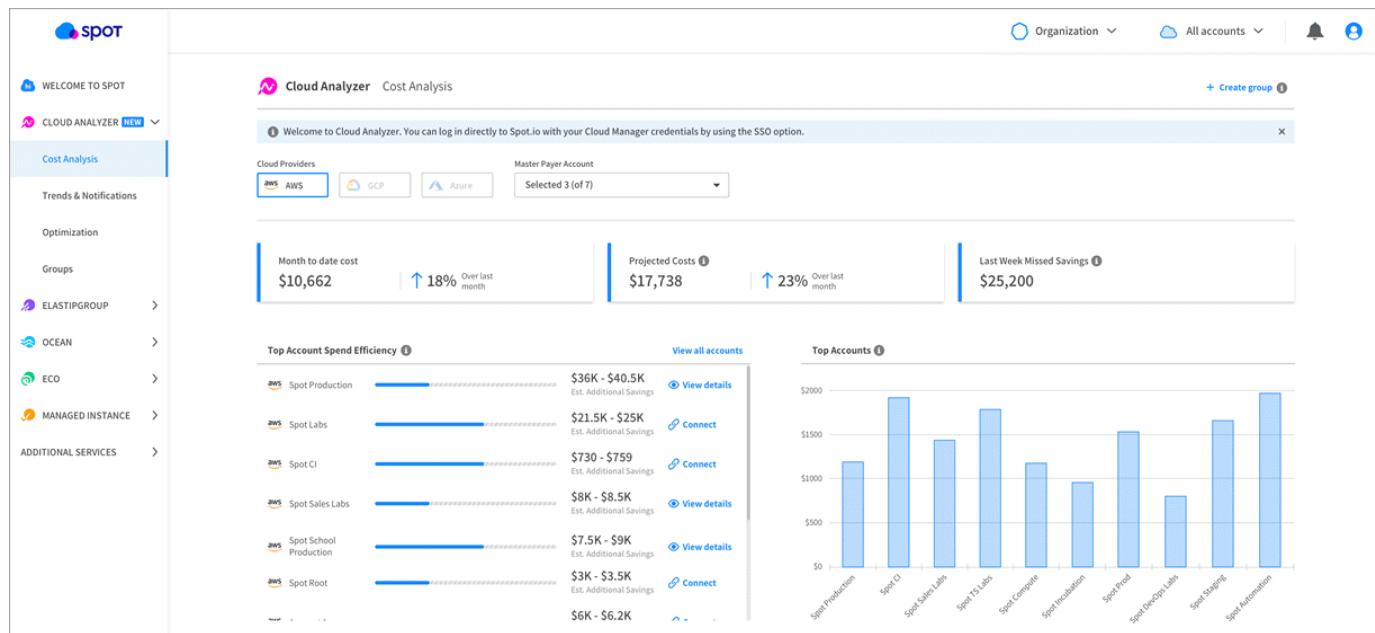
This is a bar graph showing your top ten accounts by cost. The graph is based on the last 30 days of spend activity.

[Learn more about the Cost Analysis page that's available in Spot's Cloud Analyzer.](#)

Go to Cloud Analyzer for more analysis and recommendations

Click **Get Full Analysis** at any time to access more charts and analysis, in-depth recommendations, a use case optimization breakdown (containers, ElasticApps, and reservations), and more.

Here's an example of what you'll see in Cloud Analyzer:



- [View the product page for Cloud Analyzer to learn more about its capabilities.](#)
- [View the documentation for Spot to get help using Cloud Analyzer.](#)

Viewing your Amazon S3 buckets

After you install a Connector in AWS, Cloud Manager can automatically discover information about the Amazon S3 buckets that reside in the AWS account where it's installed.

You can see details about your S3 buckets, including the region, access level, storage class, and whether the bucket is used with Cloud Volumes ONTAP for backups or data tiering. And you can scan the S3 buckets with Cloud Data Sense.

Steps

1. [Install a Connector](#) in the AWS account where you want to view your Amazon S3 buckets.

You should automatically see an Amazon S3 working environment shortly after.



2. Click the working environment and select an action from the right pane.

The screenshot shows the AWS Cloud Data Sense interface for the Amazon S3 service. At the top, there is a circular icon with a trash can symbol and the text "Amazon S3" followed by a green square indicating it is "On". To the right is a small "X" button. Below this, there is a section titled "INFORMATION" with two data points: "241 Buckets" and "15 Regions". Under the "SERVICES" section, there is a "Cloud Compliance" setting with a cloud and lock icon, currently set to "Off". A blue button labeled "Enable Compliance" is available to change this setting. A large blue button at the bottom left labeled "View Buckets" is also visible.

3. Click **Enable Compliance** to scan the S3 buckets for personal and sensitive data.

For more details, see [Getting started with Cloud Data Sense for Amazon S3](#).

4. Click **View Buckets** to view details about the S3 buckets in your AWS account.

Administer Cloud Manager

Manage Connectors

Finding the system ID for a Connector

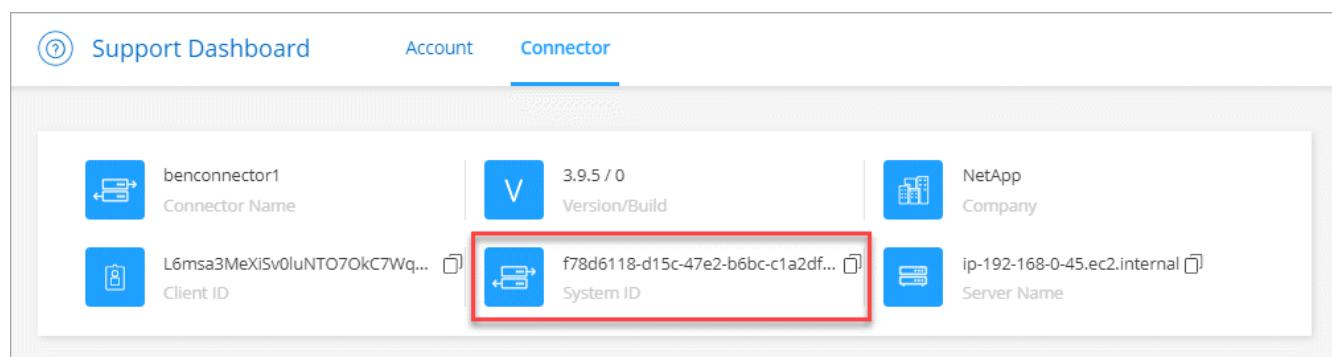
To help you get started, your NetApp representative might ask you for the system ID for a Connector. The ID is typically used for licensing and troubleshooting purposes.

Steps

1. In the upper right of the Cloud Manager console, click the Help icon.
2. Click **Support > Connector**.

The system ID appears at the top.

Example



Managing existing Connectors

After you create one or more Connectors, you can manage them by switching between Connectors, connecting to the local user interface running on a Connector, and more.

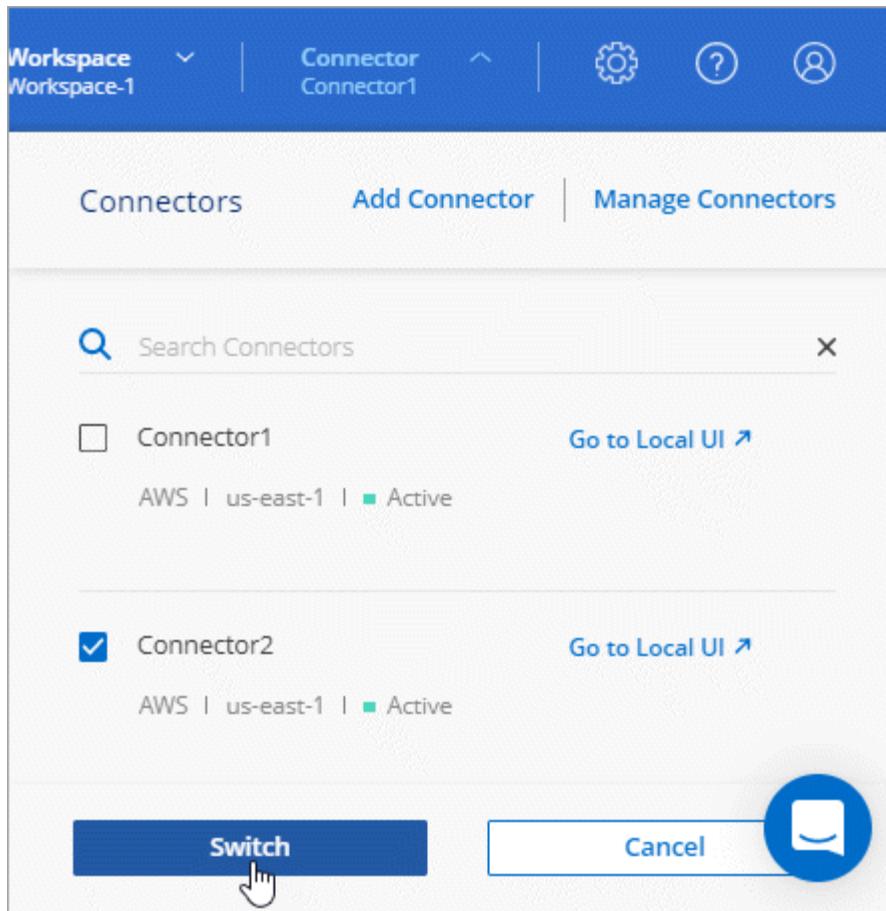
Switching between Connectors

If you have multiple Connectors, you can switch between them to see the Working Environments that are associated with a specific Connector.

For example, let's say that you're working in a multi-cloud environment. You might have one Connector in AWS and another in Google Cloud. You'd need to switch between those Connectors to manage the Cloud Volumes ONTAP systems running in those clouds.

Step

1. Click the **Connector** drop-down, select another Connector, and then click **Switch**.



Cloud Manager refreshes and shows the Working Environments associated with the selected Connector.

Accessing the local UI

While you should perform almost all tasks from the SaaS user interface, a local user interface is still available on the Connector. This interface is needed for a few tasks that need to be performed from the Connector itself:

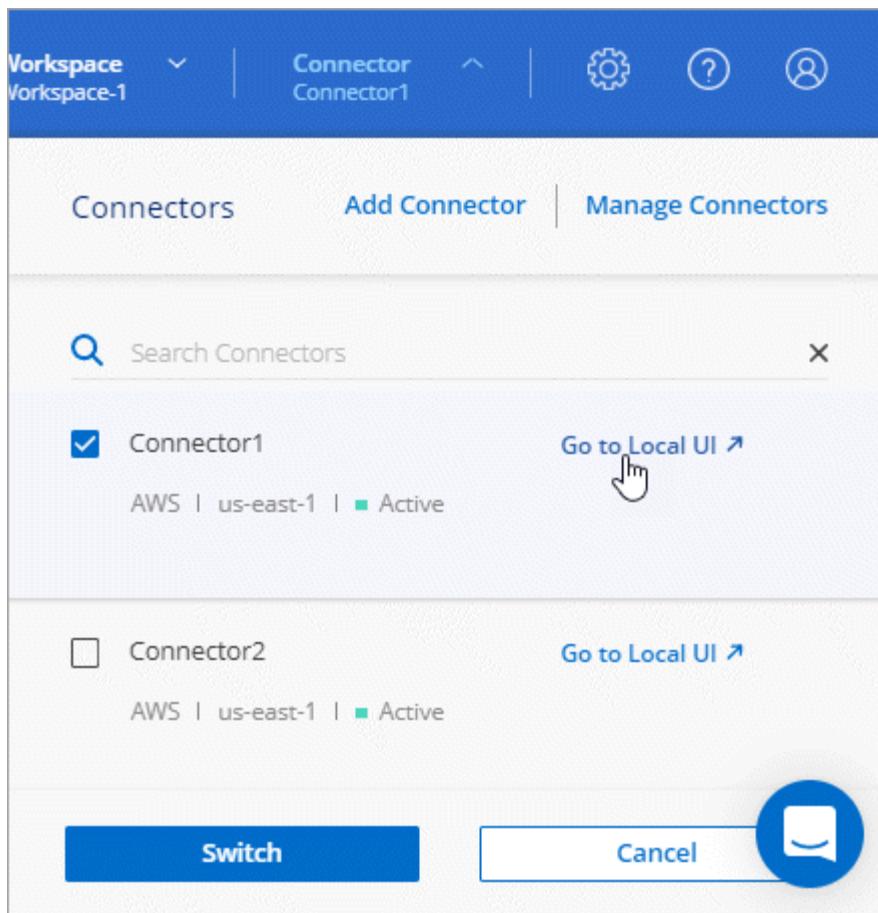
- [Setting a proxy server](#)
- Installing a patch (you'll typically work with NetApp personnel to install a patch)
- Downloading AutoSupport messages (usually directed by NetApp personnel when you have issues)

Steps

1. [Log in to the Cloud Manager SaaS interface](#) from a machine that has a network connection to the Connector instance.

If the Connector doesn't have a public IP address, you'll need a VPN connection or you'll need to connect from a jump host that's in the same network as the Connector.

2. Click the **Connector** drop-down and then click **Go to Local UI**.



The Cloud Manager interface running on the Connector loads in a new browser tab.

Editing a Connector's URIs

Add and remove the URIs for a Connector.

Steps

1. Click the **Connector** drop-down from the Cloud Manager header.
2. Click **Manage Connectors**.
3. Click the action menu for a Connector and click **Edit URIs**.
4. Add and remove URIs and then click **Apply**.

Removing Connectors from Cloud Manager

If a Connector is inactive, you can remove it from the list of Connectors in Cloud Manager. You might do this if you deleted the Connector virtual machine or if you uninstalled the Connector software.

Note the following about removing a Connector:

- This action doesn't delete the virtual machine.
- This action can't be reverted—once you remove a Connector from Cloud Manager, you can't add it back to Cloud Manager.

Steps

1. Click the **Connector** drop-down from the Cloud Manager header.
2. Click **Manage Connectors**.
3. Click the action menu for an inactive Connector and click **Remove Connector**.

Connector Name	Status	Cloud Provider	Region	Actions
Connector1	● Active	aws	US East (N. Virginia)	...
Connector2	● Inactive	aws	US East (N. Virginia)	...

Context menu for Connector2:

- Go to Local UI ↗
- Connector Id: iEMkyjIEyG4U5fXpmpmsnZS... 📁
- Edit URIs
- Remove Connector** (highlighted)

4. Enter the name of the Connector to confirm and then click Remove.

Result

Cloud Manager removes the Connector from its records.

Uninstalling the Connector software

The Connector includes an uninstallation script that you can use to uninstall the software to troubleshoot issues or to permanently remove the software from the host.

Step

1. From the Linux host, run the uninstallation script:

```
/opt/application/netapp/cloudmanager/bin/uninstall.sh [silent]
```

silent runs the script without prompting you for confirmation.

What about software upgrades?

The Connector automatically updates its software to the latest version, as long as it has [outbound internet access](#) to obtain the software update.

More ways to create Connectors

Connector host requirements

The Connector software must run on a host that meets specific operating system requirements, RAM requirements, port requirements, and so on.

A dedicated host is required

The Connector is not supported on a host that is shared with other applications. The host must be a dedicated host.

CPU

4 cores or 4 vCPUs

RAM

14 GB

AWS EC2 instance type

An instance type that meets the CPU and RAM requirements above. We recommend t3.xlarge and use that instance type when you deploy the Connector directly from Cloud Manager.

Azure VM size

An instance type that meets the CPU and RAM requirements above. We recommend DS3 v2 and use that VM size when you deploy the Connector directly from Cloud Manager.

GCP machine type

An instance type that meets the CPU and RAM requirements above. We recommend n1-standard-4 and use that machine type when you deploy the Connector directly from Cloud Manager.

Supported operating systems

- CentOS 7.6
- CentOS 7.7
- CentOS 7.8
- CentOS 7.9
- Red Hat Enterprise Linux 7.6
- Red Hat Enterprise Linux 7.7
- Red Hat Enterprise Linux 7.8
- Red Hat Enterprise Linux 7.9

The Red Hat Enterprise Linux system must be registered with Red Hat Subscription Management. If it is not registered, the system cannot access repositories to update required 3rd party software during Connector installation.

The Connector is supported on English-language versions of these operating systems.

Hypervisor

A bare metal or hosted hypervisor that is certified to run CentOS or Red Hat Enterprise Linux

[Red Hat Solution: Which hypervisors are certified to run Red Hat Enterprise Linux?](#)

Disk space in /opt

100 GiB of space must be available

Disk space in /var

20 GiB of space must be available

Outbound internet access

Outbound internet access is required to install the Connector and for the Connector to manage resources and processes within your public cloud environment. For a list of endpoints, see [Networking requirements for the Connector](#).

Creating a Connector from the AWS Marketplace

It's best to create a Connector directly from Cloud Manager, but you can launch a Connector from the AWS Marketplace, if you'd rather not specify AWS access keys. After you create and set up the Connector, Cloud Manager will automatically use it when you create new working environments.

Steps

1. Create an IAM policy and role for the EC2 instance:
 - a. Download the Cloud Manager IAM policy from the following location:
[NetApp Cloud Manager: AWS, Azure, and GCP Policies](#)
 - b. From the IAM console, create your own policy by copying and pasting the text from the Cloud Manager IAM policy.
 - c. Create an IAM role with the role type Amazon EC2 and attach the policy that you created in the previous step to the role.
2. Now go to the [Cloud Manager page on the AWS Marketplace](#) to deploy Cloud Manager from an AMI.

The IAM user must have AWS Marketplace permissions to subscribe and unsubscribe.

3. On the Marketplace page, click **Continue to Subscribe** and then click **Continue to Configuration**.

a

Cloud Manager - Manual Installation without access keys

By: NetApp, Inc. Latest Version: 3.8.4

Read below for instructions on how to deploy Cloud Volumes ONTAP.

Linux/Unix ★★★★★ 6 AWS reviews

Overview Pricing Usage Support Reviews

Product Overview

Do NOT subscribe on this page unless instructed by NetApp or redirected here from the NetApp website.

This listing lets you manually launch a Cloud Manager instance without providing your AWS credentials. After launching the Cloud Manager software in AWS, you can access it by entering the instance's IP address in a web browser. If you subscribe here, you still need to subscribe on the listing below for PAYGO charges.

Highlights

- See Product Overview for instructions on how to deploy NetApp Cloud Manager.

b

Cloud Manager - Manual Installation without access keys

< Product Detail [Subscribe](#)

Subscribe to this software

You're subscribed to this software. Please see the terms and pricing details below or click the button above to configure your software.

Terms and Conditions

NetApp, Inc. Offer

You have subscribed to this software and agreed that your use of this software is subject to the pricing terms and the seller's [End User License Agreement \(EULA\)](#). You agreed that AWS may share information about this transaction (including your payment terms) with the respective seller, reseller or underlying provider, as applicable, in accordance with the [AWS Privacy Notice](#). Your use of AWS services remains subject to the [AWS Customer Agreement](#) or other agreement with AWS governing your use of such services.

4. Change any of the default options and click **Continue to Launch**.

5. Under **Choose Action**, select **Launch through EC2** and then click **Launch**.

These steps describe how to launch the instance from the EC2 Console because the console enables you to attach an IAM role to the Cloud Manager instance. This isn't possible using the **Launch from Website** action.

6. Follow the prompts to configure and deploy the instance:

- **Choose Instance Type:** Depending on region availability, choose one of the supported instance types (t3.xlarge is recommended).

[Review the instance requirements](#).

- **Configure Instance:** Select a VPC and subnet, choose the IAM role that you created in step 1, enable termination protection (recommended), and choose any other configuration options that meet your requirements.

Number of instances	<input type="text" value="1"/>	Launch into Auto Scaling Group
Purchasing option	<input type="checkbox"/> Request Spot instances	
Network	vpc-a76d91c2 VPC4QA (default)	<input type="button" value="Create new VPC"/>
Subnet	subnet-39536c13 QASubnet1 us-east-1b 155 IP Addresses available	<input type="button" value="Create new subnet"/>
Auto-assign Public IP	Enable	
Placement group	<input type="checkbox"/> Add instance to placement group	
Capacity Reservation	Open	
IAM role	Cloud_Manager	<input type="button" value="Create new IAM role"/>
CPU options	<input type="checkbox"/> Specify CPU options	
Shutdown behavior	Stop	
Enable termination protection	<input checked="" type="checkbox"/> Protect against accidental termination	
Monitoring	<input type="checkbox"/> Enable CloudWatch detailed monitoring <small>Additional charges apply.</small>	

- **Add Storage:** Keep the default storage options.
- **Add Tags:** Enter tags for the instance, if desired.
- **Configure Security Group:** Specify the required connection methods for the Connector instance: SSH, HTTP, and HTTPS.
- **Review:** Review your selections and click **Launch**.

AWS launches the software with the specified settings. The Connector instance and software should be running in approximately five minutes.

7. Open a web browser from a host that has a connection to the Connector instance and enter the following URL:

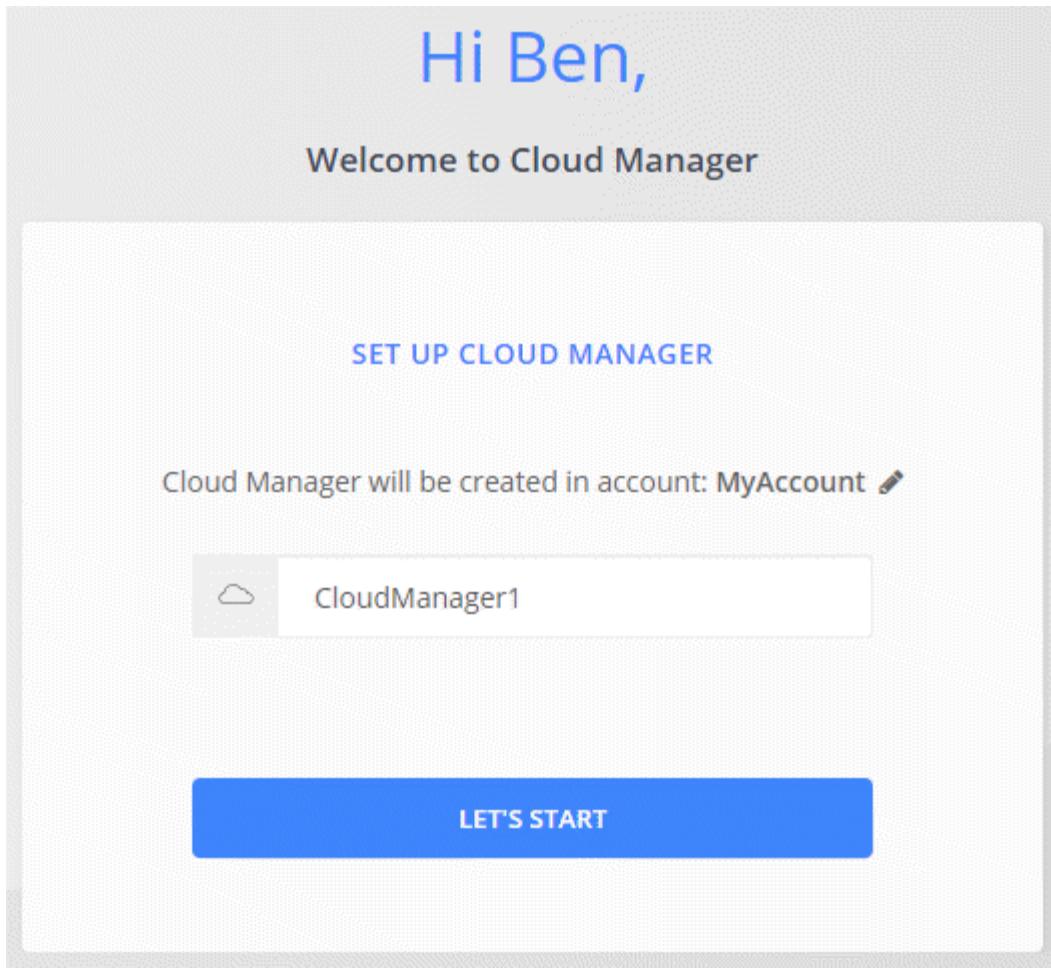
<http://ipaddress:80>

8. After you log in, set up the Connector:

- a. Specify the Cloud Central account to associate with the Connector.

[Learn about Cloud Central accounts.](#)

- b. Enter a name for the system.



Result

The Connector is now installed and set up with your Cloud Central account. Cloud Manager will automatically use this Connector when you create new working environments. But if you have more than one Connector, you'll need to [switch between them](#).

Creating a Connector from the Azure Marketplace

It's best to create a Connector directly from Cloud Manager, but you can launch a Connector from the Azure Marketplace, if you prefer. After you create and set up the Connector, Cloud Manager will automatically use it when you create new working environments.

Creating a Connector in Azure

Deploy the Connector in Azure using the image in the Azure Marketplace and then log in to the Connector to specify your Cloud Central account.

Steps

1. [Go to the Azure Marketplace page for Cloud Manager](#).
2. Click **Get it now** and then click **Continue**.
3. From the Azure portal, click **Create** and follow the steps to configure the virtual machine.

Note the following as you configure the VM:

- Cloud Manager can perform optimally with either HDD or SSD disks.
- Choose a VM size that meets CPU and RAM requirements. We recommend DS3 v2.

[Review the VM requirements.](#)

- For the network security group, the Connector requires inbound connections using SSH, HTTP, and HTTPS.

[Learn more about security group rules for the Connector.](#)

- Under **Management**, enable **System assigned managed identity** for the Connector by selecting **On**.

This setting is important because a managed identity allows the Connector virtual machine to identify itself to Azure Active Directory without providing any credentials. [Learn more about managed identities for Azure resources](#).

4. On the **Review + create** page, review your selections and click **Create** to start the deployment.

Azure deploys the virtual machine with the specified settings. The virtual machine and Connector software should be running in approximately five minutes.

5. Open a web browser from a host that has a connection to the Connector virtual machine and enter the following URL:

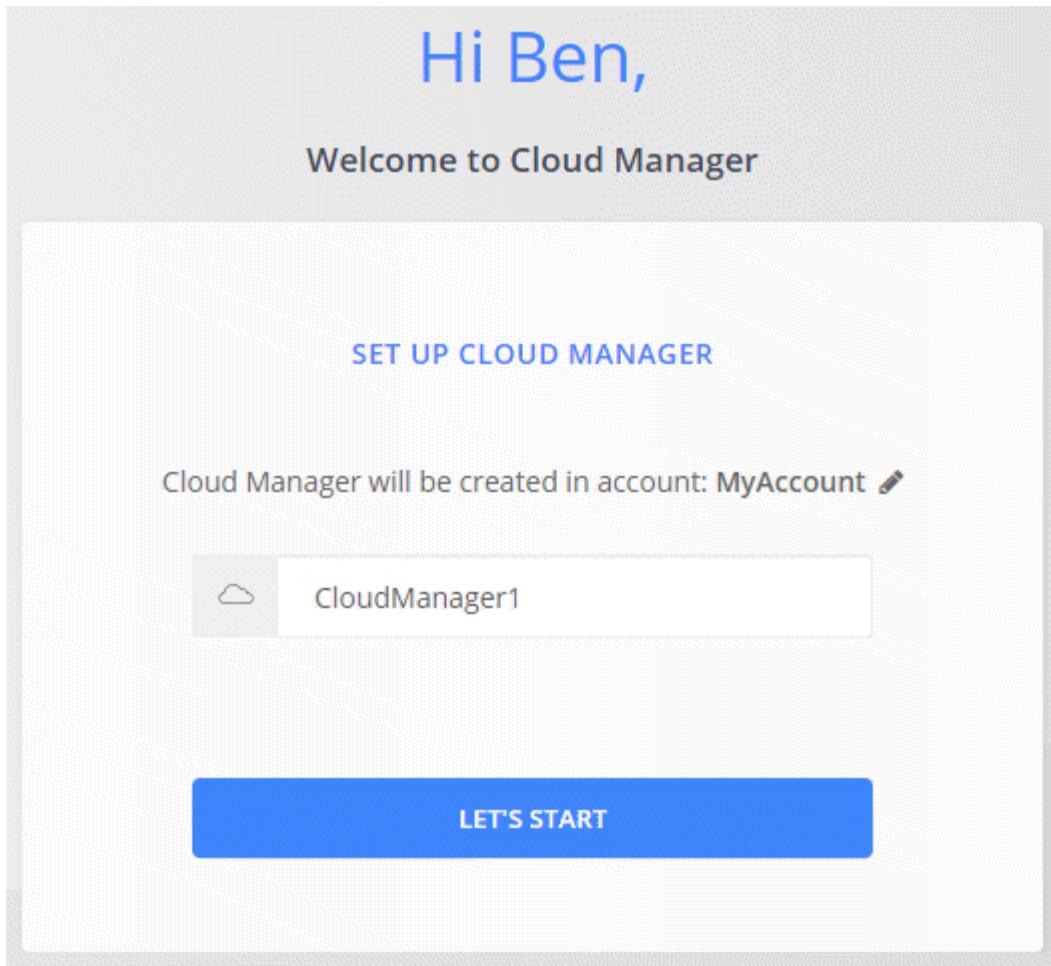
`http://ipaddress:80`

6. After you log in, set up the Connector:

- a. Specify the Cloud Central account to associate with the Connector.

[Learn about Cloud Central accounts.](#)

- b. Enter a name for the system.



Result

The Connector is now installed and set up. You must grant Azure permissions before users can deploy Cloud Volumes ONTAP in Azure.

Granting Azure permissions

When you deployed the Connector in Azure, you should have enabled a [system-assigned managed identity](#). You must now grant the required Azure permissions by creating a custom role and then by assigning the role to the Connector virtual machine for one or more subscriptions.

Steps

1. Create a custom role using the Cloud Manager policy:
 - a. Download the [Cloud Manager Azure policy](#).
 - b. Modify the JSON file by adding Azure subscription IDs to the assignable scope.

You should add the ID for each Azure subscription from which users will create Cloud Volumes ONTAP systems.

Example

```
"AssignableScopes": [  
  "/subscriptions/d333af45-0d07-4154-943d-c25fbzzzzzz",  
  "/subscriptions/54b91999-b3e6-4599-908e-416e0zzzzzz",  
  "/subscriptions/398e471c-3b42-4ae7-9b59-ce5bbzzzzzz"
```

- c. Use the JSON file to create a custom role in Azure.

The following example shows how to create a custom role using the Azure CLI 2.0:

```
az role definition create --role-definition  
C:\Policy_for_Cloud_Manager_Azure_3.9.8.json
```

You should now have a custom role called Cloud Manager Operator that you can assign to the Connector virtual machine.

2. Assign the role to the Connector virtual machine for one or more subscriptions:

- a. Open the **Subscriptions** service and then select the subscription in which you want to deploy Cloud Volumes ONTAP systems.
- b. Click **Access control (IAM)**.
- c. Click **Add > Add role assignment** and then add the permissions:
 - Select the **Cloud Manager Operator** role.



Cloud Manager Operator is the default name provided in the [Cloud Manager policy](#). If you chose a different name for the role, then select that name instead.

- d. Assign access to a **Virtual Machine**.
 - Select the subscription in which the Connector virtual machine was created.
 - Select the Connector virtual machine.
 - Click **Save**.
- d. If you want to deploy Cloud Volumes ONTAP from additional subscriptions, switch to that subscription and then repeat these steps.

Result

The Connector now has the permissions that it needs to manage resources and processes within your public cloud environment. Cloud Manager will automatically use this Connector when you create new working environments. But if you have more than one Connector, you'll need to [switch between them](#).

Installing the Connector software on an existing Linux host

The most common way to create a Connector is directly from Cloud Manager or from a cloud provider's marketplace. But you have the option to download and install the Connector software on an existing Linux host in your network or in the cloud.



If you want to create a Cloud Volumes ONTAP system in Google Cloud, then you must have a Connector running in Google Cloud, as well. You can't use a Connector that's running in another location.

What you'll need

- The host must meet [requirements for the Connector](#).
- A Red Hat Enterprise Linux system must be registered with Red Hat Subscription Management. If it is not registered, the system cannot access repositories to update required 3rd party software during installation.
- The Connector installer accesses several URLs during the installation process. You must ensure that outbound internet access is allowed to these endpoints:

- <http://dev.mysql.com/get/mysql-community-release-el7-5.noarch.rpm>
- <https://dl.fedoraproject.org/pub/epel/epel-release-latest-7.noarch.rpm>
- <https://s3.amazonaws.com/aws-cli/awscli-bundle.zip>

The host might try to update operating system packages during installation. The host can contact different mirroring sites for these OS packages.

- After you install the Connector, you'll need to set up permissions so Cloud Manager can manage resources and processes within your public cloud environment. See the steps below for more details.

About this task

- Root privileges are not required to install the Connector.
- The installation installs the AWS command line tools (awscli) to enable recovery procedures from NetApp support.

If you receive a message that installing the awscli failed, you can safely ignore the message. The Connector can operate successfully without the tools.

- The installer that is available on the NetApp Support Site might be an earlier version. After installation, the Connector automatically updates itself if a new version is available.

Steps

1. Download the Cloud Manager software from the [NetApp Support Site](#), and then copy it to the Linux host.

For help with connecting and copying the file to an EC2 instance in AWS, see [AWS Documentation: Connecting to Your Linux Instance Using SSH](#).

2. Assign permissions to run the script.

Example

```
chmod +x OnCommandCloudManager-V3.8.9.sh
```

3. Run the installation script:

```
./OnCommandCloudManager-V3.8.9.sh [silent] [proxy=ipaddress]
[proxyport=port] [proxyuser=user_name] [proxypwd=password]
```

silent runs the installation without prompting you for information.

proxy is required if the host is behind a proxy server.

proxyport is the port for the proxy server.

proxyuser is the user name for the proxy server, if basic authentication is required.

proxypwd is the password for the user name that you specified.

4. Unless you specified the silent parameter, type **Y** to continue the script, and then enter the HTTP and HTTPS ports when prompted.

Cloud Manager is now installed. At the end of the installation, the Cloud Manager service (occm) restarts twice if you specified a proxy server.

5. Open a web browser and enter the following URL:

`https://ipaddress:port`

ipaddress can be localhost, a private IP address, or a public IP address, depending on the configuration of the host. For example, if the Connector is in the public cloud without a public IP address, you must enter a private IP address from a host that has a connection to the Connector host.

port is required if you changed the default HTTP (80) or HTTPS (443) ports. For example, if the HTTPS port was changed to 8443, you would enter `https://ipaddress:8443`

6. Sign up at NetApp Cloud Central or log in.

7. If you installed the Connector in Google Cloud, set up a service account that has the permissions that Cloud Manager needs to create and manage Cloud Volumes ONTAP systems in projects.

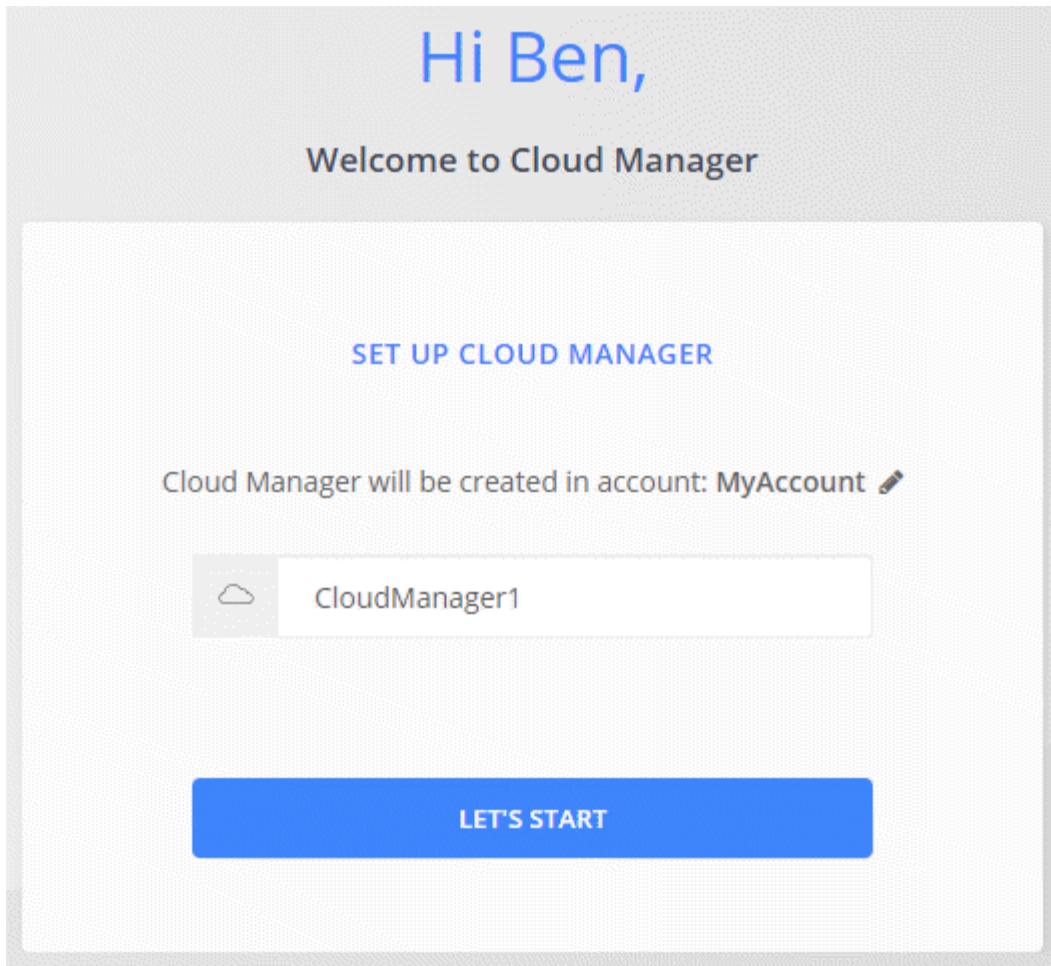
- Create a role in GCP that includes the permissions defined in the [Cloud Manager policy for GCP](#).
- Create a GCP service account and apply the custom role that you just created.
- Associate this service account with the Connector VM.
- If you want to deploy Cloud Volumes ONTAP in other projects, [grant access by adding the service account with the Cloud Manager role to that project](#). You'll need to repeat this step for each project.

8. After you log in, set up Cloud Manager:

- Specify the Cloud Central account to associate with the Connector.

[Learn about Cloud Central accounts](#).

- Enter a name for the system.



Result

The Connector is now installed and set up with your Cloud Central account. Cloud Manager will automatically use this Connector when you create new working environments.

After you finish

Set up permissions so Cloud Manager can manage resources and processes within your public cloud environment:

- AWS: [Set up an AWS account and then add it to Cloud Manager](#).
- Azure: [Set up an Azure account and then add it to Cloud Manager](#).
- Google Cloud: See step 7 above.

Default configuration for the Connector

If you need to troubleshoot the Connector, it might help to understand how it's configured.

- If you deployed the Connector from Cloud Manager (or directly from a cloud provider's marketplace), note the following:
 - In AWS, the user name for the EC2 Linux instance is ec2-user.
 - The operating system for the image is as follows:
 - AWS: Red Hat Enterprise Linux 7.6 (HVM)

- Azure: CentOS 7.6
- GCP: CentOS 7.9

The operating system does not include a GUI. You must use a terminal to access the system.

- The Connector installation folder resides in the following location:

/opt/application/netapp/cloudmanager

- Log files are contained in the following folders:

- /opt/application/netapp/cloudmanager/log

The logs in this folder provide details about the Connector and docker images.

- /opt/application/netapp/cloudmanager/docker_occm/data/log

The logs in this folder provide details about cloud services and the Cloud Manager service that runs on the Connector.

- The Cloud Manager service is named occm.
- The occm service is dependent on the MySQL service.

If the MySQL service is down, then the occm service is down too.

- Cloud Manager installs the following packages on the Linux host, if they are not already installed:

- 7Zip
 - AWSCLI
 - Docker
 - Java
 - Kubectl
 - MySQL
 - Tridentctl
 - Pull
 - Wget

- The Connector uses the following ports on the Linux host:
 - 80 for HTTP access
 - 443 for HTTPS access
 - 3306 for the Cloud Manager database
 - 8080 for the Cloud Manager API proxy
 - 8666 for the Service Manager API
 - 8777 for the Health-Checker Container Service API

Manage cloud provider credentials

AWS

AWS credentials and permissions

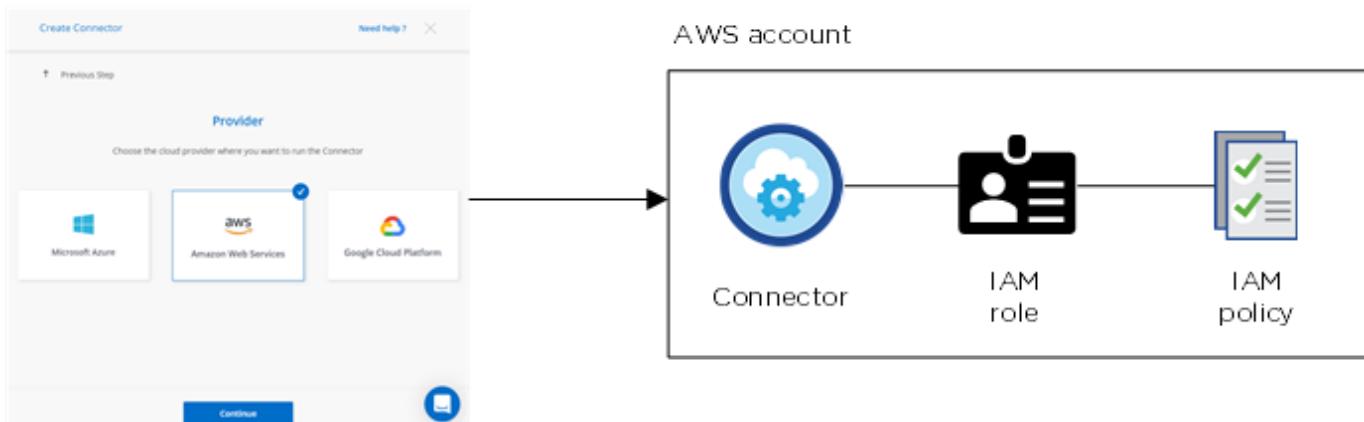
Cloud Manager enables you to choose the AWS credentials to use when deploying Cloud Volumes ONTAP. You can deploy all of your Cloud Volumes ONTAP systems using the initial AWS credentials, or you can add additional credentials.

Initial AWS credentials

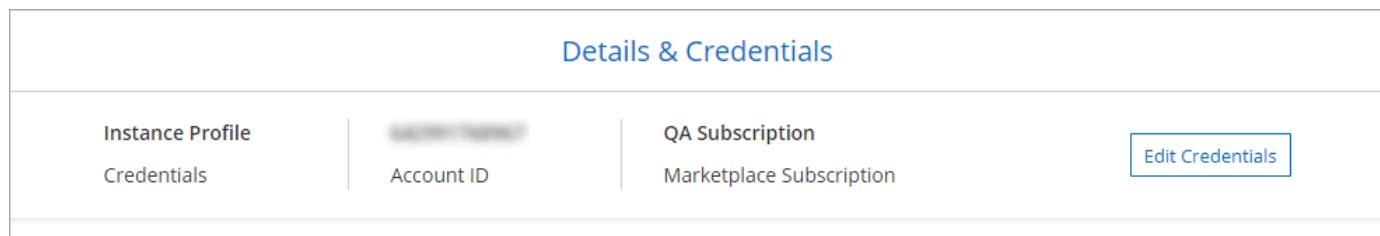
When you deploy a Connector from Cloud Manager, you need to use an AWS account that has permissions to launch the Connector instance. The required permissions are listed in the [Connector deployment policy for AWS](#).

When Cloud Manager launches the Connector instance in AWS, it creates an IAM role and an instance profile for the instance. It also attaches a policy that provides Cloud Manager with permissions to manage resources and processes within that AWS account. [Review how Cloud Manager uses the permissions](#).

Cloud Manager

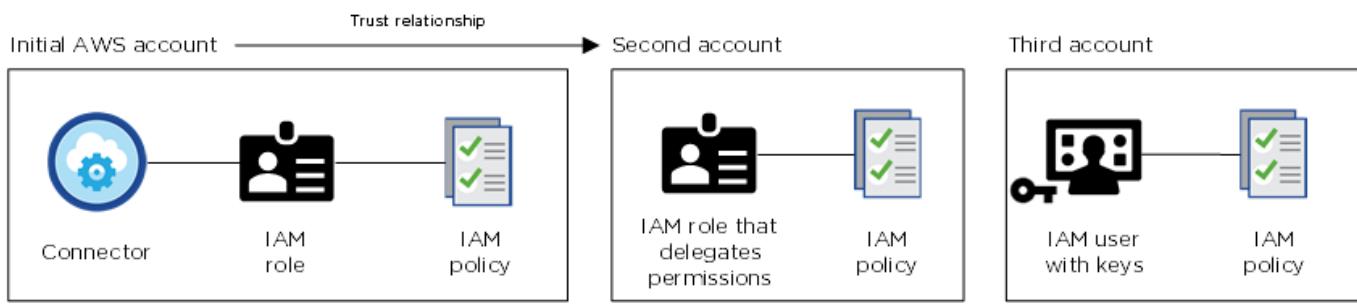


Cloud Manager selects these AWS credentials by default when you create a new working environment for Cloud Volumes ONTAP:



Additional AWS credentials

If you want to launch Cloud Volumes ONTAP in different AWS accounts, then you can either [provide AWS keys for an IAM user or the ARN of a role in a trusted account](#). The following image shows two additional accounts, one providing permissions through an IAM role in a trusted account and another through the AWS keys of an IAM user:



You would then [add the account credentials to Cloud Manager](#) by specifying the Amazon Resource Name (ARN) of the IAM role, or the AWS keys for the IAM user.

After you add another set of credentials, you can switch to them when creating a new working environment:

Edit Account & Add Subscription

Credentials

[Keys | Account ID: \[REDACTED\]](#)
[Hand cursor icon]

[Instance Profile | Account ID: \[REDACTED\]](#)

● QA Subscription

Associate Subscription to Credentials

To create a pay-as-you-go Cloud Volumes ONTAP system, you need to select AWS credentials that are associated with a subscription to Cloud Volumes ONTAP from the AWS Marketplace.

[+ Add Subscription](#)

[Apply](#)

[Cancel](#)

What about Marketplace deployments and on-prem deployments?

The sections above describe the recommended deployment method for the Connector, which is from Cloud Manager. You can also deploy a Connector in AWS from the [AWS Marketplace](#) and you can [install the Connector on-premises](#).

If you use the Marketplace, permissions are provided in the same way. You just need to manually create and set up the IAM role, and then provide permissions for any additional accounts.

For on-premises deployments, you can't set up an IAM role for the Cloud Manager system, but you can provide permissions just like you would for additional AWS accounts.

How can I securely rotate my AWS credentials?

As described above, Cloud Manager enables you to provide AWS credentials in a few ways: an IAM role associated with the Connector instance, by assuming an IAM role in a trusted account, or by providing AWS access keys.

With the first two options, Cloud Manager uses the AWS Security Token Service to obtain temporary credentials that rotate constantly. This process is the best practice—it's automatic and it's secure.

If you provide Cloud Manager with AWS access keys, you should rotate the keys by updating them in Cloud Manager at a regular interval. This is a completely manual process.

Managing AWS credentials and subscriptions for Cloud Manager

When you create a Cloud Volumes ONTAP system, you need to select the AWS credentials and subscription to use with that system. If you manage multiple AWS subscriptions, you can assign each one of them to different AWS credentials from the Credentials page.

Before you add AWS credentials to Cloud Manager, you need to provide the required permissions to that account. The permissions enable Cloud Manager to manage resources and processes within that AWS account. How you provide the permissions depends on whether you want to provide Cloud Manager with AWS keys or the ARN of a role in a trusted account.

 When you deployed a Connector from Cloud Manager, Cloud Manager automatically added AWS credentials for the account in which you deployed the Connector. This initial account is not added if you manually installed the Connector software on an existing system. [Learn about AWS credentials and permissions](#).

Choices

- [Granting permissions by providing AWS keys](#)
- [Granting permissions by assuming IAM roles in other accounts](#)

How can I securely rotate my AWS credentials?

Cloud Manager enables you to provide AWS credentials in a few ways: an IAM role associated with the Connector instance, by assuming an IAM role in a trusted account, or by providing AWS access keys. [Learn more about AWS credentials and permissions.](#)

With the first two options, Cloud Manager uses the AWS Security Token Service to obtain temporary credentials that rotate constantly. This process is the best practice, it's automatic and it's secure.

If you provide Cloud Manager with AWS access keys, you should rotate the keys by updating them in Cloud Manager at a regular interval. This is a completely manual process.

Granting permissions by providing AWS keys

If you want to provide Cloud Manager with AWS keys for an IAM user, then you need to grant the required permissions to that user. The Cloud Manager IAM policy defines the AWS actions and resources that Cloud Manager is allowed to use.

Steps

1. Download the Cloud Manager IAM policy from the [Cloud Manager Policies page](#).
2. From the IAM console, create your own policy by copying and pasting the text from the Cloud Manager IAM policy.
[AWS Documentation: Creating IAM Policies](#)
3. Attach the policy to an IAM role or an IAM user.
 - [AWS Documentation: Creating IAM Roles](#)
 - [AWS Documentation: Adding and Removing IAM Policies](#)

Result

The account now has the required permissions. [You can now add it to Cloud Manager.](#)

Granting permissions by assuming IAM roles in other accounts

You can set up a trust relationship between the source AWS account in which you deployed the Connector instance and other AWS accounts by using IAM roles. You would then provide Cloud Manager with the ARN of the IAM roles from the trusted accounts.

Steps

1. Go to the target account where you want to deploy Cloud Volumes ONTAP and create an IAM role by selecting **Another AWS account**.

Be sure to do the following:

- Enter the ID of the account where the Connector instance resides.
 - Attach the Cloud Manager IAM policy, which is available from the [Cloud Manager Policies page](#).
-
2. Go to the source account where the Connector instance resides and select the IAM role that is attached to the instance.

- a. Click **Attach policies** and then click **Create policy**.
- b. Create a policy that includes the "sts:AssumeRole" action and the ARN of the role that you created in the target account.

Example

```
{  
    "Version": "2012-10-17",  
    "Statement": {  
        "Effect": "Allow",  
        "Action": "sts:AssumeRole",  
        "Resource": "arn:aws:iam::ACCOUNT-B-ID:role/ACCOUNT-B-ROLENNAME"  
    }  
}
```

Result

The account now has the required permissions. [You can now add it to Cloud Manager](#).

Adding AWS credentials to Cloud Manager

After you provide an AWS account with the required permissions, you can add the credentials for that account to Cloud Manager. This enables you to launch Cloud Volumes ONTAP systems in that account.

Before you get started

If you just created these credentials in your cloud provider, it might take a few minutes until they are available for use. Wait a few minutes before you add the credentials to Cloud Manager.

Steps

1. In the upper right of the Cloud Manager console, click the Settings icon, and select **Credentials**.



2. Click **Add Credentials** and select **AWS**.
3. Provide AWS keys or the ARN of a trusted IAM role.
4. Confirm that the policy requirements have been met and click **Continue**.
5. Choose the subscription that you want to associate with the credentials, or click **Add Subscription** if you don't have one yet.

To pay for Cloud Volumes ONTAP at an hourly rate (PAYGO) or with an annual contract, AWS credentials must be associated with a subscription to Cloud Volumes ONTAP from the AWS Marketplace.

6. Click **Add**.

Result

You can now switch to a different set of credentials from the Details and Credentials page when creating a new

working environment:

Edit Account & Add Subscription

Credentials

Keys | Account ID: [REDACTED]

Instance Profile | Account ID: [REDACTED]

● QA Subscription

Associate Subscription to Credentials

To create a pay-as-you-go Cloud Volumes ONTAP system, you need to select AWS credentials that are associated with a subscription to Cloud Volumes ONTAP from the AWS Marketplace.

+ Add Subscription

Apply Cancel

Associating an AWS subscription to credentials

After you add your AWS credentials to Cloud Manager, you can associate an AWS Marketplace subscription with those credentials. The subscription enables you to pay for Cloud Volumes ONTAP at an hourly rate (PAYGO) or using an annual contract, and to use other NetApp cloud services.

There are two scenarios in which you might associate an AWS Marketplace subscription after you've already added the credentials to Cloud Manager:

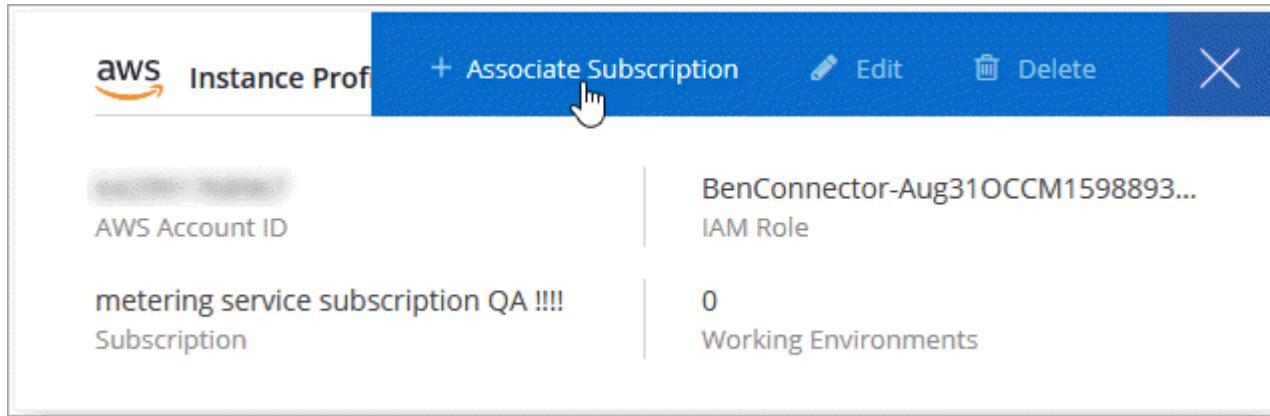
- You didn't associate a subscription when you initially added the credentials to Cloud Manager.
- You want to replace an existing AWS Marketplace subscription with a new subscription.

What you'll need

You need to create a Connector before you can change Cloud Manager settings. [Learn how](#).

Steps

1. In the upper right of the Cloud Manager console, click the Settings icon, and select **Credentials**.
2. Hover over a set of credentials and click the action menu.
3. From the menu, click **Associate Subscription**.



4. Select a subscription from the down-down list or click **Add Subscription** and follow the steps to create a new subscription.

► https://docs.netapp.com/us-en/occm//media/video_subscribing_aws.mp4 (video)

Azure

Azure credentials and permissions

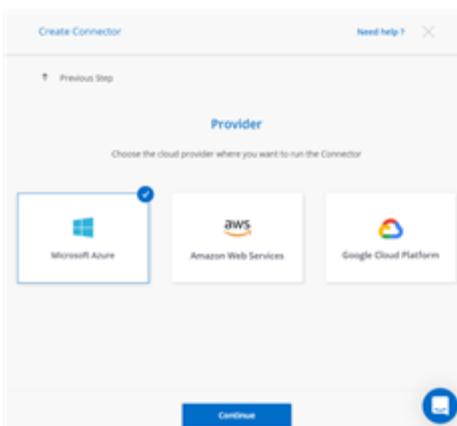
Cloud Manager enables you to choose the Azure credentials to use when deploying Cloud Volumes ONTAP. You can deploy all of your Cloud Volumes ONTAP systems using the initial Azure credentials, or you can add additional credentials.

Initial Azure credentials

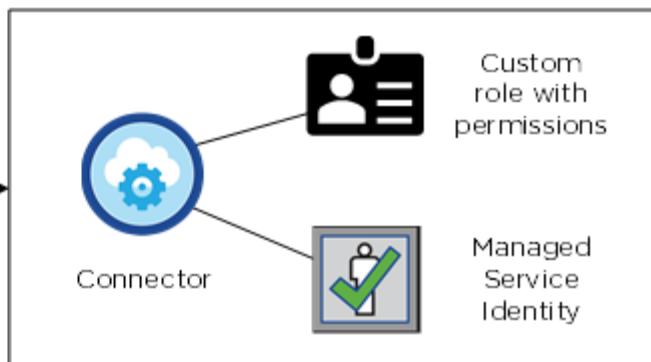
When you deploy a Connector from Cloud Manager, you need to use an Azure account that has permissions to deploy the Connector virtual machine. The required permissions are listed in the [Connector deployment policy for Azure](#).

When Cloud Manager deploys the Connector virtual machine in Azure, it enables a [system-assigned managed identity](#) on virtual machine, creates a custom role, and assigns it to the virtual machine. The role provides Cloud Manager with permissions to manage resources and processes within that Azure subscription. [Review how Cloud Manager uses the permissions](#).

Cloud Manager



Azure account



Cloud Manager selects these Azure credentials by default when you create a new working environment for Cloud Volumes ONTAP:

This screenshot shows the 'Details & Credentials' section. It displays the following information:

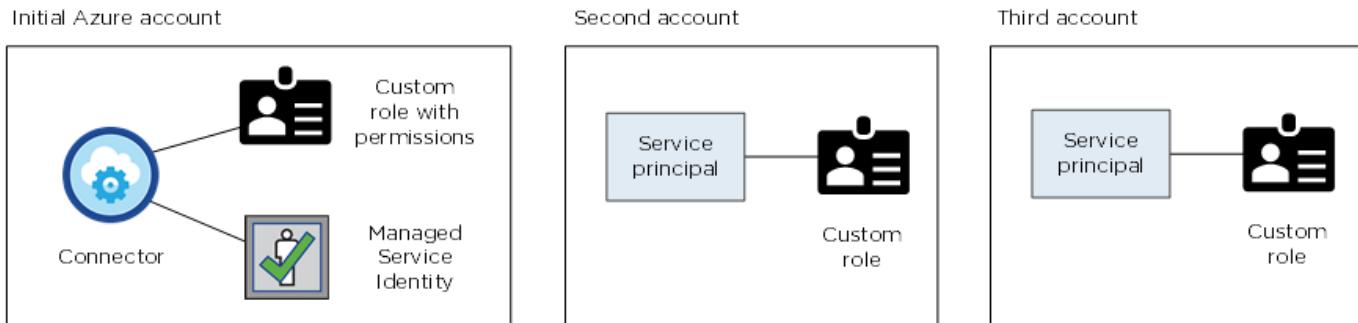
Managed Service Ide...	OCCM QA1	 ⓘ No subscription is associated	Edit Credentials
Credential Name	Azure Subscription	Marketplace Subscription	

Additional Azure subscriptions for a managed identity

The managed identity is associated with the subscription in which you launched the Connector. If you want to select a different Azure subscription, then you need to [associate the managed identity with those subscriptions](#).

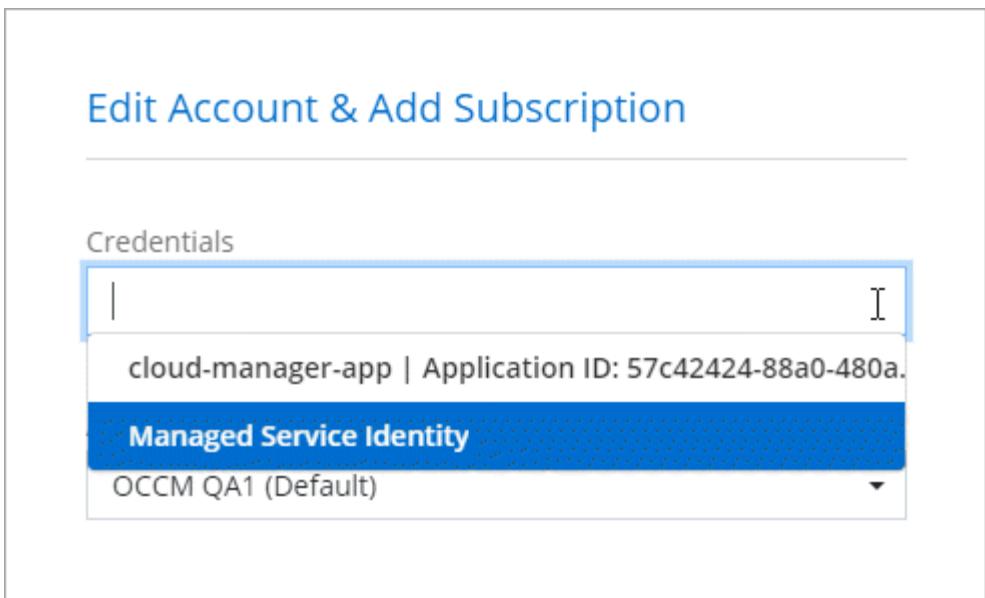
Additional Azure credentials

If you want to deploy Cloud Volumes ONTAP using different Azure credentials, then you must grant the required permissions by [creating and setting up a service principal in Azure Active Directory](#) for each Azure account. The following image shows two additional accounts, each set up with a service principal and custom role that provides permissions:



You would then [add the account credentials to Cloud Manager](#) by providing details about the AD service principal.

After you add another set of credentials, you can switch to them when creating a new working environment:



What about Marketplace deployments and on-prem deployments?

The sections above describe the recommended deployment method for the Connector, which is from NetApp Cloud Central. You can also deploy a Connector in Azure from the [Azure Marketplace](#), and you can [install the Connector on-premises](#).

If you use the Marketplace, permissions are provided in the same way. You just need to manually create and set up the managed identity for the Connector, and then provide permissions for any additional accounts.

For on-premises deployments, you can't set up a managed identity for the Connector, but you can provide permissions just like you would for additional accounts by using a service principal.

Managing Azure credentials and subscriptions for Cloud Manager

When you create a Cloud Volumes ONTAP system, you need to select the Azure credentials to use with that system. You also need to choose a Marketplace subscription, if you're using pay-as-you-go licensing. Follow the steps on this page if you need to use multiple Azure credentials or multiple Azure Marketplace subscriptions for Cloud Volumes ONTAP.

There are two ways to manage Azure credentials in Cloud Manager. First, if you want to deploy Cloud Volumes ONTAP using different Azure credentials, then you need to provide the required permissions and add the credentials to Cloud Manager. The second way is to associate additional subscriptions with the Azure managed identity.

Adding additional Azure credentials to Cloud Manager

When you deploy a Connector from Cloud Manager, Cloud Manager enables a system-assigned managed identity on the virtual machine that has the required permissions. Cloud Manager selects these Azure credentials by default when you create a new working environment for Cloud Volumes ONTAP.



An initial set of credentials isn't added if you manually installed the Connector software on an existing system. [Learn about Azure credentials and permissions](#).

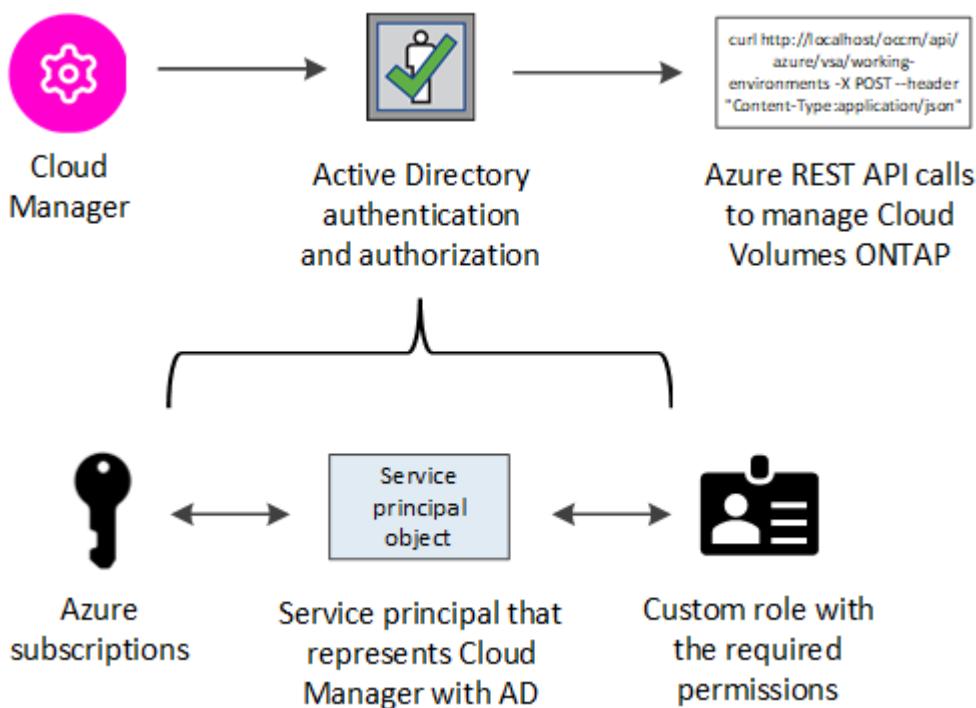
If you want to deploy Cloud Volumes ONTAP using *different* Azure credentials, then you must grant the required permissions by creating and setting up a service principal in Azure Active Directory for each Azure account. You can then add the new credentials to Cloud Manager.

Granting Azure permissions using a service principal

Cloud Manager needs permissions to perform actions in Azure. You can grant the required permissions to an Azure account by creating and setting up a service principal in Azure Active Directory and by obtaining the Azure credentials that Cloud Manager needs.

About this task

The following image depicts how Cloud Manager obtains permissions to perform operations in Azure. A service principal object, which is tied to one or more Azure subscriptions, represents Cloud Manager in Azure Active Directory and is assigned to a custom role that allows the required permissions.



Steps

1. [Create an Azure Active Directory application](#).
2. [Assign the application to a role](#).
3. [Add Windows Azure Service Management API permissions](#).
4. [Get the application ID and directory ID](#).
5. [Create a client secret](#).

Creating an Azure Active Directory application

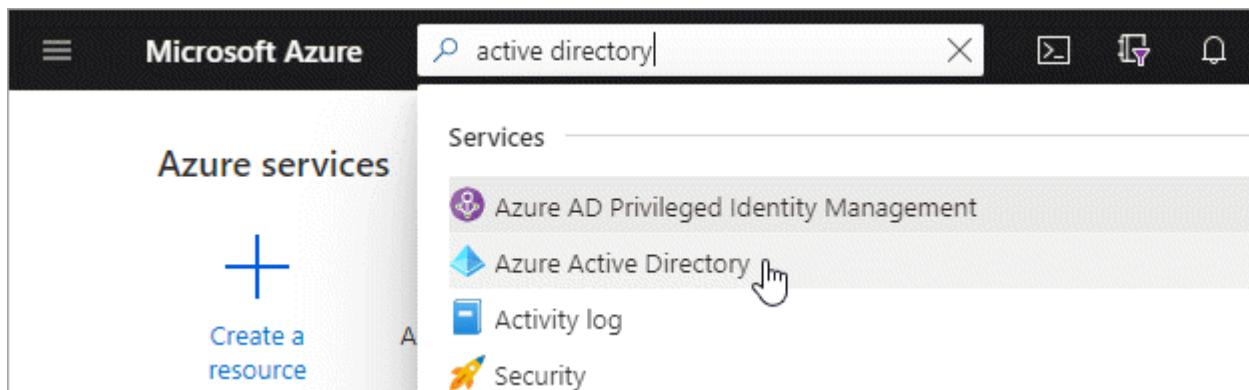
Create an Azure Active Directory (AD) application and service principal that Cloud Manager can use for role-based access control.

Before you begin

You must have the right permissions in Azure to create an Active Directory application and to assign the application to a role. For details, refer to [Microsoft Azure Documentation: Required permissions](#).

Steps

1. From the Azure portal, open the **Azure Active Directory** service.



2. In the menu, click **App registrations**.
3. Click **New registration**.
4. Specify details about the application:
 - **Name:** Enter a name for the application.
 - **Account type:** Select an account type (any will work with Cloud Manager).
 - **Redirect URI:** You can leave this field blank.
5. Click **Register**.

Result

You've created the AD application and service principal.

Assigning the application to a role

You must bind the service principal to one or more Azure subscriptions and assign it the custom "OnCommand Cloud Manager Operator" role so Cloud Manager has permissions in Azure.

Steps

1. Create a custom role:
 - a. Download the [Cloud Manager Azure policy](#).
 - b. Modify the JSON file by adding Azure subscription IDs to the assignable scope.

You should add the ID for each Azure subscription from which users will create Cloud Volumes ONTAP systems.

Example

```
"AssignableScopes": [  
    "/subscriptions/d333af45-0d07-4154-943d-c25fbzzzzzz",  
    "/subscriptions/54b91999-b3e6-4599-908e-416e0zzzzzz",  
    "/subscriptions/398e471c-3b42-4ae7-9b59-ce5bbzzzzzz"
```

- c. Use the JSON file to create a custom role in Azure.

The following example shows how to create a custom role using the Azure CLI 2.0:

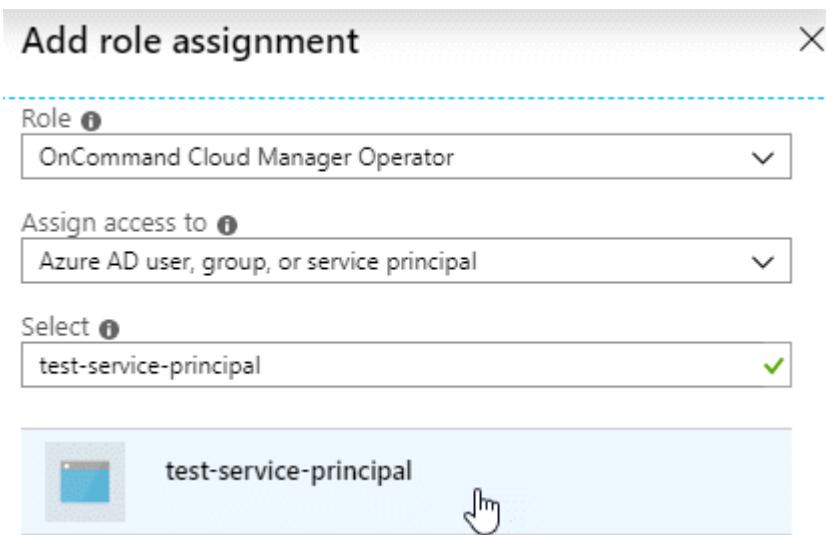
```
az role definition create --role-definition  
C:\Policy_for_cloud_Manager_Azure_3.9.8.json
```

You should now have a custom role called *Cloud Manager Operator*.

2. Assign the application to the role:

- a. From the Azure portal, open the **Subscriptions** service.
- b. Select the subscription.
- c. Click **Access control (IAM) > Add > Add role assignment**.
- d. Select the **Cloud Manager Operator** role.
- e. Keep **Azure AD user, group, or service principal** selected.
- f. Search for the name of the application (you can't find it in the list by scrolling).

Here's an example:



- g. Select the application and click **Save**.

The service principal for Cloud Manager now has the required Azure permissions for that subscription.

If you want to deploy Cloud Volumes ONTAP from multiple Azure subscriptions, then you must bind the service principal to each of those subscriptions. Cloud Manager enables you to select the subscription that you want to use when deploying Cloud Volumes ONTAP.

Adding Windows Azure Service Management API permissions

The service principal must have "Windows Azure Service Management API" permissions.

Steps

1. In the **Azure Active Directory** service, click **App registrations** and select the application.
2. Click **API permissions > Add a permission**.
3. Under **Microsoft APIs**, select **Azure Service Management**.

Request API permissions

Select an API

[Microsoft APIs](#) [APIs my organization uses](#) [My APIs](#)

Commonly used Microsoft APIs

Microsoft Graph Take advantage of the tremendous amount of data in Office 365, Enterprise Mobility + Security, and Windows 10. Access Azure AD, Excel, Intune, Outlook/Exchange, OneDrive, OneNote, SharePoint, Planner, and more through a single endpoint. 		
Azure Batch Schedule large-scale parallel and HPC applications in the cloud	Azure Data Catalog Programmatic access to Data Catalog resources to register, annotate and search data assets	Azure Data Explorer Perform ad-hoc queries on terabytes of data to build near real-time and complex analytics solutions
Azure Data Lake Access to storage and compute for big data analytic scenarios	Azure DevOps Integrate with Azure DevOps and Azure DevOps server	Azure Import/Export Programmatic control of import/export jobs
Azure Key Vault Manage your key vaults as well as the keys, secrets, and certificates within your Key Vaults	Azure Rights Management Services Allow validated users to read and write protected content	Azure Service Management Programmatic access to much of the functionality available through the Azure portal
Azure Storage Secure, massively scalable object and data lake storage for unstructured and semi-structured data	Customer Insights Create profile and interaction models for your products	Data Export Service for Microsoft Dynamics 365 Export data from Microsoft Dynamics CRM organization to an external destination

4. Click **Access Azure Service Management as organization users** and then click **Add permissions**.

Request API permissions

>

[All APIs](#)

What type of permissions does your application require?

Delegated permissions

Your application needs to access the API as the signed-in user.

Application permissions

Your application runs as a background service or daemon without a signed-in user.

Select permissions

[expand all](#)

Type to search	
PERMISSION	ADMIN CONSENT REQUIRED
<input checked="" type="checkbox"/> user_impersonation Access Azure Service Management as organization users (preview) <small> ⓘ</small>	-

Getting the application ID and directory ID

When you add the Azure account to Cloud Manager, you need to provide the application (client) ID and the directory (tenant) ID for the application. Cloud Manager uses the IDs to programmatically sign in.

Steps

1. In the **Azure Active Directory** service, click **App registrations** and select the application.
2. Copy the **Application (client) ID** and the **Directory (tenant) ID**.

The screenshot shows the Azure App Registrations page. At the top, there are 'Delete' and 'Endpoints' buttons. Below that is a message: 'Welcome to the new and improved App registrations. Looking to learn'. The page lists several application details:

- Display name : test-service-principal
- Application (client) ID : 73de25f9-99be-4ae0-8b24-538ca787a6b3 (highlighted with a red box)
- Directory (tenant) ID : 4b0911a0-929b-4715-944b-c03745165b3a (highlighted with a red box)
- Object ID : b37489a9-379f-49c2-b27c-e630514106a5

Creating a client secret

You need to create a client secret and then provide Cloud Manager with the value of the secret so Cloud Manager can use it to authenticate with Azure AD.



When you add the account to Cloud Manager, Cloud Manager refers to the client secret as the Application Key.

Steps

1. Open the **Azure Active Directory** service.
2. Click **App registrations** and select your application.
3. Click **Certificates & secrets > New client secret**.
4. Provide a description of the secret and a duration.
5. Click **Add**.
6. Copy the value of the client secret.

Client secrets

A secret string that the application uses to prove its identity when requesting a token. Also can be referred to as application password.

New client secret			Copy to clipboard
DESCRIPTION	EXPIRES	VALUE	
test secret	8/16/2020	*sZ1jSe2By:D*-ZRov4NLfdAcY7:+0vA	

Result

Your service principal is now setup and you should have copied the application (client) ID, the directory (tenant) ID, and the value of the client secret. You need to enter this information in Cloud Manager when you add an Azure account.

Adding the credentials to Cloud Manager

After you provide an Azure account with the required permissions, you can add the credentials for that account to Cloud Manager. Completing this step enables you to launch Cloud Volumes ONTAP using different Azure credentials.

Before you get started

If you just created these credentials in your cloud provider, it might take a few minutes until they are available for use. Wait a few minutes before you add the credentials to Cloud Manager.

What you'll need

You need to create a Connector before you can change Cloud Manager settings. [Learn how](#).

Steps

1. In the upper right of the Cloud Manager console, click the Settings icon, and select **Credentials**.



2. Click **Add Credentials** and select **Microsoft Azure**.
3. Enter information about the Azure Active Directory service principal that grants the required permissions:
 - Application (client) ID: See [Getting the application ID and directory ID](#).
 - Directory (tenant) ID: See [Getting the application ID and directory ID](#).
 - Client Secret: See [Creating a client secret](#).

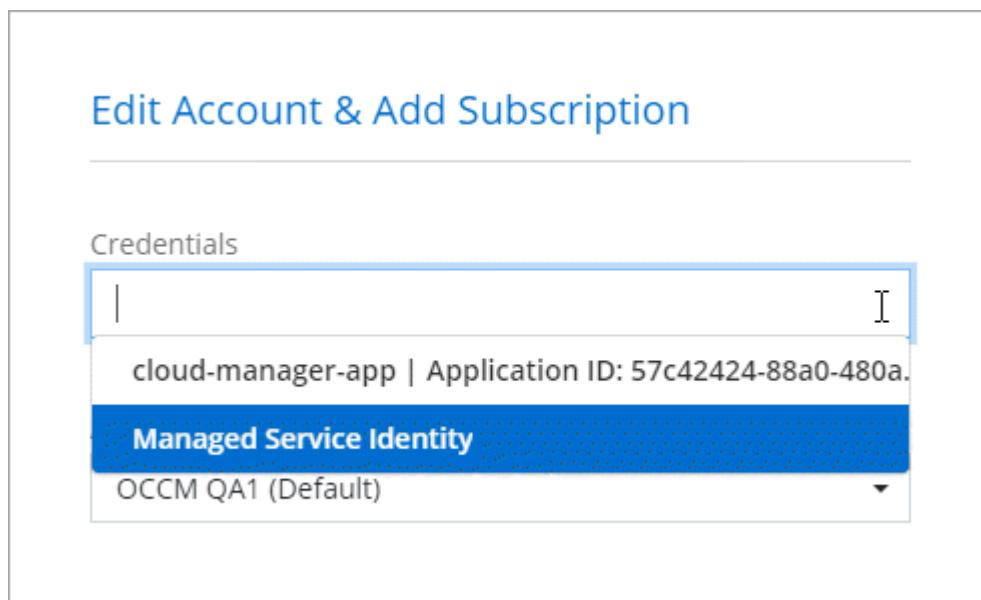
4. Confirm that the policy requirements have been met and then click **Continue**.
5. Choose the pay-as-you-go subscription that you want to associate with the credentials, or click **Add Subscription** if you don't have one yet.

To create a pay-as-you-go Cloud Volumes ONTAP system, Azure credentials must be associated with a subscription to Cloud Volumes ONTAP from the Azure Marketplace.

6. Click **Add**.

Result

You can now switch to different set of credentials from the Details and Credentials page [when creating a new working environment](#):



Associating an Azure Marketplace subscription to credentials

After you add your Azure credentials to Cloud Manager, you can associate an Azure Marketplace subscription to those credentials. The subscription enables you to create a pay-as-you-go Cloud Volumes ONTAP system, and to use other NetApp cloud services.

There are two scenarios in which you might associate an Azure Marketplace subscription after you've already added the credentials to Cloud Manager:

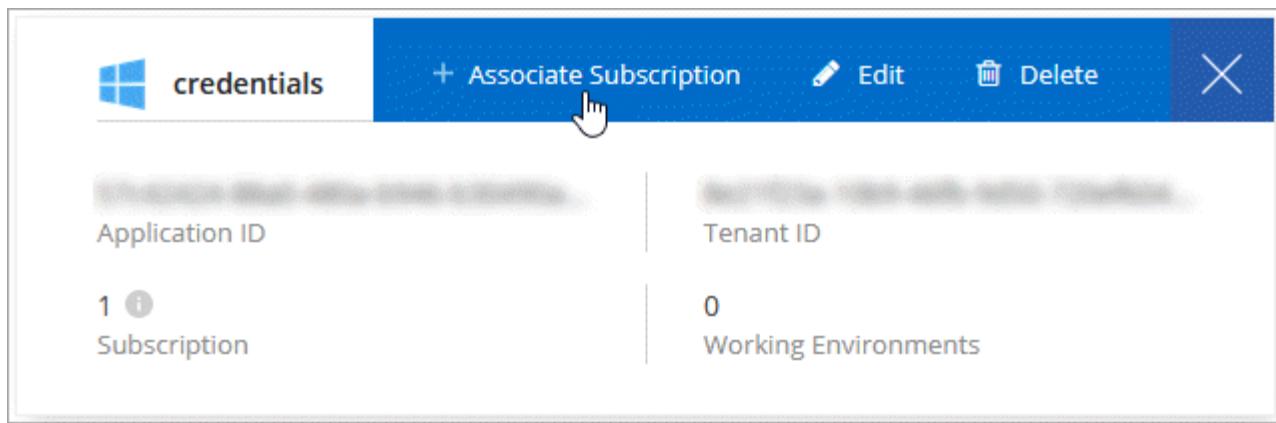
- You didn't associate a subscription when you initially added the credentials to Cloud Manager.
- You want to replace an existing Azure Marketplace subscription with a new subscription.

What you'll need

You need to create a Connector before you can change Cloud Manager settings. [Learn how](#).

Steps

1. In the upper right of the Cloud Manager console, click the Settings icon, and select **Credentials**.
2. Hover over a set of credentials and click the action menu.
3. From the menu, click **Associate Subscription**.



4. Select a subscription from the down-down list or click **Add Subscription** and follow the steps to create a new subscription.

The following video starts from the context of the working environment wizard, but shows you the same workflow after you click **Add Subscription**:

► https://docs.netapp.com/us-en/occm//media/video_subscribing_azure.mp4 (video)

Associating additional Azure subscriptions with a managed identity

Cloud Manager enables you to choose the Azure credentials and Azure subscription in which you want to deploy Cloud Volumes ONTAP. You can't select a different Azure subscription for the managed identity profile unless you associate the [managed identity](#) with those subscriptions.

About this task

A managed identity is [the initial Azure account](#) when you deploy a Connector from Cloud Manager. When you deployed the Connector, Cloud Manager created the Cloud Manager Operator role and assigned it to the Connector virtual machine.

Steps

1. Log in to the Azure portal.
2. Open the **Subscriptions** service and then select the subscription in which you want to deploy Cloud Volumes ONTAP.
3. Click **Access control (IAM)**.

- a. Click **Add > Add role assignment** and then add the permissions:

- Select the **Cloud Manager Operator** role.



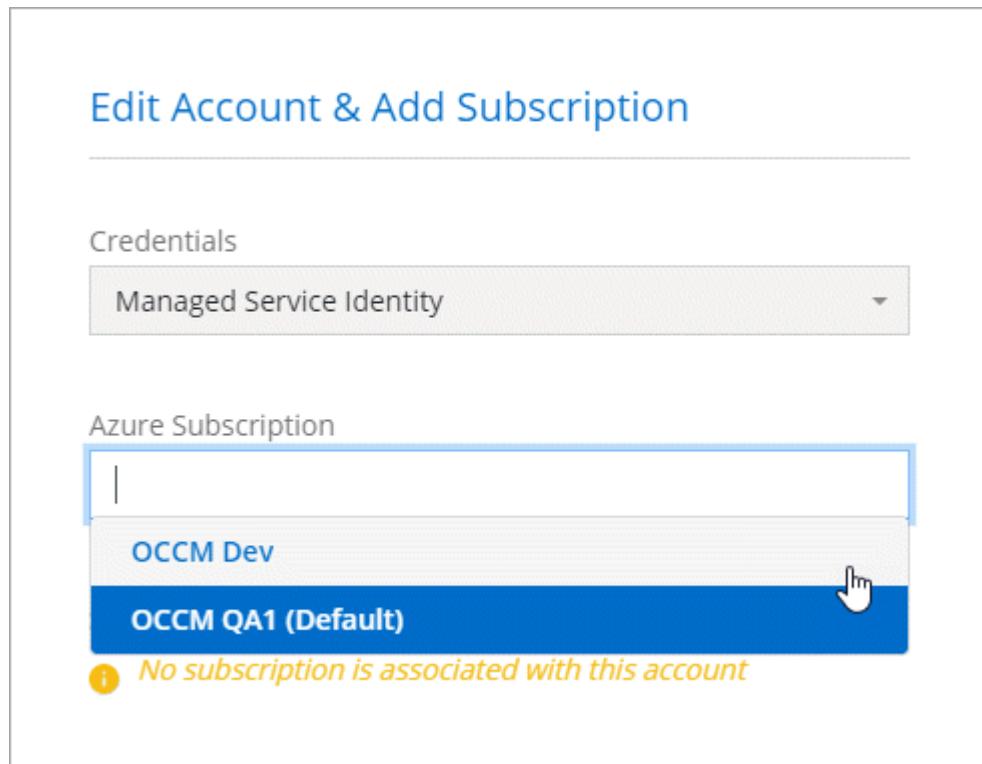
Cloud Manager Operator is the default name provided in the [Cloud Manager policy](#). If you chose a different name for the role, then select that name instead.

- Assign access to a **Virtual Machine**.
- Select the subscription in which the Connector virtual machine was created.
- Select the Connector virtual machine.
- Click **Save**.

4. Repeat these steps for additional subscriptions.

Result

When you create a new working environment, you should now have the ability to select from multiple Azure subscriptions for the managed identity profile.



GCP

Google Cloud projects, permissions, and accounts

A service account provides Cloud Manager with permissions to deploy and manage Cloud Volumes ONTAP systems that are in the same project as the Connector, or in different projects.

Project and permissions for Cloud Manager

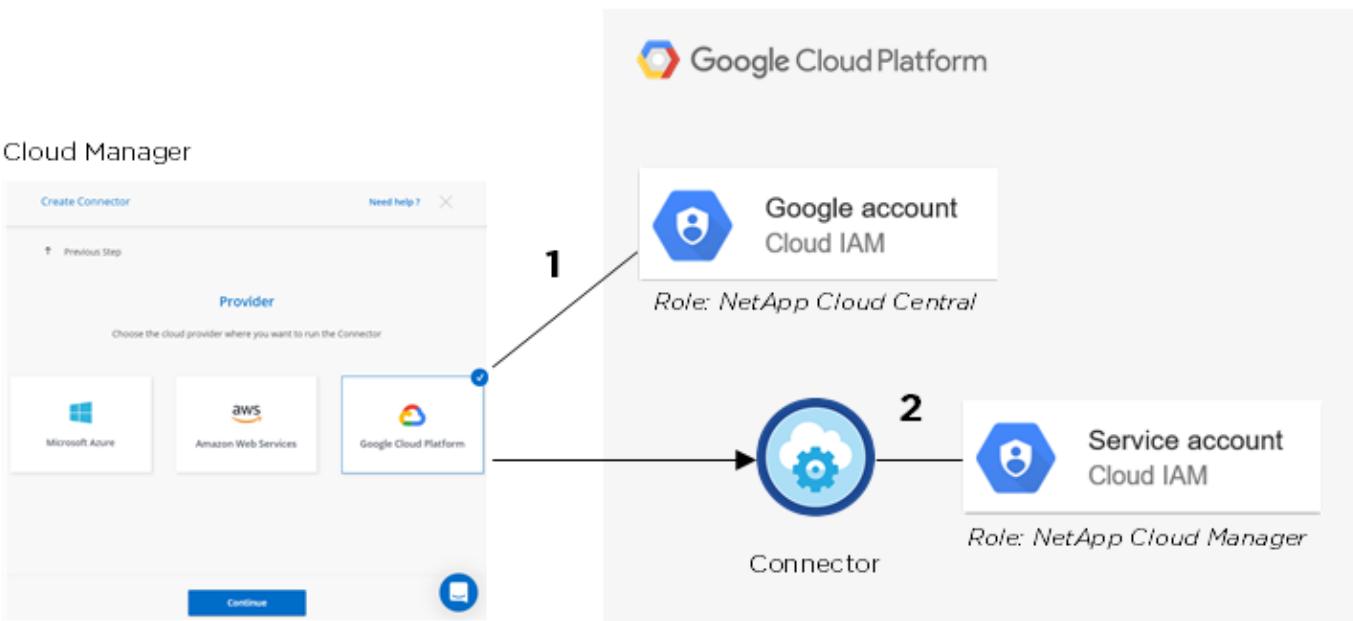
Before you can deploy Cloud Volumes ONTAP in Google Cloud, you must first deploy a Connector in a Google Cloud project. The Connector can't be running on your premises, or in a different cloud provider.

Two sets of permissions must be in place before you deploy a Connector directly from Cloud Manager:

1. You need to deploy a Connector using a Google account that has permissions to launch the Connector VM instance from Cloud Manager.
2. When deploying the Connector, you are prompted to select a [service account](#) for the VM instance. Cloud Manager gets permissions from the service account to create and manage Cloud Volumes ONTAP systems on your behalf. Permissions are provided by attaching a custom role to the service account.

We have set up two YAML files that include the required permissions for the user and the service account. [Learn how to use the YAML files to set up permissions.](#)

The following image depicts the permission requirements described in numbers 1 and 2 above:



Project for Cloud Volumes ONTAP

Cloud Volumes ONTAP can reside in the same project as the Connector, or in a different project. To deploy Cloud Volumes ONTAP in a different project, you need to first add the Connector service account and role to that project.

- [Learn how to set up service account \(see step 2\).](#)
- [Learn how to deploy Cloud Volumes ONTAP in GCP and select a project.](#)

Account for data tiering



Cloud Manager requires a GCP account for Cloud Volumes ONTAP 9.6, but not for 9.7 and later. If you want to use data tiering with Cloud Volumes ONTAP 9.7 or later, then follow step 4 in [Getting started with Cloud Volumes ONTAP in Google Cloud Platform](#).

Adding a Google Cloud account to Cloud Manager is required to enable data tiering on a Cloud Volumes ONTAP 9.6 system. Data tiering automatically tiers cold data to low-cost object storage, enabling you to reclaim space on your primary storage and shrink secondary storage.

When you add the account, you need to provide Cloud Manager with a storage access key for a service account that has Storage Admin permissions. Cloud Manager uses the access keys to set up and manage a Cloud Storage bucket for data tiering.

After you add a Google Cloud account, you can then enable data tiering on individual volumes when you create, modify, or replicate them.

- [Learn how to set up and add GCP accounts to Cloud Manager.](#)
- [Learn how to tier inactive data to low-cost object storage.](#)

Managing GCP credentials and subscriptions for Cloud Manager

You can manage two types of Google Cloud Platform credentials from Cloud Manager: the credentials that are associated with the Connector VM instance and storage access keys used with a Cloud Volumes ONTAP 9.6 system for [data tiering](#).

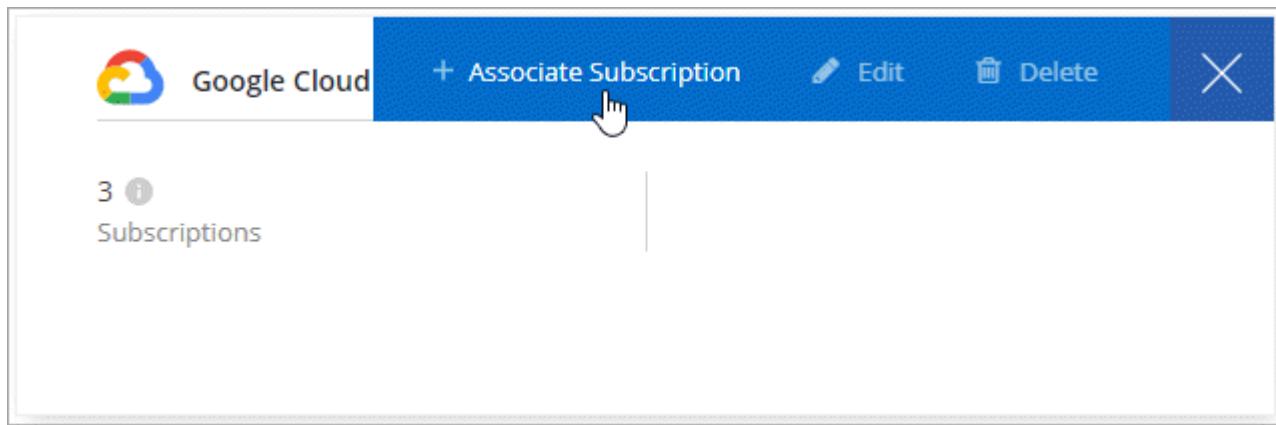
Associating a Marketplace subscription with GCP credentials

When you deploy a Connector in GCP, Cloud Manager creates a default set of credentials that are associated with the Connector VM instance. These are the credentials that Cloud Manager uses to deploy Cloud Volumes ONTAP.

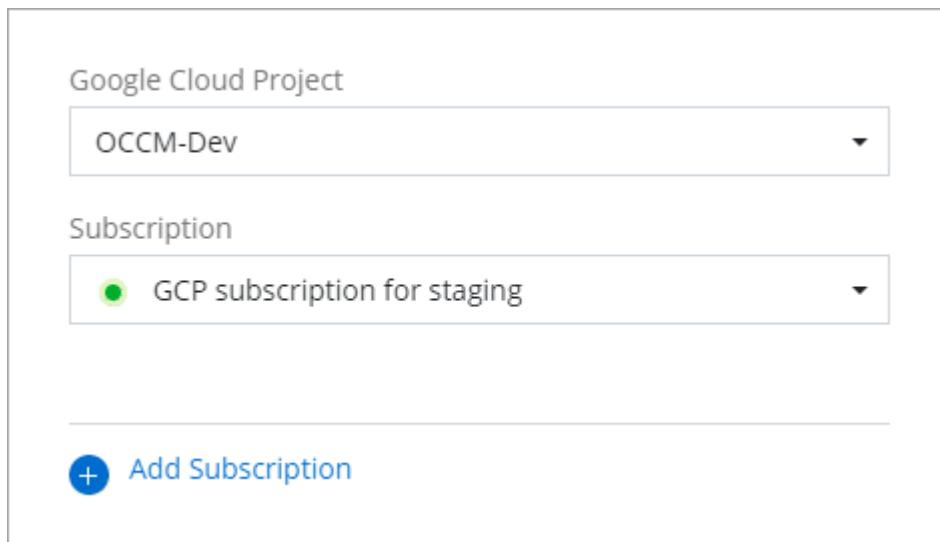
At any time, you can change the Marketplace subscription that's associated with these credentials. The subscription enables you to create a pay-as-you-go Cloud Volumes ONTAP system, and to use other NetApp cloud services.

Steps

1. In the upper right of the Cloud Manager console, click the Settings icon, and select **Credentials**.
2. Hover over a set of credentials and click the action menu.
3. From the menu, click **Associate Subscription**.



4. Select a Google Cloud project and subscription from the down-down list or click **Add Subscription** and follow the steps to create a new subscription.



5. Click **Associate**.

Setting up and adding GCP accounts for data tiering with Cloud Volumes ONTAP 9.6

If you want to enable a Cloud Volumes ONTAP 9.6 system for [data tiering](#), you need to provide Cloud Manager with a storage access key for a service account that has Storage Admin permissions. Cloud Manager uses the

access keys to set up and manage a Cloud Storage bucket for data tiering.



If you want to use data tiering with Cloud Volumes ONTAP 9.7 or later, then follow step 4 in [Getting started with Cloud Volumes ONTAP in Google Cloud Platform](#).

Setting up a service account and access keys for Google Cloud Storage

A service account enables Cloud Manager to authenticate and access Cloud Storage buckets used for data tiering. The keys are required so that Google Cloud Storage knows who is making the request.

Steps

1. Open the GCP IAM console and [create a service account](#) that has the Storage Admin role.

The screenshot shows the 'Create service account' dialog in the GCP IAM console. It's divided into three main sections: 'Service account details' (step 1), 'Grant this service account access to project (optional)' (step 2), and a 'Select a role' dropdown (step 3). The 'storage admin' role is selected in the dropdown. The 'Storage Admin' role description is visible, stating: 'Full control of GCS resources.' A cursor icon is hovering over the 'Storage Admin' role entry. At the bottom left of the dialog, there is a 'DONE' button.

2. Go to [GCP Storage Settings](#).
3. If you're prompted, select a project.
4. Click the **Interoperability** tab.
5. If you haven't already done so, click **Enable interoperability access**.

6. Under **Access keys for service accounts**, click **Create a key for a service account**.
7. Select the service account that you created in step 1.

Select a service account

Search by prefix...

Email	Name	Keys
<input checked="" type="radio"/> data-tiering-for-netapp@top-monitor-250116.iam.gserviceaccount.com	data tiering for netapp	-

CANCEL CREATE KEY | CREATE NEW ACCOUNT

8. Click **Create Key**.
9. Copy the access key and secret.

You'll need to enter this information in Cloud Manager when you add the GCP account for data tiering.

Adding a GCP account to Cloud Manager

Now that you have an access key for a service account, you can add it to Cloud Manager.

What you'll need

You need to create a Connector before you can change Cloud Manager settings. [Learn how](#).

Steps

1. In the upper right of the Cloud Manager console, click the Settings icon, and select **Credentials**.



2. Click **Add Credentials** and select **Google Cloud**.
3. Enter the access key and secret for the service account.

The keys enable Cloud Manager to set up a Cloud Storage bucket for data tiering.

4. Confirm that the policy requirements have been met and then click **Create Account**.

What's next?

You can now enable data tiering on individual volumes on a Cloud Volumes ONTAP 9.6 system when you create, modify, or replicate them. For details, see [Tiering inactive data to low-cost object storage](#).

But before you do, be sure that the subnet in which Cloud Volumes ONTAP resides is configured for Private Google Access. For instructions, refer to [Google Cloud Documentation: Configuring Private Google Access](#).

Adding NetApp Support Site accounts to Cloud Manager

Adding your NetApp Support Site account to Cloud Manager is required to deploy a node-based BYOL system or to enter the serial number for a capacity-based license. It's also required to register pay-as-you-go systems and to upgrade ONTAP software.

Steps

1. If you don't have a NetApp Support Site account yet, [register for one](#).
2. In the upper right of the Cloud Manager console, click the Help icon, and select **Support**.
3. Click **NSS Management > Add NSS Account**.
4. Enter the user name and password for your account.
 - The account must be a customer-level account (not a guest or temp account).
 - If you plan to deploy a node-based BYOL system:
 - The account must be authorized to access the serial numbers of the BYOL systems.
 - If you purchased a secure BYOL subscription, then a secure NSS account is required.
5. Click **Register**.

What's next?

Users can now select the account when creating new Cloud Volumes ONTAP systems and when registering existing systems.

- [Launching Cloud Volumes ONTAP in AWS](#)
- [Launching Cloud Volumes ONTAP in Azure](#)
- [Registering pay-as-you-go systems](#)
- [Learn how Cloud Manager manages license files](#)

Managing your account

After you perform initial setup, you can administer your account settings later by managing users, service accounts, workspaces, Connectors, and subscriptions.

[Learn more about how Cloud Central accounts work](#).

Managing your account with the Tenancy API

If you want to manage your account settings by sending API requests, then you'll need to use the *Tenancy API*. This API is different than the Cloud Manager API, which you use to create and manage working environments.

[View endpoints for the Tenancy API](#).

Creating and managing users

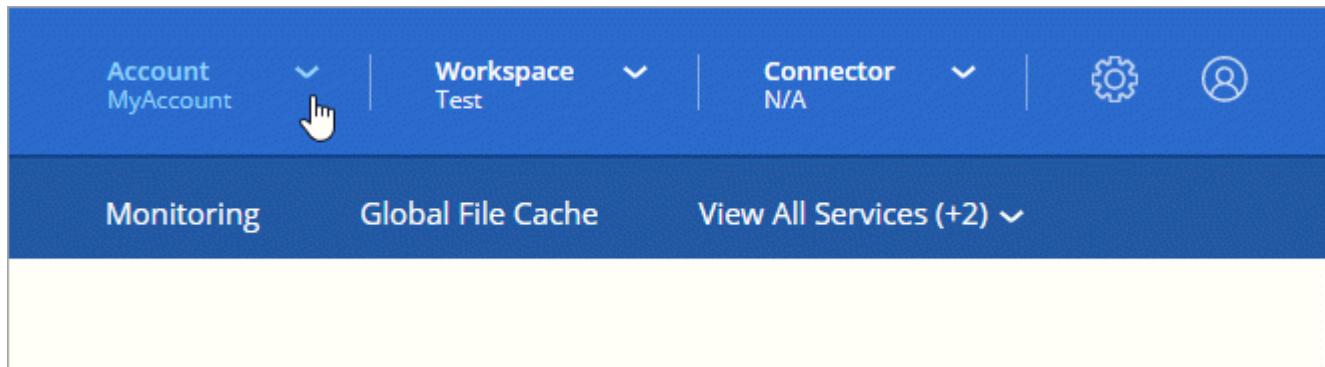
The user's in your account can access the manage the resources in your account's workspaces.

Adding users

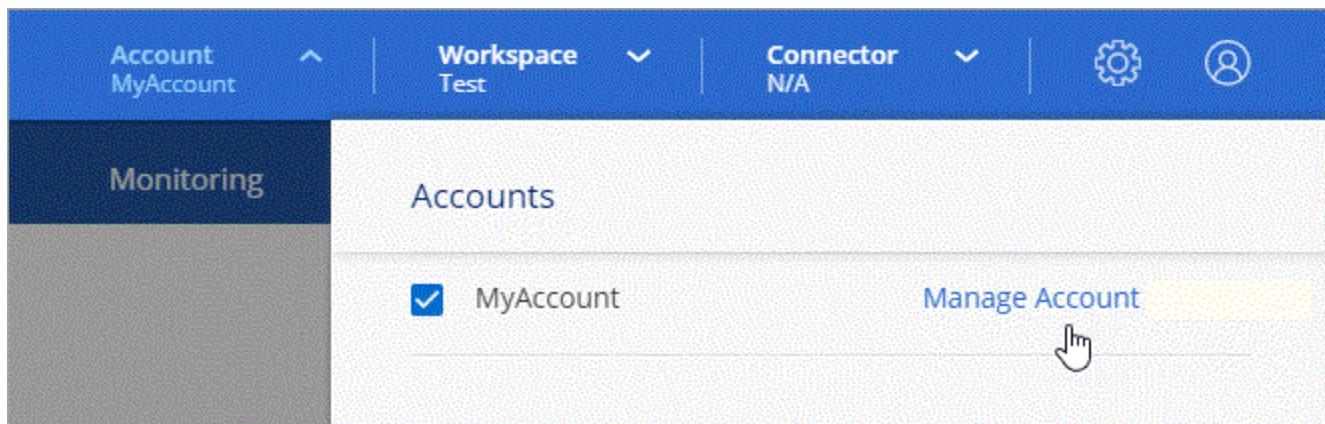
Associate Cloud Central users with the Cloud Central account so those users can create and manage working environments in Cloud Manager.

Steps

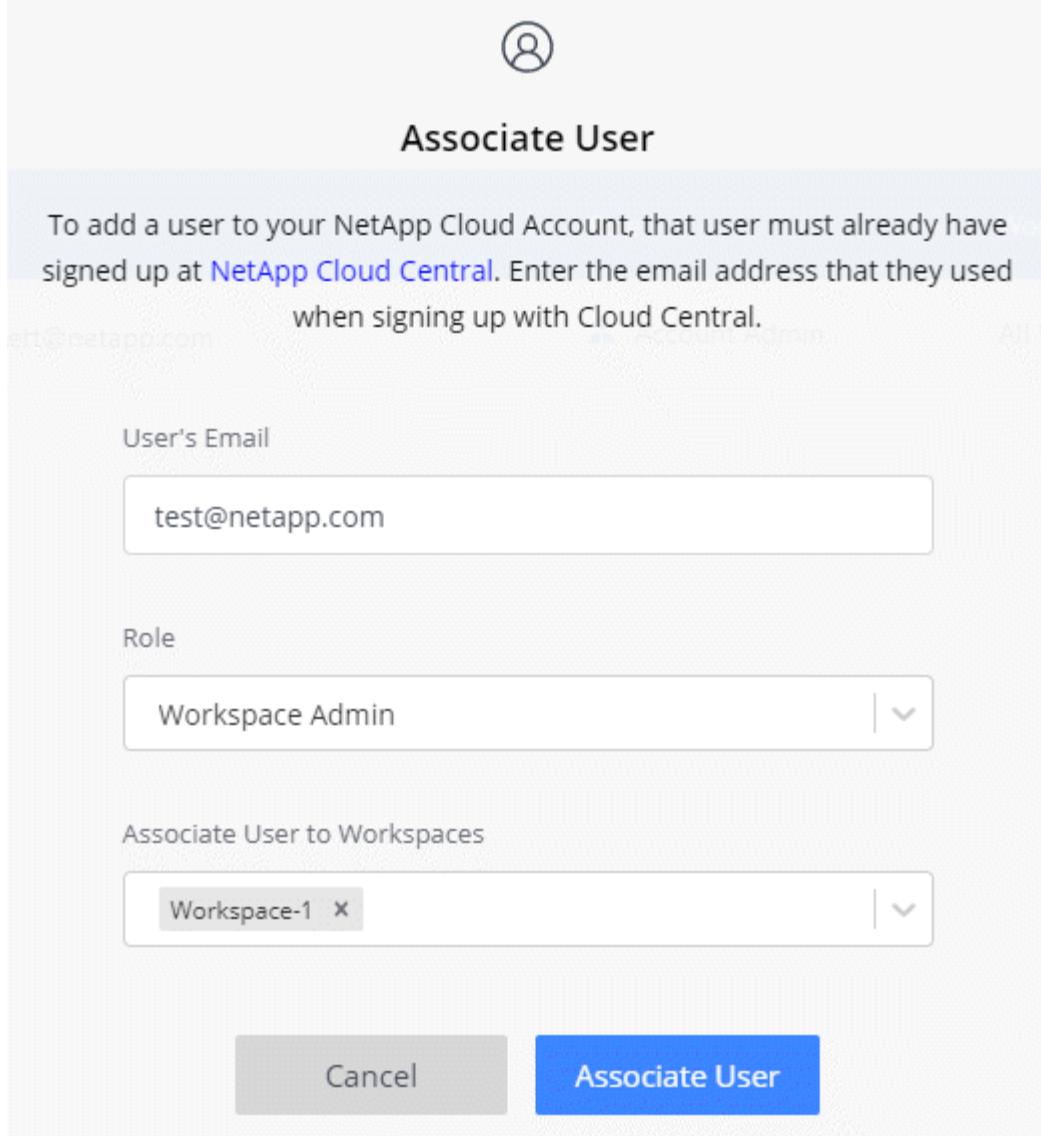
1. If the user hasn't already done so, ask the user to go to [NetApp Cloud Central](#) and sign up.
2. From the top of Cloud Manager, click the **Account** drop-down.



3. Click **Manage Account** next to the currently selected account.



4. From the Members tab, click **Associate User**.
5. Enter the user's email address and select a role for the user:
 - **Account Admin:** Can perform any action in Cloud Manager.
 - **Workspace Admin:** Can create and manage resources in assigned workspaces.
 - **Compliance Viewer:** Can only view Cloud Data Sense compliance information and generate reports for workspaces that they have permission to access.
 - **SnapCenter Admin:** Can use the SnapCenter Service to create application consistent backups and restore data using those backups. *This service is currently in Beta.*
6. If you selected Workspace Admin or Compliance Viewer, select one or more workspaces to associate with that user.



The screenshot shows a 'Associate User' dialog box. At the top is a user icon. Below it is the title 'Associate User'. A descriptive text explains that the user must already be signed up at NetApp Cloud Central and provides a link to the service. The main area contains fields for 'User's Email' (containing 'test@netapp.com') and 'Role' (set to 'Workspace Admin'). Below these is a section titled 'Associate User to Workspaces' with a dropdown menu showing 'Workspace-1' and an 'X' button. At the bottom are two buttons: 'Cancel' and a blue 'Associate User' button.

To add a user to your NetApp Cloud Account, that user must already have signed up at [NetApp Cloud Central](#). Enter the email address that they used when signing up with Cloud Central.

User's Email

test@netapp.com

Role

Workspace Admin

Associate User to Workspaces

Workspace-1 X

Cancel Associate User

7. Click **Associate**.

Result

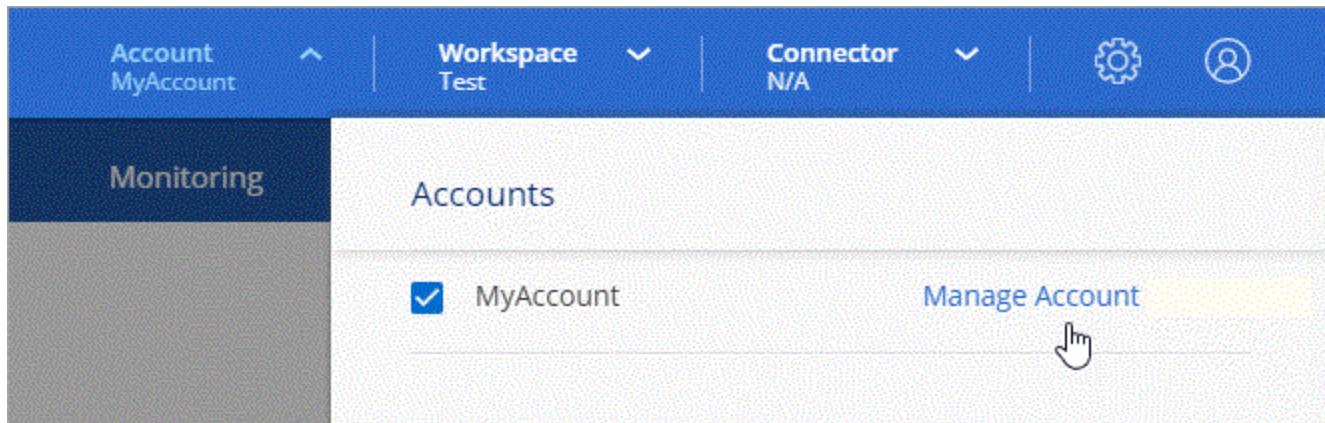
The user should receive an email from NetApp Cloud Central titled "Account Association." The email includes the information needed to access Cloud Manager.

Removing users

Disassociating a user makes it so they can no longer access the resources in a Cloud Central account.

Steps

1. From the top of Cloud Manager, click the **Account** drop-down and click **Manage Account**.



- From the Members tab, click the action menu in the row that corresponds to the user.

Type	Name	Email	Role	Workspace	Action
👤	Ben	[REDACTED]	⭐ Account Admin	All Workspaces	...
👤	Tom	[REDACTED]	⭐ Account Admin	All Workspaces	...
👤	Ben	[REDACTED]	Workspace Admin	Newone	

- Click **Disassociate User** and click **Disassociate** to confirm.

Result

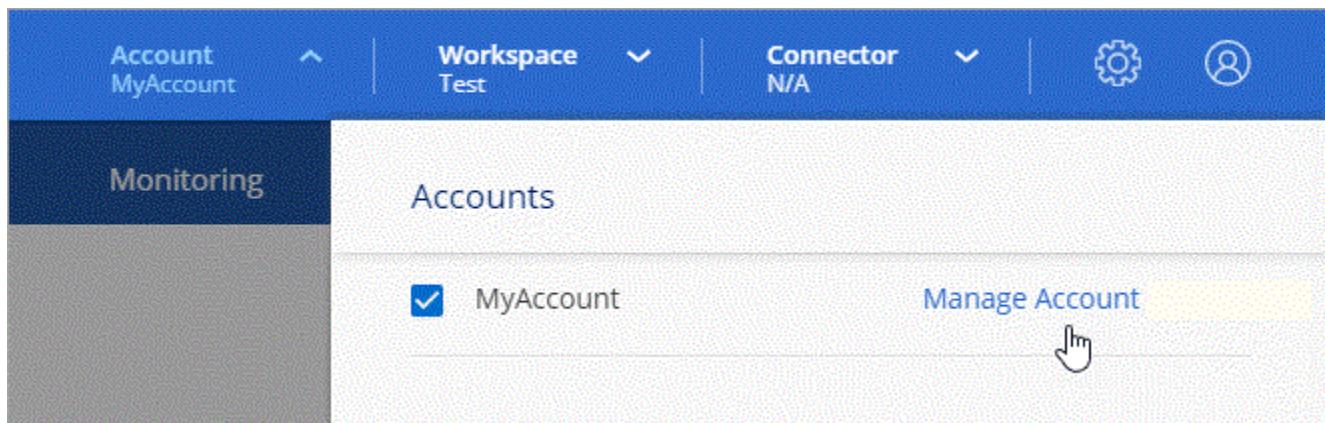
The user can no longer access the resources in this Cloud Central account.

Managing a Workspace Admin's workspaces

You can associate and disassociate Workspace Admins with workspaces at any time. Associating the user enables them to create and view the working environments in that workspace.

Steps

- From the top of Cloud Manager, click the **Account** drop-down and click **Manage Account**.



- From the Members tab, click the action menu in the row that corresponds to the user.

Type	Name	Email	Role	Workspace	
👤	Ben	[REDACTED]	⭐ Account Admin	All Workspaces	...
👤	Tom	[REDACTED]	⭐ Account Admin	All Workspaces	...
👤	Ben	[REDACTED]	Workspace Admin	Newone	

3. Click **Manage Workspaces**.

4. Select the workspaces to associate with the user and click **Apply**.

Result

The user can now access those workspaces from Cloud Manager, as long as the Connector was also associated with the workspaces.

Creating and managing service accounts

A service account acts as a "user" that can make authorized API calls to Cloud Manager for automation purposes. This makes it easier to manage automation because you don't need to build automation scripts based on a real person's user account who can leave the company at any time. And if you're using federation, you can create a token without generating a refresh token from the cloud.

You give permissions to a service account by assigning it a role, just like any other Cloud Manager user. You can also associate the service account with specific workspaces in order to control the working environments (resources) that the service can access.

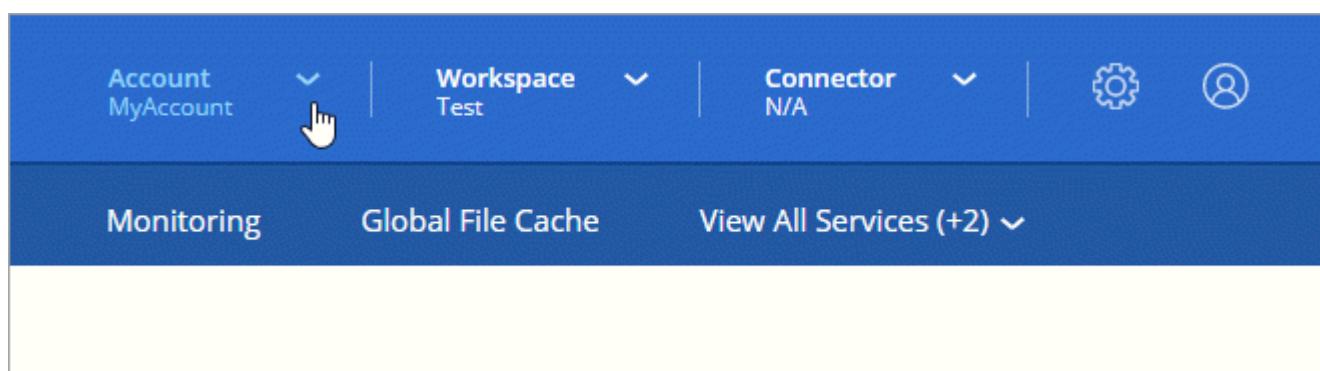
When you create the service account, Cloud Manager enables you to copy or download a client ID and client secret for the service account. This key pair is used for authentication with Cloud Manager.

Creating a service account

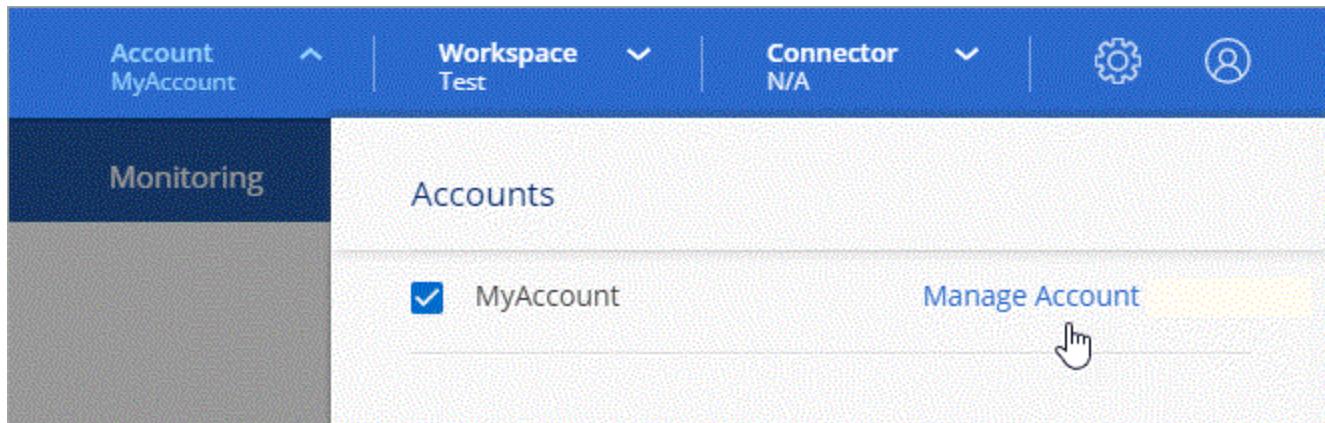
Create as many service accounts as you need to manage the resources in your working environments.

Steps

1. From the top of Cloud Manager, click the **Account** drop-down.



2. Click **Manage Account** next to the currently selected account.



3. From the Members tab, click **Create Service Account**.
4. Enter a name and select a role. If you chose a role other than Account Admin, choose the workspace to associate with this service account.
5. Click **Create**.
6. Copy or download the client ID and client secret.

The client secret is visible only once and is not stored anywhere by Cloud Manager. Copy or download the secret and store it safely.

7. Click **Close**.

Obtaining a bearer token for a service account

In order to make API calls to the [Tenancy API](#), you'll need to obtain a bearer token for a service account.

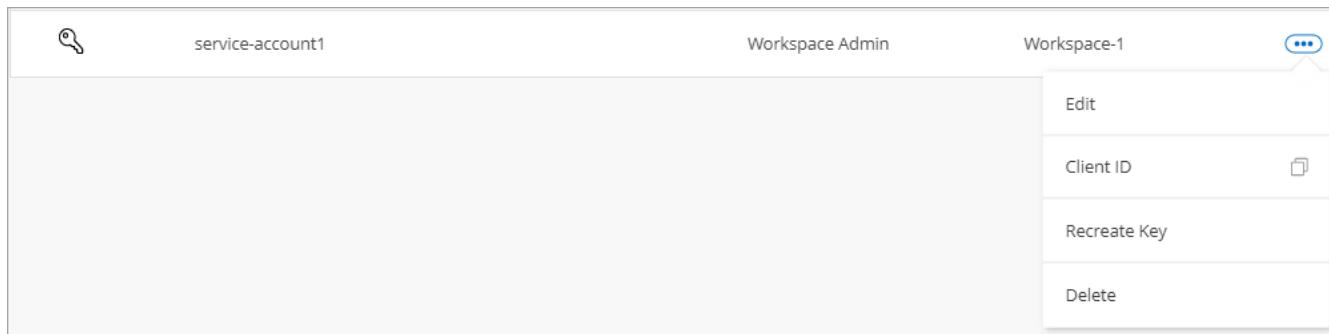
```
curl --location --request POST 'https://netapp-cloud-
account.auth0.com/oauth/token' \
--header 'Content-Type: application/json' \
--data-raw '{
  "grant_type": "client_credentials",
  "client_secret": "<client secret>",
  "audience": "https://api.cloud.netapp.com",
  "client_id": "<client id>"
}'
```

Copying the client ID

You can copy a service account's client ID at any time.

Steps

1. From the Members tab, click the action menu in the row that corresponds to the service account.



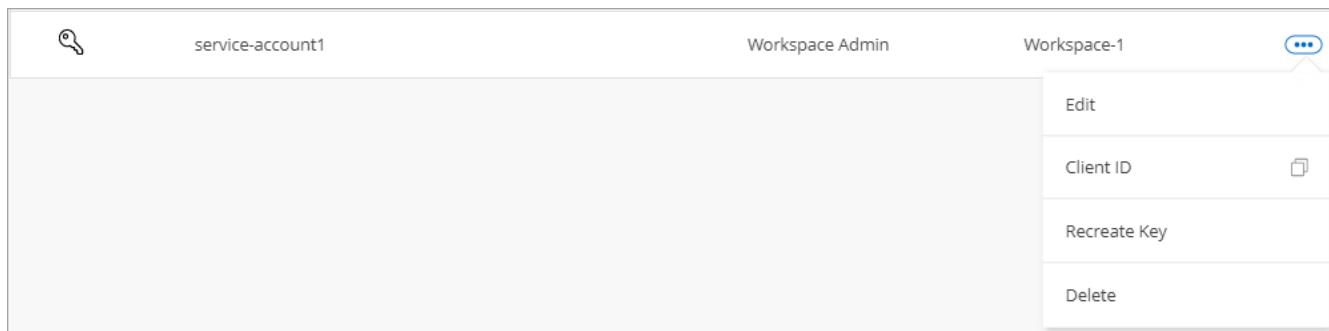
2. Click **Client ID**.
3. The ID is copied to your clipboard.

Recreating keys

Recreating the key will delete the existing key for this service account and then create a new key. You won't be able to use the previous key.

Steps

1. From the Members tab, click the action menu in the row that corresponds to the service account.



2. Click **Recreate Key**.
3. Click **Recreate** to confirm.
4. Copy or download the client ID and client secret.

The client secret is visible only once and is not stored anywhere by Cloud Manager. Copy or download the secret and store it safely.

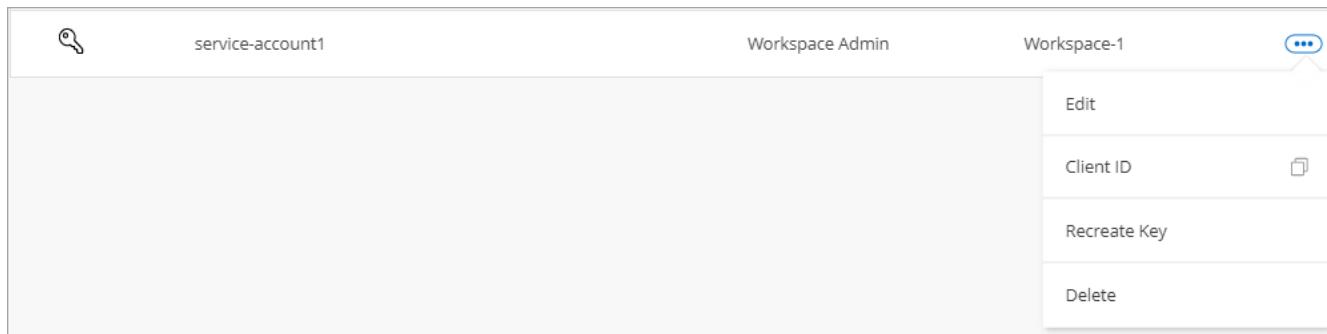
5. Click **Close**.

Deleting a service account

Delete a service account if you no longer need to use it.

Steps

1. From the Members tab, click the action menu in the row that corresponds to the service account.



2. Click **Delete**.
3. Click **Delete** again to confirm.

Managing workspaces

Manage your workspaces by creating, renaming, and deleting them. Note that you can't delete a workspace if it contains any resources. It must be empty.

Steps

1. From the top of Cloud Manager, click the **Account** drop-down and click **Manage Account**.
2. Click **Workspaces**.
3. Choose one of the following options:
 - Click **Add New Workspace** to create a new workspace.
 - Click **Rename** to rename the workspace.
 - Click **Delete** to delete the workspace.

Managing a Connector's workspaces

You need to associate the Connector with workspaces so Workspace Admins can access those workspaces from Cloud Manager.

If you only have Account Admins, then associating the Connector with workspaces isn't required. Account Admins have the ability to access all workspaces in Cloud Manager by default.

[Learn more about users, workspaces, and Connectors.](#)

Steps

1. From the top of Cloud Manager, click the **Account** drop-down and click **Manage Account**.
2. Click **Connector**.
3. Click **Manage Workspaces** for the Connector that you want to associate.
4. Select the workspaces to associate with the Connector and click **Apply**.

Managing subscriptions

After you subscribe from a cloud provider's marketplace, each subscription is available from the Account Settings widget. You have the option to rename a subscription and to disassociate the subscription from one or more accounts.

For example, let's say that you have two accounts and each is billed through separate subscriptions. You might

disassociate a subscription from one of the accounts so the users in that account don't accidentally choose the wrong subscription when creating a Cloud Volume ONTAP working environment.

[Learn more about subscriptions.](#)

Steps

1. From the top of Cloud Manager, click the **Account** drop-down and click **Manage Account**.
2. Click **Subscriptions**.

You'll only see the subscriptions that are associated with the account that you're currently viewing.

3. Click the action menu in the row that corresponds to the subscription that you want to manage.

The screenshot shows a table with two rows of subscription data. The columns are Name, Service, Cloud Provider, and Status. The first row has a blue header bar above it. The second row has a grey header bar above it. A context menu is open over the second row, with the following options: 'Rename Subscription' and 'Manage Accounts'. The 'Cloud Provider' column for both rows shows the AWS logo.

Name	Service	Cloud Provider	Status
QA Subscription	test-service	aws	Unsubscribed
metering service subscription QA !!!!	cloud-volumes-ontap	aws	Subscribed

4. Choose to rename the subscription or to manage the accounts that are associated with the subscription.

Changing your account name

Change your account name at any time to change it to something meaningful for you.

Steps

1. From the top of Cloud Manager, click the **Account** drop-down and click **Manage Account**.
2. In the **Overview** tab, click the edit icon next to the account name.
3. Type a new account name and click **Save**.

Allowing private previews

Allow private previews in your account to get access to new NetApp cloud services that are made available as a preview in Cloud Manager.

Services in private preview are not guaranteed to behave as expected and might sustain outages and be missing functionality.

Steps

1. From the top of Cloud Manager, click the **Account** drop-down and click **Manage Account**.
2. In the **Overview** tab, enable the **Allow Private Preview** setting.

Allowing third-party services

Allow third-party services in your account to get access to third-party services that are available in Cloud Manager. Third-party services are cloud services similar to the services that NetApp offers, but they're

managed and supported by third-party companies.

Steps

1. From the top of Cloud Manager, click the **Account** drop-down and click **Manage Account**.
2. In the **Overview** tab, enable the **Allow Third Party Services** setting.

Disabling the SaaS platform

We don't recommend disabling the SaaS platform unless you need to in order to comply with your company's security policies. Disabling the SaaS platform limits your ability to use NetApp's integrated cloud services.

The following services aren't available from Cloud Manager if you disable the SaaS platform:

- Cloud Data Sense
- Kubernetes
- Cloud Tiering
- Global File Cache

If you do disable the SaaS platform, you'll need to perform all tasks from [the local user interface that is available on a Connector](#).

This is an irreversible action that will prevent you from using the Cloud Manager SaaS platform. You'll need to perform actions from the local Connector. You won't have the ability to use many of NetApp's integrated cloud services, and re-enabling the SaaS platform will require the help of NetApp support.

Steps

1. From the top of Cloud Manager, click the **Account** drop-down and click **Manage Account**.
2. In the Overview tab, toggle the option to disable use of the SaaS platform.

Monitoring operations in your account

You can monitor the status of the operations that Cloud Manager is performing to see if there are any issues that you need to address. You can view the status in the Notification Center or in the Timeline.

This table provides a comparison of the Notification Center and the Timeline so you can understand what each has to offer.

Notification Center	Timeline
Shows high level status for events and actions	Provides details for each event or action for further investigation
Shows status for the current login session - the information won't appear in the Notification Center after you log off	Retains status for up to the last month
Shows only actions initiated in the user interface	Shows all actions from the UI or APIs

Notification Center	Timeline
Shows user-initiated actions	Shows all actions, whether user-initiated or system-initiated
Filter results by importance	Filter by service, action, user, status, and more

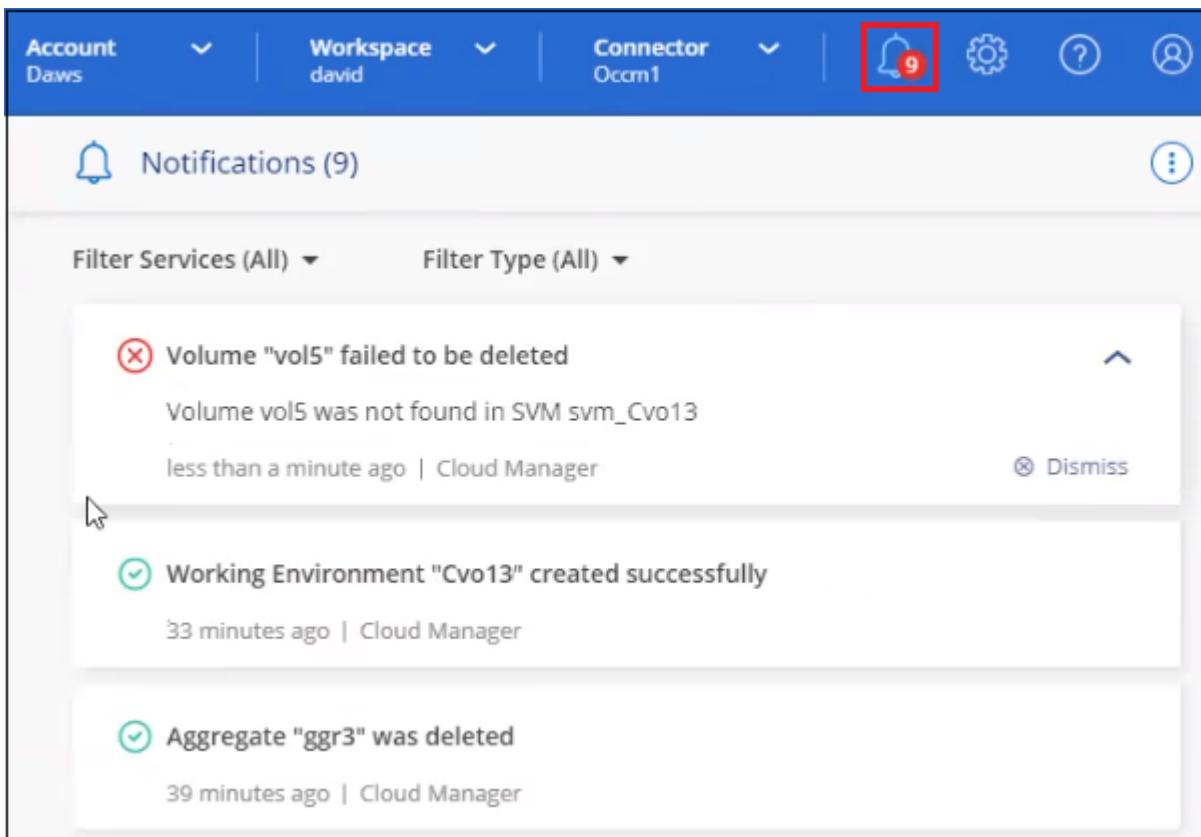
Monitoring operations status using the Notification Center

Notifications are like events where they track the progress of operations that you've initiated in Cloud Manager so you can verify whether the operation was successful, or if it failed. They enable you to view the status for Cloud Manager operations (and cloud services operations in the future) that you initiated during your current login session.

At this time, only notifications for creating and deleting the following Cloud Volumes ONTAP objects are supported:

- working environments
- aggregates
- volumes

You display the notifications by clicking the notification bell () in the menu bar. The color of the little bubble in the bell indicates the highest level severity notification that is active. So if you see a red bubble, it means there's an important notification that you should look at.

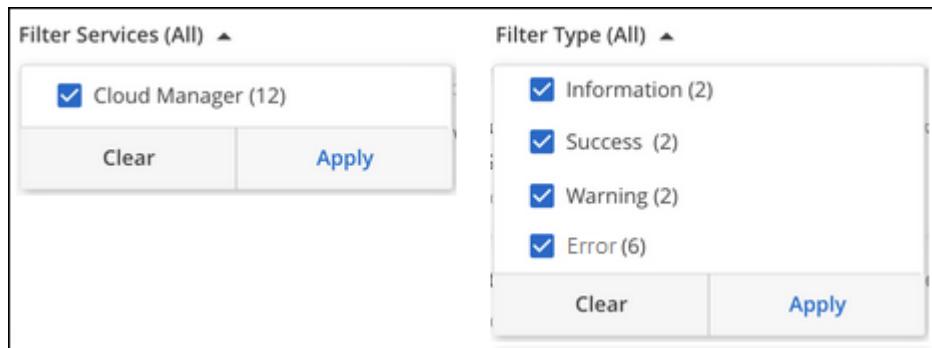


The screenshot shows the Cloud Manager interface with the Notifications panel open. The top navigation bar includes Account (Daws), Workspace (david), Connector (Ocm1), and a notification bell icon with a red bubble containing the number 9. The Notifications panel displays three notifications:

- Volume "vol5" failed to be deleted**: Volume vol5 was not found in SVM svm_Cvo13. (less than a minute ago | Cloud Manager)
- Working Environment "Cvo13" created successfully**: 33 minutes ago | Cloud Manager
- Aggregate "ggr3" was deleted**: 39 minutes ago | Cloud Manager

Filtering notifications

By default you'll see all notifications. You can filter the notifications that you see in the Notification Center to show only those notifications that are important to you. You can filter by Cloud Manager "Service" and by notification "Type".

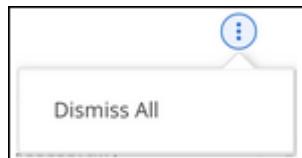


For example, if you want to see only "Error" and "Warning" notifications for Cloud Manager operations, select those entries and you'll see only those types of notifications.

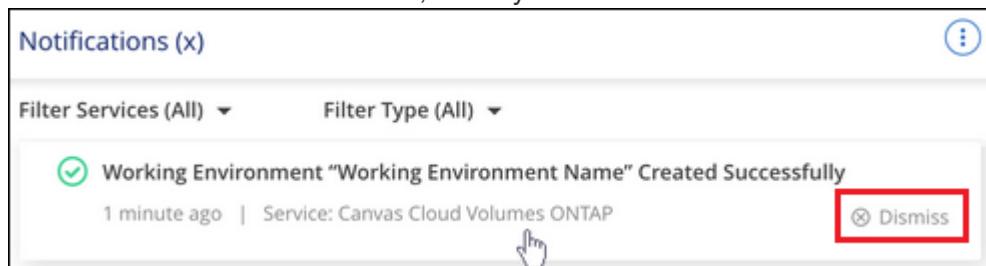
Dismissing notifications

You can remove notifications from the page if you no longer need to see them. You can dismiss all notifications at once, or you can dismiss individual notifications.

To dismiss all notifications, in the Notification Center, click : and select **Dismiss All**.



To dismiss individual notifications, hover your cursor over the notification and click **Dismiss**.



Auditing user activity in your account

The Timeline in Cloud Manager shows the actions that users completed to manage your account. This includes management actions such as associating users, creating workspaces, creating Connectors, and more.

Checking the Timeline can be helpful if you need to identify who performed a specific action, or if you need to identify the status of an action.

Steps

1. Click **All Services > Timeline**.
2. Under the Filters, click **Service**, enable **Tenancy**, and click **Apply**.

Result

The Timeline updates to show your account management actions.

Managing an HTTPS certificate for secure access

By default, Cloud Manager uses a self-signed certificate for HTTPS access to the web console. You can install a certificate signed by a certificate authority (CA), which provides better security protection than a self-signed certificate.

Before you get started

You need to create a Connector before you can change Cloud Manager settings. [Learn how](#).

Installing an HTTPS certificate

Install a certificate signed by a CA for secure access.

Steps

1. In the upper right of the Cloud Manager console, click the Settings icon, and select **HTTPS Setup**.



2. In the HTTPS Setup page, install a certificate by generating a certificate signing request (CSR) or by installing your own CA-signed certificate:

Option	Description
Generate a CSR	<ol style="list-style-type: none">a. Enter the host name or DNS of the Connector host (its Common Name), and then click Generate CSR. Cloud Manager displays a certificate signing request.b. Use the CSR to submit an SSL certificate request to a CA. The certificate must use the Privacy Enhanced Mail (PEM) Base-64 encoded X.509 format.c. Upload the certificate file and then click Install.
Install your own CA-signed certificate	<ol style="list-style-type: none">a. Select Install CA-signed certificate.b. Load both the certificate file and the private key and then click Install. The certificate must use the Privacy Enhanced Mail (PEM) Base-64 encoded X.509 format.

Result

Cloud Manager now uses the CA-signed certificate to provide secure HTTPS access. The following image shows a Cloud Manager system that is configured for secure access:

The screenshot shows the 'HTTPS Certificate' page with the following details:

- HTTPS Setup is active** (indicated by a green checkmark)
- Expiration:** Aug 15, 2029 10:09:01 am
- Issuer:** C=IL, ST=Israel, L=Tel Aviv, O=NetApp, OU=Dev, CN= Localhost, E=Admin@netapp.com
- Subject:** C=IL, ST=Israel, L=Tel Aviv, O=NetApp, OU=Dev, CN= Localhost, E=Admin@netapp.com
- Certificate:** View CSR (button)

Renewing the Cloud Manager HTTPS certificate

You should renew the Cloud Manager HTTPS certificate before it expires to ensure secure access to the Cloud Manager web console. If you do not renew the certificate before it expires, a warning appears when users access the web console using HTTPS.

Steps

1. In the upper right of the Cloud Manager console, click the Settings icon, and select **HTTPS Setup**.
Details about the Cloud Manager certificate displays, including the expiration date.
2. Click **Change Certificate** and follow the steps to generate a CSR or install your own CA-signed certificate.

Result

Cloud Manager uses the new CA-signed certificate to provide secure HTTPS access.

Removing Cloud Volumes ONTAP working environments

The Account Admin can remove a Cloud Volumes ONTAP working environment to move it to another system or to troubleshoot discovery issues.

About this task

Removing a Cloud Volumes ONTAP working environment removes it from Cloud Manager. It does not delete the Cloud Volumes ONTAP system. You can later rediscover the working environment.

Removing a working environment from Cloud Manager enables you to do the following:

- Rediscover it in another workspace
- Rediscover it from another Cloud Manager system
- Rediscover it if you had problems during the initial discovery

Steps

1. In the upper right of the Cloud Manager console, click the Settings icon, and select **Tools**.



2. From the Tools page, click **Launch**.
3. Select the Cloud Volumes ONTAP working environment that you want to remove.
4. On the Review and Approve page, click **Go**.

Result

Cloud Manager removes the working environment. Users can rediscover this working environment from the Canvas page at any time.

Configuring a Connector to use a proxy server

If your corporate policies dictate that you use a proxy server for all HTTP communication to the internet, then you must configure your Connectors to use that proxy server. The proxy server can be in the cloud or in your network.

Enable a proxy on a Connector

When you configure a Connector to use a proxy server, that Connector and the Cloud Volumes ONTAP systems that it manages (including any HA mediators), all use the proxy server.

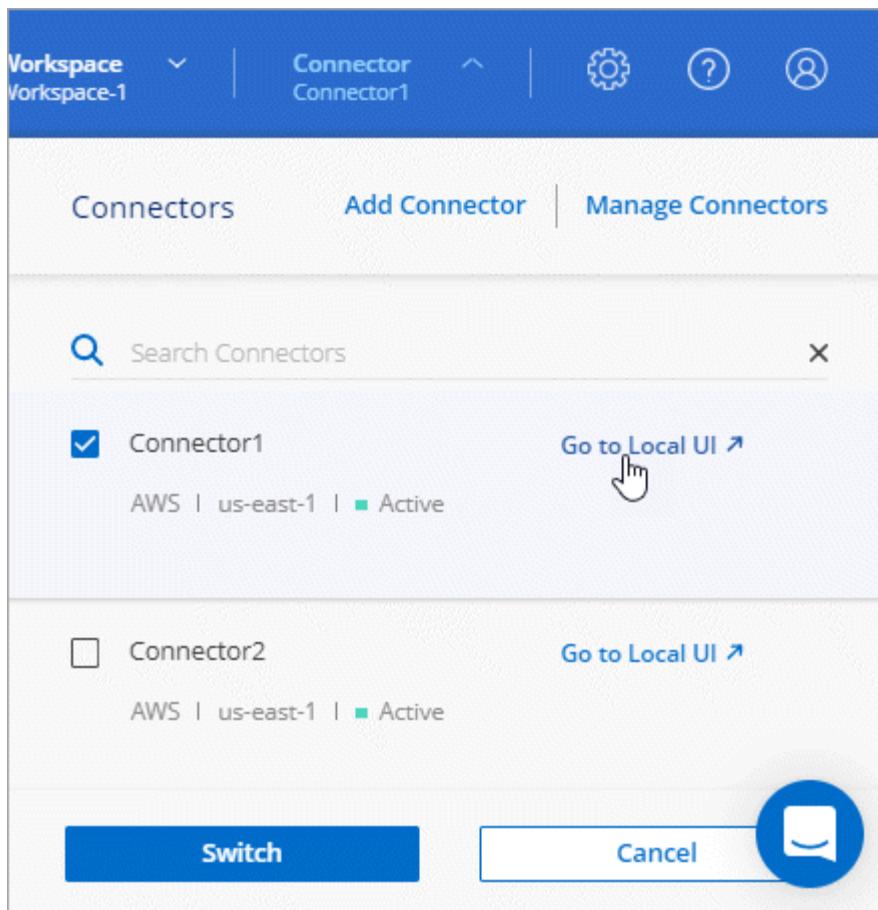
Note that this operation restarts the Connector. Ensure that the Connector isn't performing any operations before you proceed.

Steps

1. [Log in to the Cloud Manager SaaS interface](#) from a machine that has a network connection to the Connector instance.

If the Connector doesn't have a public IP address, you'll need a VPN connection or you'll need to connect from a jump host that's in the same network as the Connector.

2. Click the **Connector** drop-down and then click **Go to local UI** for a specific Connector.



The Cloud Manager interface running on the Connector loads in a new browser tab.

3. In the upper right of the Cloud Manager console, click the Settings icon, and select **Connector Settings**.



4. Under **General**, click **HTTP Proxy Configuration**.

5. Set up the proxy:

- a. Click **Enable Proxy**.
- b. Specify the server using the syntax `http://address:port`
- c. Specify a user name and password if basic authentication is required for the server
- d. Click **Save**.



Cloud Manager doesn't support passwords that include the @ character.

Result

After you specify the proxy server, new Cloud Volumes ONTAP systems are automatically configured to use the proxy server when sending AutoSupport messages. If you didn't specify the proxy server before users create Cloud Volumes ONTAP systems, then they must use System Manager to manually set the proxy server in the AutoSupport options for each system.

Enable direct API traffic

If you configured a proxy server, you can send API calls directly to Cloud Manager without going through the proxy. This option is supported with Connectors that are running in AWS, in Azure, or in Google Cloud.

Steps

1. In the upper right of the Cloud Manager console, click the Settings icon, and select **Connector Settings**.



2. Under **General**, click **Support Direct API Traffic**.
3. Click the checkbox to enable the option and then click **Save**.

Reference

Roles

The Account Admin, Workspace Admin, Compliance Viewer, and SnapCenter Admin roles provide specific permissions to users.

The Compliance Viewer role is for read-only Cloud Data Sense access.

Task	Account Admin	Workspace Admin	Compliance Viewer	SnapCenter Admin
Manage working environments	Yes	Yes	No	No
Enable services on working environments	Yes	Yes	No	No
View data replication status	Yes	Yes	No	No
View the timeline	Yes	Yes	No	No
Switch between workspaces	Yes	Yes	Yes	No
View Data Sense scan results	Yes	Yes	Yes	No
Delete working environments	Yes	No	No	No
Connect Kubernetes clusters to working environments	Yes	No	No	No
Receive the Cloud Volumes ONTAP report	Yes	No	No	No
Create Connectors	Yes	No	No	No

Task	Account Admin	Workspace Admin	Compliance Viewer	SnapCenter Admin
Manage Cloud Central accounts	Yes	No	No	No
Manage credentials	Yes	No	No	No
Modify Cloud Manager settings	Yes	No	No	No
View and manage the Support Dashboard	Yes	No	No	No
Remove working environments from Cloud Manager	Yes	No	No	No
Install an HTTPS certificate	Yes	No	No	No
Use the SnapCenter Service	Yes	Yes	No	Yes

Related links

- [Setting up workspaces and users in the Cloud Central account](#)
- [Managing workspaces and users in the Cloud Central account](#)

How Cloud Manager uses cloud provider permissions

Cloud Manager requires permissions to perform actions in your cloud provider. These permissions are included in [the policies provided by NetApp](#). You might want to understand what Cloud Manager does with these permissions.

What Cloud Manager does with AWS permissions

Cloud Manager uses an AWS account to make API calls to several AWS services, including EC2, S3, CloudFormation, IAM, the Security Token Service (STS), and the Key Management Service (KMS).

Actions	Purpose
"ec2:StartInstances", "ec2:StopInstances", "ec2:DescribeInstances", "ec2:DescribeInstanceStatus", "ec2:RunInstances", "ec2:TerminateInstances", "ec2:ModifyInstanceAttribute",	Launches a Cloud Volumes ONTAP instance and stops, starts, and monitors the instance.
"ec2:DescribeInstanceAttribute",	Verifies that enhanced networking is enabled for supported instance types.
"ec2:DescribeRouteTables", "ec2:DescribeImages",	Launches a Cloud Volumes ONTAP HA configuration.

Actions	Purpose
"ec2:CreateTags",	Tags every resource that Cloud Manager creates with the "WorkingEnvironment" and "WorkingEnvironmentId" tags. Cloud Manager uses these tags for maintenance and cost allocation.
"ec2>CreateVolume", "ec2:DescribeVolumes", "ec2:ModifyVolumeAttribute", "ec2:AttachVolume", "ec2:DeleteVolume", "ec2:DetachVolume",	Manages the EBS volumes that Cloud Volumes ONTAP uses as back-end storage.
"ec2:CreateSecurityGroup", "ec2:DeleteSecurityGroup", "ec2:DescribeSecurityGroups", "ec2:RevokeSecurityGroupEgress", "ec2:AuthorizeSecurityGroupEgress", "ec2:AuthorizeSecurityGroupIngress", "ec2:RevokeSecurityGroupIngress",	Creates predefined security groups for Cloud Volumes ONTAP.
"ec2:CreateNetworkInterface", "ec2:DescribeNetworkInterfaces", "ec2:DeleteNetworkInterface", "ec2:ModifyNetworkInterfaceAttribute",	Creates and manages network interfaces for Cloud Volumes ONTAP in the target subnet.
"ec2:DescribeSubnets", "ec2:DescribeVpcs",	Gets the list of destination subnets and security groups, which is needed when creating a new working environment for Cloud Volumes ONTAP.
"ec2:DescribeDhcpOptions",	Determines DNS servers and the default domain name when launching Cloud Volumes ONTAP instances.
"ec2:CreateSnapshot", "ec2:DeleteSnapshot", "ec2:DescribeSnapshots",	Takes snapshots of EBS volumes during initial setup and whenever a Cloud Volumes ONTAP instance is stopped.
"ec2:GetConsoleOutput",	Captures the Cloud Volumes ONTAP console, which is attached to AutoSupport messages.
"ec2:DescribeKeyPairs",	Obtains the list of available key pairs when launching instances.
"ec2:DescribeRegions",	Gets a list of available AWS regions.
"ec2:DeleteTags", "ec2:DescribeTags",	Manages tags for resources associated with Cloud Volumes ONTAP instances.
"cloudformation>CreateStack", "cloudformation>DeleteStack", "cloudformation:DescribeStacks", "cloudformation:DescribeStackEvents", "cloudformation:ValidateTemplate",	Launches Cloud Volumes ONTAP instances.

Actions	Purpose
"iam:PassRole", "iam>CreateRole", "iam>DeleteRole", "iam:PutRolePolicy", "iam:CreateInstanceProfile", "iam>DeleteRolePolicy", "iam:AddRoleToInstanceProfile", "iam:RemoveRoleFromInstanceProfile", "iam:DeleteInstanceProfile",	Launches a Cloud Volumes ONTAP HA configuration.
"iam>ListInstanceProfiles", "sts:DecodeAuthorizationMessage", "ec2:AssociateIamInstanceProfile", "ec2:DescribeIamInstanceProfileAssociations", "ec2:DisassociateIamInstanceProfile",	Manages instance profiles for Cloud Volumes ONTAP instances.
"s3:GetBucketTagging", "s3:GetBucketLocation", "s3>ListAllMyBuckets", "s3>ListBucket"	Obtains information about AWS S3 buckets so Cloud Manager can integrate with the NetApp Data Fabric Cloud Sync service.
"s3>CreateBucket", "s3>DeleteBucket", "s3:GetLifecycleConfiguration", "s3:PutLifecycleConfiguration", "s3:PutBucketTagging", "s3>ListBucketVersions", "s3:GetBucketPolicyStatus", "s3:GetBucketPublicAccessBlock", "s3:GetBucketAcl", "s3:GetBucketPolicy", "s3:PutBucketPublicAccessBlock"	Manages the S3 bucket that a Cloud Volumes ONTAP system uses as a capacity tier for data tiering.
"kms>List*", "kms:ReEncrypt*", "kms:Describe*", "kms>CreateGrant",	Enables data encryption of Cloud Volumes ONTAP using the AWS Key Management Service (KMS).
"ce:GetReservationUtilization", "ce:GetDimensionValues", "ce:GetCostAndUsage", "ce:GetTags"	Obtains AWS cost data for Cloud Volumes ONTAP.
"ec2>CreatePlacementGroup", "ec2>DeletePlacementGroup"	When you deploy an HA configuration in a single AWS Availability Zone, Cloud Manager launches the two HA nodes and the mediator in an AWS spread placement group.
"ec2:DescribeReservedInstancesOfferings"	Cloud Manager uses the permission as part of Cloud Data Sense deployment to choose which instance type to use.

Actions	Purpose
"ec2:CreateTags", "ec2>DeleteTags", "ec2:DescribeTags", >tag:getResources", >tag:getTagKeys", >tag:getTagValues", >tag:TagResources", >tag:UntagResources"	Enables you to manage tags on your AWS resources using the Cloud Manager Tagging service.
"s3:DeleteBucket", "s3:GetLifecycleConfiguration", "s3:PutLifecycleConfiguration", "s3:PutBucketTagging", "s3>ListBucketVersions", "s3:GetObject", "s3>ListBucket", "s3>ListAllMyBuckets", "s3:GetBucketTagging", "s3:GetBucketLocation", "s3:GetBucketPolicyStatus", "s3:GetBucketPublicAccessBlock", "s3:GetBucketAcl", "s3:GetBucketPolicy", "s3:PutBucketPublicAccessBlock"	Cloud Manager uses these permissions when you enable the Backup to S3 service.

What Cloud Manager does with Azure permissions

The Cloud Manager Azure policy includes the permissions that Cloud Manager needs to deploy and manage Cloud Volumes ONTAP in Azure.

Actions	Purpose
"Microsoft.Compute/locations/operations/read", "Microsoft.Compute/locations/vmSizes/read", "Microsoft.Compute/operations/read", "Microsoft.Compute/virtualMachines/instanceView/read", "Microsoft.Compute/virtualMachines/powerOff/action", "Microsoft.Compute/virtualMachines/read", "Microsoft.Compute/virtualMachines/restart/action", "Microsoft.Compute/virtualMachines/start/action", "Microsoft.Compute/virtualMachines/deallocate/action", ", "Microsoft.Compute/virtualMachines/vmSizes/read", "Microsoft.Compute/virtualMachines/write",	Creates Cloud Volumes ONTAP and stops, starts, deletes, and obtains the status of the system.
"Microsoft.Compute/images/write", "Microsoft.Compute/images/read",	Enables Cloud Volumes ONTAP deployment from a VHD.

Actions	Purpose
"Microsoft.Compute/disks/delete", "Microsoft.Compute/disks/read", "Microsoft.Compute/disks/write", "Microsoft.Storage/checknameavailability/read", "Microsoft.Storage/operations/read", "Microsoft.Storage/storageAccounts/listkeys/action", "Microsoft.Storage/storageAccounts/read", "Microsoft.Storage/storageAccounts/regeneratekey/action", "Microsoft.Storage/storageAccounts/write" "Microsoft.Storage/storageAccounts/delete", "Microsoft.Storage/usages/read",	Manages Azure storage accounts and disks, and attaches the disks to Cloud Volumes ONTAP.
"Microsoft.Storage/storageAccounts/blobServices/containers/read", "Microsoft.KeyVault/vaults/read", "Microsoft.KeyVault/vaults/accessPolicies/write"	Enables backups to Azure Blob storage and encryption of storage accounts
"Microsoft.Network/networkInterfaces/read", "Microsoft.Network/networkInterfaces/write", "Microsoft.Network/networkInterfaces/join/action",	Creates and manages network interfaces for Cloud Volumes ONTAP in the target subnet.
"Microsoft.Network/networkSecurityGroups/read", "Microsoft.Network/networkSecurityGroups/write", "Microsoft.Network/networkSecurityGroups/join/action",	Creates predefined network security groups for Cloud Volumes ONTAP.
"Microsoft.Resources/subscriptions/locations/read", "Microsoft.Network/locations/operationResults/read", "Microsoft.Network/locations/operations/read", "Microsoft.Network/virtualNetworks/read", "Microsoft.Network/virtualNetworks/checkIpAddressAvailability/read", "Microsoft.Network/virtualNetworks/subnets/read", "Microsoft.Network/virtualNetworks/subnets/virtualMachines/read", "Microsoft.Network/virtualNetworks/virtualMachines/read", "Microsoft.Network/virtualNetworks/subnets/join/action"	Gets network information about regions, the target VNet and subnet, and adds Cloud Volumes ONTAP to VNets.
"Microsoft.Network/virtualNetworks/subnets/write", "Microsoft.Network/routeTables/join/action",	Enables VNet service endpoints for data tiering.
"Microsoft.Resources/deployments/operations/read", "Microsoft.Resources/deployments/read", "Microsoft.Resources/deployments/write",	Deploys Cloud Volumes ONTAP from a template.

Actions	Purpose
"Microsoft.Resources/deployments/operations/read", "Microsoft.Resources/deployments/read", "Microsoft.Resources/deployments/write", "Microsoft.Resources/resources/read", "Microsoft.Resources/subscriptions/operationresults/read", "Microsoft.Resources/subscriptions/resourceGroups/delete", "Microsoft.Resources/subscriptions/resourceGroups/read", "Microsoft.Resources/subscriptions/resourcegroups/resources/read", "Microsoft.Resources/subscriptions/resourceGroups/write",	Creates and manages resource groups for Cloud Volumes ONTAP.
"Microsoft.Compute/snapshots/write", "Microsoft.Compute/snapshots/read", "Microsoft.Compute/snapshots/delete", "Microsoft.Compute/disks/beginGetAccess/action",	Creates and manages Azure managed snapshots.
"Microsoft.Compute/availabilitySets/write", "Microsoft.Compute/availabilitySets/read",	Creates and manages availability sets for Cloud Volumes ONTAP.
"Microsoft.MarketplaceOrdering/offertypes/publishers/offers/plans/agreements/read", "Microsoft.MarketplaceOrdering/offertypes/publishers/offers/plans/agreements/write",	Enables programmatic deployments from the Azure Marketplace.
"Microsoft.Network/loadBalancers/read", "Microsoft.Network/loadBalancers/write", "Microsoft.Network/loadBalancers/delete", "Microsoft.Network/loadBalancers/backendAddressPools/read", "Microsoft.Network/loadBalancers/backendAddressPools/join/action", "Microsoft.Network/loadBalancers/frontendIPConfigurations/read", "Microsoft.Network/loadBalancers/loadBalancingRules/read", "Microsoft.Network/loadBalancers/probes/read", "Microsoft.Network/loadBalancers/probes/join/action",	Manages an Azure load balancer for HA pairs.
"Microsoft.Authorization/locks/*",	Enables management of locks on Azure disks.
"Microsoft.Authorization/roleDefinitions/write", "Microsoft.Authorization/roleAssignments/write", "Microsoft.Web/sites/*"	Manages failover for HA pairs.

Actions	Purpose
"Microsoft.Network/privateEndpoints/write", "Microsoft.Storage/storageAccounts/PrivateEndpointConnectionsApproval/action", "Microsoft.Storage/storageAccounts/privateEndpointConnections/read", "Microsoft.Network/privateEndpoints/read", "Microsoft.Network/privateDnsZones/write", "Microsoft.Network/privateDnsZones/virtualNetworkLinks/write", "Microsoft.Network/virtualNetworks/join/action", "Microsoft.Network/privateDnsZones/A/write", "Microsoft.Network/privateDnsZones/read", "Microsoft.Network/privateDnsZones/virtualNetworkLinks/read",	Enables the management of private endpoints. Private endpoints are used when connectivity isn't provided to outside the subnet. Cloud Manager creates the storage account for HA with only internal connectivity within the subnet.
"Microsoft.NetApp/netAppAccounts/capacityPools/volumes/delete",	Enables Cloud Manager to delete volumes for Azure NetApp Files.
"Microsoft.Resources/deployments/operationStatuses/read"	Azure requires this permission for some virtual machine deployments (it depends on the underlying physical hardware that's used during deployment).
"Microsoft.Resources/deployments/operationStatuses/read", "Microsoft.Insights/Metrics/Read", "Microsoft.Compute/virtualMachines/extensions/write", "Microsoft.Compute/virtualMachines/extensions/read", "Microsoft.Compute/virtualMachines/extensions/delete", "Microsoft.Compute/virtualMachines/delete", "Microsoft.Network/networkInterfaces/delete", "Microsoft.Network/networkSecurityGroups/delete", "Microsoft.Resources/deployments/delete",	Enables you to use Global File Cache.
"Microsoft.Network/privateEndpoints/delete", "Microsoft.Compute/availabilitySets/delete",	Enables Cloud Manager to remove resources from a resource group that belong to Cloud Volumes ONTAP in case of deployment failure or deletion.
"Microsoft.Compute/diskEncryptionSets/read" "Microsoft.Compute/diskEncryptionSets/write", "Microsoft.Compute/diskEncryptionSets/delete" "Microsoft.KeyVault/vaults/deploy/action", "Microsoft.KeyVault/vaults/read", "Microsoft.KeyVault/vaults/accessPolicies/write"	Enables use of customer-managed encryption keys with Cloud Volumes ONTAP. This feature is supported using APIs.
"Microsoft.Resources/tags/read", "Microsoft.Resources/tags/write", "Microsoft.Resources/tags/delete"	Enables you to manage tags on your Azure resources using the Cloud Manager Tagging service.

What Cloud Manager does with GCP permissions

The Cloud Manager policy for GCP includes the permissions that Cloud Manager needs to deploy and manage Cloud Volumes ONTAP.

Actions	Purpose
<ul style="list-style-type: none"> - compute.disks.create - compute.disks.createSnapshot - compute.disks.delete - compute.disks.get - compute.disks.list - compute.disks.setLabels - compute.disks.use 	To create and manage disks for Cloud Volumes ONTAP.
<ul style="list-style-type: none"> - compute.firewalls.create - compute.firewalls.delete - compute.firewalls.get - compute.firewalls.list 	To create firewall rules for Cloud Volumes ONTAP.
<ul style="list-style-type: none"> - compute.globalOperations.get 	To get the status of operations.
<ul style="list-style-type: none"> - compute.images.get - compute.images.getFromFamily - compute.images.list - compute.images.useReadOnly 	To get images for VM instances.
<ul style="list-style-type: none"> - compute.instances.attachDisk - compute.instances.detachDisk 	To attach and detach disks to Cloud Volumes ONTAP.
<ul style="list-style-type: none"> - compute.instances.create - compute.instances.delete 	To create and delete Cloud Volumes ONTAP VM instances.
<ul style="list-style-type: none"> - compute.instances.get 	To list VM instances.
<ul style="list-style-type: none"> - compute.instances.getSerialPortOutput 	To get console logs.
<ul style="list-style-type: none"> - compute.instances.list 	To retrieve the list of instances in a zone.
<ul style="list-style-type: none"> - compute.instances.setDeletionProtection 	To set deletion protection on the instance.
<ul style="list-style-type: none"> - compute.instances.setLabels 	To add labels.
<ul style="list-style-type: none"> - compute.instances.setMachineType - compute.instances.setMinCpuPlatform 	To change the machine type for Cloud Volumes ONTAP.
<ul style="list-style-type: none"> - compute.instances.setMetadata 	To add metadata.
<ul style="list-style-type: none"> - compute.instances.setTags 	To add tags for firewall rules.
<ul style="list-style-type: none"> - compute.instances.start - compute.instances.stop - compute.instances.updateDisplayDevice 	To start and stop Cloud Volumes ONTAP.
<ul style="list-style-type: none"> - compute.machineTypes.get 	To get the numbers of cores to check quotas.
<ul style="list-style-type: none"> - compute.projects.get 	To support multi-projects.
<ul style="list-style-type: none"> - compute.snapshots.create - compute.snapshots.delete - compute.snapshots.get - compute.snapshots.list - compute.snapshots.setLabels 	To create and manage persistent disk snapshots.

Actions	Purpose
<ul style="list-style-type: none"> - compute.networks.get - compute.networks.list - compute.regions.get - compute.regions.list - compute.subnetworks.get - compute.subnetworks.list - compute.zoneOperations.get - compute.zones.get - compute.zones.list 	To get the networking information needed to create a new Cloud Volumes ONTAP virtual machine instance.
<ul style="list-style-type: none"> - deploymentmanager compositeTypes.get - deploymentmanager compositeTypes.list - deploymentmanager deployments.create - deploymentmanager deployments.delete - deploymentmanager deployments.get - deploymentmanager deployments.list - deploymentmanager manifests.get - deploymentmanager manifests.list - deploymentmanager operations.get - deploymentmanager operations.list - deploymentmanager resources.get - deploymentmanager resources.list - deploymentmanager typeProviders.get - deploymentmanager typeProviders.list - deploymentmanager types.get - deploymentmanager types.list 	To deploy the Cloud Volumes ONTAP virtual machine instance using Google Cloud Deployment Manager.
<ul style="list-style-type: none"> - logging logEntries.list - logging privateLogEntries.list 	To get stack log drives.
<ul style="list-style-type: none"> - resourcemanager projects.get 	To support multi-projects.
<ul style="list-style-type: none"> - storage buckets.create - storage buckets.delete - storage buckets.get - storage buckets.list - storage buckets.update 	To create and manage a Google Cloud Storage bucket for data tiering.
<ul style="list-style-type: none"> - cloudkms cryptoKeyVersions useToEncrypt - cloudkms cryptoKeys.get - cloudkms cryptoKeys.list - cloudkms keyRings.list 	To use customer-managed encryption keys from the Cloud Key Management Service with Cloud Volumes ONTAP.
<ul style="list-style-type: none"> - compute instances setServiceAccount - iam serviceAccounts actAs - iam serviceAccounts getIamPolicy - iam serviceAccounts.list - storage objects.get - storage objects.list 	To set a service account on the Cloud Volumes ONTAP instance. This service account provides permissions for data tiering to a Google Cloud Storage bucket.

Actions	Purpose
<ul style="list-style-type: none"> - compute.addresses.list - compute.backendServices.create - compute.networks.updatePolicy - compute.regionBackendServices.create - compute.regionBackendServices.get - compute.regionBackendServices.list 	To deploy HA pairs.
<ul style="list-style-type: none"> - compute.subnetworks.use - compute.subnetworks.useExternalIp - compute.instances.addAccessConfig 	To enable Cloud Data Sense.

AWS Marketplace pages for Cloud Manager and Cloud Volumes ONTAP

Several offerings are available in the AWS Marketplace for Cloud Manager and Cloud Volumes ONTAP. If need help understanding the purpose of each page, read the descriptions below.

In all cases, remember that you can't launch Cloud Volumes ONTAP in AWS from the AWS Marketplace. You need to launch it directly from Cloud Manager.

Goal	AWS Marketplace page to use	More information
Enable the use of Cloud Volumes ONTAP PAYGO, Cloud Tiering, Cloud Data Sense (Cloud Compliance), and other add-on services	Cloud Manager - Deploy & Manage NetApp Cloud Data Services	<p>This subscription enables charging for the PAYGO version of Cloud Volumes ONTAP 9.6 and later. It also enables charging for Cloud Tiering, Cloud Data Sense, and other add-on services.</p> <p>You should subscribe to this offering when Cloud Manager prompts you and redirects you to the page. Cloud Manager prompts you in the Working Environment wizard or when you add new credentials in the Settings.</p> <p>This page doesn't enable you to launch Cloud Manager in AWS. That should be done from NetApp Cloud Central, or alternatively using the AMI listed in row 3 of this table.</p>
Enable the use of Cloud Volumes ONTAP PAYGO, Cloud Tiering, Cloud Data Sense, and other add-on services <i>using an annual contract</i>	Cloud Manager (Contracts) - Deploy & Manage NetApp Cloud Data Services	This subscription is an alternative to the subscription in the first row. It enables you to get an annual upfront payment for the listings. It's mostly for NetApp partners.

Goal	AWS Marketplace page to use	More information
Deploy Cloud Manager from the AWS Marketplace using an AMI	Cloud Manager - Manual installation without access keys	We recommend that you launch Cloud Manager in AWS from NetApp Cloud Central , but you can launch it from this AWS Marketplace page, if you prefer.
Enable deployment of Cloud Volumes ONTAP PAYGO (9.5 or earlier)	<ul style="list-style-type: none"> • Cloud Volumes ONTAP for AWS • Cloud Volumes ONTAP for AWS - High Availability 	<p>These AWS Marketplace pages enable you to subscribe to the single node or HA versions of Cloud Volumes ONTAP PAYGO for versions 9.5 and earlier.</p> <p>Starting with version 9.6, you need to subscribe through the AWS Marketplace page listed in row 1 of this table for PAYGO deployments.</p>

Automate with the REST API

Automation using the Cloud Manager REST API platform

In addition to the web browser interface, Cloud Manager has a REST API that provides software developers with direct access to the Cloud Manager functionality through the SaaS interface. The Cloud Manager service consists of several distinct components that collectively form an extensible development platform.

To learn more, go to the [Cloud Manager automation docs](#).

Automation resources for infrastructure as code

Use the resources on this page to get help integrating Cloud Manager and Cloud Volumes ONTAP with your [infrastructure as code](#) (IaC).

DevOps teams use a variety of tools to automate the setup of new environments, which allows them to treat infrastructure as code. One such tool is Terraform. We have developed a Terraform provider that DevOps teams can use with Cloud Manager to automate and integrate Cloud Volumes ONTAP with infrastructure as code.

[View the netapp-cloudmanager provider](#).

Related links

- [NetApp Cloud Blog: Using Cloud Manager REST APIs with Federated Access](#)
- [NetApp Cloud Blog: Cloud Automation with Cloud Volumes ONTAP and REST](#)
- [NetApp Cloud Blog: Automated Data Cloning for Cloud-Based Testing of Software Applications](#)
- [NetApp Blog: Infrastructure-As-Code \(IaC\) Accelerated with Ansible + NetApp](#)
- [NetApp thePub: Configuration Management & Automation with Ansible](#)
- [NetApp thePub: Roles for Ansible ONTAP use](#)

Where to get help and find more information

You can get help and find more information about Cloud Manager and Cloud Volumes ONTAP through various resources, including videos, forums, and support.

- [NetApp Cloud Volumes ONTAP Support](#)

Access support resources to get help and troubleshoot issues with Cloud Volumes ONTAP.

- [Videos for Cloud Manager and Cloud Volumes ONTAP](#)

Watch videos that show you how to deploy and manage Cloud Volumes ONTAP and how to replicate data across your hybrid cloud.

- [Policies for Cloud Manager](#)

Download JSON files that include the permissions that Cloud Manager needs to perform actions in a cloud provider.

- [Cloud Manager API documentation](#)

Read an overview of the APIs, examples of how to use them, and an API reference.

- Training for Cloud Volumes ONTAP

- [Cloud Volumes ONTAP Fundamentals](#)
- [Cloud Volumes ONTAP Deployment and Management for Azure](#)
- [Cloud Volumes ONTAP Deployment and Management for AWS](#)

- Technical reports

- [NetApp Technical Report 4383: Performance Characterization of Cloud Volumes ONTAP in Amazon Web Services with Application Workloads](#)
- [NetApp Technical Report 4671: Performance Characterization of Cloud Volumes ONTAP in Azure with Application Workloads](#)
- [NetApp Technical Report 4816: Performance Characterization of Cloud Volumes ONTAP for Google Cloud](#)

- SVM disaster recovery

SVM disaster recovery is the asynchronous mirroring of SVM data and configuration from a source SVM to a destination SVM. You can quickly activate a destination SVM for data access if the source SVM is no longer available.

- [Cloud Volumes ONTAP 9 SVM Disaster Recovery Preparation Express Guide](#)

Describes how to quickly configure a destination SVM in preparation for disaster recovery.

- [Cloud Volumes ONTAP 9 SVM Disaster Recovery Express Guide](#)

Describes how to quickly activate a destination SVM after a disaster, and then reactivate the source SVM.

- [FlexCache Volumes for Faster Data Access Power Guide](#)

Describes how to create and manage FlexCache volumes in the same cluster or different cluster as the origin volume for accelerating data access.

- [Security advisories](#)

Identify known vulnerabilities (CVEs) for NetApp products, including ONTAP. Note that you can remediate security vulnerabilities for Cloud Volumes ONTAP by following ONTAP documentation.

- [ONTAP 9 Documentation Center](#)

Access product documentation for ONTAP, which can help you as you use Cloud Volumes ONTAP.

- [NetApp Community: Cloud Data Services](#)

Connect with peers, ask questions, exchange ideas, find resources, and share best practices.

- [NetApp Cloud Central](#)

Find information about additional NetApp products and solutions for the cloud.

- [NetApp Product Documentation](#)

Search NetApp product documentation for instructions, resources, and answers.

Earlier versions of Cloud Manager documentation

Documentation for previous releases of Cloud Manager is available in case you're not running the latest version.

- [Cloud Manager 3.8](#)
- [Cloud Manager 3.7](#)
- [Cloud Manager 3.6](#)
- [Cloud Manager 3.5](#)
- [Cloud Manager 3.4](#)
- [Cloud Manager 3.3](#)
- [Cloud Manager 3.2](#)

Legal notices

Legal notices provide access to copyright statements, trademarks, patents, and more.

Copyright

<http://www.netapp.com/us/legal/copyright.aspx>

Trademarks

NETAPP, the NETAPP logo, and the marks listed on the NetApp Trademarks page are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.

<http://www.netapp.com/us/legal/netapptmlist.aspx>

Patents

A current list of NetApp owned patents can be found at:

<https://www.netapp.com/us/media/patents-page.pdf>

Privacy policy

<https://www.netapp.com/us/legal/privacypolicy/index.aspx>

Open source

Notice files provide information about third-party copyright and licenses used in NetApp software.

- [Notice for Cloud Manager 3.9](#)
- [Notice for the Cloud Backup](#)
- [Notice for Single File Restore](#)
- [Notice for Global File Cache](#)
- [Notice for Cloud Sync](#)
- [Notice for Cloud Tiering](#)
- [Notice for Cloud Data Sense](#)
- [Notice for Application Templates](#)

Copyright Information

Copyright © 2021 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system-without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP "AS IS" AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark Information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.